

1 Main Page	1
1.0.0.1 Modules	1
2 Installation Instructions	3
2.0.1 Building	3
2.0.2 Building Ginkgo in Windows	5
2.0.3 Building Ginkgo with HIP support	5
2.0.3.1 Changing the paths to search for HIP and other packages	6
2.0.3.2 HIP platform detection of AMD and NVIDIA	6
2.0.3.3 Setting platform specific compilation flags	6
2.0.4 Third party libraries and packages	7
2.0.5 Installing Ginkgo	7
3 Testing Instructions	9
3.0.1 Running the unit tests	9
3.0.1.1 Using make test	9
3.0.1.2 Using make quick_test	9
3.0.1.3 Using CTest	9
4 Running the benchmarks	11
4.0.1 1: Ginkgo setup and best practice guidelines	11
4.0.2 2: Using ssget to fetch the matrices	12
4.0.3 3: Benchmarking overview	12
4.0.4 4: Publishing the results on Github and analyze the results with the GPE (optional)	13
4.0.5 5: Detailed performance analysis and debugging	14
4.0.6 6: Available benchmark options	14
5 Contributing guidelines	17
5.1 Table of Contents	17
5.2 Most important stuff (A TL;DR)	18
5.3 Project structure	18
5.3.1 Extended header files	18
5.3.2 Using library classes	19
5.4 Git related	19
5.4.1 Our git workflow	19
5.4.2 Writing good commit messages	19
5.4.2.1 Attributing credit	19
5.4.3 Creating, Reviewing and Merging Pull Requests	20
5.5 Code style	20
5.5.1 Automatic code formatting	20
5.5.2 Naming scheme	20
5.5.2.1 Filenames	20
5.5.2.2 Macros	21
5.5.2.3 Variables	21

7 Example programs	31
6.0.4 On SpMV or solvers performance	30
6.0.3 On Software Sustainability	30
6.0.2 On Portability	29
6.0.1 The Ginkgo Software	29
6 Citing Ginkgo	29
5.9.3 Avoiding circular dependencies	27
5.9.2 Warnings	27
5.9.1 C++ standard stream objects	27
5.9 Other programming comments	27
5.8.3 Documenting examples	27
5.8.2.1 After named tags such as <tt>@param foo</tt>	26
5.8.2 Whitespaces	26
5.8.1 Developer targeted notes	26
5.8 Documentation style	_
5.7.3 Writing tests for kernels	26
5.7.2 Some general rules	26
5.7.1 Testing know-how	26
5.7 Writing Tests	26
5.6.2 Converting CUDA code to HIP code	25
5.6.1 Create a new algorithm	25 25
5.6 Helper scripts	
5.5.6.3 Naming style	25
5.5.6.2 Use of macros vs functions	
5.5.6.1 Whitespaces	25 25
5.5.6 CMake coding style	24 25
5.5.5.1 Control flow constructs	24 24
5.5.5 Other Code Formatting not handled by ClangFormat	24 24
5.5.4.3 Automatic header arrangement	
5.5.4.2 Some general comments	
5.5.4.1 Main header	23
5.5.4 Include statement grouping	23
5.5.3 Whitespace	22
5.5.2.9 Template parameters	22
5.5.2.8 Namespaces	22
5.5.2.7 Members	21
5.5.2.6 Structures and classes	
5.5.2.5 Functions	21
5.5.2.4 Constants	21

8 The adaptiveprecision-blockjacobi program	35
9 The cb-gmres program	39
10 The custom-logger program	45
11 The custom-matrix-format program	55
12 The custom-stopping-criterion program	63
13 The distributed-solver program	69
14 The external-lib-interfacing program	77
15 The ginkgo-overhead program	99
16 The ginkgo-ranges program	103
17 The heat-equation program	107
18 The ilu-preconditioned-solver program	113
19 The inverse-iteration program	117
20 The ir-ilu-preconditioned-solver program	123
21 The iterative-refinement program	129
22 The kokkos_assembly program	133
23 The minimal-cuda-solver program	139
24 The mixed-multigrid-preconditioned-solver program	141
25 The mixed-multigrid-solver program	149
26 The mixed-precision-ir program	155
27 The mixed-spmv program	161
28 The multigrid-preconditioned-solver program	169
29 The multigrid-preconditioned-solver-customized program	175
30 The nine-pt-stencil-solver program	181
31 The papi-logging program	191
32 The par-ilu-convergence program	197
33 The performance-debugging program	203

34 The poisson-solver program	215
35 The preconditioned-solver program	221
36 The preconditioner-export program	225
37 The schroedinger-splitting program	233
38 The simple-solver program	239
39 The simple-solver-logging program	245
40 The three-pt-stencil-solver program	271
41 Module Documentation	279
41.1 CUDA Executor	279
41.1.1 Detailed Description	279
41.2 DPC++ Executor	280
41.2.1 Detailed Description	280
41.3 Executors	281
41.3.1 Detailed Description	281
41.3.2 Executors in Ginkgo.	282
41.3.3 Macro Definition Documentation	282
41.3.3.1 GKO_REGISTER_HOST_OPERATION	282
41.3.3.2 Example	283
41.3.3.3 GKO_REGISTER_OPERATION	283
41.3.3.4 Example	283
41.4 Factorizations	285
41.4.1 Detailed Description	285
41.5 HIP Executor	286
41.5.1 Detailed Description	286
41.6 Jacobi Preconditioner	287
41.6.1 Detailed Description	287
41.7 Linear Operators	288
41.7.1 Detailed Description	292
41.7.2 Advantages of this approach and usage	292
41.7.3 Linear operator as a concept	292
41.7.4 Macro Definition Documentation	293
41.7.4.1 GKO_CREATE_FACTORY_PARAMETERS	293
41.7.4.2 GKO_ENABLE_BUILD_METHOD	294
41.7.4.3 GKO_ENABLE_LIN_OP_FACTORY	294
41.7.4.4 GKO_FACTORY_PARAMETER	
41.7.4.5 GKO_FACTORY_PARAMETER_SCALAR	295
41.7.4.6 GKO_FACTORY_PARAMETER_VECTOR	
41.7.5 Typedef Documentation	

41.7.5.1 EnableDefaultLinOpFactory	. 296
41.7.6 Function Documentation	. 297
41.7.6.1 initialize() [1/4]	. 297
41.7.6.2 initialize() [2/4]	. 298
41.7.6.3 initialize() [3/4]	. 298
41.7.6.4 initialize() [4/4]	. 299
41.8 Logging	. 300
41.8.1 Detailed Description	. 300
41.9 SpMV employing different Matrix formats	. 301
41.9.1 Detailed Description	. 301
41.10 OpenMP Executor	. 302
41.10.1 Detailed Description	. 302
41.11 Preconditioners	. 303
41.11.1 Detailed Description	. 303
41.12 Reference Executor	. 304
41.12.1 Detailed Description	. 304
41.13 Solvers	. 305
41.13.1 Detailed Description	. 305
41.14 Stopping criteria	. 306
41.14.1 Detailed Description	. 307
41.14.2 Enumeration Type Documentation	. 307
41.14.2.1 mode	. 307
41.14.3 Function Documentation	. 307
41.14.3.1 combine()	. 307
42 Namespace Documentation	309
42.1 gko Namespace Reference	
42.1.1 Detailed Description	
42.1.2 Typedef Documentation	
42.1.2.1 highest_precision	
42.1.2.2 is_complex_or_scalar_s	
42.1.2.3 is_complex_s	
42.1.2.4 previous_precision	
42.1.2.5 remove_complex	
42.1.2.6 to_complex	
42.1.2.7 to_real	
42.1.3 Enumeration Type Documentation	
42.1.3.1 allocation_mode	
42.1.3.2 layout_type	
42.1.3.3 log_propagation_mode	
42.1.4 Function Documentation	
42.1.4.1 abs()	. 324

42.1.4.2 as() [1/7]
42.1.4.3 as() [2/7]
42.1.4.4 as() [3/7]
42.1.4.5 as() [4/7]
42.1.4.6 as() [5/7]
42.1.4.7 as() [6/7]
42.1.4.8 as() [7/7]
42.1.4.9 ceildiv()
42.1.4.10 clone() [1/2]
42.1.4.11 clone() [2/2]
42.1.4.12 conj()
42.1.4.13 copy_and_convert_to() [1/4]
42.1.4.14 copy_and_convert_to() [2/4]
42.1.4.15 copy_and_convert_to() [3/4]
42.1.4.16 copy_and_convert_to() [4/4]
42.1.4.17 get_significant_bit()
42.1.4.18 get_superior_power()
42.1.4.19 give()
42.1.4.20 imag()
42.1.4.21 is_complex()
42.1.4.22 is_complex_or_scalar()
42.1.4.23 is_finite() [1/2]
42.1.4.24 is_finite() [2/2]
42.1.4.25 is_nan() [1/2]
42.1.4.26 is_nan() [2/2]
42.1.4.27 is_nonzero()
42.1.4.28 is_zero()
42.1.4.29 lend() [1/2]
42.1.4.30 lend() [2/2]
42.1.4.31 make_array_view()
42.1.4.32 make_const_array_view()
42.1.4.33 make_const_dense_view()
42.1.4.34 make_dense_view()
42.1.4.35 make_temporary_clone()
42.1.4.36 make_temporary_conversion()
42.1.4.37 make_temporary_output_clone()
42.1.4.38 max()
42.1.4.39 min()
42.1.4.40 mixed_precision_dispatch()
42.1.4.41 mixed_precision_dispatch_real_complex()
42.1.4.42 nan() [1/2]
42.1.4.43 nan() [2/2]

42.1.4.44 one() [1/2]	47
42.1.4.45 one() [2/2]	47
42.1.4.46 operator"!=() [1/3]	48
42.1.4.47 operator"!=() [2/3]	48
42.1.4.48 operator"!=() [3/3]	49
42.1.4.49 operator<<() [1/2]	49
42.1.4.50 operator<<() [2/2]	49
42.1.4.51 operator==() [1/2]	
42.1.4.52 operator==() [2/2]	51
42.1.4.53 pi()	52
42.1.4.54 precision_dispatch()	52
42.1.4.55 precision_dispatch_real_complex() [1/3]	
42.1.4.56 precision_dispatch_real_complex() [2/3]	
42.1.4.57 precision_dispatch_real_complex() [3/3]	
42.1.4.58 read()	53
42.1.4.59 read_binary()	54
42.1.4.60 read_binary_raw()	55
42.1.4.61 read_generic()	
42.1.4.62 read_generic_raw()	
42.1.4.63 read_raw()	
42.1.4.64 real()	
42.1.4.65 reduce_add() [1/2]	
42.1.4.66 reduce_add() [2/2]	
42.1.4.67 round_down()	
42.1.4.68 round_up()	
42.1.4.69 safe_divide()	
42.1.4.70 share()	
42.1.4.71 squared_norm()	
42.1.4.72 transpose()	
42.1.4.73 unit_root()	
42.1.4.74 with_matrix_type()	
42.1.4.75 write()	
42.1.4.76 write_binary()	
42.1.4.77 write_binary_raw()	
42.1.4.78 write_raw()	
42.1.4.79 zero() [1/2]	
42.1.4.80 zero() [2/2]	
42.2 gko::accessor Namespace Reference	
42.2.1 Detailed Description	
42.3 gko::experimental::distributed Namespace Reference	
42.3.1 Detailed Description	
42.3.2 Typedef Documentation	67

42.3.2.1 comm_index_type
42.3.3 Function Documentation
42.3.3.1 make_temporary_conversion() [1/2]
42.3.3.2 make_temporary_conversion() [2/2]
42.3.3.3 precision_dispatch()
42.3.3.4 precision_dispatch_real_complex() [1/3]
42.3.3.5 precision_dispatch_real_complex() [2/3]
42.3.3.6 precision_dispatch_real_complex() [3/3]
42.4 gko::experimental::distributed::preconditioner Namespace Reference
42.4.1 Detailed Description
42.5 gko::experimental::mpi Namespace Reference
42.5.1 Detailed Description
42.5.2 Function Documentation
42.5.2.1 get_walltime()
42.5.2.2 map_rank_to_device_id()
42.5.2.3 wait_all()
42.6 gko::experimental::reorder Namespace Reference
42.6.1 Detailed Description
42.7 gko::factorization Namespace Reference
42.7.1 Detailed Description
42.8 gko::log Namespace Reference
42.8.1 Detailed Description
42.8.2 Enumeration Type Documentation
42.8.2.1 profile_event_category
42.9 gko::matrix Namespace Reference
42.9.1 Detailed Description
42.10 gko::multigrid Namespace Reference
42.10.1 Detailed Description
42.11 gko::name_demangling Namespace Reference
42.11.1 Detailed Description
42.11.2 Function Documentation
42.11.2.1 get_dynamic_type()
42.11.2.2 get_static_type()
42.12 gko::preconditioner Namespace Reference
42.12.1 Detailed Description
42.12.2 Enumeration Type Documentation
42.12.2.1 isai_type
42.13 gko::reorder Namespace Reference
42.13.1 Detailed Description
42.13.2 Typedef Documentation
42.13.2.1 EnableDefaultReorderingBaseFactory
42.14 gko::solver Namespace Reference

42.14.1 Detailed Description	382
42.14.2 Enumeration Type Documentation	382
42.14.2.1 initial_guess_mode	382
42.14.2.2 trisolve_algorithm	383
42.14.3 Function Documentation	383
42.14.3.1 build_smoother() [1/2]	383
42.14.3.2 build_smoother() [2/2]	384
42.15 gko::solver::multigrid Namespace Reference	384
42.15.1 Detailed Description	384
42.15.2 Enumeration Type Documentation	384
42.15.2.1 cycle	384
42.15.2.2 mid_smooth_type	385
42.16 gko::stop Namespace Reference	385
42.16.1 Detailed Description	386
42.16.2 Typedef Documentation	386
42.16.2.1 EnableDefaultCriterionFactory	386
42.17 gko::syn Namespace Reference	387
42.17.1 Detailed Description	387
42.17.2 Typedef Documentation	387
42.17.2.1 as_list	387
42.17.2.2 concatenate	388
42.17.3 Function Documentation	388
42.17.3.1 as_array()	388
42.18 gko::xstd Namespace Reference	389
42.18.1 Detailed Description	389
43 Class Documentation	391
43.1 gko::AbsoluteComputable Class Reference	
43.1.1 Detailed Description	
·	391
43.1.2.1 compute_absolute_linop()	
. – –	392
	392
·	392
	392
·	393
	393
	393
	394
·	394
43.4.2.1 AllocationError()	394
43.5 gko::experimental::reorder::Amd< IndexType > Class Template Reference	

43.5.1 Detailed Description
43.5.2 Member Function Documentation
43.5.2.1 generate()
43.6 gko::amd_device Class Reference
43.6.1 Detailed Description
43.7 gko::solver::ApplyWithInitialGuess Class Reference
43.7.1 Detailed Description
43.8 gko::are_all_integral < Args > Struct Template Reference
43.8.1 Detailed Description
43.9 gko::array< ValueType > Class Template Reference
43.9.1 Detailed Description
43.9.2 Constructor & Destructor Documentation
43.9.2.1 array() [1/11]
43.9.2.2 array() [2/11]
43.9.2.3 array() [3/11]
43.9.2.4 array() [4/11]
43.9.2.5 array() [5/11]
43.9.2.6 array() [6/11]
43.9.2.7 array() [7/11]
43.9.2.8 array() [8/11]
43.9.2.9 array() [9/11]
43.9.2.10 array() [10/11]
43.9.2.11 array() [11/11]
43.9.3 Member Function Documentation
43.9.3.1 as_const_view()
43.9.3.2 as_view()
43.9.3.3 clear()
43.9.3.4 const_view()
43.9.3.5 fill()
43.9.3.6 get_const_data()
43.9.3.7 get_data()
43.9.3.8 get_executor()
43.9.3.9 get_num_elems()
43.9.3.10 is_owning()
43.9.3.11 operator=() [1/3]
43.9.3.12 operator=() [2/3]
43.9.3.13 operator=() [3/3]
43.9.3.14 resize_and_reset()
43.9.3.15 set_executor()
43.9.3.16 view()
43.10 gko::device_matrix_data< ValueType, IndexType >::arrays Struct Reference
43.10.1 Detailed Description

43.11 gko::matrix::Hybrid< ValueType, IndexType >::automatic Class Reference	411
43.11.1 Detailed Description	411
43.11.2 Member Function Documentation	411
43.11.2.1 compute_ell_num_stored_elements_per_row()	411
43.12 gko::BadDimension Class Reference	412
43.12.1 Detailed Description	412
43.12.2 Constructor & Destructor Documentation	412
43.12.2.1 BadDimension()	412
43.13 gko::solver::Bicg ValueType > Class Template Reference	413
43.13.1 Detailed Description	413
43.13.2 Member Function Documentation	413
43.13.2.1 apply_uses_initial_guess()	413
43.13.2.2 conj_transpose()	414
43.13.2.3 transpose()	414
43.14 gko::solver::Bicgstab < ValueType > Class Template Reference	414
43.14.1 Detailed Description	415
43.14.2 Member Function Documentation	415
43.14.2.1 apply_uses_initial_guess()	415
43.14.2.2 conj_transpose()	416
43.14.2.3 transpose()	416
43.15 gko::preconditioner::block_interleaved_storage_scheme < IndexType > Struct Template Reference	416
43.15.1 Detailed Description	417
43.15.2 Member Function Documentation	417
43.15.2.1 compute_storage_space()	417
43.15.2.2 get_block_offset()	418
43.15.2.3 get_global_block_offset()	418
43.15.2.4 get_group_offset()	419
43.15.2.5 get_group_size()	419
43.15.2.6 get_stride()	419
43.15.3 Member Data Documentation	420
43.15.3.1 group_power	420
43.16~gko::BlockSizeError < IndexType > Class~Template~Reference~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.	420
43.16.1 Detailed Description	420
43.16.2 Constructor & Destructor Documentation	420
43.16.2.1 BlockSizeError()	421
43.17 gko::solver::CbGmres < ValueType > Class Template Reference	421
43.17.1 Detailed Description	421
43.17.2 Member Function Documentation	422
43.17.2.1 get_krylov_dim()	422
43.17.2.2 get_storage_precision()	422
43.17.2.3 set_krylov_dim()	422
43.18 gkg···solver··Cg< ValueType > Class Template Reference	423

43.18.1 Detailed Description	423
43.18.2 Member Function Documentation	423
43.18.2.1 apply_uses_initial_guess()	423
43.18.2.2 conj_transpose()	424
43.18.2.3 transpose()	424
43.19 gko::solver::Cgs< ValueType > Class Template Reference	424
43.19.1 Detailed Description	425
43.19.2 Member Function Documentation	425
43.19.2.1 apply_uses_initial_guess()	425
43.19.2.2 conj_transpose()	426
43.19.2.3 transpose()	426
43.20~gko:: experimental:: factorization:: Cholesky < Value Type, Index Type > Class~Template~Reference~~.~~	426
43.20.1 Detailed Description	426
43.20.2 Member Function Documentation	427
43.20.2.1 generate()	427
43.21 gko::matrix::Csr< ValueType, IndexType >::classical Class Reference	427
43.21.1 Detailed Description	428
43.21.2 Member Function Documentation	428
43.21.2.1 clac_size()	428
43.21.2.2 copy()	428
43.21.2.3 process()	429
43.22 gko::matrix::Hybrid< ValueType, IndexType >::column_limit Class Reference	429
43.22.1 Detailed Description	429
43.22.2 Constructor & Destructor Documentation	430
43.22.2.1 column_limit()	430
43.22.3 Member Function Documentation	430
43.22.3.1 compute_ell_num_stored_elements_per_row()	430
43.22.3.2 get_num_columns()	430
43.23 gko::Combination < ValueType > Class Template Reference	431
43.23.1 Detailed Description	431
43.23.2 Constructor & Destructor Documentation	432
43.23.2.1 Combination() [1/2]	432
43.23.2.2 Combination() [2/2]	432
43.23.3 Member Function Documentation	432
43.23.3.1 conj_transpose()	432
43.23.3.2 get_coefficients()	433
43.23.3.3 get_operators()	433
43.23.3.4 operator=() [1/2]	433
43.23.3.5 operator=() [2/2]	433
43.23.3.6 transpose()	434
43.24 gko::stop::Combined Class Reference	434
43.24.1 Detailed Description	434

43.25 gko::experimental::mpi::communicator Class Reference	34
43.25.1 Detailed Description	37
43.25.2 Constructor & Destructor Documentation	38
43.25.2.1 communicator() [1/3]	38
43.25.2.2 communicator() [2/3]	38
43.25.2.3 communicator() [3/3]	38
43.25.3 Member Function Documentation	39
43.25.3.1 all_gather()	39
43.25.3.2 all_reduce() [1/2]	40
43.25.3.3 all_reduce() [2/2]	40
43.25.3.4 all_to_all() [1/2]	41
43.25.3.5 all_to_all() [2/2]	41
43.25.3.6 all_to_all_v() [1/2]4	42
43.25.3.7 all_to_all_v() [2/2]4	43
43.25.3.8 broadcast()	43
43.25.3.9 gather()	44
43.25.3.10 gather_v()	44
43.25.3.11 get()	45
43.25.3.12 i_all_gather()	45
43.25.3.13 i_all_reduce() [1/2] 4	46
43.25.3.14 i_all_reduce() [2/2] 4	47
43.25.3.15 i_all_to_all() [1/2] 44	47
43.25.3.16 i_all_to_all() [2/2] 44	48
43.25.3.17 i_all_to_all_v() [1/2]	49
43.25.3.18 i_all_to_all_v() [2/2]	49
43.25.3.19 i_broadcast()	50
43.25.3.20 i_gather()	51
43.25.3.21 i_gather_v()	52
43.25.3.22 i_recv()	52
43.25.3.23 i_reduce()	53
43.25.3.24 i_scan()	54
43.25.3.25 i_scatter()	55
43.25.3.26 i_scatter_v()	55
43.25.3.27 i_send()	56
43.25.3.28 node_local_rank()	57
43.25.3.29 operator"!=()	57
43.25.3.30 operator==()	57
43.25.3.31 rank()	58
43.25.3.32 recv()	58
43.25.3.33 reduce()	58
43.25.3.34 scan()	59
43.25.3.35 scatter()	60

43.25.3.36 scatter_v()	60
43.25.3.37 send()	62
43.25.3.38 size()	63
43.25.3.39 synchronize()	63
43.26 gko::Composition < ValueType > Class Template Reference	63
43.26.1 Detailed Description	64
43.26.2 Constructor & Destructor Documentation	64
43.26.2.1 Composition() [1/2]	64
43.26.2.2 Composition() [2/2]	64
43.26.3 Member Function Documentation	65
43.26.3.1 conj_transpose()	65
43.26.3.2 get_operators()	65
43.26.3.3 operator=() [1/2]	65
43.26.3.4 operator=() [2/2]	66
43.26.3.5 transpose()	66
43.27 gko::experimental::mpi::contiguous_type Class Reference	66
43.27.1 Detailed Description	67
43.27.2 Constructor & Destructor Documentation	67
43.27.2.1 contiguous_type() [1/2]	67
43.27.2.2 contiguous_type() [2/2]	67
43.27.3 Member Function Documentation	67
43.27.3.1 get()	67
43.27.3.2 operator=()	68
43.28 gko::log::Convergence < ValueType > Class Template Reference	68
43.28.1 Detailed Description	69
43.28.2 Member Function Documentation	69
43.28.2.1 create() [1/2]	69
43.28.2.2 create() [2/2]4	70
43.28.2.3 get_implicit_sq_resnorm()	70
43.28.2.4 get_num_iterations()	70
43.28.2.5 get_residual()	71
43.28.2.6 get_residual_norm()	71
43.28.2.7 has_converged()	71
43.29 gko::ConvertibleTo< ResultType > Class Template Reference	71
43.29.1 Detailed Description	72
43.29.2 Member Function Documentation	72
43.29.2.1 convert_to()	72
43.29.2.2 move_to()	73
43.30 gko::matrix::Coo< ValueType, IndexType > Class Template Reference	74
43.30.1 Detailed Description	75
43.30.2 Member Function Documentation	75
43.30.2.1 apply2() [1/4] 4	75

43.30.2.2 apply2() [2/4]
43.30.2.3 apply2() [3/4]
43.30.2.4 apply2() [4/4]
43.30.2.5 compute_absolute()
43.30.2.6 create_const()
43.30.2.7 extract_diagonal()
43.30.2.8 get_col_idxs()
43.30.2.9 get_const_col_idxs()
43.30.2.10 get_const_row_idxs()
43.30.2.11 get_const_values()
43.30.2.12 get_num_stored_elements()
43.30.2.13 get_row_idxs()
43.30.2.14 get_values()
43.30.2.15 read() [1/3]
43.30.2.16 read() [2/3]
43.30.2.17 read() [3/3]
43.30.2.18 write()
43.31 gko::CpuTimer Class Reference
43.31.1 Detailed Description
43.31.2 Member Function Documentation
43.31.2.1 difference_async()
43.32 gko::cpx_real_type< T > Struct Template Reference
43.32.1 Detailed Description
43.32.2 Member Typedef Documentation
43.32.2.1 type
43.33 gko::stop::Criterion Class Reference
43.33.1 Detailed Description
43.33.2 Member Function Documentation
43.33.2.1 check()
43.33.2.2 update()
43.34 gko::log::criterion_data Struct Reference
43.34.1 Detailed Description
43.35 gko::stop::CriterionArgs Struct Reference
43.35.1 Detailed Description
43.36 gko::matrix::Csr< ValueType, IndexType > Class Template Reference
43.36.1 Detailed Description
43.36.2 Constructor & Destructor Documentation
43.36.2.1 Csr() [1/2]
43.36.2.2 Csr() [2/2]
43.36.3 Member Function Documentation
43.36.3.1 column_permute()
43.36.3.2 compute absolute()

43.36.3.3 conj_transpose()	 490
43.36.3.4 create_const()	 491
43.36.3.5 create_submatrix() [1/2]	 491
43.36.3.6 create_submatrix() [2/2]	 492
43.36.3.7 extract_diagonal()	 492
43.36.3.8 get_col_idxs()	 493
43.36.3.9 get_const_col_idxs()	 493
43.36.3.10 get_const_row_ptrs()	 493
43.36.3.11 get_const_srow()	 494
43.36.3.12 get_const_values()	 494
43.36.3.13 get_num_srow_elements()	 494
43.36.3.14 get_num_stored_elements()	 495
43.36.3.15 get_row_ptrs()	 495
43.36.3.16 get_srow()	 495
43.36.3.17 get_strategy()	 496
43.36.3.18 get_values()	 496
43.36.3.19 inv_scale()	 496
43.36.3.20 inverse_column_permute()	 496
43.36.3.21 inverse_permute()	 497
43.36.3.22 inverse_row_permute()	 497
43.36.3.23 operator=() [1/2]	 498
43.36.3.24 operator=() [2/2]	 498
43.36.3.25 permute()	 498
43.36.3.26 read() [1/3]	 499
43.36.3.27 read() [2/3]	 499
43.36.3.28 read() [3/3]	 499
43.36.3.29 row_permute()	 500
43.36.3.30 scale()	 500
43.36.3.31 set_strategy()	 500
43.36.3.32 transpose()	 501
43.36.3.33 write()	 501
43.37 gko::CublasError Class Reference	 501
43.37.1 Detailed Description	 502
43.37.2 Constructor & Destructor Documentation	 502
43.37.2.1 CublasError()	 502
43.38 gko::cuda_stream Class Reference	 502
43.38.1 Detailed Description	 503
43.38.2 Member Function Documentation	 503
43.38.2.1 get()	 503
43.39 gko::CudaError Class Reference	 503
43.39.1 Detailed Description	 503
43.39.2 Constructor & Destructor Documentation	 504

43.39.2.1 CudaError()	04
43.40 gko::CudaExecutor Class Reference	04
43.40.1 Detailed Description	05
43.40.2 Member Function Documentation	05
43.40.2.1 create()	05
43.40.2.2 get_closest_numa()	06
43.40.2.3 get_closest_pus()	06
43.40.2.4 get_cublas_handle()	06
43.40.2.5 get_cusparse_handle()	06
43.40.2.6 get_master() [1/2] 5	07
43.40.2.7 get_master() [2/2] 5	07
43.40.2.8 get_stream()	07
43.41 gko::CudaTimer Class Reference	07
43.41.1 Detailed Description	80
43.41.2 Member Function Documentation	80
43.41.2.1 difference_async()	80
43.42 gko::CufftError Class Reference	09
43.42.1 Detailed Description	09
43.42.2 Constructor & Destructor Documentation	09
43.42.2.1 CufftError()	09
43.43 gko::CurandError Class Reference	09
43.43.1 Detailed Description	10
43.43.2 Constructor & Destructor Documentation	10
43.43.2.1 CurandError()	10
43.44 gko::matrix::Csr< ValueType, IndexType >::cusparse Class Reference	10
43.44.1 Detailed Description	11
43.44.2 Member Function Documentation	11
43.44.2.1 clac_size()	11
43.44.2.2 copy()	12
43.44.2.3 process()	
43.45 gko::CusparseError Class Reference	
43.45.1 Detailed Description	
43.45.2 Constructor & Destructor Documentation	
43.45.2.1 CusparseError()	
43.46 gko::default_converter< S, R > Struct Template Reference	
43.46.1 Detailed Description	
43.46.2 Member Function Documentation	
43.46.2.1 operator()()	
43.47 gko::matrix::Dense< ValueType > Class Template Reference	
43.47.1 Detailed Description	
43.47.2 Constructor & Destructor Documentation	
43.47.2.1 Dense() [1/2]	
10.1.1.2.1.201.00 _{(/ [1/2]}	. J

43.47.2.2 Dense() [2/2]	518
43.47.3 Member Function Documentation	518
43.47.3.1 add_scaled()	518
43.47.3.2 at() [1/4]	519
43.47.3.3 at() [2/4]	519
43.47.3.4 at() [3/4]	
43.47.3.5 at() [4/4]	
43.47.3.6 column_permute() [1/4]	521
43.47.3.7 column_permute() [2/4]	521
43.47.3.8 column_permute() [3/4]	521
43.47.3.9 column_permute() [4/4]	
43.47.3.10 compute_absolute() [1/2]	
43.47.3.11 compute_absolute() [2/2]	
43.47.3.12 compute_conj_dot() [1/2]	
43.47.3.13 compute_conj_dot() [2/2]	
43.47.3.14 compute_dot() [1/2]	
43.47.3.15 compute_dot() [2/2]	
43.47.3.16 compute_norm1() [1/2]	
43.47.3.17 compute_norm1() [2/2]	
43.47.3.18 compute_norm2() [1/2]	
43.47.3.19 compute_norm2() [2/2]	
43.47.3.20 compute_squared_norm2() [1/2]	
43.47.3.21 compute_squared_norm2() [2/2]	
43.47.3.22 conj_transpose() [1/2]	
43.47.3.23 conj_transpose() [2/2]	
43.47.3.24 create_const()	
43.47.3.25 create_const_view_of()	
43.47.3.26 create_real_view() [1/2]	
43.47.3.27 create_real_view() [2/2]	
43.47.3.28 create_submatrix() [1/2]	
43.47.3.29 create_submatrix() [2/2]	
43.47.3.30 create_view_of()	530
43.47.3.31 create_with_config_of()	530
43.47.3.32 create_with_type_of() [1/3]	530
43.47.3.33 create_with_type_of() [2/3]	531
43.47.3.34 create_with_type_of() [3/3]	531
43.47.3.35 extract_diagonal() [1/2]	532
43.47.3.36 extract_diagonal() [2/2]	532
43.47.3.37 fill()	
43.47.3.38 get_const_values()	533
43.47.3.39 get_num_stored_elements()	533
43.47.3.40 get_stride()	533

43.47.3.41 get_values()	534
43.47.3.42 inv_scale()	534
43.47.3.43 inverse_column_permute() [1/4]	534
43.47.3.44 inverse_column_permute() [2/4]	535
43.47.3.45 inverse_column_permute() [3/4]	535
43.47.3.46 inverse_column_permute() [4/4]	535
43.47.3.47 inverse_permute() [1/4]	536
43.47.3.48 inverse_permute() [2/4]	536
43.47.3.49 inverse_permute() [3/4]	536
43.47.3.50 inverse_permute() [4/4]	537
43.47.3.51 inverse_row_permute() [1/4]	537
43.47.3.52 inverse_row_permute() [2/4]	537
43.47.3.53 inverse_row_permute() [3/4]	
43.47.3.54 inverse_row_permute() [4/4]	538
43.47.3.55 make_complex() [1/2]	
43.47.3.56 make_complex() [2/2]	539
43.47.3.57 operator=() [1/2]	539
43.47.3.58 operator=() [2/2]	539
43.47.3.59 permute() [1/4]	539
43.47.3.60 permute() [2/4]	540
43.47.3.61 permute() [3/4]	
43.47.3.62 permute() [4/4]	
43.47.3.63 row_gather() [1/6]	
43.47.3.64 row_gather() [2/6]	541
43.47.3.65 row_gather() [3/6]	
43.47.3.66 row_gather() [4/6]	
43.47.3.67 row_gather() [5/6]	542
43.47.3.68 row_gather() [6/6]	543
43.47.3.69 row_permute() [1/4]	543
43.47.3.70 row_permute() [2/4]	543
43.47.3.71 row_permute() [3/4]	544
43.47.3.72 row_permute() [4/4]	544
43.47.3.73 scale()	544
43.47.3.74 sub_scaled()	
43.47.3.75 transpose() [1/2]	
43.47.3.76 transpose() [2/2]	545
43.48 gko::device_matrix_data < ValueType, IndexType > Class Template Reference	546
43.48.1 Detailed Description	
43.48.2 Constructor & Destructor Documentation	547
43.48.2 Constructor & Destructor Documentation	547 547
43.48.2 Constructor & Destructor Documentation	547 547 548

43.48.3 Member Function Documentation	49
43.48.3.1 copy_to_host()	49
43.48.3.2 create_from_host()	49
43.48.3.3 empty_out()	50
43.48.3.4 get_col_idxs()	50
43.48.3.5 get_const_col_idxs()	50
43.48.3.6 get_const_row_idxs()	51
43.48.3.7 get_const_values()	51
43.48.3.8 get_executor()	51
43.48.3.9 get_num_elems()	52
43.48.3.10 get_row_idxs()	52
43.48.3.11 get_size()	52
43.48.3.12 get_values()	52
43.48.3.13 remove_zeros()	53
43.48.3.14 resize_and_reset() [1/2]	53
43.48.3.15 resize_and_reset() [2/2]	53
43.48.3.16 sum_duplicates()	53
43.49 gko::matrix::Diagonal < ValueType > Class Template Reference	54
43.49.1 Detailed Description	54
43.49.2 Member Function Documentation	55
43.49.2.1 compute_absolute()	55
43.49.2.2 conj_transpose()	55
43.49.2.3 create_const()	55
43.49.2.4 get_const_values()	56
43.49.2.5 get_values()	56
43.49.2.6 inverse_apply()	56
43.49.2.7 rapply()	57
43.49.2.8 transpose()	57
43.50 gko::DiagonalExtractable < ValueType > Class Template Reference	58
43.50.1 Detailed Description	58
43.50.2 Member Function Documentation	58
43.50.2.1 extract_diagonal()	58
43.50.2.2 extract_diagonal_linop()	59
43.51 gko::DiagonalLinOpExtractable Class Reference	59
43.51.1 Detailed Description	59
43.51.2 Member Function Documentation	59
43.51.2.1 extract_diagonal_linop()	60
43.52 gko::dim< Dimensionality, DimensionType > Struct Template Reference	60
43.52.1 Detailed Description	60
43.52.2 Constructor & Destructor Documentation	61
43.52.2.1 dim() [1/2]	61
43.52.2.2 dim() [2/2]	61

43.52.3 Member Function Documentation
43.52.3.1 operator bool()
43.52.3.2 operator[]() [1/2]
43.52.3.3 operator[]() [2/2]
43.52.4 Friends And Related Function Documentation
43.52.4.1 operator*
43.52.4.2 operator <<
43.52.4.3 operator==
43.53 gko::DimensionMismatch Class Reference
43.53.1 Detailed Description
43.53.2 Constructor & Destructor Documentation
43.53.2.1 DimensionMismatch()
43.54 gko::experimental::solver::Direct< ValueType, IndexType > Class Template Reference
43.54.1 Detailed Description
43.54.2 Member Function Documentation
43.54.2.1 conj_transpose()
43.54.2.2 transpose()
43.55 gko::experimental::distributed::DistributedBase Class Reference
43.55.1 Detailed Description
43.55.2 Member Function Documentation
43.55.2.1 get_communicator()
43.55.2.2 operator=() [1/2]
43.55.2.3 operator=() [2/2] 568
43.56 gko::DpcppExecutor Class Reference
43.56.1 Detailed Description
43.56.2 Member Function Documentation
43.56.2.1 create()
43.56.2.2 get_device_id()
43.56.2.3 get_device_type()
43.56.2.4 get_master() [1/2]
43.56.2.5 get_master() [2/2] 570
43.56.2.6 get_max_subgroup_size()
43.56.2.7 get_max_workgroup_size()
43.56.2.8 get_max_workitem_sizes()
43.56.2.9 get_num_computing_units()
43.56.2.10 get_num_devices()
43.56.2.11 get_subgroup_sizes()
43.57 gko::DpcppTimer Class Reference
43.57.1 Detailed Description
43.57.2 Member Function Documentation
43.57.2.1 difference_async()
43.58 gko::matrix::Ell< ValueType, IndexType > Class Template Reference

45.56.1 Detailed Description	5/4
43.58.2 Constructor & Destructor Documentation	575
43.58.2.1 Ell() [1/2]	575
43.58.2.2 Ell() [2/2]	575
43.58.3 Member Function Documentation	575
43.58.3.1 col_at() [1/2]	575
43.58.3.2 col_at() [2/2]	576
43.58.3.3 compute_absolute()	576
43.58.3.4 create_const()	576
43.58.3.5 extract_diagonal()	577
43.58.3.6 get_col_idxs()	577
43.58.3.7 get_const_col_idxs()	578
43.58.3.8 get_const_values()	578
43.58.3.9 get_num_stored_elements()	578
43.58.3.10 get_num_stored_elements_per_row()	579
43.58.3.11 get_stride()	579
43.58.3.12 get_values()	579
43.58.3.13 operator=() [1/2]	579
43.58.3.14 operator=() [2/2]	580
43.58.3.15 read() [1/3]	580
43.58.3.16 read() [2/3]	580
43.58.3.17 read() [3/3]	580
43.58.3.18 val_at() [1/2]	581
43.58.3.19 val_at() [2/2]	581
43.58.3.20 write()	582
$43.59~gko::enable_parameters_type < Concrete Parameters Type,~Factory > Class~Template~Reference~~.$	582
43.59.1 Detailed Description	582
43.59.2 Member Function Documentation	583
43.59.2.1 on()	583
$43.60~gko:: Enable Absolute Computation < Absolute Lin Op > Class~Template~Reference ~~\dots ~~\dots ~~.$	583
43.60.1 Detailed Description	584
43.60.2 Member Function Documentation	584
43.60.2.1 compute_absolute()	584
43.60.2.2 compute_absolute_linop()	584
43.61 gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase > Class Template Reference	585
43.61.1 Detailed Description	585
43.62 gko::solver::EnableApplyWithInitialGuess< DerivedType > Class Template Reference	585
43.62.1 Detailed Description	585
43.63 gko::EnableCreateMethod< ConcreteType > Class Template Reference	
43.63.1 Detailed Description	
43.64 gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >	
Class Template Reference	586

43.64.1 Detailed Description	587
43.64.2 Member Function Documentation	587
43.64.2.1 create()	587
43.64.2.2 get_parameters()	588
43.65 gko::experimental::EnableDistributedLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference	588
43.65.1 Detailed Description	588
43.66 gko::experimental::EnableDistributedPolymorphicObject< ConcreteObject, PolymorphicBase > Class Template Reference	589
43.66.1 Detailed Description	
43.67 gko::solver::EnableIterativeBase< DerivedType > Class Template Reference	
43.67.1 Detailed Description	
43.67.2 Constructor & Destructor Documentation	
43.67.2.1 EnableIterativeBase()	
43.67.3 Member Function Documentation	
43.67.3.1 operator=()	
43.67.3.2 set_stop_criterion_factory()	
43.68 gko::EnableLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference	
43.68.1 Detailed Description	
43.69 gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase > Class Template Reference	
43.69.1 Detailed Description	
43.70 gko::multigrid::EnableMultigridLevel< ValueType > Class Template Reference	593
43.70.1 Detailed Description	593
43.70.2 Member Function Documentation	593
43.70.2.1 get_coarse_op()	594
43.70.2.2 get_fine_op()	594
43.70.2.3 get_prolong_op()	594
43.70.2.4 get_restrict_op()	595
43.71 gko::EnablePolymorphicAssignment< ConcreteType, ResultType > Class Template Reference	595
43.71.1 Detailed Description	595
43.71.2 Member Function Documentation	596
43.71.2.1 convert_to()	596
43.71.2.2 move_to()	596
43.72 gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase > Class Template Reference	597
43.72.1 Detailed Description	597
43.73 gko::solver::EnablePreconditionable< DerivedType > Class Template Reference	597
43.73.1 Detailed Description	598
43.73.2 Constructor & Destructor Documentation	598
43.73.2.1 EnablePreconditionable()	598
43.73.3 Member Function Documentation	598
43.73.3.1 operator=()	599
43.73.3.2 set_preconditioner()	599

43.74 gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType > Class Template Ref	-00
erence	
43.74.1 Detailed Description	
43.75.1 Detailed Description	
·	
43.75.2 Constructor & Destructor Documentation	
43.75.2.1 EnableSolverBase()	
43.75.3 Member Function Documentation	
43.75.3.1 operator=()	
43.76 gko::experimental::mpi::environment Class Reference	
43.76.1 Detailed Description	
43.76.2 Constructor & Destructor Documentation	
43.76.2.1 environment()	
43.76.3 Member Function Documentation	
43.76.3.1 get_provided_thread_support()	
43.77 gko::Error Class Reference	
43.77.1 Detailed Description	
43.77.2 Constructor & Destructor Documentation	
43.77.2.1 Error()	
43.78 gko::Executor Class Reference	
43.78.1 Detailed Description	
43.78.2 Member Function Documentation	
43.78.2.1 add_logger()	
43.78.2.2 alloc()	
43.78.2.3 copy()	
43.78.2.4 copy_from()	
43.78.2.5 copy_val_to_host()	308
43.78.2.6 free()	308
43.78.2.7 get_master() [1/2]	308
43.78.2.8 get_master() [2/2]	309
43.78.2.9 memory_accessible()	309
43.78.2.10 remove_logger()	309
43.78.2.11 run() [1/2]	310
43.78.2.12 run() [2/2]	310
43.78.2.13 set_log_propagation_mode()	310
43.78.2.14 should_propagate_log()	311
43.79 gko::log::executor_data Struct Reference	311
43.79.1 Detailed Description	311
43.80 gko::executor_deleter< T > Class Template Reference	311
43.80.1 Detailed Description	312
43.80.2 Constructor & Destructor Documentation	312
43.80.2.1 executor_deleter()	812

43.80.3 Member Function Documentation	612
43.80.3.1 operator()()	612
43.81 gko::experimental::factorization::Factorization< ValueType, IndexType > Class Template Reference	613
43.81.1 Detailed Description	614
43.81.2 Member Function Documentation	614
43.81.2.1 create_from_combined_lu()	614
43.81.2.2 create_from_composition()	614
43.81.2.3 create_from_symm_composition()	616
43.81.2.4 unpack()	616
43.82 gko::matrix::Fbcsr< ValueType, IndexType > Class Template Reference	617
43.82.1 Detailed Description	618
43.82.2 Constructor & Destructor Documentation	618
43.82.2.1 Fbcsr() [1/2]	619
43.82.2.2 Fbcsr() [2/2]	619
43.82.3 Member Function Documentation	619
43.82.3.1 compute_absolute()	619
43.82.3.2 conj_transpose()	619
43.82.3.3 convert_to() [1/2]	620
43.82.3.4 convert_to() [2/2]	620
43.82.3.5 create_const()	620
43.82.3.6 extract_diagonal()	621
43.82.3.7 get_block_size()	621
43.82.3.8 get_col_idxs()	621
43.82.3.9 get_const_col_idxs()	622
43.82.3.10 get_const_row_ptrs()	622
43.82.3.11 get_const_values()	622
43.82.3.12 get_num_block_cols()	623
43.82.3.13 get_num_block_rows()	623
43.82.3.14 get_num_stored_blocks()	623
43.82.3.15 get_num_stored_elements()	623
43.82.3.16 get_row_ptrs()	624
43.82.3.17 get_values()	624
43.82.3.18 is_sorted_by_column_index()	624
43.82.3.19 operator=() [1/2]	624
43.82.3.20 operator=() [2/2]	625
43.82.3.21 read() [1/3]	625
43.82.3.22 read() [2/3]	625
43.82.3.23 read() [3/3]	625
43.82.3.24 transpose()	626
43.82.3.25 write()	626
43.83 gko::solver::Fcg< ValueType > Class Template Reference	626
43.83.1 Detailed Description	627

43.83.2 Member Function Documentation	27
43.83.2.1 apply_uses_initial_guess()	27
43.83.2.2 conj_transpose()	28
43.83.2.3 transpose()	28
43.84 gko::matrix::Fft Class Reference	28
43.84.1 Detailed Description	29
43.84.2 Member Function Documentation	29
43.84.2.1 conj_transpose()	29
43.84.2.2 transpose()	29
43.84.2.3 write() [1/4]	29
43.84.2.4 write() [2/4]	30
43.84.2.5 write() [3/4]	30
43.84.2.6 write() [4/4]	30
43.85 gko::matrix::Fft2 Class Reference	31
43.85.1 Detailed Description	31
43.85.2 Member Function Documentation	32
43.85.2.1 conj_transpose()	32
43.85.2.2 transpose()	32
43.85.2.3 write() [1/4]	32
43.85.2.4 write() [2/4]	33
43.85.2.5 write() [3/4]	33
43.85.2.6 write() [4/4]	33
43.86 gko::matrix::Fft3 Class Reference	34
43.86.1 Detailed Description	34
43.86.2 Member Function Documentation	34
43.86.2.1 conj_transpose()	35
43.86.2.2 transpose()	35
43.86.2.3 write() [1/4]	35
43.86.2.4 write() [2/4]	35
43.86.2.5 write() [3/4]	36
43.86.2.6 write() [4/4]	36
43.87 gko::multigrid::FixedCoarsening< ValueType, IndexType > Class Template Reference 63	36
43.87.1 Detailed Description	37
43.87.2 Member Function Documentation	37
43.87.2.1 get_system_matrix()	37
43.88 gko::solver::Gcr< ValueType > Class Template Reference	38
43.88.1 Detailed Description	38
43.88.2 Member Function Documentation	38
43.88.2.1 apply_uses_initial_guess()	38
43.88.2.2 conj_transpose()	39
43.88.2.3 get_krylov_dim()	39
43.88.2.4 set_krylov_dim()	39

43.88.2.5 transpose()	640
43.89 gko::solver::Gmres < ValueType > Class Template Reference	640
43.89.1 Detailed Description	640
43.89.2 Member Function Documentation	641
43.89.2.1 apply_uses_initial_guess()	641
43.89.2.2 conj_transpose()	641
43.89.2.3 get_krylov_dim()	641
43.89.2.4 set_krylov_dim()	641
43.89.2.5 transpose()	642
43.90 gko::solver::has_with_criteria< SolverType, typename > Struct Template Reference	642
43.90.1 Detailed Description	642
43.91 gko::solver::has_with_criteria< SolverType, xstd::void_t< decltype(SolverType::build().with_ \hookleftarrow criteria(std::shared_ptr< const stop::CriterionFactory >()))> > Struct Template Reference	643
43.91.1 Detailed Description	643
43.92 gko::hip_stream Class Reference	643
43.92.1 Detailed Description	644
43.92.2 Member Function Documentation	644
43.92.2.1 get()	644
43.93 gko::HipblasError Class Reference	644
43.93.1 Detailed Description	644
43.93.2 Constructor & Destructor Documentation	644
43.93.2.1 HipblasError()	644
43.94 gko::HipError Class Reference	645
43.94.1 Detailed Description	645
43.94.2 Constructor & Destructor Documentation	645
43.94.2.1 HipError()	645
43.95 gko::HipExecutor Class Reference	646
43.95.1 Detailed Description	647
43.95.2 Member Function Documentation	647
43.95.2.1 create()	647
43.95.2.2 get_closest_numa()	647
43.95.2.3 get_closest_pus()	647
43.95.2.4 get_hipblas_handle()	648
43.95.2.5 get_hipsparse_handle()	648
43.95.2.6 get_master() [1/2]	648
43.95.2.7 get_master() [2/2]	648
43.96 gko::HipfftError Class Reference	649
43.96.1 Detailed Description	649
43.96.2 Constructor & Destructor Documentation	649
43.96.2.1 HipfftError()	649
43.97 gko::HiprandError Class Reference	649
43.97.1 Detailed Description	650

43.97.2 Constructor & Destructor Documentation	650
43.97.2.1 HiprandError()	650
43.98 gko::HipsparseError Class Reference	650
43.98.1 Detailed Description	651
43.98.2 Constructor & Destructor Documentation	651
43.98.2.1 HipsparseError()	651
43.99 gko::HipTimer Class Reference	651
43.99.1 Detailed Description	652
43.99.2 Member Function Documentation	652
43.99.2.1 difference_async()	652
43.100 gko::matrix::Hybrid< ValueType, IndexType > Class Template Reference	652
43.100.1 Detailed Description	654
43.100.2 Constructor & Destructor Documentation	655
43.100.2.1 Hybrid() [1/2]	655
43.100.2.2 Hybrid() [2/2]	655
43.100.3 Member Function Documentation	655
43.100.3.1 compute_absolute()	655
43.100.3.2 ell_col_at() [1/2]	655
43.100.3.3 ell_col_at() [2/2]	656
43.100.3.4 ell_val_at() [1/2]	656
43.100.3.5 ell_val_at() [2/2]	657
43.100.3.6 extract_diagonal()	657
43.100.3.7 get_const_coo_col_idxs()	657
43.100.3.8 get_const_coo_row_idxs()	658
43.100.3.9 get_const_coo_values()	658
43.100.3.10 get_const_ell_col_idxs()	659
43.100.3.11 get_const_ell_values()	659
43.100.3.12 get_coo()	659
43.100.3.13 get_coo_col_idxs()	660
43.100.3.14 get_coo_num_stored_elements()	660
43.100.3.15 get_coo_row_idxs()	660
43.100.3.16 get_coo_values()	660
43.100.3.17 get_ell()	661
43.100.3.18 get_ell_col_idxs()	661
43.100.3.19 get_ell_num_stored_elements()	661
43.100.3.20 get_ell_num_stored_elements_per_row()	661
43.100.3.21 get_ell_stride()	662
43.100.3.22 get_ell_values()	662
43.100.3.23 get_num_stored_elements()	662
43.100.3.24 get_strategy() [1/2]	662
43.100.3.25 get_strategy() [2/2]	663
43.100.3.26 operator=() [1/2]	663

43.100.3.27 operator=() [2/2]	63
43.100.3.28 read() [1/3]	63
43.100.3.29 read() [2/3]	64
43.100.3.30 read() [3/3]	64
43.100.3.31 write()	64
43.101 gko::factorization::lc< ValueType, IndexType > Class Template Reference	65
43.101.1 Detailed Description	65
43.102 gko::preconditioner::lc< LSolverType, IndexType > Class Template Reference	65
43.102.1 Detailed Description	66
43.102.2 Constructor & Destructor Documentation	67
43.102.2.1 lc() [1/2]	67
43.102.2.2 lc() [2/2]	67
43.102.3 Member Function Documentation	67
43.102.3.1 conj_transpose()	68
43.102.3.2 get_l_solver()	68
43.102.3.3 get_lh_solver()	68
43.102.3.4 operator=() [1/2]	69
43.102.3.5 operator=() [2/2]	69
43.102.3.6 transpose()	69
43.103 gko::matrix::Identity < ValueType > Class Template Reference 6	70
43.103.1 Detailed Description	70
43.103.2 Member Function Documentation	70
43.103.2.1 conj_transpose()	70
43.103.2.2 transpose()	71
43.104 gko::matrix::IdentityFactory< ValueType > Class Template Reference 6	71
43.104.1 Detailed Description	71
43.104.2 Member Function Documentation	72
43.104.2.1 create()	72
43.105 gko::solver::ldr< ValueType > Class Template Reference 6	72
43.105.1 Detailed Description	73
43.105.2 Member Function Documentation	
43.105.2.1 apply_uses_initial_guess()	73
43.105.2.2 conj_transpose()	74
43.105.2.3 get_complex_subspace()	74
43.105.2.4 get_deterministic()	
43.105.2.5 get_kappa()	
43.105.2.6 get_subspace_dim()	
43.105.2.7 set_complex_subpsace()	75
43.105.2.8 set_deterministic()	
43.105.2.9 set_kappa()	
43.105.2.10 set_subspace_dim()	76
43.105.2.11 transpose()	76

43.106~gko:: factorization:: Ilu < Value Type, ~Index Type > Class~Template~Reference~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.	676
43.106.1 Detailed Description	677
43.107 gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType > Class Template Reference	
43.107.1 Detailed Description	678
43.107.2 Constructor & Destructor Documentation	679
43.107.2.1 llu() [1/2]	679
43.107.2.2 llu() [2/2]	680
43.107.3 Member Function Documentation	680
43.107.3.1 conj_transpose()	680
43.107.3.2 get_l_solver()	680
43.107.3.3 get_u_solver()	681
43.107.3.4 operator=() [1/2]	681
43.107.3.5 operator=() [2/2]	681
43.107.3.6 transpose()	
43.108 gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit Class Reference	
43.108.1 Detailed Description	682
43.108.2 Member Function Documentation	
43.108.2.1 compute_ell_num_stored_elements_per_row()	683
43.108.2.2 get_percentage()	683
43.108.2.3 get_ratio()	683
43.109 gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit Class Reference	
43.109.1 Detailed Description	684
43.109.2 Constructor & Destructor Documentation	684
43.109.2.1 imbalance_limit()	684
43.109.3 Member Function Documentation	685
43.109.3.1 compute_ell_num_stored_elements_per_row()	685
43.109.3.2 get_percentage()	685
43.110 gko::stop::ImplicitResidualNorm< ValueType > Class Template Reference	686
43.110.1 Detailed Description	686
43.111 gko::index_set< IndexType > Class Template Reference	686
43.111.1 Detailed Description	688
43.111.2 Constructor & Destructor Documentation	688
43.111.2.1 index_set() [1/7]	688
43.111.2.2 index_set() [2/7]	689
43.111.2.3 index_set() [3/7]	689
43.111.2.4 index_set() [4/7]	690
43.111.2.5 index_set() [5/7]	690
43.111.2.6 index_set() [6/7]	690
43.111.2.7 index_set() [7/7]	690
43.111.3 Member Function Documentation	691
43.111.3.1 clear()	691

43.111.3.2 contains() [1/2]	1
43.111.3.3 contains() [2/2]	2
43.111.3.4 get_executor()	2
43.111.3.5 get_global_index()	2
43.111.3.6 get_local_index()	3
43.111.3.7 get_num_elems()	3
43.111.3.8 get_num_subsets()	4
43.111.3.9 get_size()	4
43.111.3.10 get_subsets_begin()	4
43.111.3.11 get_subsets_end()	5
43.111.3.12 get_superset_indices()	5
43.111.3.13 is_contiguous()	5
43.111.3.14 map_global_to_local()	5
43.111.3.15 map_local_to_global()	6
43.111.3.16 operator=() [1/2]	6
43.111.3.17 operator=() [2/2]	7
43.111.3.18 to_global_indices()	7
43.112 gko::solver::Ir < ValueType > Class Template Reference	8
43.112.1 Detailed Description	8
43.112.2 Constructor & Destructor Documentation	9
43.112.2.1 lr() [1/2]	9
43.112.2.2 lr() [2/2]	9
43.112.3 Member Function Documentation	9
43.112.3.1 apply_uses_initial_guess()	0
43.112.3.2 conj_transpose()	0
43.112.3.3 get_solver()	0
43.112.3.4 operator=() [1/2]	0
43.112.3.5 operator=() [2/2]	1
43.112.3.6 set_solver()	1
43.112.3.7 transpose()	1
43.113 gko::preconditioner::lsai< lsaiType, ValueType, IndexType > Class Template Reference 70	1
43.113.1 Detailed Description	2
43.113.2 Constructor & Destructor Documentation	3
43.113.2.1 Isai() [1/2]	3
43.113.2.2 lsai() [2/2]	3
43.113.3 Member Function Documentation	3
43.113.3.1 conj_transpose()	3
43.113.3.2 get_approximate_inverse()	4
43.113.3.3 operator=() [1/2]	4
43.113.3.4 operator=() [2/2]	4
43.113.3.5 transpose()	4
43.114 gko::stop::/teration Class Reference	15

43.114.1 Detailed Description
43.115 gko::log::iteration_complete_data Struct Reference
43.115.1 Detailed Description
43.116 gko::solver::IterativeBase Class Reference
43.116.1 Detailed Description
43.116.2 Member Function Documentation
43.116.2.1 get_stop_criterion_factory()
43.116.2.2 set_stop_criterion_factory()
43.117 gko::preconditioner::Jacobi < ValueType, IndexType > Class Template Reference
43.117.1 Detailed Description
43.117.2 Constructor & Destructor Documentation
43.117.2.1 Jacobi() [1/2]
43.117.2.2 Jacobi() [2/2]
43.117.3 Member Function Documentation
43.117.3.1 conj_transpose()
43.117.3.2 convert_to()
43.117.3.3 get_blocks()
43.117.3.4 get_conditioning()
43.117.3.5 get_num_blocks()
43.117.3.6 get_num_stored_elements()
43.117.3.7 get_storage_scheme()
43.117.3.8 move_to()
43.117.3.9 operator=() [1/2]
43.117.3.10 operator=() [2/2]
43.117.3.11 transpose()
43.117.3.12 write()
43.118 gko::KernelNotFound Class Reference
43.118.1 Detailed Description
43.118.2 Constructor & Destructor Documentation
43.118.2.1 KernelNotFound()
43.119 gko::liog::linop_data Struct Reference
43.119.1 Detailed Description
43.120 gko::liog::linop_factory_data Struct Reference
43.120.1 Detailed Description
43.121 gko::LinOpFactory Class Reference
43.121.1 Detailed Description
43.121.1.1 Example: using CG in Ginkgo
43.122 gko::matrix::Csr< ValueType, IndexType >::load_balance Class Reference
43.122.1 Detailed Description
43.122.2 Constructor & Destructor Documentation
43.122.2.1 load_balance() [1/5]
43.122.2.2 load_balance() [2/5]

43.122.2.3 load_balance() [3/5]
43.122.2.4 load_balance() [4/5]
43.122.2.5 load_balance() [5/5]
43.122.3 Member Function Documentation
43.122.3.1 clac_size()
43.122.3.2 copy()
43.122.3.3 process()
43.123 gko::log::Loggable Class Reference
43.123.1 Detailed Description
43.123.2 Member Function Documentation
43.123.2.1 add_logger()
43.123.2.2 get_loggers()
43.123.2.3 remove_logger()
43.124 gko::log::Record::logged_data Struct Reference
43.124.1 Detailed Description
43.125 gko::solver::LowerTrs< ValueType, IndexType > Class Template Reference
43.125.1 Detailed Description
43.125.2 Constructor & Destructor Documentation
43.125.2.1 LowerTrs() [1/2]
43.125.2.2 LowerTrs() [2/2]
43.125.3 Member Function Documentation
43.125.3.1 conj_transpose()
43.125.3.2 operator=() [1/2]
43.125.3.3 operator=() [2/2]
43.125.3.4 transpose()
43.126 gko::experimental::factorization::Lu< ValueType, IndexType > Class Template Reference 72
43.126.1 Detailed Description
43.126.2 Member Function Documentation
43.126.2.1 generate()
43.127 gko::machine_topology Class Reference
43.127.1 Detailed Description
43.127.2 Member Function Documentation
43.127.2.1 bind_to_core()
43.127.2.2 bind_to_cores()
43.127.2.3 bind_to_pu()
43.127.2.4 bind_to_pus()
43.127.2.5 get_core()
43.127.2.6 get_instance()
43.127.2.7 get_num_cores()
43.127.2.8 get_num_numas()
43.127.2.9 get_num_pci_devices()
43.127.2.10 get_num_pus()

43.127.2.11 get_pci_device() [1/2]	728
43.127.2.12 get_pci_device() [2/2]	728
43.127.2.13 get_pu()	729
43.128 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType > Class	
Template Reference	
43.128.1 Detailed Description	
43.128.2 Constructor & Destructor Documentation	
43.128.2.1 Matrix() [1/2]	
43.128.2.2 Matrix() [2/2]	
43.128.3 Member Function Documentation	
43.128.3.1 get_local_matrix()	732
43.128.3.2 get_non_local_matrix()	732
43.128.3.3 operator=() [1/2]	733
43.128.3.4 operator=() [2/2]	733
43.128.3.5 read_distributed() [1/4]	733
43.128.3.6 read_distributed() [2/4]	734
43.128.3.7 read_distributed() [3/4]	734
43.128.3.8 read_distributed() [4/4]	735
$43.129~gko::matrix_assembly_data < ValueType,~IndexType > Class~Template~Reference~~.~~.~.~.~.$	735
43.129.1 Detailed Description	736
43.129.2 Member Function Documentation	736
43.129.2.1 add_value()	736
43.129.2.2 contains()	736
43.129.2.3 get_num_stored_elements()	737
43.129.2.4 get_ordered_data()	737
43.129.2.5 get_size()	737
43.129.2.6 get_value()	738
43.129.2.7 set_value()	738
43.130 gko::matrix_data< ValueType, IndexType > Struct Template Reference	738
43.130.1 Detailed Description	740
43.130.2 Constructor & Destructor Documentation	740
43.130.2.1 matrix_data() [1/6]	740
43.130.2.2 matrix_data() [2/6]	741
43.130.2.3 matrix_data() [3/6]	741
43.130.2.4 matrix_data() [4/6]	741
43.130.2.5 matrix_data() [5/6]	742
43.130.2.6 matrix_data() [6/6]	742
43.130.3 Member Function Documentation	743
43.130.3.1 cond() [1/2]	743
43.130.3.2 cond() [2/2]	743
43.130.3.3 diag() [1/5]	744
43.130.3.4 diag() [2/5]	

43.130.3.5 diag() [3/5]	745
43.130.3.6 diag() [4/5]	746
43.130.3.7 diag() [5/5]	746
43.130.3.8 sum_duplicates()	747
43.130.4 Member Data Documentation	747
43.130.4.1 nonzeros	747
43.131 gko::matrix_data_entry< ValueType, IndexType > Struct Template Reference	747
43.131.1 Detailed Description	748
43.132 gko::matrix::Csr< ValueType, IndexType >::merge_path Class Reference	748
43.132.1 Detailed Description	748
43.132.2 Member Function Documentation	748
43.132.2.1 clac_size()	748
43.132.2.2 copy()	749
43.132.2.3 process()	749
43.133 gko::MetisError Class Reference	750
43.133.1 Detailed Description	750
43.133.2 Constructor & Destructor Documentation	750
43.133.2.1 MetisError()	750
43.134 gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit Class Reference	750
43.134.1 Detailed Description	751
43.134.2 Member Function Documentation	751
43.134.2.1 compute_ell_num_stored_elements_per_row()	751
43.134.2.2 get_percentage()	752
43.135 gko::MpiError Class Reference	752
43.135.1 Detailed Description	752
43.135.2 Constructor & Destructor Documentation	752
43.135.2.1 MpiError()	752
43.136 gko::solver::Multigrid Class Reference	753
43.136.1 Detailed Description	753
43.136.2 Member Function Documentation	754
43.136.2.1 apply_uses_initial_guess()	754
43.136.2.2 get_coarsest_solver()	754
43.136.2.3 get_cycle()	755
43.136.2.4 get_mg_level_list()	755
43.136.2.5 get_mid_smoother_list()	755
43.136.2.6 get_post_smoother_list()	755
43.136.2.7 get_pre_smoother_list()	756
43.136.2.8 set_cycle()	756
43.137 gko::multigrid::MultigridLevel Class Reference	756
43.137.1 Detailed Description	757
43.137.2 Member Function Documentation	757
	757

43.137.2.2 get_fine_op()	757
43.137.2.3 get_prolong_op()	757
43.137.2.4 get_restrict_op()	758
43.138 gko::log::ProfilerHook::NestedSummaryWriter Class Reference	758
43.138.1 Detailed Description	758
43.138.2 Member Function Documentation	758
43.138.2.1 write_nested()	758
43.139 gko::NotCompiled Class Reference	759
43.139.1 Detailed Description	759
43.139.2 Constructor & Destructor Documentation	759
43.139.2.1 NotCompiled()	759
43.140 gko::NotImplemented Class Reference	760
43.140.1 Detailed Description	760
43.140.2 Constructor & Destructor Documentation	760
43.140.2.1 NotImplemented()	760
43.141 gko::NotSupported Class Reference	761
43.141.1 Detailed Description	761
43.141.2 Constructor & Destructor Documentation	761
43.141.2.1 NotSupported()	761
43.142 gko::null_deleter< T > Class Template Reference	761
43.142.1 Detailed Description	762
43.142.2 Member Function Documentation	762
43.142.2.1 operator()()	762
43.143 gko::nvidia_device Class Reference	762
43.143.1 Detailed Description	762
43.144 gko::OmpExecutor Class Reference	763
43.144.1 Detailed Description	763
43.144.2 Member Function Documentation	763
43.144.2.1 get_master() [1/2]	763
43.144.2.2 get_master() [2/2]	764
43.145 gko::Operation Class Reference	764
43.145.1 Detailed Description	764
43.145.2 Member Function Documentation	765
43.145.2.1 get_name()	765
43.146 gko::log::operation_data Struct Reference	766
43.146.1 Detailed Description	766
43.147 gko::OutOfBoundsError Class Reference	766
43.147.1 Detailed Description	766
43.147.2 Constructor & Destructor Documentation	766
43.147.2.1 OutOfBoundsError()	766
43.148 gko::OverflowError Class Reference	767
43 148 1 Detailed Description	767

43.148.2 Constructor & Destructor Documentation	767
43.148.2.1 OverflowError()	767
43.149 gko::log::Papi< ValueType > Class Template Reference	768
43.149.1 Detailed Description	768
43.149.2 Member Function Documentation	769
43.149.2.1 create() [1/2]	769
43.149.2.2 create() [2/2]	769
43.149.2.3 get_handle_name()	769
43.150 gko::factorization::ParIc< ValueType, IndexType > Class Template Reference	770
43.150.1 Detailed Description	770
43.151 gko::factorization::ParIct< ValueType, IndexType $>$ Class Template Reference	771
43.151.1 Detailed Description	771
43.152 gko::factorization::Parllu< ValueType, IndexType > Class Template Reference	772
43.152.1 Detailed Description	772
43.153~gko:: factorization:: Parllut < Value Type, Index Type > Class~Template~Reference~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.~.	773
43.153.1 Detailed Description	773
43.154 gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType > Class Template	
Reference	
43.154.1 Detailed Description	
43.154.2 Member Function Documentation	
43.154.2.1 build_from_contiguous()	
43.154.2.2 build_from_global_size_uniform()	
43.154.2.3 build_from_mapping()	
43.154.2.4 get_num_empty_parts()	
43.154.2.5 get_num_parts()	
43.154.2.6 get_num_ranges()	
43.154.2.7 get_part_ids()	
43.154.2.8 get_part_size()	
43.154.2.9 get_part_sizes()	
43.154.2.10 get_range_bounds()	
43.154.2.11 get_range_starting_indices()	
43.154.2.12 get_size()	
43.154.2.13 has_connected_parts()	
43.154.2.14 has_ordered_parts()	
43.155 gko::log::PerformanceHint Class Reference	
43.155.1 Detailed Description	
43.155.2 Member Function Documentation	
43.155.2.1 create()	
43.156 gko::Permutable < IndexType > Class Template Reference	
43.156.1 Detailed Description	
43.156.1.1 Example: Permuting a Csr matrix:	
43.156.2 Member Function Documentation	783

43.156.2.1 column_permute()	783
43.156.2.2 inverse_column_permute()	784
43.156.2.3 inverse_permute()	784
43.156.2.4 inverse_row_permute()	784
43.156.2.5 permute()	786
43.156.2.6 row_permute()	786
43.157 gko::matrix::Permutation < IndexType > Class Template Reference	787
43.157.1 Detailed Description	787
43.157.2 Member Function Documentation	788
43.157.2.1 create_const()	788
43.157.2.2 get_const_permutation()	788
43.157.2.3 get_permutation()	789
43.157.2.4 get_permutation_size()	789
43.157.2.5 get_permute_mask()	789
43.157.2.6 set_permute_mask()	789
43.158 gko::Perturbation < ValueType > Class Template Reference	790
43.158.1 Detailed Description	790
43.158.2 Member Function Documentation	791
43.158.2.1 get_basis()	791
43.158.2.2 get_projector()	791
43.158.2.3 get_scalar()	791
10.1.00.2.0 goz_00.0.()	
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792 792
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792 792 793
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792 792 793 793
43.159 gko::multigrid::Pgm ValueType, IndexType > Class Template Reference 7 43.159.1 Detailed Description 7 43.159.2 Member Function Documentation 7 43.159.2.1 get_agg() 7 43.159.2.2 get_const_agg() 7 43.159.2.3 get_system_matrix() 7	792 792 793 793 794
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description	792 792 793 793 794 794
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792 793 793 794 794
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference	792 792 793 793 794 794 795
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description	792 792 793 793 794 794 795
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation	792 792 793 793 794 794 795 795
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation 43.161.2.1 clear()	792 792 793 793 794 794 795 795 795
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation 43.161.2.1 clear() 43.161.2.2 clone() [1/2]	792 792 793 793 794 794 795 795 795
43.159 gko::multigrid::Pgm ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 7 43.159.2 Member Function Documentation 7 43.159.2.1 get_agg() 7 43.159.2.2 get_const_agg() 7 43.160 gko::log::polymorphic_object_data Struct Reference 7 43.160.1 Detailed Description 7 43.161.1 Detailed Description 7 43.161.2 Member Function Documentation 7 43.161.2.1 clear() 7 43.161.2.2 clone() [1/2] 43.161.2.3 clone() [2/2]	792 792 793 793 794 794 795 795 796 796
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation 43.161.2.1 clear() 43.161.2.2 clone() [1/2] 43.161.2.3 clone() [2/2] 43.161.2.4 copy_from() [1/4] 43.161.2.5 copy_from() [2/4]	792 792 793 793 794 794 795 795 796 796
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation 43.161.2.1 clear() 43.161.2.2 clone() [1/2] 43.161.2.3 clone() [2/2] 43.161.2.4 copy_from() [1/4] 43.161.2.5 copy_from() [2/4]	792 792 793 793 794 794 795 795 796 796 797
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	792 792 793 793 794 794 795 795 796 796 797 797
43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference 43.159.1 Detailed Description 43.159.2 Member Function Documentation 43.159.2.1 get_agg() 43.159.2.2 get_const_agg() 43.159.2.3 get_system_matrix() 43.160 gko::log::polymorphic_object_data Struct Reference 43.160.1 Detailed Description 43.161 gko::PolymorphicObject Class Reference 43.161.1 Detailed Description 43.161.2 Member Function Documentation 43.161.2.1 clear() 43.161.2.2 clone() [1/2] 43.161.2.3 clone() [2/2] 43.161.2.5 copy_from() [1/4] 43.161.2.5 copy_from() [3/4] 43.161.2.6 copy_from() [3/4] 43.161.2.7 copy_from() [4/4]	792 792 793 793 794 794 795 795 796 796 797 797 798

43.161.2.11 move_from()	0
43.162 gko::precision_reduction Class Reference	0
43.162.1 Detailed Description	1
43.162.2 Constructor & Destructor Documentation	2
43.162.2.1 precision_reduction() [1/2]	2
43.162.2.2 precision_reduction() [2/2]	2
43.162.3 Member Function Documentation	2
43.162.3.1 autodetect()	2
43.162.3.2 common()	3
43.162.3.3 get_nonpreserving()	3
43.162.3.4 get_preserving()	3
43.162.3.5 operator storage_type()	4
43.163 gko::Preconditionable Class Reference	
43.163.1 Detailed Description	4
43.163.2 Member Function Documentation	4
43.163.2.1 get_preconditioner()	4
43.163.2.2 set_preconditioner()	4
43.164 gko::log::ProfilerHook Class Reference	5
43.164.1 Detailed Description	6
43.164.2 Member Function Documentation	6
43.164.2.1 create_nested_summary()	6
43.164.2.2 create_nvtx()	7
43.164.2.3 create_summary()	7
43.164.2.4 create_tau()	8
43.164.2.5 set_object_name()	8
43.164.2.6 user_range()	8
43.165 gko::log::profiling_scope_guard Class Reference	9
43.165.1 Detailed Description	9
43.165.2 Constructor & Destructor Documentation	9
43.165.2.1 profiling_scope_guard()	9
43.166 gko::ptr_param< T > Class Template Reference	0
43.166.1 Detailed Description	0
43.166.2 Member Function Documentation	0
43.166.2.1 get()	0
43.166.2.2 operator bool()	1
43.166.2.3 operator*()	
43.166.2.4 operator->()	1
43.167 gko::syn::range < Start, End, Step > Struct Template Reference	
43.167.1 Detailed Description	
43.168 gko::range< Accessor > Class Template Reference	
43.168.1 Detailed Description	3
43.168.1.1 Range operations	3

43.168.1.2 Compound operations	814
43.168.1.3 Caveats	814
43.168.1.4 Examples	814
43.168.2 Constructor & Destructor Documentation	814
43.168.2.1 range()	814
43.168.3 Member Function Documentation	815
43.168.3.1 get_accessor()	815
43.168.3.2 length()	815
43.168.3.3 operator()()	816
43.168.3.4 operator->()	816
43.168.3.5 operator=() [1/2]	816
43.168.3.6 operator=() [2/2]	817
43.169 gko::reorder::Rcm< ValueType, IndexType $>$ Class Template Reference	817
43.169.1 Detailed Description	817
43.169.2 Member Function Documentation	818
43.169.2.1 get_inverse_permutation()	818
43.169.2.2 get_permutation()	818
43.170~gko:: Readable From Matrix Data < Value Type,~Index Type > Class~Template~Reference~.~.~.~.~.	819
43.170.1 Detailed Description	819
43.170.2 Member Function Documentation	819
43.170.2.1 read() [1/4]	819
43.170.2.2 read() [2/4]	820
43.170.2.3 read() [3/4]	820
43.170.2.4 read() [4/4]	820
43.171 gko::log::Record Class Reference	821
43.171.1 Detailed Description	821
43.171.2 Member Function Documentation	821
43.171.2.1 create() [1/2]	821
43.171.2.2 create() [2/2]	822
43.171.2.3 get() [1/2]	822
43.171.2.4 get() [2/2]	823
43.172 gko::ReferenceExecutor Class Reference	823
43.172.1 Detailed Description	823
43.172.2 Member Function Documentation	823
43.172.2.1 run()	823
$43.173~gko::stop::Relative Residual Norm < Value Type > Class~Template~Reference ~~\dots ~~\dots ~~.$	824
43.173.1 Detailed Description	824
43.174 gko::reorder::ReorderingBase< IndexType > Class Template Reference	824
43.174.1 Detailed Description	824
43.175 gko::reorder::ReorderingBaseArgs Struct Reference	825
43.175.1 Detailed Description	825
43 176 ako: experimental: mpi: request Class Reference	825

43.176.1 Detailed Description	325
43.176.2 Constructor & Destructor Documentation	325
43.176.2.1 request()	325
43.176.3 Member Function Documentation	326
43.176.3.1 get()	326
43.176.3.2 wait()	326
43.177 gko::stop::ResidualNorm< ValueType > Class Template Reference	326
43.177.1 Detailed Description	327
43.178 gko::stop::ResidualNormBase< ValueType > Class Template Reference	327
43.178.1 Detailed Description	327
43.179 gko::stop::ResidualNormReduction< ValueType > Class Template Reference	328
43.179.1 Detailed Description	328
43.180 gko::accessor::row_major< ValueType, Dimensionality > Class Template Reference 8	328
43.180.1 Detailed Description	329
43.180.2 Member Function Documentation	330
43.180.2.1 copy_from()	330
43.180.2.2 length()	330
43.180.2.3 operator()() [1/2]	331
43.180.2.4 operator()() [2/2]	331
43.181 gko::matrix::RowGatherer< IndexType > Class Template Reference	332
43.181.1 Detailed Description	332
43.181.2 Member Function Documentation	333
43.181.2.1 create_const()	333
43.181.2.2 get_const_row_idxs()	333
43.181.2.3 get_row_idxs()	334
43.182 gko::ScaledIdentityAddable Class Reference	334
43.182.1 Detailed Description	334
43.182.2 Member Function Documentation	334
43.182.2.1 add_scaled_identity()	334
43.183 gko::experimental::reorder::ScaledReordered< ValueType, IndexType > Class Template Reference8	335
43.183.1 Detailed Description	335
43.184 gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, Global ←	
IndexType > Class Template Reference	
43.184.1 Detailed Description	
43.185 gko::scoped_device_id_guard Class Reference	
43.185.1 Detailed Description	
43.185.2 Constructor & Destructor Documentation	
43.185.2.1 scoped_device_id_guard() [1/5]	
43.185.2.2 scoped_device_id_guard() [2/5]	
43.185.2.3 scoped_device_id_guard() [3/5]	
43.185.2.4 scoped_device_id_guard() [4/5]	
43.185.2.5 scoped device id quard() [5/5]	339

43.186 gko::matrix::Sellp< ValueType, IndexType > Class Template Reference	39
43.186.1 Detailed Description	40
43.186.2 Constructor & Destructor Documentation	41
43.186.2.1 Sellp() [1/2]	41
43.186.2.2 Sellp() [2/2]	41
43.186.3 Member Function Documentation	41
43.186.3.1 col_at() [1/2]	41
43.186.3.2 col_at() [2/2]	42
43.186.3.3 compute_absolute()	42
43.186.3.4 extract_diagonal()	42
43.186.3.5 get_col_idxs()	43
43.186.3.6 get_const_col_idxs()	43
43.186.3.7 get_const_slice_lengths()	44
43.186.3.8 get_const_slice_sets()	44
43.186.3.9 get_const_values()	44
43.186.3.10 get_num_stored_elements()	45
43.186.3.11 get_slice_lengths()	45
43.186.3.12 get_slice_sets()	45
43.186.3.13 get_slice_size()	46
43.186.3.14 get_stride_factor()	46
43.186.3.15 get_total_cols()	46
43.186.3.16 get_values()	46
43.186.3.17 operator=() [1/2]	47
43.186.3.18 operator=() [2/2]	47
43.186.3.19 read() [1/3]	47
43.186.3.20 read() [2/3]	47
43.186.3.21 read() [3/3]	48
43.186.3.22 val_at() [1/2] 8	48
43.186.3.23 val_at() [2/2] 8	49
43.186.3.24 write()	49
43.187 gko::span Struct Reference	49
43.187.1 Detailed Description	50
43.187.2 Constructor & Destructor Documentation	50
43.187.2.1 span() [1/2]	50
43.187.2.2 span() [2/2]	51
43.187.3 Member Function Documentation	51
43.187.3.1 is_valid()	51
43.187.3.2 length()	51
43.188 gko::matrix::Csr< ValueType, IndexType >::sparselib Class Reference	52
43.188.1 Detailed Description	52
43.188.2 Member Function Documentation	52
43.188.2.1 clac_size()	52

43.188.2.2 copy()	353
43.188.2.3 process()	353
43.189 gko::matrix::SparsityCsr< ValueType, IndexType > Class Template Reference	353
43.189.1 Detailed Description	
43.189.2 Constructor & Destructor Documentation	355
43.189.2.1 SparsityCsr() [1/2]	355
43.189.2.2 SparsityCsr() [2/2]	355
43.189.3 Member Function Documentation	355
43.189.3.1 conj_transpose()	356
43.189.3.2 create_const()	356
43.189.3.3 get_col_idxs()	357
43.189.3.4 get_const_col_idxs()	357
43.189.3.5 get_const_row_ptrs()	357
43.189.3.6 get_const_value()	358
43.189.3.7 get_num_nonzeros()	358
43.189.3.8 get_row_ptrs()	358
43.189.3.9 get_value()	359
43.189.3.10 operator=() [1/2]	359
43.189.3.11 operator=() [2/2]	359
43.189.3.12 read() [1/3]	359
43.189.3.13 read() [2/3]	360
43.189.3.14 read() [3/3]	360
43.189.3.15 to_adjacency_matrix()	360
43.189.3.16 transpose()	361
43.189.3.17 write()	361
43.190 gko::experimental::mpi::status Struct Reference	361
43.190.1 Detailed Description	862
43.190.2 Constructor & Destructor Documentation	362
43.190.2.1 status()	362
43.190.3 Member Function Documentation	362
43.190.3.1 get()	362
43.190.3.2 get_count()	362
43.191 gko::stopping_status Class Reference	363
43.191.1 Detailed Description	363
43.191.2 Member Function Documentation	364
43.191.2.1 converge()	364
43.191.2.2 get_id()	364
43.191.2.3 has_converged()	364
43.191.2.4 has_stopped()	365
43.191.2.5 is_finalized()	365
43.191.2.6 stop()	365
43.191.3 Friends And Related Function Documentation	865

43.191.3.1 operator"!=
43.191.3.2 operator==
43.192 gko::matrix::Hybrid< ValueType, IndexType >::strategy_type Class Reference
43.192.1 Detailed Description
43.192.2 Member Function Documentation
43.192.2.1 compute_ell_num_stored_elements_per_row()
43.192.2.2 compute_hybrid_config()
43.192.2.3 get_coo_nnz()
43.192.2.4 get_ell_num_stored_elements_per_row()
43.193 gko::matrix::Csr< ValueType, IndexType >::strategy_type Class Reference
43.193.1 Detailed Description
43.193.2 Constructor & Destructor Documentation
43.193.2.1 strategy_type()
43.193.3 Member Function Documentation
43.193.3.1 clac_size()
43.193.3.2 copy()
43.193.3.3 get_name()
43.193.3.4 process()
43.194 gko::log::Stream< ValueType > Class Template Reference
43.194.1 Detailed Description
43.194.2 Member Function Documentation
43.194.2.1 create() [1/2]
43.194.2.2 create() [2/2]
43.195 gko::StreamError Class Reference
43.195.1 Detailed Description
43.195.2 Constructor & Destructor Documentation
43.195.2.1 StreamError()
43.196 gko::log::ProfilerHook::SummaryWriter Class Reference
43.196.1 Detailed Description
43.196.2 Member Function Documentation
43.196.2.1 write()
43.197 gko::log::ProfilerHook::TableSummaryWriter Class Reference
43.197.1 Detailed Description
43.197.2 Constructor & Destructor Documentation
43.197.2.1 TableSummaryWriter()
43.197.3 Member Function Documentation
43.197.3.1 write()
43.197.3.2 write_nested()
43.198 gko::stop::Time Class Reference
43.198.1 Detailed Description
43.199 gko::time_point Class Reference
43.199.1 Detailed Description

43.200 gko::Timer Class Reference
43.200.1 Detailed Description
43.200.2 Member Function Documentation
43.200.2.1 create_for_executor()
43.200.2.2 create_time_point()
43.200.2.3 difference()
43.200.2.4 difference_async()
43.201 gko::Transposable Class Reference
43.201.1 Detailed Description
43.201.1.1 Example: Transposing a Csr matrix:
43.201.2 Member Function Documentation
43.201.2.1 conj_transpose()
43.201.2.2 transpose()
$43.202~gko::experimental::mpi::type_impl < T > Struct~Template~Reference~.~.~.~.~.~.~.~.~.~88 = 1.000~MeV~MeV~MeV~MeV~MeV~MeV~MeV~MeV~MeV~MeV$
43.202.1 Detailed Description
43.203 gko::syn::type_list< Types > Struct Template Reference
43.203.1 Detailed Description
43.204 gko::UnsupportedMatrixProperty Class Reference
43.204.1 Detailed Description
43.204.2 Constructor & Destructor Documentation
43.204.2.1 UnsupportedMatrixProperty()
43.205 gko::stop::Criterion::Updater Class Reference
43.205.1 Detailed Description
43.205.2 Member Function Documentation
43.205.2.1 check()
43.206 gko::solver::UpperTrs< ValueType, IndexType > Class Template Reference
43.206.1 Detailed Description
43.206.2 Constructor & Destructor Documentation
43.206.2.1 UpperTrs() [1/2]
43.206.2.2 UpperTrs() [2/2]
43.206.3 Member Function Documentation
43.206.3.1 conj_transpose()
43.206.3.2 operator=() [1/2]
43.206.3.3 operator=() [2/2]
43.206.3.4 transpose()
43.207 gko::UseComposition< ValueType > Class Template Reference
43.207.1 Detailed Description
43.207.2 Member Function Documentation
43.207.2.1 get_composition()
43.207.2.2 get_operator_at()
43.208 gko::syn::value_list< T, Values > Struct Template Reference
43 208 1 Detailed Description 888

43.209 gko::ValueMismatch Class Reference
43.209.1 Detailed Description
43.209.2 Constructor & Destructor Documentation
43.209.2.1 ValueMismatch()
43.210 gko::experimental::distributed::Vector< ValueType > Class Template Reference
43.210.1 Detailed Description
43.210.2 Member Function Documentation
43.210.2.1 add_scaled()
43.210.2.2 at_local() [1/4]
43.210.2.3 at_local() [2/4]
43.210.2.4 at_local() [3/4]
43.210.2.5 at_local() [4/4]
43.210.2.6 compute_absolute()
43.210.2.7 compute_conj_dot() [1/2]
43.210.2.8 compute_conj_dot() [2/2]
43.210.2.9 compute_dot() [1/2]
43.210.2.10 compute_dot() [2/2]
43.210.2.11 compute_norm1() [1/2]
43.210.2.12 compute_norm1() [2/2]89
43.210.2.13 compute_norm2() [1/2]
43.210.2.14 compute_norm2() [2/2]89
43.210.2.15 compute_squared_norm2() [1/2]
43.210.2.16 compute_squared_norm2() [2/2]
43.210.2.17 create_const() [1/2]
43.210.2.18 create_const() [2/2]
43.210.2.19 create_real_view() [1/2]
43.210.2.20 create_real_view() [2/2]
43.210.2.21 create_with_config_of()
43.210.2.22 create_with_type_of() [1/2]
43.210.2.23 create_with_type_of() [2/2]
43.210.2.24 fill()
43.210.2.25 get_const_local_values()
43.210.2.26 get_local_values()
43.210.2.27 get_local_vector()
43.210.2.28 inv_scale()
43.210.2.29 make_complex() [1/2] 90
43.210.2.30 make_complex() [2/2] 90
43.210.2.31 read_distributed() [1/2]
43.210.2.32 read_distributed() [2/2]
43.210.2.33 scale()
43.210.2.34 sub_scaled()
43.211 ako::version Struct Reference

43.211.1 Detailed Description	04
43.211.2 Member Data Documentation	04
43.211.2.1 tag	04
43.212 gko::version_info Class Reference	04
43.212.1 Detailed Description	05
43.212.2 Member Function Documentation	05
43.212.2.1 get()	05
43.212.3 Member Data Documentation	05
43.212.3.1 core_version	05
43.212.3.2 cuda_version	06
43.212.3.3 dpcpp_version	06
43.212.3.4 hip_version	06
43.212.3.5 omp_version	06
43.212.3.6 reference_version	06
43.213 gko::experimental::mpi::window< ValueType > Class Template Reference	06
43.213.1 Detailed Description	80
43.213.2 Constructor & Destructor Documentation	80
43.213.2.1 window() [1/3]	09
43.213.2.2 window() [2/3]	09
43.213.2.3 window() [3/3]	09
43.213.3 Member Function Documentation	10
43.213.3.1 accumulate()	10
43.213.3.2 fence()	10
43.213.3.3 fetch_and_op()	11
43.213.3.4 flush()	11
43.213.3.5 flush_local()	11
43.213.3.6 get()	12
43.213.3.7 get_accumulate()	12
43.213.3.8 get_window()	13
43.213.3.9 lock()	13
43.213.3.10 lock_all()	14
43.213.3.11 operator=()	14
43.213.3.12 put()	14
43.213.3.13 r_accumulate()	15
43.213.3.14 r_get()	15
43.213.3.15 r_get_accumulate()	16
43.213.3.16 r_put()	17
43.213.3.17 unlock()	17
43.214 gko::solver::workspace_traits< Solver > Struct Template Reference	18
43.214.1 Detailed Description	18
43.215 gko::WritableToMatrixData< ValueType, IndexType > Class Template Reference	18
43.215.1 Detailed Description	18

dex									921
	43.215.2.1 write()	 	 	 		 		 	919
	43.215.2 Member Function Documentation	 	 	 		 		 	919

Main Page

This is the main page for the Ginkgo library pdf documentation. The repository is hosted on github. Documentation on aspects such as the build system, can be found at the Installation Instructions page. The Example programs can help you get started with using Ginkgo.

1.0.0.1 Modules

The Ginkgo library can be grouped into modules and these modules form the basic building blocks of Ginkgo. The modules can be summarized as follows:

- Executors: Where do you want your code to be executed?
- · Linear Operators: What kind of operation do you want Ginkgo to perform?
 - Solvers : Solve a linear system for a given matrix.
 - Preconditioners: Precondition a system for a solve.
 - SpMV employing different Matrix formats: Perform a sparse matrix vector multiplication with a particular matrix format.
- Logging : Monitor your code execution.
- Stopping criteria: Manage your iteration stopping criteria.

2 Main Page

Installation Instructions

2.0.1 Building

Use the standard CMake build procedure:

```
mkdir build; cd build cmake -G "Unix Makefiles" [OPTIONS] .. && make
```

Use cmake --build . in some systems like MinGW or Microsoft Visual Studio which do not use make.

For Microsoft Visual Studio, use cmake --build . --config <build_type> to decide the build type. The possible options are Debug, Release, RelWithDebInfo and MinSizeRel.

Replace <code>[OPTIONS]</code> with desired cmake options for your build. Ginkgo adds the following additional switches to control what is being built:

- -DGINKGO_DEVEL_TOOLS={ON, OFF} sets up the build system for development (requires clang-format, will also download git-cmake-format), default is OFF. The default behavior installs a pre-commit hook, which disables git commits. If it is set to ON, a new pre-commit hook for formatting will be installed (enabling commits again). In both cases the hook may overwrite a user defined pre-commit hook when Ginkgo is used as a submodule.
- -DGINKGO_MIXED_PRECISION={ON, OFF} compiles true mixed-precision kernels instead of converting data on the fly, default is OFF. Enabling this flag increases the library size, but improves performance of mixed-precision kernels.
- -DGINKGO_BUILD_TESTS={ON, OFF} builds Ginkgo's tests (will download googletest), default is ON.
- -DGINKGO_FAST_TESTS={ON, OFF} reduces the input sizes for a few slow tests to speed them up, default is OFF.
- -DGINKGO_BUILD_BENCHMARKS={ON, OFF} builds Ginkgo's benchmarks (will download gflags and rapidjson), default is ON.
- -DGINKGO_BUILD_EXAMPLES={ON, OFF} builds Ginkgo's examples, default is ON
- -DGINKGO_BUILD_EXTLIB_EXAMPLE={ON, OFF} builds the interfacing example with deal.II, default is OFF.
- -DGINKGO_BUILD_REFERENCE={ON, OFF} build reference implementations of the kernels, useful for testing, default is ON
- -DGINKGO_BUILD_OMP={ON, OFF} builds optimized OpenMP versions of the kernels, default is ON if the selected C++ compiler supports OpenMP, OFF otherwise.

4 Installation Instructions

• -DGINKGO_BUILD_CUDA={ON, OFF} builds optimized cuda versions of the kernels (requires CUDA), default is ON if a CUDA compiler could be detected, OFF otherwise.

- -DGINKGO_BUILD_DPCPP={ON, OFF} builds optimized DPC++ versions of the kernels (requires C← MAKE_CXX_COMPILER to be set to the dpcpp compiler). The default is ON if CMAKE_CXX_COMPILER is a DPC++ compiler, OFF otherwise.
- -DGINKGO_BUILD_HIP={ON, OFF} builds optimized HIP versions of the kernels (requires HIP), default is ON if an installation of HIP could be detected, OFF otherwise.
- -DGINKGO_HIP_AMDGPU="gpuarch1; gpuarch2" the amdgpu_target(s) variable passed to hipcc for the hcc HIP backend. The default is none (auto).
- -DGINKGO_BUILD_HWLOC={ON, OFF} builds Ginkgo with HWLOC. If system HWLOC is not found, Ginkgo will try to build it. Default is ON on Linux. Ginkgo does not support HWLOC on Windows/MacOS, so the default is OFF on Windows/MacOS.
- -DGINKGO_BUILD_DOC={ON, OFF} creates an HTML version of Ginkgo's documentation from inline comments in the code. The default is OFF.
- -DGINKGO_DOC_GENERATE_EXAMPLES={ON, OFF} generates the documentation of examples in Ginkgo. The default is ON.
- -DGINKGO_DOC_GENERATE_PDF={ON, OFF} generates a PDF version of Ginkgo's documentation from inline comments in the code. The default is OFF.
- -DGINKGO_DOC_GENERATE_DEV={ON, OFF} generates the developer version of Ginkgo's documentation. The default is OFF.
- -DGINKGO_EXPORT_BUILD_DIR={ON, OFF} adds the Ginkgo build directory to the CMake package registry. The default is OFF.
- -DGINKGO_WITH_CLANG_TIDY={ON, OFF} makes Ginkgo call clang-tidy to find programming issues. The path can be manually controlled with the CMake variable -DGINKGO_CLANG_TIDY_PA← TH=<path>. The default is OFF.
- -DGINKGO_WITH_IWYU={ON, OFF} makes Ginkgo call iwyu to find include issues. The path can be manually controlled with the CMake variable -DGINKGO_IWYU_PATH=<path>. The default is OFF.
- -DGINKGO_CHECK_CIRCULAR_DEPS={ON, OFF} enables compile-time checks for circular dependencies between different Ginkgo libraries and self-sufficient headers. Should only be used for development purposes. The default is OFF.
- -DGINKGO_VERBOSE_LEVEL=integer sets the verbosity of Ginkgo.
 - 0 disables all output in the main libraries,
 - 1 enables a few important messages related to unexpected behavior (default).
- GINKGO_INSTALL_RPATH allows setting any RPATH information when installing the Ginkgo libraries. If this is OFF, the behavior is the same as if all other RPATH flags are set to OFF as well. The default is ON.
- GINKGO_INSTALL_RPATH_ORIGIN adds \$ORIGIN (Linux) or @loader_path (MacOS) to the installation RPATH. The default is ON.
- GINKGO_INSTALL_RPATH_DEPENDENCIES adds the dependencies to the installation RPATH. The default is OFF.
- -DCMAKE_INSTALL_PREFIX=path sets the installation path for make install. The default value is usually something like /usr/local.
- -DCMAKE_BUILD_TYPE=type specifies which configuration will be used for this build of Ginkgo. The
 default is RELEASE. Supported values are CMake's standard build types such as DEBUG and RELEASE and
 the Ginkgo specific COVERAGE, ASAN (AddressSanitizer), LSAN (LeakSanitizer), TSAN (ThreadSanitizer)
 and UBSAN (undefined behavior sanitizer) types.

- -DBUILD_SHARED_LIBS={ON, OFF} builds ginkgo as shared libraries (OFF) or as dynamic libraries (ON), default is ON.
- -DGINKGO_JACOBI_FULL_OPTIMIZATIONS={ON, OFF} use all the optimizations for the CUDA Jacobi algorithm. OFF by default. Setting this option to ON may lead to very slow compile time (>20 minutes) for the jacobi_generate_kernels.cu file and high memory usage.
- -DCMAKE_CUDA_HOST_COMPILER=path instructs the build system to explicitly set CUDA's host compiler to the path given as argument. By default, CUDA uses its toolchain's host compiler. Setting this option may help if you're experiencing linking errors due to ABI incompatibilities. This option is supported since CMake 3.8 but documented starting from 3.10.
- -DGINKGO_CUDA_ARCHITECTURES=<list> where <list> is a semicolon (;) separated list of architectures. Supported values are:
 - Auto
 - Kepler, Maxwell, Pascal, Volta, Turing, Ampere
 - CODE, CODE (COMPUTE), (COMPUTE)

Auto will automatically detect the present CUDA-enabled GPU architectures in the system. Kepler, Maxwell, Pascal, Volta and Ampere will add flags for all architectures of that particular NVIDIA GPU generation. COMPUTE and CODE are placeholders that should be replaced with compute and code numbers (e.g. for compute_70 and sm_70 COMPUTE and CODE should be replaced with 70. Default is Auto. For a more detailed explanation of this option see the ARCHITECTURES specification list section in the documentation of the CudaArchitectureSelector CMake module.

```
For example, to build everything (in debug mode), use:
```

```
cmake -G "Unix Makefiles" -H. -BDebug -DCMAKE_BUILD_TYPE=Debug -DGINKGO_DEVEL_TOOLS=ON \
    -DGINKGO_BUILD_TESTS=ON -DGINKGO_BUILD_REFERENCE=ON -DGINKGO_BUILD_OMP=ON \
    -DGINKGO_BUILD_CUDA=ON -DGINKGO_BUILD_HIP=ON
cmake --build Debug
```

NOTE: Ginkgo is known to work with the Unix Makefiles, Ninja, MinGW Makefiles and Visual Studio 16 2019 based generators. Other CMake generators are untested.

2.0.2 Building Ginkgo in Windows

Depending on the configuration settings, some manual work might be required:

- Build Ginkgo with Debug mode: Some Debug build specific issues can appear depending on the machine and environment: When you encounter the error message ld: error: export ordinal too large, add the compilation flag -O1 by adding -DCMAKE_CXX_FLAGS=-O1 to the CMake invocation.
- Build Ginkgo in *MinGW*:\ If encountering the issue cclplus.exe: out of memory allocating 65536 bytes, please follow the workaround in reference, or trying to compile ginkgo again might work.

2.0.3 Building Ginkgo with HIP support

Ginkgo provides a HIP backend. This allows to compile optimized versions of the kernels for either AMD or NV \leftarrow IDIA GPUs. The CMake configuration step will try to auto-detect the presence of HIP either at /opt/rocm/hip or at the path specified by HIP_PATH as a CMake parameter (-DHIP_PATH=) or environment variable (export HIP_PATH=), unless -DGINKGO_BUILD_HIP=ON/OFF is set explicitly.

6 Installation Instructions

2.0.3.1 Changing the paths to search for HIP and other packages

All HIP installation paths can be configured through the use of environment variables or CMake variables. This way of configuring the paths is currently imposed by the HIP tool suite. The variables are the following:

- CMake -DROCM_PATH= or environment export ROCM_PATH=: sets the ROCM installation path. The default value is /opt/rocm/.
- CMake -DHIP_CLANG_PATH or environment export HIP_CLANG_PATH=: sets the HIP compatible clang binary path. The default value is \${ROCM PATH}/llvm/bin.
- CMake -DHIP_PATH= or environment export HIP_PATH=: sets the HIP installation path. The default value is \${ROCM_PATH}/hip.
- CMake -DHIPBLAS_PATH= or environment export HIPBLAS_PATH=: sets the hipBLAS installation path. The default value is \$ {ROCM_PATH} / hipblas.
- CMake -DHIPSPARSE_PATH= or environment export HIPSPARSE_PATH=: sets the hipSPARSE installation path. The default value is \${ROCM_PATH}/hipsparse.
- CMake -DHIPFFT_PATH= or environment export HIPFFT_PATH=: sets the hipFFT installation path. The default value is \${ROCM_PATH}/hipfft.
- CMake -DROCRAND_PATH= or environment export ROCRAND_PATH=: sets the rocRAND installation path. The default value is \${ROCM_PATH}/rocrand.
- CMake -DHIPRAND_PATH= or environment export HIPRAND_PATH=: sets the hipRAND installation path. The default value is \${ROCM_PATH}/hiprand.
- environment export CUDA_PATH=: where hipcc can find CUDA if it is not in the default /usr/local/cuda path.

2.0.3.2 HIP platform detection of AMD and NVIDIA

By default, Ginkgo uses the output of /opt/rocm/hip/bin/hipconfig --platform to select the backend. The accepted values are either hcc (amd with ROCM >= 4.1) or nvcc (nvidia with ROCM >= 4.1). When on an AMD or NVIDIA system, this should output the correct platform by default. When on a system without GPUs, this should output hcc by default. To change this value, export the environment variable HIP_PLATFORM like so: export HIP_PLATFORM=nvcc # or nvidia for ROCM >= 4.1

2.0.3.3 Setting platform specific compilation flags

Platform specific compilation flags can be given through the following CMake variables:

- -DGINKGO_HIP_COMPILER_FLAGS=: compilation flags given to all platforms.
- -DGINKGO_HIP_NVCC_COMPILER_FLAGS=: compilation flags given to NVIDIA platforms.
- -DGINKGO_HIP_CLANG_COMPILER_FLAGS=: compilation flags given to AMD clang compiler.

2.0.4 Third party libraries and packages

Ginkgo relies on third party packages in different cases. These third party packages can be turned off by disabling the relevant options.

- GINKGO_BUILD_TESTS=ON: Our tests are implemented with Google Test;
- GINKGO_BUILD_BENCHMARKS=ON: For argument management we use gflags and for JSON parsing we use RapidJSON;
- GINKGO_DEVEL_TOOLS=ON: git-cmake-format is our CMake helper for code formatting.
- GINKGO BUILD HWLOC=ON: hwloc to detect and control cores and devices.
- GINKGO_BUILD_HWLOC=ON and GINKGO_BUILD_TESTS=ON: libnuma is required when testing the functions provided through MachineTopology.
- GINKGO_BUILD_EXAMPLES=ON: OpenCV is required for some examples, they are disabled when OpenCV is not available.
- GINKGO_BUILD_DOC=ON: doxygen is required to build the documentation and additionally graphviz is required to build the class hierarchy graphs.

Ginkgo attempts to use pre-installed versions of these package if they match version requirements using find—package. Otherwise, the configuration step will download the files for each of the packages GTest, gflags, RapidJSON and hwloc and build them internally.

Note that, if the external packages were not installed to the default location, the CMake option $-DCMAKE_{\leftarrow}$ PREFIX_PATH=<path-list> needs to be set to the semicolon (;) separated list of install paths of these external packages. For more Information, see the CMake documentation for CMAKE_PREFIX_PATH for details.

For convenience, the options GINKGO_INSTALL_RPATH [_.*] can be used to bind the installed Ginkgo shared libraries to the path of its dependencies.

2.0.5 Installing Ginkgo

To install Ginkgo into the specified folder, execute the following command in the build folder ${\tt make\ install}$

If the installation prefix (see CMAKE_INSTALL_PREFIX) is not writable for your user, e.g. when installing Ginkgo system-wide, it might be necessary to prefix the call with sudo.

After the installation, CMake can find ginkgo with $find_package$ (Ginkgo). An example can be found in the test_install.

8 Installation Instructions

Testing Instructions

3.0.1 Running the unit tests

You need to compile ginkgo with -DGINKGO_BUILD_TESTS=ON option to be able to run the tests.

3.0.1.1 Using make test

After configuring Ginkgo, use the following command inside the build folder to run all tests:

The output should contain several lines of the form:

To run only a specific test and see more details results (e.g. if a test failed) run the following from the build folder:

where path/to/test is the path returned by make test.

3.0.1.2 Using make quick_test

After compiling Ginkgo, use the following command inside the build folder to run a small subset of tests that should execute quickly:

make quick_test

These tests do not use GPU features except for a few device property queries, so they may still fail if Ginkgo was compiled with GPU support, but no such GPU is available. The output is equivalent to make test.

3.0.1.3 Using CTest

The tests can also be ran through CTest from the command line, for example when in a configured build directory: CTEST - T start -T build -T test -T submit

Will start a new test campaign (usually in Experimental mode), build Ginkgo with the set configuration, run the tests and submit the results to our CDash dashboard.

Another option is to use Ginkgo's CTest script which is configured to build Ginkgo with default settings, runs the tests and submits the test to our CDash dashboard automatically.

To run the script, use the following command:

```
ctest -S cmake/CTestScript.cmake
```

The default settings are for our own CI system. Feel free to configure the script before launching it through variables or by directly changing its values. A documentation can be found in the script itself.

10 Testing Instructions

Running the benchmarks

In addition to the unit tests designed to verify correctness, Ginkgo also includes an extensive benchmark suite for checking its performance on all Ginkgo supported systems. The purpose of Ginkgo's benchmarking suite is to allow easy and complete reproduction of Ginkgo's performance, and to facilitate performance debugging as well. Most results published in Ginkgo papers are generated thanks to this benchmarking suite and are accessible online under the ginkgo-data repository. These results can also be used for performance comparison in order to ensure that you get similar performance as what is published on this repository.

To compile the benchmarks, the flag <code>-GINKGO_BUILD_BENCHMARKS=ON</code> has to be set during the <code>cmake</code> step. In addition, the <code>ssget command-line utility</code> has to be installed on the system. The purpose of this file is to explain in detail the capacities of this benchmarking suite as well as how to properly setup everything.

Here is a short description of the content of this file:

- 1. Ginkgo setup and best practice guidelines
- 2. Installing and using the ssget tool to fetch the SuiteSparse matrices.
- 3. Benchmarking overview and how to run them in a simple way.
- 4. How to publish the benchmark results online and use the Ginkgo Performance Explorer (GPE) for performance analysis (optional).
- 5. Using the benchmark suite for performance debugging thanks to the loggers.
- 6. All available benchmark customization options.

4.0.1 1: Ginkgo setup and best practice guidelines

Before benchmarking Ginkgo, make sure that you follow the general guidelines in order to ensure best performance.

- 1. The code should be compiled in Release mode.
- 2. Make sure the machine has no competing jobs. On a Linux machine multiple commands can be used, last shows the currently opened sessions, top or htop allows to show the current machine load, and if considering using specific GPUs, nvidia-smi or room-smi can be used to check their load.
- 3. By default, Ginkgo's benchmarks will always do at least one warm-up run. For better accuracy, every benchmark is also averaged over 10 runs, except for the solver benchmark which are usually fairly long. These parameters can be tuned at the command line to either shorten benchmarking time or improve benchmarking accuracy.

In addition, the following specific options can be considered:

- 1. When specifically using the adaptive block jacobi preconditioner, enable the GINKGO_JACOBI_FULL_O← PTIMIZATIONS CMake flag. Be careful that this will use much more memory and time for the compilation due to compiler performance issues with register optimizations, in particular.
- 2. The current benchmarking setup also allows to benchmark only the overhead by using as either (or for all) preconditioner/spmv/solver, the special overhead LinOp. If your purpose is to check Ginkgo's overhead, make sure to try this mode.

4.0.2 2: Using ssget to fetch the matrices

The benchmark suite tests Ginkgo's performance using the SuiteSparse matrix collection and artificially generated matrices. The suite sparse collection will be downloaded automatically when the benchmarks are run. This is done thanks to the ssget command—line utility.

To install ssget, access the repository and copy the file ssget into a directory present in your PATH variable as per the tool's README.md instructions. The tool can be installed either in a global system path or a local directory such as \$HOME/.local/bin. After installing the tool, it is important to review the ssget script and configure as needed the variable ARCHIVE LOCATION on line 39. This is where the matrices will be stored into.

The Ginkgo benchmark can be set to run on only a portion of the SuiteSparse matrix collection as we will see in the following section. Please note that the entire collection requires roughly 100GB of disk storage in its compressed format, and roughly 25GB of additional disk space for intermediate data (such us uncompressing the archive). Additionally, the benchmark runs usually take a long time (SpMV benchmarks on the complete collection take roughly 24h using the K20 GPU), and will stress the system.

Before proceeding, it can be useful in order to save time to download the matrices as preparation. This can be done by using the ssget -f -i i command where i is the ID of the matrix to be downloaded. The following loop allows to download the full SuiteSparse matrix collection:

```
for i in $(seq 0 $(ssget -n)); do
    ssget -f -i ${i}
done
```

Note that ssget can also be used to query properties of the matrix and filter the matrices which are downloaded. For example, the following will download only positive definite matrices with less than 500M non zero elements and 10M columns. Please refer to the ssget documentation for more information.

```
for i in $(seq 0 $(ssget -n)); do
    posdef=$(ssget -p posdef -i ${i})
    cols=$(ssget -p cols -i ${i})
    nnz=$(ssget -p nonzeros -i ${i})
    if [ "$posdef" -eq 1 -a "$cols" -lt 10000000 -a "$nnz" -lt 500000000 ]; then
        ssget -f -i ${i}
    fi
done
```

4.0.3 3: Benchmarking overview

The benchmark suite is invoked using the <code>make benchmark command</code> in the build directory. Under the hood, this command simply calls the script <code>benchmark/run_all_benchmarks.sh</code> so it is possible to manually launch this script as well. The behavior of the suite can be modified using environment variables. Assuming the <code>bash shell</code> is used, these can either be specified via the <code>export command</code> to persist between multiple runs:

<code>export VARIABLE="value"</code>

```
... make benchmark
```

or specified on the fly, on the same line as the make benchmark command:

```
VARIABLE="value" ... make benchmark
```

Since make sets any variables passed to it as temporary environment variables, the following shorthand can also be used:

```
make benchmark VARIABLE="value" ...
```

A combination of the above approaches is also possible (e.g. it may be useful to <code>export</code> the <code>SYSTEM_NAME</code> variable, and specify the others at every benchmark run).

The benchmark suite can take a number of configuration parameters. Benchmarks can be run only for sparse matrix vector products (spmv), for full solvers (with or without preconditioners), or for preconditioners only when supported. The benchmark suite also allows to target a sub-part of the SuiteSparse matrix collection. For details, see the available benchmark options. Here are the most important options:

- BENCHMARK={spmv, solver, preconditioner} allows to select the type of benchmark to be ran.
- EXECUTOR={reference, cuda, hip, omp, dpcpp} select the executor and platform the benchmarks should be ran on.
- SYSTEM_NAME=<name> a name which will be used to designate this platform (e.g. V100, RadeonVII, ...).
- SEGMENTS=<N> Split the benchmarked matrix space into <N> segments. If specified, SEGMENT_ID also has to be set.
- SEGMENT_ID=<I> used in combination with the SEGMENTS variable. <I> should be an integer between 1 and <N>, the number of SEGMENTS. If specified, only the <I>-th segment of the benchmark suite will be run.
- BENCHMARK_PRECISION defines the precision the benchmarks are run in. Supported values are ←: "double" (default), "single", "dcomplex" and "scomplex"
- MATRIX_LIST_FILE=/path/to/matrix_list.file allows to list SuiteSparse matrix id or name to benchmark. As an example, a matrix list file containing the following will ensure that benchmarks are ran for only those three matrices:
 - "1903 Freescale/circuit5M thermal2"

4.0.4 4: Publishing the results on Github and analyze the results with the GPE (optional)

The previous experiments generated json files for each matrices, each containing timing, iteration count, achieved precision, ... depending on the type of benchmark run. These files are available in the directory \${ginkgo} _build_dir}/benchmark/results/. These files can be analyzed and processed through any tool (e.g. python). In this section, we describe how to generate the plots by using Ginkgo's GPE tool. First, we need to publish the experiments into a Github repository which will be then linked as source input to the GPE. For this, we can simply fork the ginkgo-data repository. To do so, we can go to the github repository and use the forking interface: https://github.com/ginkgo-project/ginkgo-data/

Once it's done, we want to clone the repository locally, put all results online and access the GPE for plotting the results. Here are the detailed steps:

```
git clone https://github.com/<username>/ginkgo-data.git $HOME/ginkgo_benchmark/ginkgo-data # send the benchmarked data to the ginkgo-data repository
# If needed, remove the old data so that no previous data is left.
# rm -r ${HOME}/ginkgo_benchmark/ginkgo-data/data/${SYSTEM_NAME}
rsync -rtv ${ginkgo_benchmark/ginkgo-data/data/$HOME/ginkgo_benchmark/ginkgo-data/data/
cd ${HOME}/ginkgo_benchmark/ginkgo-data/data/
# The following updates the main `.json` files with the list of data.
# Ensure a python 3 installation is available.
./build-list . > list.json
./agregate < list.json > agregate.json
./represent . > represent.json
git config --local user.name "<Name>"
git config --local user.email "<email>"
git commit -am "Ginkgo benchmark ${BENCHMARK} of ${SYSTEM_NAME}..."
```

git push

Note that depending on what data is of interest, you may need to update the scripts build-list or agregate to change which files you want to agglomerate and summarize (depending on the system name), or which data you want to select (solver results, spmv results, ...).

For the generating the plots in the GPE, here are the steps to go through:

- 1. Access the GPE: https://ginkgo-project.github.io/gpe/
- 2. Update data root URL, from https://raw.githubusercontent.com/ginkgo-project/ginkgo-data/mas to https://raw.githubusercontent.com/<username>/ginkgo-data/
branch>/data
- 3. Click on the arrow to load the data, select the Result Summary entry above.
- 4. Click on select an example to choose a plotting script. Multiple scripts are available by default in different branches. You can use the jsonata and chart js languages to develop your own as well.
- 5. The results should be available in the tab "plot" on the right side. Other tabs allow to access the result of the processed data after invoking the processing script.

4.0.5 5: Detailed performance analysis and debugging

Detailed performance analysis can be ran by passing the environment variable DETAILED=1 to the benchmarking script. This detailed run is available for solvers and allows to log the internal residual after every iteration as well as log the time taken by all operations. These features are also available in the performance-debugging example which can be used instead and modified as needed to analyze Ginkgo's performance.

These features are implemented thanks to the loggers located in the file \${ginkgo_src_dir}/benchmark/utils/loggers
Ginkgo possesses hooks at all important code location points which can be inspected thanks to the logger. In this
fashion, it is easy to use these loggers also for tracking memory allocation sizes and other important library aspects.

4.0.6 6: Available benchmark options

There are a set amount of options available for benchmarking. Most important options can be configured through the benchmarking script itself thanks to environment variables. Otherwise, some specific options are not available through the benchmarking scripts but can be directly configured when running the benchmarking program itself. For a list of all options, run for example \${ginkgo_build_dir}/benchmark/solver/solver --help.

The supported environment variables are described in the following list:

- BENCHMARK={spmv, solver, preconditioner} allows to select the type of benchmark to be ran. Default is spmv.
 - spmv Runs the sparse matrix-vector product benchmarks on the SuiteSparse collection.
 - solver Runs the solver benchmarks on the SuiteSparse collection. The matrix format is determined
 by running the spmv benchmarks first, and using the fastest format determined by that benchmark.
 - preconditioner Runs the preconditioner benchmarks on artificially generated block-diagonal matrices.
- EXECUTOR={reference, cuda, hip, omp, dpcpp} select the executor and platform the benchmarks should be ran on. Default is cuda.
- SYSTEM_NAME=<name> a name which will be used to designate this platform (e.g. V100, RadeonVII, ...) and not overwrite previous results. Default is unknown.

- SEGMENTS=<N> Split the benchmarked matrix space into <N> segments. If specified, SEGMENT_ID also has to be set. Default is 1.
- SEGMENT_ID=<I> used in combination with the SEGMENTS variable. <I> should be an integer between 1 and <N>, the number of SEGMENTS. If specified, only the <I>-th segment of the benchmark suite will be run. Default is 1.
- MATRIX_LIST_FILE=/path/to/matrix_list.file allows to list SuiteSparse matrix id or name to benchmark. As an example, a matrix list file containing the following will ensure that benchmarks are ran for only those three matrices: ``` 1903 Freescale/circuit5M thermal2 ``*DEVICE_ID- the accelerator device ID to target for the benchmark. The default is0. *DRY_RUN={true, false}- If set totrue`, prepares the system for the benchmark runs (downloads the collections, creates the result structure, etc.) and outputs the list of commands that would normally be run, but does not run the benchmarks themselves. Default is false.
- PRECONDS={ jacobi, ic, ilu, paric, parict, parilu, parilut, ic-isai, ilu-isai, paric-isai, parithe preconditioners to use for either solver or preconditioner benchmarks. Multiple options can be passed to this variable. Default is none.
- FORMATS={csr,coo,ell,hybrid,sellp,hybridxx,cusparse_xx,hipsparse_xx} the matrix formats to benchmark for the spmv phase of the benchmark. Run \${ginkgo_build_← dir}/benchmark/spmv/spmv --help for a full list. If needed, multiple options for hybrid with different optimization parameters are available. Depending on the libraries available at build time, vendor library formats (cuSPARSE with cusparse_prefix or hipSPARSE with hipsparse_prefix) can be used as well. Multiple options can be passed. The default is csr,coo,ell,hybrid,sellp.
- SOLVERS={bicgstab,bicg,cg,cgs,fcg,gmres,cb_gmres_{keep,reduce1,reduce2,integer,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce1,reduce2,ireduce3,ireduce1,reduce2,ireduce3,ireduc
 - the solvers which should be benchmarked. Multiple options can be passed. The default is bicgstab,cg,cgs,fcg,gmres,idr. Note that lower/upper_trs by default don't use a preconditioner, as they are by default exact direct solvers.
- SOLVERS_PRECISION=cision> the minimal residual reduction before which the solver should stop. The default is 1e-6.
- SOLVERS_MAX_ITERATIONS=<number> the maximum number of iterations with which a solver should be ran. The default is 10000.
- SOLVERS_RHS={1, random, sinus} whether to use a vector of all ones, random values or b = A * (s / |s|)\$ with s(idx) = sin(idx) (for complex numbers, s(idx) = sin(2*idx) + i * sin(2*idx+1)) as the right-hand side in solver benchmarks. Default is 1.
- SOLVERS_INITIAL_GUESS={rhs, 0, random} the initial guess generation of the solvers. rhs uses the right-hand side, 0 uses a zero vector and random generates a random vector as the initial guess.
- DETAILED={0,1} selects whether detailed benchmarks should be ran for the solver benchmarks, can be either 0 (off) or 1 (on). The default is 0.
- GPU_TIMER={true, false} If set to true, use the gpu timer, which is valid for cuda/hip executor, to measure the timing. Default is false.
- SOLVERS_JACOBI_MAX_BS sets the maximum block size for the Jacobi preconditioner (if used, otherwise, it does nothing) in the solvers benchmark. The default is '32'.
- SOLVERS_GMRES_RESTART the maximum dimension of the Krylov space to use in GMRES. The default is 100

Contributing guidelines

We are glad that you are interested in contributing to Ginkgo. Please have a look at our coding guidelines before proposing a pull request.

5.1 Table of Contents

Most Important stuff

Project Structure

- Extended header files
- Using library classes

Git related

- Our git Workflow
- Writing good commit messages
- Creating, Reviewing and Merging Pull Requests

Code Style

- Automatic code formatting
- · Naming Scheme
- Whitespace
- Include statement grouping
- Other Code Formatting not handled by ClangFormat
- CMake coding style

Helper Scripts

- · Create a new algorithm
- Converting CUDA code to HIP code

Writing Tests

- Testing know-how
- Some general rules
- Writing tests for kernels

Documentation style

- Developer targeted notes
- · Whitespaces
- Documenting examples

Other programming comments

- C++ standard stream objects
- Warnings
- · Avoiding circular dependencies

5.2 Most important stuff (A TL;DR)

- GINKGO_DEVEL_TOOLS needs to be set to on to commit. This requires clang-format to be installed. See Automatic code formatting for more details. Once installed, you can run make format in your build/ folder to automatically format your modified files. As make format unstages your files post-formatting, you must stage the files again once you have verified that make format has done the appropriate formatting, before committing the files.
- See Our git workflow to get a quick overview of our workflow.
- See Creating, Reviewing and Merging Pull Requests on how to create a Pull request.

5.3 Project structure

Ginkgo is divided into a core module with common functionalities independent of the architecture, and several kernel modules (reference, omp, cuda, hip, dpcpp) which contain low-level computational routines for each supported architecture.

5.3.1 Extended header files

Some header files from the core module have to be extended to include special functionality for specific architectures. An example of this is core/base/math.hpp, which has a GPU counterpart in cuda/base/math.co
hpp. For such files you should always include the version from the module you are working on, and this file will internally include its core counterpart.

5.4 Git related 19

5.3.2 Using library classes

You can use and call functions of existing classes inside a kernel (that are defined and not just declared in a header file), however, you are not allowed to create new instances of a polymorphic class inside a kernel (or in general inside any kernel module like cuda/hip/omp/reference) as this creates circular dependencies between the core and the backend library. With this in mind, our CI contains a job which checks if such a circular dependency exists. These checks can be run manually using the <code>-DGINKGO_CHECK_CIRCULAR_DEPS=ON</code> option in the CMake configuration.

For example, when creating a new matrix class AB by combining existing classes A and B, the AB::apply() function composed of invocations to A::apply() and B::apply() can only be defined in the core module, it is not possible to create instances of A and B inside the AB kernel files. This is to avoid the aforementioned circular dependency issue. An example for such a class is the Hybrid matrix format, which uses the apply() of the Ell and Coo matrix formats. Nevertheless, it is possible to call the kernels themselves directly within the same executor. For example, cuda::dense::add_scaled() can be called from any other cuda kernel.

5.4 Git related

Ginkgo uses git, the distributed version control system to track code changes and coordinate work among its developers. A general guide to git can be found in its extensive documentation.

5.4.1 Our git workflow

In Ginkgo, we prioritize keeping a clean history over accurate tracking of commits. <code>git rebase</code> is hence our command of choice to make sure that we have a nice and linear history, especially for pulling the latest changes from the <code>develop</code> branch. More importantly, rebasing upon develop is **required** before the commits of the PR are merged into the <code>develop</code> branch.

5.4.2 Writing good commit messages

With software sustainability and maintainability in mind, it is important to write commit messages that are short, clear and informative. Ideally, this would be the format to prefer:

```
Summary of the changes in a sentence, max 50 chars.

More detailed comments:
+ Changes that have been added.
- Changes that been removed.

Related PR: https://github.com/ginkgo-project/ginkgo/pull/<PR-number>
```

You can refer to this informative guide for more details.

5.4.2.1 Attributing credit

Git has a nice feature where it allows you to add a co-author for your commit, if you would like to attribute credits for the changes made in the commit. This can be done by:

```
Commit message.
Co-authored-by: Name <email@domain>
```

In the Ginkgo commit history, this is most common associated with suggested improvements from code reviews.

5.4.3 Creating, Reviewing and Merging Pull Requests

- The develop branch is the default branch to submit PR's to. From time to time, we merge the develop branch to the master branch and create tags on the master to create new releases of Ginkgo. Therefore, all pull requests must be merged into develop.
- Please have a look at the labels and make sure to add the relevant labels.
- You can mark the PR as a WIP if you are still working on it, Ready for Review when it is ready for others to review it.
- Assignees to the PR should be the ones responsible for merging that PR. Currently, it is only possible to assign members within the ginkgo-project.
- · Each pull request requires at least two approvals before merging.
- PR's created from within the repository will automatically trigger two CI pipelines on pushing to the branch from the which the PR has been created. The Github Actions pipeline tests our framework on Mac OSX and on Windows platforms. Another comprehensive Linux based pipeline is run from a mirror on gitlab and contains additional checks like static analysis and test coverage.
- Once a PR has been approved and the build has passed, one of the reviewers can mark the PR as READY TO MERGE. At this point the creator/assignee of the PR needs to verify that the branch is up to date with develop and rebase it on develop if it is not.

5.5 Code style

5.5.1 Automatic code formatting

Ginkgo uses ClangFormat (executable is usually named clang-format) and a custom .clang-format configuration file (mostly based on ClangFormat's *Google* style) to automatically format your code. **Make sure you have ClangFormat set up and running properly** (you should be able to run make format from Ginkgo's build directory) before committing anything that will end up in a pull request against ginkgo-project/ginkgo repository. In addition, you should **never** modify the .clang-format configuration file shipped with Ginkgo. E.g. if ClangFormat has trouble reading this file on your system, you should install a newer version of ClangFormat, and avoid commenting out parts of the configuration file.

ClangFormat is the primary tool that helps us achieve a uniform look of Ginkgo's codebase, while reducing the learning curve of potential contributors. However, ClangFormat configuration is not expressive enough to incorporate the entire coding style, so there are several additional rules that all contributed code should follow.

 $\it Note$: To learn more about how ClangFormat will format your code, see existing files in Ginkgo, .clang-format configuration file shipped with Ginkgo, and ClangFormat's documentation.

5.5.2 Naming scheme

5.5.2.1 Filenames

Filenames use snake_case and use the following extensions:

• C++ source files: .cpp

C++ header files: .hpp

5.5 Code style 21

- CUDA source files: .cu
- CUDA header files: .cuh
- HIP source files: .hip.cpp
- HIP header files: .hip.hpp
- Common source files used by both CUDA and HIP: .hpp.inc
- CMake utility files: .cmake
- Shell scripts: .sh

Note: A C++ source/header file is considered a CUDA file if it contains CUDA code that is not guarded with #if guards that disable this code in non-CUDA compilers. I.e. if a file can be compiled by a general C++ compiler, it is not considered a CUDA file.

5.5.2.2 Macros

Macros (both object-like and function-like macros) use CAPITAL_CASE. They have to start with GKO_ to avoid name clashes (even if they are #undef-ed in the same file!).

5.5.2.3 Variables

Variables use snake_case.

5.5.2.4 Constants

Constants use snake_case.

5.5.2.5 Functions

Functions use snake_case.

5.5.2.6 Structures and classes

Structures and classes which do not experience polymorphic behavior (i.e. do not contain virtual methods, nor members which experience polymorphic behavior) use snake case.

All other structures and classes use CamelCase.

5.5.2.7 **Members**

All structure / class members use the same naming scheme as they would if they were not members:

- · methods use the naming scheme for functions
- · data members the naming scheme for variables or constants
- type members for classes / structures

Additionally, non-public data members end with an underscore (_).

5.5.2.8 Namespaces

Namespaces use snake_case.

5.5.2.9 Template parameters

- Type template parameters use CamelCase, for example ValueType.
- Non-type template parameters use snake_case, for example subwarp_size.

5.5.3 Whitespace

Spaces and tabs are handled by ClangFormat, but blank lines are only partially handled (the current configuration doesn't allow for more than 2 blank lines). Thus, contributors should be aware of the following rules for blank lines:

- Top-level statements and statements directly within namespaces are separated with 2 blank lines. The first
 / last statement of a namespace is separated by two blank lines from the opening / closing brace of the
 namespace.
 - (a) exception: if the first **or** the last statement in the namespace is another namespace, then no blank lines are required example: ```c++ namespace foo {

```
struct x {
    };
    } // namespace foo
    namespace bar {
    namespace baz {
    void f();
       // namespace baz
    } // namespace bar
2. _exception_: in header files whose only purpose is to _declare_ a bunch
    of functions (e.g. the '*_kernel.hpp' files) these declarations can be
    separated by only 1 blank line (note: standard rules apply for all other
    statements that might be present in that file)
3. _exception_: "related" statement can have 1 blank line between them.

"Related" is not a strictly defined adjective in this sense, but is in
    general one of:
    1. overload of a same function,
        function / class template and it's specializations,
    3. macro that modifies the meaning or adds functionality to the % \left( 1\right) =\left( 1\right) \left( 1\right) 
         previous / following statement.
    However, simply calling function 'f' from function 'g' does not imply
    that 'f' and 'g' are "related".
```

1. Statements within structures / classes are separated with 1 blank line. There are no blank lines between the first / last statement in the structure / class.

5.5 Code style 23

```
(a) exception: there is no blank line between an access modifier (private, protected, public) and
the following statement. example: ```c++ class foo { public: int get_x() const noexcept { return x_; }
int &get_x() noexcept { return x_; }
private: int x_; }; ```
```

- 2. Function bodies cannot have multiple consecutive blank lines, and a single blank line can only appear between two logical sections of the function.
- 3. Unit tests should follow the AAA pattern, and a single blank line must appear between consecutive "A" sections. No other blank lines are allowed in unit tests.
- 4. Enumeration definitions should have no blank lines between consecutive enumerators.

5.5.4 Include statement grouping

In general, all include statements should be present on the top of the file, ordered in the following groups, with two blank lines between each group:

- Related header file (e.g. core/foo/bar.hpp included in core/foo/bar.cpp, or in the unit testcore/test/foo/bar.cpp)
- 2. Standard library headers (e.g. vector)
- 3. Executor specific library headers (e.g. omp.h)
- 4. System third-party library headers (e.g. papi.h)
- 5. Local third-party library headers
- 6. Public Ginkgo headers
- 7. Private Ginkgo headers

Example: A file core/base/my_file.cpp might have an include list like this:

```
{c++}
#include <ginkgo/core/base/my_file.hpp>
#include <algorithm>
#include <vector>
#include <omp.h>
#include <omp.h>
#include "third_party/blas/cblas.hpp"
#include "third_party/lapack/lapack.hpp"
#include <ginkgo/core/base/executor.hpp>
#include <ginkgo/core/base/lin_op.hpp>
#include <ginkgo/core/base/types.hpp>
#include <core/base/my_file_kernels.hpp"</pre>
```

5.5.4.1 Main header

This section presents general rules used to define the main header attributed to the file. In the previous example, this would be $\#include < ginkgo/core/base/my_file.hpp>$.

General rules:

- 1. Some fixed main header.
- 2. components:
 - with _kernel suffix looks for the header in the same folder.

- without _kernel suffix looks for the header in core.
- 3. test/utils: looks for the header in core
- 4. core: looks for the header in ginkgo
- 5. test or base: looks for the header in ginkgo/core
- 6. others: looks for the header in core

Note: Please see the detail in the dev_tools/scripts/config.

5.5.4.2 Some general comments.

- 1. Private headers of Ginkgo should not be included within the public Ginkgo header.
- 2. It is a good idea to keep the headers self-sufficient, See Google Style guide for reasoning. When compiling with GINKGO_CHECK_CIRCULAR_DEPS enabled, this property is explicitly checked.
- 3. The recommendations of the <code>iwyu</code> (Include what you use) tool can be used to make sure that the headers are self-sufficient and that the compiled files (<code>.cu,.cpp,.hip.cpp</code>) include only what they use. A <code>CI pipeline</code> is available that runs with the <code>iwyu</code> tool. Please be aware that this tool can be incorrect in some cases.

5.5.4.3 Automatic header arrangement

- 1. dev_tools/script/format_header.sh will take care of the group/sorting of headers according to this guideline.
- 2. make format_header arranges the header of the modified files in the branch.
- 3. make format_header_all arranges the header of all files.

5.5.5 Other Code Formatting not handled by ClangFormat

5.5.5.1 Control flow constructs

Single line statements should be avoided in all cases. Use of brackets is mandatory for all control flow constructs (e.g. if, for, while, ...).

5.5.5.2 Variable declarations

C++ supports declaring / defining multiple variables using a single *type-specifier*. However, this is often very confusing as references and pointers exhibit strange behavior:

For this reason, always declare each variable on a separate line, with its own type-specifier.

5.6 Helper scripts 25

5.5.6 CMake coding style

5.5.6.1 Whitespaces

All alignment in CMake files should use four spaces.

5.5.6.2 Use of macros vs functions

Macros in CMake do not have a scope. This means that any variable set in this macro will be available to the whole project. In contrast, functions in CMake have local scope and therefore all set variables are local only. In general, wrap all piece of algorithms using temporary variables in a function and use macros to propagate variables to the whole project.

5.5.6.3 Naming style

All Ginkgo specific variables should be prefixed with a GINKGO_ and all functions by ginkgo_.

5.6 Helper scripts

To facilitate easy development within Ginkgo and to encourage coders and scientists who do not want get bogged down by the details of the Ginkgo library, but rather focus on writing the algorithms and the kernels, Ginkgo provides the developers with a few helper scripts.

5.6.1 Create a new algorithm

A create_new_algorithm.sh script is available for developers to facilitate easy addition of new algorithms. The options it provides can be queried with ./create_new_algorithm.sh --help

The main objective of this script is to add files and boiler plate code for the new algorithm using a model and an instance of that model. For example, models can be any one of factorization, matrix, preconditioner or solver. For example to create a new solver named my_solver similar to gmres, you would set the ModelType to solver and set the ModelName to gmres. This would duplicate the core algorithm and kernels of the gmres algorithm and replace the naming to my_solver. Additionally, all the kernels of the new my_\circ solver are marked as GKO_NOT_IMPLEMENTED. For easy navigation and .txt file is created in the folder where the script is run, which lists all the TODO's. These TODO's can also be found in the corresponding files.

5.6.2 Converting CUDA code to HIP code

We provide a cuda2hip script that converts cuda kernel code into hip kernel code. Internally, this script calls the hipify script provided by HIP, converting the CUDA syntax to HIP syntax. Additionally, it also automatically replaces the instances of CUDA with HIP as appropriate. Hence, this script can be called on a Ginkgo CUDA file. You can find this script in the dev_tools/scripts/ folder.

5.7 Writing Tests

Ginkgo uses the GTest framework for the unit test framework within Ginkgo. Writing good tests are extremely important to verify the functionality of the new code and to make sure that none of the existing code has been broken.

5.7.1 Testing know-how

- GTest provides a comprehensive documentation of the functionality available within Gtest.
- Reduce code duplication with Testing Fixtures, TEST_F
- Write templated tests using TYPED_TEST.

5.7.2 Some general rules.

- Unit tests must follow the KISS principle.
- Unit tests must follow the AAA pattern, and a single blank line must appear between consecutive "A" sections.

5.7.3 Writing tests for kernels

- Reference kernels, kernels on the ReferenceExecutor, are meant to be single threaded reference implementations. Therefore, tests for reference kernels need to be performed with data that can be as small as possible. For example, matrices lesser than 5x5 are acceptable. This allows the reviewers to verify the results for exactness with tools such as MATLAB.
- OpenMP, CUDA, HIP and DPC++ kernels have to be tested against the reference kernels. Hence data for the tests of these kernels can be generated in the test files using helper functions or by using external files to be read through the standard input. In particular for CUDA, HIP and DPC++ the data size should be at least bigger than the architecture's warp size to ensure there is no corner case in the kernels.

5.8 Documentation style

Documentation uses standard Doxygen.

5.8.1 Developer targeted notes

Make use of @internal doxygen tag. This can be used for any comment which is not intended for users, but is useful to better understand a piece of code.

5.8.2 Whitespaces

5.8.2.1 After named tags such as <tt>@param foo</tt>

The documentation tags which use an additional name should be followed by two spaces in order to better distinguish the text from the doxygen tag. It is also possible to use a line break instead.

5.8.3 Documenting examples

There are two main steps:

- 1. First, you can just copy over the doc/ folder (you can copy it from the example most relevant to you) and adapt your example names and such, then you can modify the actual documentation.
- In tooltip: A short description of the example.
- In short-intro: The name of the example.
- In results.dox: Run the example and write the output you get.
- In kind: The kind of the example. For different kinds see the documentation. Examples can be of basic, techniques, logging, stopping_criteria or preconditioners. If your example does not fit any of these categories, feel free to create one.
- In intro.dox: You write an explanation of your code with some introduction similar to what you see in an existing example most relevant to you.
- In builds-on: You write the examples it builds on.
- 1. You also need to modify the examples.hpp.in file. You add the name of the example in the main section and in the section that you specified in the doc/kind file in the example documentation.

5.9 Other programming comments

5.9.1 C++ standard stream objects

These are global objects and are shared inside the same translation unit. Therefore, whenever its state or formatting is changed (e.g. using std::hex or floating point formatting) inside library code, make sure to restore the state before returning the control to the user. See this stackoverflow question for examples on how to do it correctly. This is extremely important for header files.

5.9.2 Warnings

By default, the <code>-DGINKGO_COMPILER_FLAGS</code> is set to <code>-Wpedantic</code> and hence pedantic warnings are emitted by default. Some of these warnings are false positives and a complete list of the resolved warnings and their solutions is listed in <code>Issue 174</code>. Specifically, when macros are being used, we have the issue of having <code>extra</code>; warnings, which is resolved by adding a <code>static_assert()</code>. The CI system additionally also has a step where it compiles for pedantic warnings to be errors.

5.9.3 Avoiding circular dependencies

To facilitate finding circular dependencies issues (see Using library classes for more details), a CI step no-circular-deps was created. For more details on its usage, see this pipeline, where Ginkgo did not abide to this policy and PR #278 which fixed this. Note that doing so is not enough to guarantee with 100% accuracy that no circular dependency is present. For an example of such a case, take a look at this pipeline where one of the compiler setups detected an incorrect dependency of the cuda module (due to jacobi) on the core module.

Citing Ginkgo

The main Ginkgo paper describing Ginkgo's purpose, design and interface is available through the following reference:

```
@article{ginkgo-toms-2022,
title = {{Ginkgo: A Modern Linear Operator Algebra Framework for High Performance Computing}},
volume = {48},
copyright = {All rights reserved},
issn = {0098-3500},
shorttitle = {Ginkgo},
url = {https://doi.org/10.1145/3480935},
doi = {10.1145/3480935},
number = {1},
urldate = {2022-02-17},
journal = {ACM Transactions on Mathematical Software},
author = {Anzt, Hartwig and Cojean, Terry and Flegar, Goran and Göbel, Fritz and Grützmacher, Thomas and
Nayak, Pratik and Ribizel, Tobias and Tsai, Yuhsiang Mike and Quintana-Ortí, Enrique S.},
month = feb,
year = {2022},
keywords = {ginkgo, healthy software lifecycle, High performance computing, multi-core and manycore
architectures},
pages = {2:1--2:33}
}
```

Multiple topical papers exist on Ginkgo and its algorithms. The following papers can be used to cite specific aspects of the Ginkgo project.

6.0.1 The Ginkgo Software

The Ginkgo software itself was reviewed and has a paper published in the Journal of Open Source Software, which can be cited with the following reference:

```
@article{GinkgoJoss2020,
    doi = {10.21105/joss.02260},
    url = {https://doi.org/10.21105/joss.02260},
    year = {2020},
    publisher = {The Open Journal},
    volume = {5},
    number = {52},
    pages = {2260},
    author = {Hartwig Anzt and Terry Cojean and Yen-Chen Chen and Goran Flegar and Fritz G\"{0}bel and Thomas
        Gr\"{u}tzmacher and Pratik Nayak and Tobias Ribizel and Yu-Hsiang Tsai},
    title = {Ginkgo: A high performance numerical linear algebra library},
    journal = {Journal of Open Source Software}
}
```

6.0.2 On Portability

```
@misc{tsai2020amdportability,
    title={Preparing Ginkgo for AMD GPUs -- A Testimonial on Porting CUDA Code to HIP},
    author={Yuhsiang M. Tsai and Terry Cojean and Tobias Ribizel and Hartwig Anzt},
    year={2020},
    eprint={2006.14290},
    archivePrefix={arXiv},
    primaryClass={cs.MS}
}
```

30 Citing Ginkgo

6.0.3 On Software Sustainability

```
@inproceedings{anzt2019pasccb,
author = {Anzt, Hartwig and Chen, Yen-Chen and Cojean, Terry and Dongarra, Jack and Flegar, Goran and Nayak,
      Pratik and Quintana-Ort\'{\i}, Enrique S. and Tsai, Yuhsiang M. and Wang, Weichung},
title = {Towards Continuous Benchmarking: An Automated Performance Evaluation Framework for High
      Performance Software},
year = {2019},
isbn = {9781450367707},
publisher = {Association for Computing Machinery},
address = {New York, NY, USA},
url = {https://doi.org/10.1145/3324989.3325719},
doi = \{10.1145/3324989.3325719\},
booktitle = {Proceedings of the Platform for Advanced Scientific Computing Conference},
articleno = {9},
numpages = \{11\}.
keywords = {interactive performance visualization, healthy software lifecycle, continuous integration,
     automated performance benchmarking},
location = {Zurich, Switzerland},
series = {PASC '19}
```

6.0.4 On SpMV or solvers performance

```
@InProceedings{tsai2020amdspmv,
author="Tsai, Yuhsiang M.
and Cojean, Terry
and Anzt, Hartwig",
editor="Sadayappan, Ponnuswamy
and Chamberlain, Bradford L.
and Juckeland, Guido
and Ltaief, Hatem",
title="Sparse Linear Algebra on AMD and  NVIDIA GPUs -- The Race Is On",
booktitle="High Performance Computing",
year="2020",
publisher="Springer International Publishing",
address="Cham",
pages="309--327"
abstract="Efficiently processing sparse matrices is a central and performance-critical part of many
      scientific simulation codes. Recognizing the adoption of manycore accelerators in HPC, we evaluate in this paper the performance of the currently best sparse matrix-vector product (SpMV) implementations
      on high-end GPUs from AMD and NVIDIA. Specifically, we optimize SpMV kernels for the CSR, COO, ELL,
      and HYB format taking the hardware characteristics of the latest GPU technologies into account. We
      compare for 2,800 test matrices the performance of our kernels against AMD's hipSPARSE library and
      NVIDIA's cuSPARSE library, and ultimately assess how the GPU technologies from AMD and NVIDIA compare
in terms of SpMV performance.", isbn="978-3-030-50743-5"
@article{anzt2020spmv,
author = {Anzt, Hartwig and Cojean, Terry and Yen-Chen, Chen and Dongarra, Jack and Flegar, Goran and Nayak,
      Pratik and Tomov, Stanimire and Tsai, Yuhsiang M. and Wang, Weichung},
title = {Load-Balancing Sparse Matrix Vector Product Kernels on GPUs},
year = {2020},
issue_date = {March 2020},
publisher = {Association for Computing Machinery},
address = {New York, NY, USA},
volume = \{7\},
number = \{1\},
issn = \{2329-4949\},
url = {https://doi.org/10.1145/3380930},
doi = {10.1145/3380930},
journal = {ACM Trans. Parallel Comput.},
month = mar,
articleno = {2},
numpages = \{26\},
keywords = {irregular matrices, GPUs, Sparse Matrix Vector Product (SpMV)}
    title={Evaluating the Performance of NVIDIA's A100 Ampere GPU for Sparse Linear Algebra Computations},
    author={Yuhsiang Mike Tsai and Terry Cojean and Hartwig Anzt},
    vear={2020},
    eprint={2008.08478}.
    archivePrefix={arXiv},
    primaryClass={cs.MS}
```

Example programs

Here you can find example programs that demonstrate the usage of Ginkgo. Some examples are built on one another and some are stand-alone and demonstrate a concept of Ginkgo, which can be used in your own code.

You can browse the available example programs

- 1. as a graph that shows how example programs build upon each other.
- 2. as a list that provides a short synopsis of each program.
- 3. or grouped by topic.

By default, all Ginkgo examples are built using CMake.

An example for building the examples and using Ginkgo as an external library without CMake can be found in the script provided for each example, which should be called with the form: ./build.sh PATH_TO_GINKGO_B UILD_DIR

By default, Ginkgo is compiled with at least <code>-DGINKGO_BUILD_REFERENCE=ON</code>. Ginkgo also tries to detect your environment setup (presence of CUDA, ...) to enable the relevant accelerator modules. If you want to target a specific GPU, make sure that Ginkgo is compiled with the accelerator specific module enabled, such as:

- 1. -DGINKGO_BUILD_CUDA=ON option for NVIDIA GPUs.
- 2. -DGINKGO_BUILD_HIP=ON option for AMD or NVIDIA GPUs.
- 3. -DGINKGO_BUILD_DPCPP=ON option for Intel GPUs (and possibly any other platform).

Connections between example programs

The following graph shows the connections between example programs and how they build on each other. Click on any of the boxes to go to one of the programs. If you hover your mouse pointer over a box, a brief description of the program should appear.



32 Example programs

Legend:



Example programs

VIDIA GPU's. The poisson-solver program Solve an actual physically relevant problem, the poisson problem. The matrix is generated within Ginkgo. The preconditioned-solver program Using a Jacobi preconditioner to solve a linear system. The ilu-preconditioned-solver program Using an ILU preconditioner to solve a linear system. The performance-debugging program Using Loggers to debug the performance within Ginkgo. The three-pt-stencil-solver program Using a three point stencil to solve the poisson equation with array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process.	The simple-solver program	A minimal CG solver in Ginkgo, which reads a matrix from a file.		
lem. The matrix is generated within Ginkgo. The preconditioned-solver program Using a Jacobi preconditioner to solve a linear system. The ilu-preconditioned-solver program Using an ILU preconditioner to solve a linear system. The performance-debugging program Using Loggers to debug the performance within Ginkgo. The three-pt-stencil-solver program Using a three point stencil to solve the poisson equation with array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The minimal-cuda-solver program	A minimal solver on the CUDA executor than can be run on N $_{\!$		
The ilu-preconditioned-solver program The performance-debugging program Using Loggers to debug the performance within Ginkgo. The three-pt-stencil-solver program Using a three point stencil to solve the poisson equation with array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The poisson-solver program	Solve an actual physically relevant problem, the poisson problem. The matrix is generated within Ginkgo.		
The performance-debugging program Using Loggers to debug the performance within Ginkgo. The three-pt-stencil-solver program Using a three point stencil to solve the poisson equation with array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The preconditioned-solver program	Using a Jacobi preconditioner to solve a linear system.		
The three-pt-stencil-solver program Using a three point stencil to solve the poisson equation with array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The ilu-preconditioned-solver program	Using an ILU preconditioner to solve a linear system.		
array views. The nine-pt-stencil-solver program Using a nine point 2D stencil to solve the poisson equation with array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The performance-debugging program	Using Loggers to debug the performance within Ginkgo.		
array views. The external-lib-interfacing program Using Ginkgo's solver with the external library deal.II. The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The three-pt-stencil-solver program	Using a three point stencil to solve the poisson equation with array views.		
The custom-logger program Creating a custom logger specifically for comparing the recurrent and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The nine-pt-stencil-solver program	Using a nine point 2D stencil to solve the poisson equation with array views.		
and the real residual norms. The custom-matrix-format program Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The external-lib-interfacing program	Using Ginkgo's solver with the external library deal.II.		
methods to build your own custom matrix format. The inverse-iteration program Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The custom-logger program	Creating a custom logger specifically for comparing the recurrent and the real residual norms.		
verse iteration method. The simple-solver-logging program Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The custom-matrix-format program	Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format.		
information to diagnose and debug your code. The papi-logging program Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The inverse-iteration program	Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method.		
mation about your code and its behaviour. The ginkgo-overhead program Measuring the overhead of the Ginkgo library. The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The simple-solver-logging program	Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code.		
The custom-stopping-criterion program Creating a custom stopping criterion for the iterative solution process. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The papi-logging program	Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour.		
cess. The ginkgo-ranges program Using the ranges concept to factorize a matrix with the LU fac-	The ginkgo-overhead program	Measuring the overhead of the Ginkgo library.		
	The custom-stopping-criterion program	Creating a custom stopping criterion for the iterative solution process.		
	The ginkgo-ranges program	Using the ranges concept to factorize a matrix with the LU factorization.		

The mixed-spmv program	Shows the Ginkgo mixed precision spmv functionality.	
The mixed-precision-ir program	Manual implementation of a Mixed Precision Iterative Refinement (MPIR) solver.	
The adaptiveprecision-blockjacobi program	Shows how to use the adaptive precision block-Jacobi preconditioner.	
The cb-gmres program	Using the Ginkgo CB-GMRES solver (Compressed Basis GM↔ RES).	
The heat-equation program	Solving a 2D heat equation and showing matrix assembly, vector initalization and solver setup in a more complex setting with output visualization.	
The iterative-refinement program	Using a low accuracy CG solver as an inner solver to an iterative refinement (IR) method which solves a linear system.	
The ir-ilu-preconditioned-solver program	Combining iterative refinement with the adaptive precision block- Jacobi preconditioner to approximate triangular systems occur- ring in ILU preconditioning.	
The par-ilu-convergence program	Convergence analysis at the examples of parallel incomplete factorization solver.	
The preconditioner-export program	Explicit generation and storage of preconditioners for given matrices.	
The multigrid-preconditioned-solver program	Use multigrid as preconditioner to a solver.	
The mixed-multigrid-solver program	Use multigrid with different precision multigrid_level as a solver.	
The distributed-solver program	Use a distributed solver to solve a 1D Laplace equation.	

Example programs grouped by topics

The simple-solver program
The performance-debugging program
The preconditioner-export program
The minimal-cuda-solver program
The preconditioned-solver program,
The ilu-preconditioned-solver program,
The ir-ilu-preconditioned-solver program,
The adaptiveprecision-blockjacobi program,
The par-ilu-convergence program,
The preconditioner-export program
The multigrid-preconditioned-solver program
The iterative-refinement program,
The mixed-precision-ir program,
The ir-ilu-preconditioned-solver program
The poisson-solver program,
The three-pt-stencil-solver program,
The nine-pt-stencil-solver program,
The custom-matrix-format program

34 Example programs

Reading in a matrix and right hand side from a file	The simple-solver program,	
	The minimal-cuda-solver program,	
	The preconditioned-solver program,	
	The ilu-preconditioned-solver program,	
	The inverse-iteration program,	
	The simple-solver-logging program,	
	The papi-logging program,	
	The custom-stopping-criterion program,	
	The custom-logger program	

Basic techniques

Using Ginkgo with external libraries	The external-lib-interfacing program	
Customizing Ginkgo	The custom-logger program,	
	The custom-stopping-criterion program,	
	The custom-matrix-format program	
Writing your own matrix format	The custom-matrix-format program	
Using Ginkgo to construct more complex linear algebra routines	The inverse-iteration program	
Logging within Ginkgo	The simple-solver-logging program,	
	The papi-logging program,	
	The performance-debugging program	
	The custom-logger program	
Constructing your own stopping criterion	The custom-stopping-criterion program	
Using ranges in Ginkgo	The ginkgo-ranges program	
Mixed precision	The mixed-spmv program,	
	The mixed-precision-ir program,	
	The adaptiveprecision-blockjacobi program	
	The mixed-multigrid-solver program	
Multigrid	The multigrid-preconditioned-solver program	
	The mixed-multigrid-solver program	
Distributed	The distributed-solver program	

Advanced techniques

The adaptive precision-block jacobi program

The preconditioned solver example..

This example depends on preconditioned-solver.

This example shows how to use the adaptive precision block-Jacobi preconditioner.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. The preconditioned CG solver is enhanced with a block-Jacobi preconditioner that optimizes the storage format for the distinct inverted diagonal blocks to the numerical requirements. The example features the iteration count and runtime of the CG solver.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
   exec_map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
```

```
{"cuda",
         [] {
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         { "dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                  gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS and initial guess as 1
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
    host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_x);
Create solver factory
const RealValueType reduction_factor = 1e-7;
auto solver gen =
    cq::build()
        .with criteria(
             gko::stop::Iteration::build().with_max_iters(10000u).on(exec),
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_reduction_factor(reduction_factor)
                 .on(exec))
Add preconditioner, these 2 lines are the only difference from the simple solver example
.with_preconditioner(bj::build()
                           .with_max_block_size(16u)
                           . \verb|with_storage_optimization|| \\
                               gko::precision_reduction::autodetect())
                           .on(exec))
.on(exec);
Create solver
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
solver_gen->add_logger(logger);
auto solver = solver_gen->generate(A);
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
Get residual
auto res = gko::as<real_vec>(logger->get_residual_norm());
auto impl_res = gko::as<real_vec>(logger->get_implicit_sq_resnorm());
std::cout « "Initial residual norm sqrt(r^T r):\n";
```

Results

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
194.679
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
5.69384e-06
Implicit residual norm squared (r^2):
%%MatrixMarket matrix array real general
1 1
1.27043e-15
CG iteration count: 5
CG execution time [ms]: 0.080041
```

Comments about programming and debugging

The plain program

```
************GINKGO LICENSE>******************
 Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
 1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the \frac{1}{2}
documentation and/or other materials provided with the distribution.
 3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
  IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
  TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
 SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
 DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
  #include <ginkgo/ginkgo.hpp>
  #include <fstream>
  #include <iomanip>
  #include <iostream>
  #include <map>
  #include <string>
 int main(int argc, char* argv[])
```

}

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec_map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
         {"cuda",
          [] {
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         {"dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                    gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec-yet_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
    host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_x);
const RealValueType reduction_factor = 1e-7;
auto solver_gen =
    cg::build()
         .with_criteria(
             gko::stop::Iteration::build().with_max_iters(10000u).on(exec),
             gko::stop::ResidualNorm<ValueType>::build()
                  .with_reduction_factor(reduction_factor)
                  .on(exec))
         .with_preconditioner(bj::build()
                                     .with_max_block_size(16u)
                                     .with_storage_optimization(
                                        gko::precision_reduction::autodetect())
                                     .on(exec))
         .on(exec);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
solver_gen->add_logger(logger);
auto solver = solver_gen->generate(A);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::as<real_vec>(logger->get_residual_norm());
auto impl_res = gko::as<real_vec>(logger->get_implicit_sq_resnorm());
std::cout « "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
std::cout « "Implicit residual norm squared (r^2):\n";
write(std::cout, impl_res);
std::cout « "CG iteration count:
                                               " « logger->get_num_iterations()
          « std::endl;
std::cout « "CG execution time [ms]: "
           « static_cast<double>(time.count()) / 1000000.0 « std::endl;
```

The cb-gmres program

The CB-GMRES solver example..

Introduction

About the example

This example showcases the usage of the Ginkgo solver CB-GMRES (Compressed Basis GMRES). A small system is solved with two un-preconditioned CB-GMRES solvers:

- 1. without compressing the krylov basis; it uses double precision for both the matrix and the krylov basis, and
- 2. with a compression of the krylov basis; it uses double precision for the matrix and all arithmetic operations, while using single precision for the storage of the krylov basis

Both solves are timed and the residual norm of each solution is computed to show that both solutions are correct.

The commented program

This is the main ginkgo header file.

```
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <cmath>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
```

Helper function which measures the time of solver->apply(b, x) in seconds To get an accurate result, the solve is repeated multiple times (while ensuring the initial guess is always the same). The result of the solve will be written to x.

Make a copy of x, so we can re-use the same initial guess multiple times auto $x_{copy} = clone(x)$;

```
for (int i = 0; i < repeats; ++i) {</pre>
```

No need to copy it in the first iteration

```
if (i != 0) {
    x_copy->copy_from(x);
}
```

Make sure all previous executor operations have finished before starting the time

```
exec->synchronize();
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x_copy);
```

Make sure all computations are done before stopping the time

```
exec->synchronize();
  auto tac = std::chrono::steady_clock::now();
  duration += std::chrono::duration<double>(tac - tic).count();
```

Copy the solution back to x, so the caller has the result

```
x->copy_from(x_copy);
return duration / static_cast<double>(repeats);
}
int main(int argc, char* argv[])
{
```

Use some shortcuts. In Ginkgo, vectors are seen as a gko::matrix::Dense with one column/one row. The advantage of this concept is that using multiple vectors is a now a natural extension of adding columns/rows are necessary.

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
```

The gko::matrix::Csr class is used here, but any other matrix class such as gko::matrix::Coo, gko::matrix::Hybrid, gko::matrix::Ell or gko::matrix::Sellp could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The gko::solver::CbGmres is used here, but any other solver class can also be used.

```
using cb_gmres = gko::solver::CbGmres<ValueType>;
```

```
Print the ginkgo version information.
```

```
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
    std::exit(-1);
}
```

Map which generates the appropriate executor

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    \verb"exec_map" \{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                              true);
        { "dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Note: this matrix is copied from "SOURCE_DIR/matrices" instead of from the local directory. For details, see "examples/cb-gmres/CMakeLists.txt"

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create a uniform right-hand side with a norm2 of 1 on the host (norm2(b) == 1), followed by copying it to the actual executor (to make sure it also works for GPUs)

```
const auto A_size = A->get_size();
auto b_host = vec::create(exec->get_master(), gko::dim<2>{A_size[0], 1});
for (gko::size_type i = 0; i < A_size[0]; ++i) {
    b_host->at(i, 0) =
        ValueType{1} / std::sqrt(static_cast<ValueType>(A_size[0]));
}
auto b_norm = gko::initialize<real_vec>({0.0}, exec);
b_host->compute_norm2(b_norm);
auto b = clone(exec, b_host);
```

As an initial guess, use the right-hand side

```
auto x_keep = clone(b);
auto x_reduce = clone(x_keep);
const RealValueType reduction_factor{le-6};
```

Generate two solver factories: _keep uses the same precision for the krylov basis as the matrix, and _reduce uses one precision below it. If ValueType is double, then _reduce uses float as the krylov basis storage type

```
auto solver_gen_keep :
    cb_gmres::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1000u).on(exec),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_baseline(gko::stop::mode::rhs_norm)
                .with reduction factor(reduction factor)
                .on(exec))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb_gmres::storage_precision::keep)
        .on(exec);
auto solver_gen_reduce
    cb_gmres::build()
        .with criteria (
            gko::stop::Iteration::build().with_max_iters(1000u).on(exec),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_baseline(gko::stop::mode::rhs_norm)
                .with_reduction_factor(reduction_factor)
                .on(exec))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb_gmres::storage_precision::reduce1)
        .on(exec);
```

Generate the actual solver from the factory and the matrix.

```
auto solver_keep = solver_gen_keep->generate(A);
auto solver_reduce = solver_gen_reduce->generate(A);
```

Solve both system and measure the time for each.

Make sure the output is in scientific notation for easier comparison

```
std::cout « std::scientific;
```

Note: The time might not be significantly different since the matrix is quite small

To measure if your solution has actually converged, the error of the solution is measured. one, neg_one are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, the (advanced) apply method is used.

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res_norm_keep = gko::initialize<real_vec>({0.0}, exec);
auto res_norm_reduce = gko::initialize<real_vec>({0.0}, exec);
auto tmp = gko::clone(b);

tmp = Ax - tmp
    A->apply(one, x_keep, neg_one, tmp);
    tmp->compute_norm2(res_norm_keep);
    std::cout « "\nResidual norm without compression:\n";
    write(std::cout, res_norm_keep);
    tmp->copy_from(b);
    A->apply(one, x_reduce, neg_one, tmp);
    tmp->compute_norm2(res_norm_reduce);
    std::cout « "\nResidual norm with compression:\n";
    write(std::cout, res_norm_reduce);
}
```

Results

The following is the expected result: Solve time without compression: 1.842690e-04 s Solve time with compression: 1.589936e-04 s Residual norm without compression: %*MatrixMarket matrix array real general 1 1 2.430544e-07 Residual norm with compression: %*MatrixMarket matrix array real general 1 1 3.437257e-07

Comments about programming and debugging

The plain program

```
************* GINKGO LICENSE>******************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <cmath>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
double measure_solve_time_in_s(std::shared_ptr<const gko::Executor> exec,
                                 gko::LinOp* solver, const gko::LinOp* b,
                                 gko::LinOp* x)
    constexpr int repeats{5};
    double duration {0};
    auto x_copy = clone(x);
for (int i = 0; i < repeats; ++i) {</pre>
        if (i != 0) {
            x_copy->copy_from(x);
         exec->synchronize();
        auto tic = std::chrono::steady_clock::now();
        solver->apply(b, x_copy);
        exec->synchronize();
        auto tac = std::chrono::steady_clock::now();
        duration += std::chrono::duration<double>(tac - tic).count();
    x->copy_from(x_copy);
    return duration / static_cast<double>(repeats);
int main(int argc, char* argv[])
```

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cb_gmres = gko::solver::CbGmres<ValueType>;
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                 true);
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         }},
        { "dpcpp",
              return gko::DpcppExecutor::create(0,
                                                  gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
const auto A_size = A->get_size();
\verb"auto b_host = vec::create(exec->get_master(), gko::dim<2>{A_size[0], 1});
for (gko::size_type i = 0; i < A_size[0]; ++i) {
b_host->at(i, 0) =
        ValueType{1} / std::sqrt(static_cast<ValueType>(A_size[0]));
auto b_norm = gko::initialize<real_vec>({0.0}, exec);
b_host->compute_norm2(b_norm);
auto b = clone(exec, b_host);
auto x_keep = clone(b);
auto x_reduce = clone(x_keep);
const RealValueType reduction_factor{1e-6};
auto solver_gen_keep =
    cb_gmres::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1000u).on(exec),
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_baseline(gko::stop::mode::rhs_norm)
                 .with_reduction_factor(reduction_factor)
                 .on(exec))
        .with_krylov_dim(100u)
        .with storage precision(
            gko::solver::cb_gmres::storage_precision::keep)
        .on(exec);
auto solver_gen_reduce =
    cb_gmres::build()
        .with criteria(
            gko::stop::Iteration::build().with_max_iters(1000u).on(exec),
            gko::stop::ResidualNorm<ValueType>::build()
                 .with_baseline(gko::stop::mode::rhs_norm)
                 .with_reduction_factor(reduction_factor)
                 .on(exec))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb gmres::storage precision::reduce1)
        .on(exec);
auto solver_keep = solver_gen_keep->generate(A);
auto solver_reduce = solver_gen_reduce->generate(A);
auto time_keep =
std::cout « std::scientific;
std::cout « "Solve time without compression: " « time_keep « " s\n" « "Solve time with compression: " « time_reduce « "
                                                     " « time_reduce « " s\n";
auto one = gko::initialize<vec>((1.0), exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto res_norm_reduce = gko::initialize<real_vec>({0.0}, exec);
auto res_norm_reduce = gko::initialize<real_vec>({0.0}, exec);
auto tmp = gko::clone(b);
A->apply(one, x_keep, neg_one, tmp);
tmp->compute_norm2(res_norm_keep);
std::cout « "\nResidual norm without compression:\n";
```

```
write(std::cout, res_norm_keep);
tmp->copy_from(b);
A->apply(one, x_reduce, neg_one, tmp);
tmp->compute_norm2(res_norm_reduce);
std::cout « "\nResidual norm with compression:\n";
write(std::cout, res_norm_reduce);
```

The custom-logger program

The simple solver with a custom logger example..

This example depends on simple-solver, simple-solver-logging, minimal-cuda-solver.

Introduction

The custom-logger example shows how Ginkgo's API can be leveraged to implement application-specific callbacks for Ginkgo's events. This is the most basic way of extending Ginkgo and a good first step for any application developer who wants to adapt Ginkgo to his specific needs.

Ginkgo's gko::log::Logger abstraction provides hooks to the events that happen during the library execution. These hooks concern any low-level event such as memory allocations, deallocations, copies and kernel launches up to high-level events such as linear operator applications and completion of solver iterations.

In this example, a simple logger is implemented to track the solver's recurrent residual norm and compute the true residual norm. At the end of the solver execution, a comparison table is shown on-screen.

About the example

Each example has the following sections:

- 1. **Introduction:**This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
- 2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
- 3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
- 4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

Include files

This is the main ginkgo header file.

```
#include <ginkgo/ginkgo.hpp>
```

Add the fstream header to read from data from files.

```
#include <fstream>
```

Add the map header for storing the executor map.

```
#include <man>
```

Add the C++ iomanip header to prettify the output.

```
#include <iomanip>
```

Add formatting flag modification capabilities.

```
#include <ios>
```

Add the C++ iostream header to output information to the console.

```
#include <iostream>
```

Add the string manipulation header to handle strings.

```
#include <string>
```

Add the vector header for storing the logger's data

```
#include <vector>
```

Utility function which returns the first element (position [0, 0]) from a given gko::matrix::Dense matrix / vector.

```
template <typename ValueType>
ValueType get_first_element(const gko::matrix::Dense<ValueType>* mtx)
{
```

Copy the matrix / vector to the host device before accessing the value in case it is stored in a GPU.

```
return mtx->get_executor()->copy_val_to_host(mtx->get_const_values());
```

Utility function which computes the norm of a Ginkgo gko::matrix::Dense vector.

```
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(
    const gko::matrix::Dense<ValueType>* b)
{
```

Get the executor of the vector

```
auto exec = b->get_executor();
```

Initialize a result scalar containing the value 0.0.

Use the dense <code>compute_norm2</code> function to compute the norm.

```
b->compute_norm2(b_norm);
```

Use the other utility function to return the norm contained in b_norm

```
return get_first_element(b_norm.get());
}
```

Custom logger class which intercepts the residual norm scalar and solution vector in order to print a table of real vs recurrent (internal to the solvers) residual norms.

```
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    using RealValueType = gko::remove_complex<ValueType>;
```

Output the logger's data in a table format

```
void write()const
```

```
Print a header for the table
```

Print a separation line. Note that for creating 10 characters std::setw() should be set to 11.

Print the data one by one in the form

std::defaultfloat could be used here but some compilers do not support it properly, e.g. the Intel compiler std::cout.unsetf(std::ios_base::floatfield);

Print a separation line

Customize the logging hook which is called everytime an iteration is completed

If the solver shares a residual norm, log its value

```
if (residual_norm) {
   auto dense_norm = gko::as<gko_real_dense>(residual_norm);
```

Add the norm to the recurrent_norms vector

```
recurrent_norms.push_back(get_first_element(dense_norm));
```

Otherwise, use the recurrent residual vector

```
} else {
   auto dense_residual = gko::as<gko_dense>(residual);
```

Compute the residual vector's norm

```
auto norm = compute_norm(dense_residual);
```

Add the computed norm to the recurrent_norms vector

```
recurrent_norms.push_back(norm);
```

If the solver shares the current solution vector

```
if (solution) {
```

Extract the matrix from the solver

Store the matrix's executor

```
auto exec = matrix->get_executor();
```

Create a scalar containing the value 1.0

```
auto one = gko::initialize<gko_dense>({1.0}, exec);
```

Create a scalar containing the value -1.0 auto neg_one = gko::initialize<gko_dense>({-1.0}, exec); Instantiate a temporary result variable auto res = gko::as<gko_dense>(gko::clone(b)); Compute the real residual vector by calling apply on the system matrix matrix->apply(one, solution, neg_one, res); Compute the norm of the residual vector and add it to the real norms vector real_norms.push_back(compute_norm(res.get())); } else { Add to the real_norms vector the value -1.0 if it could not be computed real_norms.push_back(-1.0); if (implicit_sq_residual_norm) { auto dense_norm = gko::as<gko_real_dense>(implicit_sq_residual_norm); Add the norm to the implicit_norms vector implicit_norms.push_back(std::sqrt(get_first_element(dense_norm))); } else { Add to the implicit_norms vector the value -1.0 if it could not be computed implicit_norms.push_back(-1.0); Add the current iteration number to the iterations vector iterations.push_back(iteration); Construct the logger ResidualLogger() : gko::log::Logger(gko::log::Logger::iteration complete mask) {} private: Vector which stores all the recurrent residual norms mutable std::vector<RealValueType> recurrent_norms{}; Vector which stores all the real residual norms mutable std::vector<RealValueType> real_norms{};

Vector which stores all the implicit residual norms

mutable std::vector<RealValueType> implicit_norms{};

Vector which stores all the iteration numbers

```
mutable std::vector<std::size_t> iterations{};
int main(int argc, char* argv[])
```

Use some shortcuts. In Ginkgo, vectors are seen as a gko::matrix::Dense with one column/one row. The advantage of this concept is that using multiple vectors is a now a natural extension of adding columns/rows are necessary.

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
```

The gko::matrix::Csr class is used here, but any other matrix class such as gko::matrix::Coo, gko::matrix::Hybrid, gko::matrix::Ell or gko::matrix::Sellp could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The gko::solver::Cg is used here, but any other solver class can also be used.

```
using cg = gko::solver::Cg<ValueType>;
```

Print the ginkgo version information.

```
std::cout « gko::version_info::get() « std::endl;
```

Where do you want to run your solver?

The gko::Executor class is one of the cornerstones of Ginkgo. Currently, we have support for an gko::OmpExecutor, which uses OpenMP multi-threading in most of its kernels, a gko::ReferenceExecutor, a single threaded specialization of the OpenMP executor and a gko::CudaExecutor which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/routines within Ginkgo should automatically work and run on the executor with any other changes.

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
         [] {
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                  true);
         }},
         {"hip",
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                 true);
         {"dpcpp",
         [] {
              return gko::DpcppExecutor::create(0,
                                                   gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Reading your data and transfer to the proper device.

Read the matrix, right hand side and the initial solution using the read function.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. gko::share and gko::give are the functions that you would need to use.

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor = 1e-7;
```

Creating the solver

Generate the gko::solver factory. Ginkgo uses the concept of Factories to build solvers with certain properties. Observe the Fluent interface used here. Here a cg solver is generated with a stopping criteria of maximum iterations of 20 and a residual norm reduction of 1e-15. You also observe that the stopping criteria(gko::stop) are also generated from factories using their build methods. You need to specify the executors which each of the object needs to be built on.

```
auto solver_gen 
cg::build()
```

```
.with_criteria(
    gko::stop::Iteration::build().with_max_iters(20u).on(exec),
    gko::stop::ResidualNorm<ValueType>::build()
        with_reduction_factor(reduction_factor)
        on(exec))
.on(exec):
```

Instantiate a ResidualLogger logger.

```
auto logger = std::make_shared<ResidualLogger<ValueType»();</pre>
```

Add the previously created logger to the solver factory. The logger will be automatically propagated to all solvers created from this factory.

```
solver_gen->add_logger(logger);
```

Generate the solver from the matrix. The solver factory built in the previous step takes a "matrix" (a gko::LinOp to be more general) as an input. In this case we provide it with a full matrix that we previously read, but as the solver only effectively uses the apply() method within the provided "matrix" object, you can effectively create a gko::LinOp class with your own apply implementation to accomplish more tasks. We will see an example of how this can be done in the custom-matrix-format example

```
auto solver = solver_gen->generate(A);
```

Finally, solve the system. The solver, being a gko::LinOp, can be applied to a right hand side, b to obtain the solution, x.

```
solver->apply(b, x);
```

Print the solution to the command line.

```
std::cout « "Solution (x):\n";
write(std::cout, x);
```

Print the table of the residuals obtained from the logger

logger->write();

To measure if your solution has actually converged, you can measure the error of the solution. one, neg_one are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, all you need to do is call the apply method, which in this case is an spmv and equivalent to the LAPACK z_spmv routine. Finally, you compute the euclidean 2-norm with the compute_norm2 function.

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

Results

The following is the expected result:

```
Solution (x):
%%MatrixMarket matrix array real general
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
```

Iteration	Recurrent Residual Norm	True Residual Norm	Implicit Residual Norm
01	4.358899e+00	4.358899e+00	4.358899e+001
11	·	2.304548e+00	
2	1.467706e+00		1.467706e+001
3	9.848751e-01	9.848751e-01	9.848751e-01
4	7.418330e-01	7.418330e-01	7.418330e-01
51	5.136231e-01	5.136231e-01	5.136231e-011
61	3.841650e-01	3.841650e-01	3.841650e-01
7	3.164394e-01	3.164394e-01	3.164394e-01
8	2.277088e-01		
9	1.703121e-01		1.703121e-01
10	9.737220e-02	9.737220e-02	9.737220e-02
11	6.168306e-02	6.168306e-02	6.168306e-02
12	4.541231e-02	4.541231e-02	4.541231e-02
13	3.195304e-02	3.195304e-02	3.195304e-02
14	1.616058e-02		
15	6.570152e-03	6.570152e-03	6.570152e-03
16	2.643669e-03	2.643669e-03	2.643669e-03
17	8.588089e-04	8.588089e-04	8.588089e-04
18	2.864613e-04	2.864613e-04	2.864613e-04
19	1.641952e-15	2.107881e-15	1.641952e-15
		-	
Residual nor	m sqrt(r^T r):		
%MatrixMark	et matrix array real gener	al	
L 1	-		
2.10788e-15			

Comments about programming and debugging

The plain program

```
*************************************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
   Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <map>
#include <iomanip>
#include <ios>
#include <iostream>
#include <string>
#include <vector>
template <typename ValueType>
ValueType get_first_element(const gko::matrix::Dense<ValueType>* mtx)
    return mtx->get_executor()->copy_val_to_host(mtx->get_const_values());
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(
    const gko::matrix::Dense<ValueType>* b)
```

```
auto exec = b->get_executor();
    auto b_norm =
        gko::initialize<gko::matrix::Dense<gko::remove_complex<ValueType>>
           {0.0}, exec);
    b->compute_norm2(b_norm);
    return get_first_element(b_norm.get());
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    using RealValueType = gko::remove_complex<ValueType>;
    void write()const
        std::cout « "Recurrent vs true vs implicit residual norm:"
                  « std::endl;
        std::cout « '|' « std::setw(10) « "Iteration" « '|' « std::setw(25)
                  « "Recurrent Residual Norm" « '|' « std::setw(25)
« "True Residual Norm" « '|' « std::setw(25)
« "Implicit Residual Norm" « '|' « std::endl;
        std::cout « ' | ' « std::setfill('-') « std::setw(11) «
                  std::cout « std::scientific;
        for (std::size_t i = 0; i < iterations.size(); i++) {
   std::cout « '|' « std::setw(10) « iterations[i] « '|'</pre>
                      « implicit_norms[i] « '|' « std::endl;
        std::cout.unsetf(std::ios_base::floatfield);
        using gko_dense = gko::matrix::Dense<ValueType>;
    using gko_real_dense = gko::matrix::Dense<RealValueType>;
    void on_iteration_complete(const gko::LinOp* solver, const gko::LinOp* b,
                                const gko::LinOp* solution,
const gko::size_type& iteration,
                                const gko::LinOp* residual,
                                const gko::LinOp* residual_norm,
                                const gko::LinOp* implicit_sq_residual_norm,
                                const gko::array<gko::stopping_status>*,
                                bool)const override
{
        if (residual_norm) {
            auto dense_norm = gko::as<gko_real_dense>(residual_norm);
            recurrent_norms.push_back(get_first_element(dense_norm));
        } else {
            auto dense_residual = gko::as<gko_dense>(residual);
            auto norm = compute norm(dense residual);
            recurrent_norms.push_back(norm);
        if (solution) {
            auto matrix = gko::as<gko::solver::detail::SolverBaseLinOp>(solver)
                               ->get_system_matrix();
            auto exec = matrix->get_executor();
            auto one = gko::initialize<gko_dense>({1.0}, exec);
            auto neg_one = gko::initialize<gko_dense>({-1.0}, exec);
            auto res = gko::as<gko_dense>(gko::clone(b));
            matrix->apply(one, solution, neg_one, res);
            real_norms.push_back(compute_norm(res.get()));
        } else {
            real_norms.push_back(-1.0);
        if (implicit_sq_residual_norm) {
            auto dense_norm =
                gko::as<gko_real_dense>(implicit_sq_residual_norm);
            implicit_norms.push_back(std::sqrt(get_first_element(dense_norm)));
        } else {
            implicit_norms.push_back(-1.0);
        iterations.push_back(iteration);
    ResidualLogger()
        : gko::log::Logger(gko::log::Logger::iteration complete mask)
    { }
    mutable std::vector<RealValueType> recurrent_norms{};
   mutable std::vector<RealValueType> real_norms{};
mutable std::vector<RealValueType> implicit_norms{};
   mutable std::vector<std::size t> iterations{};
int main(int argc, char* argv[])
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
```

```
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec_map{
         ("omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
          [] {
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                      true);
         {"hip",
          [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                       gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor = 1e-7;
auto solver_gen =
    cg::build()
         .with_criteria(
              gko::stop::Iteration::build().with_max_iters(20u).on(exec),
              gko::stop::ResidualNorm<ValueType>::build()
                  .with_reduction_factor(reduction_factor)
                   .on(exec))
         .on(exec);
auto logger = std::make_shared<ResidualLogger<ValueType»();</pre>
solver_gen->add_logger(logger);
auto solver = solver_gen->generate(A);
solver->apply(b, x);
std::cout « "Solution (x):\n";
write(std::cout, x);
logger->write();
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

The custom-matrix-format program

The custom matrix format example..

This example depends on simple-solver, poisson-solver, three-pt-stencil-solver, .

Introduction

This example solves a 1D Poisson equation:

$$\begin{aligned} u:[0,1] &\to R \\ u'' &= f \\ u(0) &= u0 \\ u(1) &= u1 \end{aligned}$$

using a finite difference method on an equidistant grid with ${\tt K}$ discretization points (${\tt K}$ can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

For an equidistant grid with K "inner" discretization points x1,...,xk,and step size h=1/(K+1), the formula produces a system of linear equations

$$2u_1 - u_2 = -f_1h^2 + u0$$

- $u(k-1) + 2u_k - u(k+1) = -f_kh^2, k = 2, ..., K-1$
- $u(K-1) + 2u_K = -f_Kh^2 + u1$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f'is set to 'f(x) = 6x' (making the solution ' $u(x) = x^3$ '), but that can be changed in the main function.

The intention of this example is to show how a custom linear operator can be created and integrated into Ginkgo to achieve performance benefits.

About the example

The commented program

```
#include <iostream>
#include <map>
#include <string>
#include <omp.h>
#include <qinkgo/ginkgo.hpp>
```

A CUDA kernel implementing the stencil, which will be used if running on the CUDA executor. Unfortunately, NVCC has serious problems interpreting some parts of Ginkgo's code, so the kernel has to be compiled separately.

A stencil matrix class representing the 3pt stencil linear operator. We include the gko::EnableLinOp mixin which implements the entire LinOp interface, except the two apply_impl methods, which get called inside the default implementation of apply (after argument verification) to perform the actual application of the linear operator. In addition, it includes the implementation of the entire PolymorphicObject interface.

It also includes the gko::EnableCreateMethod mixin which provides a default implementation of the static create method. This method will forward all its arguments to the constructor to create the object, and return an stdc::unique_ptr to the created object.

This constructor will be called by the create method. Here we initialize the coefficients of the stencil.

Here we implement the application of the linear operator, x = A * b. apply_impl will be called by the apply method, after the arguments have been moved to the correct executor and the operators checked for conforming sizes.

For simplicity, we assume that there is always only one right hand side and the stride of consecutive elements in the vectors is 1 (both of these are always true in this example).

```
void apply_impl(const gko::LinOp* b, gko::LinOp* x)const override
{
```

we only implement the operator for dense RHS. gko::as will throw an exception if its argument is not Dense.

```
auto dense_b = gko::as<vec>(b);
auto dense_x = gko::as<vec>(x);
```

we need separate implementations depending on the executor, so we create an operation which maps the call to the correct implementation

OpenMP implementation

```
void run(std::shared_ptr<const gko::OmpExecutor>)const override
{
    auto b_values = b->get_const_values();
    auto x_values = x->get_values();

#pragma omp parallel for
    for (std::size_t i = 0; i < x->get_size()[0]; ++i) {
        auto coefs = coefficients.get_const_data();
        auto result = coefs[1] * b_values[i];
        if (i > 0) {
            result += coefs[0] * b_values[i - 1];
        }
}
```

```
}
if (i < x->get_size()[0] - 1) {
    result += coefs[2] * b_values[i + 1];
}
x_values[i] = result;
}
```

CUDA implementation

We do not provide an implementation for reference executor. If not provided, Ginkgo will use the implementation for the OpenMP executor when calling it in the reference executor.

```
const coef_type& coefficients;
const vec* b;
vec* x;
);
this->get_executor()->run(
   stencil_operation(coefficients, dense_b, dense_x));
```

There is also a version of the apply function which does the operation x = alpha * A * b + beta * x. This function is commonly used and can often be better optimized than implementing it using x = A * b. However, for simplicity, we will implement it exactly like that in this example.

Creates a stencil matrix in CSR format for the given number of discretization points.

Generates the RHS vector given f and the boundary conditions.

```
Prints the solution u.
template <typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                                               const gko::matrix::Dense<ValueType>* u)
         std::cout « u0 « '\n';
for (int i = 0; i < u->get_size()[0]; ++i) {
    std::cout « u->get_const_values()[i] « '\n';
          std::cout « u1 « std::endl;
Computes the 1-norm of the error given the computed u and the correct solution function correct_u.
template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                                                      const gko::matrix::Dense<ValueType>* u,
                                                      Closure correct_u)
          const auto h = 1.0 / (discretization_points + 1);
          auto error = 0.0;
for (int i = 0; i < discretization_points; ++i) {</pre>
                 using std::abs;
const auto xi = (i + 1) * h;
                   error +=
                          abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
          return error;
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
Figure out where to run the code
       (argc == 2 && (std::string(argv[1]) == "--help")) {
  std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
          std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int discretization_points =
          argc >= 3 ? std::atoi(argv[2]) : 100u;
\verb|std::map| < \verb|std::string|, std::function| < \verb|std::shared_ptr| < gko::Executor| > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ > () \\ >
          exec_map{
                    {"omp",
                                    [] { return gko::OmpExecutor::create(); }},
                    {"cuda",
                     [] {
                               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                                                                             true);
                    {"hip",
                      [] {
                              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                    {"dpcpp",
                      [] {
                               return gko::DpcppExecutor::create(0,
                                                                                                               gko::OmpExecutor::create());
                   {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
executor used by the application
const auto app_exec = exec->get_master();
problem:
auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType{6} * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
```

```
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
const RealValueType reduction factor{1e-7};
Generate solver and solve the system
    .with_criteria(gko::stop::Iteration::build()
                         .with_max_iters(discretization_points)
                        .on(exec),
                    gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor)
                        .on(exec))
    .on(exec)
notice how our custom StencilMatrix can be used in the same way as any built-in type
        ->generate(StencilMatrix<ValueType>::create(exec, discretization_points, -1, 2, -1))
        ->apply(rhs, u);
    std::cout « "\nSolve complete."

« "\nThe average relative error is "
               « calculate_error(discretization_points, u.get(), correct_u) /
                      {\tt discretization\_points}
               « std::endl;
```

Results

Comments about programming and debugging

The plain program

```
******GINKGO LICENSE>***************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
   Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <iostream>
#include <map>
#include <string>
#include <omp.h>
#include <ginkgo/ginkgo.hpp>
template <typename ValueType>
void stencil_kernel(std::size_t size, const ValueType* coefs,
                    const ValueType* b, ValueType* x);
template <typename ValueType>
class StencilMatrix : public gko::EnableLinOp<StencilMatrix<ValueType»,</pre>
                      public gko::EnableCreateMethod<StencilMatrix<ValueType» {</pre>
```

```
public:
    StencilMatrix(std::shared_ptr<const gko::Executor> exec,
                   gko::size_type size = 0, ValueType left = -1.0,
ValueType center = 2.0, ValueType right = -1.0)
         : gko::EnableLinOp<StencilMatrix>(exec, gko::dim<2>{size}),
          coefficients(exec, {left, center, right})
    { }
protected:
    using vec = gko::matrix::Dense<ValueType>;
    using coef_type = gko::array<ValueType>;
    void apply_impl(const gko::LinOp* b, gko::LinOp* x)const override
         auto dense_b = gko::as<vec>(b);
auto dense_x = gko::as<vec>(x);
         struct stencil_operation: gko::Operation {
    stencil_operation(const coef_type& coefficients, const vec* b,
                                 vec* x)
                  : coefficients{coefficients}, b{b}, x{x}
             { }
             void run(std::shared_ptr<const gko::OmpExecutor>)const override
                  auto b_values = b->get_const_values();
                  auto x_values = x->get_values();
#pragma omp parallel for
                  for (std::size_t i = 0; i < x->get_size()[0]; ++i) {
                      auto coefs = coefficients.get_const_data();
                       auto result = coefs[1] * b_values[i];
                       if (i > 0) {
                           result += coefs[0] * b_values[i - 1];
                       if (i < x->get_size()[0] - 1) {
                           result += coefs[2] * b_values[i + 1];
                       x_values[i] = result;
                  }
             void run(std::shared ptr<const gko::CudaExecutor>)const override
                  stencil_kernel(x->get_size()[0], coefficients.get_const_data(),
                                  b->get_const_values(), x->get_values());
             const coef_type& coefficients;
             const vec* b;
             vec* x;
         this->get_executor()->run(
             stencil_operation(coefficients, dense_b, dense_x));
    {
         auto dense_b = gko::as<vec>(b);
         auto dense_x = gko::as<vec>(x);
         auto tmp_x = dense_x->clone();
         this->apply_impl(b, tmp_x.get());
         dense_x->scale(beta);
         dense_x->add_scaled(alpha, tmp_x);
private:
    coef_type coefficients;
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
    const auto discretization_points = matrix->get_size()[0];
    auto row_ptrs = matrix->get_row_ptrs();
auto col_idxs = matrix->get_col_idxs();
    auto values = matrix->get_values();
    IndexType pos = 0;
    const ValueType coefs[] = \{-1, 2, -1\};
    row_ptrs[0] = pos;
for (int i = 0; i < discretization_points; ++i) {</pre>
         for (auto ofs : {-1, 0, 1}) {
   if (0 <= i + ofs && i + ofs < discretization_points) {
      values[pos] = coefs[ofs + 1];</pre>
                  col_idxs[pos] = i + ofs;
                  ++pos;
             }
         row_ptrs[i + 1] = pos;
template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                   gko::matrix::Dense<ValueType>* rhs)
    const auto discretization_points = rhs->get_size()[0];
```

```
auto values = rhs->get_values();
    const ValueType h = 1.0 / (discretization_points + 1);
for (int i = 0; i < discretization_points; ++i) {
   const ValueType xi = ValueType(i + 1) * h;
   valueS[i] = -f(xi) * h * h;
}</pre>
    values[0] += u0;
    values[discretization_points - 1] += u1;
template <typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                     const gko::matrix::Dense<ValueType>* u)
    std::cout « u0 « '\n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout « u->get_const_values()[i] « '\n';
    std::cout « u1 « std::endl;
template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                         const gko::matrix::Dense<ValueType>* u,
                         Closure correct_u)
    const auto h = 1.0 / (discretization_points + 1);
    for (int i = 0; i < discretization_points; ++i) {</pre>
        using std::abs;
        const auto xi = (i + 1) * h;
        error +=
            abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    return error;
int main(int argc, char* argv[])
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
        std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const unsigned int discretization_points
        argc >= 3 ? std::atoi(argv[2]) : 100u;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
        exec_map{
             {"omp", [] { return gko::OmpExecutor::create(); }},
             {"cuda",
              [] {
                  return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                       true);
              } } ,
             {"hip",
                  return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
              }},
             { "dpcpp",
                  return gko::DpcppExecutor::create(0,
                                                        gko::OmpExecutor::create());
             {"reference", [] { return gko::ReferenceExecutor::create(); }};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
    auto correct_u = [](ValueType x) { return x * x * x; };
    auto f = [](ValueType x) { return ValueType{6} * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
    generate_rhs(f, u0, u1, rhs.get());
    auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
    for (int i = 0; i < u->get_size()[0]; ++i) {
        u \rightarrow get_values()[i] = 0.0;
    const RealValueType reduction_factor{1e-7};
    cg::build()
         .with_criteria(gko::stop::Iteration::build()
                              .with_max_iters(discretization_points)
                              .on(exec),
                         gko::stop::ResidualNorm<ValueType>::build()
                              .with_reduction_factor(reduction_factor)
```

Chapter 12

The custom-stopping-criterion program

The custom stopping criterion creation example..

This example depends on simple-solver, minimal-cuda-solver.

Introduction

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <thread>
* The ByInteraction class is a criterion which asks for user input to stop

* the iteration process. Using this criterion is slightly more complex than the

* other ones, because it is asynchronous therefore requires the use of threads.
class ByInteraction
    : public gko::EnablePolymorphicObject<ByInteraction, gko::stop::Criterion> {
    friend class gko::EnablePolymorphicObject<ByInteraction,
                                                    gko::stop::Criterion>;
    using Criterion = gko::stop::Criterion;
public:
    GKO_CREATE_FACTORY_PARAMETERS(parameters, Factory)
         * Boolean set by the user to stop the iteration process
         std::add_pointer<volatile bool>::type GKO_FACTORY_PARAMETER_SCALAR(
             stop_iteration_process, nullptr);
    GKO_ENABLE_CRITERION_FACTORY(ByInteraction, parameters, Factory);
    GKO_ENABLE_BUILD_METHOD (Factory);
protected:
    bool check_impl(gko::uint8 stoppingId, bool setFinalized,
                      gko::array<gko::stopping_status>* stop_status,
                      bool* one_changed, const Criterion::Updater&)override
{
         bool result = *(parameters_.stop_iteration_process);
             this->set_all_statuses(stoppingId, setFinalized, stop_status);
             *one_changed = true;
         return result;
    explicit ByInteraction(std::shared_ptr<const gko::Executor> exec)
         : EnablePolymorphicObject<ByInteraction, Criterion>(std::move(exec))
```

```
{ }
    explicit ByInteraction(const Factory* factory,
                             const gko::stop::CriterionArgs& args)
         : EnablePolymorphicObject<ByInteraction, Criterion>(
               factory->get_executor()),
          parameters_{factory->get_parameters()}
    { }
void run_solver(volatile bool* stop_iteration_process,
                 std::shared_ptr<gko::Executor> exec)
{
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using bicg = gko::solver::Bicgstab<ValueType>;
Read Data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read < vec > (std::ifstream("data/x0.mtx"), exec);
Create solver factory and solve system
auto solver = bicg::build()
                    .with_criteria(ByInteraction::build()
                                        .with_stop_iteration_process(
                                           stop_iteration_process)
                                        .on(exec))
                   .on(exec)
                    ->generate(A);
solver->add_logger(gko::log::Stream<ValueType>::create(
   gko::log::Logger::iteration_complete_mask, std::cout, true));
solver->apply(b, x);
std::cout « "Solver stopped" « std::endl;
Print solution
std::cout \leftarrow "Solution (x): n";
write(std::cout, x);
Calculate residual
    auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(res);
    std::cout « "Residual norm sqrt(r^T r): n";
    write(std::cout, res);
int main(int argc, char* argv[])
Print version information
std::cout « gko::version_info::get() « std::endl;
Figure out where to run the code
   (argc = 2 && (std::string(argv[1]) == "--help")) {
  std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
Figure out where to run the code
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec map{
         {"omp",
                 [] { return gko::OmpExecutor::create(); }},
         {"cuda",
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
```

```
true);
        { "dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                 gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Declare a user controled boolean for the iteration process
volatile bool stop_iteration_process{};
Create a new a thread to launch the solver
std::thread t(run_solver, &stop_iteration_process, exec);
Look for an input command "stop" in the console, which sets the boolean to true
    std::cout « "Type 'stop' to stop the iteration process" « std::endl;
std::string command;
    while (std::cin » command) {
        if (command == "stop") {
            break;
        } else {
           std::cout « "Unknown command" « std::endl;
    std::cout « "User input command 'stop' - The solver will stop!"
              « std::endl;
    stop\_iteration\_process = true;
    t.join();
```

Results

```
This is the expected output:
```

```
[LOG] >> iteration 22516 completed with solver LinOp[gko::solver::Bicgstab<double>,0x7fe6a4003710] with
       residual LinOp[gko::matrix::Dense<double>,0x7fe6a40050b0], solution
       LinOp[gko::matrix::Dense<double>,0x7fe6a40048e0] and residual_norm LinOp[gko::LinOp const*,0]
LinOp[gko::matrix::Dense<double>,0x7fe6a40050b0][
    5.17803e-164
    -7.6865e-165
    -2.06149e-164
    -4.84737e-165
    -3.36597e-164
    2.22353e-164
    1.47594e-165
    -1.78592e-165
    -6.17274e-166
    -3.02681e-166
    7.82009e-166
    8.57102e-165
    -1.28879e-164
    -2.62076e-165
    2.55329e-165
    -5.95988e-166
    -5.79273e-166
    -5.20172e-166
    -6.79458e-166
// Typing 'stop' stops the solver.
User input command 'stop' - The solver will stop
LinOp[gko::matrix::Dense<double>,0x7fe6a40048e0][
    0.252218
    0.108645
    0.0662811
    0.0630433
    0.0384088
    0.0396536
    0.0402648
```

```
0.0338935
    0.0193098
    0.0234653
    0.0211499
    0.0196413
    0.0199151
    0.0181674
    0.0162722
    0.0150714
    0.0107016
    0.0121141
    0.0123025
Solver stopped
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
6.50306e-16
```

Comments about programming and debugging

The plain program

```
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
```

```
#include <thread>
class ByInteraction
       public gko::EnablePolymorphicObject<ByInteraction, gko::stop::Criterion> {
    friend \ class \ gko:: Enable Polymorphic Object < By Interaction,
                                               gko::stop::Criterion>;
    using Criterion = gko::stop::Criterion;
public:
    GKO_CREATE_FACTORY_PARAMETERS (parameters, Factory)
        std::add_pointer<volatile bool>::type GKO_FACTORY_PARAMETER_SCALAR(
            stop_iteration_process, nullptr);
    };
    GKO_ENABLE_CRITERION_FACTORY(ByInteraction, parameters, Factory);
    GKO_ENABLE_BUILD_METHOD (Factory);
protected:
    bool check_impl(gko::uint8 stoppingId, bool setFinalized,
                     gko::array<gko::stopping_status>* stop_status,
                    bool* one_changed, const Criterion::Updater&)override
{
        bool result = *(parameters_.stop_iteration_process);
        if (result) {
            this->set_all_statuses(stoppingId, setFinalized, stop_status);
            *one_changed = true;
        return result;
    explicit ByInteraction(std::shared_ptr<const gko::Executor> exec)
        : EnablePolymorphicObject<ByInteraction, Criterion>(std::move(exec))
    explicit ByInteraction(const Factory* factory,
                           const gko::stop::CriterionArgs& args)
        : EnablePolymorphicObject<ByInteraction, Criterion>(
              factory->get_executor()),
          parameters_{factory->get_parameters()}
    {}
};
void run_solver(volatile bool* stop_iteration_process,
                std::shared_ptr<gko::Executor> exec)
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using bicg = gko::solver::Bicgstab<ValueType>;
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    auto solver = bicg::build()
                      .with_criteria(ByInteraction::build()
                                          .with_stop_iteration_process(
                                              stop_iteration_process)
                                          .on(exec))
                       .on(exec)
                       ->generate(A);
    solver->add_logger(gko::log::Stream<ValueType>::create(
        gko::log::Logger::iteration_complete_mask, std::cout, true));
    solver->apply(b, x);
std::cout « "Solver stopped" « std::endl;
    std::cout « "Solution (x): \n";
    write(std::cout, x);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(res);
    std::cout « "Residual norm sqrt(r^T r): \n";
    write(std::cout, res);
int main(int argc, char* argv[])
    std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
        std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    exec map{
            { "omp",
                    [] { return gko::OmpExecutor::create(); }},
            {"cuda",
             [] {
                 return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
            }},
{"hip",
```

Chapter 13

The distributed-solver program

The distributed solver example..

This example depends on simple-solver, three-pt-stencil-solver.

Introduction

This distributed solver example should help you understand the basics of using Ginkgo in a distributed setting. The example will solve a simple 1D Laplace equation where the system can be distributed row-wise to multiple processes. To run the solver with multiple processes, use mpirun -n NUM_PROCS ./distributed-solver [executor] [num_grid_points] [num_iterations].

If you are using GPU devices, please make sure that you run this example with at most as many processes as you have GPU devices available.

The commented program

Include files

This is the main ginkgo header file.

```
#include <ginkgo/ginkgo.hpp>
```

Add the C++ iostream header to output information to the console.

#include <iostream>

Add the STL map header for the executor selection

#include <map>

Add the string manipulation header to handle strings.

```
#include <string>
int main(int argc, char* argv[])
{
```

Initialize the MPI environment

Since this is an MPI program, we need to initialize and finalize MPI at the begin and end respectively of our program. This can be easily done with the following helper construct that uses RAII to automate the initialization and finalization.

```
const gko::experimental::mpi::environment env(argc, argv);
```

Type Definitiions

Define the needed types. In a parallel program we need to differentiate beweeen global and local indices, thus we have two index types.

```
using GlobalIndexType = gko::int64;
using LocalIndexType = gko::int32;
```

The underlying value type.

```
using ValueType = double;
```

As vector type we use the following, which implements a subset of gko::matrix::Dense.

```
using dist_vec = gko::experimental::distributed::Vector<ValueType>;
```

As matrix type we simply use the following type, which can read distributed data and be applied to a distributed vector.

We still need a localized vector type to be used as scalars in the advanced apply operations.

```
using vec = gko::matrix::Dense<ValueType>;
```

The partition type describes how the rows of the matrices are distributed.

We can use here the same solver type as you would use in a non-distributed program. Please note that not all solvers support distributed systems at the moment.

```
using solver = gko::solver::Cg<ValueType>;
using schwarz = gko::experimental::distributed::preconditioner::Schwarz<
    ValueType, LocalIndexType, GlobalIndexType>;
using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
```

Create an MPI communicator get the rank of the calling process.

```
const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);
const auto rank = comm.rank();
```

User Input Handling

User input settings:

- · The executor, defaults to reference.
- The number of grid points, defaults to 100.
- The number of iterations, defaults to 1000.

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   if (rank == 0) {
       « std::endl;
   std::exit(-1);
ValueType t_init = gko::experimental::mpi::get_walltime();
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const auto grid_dim =
   static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
const auto num_iters =
   static_cast<gko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
const std::map<std::string,
              std::function<std::shared_ptr<gko::Executor>(MPI_Comm)»
   executor_factory_mpi{
       {"reference",
        [](MPI_Comm) { return gko::ReferenceExecutor::create(); }},
       {"omp", [](MPI_Comm) { return gko::OmpExecutor::create(); }},
{"cuda",
        [](MPI_Comm comm) {
```

```
int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::CudaExecutor::get_num_devices());
             return gko::CudaExecutor::create(
                device_id, gko::ReferenceExecutor::create(), false,
                gko::allocation_mode::device);
         }},
        {"hip",
            int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::HipExecutor::get_num_devices());
             return gko::HipExecutor::create(
                device_id, gko::ReferenceExecutor::create(), true);
        {"dpcpp", [](MPI_Comm comm) {
             int device_id = 0;
             if (gko::DpcppExecutor::get_num_devices("gpu")) {
                device_id = gko::experimental::mpi::map_rank_to_device_id(
                    comm, gko::DpcppExecutor::get_num_devices("gpu"));
             } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
                device_id = gko::experimental::mpi::map_rank_to_device_id(
                     comm, gko::DpcppExecutor::get_num_devices("cpu"));
             } else {
                throw std::runtime_error("No suitable DPC++ devices");
             return gko::DpcppExecutor::create(
                device_id, gko::ReferenceExecutor::create());
auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
```

Creating the Distributed Matrix and Vectors

As a first step, we create a partition of the rows. The partition consists of ranges of consecutive rows which are assigned a part-id. These part-ids will be used for the distributed data structures to determine which rows will be stored locally. In this example each rank has (nearly) the same number of rows, so we can use the following specialized constructor. See gko::distributed::Partition for other modes of creating a partition.

```
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
```

Assemble the matrix using a 3-pt stencil and fill the right-hand-side with a sine value. The distributed matrix supports only constructing an empty matrix of zero size and filling in the values with gko::experimental::distributed::Matrix::read_distributed. Only the data that belongs to the rows by this rank will be assembled.

```
gko::matrix_data<ValueType, GlobalIndexType> A_data;
gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {</pre>
    if (i > 0) {
        A data.nonzeros.emplace back(i, i - 1, -1);
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid dim -
                         1) {
        A_data.nonzeros.emplace_back(i, i + 1, -1);
    b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
}
```

Take timings.

```
comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();
```

Read the matrix data, currently this is only supported on CPU executors. This will also set up the communication pattern needed for the distributed matrix-vector multiplication.

```
auto A_host = gko::share(dist_mtx::create(exec->get_master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);
```

After reading, the matrix and vector can be moved to the chosen executor, since the distributed matrix supports SpMV also on devices.

```
auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);
Take timings.
comm.synchronize();
ValueType t_read_setup_end = gko::experimental::mpi::get_walltime();
```

Solve the Distributed System

Generate the solver, this is the same as in the non-distributed case. with a local block diagonal preconditioner.

```
Setup the local block diagonal solver factory.
```

```
auto local_solver = gko::share(bj::build().on(exec));
```

```
Setup the stopping criterion and logger
```

Add logger to the generated solver to log the iteration count and residual norm

```
Ainv->add_logger(logger);
```

```
Take timings.
```

```
comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();
```

Apply the distributed solver, this is the same as in the non-distributed case.

```
Ainv->apply(b, x);
```

Take timings.

```
comm.synchronize();
ValueType t_end = gko::experimental::mpi::get_walltime();
```

Get the residual.

```
auto res_norm = gko::as<vec>(logger->get_residual_norm());
```

Printing Results

Print the achieved residual norm and timings on rank 0.

```
if (comm.rank() == 0) {
```

```
clang-format off
```

```
std::cout « "\nNum rows in matrix: " « num_rows

« "\nNum ranks: " « comm.size()

« "\nFinal Res norm: " « *res_norm->get_const_values()

« "\nIteration count: " « logger->get_num_iterations()

« "\nInit time: " « t_init_end - t_init

« "\nRead time: " « t_read_setup_end - t_init

« "\nSolver generate time: " « t_solver_generate_end - t_read_setup_end

« "\nSolver apply time: " « t_end - t_solver_generate_end

« "\nTotal time: " « t_end - t_init

« std::endl;
```

clang-format on

Results

```
This is the expected output for mpirun -n 4 ./distributed-solver:
Num rows in matrix: 100
Num ranks: 4
Final Res norm: 5.58392e-12
Iteration count: 7
Init time: 0.0663887
Read time: 0.0729806
Solver generate time: 7.6348e-05
Solver apply time: 0.0680783
Total time: 0.141351
```

The timings may vary depending on the machine.

The plain program

```
***********GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the \frac{1}{2}
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its contributors may be used to endorse or promote products derived from \frac{1}{2}
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    const gko::experimental::mpi::environment env(argc, argv);
    using GlobalIndexType = gko::int64;
using LocalIndexType = gko::int32;
    using ValueType = double;
    using dist_vec = gko::experimental::distributed::Vector<ValueType>;
    using dist mtx =
        gko::experimental::distributed::Matrix<ValueType, LocalIndexType,</pre>
    using vec = gko::matrix::Dense<ValueType>;
    using part_type =
         gko::experimental::distributed::Partition<LocalIndexType,</pre>
                                                       GlobalIndexType>;
    using solver = gko::solver::Cg<ValueType>;
    using schwarz = gko::experimental::distributed::preconditioner::Schwarz<
        ValueType, LocalIndexType, GlobalIndexType>;
    using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
    const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);
const auto rank = comm.rank();
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
         if (rank == 0) {
             « std::endl;
        std::exit(-1);
```

```
ValueType t_init = gko::experimental::mpi::get_walltime();
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const auto grid_dim =
   static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
const auto num iters =
   static_cast<qko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
const std::map<std::string,
               std::function<std::shared_ptr<gko::Executor>(MPI_Comm)»
    executor_factory_mpi{
        {"reference".
         [](MPI_Comm) { return gko::ReferenceExecutor::create(); }},
        """ | (MPI_Comm) { return gko::OmpExecutor::create(); }},
""" cuda",
         [](MPI_Comm comm) {
             int device_id = gko::experimental::mpi::map_rank_to_device_id(
                 comm, gko::CudaExecutor::get_num_devices());
             return gko::CudaExecutor::create(
                 device id, gko::ReferenceExecutor::create(), false,
                 gko::allocation_mode::device);
         }},
        {"hip",
         [] (MPI_Comm comm) {
             int device_id = gko::experimental::mpi::map_rank_to_device_id(
    comm, gko::HipExecutor::get_num_devices());
             return gko::HipExecutor::create(
                 device_id, gko::ReferenceExecutor::create(), true);
        {"dpcpp", [](MPI_Comm comm) {
   int device_id = 0;
             if (gko::DpcppExecutor::get_num_devices("gpu")) {
                 device_id = gko::experimental::mpi::map_rank_to_device_id(
                    comm, gko::DpcppExecutor::get_num_devices("gpu"));
             } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
                 device_id = gko::experimental::mpi::map_rank_to_device_id(
                     comm, gko::DpcppExecutor::get_num_devices("cpu"));
             } else {
                 throw std::runtime error("No suitable DPC++ devices");
             return gko::DpcppExecutor::create(
                 device_id, gko::ReferenceExecutor::create());
auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
gko::matrix_data<ValueType, GlobalIndexType> A_data;
gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {</pre>
    if (i > 0) {
        A_data.nonzeros.emplace_back(i, i - 1, -1);
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid_dim - 1) {</pre>
        A_data.nonzeros.emplace_back(i, i + 1, -1);
   b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
    x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();
auto A host = gko::share(dist mtx::create(exec->get master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);
auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);
comm.synchronize():
ValueType t read setup end = gko::experimental::mpi::get walltime();
auto local_solver = gko::share(bj::build().on(exec));
const gko::remove_complex<ValueType> reduction_factor{1e-8};
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
   gko::log::Convergence<ValueType>::create();
auto Ainv =
   solver::build()
```

```
.with_preconditioner(schwarz::build()
                               .with_local_solver_factory(local_solver)
                                .on(exec))
        .with_criteria(
           gko::stop::Iteration::build().with_max_iters(num_iters).on(
               exec),
           gko::stop::ResidualNorm<ValueType>::build()
               .with_reduction_factor(reduction_factor)
                .on(exec))
        .on(exec)
        ->generate(A);
Ainv->add_logger(logger);
comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();
Ainv->apply(b, x);
comm.synchronize();
valueType t_end = gko::experimental::mpi::get_walltime();
auto res_norm = gko::as<vec>(logger->get_residual_norm());
if (comm.rank() == 0) {
   « std::endl;
}
```

Chapter 14

The external-lib-interfacing program

The external library(deal.II) interfacing example..

Introduction

About the example

The commented program

```
#include <deal.II/base/function.h>
#include <deal.II/base/logstream.h>
#include <deal.II/base/quadrature_lib.h>
#include <deal.II/dofs/dof_accessor.h>
#include <deal.II/dofs/dof_handler.h>
#include <deal.II/dofs/dof_tools.h>
#include <deal.II/fe/fe_q.h>
#include <deal.II/fe/fe_values.h>
#include <deal.II/grid/grid_generator.h>
#include <deal.II/grid/grid_out.h>
#include <deal.II/grid/grid_refinement.h>
#include <deal.II/grid/tria.h>
#include <deal.II/grid/tria_accessor.h>
#include <deal.II/grid/tria_iterator.h>
#include <deal.II/lac/constraint_matrix.h>
#include <deal.II/lac/dynamic_sparsity_pattern.h>
#include <deal.II/lac/full_matrix.h>
#include <deal.II/lac/precondition.h>
#include <deal.II/lac/solver_bicgstab.h>
#include <deal.II/lac/sparse_matrix.h>
#include <deal.II/lac/vector.h>
#include <deal.II/numerics/data_out.h>
#include <deal.II/numerics/matrix_tools.h>
#include <deal.II/numerics/vector_tools.h>
```

The following two files provide classes and information for multithreaded programs. In the first one, the classes and functions are declared which we need to do assembly in parallel (i.e. the WorkStream namespace). The second file has a class MultithreadInfo which can be used to query the number of processors in your system, which is often useful when deciding how many threads to start in parallel.

```
#include <deal.II/base/multithread_info.h>
#include <deal.II/base/work_stream.h>
```

The next new include file declares a base class <code>TensorFunction</code> not unlike the <code>Function</code> class, but with the difference that the return value is tensor-valued rather than scalar of vector-valued.

```
#include <deal.II/base/tensor_function.h>
#include <deal.II/numerics/error_estimator.h>
```

Ginkgo's header file

#include <ginkgo/ginkgo.hpp>

This is C++, as we want to write some output to disk:

```
#include <fstream>
#include <iostream>
```

The last step is as in previous programs:

```
namespace Step9 {
using namespace dealii;
```

AdvectionProblem class declaration

Following we declare the main class of this program. It is very much like the main classes of previous examples, so we again only comment on the differences.

```
template <int dim>
class AdvectionProblem {
public:
   AdvectionProblem();
   AdvectionProblem();
   void run();
private:
   void setup system();
```

The next set of functions will be used to assemble the matrix. However, unlike in the previous examples, the <code>assemble_system()</code> function will not do the work itself, but rather will delegate the actual assembly to helper functions <code>assemble_local_system()</code> and <code>copy_local_to_global()</code>. The rationale is that matrix assembly can be parallelized quite well, as the computation of the local contributions on each cell is entirely independent of other cells, and we only have to synchronize when we add the contribution of a cell to the global matrix.

The strategy for parallelization we choose here is one of the possibilities mentioned in detail in the threads module in the documentation. Specifically, we will use the WorkStream approach discussed there. Since there is so much documentation in this module, we will not repeat the rationale for the design choices here (for example, if you read through the module mentioned above, you will understand what the purpose of the <code>AssemblyScratchData</code> and <code>AssemblyCopyData</code> structures is). Rather, we will only discuss the specific implementation.

If you read the page mentioned above, you will find that in order to parallelize assembly, we need two data structures – one that corresponds to data that we need during local integration ("scratch data", i.e., things we only need as temporary storage), and one that carries information from the local integration to the function that then adds the local contributions to the corresponding elements of the global matrix. The former of these typically contains the FEValues and FEFaceValues objects, whereas the latter has the local matrix, local right hand side, and information about which degrees of freedom live on the cell for which we are assembling a local contribution. With this information, the following should be relatively self-explanatory:

```
struct AssemblyScratchData {
    AssemblyScratchData(const FiniteElement<dim>& fe);
    AssemblyScratchData(const AssemblyScratchData& scratch_data);
    FEValues<dim> fe_values;
    FEFaceValues<dim> fe_face_values;
};
struct AssemblyCopyData {
    FullMatrix<double> cell_matrix;
    Vector<double> cell_rhs;
    std::vector<types::global_dof_index> local_dof_indices;
};
void assemble_system();
void local_assemble_system(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    AssemblyScratchData& scratch, AssemblyCopyData& copy_data);
void copy_local_to_global(const AssemblyCopyData& copy_data);
```

The following functions again are as in previous examples, as are the subsequent variables.

```
void solve();
void refine_grid();
void output_results(const unsigned int cycle) const;
Triangulation<dim> triangulation;
DoFHandler<dim> dof_handler;
FE_Q<dim> fe;
ConstraintMatrix hanging_node_constraints;
SparsityPattern sparsity_pattern;
SparseMatrix<double> system_matrix;
Vector<double> solution;
Vector<double> system_rhs;
```

Equation data declaration

};

Next we declare a class that describes the advection field. This, of course, is a vector field with as many components as there are space dimensions. One could now use a class derived from the Function base class, as we have done for boundary values and coefficients in previous examples, but there is another possibility in the library, namely

a base class that describes tensor valued functions. In contrast to the usual Function objects, we provide the compiler with knowledge on the size of the objects of the return type. This enables the compiler to generate efficient code, which is not so simple for usual vector-valued functions where memory has to be allocated on the heap (thus, the Function::vector_value function has to be given the address of an object into which the result is to be written, in order to avoid copying and memory allocation and deallocation on the heap). In addition to the known size, it is possible not only to return vectors, but also tensors of higher rank; however, this is not very often requested by applications, to be honest...

The interface of the TensorFunction class is relatively close to that of the Function class, so there is probably no need to comment in detail the following declaration:

In previous examples, we have used assertions that throw exceptions in several places. However, we have never seen how such exceptions are declared. This can be done as follows:

The syntax may look a little strange, but is reasonable. The format is basically as follows: use the name of one of the macros DeclexceptionN, where N denotes the number of additional parameters which the exception object shall take. In this case, as we want to throw the exception when the sizes of two vectors differ, we need two arguments, so we use Declexception2. The first parameter then describes the name of the exception, while the following declare the data types of the parameters. The last argument is a sequence of output directives that will be piped into the std:cerr object, thus the strange format with the leading << operator and the like. Note that we can access the parameters which are passed to the exception upon construction (i.e. within the Assert call) by using the names arg1 through argN, where N is the number of arguments as defined by the use of the respective macro DeclexceptionN.

To learn how the preprocessor expands this macro into actual code, please refer to the documentation of the exception classes in the base library. Suffice it to say that by this macro call, the respective exception class is declared, which also has error output functions already implemented.

The following two functions implement the interface described above. The first simply implements the function as described in the introduction, while the second uses the same trick to avoid calling a virtual function as has already been introduced in the previous example program. Note the check for the right sizes of the arguments in the second function, which should always be present in such functions; it is our experience that many if not most programming errors result from incorrectly initialized arrays, incompatible parameters to functions and the like; using assertion as in this case can eliminate many of these problems.

Besides the advection field, we need two functions describing the source terms (right hand side) and the boundary values. First for the right hand side, which follows the same pattern as in previous examples. As described in the introduction, the source is a constant function in the vicinity of a source point, which we denote by the constant static variable center_point. We set the values of this center using the same template tricks as we have shown

in the step-7 example program. The rest is simple and has been shown previously, including the way to avoid virtual function calls in the value list function.

The only new thing here is that we check for the value of the component parameter. As this is a scalar function, it is obvious that it only makes sense if the desired component has the index zero, so we assert that this is indeed the case. ExcIndexRange is a global predefined exception (probably the one most often used, we therefore made it global instead of local to some class), that takes three parameters: the index that is outside the allowed range, the first element of the valid range and the one past the last (i.e. again the half-open interval so often used in the C++ standard library):

Finally for the boundary values, which is just another class derived from the Function base class:

```
template <int dim>
class BoundaryValues : public Function<dim> {
public:
    BoundaryValues() : Function<dim>() {}
    virtual double value(const Point < dim > & p,
                         const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                            std::vector<double>& values,
                            const unsigned int component = 0) const;
template <int dim>
double BoundaryValues<dim>::value(const Point<dim>& p,
                                  const unsigned int component) const
    (void) component;
    Assert (component == 0, ExcIndexRange (component, 0, 1));
    const double sine term =
       std::sin(16 * numbers::PI * std::sqrt(p.norm_square()));
    const double weight = std::exp(-5 * p.norm_square()) / std::exp(-5.);
    return sine_term * weight;
template <int dim>
void BoundaryValues<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<double>& values,
                                     const unsigned int component) const
    Assert(values.size() == points.size(),
          ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)</pre>
        values[i] = BoundaryValues<dim>::value(points[i], component);
```

GradientEstimation class declaration

Now, finally, here comes the class that will compute the difference approximation of the gradient on each cell and weighs that with a power of the mesh size, as described in the introduction. This class is a simple version of the DerivativeApproximation class in the library, that uses similar techniques to obtain finite difference approximations of the gradient of a finite element field, or of higher derivatives.

The class has one public static function <code>estimate</code> that is called to compute a vector of error indicators, and a few private functions that do the actual work on all active cells. As in other parts of the library, we follow an informal convention to use vectors of floats for error indicators rather than the common vectors of doubles, as the additional accuracy is not necessary for estimated values.

In addition to these two functions, the class declares two exceptions which are raised when a cell has no neighbors in each of the space directions (in which case the matrix described in the introduction would be singular and can't be inverted), while the other one is used in the more common case of invalid parameters to a function, namely a vector of wrong size.

Two other comments: first, the class has no non-static member functions or variables, so this is not really a class, but rather serves the purpose of a namespace in C++. The reason that we chose a class over a namespace is that this way we can declare functions that are private. This can be done with namespaces as well, if one declares some functions in header files in the namespace and implements these and other functions in the implementation file. The functions not declared in the header file are still in the namespace but are not callable from outside. However, as we have only one file here, it is not possible to hide functions in the present case.

The second comment is that the dimension template parameter is attached to the function rather than to the class itself. This way, you don't have to specify the template parameter yourself as in most other cases, but the compiler can figure its value out itself from the dimension of the DoF handler object that one passes as first argument.

Before jumping into the fray with the implementation, let us also comment on the parallelization strategy. We have already introduced the necessary framework for using the WorkStream concept in the declaration of the main class of this program above. We will use it again here. In the current context, this means that we have to define (i) classes for scratch and copy objects, (ii) a function that does the local computation on one cell, and (iii) a function that copies the local result into a global object. Given this general framework, we will, however, deviate from it a bit. In particular, WorkStream was generally invented for cases where each local computation on a cell adds to a global object - for example, when assembling linear systems where we add local contributions into a global matrix and right hand side. WorkStream is designed to handle the potential conflict of multiple threads trying to do this addition at the same time, and consequently has to provide for some way to ensure that only thread gets to do this at a time. Here, however, the situation is slightly different: we compute contributions from every cell individually, but then all we need to do is put them into an element of an output vector that is unique to each cell. Consequently, there is no risk that the write operations from two cells might conflict, and the elaborate machinery of WorkStream to avoid conflicting writes is not necessary. Consequently, what we will do is this: We still need a scratch object that holds, for example, the FEValues object. However, we only create a fake, empty copy data structure. Likewise, we do need the function that computes local contributions, but since it can already put the result into its final location, we do not need a copy-local-to-global function and will instead give the WorkStream::run() function an empty function object - the equivalent to a NULL function pointer.

```
class GradientEstimation {
    template <int dim>
    static void estimate(const DoFHandler<dim>& dof,
                          const Vector<double>& solution.
                          Vector<float>& error per cell);
    DeclException2(ExcInvalidVectorLength, int, int, "Vector has length " « arg1 « ", but should have "
                    « arg2);
    DeclException0(ExcInsufficientDirections);
private:
    template <int dim>
    struct EstimateScratchData {
        EstimateScratchData(const FiniteElement<dim>& fe,
                              const Vector<double>& solution,
                              Vector<float>& error_per_cell);
        EstimateScratchData(const EstimateScratchData& data);
        FEValues<dim> fe midpoint value;
        const Vector<double>& solution;
        Vector<float>& error_per_cell;
```

```
};
struct EstimateCopyData {};
template <int dim>
static void estimate_cell(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    EstimateScratchData<dim>& scratch_data,
    const EstimateCopyData& copy_data);
```

AdvectionProblem class implementation

Now for the implementation of the main class. Constructor, destructor and the function <code>setup_system</code> follow the same pattern that was used previously, so we need not comment on these three function:

```
template <int dim>
AdvectionProblem<dim>::AdvectionProblem() : dof_handler(triangulation), fe(1)
template <int dim>
AdvectionProblem < dim>:: AdvectionProblem ()
   dof handler.clear();
template <int dim>
void AdvectionProblem < dim > :: setup system ()
    dof_handler.distribute_dofs(fe);
    hanging node constraints.clear();
    DoFTools::make hanging node constraints(dof handler.
                                             hanging_node_constraints);
    hanging_node_constraints.close();
    DynamicSparsityPattern dsp(dof_handler.n_dofs(), dof_handler.n_dofs());
    DoFTools::make_sparsity_pattern(dof_handler, dsp, hanging_node_constraints,
                                      *keep_constrained_dofs = * / true);
    sparsity pattern.copy from(dsp);
    system matrix.reinit(sparsity pattern);
    solution.reinit(dof_handler.n_dofs());
    system_rhs.reinit(dof_handler.n_dofs());
```

In the following function, the matrix and right hand side are assembled. As stated in the documentation of the main class above, it does not do this itself, but rather delegates to the function following next, utilizing the WorkStream concept discussed in threads .

If you have looked through the threads module, you will have seen that assembling in parallel does not take an incredible amount of extra code as long as you diligently describe what the scratch and copy data objects are, and if you define suitable functions for the local assembly and the copy operation from local contributions to global objects. This done, the following will do all the heavy lifting to get these operations done on multiple threads on as many cores as you have in your system:

After the matrix has been assembled in parallel, we still have to eliminate hanging node constraints. This is something that can't be done on each of the threads separately, so we have to do it now. Note also, that unlike in previous examples, there are no boundary conditions to be applied to the system of equations. This, of course, is due to the fact that we have included them into the weak formulation of the problem.

```
hanging_node_constraints.condense(system_matrix);
hanging_node_constraints.condense(system_rhs);
```

As already mentioned above, we need to have scratch objects for the parallel computation of local contributions. These objects contain FEValues and FEFaceValues objects, and so we will need to have constructors and copy constructors that allow us to create them. In initializing them, note first that we use bilinear elements, soGauss formulae with two points in each space direction are sufficient. For the cell terms we need the values and gradients of the shape functions, the quadrature points in order to determine the source density and the advection field at a given point, and the weights of the quadrature points times the determinant of the Jacobian at these points. In contrast, for the boundary integrals, we don't need the gradients, but rather the normal vectors to the cells. This determines which update flags we will have to pass to the constructors of the members of the class:

```
template <int dim>
{\tt AdvectionProblem < dim > :: Assembly Scratch Data:: Assembly Scratch Data()}
    const FiniteElement<dim>& fe)
    : fe_values(fe, QGauss<dim>(2),
               update_values | update_gradients | update_quadrature_points |
                    update JxW values),
      fe_face_values(fe, QGauss<dim - 1>(2),
                     update_values | update_quadrature_points |
                         update_JxW_values | update_normal_vectors)
template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const AssemblyScratchData& scratch data)
    : fe_values(scratch_data.fe_values.get_fe(),
                scratch_data.fe_values.get_quadrature(),
                update_values | update_gradients | update_quadrature_points |
                    update_JxW_values),
      fe_face_values(scratch_data.fe_face_values.get_fe(),
                     scratch_data.fe_face_values.get_quadrature(),
                     update_values | update_quadrature_points |
                          update_JxW_values | update_normal_vectors)
{ }
```

Now, this is the function that does the actual work. It is not very different from the assemble_system functions of previous example programs, so we will again only comment on the differences. The mathematical stuff follows closely what we have said in the introduction.

There are a number of points worth mentioning here, though. The first one is that we have moved the FEValues and FEFaceValues objects into the ScratchData object. We have done so because the alternative would have been to simply create one every time we get into this function – i.e., on every cell. It now turns out that the FEValues classes were written with the explicit goal of moving everything that remains the same from cell to cell into the construction of the object, and only do as little work as possible in FEValues::reinit() whenever we move to a new cell. What this means is that it would be very expensive to create a new object of this kind in this function as we would have to do it for every cell – exactly the thing we wanted to avoid with the FEValues class. Instead, what we do is create it only once (or a small number of times) in the scratch objects and then re-use it as often as we can.

This begs the question of whether there are other objects we create in this function whose creation is expensive compared to its use. Indeed, at the top of the function, we declare all sorts of objects. The <code>AdvectionField</code>, <code>RightHandSide</code> and <code>BoundaryValues</code> do not cost much to create, so there is no harm here. However, allocating memory in creating the <code>rhs_values</code> and similar variables below typically costs a significant amount of time, compared to just accessing the (temporary) values we store in them. Consequently, these would be candidates for moving into the <code>AssemblyScratchData</code> class. We will leave this as an exercise.

```
template <int dim>
void AdvectionProblem<dim>::local_assemble_system(
   const typename DoFHandler<dim>::active_cell_iterator& cell,
   AssemblyScratchData& scratch_data, AssemblyCopyData& copy_data)
{
```

First of all, we will need some objects that describe boundary values, right hand side function and the advection field. As we will only perform actions on these objects that do not change them, we declare them as constant, which can enable the compiler in some cases to perform additional optimizations.

```
const AdvectionField<dim> advection_field;
const RightHandSide<dim> right_hand_side;
const BoundaryValues<dim> boundary_values;
```

Then we define some abbreviations to avoid unnecessarily long lines:

```
const unsigned int dofs_per_cell = fe.dofs_per_cell;
const unsigned int n_q_points =
    scratch_data.fe_values.get_quadrature().size();
const unsigned int n_face_q_points =
    scratch_data.fe_face_values.get_quadrature().size();
```

We declare cell matrix and cell right hand side...

```
copy_data.cell_matrix.reinit(dofs_per_cell, dofs_per_cell);
copy_data.cell_rhs.reinit(dofs_per_cell);
```

... an array to hold the global indices of the degrees of freedom of the cell on which we are presently working...

... and array in which the values of right hand side, advection direction, and boundary values will be stored, for cell and face integrals respectively:

... and assemble the local contributions to the system matrix and right hand side as also discussed above:

Besides the cell terms which we have built up now, the bilinear form of the present problem also contains terms on the boundary of the domain. Therefore, we have to check whether any of the faces of this cell are on the boundary of the domain, and if so assemble the contributions of this face as well. Of course, the bilinear form only contains contributions from the inflow part of the boundary, but to find out whether a certain part of a face of the present cell is part of the inflow boundary, we have to have information on the exact location of the quadrature points and on the direction of flow at this point; we obtain this information using the FEFaceValues object and only decide within the main loop whether a quadrature point is on the inflow boundary.

Ok, this face of the present cell is on the boundary of the domain. Just as for the usual FEValues object which we have used in previous examples and also above, we have to reinitialize the FEFaceValues object for the present face:

```
scratch_data.fe_face_values.reinit(cell, face);
```

For the quadrature points at hand, we ask for the values of the inflow function and for the direction of flow:

```
boundary_values.value_list(
    scratch_data.fe_face_values.get_quadrature_points(),
    face_boundary_values);
advection_field.value_list(
    scratch_data.fe_face_values.get_quadrature_points(),
    face_advection_directions);
```

Now loop over all quadrature points and see whether it is on the inflow or outflow part of the boundary. This is determined by a test whether the advection direction points inwards or outwards of the domain (note that the normal vector points outwards of the cell, and since the cell is at the boundary, the normal vector points outward of the domain, so if the advection direction points into the domain, its scalar product with the normal vector must be negative):

If the is part of the inflow boundary, then compute the contributions of this face to the global matrix and right hand side, using the values obtained from the FEFaceValues object and the formulae discussed in the introduction:

```
for (unsigned int i = 0; i < dofs_per_cell; ++i) {</pre>
    for (unsigned int j = 0; j < dofs_per_cell; ++j)</pre>
         copy_data.cell_matrix(i, j) -
              (\texttt{face\_advection\_directions} \, [\, \underline{q} \, \underline{\hspace{0.1cm}} \, point \, ] \  \, \star \\
               scratch_data.fe_face_values.normal_vector(
                   a point) *
               scratch_data.fe_face_values.shape_value(
                   i, q_point)
               scratch_data.fe_face_values.shape_value(
                   j, q_point)
               scratch_data.fe_face_values.JxW(q_point));
    copy_data.cell_rhs(i) -=
         (face advection directions[g point] *
          scratch_data.fe_face_values.normal_vector(
              q_point) *
          face_boundary_values[q_point] *
          scratch_data.fe_face_values.shape_value(i,
                                                         q_point) *
          scratch_data.fe_face_values.JxW(q_point));
```

Now go on by transferring the local contributions to the system of equations into the global objects. The first step was to obtain the global indices of the degrees of freedom on this cell.

```
cell->get_dof_indices(copy_data.local_dof_indices);
```

The second function we needed to write was the one that copies the local contributions the previous function has computed and put into the copy data object, into the global matrix and right hand side vector objects. This is essentially what we always had as the last block of code when assembling something on every cell. The following should therefore be pretty obvious:

Following is the function that solves the linear system of equations. As the system is no more symmetric positive definite as in all the previous examples, we can't use the Conjugate Gradients method anymore. Rather, we use a solver that is tailored to nonsymmetric systems like the one at hand, the BiCGStab method. As preconditioner, we use the Block Jacobi method.

```
template <int dim>
void AdvectionProblem<dim>::solve()
{
```

Assert that the system be symmetric.

```
Assert(system_matrix.m() == system_matrix.n(), ExcNotQuadratic());
auto num_rows = system_matrix.m();
```

Make a copy of the rhs to use with Ginkgo.

```
std::vector<double> rhs(num_rows);
std::copy(system_rhs.begin(), system_rhs.begin() + num_rows, rhs.begin());
```

Ginkgo setup Some shortcuts: A vector is a Dense matrix with co-dimension 1. The matrix is setup in CSR. But various formats can be used. Look at Ginkgo's documentation.

```
using vec = gko::matrix::Dense<>;
using mtx = gko::matrix::Csr<>;
using bicgstab = gko::solver::Bicgstab<>;
using bj = gko::preconditioner::Jacobi<>;
using val_array = gko::array<double>;
```

Where the code is to be executed. Can be changed to omp or cuda to run on multiple threads or on gpu's std::shared_ptr<gko::Executor> exec = gko::ReferenceExecutor::create();

Setup Ginkgo's data structures

```
system_matrix.n_nonzero_elements());
mtx::value_type* values = A->get_values();
mtx::index_type* row_ptr = A->get_row_ptrs();
mtx::index_type* col_idx = A->get_col_idxs();
```

Convert to standard CSR format As deal.ii does not expose its system matrix pointers, we construct them individually

write entry into the first free one for this row

```
col_idx[ptrs[row]] = p->column();
values[ptrs[row]] = p->value();
```

then move pointer ahead

```
++ptrs[row];
}
```

Ginkgo solve The stopping criteria is set at maximum iterations of 1000 and a reduction factor of 1e-12. For other options, refer to Ginkgo's documentation.

Solve system

solver->apply(b, x);

Copy the solution vector back to deal.ii's data structures.

Distribute the possible hanging nodes constraints

```
hanging_node_constraints.distribute(solution);
```

The following function refines the grid according to the quantity described in the introduction. The respective computations are made in the class <code>GradientEstimation</code>. The only difference to previous examples is that we refine a little more aggressively (0.5 instead of 0.3 of the number of cells).

Writing output to disk is done in the same way as in the previous examples. Indeed, the function is identical to the one in step-6.

```
template <int dim>
```

```
void AdvectionProblem < dim > :: output_results (const unsigned int cycle) const
        GridOut grid out;
        std::ofstream output("grid-" + std::to_string(cycle) + ".eps");
        grid_out.write_eps(triangulation, output);
        DataOut<dim> data_out;
         data_out.attach_dof_handler(dof_handler);
         data_out.add_data_vector(solution, "solution");
        data_out.build_patches();
         std::ofstream output("solution-" + std::to_string(cycle) + ".vtk");
        data_out.write_vtk(output);
}
... as is the main loop (setup – solve – refine)
template <int dim>
void AdvectionProblem < dim >:: run ()
    for (unsigned int cycle = 0; cycle < 6; ++cycle) {
   std::cout « "Cycle " « cycle « ':' « std::endl;
   if (cycle == 0) {</pre>
             GridGenerator::hyper_cube(triangulation, -1, 1);
             triangulation.refine_global(4);
             refine_grid();
         std::cout « " Number of active cells:
                    « triangulation.n_active_cells() « std::endl;
         setup_system();
        std::cout « "
                          Number of degrees of freedom: " « dof_handler.n_dofs()
                    « std::endl;
         assemble_system();
        solve();
        output_results(cycle);
```

GradientEstimation class implementation

Now for the implementation of the GradientEstimation class. Let us start by defining constructors for the EstimateScratchData class used by the estimate_cell() function:

```
template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const FiniteElement<dim>& fe, const Vector<double>& solution,
    Vector<float>& error_per_cell)
    : fe_midpoint_value(fe, QMidpoint<dim>(),
                        update_values | update_quadrature_points),
      solution(solution),
      error_per_cell(error_per_cell)
{ }
template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const EstimateScratchData& scratch_data)
    : fe_midpoint_value(scratch_data.fe_midpoint_value.get_fe(),
                        scratch_data.fe_midpoint_value.get_quadrature(),
                        update_values | update_quadrature_points),
      solution(scratch_data.solution),
      error_per_cell(scratch_data.error_per_cell)
{ }
```

Next for the implementation of the GradientEstimation class. The first function does not much except for delegating work to the other function, but there is a bit of setup at the top.

Before starting with the work, we check that the vector into which the results are written has the right size. Programming mistakes in which one forgets to size arguments correctly at the calling site are quite common. Because the resulting damage from not catching such errors is often subtle (e.g., corruption of data somewhere in memory, or non-reproducible results), it is well worth the effort to check for such things.

Following now the function that actually computes the finite difference approximation to the gradient. The general outline of the function is to first compute the list of active neighbors of the present cell and then compute the quantities described in the introduction for each of the neighbors. The reason for this order is that it is not a one-liner to find a given neighbor with locally refined meshes. In principle, an optimized implementation would find neighbors and the quantities depending on them in one step, rather than first building a list of neighbors and in a second step their contributions but we will gladly leave this as an exercise. As discussed before, the worker function passed to WorkStream::run works on "scratch" objects that keep all temporary objects. This way, we do not need to create and initialize objects that are expensive to initialize within the function that does the work, every time it is called for a given cell. Such an argument is passed as the second argument. The third argument would be a "copy-data" object (see threads for more information) but we do not actually use any of these here. Because WorkStream::run() insists on passing three arguments, we declare this function with three arguments, but simply ignore the last one.

(This is unsatisfactory from an esthetic perspective. It can be avoided, at the cost of some other trickery. If you allow, let us here show how. First, assume that we had declared this function to only take two arguments by omitting the unused last one. Now, WorkStream::run still wants to call this function with three arguments, so we need to find a way to "forget" the third argument in the call. Simply passing WorkStream::run the pointer to the function as we do above will not do this – the compiler will complain that a function declared to have two arguments is called with three arguments. However, we can do this by passing the following as the third argument when calling WorkStream::run() above:

This creates a function object taking three arguments, but when it calls the underlying function object, it simply only uses the first and second argument – we simply "forget" to use the third argument :-) In the end, this isn't completely obvious either, and so we didn't implement it, but hey – it can be done!)

Now for the details:

We need space for the tensor Y, which is the sum of outer products of the y-vectors.

Tensor<2, dim> Y;

Then we allocate a vector to hold iterators to all active neighbors of a cell. We reserve the maximal number of active neighbors in order to avoid later reallocations. Note how this maximal number of active neighbors is computed here.

First initialize the FEValues object, as well as the Y tensor:

```
scratch_data.fe_midpoint_value.reinit(cell);
```

Then allocate the vector that will be the sum over the y-vectors times the approximate directional derivative: Tensor<1, dim> projected gradient;

Now before going on first compute a list of all active neighbors of the present cell. We do so by first looping over all faces and see whether the neighbor there is active, which would be the case if it is on the same level as the present

cell or one level coarser (note that a neighbor can only be once coarser than the present cell, as we only allow a maximal difference of one refinement over a face in deal.II). Alternatively, the neighbor could be on the same level and be further refined; then we have to find which of its children are next to the present cell and select these (note that if a child of a neighbor of an active cell that is next to this active cell, needs necessarily be active itself, due to the one-refinement rule cited above).

Things are slightly different in one space dimension, as there the one-refinement rule does not exist: neighboring active cells may differ in as many refinement levels as they like. In this case, the computation becomes a little more difficult, but we will explain this below.

Before starting the loop over all neighbors of the present cell, we have to clear the array storing the iterators to the active neighbors, of course.

First define an abbreviation for the iterator to the face and the neighbor

```
const typename DoFHandler<dim>::face_iterator face =
   cell->face(face_no);
const typename DoFHandler<dim>::cell_iterator neighbor =
   cell->neighbor(face_no);
```

Then check whether the neighbor is active. If it is, then it is on the same level or one level coarser (if we are not in 1D), and we are interested in it in any case.

```
if (neighbor->active())
    active_neighbors.push_back(neighbor);
else {
```

If the neighbor is not active, then check its children.

```
if (dim == 1) {
```

To find the child of the neighbor which bounds to the present cell, successively go to its right child if we are left of the present cell (n==0), or go to the left child if we are on the right (n==1), until we find an active cell.

```
typename DoFHandler<dim>::cell_iterator neighbor_child =
    neighbor;
while (neighbor_child->has_children())
    neighbor_child =
        neighbor_child->child(face_no == 0 ? 1 : 0);
```

As this used some non-trivial geometrical intuition, we might want to check whether we did it right, i.e. check whether the neighbor of the cell we found is indeed the cell we are presently working on. Checks like this are often useful and have frequently uncovered errors both in algorithms like the line above (where it is simple to involuntarily exchange n=1 for n=0 or the like) and in the library (the assumptions underlying the algorithm above could either be wrong, wrongly documented, or are violated due to an error in the library). One could in principle remove such checks after the program works for some time, but it might be a good things to leave it in anyway to check for changes in the library or in the algorithm above.

Note that if this check fails, then this is certainly an error that is irrecoverable and probably qualifies as an internal error. We therefore use a predefined exception class to throw here.

```
Assert(
  neighbor_child->neighbor(face_no == 0 ? 1 : 0) == cell,
  ExcInternalError());
```

If the check succeeded, we push the active neighbor we just found to the stack we keep:

```
active_neighbors.push_back(neighbor_child);
} else
```

If we are not in 1d, we collect all neighbor children 'behind' the subfaces of the current face

OK, now that we have all the neighbors, lets start the computation on each of them. First we do some preliminaries: find out about the center of the present cell and the solution at this point. The latter is obtained as a vector of

function values at the quadrature points, of which there are only one, of course. Likewise, the position of the center is the position of the first (and only) quadrature point in real space.

Now loop over all active neighbors and collect the data we need. Allocate a vector just like this_midpoint_ \leftarrow value which we will use to store the value of the solution in the midpoint of the neighbor cell. We allocate it here already, since that way we don't have to allocate memory repeatedly in each iteration of this inner loop (memory allocation is a rather expensive operation):

```
std::vector<double> neighbor_midpoint_value(1);
typename std::vector<typename DoFHandler<dim>::active_cell_iterator>::
    const_iterator neighbor_ptr = active_neighbors.begin();
for (; neighbor_ptr != active_neighbors.end(); ++neighbor_ptr) {
```

First define an abbreviation for the iterator to the active neighbor cell:

```
const typename DoFHandler<dim>::active_cell_iterator neighbor =
   *neighbor_ptr;
```

Then get the center of the neighbor cell and the value of the finite element function thereon. Note that for this information we have to reinitialize the FEValues object for the neighbor cell.

```
scratch_data.fe_midpoint_value.reinit(neighbor);
const Point<dim> neighbor_center =
    scratch_data.fe_midpoint_value.quadrature_point(0);
scratch_data.fe_midpoint_value.get_function_values(
    scratch_data.solution, neighbor_midpoint_value);
```

Compute the vector y connecting the centers of the two cells. Note that as opposed to the introduction, we denote by y the normalized difference vector, as this is the quantity used everywhere in the computations.

Tensor<1, dim> y = neighbor_center - this_center;

If now, after collecting all the information from the neighbors, we can determine an approximation of the gradient for the present cell, then we need to have passed over vectors y which span the whole space, otherwise we would not have all components of the gradient. This is indicated by the invertibility of the matrix.

If the matrix should not be invertible, this means that the present cell had an insufficient number of active neighbors. In contrast to all previous cases, where we raised exceptions, this is, however, not a programming error: it is a runtime error that can happen in optimized mode even if it ran well in debug mode, so it is reasonable to try to catch this error also in optimized mode. For this case, there is the AssertThrow macro: it checks the condition like the Assert macro, but not only in debug mode; it then outputs an error message, but instead of terminating the program as in the case of the Assert macro, the exception is thrown using the throw command of C++. This way, one has the possibility to catch this error and take reasonable counter actions. One such measure would be to refine the grid globally, as the case of insufficient directions can not occur if every cell of the initial grid has been refined at least once.

```
AssertThrow(determinant(Y) != 0, ExcInsufficientDirections());
```

If, on the other hand the matrix is invertible, then invert it, multiply the other quantity with it and compute the estimated error using this quantity and the right powers of the mesh width:

```
const Tensor<2, dim> Y_inverse = invert(Y);
Tensor<1, dim> gradient = Y_inverse * projected_gradient;
```

The last part of this function is the one where we write into the element of the output vector what we have just computed. The address of this vector has been stored in the scratch data object, and all we have to do is know how to get at the correct element inside this vector – but we can ask the cell we're on the how-manyth active cell it is for this:

```
scratch_data.error_per_cell(cell->active_cell_index()) =
    (std::pow(cell->diameter(), 1 + 1.0 * dim / 2) *
        std::sqrt(gradient.norm_square()));
}
// namespace Step9
```

Main function

The main function is similar to the previous examples. The main difference is that we use MultithreadInfo to set the maximum number of threads (see Parallel computing with multiple processors accessing shared memory" documentation module for more explanation). The number of threads used is the minimum of the environment variable DEAL_II_NUM_THREADS and the parameter of set_thread_limit. If no value is given to set thread_limit, the default value from the Intel Threading Building Blocks (TBB) library is used. If the call to set_thread_limit is omitted, the number of threads will be chosen by TBB indepently of DEAL_II_NUM_T HREADS.

```
int main()
       dealii::MultithreadInfo::set thread limit();
       Step9::AdvectionProblem<2> advection_problem_2d;
       advection_problem_2d.run();
   } catch (std::exception& exc) {
       std::cerr « std::endl
                 « std::endl
« "-----
                 « std::endl;
       std::cerr « "Exception on processing: " « std::endl
                 « exc.what() « std::endl
« "Aborting!" « std::endl
                 « "-
                 « std::endl;
       return 1:
   } catch (...)
       std::cerr « std::endl
                 « std::endl
                 « std::endl;
       « std::endl;
       return 1;
   return 0:
```

Results

Comments about programming and debugging

The plain program

```
/*

* Copyright (C) 2000 - 2018 by the deal.II authors

*

* This file is part of the deal.II library.

* The deal.II library is free software; you can use it, redistribute

* it, and/or modify it under the terms of the GNU Lesser General

* Public License as published by the Free Software Foundation; either

* version 2.1 of the License, or (at your option) any later version.

* The full text of the license can be found in the file LICENSE at

* the top level of the deal.II distribution.

*

* Author: Wolfgang Bangerth, University of Heidelberg, 2000

*/

/*

* This file has been taken verbatim from the deal.ii (version 9.0)

* examples directory and modified.

* This example aims to demonstrate the ease with which Ginkgo can

* be interfaced with other libraries. The only modification/ addition

* has been to the AdvectionProblem::solve () function.
```

```
#include <deal.II/base/function.h>
#include <deal.II/base/logstream.h>
#include <deal.II/base/quadrature_lib.h>
#include <deal.II/dofs/dof_accessor.h>
#include <deal.II/dofs/dof_handler.h>
#include <deal.II/dofs/dof_tools.h>
#include <deal.II/fe/fe_q.h>
#include <deal.II/fe/fe_values.h>
#include <deal.II/grid/grid_generator.h>
#include <deal.II/grid/grid_out.h>
#include <deal.II/grid/grid_refinement.h>
#include <deal.II/grid/tria.h>
#include <deal.II/grid/tria_accessor.h>
#include <deal.II/grid/tria_iterator.h>
#include <deal.II/lac/constraint_matrix.h>
#include <deal.II/lac/dynamic_sparsity_pattern.h>
#include <deal.II/lac/full_matrix.h>
#include <deal.II/lac/precondition.h>
#include <deal.II/lac/solver_bicgstab.h>
#include <deal.II/lac/sparse_matrix.h>
#include <deal.II/lac/vector.h>
#include <deal.II/numerics/data_out.h>
#include <deal.II/numerics/matrix tools.h>
#include <deal.II/numerics/vector_tools.h>
#include <deal.II/base/multithread_info.h>
#include <deal.II/base/work_stream.h>
#include <deal.II/base/tensor_function.h>
#include <deal.II/numerics/error_estimator.h>
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
namespace Step9 {
using namespace dealii;
template <int dim>
class AdvectionProblem {
public:
    AdvectionProblem();
     AdvectionProblem();
    void run();
private:
    void setup_system();
    struct AssemblyScratchData {
        AssemblyScratchData(const FiniteElement<dim>& fe);
        AssemblyScratchData(const AssemblyScratchData& scratch_data);
        FEValues<dim> fe_values;
        FEFaceValues<dim> fe_face_values;
    struct AssemblyCopyData {
        FullMatrix<double> cell_matrix;
        Vector<double> cell_rhs;
        std::vector<types::global_dof_index> local_dof_indices;
    void assemble_system();
    void local_assemble_system(
        const typename DoFHandler<dim>::active_cell_iterator& cell,
        AssemblyScratchData& scratch, AssemblyCopyData& copy_data);
    void copy_local_to_global(const AssemblyCopyData& copy_data);
    void solve();
    void refine_grid();
    void output_results(const unsigned int cycle) const;
    Triangulation<dim> triangulation;
    DoFHandler<dim> dof_handler;
    FE_Q<dim> fe;
    ConstraintMatrix hanging_node_constraints;
    SparsityPattern sparsity_pattern;
    SparseMatrix<double> system_matrix;
    Vector<double> solution:
    Vector<double> system_rhs;
};
template <int dim>
class AdvectionField : public TensorFunction<1, dim> {
public:
    AdvectionField() : TensorFunction<1, dim>() {}
    virtual Tensor<1, dim> value(const Point<dim>& p) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                           std::vector<Tensor<1, dim>& values) const;
    template <int dim>
Tensor<1, dim> AdvectionField<dim>::value(const Point<dim>& p)const
    Point<dim> value;
value[0] = 2;
    for (unsigned int i = 1; i < dim; ++i)</pre>
```

```
value[i] = 1 + 0.8 * std::sin(8 * numbers::PI * p[0]);
template <int dim>
void AdvectionField<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<Tensor<1, dim>& values)const
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)</pre>
       values[i] = AdvectionField<dim>::value(points[i]);
template <int dim>
class RightHandSide : public Function<dim> {
public:
    RightHandSide() : Function<dim>() {}
    virtual double value(const Point<dim>& p,
                         const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                            std::vector<double>& values,
                            const unsigned int component = 0) const;
private:
    static const Point<dim> center_point;
template <>
const Point<1> RightHandSide<1>::center_point = Point<1>(-0.75);
const Point<2> RightHandSide<2>::center_point = Point<2>(-0.75, -0.75);
template <>
const Point<3> RightHandSide<3>::center_point = Point<3>(-0.75, -0.75, -0.75);
template <int dim>
double RightHandSide<dim>::value(const Point<dim>& p,
                                 const unsigned int component) const
    (void) component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
const double diameter = 0.1;
    return ((p - center_point).norm_square() < diameter * diameter</pre>
                ? .1 / std::pow(diameter, dim) : 0);
template <int dim>
void RightHandSide<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<double>& values,
                                     const unsigned int component) const
    Assert(values.size() == points.size(),
          ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)</pre>
       values[i] = RightHandSide<dim>::value(points[i], component);
template <int dim>
class BoundaryValues : public Function<dim> {
public:
    BoundaryValues() : Function<dim>() {}
    virtual void value_list(const std::vector<Point<dim>& points,
                            std::vector<double>& values,
                            const unsigned int component = 0) const;
template <int dim>
double BoundaryValues<dim>::value(const Point<dim>& p,
                                  const unsigned int component) const
    (void) component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
    const double sine_term =
       std::sin(16 * numbers::PI * std::sqrt(p.norm_square()));
    const double weight = std::exp(-5 * p.norm_square()) / std::exp(-5.);
    return sine_term * weight;
template <int dim>
void BoundaryValues<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<double>& values,
                                     const unsigned int component) const
    Assert(values.size() == points.size(),
          ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
  values[i] = BoundaryValues<dim>::value(points[i], component);
class GradientEstimation {
public:
    template <int dim>
    static void estimate(const DoFHandler<dim>& dof,
                         const Vector<double>& solution.
```

```
Vector<float>& error_per_cell);
    DeclException2(ExcInvalidVectorLength, int, int, "Vector has length " « arg1 « ", but should have "
                   « arg2);
    DeclException0 (ExcInsufficientDirections);
private:
    template <int dim>
    struct EstimateScratchData {
        EstimateScratchData(const FiniteElement<dim>& fe,
                            const Vector<double>& solution,
                            Vector<float>& error_per_cell);
        EstimateScratchData(const EstimateScratchData& data):
        FEValues<dim> fe_midpoint_value;
        const Vector<double>& solution;
        Vector<float>& error_per_cell;
    };
    struct EstimateCopyData {};
    template <int dim>
    static void estimate_cell(
        const typename DoFHandler<dim>::active_cell_iterator& cell,
        EstimateScratchData<dim>& scratch_data,
        const EstimateCopyData& copy_data);
template <int dim>
AdvectionProblem<dim>::AdvectionProblem() : dof_handler(triangulation), fe(1)
{ }
template <int dim>
AdvectionProblem < dim > :: AdvectionProblem ()
    dof handler.clear();
template <int dim>
void AdvectionProblem<dim>::setup_system()
    dof_handler.distribute_dofs(fe);
    hanging_node_constraints.clear();
    DoFTools::make_hanging_node_constraints(dof_handler,
                                             hanging_node_constraints);
    hanging_node_constraints.close();
    DynamicSparsityPattern dsp(dof_handler.n_dofs(), dof_handler.n_dofs());
    DoFTools::make_sparsity_pattern(dof_handler, dsp, hanging_node_constraints,
                                    /*keep_constrained_dofs = */ true);
    sparsity_pattern.copy_from(dsp);
    system_matrix.reinit(sparsity_pattern);
    solution.reinit(dof_handler.n_dofs());
    system_rhs.reinit(dof_handler.n_dofs());
template <int dim>
void AdvectionProblem<dim>::assemble_system()
    WorkStream::run(dof_handler.begin_active(), dof_handler.end(), *this,
                    &AdvectionProblem::local_assemble_system,
                    &AdvectionProblem::copy_local_to_global,
                    AssemblyScratchData(fe), AssemblyCopyData());
    hanging_node_constraints.condense(system_matrix);
    hanging_node_constraints.condense(system_rhs);
template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const FiniteElement<dim>& fe)
    : fe_values(fe, QGauss<dim>(2),
                update_values | update_gradients | update_quadrature_points |
      update_JxW_values),
fe_face_values(fe, QGauss<dim - 1>(2),
                     update_values | update_quadrature_points |
                         update_JxW_values | update_normal_vectors)
template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const AssemblyScratchData& scratch_data)
    : fe_values(scratch_data.fe_values.get_fe(),
                scratch_data.fe_values.get_quadrature(),
                fe_face_values(scratch_data.fe_face_values.get_fe(),
                     scratch_data.fe_face_values.get_quadrature(),
                     update_values | update_quadrature_points |
                         update_JxW_values | update_normal_vectors)
template <int dim>
void AdvectionProblem<dim>::local_assemble_system(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    AssemblyScratchData& scratch_data, AssemblyCopyData& copy_data)
    const AdvectionField<dim> advection_field;
    const RightHandSide<dim> right_hand_side;
const BoundaryValues<dim> boundary_values;
```

```
const unsigned int dofs_per_cell = fe.dofs_per_cell;
    const unsigned int n_q_points =
        scratch_data.fe_values.get_quadrature().size();
    const unsigned int n_face_q_points =
       scratch_data.fe_face_values.get_quadrature().size();
    copy_data.cell_matrix.reinit(dofs_per_cell, dofs_per_cell);
    copy_data.cell_rhs.reinit(dofs_per_cell);
    copy_data.local_dof_indices.resize(dofs_per_cell);
    std::vector<double> rhs_values(n_q_points);
    std::vector<Tensor<1, dim» advection_directions(n_q_points);
std::vector<double> face_boundary_values(n_face_q_points);
    std::vector<Tensor<1, dim» face_advection_directions(n_face_q_points);
    scratch_data.fe_values.reinit(cell);
    advection_field.value_list(scratch_data.fe_values.get_quadrature_points(),
                                advection_directions);
    right_hand_side.value_list(scratch_data.fe_values.get_quadrature_points(),
                                rhs_values);
    const double delta = 0.1 * cell->diameter();
    for (unsigned int q_point = 0; q_point < n_q_points; ++q_point)</pre>
        for (unsigned int i = 0; i < dofs_per_cell; ++i) {</pre>
            for (unsigned int j = 0; j < dofs_per_cell; ++j)</pre>
                copy_data.cell_matrix(i, j) +=
                     ((advection\_directions[q\_point] *
                       \verb|scratch_data.fe_values.shape_grad(j, q_point)| *
                       (scratch_data.fe_values.shape_value(i, q_point) +
                       delta *
                            (advection_directions[q_point] +
                             scratch_data.fe_values.shape_grad(i, q_point)))) *
                     scratch_data.fe_values.JxW(q_point));
            copy_data.cell_rhs(i) +=
                delta * (advection_directions[q_point] *
                            scratch_data.fe_values.shape_grad(i, q_point))) *
                 rhs_values[q_point] * scratch_data.fe_values.JxW(q_point));
    for (unsigned int face = 0; face < GeometryInfo<dim>::faces_per_cell;
         ++face)
        if (cell->face(face)->at_boundary()) {
            scratch_data.fe_face_values.reinit(cell, face);
            boundary_values.value_list(
                scratch_data.fe_face_values.get_quadrature_points(),
                face_boundary_values);
            advection field.value list (
                scratch_data.fe_face_values.get_quadrature_points(),
                face_advection_directions);
            for (unsigned int q_point = 0; q_point < n_face_q_points; ++q_point)</pre>
                if (scratch_data.fe_face_values.normal_vector(q_point) *
                         face_advection_directions[q_point] <</pre>
                     for (unsigned int i = 0; i < dofs_per_cell; ++i) {</pre>
                         for (unsigned int j = 0; j < dofs_per_cell; ++j)</pre>
                             copy_data.cell_matrix(i, j) -=
                                 (face_advection_directions[q_point] *
                                  scratch_data.fe_face_values.normal_vector(
                                      q_point) *
                                  scratch data.fe face values.shape value(
                                      i, q_point) *
                                  scratch_data.fe_face_values.shape_value(
                                      j, q_point) *
                                  \verb|scratch_data.fe_face_values.JxW(q_point)|;\\
                         copy_data.cell_rhs(i) -=
                             (face_advection_directions[q_point] *
                              scratch_data.fe_face_values.normal_vector(
                                  q_point) *
                              face_boundary_values[q_point] *
                              scratch_data.fe_face_values.shape_value(i,
                                                                        q_point) *
                              scratch_data.fe_face_values.JxW(q_point));
    cell->get_dof_indices(copy_data.local_dof_indices);
template <int dim>
void AdvectionProblem<dim>::copy_local_to_global(
    const AssemblyCopyData& copy_data)
    for (unsigned int i = 0; i < copy_data.local_dof_indices.size(); ++i) {</pre>
        for (unsigned int j = 0; j < copy_data.local_dof_indices.size(); ++j)</pre>
            system_matrix.add(copy_data.local_dof_indices[i],
                               copy_data.local_dof_indices[j],
        copy_data.cell_matrix(i, j));
system_rhs(copy_data.local_dof_indices[i]) += copy_data.cell_rhs(i);
template <int dim>
void AdvectionProblem<dim>::solve()
```

```
Assert(system_matrix.m() == system_matrix.n(), ExcNotQuadratic());
    auto num_rows = system_matrix.m();
    std::vector<double> rhs(num_rows);
    std::copy(system_rhs.begin(), system_rhs.begin() + num_rows, rhs.begin());
    using vec = gko::matrix::Dense<>;
using mtx = gko::matrix::Csr<>;
    using bicgstab = gko::solver::Bicgstab<>;
    using bj = gko::preconditioner::Jacobi<>;
    using val_array = gko::array<double>;
    std::shared_ptr<gko::Executor> exec = gko::ReferenceExecutor::create();
    auto b = vec::create(exec, gko::dim<2>(num_rows, 1),
                         val_array::view(exec, num_rows, rhs.data()), 1);
    auto x = vec::create(exec, gko::dim<2>(num_rows, 1));
auto A = mtx::create(exec, gko::dim<2>(num_rows),
                          system_matrix.n_nonzero_elements());
    mtx::value_type* values = A->get_values();
mtx::index_type* row_ptr = A->get_row_ptrs();
    mtx::index_type* col_idx = A->get_col_idxs();
    row_ptr[0] = 0;
    for (auto row = 1; row <= num_rows; ++row) {</pre>
        row_ptr[row] = row_ptr[row - 1] + system_matrix.get_row_length(row - 1);
    std::vector<mtx::index_type> ptrs(num_rows + 1);
    \verb|std::copy(A->get_row_ptrs(), A->get_row_ptrs() + num_rows + 1|,\\
              ptrs.begin());
    for (auto row = 0; row < system_matrix.m(); ++row) {</pre>
        for (auto p = system_matrix.begin(row); p != system_matrix.end(row);
             col_idx[ptrs[row]] = p->column();
            values[ptrs[row]] = p->value();
             ++ptrs[row];
        }
    auto solver_gen =
        bicgstab::build()
             .with criteria(
                 gko::stop::Iteration::build().with_max_iters(1000).on(exec),
                 gko::stop::ResidualNorm<>::build()
                     .with_reduction_factor(1e-12)
                     .on(exec))
             .with_preconditioner(bj::build().on(exec))
             .on(exec);
    auto solver = solver_gen->generate(gko::give(A));
    solver->apply(b, x);
    std::copy(x->get_values(), x->get_values() + num_rows, solution.begin());
/**************
* deal.ii internal solver. Here for reference.
SolverControl
                        solver_control (1000, 1e-12);
SolverBicgstab<>
                        bicgstab (solver_control);
PreconditionJacobi<> preconditioner;
preconditioner.initialize(system_matrix, 1.0);
bicgstab.solve (system_matrix, solution, system_rhs,
preconditioner);
    hanging_node_constraints.distribute(solution);
template <int dim>
void AdvectionProblem<dim>::refine_grid()
    Vector<float> estimated_error_per_cell(triangulation.n_active_cells());
    GradientEstimation::estimate(dof_handler, solution,
                                  estimated_error_per_cell);
    GridRefinement::refine_and_coarsen_fixed_number(
        triangulation, estimated_error_per_cell, 0.5, 0.03);
    triangulation.execute_coarsening_and_refinement();
template <int dim>
void AdvectionProblem<dim>::output_results(const unsigned int cycle)const
        GridOut grid_out;
        std::ofstream output("grid-" + std::to_string(cycle) + ".eps");
        grid_out.write_eps(triangulation, output);
        DataOut<dim> data_out;
        data_out.attach_dof_handler(dof_handler);
        data_out.add_data_vector(solution, "solution");
        data_out.build_patches();
        std::ofstream output("solution-" + std::to_string(cycle) + ".vtk");
        data_out.write_vtk(output);
template <int dim>
void AdvectionProblem < dim > :: run ()
```

```
for (unsigned int cycle = 0; cycle < 6; ++cycle) {
   std::cout w "Cycle " w cycle w ':' w std::endl;
   if (cycle == 0) {</pre>
                     {\tt GridGenerator::hyper\_cube} \ ({\tt triangulation, -1, 1}) \ ;
                     triangulation.refine_global(4);
              } else {
                    refine_grid();
              std::cout « " Number of active cells:
                                « triangulation.n_active_cells() « std::endl;
              setup_system();
std::cout « " Number of degrees of freedom: " « dof_handler.n_dofs()
              assemble_system();
              solve();
              output_results(cycle);
template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
       const FiniteElement<dim>& fe, const Vector<double>& solution,
       Vector<float>& error_per_cell)
       : fe_midpoint_value(fe, QMidpoint<dim>(),
                                          update_values | update_quadrature_points),
          solution(solution),
           error_per_cell(error_per_cell)
template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
       const EstimateScratchData& scratch data)
       : fe_midpoint_value(scratch_data.fe_midpoint_value.get_fe(),
                                          scratch_data.fe_midpoint_value.get_quadrature(),
                                           update_values | update_quadrature_points),
           solution(scratch_data.solution),
          error_per_cell(scratch_data.error_per_cell)
{ }
template <int dim>
void GradientEstimation::estimate(const DoFHandler<dim>& dof_handler,
                                                             const Vector<double>& solution,
                                                             Vector<float>& error_per_cell)
       Assert (error per cell.size() ==
                           dof_handler.get_triangulation().n_active_cells(),
                    ExcInvalidVectorLength(
                           error_per_cell.size(),
                           dof_handler.get_triangulation().n_active_cells()));
       \label{thm:workStream::run(dof_handler.begin\_active(), dof\_handler.end(),} WorkStream::run(dof\_handler.begin\_active(), dof\_handler.end(), dof\_ha
                                    &GradientEstimation::template estimate_cell<dim>,
                                    std::function<void(const EstimateCopyData&)>(),
                                    EstimateScratchData<dim>(dof_handler.get_fe(), solution,
                                                                                 error_per_cell),
                                    EstimateCopyData());
template <int dim>
void GradientEstimation::estimate_cell(
       const typename DoFHandler<dim>::active_cell_iterator& cell,
       EstimateScratchData<dim>& scratch_data, const EstimateCopyData&)
      Tensor<2, dim> Y:
       std::vector<typename DoFHandler<dim>::active cell iterator>
             active neighbors;
       active_neighbors.reserve(GeometryInfo<dim>::faces_per_cell *
                                                   GeometryInfo<dim>::max_children_per_face);
       scratch_data.fe_midpoint_value.reinit(cell);
       Tensor<1, dim> projected_gradient;
       active_neighbors.clear();
       for (unsigned int face_no = 0; face_no < GeometryInfo<dim>::faces_per_cell;
                ++face no)
              if (!cell->at_boundary(face_no)) {
                      const typename DoFHandler<dim>::face_iterator face =
                            cell->face(face_no);
                      const typename DoFHandler<dim>::cell_iterator neighbor =
                            cell->neighbor(face_no);
                      if (neighbor->active())
                            active_neighbors.push_back(neighbor);
                      else {
                             if (dim == 1) {
                                    typename DoFHandler<dim>::cell_iterator neighbor_child =
                                          neighbor:
                                    while (neighbor_child->has_children())
                                          neighbor_child =
                                                 neighbor_child->child(face_no == 0 ? 1 : 0);
                                    Assert (
                                           neighbor\_child->neighbor(face\_no == 0 ? 1 : 0) == cell,
                                           ExcInternalError());
                                    active_neighbors.push_back(neighbor_child);
```

```
} else
                       for (unsigned int subface_no = 0;
                             subface_no < face->n_children(); ++subface_no)
                            \verb"active_neighbors.push_back" (
                                cell->neighbor_child_on_subface(face_no,
                                                                      subface no));
    const Point<dim> this_center =
         scratch_data.fe_midpoint_value.quadrature_point(0);
    std::vector<double> this_midpoint_value(1);
    scratch_data.fe_midpoint_value.get_function_values(scratch_data.solution,
                                                                this_midpoint_value);
    std::vector<double> neighbor_midpoint_value(1);
    typename std::vector<typename DoFHandler<dim>::active_cell_iterator>::
         {\tt const\_iterator\ neighbor\_ptr = active\_neighbors.begin();}
    for (; neighbor_ptr != active_neighbors.end(); ++neighbor_ptr) {
    const typename DoFHandler<dim>::active_cell_iterator neighbor =
            *neighbor_ptr;
         scratch_data.fe_midpoint_value.reinit(neighbor);
         const Point<dim> neighbor_center =
             scratch_data.fe_midpoint_value.quadrature_point(0);
         {\tt scratch\_data.fe\_midpoint\_value.get\_function\_values} (
         scratch_data.solution, neighbor_midpoint_value);
Tensor<1, dim> y = neighbor_center - this_center;
const double distance = y.norm();
         y /= distance;
         for (unsigned int i = 0; i < dim; ++i)
    for (unsigned int j = 0; j < dim; ++j) Y[i][j] += y[i] * y[j];</pre>
         projected_gradient +=
              (neighbor_midpoint_value[0] - this_midpoint_value[0]) / distance *
    AssertThrow(determinant(Y) != 0, ExcInsufficientDirections());
    const Tensor<2, dim> Y_inverse = invert(Y);
    Tensor<1, dim> gradient = Y_inverse * projected_gradient;
    scratch_data.error_per_cell(cell->active_cell_index()) =
   (std::pow(cell->diameter(), 1 + 1.0 * dim / 2) *
          std::sqrt(gradient.norm_square()));
   // namespace Step9
int main()
         dealii::MultithreadInfo::set_thread_limit();
         Step9::AdvectionProblem<2> advection_problem_2d;
         advection_problem_2d.run();
    } catch (std::exception& exc) {
         std::cerr « std::endl
                    « std::endl
                     « std::endl;
         std::cerr « "Exception on processing: " « std::endl
                    « exc.what() « std::endl
« "Aborting!" « std::endl
                     « std::endl;
         return 1;
    } catch (...)
         std::cerr « std::endl
                    « std::endl
                    « "-
                     « std::endl;
         std::cerr « "Unknown exception!" « std::endl
                     « "Aborting!" « std::endl
                    « "---
                    « std::endl;
         return 1;
    return 0;
```

The ginkgo-overhead program

The ginkgo overhead measurement example..

Introduction

About the example

```
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <cmath>
#include <iostream>
[[noreturn]] void print_usage_and_exit(const char* name)
    std::cerr « "Usage: " « name « " [NUM_ITERS]" « std::endl;
    std::exit(-1);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
long unsigned num_iters = 1000000;
    if (argc > 2) {
        print_usage_and_exit(argv[0]);
    if (argc == 2) {
    num_iters = std::atol(argv[1]);
         if (num iters == 0) {
            print_usage_and_exit(argv[0]);
    std::cout « gko::version_info::get() « std::endl;
    auto exec = gko::ReferenceExecutor::create();
    auto cg_factory =
        cg::build()
             .with_criteria(
                 gko::stop::Iteration::build().with_max_iters(num_iters).on(
                      exec))
             .on(exec);
    auto A = gko::initialize<mtx>({1.0}, exec);
auto b = gko::initialize<vec>({std::nan("")}, exec);
    auto x = gko::initialize<vec>({0.0}, exec);
auto tic = std::chrono::steady_clock::now();
    auto solver = cg_factory->generate(gko::give(A));
    solver->apply(x, b);
    exec->synchronize();
    auto tac = std::chrono::steady_clock::now();
    auto time = std::chrono::duration_cast<std::chrono::nanoseconds>(tac - tic);
    std::cout « "Running " « num_iters
```

This is the expected output:

```
Running 1000000 iterations of the CG solver took a total of 1.60337 seconds.

Average library overhead: 1603.37 [nanoseconds / iteration]
```

Comments about programming and debugging

```
**********GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
    Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <cmath>
#include <iostream>
[[noreturn]] void print_usage_and_exit(const char* name)
    std::cerr « "Usage: " « name « " [NUM_ITERS]" « std::endl;
    std::exit(-1);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    long unsigned num_iters = 1000000;
    if (argc > 2) {
        print_usage_and_exit(argv[0]);
    if (argc == 2) {
   num_iters = std::atol(argv[1]);
        <u>if</u> (num_iters == 0) {
            print_usage_and_exit(argv[0]);
```

```
}
     std::cout « gko::version_info::get() « std::endl;
     auto exec = gko::ReferenceExecutor::create();
     auto cg_factory =
         cg::build()
              .with_criteria(
                    gko::stop::Iteration::build().with_max_iters(num_iters).on(
                        exec))
               .on(exec);
    auto A = gko::initialize<mtx>({1.0}, exec);
auto b = gko::initialize<vec>({std::nan("")}, exec);
    auto x = gko::initialize<vec>({su.:nam(),})
auto x = gko::initialize<vec>({su.:nam(),})
auto tic = std::chrono::steady_clock::now();
     auto solver = cg_factory->generate(gko::give(A));
     solver->apply(x, b);
     exec->synchronize();
    auto tame = std::chrono::duration_cast<std::chrono::nanoseconds>(tac - tic);
     std::cout « "Running " « num_iters
                 " iterations of the CG solver took a total of "
« static_cast<double>(time.count()) /
                  static_cast<double>(std::nano::den)
« " seconds." « std::endl
                  "\tAverage library overhead:
« static_cast<double>(time.count()) /
                          static_cast<double>(num_iters)
                  « " [nanoseconds / iteration]" « std::endl;
}
```

The ginkgo-ranges program

The ranges and accessor example..

Introduction

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <iomanip>
#include <iostream>
```

LU factorization implementation using Ginkgo ranges For simplicity, we only consider square matrices, and no pivoting.

```
template <typename Accessor>
void factorize(const gko::range<Accessor>& A)
```

note: const means that the range (i.e. the data handler) is constant, not that the underlying data is constant! $^{\{}$

```
using gko::span;
assert(A.length(0) == A.length(1));
for (gko::size_type i = 0; i < A.length(0) - 1; ++i) {
   const auto trail = span{i + 1, A.length(0)};
```

note: neither of the lines below need additional memory to store intermediate arrays, all computation is done at the point of assignment

```
A(trail, i) = A(trail, i) / A(i, i);
```

a utility function for printing the factorization on screen

```
template <typename Accessor>
void print_lu(const gko::range<Accessor>& A)
{
    std::cout « std::setprecision(2) « std::fixed;
    std::cout « "L = [";
    for (int i = 0; i < A.length(0); ++i) {
        std::cout « "\n ";
        for (int j = 0; j < A.length(1); ++j) {
            std::cout « (i > j ? A(i, j) : (i == j) * 1.) « " ";
        }
    }
    std::cout « "\n]\n\nU = [";
```

```
for (int i = 0; i < A.length(0); ++i) {
    std::cout « "\n ";
    for (int j = 0; j < A.length(1); ++j) {
        std::cout « (i <= j ? A(i, j) : 0.) « " ";
    }
} std::cout « "\n]" « std::endl;
}
int main(int argc, char* argv[]) {
    using ValueType = double;
    using IndexType = int;</pre>
```

Print version information

std::cout « gko::version_info::get() « std::endl;

Create some test data, add some padding just to demonstrate how to use it with ranges. clang-format off

```
ValueType data[] = {
    2., 4., 5., -1.0,
    4., 11., 12., -1.0,
    6., 24., 24., -1.0
}:
```

clang-format on

Create a 3-by-3 range, with a 2D row-major accessor using data as the underlying storage. Set the stride (a.k.a. "LDA") to 4.

```
auto A =
    gko::range<gko::accessor::row_major<ValueType, 2>>(data, 3u, 3u, 4u);
```

use the LU factorization routine defined above to factorize the matrix

factorize(A);

```
print the factorization on screen
```

```
print_lu(A);
```

Results

This is the expected output:

```
L = [
    1.00 0.00 0.00
    2.00 1.00 0.00
]
U = [
    2.00 4.00 5.00
    0.00 3.00 2.00
    0.00 0.00 1.00
```

Comments about programming and debugging

```
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <iomanip>
#include <iostream>
template <typename Accessor>
void factorize(const gko::range<Accessor>& A)
    using gko::span;
    assert (A.length(0) == A.length(1));
    for (gko::size_type i = 0; i < A.length(0) - 1; ++i) {
         const auto trail = span{i + 1, A.length(0)};
A(trail, i) = A(trail, i) / A(i, i);
         A(trail, trail) = A(trail, trail) - mmul(A(trail, i), A(i, trail));
template <typename Accessor>
void print_lu(const gko::range<Accessor>& A)
    std::cout « std::setprecision(2) « std::fixed;
std::cout « "L = [";
for (int i = 0; i < A.length(0); ++i) {
    std::cout « "\n ";</pre>
         for (int j = 0; j < A.length(1); ++j) {</pre>
            std::cout « (i > j ? A(i, j) : (i == j) * 1.) « " ";
    std::cout « "\n]\n\nU = [";
for (int i = 0; i < A.length(0); ++i) {
    std::cout « "\n ";</pre>
         for (int j = 0; j < A.length(1); ++j) {
             std::cout « (i <= j ? A(i, j) : 0.) « " ";
    std::cout « "\n]" « std::endl;
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    std::cout « gko::version_info::get() « std::endl;
    ValueType data[] = {
        2., 4., 5., -1.0,
4., 11., 12., -1.0,
6., 24., 24., -1.0
    }:
    auto A =
        gko::range<gko::accessor::row_major<ValueType, 2>>(data, 3u, 3u, 4u);
    factorize(A);
    print_lu(A);
```

The heat-equation program

The heat equation example..

This example depends on simple-solver, three-pt-stencil-solver.

Introduction

This example solves a 2D heat conduction equation

$$u: [0, d]^2 \to R$$

 $\partial_t u = \delta u + f$

with Dirichlet boundary conditions and given initial condition and constant-in-time source function f.

The partial differential equation (PDE) is solved with a finite difference spatial discretization on an equidistant grid: For n grid points, and grid distance h=1/n we write

$$u_{i,j}' = \alpha \frac{u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1} - 4u_{i,j}}{h^2} + f_{i,j}$$

We then build an implicit Euler integrator by discretizing with time step au

$$\frac{u_{i,j}^{k+1} - u_{i,j}^k}{\tau} = \alpha \frac{u_{i-1,j}^{k+1} + u_{i+1,j}^{k+1} - u_{i,j-1}^{k+1} - u_{i,j+1}^{k+1} + 4u_{i,j}^{k+1}}{h^2} + f_{i,j}$$

and solve the resulting linear system for u^{k+1} using Ginkgo's CG solver preconditioned with an incomplete Cholesky factorization for each time step, occasionally writing the resulting grid values into a video file using OpenCV and a custom color mapping.

The intention of this example is to provide a mini-app showing matrix assembly, vector initialization, solver setup and the use of Ginkgo in a more complex setting.

About the example

```
This example solves a 2D heat conduction equation
    u: [0, d]^2 \rightarrow R\\
\partial_t u = \delta u + f
with Dirichlet boundary conditions and given initial condition and
constant-in-time source function f.
The partial differential equation (PDE) is solved with a finite difference
spatial discretization on an equidistant grid: For 'n' grid points,
and grid distance h = 1/n we write
    u_{i,j}' = \alpha_{i,j}' = \alpha_{i,j+1} + u_{i,j+1} + u_{i,j+1} + u_{i,j+1}
 4 u_{i,j}) / h^2
+ f_{i,j}
We then build an implicit Euler integrator by discretizing with time step \tau
(u_{i,j}^{k+1} - u_{i,j}^k) / tau =
+ f_{i,j}
and solve the resulting linear system for u_{\langle cdot \rangle^{k+1}} using Ginkgo's CG solver preconditioned with an incomplete Cholesky factorization for each time
step, occasionally writing the resulting grid values into a video file using
OpenCV and a custom color mapping.
The intention of this example is to provide a mini-app showing matrix assembly,
vector initialization, solver setup and the use of Ginkgo in a more complex
setting.
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <fstream>
#include <iostream>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
This function implements a simple Ginkgo-themed clamped color mapping for values in the range [0,5].
void set_val(unsigned char* data, double value)
RGB values for the 6 colors used for values 0, 1, ..., 5 We will interpolate linearly between these values.
double col_r[] = \{255, 221, 129, 201, 249, 255\}; double col_g[] = \{255, 220, 130, 161, 158, 204\};
double col_b[] = {255, 220, 133, 93, 24, 8};
value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
auto d = std::max(0.0, std::min(1.0, value - i));
OpenCV uses BGR instead of RGB by default, revert indices
    data[0] = static_cast < unsigned char > (col_b[i + 1] * d + col_b[i] * (1 - d));
Initialize video output with given dimension and FPS (frames per seconds)
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
    cv::Size videosize{n, n};
    auto output =
       std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', '1');
    output.first.open("heat.mp4", fourcc, fps, videosize);
    return output;
}
Write the current frame to video output using the above color mapping
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                     const double* data)
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
        for (int j = 0; j < n; j++) {
    set_val(&row[3 * j], data[i * n + j]);</pre>
    output.first.write(output.second);
```

```
int main(int argc, char* argv[])
    using mtx = gko::matrix::Csr<>;
    using vec = gko::matrix::Dense<>;
Problem parameters: simulation length
diffusion factor
auto diffusion = 0.0005;
scaling factor for heat source
auto source_scale = 2.5;
Simulation parameters: inner grid points per discretization direction
auto n = 256;
number of simulation steps per second
auto steps_per_sec = 500;
number of video frames per second
auto fps = 25;
number of grid points
auto n2 = n * n;
grid point distance (ignoring boundary points)
auto h = 1.0 / (n + 1);
auto h2 = h * h;
time step size for the simulation
auto tau = 1.0 / steps_per_sec;
create a CUDA executor with an associated OpenMP host executor
auto exec = gko::CudaExecutor::create(0, gko::OmpExecutor::create());
load heat source and initial state vectors
std::ifstream initial_stream("data/gko_logo_2d.mtx");
std::ifstream source_stream("data/gko_text_2d.mtx");
auto source = gko::read<vec>(source_stream, exec);
auto in_vector = gko::read<vec>(initial_stream, exec);
create output vector with initial guess for
auto out_vector = in_vector->clone();
create scalar for source update
auto tau_source_scalar = gko::initialize<vec>({source_scale * tau}, exec);
create stencil matrix as shared_ptr for solver
auto stencil_matrix = gko::share(mtx::create(exec));
assemble matrix
gko::matrix_data<> mtx_data{gko::dim<2>(n2, n2)};
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
  auto c = i * n + j;
  auto c_val = diffusion * tau * 4.0 / h2 + 1.0;</pre>
        auto off_val = -diffusion \star tau / h2;
for each grid point: insert 5 stencil points with Dirichlet boundary conditions, i.e. with zero boundary value
            mtx_data.nonzeros.emplace_back(c, c - n, off_val);
            mtx_data.nonzeros.emplace_back(c, c - 1, off_val);
        mtx_data.nonzeros.emplace_back(c, c, c_val);
            mtx_data.nonzeros.emplace_back(c, c + 1, off_val);
```

if (i < n - 1) {

mtx_data.nonzeros.emplace_back(c, c + n, off_val);

```
}
stencil_matrix->read(mtx_data);
prepare video output
auto output = build_output(n, fps);
build CG solver on stencil with incomplete Cholesky preconditioner stopping at 1e-10 relative accuracy
    gko::solver::Cg<>::build()
        .with_preconditioner(gko::preconditioner::Ic<>::build().on(exec))
        .with_criteria(gko::stop::ResidualNorm<>::build()
                            .with_baseline(gko::stop::mode::rhs_norm)
                            .with reduction factor(1e-10)
                            .on(exec))
         ->generate(stencil_matrix);
time stamp of the last output frame (initialized to a sentinel value)
execute implicit Euler method: for each timestep, solve stencil system
for (double t = 0; t < t0; t += tau) {
if enough time has passed, output the next video frame
if (t - last_t > 1.0 / fps)
    last_t = t;
    std::cout « t « std::endl;
    output_timestep(
        output, n,
        gko::make_temporary_clone(exec->get_master(), in_vector)
            ->get_const_values());
add heat source contribution
in_vector->add_scaled(tau_source_scalar, source);
execute Euler step
solver->apply(in_vector, out_vector);
swap input and output
        std::swap(in_vector, out_vector);
```

The program will generate a video file named heat.mp4 and output the timestamp of each generated frame.

Comments about programming and debugging

```
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES: LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
This example solves a 2D heat conduction equation
u : [0, d]^2 \rightarrow R 
\partial_t u = \beta u + f
with Dirichlet boundary conditions and given initial condition and
constant-in-time source function f.
The partial differential equation (PDE) is solved with a finite difference
spatial discretization on an equidistant grid: For 'n' grid points,
and grid distance @f$h = 1/n@f$ we write
u_{i,j}' = \alpha(u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1})
 4 u_{i,j}) / h^2
+ f_{i,j}
We then build an implicit Euler integrator by discretizing with time step @f$\tau@f$
u_{i,j-1}^{k+1} - u_{i,j+1}^{k+1} - 4 u_{i,j}^{k+1}) / h^2
and solve the resulting linear system for Qf$ u_{\cdot}^{k+1}Qf$ using Ginkgo's CG and
solver preconditioned with an incomplete Cholesky factorization for each time
step, occasionally writing the resulting grid values into a video file using
OpenCV and a custom color mapping.
The intention of this example is to provide a mini-app showing matrix assembly,
vector initialization, solver setup and the use of {\tt Ginkgo} in a more complex
settina.
#include <ginkgo/ginkgo.hpp>
#include <chrono>
#include <fstream>
#include <iostream>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
void set_val(unsigned char* data, double value)
    double col_r[] = {255, 221, 129, 201, 249, 255}; double col_g[] = {255, 220, 130, 161, 158, 204};
    double col_b[] = {255, 220, 133, 93, 24, 8};
    value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
    auto d = std::max(0.0, std::min(1.0, value - i));
    data[2] = static_cast < unsigned char > (col_r[i + 1] * d + col_r[i] * (1 - d));
    \label{eq:data} \texttt{data[1] = static\_cast} < \texttt{unsigned char} > (\texttt{col\_g[i + 1]} * \texttt{d + col\_g[i]} * (1 - \texttt{d}));
    std::pair<cv::VideoWriter, cv::Mat> build output(int n, double fps)
    cv::Size videosize{n, n};
    auto output =
    std::make_pair(cv::VideoWriter{}), cv::Mat{videosize, CV_8UC3});
auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', '1');
   output.first.open("heat.mp4", fourcc, fps, videosize);
    return output;
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                    const double* data)
    for (int i = 0; i < n; i++) {</pre>
       auto row = output.second.ptr(i);
for (int j = 0; j < n; j++) {
            set_val(&row[3 * j], data[i * n + j]);
    output.first.write(output.second);
```

```
int main(int argc, char* argv[])
    using mtx = gko::matrix::Csr<>;
using vec = gko::matrix::Dense<>;
auto t0 = 5.0;
    auto diffusion = 0.0005;
    auto source_scale = 2.5;
    auto n = 256;
    auto steps_per_sec = 500;
    auto fps = 25;
auto n2 = n * n;
    auto h = 1.0 / (n + 1);
    auto h2 = h * h;
    auto tau = 1.0 / steps_per_sec;
    auto exec = gko::CudaExecutor::create(0, gko::OmpExecutor::create());
std::ifstream initial_stream("data/gko_logo_2d.mtx");
    std::ifstream source_stream("data/gko_text_2d.mtx");
    auto source = gko::read<vec>(source_stream, exec);
    auto in_vector = gko::read<vec>(initial_stream, exec);
    auto out_vector = in_vector->clone();
    auto tau_source_scalar = gko::initialize<vec>((source_scale * tau), exec);
auto stencil_matrix = gko::share(mtx::create(exec));
    gko::matrix_data<> mtx_data{gko::dim<2>(n2, n2)};
for (int i = 0; i < n; i++) {</pre>
        for (int j = 0; j < n; j++)
auto c = i * n + j;
             auto c_val = diffusion * tau * 4.0 / h2 + 1.0;
             auto off_val = -diffusion * tau / h2;
             if (i > 0) {
                  mtx data.nonzeros.emplace back(c, c - n, off val);
                  mtx_data.nonzeros.emplace_back(c, c - 1, off_val);
             mtx_data.nonzeros.emplace_back(c, c, c_val);
             if (j < n - 1) {
                  mtx_data.nonzeros.emplace_back(c, c + 1, off_val);
             if (i < n - 1) {</pre>
                  mtx_data.nonzeros.emplace_back(c, c + n, off_val);
        }
    stencil_matrix->read(mtx_data);
    auto output = build_output(n, fps);
    auto solver =
        gko::solver::Cg<>::build()
             .with_preconditioner(gko::preconditioner::Ic<>::build().on(exec))
             .with_criteria(gko::stop::ResidualNorm<>::build()
                                   .with_baseline(gko::stop::mode::rhs_norm)
                                   .with_reduction_factor(1e-10)
                                    .on(exec))
             .on(exec)
             ->generate(stencil_matrix);
    double last_t = -t0;
for (double t = 0; t < t0; t += tau) {</pre>
         if (t - last_t > 1.0 / fps) {
             last_t = t;
             std::cout « t « std::endl;
             output timestep(
                  output, n,
gko::make_temporary_clone(exec->get_master(), in_vector)
                       ->get_const_values());
         in_vector->add_scaled(tau_source_scalar, source);
        solver->apply(in_vector, out_vector);
         std::swap(in_vector, out_vector);
```

The ilu-preconditioned-solver program

The ILU-preconditioned solver example..

This example depends on simple-solver.

Introduction

About the example

exec_map{

This example shows how to use incomplete factors generated via the ParILU algorithm to generate an incomplete factorization (ILU) preconditioner, how to specify the sparse triangular solves in the ILU preconditioner application, and how to generate an ILU-preconditioned solver and apply it to a specific problem.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <cstdlib>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using gmres = gko::solver::Gmres<ValueType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
```

{"omp", [] { return gko::OmpExecutor::create(); }},

```
{"cuda",
           [] {
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                       true):
          {"hip",
           [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          { "dpcpp",
           [] {
               return gko::DpcppExecutor::create(0,
                                                        gko::OmpExecutor::create());
          {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
Generate incomplete factors using ParILU
auto par_ilu_fact :
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
Generate concrete factorization for input matrix
```

Generate an ILU preconditioner factory by setting lower and upper triangular solver - in this case the exact triangular solves

Use incomplete factors to generate ILU preconditioner

auto par_ilu = gko::share(par_ilu_fact->generate(A));

auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));

Use preconditioner inside GMRES solver factory Generating a solver factory tied to a specific preconditioner makes sense if there are several very similar systems to solve, and the same solver+preconditioner combination is expected to be effective.

Generate preconditioned solver for a specific target system

auto ilu_gmres = ilu_gmres_factory->generate(A);

Solve system

```
ilu_gmres->apply(b, x);
```

Print solution

```
std::cout « "Solution (x):\n";
write(std::cout, x);
```

Calculate residual

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

```
This is the expected output: Solution (x):
```

```
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1.46249e-08
```

Comments about programming and debugging

```
******GINKGO LICENSE>*****************
 Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
          Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
            Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
 IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
 TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
 PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
 DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
 #include <ginkgo/ginkgo.hpp>
  #include <cstdlib>
 #include <fstream>
 #include <iostream>
 #include <map>
 #include <string>
 int main(int argc, char* argv[])
              using ValueType = double;
              using RealValueType = gko::remove_complex<ValueType>;
              using IndexType = int;
              using vec = gko::matrix::Dense<ValueType>;
              using real_vec = gko::matrix::Dense<RealValueType>;
              using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

```
using gmres = gko::solver::Gmres<ValueType>;
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec_map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
          [] {
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
         {"hip",
          [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"dpcpp",
              return gko::DpcppExecutor::create(0,
                                                       gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::readstd::ifstream("data/b.mtx"), exec);
auto x = gko::readvec>(std::ifstream("data/x0.mtx"), exec);
auto par ilu fact =
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
auto par_ilu = gko::share(par_ilu_fact->generate(A));
auto ilu_pre_factory =
    gko::preconditioner::Ilu<gko::solver::LowerTrs<ValueType, IndexType>,
                                  gko::solver::UpperTrs<ValueType, IndexType>,
                                  false>::build()
         .on(exec);
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
const RealValueType reduction_factor{1e-7};
auto ilu_gmres_factory =
    gmres::build()
         .with criteria(
              gko::stop::Iteration::build().with_max_iters(1000u).on(exec),
              gko::stop::ResidualNorm<ValueType>::build()
                   .with_reduction_factor(reduction_factor)
                   .on(exec))
         . \verb|with_generated_preconditioner|| (\verb|ilu_preconditioner||)
         .on(exec);
auto ilu_gmres = ilu_gmres_factory->generate(A);
ilu_gmres->apply(b, x);
std::cout « "Solution (x):\n";
write(std::cout, x);
auto one = gko::initialize<vec>((1.0), exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

The inverse-iteration program

The inverse iteration example..

This example depends on simple-solver, .

Introduction

This example shows how components available in Ginkgo can be used to implement higher-level numerical methods. The method used here will be the shifted inverse iteration method for eigenvalue computation which find the eigenvalue and eigenvector of A closest to z, for some scalar z. The method requires repeatedly solving the shifted linear system (A - zI)x = b, as well as performing matrix-vector products with the matrix A. Here is the complete pseudocode of the method:

```
x_0 = initial guess
for i = 0 .. max_iterations:
    solve (A - zI) y_i = x_i for y_i+1
    x_(i+1) = y_i / || y_i || # compute next eigenvector approximation
    g_(i+1) = x_(i+1)^* A x_(i+1) # approximate eigenvalue (Rayleigh quotient)
    if ||A x_(i+1) - g_(i+1)x_(i+1)|| < tol * g_(i+1): # check convergence</pre>
```

About the example

```
#include <ginkgo/ginkgo.hpp>
#include <cmath>
#include <complex>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using precision = std::complex<double>;
using real_precision = gko::remove_complex<precision>;
using vec = gko::matrix::Dense<precision>;
using real_vec = gko::matrix::Dense<real_precision>;
using mtx = gko::matrix::Csr<precision>;
using solver_type = gko::solver::Bicgstab<precision>;
using std::abs;
using std::sqrt;
```

```
Print version information
```

```
std::cout « gko::version_info::get() « std::endl;
std::cout « std::scientific « std::setprecision(8) « std::showpos;
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
         {"cuda",
          [] {
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                   true);
          }},
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         { "dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                    gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto this_exec = exec->get_master();
linear system solver parameters
auto system_max_iterations = 100u;
auto system_residual_goal = real_precision{le-16};
eigensolver parameters
auto max_iterations = 20u;
auto residual_goal = real_precision{le-8};
auto z = precision{20.0, 2.0};
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Generate shifted matrix A - zI

• we avoid duplicating memory by not storing both A and A - zI, but compute A - zI on the fly by using Ginkgo's utilities for creating linear combinations of operators

```
auto one = share(gko::initialize<vec>({precision{1.0}}), exec));
auto neg_one = share(gko::initialize<vec>({-precision{1.0}}, exec));
auto neg_z = gko::initialize<vec>((-z), exec);
auto system_matrix = share(gko::Combination<precision>::create(
    one, A, gko::initialize<vec>({-z}, exec),
    gko::matrix::Identity<precision>::create(exec, A->get_size()[0])));
Generate solver operator (A - zI)^-1
auto solver =
    solver_type::build()
        .with_criteria(gko::stop::Iteration::build()
                           .with_max_iters(system_max_iterations)
                            .on(exec),
                       gko::stop::ResidualNorm<precision>::build()
                           .with_reduction_factor(system_residual_goal)
                            .on(exec))
        .on(exec)
        ->generate(system matrix);
inverse iterations
start with guess [1, 1, ..., 1]
auto x = [\&] {
    auto work = vec::create(this_exec, gko::dim<2>{A->get_size()[0], 1});
    const auto n = work->get_size()[0];
```

```
for (int i = 0; i < n; ++i) {
         work->get_values()[i] = precision(1.0) / sqrt(n);
     return clone(exec, work);
}();
auto v = clone(x);
auto tmp = clone(x);
auto norm = gko::initialize<real_vec>({1.0}, exec);
auto inv_norm = clone(this_exec, one);
auto g = clone(one);
for (auto i = Ou; i < max_iterations; ++i) {
    std::cout « "{ ";</pre>
(A - zI)y = x
solver->apply(x, y);
system_matrix->apply(one, y, neg_one, x);
x->compute_norm2(norm);
std::cout « "\"system_residual\": "
           « clone(this_exec, norm)->get_values()[0] « ", ";
x->copy_from(y);
x = y / || y ||
x->compute_norm2(norm);
inv_norm->get_values()[0] =
     real_precision{1.0} / clone(this_exec, norm)->get_values()[0];
x->scale(clone(exec, inv_norm));
g = x^{\wedge} * A x
A->apply(x, tmp);
x->compute_dot(tmp, g);
auto g_val = clone(this_exec, g)->get_values()[0]; std::cout « "\"eigenvalue\": " « g_val « ", ";
||Ax - gx|| < tol * g
          auto v = gko::initialize<vec>({-g_val}, exec);
         tmp->add scaled(v, x);
         tmp->compute_norm2(norm);
         auto res_val = clone(exec->get_master(), norm)->get_values()[0];
std::cout « "\"residual\": " « res_val / g_val « " }," « std::e
         if (abs(res_val) < residual_goal * abs(g_val)) {</pre>
              break;
    }
```

```
This is the expected output:
```

```
+1.61736920e-14, "eigenvalue": (+2.03741410e+01,-1.17744356e-16), "residual":
{ "system_residual":
      (+2.92231055e-01,+1.68883476e-18) },
{ "system_residual": +4.98014795e-15, "eigenvalue": (+1.94878474e+01,+1.25948378e-15), "residual":
(+7.94370276e-02,-5.13395071e-18) }, { "system_residual": +3.39296916e-15, "eigenvalue": (+1.93282121e+01,-1.19329332e-15), "residual":
(+4.11149623e-02,+2.53837290e-18) }, { "system_residual": +3.35953656e-15, "eigenvalue": (+1.92638912e+01,+3.28657016e-16), "residual":
      (+2.34717040e-02,-4.00445585e-19) },
{ "system_residual": +2.91474009e-15,
                                         "eigenvalue": (+1.92409166e+01,+3.65597737e-16), "residual":
      (+1.34709547e-02,-2.55962367e-19) },
{ "system_residual": +3.09863953e-15, "eigenvalue": (+1.92331106e+01,-1.07919176e-15), "residual":
(+7.72060707e-03,+4.33212063e-19) }, { "system_residual": +2.31198069e-15, "eigenvalue": (+1.92305014e+01,-2.89755360e-16), "residual":
      (+4.42106625e-03,+6.66143651e-20) },
{ "system_residual": +3.02771202e-15, "eigenvalue": (+1.92296339e+01,+8.04259901e-16), "residual":
      (+2.53081312e-03,-1.05848687e-19) },
{ "system_residual": +2.02954523e-15, "eigenvalue": (+1.92293461e+01,+7.81834016e-16), "residual":
(+1.44862114e-03,-5.88985854e-20) }, { "system_residual": +2.31762332e-15, "eigenvalue": (+1.92292506e+01,-1.11718775e-16), "residual":
      stem_residual": +2.31702333 --, (+8.29183451e-04,+4.81741912e-21) }, (+8.29183451e-04,+4.81741912e-21) }, **residual": +8.12541038e-15, "eigenvalue": (+1.92292190e+01,-6.55606254e-16), "residual":
{ "system_residual":
      (+4.74636702e-04,+1.61823936e-20) },
{ "system_residual": +2.77259926e-15, "eigenvalue": (+1.92292085e+01,+4.30588140e-16), "residual":
      (+2.71701077e-04,-6.08403935e-21) },
(+8.90457139e-05,+2.09737561e-21) },
```

*********GINKGO LICENSE>*****************

Comments about programming and debugging

```
Copyright (c) 2017-2023, the Ginkgo authors
 All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
          Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
 contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
 IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
 TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
 DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
 #include <ginkgo/ginkgo.hpp>
  #include <cmath>
 #include <complex>
 #include <fstream>
 #include <iomanip>
 #include <iostream>
 #include <map>
 #include <string>
 int main(int argc, char* argv[])
             using precision = std::complex<double>;
            using real_precision = gko::remove_complex<precision>;
using vec = gko::matrix::Dense<precision>;
             using real_vec = gko::matrix::Dense<real_precision>;
             using mtx = gko::matrix::Csr<precision>;
             using solver_type = gko::solver::Bicgstab<precision>;
             using std::abs;
             using std::sqrt;
             std::cout « gko::version_info::get() « std::endl;
             std::cout « std::scientific « std::setprecision(8) « std::showpos;
             if (argc == 2 && (std::string(argy[1]) == "-help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
                         std::exit(-1);
             const auto executor_string = argc >= 2 ? argv[1] : "reference";
             std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
                         exec map{
                                      { "omp"
                                                            [] { return gko::OmpExecutor::create(); }},
                                      {"omp",
{"cuda",
                                         [] {
                                                     return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                                                                                                                           true);
                                         }},
```

```
{"hip",
          []
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                    true);
          11.
         { "dpcpp",
          [] {
               return gko::DpcppExecutor::create(0,
                                                      gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto this_exec = exec->get_master();
auto system_max_iterations = 100u;
auto system_residual_goal = real_precision{1e-16};
auto max_iterations = 20u;
auto residual_goal = real_precision{1e-8};
auto z = precision\{20.0, 2.0\};
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto one = share(gko::initialize<vec>({precision{1.0}}), exec));
auto neg_one = share(gko::initialize<vec>({-precision{1.0}}, exec));
auto neg_z = gko::initialize<vec>({-z}, exec);
auto system_matrix = share(gko::Combinationprecision>::create(
    one, A, gko::initialize<vec>({-z}, exec),
    gko::matrix::Identity<precision>::create(exec, A->get_size()[0])));
auto solver =
     solver_type::build()
         .with_criteria(gko::stop::Iteration::build()
                               .with_max_iters(system_max_iterations)
                               .on(exec),
                           gko::stop::ResidualNorm<precision>::build()
                               .with_reduction_factor(system_residual_goal)
         ->generate(system_matrix);
auto x = [\&]
    auto work = vec::create(this_exec, gko::dim<2>{A->get_size()[0], 1}); const auto n = work->get_size()[0];
     for (int i = 0; i < n; ++i) {
         work->get_values()[i] = precision{1.0} / sqrt(n);
    return clone (exec, work);
}():
auto y = clone(x);
auto tmp = clone(x);
auto norm = gko::initialize<real_vec>({1.0}, exec);
auto inv_norm = clone(this_exec, one);
auto g = clone(one);
for (auto i = Ou; i < max_iterations; ++i) {
    std::cout « "{ ";</pre>
     solver->apply(x, y);
     system_matrix->apply(one, y, neg_one, x);
     x->compute_norm2(norm);
     \verb|std::cout| « "\"system_residual\": "
               « clone(this_exec, norm)->get_values()[0] « ", ";
     x->copy_from(y);
     x->compute_norm2(norm);
     inv_norm->get_values()[0]
         real_precision{1.0} / clone(this_exec, norm)->get_values()[0];
     x->scale(clone(exec, inv_norm));
    A->apply(x, tmp);
     x->compute_dot(tmp, g);
     auto g_val = clone(this_exec, g)->get_values()[0];
std::cout « "\"eigenvalue\": " « g_val « ", ";
     auto v = gko::initialize<vec>({-g_val}, exec);
     tmp->add_scaled(v, x);
     tmp->compute_norm2(norm);
    auto res_val = clone(exec->get_master(), norm)->get_values()[0]; std::cout « "\"residual\": " « res_val / g_val « " }," « std::endl;
     if (abs(res_val) < residual_goal * abs(g_val)) {</pre>
}
```

The ir-ilu-preconditioned-solver program

The IR-ILU preconditioned solver example..

This example depends on ilu-preconditioned-solver, iterative-refinement.

Introduction

About the example

This example shows how to combine iterative refinement with the adaptive precision block-Jacobi preconditioner in order to approximately solve the triangular systems occurring in ILU preconditioning. Using an adaptive precision block-Jacobi preconditioner matrix as inner solver for the iterative refinement method is equivalent to doing adaptive precision block-Jacobi relaxation in the triangular solves. This example roughly approximates the triangular solves with five adaptive precision block-Jacobi sweeps with a maximum block size of 16.

This example is motivated by "Multiprecision block-Jacobi for Iterative Triangular Solves" (Göbel, Anzt, Cojean, Flegar, Quintana-Ortí, Euro-Par 2020). The theory and a detailed analysis can be found there.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <cstdlib>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using gmres = gko::solver::Gmres<ValueType>;
using ir = gko::solver::Ir<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
```

Print version information

```
std::cout « gko::version_info::get() « std::endl;
```

```
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int sweeps = argc == 3 ? std::atoi(argv[2]) : 5u;
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
          [] {
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                  true);
          }},
         {"dpcpp",
              return gko::DpcppExecutor::create(0,
                                                    gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS and initial guess as 1
gko::size_type num_rows = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(num_rows, 1));
for (gko::size_type i = 0; i < num_rows; i++) {</pre>
    host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto clone_x = gko::clone(exec, x);
Generate incomplete factors using ParILU
auto par_ilu_fact
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
```

Generate concrete factorization for input matrix

auto par_ilu = gko::share(par_ilu_fact->generate(A));

Generate an iterative refinement factory to be used as a triangular solver in the preconditioner application. The generated method is equivalent to doing five block-Jacobi sweeps with a maximum block size of 16.

Generate an ILU preconditioner factory by setting lower and upper triangular solver - in this case the previously defined iterative refinement method.

```
auto ilu_pre_factory =
    gko::preconditioner::Ilu<ir, ir>::build()
        .with_l_solver_factory(gko::clone(trisolve_factory))
        .with_u_solver_factory(gko::clone(trisolve_factory))
        .on(exec);
```

Use incomplete factors to generate ILU preconditioner

```
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
```

Create stopping criteria for Gmres

```
const RealValueType reduction_factor{1e-12};
auto iter_stop = gko::share(
```

Use preconditioner inside GMRES solver factory Generating a solver factory tied to a specific preconditioner makes sense if there are several very similar systems to solve, and the same solver+preconditioner combination is expected to be effective.

```
to be effective.
auto ilu_gmres_factory =
    gmres::build()
        .with_criteria(iter_stop, tol_stop)
         . \verb|with_generated_preconditioner| (ilu\_preconditioner)|\\
        .on(exec);
Generate preconditioned solver for a specific target system
auto ilu_gmres = ilu_gmres_factory->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
    gko::log::Convergence<ValueType>::create();
ilu_gmres->add_logger(logger);
Warmup run
ilu_gmres->apply(b, x);
Solve system 100 times and take the average time.
std::chrono::nanoseconds time(0);
for (int i = 0; i < 100; i++) {
    x->copy_from(clone_x);
    auto tic = std::chrono::high_resolution_clock::now();
    ilu_gmres->apply(b, x);
    auto toc = std::chrono::high_resolution_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
std::cout « "Using " « sweeps « " block-Jacobi sweeps.\n";
Print solution
std::cout « "Solution (x):\n";
write(std::cout, x);
Get residual
    auto res = gko::as<vec>(logger->get_residual_norm());
    std::cout « "GMRES iteration count:
                                                     ' « logger->get num iterations()
              « "\n";
    std::cout « "GMRES execution time [ms]: "
```

« static_cast<double>(time.count()) / 100000000.0 « "\n";

Results

}

This is the expected output:

write(std::cout, res);

```
Using 5 block-Jacobi sweeps. Solution (x):
%%MatrixMarket matrix array real general
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
GMRES iteration count:
GMRES execution time [ms]: 0.377673
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.65303e-12
```

std::cout « "Residual norm sqrt($r^T r$):n";

Comments about programming and debugging

```
*********GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
    Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <cstdlib>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using gmres = gko::solver::Gmres<ValueType>;
    using ir = gko::solver::Ir<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
         std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int sweeps = argc == 3 ? std::atoi(argv[2]) : 5u;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         exec_map{
             { "omp"
                     [] { return gko::OmpExecutor::create(); }},
             {"cuda",
              [] {
                  return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
              }},
             {"hip",
                  return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
             { "dpcpp",
              [] {
                  return gko::DpcppExecutor::create(0,
                                                        gko::OmpExecutor::create());
             {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type num_rows = A->get_size()[0];
    auto host x = vec::create(exec->get master(), gko::dim<2>(num rows, 1));
    for (gko::size_type i = 0; i < num_rows; i++)</pre>
        host_x->at(i, 0) = 1.;
```

```
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto clone_x = gko::clone(exec, x);
auto par_ilu_fact =
   gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
auto par_ilu = gko::share(par_ilu_fact->generate(A));
auto bj_factory = gko::share(
   bj::build()
        .with_max_block_size(16u)
        . \verb|with_storage_optimization(gko::precision_reduction::autodetect())|\\
        .on(exec));
auto trisolve_factory
    ir::build()
        .with_solver(bj_factory)
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(sweeps).on(exec))
        .on(exec);
auto ilu_pre_factory =
    gko::preconditioner::Ilu<ir, ir>::build()
        .with_l_solver_factory(gko::clone(trisolve_factory))
        .with_u_solver_factory(gko::clone(trisolve_factory))
        .on(exec);
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
const RealValueType reduction_factor{le-12};
auto iter_stop = gko::share(
    gko::stop::Iteration::build().with_max_iters(1000u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                .with_reduction_factor(reduction_factor)
                                .on(exec));
auto ilu_gmres_factory =
    gmres::build()
        .with_criteria(iter_stop, tol_stop)
        . \verb|with_generated_preconditioner|| (\verb|ilu_preconditioner||) \\
.on(exec);
auto ilu_gmres = ilu_gmres_factory->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
ilu_gmres->add_logger(logger);
ilu_gmres->apply(b, x);
std::chrono::nanoseconds time(0);
for (int i = 0; i < 100; i++) {
    x->copy_from(clone_x);
    auto tic = std::chrono::high_resolution_clock::now();
    ilu_gmres->apply(b, x);
    auto toc = std::chrono::high_resolution_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
std::cout « "Using " « sweeps « " block-Jacobi sweeps.\n";
std::cout « "Solution (x):\n";
write(std::cout, x);
auto res = gko::as<vec>(logger->get_residual_norm());
" « logger->get_num_iterations()
std::cout « "GMRES execution time [ms]: "
         « static_cast<double>(time.count()) / 100000000.0 « "\n";
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

The iterative-refinement program

The iterative refinement solver example..

This example depends on simple-solver.

This example shows how to use the iterative refinement solver.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. An inaccurate CG solver is used as the inner solver to an iterative refinement (IR) method which solves a linear system. The example features the iteration count and runtime of the IR solver.

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using ir = gko::solver::Ir<ValueType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
          {"omp", [] { return gko::OmpExecutor::create(); }},
          {"cuda",
                return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                        true);
```

```
}},
         {"hip",
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                 true):
          }}.
         {"dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                   gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS and initial guess as 1
gko::size_type size = A->get_size()[0];
auto host_x = gko::matrix::Dense<ValueType>::create(exec->get_master(),
                                                        gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
   host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_x);
Create solver factory
gko::size_type max_iters = 10000u;
RealValueType outer_reduction_factor{1e-12};
RealValueType inner_reduction_factor{1e-2};
auto solver_gen =
    ir::build()
        .with_solver(
             cq::build()
                 .with_criteria(
                      gko::stop::ResidualNorm<ValueType>::build()
                          .with_reduction_factor(inner_reduction_factor)
                          .on(exec))
                  .on(exec))
         .with_criteria(
             gko::stop::Iteration::build().with_max_iters(max_iters).on(
                 exec),
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_reduction_factor(outer_reduction_factor)
                  .on(exec))
         .on(exec);
Create solver
auto solver = solver_gen->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
Get residual
auto res = gko::as<real_vec>(logger->get_residual_norm());
std::cout « "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
Print solver statistics
    std::cout « "IR iteration count:
                                                  " « logger->get_num_iterations()
    « static_cast<double>(time.count()) / 1000000.0 « std::endl;
}
```

This is the expected output: Initial residual norm sqrt(r^T r): %%MatrixMarket matrix array real general 1 1 194.679 Final residual norm sqrt(r^T r): %%MatrixMarket matrix array real general 1 1 4.23821e-11 IR iteration count: 24 IR execution time [ms]: 0.794962

Comments about programming and debugging

```
************** GINKGO LICENSE>******************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double:
    using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
using ir = gko::solver::Ir<ValueType>;
    istd::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
         std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         exec map{
             {"omp",
{"cuda",
                     [] { return gko::OmpExecutor::create(); }},
              [] {
                   return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                       true);
              }},
```

}

```
{"hip",
         []
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                   true);
          }},
         { "dpcpp",
          [] {
              return gko::DpcppExecutor::create(0,
                                                    gko::OmpExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // thrws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type size = A->get_size()[0];
auto host_x = gko::matrix::Dense<ValueType>::create(exec->get_master(),
                                                          gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
   host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto initres = gko::initialize<real_vec>((0.0), exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_x);
gko::size_type max_iters = 10000u;
RealValueType outer_reduction_factor{1e-12};
RealValueType inner_reduction_factor{1e-2};
auto solver gen =
    ir::build()
        .with_solver(
             cg::build()
                  . \verb|with_criteria|| (
                      gko::stop::ResidualNorm<ValueType>::build()
                          .with_reduction_factor(inner_reduction_factor)
                           .on(exec))
         .with_criteria(
             gko::stop::Iteration::build().with_max_iters(max_iters).on(
                 exec),
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_reduction_factor(outer_reduction_factor)
                  .on(exec))
         .on(exec);
auto solver = solver_gen->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::as<real_vec>(logger->get_residual_norm());
std::cout « "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
std::cout « "IR iteration count:
                                               " « logger->get_num_iterations()
          « std::endl;
std::cout « "IR execution time [ms]: "
           « static_cast<double>(time.count()) / 1000000.0 « std::endl;
```

The kokkos_assembly program

The Kokkos assembly example..

This example depends on simple-solver, poisson-solver, three-pt-stencil-solver, .

Introduction

This example solves a 1D Poisson equation:

$$u:[0,1]\to Ru"=fu(0)=u0u(1)=u1$$

using a finite difference method on an equidistant grid with $\mathbb K$ discretization points ($\mathbb K$ can be controlled with a command line parameter).

The resulting CSR matrix is assembled using Kokkos kernels. This example show how to use Ginkgo data with Kokkos kernels.

Notes

If this example is built as part of Ginkgo, it is advised to configure Ginkgo with <code>-DGINKGO_WITH_CCACHE=OFF</code> to prevent incompabilities with Kokkos' compiler wrapper for <code>nvcc</code>.

The commented program

```
#include <iostream>
#include <map>
#include <string>
#include <omp.h>
#include <Kokkos_Core.hpp>
#include <ginkgo/ginkgo.hpp>
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
   auto exec = matrix->get_executor();
   const auto discretization_points = matrix->get_size()[0];
```

Over-allocate storage for the matrix elements. Each row has 3 entries, except for the first and last one. To handle each row uniformly, we allocate space for 3x the number of rows.

Create Kokkos views on Ginkgo data.

```
Kokkos::View<IndexType*> v_row_idxs(md.get_row_idxs(), md.get_num_elems());
Kokkos::View<IndexType*> v_col_idxs(md.get_col_idxs(), md.get_num_elems());
Kokkos::View<ValueType*> v_values(md.get_values(), md.get_num_elems());
```

Create the matrix entries. This also creates zero entries for the first and second row to handle all rows uniformly.

```
Kokkos::parallel_for(
   "generate_stencil_matrix", md.get_num_elems(), KOKKOS_LAMBDA(int i) {
      const ValueType coefs[] = {-1, 2, -1};
      auto ofs = static_cast<IndexType>((i % 3) - 1);
      auto row = static_cast<IndexType>(i / 3);
      auto col = row + ofs;
```

To prevent branching, a mask is used to set the entry to zero, if the column is out-of-bounds

Add up duplicate (zero) entries.

```
md.sum_duplicates();
```

Build Csr matrix.

```
matrix->read(std::move(md));
```

Generates the RHS vector given f and the boundary conditions.

Computes the 1-norm of the error given the computed u and the correct solution function correct_u.

```
template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                            const gko::matrix::Dense<ValueType>* u,
                            Closure&& correct_u)
     Kokkos::View<const ValueType*> v_u(u->get_const_values(),
                                                discretization_points);
     auto error = 0.0;
     Kokkos::parallel_reduce(
          "calculate_error", discretization_points,
KOKKOS_LAMBDA(int i, double& lsum) {
               const auto h = 1.0 / (discretization_points + 1); const auto xi = (i + 1) * h;
               lsum += Kokkos::Experimental::abs(
                    (v u(i) - correct u(xi)) /
                    Kokkos::Experimental::abs(correct_u(xi)));
          error);
     return error;
int main(int argc, char* argv[])
     Kokkos::ScopeGuard kokkos(argc, argv);
```

```
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType>;
Print help message. For details on the kokkos-options see https://kokkos.github.io/kokkos-core-wiki/←
ProgrammingGuide/Initialization.html#initialization-by-command-line-arguments
std::exit(-1);
const unsigned int discretization_points =
    argc >= 2 ? std::atoi(argv[1]) : 100u;
chooses the executor that corresponds to the Kokkos DefaultExecutionSpace
   auto exec = []() -> std::shared_ptr<gko::Executor> {
#ifdef KOKKOS_ENABLE_SERIAL
        if (std::is_same<Kokkos::DefaultExecutionSpace,</pre>
                        Kokkos::Serial>::value)
            return gko::ReferenceExecutor::create();
#endif
#ifdef KOKKOS_ENABLE_OPENMP
        if (std::is_same<Kokkos::DefaultExecutionSpace,</pre>
                        Kokkos::OpenMP>::value) {
            return gko::OmpExecutor::create();
#endif
#ifdef KOKKOS_ENABLE_CUDA
        if (std::is_same<Kokkos::DefaultExecutionSpace, Kokkos::Cuda>::value) {
            return gko::CudaExecutor::create(0,
                                             gko::ReferenceExecutor::create());
#endif
#ifdef KOKKOS_ENABLE_HIP
        if (std::is_same<Kokkos::DefaultExecutionSpace, Kokkos::HIP>::value) {
            return gko::HipExecutor::create(0,
                                            gko::ReferenceExecutor::create());
#endif
}();
problem:
auto correct_u = [] KOKKOS_FUNCTION(ValueType x) { return x * x * x; };
auto f = [] KOKKOS_FUNCTION(ValueType x) { return ValueType{6} * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
initialize vectors
auto rhs = vec::create(exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
initialize the stencil matrix
auto A = share(mtx::create(
    exec, gko::dim<2>{discretization_points, discretization_points}));
generate_stencil_matrix(A.get());
const RealValueType reduction_factor{1e-7};
Generate solver and solve the system
    cg::build()
        .with_criteria(gko::stop::Iteration::build()
                           .with max iters(discretization points)
                           .on(exec),
                       gko::stop::ResidualNorm<ValueType>::build()
                          .with_reduction_factor(reduction_factor)
                           .on(exec))
        .with_preconditioner(bj::build().on(exec))
        .on(exec)
        ->generate(A)
        ->apply(rhs, u);
    std::cout « "\nSolve complete."

« "\nThe average relative error is "
              « calculate_error(discretization_points, u.get(), correct_u) /
                    discretization_points
              « std::endl;
```

}

Example output:

```
> ./kokkos-assembly
Solve complete.
The average relative error is 1.05488e-11
```

The actual error depends on the used hardware.

```
*********GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
           Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
 IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
  #include <iostream>
 #include <map>
 #include <string>
 #include <omp.h>
 #include <Kokkos_Core.hpp>
 #include <ginkgo/ginkgo.hpp>
 template <typename ValueType, typename IndexType>
 void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
            auto exec = matrix->get_executor();
            const auto discretization_points = matrix->get_size()[0];
gko::device_matrix_data<ValueType, IndexType> md(exec, matrix->get_size(),
                                                                                                                                                             discretization_points * 3);
            Kokkos::View<IndexType*> v_row_idxs(md.get_row_idxs(), md.get_num_elems());
            Kokkos::View<IndexType*> v_col_idxs(md.get_col_idxs(), md.get_num_elems());
            Kokkos::View<ValueType*> v_values(md.get_values(), md.get_num_elems());
            Kokkos::parallel for(
                          "generate_stencil_matrix", md.get_num_elems(), KOKKOS_LAMBDA(int i) {
    const ValueType coefs[] = {-1, 2, -1};
    auto ofs = static_cast<IndexType>((i % 3) - 1);
                                    auto row = static_cast<IndexType>(i / 3);
                                    auto col = row + ofs;
                                   auto mask =
                                              static_cast<IndexType>(0 <= col && col < discretization_points);</pre>
                                    v_row_idxs[i] = mask * row;
v_col_idxs[i] = mask * col;
                                    v_values[i] = mask * coefs[ofs + 1];
                       });
            md.sum_duplicates();
            matrix->read(std::move(md));
 template <typename Closure, typename ValueType>
 void generate_rhs(Closure&& f, ValueType u0, ValueType u1,
                                                     gko::matrix::Dense<ValueType>* rhs)
            const auto discretization points = rhs->get size()[0]:
            auto values = rhs->get_values();
            Kokkos::View<ValueType*> values_view(values, discretization_points);
```

```
Kokkos::parallel_for(
        "generate_rhs", discretization_points, KOKKOS_LAMBDA(int i) {
            const ValueType h = 1.0 / (discretization_points + 1);
const ValueType xi = ValueType(i + 1) * h;
             values\_view[i] = -f(xi) * h * h;
             if (i == 0) {
                 values_view[i] += u0;
             if (i == discretization_points - 1) {
                 values_view[i] += u1;
        });
template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                         const gko::matrix::Dense<ValueType>* u,
                        Closure&& correct u)
{
    Kokkos::View<const ValueType*> v_u(u->get_const_values(),
                                         discretization_points);
    auto error = 0.0;
    Kokkos::parallel_reduce(
         "calculate_error", discretization_points,
        KOKKOS_LAMBDA(int i, double& lsum) {
  const auto h = 1.0 / (discretization_points + 1);
  const auto xi = (i + 1) * h;
             lsum += Kokkos::Experimental::abs(
                 (v_u(i) - correct_u(xi)) /
                 Kokkos::Experimental::abs(correct_u(xi)));
        }.
        error);
    return error;
int main(int argc, char* argv[])
    Kokkos::ScopeGuard kokkos(argc, argv);
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType>;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::exit(-1);
    const unsigned int discretization_points =
    argc >= 2 ? std::atoi(argv[1]) : 100u;
auto exec = []() -> std::shared_ptr<gko::Executor> {
#ifdef KOKKOS_ENABLE_SERIAL
        if (std::is_same<Kokkos::DefaultExecutionSpace,</pre>
                          Kokkos::Serial>::value) {
            return gko::ReferenceExecutor::create();
#endif
#ifdef KOKKOS_ENABLE_OPENMP
        if (std::is_same<Kokkos::DefaultExecutionSpace,</pre>
                          Kokkos::OpenMP>::value) {
            return gko::OmpExecutor::create();
        }
#endif
#ifdef KOKKOS ENABLE CUDA
        if (std::is_same<Kokkos::DefaultExecutionSpace, Kokkos::Cuda>::value) {
             return gko::CudaExecutor::create(0,
                                                gko::ReferenceExecutor::create());
#endif
#ifdef KOKKOS_ENABLE_HIP
        if (std::is_same<Kokkos::DefaultExecutionSpace, Kokkos::HIP>::value) {
             return gko::HipExecutor::create(0,
                                                gko::ReferenceExecutor::create());
#endif
    }();
    auto correct_u = [] KOKKOS_FUNCTION(ValueType x) { return x * x * x; };
    auto f = [] KOKKOS_FUNCTION(ValueType x) { return ValueType{6} * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    auto rhs = vec::create(exec, gko::dim<2>(discretization_points, 1));
    generate_rhs(f, u0, u1, rhs.get());
    auto u = vec::create(exec, gko::dim<2>(discretization_points, 1));
    for (int i = 0; i < u->get_size()[0]; ++i) {
        u->get_values()[i] = 0.0;
    auto A = share(mtx::create(
```

The minimal-cuda-solver program

The minimal CUDA solver example..

This example depends on simple-solver.

Introduction

This is a minimal example that solves a system with Ginkgo. The matrix, right hand side and initial guess are read from standard input, and the result is written to standard output. The system matrix is stored in CSR format, and the system solved using the CG method, preconditioned with the block-Jacobi preconditioner. All computations are done on the GPU.

The easiest way to use the example data from the data/ folder is to concatenate the matrix, the right hand side and the initial solution (in that exact order), and pipe the result to the minimal_solver_cuda executable:

```
cat data/A.mtx data/b.mtx data/x0.mtx \mid ./minimal-cuda-solver
```

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <iostream>
int main()
Instantiate a CUDA executor
auto gpu = gko::CudaExecutor::create(0, gko::OmpExecutor::create(), true);
auto A = gko::read<gko::matrix::Csr<»(std::cin, gpu);
auto b = gko::read<gko::matrix::Dense<»(std::cin, gpu);</pre>
auto x = gko::read<gko::matrix::Dense<>(std::cin, gpu);
Create the solver
auto solver
    gko::solver::Cg<>::build()
        .with_preconditioner(gko::preconditioner::Jacobi<>::build().on(gpu))
             gko::stop::Iteration::build().with_max_iters(20u).on(gpu),
             gko::stop::ResidualNorm<>::build()
                .with_reduction_factor(1e-15)
                  .on(gpu))
         .on(gpu);
Solve system
solver->generate(give(A))->apply(b, x);
Write result
    write(std::cout, x);
```

The following is the expected result when using the data contained in the folder data as input:

```
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
```

Comments about programming and debugging

```
**********GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <iostream>
int main()
    auto gpu = gko::CudaExecutor::create(), gko::OmpExecutor::create(), true);
    auto A = gko::read<gko::matrix::Csr<>(std::cin, gpu);
    auto b = gko::read<gko::matrix::Dense<> (std::cin, gpu);
    auto x = gko::read<gko::matrix::Dense<>(std::cin, gpu);
    auto solver =
        gko::solver::Cg<>::build()
            .with_preconditioner(gko::preconditioner::Jacobi<>::build().on(gpu))
            .with criteria(
                gko::stop::Iteration::build().with_max_iters(20u).on(gpu),
                gko::stop::ResidualNorm<>::build()
                    .with_reduction_factor(1e-15)
                     .on(gpu))
            .on(gpu);
    solver->generate(give(A))->apply(b, x);
    write(std::cout, x);
```

The mixed-multigrid-preconditioned-solver program

The preconditioned solver example..

This example depends on multigrid-preconditioned-solver, mixed-multigrid-solver.

This example shows how to use the mixed-precision multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a mixed-precision multigrid preconditioner. The example features the generating time and runtime of the CG solver.

The commented program

exec_map{

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using MixedType = float;
using IndexType = int;
using wee = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using fcg = gko::solver::Fcg<ValueType>;
using cg = gko::solver::Cg<ValueType>;
using ir = gko::solver::Ir<ValueType>;
using mg = gko::solver::Multigrid;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
using cg_f = gko::solver::Cg<MixedType>;
using ir_f = gko::solver::Ir<MixedType>;
using bj_f = gko::preconditioner::Jacobi<MixedType, IndexType>;
using pgm_f = gko::multigrid::Pgm<MixedType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
```

std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»

```
{"omp", [] { return gko::OmpExecutor::create(); }},
          []
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                    true):
          }}.
         {"hip",
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          }}.
         {"dpcpp",
          [] {
              return gko::DpcppExecutor::create(
                   0, gko::ReferenceExecutor::create());
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)();
                                                         // throws if not valid
const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout « "Using mixed precision? " « use_mixed « std::endl;
Read data
auto A = \text{share}(gko::read < mtx > (std::ifstream("data/A.mtx"). exec)):
Create RHS as 1 and initial guess as 0
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_b);
Prepare the stopping criteria
const gko::remove_complex<ValueType> tolerance = 1e-8;
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                   .with_baseline(gko::stop::mode::absolute)
                                   .with_reduction_factor(tolerance)
                                   .on(exec));
std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
Create smoother factory (ir with bj)
auto inner_solver_gen
    gko::share(bj::build().with_max_block_size(lu).on(exec));
auto inner_solver_gen_f =
    gko::share(bj_f::build().with_max_block_size(lu).on(exec));
auto smoother_gen = gko::share(
    ir::build()
         .with_solver(inner_solver_gen)
         .with_relaxation_factor(static_cast<ValueType>(0.9))
         .with_criteria(
             gko::stop::Iteration::build().with_max_iters(1u).on(exec))
         .on(exec));
auto smoother_gen_f = gko::share(
    ir_f::build()
         .with_solver(inner_solver_gen_f)
         .with_relaxation_factor(static_cast<MixedType>(0.9))
         .with_criteria(
             gko::stop::Iteration::build().with_max_iters(1u).on(exec))
```

```
.on(exec));
```

```
Create MultigridLevel factory
```

auto mg_level_gen

```
gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen_f =
    gko::share(pgm_f::build().with_deterministic(true).on(exec));
```

Create CoarsestSolver factory

Create multigrid factory

The first (index 0) level will use the first mg_level_gen, smoother_gen which are the factories with ValueType. The rest of levels (>= 1) will use the second (index 1) mg_level_gen2 and smoother_gen2 which use the MixedType. The rest of levels will use different type than the normal multigrid.

```
return level >= 1 ? 1 :
            .with_coarsest_solver(coarsest_gen_f)
            .with_default_initial_guess(
                gko::solver::initial_guess_mode::zero)
            .with criteria(
                gko::stop::Iteration::build().with_max_iters(1u).on(exec))
            .on(exec);
} else {
    multigrid_gen =
        mg::build()
            .with_max_levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen)
            .with_coarsest_solver(coarsest_gen)
            . \verb|with_default_initial_guess|| (
                gko::solver::initial_guess_mode::zero)
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(1u).on(exec))
}
```

Create solver factory

```
auto solver_gen = cg::build()
    .with_criteria(iter_stop, tol_stop)
    .with_preconditioner(multigrid_gen)
    .on(exec);
```

Create solver

```
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
```

Solve system

```
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steadv clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
Calculate residual
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
Print solver statistics
    std::cout « "CG iteration count:
                                               " « logger->get_num_iterations()
             « std::endl;
    std::cout « "CG generation time [ms]: "
             « static_cast<double>(gen_time.count()) / 1000000.0 « std::endl;
    std::cout « "CG execution time [ms]: "
              « static_cast<double>(time.count()) / 1000000.0 « std::endl;
    std::cout « "CG execution time per iteraion[ms]:
              « static_cast<double>(time.count()) / 1000000.0 /
                     logger->get_num_iterations()
              « std::endl;
```

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.3589
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.69858e-09
CG iteration count: 39
CG generation time [ms]: 2.04293
CG execution time [ms]: 22.3874
CG execution time per iteraion[ms]: 0.574036
```

Comments about programming and debugging

```
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
```

```
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double;
    using MixedType = float;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using fcg = gko::solver::Fcg<ValueType>;
    using cg = gko::solver::Cg<ValueType>;
using ir = gko::solver::Ir<ValueType>;
    using mg = gko::solver::Multigrid;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    using cg_f = gko::solver::Cg<MixedType>;
    using ir_f = gko::solver::Ir<MixedType>;
    using bj_f = gko::preconditioner::Jacobi<MixedType, IndexType>;
using pgm_f = gko::multigrid::Pgm<MixedType, IndexType>;
    std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] :
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         exec_map{
             {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
              [] {
                   return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
             {"hip",
                   return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
               }},
             {"dpcpp",
                   return gko::DpcppExecutor::create(
                       0, gko::ReferenceExecutor::create());
              {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout « "Using mixed precision? " « use_mixed « std::endl;
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {</pre>
        host_x - at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    auto x = vec::create(exec);
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>((-1.0), exec);
    auto initres = gko::initialize<vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_b);
    const gko::remove_complex<ValueType> tolerance = 1e-8;
    auto iter stop =
        gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
    auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                       .with_baseline(gko::stop::mode::absolute)
                                        .with_reduction_factor(tolerance)
                                        .on(exec)):
    std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
        gko::log::Convergence<ValueType>::create();
    iter_stop->add_logger(logger);
    tol_stop->add_logger(logger);
    auto inner_solver_gen =
        gko::share(bj::build().with_max_block_size(lu).on(exec));
    auto inner solver gen f =
```

```
gko::share(bj_f::build().with_max_block_size(1u).on(exec));
auto smoother_gen = gko::share(
    ir::build()
        . \verb|with_solver(inner_solver_gen)|\\
        .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
auto smoother_gen_f = gko::share(
    ir_f::build()
        .with_solver(inner_solver_gen_f)
        .with_relaxation_factor(static_cast<MixedType>(0.9))
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
        .on(exec));
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg level gen f =
   gko::share(pgm_f::build().with_deterministic(true).on(exec));
auto coarsest_gen = gko::share(
    ir::build()
        .with_solver(inner_solver_gen)
        .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with criteria(
            gko::stop::Iteration::build().with_max_iters(4u).on(exec))
        .on(exec));
auto coarsest_gen_f = gko::share(
    ir_f::build()
        .with_solver(inner_solver_gen_f)
        .with_relaxation_factor(static_cast<MixedType>(0.9))
        .with criteria(
            gko::stop::Iteration::build().with_max_iters(4u).on(exec))
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
   multigrid_gen =
       mq::build()
            .with_max_levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen, smoother_gen_f)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen, mg_level_gen_f)
            return level >= 1 ? 1 : 0;
            })
            .with_coarsest_solver(coarsest_gen_f)
            . \verb|with_default_initial_guess|| (
                gko::solver::initial_guess_mode::zero)
            .with criteria(
               gko::stop::Iteration::build().with_max_iters(lu).on(exec))
} else {
   multigrid_gen =
        mg::build()
            .with max levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen)
            . \verb|with_coarsest_solver(coarsest_gen)|\\
            .with_default_initial_guess(
                gko::solver::initial_guess_mode::zero)
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(1u).on(exec))
            .on(exec);
auto solver_gen = cg::build()
                      .with_criteria(iter_stop, tol_stop)
                      .with_preconditioner(multigrid_gen)
                      .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
   std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
exec->synchronize();
std::chrono::nanoseconds time(0):
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
```

The mixed-multigrid-preconditioned-solver	program

The mixed-multigrid-solver program

The mixed multigrid solver example..

This example depends on simple-solver.

This example shows how to use the mixed-precision multigrid solver.

In this example, we first read in a matrix from a file, then generate a right-hand side and an initial guess. The multigrid solver can mix different precision of MultigridLevel. The example features the generating time and runtime of the multigrid solver.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using MixedType = float;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using fcg = gko::solver::Fcg<ValueType>;
using cg = gko::solver::Cg<MixedType>;
using ir = gko::solver::Ir<ValueType>;
using ir2 = gko::solver::Ir<MixedType>;
using mg = gko::solver::Multigrid;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using bj2 = gko::preconditioner::Jacobi<MixedType, IndexType>;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
using pgm2 = gko::multigrid::Pgm<MixedType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
     exec_map{
           {"omp", [] { return gko::OmpExecutor::create(); }},
           {"cuda",
```

```
return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"hip",
          [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         { "dpcpp",
          [] {
               return gko::DpcppExecutor::create(
                   0, gko::ReferenceExecutor::create());
          }},
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
                                                           // throws if not valid
const auto exec = exec_map.at(executor_string)();
const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout « "Using mixed precision? " « use_mixed « std::endl;
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS as 1 and initial guess as 0
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
host_b->at(i, 0) = 1.;
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>((1.0), exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto initres = gko::initialize<vec>((0.0), exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_b);
Prepare the stopping criteria
const gko::remove_complex<ValueType> tolerance = 1e-12;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with max iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                   .with_baseline(gko::stop::mode::absolute)
                                    .with_reduction_factor(tolerance)
                                    .on(exec));
Create smoother factory (ir with bj)
auto smoother_gen = gko::share(
    ir::build()
         . \verb|with_solver(bj::build().with_max_block_size(1u).on(exec))|\\
         .with_relaxation_factor(static_cast<ValueType>(0.9))
         .with criteria(
             gko::stop::Iteration::build().with_max_iters(1u).on(exec))
         .on(exec));
auto smoother_gen2 = gko::share(
    ir2::build()
         .with_solver(bj2::build().with_max_block_size(1u).on(exec))
.with_relaxation_factor(static_cast<MixedType>(0.9))
         .with_criteria(
              gko::stop::Iteration::build().with_max_iters(1u).on(exec))
         .on(exec));
Create RestrictProlong factory
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen2 =
    gko::share(pgm2::build().with_deterministic(true).on(exec));
```

Create CoarsesSolver factory

```
auto coarsest_solver_gen = gko::share(
        .with_solver(bj::build().with_max_block_size(lu).on(exec))
        .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with criteria(
            gko::stop::Iteration::build().with max iters(4u).on(exec))
        .on(exec));
auto coarsest_solver_gen2 = gko::share(
    ir2::build()
        .with_solver(bj2::build().with_max_block_size(1u).on(exec))
        .with\_relaxation\_factor(static\_cast < MixedType > (0.9))\\
        .with_criteria(
            gko::stop::Iteration::build().with max iters(4u).on(exec))
        .on(exec));
Create multigrid factory
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
    multigrid_gen =
        mg::build()
            .with_max_levels(10u)
            .with min coarse rows(2u)
            .with_pre_smoother(smoother_gen, smoother_gen2)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen, mg_level_gen2)
            .with_level_selector([](const gko::size_type level,
                                     const gko::LinOp*) -> gko::size_type {
```

The first (index 0) level will use the first mg_level_gen, smoother_gen which are the factories with ValueType. The rest of levels (>= 1) will use the second (index 1) mg_level_gen2 and smoother_gen2 which use the MixedType. The rest of levels will use different type than the normal multigrid.

```
return level >= 1 ? 1 : 0;
            .with_coarsest_solver(coarsest_solver_gen2)
            .with_criteria(iter_stop, tol_stop)
            .on(exec);
} else {
   multigrid_gen = mg::build()
                         .with_max_levels(10u)
                         .with_min_coarse_rows(2u)
                         .with_pre_smoother(smoother_gen)
                         .with_post_uses_pre(true)
                         .with_mg_level(mg_level_gen)
                         .with coarsest solver(coarsest solver gen)
                         .with_criteria(iter_stop, tol_stop)
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
auto solver = multigrid_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
```

Calculate residual explicitly, because the residual is not available inside of the multigrid solver

```
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
```

Print solver statistics

Results

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.3589
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
6.31088e-14
Multigrid iteration count: 9
Multigrid generation time [ms]: 3.35361
Multigrid execution time [ms]: 10.048
Multigrid execution time per iteraion[ms]: 1.11644
```

Comments about programming and debugging

```
******GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
   Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double:
    using MixedType = float;
    using IndexType = int;
```

```
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using fcg = gko::solver::Fcg<ValueType>;
using cg = gko::solver::Cg<MixedType>;
using ir = gko::solver::Ir<ValueType>;
using ir2 = gko::solver::Ir<MixedType>;
using mg = gko::solver::Multigrid;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using bj2 = gko::preconditioner::Jacobi<MixedType, IndexType>;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
using pgm2 = gko::multigrid::Pgm<MixedType, IndexType>;
std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                  true);
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                 true);
          }},
         { "dpcpp",
              return gko::DpcppExecutor::create(
                  0, gko::ReferenceExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout « "Using mixed precision? " « use_mixed « std::endl;
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
    host_x->at(i, 0) = 0.;
host_b->at(i, 0) = 1.;
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_b);
const gko::remove_complex<ValueType> tolerance = 1e-12;
auto iter stop =
   gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                 .with_baseline(gko::stop::mode::absolute)
                                  .with_reduction_factor(tolerance)
                                 .on(exec));
auto smoother_gen = gko::share(
    ir::build()
        .with_solver(bj::build().with_max_block_size(lu).on(exec))
         .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
        .on(exec));
auto smoother_gen2 = gko::share(
    ir2::build()
        .with_solver(bj2::build().with_max_block_size(1u).on(exec))
         .with_relaxation_factor(static_cast<MixedType>(0.9))
         .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
         .on(exec));
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen2 =
    gko::share(pgm2::build().with_deterministic(true).on(exec));
auto coarsest_solver_gen = gko::share(
    ir::build()
        .with_solver(bj::build().with_max_block_size(1u).on(exec))
        .with_relaxation_factor(static_cast<ValueType>(0.9))
         .with_criteria(
            gko::stop::Iteration::build().with_max_iters(4u).on(exec))
        .on(exec));
auto coarsest solver gen2 = gko::share(
```

```
ir2::build()
        .with_solver(bj2::build().with_max_block_size(1u).on(exec))
        .with_relaxation_factor(static_cast<MixedType>(0.9))
        .with criteria(
           gko::stop::Iteration::build().with_max_iters(4u).on(exec))
        .on(exec));
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
   multigrid_gen =
       mg::build()
            .with_max_levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen, smoother_gen2)
           .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen, mg_level_gen2)
            .with_level_selector([](const gko::size_type level,
                                   const gko::LinOp*) -> gko::size_type {
                return level >= 1 ? 1 : 0;
            .with_coarsest_solver(coarsest_solver_gen2)
            .with_criteria(iter_stop, tol_stop)
            .on(exec);
} else {
   multigrid_gen = mg::build()
                        .with_max_levels(10u)
                        .with_min_coarse_rows(2u)
                        .with_pre_smoother(smoother_gen)
                        .with_post_uses_pre(true)
                        .with_mg_level(mg_level_gen)
                        .with_coarsest_solver(coarsest_solver_gen)
                        .with_criteria(iter_stop, tol_stop)
                        .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = multigrid_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
   std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
   gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
std::cout « "Multigrid iteration count:
         « logger->get_num_iterations() « std::endl;
std::cout « "Multigrid execution time [ms]: '
         « static_cast<double>(time.count()) / 1000000.0 « std::endl;
std::cout « "Multigrid execution time per iteraion[ms]:
         « static_cast<double>(time.count()) / 1000000.0 /
                logger->get_num_iterations()
          « std::endl;
```

The mixed-precision-ir program

The Mixed Precision Iterative Refinement (MPIR) solver example..

This example depends on iterative-refinement.

This example manually implements a Mixed Precision Iterative Refinement (MPIR) solver.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. An inaccurate CG solver in single precision is used as the inner solver to an iterative refinement (IR) in double precision method which solves a linear system.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using SolverType = float;
using RealSolverType = gko::remove_complex<SolverType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using solver_vec = gko::matrix::Dense<SolverType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using solver_mtx = gko::matrix::Csr<SolverType, IndexType>;
using cg = gko::solver::Cg<SolverType>;
gko::size_type max_outer_iters = 100u;
gko::size_type max_inner_iters = 100u;
RealValueType outer_reduction_factor{1e-12};
RealSolverType inner_reduction_factor{1e-2};
Print version information
std::cout « gko::version_info::get() « std::endl;
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
```

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
          {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
     exec_map{
          [] {
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                      true);
          {"hip",
           [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          { "dpcpp",
           [] {
               return gko::DpcppExecutor::create(
                    0, gko::ReferenceExecutor::create());
           }},
          {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS and initial guess as 1
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 1.;
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>((1.0), exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto initres_vec = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres_vec);
Build lower-precision system matrix and residual
auto solver_A = solver_mtx::create(exec);
auto inner_residual = solver_vec::create(exec);
auto outer_residual = vec::create(exec);
A->convert_to(solver_A);
b->convert_to(outer_residual);
restore b
b->copy_from(host_x);
Create inner solver
auto inner solver =
     cg::build()
         .with_criteria(gko::stop::ResidualNorm<SolverType>::build()
                                .with_reduction_factor(inner_reduction_factor)
                                .on(exec),
                           gko::stop::Iteration::build()
                                .with_max_iters(max_inner_iters)
                                .on(exec))
          ->generate(give(solver_A));
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto res_vec = gko::initialize<real_vec>((0.0), exec);
auto initres = exec->copy_val_to_host(initres_vec->get_const_values());
auto inner_solution = solver_vec::create(exec);
auto outer_delta = vec::create(exec);
auto tic = std::chrono::steady_clock::now();
int iter = -1;
while (true) {
     ++iter;
```

convert residual to inner precision

```
outer_residual->convert_to(inner_residual);
outer_residual->compute_norm2(res_vec);
auto res = exec->copy_val_to_host(res_vec->get_const_values());
break if we exceed the number of iterations or have converged
if (iter > max_outer_iters || res / initres < outer_reduction_factor) {</pre>
Use the inner solver to solve A * inner_solution = inner_residual with residual as initial guess.
inner_solution->copy_from(inner_residual);
inner_solver->apply(inner_residual, inner_solution);
convert inner solution to outer precision
inner_solution->convert_to(outer_delta);
x = x + inner solution
x->add_scaled(one, outer_delta);
residual = b - A * x
    outer_residual->copy_from(b);
    A->apply(neg_one, x, one, outer_residual);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
Calculate residual
A->apply(one, x, neg_one, b);
b->compute_norm2(res_vec);
std::cout « "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres_vec);
std::cout « "Final residual norm sqrt(r^T r):\n";
write(std::cout, res_vec);
Print solver statistics
    std::cout « "MPIR iteration count:
std::cout « "MPIR execution time [ms]: "
                                                     " « iter « std::endl;
               « static_cast<double>(time.count()) / 1000000.0 « std::endl;
```

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
194.679
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.22728e-10
MPIR iteration count: 25
MPIR execution time [ms]: 0.846559
```

Comments about programming and debugging

```
Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
     using ValueType = double;
     using RealValueType = gko::remove_complex<ValueType>;
using SolverType = float;
     using RealSolverType = gko::remove_complex<SolverType>;
     using IndexType = int;
     using vec = gko::matrix::Dense<ValueType>;
     using real_vec = gko::matrix::Dense<RealValueType>;
     using solver_vec = gko::matrix::Dense<SolverType>;
     using mtx = gko::matrix::Csr<ValueType, IndexType>;
     using solver_mtx = gko::matrix::Csr<SolverType, IndexType>;
     using cg = gko::solver::Cg<SolverType>;
     gko::size_type max_outer_iters = 100u;
gko::size_type max_inner_iters = 100u;
     RealValueType outer_reduction_factor{1e-12};
     RealSolverType inner_reduction_factor{1e-2};
     std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
          std::exit(-1);
     const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
          exec_map{
                        [] { return gko::OmpExecutor::create(); }},
               { "omp",
               {"cuda",
                [] {
                     return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                              true);
                } } ,
               {"hip",
                     return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                }},
               { "dpcpp",
                     return gko::DpcppExecutor::create(
                          0, gko::ReferenceExecutor::create());
               {"reference", [] { return gko::ReferenceExecutor::create(); }}};
     const auto exec = exec_map.at(executor_string)(); // throws if not valid
     auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
     gko::size_type size = A->get_size()[0];
     auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
     for (auto i = 0; i < size; i++) {
   host_x->at(i, 0) = 1.;
     auto x = gko::clone(exec, host_x);
     auto b = gko::clone(exec, host_x);
     auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
     auto initres_vec = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
     b->compute_norm2(initres_vec);
     auto solver_A = solver_mtx::create(exec);
     auto inner_residual = solver_vec::create(exec);
     auto outer_residual = vec::create(exec);
     A->convert to(solver A);
     b->convert to(outer residual);
```

```
b->copy_from(host_x);
auto inner_solver =
    cg::build()
        .with_criteria(gko::stop::ResidualNorm<SolverType>::build()
                             .with_reduction_factor(inner_reduction_factor)
                             .on(exec).
                        gko::stop::Iteration::build()
                             .with_max_iters(max_inner_iters)
                             .on(exec))
        .on(exec)
        ->generate(give(solver_A));
exec->synchronize();
std::chrono::nanoseconds time(0);
auto res_vec = gko::initialize<real_vec>({0.0}, exec);
auto initres = exec->copy_val_to_host(initres_vec->get_const_values());
auto inner_solution = solver_vec::create(exec);
auto outer_delta = vec::create(exec);
auto tic = std::chrono::steady_clock::now();
int iter = -1;
while (true) {
    ++iter;
    outer_residual->convert_to(inner_residual);
    outer_residual->compute_norm2(res_vec);
    auto res = exec->copy_val_to_host(res_vec->get_const_values());
    if (iter > max_outer_iters || res / initres < outer_reduction_factor) {
        break;
    inner_solution->copy_from(inner_residual);
    inner_solver->apply(inner_residual, inner_solution);
    inner_solution->convert_to(outer_delta);
    x->add scaled(one, outer delta);
    outer_residual->copy_from(b);
A->apply(neg_one, x, one, outer_residual);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
A->apply(one, x, neg_one, b);
b->compute_norm2(res_vec);
std::cout « "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres_vec);
std::cout « "Final residual norm sqrt(r^T r):\n";
write(std::cout, res_vec);
std::cout « "MPIR iteration count:
std::cout « "MPIR execution time [ms]: " « iter « std::endl;
          « static_cast<double>(time.count()) / 1000000.0 « std::endl;
```

The mixed-spmv program

The mixed spmv example..

Introduction

This mixed spmv example should give the usage of Ginkgo mixed precision. This example is meant for you to understand how Ginkgo works with different precision of data. We encourage you to play with the code, change the parameters and see what is best suited for your purposes.

About the example

Each example has the following sections:

- 1. **Introduction:**This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
- 2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
- 3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
- 4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

Include files

This is the main ginkgo header file. #include <ginkgo/ginkgo.hpp>

Add the fstream header to read from data from files.

#include <fstream>

Add the C++ iostream header to output information to the console.

#include <iostream

Add the STL map header for the executor selection.

#include <map>

Add the string manipulation header to handle strings.

#include <string>

Add the timing header for timing.

#include <chrono>

Add the random header to generate random vectors.

```
#include <random
namespace {
 * Generate a random value.
 * @tparam ValueType valuetype of the value
 * Otparam Valuelyser varietyse of the value distribution 
* Otparam Engine type of random engine
 * @param value_dist distribution of array values
 * @param engine a random engine
 * @return ValueType
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<!gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
    return value dist (gen);
 \star Specialization for complex types.
 * @copydoc get_rand_value
template <typename ValueType, typename ValueDistribution, typename Engine>typename std::enable_if<gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
    return ValueType(value_dist(gen), value_dist(gen));
}
 \star timing the apply operation A->apply(b, x). It will runs 2 warmup and get
 * average time among 10 times.
 * @return seconds
double timing(std::shared_ptr<const gko::Executor> exec,
                std::shared_ptr<const gko::LinOp> A,
                std::shared_ptr<const gko::LinOp> b,
                std::shared_ptr<gko::LinOp> x)
    int warmup = 2;
    int rep = 10;
for (int i = 0; i < warmup; i++) {</pre>
        A->apply(b, x);
    double total_sec = 0;
for (int i = 0; i < rep; i++) {</pre>
```

always clone the x in each apply

auto xx = x -> clone();

synchronize to make sure data is already on device

```
exec->synchronize();
auto start = std::chrono::steady_clock::now();
A->apply(b, xx);
```

synchronize to make sure the operation is done

```
exec->synchronize();
auto stop = std::chrono::steady_clock::now();
```

get the duration in seconds

```
std::chrono::duration<double> duration_time = stop - start;
total_sec += duration_time.count();
if (i + 1 == rep) {
```

copy the result back to x

```
x->copy_from(xx);
}

return total_sec / rep;
}
} // namespace
int main(int argc, char* argv[])
{
```

Use some shortcuts. In Ginkgo, vectors are seen as a gko::matrix::Dense with one column/one row. The advantage of this concept is that using multiple vectors is a now a natural extension of adding columns/rows are necessary.

```
using HighPrecision = double;
using RealValueType = gko::remove_complex<HighPrecision>;
using LowPrecision = float;
using IndexType = int;
using hp_vec = gko::matrix::Dense<HighPrecision>;
using lp_vec = gko::matrix::Dense<LowPrecision>;
using real_vec = gko::matrix::Dense<RealValueType>;
```

The gko::matrix::Ell class is used here, but any other matrix class such as gko::matrix::Coo, gko::matrix::Hybrid, gko::matrix::Csr or gko::matrix::Sellp could also be used. Note. the behavior will depends GINKGO_MIXED_PR← ECISION flags and the actual implementation from different matrices.

```
using hp_mtx = gko::matrix::Ell<HighPrecision, IndexType>;
using lp_mtx = gko::matrix::Ell<LowPrecision, IndexType>;
```

Print the ginkgo version information.

```
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
    std::exit(-1);
```

Where do you want to run your operation?

The gko::Executor class is one of the cornerstones of Ginkgo. Currently, we have support for an gko::OmpExecutor, which uses OpenMP multi-threading in most of its kernels, a gko::ReferenceExecutor, a single threaded specialization of the OpenMP executor and a gko::CudaExecutor which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/routines within Ginkgo should automatically work and run on the executor with any other changes.

```
const auto executor string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared ptr<qko::Executor>()>
    exec_map{
        {"omp",
               [] { return gko::OmpExecutor::create(); }},
        {"cuda",
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                              true):
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                             true);
         11.
        { "dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                               gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Preparing your data and transfer to the proper device.

Read the matrix using the read function and set the right hand side randomly.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. gko::share and gko::give are the functions that you would need to use.

```
read the matrix into HighPrecision and LowPrecision.
```

```
auto hp_A = share(gko::read<hp_mtx>(std::ifstream("data/A.mtx"), exec));
auto lp_A = share(gko::read<lp_mtx>(std::ifstream("data/A.mtx"), exec));
Set the shortcut for each dimension
auto A dim = hp A->get size();
auto b_dim = gko::dim<2>{A_dim[1], 1};
auto x_dim = gko::dim<2>{A_dim[0], b_dim[1]};
auto host_b = hp_vec::create(exec->get_master(), b_dim);
fill the b vector with some random data
std::default_random_engine rand_engine(32);
auto dist = std::uniform_real_distribution<RealValueType>(0.0, 1.0);
for (int i = 0; i < host_b->get_size()[0]; i++)
    host_b->at(i, 0) = get_rand_value<HighPrecision>(dist, rand_engine);
copy the data from host to device
auto hp_b = share(gko::clone(exec, host_b));
auto lp_b = share(lp_vec::create(exec));
lp_b->copy_from(hp_b);
create several result x vector in different precision
auto hp_x = share(hp_vec::create(exec, x_dim));
auto lp_x = share(lp_vec::create(exec, x_dim));
auto hplp_x = share(hp_x->clone());
auto lplp_x = share(hp_x->clone());
```

Measure the time of apply

auto lphp_x = share(hp_x->clone());

We measure the time among different combination of apply operation.

```
Hp * Hp -> Hp
auto hp_sec = timing(exec, hp_A, hp_b, hp_x);

Lp * Lp -> Lp
auto lp_sec = timing(exec, lp_A, lp_b, lp_x);

Hp * Lp -> Hp
auto hplp_sec = timing(exec, hp_A, lp_b, hplp_x);

Lp * Lp -> Hp
auto lplp_sec = timing(exec, lp_A, lp_b, lplp_x);
Lp * Hp -> Hp
auto lphp_sec = timing(exec, lp_A, hp_b, lphp_x);
```

To measure error of result. neg_one is an object that represent the number -1.0 which allows for a uniform interface when computing on any device. To compute the residual, all you need to do is call the add_scaled method, which in this case is an axpy and equivalent to the LAPACK axpy routine. Finally, you compute the euclidean 2-norm with the compute_norm2 function.

```
auto neg_one = gko::initialize<hp_vec>({-1.0}, exec);
auto hp_x_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
```

```
auto hplp_diff_norm = gko::initialize<real_vec>((0.0), exec->get_master());
auto lplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lphp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp\_diff = hp\_x -> clone();
auto hplp\_diff = hp\_x -> clone();
auto lplp_diff = hp_x->clone();
auto lphp_diff = hp_x->clone();
auto lphp_diff = hp_x->clone();
hp_x->compute_norm2(hp_x_norm);
lp_diff->add_scaled(neg_one, lp_x);
lp_diff->compute_norm2(lp_diff_norm);
hplp_diff->add_scaled(neg_one, hplp_x);
hplp_diff->compute_norm2(hplp_diff_norm);
lplp_diff->add_scaled(neg_one, lplp_x);
lplp_diff->compute_norm2(lplp_diff_norm);
lphp_diff->add_scaled(neg_one, lphp_x);
lphp_diff->compute_norm2(lphp_diff_norm);
exec->synchronize();
std::cout.precision(10);
std::cout « std::scientific;
std::cout « "High Precision time(s): " « hp_sec « std::endl;
std::cout « "High Precision result norm: " « hp_x_norm->at(0)
          « std::endl;
std::cout « "Low Precision time(s): " « lp_sec « std::endl;
std::cout « "Hp * Lp -> Hp relative error:
" lplp_diff_norm->at(0) / hp_x_norm->at(0) « "\n";
" Lp * Hp -> Hp time(s): " « lplp_sec « std::endl;
std::cout « "Lp * Hp -> Hp time(s): " « lp std::cout « "Lp * Hp -> Hp relative error:
          « lphp_diff_norm->at(0) / hp_x_norm->at(0) « "\n";
```

The following is the expected result (omp):

```
High Precision time(s): 2.0568800000e-05
High Precision result norm: 1.7725534898e+05
Low Precision time(s): 2.0955600000e-05
Low Precision relative error: 9.1052887738e-08
Hp * Lp -> Hp time(s): 2.1186100000e-05
Hp * Lp -> Hp relative error: 3.7799774251e-08
Lp * Lp -> Hp time(s): 2.0312300000e-05
Lp * Lp -> Hp relative error: 5.7910008031e-08
Lp * Hp -> Hp time(s): 2.0312300000e-05
Lp * Hp -> Hp relative error: 3.7173133506e-08
```

Comments about programming and debugging

```
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <chrono>
#include <random>
namespace {
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<!gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
    return value_dist(gen);
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
    return ValueType(value_dist(gen), value_dist(gen));
double timing(std::shared_ptr<const gko::Executor> exec,
               std::shared_ptr<const gko::LinOp> A,
               std::shared_ptr<const gko::LinOp> b,
               std::shared_ptr<gko::LinOp> x)
    int warmup = 2;
    int rep = 10;

for (int i = 0; i < warmup; i++) {
        A->apply(b, x);
    double total_sec = 0;
    for (int i = 0; i < rep; i++) {</pre>
        auto xx = x - clone();
        exec->synchronize();
        auto start = std::chrono::steady_clock::now();
        A->apply(b, xx);
        exec->synchronize();
        auto stop = std::chrono::steady_clock::now();
        std::chrono::duration<double> duration_time = stop - start;
        total sec += duration time.count();
        if (i + 1 == rep) {
             x->copy_from(xx);
    return total_sec / rep;
   // namespace
int main(int argc, char* argv[])
    using HighPrecision = double;
    using RealValueType = gko::remove_complex<HighPrecision>;
    using LowPrecision = float;
    using IndexType = int;
    using hp_vec = gko::matrix::Dense<HighPrecision>;
using lp_vec = gko::matrix::Dense<LowPrecision>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using hp_mtx = gko::matrix::Ell<HighPrecision, IndexType>;
using lp_mtx = gko::matrix::Ell<LowPrecision, IndexType>;
    std::cout « gko::version_info::get() « std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
        std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
        exec_map{
                     [] { return gko::OmpExecutor::create(); }},
             { "omp",
             {"cuda",
              [] {
                  return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                       true);
              }},
                  return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                      true):
              }},
```

```
{"dpcpp",
             return gko::DpcppExecutor::create(0,
                                                 gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto hp_A = share(gko::read<hp_mtx>(std::ifstream("data/A.mtx"), exec));
auto lp_A = share(gko::read<lp_mtx>(std::ifstream("data/A.mtx"), exec));
auto A_dim = hp_A->get_size();
auto b_dim = gko::dim<2>{A_dim[1], 1};
auto x_dim = gko::dim<2>{A_dim[0], b_dim[1]};
auto host_b = hp_vec::create(exec->get_master(), b_dim);
std::default_random_engine rand_engine(32);
auto dist = std::uniform_real_distribution<RealValueType>(0.0, 1.0);
for (int i = 0; i < host_b->get_size()[0]; i++) {
   host_b->at(i, 0) = get_rand_value<HighPrecision>(dist, rand_engine);
auto hp_b = share(gko::clone(exec, host_b));
auto lp_b = share(lp_vec::create(exec));
lp_b->copy_from(hp_b);
auto hp_x = share(hp_vec::create(exec, x_dim));
auto lp_x = share(lp_vec::create(exec, x_dim));
auto hplp_x = share(hp_x->clone());
auto lplp_x = share(hp_x->clone());
auto lphp_x = share(hp_x->clone());
auto hp_sec = timing(exec, hp_A, hp_b, hp_x);
auto lp_sec = timing(exec, lp_A, lp_b, lp_x);
auto hplp_sec = timing(exec, hp_A, lp_b, hplp_x);
auto lphp_sec = timing(exec, lp_A, lp_b, lphp_x);
auto lphp_sec = timing(exec, lp_A, hp_b, lphp_x);
auto neg_one = gko::initialize<hp_vec>({-1.0}, exec);
auto hp_x_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto hplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lphp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff = hp_x->clone();
auto hplp_diff = hp_x->clone();
auto lplp_diff = hp_x->clone();
auto lphp_diff = hp_x->clone();
hp_x->compute_norm2(hp_x_norm);
lp diff->add scaled(neg one, lp x);
lp_diff->compute_norm2(lp_diff_norm);
hplp_diff->add_scaled(neg_one, hplp_x);
hplp_diff->compute_norm2(hplp_diff_norm);
lplp_diff->add_scaled(neg_one, lplp_x);
lplp_diff->compute_norm2(lplp_diff_norm);
lphp_diff->add_scaled(neg_one, lphp_x);
lphp_diff->compute_norm2(lphp_diff_norm);
exec->synchronize();
std::cout.precision(10);
std::cout « std::scientific;
std::cout « "High Precision time(s): " « hp_sec « std::endl; std::cout « "High Precision result norm: " « hp_x_norm->at(0)
          « std::endl;
std::cout « "Low Precision time(s): " « lp_sec « std::endl;
std::cout « "Low Precision relative error:
< lplp_diff_norm->at(0) / hp_x_norm->at(0) < "\n";
```

The multigrid-preconditioned-solver program

The preconditioned solver example..

This example depends on preconditioned-solver.

This example shows how to use the multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a multigrid preconditioner. The example features the generating time and runtime of the CG solver.

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using mg = gko::solver::Multigrid;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
         {"omp",
{"cuda",
                 [] { return gko::OmpExecutor::create(); }},
              return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          }},
         {"dpcpp",
```

```
return gko::DpcppExecutor::create(
                  0, gko::ReferenceExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read data
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS as 1 and initial guess as 0
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_b);
Create multigrid factory
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid gen =
    mg::build()
        .with_mg_level(pgm::build().with_deterministic(true).on(exec))
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
        .on(exec);
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto solver_gen =
    cg::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(100u).on(exec),
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_baseline(gko::stop::mode::absolute)
                 .with_reduction_factor(tolerance)
                 .on(exec))
        .with_preconditioner(multigrid_gen)
        .on(exec);
Create solver
std::chrono::nanoseconds gen_time(0);
auto gen tic = std::chrono::steady clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
Add logger
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->svnchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
```

Calculate residual

```
auto res = gko::as<vec>(logger->get_residual_norm());
std::cout « "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
Print solver statistics
    std::cout « "CG iteration count:
                                              " « logger->get num iterations()
              « std::endl;
    std::cout « "CG generation time [ms]: "
             « static_cast<double>(gen_time.count()) / 1000000.0 « std::endl;
    std::cout « "CG execution time [ms]: "
             « static_cast<double>(time.count()) / 1000000.0 « std::endl;
    std::cout « "CG execution time per iteraion[ms]: "
             « static_cast<double>(time.count()) / 1000000.0 /
                    logger->get_num_iterations()
              « std::endl;
```

Results

This is the expected output:

Comments about programming and debugging

```
************GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
```

```
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using mg = gko::solver::Multigrid;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
        exec_map{
            {"omp", [] { return gko::OmpExecutor::create(); }},
            {"cuda",
             [] {
                 return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
             }},
            { "hip",
             [] {
                 return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
            {"dpcpp",
             [] {
                 return gko::DpcppExecutor::create(
                     0, gko::ReferenceExecutor::create());
            {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    auto x = vec::create(exec):
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_b);
    std::shared_ptr<gko::LinOpFactory> multigrid_gen;
    multigrid_gen =
        mg::build()
            .with_mg_level(pgm::build().with_deterministic(true).on(exec))
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(1u).on(exec))
            .on(exec);
    const gko::remove_complex<ValueType> tolerance = 1e-8;
    auto solver_gen =
        cq::build()
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(100u).on(exec),
                gko::stop::ResidualNorm<ValueType>::build()
                     .with_baseline(gko::stop::mode::absolute)
                     .with_reduction_factor(tolerance)
                     .on(exec))
            .with_preconditioner(multigrid_gen)
            .on(exec);
    std::chrono::nanoseconds gen_time(0);
    auto gen_tic = std::chrono::steady_clock::now();
    auto solver = solver_gen->generate(A);
    exec->svnchronize();
    auto gen_toc = std::chrono::steady_clock::now();
    gen_time +=
        std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
    std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
        gko::log::Convergence<ValueType>::create();
    solver->add_logger(logger);
    exec->synchronize();
    std::chrono::nanoseconds time(0);
    auto tic = std::chrono::steady_clock::now();
    solver->apply(b, x);
    exec->synchronize();
    auto toc = std::chrono::steady_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
```

The multigrid-preconditioned-solver-customized program

The customized multigrid preconditioned solver example..

This example depends on multigrid-preconditioned-solver.

This example shows how to customize the multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a multigrid preconditioner. Several non-default options are used to create this preconditioner. The example features the generating time and runtime of the CG solver.

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using ir = gko::solver::Ir<ValueType>;
using mg = gko::solver::Multigrid;
using ic = gko::preconditioner::Ic<gko::solver::LowerTrs<ValueType>>;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
```

```
true);
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
         {"dpcpp",
          [] {
              return gko::DpcppExecutor::create(
                  0, gko::ReferenceExecutor::create());
          }},
         {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
Create RHS as 1 and initial guess as 0
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {</pre>
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
Calculate initial residual by overwriting b
auto one = gko::initialize<vec>((1.0), exec);
auto neg_one = gko::initialize<vec>((-1.0), exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
copy b again
b->copy_from(host_b);
Prepare the stopping criteria
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                 .with_baseline(gko::stop::mode::absolute)
                                  .with_reduction_factor(tolerance)
                                  .on(exec));
auto exact tol stop =
    gko::share(gko::stop::ResidualNorm<ValueType>::build()
                    .with_baseline(gko::stop::mode::rhs_norm)
                    .with_reduction_factor(1e-14)
                     .on(exec));
std::shared_ptr<const gko::log::Convergence<ValueType» logger =
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
```

Now we customize some settings of the multigrid preconditioner. First we choose a smoother. Since the input matrix is spd, we use iterative refinement with two iterations and an Ic solver.

```
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
```

Next we select a CG solver for the coarsest level. Again, since the input matrix is known to be spd, and the Pgm restriction preserves this characteristic, we can safely choose the CG. We reuse the lc factory here to generate an lc preconditioner. It is important to solve until machine precision here to get a good convergence rate.

```
auto coarsest_gen = gko::share(cg::build()
                                   .with_preconditioner(ic_gen)
                                   .with_criteria(iter_stop, exact_tol_stop)
                                   .on(exec));
Here we put the customized options together and create the multigrid factory.
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid_gen =
   mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(32u)
        .with_pre_smoother(smoother_gen)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen)
        .with_coarsest_solver(coarsest_gen)
        .with\_default\_initial\_guess (gko::solver::initial\_guess\_mode::zero) \\
        .with criteria(
           gko::stop::Iteration::build().with_max_iters(1u).on(exec))
        .on(exec);
Create solver factory
auto solver_gen = cg::build()
                      .with_criteria(iter_stop, tol_stop)
                      .with_preconditioner(multigrid_gen)
                      .on(exec);
Create solver
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
Solve system
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
Calculate residual
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Initial residual norm sqrt(r^T r): n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
Print solver statistics
    std::cout « "CG iteration count:
                                              " « logger->get_num_iterations()
              « std::endl;
    std::cout « "CG generation time [ms]: "
    « static_cast<double>(time.count()) / 1000000.0 « std::endl;
    std::cout « "CG execution time per iteraion[ms]: "
              « static_cast<double>(time.count()) / 1000000.0 /
                    logger->get_num_iterations()
              « std::endl;
}
Results
```

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%*MatrixMarket matrix array real general
1 1
25.9808
Final residual norm sqrt(r^T r):
%*MatrixMarket matrix array real general
1 1
5.81328e-09
GG iteration count: 12
CG generation time [ms]: 1.41642
CG execution time [ms]: 6.59244
CG execution time per iteraion[ms]: 0.54937
```

Comments about programming and debugging

```
*********GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
   Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using ir = gko::solver::Ir<ValueType>;
    using mg = gko::solver::Multigrid;
    using ic = gko::preconditioner::Ic<gko::solver::LowerTrs<ValueType>>;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    std::cout « gko::version_info::get() « std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
        exec_map{
             {"omp", [] { return gko::OmpExecutor::create(); }},
             {"cuda",
             [] {
                  return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
             {"hip",
             [] {
                  return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
             } } ,
             {"dpcpp",
                  return gko::DpcppExecutor::create(
                      0, gko::ReferenceExecutor::create());
             {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size type size = A->get size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {</pre>
        host_x->at(i, 0) = 0.;
host_b->at(i, 0) = 1.;
    auto x = vec::create(exec);
    auto b = vec::create(exec);
```

```
x->copy_from(host_x);
b->copy_from(host_b);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_b);
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
   gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
                                .with_baseline(gko::stop::mode::absolute)
                                .with_reduction_factor(tolerance)
                                 .on(exec));
auto exact_tol_stop =
   gko::share(gko::stop::ResidualNorm<ValueType>::build()
                   .with_baseline(gko::stop::mode::rhs_norm)
                   .with_reduction_factor(1e-14)
                    .on(exec));
std::shared_ptr<const gko::log::Convergence<ValueType» logger =</pre>
   gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
auto ic_gen = gko::share(
   ic::build()
        .with_factorization_factory(
            gko::factorization::Ic<ValueType, int>::build().on(exec))
        .on(exec));
auto smoother_gen = gko::share(
   gko::solver::build_smoother(ic_gen, 2u, static_cast<ValueType>(0.9)));
auto mg level gen =
   gko::share(pgm::build().with_deterministic(true).on(exec));
auto coarsest_gen = gko::share(cg::build()
                                     . \verb|with_preconditioner(ic_gen)|\\
                                     .with_criteria(iter_stop, exact_tol_stop)
                                     .on(exec));
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid_gen
   mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(32u)
        .with pre smoother(smoother gen)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen)
        .with_coarsest_solver(coarsest_gen)
        .with_default_initial_guess(gko::solver::initial_guess_mode::zero)
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1u).on(exec))
        .on(exec);
auto solver_gen = cg::build()
                       .with_criteria(iter_stop, tol_stop)
                       .with_preconditioner(multigrid_gen)
                       .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->svnchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
   std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Initial residual norm sqrt(r^T r): n";
write(std::cout, initres);
std::cout « "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);
std::cout « "CG iteration count:
                                            " « logger->get_num_iterations()
          « std::endl;
std::cout « "CG generation time [ms]: "
          « static_cast<double>(gen_time.count()) / 1000000.0 « std::endl;
std::cout « "CG execution time [ms]:
          « static_cast<double>(time.count()) / 1000000.0 « std::endl;
std::cout « "CG execution time per iteraion[ms]:
          « static_cast<double>(time.count()) / 1000000.0 /
                 logger->get_num_iterations()
          « std::endl;
```

The multigrid-preconditioned-solver-customized program	

180

The nine-pt-stencil-solver program

The 9-point stencil example..

This example depends on simple-solver, three-pt-stencil-solver, poisson-solver.

Introduction

This example solves a 2D Poisson equation:

[$\Omega = (0,1)^2 \ D = [0,1]^2 \ U = (0,1)^2 \ U = (0,1)^2$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization may be done by any order Taylor polynomial. For an equidistant grid with K "inner" discretization points ((x1,y1), \ldots, (xk,y1),(x1,y2), \ldots, (xk,yk,z1)) step size (h = 1 / (K + 1)) and a stencil (\in \mathb{R}^{3} \times 3), the formula produces a system of linear equations

 $(\sum_{a,b=-1}^{1} stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2)$, on any inner node with a neighborhood of inner nodes

On any node, where neighbor is on the border, the neighbor is replaced with a (-stencil(a,b) * u_{i+a,j+b}) and added to the right hand side vector. For example a node with a neighborhood of only edge nodes may look like this

```
[\sum_{a,b=-1}^{(1,0)} stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2 - \sum_{a=-1}^{1} stencil(a,1) * u_{(i+a,j+1)}]
```

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function f is set to (f(x,y) = 6x + 6y) (making the solution $(u(x,y) = x^3)$

• y³)), but that can be changed in the main function. Also the stencil values for the core, the faces, the edge and the corners can be changed when passing additional parameters.

The intention of this is to show how generation of stencil values and the right hand side vector changes when increasing the dimension.

About the example

The commented program

```
This example solves a 2D Poisson equation:
    \comega_b = [0,1]^2 (with boundary) \partial\Omega = \Omega_b \backslash \Omega u : \Omega_b -> R u" = f in \Omega u = u_D on \no."
     u = u_D on \gamma (u)
using a finite difference method on an equidistant grid with \ 'K' discretization
points ('K' can be controlled with a command line parameter). The discretization
may be done by any order Taylor polynomial.
For an equidistant grid with K "inner" discretization points (x1,y1),
(xk,y1), (x1,y2), ..., (xk,yk) step size h=1 / (K+1) and a stencil \in
\R^{3} x 3}, the formula produces a system of linear equations
\sum_{a,b=-1}^1 \text{stencil}(a,b) * u_{(i+a,j+b)} = -f_k h^2,
a neighborhood of inner nodes
On any node, where neighbor is on the border, the neighbor is replaced with a '-stencil(a,b) \star u_{i+a,j+b}' and added to the right hand side vector. For
example a node with a neighborhood of only edge nodes may look like this
\sum_{a,b=-1}^{(1,0)} stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2 - \sum_{a=-1}^1 stance{-1}
stencil(a,1) * u_{(i+a,j+1)
which is then solved using Ginkgo's implementation of the CG method
preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x,y) = 6x + 6y' (making the solution 'u(x,y) = x^3 + y^3'), but that can be changed in the 'main' function. Also the stencil values
for the core, the faces, the edge and the corners can be changed when passing
additional parameters.
The intention of this is to show how generation of stencil values and the right
hand side vector changes when increasing the dimension.
        #include <array>
#include <chrono>
#include <ainkao/ainkao.hpp>
#include <iostream>
#include <map>
#include <string>
#include <vector>
```

Stencil values. Ordering can be seen in the main function Can also be changed by passing additional parameter when executing

```
constexpr double default_alpha = 10.0 / 3.0;
constexpr double default_beta = -2.0 / 3.0;
constexpr double default_gamma = -1.0 / 6.0;
/ * Possible alternative default values are
* default_alpha = 8.0;
* default_beta = -1.0;
* default_gamma = -1.0;
* /
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
Generates the RHS vector given f and the boundary conditions.
```

Iterating over the edges to add boundary values and adding the overlapping 3x1 to the rhs

```
for (size_t i = 0; i < dp; ++i) {
    const auto xi = ValueType(i + 1) * h;
    const auto index_top = i;
    const auto index_bot = i + dp * (dp - 1);
    rhs[index_top] -= u(xi - h, 0.0) * coefs[0];
    rhs[index_top] -= u(xi, 0.0) * coefs[1];
    rhs[index_top] -= u(xi + h, 0.0) * coefs[2];
    rhs[index_top] -= u(xi - h, 1.0) * coefs[6];
    rhs[index_bot] -= u(xi - h, 1.0) * coefs[7];
    rhs[index_bot] -= u(xi + h, 1.0) * coefs[7];
    rhs[index_bot] -= u(xi + h, 1.0) * coefs[8];
}
for (size_t i = 0; i < dp; ++i) {
    const auto yi = ValueType(i + 1) * h;
    const auto index_left = i * dp;
    const auto index_right = i * dp + (dp - 1);
    rhs[index_left] -= u(0.0, yi - h) * coefs[0];
    rhs[index_left] -= u(0.0, yi + h) * coefs[6];
    rhs[index_right] -= u(1.0, yi - h) * coefs[2];
    rhs[index_right] -= u(1.0, yi + h) * coefs[5];
    rhs[index_right] -= u(1.0, yi + h) * coefs[8];
}</pre>
```

remove the double corner values

```
rhs[0] += u(0.0, 0.0) * coefs[0];
rhs[(dp - 1)] += u(1.0, 0.0) * coefs[2];
rhs[(dp - 1) * dp] += u(0.0, 1.0) * coefs[6];
rhs[dp * dp - 1] += u(1.0, 1.0) * coefs[8];
```

Prints the solution u.

```
template <typename ValueType, typename IndexType>
void print_solution(IndexType dp, const ValueType* u)
{
    for (IndexType i = 0; i < dp; ++i) {
        for (IndexType j = 0; j < dp; ++j) {
            std::cout « u[i * dp + j] « ' ';
        }
        std::cout « '\n';
    }
    std::cout « std::endl;
}</pre>
```

Computes the 1-norm of the error given the computed u and the correct solution function correct_u.

```
void solve_system(const std::string& executor_string,
                  unsigned int discretization_points, IndexType* row_ptrs,
                  IndexType* col_idxs, ValueType* values, ValueType* rhs,
                  ValueType* u, gko::remove_complex<ValueType> reduction_factor)
Some shortcuts
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto& dp = discretization_points;
const gko::size_type dp_2 = dp * dp;
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        { "omp",
                [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                 gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
executor where the application initialized the data
const auto app_exec = exec->get_master();
```

Tell Ginkgo to use the data in our application

Matrix: we have to set the executor of the matrix to the one where we want SpMVs to run (in this case exec). When creating array views, we have to specify the executor where the data is (in this case app_exec).

If the two do not match, Ginkgo will automatically create a copy of the data on exec (however, it will not copy the data back once it is done

• here this is not important since we are not modifying the matrix).

Solution: we have to be careful here - if the executors are different, once we compute the solution the array will not be automatically copied back to the original memory locations. Fortunately, whenever \mathtt{apply} is called on a linear operator (e.g. matrix, solver) the arguments automatically get copied to the executor where the operator is, and copied back once the operation is completed. Thus, in this case, we can just define the solution on $\mathtt{app_exec}$, and it will be automatically transferred to/from \mathtt{exec} if needed.

Generate solver

auto solver_gen =

```
cg::build()
        .with criteria(
            gko::stop::Iteration::build().with_max_iters(dp_2).on(exec),
             gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor)
                 .on(exec))
        .with_preconditioner(bj::build().on(exec))
         .on(exec);
auto solver = solver_gen->generate(gko::give(matrix));
Solve system
    solver->apply(b, x);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && std::string(argv[1]) == "--help") {
    std::cerr
       « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const IndexType discretization_points =
argc >= 3 ? std::atoi(argv[2]) : 100;
const ValueType alpha_c = argc >= 4 ? std::atof(argv[3]) : default_alpha;
const ValueType beta_c = argc >= 5 ? std::atof(argv[4]) : default_beta;
const ValueType gamma_c = argc >= 6 ? std::atof(argv[5]) : default_gamma;
clang-format off
std::array<ValueType, 9> coefs{
   gamma_c, beta_c, gamma_c,
       beta_c, alpha_c, beta_c,
    gamma_c, beta_c, gamma_c);
clang-format on
const auto dp = discretization_points;
const size_t dp_2 = dp * dp;
problem:
auto correct_u = [](ValueType x, ValueType y) {
    return x * x * x + y * y * y;
auto f = [](ValueType x, ValueType y) {
    return ValueType(6) * x + ValueType(6) * y;
matrix
right hand side
std::vector<ValueType> rhs(dp_2);
std::vector<ValueType> u(dp_2, 0.0);
generate_stencil_matrix(dp, row_ptrs.data(), col_idxs.data(), values.data(),
                         coefs.data());
looking for solution u = x^3: f = 6x, u(0) = 0, u(1) = 1
generate_rhs(dp, f, correct_u, rhs.data(), coefs.data());
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
auto start_time = std::chrono::steady_clock::now();
auto stop_time = std::chrono::steady_clock::now();
auto runtime_duration =
    static_cast<double>(
        std::chrono::duration_cast<std::chrono::nanoseconds>(stop_time -
                                                                start_time)
             .count()) *
    1e-6;
Uncomment to print the solution print_solution(dp, u.data());
    std::cout « "The average relative error is
              « calculate_error(dp, u.data(), correct_u) /
                      static_cast<gko::remove_complex<ValueType>> (dp_2)
              « std::endl;
    std::cout « "The runtime is " « std::to_string(runtime_duration) « " ms"
              « std::endl;
}
```

Results

The expected output should be

The average relative error is 6.35715e-06 The runtime is 167.320520 ms

Comments about programming and debugging

The plain program

3. Neither the name of the copyright holder nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

using a finite difference method on an equidistant grid with 'K' discretization points ('K' can be controlled with a command line parameter). The discretization may be done by any order Taylor polynomial. For an equidistant grid with K "inner" discretization points (x1,y1), ..., (xk,y1), (x1,y2), ..., (xk,yk) step size h=1 / (K + 1) and a stencil \in \R^{3} x 3}, the formula produces a system of linear equations

 $\label{eq:local_sum_{a,b=-1}^1 stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2, \ \ \mbox{on any inner node with a neighborhood of inner nodes}$

On any node, where neighbor is on the border, the neighbor is replaced with a '-stencil(a,b) * $u_{i+a,j+b}'$ and added to the right hand side vector. For example a node with a neighborhood of only edge nodes may look like this

 $\sum_{a,b=-1}^{1,0} \sin_{a,b} \cdot u_{(i+a,j+b)} = -f_k h^2 - \sum_{a=-1}^1 \operatorname{stencil}(a,1) \cdot u_{(i+a,j+1)}$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x,y) = 6x + 6y' (making the solution 'u(x,y) = $x^3 + y^3$ '), but that can be changed in the 'main' function. Also the stencil values for the core, the faces, the edge and the corners can be changed when passing additional parameters.

```
#include <array>
#include <chrono>
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
#include <vector>
constexpr double default_alpha = 10.0 / 3.0;
constexpr double default_beta = -2.0 / 3.0;
constexpr double default_gamma = -1.0 / 6.0;
/* Possible alternative default values are
* default_alpha = 8.0;
* default_beta = -1.0;
* default_gamma = -1.0;
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType dp, IndexType* row_ptrs,
                                            IndexType* col_idxs, ValueType* values,
                                             ValueType* coefs)
{
     IndexType pos = 0;
      const size_t dp_2 = dp * dp;
row_ptrs[0] = pos;
      for (IndexType k = 0; k < dp; ++k) {
    for (IndexType i = 0; i < dp; ++i) {
                  const size_t index = i + k * dp;
                  const size_t index = i + k * dp;
for (IndexType j = -1; j <= 1; ++j) {
  for (IndexType l = -1; l <= 1; ++l) {
    const IndexType offset = l + l + 3 * (j + l);
    if ((k + j) >= 0 && (k + j) < dp && (i + l) >= 0 &&
        (i + l) < dp) {
        values[pos] = coefs[offset];
    }
}</pre>
                                     col_idxs[pos] = index + 1 + dp * j;
                                     ++pos;
                        }
                  row_ptrs[index + 1] = pos;
      }
template <typename Closure, typename ClosureT, typename ValueType,
              typename IndexType>
void generate_rhs(IndexType dp, Closure f, ClosureT u, ValueType* rhs,
                           ValueType* coefs)
      const size_t dp_2 = dp * dp;
const ValueType h = 1.0 / (dp + 1.0);
for (IndexType i = 0; i < dp; ++i) {</pre>
            const auto yi = ValueType(i + 1) * h;
            for (IndexType j = 0; j < dp; ++j) {
                  const auto xi = ValueType(j + 1) * h;
                  const auto index = i * dp + j;
                  rhs[index] = -f(xi, yi) * h * h;
      for (size_t i = 0; i < dp; ++i) {</pre>
            const auto xi = ValueType(i + 1) * h;
            const auto index_top = i;
const auto index_bot = i + dp * (dp - 1);
            rhs[index_top] -= u(xi - h, 0.0) * coefs[0];
rhs[index_top] -= u(xi, 0.0) * coefs[1];
            rhs[index_top] -= u(xi + h, 0.0) * coefs[2];
rhs[index_bot] -= u(xi - h, 1.0) * coefs[6];
            rhs[index_bot] -= u(xi, 1.0) * coefs[7];
            rhs[index_bot] -= u(xi + h, 1.0) * coefs[8];
      for (size_t i = 0; i < dp; ++i) {</pre>
            const auto yi = ValueType(i + 1) * h;
const auto index_left = i * dp;
const auto index_right = i * dp + (dp - 1);
           const auto index_right = 1 * dp + (dp - 1);
rhs[index_left] -= u(0.0, yi - h) * coefs[0];
rhs[index_left] -= u(0.0, yi) * coefs[3];
rhs[index_left] -= u(0.0, yi + h) * coefs[6];
rhs[index_right] -= u(1.0, yi - h) * coefs[2];
rhs[index_right] -= u(1.0, yi) * coefs[5];
rhs[index_right] -= u(1.0, yi + h) * coefs[8];
      rhs[0] += u(0.0, 0.0) * coefs[0];
rhs[(dp - 1)] += u(1.0, 0.0) * coefs[2];
rhs[(dp - 1) * dp] += u(0.0, 1.0) * coefs[6];
      rhs[dp * dp - 1] += u(1.0, 1.0) * coefs[8];
template <typename ValueType, typename IndexType>
void print_solution(IndexType dp, const ValueType* u)
      for (IndexType i = 0; i < dp; ++i) {</pre>
```

```
for (IndexType j = 0; j < dp; ++j) {
    std::cout « u[i * dp + j] « ' ';</pre>
         std::cout « '\n';
    std::cout « std::endl;
template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType dp, const ValueType* u,
                                                        Closure correct u)
    const ValueType h = 1.0 / (dp + 1);
    gko::remove_complex<ValueType> error = 0.0;
     for (IndexType j = 0; j < dp; ++j) {</pre>
         const auto xi = ValueType(j + 1) * h;
         for (IndexType i = 0; i < dp; ++i) {</pre>
             using std::abs;
const auto yi = ValueType(i + 1) * h;
              error +=
                  abs(u[i * dp + j] - correct_u(xi, yi)) / abs(correct_u(xi, yi));
    return error;
template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                     unsigned int discretization_points, IndexType* row_ptrs,
                     IndexType* col_idxs, ValueType* values, ValueType* rhs,
                     ValueType* u, gko::remove_complex<ValueType> reduction_factor)
{
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto@ dp = discretization_points;
const gko::size_type dp_2 = dp * dp;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         exec_map{
              {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
               [] {
                    return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
              {"hip",
               [] {
                    return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                          true);
               }},
              {"dpcpp",
                    return gko::DpcppExecutor::create(0,
                                                            gko::OmpExecutor::create());
              {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
    auto matrix = mtx::create(
         exec, gko::dim<2>(dp_2),
         val_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), values),
idx_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), col_idxs),
idx_array::view(app_exec, dp_2 + 1, row_ptrs));
    auto b = vec::create(exec, gko::dim<2>(dp_2, 1),
                             val_array::view(app_exec, dp_2, rhs), 1);
    auto x = vec::create(app_exec, gkc::dim<2>(dp_2, 1), val_array::view(app_exec, dp_2, u), 1);
    auto solver_gen =
         cg::build()
              .with_criteria(
                   gko::stop::Iteration::build().with_max_iters(dp_2).on(exec),
                   gko::stop::ResidualNorm<ValueType>::build()
                       .with_reduction_factor(reduction_factor)
                        .on(exec))
              .with_preconditioner(bj::build().on(exec))
              .on(exec);
    auto solver = solver_gen->generate(gko::give(matrix));
    solver->apply(b, x);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && std::string(argv[1]) == "--help") {
         std::cerr
```

```
« "Usage: " « argv[0]
          « " [executor] [DISCRETIZATION_POINTS] [alpha] [beta] [gamma]"
          « std::endl;
     std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const IndexType discretization_points =
argc >= 3 ? std::atoi(argv[2]) : 100;

const ValueType alpha_c = argc >= 4 ? std::atof(argv[3]) : default_alpha;

const ValueType beta_c = argc >= 5 ? std::atof(argv[4]) : default_beta;

const ValueType gamma_c = argc >= 6 ? std::atof(argv[5]) : default_gamma;
std::array<ValueType, 9> coefs{
     gamma_c, beta_c, gamma_c, beta_c, alpha_c, beta_c,
     gamma_c, beta_c, gamma_c);
const auto dp = discretization_points;
const size_t dp_2 = dp * dp;
auto correct_u = [](ValueType x, ValueType y) {
    return x * x * x + y * y * y;
};
auto f = [](ValueType x, ValueType y) {
    return ValueType(6) * x + ValueType(6) * y;
};
std::vector<IndexType> row_ptrs(dp_2 + 1);
std::vector<IndexType> col_idxs((3 * dp - 2) * (3 * dp - 2));
std::vector<ValueType> values((3 * dp - 2) * (3 * dp - 2));
std::vector<ValueType> rhs(dp_2);
std::vector<ValueType> u(dp_2, 0.0);
generate_stencil_matrix(dp, row_ptrs.data(), col_idxs.data(), values.data(),
                              coefs.data());
generate_rhs(dp, f, correct_u, rhs.data(), coefs.data());
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
auto start_time = std::chrono::steady_clock::now();
solve_system(executor_string, dp, row_ptrs.data(), col_idxs.data(),
                values.data(), rhs.data(), u.data(), reduction_factor);
auto stop_time = std::chrono::steady_clock::now();
auto runtime duration =
     static_cast<double>(
         std::chrono::duration_cast<std::chrono::nanoseconds>(stop_time
               .count()) *
    1e-6:
std::cout « "The average relative error is "
            « calculate_error(dp, u.data(), correct_u) /
                    static_cast<gko::remove_complex<ValueType>> (dp_2)
            « std::endl;
std::cout « "The runtime is " « std::to_string(runtime_duration) « " ms"
            « std::endl;
```

The papi-logging program

The papi logging example..

This example depends on simple-solver-logging.

Introduction

About the example

```
#include <ginkgo/ginkgo.hpp>
#include <ginkgo/g
#include <papi.h>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <thread>
namespace {
void papi_add_event(const std::string& event_name, int& eventset)
     int ret_val = PAPI_event_name_to_code(event_name.c_str(), &code);
     if (PAPI_OK != ret_val) {
    std::cerr « "Error at PAPI_name_to_code()" « std::endl;
          std::exit(-1);
     ret_val = PAPI_add_event(eventset, code);
     if (PAPI_OK != ret_val) {
          std::cerr « "Error at PAPI_name_to_code()" « std::endl;
          std::exit(-1);
template <typename T>
std::string to_string(T* ptr)
     std::ostringstream os;
    os « reinterpret_cast<gko::uintptr>(ptr);
return os.str();
int init_papi_counters(std::string solver_name, std::string A_name)
Initialize PAPI, add events and start it up
     int eventset = PAPI_NULL;
int ret_val = PAPI_library_init(PAPI_VER_CURRENT);
if (ret_val != PAPI_VER_CURRENT) {
    std::cerr « "Error at PAPI_library_init()" « std::endl;
          std::exit(-1);
```

```
ret_val = PAPI_create_eventset(&eventset);
     if (PAPI_OK != ret_val) {
    std::cerr « "Error at PAPI_create_eventset()" « std::endl;
          std::exit(-1);
     ,std::string simple_apply_string("sde:::ginkgo0::linop_apply_completed::");
std::string advanced_apply_string(
          "sde:::ginkgo0::linop_advanced_apply_completed::");
     papi_add_event(simple_apply_string + solver_name, eventset);
papi_add_event(simple_apply_string + A_name, eventset);
     papi_add_event(advanced_apply_string + A_name, eventset);
     ret_val = PAPI_start(eventset);
     if (PAPI_OK != ret_val) {
   std::cerr « "Error at PAPI_start()" « std::endl;
          std::exit(-1);
     return eventset:
void print_papi_counters(int eventset)
Stop PAPI and read the linop apply completed event for all of them
long long int values[3];
int ret_val = PAPI_stop(eventset, values);
if (PAPI_OK != ret_val) {
     std::cerr « "Error at PAPI_stop()" « std::endl;
     std::exit(-1);
PAPI_shutdown();
Print all values returned from PAPI
     std::cout « "PAPI SDE counters:" « std::endl;
std::cout « "solver did " « values[0] « " applies." « std::endl;
std::cout « "A did " « values[1] « " simple applies." « std::endl;
std::cout « "A did " « values[2] « " advanced applies." « std::endl;
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argy[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1):
Figure out where to run the code
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
     exec map{
                   [] { return gko::OmpExecutor::create(); }},
           {"omp",
           {"cuda",
                 return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                           true);
           }},
           {"hip",
                 return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                          true);
           { "dpcpp",
           [] {
                return gko::DpcppExecutor::create(0,
                                                            gko::OmpExecutor::create());
           {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
Generate solver
const RealValueType reduction_factor{1e-7};
auto solver gen =
    cg::build()
       .with_criteria(
            gko::stop::Iteration::build().with_max_iters(20u).on(exec),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor)
                .on(exec))
        .on(exec);
auto solver = solver_gen->generate(A);
In this example, we split as much as possible the Ginkgo solver/logger and the PAPI interface. Note that the PAPI
ginkgo namespaces are of the form sde:::ginkgo<x> where <x> starts from 0 and is incremented with every new
PAPI logger.
int eventset
    init_papi_counters(to_string(solver.get()), to_string(A.get()));
Create a PAPI logger and add it to relevant LinOps
auto logger = gko::log::Papi<ValueType>::creat
    gko::log::Logger::linop_apply_completed_mask |
    gko::log::Logger::linop_advanced_apply_completed_mask);
solver->add_logger(logger);
A->add_logger(logger);
Solve system
solver->apply(b, x);
Stop PAPI event gathering and print the counters
print_papi_counters(eventset);
Print solution
std::cout « "Solution (x): \n";
write(std::cout, x);
Calculate residual
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto res = gko::initialize<real vec>({0.0}, exec);
    A->apply(one, x, neq_one, b);
    b->compute_norm2(res);
    std::cout « "Residual norm sqrt(r^T r): \n";
    write(std::cout, res);
```

Results

The following is the expected result:

```
PAPI SDE counters:
solver did 1 applies.
A did 20 simple applies.
A did 1 advanced applies.
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
8.87107e-16
```

Comments about programming and debugging

```
*********GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
    Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <papi.h>
#include <fstream>
#include <iostream>
#include <map>
#include <string:
#include <thread>
namespace {
void papi_add_event(const std::string& event_name, int& eventset)
    int ret_val = PAPI_event_name_to_code(event_name.c_str(), &code);
    if (PAPI_OK != ret_val) {
        std::cerr « "Error at PAPI_name_to_code()" « std::endl;
        std::exit(-1);
    ret_val = PAPI_add_event(eventset, code);
    if (PAPI_OK != ret_val) {
        std::cerr « "Error at PAPI_name_to_code()" « std::endl;
        std::exit(-1);
    }
template <typename T>
std::string to_string(T* ptr)
    std::ostringstream os;
    os « reinterpret_cast<gko::uintptr>(ptr);
    return os.str();
   // namespace
int init_papi_counters(std::string solver_name, std::string A_name)
    int eventset = PAPI_NULL;
    int ret_val = PAPI_library_init(PAPI_VER_CURRENT);
if (ret_val != PAPI_VER_CURRENT) {
        std::cerr « "Error at PAPI_library_init()" « std::endl;
        std::exit(-1);
    ret_val = PAPI_create_eventset(&eventset);
    if (PAPI_OK != ret_val) {
    std::cerr « "Error at PAPI_create_eventset()" « std::endl;
        std::exit(-1);
    std::string simple_apply_string("sde:::ginkgo0::linop_apply_completed::");
    std::string advanced_apply_string(
         "sde:::ginkgo0::linop_advanced_apply_completed::");
    papi_add_event(simple_apply_string + solver_name, eventset);
papi_add_event(simple_apply_string + A_name, eventset);
    papi_add_event(advanced_apply_string + A_name, eventset);
```

```
ret_val = PAPI_start(eventset);
    if (PAPI_OK != ret_val) {
    std::cerr « "Error at PAPI_start()" « std::endl;
         std::exit(-1);
    return eventset:
void print_papi_counters(int eventset)
    long long int values[3];
    int ret_val = PAPI_stop(eventset, values);
    if (PAPI_OK != ret_val) {
         std::cerr « "Error at PAPI_stop()" « std::endl;
         std::exit(-1);
    PAPI shutdown();
    std::cout « "PAPI SDE counters:" « std::endl;
    std::cout « "FAFI SDE Counters: « Std::endf; std::cout « "solver did " « values[0] « " applies." « std::endl; std::cout « "A did " « values[1] « " simple applies." « std::endl; std::cout « "A did " « values[2] « " advanced applies." « std::endl;
int main(int argc, char* argv[])
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
         std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
         exec_map{
              -
{"omp",
                      [] { return gko::OmpExecutor::create(); }},
              {"cuda",
               [] {
                   return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                           true):
               }},
              {"hip",
               [] {
                   return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                         true);
               }},
              {"dpcpp",
               [] {
                   return gko::DpcppExecutor::create(0,
                                                            gko::OmpExecutor::create());
              }},
{"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    const RealValueType reduction_factor{1e-7};
    auto solver_gen =
         cq::build()
              .with_criteria(
                  gko::stop::Iteration::build().with_max_iters(20u).on(exec),
                  gko::stop::ResidualNorm<ValueType>::build()
                       .with_reduction_factor(reduction_factor)
                       .on(exec))
              .on(exec);
    auto solver = solver_gen->generate(A);
    int eventset =
         init_papi_counters(to_string(solver.get()), to_string(A.get()));
    auto logger = gko::log::Papi<ValueType>::create(
         gko::log::Logger::linop_apply_completed_mask |
         gko::log::Logger::linop_advanced_apply_completed_mask);
    solver->add logger(logger);
    A->add_logger(logger);
    solver->apply(b, x);
    print_papi_counters(eventset);
std::cout « "Solution (x): \n";
    write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>((-1.0), exec);
    auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r): \n";
    write(std::cout, res);
```

The par-ilu-convergence program

The ParILU convergence example..

This example depends on simple-solver.

Introduction

About the example

This example can be used to inspect the convergence behavior of parallel incomplete factorizations. *

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <functional>
#include <iostream>
#include <memory>
#include <string>
const std::map<std::string, std::function<std::shared_ptr<gko::Executor>() >>
    executors{
        {"reference", [] { return gko::ReferenceExecutor::create(); }},
        {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create());
         }},
         {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         {"dpcpp", [] {
              return gko::DpcppExecutor::create(0, gko::OmpExecutor::create());
         } } ;
template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
    decltype(fun()) result;
        result = fun();
    } catch (const gko::Error& err) {
   std::cerr « "Error: " « err.what() « '\n';
        std::exit(-1);
    return result;
template <typename ValueType, typename IndexType>
double compute_ilu_residual_norm(
```

```
const gko::matrix::Csr<ValueType, IndexType>* residual,
    const gko::matrix::Csr<ValueType, IndexType>* mtx)
    gko::matrix_data<ValueType, IndexType> residual_data;
    gko::matrix_data<ValueType, IndexType> mtx_data;
    residual->write(residual_data);
    mtx->write(mtx_data);
    residual_data.ensure_row_major_order();
    mtx_data.ensure_row_major_order();
    auto it = mtx data.nonzeros.begin();
    double residual_norm{};
    for (auto entry : residual_data.nonzeros) {
   auto ref_row = it->row;
         auto ref_col = it->column;
         if (entry.row == ref_row && entry.column == ref_col) {
             residual_norm += gko::squared_norm(entry.value);
             ++i+:
    return std::sqrt(residual_norm);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
print usage message
   (argc < 2 || executors.find(argv[1]) == executors.end()) {</pre>
    std::cerr « "Usage: executable"
               « " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
                  "[[[cariful pariful paric|parict] [<max-iterations>] "
"[<num-repetitions>] [<fill-in-limit>]\n";
    return -1;
generate executor based on first argument
auto exec = try_generate([&] { return executors.at(argv[1])(); });
set matrix and preconditioner name with default values
std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
std::string precond = argc < 4 ? "parilu" : argv[3];</pre>
int max_iterations = argc < 5 ? 10 : std::stoi(argv[4]);
int num_repetitions = argc < 6 ? 10 : std::stoi(argv[5]);
double limit = argc < 7 ? 2 : std::stod(argv[6]);</pre>
load matrix file into Csr format
    auto mtx = gko::share(try_generate([&] {
         std::ifstream mtx_stream{matrix};
         if (!mtx_stream) {
             throw GKO_STREAM_ERROR("Unable to open matrix file");
         std::cerr « "Reading " « matrix « std::endl;
         return gko::read<gko::matrix::Csr<ValueType, IndexType»(mtx_stream,</pre>
    }));
    std::shared_ptr<gko::LinOpFactory> factory;
    std::function<void(int)> set_iterations;
    if (precond == "parilu") {
         factory =
             gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
         set iterations = [&](int it) {
             gko::as<gko::factorization::ParIlu<ValueType, IndexType>::Factory>(
                 factory)
                 ->get_parameters()
                 .iterations = it;
    } else if (precond == "paric") {
        factory =
             gko::factorization::ParIc<ValueType, IndexType>::build().on(exec);
         set_iterations = [&](int it) {
             gko::as<gko::factorization::ParIc<ValueType, IndexType>::Factory>(
                factory)
                 ->get_parameters()
                 .iterations = it;
        };
    } else if (precond == "parilut") {
         factory = gko::factorization::ParIlut<ValueType, IndexType>::build()
                        .with_fill_in_limit(limit)
                        .on(exec);
         set_iterations = [&](int it) {
             gko::as<gko::factorization::ParIlut<ValueType, IndexType>::Factory>(
                 factory)
                  ->get_parameters()
```

```
.iterations = it;
} else if (precond == "parict") {
   factory = gko::factorization::ParIct<ValueType, IndexType>::build()
                  .with_fill_in_limit(limit)
                  .on(exec);
    set_iterations = [&](int it) {
        gko::as<gko::factorization::ParIct<ValueType, IndexType>::Factory>(
           factory)
            ->get_parameters()
            .iterations = it;
   };
auto one = gko::initialize<gko::matrix::Dense<ValueType»({1.0}, exec);</pre>
    gko::initialize<gko::matrix::Dense<ValueType»({-1.0}, exec);</pre>
for (int it = 1; it <= max_iterations; ++it) {</pre>
   set_iterations(it);
    std::cout « it « ';';
    std::vector<long> times;
    std::vector<double> residuals;
    for (int rep = 0; rep < num_repetitions; ++rep) {</pre>
       auto tic = std::chrono::high_resolution_clock::now();
       auto result =
           gko::as<gko::Composition<ValueType»(factory->generate(mtx));
        exec->synchronize();
        auto toc = std::chrono::high_resolution_clock::now();
        auto residual = gko::clone(exec, mtx);
        result->get_operators()[0]->apply(one, result->get_operators()[1],
                                           minus_one, residual);
       times.push back (
            std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic)
        residuals.push_back(
            compute_ilu_residual_norm(residual.get(), mtx.get()));
    for (auto el : times) {
        std::cout « el « ';';
    for (auto el : residuals) {
        std::cout « el « ';';
   std::cout « '\n';
```

Results

This is the expected output:

```
Usage: executable <reference|omp|cuda|hip|dpcpp> [<matrix-file>] [<parilu|parilut|paric|parict] [<max-iterations>] [<num-repetitions>] [fill-in-limit]
```

When specifying an executor:

```
Reading data/A.mtx
1;71800;10300;8800;8200;8000;7700;7500;7500;7500;7500;7400;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0331e-14;1.0031e-16;4.15407e-16;4.15407e-16;4.15407e-16;4.15407e-16;4.1540
```

Comments about programming and debugging

```
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <functional>
#include <iostream>
#include <map>
#include <memory>
#include <string
const std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    executors{
         "reference", [] { return gko::ReferenceExecutor::create(); }},
"omp", [] { return gko::OmpExecutor::create(); }},
"cuda",
          [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create());
          }},
         {"hip",
          [] {
              return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         {"dpcpp", [] {
              return gko::DpcppExecutor::create(0, gko::OmpExecutor::create());
          }}};
template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
    decltype(fun()) result;
    try {
        result = fun();
    } catch (const gko::Error& err) {
  std::cerr « "Error: " « err.what() « '\n';
        std::exit(-1);
    return result;
template <typename ValueType, typename IndexType>
double compute_ilu_residual_norm(
    const gko::matrix::Csr<ValueType, IndexType>* residual,
    const gko::matrix::Csr<ValueType, IndexType>* mtx)
    gko::matrix_data<ValueType, IndexType> residual_data;
    gko::matrix_data<ValueType, IndexType> mtx_data;
    residual->write(residual data);
    mtx->write(mtx_data);
    residual_data.ensure_row_major_order();
    mtx_data.ensure_row_major_order();
    auto it = mtx_data.nonzeros.begin();
    double residual_norm{};
    for (auto entry : residual_data.nonzeros) {
   auto ref_row = it->row;
         auto ref_col = it->column;
         if (entry.row == ref_row && entry.column == ref_col) {
             residual_norm += gko::squared_norm(entry.value);
             ++it;
    return std::sqrt(residual_norm);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int:
```

```
if (argc < 2 || executors.find(argv[1]) == executors.end()) {</pre>
    std::cerr « "Usage: executable"
                « " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
                   "[<parilu|parilut|paric|parict] [<max-iterations>] "
                    "[<num-repetitions>] [<fill-in-limit>]\\n";
    return -1:
auto exec = try_generate([&] { return executors.at(argv[1])(); });
std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
std::string precond = argc < 4 ? "parilu" : argv[3];
int max_iterations = argc < 5 ? 10 : std::stoi(argv[4]);
int num_repetitions = argc < 6 ? 10 : std::stoi(argv[5]);
double limit = argc < 7 ? 2 : std::stod(argv[6]);</pre>
auto mtx = gko::share(try_generate([&] {
     std::ifstream mtx_stream{matrix};
     if (!mtx_stream) {
         throw GKO_STREAM_ERROR("Unable to open matrix file");
    std::cerr « "Reading " « matrix « std::endl;
    return gko::read<gko::matrix::Csr<ValueType, IndexType»(mtx_stream,</pre>
std::shared_ptr<gko::LinOpFactory> factory;
std::function<void(int)> set_iterations;
if (precond == "parilu") {
    factory =
        gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
     set_iterations = [&](int it) {
         gko::as<gko::factorization::ParIlu<ValueType, IndexType>::Factory>(
             factory)
             ->get parameters()
             .iterations = it;
} else if (precond == "paric") {
    factory =
    gko::factorization::ParIc<ValueType, IndexType>::build().on(exec);
    set iterations = [&](int it) {
         gko::as<gko::factorization::ParIc<ValueType, IndexType>::Factory>(
             factory)
             ->get_parameters()
             .iterations = it;
    };
} else if (precond == "parilut") {
    factory = gko::factorization::ParIlut<ValueType, IndexType>::build()
                    .with_fill_in_limit(limit)
                    .on(exec);
    set_iterations = [&](int it) {
         gko::as<gko::factorization::ParIlut<ValueType, IndexType>::Factory>(
             factory)
             ->get parameters()
             .iterations = it;
} else if (precond == "parict") {
   factory = gko::factorization::ParIct<ValueType, IndexType>::build()
                    .with_fill_in_limit(limit)
                    .on(exec);
    set_iterations = [&](int it) {
         gko::as<gko::factorization::ParIct<ValueType, IndexType>::Factory>(
             factory)
             ->get_parameters()
             .iterations = it;
    };
auto one = gko::initialize<gko::matrix::Dense<ValueType»({1.0}, exec);</pre>
    gko::initialize<gko::matrix::Dense<ValueType»({-1.0}, exec);</pre>
for (int it = 1; it <= max_iterations; ++it) {</pre>
    set_iterations(it);
    std::cout « it « ';';
    std::vector<long> times;
     std::vector<double> residuals;
     for (int rep = 0; rep < num_repetitions; ++rep) {</pre>
         auto tic = std::chrono::high_resolution_clock::now();
         auto result =
             gko::as<gko::Composition<ValueType»(factory->generate(mtx));
         exec->synchronize();
         auto toc = std::chrono::high_resolution_clock::now();
         auto residual = gko::clone(exec, mtx);
         result->get_operators()[0]->apply(one, result->get_operators()[1],
                                               minus one, residual);
         times.push back (
             std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic)
         residuals.push_back(
             compute_ilu_residual_norm(residual.get(), mtx.get()));
     for (auto el : times) {
```

```
std::cout « el « ';';
}
for (auto el : residuals) {
    std::cout « el « ';';
}
std::cout « '\n';
}
```

The performance-debugging program

The simple solver with performance debugging example..

This example depends on simple-solver-logging, minimal-cuda-solver.

Introduction

About the example

This example runs a solver on a test problem and shows how to use loggers to debug performance and convergence rate.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <algorithm>
#include <array>
#include <chrono>
#include <cstdlib>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <ostream>
#include <sstream>
#include <string>
#include <utility>
#include <vector>
template <typename ValueType>
using vec = gko::matrix::Dense<ValueType>;
template <typename ValueType>
using real_vec = gko::matrix::Dense<gko::remove_complex<ValueType>>;
namespace utils {
creates a zero vector
template <typename ValueType>
std::unique_ptr<vec<ValueType» create_vector(</pre>
    std::shared_ptr<const gko::Executor> exec, gko::size_type size,
    ValueType value)
    auto res = vec<ValueType>::create(exec);
    res->read(gko::matrix_data<ValueType>(gko::dim<2>{size, 1}, value));
```

utilities for computing norms and residuals

```
template <typename ValueType>
ValueType get_first_element(const vec<ValueType>* norm)
    return norm->get_executor()->copy_val_to_host(norm->get_const_values());
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(const vec<ValueType>* b)
    auto exec = b->get_executor();
    auto b_norm = gko::initialize<real_vec<ValueType»({0.0}, exec);</pre>
    b->compute_norm2(b_norm);
    return get_first_element(b_norm.get());
template <typename ValueType>
gko::remove_complex<ValueType> compute_residual_norm(
    const gko::LinOp* system_matrix, const vec<ValueType>* b,
    const vec<ValueType>* x)
    auto exec = system_matrix->get_executor();
    auto one = gko::initialize<vec<ValueType»({1.0}, exec);</pre>
    auto neg_one = gko::initialize<vec<ValueType»({-1.0}, exec);</pre>
    auto res = gko::clone(b);
    system_matrix->apply(one, x, neg_one, res);
    return compute_norm(res.get());
   // namespace utils
namespace loggers {
```

A logger that accumulates the time of all operations. For each operation type (allocations, free, copy, internal operations i.e. kernels), the timing is taken before and after. This can create significant overhead since to ensure proper timings, calls to synchronize are required.

```
struct OperationLogger : gko::log::Logger {
    void on_allocation_started(const gko::Executor* exec,
                               const gko::size_type&)const override
       this->start_operation(exec, "allocate");
   void on_allocation_completed(const gko::Executor* exec,
                                const gko::size type&,
                                 const gko::uintptr&)const override
       this->end_operation(exec, "allocate");
   void on_free_started(const gko::Executor* exec,
                         const gko::uintptr&)const override
       this->start_operation(exec, "free");
   void on_free_completed(const gko::Executor* exec,
                          const gko::uintptr&)const override
       this->end operation(exec, "free");
   void on_copy_started(const gko::Executor* from, const gko::Executor* to,
                        const gko::uintptr&, const gko::uintptr&,
                         const gko::size_type&)const override
{
       from->synchronize();
       this->start_operation(to, "copy");
   const gko::size_type&)const override
{
       from->synchronize();
       this->end_operation(to, "copy");
   void on_operation_launched(const gko::Executor* exec,
                               const gko::Operation* op)const override
       this->start operation(exec, op->get name());
   void on_operation_completed(const gko::Executor* exec,
                               const gko::Operation* op)const override
       this->end_operation(exec, op->get_name());
   void write_data(std::ostream& ostream)
        for (const auto& entry : total) {
            ostream « "\t" « entry.first.c_str() « ": "
                    \mbox{\tt w std::} chrono:: duration\_cast < std:: chrono:: nanoseconds > (
                          entry.second)
                           .count()
                    « std::endl;
```

```
}
private:
```

Helper which synchronizes and starts the time before every operation.

Helper to compute the end time and store the operation's time at its end. Also time nested operations.

```
void end_operation(const gko::Executor* exec, const std::string& name)const
{
    exec->synchronize();
    const auto end = std::chrono::steady_clock::now();
    const auto diff = end - start[name];
```

make sure timings for nested operations are not counted twice

```
total[name] += diff - nested.back();
nested.pop_back();
if (nested.size() > 0) {
    nested.back() += diff;
}
mutable std::map<std::string, std::chrono::steady_clock::time_point> start;
mutable std::map<std::string, std::chrono::steady_clock::duration> total;
```

the position i of this vector holds the total time spend on child operations on nesting level i

```
mutable std::vector<std::chrono::steady_clock::duration> nested;
```

This logger tracks the persistently allocated data

```
struct StorageLogger : gko::log::Logger {
```

Store amount of bytes allocated on every allocation

Reset the amount of bytes on every free

Write the data after summing the total from all allocations

```
void write_data(std::ostream& ostream)
{
    gko::size_type total{};
    for (const auto& e : storage) {
        total += e.second;
    }
    ostream « "Storage: " « total « std::endl;
}
private:
    mutable std::unordered_map<gko::uintptr, gko::size_type> storage;
};
```

Logs true and recurrent residuals of the solver

```
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
```

Depending on the available information, store the norm or compute it from the residual. If the true residual norm could not be computed, store the value -1.0.

```
rec_res_norms.push_back(utils::get_first_element(
                  gko::as<real_vec<ValueType»(residual_norm)));</pre>
         } else {
              rec_res_norms.push_back(
                 utils::compute_norm(gko::as<vec<ValueType»(residual)));
              true_res_norms.push_back(utils::compute_residual_norm(
                 matrix, b, gko::as<vec<ValueType»(solution)));</pre>
         } else {
              true_res_norms.push_back(-1.0);
    ResidualLogger(const gko::LinOp* matrix, const vec<ValueType>* b)
         : gko::log::Logger(gko::log::Logger::iteration_complete_mask),
           matrix{matrix},
           h{h}
    {}
    void write_data(std::ostream& ostream)
         ostream \ll "Recurrent Residual Norms: " \ll std::endl; ostream \ll "[" \ll std::endl;
         for (const auto& entry : rec_res_norms) {
    ostream « "\t" « entry « std::endl;
         ostream « "];" « std::endl;
         ostream « "True Residual Norms: " « std::endl;
         ostream « "[" « std::endl;
         for (const auto& entry : true_res_norms) {
  ostream « "\t" « entry « std::endl;
         ostream « "]; " « std::endl;
    }
private:
    const gko::LinOp* matrix;
    const vec<ValueType>* b;
    mutable std::vector<gko::remove_complex<ValueType> rec_res_norms;
    mutable std::vector<gko::remove_complex<ValueType> true_res_norms;
  // namespace loggers
namespace {
Print usage help
void print_usage(const char* filename)
    std::cerr « "Usage: " « filename « " [executor] [matrix file]"
                « std::endl;
    \verb|std::cerr & \verb|matrix file should be a file in matrix market format. | \verb| "
                  "The file data/A.mtx is provided as an example."
                « std::endl;
    std::exit(-1);
template <typename ValueType>
void print_vector(const gko::matrix::Dense<ValueType>* vec)
    auto elements_to_print = std::min(gko::size_type(10), vec->get_size()[0]);
    std::cout « "[" « std::endl;
for (int i = 0; i < elements_to_print; ++i) {
    std::cout « "\t" « vec->at(i) « std::endl;
    std::cout « "];" « std::endl;
   // namespace
int main(int argc, char* argv[])
Parametrize the benchmark here Pick a value type
using ValueType = double;
using IndexType = int;
Pick a matrix format
using mtx = gko::matrix::Csr<ValueType, IndexType>;
Pick a solver
using solver = gko::solver::Cg<ValueType>;
Pick a preconditioner type
using preconditioner = gko::matrix::IdentityFactory<ValueType>;
Pick a residual norm reduction value
const gko::remove_complex<ValueType> reduction_factor = 1e-12;
```

```
Pick an output file name
const auto of_name = "log.txt";
Simple shortcut
using vec = gko::matrix::Dense<ValueType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
Figure out where to run the code
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec map{
         {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                true):
         }},
        {"hip",
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                              true);
         }},
        { "dpcpp",
             return gko::DpcppExecutor::create(0,
                                                 gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
Read the input matrix file directory
std::string input_mtx = "data/A.mtx";
if (argc == 3)
    input_mtx = std::string(argv[2]);
Read data: A is read from disk Create a StorageLogger to track the size of A
auto storage_logger = std::make_shared<loggers::StorageLogger>();
Add the logger to the executor
exec->add_logger(storage_logger);
Read the matrix A from file
auto A = gko::share(gko::read<mtx>(std::ifstream(input mtx), exec));
Remove the storage logger
exec->remove_logger(storage_logger);
Pick a maximum iteration count
const auto max_iters = A->get_size()[0];
Generate b and x vectors
auto b = utils::create_vector<ValueType>(exec, A->get_size()[0], 1.0);
auto x = utils::create_vector<ValueType>(exec, A->get_size()[0], 0.0);
Declare the solver factory. The preconditioner's arguments should be adapted if needed.
auto solver factory
    solver::build()
        .with_criteria(
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor)
                 .on(exec),
            gko::stop::Iteration::build().with_max_iters(max_iters).on(
```

exec))

.with_preconditioner(preconditioner::create(exec))

```
.on(exec);
```

```
Declare the output file for all our loggers
```

```
std::ofstream output_file(of_name);
```

Do a warmup run

{

Clone x to not overwrite the original one

```
auto x_clone = gko::clone(x);
```

Generate and call apply on a solver

```
solver_factory->generate(A)->apply(b, x_clone);
exec->synchronize();
```

Do a timed run

{

Clone x to not overwrite the original one

auto x_clone = gko::clone(x);

Synchronize ensures no operation are ongoing

exec->synchronize();

Time before generate

auto g_tic = std::chrono::steady_clock::now();

Generate a solver

```
auto generated_solver = solver_factory->generate(A);
exec->synchronize();
```

-. . . .

Time after generate
auto g_tac = std::chrono::steady_clock::now();

Compute the generation time

```
auto generate_time =
    std::chrono::duration_cast<std::chrono::nanoseconds>(g_tac - g_tic);
```

Write the generate time to the output file

Similarly time the apply

```
exec->synchronize();
auto a_tic = std::chrono::steady_clock::now();
generated_solver->apply(b, x_clone);
exec->synchronize();
auto a_tac = std::chrono::steady_clock::now();
auto apply_time =
    std::chrono::duration_cast<std::chrono::nanoseconds>(a_tac - a_tic);
output_file « "Apply time (ns): " « apply_time.count() « std::endl;
```

Compute the residual norm

```
auto residual =
    utils::compute_residual_norm(A.get(), b.get(), x_clone.get());
output_file « "Residual_norm: " « residual « std::endl;
```

Log the internal operations using the OperationLogger without timing

{

Create an OperationLogger to analyze the generate step

```
auto gen_logger = std::make_shared<loggers::OperationLogger>();
```

Add the generate logger to the executor

```
exec->add_logger(gen_logger);
```

Generate a solver

auto generated_solver = solver_factory->generate(A);

Remove the generate logger from the executor

```
exec->remove_logger(gen_logger);
```

Write the data to the output file

```
output_file « "Generate operations times (ns):" « std::endl;
gen_logger->write_data(output_file);
```

Create an OperationLogger to analyze the apply step

```
auto apply_logger = std::make_shared<loggers::OperationLogger>();
exec->add_logger(apply_logger);
```

Create a ResidualLogger to log the recurent residual

```
auto res_logger = std::make_shared<loggers::ResidualLogger<ValueType»(
    A.get(), b.get());
generated_solver->add_logger(res_logger);
```

Solve the system

```
generated_solver->apply(b, x);
exec->remove_logger(apply_logger);
```

Write the data to the output file

```
output_file « "Apply operations times (ns):" « std::endl;
apply_logger->write_data(output_file);
res_logger->write_data(output_file);
```

Print solution

```
std::cout \ "Solution, first ten entries: \ "; print_vector(x.get());
```

Print output file location

Results

This is the expected standard output:

Here is a sample output in the file log.txt:

```
Generate time (ns): 861
Apply time (ns): 108144
Residual_norm: 2.10788e-15
Generate operations times (ns):
Apply operations times (ns):
allocate: 14991
cg::initialize#8: 872
cg::step_1#5: 7683
cg::step_2#7: 7756
copy: 7751
csr::advanced_spmv#5: 21819
csr::spmv#3: 20429
dense::compute_dot#3: 18043
dense::compute_norm2#2: 16726
free: 8857
residual_norm::residual_norm#9: 3614
Recurrent Residual Norms:
[
4.3589
2.30455
```

```
1.46771
    0.984875
    0.741833
    0.513623
    0.384165
    0.316439
    0.227709
    0.170312
    0.0973722
    0.0616831
    0.0454123
    0.031953
    0.0161606
    0.00657015
    0.00264367
    0.000858809
    0.000286461
    1.64195e-15
];
True Residual Norms:
    4.3589
    2.30455
    1.46771
    0.984875
    0.741833
    0.513623
    0.384165
    0.316439
    0.227709
    0.170312
    0.0973722
    0.0616831
    0.0454123
    0.031953
    0.0161606
    0.00657015
    0.00264367
    0.000858809
    0.000286461
    2.10788e-15
1;
```

Comments about programming and debugging

The plain program

```
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright \ensuremath{\text{copyright}}
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <algorithm>
#include <array>
```

************GINKGO LICENSE>*****************

```
#include <chrono>
#include <cstdlib>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <ostream>
#include <sstream>
#include <string>
#include <unordered_map>
#include <utility>
#include <vector>
template <typename ValueType>
using vec = gko::matrix::Dense<ValueType>;
template <typename ValueType>
using real_vec = gko::matrix::Dense<gko::remove_complex<ValueType>>;
namespace utils {
template <typename ValueType>
std::unique_ptr<vec<ValueType» create_vector(
    std::shared_ptr<const gko::Executor> exec, gko::size_type size,
    ValueType value)
    auto res = vec<ValueType>::create(exec);
    res->read(gko::matrix_data<ValueType>(gko::dim<2>{size, 1}, value));
    return res;
template <typename ValueType>
ValueType get_first_element(const vec<ValueType>* norm)
    return norm->get_executor()->copy_val_to_host(norm->get_const_values());
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(const vec<ValueType>* b)
    auto exec = b->get_executor();
    auto b_norm = gko::initialize<real_vec<ValueType»({0.0}, exec);</pre>
    b->compute_norm2(b_norm);
    return get_first_element(b_norm.get());
template <typename ValueType>
gko::remove_complex<ValueType> compute_residual_norm(
    const gko::LinOp* system_matrix, const vec<ValueType>* b,
    const vec<ValueType>* x)
    auto exec = system_matrix->get_executor();
    auto one = gko::initialize<vec<ValueType»({1.0}, exec);</pre>
    auto neg_one = gko::initialize<vec<ValueType»({-1.0}, exec);</pre>
    auto res = gko::clone(b);
    system_matrix->apply(one, x, neg_one, res);
    return compute_norm(res.get());
  // namespace utils
namespace loggers {
struct OperationLogger : gko::log::Logger {
    void on_allocation_started(const gko::Executor* exec,
                               const gko::size_type&)const override
{
       this->start_operation(exec, "allocate");
    void on_allocation_completed(const gko::Executor* exec,
                                 const gko::size_type&,
                                 const gko::uintptr&)const override
{
       this->end_operation(exec, "allocate");
    void on_free_started(const gko::Executor* exec,
                         const gko::uintptr&)const override
{
       this->start_operation(exec, "free");
    void on_free_completed(const gko::Executor* exec,
                           const gko::uintptr&)const override
{
        this->end_operation(exec, "free");
    void on_copy_started(const gko::Executor* from, const gko::Executor* to,
                         const gko::uintptr&, const gko::uintptr&,
                         const gko::size_type&)const override
{
        from->synchronize():
       this->start_operation(to, "copy");
    void on_copy_completed(const gko::Executor* from, const gko::Executor* to,
                           const gko::uintptr&, const gko::uintptr&,
                           const gko::size_type&)const override
        from->synchronize();
```

```
this->end_operation(to, "copy");
    void on_operation_launched(const gko::Executor* exec,
                                 const gko::Operation* op)const override
{
        this->start operation(exec, op->get name());
    void on_operation_completed(const gko::Executor* exec,
                                 const gko::Operation* op)const override
        this->end_operation(exec, op->get_name());
    void write_data(std::ostream& ostream)
        for (const auto& entry : total) {
            ostream « "\t" « entry.first.c_str() « ": "
                     « std::chrono::duration_cast<std::chrono::nanoseconds>(
                            entry.second)
                            .count()
                     « std::endl;
    }
private:
    void start_operation(const gko::Executor* exec,
                          const std::string& name)const
{
        nested.emplace_back(0);
        exec->synchronize();
        start[name] = std::chrono::steady_clock::now();
    void end operation(const gko::Executor* exec, const std::string& name)const
{
        exec->synchronize();
        const auto end = std::chrono::steady_clock::now();
const auto diff = end - start[name];
        total[name] += diff - nested.back();
        nested.pop_back();
        if (nested.size() > 0) {
            nested.back() += diff;
    mutable std::map<std::string, std::chrono::steady_clock::time_point> start;
mutable std::map<std::string, std::chrono::steady_clock::duration> total;
    mutable std::vector<std::chrono::steady_clock::duration> nested;
struct StorageLogger : gko::log::Logger {
    void on_allocation_completed(const gko::Executor*,
                                   const gko::size_type& num_bytes,
                                   const gko::uintptr& location)const override
{
        storage[location] = num_bytes;
    void on_free_completed(const gko::Executor*,
                            const gko::uintptr& location)const override
        storage[location] = 0;
    void write_data(std::ostream& ostream)
        gko::size_type total{};
        for (const auto& e : storage) {
   total += e.second;
        ostream « "Storage: " « total « std::endl;
    }
private:
   mutable std::unordered_map<gko::uintptr, gko::size_type> storage;
};
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    void on_iteration_complete(const gko::LinOp*, const gko::size_type&,
                                 const gko::LinOp* residual,
                                 const gko::LinOp* solution,
                                 const gko::LinOp* residual_norm)const override
{
        if (residual norm) {
            rec_res_norms.push_back(utils::get_first_element(
                 gko::as<real_vec<ValueType»(residual_norm)));</pre>
        } else {
            rec_res_norms.push_back(
                utils::compute_norm(gko::as<vec<ValueType»(residual)));
         if (solution) {
             true_res_norms.push_back(utils::compute_residual_norm(
                matrix, b, gko::as<vec<ValueType>(solution)));
        } else {
            true res norms.push back (-1.0);
```

```
}
     ResidualLogger(const gko::LinOp* matrix, const vec<ValueType>* b)
         : gko::log::Logger(gko::log::Logger::iteration_complete_mask),
           matrix{matrix},
           b{b}
     {}
     void write_data(std::ostream& ostream)
         ostream \ll "Recurrent Residual Norms: " \ll std::endl; ostream \ll "[" \ll std::endl;
         for (const auto& entry : rec_res_norms) {
             ostream « "\t" « entry « std::endl;
         ostream « "]; " « std::endl;
         ostream « "True Residual Norms: " « std::endl;
ostream « "[" « std::endl;
for (const auto& entry : true_res_norms) {
    ostream « "\t" « entry « std::endl;
         ostream « "]; " « std::endl;
    }
private:
    const gko::LinOp* matrix;
     const vec<ValueType>* b;
     mutable std::vector<gko::remove_complex<ValueType» rec_res_norms;</pre>
     mutable std::vector<gko::remove_complex<ValueType> true_res_norms;
} // namespace loggers
namespace {
void print_usage(const char* filename)
     std::cerr « "Usage: " « filename « " [executor] [matrix file]"
                « std::endl;
     std::cerr « "matrix file should be a file in matrix market format. "
The file data/A.mtx is provided as an example."
                « std::endl;
     std::exit(-1);
template <typename ValueType>
void print_vector(const gko::matrix::Dense<ValueType>* vec)
     auto elements_to_print = std::min(gko::size_type(10), vec->get_size()[0]);
     std::cout « "[" « std::endl;
for (int i = 0; i < elements_to_print; ++i) {
    std::cout « "\t" « vec->at(i) « std::endl;
     std::cout « "];" « std::endl;
   // namespace
int main(int argc, char* argv[])
     using ValueType = double;
     using IndexType = int;
     using mtx = gko::matrix::Csr<ValueType, IndexType>;
     using solver = gko::solver::Cg<ValueType>;
using preconditioner = gko::matrix::IdentityFactory<ValueType>;
     const gko::remove_complex<ValueType> reduction_factor = le-12;
const auto of_name = "log.txt";
     using vec = gko::matrix::Dense<ValueType>;
     std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
         std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     const auto executor_string = argc >= 2 ? argv[1] : "reference";
     exec map{
              {"omp", [] { return gko::OmpExecutor::create(); }},
              {"cuda",
               [] {
                    return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                           true);
               }},
              {"hip",
               [] {
                    return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
               }},
              { "dpcpp",
                   return gko::DpcppExecutor::create(0,
                                                            gko::OmpExecutor::create());
              {"reference", [] { return gko::ReferenceExecutor::create(); }}};
     const auto exec = exec_map.at(executor_string)(); // throws if not valid
std::string input_mtx = "data/A.mtx";
```

```
if (argc == 3) {
    input_mtx = std::string(argv[2]);
auto storage_logger = std::make_shared<loggers::StorageLogger>();
exec->add_logger(storage_logger);
auto A = gko::share(gko::read<mtx>(std::ifstream(input_mtx), exec));
exec->remove_logger(storage_logger);
const auto max_iters = A->get_size()[0];
auto b = utils::create_vector<ValueType>(exec, A->get_size()[0], 1.0);
auto x = utils::create_vector<ValueType>(exec, A->get_size()[0], 0.0);
auto solver_factory =
    solver::build()
        .with criteria(
             gko::stop::ResidualNorm<ValueType>::build()
                 .with_reduction_factor(reduction_factor)
                  .on(exec),
             gko::stop::Iteration::build().with_max_iters(max_iters).on(
                 exec))
         .with_preconditioner(preconditioner::create(exec))
         .on(exec);
std::ofstream output_file(of_name);
    auto x_clone = gko::clone(x);
solver_factory->generate(A)->apply(b, x_clone);
    exec->synchronize();
    auto x_clone = gko::clone(x);
    exec->synchronize();
    auto q_tic = std::chrono::steady_clock::now();
    auto generated solver = solver factory->generate(A):
    exec->synchronize();
    auto g_tac = std::chrono::steady_clock::now();
    auto generate_time =
    std::chrono::duration_cast<std::chrono::nanoseconds>(g_tac - g_tic);
output_file « "Generate time (ns): " « generate_time.count()
                 « std::endl;
    exec->synchronize();
    auto a_tic = std::chrono::steady_clock::now();
    generated_solver->apply(b, x_clone);
    exec->synchronize();
    auto a_tac = std::chrono::steady_clock::now();
    auto apply_time =
        std::chrono::duration_cast<std::chrono::nanoseconds>(a_tac - a_tic);
    output_file « "Apply time (ns): " « apply_time.count() « std::endl;
    auto residual =
        utils::compute_residual_norm(A.get(), b.get(), x_clone.get());
    output_file « "Residual_norm: " « residual « std::endl;
    auto gen_logger = std::make_shared<loggers::OperationLogger>();
    exec->add_logger(gen_logger);
    auto generated_solver = solver_factory->generate(A);
    exec->remove_logger(gen_logger);
output_file « "Generate operations times (ns):" « std::endl;
    gen_logger->write_data(output_file);
auto apply_logger = std::make_shared<loggers::OperationLogger>();
    exec->add_logger(apply_logger);
    auto res_logger = std::make_shared<loggers::ResidualLogger<ValueType»(</pre>
    A.get(), b.get());
generated_solver->add_logger(res_logger);
    generated_solver->apply(b, x);
    exec->remove_logger(apply_logger);
output_file « "Apply operations times (ns):" « std::endl;
    apply_logger->write_data(output_file);
    res_logger->write_data(output_file);
std::cout « "Solution, first ten entries: \n";
print vector(x.get());
std::cout « "The performance and residual data can be found in " « of_name
           « std::endl;
```

The poisson-solver program

The poisson solver example..

This example depends on simple-solver.

Introduction

This example solves a 1D Poisson equation:

$$\begin{aligned} u:[0,1] &\to R \\ u'' &= f \\ u(0) &= u0 \\ u(1) &= u1 \end{aligned}$$

using a finite difference method on an equidistant grid with ${\tt K}$ discretization points (${\tt K}$ can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

For an equidistant grid with K "inner" discretization points x1,...,xk,and step size h=1/(K+1), the formula produces a system of linear equations

$$2u_1 - u_2 = -f_1h^2 + u0$$

- $u(k-1) + 2u_k - u(k+1) = -f_kh^2, k = 2, ..., K-1$
- $u(K-1) + 2u_K = -f_Kh^2 + u1$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f'is set to 'f(x) = 6x' (making the solution ' $u(x) = x^3$ '), but that can be changed in the main function.

The intention of the example is to show how Ginkgo can be used to build an application solving a real-world problem, which includes a solution of a large, sparse linear system as a component.

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
#include <vector>
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    const auto discretization_points = matrix->get_size()[0];
    auto row_ptrs = matrix->get_row_ptrs();
    auto col_idxs = matrix->get_col_idxs();
    auto values = matrix->get_values();
    int pos = 0;
    const ValueType coefs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (int i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
            values[pos] = coefs[ofs + 1];
            col_idxs[pos] = i + ofs;
            ++pos;
        }
    }
    row_ptrs[i + 1] = pos;
}
</pre>
```

Generates the RHS vector given f and the boundary conditions.

Prints the solution u.

Computes the 1-norm of the error given the computed ${\tt u}$ and the correct solution function ${\tt correct_u}.$

```
template <typename Closure, typename ValueType>
gko::remove_complex<ValueType> calculate_error(
   int discretization_points, const gko::matrix::Dense<ValueType>* u,
   Closure correct_u)
{
   const ValueType h = 1.0 / static_cast<ValueType> (discretization_points + 1);
   gko::remove_complex<ValueType> error = 0.0;
   for (int i = 0; i < discretization_points; ++i) {
      using std::abs;
      const auto xi = static_cast<ValueType>(i + 1) * h;
      error +=
      abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
   }
   return error;
}
int main(int argc, char* argv[])
{
```

```
Some shortcuts
using ValueType = double;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
Print version information
std::exit(-1);
}
Get number of discretization points
const unsigned int discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
Get the executor string
const auto executor_string = argc >= 2 ? argv[1] : "reference";
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        {"omp",
{"cuda",
               [] { return gko::OmpExecutor::create(); }},
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                               true);
         }},
        {"hip",
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                              true);
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
executor used by the application
const auto app_exec = exec->get_master();
Set up the problem: define the exact solution, the right hand side and the Dirichlet boundary condition.
auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType(6) * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
initialize matrix and vectors
auto matrix = mtx::create(app_exec, gko::dim<2>(discretization_points),
                          3 * discretization_points - 2);
generate_stencil_matrix(matrix.get());
auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
Generate solver and solve the system
ca::build()
    .with_criteria(gko::stop::Iteration::build()
                        .with_max_iters(discretization_points)
```

gko::stop::ResidualNorm<ValueType>::build()
 .with_reduction_factor(reduction_factor)

->generate(clone(exec, matrix)) // copy the matrix to the executor

.on(exec))
.with_preconditioner(bj::build().on(exec))

.on(exec)

This is the expected output:

```
Solve complete. The average relative error is 2.52236e-11
```

Comments about programming and debugging

```
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
   Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
#include <vector>
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
    const auto discretization points = matrix->get size()[0];
    auto row_ptrs = matrix->get_row_ptrs();
    auto col_idxs = matrix->get_col_idxs();
    auto values = matrix->get_values();
    int pos = 0;
    const ValueType coefs[] = \{-1, 2, -1\};
    row_ptrs[0] = pos;
for (int i = 0; i < discretization_points; ++i) {</pre>
        for (auto ofs : {-1, 0, 1}) {
    if (0 <= i + ofs && i + ofs < discretization_points) {</pre>
                values[pos] = coefs[ofs + 1];
                col_idxs[pos] = i + ofs;
                ++pos:
            }
        }
```

```
row_ptrs[i + 1] = pos;
template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                  gko::matrix::Dense<ValueType>* rhs)
    const auto discretization_points = rhs->get_size()[0];
    auto values = rhs->get_values();
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
for (gko::size_type i = 0; i < discretization_points; ++i) {</pre>
        const auto xi = static_cast<ValueType>(i + 1) * h;
        values[i] = -f(xi) * h * h;
    values[0] += u0;
    values[discretization_points - 1] += u1;
template <typename Closure, typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                    const gko::matrix::Dense<ValueType>* u)
    std::cout « u0 « ' \n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout « u->get_const_values()[i] « '\n';
    std::cout « u1 « std::endl;
template <typename Closure, typename ValueType>
gko::remove_complex<ValueType> calculate_error(
    int discretization_points, const gko::matrix::Dense<ValueType>* u,
    Closure correct u)
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (int i = 0; i < discretization_points; ++i) {</pre>
        using std::abs;
        const auto xi = static_cast<ValueType>(i + 1) * h;
        error +=
           abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    return error;
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    if (argc == 2 && (std::string(argv[1]) == "--help"))
        std::exit(-1);
    const unsigned int discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
    exec map{
            {"omp", [] { return gko::OmpExecutor::create(); }},
            {"cuda",
             [] {
                 return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                   true);
             }},
            {"hip",
             [] {
                 return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
            {"dpcpp",
             [] {
                 return gko::DpcppExecutor::create(0,
                                                    gko::OmpExecutor::create());
            {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
    auto correct_u = [](ValueType x) { return x * x * x; };
    auto f = [](ValueType x) { return ValueType(6) * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    auto matrix = mtx::create(app_exec, gko::dim<2>(discretization_points),
                               3 * discretization_points - 2);
    generate_stencil_matrix(matrix.get());
```

```
auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
    auto fis = vec::create(app_exec, gko::dim<2>(discretization_points, 1)
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
}
     const gko::remove_complex<ValueType> reduction_factor = 1e-7;
     cg::build()
          .with_criteria(gko::stop::Iteration::build()
                                     .with_max_iters(discretization_points)
                               .on(exec),
gko::stop::ResidualNorm<ValueType>::build()
                                    .with_reduction_factor(reduction_factor)
                                     .on(exec))
           .with_preconditioner(bj::build().on(exec))
           .on(exec)
           ->generate(clone(exec, matrix)) // copy the matrix to the executor
     ->apply(rhs, u);
std::cout « "Solve complete.\nThe average relative error is "
                  « calculate_error(discretization_points, u.get(), correct_u) /
    static_cast<gko::remove_complex<ValueType>>(
                                 discretization_points)
                   « std::endl;
}
```

The preconditioned-solver program

The preconditioned solver example..

This example depends on simple-solver.

Introduction

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
Figure out where to run the code
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1);
Figure out where to run the code
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
         return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                     true);
          }},
```

```
{"hip",
          [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                      true);
          }},
          { "dpcpp",
           [] {
               return gko::DpcppExecutor::create(0,
                                                        gko::OmpExecutor::create());
          }},
{"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
Create solver factory
auto solver_gen
     cg::build()
          .with_criteria(
              gko::stop::Iteration::build().with_max_iters(20u).on(exec),
gko::stop::ResidualNorm<ValueType>::build()
                   .with_reduction_factor(reduction_factor)
                   .on(exec))
Add preconditioner, these 2 lines are the only difference from the simple solver example
.with_preconditioner(bj::build().with_max_block_size(8u).on(exec))
.on(exec);
Create solver
auto solver = solver gen->generate(A);
Solve system
solver->apply(b, x);
Print solution
std::cout « "Solution (x):\n";
write(std::cout, x);
Calculate residual
     auto one = gko::initialize<vec>({1.0}, exec);
     auto neg_one = gko::initialize<vec>({-1.0}, exec);
     auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
     write(std::cout, res);
```

This is the expected output:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
4.82005e-08
```

Comments about programming and debugging

```
*********GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
    Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
   Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    std::cout « gko::version_info::get() « std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
        std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
        exec_map{
             {"omp",
                     [] { return gko::OmpExecutor::create(); }},
             {"cuda",
                  return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
              }},
             {"hip",
                  return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
              }},
             { "dpcpp",
              [] {
                  return gko::DpcppExecutor::create(0,
                                                       gko::OmpExecutor::create());
             {"reference", [] { return qko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    const RealValueType reduction_factor{1e-7};
    auto solver gen
        cg::build()
                 gko::stop::Iteration::build().with_max_iters(20u).on(exec),
```

The preconditioner-export program

The preconditioner export example..

This example depends on simple-solver.

Introduction

About the example

This example shows how to explicitly generate and store preconditioners for a given matrix. It can also be used to inspect and debug the preconditioner generation.

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <functional>
#include <iostream>
#include <map>
#include <memory>
#include <string>
const std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    executors{{"reference", [] { return gko::ReferenceExecutor::create(); }},
                {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
                 [] {
                      return gko::CudaExecutor::create(
                         0, gko::ReferenceExecutor::create());
                      return gko::HipExecutor::create(
                          0, gko::ReferenceExecutor::create());
                 }},
                {"dpcpp", [] {
    return gko::DpcppExecutor::create(
                         0, gko::ReferenceExecutor::create());
void output(gko::ptr_param<const gko::WritableToMatrixData<double, int>> mtx,
             std::string name)
     std::ofstream stream{name};
     std::cerr « "Writing " « name « std::endl;
    gko::write(stream, mtx, gko::layout_type::coordinate);
template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
```

```
decltype(fun()) result;
    try {
        result = fun();
    } catch (const gko::Error& err) {
  std::cerr « "Error: " « err.what() « '\n';
         std::exit(-1);
    return result;
int main(int argc, char* argv[])
print usage message
 f (argc < 2 || executors.find(argv[1]) == executors.end()) {</pre>
    std::cerr « "Usage: executable"
               « " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
                  "[<jacobi|ilu|parilu|parilut|ilu-isai|parilu-isai|parilut-"
                  "isai] [conditioner args>]\n";
    std::cerr « "Jacobi parameters: [<max-block-size>] [<accuracy>] "
                  "[<storage-optimization:auto|0|1|2>]\n";
    std::cerr « "ParILU parameters: [<iteration-count>]\n";
    std::cerr
        « "ParILUT parameters: [<iteration-count>] [<fill-in-limit>]\n";
    std::cerr « "ILU-ISAI parameters: [<sparsity-power>]\n";
    std::cerr « "ParILU-ISAI parameters: [<iteration-count>]
                  "[<sparsity-power>]\n";
    std::cerr « "ParILUT-ISAI parameters: [<iteration-count>] "
                   "[<fill-in-limit>] [<sparsity-power>]\n";
    return -1;
generate executor based on first argument
auto exec = try_generate([&] { return executors.at(argv[1])(); });
set matrix and preconditioner name with default values
std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
std::string precond = argc < 4 ? "jacobi" : argv[3];
load matrix file into Csr format
auto mtx = gko::share(try_generate([&] {
    std::ifstream mtx_stream{matrix};
    if (!mtx stream) {
        throw GKO_STREAM_ERROR("Unable to open matrix file");
    std::cerr « "Reading " « matrix « std::endl;
return gko::read<gko::matrix::Csr<»(mtx_stream, exec);</pre>
}));
concatenate remaining arguments for filename
std::string output_suffix;
for (auto i = 4; i < argc; ++i) {</pre>
    output_suffix = output_suffix + "-" + argv[i];
handle different preconditioners
if (precond == "jacobi") {
jacobi: max_block_size, accuracy, storage_optimization
    auto factory = gko::preconditioner::Jacobi<>::build().on(exec);
if (argc >= 5) {
        factory->get_parameters().max_block_size = std::stoi(argv[4]);
        factory->get_parameters().accuracy = std::stod(argv[5]);
    if (argc >= 7) {
         factory->get_parameters().storage_optimization =
    std::string{argv[6]} == "auto"
                 ? gko::precision_reduction::autodetect()
                  : gko::precision_reduction(0, std::stoi(argv[6]));
auto jacobi = try_generate([&] { return factory->generate(mtx); });
  output(jacobi, matrix + ".jacobi" + output_suffix);
} else if (precond == "ilu") {
ilu: no parameters
    auto ilu = gko::as<gko::Composition<»(try_generate([&] {</pre>
        return gko::factorization::Ilu<>::build().on(exec)->generate(mtx);
    output(gko::as<gko::matrix::Csr<>>(ilu->get_operators()[0]),
```

```
matrix + ".ilu-l");
    output(gko::as<gko::matrix::Csr<>>(ilu->get_operators()[1]),
          matrix + ".ilu-u");
} else if (precond == "parilu") {
parilu: iterations
    auto factory = gko::factorization::ParIlu<>::build().on(exec);
    if (argc >= 5) {
        factory->get_parameters().iterations = std::stoi(argv[4]);
    auto ilu = gko::as<gko::Composition<>>(
    output(gko::as<gko::matrix::Csr<>>(ilu->get_operators()[1]),
          matrix + ".parilu" + output_suffix + "-u");
} else if (precond == "parilut") {
parilut: iterations, fill-in limit
    auto factory = gko::factorization::ParIlut<>::build().on(exec);
if (argc >= 5) {
       factory->get_parameters().iterations = std::stoi(argv[4]);
    if (argc >= 6) {
       factory->get_parameters().fill_in_limit = std::stod(argv[5]);
    auto ilut = gko::as<gko::Composition<»(
   try_generate([&] { return factory->generate(mtx); }));
output(gko::as<gko::matrix::Csr<>>(ilut->get_operators()[0]),
          matrix + ".parilut" + output_suffix + "-1");
ilu-isai: sparsity power
    auto fact factory =
       gko::share(gko::factorization::Ilu<>::build().on(exec));
    int sparsity_power = 1;
    if (argc >= 5) {
       sparsity_power = std::stoi(argv[4]);
    auto factory =
       gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
                                gko::preconditioner::UpperIsai<>>::build()
            .with_factorization_factory(fact_factory)
            .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                      .with_sparsity_power(sparsity_power)
                                      .on(exec))
            .with_u_solver_factory(gko::preconditioner::UpperIsai<>::build()
                                      .with_sparsity_power(sparsity_power)
                                      .on(exec))
            .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
          matrix + ".ilu-isai" + output_suffix + "-u");
} else if (precond == "parilu-isai") {
parilu-isai: iterations, sparsity power
    auto fact_factory =
       gko::share(gko::factorization::ParIlu<>::build().on(exec));
    int sparsity_power = 1;
if (argc >= 5) {
        fact_factory->get_parameters().iterations = std::stoi(argv[4]);
    if (argc >= 6) {
       sparsity_power = std::stoi(argv[5]);
    auto factory =
       gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
                                gko::preconditioner::UpperIsai<>>::build()
            .with_factorization_factory(fact_factory)
            .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                      . \verb|with_sparsity_power(sparsity_power)|\\
                                      .on(exec))
            .with_u_solver_factory(gko::preconditioner::UpperIsai<>::build()
                                      .with_sparsity_power(sparsity_power)
                                      .on(exec))
            .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
          matrix + ".parilu-isai" + output_suffix + "-1");
```

```
output(ilu_isai->get_u_solver()->get_approximate_inverse(),
matrix + ".parilu-isai" + output_suffix + "-u");
} else if (precond == "parilut-isai") {
parilut-isai: iterations, fill-in limit, sparsity power
        auto fact_factory
            gko::share(gko::factorization::ParIlut<>::build().on(exec));
        int sparsity_power = 1;
        if (argc >= 5) {
            fact factory->get parameters().iterations = std::stoi(argv[4]);
            fact_factory->get_parameters().fill_in_limit = std::stod(argv[5]);
        if (argc >= 7) {
            sparsity_power = std::stoi(argv[6]);
            gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
                                      gko::preconditioner::UpperIsai<>>::build()
                .with_factorization_factory(fact_factory)
                .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                            .with_sparsity_power(sparsity_power)
                                            .on(exec))
                .with_u_solver_factory(gko::preconditioner::UpperIsai<>::build()
                                            .with_sparsity_power(sparsity_power)
                                            .on(exec))
                .on(exec);
        auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
        output(ilu_isai->get_l_solver()->get_approximate_inverse(),
matrix + ".parilut-isai" + output_suffix + "-l");
```

This is the expected output:

When specifying an executor:

```
Reading data/A.mtx
Writing data/A.mtx.jacobi
```

Comments about programming and debugging

```
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <functional>
#include <iostream>
#include <map>
#include <memory>
#include <string>
{"omp", [] { return gko::OmpExecutor::create(); }}, {"cuda",
                [] {
                    return gko::CudaExecutor::create(
                         0, gko::ReferenceExecutor::create());
                }},
               {"hip",
                [] {
                    return gko::HipExecutor::create(
                         0, gko::ReferenceExecutor::create());
                }},
               {"dpcpp", [] {
                    return gko::DpcppExecutor::create(
                        0, gko::ReferenceExecutor::create());
                }}};
void output(gko::ptr_param<const gko::WritableToMatrixData<double, int>> mtx,
            std::string name)
    std::ofstream stream{name};
std::cerr « "Writing " « name « std::endl;
    gko::write(stream, mtx, gko::layout_type::coordinate);
template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
    decltype(fun()) result;
        result = fun();
    } catch (const gko::Error& err) {
   std::cerr « "Error: " « err.what() « '\n';
        std::exit(-1);
    return result;
int main(int argc, char* argv[])
    if (argc < 2 || executors.find(argv[1]) == executors.end()) {
   std::cerr « "Usage: executable"</pre>
                   « " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
                       "[<jacobi|ilu|parilu|parilut|ilu-isai|parilu-isai|parilut-"
                       "isai] [conditioner args>]\n";
         std::cerr « "Jacobi parameters: [<max-block-size>] [<accuracy>] "
                      "[<storage-optimization:auto|0|1|2>]\n"
         std::cerr « "ParILU parameters: [<iteration-count>]\n";
         std::cerr
             « "ParILUT parameters: [<iteration-count>] [<fill-in-limit>]\n";
         std::cerr « "ILU-ISAI parameters: [<sparsity-power>]\n";
         std::cerr « "ParILU-ISAI parameters: [<iteration-count>] "
                      "[<sparsity-power>]\n";
        return -1:
    auto exec = try_generate([&] { return executors.at(argv[1])(); });
    std::string matrix = argc < 3 ? "data/A.mtx": argv[2]; std::string precond = argc < 4 ? "jacobi": argv[3];
    auto mtx = gko::share(try_generate([&] {
        std::ifstream mtx_stream{matrix};
         if (!mtx_stream) {
             throw GKO_STREAM_ERROR("Unable to open matrix file");
         std::cerr « "Reading " « matrix « std::endl;
        return gko::read<gko::matrix::Csr<»(mtx_stream, exec);</pre>
    })):
    std::string output suffix;
```

```
for (auto i = 4; i < argc; ++i) {</pre>
   output_suffix = output_suffix + "-" + argv[i];
if (precond == "jacobi") {
   auto factory = gko::preconditioner::Jacobi<>::build().on(exec);
if (argc >= 5) {
       factory->get_parameters().max_block_size = std::stoi(argv[4]);
       factory->get_parameters().accuracy = std::stod(argv[5]);
   if (argc >= 7) {
       factory->get_parameters().storage_optimization =
           std::string{argv[6]} == "auto"
              ? gko::precision_reduction::autodetect()
               : gko::precision_reduction(0, std::stoi(argv[6]));
   auto jacobi = try_generate([&] { return factory->generate(mtx); });
output(jacobi, matrix + ".jacobi" + output_suffix);
} else if (precond == "ilu") {
   auto ilu = gko::as<gko::Composition<> (try_generate([&] {
       return gko::factorization::Ilu<>::build().on(exec)->generate(mtx);
   1)):
   } else if (precond == "parilu") {
   auto factory = gko::factorization::ParIlu<>::build().on(exec);
   if (argc >= 5) {
       factory->get parameters().iterations = std::stoi(argy[4]);
   auto ilu = gko::as<gko::Composition<»(</pre>
       try_generate([&] { return factory->generate(mtx); }));
   } else if (precond == "parilut") {
   auto factory = gko::factorization::ParIlut<>::build().on(exec);
   if (argc >= 5) {
       factory->get_parameters().iterations = std::stoi(argv[4]);
   if (argc >= 6) {
       factory->get_parameters().fill_in_limit = std::stod(argv[5]);
   auto ilut = gko::as<gko::Composition<»(
   matrix + ".parilut" + output_suffix + "-u");
} else if (precond == "ilu-isai") {
   auto fact_factory =
       gko::share(gko::factorization::Ilu<>::build().on(exec));
   int sparsity_power = 1;
   if (argc >= 5) {
       sparsity_power = std::stoi(argv[4]);
   auto factory =
       gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
                              qko::preconditioner::UpperIsai<>>::build()
           .with_factorization_factory(fact_factory)
           .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                    .with_sparsity_power(sparsity_power)
                                    .on(exec))
           .with_u_solver_factory(gko::preconditioner::UpperIsai<>::build()
                                    .with_sparsity_power(sparsity_power)
                                    .on(exec))
           .on(exec);
   auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
   output(ilu_isai->get_u_solver()->get_approximate_inverse(),
    matrix + ".ilu-isai" + output_suffix + "-u");
} else if (precond == "parilu-isai") {
   auto fact_factory =
      gko::share(gko::factorization::ParIlu<>::build().on(exec));
   int sparsity_power = 1;
   if (argc >= 5) {
       fact_factory->get_parameters().iterations = std::stoi(argv[4]);
   if (argc >= 6) {
       sparsity_power = std::stoi(argv[5]);
   auto factory =
       gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
```

```
gko::preconditioner::UpperIsai<>>::build()
             .with_factorization_factory(fact_factory)
             .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                            .with_sparsity_power(sparsity_power)
                                            .on(exec))
             .with_u_solver_factory(gko::preconditioner::UpperIsai<>::build()
                                           .with_sparsity_power(sparsity_power)
                                            .on(exec))
             .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver(),->get_approximate_inverse(), matrix + ".parilu-isai" + output_suffix + "-1");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
matrix + ".parilu-isai" + output_suffix + "-u");
} else if (precond == "parilut-isai") {
  auto fact_factory =
        gko::share(gko::factorization::ParIlut<>::build().on(exec));
    int sparsity_power = 1;
if (argc >= 5) {
        fact_factory->get_parameters().iterations = std::stoi(argv[4]);
    if (argc >= 6) {
         fact_factory->get_parameters().fill_in_limit = std::stod(argv[5]);
    if (argc >= 7) {
        sparsity_power = std::stoi(argv[6]);
    auto factory =
        gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
                                     gko::preconditioner::UpperIsai<>>::build()
             .with_factorization_factory(fact_factory)
             .with_l_solver_factory(gko::preconditioner::LowerIsai<>::build()
                                           .with_sparsity_power(sparsity_power)
                                            .on(exec))
             . with\_u\_solver\_factory (gko::preconditioner::UpperIsai<>::build()
                                            .with_sparsity_power(sparsity_power)
                                            .on(exec))
             .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
    matrix + ".parilut-isai" + output_suffix + "-l");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
            matrix + ".parilut-isai" + output_suffix + "-u");
```

The schroedinger-splitting program

The Schroedinger equation example..

This example depends on heat-equation.

Introduction

This example shows how to use the FFT and iFFT implementations in Ginkgo to solve the non-linear Schrödinger equation with a splitting method.

The non-linear Schrödinger equation (NLS) is given by

$$i\partial_t \theta = -\delta \theta + |\theta|^2 \theta$$

Here θ is the wave function of a single particle in two dimensions. Its magnitude $|\theta|^2$ describes the probability distribution of the particle's position.

This equation can be split in to its linear (1) and non-linear (2) part

$$(1) \quad i\partial_t \theta = -\delta \theta$$

(2)
$$i\partial_t \theta = |\theta|^2 \theta$$

For both of these equations, we can compute exact solutions, assuming periodic boundary conditions and using the Fourier series expansion for (1) and using the fact that $|\theta|^2$ is constant in (2):

$$(\hat{1})$$
 $\partial_t \hat{\theta}_k = -i|k|^2 \theta$

$$\begin{aligned} &(\hat{1}) & \partial_t \hat{\theta}_k = -i|k|^2 \theta \\ &(2') & \partial_t |\theta|^2 = i|\theta|^2 (\theta - \theta) = 0 \end{aligned}$$

The exact solutions are then given by

$$(\hat{1}) \quad \hat{\theta}(t) = e^{-i|k|^2 t} \hat{\theta}(0)$$

$$(2') \quad \theta(t) = e^{-i|\theta|^2 t} \theta(0)$$

These partial solutions can be used to approximate a solution to the full NLS by alternating between small time steps for (1) and (2).

For nicer visual results, we add another constant potential term V(x) \theta to the non-linear part, which turns it into the Gross-Pitaevskii equation.

About the example

The commented program

```
This example shows how to use the FFT and iFFT implementations in Ginkgo
 to solve the non-linear Schrödinger equation with a splitting method.
The non-linear Schrödinger equation (NLS) is given by
i \partial_t \theta = -\delta \theta + 1\theta |^2 \theta Here \theta is the wave function of a single particle in two dimensions.
Its magnitude |\theta|^2 describes the probability distribution of the
particle's position.
This equation can be split in to its linear (1) and non-linear (2) part
 \f{align*}{
(1) \quad i \partial_t \theta &= -\delta \theta\\
(2) \quad i \partial_t \theta &= |\theta|^2 \theta
For both of these equations, we can compute exact solutions, assuming periodic
boundary conditions and using the Fourier series expansion for (1) and using the fact that | \theta^2  is constant in (2):
 (\hat 1) \quad \quad \partial_t \hat\theta_k &= -i |k|^2 \neq 1
 \f}
The exact solutions are then given by
 \f{align*}{
           (\hat 1) \quad \hat\theta(t) &= e^{-i |k|^2 t} \hat{t} (0) 
           (2') \quad \text{(2')} \quad \text{(2')} \quad \text{(b)} \quad \text{(2')} \quad \text{(2')} \quad \text{(b)} \quad \text{(2')} \quad \text{(2')} \quad \text{(2')} \quad \text{(3')} \quad
These partial solutions can be used to approximate a solution to the full NLS
by alternating between small time steps for (1) and (2).
For nicer visual results, we add another constant potential term V\left(x\right) \theta
to the non-linear part, which turns it into the Gross-Pitaevskii equation.
 #include <ginkgo/ginkgo.hpp>
#include <algorithm>
 #include <chrono>
 #include <fstream>
 #include <iostream>
#include <utility>
 #include <opencv2/core.hpp>
 #include <opencv2/videoio.hpp>
This function implements a simple Ginkgo-themed clamped color mapping for values in the range [0,5].
void set_val(unsigned char* data, double value)
RGB values for the 6 colors used for values 0, 1, ..., 5 We will interpolate linearly between these values.
double col_r[] = {255, 221, 129, 201, 249, 255};
 double col_g[] = {255, 220, 130, 161, 158, 204};
double col_b[] = {255, 220, 133, 93, 24, 8};
value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
auto d = std::max(0.0, std::min(1.0, value - i));
OpenCV uses BGR instead of RGB by default, revert indices
          \label{eq:data_obj} \texttt{data[0]} = \texttt{static\_cast} < \texttt{unsigned char} > (\texttt{col\_b[i + 1]} \ * \ \texttt{d} \ + \ \texttt{col\_b[i]} \ * \ (\texttt{1 - d)});
}
Initialize video output with given dimension and FPS (frames per seconds)
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
           cv::Size videosize{n, n};
          auto output =
         std::make_pair(cv::VideoWriter{}), cv::Mat{videosize, CV_8UC3});
auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', '1');
output.first.open("nls.mp4", fourcc, fps, videosize);
          return output;
```

```
}
Write the current frame to video output using the above color mapping
for (int i = 0; i < n; i++) {</pre>
         auto row = output.second.ptr(i);
         for (int j = 0; j < n; j++)
             set_val(@row[3 * j], abs(data[i * n + j]));
    output.first.write(output.second);
int main(int argc, char* argv[])
    using vec = gko::matrix::Dense<std::complex<double>>;
using real_vec = gko::matrix::Dense<double>;
    using fft2 = gko::matrix::Fft2;
Problem parameters: simulation length
const auto t0 = 15.0;
scaling factor for non-linearity
const auto nonlinear_scale = 1.0;
scaling factor for potential
const auto potential_scale = 3.0;
Simulation parameters: time scaling factor
const auto time_scale = 0.25;
number of grid points in each dimension
const auto n = 256;
number of simulation steps per second
const auto steps_per_sec = 1000;
number of video frames per second
const auto fps = 25;
number of grid points
const auto n2 = n * n;
phase difference between neighboring grid points
const auto h = 2.0 * gko::pi < double > () / n; const auto h2 = h * h;
time step size for the simulation
const auto tau = 1.0 / steps_per_sec;
const auto idx = [&](int i, int j) { return i * n + j; };
create an OpenMP executor
auto exec = gko::OmpExecutor::create();
load initial state vector
std::ifstream initial_stream("data/gko_logo_2d.mtx");
std::ifstream potential_stream("data/gko_text_2d.mtx");
auto amplitude = gko::read<vec>(initial_stream, exec);
auto potential = gko::read<real_vec>(potential_stream, exec);
create vector for frequency space representation
auto frequency = vec::create(exec, amplitude->get_size());
create Fourier matrix
```

prepare video output

auto fft = fft2::create(exec, n, n);
auto ifft = fft->conj_transpose();

auto output = build_output(n, fps);

```
time stamp of the last output frame (sentinel value)
double last_t = -t0;
execute splitting method: time step in linear part, then non-linear part
for (double t = 0; t < t0; t += tau) {</pre>
if enough time has passed, output the next frame
if (t - last_t > 1.0 / fps) {
     last_t = t;
     std::cout « t « std::endl;
     output_timestep(output, n, amplitude->get_const_values());
time step in linear part
fft->apply(amplitude, frequency);
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        frequency->at(idx(i, j)) *=
            std::polar(1.0, -h2 * (i * i + j * j) * tau * time_scale);
scale by FFT*iFFT normalization factor
         frequency->at(idx(i, j)) \star= 1.0 / n2;
ifft->apply(frequency, amplitude);
time step in non-linear part
         for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
                   amplitude->at(idx(i, j)) *= std::polar(
                       1.0, -(nonlinear_scale *
                                     gko::squared_norm(amplitude->at(idx(i, j))) +
                                potential_scale * potential->at(idx(i, j))) *
                                   tau * time_scale);
         }
    }
```

The program will generate a video file named nls.mp4 and output the timestamp of each generated frame.

Comments about programming and debugging

```
************GINKGO LICENSE>*****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
are met:
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
   Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
```

```
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
This example shows how to use the FFT and iFFT implementations in Ginkgo
to solve the non-linear Schrödinger equation with a splitting method.
The non-linear Schrödinger equation (NLS) is given by
i \partial_t \theta = -\delta \theta + |\theta^2 \rangle \theta
Here @f$\theta@f$ is the wave function of a single particle in two dimensions.
Its magnitude @f$|\theta|^2@f$ describes the probability distribution of the
particle's position.
This equation can be split in to its linear (1) and non-linear (2) part
\f{align*}{
(1) \quad i \partial_t \theta &= -\delta \theta\\
(2) \quad i \partial_t \theta &= |\theta|^2 \theta
For both of these equations, we can compute exact solutions, assuming periodic
boundary conditions and using the Fourier series expansion for (1) and using the
fact that @f \theta |^2@f$ is constant in (2):
(\hat 1) \quad \quad \partial_t \hat\theta_k &= -i |k|^2 \theta \\ (2') \quad \partial_t |\theta|^2 &= i |\theta|^2 (\theta - \theta) = 0
The exact solutions are then given by
f{align*}{
(2') \quad \text{quad } \quad \text{theta(t)} \quad \text{e= e^{-i} } \quad \text{theta(0)}
These partial solutions can be used to approximate a solution to the full NLS
by alternating between small time steps for (1) and (2).
For nicer visual results, we add another constant potential term \text{V}\left(x\right) \theta
to the non-linear part, which turns it into the Gross-Pitaevskii equation.
#include <ginkgo/ginkgo.hpp>
#include <algorithm>
#include <chrono>
#include <fstream>
#include <iostream>
#include <utility>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
void set_val(unsigned char* data, double value)
    double col_r[] = {255, 221, 129, 201, 249, 255}; double col_g[] = {255, 220, 130, 161, 158, 204};
    double col_b[] = \{255, 220, 133, 93, 24, 8\};
    value = std::max(0.0, value);
    auto i = std::max(0, std::min(4, int(value)));
    auto d = std::max(0.0, std::min(1.0, value - i)); data[2] = static_cast<unsigned char>(col_r[i + 1] * d + col_r[i] * (1 - d));
    data[1] = static_cast<unsigned char>(col_g[i + 1] * d + col_g[i] * (1 - d));
    data[0] = static_cast<unsigned char>(col_b[i + 1] * d + col_b[i] * (1 - d));
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
    cv::Size videosize{n, n};
    auto output =
        std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', '1');
    output.first.open("nls.mp4", fourcc, fps, videosize);
    return output;
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                      const std::complex<double>* data)
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
for (int j = 0; j < n; j++) {
    set_val(&row[3 * j], abs(data[i * n + j]));</pre>
```

```
output.first.write(output.second);
int main(int argc, char* argv[])
     using vec = gko::matrix::Dense<std::complex<double>>;
     using real_vec = gko::matrix::Dense<double>;
     using fft2 = gko::matrix::Fft2;
     const auto t0 = 15.0;
     const auto nonlinear_scale = 1.0;
const auto potential_scale = 3.0;
     const auto time_scale = 0.25;
     const auto n = 256;
     const auto steps_per_sec = 1000;
    const auto fps = 25;

const auto n2 = n * n;

const auto h = 2.0 * gko::pi<double>() / n;

const auto h2 = h * h;
     const auto tau = 1.0 / steps_per_sec;
     const auto idx = [\&](int i, int j) { return i * n + j; };
     auto exec = gko::OmpExecutor::create();
     std::ifstream initial_stream("data/gko_logo_2d.mtx");
     std::ifstream potential_stream("data/gko_text_2d.mtx");
auto amplitude = gko::read<vec>(initial_stream, exec);
auto potential = gko::read<real_vec>(potential_stream, exec);
     auto frequency = vec::create(exec, amplitude->get_size());
     auto fft = fft2::create(exec, n, n);
auto ifft = fft->conj_transpose();
     auto output = build_output(n, fps);
     double last_t = -t0;
for (double t = 0; t < t0; t += tau) {</pre>
          if (t - last_t > 1.0 / fps) {
                last_t = t;
                std::cout « t « std::endl;
                output_timestep(output, n, amplitude->get_const_values());
          fft->apply(amplitude, frequency);
          for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
                     frequency->at(idx(i, j)) *=
    std::polar(1.0, -h2 * (i * i + j * j) * tau * time_scale);
frequency->at(idx(i, j)) *= 1.0 / n2;
          ifft->apply(frequency, amplitude);
          for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        amplitude->at(idx(i, j)) *= std::polar()
                          1.0, -(nonlinear_scale *
                                         gko::squared_norm(amplitude->at(idx(i, j))) +
                                    potential_scale * potential->at(idx(i, j)))
                                       tau * time_scale);
              }
        }
   }
```

Chapter 38

The simple-solver program

The simple solver example..

Introduction

This simple solver example should help you get started with Ginkgo. This example is meant for you to understand how Ginkgo works and how you can solve a simple linear system with Ginkgo. We encourage you to play with the code, change the parameters and see what is best suited for your purposes.

About the example

Each example has the following sections:

- 1. **Introduction:**This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
- 2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
- 3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
- 4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

Include files

This is the main ginkgo header file.

#include <ginkgo/ginkgo.hpp>

Add the fstream header to read from data from files.

#include <fstream>

Add the C++ iostream header to output information to the console.

#include <iostream

Add the STL map header for the executor selection

#include <map>

Add the string manipulation header to handle strings.

```
#include <string>
int main(int argc, char* argv[])
```

Use some shortcuts. In Ginkgo, vectors are seen as a gko::matrix::Dense with one column/one row. The advantage of this concept is that using multiple vectors is a now a natural extension of adding columns/rows are necessary.

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
```

The gko::matrix::Csr class is used here, but any other matrix class such as gko::matrix::Coo, gko::matrix::Hybrid, gko::matrix::Ell or gko::matrix::Sellp could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The gko::solver::Cg is used here, but any other solver class can also be used.

```
using cg = gko::solver::Cg<ValueType>;
```

Print the ginkgo version information.

```
std::cout « gko::version_info::get() « std::endl;
```

```
Print help on how to execute this example.
```

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
   std::exit(-1);
}
```

Where do you want to run your solver?

The gko::Executor class is one of the cornerstones of Ginkgo. Currently, we have support for an gko::OmpExecutor, which uses OpenMP multi-threading in most of its kernels, a gko::ReferenceExecutor, a single threaded specialization of the OpenMP executor and a gko::CudaExecutor which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/routines within Ginkgo should automatically work and run on the executor with any other changes.

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec map{
        { "omp",
               [] { return gko::OmpExecutor::create(); }},
        {"cuda",
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                              true);
         }},
        {"hip",
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                             true);
        { "dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                               gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Reading your data and transfer to the proper device.

Read the matrix, right hand side and the initial solution using the read function.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. gko::share and gko::give are the functions that you would need to use.

```
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
```

Creating the solver

Generate the gko::solver factory. Ginkgo uses the concept of Factories to build solvers with certain properties. Observe the Fluent interface used here. Here a cg solver is generated with a stopping criteria of maximum iterations of 20 and a residual norm reduction of 1e-7. You also observe that the stopping criteria(gko::stop) are also generated from factories using their build methods. You need to specify the executors which each of the object needs to be built on.

Generate the solver from the matrix. The solver factory built in the previous step takes a "matrix" (a gko::LinOp to be more general) as an input. In this case we provide it with a full matrix that we previously read, but as the solver only effectively uses the apply() method within the provided "matrix" object, you can effectively create a gko::LinOp class with your own apply implementation to accomplish more tasks. We will see an example of how this can be done in the custom-matrix-format example

```
auto solver = solver_gen->generate(A);
```

Finally, solve the system. The solver, being a gko::LinOp, can be applied to a right hand side, b to obtain the solution, x.

```
solver->apply(b, x);
```

Print the solution to the command line.

```
\label{eq:std:cout} $$ std::cout & "Solution (x):\n"; \\ write(std::cout, x); \\ \end{aligned}
```

To measure if your solution has actually converged, you can measure the error of the solution. one, neg_one are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, all you need to do is call the apply method, which in this case is an spmv and equivalent to the LAPACK z_spmv routine. Finally, you compute the euclidean 2-norm with the compute_norm2 function.

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

Results

The following is the expected result:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
2.10788e-15
```

Comments about programming and debugging

The plain program

```
******GINKGO LICENSE>****************
 Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
             Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
 3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
  IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
 TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
 PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
 DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
 #include <ginkgo/ginkgo.hpp>
  #include <fstream>
 #include <iostream>
 #include <map>
 #include <string>
 int main(int argc, char* argv[])
              using ValueType = double;
              using RealValueType = gko::remove_complex<ValueType>;
              using IndexType = int;
              using vec = gko::matrix::Dense<ValueType>;
              using real_vec = gko::matrix::Dense<RealValueType>;
              using mtx = gko::matrix::Csr<ValueType, IndexType>;
              using cg = gko::solver::Cg<ValueType>;
```

```
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
   std::cerr « "Usage: " « argv[0] « " [executor] " « std::endl;
     std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
     exec_map{
          {"omp", [] { return gko::OmpExecutor::create(); }},
          {"cuda",
           [] {
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                        true);
          } } ,
          {"hip",
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                                      true);
           }},
          {"dpcpp",
           [] {
               return gko::DpcppExecutor::create(0,
                                                         gko::OmpExecutor::create());
         }},
{"reference", [] { return qko::ReferenceExecutor::create(); }};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::readc>(std::ifstream("data/b.mtx"), exec);
auto x = gko::readvec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
         .with_criteria(
              gko::stop::Iteration::build().with_max_iters(20u).on(exec),
              gko::stop::ResidualNorm<ValueType>::build()
                   .with_reduction_factor(reduction_factor)
                   .on(exec))
          .on(exec);
auto solver = solver_gen->generate(A);
solver->apply(b, x);
std::cout « "Solution (x):\n";
write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

Chapter 39

The simple-solver-logging program

The simple solver with logging example..

This example depends on simple-solver, minimal-cuda-solver.

Introduction

About the example

The commented program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
namespace {
template <typename ValueType>
void print_vector(const std::string& name,
                      const gko::matrix::Dense<ValueType>* vec)
     std::cout « name « " = [" « std::endl;
     for (int i = 0; i < vec->get_size()[0]; ++i) {
    std::cout « " " « vec->at(i, 0) « std::endl;
     std::cout « "];" « std::endl;
   // namespace
int main(int argc, char* argv[])
Some shortcuts
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using real_vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
     std::exit(-1);
```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()»
    exec_map{
        { "omp",
               [] { return gko::OmpExecutor::create(); }},
        {"cuda".
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
        { "dpcpp",
             return gko::DpcppExecutor::create(0,
                                               gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read < vec > (std::ifstream("data/x0.mtx"), exec);
```

Let's declare a logger which prints to std::cout instead of printing to a file. We log all events except for all linop factory and polymorphic object events. Events masks are group of events which are provided for convenience.

```
std::shared_ptr<gko::log::Stream<ValueType» stream_logger =
    gko::log::Stream<ValueType>::create(
        gko::log::Logger::all_events_mask ^
             gko::log::Logger::linop_factory_events_mask ^
             gko::log::Logger::polymorphic_object_events_mask,
        std::cout);
```

Add stream_logger to the executor

exec->add_logger(stream_logger);

Add stream_logger only to the ResidualNorm criterion Factory Note that the logger will get automatically propagated to every criterion generated from this factory.

First we add facilities to only print to a file. It's possible to select events, using masks, e.g. only iterations mask: gko::log::Logger::iteration_complete_mask. See the documentation of Logger class for more information.

```
std::ofstream filestream("my_file.txt");
solver->add_logger(gko::log::Stream<ValueType>::create(
    gko::log::Logger::all_events_mask, filestream));
solver->add_logger(stream_logger);
```

This adds a simple logger that only reports convergence state at the end of the solver. Specifically it captures the last residual norm, the final number of iterations, and the converged or not converged status.

```
std::shared_ptr<gko::log::Convergence<ValueType» convergence_logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(convergence_logger);
```

Add another logger which puts all the data in an object, we can later retrieve this object in our code. Here we only have want Executor and criterion check completed events.

```
std::shared_ptr<gko::log::Record> record_logger =
   gko::log::Record::create(gko::log::Logger::executor_events_mask |
                            gko::log::Logger::iteration_complete_mask);
exec->add logger(record logger);
solver->add_logger(record_logger);
Solve system
solver->apply(b, x);
Print the residual of the last criterion check event (where convergence happened)
auto residual =
    record_logger->get().iteration_completed.back()->residual.get();
auto residual_d = gko::as<vec>(residual);
print_vector("Residual", residual_d);
Print solution
    std::cout « "Solution (x):\n";
    write(std::cout, x);
    std::cout « "Residual norm sqrt(r^T r):\n";
    write(std::cout, gko::as<vec>(convergence_logger->get_residual_norm()));
    std::cout « "Number of iterations "
    « convergence_logger->has_converged() « std::endl;
}
```

Results

This is the expected output:

```
[LOG] >> apply started on A LinOp[gko::solver::Cg<double>,0x2142d60] with b
LinOp[gko::matrix::Dense<double>,0x2142140] and x LinOp[gko::matrix::Dense<double>,0x2143450] [LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142280] with
         Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410] with
         Bytes[8]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0] with
         Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482f0] with
         Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0] with
         Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21486b0] with
         Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8] [LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148010] with
         Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60] with
         Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor, 0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0] with
         Bvtes[8]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40] with
         Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[1]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90] with
         Bytes[1]
[LOG] »> Operation[gko::solver::cg::initialize_operation<gko::matrix::Dense<double> const*&,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::initialize_operation<gko::matrix::Dense<double> const*&,
         gko::matrix::Dense<double>*, gko::matrix::Den
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>
const*, gko::matrix::Dense<double>*>,0x7ffd93d14aa0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
```

```
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14aa0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[2]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0] with
      Bytes[2]
[LOG] »- allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148e50] with
      Bytes[8]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0] with
      Bytes[8]
[LOG] »- Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14a20] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
gko::matrix::Dense<double>*>,0x7ffd93d14a20] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bvtes[152]
Bytes[152]
[LOG] »> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 0 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
      LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
      residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 0 with ID
       1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
       gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 0 with
      ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 \,
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149550] with
      Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149550] with
      Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149550] with
      Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730] with
      Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor, 0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149730] with
      Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149730] with
      Bytes[152]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
```

```
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
           gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 1 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
          \label{lin0p} \verb|[gko::matrix::Dense<double>, 0x2147b30]|, solution Lin0p| \verb|[gko::matrix::Dense<double>, 0x2143450]| and the control of the
          residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 1 with ID
          1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg Operation[gko::stop::residual_norm::residual_norm_operation < gko::matrix::Dense < double > const*\&, figure = 
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*&>,0x7ffd93d14b90] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 1 with
          ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor, 0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980] with
         Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149980] with
          Bytes [152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149980] with
          Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor.0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80] with
          Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b80] with
          Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b80] with
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2149550]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149550]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
 [LOG] \  \, \verb">> Operation[gko::solver::cg::step_1_operation < gko::matrix::Dense < double > \star \textit{,} \\
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
```

```
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes [152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
        »> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 2 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 2 with ID
        1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*.
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149290] with
        Bytes[152]
Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor, 0x21400d0] to
        Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x2149690] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149690] with
        Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80] [LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
         gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \, \verb">> Operation[gko::solver::cg::step_2_operation < gko::matrix::Dense < double > * \&, and the property of the pro
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
```

```
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
           Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
           Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c501 completed on
           Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> iteration 3 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
           LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
           residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 3 with ID
           1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
            gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
           gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping\_status>*&, gko::array<br/>bool>*, bool*&>,0x7ffd93d14b90] started on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
           gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
           gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \  \, \text{$>$} \  \  \, \text{check completed for stop::Criterion} \\ [gko::stop::ResidualNorm < double >, 0x2148db0] \  \  \, \text{at iteration 3 with 1} \\ [xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::xtop::
           ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
           Bytes[152]
[LOG] \gg copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
           Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0] with
           Bytes [152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149ae0] with
           Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[qko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149ae0] with
           Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149290]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149290]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
           Executor[qko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
           Executor[gko::ReferenceExecutor, 0x21400d0]
[\texttt{LOG}] \  \, \texttt{"NOP} \  \, \texttt{Operation} \\ [\texttt{gko}::solver::cg::step\_2\_operation \\ \texttt{gko}::matrix::Dense \\ \texttt{double} \\ \texttt{>} \star \&, \\ \texttt{(LOG)} \\ \texttt{(NOP)} \\ 
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
           Bytes[152]
```

```
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 4 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
      LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
      residual_norm LinOp[gko::LinOp const*,0]
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 4 with ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 4 with
      ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200] with
      Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
      Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
      Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310] with
      Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149310] with
      Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149310] with
      Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
      »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      \textbf{Executor}[\textbf{gko}:: \textbf{ReferenceExecutor}, \textbf{0x} \textbf{21400d0}] \text{ from Location}[\textbf{0x} \textbf{21480a0}] \text{ to Location}[\textbf{0x} \textbf{21482f0}] \text{ with } \textbf{0x} \textbf{11480a0}]
      Bytes[152]
[LOG] »> Operation[gko::matrix::dense::compute dot operation<gko::matrix::Dense<double> const*,
```

```
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> iteration 5 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 5 with ID
        1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 5 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 [LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
[LOG] \gg copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
        Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
        Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0] with
        Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cc0] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cc0] with
        Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        qko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
 [LOG] \  \, \hbox{$\tt w$-} \  \, {\tt Operation[gko::solver::cg::step\_2\_operation$< gko::matrix::Dense$< double>*\&, figure for the content of the c
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes [152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
```

```
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 6 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 6 with ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping\_status>*&, gko::array<br/>bool>*, bool*&>,0x7ffd93d14b90] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 6 with
        ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor, 0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
        Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0] with
        Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21494f0] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21494f0] with
        Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
 [LOG] \  \, \verb""" Operation[gko::solver::cg::step_1_operation < gko::matrix::Dense < double > \star \textit{""}, the properties of the properties o
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        \label{lem:gko::matrix::Dense<double>*>,0x7ffd93d14c50]} completed on \\ \texttt{Executor[gko::ReferenceExecutor,0x21400d0]}
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor, 0x21400d0]
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        \label{lem:gko::matrix::Dense<double>*>,0x7ffd93d14c50]} started on \\ \texttt{Executor[gko::ReferenceExecutor,0x21400d0]}
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 7 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
```

```
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 7 with ID
      1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
      gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*&, 0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 7 with
      ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 \,
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730] with
      Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149730] with
      Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149730] with
      Bytes [152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0] with
      Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21497d0] with
      Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21497d0] with
      Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2149450]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*
     gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 8 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
      LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
      residual_norm LinOp[gko::LinOp const*,0]
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 8 with ID
      1\ \mathrm{and}\ \mathrm{finalized}\ \mathrm{set}\ \mathrm{to}\ 1
```

```
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
        gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 8 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 [LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        \textbf{Executor}[\textbf{gko}:: \textbf{ReferenceExecutor}, \textbf{0x} \textbf{21400d0}] \text{ from Location}[\textbf{0x} \textbf{21480a0}] \text{ to Location}[\textbf{0x} \textbf{2149200}] \text{ with } \textbf{0x} \textbf{21490d0}]
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
        Bytes[152]
[LOG] »- allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x21492a0] with
        Bytes [152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21492a0] with
        Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2149730]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \, \text{"Normal of the construction of 
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
 [LOG] \  \, \text{``operation[gko::solver::cg::step\_2\_operation<gko::matrix::Dense<double>*&, } \\
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        \textbf{Executor}[\textbf{gko}:: \textbf{ReferenceExecutor}, \textbf{0x} \textbf{21400d0}] \text{ from Location}[\textbf{0x} \textbf{21480a0}] \text{ to Location}[\textbf{0x} \textbf{21482f0}] \text{ with } \textbf{0x} \textbf{11480a0}]
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 9 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 9 with ID
        1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&, gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
```

```
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 9 with
         ID 1 and finalized set to \hat{1}. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
        Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0] with
        Bvtes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21496c0] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21496c0] with
        Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
         qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
         Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
         gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
         gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed or
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor, 0x21400d0] to
        \texttt{Executor}[\textbf{gko}:: \textbf{ReferenceExecutor}, \textbf{0x21400d0}] \text{ from } \texttt{Location}[\textbf{0x21480a0}] \text{ to } \texttt{Location}[\textbf{0x21482f0}] \text{ with } \texttt{Location}[\textbf{0x21480a0}] \text{ to } \texttt{Location}[\textbf{0x21482f0}] \text{ with } \texttt{Location}[\textbf{0x21480a0}] \text{ to } \texttt{Location}[\textbf{0x21482f0}] \text{ with } \texttt{L
        Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 10 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
         LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
         residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 10 with
         ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
         gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
         gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
         gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        qko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
```

```
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 10 with
      ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450] with
      Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
      Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
      Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760] with
      Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149760] with
      Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149760] with
      Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      \label{lem:gko::matrix::Dense<double>*>,0x7ffd93d14c50]} started on \\ \texttt{Executor[gko::ReferenceExecutor,0x21400d0]}
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 11 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
      LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
      residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 11 with
      {\tt ID}\ 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
      Executor[qko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 11 with ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
```

```
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149860] with
        Bytes [152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149860] with
        Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149860] with
        Bytes [152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900] with
        Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x2149900] with
        Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[qko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149900] with
        Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> *>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[\texttt{LOG}] \  \, \texttt{"NOP} \  \, \texttt{Operation} \\ [\texttt{gko::solver::cg::step\_2\_operation} \\ \texttt{`gko::matrix::Dense} \\ \texttt{`double} \\ \texttt{`*\&, operation} \\ \texttt{`matrix::Dense} \\ \texttt{`double} \\ \texttt{`*\&, operation} \\ \texttt{`gko::matrix::Dense} \\ \texttt{`double} \\ \texttt{``homographical properation} 
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
         gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
qko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
[LOG] »> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 12 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 12 with
        ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 12 with
        ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
```

```
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21499a0] with
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21499a0] with
          Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21493d0] with
          Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21493d0] with
          Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor.0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x21493d0] with
          Bytes[152]
[LOG] » free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2149860]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149860]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
           Executor[gko::ReferenceExecutor,0x21400d0]
 [LOG] \  \, \text{``operation[gko::solver::cg::step\_1\_operation<gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation<gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step\_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step_1\_operation=gko::matrix::Dense<double>*, operation[gko::solver::cg::step_1\_operation=gko::step_1\_operation=gko::matrix::dense<double>*, operation[gko::solver::cg::step_1\_operation=gko::step_1\_operation=gko::matrix::dense<double>*, operation=gko
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
           Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor.0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
           gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
           Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] \begin{tabular}{l} \verb|MOG| >> Operation[gko::matrix::dense::compute\_dot\_operation < gko::matrix::Dense < double > const \star, dense < double > const \lambda < double > const \la
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> iteration 13 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
           residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 13 with
           ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
           gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
          qko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[qko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
           gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
           gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 13 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 [LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152] [LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490] with
          Bytes [152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149490] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149490] with
```

```
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor, 0x21400d0] to
         Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x2149580] with
         Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
         Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149580] with
         Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x21493d0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21493d0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
         Executor[qko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
         gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
         Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
         Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
         Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
         Bytes[152]
[LOG] »> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 14 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
         LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
         residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 14 with
         ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
         gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
         gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping\_status>*&, gko::array<br/>bool>*, bool*&>,0x7ffd93d14b90] started on
         Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
         gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
         gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
         Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 14 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 [LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b50] with
         Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
         Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149b50] with
         Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
         Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149b50] with
         Bytes [152]
[LOG] \  \, \text{">"} \  \, \text{allocation started on Executor} \\ [gko::ReferenceExecutor,0x21400d0] \  \, \text{with Bytes} \\ [152] \  \, \text{(152)} \\ [152] \  \, \text{(1
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0] with
         Bytes[152]
```

```
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x2143e90] to Location[0x21499c0] with
          Bytes [152]
\texttt{[LOG] w> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to}\\
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21499c0] with
          Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490]
 [LOG] \  \, \text{"Normalization} = \text{"Comparison} = \text{"Com
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[qko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
          \label{lem:gko::matrix::Dense<double>*>,0x7ffd93d14b80]} completed on \\ \texttt{Executor[gko::ReferenceExecutor,0x21400d0]}
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes [152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 15 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
          LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
          residual_norm LinOp[gko::LinOp const*,0]
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 15 with
          ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg Operation[gko::stop::residual_norm::residual_norm_operation < gko::matrix::Dense < double > const*\&, figure = 
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping\_status>*&, gko::array<br/>bool>*, bool*&>,0x7ffd93d14b90] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 15 with
          ID 1 and finalized set to \hat{1}. It changed one RHS \hat{0}, stopped the iteration process \hat{0}
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70] with
          Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149a70] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[qko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149a70] with
          Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor, 0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340] with
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
```

```
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2149b50]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b50]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor, 0x21400d0]
 [LOG] \  \, \text{"None of the continuous of the c
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
        Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
        gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bvtes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
        Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
        Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 16 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
        LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
        residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 16 with
        ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
        gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
        gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
        Executor[gko::ReferenceExecutor,0x21400d0]
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 16 with
       ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 \,
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970] with
        Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149970] with
        Bytes[152]
[LOG] »> copy completed from Executor[qko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149970] with
        Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10] with
        Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b10] with
        Bytes [152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
        Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b10] with
        Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340]
```

```
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b801 started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
          gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
          Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
          gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
          Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> iteration 17 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
          LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
          residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 17 with
          ID 1 and finalized set to 1
[LOG] >> Operation(gko::matrix::dense::compute norm2 operation(gko::matrix::Dense(double) const*
          gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
           gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg Operation[gko::stop::residual_norm::residual_norm_operation < gko::matrix::Dense < double > const*\&, figure = 
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
          gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
          Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \, \text{"Normalization of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::Criterion} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm} \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNorm } \\ [gko::stop::ResidualNorm < double"] \  \, \text{one of Stop::ResidualNor
          ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780] with
          Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x2149780] with
          Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149780] with
          Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
          Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149890] with
          Bytes[152]
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
          Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149890] with
          Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
          gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
```

```
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
              gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
              gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
              gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
              gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
             Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
              gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
              Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
              qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
             Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \, \text{"None of the continuous of the co
             gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
              gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
             Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
             gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*,
              gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
             Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
             Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
              Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
[LOG] \  \  \, \text{"No Operation} \  \  \, \text{"Sense::compute\_dot\_operation} \  \  \, \text{"Matrix::Dense<double> const*} \  \  \, \text{"Months of the operation} \  \  \, \text{"Months of th
              qko::matrix::Dense<double> const*, qko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
             Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
              gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 18 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
             \label{linOp} \texttt{LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and the line of the li
             residual_norm LinOp[gko::LinOp const*,0]
[LOG] »> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 18 with
             ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
              gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
              gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
              gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
              gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \gg Operation[gko::stop::residual_norm::residual_norm_operation < gko::matrix::Dense < double > const*\&, figure = 
             gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
              Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 18 with
              ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0 \,
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620] with
             Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
              Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
[LOG] \gg copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
             Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
             Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0] with
             Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
             Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cf0] with
             Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
              Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cf0] with
              Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
              gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
              gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
             Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] \  \, \verb">> Operation[gko::solver::cg::step_1_operation < gko::matrix::Dense < double > \star \textit{,} \\
             gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
```

```
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation(gko::matrix::csr::spmv operation(gko::matrix::Csr<double, int> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
      gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor, 0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes [152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
      Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> *>,0x7ffd93d14c50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> iteration 19 completed with solver Linop[gko::solver::Cg<double>,0x2142d60] with residual Linop[gko::matrix::Dense<double>,0x2147b30], solution Linop[gko::matrix::Dense<double>,0x2143450] and
      residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 19 with
      ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] »> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
      gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&, gko::array<br/>gko::array<br/>bool>*, bool*&>,0x7ffd93d14b90] completed on
      Executor[gko::ReferenceExecutor,0x21400d0]
Bvtes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
      Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
      Bytes[152]
[LOG] »> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340] with
      Bytes[152]
[LOG] »> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
      Bytes[152]
[LOG] »> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
      Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
      Bytes[152]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2148e50]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148e50]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0] [LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0]
```

```
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2148010]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148010]
[LOG] >> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x21486b0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21486b0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0]
[LOG] >> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x21482f0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor.0x21400d0] at Location[0x21482f0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2142280]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142280]
[LOG] >> apply completed on A LinOp[gko::solver::Cg<double>,0x2142d60] with b
      LinOp[gko::matrix::Dense<double>,0x2142140] and x LinOp[gko::matrix::Dense<double>,0x2143450]
            copied was of size 98 FROM executor 0x21400d0 pointer 2143e90 TO executor 0x21400d0 pointer
      2149340
Residual = [
    8.1654e-19
    -1.51449e-17
    2.23854e-17
    -1.0842e-19
    6.09864e-20
    -1.92446e-18
    1.97867e-18
    -4.58075e-18
    -1.55854e-18
    -2.64274e-17
    4.20128e-17
    -8.71427e-18
    -2.62919e-18
    -5.49947e-17
    5.51893e-17
    -1.57022e-16
    -4.2034e-17
    -8.71951e-16
    1.37837e-15
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0 0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
[LOG] \gg allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0] with
      Bvtes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] »> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870] with
      Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500] with
      Bytes[8]
[LOG] >> Operation(gko::matrix::csr::advanced spmv operation<gko::matrix::Dense<double> const*,
      gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>
      const*, gko::matrix::Dense<double>*>,0x7ffd93d14e50] started on
      Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14e50] completed on
      Executor[gko::ReferenceExecutor, 0x21400d0]
[LOG] >> Operation(gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14f70] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
      gko::matrix::Dense<double>*>,0x7ffd93d14f70] completed on Executor[gko::ReferenceExecutor,0x21400d0]
Residual norm sqrt(r^T r):
```

```
%%MatrixMarket matrix array real general
2.10788e-15
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143e90]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143e90]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143590]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143590]
[LOG] »> free started on Executor[gko::ReferenceExecutor, 0x21400d0] at Location[0x2142b10]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142b10]
[LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143c30]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143c30] [LOG] »> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143790]
[LOG] »> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143790]
```

Comments about programming and debugging

The plain program

```
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer.
           Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
 3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from % \left( 1\right) =\left( 1\right) \left( 1\right
this software without specific prior written permission.
 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
 IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
 TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
  #include <ginkgo/ginkgo.hpp>
 #include <fstream>
 #include <iomanip>
 #include <iostream>
 #include <map>
 #include <string>
namespace {
 template <typename ValueType>
 void print_vector(const std::string& name,
                                                        const gko::matrix::Dense<ValueType>* vec)
             std::cout « name « " = [" « std::endl;
             for (int i = 0; i < vec->get_size()[0]; ++i) {
    std::cout « " " « vec->at(i, 0) « std::endl;
             std::cout « "];" « std::endl;
          // namespace
 int main(int argc, char* argv[])
              using ValueType = double;
             using RealValueType = gko::remove_complex<ValueType>;
             using IndexType = int;
             using vec = gko::matrix::Dense<ValueType>;
             using real_vec = gko::matrix::Dense<RealValueType>;
             using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

```
using cg = gko::solver::Cg<ValueType>;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr « "Usage: " « argv[0] « " [executor]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
{"cuda",
         [] {
             return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
                                               true);
         }},
        {"dpcpp",
             return gko::DpcppExecutor::create(0,
                                                 gko::OmpExecutor::create());
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
std::shared_ptr<gko::log::Stream<ValueType» stream_logger =
    gko::log::Stream<ValueType>::create(
        gko::log::Logger::all_events_mask ^
            gko::log::Logger::linop_factory_events_mask ^
            gko::log::Logger::polymorphic_object_events_mask,
        std::cout);
exec->add logger(stream logger);
const RealValueType reduction_factor{1e-7};
using ResidualCriterionFactory =
    gko::stop::ResidualNorm<ValueType>::Factory;
std::shared_ptr<ResidualCriterionFactory> residual_criterion =
    ResidualCriterionFactory::create()
        .with reduction factor(reduction factor)
        .on(exec);
residual_criterion->add_logger(stream_logger);
auto solver_gen =
    ca::build()
        .with_criteria(
            residual_criterion,
            gko::stop::Iteration::build().with_max_iters(20u).on(exec))
        .on(exec);
auto solver = solver_gen->generate(A);
std::ofstream filestream("my_file.txt");
solver->add_logger(gko::log::Stream<ValueType>::create(
    gko::log::Logger::all_events_mask, filestream));
solver->add_logger(stream_logger);
std::shared_ptr<gko::log::Convergence<ValueType» convergence_logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(convergence_logger);
std::shared_ptr<gko::log::Record> record_logger =
    gko::log::Record::create(gko::log::Logger::executor_events_mask |
                             gko::log::Logger::iteration_complete_mask);
exec->add_logger(record_logger);
solver->add_logger(record_logger);
solver->apply(b, x);
auto residual =
   record_logger->get().iteration_completed.back()->residual.get();
auto residual_d = gko::as<vec>(residual);
print_vector("Residual", residual_d);
std::cout « "Solution (x):\n";
write(std::cout, x);
std::cout « "Residual norm sqrt(r^T r):\n";
write(std::cout, gko::as<vec>(convergence_logger->get_residual_norm()));
std::cout « "Number of iterations "
« convergence_logger->has_converged() « std::endl;
```

}

Chapter 40

The three-pt-stencil-solver program

The 3-point stencil example..

This example depends on simple-solver, poisson-solver.

Introduction

This example solves a 1D Poisson equation:

$$u: [0,1] \to R$$

$$u'' = f$$

$$u(0) = u0$$

$$u(1) = u1$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

For an equidistant grid with K "inner" discretization points x1,...,xk,and step size h=1/(K+1), the formula produces a system of linear equations

$$2u_1 - u_2 = -f_1h^2 + u0$$

- $u_(k-1) + 2u_k - u_(k+1) = -f_kh^2, k = 2, ..., K-1$
- $u_(K-1) + 2u_K = -f_Kh^2 + u1$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f'is set to 'f(x) = 6x' (making the solution ' $u(x) = x^3$ '), but that can be changed in the main function.

The intention of the example is to show how Ginkgo can be integrated into existing software - the <code>generate</code>—<code>stencil_matrix</code>, <code>generate_rhs</code>, <code>print_solution</code>, <code>compute_error</code> and <code>main</code> function do not reference Ginkgo at all (i.e. they could have been there before the application developer decided to use Ginkgo, and the only part where Ginkgo is introduced is inside the <code>solve_system</code> function.

About the example

The commented program

This example solves a 1D Poisson equation:

```
u: [0, 1] -> R
u" = f
      u(0) = u0
     u(1) = u1
using a finite difference method on an equidistant grid with 'K' discretization
points ('K' can be controlled with a command line parameter). The discretization
is done via the second order Taylor polynomial:
u(x + h) = u(x) - u'(x)h + 1/2 u'(x)h^2 + 0(h^3)

u(x - h) = u(x) + u'(x)h + 1/2 u''(x)h^2 + 0(h^3)
-u(x - h) + 2u(x) + -u(x + h) = -f(x)h^2 + O(h^3)
For an equidistant grid with K "inner" discretization points x1, ..., xk, and
step size h=1 / (K + 1), the formula produces a system of linear equations 2u_1 - u_2 = -f_1 h^2 + u0
-u_{k-1} + 2u_{k} - u_{k+1} = -f_{k} h^2,

-u_{k-1} + 2u_{k} = -f_{k} h^2 + u_{k+1}
                                                             k = 2, ..., K - 1
-u_(K-1) + 2u_K
-u_(x-1) + 2u_x = -1_x h^2 + u1 which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x) = 6x' (making the solution 'u(x) = x^3'), but
that can be changed in the 'main' function.
The intention of the example is to show how Ginkgo can be integrated into existing software - the 'generate_stencil_matrix', 'generate_rhs', 'print_solution', 'compute_error' and 'main' function do not reference Ginkgo at
all (i.e. they could have been there before the application developer decided to
use Ginkgo, and the only part where Ginkgo is introduced is inside the
 'solve_system' function.
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
#include <vector>
Creates a stencil matrix in CSR format for the given number of discretization points.
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType discretization_points,
                                         IndexType* row_ptrs, IndexType* col_idxs,
                                        ValueType* values)
      IndexType pos = 0;
      const ValueType coefs[] = \{-1, 2, -1\};
      row_ptrs[0] = pos;
      for (IndexType i = 0; i < discretization_points; ++i) {</pre>
           for (auto ofs: {-1, 0, 1}) {
   if (0 <= i + ofs && i + ofs < discretization_points) {
      values[pos] = coefs[ofs + 1];
      col_idxs[pos] = i + ofs;</pre>
                      ++pos:
                 }
           row_ptrs[i + 1] = pos;
      }
}
Generates the RHS vector given f and the boundary conditions.
template <typename Closure, typename ValueType, typename IndexType>
void generate_rhs(IndexType discretization_points, Closure f, ValueType u0,
                         ValueType u1, ValueType* rhs)
     const ValueType h = 1.0 / (discretization_points + 1);
for (IndexType i = 0; i < discretization_points; ++i) {
   const ValueType xi = ValueType(i + 1) * h;
   rhs[i] = -f(xi) * h * h;</pre>
      rhs[0] += u0;
      rhs[discretization_points - 1] += u1;
}
Prints the solution u.
template <typename ValueType, typename IndexType>
void print_solution(IndexType discretization_points, ValueType u0, ValueType u1,
                            const ValueType* u)
```

```
std::cout « u0 « ' \n';
for (IndexType i = 0; i < discretization_points; ++i) {
        std::cout « u[i] « '\n';
     std::cout « u1 « std::endl;
Computes the 1-norm of the error given the computed u and the correct solution function correct_u.
template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType discretization_points,
                                                       const ValueType* u,
                                                      Closure correct u)
{
     const ValueType h = 1.0 / (discretization_points + 1);
     gko::remove_complex<ValueType> error = 0.0;
     for (IndexType i = 0; i < discretization_points; ++i) {</pre>
         using std::abs;
         const ValueType xi = ValueType(i + 1) * h;
error += abs(u[i] - correct_u(xi)) / abs(correct_u(xi));
     return error;
template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                     IndexType discretization_points, IndexType* row_ptrs,
IndexType* col_idxs, ValueType* values, ValueType* rhs,
                     ValueType* u, gko::remove_complex<ValueType> reduction_factor)
Some shortcuts
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto& dp = discretization_points;
Figure out where to run the code
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
     exec_map{
          {"omp", [] { return gko::OmpExecutor::create(); }},
          {"cuda",
               return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                     true);
          11.
         {"hip",
           [] {
               return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
          }},
          { "dpcpp",
               return gko::DpcppExecutor::create(0,
                                                      gko::OmpExecutor::create());
          {"reference", [] { return gko::ReferenceExecutor::create(); }}};
executor where Ginkgo will perform the computation
const auto exec = exec_map.at(executor_string)(); // throws if not valid
executor where the application initialized the data
const auto app_exec = exec->get_master();
```

Tell Ginkgo to use the data in our application

Matrix: we have to set the executor of the matrix to the one where we want SpMVs to run (in this case exec). When creating array views, we have to specify the executor where the data is (in this case app_exec).

If the two do not match, Ginkgo will automatically create a copy of the data on exec (however, it will not copy the data back once it is done

· here this is not important since we are not modifying the matrix).

Solution: we have to be careful here - if the executors are different, once we compute the solution the array will not be automatically copied back to the original memory locations. Fortunately, whenever apply is called on a linear operator (e.g. matrix, solver) the arguments automatically get copied to the executor where the operator is, and copied back once the operation is completed. Thus, in this case, we can just define the solution on app_exec , and it will be automatically transferred to/from exec if needed

```
and it will be automatically transferred to/from exec if needed.
auto x = vec::create(app_exec, gko::dim<2>(dp, 1),
                      val_array::view(app_exec, dp, u), 1);
Generate solver
auto solver_gen =
    cq::build()
        .with_criteria(gko::stop::Iteration::build()
                            .with_max_iters(gko::size_type(dp))
                             .on(exec),
                        gko::stop::ResidualNorm<ValueType>::build()
                            .with_reduction_factor(reduction_factor)
                            .on(exec))
        .with_preconditioner(bj::build().on(exec))
auto solver = solver_gen->generate(gko::give(matrix));
Solve system
    solver->apply(b, x);
int main(int argc, char* argv[])
    using ValueType = double;
    using IndexType = int;
Print version information
std::cout « gko::version_info::get() « std::endl;
if (argc == 2 && std::string(argv[1]) == "--help") {
    std::cerr « "Usage: " « argv[0]
               « " [executor] [DISCRETIZATION_POINTS]" « std::endl;
    std::exit(-1);
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const IndexType discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
problem:
auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType(6) * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
std::vector<IndexType> row_ptrs(discretization_points + 1);
std::vector<IndexType> col_idxs(3 * discretization_points
std::vector<ValueType> values(3 * discretization_points - 2);
right hand side
std::vector<ValueType> rhs(discretization_points);
std::vector<ValueType> u(discretization_points, 0.0);
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
generate_stencil_matrix(discretization_points, row_ptrs.data(),
                         col_idxs.data(), values.data());
looking for solution u = x^3: f = 6x, u(0) = 0, u(1) = 1
reduction_factor);
Uncomment to print the solution print solution<ValueType, IndexType>(discretization points, 0, 1, u.data());
    std::cout « "The average relative error is "
               « calculate_error(discretization_points, u.data(), correct_u) /
                      discretization_points
               « std::endl;
}
```

Results

This is the expected output:

The average relative error is 2.52236e-11

Comments about programming and debugging

The plain program

```
******GINKGO LICENSE>****************
Copyright (c) 2017-2023, the Ginkgo authors
All rights reserved.
Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions
1. Redistributions of source code must retain the above copyright
notice, this list of conditions and the following disclaimer
2. Redistributions in binary form must reproduce the above copyright
notice, this list of conditions and the following disclaimer in the
documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its
contributors may be used to endorse or promote products derived from
this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS
IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED
TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A
PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
This example solves a 1D Poisson equation:
u : [0, 1] \rightarrow R
u'' = f
u(0) = u0
u(1) = u1
using a finite difference method on an equidistant grid with \ discretization
points ('K' can be controlled with a command line parameter). The discretization
is done via the second order Taylor polynomial:
u(x + h) = u(x) - u'(x)h + 1/2 u''(x)h^2 + O(h^3)
u(x - h) = u(x) + u'(x)h + 1/2 u''(x)h^2 + O(h^3)
-u(x - h) + 2u(x) + -u(x + h) = -f(x)h^2 + O(h^3)
For an equidistant grid with K "inner" discretization points x1, ..., xk, and
step size h = 1 / (K + 1), the formula produces a system of linear equations
2u_1 - u_2 = -f_1 h^2 + u^0

-u_(k-1) + 2u_k - u_(k+1) = -f_k h^2,
                                            k = 2, ..., K - 1
                         = -f_K h^2 + u1
-u_{K-1} + 2u_{K}
which is then solved using Ginkgo's implementation of the CG method
preconditioned with block-Jacobi. It is also possible to specify on which
executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x) = 6x' (making the solution 'u(x) = x^3'), but
that can be changed in the 'main' function.
The intention of the example is to show how Ginkgo can be integrated into
existing software - the 'generate_stencil_matrix', 'generate_rhs',
'print_solution', 'compute_error' and 'main' function do not reference Ginkgo at
all (i.e. they could have been there before the application developer decided to
use Ginkgo, and the only part where Ginkgo is introduced is inside the
'solve_system' function.
 #include <ginkgo/ginkgo.hpp>
```

```
#include <iostream>
#include <map>
#include <string>
#include <vector>
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType discretization_points,
                                   IndexType* row_ptrs, IndexType* col_idxs,
                                    ValueType* values)
    IndexType pos = 0;
     const ValueType coefs[] = {-1, 2, -1};
     row_ptrs[0] = pos;
for (IndexType i = 0; i < discretization_points; ++i) {</pre>
          for (auto ofs : {-1, 0, 1}) {
    if (0 <= i + ofs && i + ofs < discretization_points) {</pre>
                    values[pos] = coefs[ofs + 1];
                    col_idxs[pos] = i + ofs;
                    ++pos;
          row_ptrs[i + 1] = pos;
template <typename Closure, typename ValueType, typename IndexType>
void generate_rhs(IndexType discretization_points, Closure f, ValueType u0,
                      ValueType u1, ValueType* rhs)
     const ValueType h = 1.0 / (discretization_points + 1);
for (IndexType i = 0; i < discretization_points; ++i) {
   const ValueType xi = ValueType(i + 1) * h;
   rhs[i] = -f(xi) * h * h;</pre>
     rhs[0] += u0;
     rhs[discretization_points - 1] += u1;
std::cout « u0 « ' \ 'n'; for (IndexType i = 0; i < discretization_points; ++i) { std::cout « u[i] « ' \ 'n';
     std::cout « u1 « std::endl;
template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType discretization_points,
                                                          const ValueType* u,
                                                          Closure correct_u)
     const ValueType h = 1.0 / (discretization_points + 1);
gko::remove_complex<ValueType> error = 0.0;
     for (IndexType i = 0; i < discretization_points; ++i) {</pre>
         using std::abs;
         const ValueType xi = ValueType(i + 1) * h;
error += abs(u[i] - correct_u(xi)) / abs(correct_u(xi));
template <typename ValueType, typename IndexType>
using vec = gko::matrix::Dense<ValueType>;
     using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
     using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto@ dp = discretization_points;
     std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
          exec_map{
               {"omp", [] { return gko::OmpExecutor::create(); }},
               {"cuda",
                [] {
                     return gko::CudaExecutor::create(0, gko::OmpExecutor::create(),
                                                              true);
                11.
               {"hip",
                [] {
                     return gko::HipExecutor::create(0, gko::OmpExecutor::create(),
               {"dpcpp",
                [] {
```

```
return gko::DpcppExecutor::create(0,
                                                   gko::OmpExecutor::create());
            {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
auto matrix = mtx::create(exec, gko::dim<2>(dp),
                              val_array::view(app_exec, 3 * dp - 2, values),
                              idx_array::view(app_exec, 3 * dp - 2, col_idxs),
                              idx_array::view(app_exec, dp + 1, row_ptrs));
    auto b = vec::create(exec, gko::dim<2>(dp, 1),
   auto solver_gen =
       cg::build()
            .with_criteria(gko::stop::Iteration::build()
                               .with_max_iters(gko::size_type(dp))
                               .on(exec),
                           gko::stop::ResidualNorm<ValueType>::build()
                               .with_reduction_factor(reduction_factor)
                               .on(exec))
            .with_preconditioner(bj::build().on(exec))
            .on(exec);
    auto solver = solver_gen->generate(gko::give(matrix));
    solver->apply(b, x);
int main(int argc, char* argv[])
    using ValueType = double:
    using IndexType = int;
   std::exit(-1);
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const IndexType discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType(6) * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    std::vector<IndexType> row_ptrs(discretization_points + 1);
    std::vector<IndexType> col_idxs(3 * discretization_points - 2);
    std::vector<ValueType> values(3 * discretization_points - 2);
    std::vector<ValueType> rhs(discretization_points);
    std::vector<ValueType> u(discretization_points, 0.0);
    const gko::remove_complex<ValueType> reduction_factor = 1e-7;
    generate_stencil_matrix(discretization_points, row_ptrs.data(),
                            col_idxs.data(), values.data());
    generate_rhs(discretization_points, f, u0, u1, rhs.data());
    solve_system(executor_string, discretization_points, row_ptrs.data(),
                 col_idxs.data(), values.data(), rhs.data(), u.data(),
                 reduction_factor);
    std::cout « "The average relative error is "
             « calculate_error(discretization_points, u.data(), correct_u) /
                    discretization_points
              « std::endl;
}
```

Chapter 41

Module Documentation

41.1 CUDA Executor

A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.

Classes

• class gko::CudaExecutor

This is the Executor subclass which represents the CUDA device.

41.1.1 Detailed Description

A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.

41.2 DPC++ Executor

A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.

Classes

• class gko::DpcppExecutor

This is the Executor subclass which represents a DPC++ enhanced device.

41.2.1 Detailed Description

A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.

41.3 Executors 281

41.3 Executors

A module dedicated to the implementation and usage of the executors in Ginkgo.

Modules

CUDA Executor

A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.

DPC++ Executor

A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.

HIP Executor

A module dedicated to the implementation and usage of the HIP executor in Ginkgo.

OpenMP Executor

A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.

Reference Executor

A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

Classes

· class gko::Operation

Operations can be used to define functionalities whose implementations differ among devices.

· class gko::Executor

The first step in using the Ginkgo library consists of creating an executor.

class gko::executor_deleter< T >

This is a deleter that uses an executor's free method to deallocate the data.

class gko::OmpExecutor

This is the Executor subclass which represents the OpenMP device (typically CPU).

class gko::ReferenceExecutor

This is a specialization of the OmpExecutor, which runs the reference implementations of the kernels used for debugging purposes.

· class gko::CudaExecutor

This is the Executor subclass which represents the CUDA device.

· class gko::HipExecutor

This is the Executor subclass which represents the HIP enhanced device.

· class gko::DpcppExecutor

This is the Executor subclass which represents a DPC++ enhanced device.

Macros

#define GKO_REGISTER_OPERATION(_name, _kernel)

Binds a set of device-specific kernels to an Operation.

#define GKO_REGISTER_HOST_OPERATION(_name, _kernel)

Binds a host-side kernel (independent of executor type) to an Operation.

41.3.1 Detailed Description

A module dedicated to the implementation and usage of the executors in Ginkgo.

Below, we provide a brief introduction to executors in Ginkgo, how they have been implemented, how to best make use of them and how to add new executors.

41.3.2 Executors in Ginkgo.

The first step in using the Ginkgo library consists of creating an executor. Executors are used to specify the location for the data of linear algebra objects, and to determine where the operations will be executed. Ginkgo currently supports three different executor types:

- OpenMP Executor specifies that the data should be stored and the associated operations executed on an OpenMP-supporting device (e.g. host CPU);
- CUDA Executor specifies that the data should be stored and the operations executed on the NVIDIA GPU accelerator;
- HIP Executor uses the HIP library to compile code for either NVIDIA or AMD GPU accelerator;
- DPC++ Executor uses the DPC++ compiler for any DPC++ supported hardware (e.g. Intel CPUs, GPU, FPGAs, ...);
- Reference Executor executes a non-optimized reference implementation, which can be used to debug the library.

41.3.3 Macro Definition Documentation

41.3.3.1 GKO REGISTER HOST OPERATION

Value:

Binds a host-side kernel (independent of executor type) to an Operation.

It also defines a helper function which creates the associated operation. Any input arguments passed to the helper function are forwarded to the kernel when the operation is executed. The kernel name is searched for in the namespace where this macro is called. Host operations are used to make computations that are not part of the device kernels visible to profiling loggers and benchmarks.

Parameters

_name	operation name
_kernel	kernel which will be bound to the operation

41.3 Executors 283

41.3.3.2 Example

```
{c++}
void host_kernel(int) {
    // do some expensive computations
}
// Bind the kernels to the operation
GKO_REGISTER_HOST_OPERATION(my_op, host_kernel);
int main() {
    // create executor
    auto ref = ReferenceExecutor::create();
    // create the operation
    auto op = make_my_op(5); // x = 5
    ref->run(op); // run host kernel
}
```

41.3.3.3 GKO_REGISTER_OPERATION

Binds a set of device-specific kernels to an Operation.

It also defines a helper function which creates the associated operation. Any input arguments passed to the helper function are forwarded to the kernel when the operation is executed.

The kernels used to bind the operation are searched in kernels::DEV_TYPE namespace, where DEV_TYPE is replaced by omp, cuda, hip, dpcpp and reference.

Parameters

_name	operation name
_kernel	kernel which will be bound to the operation

41.3.3.4 Example

```
{c++}
// define the omp, cuda, hip and reference kernels which will be bound to the
// operation
namespace kernels {
namespace omp {
void my_kernel(int x) {
     // omp code
namespace cuda {
void my_kernel(int x) {
      // cuda code
namespace hip {
void my_kernel(int x) {
    // hip code
namespace dpcpp {
void my_kernel(int x) {
     // dpcpp code
namespace reference {
void my_kernel(int x) {
    // reference code
```

```
// Bind the kernels to the operation
GKO_REGISTER_OPERATION(my_op, my_kernel);
int main() {
    // create executors
    auto omp = OmpExecutor::create();
    auto cuda = CudaExecutor::create(0, omp);
    auto hip = HipExecutor::create(0, omp);
    auto dpcpp = DpcppExecutor::create(0, omp);
    auto ref = ReferenceExecutor::create();
    // create the operation
    auto op = make_my_op(5); // x = 5
    omp->run(op); // run omp kernel
    cuda->run(op); // run hip kernel
    dpcpp->run(op); // run bPC++ kernel
    ref->run(op); // run reference kernel
}
```

41.4 Factorizations 285

41.4 Factorizations

A module dedicated to the implementation and usage of the Factorizations in Ginkgo.

Namespaces

• gko::factorization

The Factorization namespace.

Classes

class gko::factorization::lc< ValueType, IndexType >

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

class gko::factorization::llu
 ValueType, IndexType

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

class gko::factorization::ParIc< ValueType, IndexType >

ParIC is an incomplete Cholesky factorization which is computed in parallel.

class gko::factorization::ParIct< ValueType, IndexType >

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

- class gko::factorization::Parllu< ValueType, IndexType >

ParILU is an incomplete LU factorization which is computed in parallel.

class gko::factorization::Parllut< ValueType, IndexType >

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

41.4.1 Detailed Description

A module dedicated to the implementation and usage of the Factorizations in Ginkgo.

41.5 HIP Executor

A module dedicated to the implementation and usage of the HIP executor in Ginkgo.

Classes

• class gko::HipExecutor

This is the Executor subclass which represents the HIP enhanced device.

41.5.1 Detailed Description

A module dedicated to the implementation and usage of the HIP executor in Ginkgo.

41.6 Jacobi Preconditioner 287

41.6 Jacobi Preconditioner

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

Classes

- struct gko::preconditioner::block_interleaved_storage_scheme < IndexType >
 Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.
- class gko::preconditioner::Jacobi< ValueType, IndexType >

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

41.6.1 Detailed Description

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

41.7 Linear Operators

A module dedicated to the implementation and usage of the Linear operators in Ginkgo.

Modules

Factorizations

A module dedicated to the implementation and usage of the Factorizations in Ginkgo.

· SpMV employing different Matrix formats

A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.

Preconditioners

A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.

Solvers

A module dedicated to the implementation and usage of the Solvers in Ginkgo.

Classes

class gko::Combination < ValueType >

The Combination class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$

class gko::Composition < ValueType >

The Composition class can be used to compose linear operators op1, op2, ..., opn and obtain the operator op1 * op2 * ...

class gko::LinOpFactory

A LinOpFactory represents a higher order mapping which transforms one linear operator into another.

class gko::ReadableFromMatrixData< ValueType, IndexType >

A LinOp implementing this interface can read its data from a matrix_data structure.

class gko::WritableToMatrixData< ValueType, IndexType >

A LinOp implementing this interface can write its data to a matrix_data structure.

· class gko::Preconditionable

A LinOp implementing this interface can be preconditioned.

· class gko::DiagonalLinOpExtractable

The diagonal of a LinOp can be extracted.

class gko::DiagonalExtractable
 ValueType >

The diagonal of a LinOp implementing this interface can be extracted.

class gko::EnableAbsoluteComputation< AbsoluteLinOp >

The EnableAbsoluteComputation mixin provides the default implementations of compute_absolute_linop and the absolute interface.

class gko::EnableLinOp
 ConcreteLinOp
 PolymorphicBase

The EnableLinOp mixin can be used to provide sensible default implementations of the majority of the LinOp and PolymorphicObject interface.

 $\bullet \ \ {\it class gko::} Perturbation < \ \ {\it ValueType} >$

The Perturbation class can be used to construct a LinOp to represent the operation (identity + scalar * basis * projector).

 $\bullet \ \, {\sf class} \ \, {\sf gko::experimental::EnableDistributedLinOp} < \ \, {\sf ConcreteLinOp}, \ \, {\sf PolymorphicBase} > \\$

This mixin does the same as EnableLinOp, but for concrete types that are derived from distributed::DistributedBase.

class gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >

A Schwarz preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

class gko::experimental::distributed::Vector< ValueType >

41.7 Linear Operators 289

Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

class gko::factorization::lc
 ValueType, IndexType >

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

class gko::factorization::Ilu
 ValueType, IndexType

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

class gko::factorization::ParIc< ValueType, IndexType >

ParIC is an incomplete Cholesky factorization which is computed in parallel.

class gko::factorization::ParIct
 ValueType, IndexType >

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

class gko::factorization::Parllu< ValueType, IndexType >

ParILU is an incomplete LU factorization which is computed in parallel.

class gko::factorization::Parllut< ValueType, IndexType >

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

class gko::matrix::Coo< ValueType, IndexType >

COO stores a matrix in the coordinate matrix format.

class gko::matrix::Csr< ValueType, IndexType >

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

class gko::matrix::Dense
 ValueType >

Dense is a matrix format which explicitly stores all values of the matrix.

class gko::matrix::Diagonal < ValueType >

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

class gko::matrix::Ell< ValueType, IndexType >

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

class gko::matrix::Fbcsr< ValueType, IndexType >

Fixed-block compressed sparse row storage matrix format.

class gko::matrix::Fft

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

class gko::matrix::Fft2

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

class gko::matrix::Fft3

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

class gko::matrix::Hybrid< ValueType, IndexType >

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

class gko::matrix::Identity
 ValueType >

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

This factory is a utility which can be used to generate Identity operators.

class gko::matrix::Permutation< IndexType >

Permutation is a matrix "format" which stores the row and column permutation arrays which can be used for reordering the rows and columns a matrix.

class gko::matrix::RowGatherer< IndexType >

RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.

class gko::matrix::Sellp< ValueType, IndexType >

SELL-P is a matrix format similar to ELL format.

class gko::matrix::SparsityCsr< ValueType, IndexType >

SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

class gko::multigrid::FixedCoarsening
 ValueType, IndexType >

FixedCoarsening is a very simple coarse grid generation algorithm.

class gko::multigrid::Pgm< ValueType, IndexType >

Parallel graph match (Pgm) is the aggregate method introduced in the paper M.

class gko::preconditioner::lc< LSolverType, IndexType >

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H*x=b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

class gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >

The Incomplete LU (ILU) preconditioner solves the equation LUx = b for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

class gko::preconditioner::lsai< lsaiType, ValueType, IndexType >

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

class gko::preconditioner::Jacobi< ValueType, IndexType >

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

class gko::solver::Bicg< ValueType >

BICG or the Biconjugate gradient method is a Krylov subspace solver.

class gko::solver::Bicgstab
 ValueType >

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

class gko::solver::CbGmres< ValueType >

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

class gko::solver::Cg< ValueType >

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

class gko::solver::Cgs< ValueType >

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

class gko::solver::Fcg< ValueType >

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

class gko::solver::Gcr< ValueType >

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

class gko::solver::Gmres < ValueType >

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

class gko::solver::ldr< ValueType >

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

class gko::solver::Ir< ValueType >

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

· class gko::solver::Multigrid

Multigrid methods have a hierarchy of many levels, whose corase level is a subset of the fine level, of the problem.

class gko::solver::EnableSolverBase
 DerivedType, MatrixType >

A LinOp deriving from this CRTP class stores a system matrix.

class gko::solver::IterativeBase

A LinOp implementing this interface stores a stopping criterion factory.

class gko::solver::EnableIterativeBase
 DerivedType >

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

class gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType >

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

41.7 Linear Operators 291

class gko::solver::LowerTrs< ValueType, IndexType >

LowerTrs is the triangular solver which solves the system L x = b, when L is a lower triangular matrix.

class gko::solver::UpperTrs< ValueType, IndexType >

UpperTrs is the triangular solver which solves the system Ux = b, when U is an upper triangular matrix.

Macros

• #define GKO CREATE FACTORY PARAMETERS(parameters name, factory name)

This Macro will generate a new type containing the parameters for the factory_factory_name.

#define GKO_ENABLE_LIN_OP_FACTORY(_lin_op, _parameters_name, _factory_name)

This macro will generate a default implementation of a LinOpFactory for the LinOp subclass it is defined in.

#define GKO ENABLE BUILD METHOD(factory name)

Defines a build method for the factory, simplifying its construction by removing the repetitive typing of factory's name.

#define GKO_FACTORY_PARAMETER(_name, ...)

Creates a factory parameter in the factory parameters structure.

Creates a scalar factory parameter in the factory parameters structure.

#define GKO_FACTORY_PARAMETER_VECTOR(_name, ...) GKO_FACTORY_PARAMETER(_name, _ ← VA_ARGS__)

Creates a vector factory parameter in the factory parameters structure.

Typedefs

template < typename ConcreteFactory , typename ConcreteLinOp , typename ParametersType , typename PolymorphicBase = Lin←
 OpFactory>

using gko::EnableDefaultLinOpFactory = EnableDefaultFactory< ConcreteFactory, ConcreteLinOp, ParametersType, PolymorphicBase >

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of LinOpFactory.

Functions

template<typename Matrix, typename... TArgs>
 std::unique_ptr< Matrix > gko::initialize (size_type stride, std::initializer_list< typename Matrix::value_type
 > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a column-vector.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > gko::initialize (std::initializer_list< typename Matrix::value_type > vals, std
 ::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a column-vector.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > gko::initialize (size_type stride, std::initializer_list< std::initializer_list< typename
 Matrix::value_type >> vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a matrix.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > gko::initialize (std::initializer_list< std::initializer_list< typename Matrix::value_
 type >> vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a matrix.

41.7.1 Detailed Description

A module dedicated to the implementation and usage of the Linear operators in Ginkgo.

Below we elaborate on one of the most important concepts of Ginkgo, the linear operator. The linear operator (LinOp) is a base class for all linear algebra objects in Ginkgo. The main benefit of having a single base class for the entire collection of linear algebra objects (as opposed to having separate hierarchies for matrices, solvers and preconditioners) is the generality it provides.

41.7.2 Advantages of this approach and usage

A common interface often allows for writing more generic code. If a user's routine requires only operations provided by the LinOp interface, the same code can be used for any kind of linear operators, independent of whether these are matrices, solvers or preconditioners. This feature is also extensively used in Ginkgo itself. For example, a preconditioner used inside a Krylov solver is a LinOp. This allows the user to supply a wide variety of preconditioners: either the ones which were designed to be used in this scenario (like ILU or block-Jacobi), a user-supplied matrix which is known to be a good preconditioner for the specific problem, or even another solver (e.g., if constructing a flexible GMRES solver).

For example, a matrix free implementation would require the user to provide an apply implementation and instead of passing the generated matrix to the solver, they would have to provide their apply implementation for all the executors needed and no other code needs to be changed. See The custom-matrix-format program example for more details.

41.7.3 Linear operator as a concept

The linear operator (LinOp) is a base class for all linear algebra objects in Ginkgo. The main benefit of having a single base class for the entire collection of linear algebra objects (as opposed to having separate hierarchies for matrices, solvers and preconditioners) is the generality it provides.

First, since all subclasses provide a common interface, the library users are exposed to a smaller set of routines. For example, a matrix-vector product, a preconditioner application, or even a system solve are just different terms given to the operation of applying a certain linear operator to a vector. As such, Ginkgo uses the same routine name, LinOp::apply() for each of these operations, where the actual operation performed depends on the type of linear operator involved in the operation.

Second, a common interface often allows for writing more generic code. If a user's routine requires only operations provided by the LinOp interface, the same code can be used for any kind of linear operators, independent of whether these are matrices, solvers or preconditioners. This feature is also extensively used in Ginkgo itself. For example, a preconditioner used inside a Krylov solver is a LinOp. This allows the user to supply a wide variety of preconditioners: either the ones which were designed to be used in this scenario (like ILU or block-Jacobi), a user-supplied matrix which is known to be a good preconditioner for the specific problem, or even another solver (e.g., if constructing a flexible GMRES solver).

A key observation for providing a unified interface for matrices, solvers, and preconditioners is that the most common operation performed on all of them can be expressed as an application of a linear operator to a vector:

- the sparse matrix-vector product with a matrix A is a linear operator application y = Ax;
- the application of a preconditioner is a linear operator application $y = M^{-1}x$, where M is an approximation of the original system matrix A (thus a preconditioner represents an "approximate inverse" operator M^{-1}).
- the system solve Ax = b can be viewed as linear operator application $x = A^{-1}b$ (it goes without saying that the implementation of linear system solves does not follow this conceptual idea), so a linear system solver can be viewed as a representation of the operator A^{-1} .

41.7 Linear Operators 293

Finally, direct manipulation of LinOp objects is rarely required in simple scenarios. As an illustrative example, one could construct a fixed-point iteration routine $x_{k+1} = Lx_k + b$ as follows:

```
std::unique_ptr<matrix::Dense<» calculate_fixed_point(
    int iters, const LinOp *L, const matrix::Dense<> *x0
    const matrix::Dense<> *b)
{
    auto x = gko::clone(x0);
    auto tmp = gko::clone(x0);
    auto one = Dense<>::create(L->get_executor(), {1.0,});
    for (int i = 0; i < iters; ++i) {
        L->apply(tmp, x);
        x->add_scaled(one, b);
        tmp->copy_from(x);
    }
    return x;
```

Here, if L is a matrix, LinOp::apply() refers to the matrix vector product, and L->apply(a, b) computes $b = L \cdot a$. x->add_scaled(one, b) is the axpy vector update x := x + b.

The interesting part of this example is the apply() routine at line 4 of the function body. Since this routine is part of the LinOp base class, the fixed-point iteration routine can calculate a fixed point not only for matrices, but for any type of linear operator.

Linear Operators

41.7.4 Macro Definition Documentation

41.7.4.1 GKO_CREATE_FACTORY_PARAMETERS

This Macro will generate a new type containing the parameters for the factory _factory_name.

For more details, see GKO_ENABLE_LIN_OP_FACTORY(). It is required to use this macro before calling the macro GKO_ENABLE_LIN_OP_FACTORY(). It is also required to use the same names for all parameters between both macros.

Parameters

_parameters_name	name of the parameters member in the class
_factory_name	name of the generated factory type

41.7.4.2 GKO_ENABLE_BUILD_METHOD

Defines a build method for the factory, simplifying its construction by removing the repetitive typing of factory's name.

Parameters

```
_factory_name the factory for which to define the method
```

41.7.4.3 GKO ENABLE LIN OP FACTORY

This macro will generate a default implementation of a LinOpFactory for the LinOp subclass it is defined in.

It is required to first call the macro GKO_CREATE_FACTORY_PARAMETERS() before this one in order to instantiate the parameters type first.

The list of parameters for the factory should be defined in a code block after the macro definition, and should contain a list of GKO_FACTORY_PARAMETER_* declarations. The class should provide a constructor with signature \leftarrow _lin_op(const _factory_name *, std::shared_ptr<const LinOp>) which the factory will use a callback to construct the object.

A minimal example of a linear operator is the following:

auto my_op = fact->generate(gko::matrix::Identity::create(exec, 2));
std::cout « my_op->get_my_parameters().my_value; // prints 0

""c++ struct MyLinOp: public EnableLinOp<MyLinOp> { GKO_ENABLE_LIN_OP_FACTORY(MyLinOp, my_parameters, Factory) { // a factory parameter named "my_value", of type int and default // value of 5 int GKO_FACTORY_PARAMETER_SCALAR(my_value, // a factory parameter named my_pair of type std::pair<int,int>// and default value {5, 5} std::pair<int, int> GKO_FACTORY_PARAMETER_VECTOR(my_pair, 5, 5); }; // constructor needed by EnableLinOp explicit MyLinOp(std::shared_ptr<const Executor> exec) { : EnableLinOp<MyLinOp>(exec) {} // constructor needed by the factory explicit MyLinOp(const Factory *factory, std::shared_ptr<const LinOp> matrix) : EnableLinOp<← MyLinOp>(factory->get_executor()), matrix->get_size()), // store factory's parameters locally my_parameters_← {factory->get_parameters()}, { int value = my_parameters_.my_value; // do something with value } MyLinOp can then be created as follows: `c++ auto exec = gko::ReferenceExecutor::create(); // create a factory with default 'my_value' parameter auto fact = MyLinOp::build().on(exec); // create a operator using the factory: auto my_op = fact->generate(gko::matrix::Identity::create(exec, 2)); std::cout « my_op->get_my_parameters().my_value; // prints 5
// create a factory with custom 'my_value' parameter auto fact = MyLinOp::build().with_my_value(0).on(exec); // create a operator using the factory:

41.7 Linear Operators 295

Note

It is possible to combine both the #GKO_CREATE_FACTORY_PARAMETER_*() macros with this one in a unique macro for class **templates** (not with regular classes). Splitting this into two distinct macros allows to use them in all contexts. See https://stackoverflow.com/q/50202718/9385966 for more details.

Parameters

_lin_op	concrete operator for which the factory is to be created [CRTP parameter]
_parameters_name	name of the parameters member in the class (its type is
	<pre><_parameters_name>_type, the protected member's name is</pre>
	<_parameters_name>_, and the public getter's name is
	<pre>get_<_parameters_name>())</pre>
_factory_name	name of the generated factory type

41.7.4.4 GKO_FACTORY_PARAMETER

Creates a factory parameter in the factory parameters structure.

Parameters

_name	name of the parameter
VA_ARGS	default value of the parameter

See also

GKO_ENABLE_LIN_OP_FACTORY for more details, and usage example

41.7.4.5 GKO_FACTORY_PARAMETER_SCALAR

Creates a scalar factory parameter in the factory parameters structure.

Scalar in this context means that the constructor for this type only takes a single parameter.

Parameters

_name	name of the parameter
_default	default value of the parameter

See also

GKO_ENABLE_LIN_OP_FACTORY for more details, and usage example

41.7.4.6 GKO_FACTORY_PARAMETER_VECTOR

Creates a vector factory parameter in the factory parameters structure.

Vector in this context means that the constructor for this type takes multiple parameters.

Parameters

_name	name of the parameter
_default	default value of the parameter

See also

GKO_ENABLE_LIN_OP_FACTORY for more details, and usage example

41.7.5 Typedef Documentation

41.7.5.1 EnableDefaultLinOpFactory

```
template<typename ConcreteFactory , typename ConcreteLinOp , typename ParametersType , typename PolymorphicBase = LinOpFactory> using gko::EnableDefaultLinOpFactory = typedef EnableDefaultFactory<ConcreteFactory, Concrete← LinOp, ParametersType, PolymorphicBase>
```

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of LinOpFactory.

41.7 Linear Operators 297

Template Parameters

ConcreteFactory	the concrete factory which is being implemented [CRTP parmeter]
ConcreteLinOp	the concrete LinOp type which this factory produces, needs to have a constructor which takes a const ConcreteFactory *, and an std::shared_ptr <const linop=""> as parameters.</const>
Parameters Type	a subclass of enable_parameters_type template which defines all of the parameters of the factory
PolymorphicBase	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of LinOpFactory

41.7.6 Function Documentation

41.7.6.1 initialize() [1/4]

Creates and initializes a matrix.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type.

Template Parameters

Ма	trix	matrix type to initialize (Dense has to implement the ConvertibleTo <matrix> interface)</matrix>
TA	rgs	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

stride	row stride for the temporary Dense matrix
vals	values used to initialize the matrix
exec	Executor associated to the matrix
create_args	additional arguments passed to Matrix::create, not including the Executor, which is passed as
	the first argument

```
1446 {
            using dense = matrix::Dense<typename Matrix::value_type>;
1447
            size_type num_rows = vals.size();
size_type num_cols = num_rows > 0 ? begin(vals)->size() : 1;
1448
1449
1450
           auto tmp =
1451
                 dense::create(exec->get_master(), dim<2>{num_rows, num_cols}, stride);
1452
            size_type ridx = 0;
           size_type ridx - 0;
for (const auto& row : vals) {
    size_type cidx = 0;
    for (const auto& elem : row) {
        tmp->at(ridx, cidx) = elem;
}
1453
1454
1455
1456
1457
                        ++cidx;
```

References gko::matrix::Dense< ValueType >::at().

41.7.6.2 initialize() [2/4]

Creates and initializes a column-vector.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type.

Template Parameters

Matrix	matrix type to initialize (Dense has to implement the ConvertibleTo <matrix> interface)</matrix>
TArgs	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

stride	row stride for the temporary Dense matrix
vals	values used to initialize the vector
exec	Executor associated to the vector
create_args	additional arguments passed to Matrix::create, not including the Executor, which is passed as the first argument

References gko::matrix::Dense< ValueType >::at().

41.7.6.3 initialize() [3/4]

Creates and initializes a matrix.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type. The stride of the intermediate Dense matrix is set to the number of columns of the initializer list.

41.7 Linear Operators 299

Template Parameters

Matrix	matrix type to initialize (Dense has to implement the ConvertibleTo <matrix> interface)</matrix>
TArgs	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

vals	values used to initialize the matrix
exec	Executor associated to the matrix
create_args	additional arguments passed to Matrix::create, not including the Executor, which is passed as the first argument

41.7.6.4 initialize() [4/4]

Creates and initializes a column-vector.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type. The stride of the intermediate Dense matrix is set to 1.

Template Parameters

Matrix	matrix type to initialize (Dense has to implement the ConvertibleTo <matrix> interface)</matrix>	
TArgs	argument types for Matrix::create method (not including the implied Executor as the first argument)]

Parameters

vals	values used to initialize the vector
exec	Executor associated to the vector
create_args	additional arguments passed to Matrix::create, not including the Executor, which is passed as the first argument

41.8 Logging

A module dedicated to the implementation and usage of the Logging in Ginkgo.

Namespaces

• gko::log

The logger namespace.

Classes

class gko::log::Convergence < ValueType >

Convergence is a Logger which logs data strictly from the criterion_check_completed event.

class gko::log::Papi< ValueType >

Papi is a Logger which logs every event to the PAPI software.

· class gko::log::PerformanceHint

PerformanceHint is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

class gko::log::Stream< ValueType >

Stream is a Logger which logs every event to a stream.

41.8.1 Detailed Description

A module dedicated to the implementation and usage of the Logging in Ginkgo.

The Logger class represents a simple Logger object. It comprises all masks and events internally. Every new logging event addition should be done here. The Logger class also provides a default implementation for most events which do nothing, therefore it is not an obligation to change all classes which derive from Logger, although it is good practice. The logger class is built using event masks to control which events should be logged, and which should not.

41.9 SpMV employing different Matrix formats

A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.

Classes

class gko::experimental::distributed::Vector< ValueType >

Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

class gko::matrix::Coo< ValueType, IndexType >

COO stores a matrix in the coordinate matrix format.

class gko::matrix::Csr< ValueType, IndexType >

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

class gko::matrix::Dense< ValueType >

Dense is a matrix format which explicitly stores all values of the matrix.

class gko::matrix::Diagonal < ValueType >

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

class gko::matrix::Ell< ValueType, IndexType >

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

class gko::matrix::Fbcsr< ValueType, IndexType >

Fixed-block compressed sparse row storage matrix format.

· class gko::matrix::Fft

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

class gko::matrix::Fft2

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

· class gko::matrix::Fft3

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

class gko::matrix::Hybrid< ValueType, IndexType >

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

class gko::matrix::Identity< ValueType >

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

class gko::matrix::IdentityFactory< ValueType >

This factory is a utility which can be used to generate Identity operators.

 $\bullet \ \, {\it class gko::matrix::Permutation}{< IndexType} >$

Permutation is a matrix "format" which stores the row and column permutation arrays which can be used for reordering the rows and columns a matrix.

class gko::matrix::Sellp< ValueType, IndexType >

SELL-P is a matrix format similar to ELL format.

class gko::matrix::SparsityCsr< ValueType, IndexType >

SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

41.9.1 Detailed Description

A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.

41.10 OpenMP Executor

A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.

Classes

• class gko::OmpExecutor

This is the Executor subclass which represents the OpenMP device (typically CPU).

41.10.1 Detailed Description

A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.

41.11 Preconditioners 303

41.11 Preconditioners

A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.

Modules

· Jacobi Preconditioner

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

Namespaces

• gko::experimental::distributed::preconditioner

The Preconditioner namespace.

· gko::preconditioner

The Preconditioner namespace.

Classes

· class gko::Preconditionable

A LinOp implementing this interface can be preconditioned.

class gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >

A Schwarz preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

class gko::preconditioner::lc< LSolverType, IndexType >

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H*x=b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

class gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >

The Incomplete LU (ILU) preconditioner solves the equation LUx = b for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

class gko::preconditioner::lsai< lsaiType, ValueType, IndexType >

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

class gko::preconditioner::Jacobi< ValueType, IndexType >

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

41.11.1 Detailed Description

A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.

41.12 Reference Executor

A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

Classes

• class gko::ReferenceExecutor

This is a specialization of the OmpExecutor, which runs the reference implementations of the kernels used for debugging purposes.

41.12.1 Detailed Description

A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

41.13 Solvers 305

41.13 Solvers

A module dedicated to the implementation and usage of the Solvers in Ginkgo.

Namespaces

· gko::solver

The ginkgo Solve namespace.

Classes

class gko::solver::Bicg
 ValueType >

BICG or the Biconjugate gradient method is a Krylov subspace solver.

class gko::solver::Bicgstab < ValueType >

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

class gko::solver::CbGmres< ValueType >

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

class gko::solver::Cg< ValueType >

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

class gko::solver::Cgs< ValueType >

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

class gko::solver::Fcg< ValueType >

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

class gko::solver::Gcr< ValueType >

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

class gko::solver::Gmres < ValueType >

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

class gko::solver::ldr< ValueType >

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

class gko::solver::Ir< ValueType >

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

· class gko::solver::Multigrid

Multigrid methods have a hierarchy of many levels, whose corase level is a subset of the fine level, of the problem.

class gko::solver::LowerTrs< ValueType, IndexType >

LowerTrs is the triangular solver which solves the system L x = b, when L is a lower triangular matrix.

class gko::solver::UpperTrs< ValueType, IndexType >

UpperTrs is the triangular solver which solves the system Ux = b, when U is an upper triangular matrix.

41.13.1 Detailed Description

A module dedicated to the implementation and usage of the Solvers in Ginkgo.

41.14 Stopping criteria

A module dedicated to the implementation and usage of the Stopping Criteria in Ginkgo.

Namespaces

gko::stop

The Stopping criterion namespace.

Classes

class gko::stop::Combined

The Combined class is used to combine multiple criterions together through an OR operation.

class gko::stop::Iteration

The Iteration class is a stopping criterion which stops the iteration process after a preset number of iterations.

class gko::stop::ResidualNormBase
 ValueType >

The ResidualNormBase class provides a framework for stopping criteria related to the residual norm.

class gko::stop::ResidualNorm< ValueType >

The ResidualNorm class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

class gko::stop::ImplicitResidualNorm< ValueType >

The ImplicitResidualNorm class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

class gko::stop::ResidualNormReduction< ValueType >

The ResidualNormReduction class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

class gko::stop::RelativeResidualNorm
 ValueType >

The RelativeResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

class gko::stop::AbsoluteResidualNorm< ValueType >

The AbsoluteResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

· class gko::stopping status

This class is used to keep track of the stopping status of one vector.

class gko::stop::Time

The Time class is a stopping criterion which stops the iteration process after a certain amout of time has passed.

Enumerations

· enum gko::stop::mode

The mode for the residual norm criterion.

Functions

template<typename FactoryContainer >
 std::shared_ptr< const CriterionFactory > gko::stop::combine (FactoryContainer &&factories)

Combines multiple criterion factories into a single combined criterion factory.

41.14 Stopping criteria 307

41.14.1 Detailed Description

A module dedicated to the implementation and usage of the Stopping Criteria in Ginkgo.

41.14.2 Enumeration Type Documentation

41.14.2.1 mode

```
enum gko::stop::mode [strong]
```

The mode for the residual norm criterion.

- absolute: Check for tolerance against residual norm. $\vert\vert r\vert\vert<\tau$
- initial_resnorm: Check for tolerance relative to the initial residual norm. $\frac{||r||}{||r_0||} < au$
- rhs_norm: Check for tolerance relative to the rhs norm. $\frac{||r||}{||b||} < \tau$

```
65 { absolute, initial_resnorm, rhs_norm };
```

41.14.3 Function Documentation

41.14.3.1 combine()

Combines multiple criterion factories into a single combined criterion factory.

This function treats a singleton container as a special case and avoids creating an additional object and just returns the input factory.

Template Parameters

FactoryContainer	a random access container type
------------------	--------------------------------

Parameters

factories a list of factories to comb	ined
---------------------------------------	------

Returns

a combined criterion factory if the input contains multiple factories or the input factory if the input contains only one factory

```
139 {
140
        switch (factories.size()) {
        case 0:
142
           GKO_NOT_SUPPORTED (nullptr);
143
        case 1:
144
            if (factories[0] == nullptr) {
                GKO_NOT_SUPPORTED (nullptr);
145
146
147
            return factories[0];
148
        default:
149
            if (factories[0] == nullptr) {
150
                // first factory must be valid to capture executor
                GKO_NOT_SUPPORTED(nullptr);
            } else {
152
               auto exec = factories[0]->get_executor();
153
                return Combined::build()
    .with_criteria(std::forward<FactoryContainer>(factories))
154
155
156
                     .on(exec);
157
            }
158
        }
159 }
```

Chapter 42

Namespace Documentation

42.1 gko Namespace Reference

The Ginkgo namespace.

Namespaces

accessor

The accessor namespace.

factorization

The Factorization namespace.

log

The logger namespace .

• matrix

The matrix namespace.

• multigrid

The multigrid components namespace.

name_demangling

The name demangling namespace.

· preconditioner

The Preconditioner namespace.

reorder

The Reorder namespace.

solver

The ginkgo Solve namespace.

• stop

The Stopping criterion namespace.

syn

The Synthesizer namespace.

xstd

The namespace for functionalities after C++14 standard.

Classes

· class AbsoluteComputable

The AbsoluteComputable is an interface that allows to get the component wise absolute of a LinOp.

class AbstractFactory

The AbstractFactory is a generic interface template that enables easy implementation of the abstract factory design pattern.

class AllocationError

AllocationError is thrown if a memory allocation fails.

· class amd_device

amd_device handles the number of executor on Amd devices and have the corresponding recursive_mutex.

struct are_all_integral

Evaluates if all template arguments Args fulfill std::is_integral.

· class array

An array is a container which encapsulates fixed-sized arrays, stored on the Executor tied to the array.

class BadDimension

BadDimension is thrown if an operation is being applied to a LinOp with bad dimensions.

· class BlockSizeError

Error that denotes issues between block sizes and matrix dimensions.

class Combination

The Combination class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$

· class Composition

The Composition class can be used to compose linear operators op1, op2, ..., opn and obtain the operator op1*op2*...

· class ConvertibleTo

Convertible To interface is used to mark that the implementer can be converted to the object of ResultType.

class CpuTimer

A timer using std::chrono::steady_clock for timing.

struct cpx_real_type

Access the underlying real type of a complex number.

class CublasError

CublasError is thrown when a cuBLAS routine throws a non-zero error code.

· class cuda_stream

An RAII wrapper for a custom CUDA stream.

class CudaError

CudaError is thrown when a CUDA routine throws a non-zero error code.

class CudaExecutor

This is the Executor subclass which represents the CUDA device.

class CudaTimer

A timer using events for timing on a CudaExecutor.

class CufftError

CufftError is thrown when a cuFFT routine throws a non-zero error code.

class CurandError

CurandError is thrown when a cuRAND routine throws a non-zero error code.

class CusparseError

CusparseError is thrown when a cuSPARSE routine throws a non-zero error code.

struct default_converter

Used to convert objects of type S to objects of type R using static_cast.

· class device matrix data

This type is a device-side equivalent to matrix_data.

class DiagonalExtractable

The diagonal of a LinOp implementing this interface can be extracted.

· class DiagonalLinOpExtractable

The diagonal of a LinOp can be extracted.

struct dim

A type representing the dimensions of a multidimensional object.

· class DimensionMismatch

DimensionMismatch is thrown if an operation is being applied to LinOps of incompatible size.

class DpcppExecutor

This is the Executor subclass which represents a DPC++ enhanced device.

class DpcppTimer

A timer using kernels for timing on a DpcppExecutor in profiling mode.

· class enable_parameters_type

The enable_parameters_type mixin is used to create a base implementation of the factory parameters structure.

class EnableAbsoluteComputation

The EnableAbsoluteComputation mixin provides the default implementations of compute_absolute_linop and the absolute interface.

class EnableAbstractPolymorphicObject

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new abstract object.

class EnableCreateMethod

This mixin implements a static create() method on ConcreteType that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an std::unique_ptr to such an object.

class EnableDefaultFactory

This mixin provides a default implementation of a concrete factory.

class EnableLinOp

The EnableLinOp mixin can be used to provide sensible default implementations of the majority of the LinOp and PolymorphicObject interface.

class EnablePolymorphicAssignment

This mixin is used to enable a default PolymorphicObject::copy_from() implementation for objects that have implemented conversions between them.

class EnablePolymorphicObject

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new concrete polymorphic object.

class Error

The Error class is used to report exceptional behaviour in library functions.

class Executor

The first step in using the Ginkgo library consists of creating an executor.

class executor_deleter

This is a deleter that uses an executor's free method to deallocate the data.

class hip_stream

An RAII wrapper for a custom HIP stream.

· class HipblasError

HipblasError is thrown when a hipBLAS routine throws a non-zero error code.

class HipError

HipError is thrown when a HIP routine throws a non-zero error code.

class HipExecutor

This is the Executor subclass which represents the HIP enhanced device.

class HipfftError

HipfftError is thrown when a hipFFT routine throws a non-zero error code.

· class HiprandError

HiprandError is thrown when a hipRAND routine throws a non-zero error code.

class HipsparseError

HipsparseError is thrown when a hipSPARSE routine throws a non-zero error code.

· class HipTimer

A timer using events for timing on a HipExecutor.

· class index set

An index set class represents an ordered set of intervals.

· class KernelNotFound

KernelNotFound is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.

class LinOpFactory

A LinOpFactory represents a higher order mapping which transforms one linear operator into another.

· class machine_topology

The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.

· class matrix assembly data

This structure is used as an intermediate type to assemble a sparse matrix.

· struct matrix data

This structure is used as an intermediate data type to store a sparse matrix.

struct matrix_data_entry

Type used to store nonzeros.

class MetisError

MetisError is thrown when METIS routine throws an error code.

class MpiError

MpiError is thrown when a MPI routine throws a non-zero error code.

class NotCompiled

NotCompiled is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.

· class NotImplemented

NotImplemented is thrown in case an operation has not yet been implemented (but will be implemented in the future).

class NotSupported

NotSupported is thrown in case it is not possible to perform the requested operation on the given object type.

· class null_deleter

This is a deleter that does not delete the object.

· class nvidia device

nvidia_device handles the number of executor on Nvidia devices and have the corresponding recursive_mutex.

class OmpExecutor

This is the Executor subclass which represents the OpenMP device (typically CPU).

class Operation

Operations can be used to define functionalities whose implementations differ among devices.

class OutOfBoundsError

OutOfBoundsError is thrown if a memory access is detected to be out-of-bounds.

class OverflowError

OverflowError is thrown when an index calculation for storage requirements overflows.

· class Permutable

Linear operators which support permutation should implement the Permutable interface.

class Perturbation

The Perturbation class can be used to construct a LinOp to represent the operation (identity + scalar * basis * projector).

· class PolymorphicObject

A PolymorphicObject is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.

· class precision_reduction

This class is used to encode storage precisions of low precision algorithms.

· class Preconditionable

A LinOp implementing this interface can be preconditioned.

class ptr_param

This class is used for function parameters in the place of raw pointers.

· class range

A range is a multidimensional view of the memory.

class ReadableFromMatrixData

A LinOp implementing this interface can read its data from a matrix_data structure.

· class ReferenceExecutor

This is a specialization of the OmpExecutor, which runs the reference implementations of the kernels used for debugging purposes.

class ScaledIdentityAddable

Adds the operation M <- a I + b M for matrix M, identity operator I and scalars a and b, where M is the calling object.

class scoped device id guard

This move-only class uses RAII to set the device id within a scoped block, if necessary.

· struct span

A span is a lightweight structure used to create sub-ranges from other ranges.

class stopping_status

This class is used to keep track of the stopping status of one vector.

class StreamError

StreamError is thrown if accessing a stream failed.

class time_point

An opaque wrapper for a time point generated by a timer.

· class Timer

Represents a generic timer that can be used to record time points and measure time differences on host or device streams

class Transposable

Linear operators which support transposition should implement the Transposable interface.

class UnsupportedMatrixProperty

Exception throws if a matrix does not have a property required by a numerical method.

class UseComposition

The UseComposition class can be used to store the composition information in LinOp.

· class ValueMismatch

ValueMismatch is thrown if two values are not equal.

struct version

This structure is used to represent versions of various Ginkgo modules.

· class version info

Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:

class WritableToMatrixData

A LinOp implementing this interface can write its data to a matrix data structure.

Typedefs

template < typename ConcreteFactory , typename ConcreteLinOp , typename ParametersType , typename PolymorphicBase = LinOp ←
Factory>

using EnableDefaultLinOpFactory = EnableDefaultFactory < ConcreteFactory, ConcreteLinOp, Parameters ← Type, PolymorphicBase >

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of LinOpFactory.

• template<typename T >

```
using is_complex_s = detail::is_complex_impl< T >
```

Allows to check if T is a complex value during compile time by accessing the value attribute of this struct.

using uintptr = std::uintptr_t

```
• template<typename T >
  using is_complex_or_scalar_s = detail::is_complex_or_scalar_impl< T >
      Allows to check if T is a complex or scalar value during compile time by accessing the value attribute of this struct.

    template<typename T >

  using remove complex = typename detail::remove complex s< T >::type
      Obtain the type which removed the complex of complex/scalar type or the template parameter of class by accessing
      the type attribute of this struct.
• template<typename T >
  using to complex = typename detail::to complex s< T >::type
      Obtain the type which adds the complex of complex/scalar type or the template parameter of class by accessing the
      type attribute of this struct.

    template<typename T >

  using to_real = remove_complex < T >
      to real is alias of remove complex
• template<typename T >
  using next precision = typename detail::next precision impl< T >::type
      Obtains the next type in the singly-linked precision list.

    template<typename T >

  using previous precision = next precision < T >
      Obtains the previous type in the singly-linked precision list.
• template<typename T >
  using reduce_precision = typename detail::reduce_precision_impl< T >::type
      Obtains the next type in the hierarchy with lower precision than T.

    template<typename T >

  using increase precision = typename detail::increase precision impl< T >::type
      Obtains the next type in the hierarchy with higher precision than T.

    template<tvpename... Ts>

  using highest_precision = typename detail::highest_precision_variadic< Ts... >::type
      Obtains the smallest arithmetic type that is able to store elements of all template parameter types exactly.

    template<typename T , size_type Limit = sizeof(uint16) * byte_size>

  using truncate_type = std::conditional_t< detail::type_size_impl< T >::value >=2 *Limit, typename detail↔
  ::truncate type impl< T >::type, T >
      Truncates the type by half (by dropping bits), but ensures that it is at least Limit bits wide.
• using size type = std::size t
      Integral type used for allocation quantities.
using int8 = std::int8_t
     8-bit signed integral type.
using int16 = std::int16 t
      16-bit signed integral type.
• using int32 = std::int32 t
      32-bit signed integral type.
• using int64 = std::int64_t
      64-bit signed integral type.
• using uint8 = std::uint8 t
      8-bit unsigned integral type.
• using uint16 = std::uint16_t
      16-bit unsigned integral type.
using uint32 = std::uint32_t
      32-bit unsigned integral type.
using uint64 = std::uint64 t
      64-bit unsigned integral type.
```

Unsigned integer type capable of holding a pointer to void.

• using float16 = half

Half precision floating point type.

• using float32 = float

Single precision floating point type.

• using float64 = double

Double precision floating point type.

• using full_precision = double

The most precise floating-point type.

• using default_precision = double

Precision used if no precision is explicitly specified.

Enumerations

enum log_propagation_mode { log_propagation_mode::never, log_propagation_mode::automatic }

How Logger events are propagated to their Executor.

· enum allocation mode

Specify the mode of allocation for CUDA/HIP GPUs.

enum layout_type { layout_type::array, layout_type::coordinate }

Specifies the layout type when writing data in matrix market format.

Functions

template < typename ValueType >
 ValueType reduce_add (const array < ValueType > &input_arr, const ValueType init_val=0)

Reduce (sum) the values in the array.

template<typename ValueType >

void reduce_add (const array< ValueType > &input_arr, array< ValueType > &result)

Reduce (sum) the values in the array.

template<typename ValueType >

array< ValueType > make_array_view (std::shared_ptr< const Executor > exec, size_type size, ValueType *data)

Helper function to create an array view deducing the value type.

template<typename ValueType >

detail::const_array_view < ValueType > make_const_array_view (std::shared_ptr< const Executor > exec, size_type size, const ValueType *data)

Helper function to create a const array view deducing the value type.

 $\bullet \ \ \text{template} {<} \text{size_type Dimensionality, typename DimensionType} >$

constexpr bool operator!= (const dim< Dimensionality, DimensionType > &x, const dim< Dimensionality, DimensionType > &y)

Checks if two dim objects are different.

 $\bullet \ \ \text{template}{<} \text{typename DimensionType}>$

 $constexpr \ dim < 2, \ Dimension Type > transpose \ (const \ dim < 2, \ Dimension Type > \& dimensions) \ no except$

Returns a dim<2> object with its dimensions swapped.

• template<typename T >

constexpr bool is_complex ()

Checks if T is a complex type.

• template<typename T >

constexpr bool is_complex_or_scalar ()

Checks if T is a complex/scalar type.

```
• template<typename T >
  constexpr reduce_precision < T > round_down (T val)
      Reduces the precision of the input parameter.
• template<typename T >
  constexpr increase precision< T > round up (T val)
      Increases the precision of the input parameter.

    constexpr int64 ceildiv (int64 num, int64 den)

      Performs integer division with rounding up.
• template<typename T >
  constexpr T zero ()
     Returns the additive identity for T.
• template<typename T >
  constexpr T zero (const T &)
      Returns the additive identity for T.

    template<typename T >

  constexpr T one ()
      Returns the multiplicative identity for T.

    template<typename T >

  constexpr T one (const T &)
      Returns the multiplicative identity for T.
• template<typename T >
  constexpr bool is_zero (T value)
      Returns true if and only if the given value is zero.
template<typename T >
  constexpr bool is_nonzero (T value)
      Returns true if and only if the given value is not zero.
template<typename T >
  constexpr T max (const T &x, const T &y)
      Returns the larger of the arguments.
• template<typename T >
  constexpr T min (const T &x, const T &y)
      Returns the smaller of the arguments.
• template<typename T >
  constexpr auto real (const T &x)
     Returns the real part of the object.
• template<typename T >
  constexpr auto imag (const T &x)
      Returns the imaginary part of the object.

    template<typename T >

  constexpr auto conj (const T &x)
      Returns the conjugate of an object.

    template<typename T >

  constexpr auto squared_norm (const T &x) -> decltype(real(conj(x) *x))
      Returns the squared norm of the object.
• template<typename T >
  constexpr xstd::enable_if_t<!is_complex_s< T >::value, T > abs (const T &x)
      Returns the absolute value of the object.
• template<typename T >
  constexpr T pi ()
      Returns the value of pi.
• template<typename T >
  constexpr std::complex < remove complex < T >> unit root (int64 n, int64 k=1)
      Returns the value of exp(2 * pi * i * k / n), i.e.
```

```
• template<typename T >
  constexpr uint32 get_significant_bit (const T &n, uint32 hint=0u) noexcept
      Returns the position of the most significant bit of the number.
template<typename T >
  constexpr T get_superior_power (const T &base, const T &limit, const T &hint=T{1}) noexcept
      Returns the smallest power of base not smaller than limit.
• template<typename T >
  std::enable if t<!is complex s< T>::value, bool > is finite (const T &value)
      Checks if a floating point number is finite, meaning it is neither +/- infinity nor NaN.

    template<typename T >

  std::enable_if_t< is_complex_s< T >::value, bool > is_finite (const T &value)
      Checks if all components of a complex value are finite, meaning they are neither +/- infinity nor NaN.

    template<typename T >

  T safe divide (T a, T b)
      Computes the quotient of the given parameters, guarding against division by zero.

    template<typename T >

  std::enable_if_t<!is_complex_s< T >::value, bool > is_nan (const T &value)
      Checks if a floating point number is NaN.
template<typename T >
  std::enable_if_t< is_complex_s< T >::value, bool > is_nan (const T &value)
      Checks if any component of a complex value is NaN.

    template<typename T >

  constexpr std::enable_if_t<!is_complex_s< T >::value, T > nan ()
      Returns a quiet NaN of the given type.

    template<typename T >

  constexpr std::enable_if_t< is_complex_s< T >::value, T > nan ()
      Returns a complex with both components quiet NaN.
• template<typename ValueType = default_precision, typename IndexType = int32>
  matrix data < ValueType, IndexType > read raw (std::istream &is)
      Reads a matrix stored in matrix market format from an input stream.

    template<typename ValueType = default_precision, typename IndexType = int32>

  matrix_data< ValueType, IndexType > read_binary_raw (std::istream &is)
      Reads a matrix stored in Ginkgo's binary matrix format from an input stream.
• template<typename ValueType = default_precision, typename IndexType = int32>
  matrix data < ValueType, IndexType > read generic raw (std::istream &is)
      Reads a matrix stored in either binary or matrix market format from an input stream.

    template<typename ValueType , typename IndexType >

  void write raw (std::ostream &os, const matrix data < ValueType, IndexType > &data, layout type
  layout=layout type::array)
      Writes a matrix_data structure to a stream in matrix market format.

    template<typename ValueType , typename IndexType >

  void write binary raw (std::ostream &os, const matrix data< ValueType, IndexType > &data)
      Writes a matrix_data structure to a stream in binary format.
• template<typename MatrixType , typename StreamType , typename... MatrixArgs>
  std::unique_ptr< MatrixType > read (StreamType &&is, MatrixArgs &&... args)
      Reads a matrix stored in matrix market format from an input stream.
• template<typename MatrixType , typename StreamType , typename... MatrixArgs>
  std::unique_ptr< MatrixType > read_binary (StreamType &&is, MatrixArgs &&... args)
      Reads a matrix stored in binary format from an input stream.
• template<typename MatrixType , typename StreamType , typename... MatrixArgs>
  std::unique_ptr< MatrixType > read_generic (StreamType &&is, MatrixArgs &&... args)
      Reads a matrix stored either in binary or matrix market format from an input stream.
```

template<typename MatrixPtrType, typename StreamType >
 void write (StreamType &&os, MatrixPtrType &&matrix, layout_type layout=detail::mtx_io_traits< std
 ::remove_cv_t< detail::pointee< MatrixPtrType >>>::default_layout)

Writes a matrix into an output stream in matrix market format.

template < typename MatrixPtrType, typename StreamType > void write binary (StreamType &&os, MatrixPtrType &&matrix)

Writes a matrix into an output stream in binary format.

template<typename R, typename T >
 std::unique_ptr< R, std::function< void(R *)> > copy_and_convert_to (std::shared_ptr< const Executor >
 exec, T *obj)

Converts the object to R and places it on Executor exec.

template<typename R, typename T >
 std::unique_ptr< const R, std::function< void(const R *)> > copy_and_convert_to (std::shared_ptr< const Executor > exec, const T *obj)

Converts the object to R and places it on Executor exec.

template<typename R, typename T >
 std::shared_ptr< R > copy_and_convert_to (std::shared_ptr< const Executor > exec, std::shared_ptr< T >
 obj)

Converts the object to R and places it on Executor exec.

- template<typename R, typename T >
 std::shared_ptr< const R > copy_and_convert_to (std::shared_ptr< const Executor > exec, std::shared_
 ptr< const T > obj)
- template<typename ValueType, typename Ptr >
 detail::temporary_conversion< std::conditional_t< std::is_const< detail::pointee< Ptr > >::value, const
 matrix::Dense< ValueType > > make_temporary_conversion (Ptr &&matrix)

Convert the given LinOp from matrix::Dense<...> to matrix::Dense< Value Type>.

template < typename ValueType , typename Function , typename... Args > void precision_dispatch (Function fn, Args *... linops)

Calls the given function with each given argument LinOp temporarily converted into matrix::Dense<ValueType> as parameters.

template<typename ValueType , typename Function >
 void precision_dispatch_real_complex (Function fn, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix::Dense< ValueType>* as parameters.

template<typename ValueType, typename Function >
void precision dispatch real complex (Function fn, const LinOp *alpha, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix:: Dense < ValueType > * as parameters.

template < typename ValueType, typename Function >
 void precision_dispatch_real_complex (Function fn, const LinOp *alpha, const LinOp *in, const LinOp *beta,
 LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix::Dense< Value Type>* as parameters.

template<typename ValueType , typename Function > void mixed_precision_dispatch (Function fn, const LinOp *in, LinOp *out)

Calls the given function with each given argument LinOp converted into matrix::Dense< ValueType> as parameters.

template<typename ValueType , typename Function , std::enable_if_t< is_complex< ValueType >()> * = nullptr> void mixed_precision_dispatch_real_complex (Function fn, const LinOp *in, LinOp *out)

 $\textit{Calls the given function with the given LinOps cast to their dynamic type \textit{matrix}::} \textbf{Dense} < \textit{ValueType} > * \textit{ as parameters}.$

Creates a temporary_clone.

Creates a uninitialized temporary_clone that will be copied back to the input afterwards.

• constexpr bool operator== (precision_reduction x, precision_reduction y) noexcept

Checks if two precision_reduction encodings are equal.

• constexpr bool operator!= (precision_reduction x, precision_reduction y) noexcept

Checks if two precision_reduction encodings are different.

template<typename IndexType >

constexpr IndexType invalid_index ()

Value for an invalid signed index type.

• template<typename Pointer >

detail::cloned_type< Pointer > clone (const Pointer &p)

Creates a unique clone of the object pointed to by p.

template<typename Pointer >

detail::cloned_type< Pointer > clone (std::shared_ptr< const Executor > exec, const Pointer &p)

Creates a unique clone of the object pointed to by p on Executor exec.

• template<typename OwningPointer >

detail::shared_type< OwningPointer > share (OwningPointer &&p)

Marks the object pointed to by p as shared.

• template<typename OwningPointer >

std::remove reference < OwningPointer >::type && give (OwningPointer &&p)

Marks that the object pointed to by p can be given to the callee.

• template<typename Pointer >

std::enable_if< detail::have_ownership_s< Pointer >::value, detail::pointee< Pointer > * >::type lend (const Pointer &p)

Returns a non-owning (plain) pointer to the object pointed to by p.

template<typename Pointer >

std::enable_if<!detail::have_ownership_s< Pointer >::value, detail::pointee< Pointer > * >::type lend (const Pointer &p)

Returns a non-owning (plain) pointer to the object pointed to by p.

• template<typename T , typename U >

```
std::decay_t< T > * as (U *obj)
```

Performs polymorphic type conversion.

• template<typename T , typename U >

const std::decay_t< T > * as (const U *obj)

Performs polymorphic type conversion.

• template<typename T , typename U >

```
std::decay_t < T > * as (ptr_param < U > obj)
```

Performs polymorphic type conversion on a ptr_param.

• template<typename T , typename U >

```
const std::decay_t < T > * as (ptr_param < const U > obj)
```

Performs polymorphic type conversion.

• template<typename T , typename U >

```
std::unique\_ptr < std::decay\_t < T >> \underbrace{as} (std::unique\_ptr < U > \&\&obj)
```

Performs polymorphic type conversion of a unique ptr.

• template<typename T , typename U >

```
std::shared\_ptr < std::decay\_t < T >> \underbrace{as} (std::shared\_ptr < U > obj)
```

Performs polymorphic type conversion of a shared_ptr.

template<typename T , typename U >

```
std::shared_ptr< const std::decay_t< T >> as (std::shared_ptr< const U > obj)
```

Performs polymorphic type conversion of a shared_ptr.

std::ostream & operator<< (std::ostream &os, const version &ver)

Prints version information to a stream.

std::ostream & operator<< (std::ostream &os, const version_info &ver_info)

Prints library version information in human-readable format to a stream.

template < template < typename, typename > class MatrixType, typename... Args > auto with _matrix_type (Args &&... create_args)

This function returns a type that delays a call to MatrixType::create.

 $\bullet \ \ \text{template}{<} \text{typename VecPtr}>$

std::unique_ptr< matrix::Dense< typename detail::pointee< VecPtr >::value_type >> make_dense_view (VecPtr &vector)

Creates a view of a given Dense vector.

• template<typename VecPtr>

std::unique_ptr< const matrix::Dense< typename detail::pointee< VecPtr >::value_type > > make_const_dense_view (VecPtr &&vector)

Creates a view of a given Dense vector.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > initialize (size_type stride, std::initializer_list< typename Matrix::value_type > vals,
 std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a column-vector.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > initialize (std::initializer_list< typename Matrix::value_type > vals, std::shared_
 ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a column-vector.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > initialize (size_type stride, std::initializer_list< std::initializer_list< typename
 Matrix::value_type >> vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a matrix.

template<typename Matrix , typename... TArgs>
 std::unique_ptr< Matrix > initialize (std::initializer_list< std::initializer_list< typename Matrix::value_type >>
 vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)

Creates and initializes a matrix.

• bool operator== (const stopping_status &x, const stopping_status &y) noexcept

Checks if two stopping statuses are equivalent.

• bool operator!= (const stopping status &x, const stopping status &y) noexcept

Checks if two stopping statuses are different.

Variables

constexpr size_type byte_size = CHAR_BIT

Number of bits in a byte.

42.1.1 Detailed Description

The Ginkgo namespace.

42.1.2 Typedef Documentation

42.1.2.1 highest_precision

```
template<typename... Ts>
using gko::highest_precision = typedef typename detail::highest_precision_variadic<Ts...>←
::type
```

Obtains the smallest arithmetic type that is able to store elements of all template parameter types exactly.

All template type parameters need to be either real or complex types, mixing them is not possible.

Formally, it computes a right-fold over the type list, with the highest precision of a pair of real arithmetic types T1, T2 computed as $decltype(T1\{\} + T2\{\})$, or $std::complex<highest_precision<remove_{} complex<T1>$, $remove_complex<T2>>> for complex types$.

42.1.2.2 is complex or scalar s

```
template<typename T >
using gko::is_complex_or_scalar_s = typedef detail::is_complex_or_scalar_impl<T>
```

Allows to check if T is a complex or scalar value during compile time by accessing the value attribute of this struct.

If value is true, T is a complex/scalar type, if it is false, T is not a complex/scalar type.

Template Parameters

T type to check

42.1.2.3 is_complex_s

```
template<typename T >
using gko::is_complex_s = typedef detail::is_complex_impl<T>
```

Allows to check if T is a complex value during compile time by accessing the value attribute of this struct.

If value is true, T is a complex type, if it is false, T is not a complex type.

Template Parameters

T type to check

42.1.2.4 previous_precision

```
template<typename T >
using gko::previous_precision = typedef next_precision<T>
```

Obtains the previous type in the singly-linked precision list.

Note

Currently our lists contains only two elements, so this is the same as next_precision.

42.1.2.5 remove_complex

```
template<typename T >
using gko::remove_complex = typedef typename detail::remove_complex_s<T>::type
```

Obtain the type which removed the complex of complex/scalar type or the template parameter of class by accessing the type attribute of this struct.

Template Parameters

T | type to remove complex

Note

remove_complex<class> can not be used in friend class declaration.

42.1.2.6 to_complex

```
template<typename T >
using gko::to_complex = typedef typename detail::to_complex_s<T>::type
```

Obtain the type which adds the complex of complex/scalar type or the template parameter of class by accessing the type attribute of this struct.

Template Parameters

T type to complex_type

Note

to_complex < class > can not be used in friend class declaration. the followings are the error message from different combination. friend to_complex < Csr >; error: can not recognize it is class correctly. friend class to_complex < Csr >; error: using alias template specialization friend class to_complex_s < Csr < ValueType, \leftarrow IndexType >>::type; error: can not recognize it is class correctly.

42.1.2.7 to_real

```
template<typename T >
using gko::to_real = typedef remove_complex<T>
```

to_real is alias of remove_complex

Template Parameters

```
T type to real
```

42.1.3 Enumeration Type Documentation

42.1.3.1 allocation_mode

```
enum gko::allocation_mode [strong]
```

Specify the mode of allocation for CUDA/HIP GPUs.

device allocates memory on the device and Unified Memory model is not used.

unified_global allocates memory on the device, but is accessible by the host through the Unified memory model.

unified_host allocates memory on the host and it is not available on devices which do not have concurrent acesses switched on, but this access can be explictly switched on, when necessary.

89 { device, unified_global, unified_host };

42.1.3.2 layout_type

```
enum gko::layout_type [strong]
```

Specifies the layout type when writing data in matrix market format.

Enumerator

array	The matrix should be written as dense matrix in column-major order.
coordinate	The matrix should be written as a sparse matrix in coordinate format.

```
121 {
125 array,
129 coordinate
130 };
```

42.1.3.3 log_propagation_mode

```
enum gko::log_propagation_mode [strong]
```

How Logger events are propagated to their Executor.

Enumerator

never	Events only get reported at loggers attached to the triggering object. (Except for allocation/free, copy and Operations, since they happen at the Executor).
automatic	Events get reported to loggers attached to the triggering object and propagating loggers (Logger::needs_propagation() return true) attached to its Executor.

42.1.4 Function Documentation

42.1.4.1 abs()

Returns the absolute value of the object.

Template Parameters

```
T the type of the object
```

Parameters

```
x the object
```

Returns

```
x >= zero<T>() ? x:-x;

1112 {
    1113          return x >= zero<T>() ? x: -x;
    1114 }
```

Referenced by is_finite().

42.1.4.2 as() [1/7]

Performs polymorphic type conversion.

This is the constant version of the function.

Template Parameters

Τ	requested result type
U	static type of the passed object

Parameters

```
obj the object which should be converted
```

Returns

If successful, returns a pointer to the subtype, otherwise throws NotSupported.

42.1.4.3 as() [2/7]

```
template<typename T , typename U > const std::decay_t<T>* gko::as ( ptr\_param< const \ U > obj \ ) \quad [inline]
```

Performs polymorphic type conversion.

This is the constant version of the function.

Template Parameters

Τ	requested result type
U	static type of the passed object

Parameters

obj	the object which should be converted
-----	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws NotSupported.

References gko::ptr_param< T >::get().

42.1.4.4 as() [3/7]

```
template<typename T , typename U > std::decay_t<T>* gko::as ( ptr\_param< U > obj ) \quad [inline]
```

Performs polymorphic type conversion on a ptr_param.

Template Parameters

T	requested result type
U	static type of the passed object

Parameters

Returns

If successful, returns a pointer to the subtype, otherwise throws NotSupported.

References gko::ptr_param< T >::get().

42.1.4.5 as() [4/7]

Performs polymorphic type conversion of a shared_ptr.

This is the constant version of the function.

Template Parameters

T	requested result type
U	static type of the passed object

Parameters

obj the shared_ptr to the object which should be converted.

Returns

If successful, returns a shared_ptr to the subtype, otherwise throws NotSupported. This pointer shares ownership with the input pointer.

42.1.4.6 as() [5/7]

```
template<typename T , typename U >  std::shared\_ptr < std::decay\_t < T > gko::as ( \\ std::shared\_ptr < U > obj ) [inline]
```

Performs polymorphic type conversion of a shared_ptr.

Template Parameters

T	requested result type
U	static type of the passed object

Parameters

obj the shared_ptr to the object which should be converted.

Returns

If successful, returns a shared_ptr to the subtype, otherwise throws NotSupported. This pointer shares ownership with the input pointer.

42.1.4.7 as() [6/7]

Performs polymorphic type conversion of a unique_ptr.

Template Parameters

	Т	requested result type
ſ	U	static type of the passed object

Parameters

obj the unique_ptr to the object which should be converted. If successful, it will be reset to a nullptr.

Returns

If successful, returns a unique_ptr to the subtype, otherwise throws NotSupported.

42.1.4.8 as() [7/7]

Performs polymorphic type conversion.

Template Parameters

T	requested result type
U	static type of the passed object

Parameters

obj	the object which should be converted
-----	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws NotSupported.

Referenced by gko::preconditioner::lsai< lsaiType, ValueType, IndexType >::get_approximate_inverse().

42.1.4.9 ceildiv()

Performs integer division with rounding up.

Parameters

num	numerator
den	denominator

Returns

returns the ceiled quotient.

Referenced by gko::matrix::Csr< ValueType, IndexType >::load_balance::clac_size(), gko::preconditioner::block \leftarrow _interleaved_storage_scheme< index_type >::compute_storage_space(), and gko::matrix::Csr< ValueType, IndexType >::load_balance::process().

42.1.4.10 clone() [1/2]

Creates a unique clone of the object pointed to by p.

The pointee (i.e. *p) needs to have a clone method that returns a std::unique_ptr in order for this method to work.

Template Parameters

Pointer	type of pointer to the object (plain or smart pointer)
---------	--

Parameters

```
p a pointer to the object
```

Note

The difference between this function and directly calling LinOp::clone() is that this one preserves the static type of the object.

Referenced by gko::preconditioner::lc< LSolverType, IndexType >::operator=(), gko::preconditioner::llu< L \leftarrow SolverType, USolverType, ReverseApply, IndexType >::operator=(), gko::solver::EnablePreconditionable< Bicg< ValueType > >::set_preconditioner(), and gko::solver::EnableIterativeBase< Bicg< ValueType > >::set_stop_ \leftarrow criterion_factory().

42.1.4.11 clone() [2/2]

Creates a unique clone of the object pointed to by p on Executor exec.

The pointee (i.e. *p) needs to have a clone method that takes an executor and returns a std::unique_ptr in order for this method to work.

Template Parameters

Pointer type of pointer to the object (plain or smart pointer)
--

exec	the executor where the cloned object should be stored
р	a pointer to the object

Note

The difference between this function and directly calling LinOp::clone() is that this one preserves the static type of the object.

42.1.4.12 conj()

Returns the conjugate of an object.

Parameters

```
x the number to conjugate
```

Returns

conjugate of the object (by default, the object itself)

Referenced by squared_norm().

42.1.4.13 copy_and_convert_to() [1/4]

Converts the object to R and places it on Executor exec.

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

	the type to which the object should be converted	
Т	the type of the input object	

ex	ec	the executor where the result should be placed
ob)j	the object that should be converted

Returns

a unique pointer (with dynamically bound deleter) to the converted object

Note

This is a version of the function which adds the const qualifier to the result if the input had the same qualifier.

```
608 {
609     return detail::copy_and_convert_to_impl<const R>(std::move(exec), obj);
610 }
```

42.1.4.14 copy_and_convert_to() [2/4]

This is the version that takes in the std::shared_ptr and returns a std::shared_ptr

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

R	the type to which the object should be converted
T	the type of the input object

Parameters

exec	the executor where the result should be placed
obj	the object that should be converted

Returns

a shared pointer to the converted object

Note

This is a version of the function which adds the const qualifier to the result if the input had the same qualifier.

42.1.4.15 copy_and_convert_to() [3/4]

```
template<typename R , typename T >
std::shared_ptr<R> gko::copy_and_convert_to (
```

```
std::shared_ptr< const Executor > exec,
std::shared_ptr< T > obj )
```

Converts the object to R and places it on Executor exec.

This is the version that takes in the std::shared_ptr and returns a std::shared_ptr

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

	the type to which the object should be converted
T	the type of the input object

Parameters

exec	the executor where the result should be placed
obj	the object that should be converted

Returns

a shared pointer to the converted object

42.1.4.16 copy_and_convert_to() [4/4]

Converts the object to R and places it on Executor exec.

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

R	the type to which the object should be converted
T	the type of the input object

exec	the executor where the result should be placed
obj	the object that should be converted

Returns

a unique pointer (with dynamically bound deleter) to the converted object

42.1.4.17 get_significant_bit()

Returns the position of the most significant bit of the number.

This is the same as the rounded down base-2 logarithm of the number.

Template Parameters

T a numeric type supporting bit shift and comparison

Parameters

n	a number
hint	a lower bound for the position o the significant bit

Returns

maximum of hint and the significant bit position of n

42.1.4.18 get_superior_power()

Returns the smallest power of base not smaller than limit.

Template Parameters

T a numeric type supporting multiplication and comparison

base	the base of the power to be returned

Parameters

limit	the lower limit on the size of the power returned
hint	a lower bound on the result, has to be a power of base

Returns

the smallest power of base not smaller than limit

42.1.4.19 give()

Marks that the object pointed to by p can be given to the callee.

Effectively calls std::move(p).

Template Parameters

OwningFointer type of pointer with ownership to the object that to be a smart pointe	OwningPointer	type of pointer with ownership to the object (has to be a smart pointer)
--	---------------	--

Parameters

```
p a pointer to the object
```

Note

The original pointer p becomes invalid after this call.

42.1.4.20 imag()

Returns the imaginary part of the object.

Template Parameters

Parameters

```
x the object
```

Returns

imaginary part of the object (by default, zero<T>())

42.1.4.21 is_complex()

```
template<typename T >
constexpr bool gko::is_complex ( ) [inline], [constexpr]
```

Checks if T is a complex type.

Template Parameters

```
T type to check
```

Returns

true if T is a complex type, false otherwise

42.1.4.22 is_complex_or_scalar()

```
template<typename T >
constexpr bool gko::is_complex_or_scalar ( ) [inline], [constexpr]
```

Checks if T is a complex/scalar type.

Template Parameters

```
T type to check
```

Returns

true if T is a complex/scalar type, false otherwise

42.1.4.23 is_finite() [1/2]

Checks if a floating point number is finite, meaning it is neither +/- infinity nor NaN.

Template Parameters

```
T type of the value to check
```

Parameters

value value to check

Returns

true if the value is finite, meaning it are neither +/- infinity nor NaN.

References abs().

Referenced by is_finite().

42.1.4.24 is_finite() [2/2]

Checks if all components of a complex value are finite, meaning they are neither +/- infinity nor NaN.

Template Parameters

T complex type of the value to check

Parameters

value	complex value to check
-------	------------------------

Returns

true if both components of the given value are finite, meaning they are neither +/- infinity nor NaN.

References is_finite().

42.1.4.25 is_nan() [1/2]

Checks if a floating point number is NaN.

Template Parameters

```
T type of the value to check
```

Parameters

```
value value to check
```

Returns

true if the value is NaN.

42.1.4.26 is_nan() [2/2]

Checks if any component of a complex value is NaN.

Template Parameters

```
T | complex type of the value to check
```

Parameters

```
value | complex value to check
```

Returns

true if any component of the given value is NaN.

42.1.4.27 is_nonzero()

Returns true if and only if the given value is not zero.

Template Parameters

T the type of the value

Parameters

value	the given value
-------	-----------------

Returns

```
true iff the given value is not zero, i.e. value != zero < T > ()
```

Referenced by gko::matrix_data< ValueType, IndexType >::diag().

42.1.4.28 is_zero()

Returns true if and only if the given value is zero.

Template Parameters

```
T the type of the value
```

Parameters

```
value the given value
```

Returns

```
true iff the given value is zero, i.e. value == zero<T>()
```

Referenced by gko::matrix_data< ValueType, IndexType >::remove_zeros().

42.1.4.29 lend() [1/2]

Returns a non-owning (plain) pointer to the object pointed to by p.

Template Parameters

Pointer	type of pointer to the object (plain or smart pointer)
---------	--

Parameters

```
p a pointer to the object
```

Note

This is the overload for owning (smart) pointers, that behaves the same as calling .get() on the smart pointer.

42.1.4.30 lend() [2/2]

Returns a non-owning (plain) pointer to the object pointed to by p.

Template Parameters

Pointer	type of pointer to the object (plain or smart pointer)	
---------	--	--

Parameters

```
p a pointer to the object
```

Note

This is the overload for non-owning (plain) pointers, that just returns p.

42.1.4.31 make_array_view()

Helper function to create an array view deducing the value type.

е	хес	the executor on which the array resides
s	ize	the number of elements for the array
a	lata	the pointer to the array we create a view on.

Template Parameters

ValueType the type of the array elements
--

Returns

```
array<ValueType>::view(exec, size, data)
762 {
763     return array<ValueType>::view(exec, size, data);
764 }
```

References make_array_view().

Referenced by make_array_view().

42.1.4.32 make_const_array_view()

Helper function to create a const array view deducing the value type.

Parameters

exec	the executor on which the array resides
size	the number of elements for the array
data	the pointer to the array we create a view on.

Template Parameters

ValueType the type of the array elements
--

Returns

```
array<ValueType>::const_view(exec, size, data)
```

References make_const_array_view().

Referenced by make_const_array_view().

42.1.4.33 make_const_dense_view()

Creates a view of a given Dense vector.

Template Parameters

<i>VecPtr</i>	a (smart or raw) pointer to the vector.
---------------	---

Parameters

vector	the vector on which to create the view
--------	--

References gko::matrix::Dense< ValueType >::create_const_view_of().

42.1.4.34 make_dense_view()

```
\label{template} $$ $ \end{template} $$ $ \e
```

Creates a view of a given Dense vector.

Template Parameters

/ecPtr a (smart or raw) pointer to the vector.
--

Parameters

References gko::matrix::Dense< ValueType >::create_view_of().

42.1.4.35 make_temporary_clone()

Creates a temporary_clone.

This is a helper function which avoids the need to explicitly specify the type of the object, as would be the case if using the constructor of temporary_clone.

ſ	ехес	the executor where the clone will be created
	ptr	a pointer to the object of which the clone will be created

Template Parameters

Ptr the (raw or smart) pointer type to be temporarily cloned

```
209 {
210    using T = detail::pointee<Ptr>
211    return detail::temporary_clone<T>(std::move(exec), std::forward<Ptr>
(ptr));
212 }
```

Referenced by gko::ScaledIdentityAddable::add_scaled_identity(), gko::matrix::Coo< ValueType, IndexType >
::apply2(), gko::matrix::Csr< ValueType, IndexType >::inv_scale(), and gko::matrix::Csr< ValueType, IndexType
>::scale().

42.1.4.36 make_temporary_conversion()

Convert the given LinOp from matrix::Dense<...> to matrix::Dense<ValueType>.

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by next—precision<...> starting from the ValueType template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

matrix the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type matrix::Dense<ValueType>, otherwise it will be converted to this type if possible.

Returns

a detail::temporary conversion pointing to the (potentially converted) object.

Exceptions

NotSupported if the input matrix cannot be converted to matrix::Dense<ValueType>

Template Parameters

ValueType the value type into whose associated matrix::Dense type to convert the input LinOp.

```
76 {
77     using Pointee = detail::pointee<Ptr>78     using Dense = matrix::Dense<ValueType>;
79     using NextDense = matrix::Dense<next_precision<ValueType»;
80     using MaybeConstDense =
81         std::conditional_t<std::is_const<Pointee>::value, const Dense, Dense>;
82     auto result = detail::temporary_conversion
83     MaybeConstDense>::template create<NextDense>(matrix);
```

```
84  if (!result) {
85      GKO_NOT_SUPPORTED(*matrix);
86  }
87  return result;
88 }
```

42.1.4.37 make_temporary_output_clone()

Creates a uninitialized temporary_clone that will be copied back to the input afterwards.

It can be used for output parameters to avoid an unnecessary copy in make_temporary_clone.

This is a helper function which avoids the need to explicitly specify the type of the object, as would be the case if using the constructor of temporary_clone.

Parameters

exec	the executor where the uninitialized clone will be created
ptr	a pointer to the object of which the clone will be created

Template Parameters

```
Ptr the (raw or smart) pointer type to be temporarily cloned
```

42.1.4.38 max()

Returns the larger of the arguments.

Template Parameters

T	type of the arguments

Χ	first argument
У	second argument

Returns

```
x >= y ? x : y
```

42.1.4.39 min()

Returns the smaller of the arguments.

Template Parameters

```
T type of the arguments
```

Parameters

Х	first argument
У	second argument

Returns

```
x \le y ? x : y
```

Referenced by gko::matrix::Csr< ValueType, IndexType >::load_balance::clac_size().

42.1.4.40 mixed precision dispatch()

Calls the given function with each given argument LinOp converted into matrix::Dense<ValueType> as parameters.

If GINKGO_MIXED_PRECISION is defined, this means that the function will be called with its dynamic type as a static type, so the (templated/generic) function will be instantiated with all pairs of Dense<ValueType> and Dense<next_precision<ValueType>> parameter types, and the appropriate overload will be called based on the dynamic type of the parameter.

If GINKGO_MIXED_PRECISION is not defined, it will behave exactly like precision_dispatch.

Parameters

fn	the given function. It will be called with one const and one non-const matrix::Dense<> parameter based on the dynamic type of the inputs (GINKGO_MIXED_PRECISION) or of type matrix::Dense <valuetype> (no GINKGO_MIXED_PRECISION).</valuetype>
in	The first parameter to be cast (GINKGO_MIXED_PRECISION) or converted (no GINKGO_MIXED_PRECISION) and used to call fn.
out	The second parameter to be cast (GINKGO_MIXED_PRECISION) or converted (no GINKGO_MIXED_PRECISION) and used to call fn.

Template Parameters

ValueType	the value type to use for the parameters of fn (no GINKGO_MIXED_PRECISION). With GINKGO_MIXED_PRECISION enabled, it only matters whether this type is complex or real.
Function	the function pointer, lambda or other functor type to call with the converted arguments.

42.1.4.41 mixed_precision_dispatch_real_complex()

Calls the given function with the given LinOps cast to their dynamic type matrix::Dense<ValueType>* as parameters.

If ValueType is real and both in and out are complex, uses matrix::Dense::get_real_view() to convert them into real matrices after precision conversion.

See also

mixed_precision_dispatch()

42.1.4.42 nan() [1/2]

```
template<typename T >
constexpr std::enable_if_t<!is_complex_s<T>::value, T> gko::nan ( ) [inline], [constexpr]
```

Returns a quiet NaN of the given type.

Template Parameters

T the type of the object

Returns

NaN.

Referenced by nan().

42.1.4.43 nan() [2/2]

```
template<typename T >
constexpr std::enable_if_t<is_complex_s<T>::value, T> gko::nan ( ) [inline], [constexpr]
```

Returns a complex with both components quiet NaN.

Template Parameters

```
T the type of the object
```

Returns

complex{NaN, NaN}.

References nan().

42.1.4.44 one() [1/2]

```
template<typename T >
constexpr T gko::one ( ) [inline], [constexpr]
```

Returns the multiplicative identity for T.

Returns

the multiplicative identity for T

Referenced by unit_root().

42.1.4.45 one() [2/2]

Returns the multiplicative identity for T.

Returns

the multiplicative identity for T

Note

This version takes an unused reference argument to avoid complicated calls like one < decltype(x) > (). Instead, it allows one(x).

42.1.4.46 operator"!=() [1/3]

Checks if two dim objects are different.

Template Parameters

Dimensionality	number of dimensions of the dim objects
DimensionType	datatype used to represent each dimension

Parameters

X	first object
У	second object

Returns

```
! (x == y)

261 {
262 return ! (x == y);
263 }
```

42.1.4.47 operator"!=() [2/3]

Checks if two stopping statuses are different.

Parameters

X	a stopping status
У	a stopping status

Returns

```
true if and only if ! (x == y)

179 {
180     return x.data_ != y.data_;
181 }
```

42.1.4.48 operator"!=() [3/3]

Checks if two precision_reduction encodings are different.

Parameters

X	an encoding
У	an encoding

Returns

true if and only if x and y are different encodings.

```
397 {
398     using st = precision_reduction::storage_type;
399     return static_cast<st>(x) != static_cast<st>(y);
400 }
```

42.1.4.49 operator << () [1/2]

Prints version information to a stream.

Parameters

os	output stream
ver	version structure

Returns

References gko::version::major, gko::version::minor, gko::version::patch, and gko::version::tag.

42.1.4.50 operator << () [2/2]

Prints library version information in human-readable format to a stream.

Parameters

os	output stream
ver_info	version information

Returns

os

42.1.4.51 operator==() [1/2]

Checks if two stopping statuses are equivalent.

Parameters

X	a stopping status
У	a stopping status

Returns

true if and only if both \boldsymbol{x} and \boldsymbol{y} have the same mask and converged and finalized state

42.1.4.52 operator==() [2/2]

Checks if two precision_reduction encodings are equal.

Parameters

Х	an encoding
У	an encoding

Returns

true if and only if \boldsymbol{x} and \boldsymbol{y} are the same encodings

42.1.4.53 pi()

```
template<typename T >
constexpr T gko::pi ( ) [inline], [constexpr]
```

Returns the value of pi.

Template Parameters

```
T \mid the value type to return
```

42.1.4.54 precision_dispatch()

Calls the given function with each given argument LinOp temporarily converted into matrix::Dense<ValueType> as parameters.

Parameters

fn	the given function. It will be passed one (potentially const) matrix::Dense <valuetype>* parameter per parameter in the parameter pack linops.</valuetype>
linops	the given arguments to be converted and passed on to fn.

Template Parameters

ſ	ValueType	the value type to use for the parameters of fn.
	Function	the function pointer, lambda or other functor type to call with the converted arguments.
	Args	the argument type list.

42.1.4.55 precision_dispatch_real_complex() [1/3]

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType>* as parameters.

If ValueType is real and both in and out are complex, uses matrix::Dense::get_real_view() to convert them into real matrices after precision conversion.

See also

```
precision_dispatch()
```

42.1.4.56 precision_dispatch_real_complex() [2/3]

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType>* as parameters.

If ValueType is real and both in and out are complex, uses matrix::Dense::get_real_view() to convert them into real matrices after precision conversion.

See also

```
precision_dispatch()
```

42.1.4.57 precision_dispatch_real_complex() [3/3]

```
template<typename ValueType , typename Function > void gko::precision_dispatch_real_complex (  Function \ fn, \\ const \ LinOp * in, \\ LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType>* as parameters.

If ValueType is real and both input vectors are complex, uses matrix::Dense::get_real_view() to convert them into real matrices after precision conversion.

See also

```
precision dispatch()
```

42.1.4.58 read()

Reads a matrix stored in matrix market format from an input stream.

Template Parameters

MatrixType	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from disk.
StreamType	type of stream used to write the data to
MatrixArgs	additional argument types passed to MatrixType constructor

Parameters

is	input stream from which to read the data
args	additional arguments passed to MatrixType constructor

Returns

A MatrixType LinOp filled with data from filename

References read_raw().

Referenced by gko::ReadableFromMatrixData< ValueType, int32 >::read().

42.1.4.59 read_binary()

Reads a matrix stored in binary format from an input stream.

Template Parameters

MatrixType	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from disk.	
StreamType	type of stream used to write the data to	
MatrixArgs	additional argument types passed to MatrixType constructor	

Parameters

is	input stream from which to read the data	
args additional arguments passed to MatrixType cons		

Returns

A MatrixType LinOp filled with data from filename

References read_binary_raw().

42.1.4.60 read_binary_raw()

Reads a matrix stored in Ginkgo's binary matrix format from an input stream.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

The binary format has the following structure (in system endianness):

- A 32 byte header consisting of 4 uint64_t values: magic = GINKGO__: The highest two bytes stand for value and index type. value type: S (float), D (double), C (complex<float>), Z(complex<double>) index type: I (int32), L (int64) num_rows: Number of rows num_cols: Number of columns num_entries: Number of (row, column, value) tuples to follow
- 2. Following are num_entries blocks of size sizeof(IndexType) * 2 + sizeof(ValueType). Each consists of a row index stored as IndexType, followed by a column index stored as IndexType and a value stored as ValueType.

Template Parameters

ValueType	type of matrix values
IndexType	type of matrix indexes

Parameters

```
is input stream from which to read the data
```

Returns

A matrix_data structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using gko::read_binary instead.

Referenced by read_binary().

42.1.4.61 read_generic()

Reads a matrix stored either in binary or matrix market format from an input stream.

Template Parameters

MatrixType	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from	
StreamType	e type of stream used to write the data to	
MatrixArgs additional argument types passed to MatrixType constructor		

Parameters

is	input stream from which to read the data
args	additional arguments passed to MatrixType constructor

Returns

A MatrixType LinOp filled with data from filename

References read_generic_raw().

42.1.4.62 read_generic_raw()

Reads a matrix stored in either binary or matrix market format from an input stream.

Template Parameters

ValueType	type of matrix values
IndexType	type of matrix indexes

Parameters

is input stream from which to read the data

Returns

A matrix_data structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using gko::read_generic instead.

Referenced by read_generic().

42.1.4.63 read_raw()

Reads a matrix stored in matrix market format from an input stream.

Template Parameters

ValueType	type of matrix values
IndexType	type of matrix indexes

Parameters

Returns

A matrix_data structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using gko::read instead.

Referenced by read().

42.1.4.64 real()

Returns the real part of the object.

Template Parameters

```
T type of the object
```

Parameters

```
x the object
```

Returns

real part of the object (by default, the object itself)

Referenced by squared_norm().

42.1.4.65 reduce_add() [1/2]

Reduce (sum) the values in the array.

Template Parameters

Parameters

in	input_arr	the input array to be reduced	
in,out	result	the reduced value. The result is written into the first entry and the value in the first	
		entry is used as the initial value for the reduce.	

42.1.4.66 reduce_add() [2/2]

Reduce (sum) the values in the array.

Template Parameters

The	type of the input data
-----	------------------------

Parameters

	in	input_arr	the input array to be reduced
ſ	in	init_val	the initial value

Returns

the reduced value

42.1.4.67 round_down()

Reduces the precision of the input parameter.

Template Parameters

```
T the original precision
```

Parameters

```
val the value to round down
```

Returns

the rounded down value

42.1.4.68 round_up()

Increases the precision of the input parameter.

Template Parameters

```
T the original precision
```

Parameters

```
val the value to round up
```

Returns

the rounded up value

42.1.4.69 safe_divide()

Computes the quotient of the given parameters, guarding against division by zero.

Template Parameters

T value type of the parameter

Parameters

а	the dividend
b	the divisor

Returns

the value of a / b if b is non-zero, zero otherwise.

42.1.4.70 share()

Marks the object pointed to by p as shared.

Effectively converts a pointer with ownership to std::shared_ptr.

Template Parameters

OwningPointer	type of pointer with ownership to the object (has to be a smart pointer)

Parameters

```
p a pointer to the object. It must be a temporary or explicitly marked movable (rvalue reference).
```

Note

The original pointer p becomes invalid after this call.

Referenced by gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose(), gko::preconditioner::lc< L \leftarrow SolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), gko::preconditioner::lc< LSolverType,

IndexType >::transpose(), and gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >
::transpose().

42.1.4.71 squared_norm()

Returns the squared norm of the object.

Template Parameters

```
T type of the object.
```

Returns

The squared norm of the object.

References conj(), and real().

42.1.4.72 transpose()

Returns a dim<2> object with its dimensions swapped.

Template Parameters

DimensionType	datatype used to represent each dimension

Parameters

dimensions original object

Returns

a dim<2> object with its dimensions swapped

Referenced by gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose(), gko::preconditioner::llu< L \leftarrow SolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), gko::preconditioner::lc< LSolverType, IndexType >::transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType > \leftarrow ::transpose().

42.1.4.73 unit_root()

Returns the value of exp(2 * pi * i * k / n), i.e.

an nth root of unity.

Parameters

n	the denominator of the argument
k	the numerator of the argument. Defaults to 1.

Template Parameters

```
T the corresponding real value type.
```

References one().

42.1.4.74 with_matrix_type()

This function returns a type that delays a call to MatrixType::create.

It can be used to set the used value and index type, as well as the executor at a later stage.

For example, the following code creates first a temporary object, which is then used later to construct an operator of the previously defined base type:

```
auto type = gko::with_matrix_type<gko::matrix::Csr>();
...
std::unique_ptr<LinOp> concrete_op
if(flag1){
    concrete_op = type.template create<double, int>(exec);
} else {
    concrete_op = type.template create<float, int>(exec);
}
```

Note

This is mainly a helper function to specify the local matrix type for a gko::experimental::distributed::Matrix more easily.

Template Parameters

MatrixType	A template type that accepts two types, the first one will be set to the value type, the second one to the index type.
Args	Types of the arguments passed to MatrixType::create.

Parameters

create_args arguments that will be forwarded to Matrix	Type::create
--	--------------

Returns

A type with a function create<value_type, index_type>(executor).

```
142 {
143     return detail::MatrixTypeBuilderFromValueAndIndex<MatrixType, Args...>{
144          std::forward_as_tuple(create_args...)};
145 }
```

42.1.4.75 write()

Writes a matrix into an output stream in matrix market format.

Template Parameters

MatrixPtrType	a (smart or raw) pointer to a WritableToMatrixData object providing data to be written.
StreamType	type of stream used to write the data to

Parameters

os	output stream where the data is to be written
matrix	the matrix to write
layout	the layout used in the output

References write_raw().

42.1.4.76 write_binary()

Writes a matrix into an output stream in binary format.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

Template Parameters

<i>MatrixPtrType</i>	a (smart or raw) pointer to a WritableToMatrixData object providing data to be written.
StreamType	type of stream used to write the data to

Parameters

os	output stream where the data is to be written
matrix	the matrix to write

References write_binary_raw().

42.1.4.77 write_binary_raw()

Writes a matrix_data structure to a stream in binary format.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

Template Parameters

ValueType	type of matrix values
IndexType	type of matrix indexes

Parameters

os	output stream where the data is to be written
data	the matrix data to write

Note

This is an advanced routine that writes the raw matrix data structure. If you are trying to write an existing matrix, consider using gko::write_binary instead.

Referenced by write_binary().

42.1.4.78 write_raw()

```
template<typename ValueType , typename IndexType >
void gko::write_raw (
```

```
std::ostream & os,
const matrix_data< ValueType, IndexType > & data,
layout_type layout = layout_type::array )
```

Writes a matrix_data structure to a stream in matrix market format.

Template Parameters

ValueType	type of matrix values
IndexType	type of matrix indexes

Parameters

os	output stream where the data is to be written
data	the matrix data to write
layout	the layout used in the output

Note

This is an advanced routine that writes the raw matrix data structure. If you are trying to write an existing matrix, consider using gko::write instead.

Referenced by write().

42.1.4.79 zero() [1/2]

```
template<typename T >
constexpr T gko::zero ( ) [inline], [constexpr]
```

Returns the additive identity for T.

Returns

additive identity for T

Referenced by gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::strategy_type().

42.1.4.80 zero() [2/2]

Returns the additive identity for T.

Returns

additive identity for T

Note

This version takes an unused reference argument to avoid complicated calls like zero < decltype(x) > (). Instead, it allows zero(x).

42.2 gko::accessor Namespace Reference

The accessor namespace.

Classes

· class row_major

A row_major accessor is a bridge between a range and the row-major memory layout.

42.2.1 Detailed Description

The accessor namespace.

42.3 gko::experimental::distributed Namespace Reference

The distributed namespace.

Namespaces

· preconditioner

The Preconditioner namespace.

Classes

· class DistributedBase

A base class for distributed objects.

class Matrix

The Matrix class defines a (MPI-)distributed matrix.

· class Partition

Represents a partition of a range of indices [0, size) into a disjoint set of parts.

· class Vector

Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

Typedefs

• using comm_index_type = int

Index type for enumerating processes in a distributed application.

Functions

template<typename ValueType >
 detail::temporary_conversion< experimental::distributed::Vector< ValueType > > make_temporary_conversion
 (LinOp *matrix)

Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector< Value ← Type>.

template<typename ValueType >
 detail::temporary_conversion< const experimental::distributed::Vector< ValueType > > make_temporary_conversion
 (const LinOp *matrix)

Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector< Value ← Type>.

template < typename ValueType, typename Function, typename... Args > void precision dispatch (Function fn, Args *... linops)

Calls the given function with each given argument LinOp temporarily converted into experimental::distributed::← Vector< Value Type> as parameters.

template < typename ValueType , typename Function >
 void precision_dispatch_real_complex (Function fn, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector< Value ← Type>* as parameters.

• template<typename ValueType, typename Function > void precision_dispatch_real_complex (Function fn, const LinOp *alpha, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector< Value

Type>* as parameters.

template<typename ValueType, typename Function >
 void precision_dispatch_real_complex (Function fn, const LinOp *alpha, const LinOp *in, const LinOp *beta,
 LinOp *out)

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector< Value ← Type>* as parameters.

42.3.1 Detailed Description

The distributed namespace.

42.3.2 Typedef Documentation

42.3.2.1 comm_index_type

using gko::experimental::distributed::comm_index_type = typedef int

Index type for enumerating processes in a distributed application.

Conforms to the MPI C interface of e.g. MPI rank or size

42.3.3 Function Documentation

42.3.3.1 make_temporary_conversion() [1/2]

Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector<Value← Type>.

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by next⇔ _precision<...> starting from the ValueType template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

matrix	the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type
	experimental::distributed::Vector <valuetype>, otherwise it will be converted to this type if possible.</valuetype>

Returns

a detail::temporary_conversion pointing to the (potentially converted) object.

Exceptions

al::distributed::Vector <valuetype></valuetype>	if the input matrix cannot be converted to experimental::c
---	--

Template Parameters

ValueType	the value type into whose associated experimental::distributed::Vector type to convert the input
	LinOp.

42.3.3.2 make_temporary_conversion() [2/2]

 $\label{lem:convert} \mbox{Convert the given LinOp from experimental::} \mbox{distributed::} \mbox{Vector} < ... > \mbox{to experimental::} \mbox{distributed::} \mbox{Vector} < \mbox{Value} \mbox{\hookrightarrow} \mbox{Type} > .$

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by next⇔_precision<...> starting from the ValueType template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

matrix	the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type	Ì
	experimental::distributed::Vector <valuetype>, otherwise it will be converted to this type if possible.</valuetype>	

Returns

a detail::temporary_conversion pointing to the (potentially converted) object.

Exceptions

NotSupported	if the input matrix cannot be converted to experimental::distributed::Vector <valuetype></valuetype>
--------------	--

Template Parameters

ValueType	the value type into whose associated experimental::distributed::Vector type to convert the input	
	LinOp.	

42.3.3.3 precision_dispatch()

Calls the given function with each given argument LinOp temporarily converted into experimental::distributed:: \leftarrow Vector<ValueType> as parameters.

Parameters

fn	the given function. It will be passed one (potentially const) experimental::distributed::Vector <valuetype>* parameter per parameter in the parameter pack linops.</valuetype>
linops	the given arguments to be converted and passed on to fn.

Template Parameters

ValueType	the value type to use for the parameters of fn.
Function	the function pointer, lambda or other functor type to call with the converted arguments.
Args	the argument type list.

42.3.3.4 precision_dispatch_real_complex() [1/3]

```
void gko::experimental::distributed::precision_dispatch_real_complex (
    Function fn,
    const LinOp * alpha,
    const LinOp * in,
    const LinOp * beta,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector<Value
Type>* as parameters.

If ValueType is real and both input vectors are complex, uses experimental::distributed::Vector::get_real_view() to convert them into real matrices after precision conversion.

See also

precision dispatch()

42.3.3.5 precision dispatch real complex() [2/3]

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector<Value \leftarrow Type>* as parameters.

If ValueType is real and both input vectors are complex, uses experimental::distributed::Vector::get_real_view() to convert them into real matrices after precision conversion.

See also

precision_dispatch()

42.3.3.6 precision_dispatch_real_complex() [3/3]

Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector<Value
Type>* as parameters.

If ValueType is real and both input vectors are complex, uses experimental::distributed::Vector::get_real_view() to convert them into real matrices after precision conversion.

See also

precision_dispatch()

42.4 gko::experimental::distributed::preconditioner Namespace Reference

The Preconditioner namespace.

Classes

· class Schwarz

A Schwarz preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

42.4.1 Detailed Description

The Preconditioner namespace.

42.5 gko::experimental::mpi Namespace Reference

The mpi namespace, contains wrapper for many MPI functions.

Classes

· class communicator

A thin wrapper of MPI_Comm that supports most MPI calls.

class contiguous_type

A move-only wrapper for a contiguous MPI_Datatype.

· class environment

Class that sets up and finalizes the MPI environment.

· class request

The request class is a light, move-only wrapper around the MPI_Request handle.

· struct status

The status struct is a light wrapper around the MPI_Status struct.

struct type_impl

A struct that is used to determine the MPI_Datatype of a specified type.

· class window

This class wraps the MPI_Window class with RAII functionality.

Enumerations

enum thread_type

This enum specifies the threading type to be used when creating an MPI environment.

Functions

• constexpr bool is_gpu_aware ()

Return if GPU aware functionality is available.

• int map_rank_to_device_id (MPI_Comm comm, int num_devices)

Maps each MPI rank to a single device id in a round robin manner.

- std::vector< status > wait_all (std::vector< request > &req)
 - Allows a rank to wait on multiple request handles.
- bool requires_host_buffer (const std::shared_ptr< const Executor > &exec, const communicator &comm)

Checks if the combination of Executor and communicator requires passing MPI buffers from the host memory.

• double get_walltime ()

Get the rank in the communicator of the calling process.

42.5.1 Detailed Description

The mpi namespace, contains wrapper for many MPI functions.

42.5.2 Function Documentation

42.5.2.1 get_walltime()

```
double gko::experimental::mpi::get_walltime () [inline]
```

Get the rank in the communicator of the calling process.

Parameters

```
comm the communicator
```

```
1522 { return MPI_Wtime(); }
```

42.5.2.2 map_rank_to_device_id()

Maps each MPI rank to a single device id in a round robin manner.

Parameters

comm	used to determine the node-local rank, if no suitable environment variable is available.
num_devices	the number of devices per node.

Returns

device id that this rank should use.

42.5.2.3 wait_all()

```
\begin{tabular}{ll} $\tt std::vector<status> gko::experimental::mpi::wait_all ( & std::vector< request > & req ) & [inline] \end{tabular}
```

Allows a rank to wait on multiple request handles.

Parameters

req The vector of request handles to be waited on.

Returns

status The vector of status objects that can be queried.

42.6 gko::experimental::reorder Namespace Reference

The Reorder namespace.

Classes

class Amd

Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.

· class ScaledReordered

Provides an interface to wrap reorderings like Rcm and diagonal scaling like equilibration around a LinOp like e.g.

42.6.1 Detailed Description

The Reorder namespace.

42.7 gko::factorization Namespace Reference

The Factorization namespace.

Classes

• class Ic

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

class Ilu

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

· class Parlc

ParIC is an incomplete Cholesky factorization which is computed in parallel.

· class Parlct

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

· class Parllu

ParILU is an incomplete LU factorization which is computed in parallel.

class Parllut

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

42.7.1 Detailed Description

The Factorization namespace.

42.8 gko::log Namespace Reference

The logger namespace.

Classes

· class Convergence

Convergence is a Logger which logs data strictly from the criterion_check_completed event.

struct criterion_data

Struct representing Criterion related data.

· class EnableLogging

EnableLogging is a mixin which should be inherited by any class which wants to enable logging.

• struct executor_data

Struct representing Executor related data.

· struct iteration_complete_data

Struct representing iteration complete related data.

struct linop_data

Struct representing LinOp related data.

struct linop_factory_data

Struct representing LinOp factory related data.

· class Loggable

Loggable class is an interface which should be implemented by classes wanting to support logging.

· struct operation_data

Struct representing Operator related data.

· class Papi

Papi is a Logger which logs every event to the PAPI software.

· class PerformanceHint

PerformanceHint is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

• struct polymorphic_object_data

Struct representing PolymorphicObject related data.

· class ProfilerHook

This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.

· class profiling_scope_guard

Scope guard that annotates its scope with the provided profiler hooks.

· class Record

Record is a Logger which logs every event to an object.

· class Stream

Stream is a Logger which logs every event to a stream.

Enumerations

```
    enum profile_event_category {
        profile_event_category::memory, profile_event_category::operation, profile_event_category::object,
        profile_event_category::linop,
        profile_event_category::factory, profile_event_category::solver, profile_event_category::criterion, profile_event_category::user,
        profile_event_category::internal }
```

Categorization of logger events.

42.8.1 Detailed Description

The logger namespace.

The Logging namespace.

Logging

42.8.2 Enumeration Type Documentation

42.8.2.1 profile_event_category

```
enum gko::log::profile_event_category [strong]
```

Categorization of logger events.

Enumerator

memory	Memory allocation.	
operation	Kernel execution and data movement.	
object	PolymorphicObject events.	
linop	LinOp events.	
factory	LinOpFactory events.	
solver	Solver events.	
criterion	Stopping criterion events.	
Generated Isomovydeser-defined events.		
internal	For development use.	

```
memory,
55
       operation,
57
       object,
59
       linop,
61
       factory.
       solver,
63
       criterion,
67
       user,
69
       internal,
70 };
```

42.9 gko::matrix Namespace Reference

The matrix namespace.

Classes

· class Coo

COO stores a matrix in the coordinate matrix format.

class Csr

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

class Dense

Dense is a matrix format which explicitly stores all values of the matrix.

class Diagonal

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

class Ell

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

class Fbcsr

Fixed-block compressed sparse row storage matrix format.

· class Fft

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

• class Fft2

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

class Fft3

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

class Hybrid

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

· class Identity

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

class IdentityFactory

This factory is a utility which can be used to generate Identity operators.

· class Permutation

Permutation is a matrix "format" which stores the row and column permutation arrays which can be used for reordering the rows and columns a matrix.

· class RowGatherer

RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix

· class Sellp

SELL-P is a matrix format similar to ELL format.

class SparsityCsr

SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

42.9.1 Detailed Description

The matrix namespace.

42.10 gko::multigrid Namespace Reference

The multigrid components namespace.

Classes

· class EnableMultigridLevel

The EnableMultigridLevel gives the default implementation of MultigridLevel with composition and provides $set_ \leftarrow multigrid_level$ function.

class FixedCoarsening

FixedCoarsening is a very simple coarse grid generation algorithm.

class MultigridLevel

This class represents two levels in a multigrid hierarchy.

class Pgm

Parallel graph match (Pgm) is the aggregate method introduced in the paper M.

42.10.1 Detailed Description

The multigrid components namespace.

42.11 gko::name demangling Namespace Reference

The name demangling namespace.

Functions

```
    template<typename T >
        std::string get_static_type (const T &)
```

This function uses name demangling facilities to get the name of the static type (T) of the object passed in arguments.

template<typename T >
 std::string get_dynamic_type (const T &t)

This function uses name demangling facilities to get the name of the dynamic type of the object passed in arguments.

42.11.1 Detailed Description

The name demangling namespace.

42.11.2 Function Documentation

42.11.2.1 get_dynamic_type()

This function uses name demangling facilities to get the name of the dynamic type of the object passed in arguments.

Template Parameters

T | the type of the object to demangle

Parameters

t the object we get the dynamic type of

```
101 {
102          return get_type_name(typeid(t));
103 }
```

42.11.2.2 get_static_type()

This function uses name demangling facilities to get the name of the static type (T) of the object passed in arguments.

Template Parameters

T | the type of the object to demangle

Parameters

unused

42.12 gko::preconditioner Namespace Reference

The Preconditioner namespace.

Classes

· struct block_interleaved_storage_scheme

Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.

• class Ic

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

· class Ilu

The Incomplete LU (ILU) preconditioner solves the equation LUx = b for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

• class Isai

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

class Jacobi

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

Enumerations

· enum isai type

This enum lists the types of the ISAI preconditioner.

42.12.1 Detailed Description

The Preconditioner namespace.

42.12.2 Enumeration Type Documentation

42.12.2.1 isai_type

```
enum gko::preconditioner::isai_type [strong]
```

This enum lists the types of the ISAI preconditioner.

ISAI can either be generated for a general square matrix, a lower triangular matrix, an upper triangular matrix or an spd matrix.

```
63 { lower, upper, general, spd };
```

42.13 gko::reorder Namespace Reference

The Reorder namespace.

Classes

class Rcm

Rcm is a reordering algorithm minimizing the bandwidth of a matrix.

· class ReorderingBase

The ReorderingBase class is a base class for all the reordering algorithms.

struct ReorderingBaseArgs

This struct is used to pass parameters to the EnableDefaultReorderingBaseFactory::generate() method.

Typedefs

- template<typename IndexType = int32>
 using ReorderingBaseFactory = AbstractFactory< ReorderingBase< IndexType >, ReorderingBaseArgs >
 Declares an Abstract Factory specialized for ReorderingBases.
- template < typename ConcreteFactory, typename ConcreteReorderingBase, typename ParametersType, typename IndexType = int32, typename PolymorphicBase = ReorderingBaseFactory < IndexType>> using EnableDefaultReorderingBaseFactory = EnableDefaultFactory
 ConcreteFactory, Concrete ← ReorderingBase, ParametersType, PolymorphicBase >

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of ReorderingBaseFactory.

42.13.1 Detailed Description

The Reorder namespace.

The reordering namespace.

42.13.2 Typedef Documentation

42.13.2.1 EnableDefaultReorderingBaseFactory

template<typename ConcreteFactory , typename ConcreteReorderingBase , typename ParametersType , typename IndexType = int32, typename PolymorphicBase = ReorderingBaseFactory<IndexType>> using gko::reorder::EnableDefaultReorderingBaseFactory = typedef EnableDefaultFactory<Concrete←Factory, ConcreteReorderingBase, ParametersType, PolymorphicBase>

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of ReorderingBaseFactory.

Template Parameters

ConcreteFactory	the concrete factory which is being implemented [CRTP parmeter]
ConcreteReorderingBase	the concrete ReorderingBase type which this factory produces, needs to have a constructor which takes a const ConcreteFactory *, and a const ReorderingBaseArgs * as parameters.
ParametersType	a subclass of enable_parameters_type template which defines all of the parameters of the factory
PolymorphicBase	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of ReorderingBaseFactory

42.14 gko::solver Namespace Reference

The ginkgo Solve namespace.

Namespaces

· multigrid

The solver multigrid namespace.

Classes

- · class ApplyWithInitialGuess
 - ApplyWithInitialGuess provides a way to give the input guess for apply function.
- class Bicg

BICG or the Biconjugate gradient method is a Krylov subspace solver.

· class Bicgstab

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

· class CbGmres

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

class Cq

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

· class Cgs

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

class EnableApplyWithInitialGuess

EnableApplyWithInitialGuess providing default operation for ApplyWithInitialGuess with correct validation and log.

class EnableIterativeBase

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

class EnablePreconditionable

Mixin providing default operation for Preconditionable with correct value semantics.

· class EnablePreconditionedIterativeSolver

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

class EnableSolverBase

A LinOp deriving from this CRTP class stores a system matrix.

· class Fcg

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

· class Gcr

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

· class Gmres

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

· struct has with criteria

Helper structure to test if the Factory of SolverType has a function with_criteria.

Helper structure to test if the Factory of SolverType has a function with_criteria.

struct has_with_criteria< SolverType, xstd::void_t< decltype(SolverType::build().with_criteria(std::shared_ptr< const stop::Criteria

• class ldr

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

· class Ir

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

· class IterativeBase

A LinOp implementing this interface stores a stopping criterion factory.

class LowerTrs

LowerTrs is the triangular solver which solves the system L x = b, when L is a lower triangular matrix.

· class Multigrid

Multigrid methods have a hierarchy of many levels, whose corase level is a subset of the fine level, of the problem.

class UpperTrs

Upper Trs is the triangular solver which solves the system Ux = b, when U is an upper triangular matrix.

struct workspace_traits

Traits class providing information on the type and location of workspace vectors inside a solver.

Enumerations

enum initial_guess_mode { initial_guess_mode::zero, initial_guess_mode::rhs, initial_guess_mode::provided }

Give a initial guess mode about the input of the apply method.

• enum trisolve_algorithm

A helper for algorithm selection in the triangular solvers.

Functions

• template<typename ValueType > auto build_smoother (std::shared_ptr< const LinOpFactory > factory, size_type iteration=1, ValueType relaxation_factor=0.9)

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relacation_factor).

template<typename ValueType >
 auto build_smoother (std::shared_ptr< const LinOp > solver, size_type iteration=1, ValueType relaxation_
 factor=0.9)

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relacation_factor).

42.14.1 Detailed Description

The ginkgo Solve namespace.

The ginkgo Solver namespace.

42.14.2 Enumeration Type Documentation

42.14.2.1 initial_guess_mode

```
enum gko::solver::initial_guess_mode [strong]
```

Give a initial guess mode about the input of the apply method.

Enumerator

zero	the input is zero
rhs	the input is right hand side
provided	the input is provided

```
67 {
71 zero,
75 rhs,
79 provided
80 }:
```

42.14.2.2 trisolve_algorithm

```
enum gko::solver::trisolve_algorithm [strong]
```

A helper for algorithm selection in the triangular solvers.

It currently only matters for the Cuda executor as there, we have a choice between the Ginkgo syncfree and cuS← PARSE implementations.

```
67 { sparselib, syncfree };
```

42.14.3 Function Documentation

42.14.3.1 build_smoother() [1/2]

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relacation_factor).

Parameters

solver	the shared pointer of solver	
iteration	the maximum number of iteraion, which default is 1	
relaxation_factor	the relaxation factor for Richardson	

Returns

the pointer of Ir(Richardson)

Note

this is the overload function for LinOp.

42.14.3.2 build_smoother() [2/2]

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relacation factor).

Parameters

factory	the shared pointer of factory	
iteration the maximum number of iteraion, which default		
relaxation_factor	ctor the relaxation factor for Richardson	

Returns

the pointer of Ir(Richardson)

42.15 gko::solver::multigrid Namespace Reference

The solver multigrid namespace.

Enumerations

· enum cycle

cycle defines which kind of multigrid cycle can be used.

• enum mid_smooth_type

mid_smooth_type gives the options to handle the middle smoother behavior between the two cycles in the same level.

42.15.1 Detailed Description

The solver multigrid namespace.

42.15.2 Enumeration Type Documentation

42.15.2.1 cycle

```
enum gko::solver::multigrid::cycle [strong]
```

cycle defines which kind of multigrid cycle can be used.

It contains V, W, and F cycle.

- V, W cycle uses the algorithm according to Briggs, Henson, and McCormick: A multigrid tutorial 2nd Edition.
- F cycle uses the algorithm according to Trottenberg, Oosterlee, and Schuller: Multigrid 1st Edition. F cycle first uses the recursive call but second uses the V-cycle call such that F-cycle is between V and W cycle.

42.15.2.2 mid_smooth_type

```
enum gko::solver::multigrid::mid_smooth_type [strong]
```

mid_smooth_type gives the options to handle the middle smoother behavior between the two cycles in the same level.

It only affects the behavior when there's no operation between the post smoother of previous cycle and the pre smoother of next cycle. Thus, it only affects W cycle and F cycle.

- both: gives the same behavior as the original algorithm, which use posts smoother from previous cycle and pre smoother from next cycle.
- post_smoother: only uses the post smoother of previous cycle in the mid smoother
- · pre_smoother: only uses the pre smoother of next cycle in the mid smoother
- · standalone: uses the defined smoother in the mid smoother

42.16 gko::stop Namespace Reference

The Stopping criterion namespace.

Classes

· class AbsoluteResidualNorm

The AbsoluteResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

· class Combined

The Combined class is used to combine multiple criterions together through an OR operation.

class Criterion

The Criterion class is a base class for all stopping criteria.

struct CriterionArgs

This struct is used to pass parameters to the EnableDefaultCriterionFactoryCriterionFactory::generate() method.

· class ImplicitResidualNorm

The ImplicitResidualNorm class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

· class Iteration

The Iteration class is a stopping criterion which stops the iteration process after a preset number of iterations.

class RelativeResidualNorm

The RelativeResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

· class ResidualNorm

The ResidualNorm class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

class ResidualNormBase

The ResidualNormBase class provides a framework for stopping criteria related to the residual norm.

• class ResidualNormReduction

The ResidualNormReduction class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

class Time

The Time class is a stopping criterion which stops the iteration process after a certain amout of time has passed.

Typedefs

- using CriterionFactory = AbstractFactory < Criterion, CriterionArgs >
 - Declares an Abstract Factory specialized for Criterions.
- template<typename ConcreteFactory , typename ConcreteCriterion , typename ParametersType , typename PolymorphicBase =
 CriterionFactory>

```
using EnableDefaultCriterionFactory = EnableDefaultFactory< ConcreteFactory, ConcreteCriterion, ParametersType, PolymorphicBase >
```

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of CriterionFactory.

Enumerations

· enum mode

The mode for the residual norm criterion.

Functions

template<typename FactoryContainer >
 std::shared_ptr< const CriterionFactory > combine (FactoryContainer &&factories)
 Combines multiple criterion factories into a single combined criterion factory.

42.16.1 Detailed Description

The Stopping criterion namespace.

Stopping criteria

42.16.2 Typedef Documentation

42.16.2.1 EnableDefaultCriterionFactory

```
template<typename ConcreteFactory , typename ConcreteCriterion , typename ParametersType ,
typename PolymorphicBase = CriterionFactory>
using gko::stop::EnableDefaultCriterionFactory = typedef EnableDefaultFactory<ConcreteFactory,
ConcreteCriterion, ParametersType, PolymorphicBase>
```

This is an alias for the EnableDefaultFactory mixin, which correctly sets the template parameters to enable a subclass of CriterionFactory.

Template Parameters

ConcreteFactory	the concrete factory which is being implemented [CRTP parmeter]	
ConcreteCriterion	the concrete Criterion type which this factory produces, needs to have a constructor which takes a const ConcreteFactory $*$, and a const CriterionArgs $*$ as parameters.	
ParametersType	a subclass of enable_parameters_type template which defines all of the parameters of the factory	
PolymorphicBase	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclaggrephted by Doxygen CriterionFactory	

42.17 gko::syn Namespace Reference

The Synthesizer namespace.

Classes

```
• struct range range records start, end, step in template
```

struct type_list

type_list records several types in template

struct value_list

value_list records several values with the same type in template.

Typedefs

```
    template<typename List1 , typename List2 >
        using concatenate = typename detail::concatenate_impl< List1, List2 >::type
        concatenate combines two value_list into one value_list.
    template<typename T >
        using as_list = typename detail::as_list_impl< T >::type
        as_list<T> gives the alias type of as_list_impl<T>::type.
```

Functions

```
    template<typename T, T... Value>
    constexpr std::array< T, sizeof...(Value)> as_array (value_list< T, Value... > vI)
    as_array<T> returns the array from value_list.
```

42.17.1 Detailed Description

The Synthesizer namespace.

42.17.2 Typedef Documentation

42.17.2.1 as_list

```
template<typename T >
using gko::syn::as_list = typedef typename detail::as_list_impl<T>::type
as_list<T> gives the alias type of as_list_impl<T>::type.
```

It gives a list (itself) if input is already a list, or generates list type from range input.

Template Parameters

```
T list or range
```

42.17.2.2 concatenate

```
template<typename List1 , typename List2 >
using gko::syn::concatenate = typedef typename detail::concatenate_impl<List1, List2>::type
```

concatenate combines two value_list into one value_list.

Template Parameters

List1	the first list	
List2	the second list	

42.17.3 Function Documentation

42.17.3.1 as_array()

as_array<T> returns the array from value_list.

It will be helpful if using for in runtime on the array.

Template Parameters

T	the type of value_list	
Value	the values of value_list	

Parameters

```
value_list the input value_list
```

Returns

std::array the std::array contains the values of value_list

```
204 {
205         return std::array<T, sizeof...(Value)>{Value...};
206 }
```

References gko::array.

42.18 gko::xstd Namespace Reference

The namespace for functionalities after C++14 standard.

42.18.1 Detailed Description

The namespace for functionalities after C++14 standard.

Chapter 43

Class Documentation

43.1 gko::AbsoluteComputable Class Reference

The AbsoluteComputable is an interface that allows to get the component wise absolute of a LinOp.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual std::unique_ptr< LinOp > compute_absolute_linop () const =0
 Gets the absolute LinOp.
- virtual void compute_absolute_inplace ()=0

 Compute absolute inplace on each element.

43.1.1 Detailed Description

The AbsoluteComputable is an interface that allows to get the component wise absolute of a LinOp.

Use EnableAbsoluteComputation<AbsoluteLinOp> to implement this interface.

43.1.2 Member Function Documentation

43.1.2.1 compute_absolute_linop()

```
virtual std::unique_ptr<LinOp> gko::AbsoluteComputable::compute_absolute_linop ( ) const
[pure virtual]
```

Gets the absolute LinOp.

Returns

a pointer to the new absolute LinOp

Implemented in gko::EnableAbsoluteComputation < AbsoluteLinOp >, gko::EnableAbsoluteComputation < remove_complex < Hybrid gko::EnableAbsoluteComputation < remove_complex < Sellp < ValueType, IndexType > > , gko::EnableAbsoluteComputation < regko::EnableAbsoluteComputation < remove_complex < Coo < ValueType, IndexType > > , gko::EnableAbsoluteComputation < remove_complex < Vector < ValueType > > , gko::EnableAbsoluteComputation < remove_complex < Vector < ValueType > > , gko::EnableAbsoluteComputation < remove_complex < Vector < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType, IndexType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType > > , and gko::EnableAbsoluteComputation < remove_complex < Csr < ValueType > > , and gko::EnableAbsoluteComputation < remove_comple

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.2 gko::stop::AbsoluteResidualNorm< ValueType > Class Template Reference

The AbsoluteResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

43.2.1 Detailed Description

```
template < typename ValueType = default_precision > class gko::stop::AbsoluteResidualNorm < ValueType >
```

The AbsoluteResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

when norm(residual) < threshold. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on b in order to get the number of right-hand sides. If this is not correctly provided, an exception ::gko::NotSupported() is thrown.

The documentation for this class was generated from the following file:

ginkgo/core/stop/residual_norm.hpp

43.3 gko::AbstractFactory< AbstractProductType, ComponentsType > Class Template Reference

The AbstractFactory is a generic interface template that enables easy implementation of the abstract factory design pattern.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

```
    template<typename... Args>
    std::unique_ptr< abstract_product_type > generate (Args &&... args) const
        Creates a new product from the given components.
```

43.3.1 Detailed Description

```
template < typename\ AbstractProductType,\ typename\ ComponentsType > \\ class\ gko:: AbstractFactory < AbstractProductType,\ ComponentsType > \\
```

The AbstractFactory is a generic interface template that enables easy implementation of the abstract factory design pattern.

The interface provides the AbstractFactory::generate() method that can produce products of type Abstract ProductType using an object of ComponentsType (which can be constructed on the fly from parameters to its constructors). The generate() method is not declared as virtual, as this allows subclasses to hide the method with a variant that preserves the compile-time type of the objects. Instead, implementers should override the generate impl() method, which is declared virtual.

Implementers of concrete factories should consider using the EnableDefaultFactory mixin to obtain default implementations of utility methods of PolymorphicObject and AbstractFactory.

Template Parameters

AbstractProductType	the type of products the factory produces
ComponentsType	the type of components the factory needs to produce the product

43.3.2 Member Function Documentation

43.3.2.1 generate()

Creates a new product from the given components.

The method will create an ComponentsType object from the arguments of this method, and pass it to the generate ← _impl() function which will create a new AbstractProductType.

Template Parameters

Parameters

```
args arguments passed to the constructor of ComponentsType
```

Returns

an instance of AbstractProductType

The documentation for this class was generated from the following file:

ginkgo/core/base/abstract_factory.hpp

43.4 gko::AllocationError Class Reference

AllocationError is thrown if a memory allocation fails.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

AllocationError (const std::string &file, int line, const std::string &device, size_type bytes)
 Initializes an allocation error.

43.4.1 Detailed Description

AllocationError is thrown if a memory allocation fails.

43.4.2 Constructor & Destructor Documentation

43.4.2.1 AllocationError()

Initializes an allocation error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
device	The device on which the error occurred	
bytes	The size of the memory block whose allocation failed.	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.5 gko::experimental::reorder::Amd < IndexType > Class Template Reference

Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.

```
#include <ginkgo/core/reorder/amd.hpp>
```

Public Member Functions

 $\bullet \ \ \text{std::unique_ptr} < \text{permutation_type} > \text{generate (std::shared_ptr} < \text{const LinOp} > \text{system_matrix) const} \\$

Static Public Member Functions

static parameters_type build ()
 Creates a new parameter type to set up the factory.

43.5.1 Detailed Description

```
template<typename IndexType = int32> class gko::experimental::reorder::Amd< IndexType >
```

Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.

Template Parameters

IndexType the type used to store sparsity pattern indices of the system matrix

43.5.2 Member Function Documentation

43.5.2.1 generate()

Note

This function overrides the default LinOpFactory::generate to return a Permutation instead of a generic LinOp, which would need to be cast to Permutation again to access its indices. It is only necessary because smart pointers aren't covariant.

The documentation for this class was generated from the following file:

ginkgo/core/reorder/amd.hpp

43.6 gko::amd_device Class Reference

amd_device handles the number of executor on Amd devices and have the corresponding recursive_mutex.

```
#include <ginkgo/core/base/device.hpp>
```

43.6.1 Detailed Description

amd device handles the number of executor on Amd devices and have the corresponding recursive mutex.

The documentation for this class was generated from the following file:

• ginkgo/core/base/device.hpp

43.7 gko::solver::ApplyWithInitialGuess Class Reference

ApplyWithInitialGuess provides a way to give the input guess for apply function.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

43.7.1 Detailed Description

ApplyWithInitialGuess provides a way to give the input guess for apply function.

All functionalities have the protected access specifier. It should only be used internally.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.8 gko::are_all_integral < Args > Struct Template Reference

Evaluates if all template arguments Args fulfill std::is integral.

```
#include <ginkgo/core/base/types.hpp>
```

43.8.1 Detailed Description

```
template<typename... Args> struct gko::are_all_integral< Args>
```

Evaluates if all template arguments Args fulfill std::is_integral.

If that is the case, this class inherits from std::true_type, otherwise, it inherits from std::false_type. If no values are passed in, std::true_type is inherited from.

Template Parameters

Arae	Arguments to test for std::is_integral	
~!us	Alduliello lo lesi loi sidis ililediai	

The documentation for this struct was generated from the following file:

ginkgo/core/base/types.hpp

43.9 gko::array < ValueType > Class Template Reference

An array is a container which encapsulates fixed-sized arrays, stored on the Executor tied to the array.

```
#include <ginkgo/core/base/array.hpp>
```

Public Types

using value type = ValueType

The type of elements stored in the array.

using default_deleter = executor_deleter < value_type[]>

The default deleter type used by array.

• using view deleter = null deleter < value type[]>

The deleter type used for views.

Public Member Functions

· array () noexcept

Creates an empty array not tied to any executor.

array (std::shared_ptr< const Executor > exec) noexcept

Creates an empty array tied to the specified Executor.

array (std::shared_ptr< const Executor > exec, size_type num_elems)

Creates an array on the specified Executor.

 $\bullet \ \ \text{template}{<} \text{typename DeleterType} >$

array (std::shared_ptr< const Executor > exec, size_type num_elems, value_type *data, DeleterType deleter)

Creates an array from existing memory.

array (std::shared_ptr< const Executor > exec, size_type num_elems, value_type *data)

Creates an array from existing memory.

 $\bullet \ \ template {<} typename \ Random Access Iterator >$

array (std::shared ptr< const Executor > exec, RandomAccessIterator begin, RandomAccessIterator end)

Creates an array on the specified Executor and initializes it with values.

• template<typename T >

```
array (std::shared ptr< const Executor > exec, std::initializer list< T > init list)
```

Creates an array on the specified Executor and initializes it with values.

array (std::shared_ptr< const Executor > exec, const array &other)

Creates a copy of another array on a different executor.

• array (const array &other)

Creates a copy of another array.

array (std::shared_ptr< const Executor > exec, array &&other)

Moves another array to a different executor.

· array (array &&other)

Moves another array.

array< ValueType > as_view ()

Returns a non-owning view of the memory owned by this array.

detail::const_array_view< ValueType > as_const_view () const

Returns a non-owning constant view of the memory owned by this array.

• array & operator= (const array &other)

Copies data from another array or view.

• array & operator= (array &&other)

Moves data from another array or view.

template<typename OtherValueType >

 $std::enable_if_t < !std::is_same < ValueType, OtherValueType >::value, array > & operator= (const array < OtherValueType > & otherValueType > &$

Copies and converts data from another array with another data type.

• void clear () noexcept

Deallocates all data used by the array.

void resize_and_reset (size_type num_elems)

Resizes the array so it is able to hold the specified number of elements.

void fill (const value_type value)

Fill the array with the given value.

size_type get_num_elems () const noexcept

Returns the number of elements in the array.

value_type * get_data () noexcept

Returns a pointer to the block of memory used to store the elements of the array.

const value_type * get_const_data () const noexcept

Returns a constant pointer to the block of memory used to store the elements of the array.

std::shared_ptr< const Executor > get_executor () const noexcept

Returns the Executor associated with the array.

void set executor (std::shared ptr< const Executor > exec)

Changes the Executor of the array, moving the allocated data to the new Executor.

bool is_owning ()

Tells whether this array owns its data or not.

Creates an array from existing memory.

Static Public Member Functions

• static array view (std::shared_ptr< const Executor > exec, size_type num_elems, value_type *data)

• static detail::const_array_view< ValueType > const_view (std::shared_ptr< const Executor > exec, size_type num_elems, const value_type *data)

Creates a constant (immutable) array from existing memory.

43.9.1 Detailed Description

```
template<typename ValueType> class gko::array< ValueType>
```

An array is a container which encapsulates fixed-sized arrays, stored on the Executor tied to the array.

The array stores and transfers its data as **raw** memory, which means that the constructors of its elements are not called when constructing, copying or moving the array. Thus, the array class is most suitable for storing POD types.

Template Parameters

ValueType the type of elements stored in the	array
--	-------

43.9.2 Constructor & Destructor Documentation

43.9.2.1 array() [1/11]

```
template<typename ValueType>
gko::array< ValueType >::array ( ) [inline], [noexcept]
```

Creates an empty array not tied to any executor.

An array without an assigned executor can only be empty. Attempts to change its size (e.g. via the resize_and_\circ} reset method) will result in an exception. If such an array is used as the right hand side of an assignment or move assignment expression, the data of the target array will be cleared, but its executor will not be modified.

The executor can later be set by using the set_executor method. If an array with no assigned executor is assigned or moved to, it will inherit the executor of the source array.

43.9.2.2 array() [2/11]

Creates an empty array tied to the specified Executor.

Parameters

```
exec the Executor where the array data is allocated
```

43.9.2.3 array() [3/11]

Creates an array on the specified Executor.

exec	the Executor where the array data will be allocated	
num_elems Generated by Doxy	the amount of memory (expressed as the number of value_type elements) allocated on the energy description of the property of the second of the s	

43.9.2.4 array() [4/11]

Creates an array from existing memory.

The memory will be managed by the array, and deallocated using the specified deleter (e.g. use std::default_delete for data allocated with new).

Template Parameters

DeleterType	type of the deleter
-------------	---------------------

Parameters

exec	executor where data is located	
num_elems	number of elements in data	
data	chunk of memory used to create the array	
deleter	the deleter used to free the memory	

See also

```
array::view() to create an array that does not deallocate memory
array(std::shared_ptr<cont Executor>, size_type, value_type*) to deallocate the memory using Executor::free()
method
```

43.9.2.5 array() [5/11]

Creates an array from existing memory.

The memory will be managed by the array, and deallocated using the Executor::free method.

exec	executor where data is located	
num_elems	number of elements in data	Generated by Doxygen
data	chunk of memory used to create the array	Generated by Doxygen

43.9.2.6 array() [6/11]

Creates an array on the specified Executor and initializes it with values.

Template Parameters

RandomAccessIterator	type of the iterators
----------------------	-----------------------

Parameters

exec	the Executor where the array data will be allocated
begin	start of range of values
end	end of range of values

43.9.2.7 array() [7/11]

Creates an array on the specified Executor and initializes it with values.

Template Parameters

T | type of values used to initialize the array (T has to be implicitly convertible to value_type)

exec	the Executor where the array data will be allocated
init_list	list of values used to initialize the array

43.9.2.8 array() [8/11]

Creates a copy of another array on a different executor.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

exec	the executor where the new array will be created
other	the array to copy from

43.9.2.9 array() [9/11]

Creates a copy of another array.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

```
other the array to copy from
```

43.9.2.10 array() [10/11]

Moves another array to a different executor.

This does not invoke the constructors of the elements, instead they are copied as POD types.

exec	the executor where the new array will be moved
other	the array to move

43.9.2.11 array() [11/11]

Moves another array.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

other	the array to move
otner	the array to move

43.9.3 Member Function Documentation

43.9.3.1 as_const_view()

```
template<typename ValueType>
detail::const_array_view<ValueType> gko::array< ValueType >::as_const_view ( ) const [inline]
```

Returns a non-owning constant view of the memory owned by this array.

It can only be used until this array gets deleted, cleared or resized.

43.9.3.2 as_view()

```
template<typename ValueType>
array<ValueType> gko::array< ValueType >::as_view ( ) [inline]
```

Returns a non-owning view of the memory owned by this array.

It can only be used until this array gets deleted, cleared or resized.

43.9.3.3 clear()

```
template<typename ValueType>
void gko::array< ValueType >::clear ( ) [inline], [noexcept]
```

Deallocates all data used by the array.

The array is left in a valid, but empty state, so the same array can be used to allocate new memory. Calls to array::get_data() will return a nullptr.

Referenced by gko::index_set< IndexType >::clear(), and gko::array< index_type >::operator=().

43.9.3.4 const_view()

Creates a constant (immutable) array from existing memory.

The array does not take ownership of the memory, and will not deallocate it once it goes out of scope. This array type cannot use the function resize_and_reset since it does not own the data it should resize.

Parameters

exec	executor where data is located
num_elems	number of elements in data
data	chunk of memory used to create the array

Returns

an array constructed from data

Referenced by gko::array< index_type >::as_const_view().

43.9.3.5 fill()

Fill the array with the given value.

Parameters

```
value the value to be filled
```

43.9.3.6 get_const_data()

```
template<typename ValueType>
const value_type* gko::array< ValueType >::get_const_data ( ) const [inline], [noexcept]
```

Returns a constant pointer to the block of memory used to store the elements of the array.

Returns

a constant pointer to the block of memory used to store the elements of the array

Referenced by gko::array< index type >::as const view(), gko::matrix::Dense< value type >::at(), gko↔ ::preconditioner::Jacobi< ValueType, IndexType >::get blocks(), gko::preconditioner::Jacobi< ValueType, Index← Type >::get_conditioning(), gko::multigrid::Pgm< ValueType, IndexType >::get_const_agg(), gko::matrix::← SparsityCsr< ValueType, IndexType >::get_const_col_idxs(), gko::matrix::Sellp< ValueType, IndexType >::get← _const_col_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_const_col_idxs(), gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs(), gko::device_matrix_data< ValueType, IndexType >::get_const_col_idxs(), gko::matrix::Fbcsr< ValueType, IndexType >::get const col idxs(), gko::matrix::Csr< ValueType, IndexType >::get_const_col_idxs(), gko::matrix::Permutation< IndexType >::get_const_permutation(), gko::matrix::Row← Gatherer< IndexType >::get_const_row_idxs(), gko::device_matrix_data< ValueType, IndexType >::get_const_← row_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_const_row_idxs(), gko::matrix::SparsityCsr< Value ← Type, IndexType >::get const row ptrs(), gko::matrix::Fbcsr< ValueType, IndexType >::get const row ptrs(), gko::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs(), gko::matrix::Sellp< ValueType, IndexType >::get const slice lengths(), gko::matrix::Sellp< ValueType, IndexType >::get const slice sets(), gko::matrix::← Csr< ValueType, IndexType >::get const srow(), gko::matrix::SparsityCsr< ValueType, IndexType >::get const ← value(), gko::matrix::Diagonal < ValueType >::get const values(), gko::matrix::Sellp < ValueType, IndexType > ← ::get_const_values(), gko::matrix::Coo< ValueType, IndexType >::get_const_values(), gko::matrix::Ell< ValueType, IndexType >::get_const_values(), gko::device_matrix_data< ValueType, IndexType >::get_const_values(), gko⇔ ::matrix::Fbcsr< ValueType, IndexType >::get const values(), gko::matrix::Dense< value type >::get const \(\lefta \) values(), gko::matrix::Csr< ValueType, IndexType >::get const values(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_ids(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get part size(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get part sizes(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get range← bounds(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get range starting ← indices(), gko::index_set< IndexType >::get_subsets_begin(), gko::index_set< IndexType >::get_subsets_end(), gko::index_set< IndexType >::get_superset_indices(), gko::matrix::Csr< ValueType, IndexType >::classical← ::process(), gko::matrix::Csr< ValueType, IndexType >::load_balance::process(), gko::matrix::Ell< ValueType, IndexType >::val_at(), and gko::matrix::Sellp< ValueType, IndexType >::val_at().

43.9.3.7 get_data()

```
template<typename ValueType>
value_type* gko::array< ValueType >::get_data ( ) [inline], [noexcept]
```

Returns a pointer to the block of memory used to store the elements of the array.

Returns

a pointer to the block of memory used to store the elements of the array

Referenced by gko::array< index_type >::as_view(), gko::matrix::Dense< value_type >::at(), gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row(), gko::multigrid::Pgm< ValueType, IndexType >::get_agg(), gko::matrix::SparsityCsr< ValueType, IndexType >::get_col_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_col_eidxs(), gko::matrix::Ell< ValueType, IndexType >::get_col_idxs(), gko::device_matrix_data< ValueType, IndexType >::get_col_idxs(), gko::matrix::Csr< ValueType, IndexType >::get_col_idxs(), gko::matrix::Csr< ValueType, IndexType >::get_col_idxs(), gko::matrix::Csr< ValueType, IndexType >::get_col_idxs(), gko::matrix::Csr< ValueType, IndexType >::get_permutation(), gko::matrix::Row Gatherer< IndexType >::get_row_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_row_idxs(), gko::matrix::SparsityCsr< ValueType, IndexType >::get_row_ptrs(), gko::matrix::Csr< ValueType, IndexType, IndexType >::get_row_ptrs(), gko::matrix::Csr< ValueType, IndexType, IndexType,

 $\label{thm:continuity:indexType} $::get_row_ptrs(), gko::matrix::Sellp< ValueType, IndexType>::get_slice_lengths(), gko::matrix:: $$ Sellp< ValueType, IndexType>::get_slice_sets(), gko::matrix::Csr< ValueType, IndexType>::get_srow(), gko::matrix::SparsityCsr< ValueType, IndexType>::get_value(), gko::matrix::Diagonal< ValueType>::get_values(), gko::matrix::Sellp< ValueType, IndexType>::get_values(), gko::matrix::Coo< ValueType, IndexType>::get_coo
<math display="block">\label{thm:coo} ValueType, IndexType>::get_values(), gko::matrix::Dense< ValueType, IndexType>::get_values(), gko::matrix::Dense< ValueType, IndexType>::get_values(), gko::matrix::Dense< ValueType>::get_values(), gko::matrix::Csr< ValueType, IndexType>::get_values(), gko::matrix::Csr< ValueType, IndexType>::get_values(), gko::matrix::Ell< ValueType, IndexType>::load_balance::process(), gko::matrix::Ell< ValueType, IndexType>::val_at(), and gko::matrix::Sellp< ValueType, IndexType>::val_at().$

43.9.3.8 get executor()

```
template<typename ValueType>
std::shared_ptr<const Executor> gko::array< ValueType >::get_executor ( ) const [inline],
[noexcept]
```

Returns the Executor associated with the array.

Returns

the Executor associated with the array

Referenced by gko::array< index_type >::as_const_view(), gko::array< index_type >::as_view(), gko::matrix:: Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config(), gko::device_matrix_data< ValueType, IndexType >::get_executor(), gko::matrix:: Csr< ValueType, IndexType >::classical::process(), and gko::matrix:: Csr< ValueType, IndexType >::load balance::process().

43.9.3.9 get_num_elems()

```
template<typename ValueType>
size_type gko::array< ValueType >::get_num_elems ( ) const [inline], [noexcept]
```

Returns the number of elements in the array.

Returns

the number of elements in the array

Referenced by gko::array< index_type >::as_const_view(), gko::array< index_type >::as_view(), gko::matrix::← Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row(), gko::matrix← ::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_stored_elements_per_row(), gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config(), gko::device_matrix_← data< ValueType, IndexType >::get_num_elems(), gko::matrix::SparsityCsr< ValueType, IndexType >::get← $_$ num $_$ nonzeros(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get $_$ num $_\leftrightarrow$ ranges(), gko::matrix::Csr< ValueType, IndexType >::get_num_srow_elements(), gko::matrix::Fbcsr< ValueType, IndexType >::get num stored blocks(), gko::matrix::Ell< ValueType, IndexType >::get num stored elements(), gko::matrix::Coo< ValueType, IndexType >::get num stored elements(), gko::matrix::Sellp< ValueType, Index← Type >::get_num_stored_elements(), gko::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored ← _elements(), gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_elements(), gko::matrix::Dense< value_type >::get_num_stored_elements(), gko::matrix::Csr< ValueType, IndexType >::get_num_stored_← elements(), gko::index_set< IndexType >::get_num_subsets(), gko::matrix::Permutation< IndexType >::get← _permutation_size(), gko::matrix::Sellp< ValueType, IndexType >::get_total_cols(), gko::index_set< IndexType >::index_set(), gko::matrix::Csr< ValueType, IndexType >::classical::process(), and gko::matrix::Csr< ValueType, IndexType >::load_balance::process().

43.9.3.10 is_owning()

```
template<typename ValueType>
bool gko::array< ValueType >::is_owning ( ) [inline]
```

Tells whether this array owns its data or not.

Views do not own their data and this has multiple implications. They cannot be resized since the data is not owned by the array which stores a view. It is also unclear whether custom deleter types are owning types as they could be a user-created view-type, therefore only proper array which use the default_deleter are considered owning types.

Returns

whether this array can be resized or not.

Referenced by gko::array< index type >::operator=().

43.9.3.11 operator=() [1/3]

Moves data from another array or view.

Only the pointer and deleter type change, a copy only happens when targeting another executor's data. This means that in the following situation:

```
gko::array<int> a; // an existing array or view
gko::array<int> b; // an existing array or view
b = std::move(a);
```

Depending on whether a and b are array or view, this happens:

- a and b are views, b becomes the only valid view of a;
- a and b are arrays, b becomes the only valid array of a;
- a is a view and b is an array, b frees its data and becomes the only valid view of a ();
- a is an array and b is a view, b becomes the only valid array of a.

In all the previous cases, a becomes empty (e.g., a nullptr).

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

other	the array to move data from
-------	-----------------------------

Returns

this

43.9.3.12 operator=() [2/3]

Copies data from another array or view.

In the case of an array target, the array is resized to match the source's size. In the case of a view target, if the dimensions are not compatible a gko::OutOfBoundsError is thrown.

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

other	the array to copy from
-------	------------------------

Returns

this

43.9.3.13 operator=() [3/3]

Copies and converts data from another array with another data type.

In the case of an array target, the array is resized to match the source's size. In the case of a view target, if the dimensions are not compatible a gko::OutOfBoundsError is thrown.

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

other	the array to copy from
-------	------------------------

Template Parameters

Returns

this

43.9.3.14 resize_and_reset()

Resizes the array so it is able to hold the specified number of elements.

For a view and other non-owning array types, this throws an exception since these types cannot be resized.

All data stored in the array will be lost.

If the array is not assigned an executor, an exception will be thrown.

Parameters

num_elems	the amount of memory (expressed as the number of value_type elements) allocated on the	
	Executor	

Referenced by gko::array< index_type >::operator=().

43.9.3.15 set_executor()

Changes the Executor of the array, moving the allocated data to the new Executor.

Parameters

exec the Executor where the data will be moved to

43.9.3.16 view()

Creates an array from existing memory.

The array does not take ownership of the memory, and will not deallocate it once it goes out of scope. This array type cannot use the function resize_and_reset since it does not own the data it should resize.

Parameters

exec	executor where data is located
num_elems	number of elements in data
data	chunk of memory used to create the array

Returns

an array constructed from data

Referenced by gko::array< index_type >::as_view().

The documentation for this class was generated from the following file:

• ginkgo/core/base/array.hpp

43.10 gko::device_matrix_data< ValueType, IndexType >::arrays Struct Reference

Stores the internal arrays of a device_matrix_data object.

```
#include <ginkgo/core/base/device_matrix_data.hpp>
```

43.10.1 Detailed Description

```
template<typename ValueType, typename IndexType>
struct gko::device_matrix_data< ValueType, IndexType >::arrays
```

Stores the internal arrays of a device_matrix_data object.

The documentation for this struct was generated from the following file:

• ginkgo/core/base/device_matrix_data.hpp

43.11 gko::matrix::Hybrid< ValueType, IndexType >::automatic Class Reference

automatic is a strategy_type which decides the number of stored elements per row of the ell part automatically.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

• automatic ()

Creates an automatic strategy.

size_type compute_ell_num_stored_elements_per_row (array < size_type > *row_nnz) const override
 Computes the number of stored elements per row of the ell part.

43.11.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Hybrid< ValueType, IndexType >::automatic
```

automatic is a strategy_type which decides the number of stored elements per row of the ell part automatically.

43.11.2 Member Function Documentation

43.11.2.1 compute_ell_num_stored_elements_per_row()

Computes the number of stored elements per row of the ell part.

Parameters

row nnz	the number of nonzeros of each row

Returns

the number of stored elements per row of the ell part

Implements gko::matrix::Hybrid< ValueType, IndexType >::strategy type.

References gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_stored ← __elements_per_row().

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/hybrid.hpp

43.12 gko::BadDimension Class Reference

BadDimension is thrown if an operation is being applied to a LinOp with bad dimensions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• BadDimension (const std::string &file, int line, const std::string &func, const std::string &op_name, size_type op_num_rows, size_type op_num_cols, const std::string &clarification)

Initializes a bad dimension error.

43.12.1 Detailed Description

BadDimension is thrown if an operation is being applied to a LinOp with bad dimensions.

43.12.2 Constructor & Destructor Documentation

43.12.2.1 BadDimension()

Initializes a bad dimension error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The function name where the error occurred
op_name	The name of the operator
op_num_rows	The row dimension of the operator
op_num_cols	The column dimension of the operator
clarification	An additional message further describing the error

The documentation for this class was generated from the following file:

ginkgo/core/base/exception.hpp

43.13 gko::solver::Bicg < ValueType > Class Template Reference

BICG or the Biconjugate gradient method is a Krylov subspace solver.

#include <ginkgo/core/solver/bicg.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
 Returns a LinOp representing the conjugate transpose of the Transposable object.
- bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

43.13.1 Detailed Description

template<typename ValueType = default_precision> class gko::solver::Bicg< ValueType >

BICG or the Biconjugate gradient method is a Krylov subspace solver.

Being a generic solver, it is capable of solving general matrices, including non-s.p.d matrices. Though, the memory and the computational requirement of the BiCG solver are higher than of its s.p.d solver counterpart, it has the capability to solve generic systems.

BiCG is based on the bi-Lanczos tridiagonalization method and in exact arithmetic should terminate in at most N iterations (2N MV's, with A and $A^{\wedge}H$). It forms the basis of many of the cheaper methods such as BiCGSTAB and CGS.

Reference: R.Fletcher, Conjugate gradient methods for indefinite systems, doi: 10.1007/BFb0080116

Template Parameters

ValueType precision of matrix elements

43.13.2 Member Function Documentation

43.13.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Bicg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
100 { return true; }
```

43.13.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.13.2.3 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/bicg.hpp

43.14 gko::solver::Bicgstab < ValueType > Class Template Reference

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

#include <ginkgo/core/solver/bicgstab.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

· bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

43.14.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::Bicgstab< ValueType >
```

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

Being a generic solver, it is capable of solving general matrices, including non-s.p.d matrices. Though, the memory and the computational requirement of the BiCGSTAB solver are higher than of its s.p.d solver counterpart, it has the capability to solve generic systems. It was developed by stabilizing the BiCG method.

Template Parameters

ValueType	precision of the elements of the system matrix.
-----------	---

43.14.2 Member Function Documentation

43.14.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Bicgstab< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
99 { return true; }
```

43.14.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicgstab< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.14.2.3 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicgstab< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/solver/bicgstab.hpp

43.15 gko::preconditioner::block_interleaved_storage_scheme < IndexType > Struct Template Reference

Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.

#include <ginkgo/core/preconditioner/jacobi.hpp>

Public Member Functions

• IndexType get_group_size () const noexcept

Returns the number of elements in the group.

size_type compute_storage_space (size_type num_blocks) const noexcept

Computes the storage space required for the requested number of blocks.

• IndexType get_group_offset (IndexType block_id) const noexcept

Returns the offset of the group belonging to the block with the given ID.

IndexType get_block_offset (IndexType block_id) const noexcept

Returns the offset of the block with the given ID within its group.

IndexType get_global_block_offset (IndexType block_id) const noexcept

Returns the offset of the block with the given ID.

• IndexType get_stride () const noexcept

Returns the stride between columns of the block.

Public Attributes

IndexType block_offset

The offset between consecutive blocks within the group.

IndexType group_offset

The offset between two block groups.

• uint32 group_power

Then base 2 power of the group.

43.15.1 Detailed Description

```
template<typename IndexType>
struct gko::preconditioner::block_interleaved_storage_scheme< IndexType >
```

Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.

Template Parameters

IndexType type used for storing indices of the matrix

43.15.2 Member Function Documentation

43.15.2.1 compute_storage_space()

Computes the storage space required for the requested number of blocks.

Parameters

- 1		
	num blocks	the total number of blocks that needs to be stored
	HUHH DIOCKS	i the total number of blocks that needs to be stored

Returns

the total memory (as the number of elements) that need to be allocated for the scheme

Note

To simplify using the method in situations where the number of blocks is not known, for a special input $size \leftarrow _type\{\} - 1$ the method returns 0 to avoid overallocation of memory.

43.15.2.2 get_block_offset()

Returns the offset of the block with the given ID within its group.

Parameters

block←	the ID of the block
_id	

Returns

the offset of the block with ID block_id within its group

Referenced by gko::preconditioner::block_interleaved_storage_scheme < index_type >::get_global_block_offset().

43.15.2.3 get_global_block_offset()

Returns the offset of the block with the given ID.

Parameters

block⊷	the ID of the block
_id	

Returns

the offset of the block with ID $block_id$

43.15.2.4 get_group_offset()

Returns the offset of the group belonging to the block with the given ID.

Parameters

block⊷	the ID of the block
_id	

Returns

the offset of the group belonging to block with ID block_id

Referenced by gko::preconditioner::block_interleaved_storage_scheme < index_type >::get_global_block_offset().

43.15.2.5 get_group_size()

```
template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_group_size (
) const [inline], [noexcept]
```

Returns the number of elements in the group.

Returns

the number of elements in the group

Referenced by $gko::preconditioner::block_interleaved_storage_scheme < index_type >::compute_storage_ < space(), and <math>gko::preconditioner::block_interleaved_storage_scheme < index_type >::get_block_offset().$

43.15.2.6 get_stride()

```
template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_stride ( )
const [inline], [noexcept]
```

Returns the stride between columns of the block.

Returns

stride between columns of the block

43.15.3 Member Data Documentation

43.15.3.1 group power

```
template<typename IndexType>
uint32 gko::preconditioner::block_interleaved_storage_scheme< IndexType >::group_power
```

Then base 2 power of the group.

I.e. the group contains 1 << group_power elements.

Referenced by gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_group_offset(), gko ::preconditioner::block_interleaved_storage_scheme< index_type >::get_group_size(), and gko::preconditioner ::block interleaved storage scheme< index type >::get stride().

The documentation for this struct was generated from the following file:

• ginkgo/core/preconditioner/jacobi.hpp

43.16 gko::BlockSizeError< IndexType > Class Template Reference

Error that denotes issues between block sizes and matrix dimensions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• BlockSizeError (const std::string &file, const int line, const int block_size, const IndexType size)

43.16.1 Detailed Description

```
\label{template} \mbox{template}{<} \mbox{typename IndexType}{>} \\ \mbox{class gko::BlockSizeError}{<} \mbox{IndexType}{>} \\
```

Error that denotes issues between block sizes and matrix dimensions.

Template Parameters

IndexType Type of index used by the linear algebra object that is incompatible with the required block size.

43.16.2 Constructor & Destructor Documentation

43.16.2.1 BlockSizeError()

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
block_size	Size of small dense blocks in a matrix
size	The size that is not exactly divided by the block size

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.17 gko::solver::CbGmres < ValueType > Class Template Reference

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

```
#include <ginkgo/core/solver/cb_gmres.hpp>
```

Public Member Functions

size_type get_krylov_dim () const

Returns the Krylov dimension.

void set_krylov_dim (size_type other)

Sets the Krylov dimension.

• cb_gmres::storage_precision get_storage_precision () const

Returns the storage precision used internally.

43.17.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::CbGmres< ValueType >
```

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CB-GMRES are merged into 2 separate steps. Classical Gram-Schmidt with reorthogonalization is used.

The Krylov basis can be stored in reduced precision (compressed) to reduce memory accesses, while all computations (including Krylov basis operations) are performed in the same arithmetic precision ValueType. By default, the Krylov basis are stored in one precision lower than ValueType.

Template Parameters

ValueType	the arithmetic precision and the precision of matrix elements
-----------	---

43.17.2 Member Function Documentation

43.17.2.1 get_krylov_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::CbGmres< ValueType >::get_krylov_dim ( ) const [inline]
```

Returns the Krylov dimension.

Returns

the Krylov dimension

```
137 { return parameters_.krylov_dim; }
```

43.17.2.2 get_storage_precision()

```
template<typename ValueType = default_precision>
cb_gmres::storage_precision gko::solver::CbGmres< ValueType >::get_storage_precision ( ) const
[inline]
```

Returns the storage precision used internally.

Returns

the storage precision used internally

43.17.2.3 set_krylov_dim()

Sets the Krylov dimension.

Parameters

other	the new Krylov dimension
-------	--------------------------

The documentation for this class was generated from the following file:

ginkgo/core/solver/cb_gmres.hpp

43.18 gko::solver::Cg< ValueType > Class Template Reference

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

#include <ginkgo/core/solver/cg.hpp>

Public Member Functions

std::unique_ptr< LinOp > transpose () const override
 Returns a LinOp representing the transpose of the Transposable object.

- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

· bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

43.18.1 Detailed Description

template<typename ValueType = default_precision> class gko::solver::Cg< ValueType >

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

Though this method performs very well for symmetric positive definite matrices, it is in general not suitable for general matrices.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CG are merged into 2 separate steps.

Template Parameters

ValueType precision of matrix elements

43.18.2 Member Function Documentation

43.18.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Cg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
93 { return true; }
```

43.18.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.18.2.3 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/cg.hpp

43.19 gko::solver::Cgs< ValueType > Class Template Reference

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

```
#include <ginkgo/core/solver/cgs.hpp>
```

Public Member Functions

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

· bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

43.19.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::Cgs< ValueType >
```

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CGS are merged into 3 separate steps.

Template Parameters

ValueType	precision of matrix elements
-----------	------------------------------

43.19.2 Member Function Documentation

43.19.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Cgs< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
91 { return true; }
```

43.19.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cgs< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.19.2.3 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cgs< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/solver/cgs.hpp

43.20 gko::experimental::factorization::Cholesky< ValueType, IndexType > Class Template Reference

Computes a Cholesky factorization of a symmetric, positive-definite sparse matrix.

```
#include <ginkgo/core/factorization/cholesky.hpp>
```

Public Member Functions

 $\bullet \ \ \mathsf{std} :: \mathsf{unique_ptr} < \mathsf{factorization_type} > \mathsf{generate} \ (\mathsf{std} :: \mathsf{shared_ptr} < \mathsf{const} \ \mathsf{LinOp} > \mathsf{system_matrix}) \ \mathsf{const} \\$

Static Public Member Functions

static parameters_type build ()
 Creates a new parameter_type to set up the factory.

43.20.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::experimental::factorization::Cholesky< ValueType, IndexType>
```

Computes a Cholesky factorization of a symmetric, positive-definite sparse matrix.

This LinOpFactory returns a Factorization storing the L and L^AH factors for the provided system matrix in matrix::Csr format. If no symbolic factorization is provided, it will be computed first.

Template Parameters

ValueType	the type used to store values of the system matrix
IndexType	the type used to store sparsity pattern indices of the system matrix

43.20.2 Member Function Documentation

43.20.2.1 generate()

Note

This function overrides the default LinOpFactory::generate to return a Factorization instead of a generic LinOp, which would need to be cast to Factorization again to access its factors. It is only necessary because smart pointers aren't covariant.

The documentation for this class was generated from the following file:

· ginkgo/core/factorization/cholesky.hpp

43.21 gko::matrix::Csr< ValueType, IndexType >::classical Class Reference

classical is a strategy_type which uses the same number of threads on each row.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

· classical ()

Creates a classical strategy.

- void process (const array < index_type > &mtx_row_ptrs, array < index_type > *mtx_srow) override
 Computes srow according to row pointers.
- int64_t clac_size (const int64_t nnz) override

Computes the srow size according to the number of nonzeros.

- $std::shared_ptr < strategy_type > copy$ () override

Copy a strategy.

43.21.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Csr< ValueType, IndexType >::classical

classical is a strategy_type which uses the same number of threads on each row.

Classical strategy uses multithreads to calculate on parts of rows and then do a reduction of these threads results. The number of threads per row depends on the max number of stored elements per row.

43.21.2 Member Function Documentation

43.21.2.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

```
nnz the number of nonzeros
```

Returns

the size of srow

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.21.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::classical::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.21.2.3 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix
mtx_srow	the srow of the matrix

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

References gko::array< ValueType >::get_const_data(), gko::array< ValueType >::get_executor(), and gko ::array< ValueType >::get_num_elems().

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/csr.hpp

43.22 gko::matrix::Hybrid< ValueType, IndexType >::column_limit Class Reference

column_limit is a strategy_type which decides the number of stored elements per row of the ell part by specifying the number of columns.

#include <ginkgo/core/matrix/hybrid.hpp>

Public Member Functions

- column_limit (size_type num_column=0)
 Creates a column_limit strategy.
- size_type compute_ell_num_stored_elements_per_row (array< size_type > *row_nnz) const override

 Computes the number of stored elements per row of the ell part.
- auto get_num_columns () const

Get the number of columns limit.

43.22.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Hybrid< ValueType, IndexType >::column_limit

column_limit is a strategy_type which decides the number of stored elements per row of the ell part by specifying the number of columns.

43.22.2 Constructor & Destructor Documentation

43.22.2.1 column limit()

Creates a column_limit strategy.

Parameters

num_column	the specified number of columns of the ell part
------------	---

43.22.3 Member Function Documentation

43.22.3.1 compute_ell_num_stored_elements_per_row()

Computes the number of stored elements per row of the ell part.

Parameters

row_nnz the numbe	r of nonzeros of each row
-------------------	---------------------------

Returns

the number of stored elements per row of the ell part

 $Implements\ gko::matrix::Hybrid< ValueType,\ IndexType>::strategy_type.$

43.22.3.2 get_num_columns()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::column_limit::get_num_columns ( ) const
[inline]
```

Get the number of columns limit.

Returns

the number of columns limit

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/hybrid.hpp

43.23 gko::Combination < ValueType > Class Template Reference

The Combination class can be used to construct a linear combination of multiple linear operators c1 * op1 + c2 * op2 + ...

#include <ginkgo/core/base/combination.hpp>

Public Member Functions

- const std::vector < std::shared_ptr < const LinOp > > & get_coefficients () const noexcept
 Returns a list of coefficients of the combination.
- const std::vector< std::shared_ptr< const LinOp > > & get_operators () const noexcept Returns a list of operators of the combination.
- std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

Combination & operator= (const Combination &)

Copy-assigns a Combination.

Combination & operator= (Combination &&)

Move-assigns a Combination.

Combination (const Combination &)

Copy-constructs a Combination.

Combination (Combination &&)

Move-constructs a Combination.

43.23.1 Detailed Description

```
\label{template} \mbox{typename ValueType = default\_precision} > \mbox{class gko::Combination} < \mbox{ValueType} > \mbox{}
```

The Combination class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$

• ck * opk.

Combination ensures that all LinOps passed to its constructor use the same executor, and if not, copies the operators to the executor of the first operator.

Template Parameters

ValueType | precision of input and result vectors

43.23.2 Constructor & Destructor Documentation

43.23.2.1 Combination() [1/2]

Copy-constructs a Combination.

This inherits the executor of the input Combination and all of its operators with shared ownership.

43.23.2.2 Combination() [2/2]

Move-constructs a Combination.

This inherits the executor of the input Combination and all of its operators. The moved-from object is empty (0x0 LinOp without operators) afterwards.

43.23.3 Member Function Documentation

43.23.3.1 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Combination< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.23.3.2 get_coefficients()

```
template<typename ValueType = default_precision> const std::vector<std::shared_ptr<const LinOp> >& gko::Combination< ValueType >::get_ \leftarrow coefficients ( ) const [inline], [noexcept]
```

Returns a list of coefficients of the combination.

Returns

a list of coefficients

```
76 {
77     return coefficients_;
78 }
```

43.23.3.3 get_operators()

```
template<typename ValueType = default_precision> const std::vector<std::shared_ptr<const LinOp> >& gko::Combination< ValueType >::get_\leftarrow operators ( ) const [inline], [noexcept]
```

Returns a list of operators of the combination.

Returns

a list of operators

43.23.3.4 operator=() [1/2]

Move-assigns a Combination.

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor, otherwise they share ownership. The moved-from object is empty (0x0 LinOp without operators) afterwards.

43.23.3.5 operator=() [2/2]

Copy-assigns a Combination.

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor.

43.23.3.6 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Combination< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/base/combination.hpp

43.24 gko::stop::Combined Class Reference

The Combined class is used to combine multiple criterions together through an OR operation.

```
#include <ginkgo/core/stop/combined.hpp>
```

43.24.1 Detailed Description

The Combined class is used to combine multiple criterions together through an OR operation.

The typical use case is to stop the iteration process if any of the criteria is fulfilled, e.g. a number of iterations, the relative residual norm has reached a threshold, etc.

The documentation for this class was generated from the following file:

• ginkgo/core/stop/combined.hpp

43.25 gko::experimental::mpi::communicator Class Reference

A thin wrapper of MPI_Comm that supports most MPI calls.

#include <ginkgo/core/base/mpi.hpp>

Public Member Functions

communicator (const MPI_Comm &comm, bool force_host_buffer=false)

Non-owning constructor for an existing communicator of type MPI Comm.

communicator (const MPI_Comm &comm, int color, int key)

Create a communicator object from an existing MPI_Comm object using color and key.

communicator (const communicator &comm, int color, int key)

Create a communicator object from an existing MPI_Comm object using color and key.

const MPI_Comm & get () const

Return the underlying MPI_Comm object.

· int size () const

Return the size of the communicator (number of ranks).

• int rank () const

Return the rank of the calling process in the communicator.

• int node local rank () const

Return the node local rank of the calling process in the communicator.

bool operator== (const communicator &rhs) const

Compare two communicator objects for equality.

bool operator!= (const communicator &rhs) const

Compare two communicator objects for non-equality.

· void synchronize () const

This function is used to synchronize the ranks in the communicator.

template<typename SendType >

void send (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, const int destination_rank, const int send_tag) const

Send (Blocking) data from calling process to destination rank.

 $\bullet \ \ \text{template}{<} \text{typename SendType} >$

request i_send (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_← count, const int destination rank, const int send tag) const

Send (Non-blocking, Immediate return) data from calling process to destination rank.

template<typename RecvType >

status recv (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count, const int source_rank, const int recv_tag) const

Receive data from source rank.

template<typename RecvType >

request i_recv (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count, const int source_rank, const int recv_tag) const

Receive (Non-blocking, Immediate return) data from source rank.

template<typename BroadcastType >

void broadcast (std::shared_ptr< const Executor > exec, BroadcastType *buffer, int count, int root_rank) const

Broadcast data from calling process to all ranks in the communicator.

template<typename BroadcastType >

request i_broadcast (std::shared_ptr< const Executor > exec, BroadcastType *buffer, int count, int root_rank) const

(Non-blocking) Broadcast data from calling process to all ranks in the communicator

template<typename ReduceType >

void reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv← _buffer, int count, MPI_Op operation, int root_rank) const

Reduce data into root from all calling processes on the same communicator.

template<typename ReduceType >

request i_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv_buffer, int count, MPI_Op operation, int root_rank) const

(Non-blocking) Reduce data into root from all calling processes on the same communicator.

template<typename ReduceType >
 void all_reduce (std::shared_ptr< const Executor > exec, ReduceType *recv_buffer, int count, MPI_Op operation) const

(In-place) Reduce data from all calling processes from all calling processes on same communicator.

template<typename ReduceType >
 request i_all_reduce (std::shared_ptr< const Executor > exec, ReduceType *recv_buffer, int count, MPI_Op
 operation) const

(In-place, non-blocking) Reduce data from all calling processes from all calling processes on same communicator.

template<typename ReduceType >
 void all_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType
 *recv buffer, int count, MPI Op operation) const

Reduce data from all calling processes from all calling processes on same communicator.

template<typename ReduceType >
 request i_all_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType
 *recv_buffer, int count, MPI_Op operation) const

Reduce data from all calling processes from all calling processes on same communicator.

template<typename SendType, typename RecvType >
 void gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count,
 RecvType *recv_buffer, const int recv_count, int root_rank) const

Gather data onto the root rank from all ranks in the communicator.

• template<typename SendType , typename RecvType > request i_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_← count, RecvType *recv_buffer, const int recv_count, int root_rank) const

(Non-blocking) Gather data onto the root rank from all ranks in the communicator.

template<typename SendType, typename RecvType >
 void gather_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count,
 RecvType *recv buffer, const int *recv counts, const int *displacements, int root rank) const

Gather data onto the root rank from all ranks in the communicator with offsets.

template<typename SendType , typename RecvType >
 request i_gather_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send
 _count, RecvType *recv_buffer, const int *recv_counts, const int *displacements, int root_rank) const

template<typename SendType , typename RecvType >
 void all_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_←
 count, RecvType *recv_buffer, const int recv_count) const

(Non-blocking) Gather data onto the root rank from all ranks in the communicator with offsets.

Gather data onto all ranks from all ranks in the communicator.

template<typename SendType, typename RecvType >
 request i_all_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count) const

(Non-blocking) Gather data onto all ranks from all ranks in the communicator.

template<typename SendType, typename RecvType >
 void scatter (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count,
 RecvType *recv_buffer, const int recv_count, int root_rank) const

Scatter data from root rank to all ranks in the communicator.

template < typename SendType , typename RecvType >
 request i_scatter (std::shared_ptr < const Executor > exec, const SendType *send_buffer, const int send_
 count, RecvType *recv buffer, const int recv count, int root rank) const

(Non-blocking) Scatter data from root rank to all ranks in the communicator.

template<typename SendType , typename RecvType >
 void scatter_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int *send_←
 counts, const int *displacements, RecvType *recv buffer, const int recv count, int root rank) const

Scatter data from root rank to all ranks in the communicator with offsets.

template<typename SendType, typename RecvType >
 request i_scatter_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int
 *send_counts, const int *displacements, RecvType *recv_buffer, const int recv_count, int root_rank) const

(Non-blocking) Scatter data from root rank to all ranks in the communicator with offsets.

template<typename RecvType >
 void all_to_all (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count) const
 (In-place) Communicate data from all ranks to all other ranks in place (MPI_Alltoall).

template<typename RecvType >
 request i_all_to_all (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count)
 const

(In-place, Non-blocking) Communicate data from all ranks to all other ranks in place (MPI_lalltoall).

template<typename SendType , typename RecvType >
 void all_to_all (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count,
 RecvType *recv_buffer, const int recv_count) const

Communicate data from all ranks to all other ranks (MPI_Alltoall).

template < typename SendType , typename RecvType >
 request i_all_to_all (std::shared_ptr < const Executor > exec, const SendType *send_buffer, const int send count, RecvType *recv_buffer, const int recv_count) const

(Non-blocking) Communicate data from all ranks to all other ranks (MPI_Ialltoall).

template<typename SendType, typename RecvType >
 void all_to_all_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int *send
 _counts, const int *send_offsets, RecvType *recv_buffer, const int *recv_counts, const int *recv_offsets)
 const

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

void all_to_all_v (std::shared_ptr< const Executor > exec, const void *send_buffer, const int *send_← counts, const int *send_offsets, MPI_Datatype send_type, void *recv_buffer, const int *recv_counts, const int *recv_offsets, MPI_Datatype recv_type) const

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

request i_all_to_all_v (std::shared_ptr< const Executor > exec, const void *send_buffer, const int *send←
 _counts, const int *send_offsets, MPI_Datatype send_type, void *recv_buffer, const int *recv_counts, const
 int *recv offsets, MPI_Datatype recv_type) const

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

template<typename SendType, typename RecvType >
 request i_all_to_all_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int
 *send_counts, const int *send_offsets, RecvType *recv_buffer, const int *recv_counts, const int *recv_counts, const int *recv_counts

Communicate data from all ranks to all other ranks with offsets (MPI_Ialltoallv).

template<typename ScanType >
 void scan (std::shared_ptr< const Executor > exec, const ScanType *send_buffer, ScanType *recv_buffer,
 int count, MPI_Op operation) const

Does a scan operation with the given operator.

template < typename ScanType >
 request i_scan (std::shared_ptr < const Executor > exec, const ScanType *send_buffer, ScanType *recv_
 buffer, int count, MPI_Op operation) const

Does a scan operation with the given operator.

43.25.1 Detailed Description

A thin wrapper of MPI Comm that supports most MPI calls.

A wrapper class that takes in the given MPI communicator. If a bare MPI_Comm is provided, the wrapper takes no ownership of the MPI_Comm. Thus the MPI_Comm must remain valid throughout the lifetime of the communicator. If the communicator was created through splitting, the wrapper takes ownership of the MPI_Comm. In this case, as the class or object goes out of scope, the underlying MPI_Comm is freed.

Note

All MPI calls that work on a buffer take in an Executor as an additional argument. This argument specifies the memory space the buffer lives in.

43.25.2 Constructor & Destructor Documentation

43.25.2.1 communicator() [1/3]

Non-owning constructor for an existing communicator of type MPI_Comm.

The MPI_Comm object will not be deleted after the communicator object has been freed and an explicit MPI_← Comm_free needs to be called on the original MPI_Comm object.

Parameters

comm	The input MPI_Comm object.
force_host_buffer	If set to true, always communicates through host memory

43.25.2.2 communicator() [2/3]

Create a communicator object from an existing MPI_Comm object using color and key.

Parameters

comm	The input MPI_Comm object.
color	The color to split the original comm object
key	The key to split the comm object

43.25.2.3 communicator() [3/3]

```
int color,
int key ) [inline]
```

Create a communicator object from an existing MPI_Comm object using color and key.

Parameters

comm	The input communicator object.
color	The color to split the original comm object
key	The key to split the comm object

References get().

43.25.3 Member Function Documentation

43.25.3.1 all_gather()

Gather data onto all ranks from all ranks in the communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

References get().

43.25.3.2 all_reduce() [1/2]

```
template<typename ReduceType >
void gko::experimental::mpi::communicator::all_reduce (
    std::shared_ptr< const Executor > exec,
    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Reduce data from all calling processes from all calling processes on same communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the data to reduce
recv_buffer	the reduced result
count	the number of elements to reduce
operation	the reduce operation. See @MPI_Op

Template Parameters

ReduceType	the type of the data to send. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

References get().

43.25.3.3 all_reduce() [2/2]

(In-place) Reduce data from all calling processes from all calling processes on same communicator.

Parameters

exec	The executor, on which the message buffer is located.
recv_buffer	the data to reduce and the reduced result
count	the number of elements to reduce
operation	the MPI_Op type reduce operation.

Template Parameters

ReduceType	the type of the data to send. Has to be a type which has a specialization of type_impl that	1
	defines its MPI_Datatype.	

References get().

43.25.3.4 all_to_all() [1/2]

Communicate data from all ranks to all other ranks (MPI_Alltoall).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
recv_buffer	the buffer to receive
recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

References get().

43.25.3.5 all_to_all() [2/2]

(In-place) Communicate data from all ranks to all other ranks in place (MPI_Alltoall).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffer is located.
buffer	the buffer to send and the buffer receive
Generated by Doxy	the number of elements to receive
comm	the communicator

Template Parameters

RecvType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines	
	its MPI_Datatype.	

Note

This overload uses MPI_IN_PLACE and the source and destination buffers are the same.

References get().

43.25.3.6 all_to_all_v() [1/2]

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
send_offsets	the offsets for the send buffer
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
recv_offsets	the offsets for the recv buffer
comm	the communicator

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

43.25.3.7 all_to_all_v() [2/2]

```
void gko::experimental::mpi::communicator::all_to_all_v (
    std::shared_ptr< const Executor > exec,
    const void * send_buffer,
    const int * send_counts,
    const int * send_offsets,
    MPI_Datatype send_type,
    void * recv_buffer,
    const int * recv_counts,
    const int * recv_counts,
    const int * recv_offsets,
    MPI_Datatype recv_type ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

See MPI documentation for more details.

Parameters

	The average of the second seco
exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
send_offsets	the offsets for the send buffer
send_type	the MPI_Datatype for the send buffer
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
recv_offsets	the offsets for the recv buffer
recv_type	the MPI_Datatype for the recv buffer
comm	the communicator

References get().

43.25.3.8 broadcast()

Broadcast data from calling process to all ranks in the communicator.

Parameters

exec	The executor, on which the message buffer is located.
buffer	the buffer to broadcsat
count	the number of elements to broadcast
root rank	the rank to broadcast from

Template Parameters

BroadcastType	the type of the data to broadcast. Has to be a type which has a specialization of type_impl
	that defines its MPI_Datatype.

References get().

43.25.3.9 gather()

Gather data onto the root rank from all ranks in the communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
root_rank	the rank to gather into

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

References get().

43.25.3.10 gather_v()

```
const int * recv_counts,
const int * displacements,
int root_rank ) const [inline]
```

Gather data onto the root rank from all ranks in the communicator with offsets.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
displacements	the offsets for the buffer
root_rank	the rank to gather into

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

References get().

43.25.3.11 get()

```
const MPI_Comm& gko::experimental::mpi::communicator::get ( ) const [inline]
```

Return the underlying MPI_Comm object.

Returns

the MPI_Comm object

Referenced by all_gather(), all_reduce(), all_to_all(), all_to_all_v(), broadcast(), communicator(), gather(), _v(), i_all_gather(), i_all_reduce(), i_all_to_all(), i_all_to_all_v(), i_broadcast(), i_gather(), i_gather_v(), i_recv(), i_reduce(), i_scan(), i_scatter(), i_scatter_v(), i_send(), recv(), reduce(), scan(), scatter(), scatter_v(), send(), synchronize(), and gko::experimental::mpi::window< ValueType >::window().

43.25.3.12 i_all_gather()

(Non-blocking) Gather data onto all ranks from all ranks in the communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.13 i_all_reduce() [1/2]

```
template<typename ReduceType >
request gko::experimental::mpi::communicator::i_all_reduce (
    std::shared_ptr< const Executor > exec,
    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Reduce data from all calling processes from all calling processes on same communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the data to reduce
recv_buffer	the reduced result
count	the number of elements to reduce
operation	the reduce operation. See @MPI_Op

Template Parameters

ReduceType	the type of the data to reduce. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.14 i_all_reduce() [2/2]

(In-place, non-blocking) Reduce data from all calling processes from all calling processes on same communicator.

Parameters

exec	The executor, on which the message buffer is located.
recv_buffer	the data to reduce and the reduced result
count	the number of elements to reduce
operation	the reduce operation. See @MPI_Op

Template Parameters

ReduceType	the type of the data to reduce. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.15 i_all_to_all() [1/2]

(Non-blocking) Communicate data from all ranks to all other ranks (MPI_lalltoall).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
recv_buffer	the buffer to receive
recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.16 i_all_to_all() [2/2]

(In-place, Non-blocking) Communicate data from all ranks to all other ranks in place (MPI_Ialltoall).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffer is located.
buffer	the buffer to send and the buffer receive
recv_count	the number of elements to receive
comm	the communicator

Template Parameters

RecvType	the type of the data to receive. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

Returns

the request handle for the call

Note

This overload uses MPI_IN_PLACE and the source and destination buffers are the same.

References gko::experimental::mpi::request::get(), and get().

43.25.3.17 i_all_to_all_v() [1/2]

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
send_offsets	the offsets for the send buffer
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
recv_offsets	the offsets for the recv buffer

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References i_all_to_all_v().

43.25.3.18 i_all_to_all_v() [2/2]

```
const void * send_buffer,
const int * send_counts,
const int * send_offsets,
MPI_Datatype send_type,
void * recv_buffer,
const int * recv_counts,
const int * recv_offsets,
MPI_Datatype recv_type ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_Ialltoallv).

See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to send
send_count	the number of elements to send
send_offsets	the offsets for the send buffer
send_type	the MPI_Datatype for the send buffer
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
recv_offsets	the offsets for the recv buffer
recv_type	the MPI_Datatype for the recv buffer

Returns

the request handle for the call

Note

This overload allows specifying the MPI_Datatype for both the send and received data.

References gko::experimental::mpi::request::get(), and get().

Referenced by i_all_to_all_v().

43.25.3.19 i_broadcast()

(Non-blocking) Broadcast data from calling process to all ranks in the communicator

Parameters

exec	The executor, on which the message buffer is located.
buffer	the buffer to broadcsat
count	the number of elements to broadcast
root_rank	the rank to broadcast from

Template Parameters

BroadcastType	the type of the data to broadcast. Has to be a type which has a specialization of type_impl
	that defines its MPI_Datatype.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.20 i_gather()

(Non-blocking) Gather data onto the root rank from all ranks in the communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
root_rank	the rank to gather into

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.21 i_gather_v()

(Non-blocking) Gather data onto the root rank from all ranks in the communicator with offsets.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
displacements	the offsets for the buffer
root_rank	the rank to gather into

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.22 i_recv()

```
template<typename RecvType >
request gko::experimental::mpi::communicator::i_recv (
```

```
std::shared_ptr< const Executor > exec,
RecvType * recv_buffer,
const int recv_count,
const int source_rank,
const int recv_tag ) const [inline]
```

Receive (Non-blocking, Immediate return) data from source rank.

Parameters

exec	The executor, on which the message buffer is located.
recv_buffer	the buffer to send
recv_count	the number of elements to receive
source_rank	the rank to receive the data from
recv_tag	the tag for the recv call

Template Parameters

RecvType	the type of the data to receive. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

Returns

the request handle for the recv call

References gko::experimental::mpi::request::get(), and get().

43.25.3.23 i_reduce()

(Non-blocking) Reduce data into root from all calling processes on the same communicator.

Parameters

exec	The executor, on which the message buffer is located.
send_buffer	the buffer to reduce
recv_buffer	the reduced result
count	the number of elements to reduce
operation	the MPI_Op type reduce operation.

Template Parameters

ReduceType	the type of the data to reduce. Has to be a type which has a specialization of type_impl that	1
	defines its MPI_Datatype.	

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.24 i_scan()

Does a scan operation with the given operator.

(MPI_Iscan). See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to scan from
recv_buffer	the result buffer
recv_count	the number of elements to scan
operation	the operation type to be used for the scan. See @MPI_Op

Template Parameters

ScanType	the type of the data to scan. Has to be a type which has a specialization of type_impl that defines
	its MPI_Datatype.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.25 i_scatter()

(Non-blocking) Scatter data from root rank to all ranks in the communicator.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.26 i_scatter_v()

(Non-blocking) Scatter data from root rank to all ranks in the communicator with offsets.

Parameters

<i>exec</i> The executor, on which the message buffers are located.

Parameters

send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
displacements	the offsets for the buffer
comm	the communicator

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

43.25.3.27 i_send()

Send (Non-blocking, Immediate return) data from calling process to destination rank.

Parameters

	-
exec	The executor, on which the message buffer is located.
send buffer	the buffer to send
	4b
send_count	the number of elements to send
destination_rank	the rank to send the data to
send_tag	the tag for the send call

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that define	
	its MPI_Datatype.	

Returns

the request handle for the send call

References gko::experimental::mpi::request::get(), and get().

43.25.3.28 node_local_rank()

```
int gko::experimental::mpi::communicator::node_local_rank ( ) const [inline]
```

Return the node local rank of the calling process in the communicator.

Returns

the node local rank

43.25.3.29 operator"!=()

Compare two communicator objects for non-equality.

Returns

if the two comm objects are not equal

43.25.3.30 operator==()

Compare two communicator objects for equality.

Returns

if the two comm objects are equal

43.25.3.31 rank()

```
int gko::experimental::mpi::communicator::rank ( ) const [inline]
```

Return the rank of the calling process in the communicator.

Returns

the rank

43.25.3.32 recv()

Receive data from source rank.

Parameters

exec	The executor, on which the message buffer is located.
recv_buffer	the buffer to receive
recv_count	the number of elements to receive
source_rank	the rank to receive the data from
recv_tag	the tag for the recv call

Template Parameters

RecvType	the type of the data to receive. Has to be a type which has a specialization of type_impl that
	defines its MPI_Datatype.

Returns

the status of completion of this call

 $References\ gko::experimental::mpi::status::get(),\ and\ get().$

43.25.3.33 reduce()

```
template<typename ReduceType >
void gko::experimental::mpi::communicator::reduce (
```

```
std::shared_ptr< const Executor > exec,
const ReduceType * send_buffer,
ReduceType * recv_buffer,
int count,
MPI_Op operation,
int root_rank ) const [inline]
```

Reduce data into root from all calling processes on the same communicator.

Parameters

exec	The executor, on which the message buffer is located.
send_buffer	the buffer to reduce
recv_buffer	the reduced result
count	the number of elements to reduce
operation	the MPI_Op type reduce operation.

Template Parameters

ReduceType	the type of the data to reduce. Has to be a type which has a specialization of type_impl that]
	defines its MPI_Datatype.	

References get().

43.25.3.34 scan()

```
template<typename ScanType >
void gko::experimental::mpi::communicator::scan (
    std::shared_ptr< const Executor > exec,
    const ScanType * send_buffer,
    ScanType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Does a scan operation with the given operator.

(MPI_Scan). See MPI documentation for more details.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to scan from
recv_buffer	the result buffer
recv_count	the number of elements to scan
operation	the operation type to be used for the scan. See @MPI_Op

Template Parameters

ScanType	the type of the data to scan. Has to be a type which has a specialization of type_impl that defines	
	its MPI_Datatype.	

References get().

43.25.3.35 scatter()

Scatter data from root rank to all ranks in the communicator.

Parameters

	exec	The executor, on which the message buffers are located.
	send_buffer	the buffer to gather from
	send_count	the number of elements to send
recv_buffer the buffer to gather into		the buffer to gather into
	recv_count	the number of elements to receive

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.	
RecvType	the type of the data to receive. The same restrictions as for SendType apply.	

References get().

43.25.3.36 scatter_v()

```
const int recv_count,
int root_rank ) const [inline]
```

Scatter data from root rank to all ranks in the communicator with offsets.

Parameters

exec	The executor, on which the message buffers are located.
send_buffer	the buffer to gather from
send_count	the number of elements to send
recv_buffer	the buffer to gather into
recv_count	the number of elements to receive
displacements	the offsets for the buffer
comm	the communicator

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
RecvType	the type of the data to receive. The same restrictions as for SendType apply.

References get().

43.25.3.37 send()

```
template<typename SendType >
void gko::experimental::mpi::communicator::send (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    const int destination_rank,
    const int send_tag ) const [inline]
```

Send (Blocking) data from calling process to destination rank.

Parameters

exec	The executor, on which the message buffer is located.
send_buffer	the buffer to send
send_count	the number of elements to send
destination_rank	the rank to send the data to
send_tag	the tag for the send call

Template Parameters

SendType	the type of the data to send. Has to be a type which has a specialization of type_impl that defines	
	its MPI_Datatype.	

References get().

43.25.3.38 size()

```
int gko::experimental::mpi::communicator::size ( ) const [inline]
```

Return the size of the communicator (number of ranks).

Returns

the size

43.25.3.39 synchronize()

```
void gko::experimental::mpi::communicator::synchronize ( ) const [inline]
```

This function is used to synchronize the ranks in the communicator.

Calls MPI_Barrier

References get().

The documentation for this class was generated from the following file:

· ginkgo/core/base/mpi.hpp

43.26 gko::Composition < ValueType > Class Template Reference

```
The Composition class can be used to compose linear operators op1, op2, ..., opn and obtain the operator op1*op2*...
```

#include <ginkgo/core/base/composition.hpp>

Public Member Functions

- const std::vector< std::shared_ptr< const LinOp > > & get_operators () const noexcept Returns a list of operators of the composition.
- std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

Composition & operator= (const Composition &)

Copy-assigns a Composition.

· Composition & operator= (Composition &&)

Move-assigns a Composition.

• Composition (const Composition &)

Copy-constructs a Composition.

Composition (Composition &&)

Move-constructs a Composition.

43.26.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::Composition< ValueType >
```

The Composition class can be used to compose linear operators op1, op2, ..., opn and obtain the operator op1 * op2 * ...

· opn.

All LinOps of the Composition must operate on Dense inputs. For an operator op_k that require an initial guess for their apply, Composition provides either

- the output of the previous op_{k+1}->apply if op_k has square dimension
- zero if op_k is rectangular as an initial guess.

Composition ensures that all LinOps passed to its constructor use the same executor, and if not, copies the operators to the executor of the first operator.

Template Parameters

```
ValueType precision of input and result vectors
```

43.26.2 Constructor & Destructor Documentation

43.26.2.1 Composition() [1/2]

Copy-constructs a Composition.

This inherits the executor of the input Composition and all of its operators with shared ownership.

43.26.2.2 Composition() [2/2]

Move-constructs a Composition.

This inherits the executor of the input Composition and all of its operators. The moved-from object is empty (0x0 LinOp without operators) afterwards.

43.26.3 Member Function Documentation

43.26.3.1 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Composition< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.26.3.2 get operators()

```
template<typename ValueType = default_precision>
const std::vector<std::shared_ptr<const LinOp> >& gko::Composition< ValueType >::get_
operators ( ) const [inline], [noexcept]
```

Returns a list of operators of the composition.

Returns

a list of operators

```
84 {
85          return operators_;
86 }
```

43.26.3.3 operator=() [1/2]

Move-assigns a Composition.

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor, otherwise they share ownership. The moved-from object is empty (0x0 LinOp without operators) afterwards.

43.26.3.4 operator=() [2/2]

Copy-assigns a Composition.

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor.

43.26.3.5 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Composition< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/base/composition.hpp

43.27 gko::experimental::mpi::contiguous type Class Reference

A move-only wrapper for a contiguous MPI_Datatype.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

• contiguous_type (int count, MPI_Datatype old_type)

Constructs a wrapper for a contiguous MPI_Datatype.

• contiguous_type ()

Constructs empty wrapper with MPI_DATATYPE_NULL.

contiguous_type (const contiguous_type &)=delete

Disallow copying of wrapper type.

contiguous_type & operator= (const contiguous_type &)=delete

Disallow copying of wrapper type.

• contiguous_type (contiguous_type &&other) noexcept

Move constructor, leaves other with MPI_DATATYPE_NULL.

• contiguous_type & operator= (contiguous_type &&other) noexcept

Move assignment, leaves other with MPI_DATATYPE_NULL.

~contiguous_type ()

Destructs object by freeing wrapped MPI Datatype.

MPI_Datatype get () const

Access the underlying MPI_Datatype.

43.27.1 Detailed Description

A move-only wrapper for a contiguous MPI_Datatype.

The underlying MPI_Datatype is automatically created and committed when an object of this type is constructed, and freed when it is destructed.

43.27.2 Constructor & Destructor Documentation

43.27.2.1 contiguous_type() [1/2]

Constructs a wrapper for a contiguous MPI_Datatype.

Parameters

count the number of old_type elements the new datatype		the number of old_type elements the new datatype contains.
	old_type	the MPI_Datatype that is contained.

43.27.2.2 contiguous_type() [2/2]

Move constructor, leaves other with MPI DATATYPE NULL.

Parameters

```
other to be moved from object.
```

43.27.3 Member Function Documentation

43.27.3.1 get()

```
MPI_Datatype gko::experimental::mpi::contiguous_type::get ( ) const [inline]
```

Access the underlying MPI_Datatype.

Returns

the underlying MPI_Datatype.

43.27.3.2 operator=()

Move assignment, leaves other with MPI DATATYPE NULL.

Parameters

Returns

this object.

The documentation for this class was generated from the following file:

· ginkgo/core/base/mpi.hpp

43.28 gko::log::Convergence< ValueType > Class Template Reference

Convergence is a Logger which logs data strictly from the <code>criterion_check_completed</code> event.

```
#include <ginkgo/core/log/convergence.hpp>
```

Public Member Functions

· bool has converged () const noexcept

Returns true if the solver has converged.

• void reset_convergence_status ()

Resets the convergence status to false.

const size_type & get_num_iterations () const noexcept

Returns the number of iterations.

- const LinOp * get_residual () const noexcept

Returns the residual.

const LinOp * get_residual_norm () const noexcept

Returns the residual norm.

const LinOp * get_implicit_sq_resnorm () const noexcept

Returns the implicit squared residual norm.

Static Public Member Functions

 static std::unique_ptr< Convergence > create (std::shared_ptr< const Executor >, const mask_type &enabled_events=Logger::criterion_events_mask|Logger::iteration_complete_mask)

Creates a convergence logger.

static std::unique_ptr< Convergence > create (const mask_type &enabled_events=Logger::criterion_
 events_mask|Logger::iteration_complete_mask)

Creates a convergence logger.

43.28.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::log::Convergence< ValueType >
```

Convergence is a Logger which logs data strictly from the criterion_check_completed event.

The purpose of this logger is to give a simple access to standard data generated by the solver once it has stopped with minimal overhead.

This logger also computes the residual norm from the residual when the residual norm was not available. This can add some slight overhead.

43.28.2 Member Function Documentation

43.28.2.1 create() [1/2]

Creates a convergence logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

```
enabled_events the events enabled for this logger. By default all events.
```

Returns

an std::unique ptr to the the constructed object

```
130 {
131     return std::unique_ptr<Convergence>(new Convergence(enabled_events));
132 }
```

43.28.2.2 create() [2/2]

Creates a convergence logger.

This dynamically allocates the memory, constructs the object and returns an std::unique ptr to this object.

Parameters

exec	the executor
enabled_events	the events enabled for this logger. By default all events.

Returns

an std::unique_ptr to the the constructed object

43.28.2.3 get_implicit_sq_resnorm()

```
template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_implicit_sq_resnorm ( ) const [inline],
[noexcept]
```

Returns the implicit squared residual norm.

Returns

the implicit squared residual norm

43.28.2.4 get_num_iterations()

```
template<typename ValueType = default_precision>
const size_type& gko::log::Convergence< ValueType >::get_num_iterations ( ) const [inline],
[noexcept]
```

Returns the number of iterations.

Returns

the number of iterations

43.28.2.5 get_residual()

```
template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_residual ( ) const [inline], [noexcept]
```

Returns the residual.

Returns

the residual

43.28.2.6 get_residual_norm()

```
template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_residual_norm ( ) const [inline], [noexcept]
```

Returns the residual norm.

Returns

the residual norm

43.28.2.7 has_converged()

```
template<typename ValueType = default_precision>
bool gko::log::Convergence< ValueType >::has_converged ( ) const [inline], [noexcept]
```

Returns true if the solver has converged.

Returns

the bool flag for convergence status

The documentation for this class was generated from the following file:

• ginkgo/core/log/convergence.hpp

43.29 gko::ConvertibleTo< ResultType > Class Template Reference

ConvertibleTo interface is used to mark that the implementer can be converted to the object of ResultType.

#include <ginkgo/core/base/polymorphic_object.hpp>

Public Member Functions

- virtual void convert_to (result_type *result) const =0
 - Converts the implementer to an object of type result_type.
- virtual void move_to (result_type *result)=0

Converts the implementer to an object of type result_type by moving data from this object.

43.29.1 Detailed Description

```
template<typename ResultType> class gko::ConvertibleTo< ResultType >
```

Convertible To interface is used to mark that the implementer can be converted to the object of Result Type.

This interface is used to enable conversions between polymorphic objects. To mark that an object of type $\tt U$ can be converted to an object of type $\tt V$, $\tt U$ should implement ConvertibleTo<V>. Then, the implementation of PolymorphicObject::copy_from automatically generated by EnablePolymorphicObject mixin will use RTTI to figure out that $\tt U$ implements the interface and convert it using the convert_to / move_to methods of the interface.

As an example, the following function:

```
{c++}
void my_function(const U *u, V *v) {
  v->copy_from(u);
}
```

will convert object u to object v by checking that u can be dynamically casted to ConvertibleTo\<V>, and calling ConvertibleTo<V>::convert_to(V*)` to do the actual conversion.

In case u is passed as a unique_ptr, call to <code>convert_to</code> will be replaced by a call to <code>move_to</code> and trigger move semantics.

Template Parameters

ResultType the type to which the implementer can be converted to, has to be a subclass of PolymorphicObject

43.29.2 Member Function Documentation

43.29.2.1 convert_to()

Converts the implementer to an object of type result type.

Parameters

result	the object used to store the result of the conversion

Implemented in gko::EnablePolymorphicAssignment < ConcreteType, ResultType >, gko::EnablePolymorphicAssignment < Factoriza gko::EnablePolymorphicAssignment< Lu< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Isai< IsaiType, Value gko::EnablePolymorphicAssignment< SparsityCsr< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Vector< ValueType gko::EnablePolymorphicAssignment< Diagonal< ValueType > >, gko::EnablePolymorphicAssignment< Pgm< ValueType, IndexTy gko::EnablePolymorphicAssignment< Dense< ValueType > >, gko::EnablePolymorphicAssignment< UpperTrs< ValueType, Index gko::EnablePolymorphicAssignment< Amd< IndexType > >, gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndex gko::EnablePolymorphicAssignment< Hybrid< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Identity< ValueType gko::EnablePolymorphicAssignment< ConcreteLinOp >, gko::EnablePolymorphicAssignment< Fft3 >, gko::EnablePolymorphicAssignment gko::EnablePolymorphicAssignment< Composition< ValueType > >, gko::EnablePolymorphicAssignment< RowGatherer< IndexType gko::EnablePolymorphicAssignment< Cholesky< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fft2 >, gko::EnablePolymorphicAssignment< FixedCoarsening< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fbcsr< ${\sf gko::} Enable {\sf PolymorphicAssignment} {\sf < Bicgstab} {\sf < ValueType>>}, {\sf gko::} Enable {\sf PolymorphicAssignment} {\sf < LowerTrs} {\sf < ValueType}, {\sf Index of the property of the prop$ gko::EnablePolymorphicAssignment< Combination< ValueType > >, gko::EnablePolymorphicAssignment< Multigrid >, gko::EnablePolymorphicAssignment< Gmres< ValueType>>, gko::EnablePolymorphicAssignment< CbGmres< ValueType>>, ${\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Csr} < {\sf ValueType}, \\ {\sf IndexType} > >, \\ {\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Ir} < {\sf ValueType} > >, \\ {\sf valueType} > >, \\ {\sf pko::} Enable {\sf PolymorphicAssignment} < {\sf Ir} < {\sf ValueType} > >, \\ {\sf valueType} >$ ${\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Coo} < {\sf ValueType}, \ {\sf IndexType} > >, \\ {\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Fcg} < {\sf ValueType} > >, \\ {\sf valueType} > >, \\ {\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Fcg} < {\sf ValueType} > >, \\ {\sf valueType} > >, \\ {\sf gko::} Enable {\sf PolymorphicAssignment} < {\sf Fcg} < {\sf ValueType} > >, \\ {\sf valueType} > , \\ {\sf valueType} >$ gko::EnablePolymorphicAssignment< Rcm< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fft >, gko::EnablePolymorphicAssignment< Direct< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Idr< ValueType > gko::EnablePolymorphicAssignment< ConcreteFactory >, gko::EnablePolymorphicAssignment< Cgs< ValueType > >, gko::EnablePolymorphicAssignment< Ell< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Ilu< LSolverType, US gko::EnablePolymorphicAssignment< Ic< LSolverType, IndexType > , gko::EnablePolymorphicAssignment< Permutation< Index gko::EnablePolymorphicAssignment< Jacobi< ValueType, IndexType > , gko::EnablePolymorphicAssignment< Gcr< ValueType : gko::EnablePolymorphicAssignment< Schwarz< ValueType, LocalIndexType, GlobalIndexType > >, gko::EnablePolymorphicAssignment gko::EnablePolymorphicAssignment< ScaledReordered< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Sellp< gko::EnablePolymorphicAssignment< Bicg< ValueType > >, gko::EnablePolymorphicAssignment< Perturbation< ValueType > >, and gko::preconditioner::Jacobi < ValueType, IndexType >.

43.29.2.2 move_to()

Converts the implementer to an object of type result_type by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

result the object used to emplace the result of the conversion

Note

ConvertibleTo::move_to can be implemented by simply calling ConvertibleTo::convert_to. However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implemented in gko::EnablePolymorphicAssignment< ConcreteType, ResultType >, gko::EnablePolymorphicAssignment< Factoriza gko::EnablePolymorphicAssignment< Lu< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Isai< IsaiType, Value gko::EnablePolymorphicAssignment< SparsityCsr< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Vector< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Diagonal< ValueType > >, gko::EnablePolymorphicAssignment< UpperTrs< ValueType, IndexTygko::EnablePolymorphicAssignment< UpperTrs< ValueType, IndexTygko::EnablePolymorphicAssignment< Amd< IndexType > >, gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndexType > >, gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndexType > >, gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndexType

gko::EnablePolymorphicAssignment< Hybrid< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Identity< ValueType gko::EnablePolymorphicAssignment< ConcreteLinOp >, gko::EnablePolymorphicAssignment< Fft3 >, gko::EnablePolymorphicAssignment gko::EnablePolymorphicAssignment< Composition< ValueType > >, gko::EnablePolymorphicAssignment< RowGatherer< IndexType gko::EnablePolymorphicAssignment< Cholesky< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fft2 >, gko::EnablePolymorphicAssignment< FixedCoarsening< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fbcsr< gko::EnablePolymorphicAssignment< Bicgstab< ValueType > >, gko::EnablePolymorphicAssignment< LowerTrs< ValueType, Index gko::EnablePolymorphicAssignment< Combination< ValueType > >, gko::EnablePolymorphicAssignment< Multigrid >, gko::EnablePolymorphicAssignment< Gmres< ValueType > >, gko::EnablePolymorphicAssignment< CbGmres< ValueType > >, gko::EnablePolymorphicAssignment< Csr< ValueType, IndexType >>, gko::EnablePolymorphicAssignment< Ir< ValueType >>, gko::EnablePolymorphicAssignment< Coo< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fcg< ValueType > gko::EnablePolymorphicAssignment< Rcm< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Fft >, ${\tt gko::} Enable {\tt PolymorphicAssignment<Direct<Value Type, Index Type>>, {\tt gko::} Enable {\tt PolymorphicAssignment<Idr<Value Type>>, {\tt gko::} Enable {\tt gko$ gko::EnablePolymorphicAssignment< ConcreteFactory >, gko::EnablePolymorphicAssignment< Cgs< ValueType > >, gko::EnablePolymorphicAssignment< Ell< ValueType, IndexType > , gko::EnablePolymorphicAssignment< Ilu< LSolverType, US gko::EnablePolymorphicAssignment< lc< LSolverType, IndexType > >, gko::EnablePolymorphicAssignment< Permutation< Index gko::EnablePolymorphicAssignment< Jacobi< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Gcr< ValueType gko::EnablePolymorphicAssignment< Schwarz< ValueType, LocalIndexType, GlobalIndexType > >, gko::EnablePolymorphicAssign gko::EnablePolymorphicAssignment< ScaledReordered< ValueType, IndexType > >, gko::EnablePolymorphicAssignment< Sellp< gko::EnablePolymorphicAssignment< Bicg< ValueType > >, gko::EnablePolymorphicAssignment< Perturbation< ValueType > >, and gko::preconditioner::Jacobi < ValueType, IndexType >.

The documentation for this class was generated from the following file:

• ginkgo/core/base/polymorphic_object.hpp

43.30 gko::matrix::Coo< ValueType, IndexType > Class Template Reference

COO stores a matrix in the coordinate matrix format.

#include <ginkgo/core/matrix/coo.hpp>

Public Member Functions

· void read (const mat data &data) override

Reads a matrix from a matrix_data structure.

void read (const device_mat_data &data) override

Reads a matrix from a device matrix data structure.

void read (device_mat_data &&data) override

Reads a matrix from a device_matrix_data structure.

· void write (mat data &data) const override

Writes a matrix to a matrix_data structure.

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

- $std::unique_ptr < absolute_type > compute_absolute$ () const override

Gets the AbsoluteLinOp.

· void compute absolute inplace () override

Compute absolute inplace on each element.

value type * get values () noexcept

Returns the values of the matrix.

const value_type * get_const_values () const noexcept

Returns the values of the matrix.

index_type * get_col_idxs () noexcept

Returns the column indexes of the matrix.

const index_type * get_const_col_idxs () const noexcept

Returns the column indexes of the matrix.

• index_type * get_row_idxs () noexcept

Returns the row indexes of the matrix.

- const index_type * get_const_row_idxs () const noexcept
- size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

LinOp * apply2 (ptr_param < const LinOp > b, ptr_param < LinOp > x)

Applies Coo matrix axpy to a vector (or a sequence of vectors).

- const LinOp * apply2 (ptr_param < const LinOp > b, ptr_param < LinOp > x) const
- LinOp * apply2 (ptr_param < const LinOp > alpha, ptr_param < const LinOp > b, ptr_param < LinOp > x)
 Performs the operation x = alpha * Coo * b + x.
- const LinOp * apply2 (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b, ptr_param< LinOp > x) const

Static Public Member Functions

static std::unique_ptr< const Coo > create_const (std::shared_ptr< const Executor > exec, const dim< 2

 &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType
 &&col_idxs, gko::detail::const_array_view< IndexType > &&row_idxs)

Creates a constant (immutable) Coo matrix from a set of constant arrays.

43.30.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Coo< ValueType, IndexType >

COO stores a matrix in the coordinate matrix format.

The nonzero elements are stored in an array row-wise (but not neccessarily sorted by column index within a row). Two extra arrays contain the row and column indexes of each nonzero element of the matrix.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.30.2 Member Function Documentation

43.30.2.1 apply2() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
LinOp* gko::matrix::Coo< ValueType, IndexType >::apply2 (
```

```
ptr_param< const LinOp > alpha,
ptr_param< const LinOp > b,
ptr_param< LinOp > x ) [inline]
```

Performs the operation x = alpha * Coo * b + x.

Parameters

alpha	scaling of the result of Coo * b
b	vector(s) on which the operator is applied
Х	output vector(s)

Returns

this

References gko::ptr_param< T >::get(), gko::PolymorphicObject::get_executor(), and gko::make_temporary_ \leftarrow clone().

43.30.2.2 apply2() [2/4]

References gko::ptr_param< T >::get(), gko::PolymorphicObject::get_executor(), and gko::make_temporary_ \leftarrow clone().

43.30.2.3 apply2() [3/4]

Applies Coo matrix axpy to a vector (or a sequence of vectors).

Performs the operation x = Coo * b + x

Parameters

b	the input vector(s) on which the operator is applied
Х	the output vector(s) where the result is stored

Returns

this

References gko::ptr_param< T >::get(), gko::PolymorphicObject::get_executor(), and gko::make_temporary_ \leftarrow clone().

43.30.2.4 apply2() [4/4]

References gko::ptr_param< T >::get(), gko::PolymorphicObject::get_executor(), and gko::make_temporary_clone().

43.30.2.5 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Coo< ValueType, IndexType >::compute_absolute ()
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove_complex< Coo< ValueType, IndexType >> >.

43.30.2.6 create_const()

Creates a constant (immutable) Coo matrix from a set of constant arrays.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
values	the value array of the matrix
col_idxs	the column index array of the matrix
row_ptrs	the row index array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

43.30.2.7 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Coo< ValueType, IndexType >::extract_\(\chi\)
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag	the vector into which the diagonal will be written
------	--

Implements gko::DiagonalExtractable < ValueType >.

43.30.2.8 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Coo< ValueType, IndexType >::get_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References gko::array< ValueType >::get_data().

43.30.2.9 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Coo< ValueType, IndexType >::get_const_col_idxs ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get const data().

43.30.2.10 get_const_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Coo< ValueType, IndexType >::get_const_row_idxs ( ) const [inline],
[noexcept]
```

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.30.2.11 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Coo< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.30.2.12 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Coo< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.30.2.13 get_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Coo< ValueType, IndexType >::get_row_idxs () [inline], [noexcept]
```

Returns the row indexes of the matrix.

Returns

the row indexes of the matrix.

References gko::array< ValueType >::get_data().

43.30.2.14 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Coo< ValueType, IndexType >::get_values () [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References gko::array< ValueType >::get_data().

43.30.2.15 read() [1/3]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.30.2.16 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.30.2.17 read() [3/3]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.30.2.18 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/coo.hpp

43.31 gko::CpuTimer Class Reference

A timer using std::chrono::steady_clock for timing.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void record (time_point &time) override
 Records a time point at the current time.
- void wait (time_point &time) override

Waits until all kernels in-process when recording the time point are finished.

• std::chrono::nanoseconds difference_async (const time_point &start, const time_point &stop) override Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

43.31.1 Detailed Description

A timer using std::chrono::steady_clock for timing.

43.31.2 Member Function Documentation

43.31.2.1 difference_async()

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. timer->wait(stop) needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements gko::Timer.

The documentation for this class was generated from the following file:

• ginkgo/core/base/timer.hpp

43.32 gko::cpx_real_type< T > Struct Template Reference

Access the underlying real type of a complex number.

#include <ginkgo/core/base/math.hpp>

Public Types

```
• using type = T

The type.
```

43.32.1 Detailed Description

```
template < typename T> struct gko::cpx_real_type < T>
```

Access the underlying real type of a complex number.

Template Parameters

```
T the type being checked.
```

43.32.2 Member Typedef Documentation

43.32.2.1 type

```
template<typename T >
using gko::cpx_real_type< T >::type = T
```

The type.

When the type is not complex, return the type itself.

The documentation for this struct was generated from the following file:

· ginkgo/core/base/math.hpp

43.33 gko::stop::Criterion Class Reference

The Criterion class is a base class for all stopping criteria.

```
#include <ginkgo/core/stop/criterion.hpp>
```

Classes

class Updater

The Updater class serves for convenient argument passing to the Criterion's check function.

Public Member Functions

• Updater update ()

Returns the updater object.

bool check (uint8 stopping_id, bool set_finalized, array< stopping_status > *stop_status, bool *one_← changed, const Updater &updater)

This checks whether convergence was reached for a certain criterion.

43.33.1 Detailed Description

The Criterion class is a base class for all stopping criteria.

It contains a factory to instantiate criteria. It is up to each specific stopping criterion to decide what to do with the data that is passed to it.

Note that depending on the criterion, convergence may not have happened after stopping.

43.33.2 Member Function Documentation

43.33.2.1 check()

```
bool gko::stop::Criterion::check (
          uint8 stopping_id,
          bool set_finalized,
          array< stopping_status > * stop_status,
          bool * one_changed,
          const Updater & updater ) [inline]
```

This checks whether convergence was reached for a certain criterion.

The actual implantation of the criterion goes here.

Parameters

stopping_id	id of the stopping criterion
set_finalized	Controls if the current version should count as finalized or not
stop_status	status of the stopping criterion
one_changed	indicates if the status of a vector has changed
updater	the Updater object containing all the information

Returns

whether convergence was completely reached

```
168
              this->template log<log::Logger::criterion_check_started>(
170
                 this, updater.num_iterations_, updater.residual_,
171
                   updater.residual_norm_, updater.solution_, stopping_id,
             set_finalized);
auto all_converged = this->check_impl(
172
173
              stopping_id, set_finalized, stop_status, one_changed, updater);
this->template log<log::Logger::criterion_check_completed>(
174
175
                   this, updater.num_iterations_, updater.residual_,
177
                   updater.residual_norm_, updater.implicit_sq_residual_norm_,
              updater.solution_, stopping_id, set_finalized, stop_status,
    *one_changed, all_converged);
return all_converged;
178
179
180
```

Referenced by gko::stop::Criterion::Updater::check().

43.33.2.2 update()

```
Updater gko::stop::Criterion::update ( ) [inline]
```

Returns the updater object.

Returns

the updater object

The documentation for this class was generated from the following file:

· ginkgo/core/stop/criterion.hpp

43.34 gko::log::criterion_data Struct Reference

Struct representing Criterion related data.

#include <ginkgo/core/log/record.hpp>

43.34.1 Detailed Description

Struct representing Criterion related data.

The documentation for this struct was generated from the following file:

· ginkgo/core/log/record.hpp

43.35 gko::stop::CriterionArgs Struct Reference

This struct is used to pass parameters to the EnableDefaultCriterionFactoryCriterionFactory::generate() method.

```
#include <ginkgo/core/stop/criterion.hpp>
```

43.35.1 Detailed Description

This struct is used to pass parameters to the EnableDefaultCriterionFactoryCriterionFactory::generate() method. It is the ComponentsType of CriterionFactory.

Note

Dependly on the use case, some of these parameters can be nullptr as only some stopping criterion require them to be set. An example is the ResidualNormReduction which really requires the initial — residual to be set.

The documentation for this struct was generated from the following file:

• ginkgo/core/stop/criterion.hpp

43.36 gko::matrix::Csr< ValueType, IndexType > Class Template Reference

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

```
#include <ginkgo/core/matrix/csr.hpp>
```

Classes

· class classical

classical is a strategy_type which uses the same number of threads on each row.

class cusparse

cusparse is a strategy_type which uses the sparselib csr.

class load_balance

load_balance is a strategy_type which uses the load balance algorithm.

class merge_path

merge_path is a strategy_type which uses the merge_path algorithm.

· class sparselib

sparselib is a strategy_type which uses the sparselib csr.

· class strategy_type

strategy_type is to decide how to set the csr algorithm.

Public Member Functions

void read (const mat_data &data) override

Reads a matrix from a matrix data structure.

void read (const device_mat_data &data) override

Reads a matrix from a device_matrix_data structure.

· void read (device_mat_data &&data) override

Reads a matrix from a device_matrix_data structure.

· void write (mat_data &data) const override

Writes a matrix to a matrix_data structure.

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique ptr< LinOp > conj transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

std::unique_ptr< LinOp > permute (const array< IndexType > *permutation_indices) const override

Returns a LinOp representing the symmetric row and column permutation of the Permutable object.

std::unique_ptr< LinOp > inverse_permute (const array< IndexType > *inverse_permutation_indices) const override

Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.

- std::unique_ptr< LinOp > row_permute (const array< IndexType > *permutation_indices) const override

 Returns a LinOp representing the row permutation of the Permutable object.
- std::unique_ptr< LinOp > column_permute (const array< IndexType > *permutation_indices) const override

 Returns a LinOp representing the column permutation of the Permutable object.
- std::unique_ptr< LinOp > inverse_row_permute (const array< IndexType > *inverse_permutation_indices) const override

Returns a LinOp representing the row permutation of the inverse permuted object.

std::unique_ptr< LinOp > inverse_column_permute (const array< IndexType > *inverse_permutation_
indices) const override

Returns a LinOp representing the row permutation of the inverse permuted object.

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

• std::unique_ptr< absolute_type > compute_absolute () const override

Gets the AbsoluteLinOp.

void compute_absolute_inplace () override

Compute absolute inplace on each element.

void sort_by_column_index ()

Sorts all (value, col_idx) pairs in each row by column index.

value_type * get_values () noexcept

Returns the values of the matrix.

const value type * get const values () const noexcept

Returns the values of the matrix.

• index_type * get_col_idxs () noexcept

Returns the column indexes of the matrix.

const index_type * get_const_col_idxs () const noexcept

Returns the column indexes of the matrix.

• index_type * get_row_ptrs () noexcept

Returns the row pointers of the matrix.

const index_type * get_const_row_ptrs () const noexcept

Returns the row pointers of the matrix.

• index_type * get_srow () noexcept

Returns the starting rows.

const index_type * get_const_srow () const noexcept
 Returns the starting rows.

• size_type get_num_srow_elements () const noexcept

Returns the number of the srow stored elements (involved warps)

size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

std::shared_ptr< strategy_type > get_strategy () const noexcept

Returns the strategy.

void set_strategy (std::shared_ptr< strategy_type > strategy)

Set the strategy.

void scale (ptr_param < const LinOp > alpha)

Scales the matrix with a scalar.

void inv_scale (ptr_param< const LinOp > alpha)

Scales the matrix with the inverse of a scalar.

std::unique_ptr< Csr< ValueType, IndexType > > create_submatrix (const index_set< IndexType > &row
 —
 index_set, const index_set< IndexType > &column_index_set) const

Creates a submatrix from this Csr matrix given row and column index_set objects.

std::unique_ptr< Csr< ValueType, IndexType > > create_submatrix (const span &row_span, const span &column_span) const

Creates a submatrix from this Csr matrix given row and column spans.

• Csr & operator= (const Csr &)

Copy-assigns a Csr matrix.

Csr & operator= (Csr &&)

Move-assigns a Csr matrix.

Csr (const Csr &)

Copy-constructs a Csr matrix.

Csr (Csr &&)

Move-constructs a Csr matrix.

Static Public Member Functions

static std::unique_ptr< const Csr > create_const (std::shared_ptr< const Executor > exec, const dim< 2

 &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType
 &&col_idxs, gko::detail::const_array_view< IndexType > &&row_ptrs, std::shared_ptr< strategy_type > strategy)

Creates a constant (immutable) Csr matrix from a set of constant arrays.

static std::unique_ptr< const Csr > create_const (std::shared_ptr< const Executor > exec, const dim< 2

 &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType
 &&col_idxs, gko::detail::const_array_view< IndexType > &&row_ptrs)

This is version of create_const with a default strategy.

43.36.1 Detailed Description

template < typename ValueType = default_precision, typename IndexType = int32 > class gko::matrix::Csr < ValueType, IndexType >

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

The nonzero elements are stored in a 1D array row-wise, and accompanied with a row pointer array which stores the starting index of each row. An additional column index array is used to identify the column of each nonzero element.

The Csr LinOp supports different operations:

Both the SpGEMM and SpGEAM operation require the input matrices to be sorted by column index, otherwise the algorithms will produce incorrect results.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.36.2 Constructor & Destructor Documentation

43.36.2.1 Csr() [1/2]

Copy-constructs a Csr matrix.

Inherits executor, strategy and data.

43.36.2.2 Csr() [2/2]

Move-constructs a Csr matrix.

Inherits executor and strategy, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor and strategy, no nonzeros and valid row pointers).

43.36.3 Member Function Documentation

43.36.3.1 column_permute()

Returns a LinOp representing the column permutation of the Permutable object.

In the resulting LinOp, the column i contains the input column perm[i].

Parameters

	permutation_indices	the array of indices containing the permutation order ${\tt perm.}$	
--	---------------------	---	--

Returns

a pointer to the new column permuted object

Implements gko::Permutable < IndexType >.

43.36.3.2 compute absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Csr< ValueType, IndexType >::compute_absolute ()
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove complex< Csr< ValueType, IndexType >> >.

43.36.3.3 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::conj_transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.36.3.4 create_const()

Creates a constant (immutable) Csr matrix from a set of constant arrays.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
values	the value array of the matrix
col_idxs	the column index array of the matrix
row_ptrs	the row pointer array of the matrix
strategy	the strategy the matrix uses for SpMV operations

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

Referenced by gko::matrix::Csr< ValueType, IndexType >::create const().

43.36.3.5 create_submatrix() [1/2]

Creates a submatrix from this Csr matrix given row and column index_set objects.

Parameters

row_index_set	the row index set containing the set of rows to be in the submatrix.
column_index_set	the col index set containing the set of columns to be in the submatrix.

Returns

A new CSR matrix with the elements that belong to the row and columns of this matrix as specified by the index sets.

Note

This is not a view but creates a new, separate CSR matrix.

43.36.3.6 create_submatrix() [2/2]

Creates a submatrix from this Csr matrix given row and column spans.

Parameters

row_span	the row span containing the contiguous set of rows to be in the submatrix.
column_span	the column span containing the contiguous set of columns to be in the submatrix.

Returns

A new CSR matrix with the elements that belong to the row and columns of this matrix as specified by the index sets.

Note

This is not a view but creates a new, separate CSR matrix.

43.36.3.7 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Csr< ValueType, IndexType >::extract_
diagonal () const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag	the vector into which the diagonal will be written
------	--

Implements gko::DiagonalExtractable < ValueType >.

43.36.3.8 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References gko::array< ValueType >::get_data().

43.36.3.9 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_col_idxs ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.36.3.10 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs ( ) const [inline],
[noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.36.3.11 get_const_srow()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_srow ( ) const [inline],
[noexcept]
```

Returns the starting rows.

Returns

the starting rows.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.36.3.12 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Csr< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.36.3.13 get_num_srow_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Csr< ValueType, IndexType >::get_num_srow_elements ( ) const [inline],
[noexcept]
```

Returns the number of the srow stored elements (involved warps)

Returns

the number of the srow stored elements (involved warps)

References gko::array< ValueType >::get_num_elems().

43.36.3.14 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Csr< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.36.3.15 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_row_ptrs ( ) [inline], [noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

References gko::array< ValueType >::get_data().

43.36.3.16 get_srow()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_srow () [inline], [noexcept]
```

Returns the starting rows.

Returns

the starting rows.

References gko::array< ValueType >::get data().

43.36.3.17 get_strategy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::get_strategy ( )
const [inline], [noexcept]
```

Returns the strategy.

Returns

the strategy

43.36.3.18 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Csr< ValueType, IndexType >::get_values () [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References gko::array< ValueType >::get_data().

43.36.3.19 inv_scale()

Scales the matrix with the inverse of a scalar.

Parameters

```
alpha The entire matrix is scaled by 1 / alpha. alpha has to be a 1x1 Dense matrix.
```

References gko::PolymorphicObject::get_executor(), and gko::make_temporary_clone().

43.36.3.20 inverse_column_permute()

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column perm[i] contains the input column i.

Parameters

```
permutation_indices the array of indices containing the permutation order perm.
```

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < IndexType >.

43.36.3.21 inverse permute()

Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (perm[i], perm[j]) contains the input value (i, j).

Parameters

```
permutation_indices the array of indices containing the permutation order.
```

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < IndexType >.

43.36.3.22 inverse_row_permute()

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row perm[i] contains the input row i.

Parameters

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < IndexType >.

43.36.3.23 operator=() [1/2]

Copy-assigns a Csr matrix.

Preserves executor, copies everything else.

43.36.3.24 operator=() [2/2]

Move-assigns a Csr matrix.

Preserves executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor and strategy, no nonzeros and valid row pointers).

43.36.3.25 permute()

Returns a LinOp representing the symmetric row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (i, j) contains the input value (perm[i], perm[j]).

Parameters

permutation_indices	the array of indices containing the permutation order.
---------------------	--

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < IndexType >.

43.36.3.26 read() [1/3]

Reads a matrix from a device matrix data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.36.3.27 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.36.3.28 read() [3/3]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device matrix data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.36.3.29 row_permute()

Returns a LinOp representing the row permutation of the Permutable object.

In the resulting LinOp, the row i contains the input row perm[i].

Parameters

permutation_indices the array of indices containing the permutat	rder.
--	-------

Returns

a pointer to the new permuted object

Implements gko::Permutable < IndexType >.

43.36.3.30 scale()

Scales the matrix with a scalar.

Parameters

```
alpha The entire matrix is scaled by alpha. alpha has to be a 1x1 Dense matrix.
```

References gko::PolymorphicObject::get_executor(), and gko::make_temporary_clone().

43.36.3.31 set_strategy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
```

Set the strategy.

Parameters

```
strategy the csr strategy
```

43.36.3.32 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.36.3.33 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- · ginkgo/core/distributed/matrix.hpp
- ginkgo/core/matrix/csr.hpp

43.37 gko::CublasError Class Reference

CublasError is thrown when a cuBLAS routine throws a non-zero error code.

#include <ginkgo/core/base/exception.hpp>

Public Member Functions

CublasError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a cuBLAS error.

43.37.1 Detailed Description

CublasError is thrown when a cuBLAS routine throws a non-zero error code.

43.37.2 Constructor & Destructor Documentation

43.37.2.1 CublasError()

Initializes a cuBLAS error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the cuBLAS routine that failed
error_code	The resulting cuBLAS error code

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.38 gko::cuda_stream Class Reference

An RAII wrapper for a custom CUDA stream.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- cuda_stream (int device_id=0)
 Creates a new custom CUDA stream.
- ~cuda_stream ()

Destroys the custom CUDA stream, if it wasn't moved-from already.

cuda_stream (cuda_stream &&)

Move-constructs from an existing stream, which will be emptied.

cuda_stream & operator= (cuda_stream &&)=delete

Move-assigns from an existing stream, which will be emptied.

• CUstream_st * get () const

Returns the native CUDA stream handle.

43.38.1 Detailed Description

An RAII wrapper for a custom CUDA stream.

The stream will be created on construction and destroyed when the lifetime of the wrapper ends.

43.38.2 Member Function Documentation

43.38.2.1 get()

```
CUstream_st* gko::cuda_stream::get ( ) const
```

Returns the native CUDA stream handle.

In a moved-from cuda_stream, this will return nullptr.

The documentation for this class was generated from the following file:

• ginkgo/core/base/executor.hpp

43.39 gko::CudaError Class Reference

CudaError is thrown when a CUDA routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

CudaError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a CUDA error.

43.39.1 Detailed Description

CudaError is thrown when a CUDA routine throws a non-zero error code.

43.39.2 Constructor & Destructor Documentation

43.39.2.1 CudaError()

Initializes a CUDA error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the CUDA routine that failed
error_code	The resulting CUDA error code

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.40 gko::CudaExecutor Class Reference

This is the Executor subclass which represents the CUDA device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- std::shared_ptr< Executor > get_master () noexcept override
 - ${\it Returns\ the\ master\ OmpExecutor\ of\ this\ Executor.}$
- std::shared_ptr< const Executor > get_master () const noexcept override
 - Returns the master OmpExecutor of this Executor.
- · void synchronize () const override
 - Synchronize the operations launched on the executor with its master.
- int get_device_id () const noexcept
 - Get the CUDA device id of the device associated to this executor.
- int get_num_warps_per_sm () const noexcept
 - Get the number of warps per SM of this executor.
- int get_num_multiprocessor () const noexcept
 - Get the number of multiprocessor of this executor.
- int get_num_warps () const noexcept
 - Get the number of warps of this executor.

int get_warp_size () const noexcept

Get the warp size of this executor.

• int get_major_version () const noexcept

Get the major verion of compute capability.

int get_minor_version () const noexcept

Get the minor verion of compute capability.

cublasContext * get_cublas_handle () const

Get the cubias handle for this executor.

cusparseContext * get_cusparse_handle () const

Get the cusparse handle for this executor.

std::vector< int > get_closest_pus () const

Get the closest PUs.

• int get_closest_numa () const

Get the closest NUMA node.

CUstream_st * get_stream () const

Returns the CUDA stream used by this executor.

Static Public Member Functions

• static std::shared_ptr< CudaExecutor > create (int device_id, std::shared_ptr< Executor > master, bool device_reset=false, allocation_mode alloc_mode=default_cuda_alloc_mode, CUstream_st *stream=nullptr)

Creates a new CudaExecutor.

static int get_num_devices ()

Get the number of devices present on the system.

43.40.1 Detailed Description

This is the Executor subclass which represents the CUDA device.

43.40.2 Member Function Documentation

43.40.2.1 create()

```
static std::shared_ptr<CudaExecutor> gko::CudaExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    bool device_reset = false,
    allocation_mode alloc_mode = default_cuda_alloc_mode,
    CUstream_st * stream = nullptr ) [static]
```

Creates a new CudaExecutor.

Parameters

device_id	the CUDA device id of this device
master	an executor on the host that is used to invoke the device kernels
device_reset	whether to reset the device after the object exits the scope.
Gentland photograph	The allocation mode that the executor should operate on. See @allocation_mode for more details

43.40.2.2 get_closest_numa()

```
int gko::CudaExecutor::get_closest_numa ( ) const [inline]
```

Get the closest NUMA node.

Returns

the closest NUMA node closest to this device

43.40.2.3 get_closest_pus()

```
std::vector<int> gko::CudaExecutor::get_closest_pus ( ) const [inline]
```

Get the closest PUs.

Returns

the array of PUs closest to this device

43.40.2.4 get_cublas_handle()

```
cublasContext* gko::CudaExecutor::get_cublas_handle ( ) const [inline]
```

Get the cubias handle for this executor.

Returns

the cubias handle (cubiasContext*) for this executor

43.40.2.5 get_cusparse_handle()

```
cusparseContext* gko::CudaExecutor::get_cusparse_handle ( ) const [inline]
```

Get the cusparse handle for this executor.

Returns

the cusparse handle (cusparseContext*) for this executor

43.40.2.6 get_master() [1/2]

```
std::shared_ptr<const Executor> gko::CudaExecutor::get_master ( ) const [override], [virtual],
[noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.40.2.7 get_master() [2/2]

```
std::shared_ptr<Executor> gko::CudaExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.40.2.8 get_stream()

```
CUstream_st* gko::CudaExecutor::get_stream ( ) const [inline]
```

Returns the CUDA stream used by this executor.

Can be nullptr for the default stream.

Returns

the stream used to execute kernels and memory operations.

The documentation for this class was generated from the following file:

· ginkgo/core/base/executor.hpp

43.41 gko::CudaTimer Class Reference

A timer using events for timing on a CudaExecutor.

#include <ginkgo/core/base/timer.hpp>

Public Member Functions

· void record (time_point &time) override

Records a time point at the current time.

· void wait (time_point &time) override

Waits until all kernels in-process when recording the time point are finished.

• std::chrono::nanoseconds difference_async (const time_point &start, const time_point &stop) override Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

43.41.1 Detailed Description

A timer using events for timing on a CudaExecutor.

43.41.2 Member Function Documentation

43.41.2.1 difference_async()

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. timer->wait(stop) needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements gko::Timer.

The documentation for this class was generated from the following file:

• ginkgo/core/base/timer.hpp

43.42 gko::CufftError Class Reference

CufftError is thrown when a cuFFT routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

CufftError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a cuFFT error.

43.42.1 Detailed Description

CufftError is thrown when a cuFFT routine throws a non-zero error code.

43.42.2 Constructor & Destructor Documentation

43.42.2.1 CufftError()

Initializes a cuFFT error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the cuFFT routine that failed
error_code	The resulting cuFFT error code

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.43 gko::CurandError Class Reference

CurandError is thrown when a cuRAND routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• CurandError (const std::string &file, int line, const std::string &func, int64 error_code)

Initializes a cuRAND error.

43.43.1 Detailed Description

CurandError is thrown when a cuRAND routine throws a non-zero error code.

43.43.2 Constructor & Destructor Documentation

43.43.2.1 CurandError()

Initializes a cuRAND error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the cuRAND routine that failed
error_code	The resulting cuRAND error code

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.44 gko::matrix::Csr< ValueType, IndexType >::cusparse Class Reference

cusparse is a strategy_type which uses the sparselib csr.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

• cusparse ()

Creates a cusparse strategy.

- void process (const array < index_type > &mtx_row_ptrs, array < index_type > *mtx_srow) override
 Computes srow according to row pointers.
- int64_t clac_size (const int64_t nnz) override

Computes the srow size according to the number of nonzeros.

 std::shared_ptr< strategy_type > copy () override Copy a strategy.

43.44.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Csr< ValueType, IndexType >::cusparse
```

cusparse is a strategy_type which uses the sparselib csr.

Note

cusparse is also known to the hip executor which converts between cuda and hip.

43.44.2 Member Function Documentation

43.44.2.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

```
nnz the number of nonzeros
```

Returns

the size of srow

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.44.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::cusparse::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.44.2.3 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix
mtx_srow	the srow of the matrix

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/csr.hpp

43.45 gko::CusparseError Class Reference

CusparseError is thrown when a cuSPARSE routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

CusparseError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a cuSPARSE error.

43.45.1 Detailed Description

CusparseError is thrown when a cuSPARSE routine throws a non-zero error code.

43.45.2 Constructor & Destructor Documentation

43.45.2.1 CusparseError()

Initializes a cuSPARSE error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the cuSPARSE routine that failed
error_code	The resulting cuSPARSE error code

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.46 gko::default_converter < S, R > Struct Template Reference

Used to convert objects of type S to objects of type R using static_cast.

```
#include <ginkgo/core/base/math.hpp>
```

Public Member Functions

R operator() (S val)
 Converts the object to result type.

43.46.1 Detailed Description

```
template < typename S, typename R> struct gko::default_converter < S, R >
```

Used to convert objects of type ${\mathbb S}$ to objects of type ${\mathbb R}$ using static_cast.

Template Parameters

S	source type
R	result type

Generated by Doxygen

43.46.2 Member Function Documentation

43.46.2.1 operator()()

Converts the object to result type.

Parameters

```
val the object to convert
```

Returns

the converted object

The documentation for this struct was generated from the following file:

· ginkgo/core/base/math.hpp

43.47 gko::matrix::Dense< ValueType > Class Template Reference

Dense is a matrix format which explicitly stores all values of the matrix.

```
#include <ginkgo/core/matrix/dense.hpp>
```

Public Member Functions

std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

void transpose (ptr_param < Dense > output) const

Writes the transposed matrix into the given output matrix.

void conj_transpose (ptr_param< Dense > output) const

Writes the conjugate-transposed matrix into the given output matrix.

void fill (const ValueType value)

Fill the dense matrix with a given value.

- $\bullet \ \, \text{std::unique_ptr} < \text{LinOp} > \text{permute (const array} < \text{int32} > *\text{permutation_indices) const override} \\$
 - Returns a LinOp representing the symmetric row and column permutation of the Permutable object.
- std::unique_ptr< LinOp > permute (const array< int64 > *permutation_indices) const override
 - Returns a LinOp representing the symmetric row and column permutation of the Permutable object.
- void permute (const array< int32 > *permutation_indices, ptr_param< Dense > output) const

Writes the symmetrically permuted matrix into the given output matrix.

- void permute (const array < int64 > *permutation_indices, ptr_param < Dense > output) const
- std::unique_ptr< LinOp > inverse_permute (const array< int32 > *permutation_indices) const override

 Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.
- std::unique_ptr< LinOp > inverse_permute (const array< int64 > *permutation_indices) const override

 Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.
- void inverse_permute (const array< int32 > *permutation_indices, ptr_param< Dense > output) const Writes the inverse symmetrically permuted matrix into the given output matrix.
- void inverse_permute (const array< int64 > *permutation_indices, ptr_param< Dense > output) const
- std::unique_ptr< LinOp > row_permute (const array< int32 > *permutation_indices) const override

 Returns a LinOp representing the row permutation of the Permutable object.
- std::unique_ptr< LinOp > row_permute (const array< int64 > *permutation_indices) const override

 Returns a LinOp representing the row permutation of the Permutable object.
- void row_permute (const array< int32 > *permutation_indices, ptr_param< Dense > output) const Writes the row-permuted matrix into the given output matrix.
- void row_permute (const array< int64 > *permutation_indices, ptr_param< Dense > output) const
- std::unique_ptr< Dense > row_gather (const array< int32 > *gather_indices) const
 Create a Dense matrix consisting of the given rows from this matrix.
- std::unique ptr< Dense > row gather (const array< int64 > *gather indices) const

Create a Dense matrix consisting of the given rows from this matrix.

- void row_gather (const array< int32 > *gather_indices, ptr_param< LinOp > row_collection) const Copies the given rows from this matrix into row_collection
- void row gather (const array< int64 > *gather indices, ptr param< LinOp > row collection) const
- void row_gather (ptr_param< const LinOp > alpha, const array< int32 > *gather_indices, ptr_param< const LinOp > beta, ptr_param< LinOp > row_collection) const

Copies the given rows from this matrix into row_collection with scaling.

- void row_gather (ptr_param < const LinOp > alpha, const array < int64 > *gather_indices, ptr_param < const LinOp > beta, ptr_param < LinOp > row_collection) const
- std::unique_ptr< LinOp > column_permute (const array< int32 > *permutation_indices) const override
 Returns a LinOp representing the column permutation of the Permutable object.
- std::unique_ptr< LinOp > column_permute (const array< int64 > *permutation_indices) const override
 Returns a LinOp representing the column permutation of the Permutable object.
- void column_permute (const array< int32 > *permutation_indices, ptr_param< Dense > output) const Writes the column-permuted matrix into the given output matrix.
- $\bullet \ \ void\ column_permute\ (const\ array < int64 > *permutation_indices,\ ptr_param < Dense > output)\ const$
- std::unique_ptr< LinOp > inverse_row_permute (const array< int32 > *permutation_indices) const override

 Returns a LinOp representing the row permutation of the inverse permuted object.
- std::unique_ptr< LinOp > inverse_row_permute (const array< int64 > *permutation_indices) const override

 Returns a LinOp representing the row permutation of the inverse permuted object.
- void inverse_row_permute (const array< int32 > *permutation_indices, ptr_param< Dense > output) const Writes the inverse row-permuted matrix into the given output matrix.
- $\bullet \ \ void \ inverse_row_permute \ (const \ array < int64 > *permutation_indices, \ ptr_param < Dense > output) \ const$
- std::unique_ptr< LinOp > inverse_column_permute (const array< int32 > *permutation_indices) const override

Returns a LinOp representing the row permutation of the inverse permuted object.

 std::unique_ptr< LinOp > inverse_column_permute (const array< int64 > *permutation_indices) const override

Returns a LinOp representing the row permutation of the inverse permuted object.

void inverse_column_permute (const array< int32 > *permutation_indices, ptr_param< Dense > output)
 const

Writes the inverse column-permuted matrix into the given output matrix.

void inverse_column_permute (const array< int64 > *permutation_indices, ptr_param< Dense > output)

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

void extract_diagonal (ptr_param< Diagonal< ValueType >> output) const

Writes the diagonal of this matrix into an existing diagonal matrix.

• std::unique_ptr< absolute_type > compute_absolute () const override

Gets the AbsoluteLinOp.

void compute_absolute (ptr_param< absolute_type > output) const

Writes the absolute values of this matrix into an existing matrix.

• void compute_absolute_inplace () override

Compute absolute inplace on each element.

std::unique_ptr< complex_type > make_complex () const

Creates a complex copy of the original matrix.

void make complex (ptr param < complex type > result) const

Writes a complex copy of the original matrix to a given complex matrix.

std::unique_ptr< real_type > get_real () const

Creates a new real matrix and extracts the real part of the original matrix into that.

void get real (ptr param< real type > result) const

Extracts the real part of the original matrix into a given real matrix.

std::unique_ptr< real_type > get_imag () const

Creates a new real matrix and extracts the imaginary part of the original matrix into that.

void get_imag (ptr_param < real_type > result) const

Extracts the imaginary part of the original matrix into a given real matrix.

value_type * get_values () noexcept

Returns a pointer to the array of values of the matrix.

const value_type * get_const_values () const noexcept

Returns a pointer to the array of values of the matrix.

size_type get_stride () const noexcept

Returns the stride of the matrix.

· size type get num stored elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

value_type & at (size_type row, size_type col) noexcept

Returns a single element of the matrix.

value_type at (size_type row, size_type col) const noexcept

Returns a single element of the matrix.

ValueType & at (size_type idx) noexcept

Returns a single element of the matrix.

ValueType at (size_type idx) const noexcept

Returns a single element of the matrix.

void scale (ptr param< const LinOp > alpha)

Scales the matrix with a scalar (aka: BLAS scal).

void inv_scale (ptr_param< const LinOp > alpha)

Scales the matrix with the inverse of a scalar.

• void add scaled (ptr param < const LinOp > alpha, ptr param < const LinOp > b)

Adds b scaled by alpha to the matrix (aka: BLAS axpy).

void sub_scaled (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)

Subtracts b scaled by alpha fron the matrix (aka: BLAS axpy).

void compute dot (ptr param< const LinOp > b, ptr param< LinOp > result) const

Computes the column-wise dot product of this matrix and b.

• void compute_dot (ptr_param < const LinOp > b, ptr_param < LinOp > result, array < char > &tmp) const

Computes the column-wise dot product of this matrix and b.

void compute_conj_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result) const

Computes the column-wise dot product of conj(this matrix) and b.

void compute_conj_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp)

Computes the column-wise dot product of conj(this matrix) and b.

void compute_norm2 (ptr_param < LinOp > result) const

Computes the column-wise Euclidian (L^{\wedge} 2) norm of this matrix.

void compute_norm2 (ptr_param < LinOp > result, array < char > &tmp) const

Computes the column-wise Euclidian (L^2) norm of this matrix.

void compute_norm1 (ptr_param< LinOp > result) const

Computes the column-wise (L^{\wedge} 1) norm of this matrix.

void compute_norm1 (ptr_param < LinOp > result, array < char > &tmp) const

Computes the column-wise (L^{\wedge} 1) norm of this matrix.

void compute_squared_norm2 (ptr_param< LinOp > result) const

Computes the square of the column-wise Euclidian (L^{\wedge} 2) norm of this matrix.

void compute_squared_norm2 (ptr_param< LinOp > result, array< char > &tmp) const

Computes the square of the column-wise Euclidian (L^2) norm of this matrix.

std::unique_ptr< Dense > create_submatrix (const span &rows, const span &columns, const size_type stride)

Create a submatrix from the original matrix.

• std::unique ptr< Dense > create submatrix (const span &rows, const span &columns)

Create a submatrix from the original matrix.

std::unique_ptr< real_type > create_real_view ()

Create a real view of the (potentially) complex original matrix.

std::unique_ptr< const real_type > create_real_view () const

Create a real view of the (potentially) complex original matrix.

Dense & operator= (const Dense &)

Copy-assigns a Dense matrix.

• Dense & operator= (Dense &&)

Move-assigns a Dense matrix.

• Dense (const Dense &)

Copy-constructs a Dense matrix.

• Dense (Dense &&)

Move-constructs a Dense matrix.

Static Public Member Functions

• static std::unique ptr< Dense > create with config of (ptr param< const Dense > other)

Creates a Dense matrix with the same size and stride as another Dense matrix.

static std::unique_ptr< Dense > create_with_type_of (ptr_param< const Dense > other, std::shared_ptr< const Executor > exec, const dim< 2 > &size=dim< 2 >{})

Creates a Dense matrix with the same type as another Dense matrix but on a different executor and with a different size.

- static std::unique_ptr< Dense > create_with_type_of (ptr_param< const Dense > other, std::shared_ptr
 const Executor > exec, const dim< 2 > &size, size type stride)
- static std::unique_ptr< Dense > create_with_type_of (ptr_param< const Dense > other, std::shared_ptr< const Executor > exec, const dim< 2 > &size, const dim< 2 > &local_size, size_type stride)
- static std::unique_ptr< Dense > create_view_of (ptr_param< Dense > other)

Creates a Dense matrix, where the underlying array is a view of another Dense matrix' array.

static std::unique_ptr< const Dense > create_const_view_of (ptr_param< const Dense > other)

Creates a immutable Dense matrix, where the underlying array is a view of another Dense matrix' array.

static std::unique_ptr< const Dense > create_const (std::shared_ptr< const Executor > exec, const dim
 2 > &size, gko::detail::const_array_view< ValueType > &&values, size_type stride)

Creates a constant (immutable) Dense matrix from a constant array.

43.47.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::matrix::Dense< ValueType >
```

Dense is a matrix format which explicitly stores all values of the matrix.

The values are stored in row-major format (values belonging to the same row appear consecutive in the memory). Optionally, rows can be padded for better memory access.

Template Parameters

ValueType	precision of matrix elements
-----------	------------------------------

Note

While this format is not very useful for storing sparse matrices, it is often suitable to store vectors, and sets of vectors.

43.47.2 Constructor & Destructor Documentation

43.47.2.1 Dense() [1/2]

Copy-constructs a Dense matrix.

Inherits executor and dimensions, but copies data without padding.

43.47.2.2 Dense() [2/2]

Move-constructs a Dense matrix.

Inherits executor, dimensions and data with padding. The moved-from object is empty (0x0 with empty Array).

43.47.3 Member Function Documentation

43.47.3.1 add_scaled()

Adds b scaled by alpha to the matrix (aka: BLAS axpy).

Parameters

alpha	If alpha is 1x1 Dense matrix, the entire matrix is scaled by alpha. If it is a Dense row vector of values,	
	then i-th column of the matrix is scaled with the i-th element of alpha (the number of columns of alpha	
	has to match the number of columns of the matrix).	
b	a matrix of the same dimension as this	

43.47.3.2 at() [1/4]

Returns a single element of the matrix.

Useful for iterating across all elements of the matrix. However, it is less efficient than the two-parameter variant of this method.

Parameters

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.47.3.3 at() [2/4]

Returns a single element of the matrix.

Useful for iterating across all elements of the matrix. However, it is less efficient than the two-parameter variant of this method.

Parameters

```
idx a linear index of the requested element (ignoring the stride)
```

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.47.3.4 at() [3/4]

Returns a single element of the matrix.

Parameters

row	the row of the requested element
col	the column of the requested element

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.47.3.5 at() [4/4]

Returns a single element of the matrix.

Parameters

row	the row of the requested element
col	the column of the requested element

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

Referenced by gko::initialize().

43.47.3.6 column_permute() [1/4]

Returns a LinOp representing the column permutation of the Permutable object.

In the resulting LinOp, the column i contains the input column perm[i].

Parameters

	permutation_indices	the array of indices containing the permutation order ${\tt perm.}$	
--	---------------------	---	--

Returns

a pointer to the new column permuted object

Implements gko::Permutable < int32 >.

43.47.3.7 column_permute() [2/4]

Writes the column-permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[1]
	elements.
output	The output matrix. It must have the dimensions this->get_size()

See also

Dense::column permute(const array<int32>*)

43.47.3.8 column_permute() [3/4]

Returns a LinOp representing the column permutation of the Permutable object.

In the resulting LinOp, the column i contains the input column perm[i].

Parameters

permutation indices	the array of indices containing the permutation order perm.

Returns

a pointer to the new column permuted object

Implements gko::Permutable < int64 >.

43.47.3.9 column_permute() [4/4]

43.47.3.10 compute_absolute() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<absolute_type> gko::matrix::Dense< ValueType >::compute_absolute ( ) const
[override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove complex< Dense< ValueType >> >.

43.47.3.11 compute_absolute() [2/2]

Writes the absolute values of this matrix into an existing matrix.

Parameters

output	The output matrix. Its size must match the size of this matrix.

See also

Dense::compute_absolute()

43.47.3.12 compute_conj_dot() [1/2]

Computes the column-wise dot product of conj (this matrix) and b.

Parameters

b	a Dense matrix of same dimension as this
result	a Dense row vector, used to store the dot product (the number of column in the vector must match the
	number of columns of this)

43.47.3.13 compute_conj_dot() [2/2]

Computes the column-wise dot product of conj(this matrix) and b.

Parameters

b	a Dense matrix of same dimension as this
result	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

43.47.3.14 compute_dot() [1/2]

Computes the column-wise dot product of this matrix and ${\tt b.}$

Parameters

b	a Dense matrix of same dimension as this
result	a Dense row vector, used to store the dot product (the number of column in the vector must match the
	number of columns of this)

43.47.3.15 compute_dot() [2/2]

Computes the column-wise dot product of this matrix and b.

Parameters

b	a Dense matrix of same dimension as this
result	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

43.47.3.16 compute_norm1() [1/2]

Computes the column-wise ($L^{\wedge}1$) norm of this matrix.

Parameters

```
result a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
```

43.47.3.17 compute_norm1() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_norm1 (
```

```
ptr_param< LinOp > result,
array< char > & tmp ) const
```

Computes the column-wise ($L^{\wedge}1$) norm of this matrix.

Parameters

result	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized
	and/or reset to the correct executor.

43.47.3.18 compute_norm2() [1/2]

Computes the column-wise Euclidian (L^2) norm of this matrix.

Parameters

result

a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)

43.47.3.19 compute_norm2() [2/2]

Computes the column-wise Euclidian ($L^{\wedge}2$) norm of this matrix.

Parameters

result	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized
	and/or reset to the correct executor.

43.47.3.20 compute_squared_norm2() [1/2]

Computes the square of the column-wise Euclidian (L^2) norm of this matrix.

Parameters

result

a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)

43.47.3.21 compute_squared_norm2() [2/2]

Computes the square of the column-wise Euclidian ($L^{\wedge}2$) norm of this matrix.

Parameters

result	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized
	and/or reset to the correct executor.

43.47.3.22 conj_transpose() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.47.3.23 conj_transpose() [2/2]

Writes the conjugate-transposed matrix into the given output matrix.

Parameters

output	The output matrix. It must have the dimensions gko::transpose(this->get_size())	1
--------	---	---

43.47.3.24 create_const()

Creates a constant (immutable) Dense matrix from a constant array.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
values	the value array of the matrix
stride	the row-stride of the matrix

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

43.47.3.25 create_const_view_of()

Creates a immutable Dense matrix, where the underlying array is a view of another Dense matrix' array.

Parameters

other	The other matrix on which to create the view

Returns

A immutable Dense matrix that is a view of other

Referenced by gko::make_const_dense_view().

43.47.3.26 create_real_view() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<real_type> gko::matrix::Dense< ValueType >::create_real_view ( )
```

Create a real view of the (potentially) complex original matrix.

If the original matrix is real, nothing changes. If the original matrix is complex, the result is created by viewing the complex matrix with as real with a reinterpret cast with twice the number of columns and double the stride.

43.47.3.27 create_real_view() [2/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<const real_type> gko::matrix::Dense< ValueType >::create_real_view ( ) const
```

Create a real view of the (potentially) complex original matrix.

If the original matrix is real, nothing changes. If the original matrix is complex, the result is created by viewing the complex matrix with as real with a reinterpret cast with twice the number of columns and double the stride.

43.47.3.28 create_submatrix() [1/2]

Create a submatrix from the original matrix.

Parameters

rows	row span
columns	column span

43.47.3.29 create_submatrix() [2/2]

Create a submatrix from the original matrix.

Warning: defining stride for this create_submatrix method might cause wrong memory access. Better use the create_submatrix(rows, columns) method instead.

Parameters

rows	row span
columns	column span
stride	stride of the new submatrix.

Referenced by gko::matrix::Dense< value_type >::create_submatrix().

43.47.3.30 create_view_of()

Creates a Dense matrix, where the underlying array is a view of another Dense matrix' array.

Parameters

The other matrix on which to crea	ate the view
-----------------------------------	--------------

Returns

A Dense matrix that is a view of other

Referenced by gko::make_dense_view().

43.47.3.31 create_with_config_of()

Creates a Dense matrix with the same size and stride as another Dense matrix.

Parameters

other The other matrix whose configuration needs to copied	d.
--	----

43.47.3.32 create_with_type_of() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_with_type_of (
```

```
ptr_param< const Dense< ValueType >> other,
std::shared_ptr< const Executor > exec,
const dim< 2 > & size,
const dim< 2 > & local_size,
size_type stride ) [inline], [static]
```

Parameters

local_size	Unused
stride	The stride of the new matrix.

Note

This is an overload to stay consistent with gko::experimental::distributed::Vector

43.47.3.33 create_with_type_of() [2/3]

Parameters

s	tride	The stride of the new matrix.
---	-------	-------------------------------

Note

This is an overload which allows full parameter specification.

43.47.3.34 create_with_type_of() [3/3]

Creates a Dense matrix with the same type as another Dense matrix but on a different executor and with a different size.

Parameters

other	The other matrix whose type we target.
exec	The executor of the new matrix.
size	The size of the new matrix.
Generated by Doxygen. Stride The Stride of the new matrix.	

Returns

a Dense matrix with the type of other.

43.47.3.35 extract_diagonal() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Dense< ValueType >::extract_diagonal ( )
const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag the vector into which the diagonal will be written

Implements gko::DiagonalExtractable < ValueType >.

43.47.3.36 extract diagonal() [2/2]

Writes the diagonal of this matrix into an existing diagonal matrix.

Parameters

output The output matrix. Its size must match the size of this matrix's diagonal.

See also

Dense::extract_diagonal()

43.47.3.37 fill()

Fill the dense matrix with a given value.

Parameters

value	the value to be filled

43.47.3.38 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::matrix::Dense< ValueType >::get_const_values ( ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.47.3.39 get_num_stored_elements()

```
template<typename ValueType = default_precision>
size_type gko::matrix::Dense< ValueType >::get_num_stored_elements ( ) const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

43.47.3.40 get_stride()

```
template<typename ValueType = default_precision>
size_type gko::matrix::Dense< ValueType >::get_stride ( ) const [inline], [noexcept]
```

Returns the stride of the matrix.

Returns

the stride of the matrix.

Referenced by gko::matrix::Dense< value_type >::create_submatrix().

43.47.3.41 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::matrix::Dense< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

43.47.3.42 inv_scale()

Scales the matrix with the inverse of a scalar.

Parameters

alpha

If alpha is 1x1 Dense matrix, the entire matrix is scaled by 1 / alpha. If it is a Dense row vector of values, then i-th column of the matrix is scaled with the inverse of the i-th element of alpha (the number of columns of alpha has to match the number of columns of the matrix).

43.47.3.43 inverse_column_permute() [1/4]

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column perm[i] contains the input column i.

Parameters

permutation_indices	the array of indices containing the permutation order perm.
---------------------	---

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < int32 >.

43.47.3.44 inverse_column_permute() [2/4]

Writes the inverse column-permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[1]	
	elements.	
output	The output matrix. It must have the dimensions this->get_size()	

See also

Dense::inverse_column_permute(const array<int32>*)

43.47.3.45 inverse_column_permute() [3/4]

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column perm[i] contains the input column i.

Parameters

permutation_indices	the array of indices containing the permutation order perm.
---------------------	---

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < int64 >.

43.47.3.46 inverse_column_permute() [4/4]

43.47.3.47 inverse_permute() [1/4]

Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (perm[i], perm[j]) contains the input value (i, j).

Parameters

```
permutation_indices the array of indices containing the permutation order.
```

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < int32 >.

43.47.3.48 inverse permute() [2/4]

Writes the inverse symmetrically permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[0] elements.
output	The output matrix. It must have the dimensions this->get_size()

See also

Dense::inverse permute(const array<int32>*)

43.47.3.49 inverse_permute() [3/4]

Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (perm[i], perm[j]) contains the input value (i, j).

Parameters

permutation_indices	the array of indices containing the permutation order.
---------------------	--

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < int64 >.

43.47.3.50 inverse_permute() [4/4]

43.47.3.51 inverse_row_permute() [1/4]

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row perm[i] contains the input row i.

Parameters

permutation_indices	the array of indices containing the permutation order perm.
---------------------	---

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < int32 >.

43.47.3.52 inverse_row_permute() [2/4]

Writes the inverse row-permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[0]
	elements.
output	The output matrix. It must have the dimensions this->get_size()

See also

Dense::inverse_row_permute(const array<int32>*)

43.47.3.53 inverse_row_permute() [3/4]

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row perm[i] contains the input row i.

Parameters

permutation_indices	the array of indices containing the permutation order perm.
---------------------	---

Returns

a pointer to the new inverse permuted object

Implements gko::Permutable < int64 >.

43.47.3.54 inverse_row_permute() [4/4]

43.47.3.55 make_complex() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<complex_type> gko::matrix::Dense< ValueType >::make_complex ( ) const
```

Creates a complex copy of the original matrix.

If the original matrix was real, the imaginary part of the result will be zero.

43.47.3.56 make_complex() [2/2]

Writes a complex copy of the original matrix to a given complex matrix.

If the original matrix was real, the imaginary part of the result will be zero.

43.47.3.57 operator=() [1/2]

Copy-assigns a Dense matrix.

Preserves the executor, reallocates the matrix with minimal stride if the dimensions don't match, then copies the data over, ignoring padding.

43.47.3.58 operator=() [2/2]

Move-assigns a Dense matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with empty Array).

43.47.3.59 permute() [1/4]

Returns a LinOp representing the symmetric row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (i, j) contains the input value (perm[i], perm[j]).

Parameters

```
permutation_indices the array of indices containing the permutation order.
```

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < int32 >.

43.47.3.60 permute() [2/4]

Writes the symmetrically permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[0]
	elements.
output	The output matrix. It must have the dimensions this->get_size()

See also

Dense::permute(const array<int32>*)

43.47.3.61 permute() [3/4]

Returns a LinOp representing the symmetric row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (i, j) contains the input value (perm[i], perm[j]).

Parameters

permutation_indices	the array of indices containing the permutation order.

Returns

a pointer to the new permuted object

Reimplemented from gko::Permutable < int64 >.

43.47.3.62 permute() [4/4]

43.47.3.63 row_gather() [1/6]

Create a Dense matrix consisting of the given rows from this matrix.

Parameters

Returns

Dense matrix on the same executor with the same number of columns and gather_indices->get_← num_elems() rows containing the gathered rows from this matrix: output(i, j) = input(gather← _indices(i), j)

43.47.3.64 row_gather() [2/6]

Copies the given rows from this matrix into row_collection

Parameters

gather_indices	pointer to an array containing row indices from this matrix. It may contain duplicates.	
row_collection	<pre>pointer to a LinOp that will store the gathered rows: row_collection(i, j) =</pre>	
	<pre>input(gather_indices(i), j) It must have the same number of columns as this</pre>	
	<pre>matrix and gather_indices->get_num_elems() rows.</pre>	

43.47.3.65 row_gather() [3/6]

```
template<typename ValueType = default_precision>
```

```
\label{lem:std::unique_ptr<Dense} $$gko::matrix::Dense< ValueType >::row_gather ($$ const array< int64 > * $$ gather_indices ) $$ const
```

Create a Dense matrix consisting of the given rows from this matrix.

Parameters

gather_indices	pointer to an array containing row indices from this matrix. It may contain duplicates.
----------------	---

Returns

Dense matrix on the same executor with the same number of columns and gather_indices->get_← num_elems() rows containing the gathered rows from this matrix: output(i,j) = input(gather← _indices(i), j)

43.47.3.66 row_gather() [4/6]

43.47.3.67 row_gather() [5/6]

Copies the given rows from this matrix into row_collection with scaling.

Parameters

alpha	scaling the result of row gathering	
gather_indices	pointer to an array containing row indices from this matrix. It may contain duplicates.	
beta	scaling the input row_collection	
row_collection	pointer to a LinOp that will store the gathered rows: row_collection(i, j) =	
	<pre>input(gather_indices(i), j) It must have the same number of columns as this</pre>	
	<pre>matrix and gather_indices->get_num_elems() rows.</pre>	

43.47.3.68 row_gather() [6/6]

43.47.3.69 row_permute() [1/4]

Returns a LinOp representing the row permutation of the Permutable object.

In the resulting LinOp, the row i contains the input row perm[i].

Parameters

permutation_indices the array of indices containing the permutation order.
--

Returns

a pointer to the new permuted object

Implements gko::Permutable < int32 >.

43.47.3.70 row_permute() [2/4]

Writes the row-permuted matrix into the given output matrix.

Parameters

permutation_indices	The array containing permutation indices. It must have this->get_size()[0] elements.	
output	The output matrix. It must have the dimensions this->get_size()	

See also

Dense::row_permute(const array<int32>*)

43.47.3.71 row_permute() [3/4]

Returns a LinOp representing the row permutation of the Permutable object.

In the resulting LinOp, the row i contains the input row perm[i].

Parameters

permutation_indices the array of indices containing the permutation order.

Returns

a pointer to the new permuted object

Implements gko::Permutable < int64 >.

43.47.3.72 row_permute() [4/4]

43.47.3.73 scale()

Scales the matrix with a scalar (aka: BLAS scal).

Parameters

alpha

If alpha is 1x1 Dense matrix, the entire matrix is scaled by alpha. If it is a Dense row vector of values, then i-th column of the matrix is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of columns of the matrix).

43.47.3.74 sub scaled()

Subtracts b scaled by alpha fron the matrix (aka: BLAS axpy).

Parameters

alpha	If alpha is 1x1 Dense matrix, b is scaled by alpha. If it is a Dense row vector of values, then i-th column of b is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of columns of the matrix).
b	a matrix of the same dimension as this

43.47.3.75 transpose() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.47.3.76 transpose() [2/2]

Writes the transposed matrix into the given output matrix.

Parameters

```
output | The output matrix. It must have the dimensions gko::transpose(this->get_size())
```

The documentation for this class was generated from the following files:

- ginkgo/core/base/dense_cache.hpp
- ginkgo/core/matrix/dense.hpp

43.48 gko::device_matrix_data< ValueType, IndexType > Class Template Reference

This type is a device-side equivalent to matrix_data.

#include <qinkqo/core/base/device_matrix_data.hpp>

Classes

· struct arrays

Stores the internal arrays of a device_matrix_data object.

Public Member Functions

- device_matrix_data (std::shared_ptr< const Executor > exec, dim< 2 > size={}, size_type num_entries=0)
 Initializes a new device_matrix_data object.
- device_matrix_data (std::shared_ptr< const Executor > exec, const device_matrix_data &data)

Initializes a device_matrix_data object by copying an existing object on another executor.

template<typename ValueArray , typename RowlndexArray , typename CollndexArray >
 device_matrix_data (std::shared_ptr< const Executor > exec, dim< 2 > size, RowlndexArray &&row_idxs,
 CollndexArray &&col_idxs, ValueArray &&values)

Initializes a new device_matrix_data object from existing data.

· host type copy to host () const

Copies the device_matrix_data entries to the host to return a regular matrix_data object with the same dimensions and entries.

void sort_row_major ()

Sorts the matrix entries in row-major order This means that they will be sorted by row index first, and then by column index inside each row.

• void remove_zeros ()

Removes all zero entries from the storage.

• void sum_duplicates ()

Sums up all duplicate entries pointing to the same non-zero location.

std::shared_ptr< const Executor > get_executor () const

Returns the executor used to store the device_matrix_data entries.

• dim< 2 > get_size () const

Returns the dimensions of the matrix.

• size_type get_num_elems () const

Returns the number of stored elements of the matrix.

• index type * get row idxs ()

Returns a pointer to the row index array.

const index_type * get_const_row_idxs () const

Returns a pointer to the constant row index array.

index type * get col idxs ()

Returns a pointer to the column index array.

const index_type * get_const_col_idxs () const

Returns a pointer to the constant column index array.

value_type * get_values ()

Returns a pointer to the value array.

const value_type * get_const_values () const

Returns a pointer to the constant value array.

void resize_and_reset (size_type new_num_entries)

Resizes the internal storage to the given number of stored matrix entries.

void resize_and_reset (dim< 2 > new_size, size_type new_num_entries)

Resizes the matrix and internal storage to the given dimensions.

arrays empty_out ()

Moves out the internal arrays of the device_matrix_data object and resets it to an empty 0x0 matrix.

Static Public Member Functions

 static device_matrix_data create_from_host (std::shared_ptr< const Executor > exec, const host_type &data)

Creates a device_matrix_data object from the given host data on the given executor.

43.48.1 Detailed Description

```
template<typename ValueType, typename IndexType> class gko::device_matrix_data< ValueType, IndexType >
```

This type is a device-side equivalent to matrix_data.

It stores the data necessary to initialize any matrix format in Ginkgo in individual value, column and row index arrays together with associated matrix dimensions. matrix_data uses array-of-Structs storage (AoS), while device matrix data uses Struct-of-Arrays (SoA).

Note

To be used with a Ginkgo matrix type, the entry array must be sorted in row-major order, i.e. by row index, then by column index within rows. This can be achieved by calling the sort_row_major function.

The data must not contain any duplicate (row, column) pairs.

Template Parameters

ValueType	the type used to store matrix values
IndexType	the type used to store matrix row and column indices

43.48.2 Constructor & Destructor Documentation

43.48.2.1 device_matrix_data() [1/3]

```
template<trypename ValueType, typename IndexType>
gko::device_matrix_data< ValueType, IndexType >::device_matrix_data (
```

```
std::shared_ptr< const Executor > exec,
dim< 2 > size = {},
size_type num_entries = 0 ) [explicit]
```

Initializes a new device_matrix_data object.

It uses the given executor to allocate storage for the given number of entries and matrix dimensions.

Parameters

exec	the executor to be used to store the matrix entries
size	the matrix dimensions
num_entries	the number of entries to be stored

43.48.2.2 device_matrix_data() [2/3]

Initializes a device_matrix_data object by copying an existing object on another executor.

Parameters

exec	the executor to be used to store the matrix entries	
data	the device_matrix data object to copy, potentially stored on another executor.	

43.48.2.3 device_matrix_data() [3/3]

Initializes a new device_matrix_data object from existing data.

Parameters

size	the matrix dimensions	
values	the array containing the matrix values	
col_idxs	ol_idxs the array containing the matrix column indice	
row_idxs	the array containing the matrix row indices	

43.48.3 Member Function Documentation

43.48.3.1 copy_to_host()

```
template<typename ValueType, typename IndexType>
host_type gko::device_matrix_data< ValueType, IndexType >::copy_to_host ( ) const
```

Copies the device_matrix_data entries to the host to return a regular matrix_data object with the same dimensions and entries.

Returns

a matrix_data object with the same dimensions and entries.

Referenced by gko::ReadableFromMatrixData< ValueType, int32 >::read().

43.48.3.2 create_from_host()

Creates a device_matrix_data object from the given host data on the given executor.

Parameters

exec	the executor to create the device_matrix_data on.
data	the data to be wrapped or copied into a device_matrix_data.

Returns

a device_matrix_data object with the same size and entries as data copied to the device executor.

43.48.3.3 empty_out()

```
template<typename ValueType, typename IndexType>
arrays gko::device_matrix_data< ValueType, IndexType >::empty_out ( )
```

Moves out the internal arrays of the device_matrix_data object and resets it to an empty 0x0 matrix.

Returns

a struct containing the internal arrays.

43.48.3.4 get_col_idxs()

```
template<typename ValueType, typename IndexType>
index_type* gko::device_matrix_data< ValueType, IndexType >::get_col_idxs ( ) [inline]
```

Returns a pointer to the column index array.

Returns

a pointer to the column index array

References gko::array< ValueType >::get_data().

43.48.3.5 get_const_col_idxs()

```
template<typename ValueType, typename IndexType>
const index_type* gko::device_matrix_data< ValueType, IndexType >::get_const_col_idxs ( )
const [inline]
```

Returns a pointer to the constant column index array.

Returns

a pointer to the constant column index array

References gko::array< ValueType >::get_const_data().

43.48.3.6 get_const_row_idxs()

```
template<typename ValueType, typename IndexType>
const index_type* gko::device_matrix_data< ValueType, IndexType >::get_const_row_idxs ( )
const [inline]
```

Returns a pointer to the constant row index array.

Returns

a pointer to the constant row index array

References gko::array< ValueType >::get_const_data().

43.48.3.7 get_const_values()

```
template<typename ValueType, typename IndexType>
const value_type* gko::device_matrix_data< ValueType, IndexType >::get_const_values ( ) const
[inline]
```

Returns a pointer to the constant value array.

Returns

a pointer to the constant value array

References gko::array< ValueType >::get_const_data().

43.48.3.8 get executor()

```
template<typename ValueType, typename IndexType>
std::shared_ptr<const Executor> gko::device_matrix_data< ValueType, IndexType >::get_executor
( ) const [inline]
```

Returns the executor used to store the device matrix data entries.

Returns

the executor used to store the device_matrix_data entries.

References gko::array< ValueType >::get_executor().

43.48.3.9 get_num_elems()

```
template<typename ValueType, typename IndexType>
size_type gko::device_matrix_data< ValueType, IndexType >::get_num_elems ( ) const [inline]
```

Returns the number of stored elements of the matrix.

Returns

the number of stored elements of the matrix.

References gko::array< ValueType >::get_num_elems().

43.48.3.10 get_row_idxs()

```
template<typename ValueType, typename IndexType>
index_type* gko::device_matrix_data< ValueType, IndexType >::get_row_idxs ( ) [inline]
```

Returns a pointer to the row index array.

Returns

a pointer to the row index array

References gko::array< ValueType >::get_data().

43.48.3.11 get_size()

```
template<typename ValueType, typename IndexType>
dim<2> gko::device_matrix_data< ValueType, IndexType >::get_size ( ) const [inline]
```

Returns the dimensions of the matrix.

Returns

the dimensions of the matrix.

43.48.3.12 get_values()

```
template<typename ValueType, typename IndexType>
value_type* gko::device_matrix_data< ValueType, IndexType >::get_values ( ) [inline]
```

Returns a pointer to the value array.

Returns

a pointer to the value array

References gko::array< ValueType >::get_data().

43.48.3.13 remove_zeros()

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::remove_zeros ( )
```

Removes all zero entries from the storage.

This does not modify the storage if there are no zero entries, and keeps the relative order of nonzero entries otherwise.

43.48.3.14 resize_and_reset() [1/2]

Resizes the matrix and internal storage to the given dimensions.

The resulting storage should be assumed uninitialized.

Parameters

new_size	the new matrix dimensions.
new_num_entries	the new number of stored matrix entries.

43.48.3.15 resize_and_reset() [2/2]

Resizes the internal storage to the given number of stored matrix entries.

The resulting storage should be assumed uninitialized.

Parameters

```
new_num_entries the new number of stored matrix entries.
```

43.48.3.16 sum_duplicates()

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::sum_duplicates ( )
```

Sums up all duplicate entries pointing to the same non-zero location.

The output will be sorted in row-major order, and it will only reallocate if duplicates exist.

The documentation for this class was generated from the following file:

ginkgo/core/base/device_matrix_data.hpp

43.49 gko::matrix::Diagonal < ValueType > Class Template Reference

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

#include <ginkgo/core/matrix/diagonal.hpp>

Public Member Functions

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

std::unique_ptr< absolute_type > compute_absolute () const override

Gets the AbsoluteLinOp.

· void compute absolute inplace () override

Compute absolute inplace on each element.

value_type * get_values () noexcept

Returns a pointer to the array of values of the matrix.

const value_type * get_const_values () const noexcept

Returns a pointer to the array of values of the matrix.

void rapply (ptr_param< const LinOp > b, ptr_param< LinOp > x) const

Applies the diagonal matrix from the right side to a matrix b, which means scales the columns of b with the according diagonal entries.

void inverse_apply (ptr_param< const LinOp > b, ptr_param< LinOp > x) const

Applies the inverse of the diagonal matrix to a matrix b, which means scales the columns of b with the inverse of the according diagonal entries.

Static Public Member Functions

static std::unique_ptr< const Diagonal > create_const (std::shared_ptr< const Executor > exec, size_type size, gko::detail::const_array_view< ValueType > &&values)

Creates a constant (immutable) Diagonal matrix from a constant array.

43.49.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::matrix::Diagonal< ValueType >
```

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

Objects of the Diagonal class always represent a square matrix, and require one array to store their values.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes of a CSR matrix the diagonal is applied or converted to.

43.49.2 Member Function Documentation

43.49.2.1 compute_absolute()

```
template<typename ValueType = default_precision>
std::unique_ptr<absolute_type> gko::matrix::Diagonal< ValueType >::compute_absolute ( ) const
[override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove_complex< Diagonal< ValueType >> >.

43.49.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Diagonal< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.49.2.3 create_const()

Creates a constant (immutable) Diagonal matrix from a constant array.

Parameters

exec	the executor to create the matrix on
size	the size of the square matrix
values	the value array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

43.49.2.4 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::matrix::Diagonal< ValueType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.49.2.5 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::matrix::Diagonal< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data().

43.49.2.6 inverse_apply()

Applies the inverse of the diagonal matrix to a matrix b, which means scales the columns of b with the inverse of the according diagonal entries.

Parameters

b	the input vector(s) on which the inverse of the diagonal matrix is applied
X	the output vector(s) where the result is stored

References gko::ptr_param< T >::get().

43.49.2.7 rapply()

Applies the diagonal matrix from the right side to a matrix b, which means scales the columns of b with the according diagonal entries.

Parameters

b	the input vector(s) on which the diagonal matrix is applied
Χ	the output vector(s) where the result is stored

References gko::ptr_param< T >::get().

43.49.2.8 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Diagonal< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following files:

- ginkgo/core/base/lin_op.hpp
- ginkgo/core/matrix/diagonal.hpp

43.50 gko::DiagonalExtractable < ValueType > Class Template Reference

The diagonal of a LinOp implementing this interface can be extracted.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > extract_diagonal_linop () const override
 Extracts the diagonal entries of the matrix into a vector.
- virtual std::unique_ptr< matrix::Diagonal< ValueType >> extract_diagonal () const =0
 Extracts the diagonal entries of the matrix into a vector.

43.50.1 Detailed Description

```
template<typename ValueType> class gko::DiagonalExtractable< ValueType >
```

The diagonal of a LinOp implementing this interface can be extracted.

extract_diagonal extracts the elements whose col and row index are the same and stores the result in a min(nrows, ncols) x 1 dense matrix.

43.50.2 Member Function Documentation

43.50.2.1 extract_diagonal()

```
template<typename ValueType >
virtual std::unique_ptr<matrix::Diagonal<ValueType> > gko::DiagonalExtractable< ValueType
>::extract_diagonal ( ) const [pure virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

```
diag the vector into which the diagonal will be written
```

Implemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType, IndexType >, gko::matrix::Hybrid< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Ell< ValueType, IndexType gko::matrix::Coo< ValueType, IndexType >, and gko::matrix::Sellp< ValueType, IndexType >.

43.50.2.2 extract_diagonal_linop()

```
template<typename ValueType >
std::unique_ptr<LinOp> gko::DiagonalExtractable< ValueType >::extract_diagonal_linop ( )
const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Returns

linop the linop of diagonal format

Implements gko::DiagonalLinOpExtractable.

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.51 gko::DiagonalLinOpExtractable Class Reference

The diagonal of a LinOp can be extracted.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

virtual std::unique_ptr< LinOp > extract_diagonal_linop () const =0
 Extracts the diagonal entries of the matrix into a vector.

43.51.1 Detailed Description

The diagonal of a LinOp can be extracted.

It will be implemented by DiagonalExtractable<ValueType>, so the class does not need to implement it. extract — _diagonal_linop returns a linop which extracts the elements whose col and row index are the same and stores the result in a min(nrows, ncols) x 1 dense matrix.

43.51.2 Member Function Documentation

43.51.2.1 extract_diagonal_linop()

virtual std::unique_ptr<LinOp> gko::DiagonalLinOpExtractable::extract_diagonal_linop () const
[pure virtual]

Extracts the diagonal entries of the matrix into a vector.

Returns

linop the linop of diagonal format

Implemented in gko::DiagonalExtractable < ValueType >.

The documentation for this class was generated from the following file:

· ginkgo/core/base/lin op.hpp

43.52 gko::dim< Dimensionality, DimensionType > Struct Template Reference

A type representing the dimensions of a multidimensional object.

#include <ginkgo/core/base/dim.hpp>

Public Member Functions

- constexpr dim (const dimension type &size=dimension type{})
 - Creates a dimension object with all dimensions set to the same value.
- template<typename... Rest>

constexpr dim (const dimension_type &first, const Rest &... rest)

Creates a dimension object with the specified dimensions.

- constexpr const dimension_type & operator[] (const size_type &dimension) const noexcept Returns the requested dimension.
- dimension_type & operator[] (const size_type &dimension) noexcept
- · constexpr operator bool () const

Checks if all dimensions evaluate to true.

Friends

- constexpr friend bool operator== (const dim &x, const dim &y)
 - Checks if two dim objects are equal.
- constexpr friend dim operator* (const dim &x, const dim &y)

Multiplies two dim objects.

std::ostream & operator<< (std::ostream &os, const dim &x)

A stream operator overload for dim.

43.52.1 Detailed Description

template < size_type Dimensionality, typename DimensionType = size_type> struct gko::dim < Dimensionality, DimensionType >

A type representing the dimensions of a multidimensional object.

Template Parameters

Dimensionality	number of dimensions of the object
DimensionType	datatype used to represent each dimension

43.52.2 Constructor & Destructor Documentation

43.52.2.1 dim() [1/2]

Creates a dimension object with all dimensions set to the same value.

Parameters

size	the size of each dimension
------	----------------------------

43.52.2.2 dim() [2/2]

Creates a dimension object with the specified dimensions.

If the number of dimensions given is less than the dimensionality of the object, the remaining dimensions are set to the same value as the last value given.

For example, in the context of matrices $dim<2>\{2, 3\}$ creates the dimensions for a 2-by-3 matrix.

Parameters

first	first dimension
rest	other dimensions

43.52.3 Member Function Documentation

43.52.3.1 operator bool()

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr gko::dim< Dimensionality, DimensionType >::operator bool ( ) const [inline], [explicit],
[constexpr]
```

Checks if all dimensions evaluate to true.

For standard arithmetic types, this is equivalent to all dimensions being different than zero.

Returns

true if and only if all dimensions evaluate to true

Note

This operator is explicit to avoid implicit dim-to-int casts. It will still be used in contextual conversions (if, &&, ||, ||)

43.52.3.2 operator[]() [1/2]

Returns the requested dimension.

For example, if d is a dim<2> object representing matrix dimensions, d [0] returns the number of rows, and d [1] returns the number of columns.

Parameters

dimension the requested dimer

Returns

the dimension-th dimension

43.52.3.3 operator[]() [2/2]

43.52.4 Friends And Related Function Documentation

43.52.4.1 operator*

Multiplies two dim objects.

Parameters

X	first object	
У	second object	

Returns

a dim object representing the size of the tensor product x * y

43.52.4.2 operator <<

A stream operator overload for dim.

Parameters

os	stream object	
Χ	dim object	

Returns

a stream object appended with the dim output

43.52.4.3 operator==

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr friend bool operator== (
```

```
const dim< Dimensionality, DimensionType > & x, const dim< Dimensionality, DimensionType > & y) [friend]
```

Checks if two dim objects are equal.

Parameters

X	first object
У	second object

Returns

true if and only if all dimensions of both objects are equal.

The documentation for this struct was generated from the following file:

· ginkgo/core/base/dim.hpp

43.53 gko::DimensionMismatch Class Reference

DimensionMismatch is thrown if an operation is being applied to LinOps of incompatible size.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• DimensionMismatch (const std::string &file, int line, const std::string &func, const std::string &first_name, size_type first_rows, size_type first_cols, const std::string &second_name, size_type second_rows, size_type second_cols, const std::string &clarification)

Initializes a dimension mismatch error.

43.53.1 Detailed Description

DimensionMismatch is thrown if an operation is being applied to LinOps of incompatible size.

43.53.2 Constructor & Destructor Documentation

43.53.2.1 DimensionMismatch()

Initializes a dimension mismatch error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The function name where the error occurred	
first_name	The name of the first operator	
first_rows	The output dimension of the first operator	
first_cols	The input dimension of the first operator	
second_name	The name of the second operator	
second_rows	The output dimension of the second operator	
second_cols	The input dimension of the second operator	
clarification	An additional message describing the error further	

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.54 gko::experimental::solver::Direct< ValueType, IndexType > Class Template Reference

A direct solver based on a factorization into lower and upper triangular factors (with an optional diagonal scaling).

#include <ginkgo/core/solver/direct.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
- Returns a LinOp representing the conjugate transpose of the Transposable object.
- Direct (const Direct &)

Creates a copy of the solver.

Direct (Direct &&)

Moves from the given solver, leaving it empty.

43.54.1 Detailed Description

 $\label{template} \begin{tabular}{ll} template < typename \ ValueType, typename \ IndexType > \\ class \ gko::experimental::solver::Direct < ValueType, IndexType > \\ \end{tabular}$

A direct solver based on a factorization into lower and upper triangular factors (with an optional diagonal scaling).

The solver is built from the Factorization returned by the provided LinOpFactory.

Template Parameters

ValueType	the type used to store values of the system matrix
IndexType	the type used to store sparsity pattern indices of the system matrix

43.54.2 Member Function Documentation

43.54.2.1 conj_transpose()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::experimental::solver::Direct< ValueType, IndexType >::conj_
transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.54.2.2 transpose()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::experimental::solver::Direct< ValueType, IndexType >::transpose (
) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/direct.hpp

43.55 gko::experimental::distributed::DistributedBase Class Reference

A base class for distributed objects.

#include <ginkgo/core/distributed/base.hpp>

Public Member Functions

- DistributedBase & operator= (const DistributedBase &)
 - Copy assignment that doesn't change the used mpi::communicator.
- DistributedBase & operator= (DistributedBase &&) noexcept
 - Move assignment that doesn't change the used mpi::communicator.
- mpi::communicator get_communicator () const

Access the used mpi::communicator.

43.55.1 Detailed Description

A base class for distributed objects.

This class stores and gives access to the used mpi::communicator object.

Note

The communicator is not changed on assignment.

43.55.2 Member Function Documentation

43.55.2.1 get_communicator()

mpi::communicator gko::experimental::distributed::DistributedBase::get_communicator () const
[inline]

Access the used mpi::communicator.

Returns

used mpi::communicator

```
84 { return comm_; }
```

43.55.2.2 operator=() [1/2]

Copy assignment that doesn't change the used mpi::communicator.

Returns

unmodified *this

43.55.2.3 operator=() [2/2]

Move assignment that doesn't change the used mpi::communicator.

Returns

unmodified *this

The documentation for this class was generated from the following file:

ginkgo/core/distributed/base.hpp

43.56 gko::DpcppExecutor Class Reference

This is the Executor subclass which represents a DPC++ enhanced device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- std::shared_ptr< Executor > get_master () noexcept override
 - Returns the master OmpExecutor of this Executor.
- std::shared_ptr< const Executor > get_master () const noexcept override

Returns the master OmpExecutor of this Executor.

· void synchronize () const override

Synchronize the operations launched on the executor with its master.

• int get_device_id () const noexcept

Get the DPCPP device id of the device associated to this executor.

const std::vector< int > & get_subgroup_sizes () const noexcept

Get the available subgroup sizes for this device.

int get_num_computing_units () const noexcept

Get the number of Computing Units of this executor.

• int get_num_subgroups () const noexcept

Get the number of subgroups of this executor.

- const std::vector< int > & get_max_workitem_sizes () const noexcept

Get the maximum work item sizes.

• int get_max_workgroup_size () const noexcept

Get the maximum workgroup size.

int get_max_subgroup_size () const noexcept

Get the maximum subgroup size.

• std::string get_device_type () const noexcept

Get a string representing the device type.

Static Public Member Functions

• static std::shared_ptr< DpcppExecutor > create (int device_id, std::shared_ptr< Executor > master, std ← ::string device_type="all", dpcpp_queue_property property=dpcpp_queue_property::in_order)

Creates a new DpcppExecutor.

static int get_num_devices (std::string device_type)

Get the number of devices present on the system.

43.56.1 Detailed Description

This is the Executor subclass which represents a DPC++ enhanced device.

43.56.2 Member Function Documentation

43.56.2.1 create()

```
static std::shared_ptr<DpcppExecutor> gko::DpcppExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    std::string device_type = "all",
    dpcpp_queue_property property = dpcpp_queue_property::in_order ) [static]
```

Creates a new DpcppExecutor.

Parameters

device_id	the DPCPP device id of this device
master	an executor on the host that is used to invoke the device kernels
device_type	a string representing the type of device to consider (accelerator, cpu, gpu or all).

43.56.2.2 get_device_id()

```
int gko::DpcppExecutor::get_device_id ( ) const [inline], [noexcept]
```

Get the DPCPP device id of the device associated to this executor.

Returns

the DPCPP device id of the device associated to this executor

43.56.2.3 get_device_type()

```
std::string gko::DpcppExecutor::get_device_type ( ) const [inline], [noexcept]
```

Get a string representing the device type.

Returns

a string representing the device type

43.56.2.4 get_master() [1/2]

```
std::shared_ptr<const Executor> gko::DpcppExecutor::get_master ( ) const [override], [virtual],
[noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.56.2.5 get_master() [2/2]

```
std::shared_ptr<Executor> gko::DpcppExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.56.2.6 get_max_subgroup_size()

```
int gko::DpcppExecutor::get_max_subgroup_size ( ) const [inline], [noexcept]
```

Get the maximum subgroup size.

Returns

the maximum subgroup size

43.56.2.7 get_max_workgroup_size()

```
int gko::DpcppExecutor::get_max_workgroup_size ( ) const [inline], [noexcept]
```

Get the maximum workgroup size.

Returns

the maximum workgroup size

43.56.2.8 get_max_workitem_sizes()

```
const std::vector<int>& gko::DpcppExecutor::get_max_workitem_sizes ( ) const [inline], [noexcept]
```

Get the maximum work item sizes.

Returns

the maximum work item sizes

43.56.2.9 get_num_computing_units()

```
int gko::DpcppExecutor::get_num_computing_units ( ) const [inline], [noexcept]
```

Get the number of Computing Units of this executor.

Returns

the number of Computing Units of this executor

43.56.2.10 get_num_devices()

Get the number of devices present on the system.

Parameters

device type	a string representing the device type
acvice type	a string representing the acvice type

Returns

the number of devices present on the system

43.56.2.11 get_subgroup_sizes()

```
const std::vector<iint>& gko::DpcppExecutor::get_subgroup_sizes ( ) const [inline], [noexcept]
```

Get the available subgroup sizes for this device.

Returns

the available subgroup sizes for this device

The documentation for this class was generated from the following file:

· ginkgo/core/base/executor.hpp

43.57 gko::DpcppTimer Class Reference

A timer using kernels for timing on a DpcppExecutor in profiling mode.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void record (time_point &time) override
 - Records a time point at the current time.
- void wait (time_point &time) override
 - Waits until all kernels in-process when recording the time point are finished.
- std::chrono::nanoseconds difference_async (const time_point &start, const time_point &stop) override Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

43.57.1 Detailed Description

A timer using kernels for timing on a DpcppExecutor in profiling mode.

43.57.2 Member Function Documentation

43.57.2.1 difference_async()

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. timer->wait(stop) needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements gko::Timer.

The documentation for this class was generated from the following file:

ginkgo/core/base/timer.hpp

43.58 gko::matrix::Ell< ValueType, IndexType > Class Template Reference

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

#include <ginkgo/core/matrix/ell.hpp>

Public Member Functions

• void read (const mat_data &data) override

Reads a matrix from a matrix data structure.

void read (const device_mat_data &data) override

Reads a matrix from a device_matrix_data structure.

void read (device_mat_data &&data) override

Reads a matrix from a device_matrix_data structure.

void write (mat_data &data) const override

Writes a matrix to a matrix_data structure.

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

• std::unique_ptr< absolute_type > compute_absolute () const override

Gets the AbsoluteLinOp.

• void compute_absolute_inplace () override

Compute absolute inplace on each element.

value_type * get_values () noexcept

Returns the values of the matrix.

• const value_type * get_const_values () const noexcept

Returns the values of the matrix.

index_type * get_col_idxs () noexcept

Returns the column indexes of the matrix.

const index type * get const col idxs () const noexcept

Returns the column indexes of the matrix.

size_type get_num_stored_elements_per_row () const noexcept

Returns the number of stored elements per row.

• size_type get_stride () const noexcept

Returns the stride of the matrix.

size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

value_type & val_at (size_type row, size_type idx) noexcept

Returns the idx-th non-zero element of the row-th row.

value_type val_at (size_type row, size_type idx) const noexcept

Returns the idx-th non-zero element of the row-th row .

• index_type & col_at (size_type row, size_type idx) noexcept

Returns the idx-th column index of the row-th row.

index_type col_at (size_type row, size_type idx) const noexcept

Returns the idx-th column index of the row-th row.

• Ell & operator= (const Ell &)

Copy-assigns an Ell matrix.

Ell & operator= (Ell &&)

Move-assigns an Ell matrix.

• Ell (const Ell &)

Copy-constructs an Ell matrix.

• Ell (Ell &&)

Move-constructs an Ell matrix.

Static Public Member Functions

static std::unique_ptr< const Ell > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&col_idxs, size_type num_stored_elements_per_row, size_type stride)

Creates a constant (immutable) Ell matrix from a set of constant arrays.

43.58.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::EII< ValueType, IndexType >

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

The number of elements stored in each row is the largest number of nonzero elements in any of the rows (obtainable through get_num_stored_elements_per_row() method). This removes the need of a row pointer like in the CSR format, and allows for SIMD processing of the distinct rows. For efficient processing, the nonzero elements and the corresponding column indices are stored in column-major fashion. The columns are padded to the length by user-defined stride parameter whose default value is the number of rows of the matrix.

This implementation uses the column index value invalid_index<IndexType>() to mark padding entries that are not part of the sparsity pattern.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.58.2 Constructor & Destructor Documentation

43.58.2.1 EII() [1/2]

Copy-constructs an Ell matrix.

Inherits executor and dimensions, but copies data without padding.

43.58.2.2 EII() [2/2]

Move-constructs an Ell matrix.

Inherits executor, dimensions and data with padding. The moved-from object is empty (0x0 with empty Array).

43.58.3 Member Function Documentation

43.58.3.1 col_at() [1/2]

Returns the $\mathtt{id}x\text{-th}$ column index of the $\mathtt{row}\text{-th}$ row .

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

```
250 {
251     return this->get_const_col_idxs()[this->linearize_index(row, idx)];
```

```
252
```

References gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs().

43.58.3.2 col_at() [2/2]

Returns the idx-th column index of the row-th row.

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References gko::matrix::Ell< ValueType, IndexType >::get_col_idxs().

43.58.3.3 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Ell< ValueType, IndexType >::compute_absolute ()
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove_complex< Ell< ValueType, IndexType >> >.

43.58.3.4 create const()

Creates a constant (immutable) Ell matrix from a set of constant arrays.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
values	the value array of the matrix
col_idxs	the column index array of the matrix
num_stored_elements_per_row	the number of stored nonzeros per row
stride	the column-stride of the value and column index array

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

43.58.3.5 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Ell< ValueType, IndexType >::extract_
diagonal () const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag	the vector into which the diagonal will be written
3	1

 $Implements\ gko:: Diagonal Extractable < Value Type >.$

43.58.3.6 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Ell< ValueType, IndexType >::get_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References gko::array< ValueType >::get_data().

Referenced by gko::matrix::Ell< ValueType, IndexType >::col_at().

43.58.3.7 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

Referenced by gko::matrix::Ell< ValueType, IndexType >::col_at().

43.58.3.8 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Ell< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.58.3.9 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.58.3.10 get_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_num_stored_elements_per_row ( ) const
[inline], [noexcept]
```

Returns the number of stored elements per row.

Returns

the number of stored elements per row.

43.58.3.11 get stride()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_stride ( ) const [inline], [noexcept]
```

Returns the stride of the matrix.

Returns

the stride of the matrix.

43.58.3.12 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Ell< ValueType, IndexType >::get_values () [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References gko::array< ValueType >::get_data().

43.58.3.13 operator=() [1/2]

Copy-assigns an Ell matrix.

Preserves the executor, reallocates the matrix with minimal stride if the dimensions don't match, then copies the data over, ignoring padding.

43.58.3.14 operator=() [2/2]

Move-assigns an Ell matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with empty Array).

43.58.3.15 read() [1/3]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.58.3.16 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.58.3.17 read() [3/3]

Reads a matrix from a device matrix data structure.

The structure may be emptied by this function.

Parameters

data the device_matrix_data structure.	
--	--

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.58.3.18 val_at() [1/2]

Returns the ${\tt idx}\text{-th}$ non-zero element of the ${\tt row}\text{-th}$ row .

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References gko::array< ValueType >::get_const_data().

43.58.3.19 val_at() [2/2]

Returns the ${\tt idx}\text{-th}$ non-zero element of the ${\tt row}\text{-th}$ row .

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References gko::array< ValueType >::get_data().

43.58.3.20 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- · ginkgo/core/matrix/csr.hpp
- · ginkgo/core/matrix/ell.hpp

43.59 gko::enable_parameters_type< ConcreteParametersType, Factory > Class Template Reference

The enable_parameters_type mixin is used to create a base implementation of the factory parameters structure.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

- template<typename... Args>
 ConcreteParametersType & with_loggers (Args &&... _value)
 - Provides the loggers to be added to the factory and its generated objects in a fluent interface.
- std::unique_ptr< Factory > on (std::shared_ptr< const Executor > exec) const

Creates a new factory on the specified executor.

43.59.1 Detailed Description

```
template < typename \ Concrete Parameters Type, \ typename \ Factory > \\ class \ gko::enable\_parameters\_type < Concrete Parameters Type, \ Factory > \\
```

The enable_parameters_type mixin is used to create a base implementation of the factory parameters structure.

It provides only the on() method which can be used to instantiate the factory give the parameters stored in the structure.

Template Parameters

ConcreteParametersType	the concrete parameters type which is being implemented [CRTP parameter]
Factory	the concrete factory for which these parameters are being used

43.59.2 Member Function Documentation

43.59.2.1 on()

Creates a new factory on the specified executor.

Parameters

the executor where the factory will be created	ated
--	------

Returns

a new factory instance

The documentation for this class was generated from the following file:

ginkgo/core/base/abstract_factory.hpp

43.60 gko::EnableAbsoluteComputation< AbsoluteLinOp > Class Template Reference

The EnableAbsoluteComputation mixin provides the default implementations of compute_absolute_linop and the absolute interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > compute_absolute_linop () const override Gets the absolute LinOp.
- virtual std::unique_ptr< absolute_type > compute_absolute () const =0
 Gets the AbsoluteLinOp.

43.60.1 Detailed Description

```
template<typename AbsoluteLinOp> class gko::EnableAbsoluteComputation< AbsoluteLinOp>
```

The EnableAbsoluteComputation mixin provides the default implementations of compute_absolute_linop and the absolute interface.

compute_absolute gets a new AbsoluteLinOp. compute_absolute_inplace applies absolute inplace, so it still keeps the value type of the class.

Template Parameters

AbsoluteLinOp | the absolute LinOp which is being returned [CRTP parameter]

43.60.2 Member Function Documentation

43.60.2.1 compute_absolute()

```
template<typename AbsoluteLinOp>
virtual std::unique_ptr<absolute_type> gko::EnableAbsoluteComputation< AbsoluteLinOp >←
::compute_absolute ( ) const [pure virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< value_type >, gko::matrix::Dense< ValueType, IndexType >, gko::matrix::Hybrid< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::experimental::distributed::Vector< gko::matrix::Ell< ValueType, IndexType >, gko::matrix::Sellp< ValueType, IndexType >, gko::matrix::Sellp< ValueType, IndexType >, and gko::matrix::Diagonal< ValueType >.

Referenced by gko::EnableAbsoluteComputation< remove_complex< Ell< ValueType, IndexType $> > \leftrightarrow$::compute_absolute_linop().

43.60.2.2 compute_absolute_linop()

```
template<typename AbsoluteLinOp>
std::unique_ptr<LinOp> gko::EnableAbsoluteComputation< AbsoluteLinOp >::compute_absolute_\LinOp
linop ( ) const [inline], [override], [virtual]
```

Gets the absolute LinOp.

Returns

a pointer to the new absolute LinOp

```
Implements gko::AbsoluteComputable.
```

```
812 {
813     return this->compute_absolute();
814 }
```

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.61 gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase > Class Template Reference

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new abstract object.

#include <ginkgo/core/base/polymorphic_object.hpp>

43.61.1 Detailed Description

template<typename AbstractObject, typename PolymorphicBase = PolymorphicObject> class gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase >

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new abstract object.

It uses method hiding to update the parameter and return types from PolymorphicObject toAbstractObject` wherever it makes sense. As opposed to EnablePolymorphicObject, it does not implement PolymorphicObject's virtual methods.

Template Parameters

AbstractObject	the abstract class which is being implemented [CRTP parameter]
PolymorphicBase	parent of AbstractObject in the polymorphic hierarchy, has to be a subclass of polymorphic object

See also

EnablePolymorphicObject for creating a concrete subclass of PolymorphicObject.

The documentation for this class was generated from the following file:

• ginkgo/core/base/polymorphic object.hpp

43.62 gko::solver::EnableApplyWithInitialGuess< DerivedType > Class Template Reference

EnableApplyWithInitialGuess providing default operation for ApplyWithInitialGuess with correct validation and log.

#include <ginkgo/core/solver/solver_base.hpp>

43.62.1 Detailed Description

template < typename DerivedType > class gko::solver::EnableApplyWithInitialGuess < DerivedType >

EnableApplyWithInitialGuess providing default operation for ApplyWithInitialGuess with correct validation and log.

It ensures that vectors of apply_with_initial_guess will always have the same executor as the object this mixin is used in, creating a clone on the correct executor if necessary.

Template Parameters

DerivedType	The type that this Mixin is used in. It must provide get_size() and get_executor() functions that
	return correctly initialized values and the logger functionality.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.63 gko::EnableCreateMethod< ConcreteType > Class Template Reference

This mixin implements a static <code>create()</code> method on <code>ConcreteType</code> that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an std::unique_ptr to such an object.

#include <ginkgo/core/base/polymorphic_object.hpp>

43.63.1 Detailed Description

template<typename ConcreteType>
class gko::EnableCreateMethod< ConcreteType>

This mixin implements a static <code>create()</code> method on <code>ConcreteType</code> that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an std::unique_ptr to such an object.

Template Parameters

ConcreteObject the concrete type for which create() is being implemented [CRTP parameter]

The documentation for this class was generated from the following file:

• ginkgo/core/base/polymorphic_object.hpp

43.64 gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase > Class Template Reference

This mixin provides a default implementation of a concrete factory.

#include <ginkgo/core/base/abstract_factory.hpp>

Public Member Functions

const parameters_type & get_parameters () const noexcept
 Returns the parameters of the factory.

Static Public Member Functions

• static parameters_type create ()

Creates a new ParametersType object which can be used to instantiate a new ConcreteFactory.

43.64.1 Detailed Description

template<typename ConcreteFactory, typename ProductType, typename ParametersType, typename PolymorphicBase> class gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >

This mixin provides a default implementation of a concrete factory.

It implements all the methods of AbstractFactory and PolymorphicObject. Its implementation of the generate — impl() method delegates the creation of the product by calling the ProductType::ProductType(const ConcreteFactory *, const components_type &) constructor. The factory also supports parameters by using the ParametersType structure, which is defined by the user.

For a simple example, see IntFactory in core/test/base/abstract_factory.cpp.

Template Parameters

ConcreteFactory	the concrete factory which is being implemented [CRTP parameter]
ProductType	the concrete type of products which this factory produces, has to be a subclass of PolymorphicBase::abstract_product_type
ParametersType	a type representing the parameters of the factory, has to inherit from the enable_parameters_type mixin
PolymorphicBase	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of AbstractFactory

43.64.2 Member Function Documentation

43.64.2.1 create()

```
template<typename ConcreteFactory , typename ProductType , typename ParametersType , typename PolymorphicBase > static parameters_type gko::EnableDefaultFactory< ConcreteFactory, ProductType, Parameters← Type, PolymorphicBase >::create ( ) [inline], [static]
```

Creates a new ParametersType object which can be used to instantiate a new ConcreteFactory.

This method does not construct the factory directly, but returns a new parameters_type object, which can be used to set the parameters of the factory. Once the parameters have been set, the parameters_type::on() method can be used to obtain an instance of the factory with those parameters.

Returns

a default parameters_type object

43.64.2.2 get_parameters()

```
template<typename ConcreteFactory , typename ProductType , typename ParametersType , typename PolymorphicBase >
const parameters_type& gko::EnableDefaultFactory< ConcreteFactory, ProductType, Parameters←
Type, PolymorphicBase >::get_parameters ( ) const [inline], [noexcept]
```

Returns the parameters of the factory.

Returns

the parameters of the factory

The documentation for this class was generated from the following file:

• ginkgo/core/base/abstract_factory.hpp

43.65 gko::experimental::EnableDistributedLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference

This mixin does the same as EnableLinOp, but for concrete types that are derived from distributed::DistributedBase.

```
#include <ginkgo/core/distributed/lin_op.hpp>
```

Additional Inherited Members

43.65.1 Detailed Description

template < typename ConcreteLinOp, typename PolymorphicBase = LinOp > class gko::experimental::EnableDistributedLinOp < ConcreteLinOp, PolymorphicBase >

This mixin does the same as EnableLinOp, but for concrete types that are derived from distributed::DistributedBase.

See also

EnableLinOp.

Template Parameters

ConcreteLinOp	the concrete LinOp which is being implemented that is derived from
	distributed::DistributedBase [CRTP parameter]
PolymorphicBase	parent of ConcreteLinOp in the polymorphic hierarchy, has to be a subclass of LinOp

The documentation for this class was generated from the following file:

ginkgo/core/distributed/lin_op.hpp

43.66 gko::experimental::EnableDistributedPolymorphicObject < ConcreteObject, PolymorphicBase > Class Template Reference

This mixin does the same as EnablePolymorphicObject, but for concrete types that are derived from distributed::DistributedBase.

#include <ginkgo/core/distributed/polymorphic_object.hpp>

43.66.1 Detailed Description

template < typename ConcreteObject, typename PolymorphicBase = PolymorphicObject > class gko::experimental::EnableDistributedPolymorphicObject < ConcreteObject, PolymorphicBase >

This mixin does the same as EnablePolymorphicObject, but for concrete types that are derived from distributed::DistributedBase.

See also

EnablePolymporphicObject.

The following is a minimal example of a distributed PolymorphicObject:

Template Parameters

ConcreteObject	the concrete type which is being implemented that is derived from distributed::DistributedBase [CRTP parameter]
PolymorphicBase	parent of ConcreteObject in the polymorphic hierarchy, has to be a subclass of polymorphic object

The documentation for this class was generated from the following file:

• ginkgo/core/distributed/polymorphic_object.hpp

43.67 gko::solver::EnableIterativeBase< DerivedType > Class Template Reference

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

#include <ginkgo/core/solver/solver_base.hpp>

Public Member Functions

• EnableIterativeBase & operator= (const EnableIterativeBase &other)

Creates a shallow copy of the provided stopping criterion, clones it onto this executor if executors don't match.

• EnableIterativeBase & operator= (EnableIterativeBase &&other)

Moves the provided stopping criterion, clones it onto this executor if executors don't match.

• EnableIterativeBase (const EnableIterativeBase &other)

Creates a shallow copy of the provided stopping criterion.

• EnableIterativeBase (EnableIterativeBase &&other)

Moves the provided stopping criterion.

• void set_stop_criterion_factory (std::shared_ptr< const stop::CriterionFactory > new_stop_factory) override Sets the stopping criterion of the solver.

43.67.1 Detailed Description

```
template<typename DerivedType>
class gko::solver::EnableIterativeBase< DerivedType >
```

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

Template Parameters

```
DerivedType the CRTP type that derives from this
```

43.67.2 Constructor & Destructor Documentation

43.67.2.1 EnableIterativeBase()

Moves the provided stopping criterion.

The moved-from object has a nullptr stopping criterion.

43.67.3 Member Function Documentation

43.67.3.1 operator=()

Moves the provided stopping criterion, clones it onto this executor if executors don't match.

The moved-from object has a nullptr stopping criterion.

43.67.3.2 set_stop_criterion_factory()

Sets the stopping criterion of the solver.

Parameters

other	the new stopping criterion factory
-------	------------------------------------

Reimplemented from gko::solver::lterativeBase.

Referenced by gko::solver::EnableIterativeBase< Bicg< ValueType > >::operator=().

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.68 gko::EnableLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference

The EnableLinOp mixin can be used to provide sensible default implementations of the majority of the LinOp and PolymorphicObject interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Additional Inherited Members

43.68.1 Detailed Description

```
template<typename ConcreteLinOp, typename PolymorphicBase = LinOp> class gko::EnableLinOp< ConcreteLinOp, PolymorphicBase >
```

The EnableLinOp mixin can be used to provide sensible default implementations of the majority of the LinOp and PolymorphicObject interface.

The goal of the mixin is to facilitate the development of new LinOp, by enabling the implementers to focus on the important parts of their operator, while the library takes care of generating the trivial utility functions. The mixin will provide default implementations for the entire PolymorphicObject interface, including a default implementation of copy_from between objects of the new LinOp type. It will also hide the default LinOp::apply() methods with versions that preserve the static type of the object.

Implementers of new LinOps are required to specify only the following aspects:

- Creation of the LinOp: This can be facilitated via either EnableCreateMethod mixin (used mostly for matrix formats), or GKO_ENABLE_LIN_OP_FACTORY macro (used for operators created from other operators, like preconditioners and solvers).
- 2. Application of the LinOp: Implementers have to override the two overloads of the LinOp::apply_impl() virtual methods.

Note

This mixin can't be used with concrete types that derive from experimental::distributed::DistributedBase. In that case use experimental::EnableDistributedLinOp instead.

Template Parameters

ConcreteLinOp	the concrete LinOp which is being implemented [CRTP parameter]
PolymorphicBase	parent of ConcreteLinOp in the polymorphic hierarchy, has to be a subclass of LinOp

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.69 gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase > Class Template Reference

EnableLogging is a mixin which should be inherited by any class which wants to enable logging.

```
#include <ginkgo/core/log/logger.hpp>
```

43.69.1 Detailed Description

template < typename ConcreteLoggable, typename PolymorphicBase = Loggable > class gko::log::EnableLogging < ConcreteLoggable, PolymorphicBase >

EnableLogging is a mixin which should be inherited by any class which wants to enable logging.

All the received events are passed to the loggers this class contains.

Template Parameters

ConcreteLoggable	the object being logged [CRTP parameter]
PolymorphicBase	the polymorphic base of this class. By default it is Loggable. Change it if you want to use
	a new superclass of Loggable as polymorphic base of this class. Generated by Doxygen

The documentation for this class was generated from the following file:

· ginkgo/core/log/logger.hpp

43.70 gko::multigrid::EnableMultigridLevel< ValueType > Class Template Reference

The EnableMultigridLevel gives the default implementation of MultigridLevel with composition and provides set ← _multigrid_level function.

#include <ginkgo/core/multigrid/multigrid_level.hpp>

Public Member Functions

- std::shared_ptr< const LinOp > get_fine_op () const override
 Returns the operator on fine level.
- std::shared_ptr< const LinOp > get_restrict_op () const override
 Returns the restrict operator.
- std::shared_ptr< const LinOp > get_coarse_op () const override
 Returns the operator on coarse level.
- std::shared_ptr< const LinOp > get_prolong_op () const override
 Returns the prolong operator.

43.70.1 Detailed Description

template<typename ValueType>
class gko::multigrid::EnableMultigridLevel< ValueType >

The EnableMultigridLevel gives the default implementation of MultigridLevel with composition and provides set ← _multigrid_level function.

A class inherit from EnableMultigridLevel should use the this->get_compositions()->apply(...) as its own apply, which represents op(b) = prolong(coarse(restrict(b))).

43.70.2 Member Function Documentation

43.70.2.1 get_coarse_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_coarse_op
( ) const [inline], [override], [virtual]
```

Returns the operator on coarse level.

Returns

the operator on coarse level.

Implements gko::multigrid::MultigridLevel.

References gko::UseComposition < ValueType >::get_operator_at().

43.70.2.2 get_fine_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_fine_op (
) const [inline], [override], [virtual]
```

Returns the operator on fine level.

Returns

the operator on fine level.

Implements gko::multigrid::MultigridLevel.

43.70.2.3 get_prolong_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_prolong←
_op ( ) const [inline], [override], [virtual]
```

Returns the prolong operator.

Returns

the prolong operator.

Implements gko::multigrid::MultigridLevel.

 $References\ gko:: Use Composition < Value Type >:: get_operator_at().$

43.70.2.4 get_restrict_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_restrict←
_op ( ) const [inline], [override], [virtual]
```

Returns the restrict operator.

Returns

the restrict operator.

Implements gko::multigrid::MultigridLevel.

References gko::UseComposition < ValueType >::get_operator_at().

The documentation for this class was generated from the following file:

• ginkgo/core/multigrid/multigrid level.hpp

43.71 gko::EnablePolymorphicAssignment< ConcreteType, ResultType > Class Template Reference

This mixin is used to enable a default PolymorphicObject::copy_from() implementation for objects that have implemented conversions between them.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

Public Member Functions

- void convert_to (result_type *result) const override
 Converts the implementer to an object of type result_type.
- void move_to (result_type *result) override

Converts the implementer to an object of type result_type by moving data from this object.

43.71.1 Detailed Description

```
\label{template} \textbf{typename ConcreteType}, \textbf{typename ResultType = ConcreteType} \\ \textbf{class gko::EnablePolymorphicAssignment} \\ \textbf{< ConcreteType}, \textbf{ResultType} \\ \textbf{>} \\ \\ \textbf{< ConcreteType}, \textbf{< ResultType} \\ \textbf{>} \\ \textbf{< ConcreteType}, \textbf{< C
```

This mixin is used to enable a default PolymorphicObject::copy_from() implementation for objects that have implemented conversions between them.

The requirement is that there is either a conversion constructor from ConcreteType in ResultType, or a conversion operator to ResultType in ConcreteType.

Template Parameters

ConcreteType	the concrete type from which the copy_from is being enabled [CRTP parameter]
ResultType	the type to which copy_from is being enabled

43.71.2 Member Function Documentation

43.71.2.1 convert_to()

Converts the implementer to an object of type result_type.

Parameters

	result	the object used to store the result of the conversion
--	--------	---

Implements gko::ConvertibleTo < ResultType >.

43.71.2.2 move to()

Converts the implementer to an object of type result_type by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

result	the object used to emplace the result of the conversion
	,

Note

ConvertibleTo::move_to can be implemented by simply calling ConvertibleTo::convert_to. However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implements gko::ConvertibleTo< ResultType >.

The documentation for this class was generated from the following file:

• ginkgo/core/base/polymorphic_object.hpp

43.72 gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase > Class Template Reference

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new concrete polymorphic object.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

43.72.1 Detailed Description

template<typename ConcreteObject, typename PolymorphicBase = PolymorphicObject> class gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase >

This mixin inherits from (a subclass of) PolymorphicObject and provides a base implementation of a new concrete polymorphic object.

The mixin changes parameter and return types of appropriate public methods of PolymorphicObject in the same way EnableAbstractPolymorphicObject does. In addition, it also provides default implementations of PolymorphicObject's vritual methods by using the *executor default constructor* and the assignment operator of ConcreteObject. Consequently, the following is a minimal example of PolymorphicObject:

In a way, this mixin can be viewed as an extension of default constructor/destructor/assignment operators.

Note

This mixin does not enable copying the polymorphic object to the object of the same type (i.e. it does not implement the ConvertibleTo<ConcreteObject> interface). To enable a default implementation of this interface see the EnablePolymorphicAssignment mixin.

This mixin can't be used with concrete types that derive from experimental::distributed::DistributedBase. In that case use experimental::EnableDistributedPolymorphicObject instead.

Template Parameters

ConcreteObject	the concrete type which is being implemented [CRTP parameter]
PolymorphicBase	parent of ConcreteObject in the polymorphic hierarchy, has to be a subclass of polymorphic object

The documentation for this class was generated from the following file:

· ginkgo/core/base/polymorphic object.hpp

43.73 gko::solver::EnablePreconditionable< DerivedType > Class Template Reference

Mixin providing default operation for Preconditionable with correct value semantics.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

• void set_preconditioner (std::shared_ptr< const LinOp > new_precond) override

Sets the preconditioner operator used by the Preconditionable.

EnablePreconditionable & operator= (const EnablePreconditionable & other)

Creates a shallow copy of the provided preconditioner, clones it onto this executor if executors don't match.

• EnablePreconditionable & operator= (EnablePreconditionable &&other)

Moves the provided preconditioner, clones it onto this executor if executors don't match.

• EnablePreconditionable (const EnablePreconditionable &other)

Creates a shallow copy of the provided preconditioner.

• EnablePreconditionable (EnablePreconditionable &&other)

Moves the provided preconditioner.

43.73.1 Detailed Description

```
template<typename DerivedType> class gko::solver::EnablePreconditionable< DerivedType >
```

Mixin providing default operation for Preconditionable with correct value semantics.

It ensures that the preconditioner stored in this class will always have the same executor as the object this mixin is used in, creating a clone on the correct executor if necessary.

Template Parameters

DerivedType	The type that this Mixin is used in. It must provide get_size() and get_executor() functions that
	return correctly initialized values when the EnablePreconditionable constructor is called, i.e. the
	constructor must be provided by a base class added before EnablePreconditionable, since the
	member initialization order also applying to multiple inheritance.

43.73.2 Constructor & Destructor Documentation

43.73.2.1 EnablePreconditionable()

Moves the provided preconditioner.

The moved-from object has a nullptr preconditioner.

43.73.3 Member Function Documentation

43.73.3.1 operator=()

Moves the provided preconditioner, clones it onto this executor if executors don't match.

The moved-from object has a nullptr preconditioner.

43.73.3.2 set_preconditioner()

Sets the preconditioner operator used by the Preconditionable.

Parameters

new_precond the new preconditioner operator used by the Preconditionable

Reimplemented from gko::Preconditionable.

Referenced by gko::solver::EnablePreconditionable
Sicg
ValueType
> ::operator=().

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.74 gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType > Class Template Reference

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Additional Inherited Members

43.74.1 Detailed Description

```
template<typename ValueType, typename DerivedType> class gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType >
```

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

Template Parameters

ValueType	the value type that iterative solver uses for its vectors	
DerivedType	the CRTP type that derives from this	

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.75 gko::solver::EnableSolverBase< DerivedType, MatrixType > Class Template Reference

A LinOp deriving from this CRTP class stores a system matrix.

#include <ginkgo/core/solver/solver_base.hpp>

Public Member Functions

EnableSolverBase & operator= (const EnableSolverBase & other)

Creates a shallow copy of the provided system matrix, clones it onto this executor if executors don't match.

EnableSolverBase & operator= (EnableSolverBase &&other)

Moves the provided system matrix, clones it onto this executor if executors don't match.

EnableSolverBase (const EnableSolverBase & other)

Creates a shallow copy of the provided system matrix.

EnableSolverBase (EnableSolverBase &&other)

Moves the provided system matrix.

std::vector< int > get_workspace_scalars () const override

Returns the IDs of all scalars (workspace vectors with system dimension-independent size, usually 1 x num rhs).

std::vector< int > get_workspace_vectors () const override

Returns the IDs of all vectors (workspace vectors with system dimension-dependent size, usually system_matrix_size x num_rhs).

43.75.1 Detailed Description

template<typename DerivedType, typename MatrixType = LinOp> class gko::solver::EnableSolverBase< DerivedType, MatrixType >

A LinOp deriving from this CRTP class stores a system matrix.

Template Parameters

DerivedType	the CRTP type that derives from this	
MatrixType the concrete matrix type to be stored as system_mat		

43.75.2 Constructor & Destructor Documentation

43.75.2.1 EnableSolverBase()

Moves the provided system matrix.

The moved-from object has a nullptr system matrix.

43.75.3 Member Function Documentation

43.75.3.1 operator=()

Moves the provided system matrix, clones it onto this executor if executors don't match.

The moved-from object has a nullptr system matrix.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.76 gko::experimental::mpi::environment Class Reference

Class that sets up and finalizes the MPI environment.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

• int get_provided_thread_support () const

Return the provided thread support.

- environment (int &argc, char **&argv, const thread_type thread_t=thread_type::serialized)
 - Call MPI_Init_thread and initialize the MPI environment.
- ∼environment ()

Call MPI_Finalize at the end of the scope of this class.

43.76.1 Detailed Description

Class that sets up and finalizes the MPI environment.

This class is a simple RAII wrapper to MPI_Init and MPI_Finalize.

MPI Init must have been called before calling any MPI functions.

Note

If MPI_Init has already been called, then this class should not be used.

43.76.2 Constructor & Destructor Documentation

43.76.2.1 environment()

Call MPI_Init_thread and initialize the MPI environment.

Parameters

argc	the number of arguments to the main function.	
argv	the arguments provided to the main function.	
thread← the type of threading for initialization. See @thread_t		
t		

43.76.3 Member Function Documentation

43.76.3.1 get provided thread support()

```
int gko::experimental::mpi::environment::get_provided_thread_support ( ) const [inline]
```

Return the provided thread support.

Returns

the provided thread support

The documentation for this class was generated from the following file:

ginkgo/core/base/mpi.hpp

43.77 gko::Error Class Reference

The Error class is used to report exceptional behaviour in library functions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- Error (const std::string &file, int line, const std::string &what)

 Initializes an error.
- virtual const char * what () const noexcept override
 Returns a human-readable string with a more detailed description of the error.

43.77.1 Detailed Description

The Error class is used to report exceptional behaviour in library functions.

Ginkgo uses C++ exception mechanism to this end, and the Error class represents a base class for all types of errors. The exact list of errors which could occur during the execution of a certain library routine is provided in the documentation of that routine, along with a short description of the situation when that error can occur. During runtime, these errors can be detected by using standard C++ try-catch blocks, and a human-readable error description can be obtained by calling the Error::what() method.

As an example, trying to compute a matrix-vector product with arguments of incompatible size will result in a DimensionMismatch error, which is demonstrated in the following program.

```
#include <ginkgo.h>
#include <iostream>
using namespace gko;
int main()
{
    auto omp = create<OmpExecutor>();
    auto A = randn_fill<matrix::Csr<float>(5, 5, 0f, 1f, omp);
    auto x = fill<matrix::Dense<float>(6, 1, 1f, omp);
    try {
        auto y = apply(A, x);
    } catch(Error e) {
        // an error occured, write the message to screen and exit std::cout « e.what() « std::endl;
        return -1;
    }
    return 0;
}
```

43.77.2 Constructor & Destructor Documentation

43.77.2.1 Error()

Initializes an error.

Parameters

file The name of the offending source file		The name of the offending source file	
	line	The source code line number where the error occurred	
	what	The error message	

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.78 gko::Executor Class Reference

The first step in using the Ginkgo library consists of creating an executor.

#include <ginkgo/core/base/executor.hpp>

Public Member Functions

virtual void run (const Operation &op) const =0

Runs the specified Operation using this Executor.

template<typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename ClosureDpcpp > void run (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const

Runs one of the passed in functors, depending on the Executor type.

• template<typename T >

T * alloc (size_type num_elems) const

Allocates memory in this Executor.

void free (void *ptr) const noexcept

Frees memory previously allocated with Executor::alloc().

• template<typename T >

void copy_from (ptr_param< const Executor > src_exec, size_type num_elems, const T *src_ptr, T *dest ← _ _ptr) const

Copies data from another Executor.

 $\bullet \ \ template {<} typename \ T >$

void copy (size_type num_elems, const T *src_ptr, T *dest_ptr) const

Copies data within this Executor.

• template<typename T >

T copy_val_to_host (const T *ptr) const

Retrieves a single element at the given location from executor memory.

virtual std::shared_ptr< Executor > get_master () noexcept=0

Returns the master OmpExecutor of this Executor.

virtual std::shared_ptr< const Executor > get_master () const noexcept=0

Returns the master OmpExecutor of this Executor.

virtual void synchronize () const =0

Synchronize the operations launched on the executor with its master.

- void add_logger (std::shared_ptr< const log::Logger > logger) override
- void remove_logger (const log::Logger *logger) override
- void set_log_propagation_mode (log_propagation_mode mode)

Sets the logger event propagation mode for the executor.

· bool should_propagate_log () const

Returns true iff events occurring at an object created on this executor should be logged at propagating loggers attached to this executor, and there is at least one such propagating logger.

• bool memory_accessible (const std::shared_ptr< const Executor > &other) const

Verifies whether the executors share the same memory.

43.78.1 Detailed Description

The first step in using the Ginkgo library consists of creating an executor.

Executors are used to specify the location for the data of linear algebra objects, and to determine where the operations will be executed. Ginkgo currently supports five different executor types:

- OmpExecutor specifies that the data should be stored and the associated operations executed on an Open
 —
 MP-supporting device (e.g. host CPU);
- CudaExecutor specifies that the data should be stored and the operations executed on the NVIDIA GPU accelerator;
- HipExecutor specifies that the data should be stored and the operations executed on either an NVIDIA or AMD GPU accelerator;
- DpcppExecutor specifies that the data should be stored and the operations executed on an hardware supporting DPC++;
- ReferenceExecutor executes a non-optimized reference implementation, which can be used to debug the library.

The following code snippet demonstrates the simplest possible use of the Ginkgo library:

```
auto omp = gko::create<gko::OmpExecutor>();
auto A = gko::read_from_mtx<gko::matrix::Csr<float>("A.mtx", omp);
```

First, we create a OMP executor, which will be used in the next line to specify where we want the data for the matrix A to be stored. The second line will read a matrix from the matrix market file 'A.mtx', and store the data on the CPU in CSR format (gko::matrix::Csr is a Ginkgo matrix class which stores its data in CSR format). At this point, matrix A is bound to the CPU, and any routines called on it will be performed on the CPU. This approach is usually desired in sparse linear algebra, as the cost of individual operations is several orders of magnitude lower than the cost of copying the matrix to the GPU.

If matrix A is going to be reused multiple times, it could be beneficial to copy it over to the accelerator, and perform the operations there, as demonstrated by the next code snippet:

```
auto cuda = gko::create<gko::CudaExecutor>(0, omp);
auto dA = gko::copy_to<gko::matrix::Csr<float»(A.get(), cuda);</pre>
```

The first line of the snippet creates a new CUDA executor. Since there may be multiple NVIDIA GPUs present on the system, the first parameter instructs the library to use the first device (i.e. the one with device ID zero, as in cudaSetDevice() routine from the CUDA runtime API). In addition, since GPUs are not stand-alone processors, it is required to pass a "master" OmpExecutor which will be used to schedule the requested CUDA kernels on the accelerator.

The second command creates a copy of the matrix A on the GPU. Notice the use of the get() method. As Ginkgo aims to provide automatic memory management of its objects, the result of calling gko::read_from_mtx() is a smart pointer (std::unique_ptr) to the created object. On the other hand, as the library will not hold a reference to A once the copy is completed, the input parameter for gko::copy_to() is a plain pointer. Thus, the get() method is used to convert from a std::unique_ptr to a plain pointer, as expected by gko::copy_to().

As a side note, the gko::copy_to routine is far more powerful than just copying data between different devices. It can also be used to convert data between different formats. For example, if the above code used gko::matrix::Ell as the template parameter, dA would be stored on the GPU, in ELLPACK format.

Finally, if all the processing of the matrix is supposed to be done on the GPU, and a CPU copy of the matrix is not required, we could have read the matrix to the GPU directly:

```
auto omp = gko::create<gko::OmpExecutor>();
auto cuda = gko::create<gko::CudaExecutor>(0, omp);
auto dA = gko::read_from_mtx<gko::matrix::Csr<float>("A.mtx", cuda);
```

Notice that even though reading the matrix directly from a file to the accelerator is not supported, the library is designed to abstract away the intermediate step of reading the matrix to the CPU memory. This is a general design approach taken by the library: in case an operation is not supported by the device, the data will be copied to the CPU, the operation performed there, and finally the results copied back to the device. This approach makes using the library more concise, as explicit copies are not required by the user. Nevertheless, this feature should be taken into account when considering performance implications of using such operations.

43.78.2 Member Function Documentation

43.78.2.1 add_logger()

Note

This specialization keeps track of whether any propagating loggers were attached to the executor.

See also

Logger::needs_propagation()

Implements gko::log::Loggable.

43.78.2.2 alloc()

Allocates memory in this Executor.

Template Parameters

T datatype to allocate

Parameters

Exceptions

AllocationError	if the allocation failed
-----------------	--------------------------

Returns

pointer to allocated memory

43.78.2.3 copy()

Copies data within this Executor.

Template Parameters

```
T datatype to copy
```

Parameters

num_elems	num_elems number of elements of type T to copy	
src_ptr pointer to a block of memory containing the data to be copied		
dest_ptr pointer to an allocated block of memory where the data will be copied		

References copy_from().

43.78.2.4 copy_from()

Copies data from another Executor.

Template Parameters

```
T datatype to copy
```

Parameters

src_exec	Executor from which the memory will be copied	
num_elems	number of elements of type T to copy	
src_ptr	src_ptr pointer to a block of memory containing the data to be copied	
dest_ptr	pointer to an allocated block of memory where the data will be copied to	

References gko::ptr_param< T >::get().

Referenced by copy().

43.78.2.5 copy_val_to_host()

Retrieves a single element at the given location from executor memory.

Template Parameters

```
T datatype to copy
```

Parameters

ptr the pointer to the element to be copied

Returns

the value stored at ptr

References get_master().

43.78.2.6 free()

Frees memory previously allocated with Executor::alloc().

If ptr is a nullptr, the function has no effect.

Parameters

ptr pointer to the allocated memory block

43.78.2.7 get_master() [1/2]

```
virtual std::shared_ptr<const Executor> gko::Executor::get_master ( ) const [pure virtual],
[noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implemented in gko::DpcppExecutor, gko::HipExecutor, gko::CudaExecutor, and gko::OmpExecutor.

43.78.2.8 get_master() [2/2]

```
virtual std::shared_ptr<Executor> gko::Executor::get_master ( ) [pure virtual], [noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implemented in gko::DpcppExecutor, gko::HipExecutor, gko::CudaExecutor, and gko::OmpExecutor.

Referenced by copy_val_to_host().

43.78.2.9 memory_accessible()

Verifies whether the executors share the same memory.

Parameters

```
other the other Executor to compare against
```

Returns

whether the executors this and other share the same memory.

43.78.2.10 remove_logger()

Note

This specialization keeps track of whether any propagating loggers were attached to the executor.

See also

Logger::needs_propagation()

Implements gko::log::Loggable.

43.78.2.11 run() [1/2]

Runs one of the passed in functors, depending on the Executor type.

Template Parameters

ClosureOmp	type of op_omp
ClosureCuda	type of op_cuda
ClosureHip	type of op_hip
ClosureDpcpp	type of op_dpcpp

Parameters

op_omp	functor to run in case of a OmpExecutor or ReferenceExecutor	
op_cuda	functor to run in case of a CudaExecutor	
op_hip	functor to run in case of a HipExecutor	
op_dpcpp	functor to run in case of a DpcppExecutor	

References run().

43.78.2.12 run() [2/2]

Runs the specified Operation using this Executor.

Parameters

```
op the operation to run
```

Implemented in gko::ReferenceExecutor.

Referenced by run().

43.78.2.13 set_log_propagation_mode()

Sets the logger event propagation mode for the executor.

This controls whether events that happen at objects created on this executor will also be logged at propagating loggers attached to the executor.

See also

Logger::needs propagation()

43.78.2.14 should_propagate_log()

```
bool gko::Executor::should_propagate_log ( ) const [inline]
```

Returns true iff events occurring at an object created on this executor should be logged at propagating loggers attached to this executor, and there is at least one such propagating logger.

See also

Logger::needs_propagation()

Executor::set_log_propagation_mode(log_propagation_mode)

References gko::automatic.

The documentation for this class was generated from the following file:

• ginkgo/core/base/executor.hpp

43.79 gko::log::executor_data Struct Reference

Struct representing Executor related data.

```
#include <ginkgo/core/log/record.hpp>
```

43.79.1 Detailed Description

Struct representing Executor related data.

The documentation for this struct was generated from the following file:

• ginkgo/core/log/record.hpp

43.80 gko::executor_deleter< T > Class Template Reference

This is a deleter that uses an executor's free method to deallocate the data.

#include <ginkgo/core/base/executor.hpp>

Public Member Functions

```
    executor_deleter (std::shared_ptr< const Executor > exec)
```

Creates a new deleter.

• void operator() (pointer ptr) const

Deletes the object.

43.80.1 Detailed Description

```
\label{template} \mbox{template} < \mbox{typename T} > \\ \mbox{class gko::executor\_deleter} < \mbox{T} > \\
```

This is a deleter that uses an executor's free method to deallocate the data.

Template Parameters

```
T | the type of object being deleted
```

43.80.2 Constructor & Destructor Documentation

43.80.2.1 executor_deleter()

Creates a new deleter.

Parameters

```
exec the executor used to free the data
```

43.80.3 Member Function Documentation

43.80.3.1 operator()()

Deletes the object.

Parameters

ptr | pointer to the object being deleted

The documentation for this class was generated from the following file:

· ginkgo/core/base/executor.hpp

43.81 gko::experimental::factorization::Factorization < ValueType, IndexType > Class Template Reference

Represents a generic factorization consisting of two triangular factors (upper and lower) and an optional diagonal scaling matrix.

#include <ginkgo/core/factorization/factorization.hpp>

Public Member Functions

std::unique_ptr< Factorization > unpack () const

Transforms the factorization from a compact representation suitable only for triangular solves to a composition representation that can also be used to access individual factors and multiply with the factorization.

storage_type get_storage_type () const

Returns the storage type used by this factorization.

std::shared_ptr< const matrix_type > get_lower_factor () const

Returns the lower triangular factor of the factorization, if available, nullptr otherwise.

std::shared_ptr< const diag_type > get_diagonal () const

Returns the diagonal scaling matrix of the factorization, if available, nullptr otherwise.

std::shared_ptr< const matrix_type > get_upper_factor () const

Returns the upper triangular factor of the factorization, if available, nullptr otherwise.

std::shared_ptr< const matrix_type > get_combined () const

Returns the matrix storing a compact representation of the factorization, if available, nullptr otherwise.

· Factorization (const Factorization &)

Creates a deep copy of the factorization.

Factorization (Factorization &&)

Moves from the given factorization, leaving it empty.

Static Public Member Functions

static std::unique_ptr< Factorization > create_from_composition (std::unique_ptr< composition_type > composition)

Creates a Factorization from an existing composition.

• static std::unique_ptr< Factorization > create_from_symm_composition (std::unique_ptr< composition_type > composition)

Creates a Factorization from an existing symmetric composition.

• static std::unique_ptr< Factorization > create_from_combined_lu (std::unique_ptr< matrix_type > matrix)

Creates a Factorization from an existing combined representation of an LU factorization.

43.81.1 Detailed Description

```
template < typename\ ValueType,\ typename\ IndexType > \\ class\ gko::experimental::factorization::Factorization < ValueType,\ IndexType > \\
```

Represents a generic factorization consisting of two triangular factors (upper and lower) and an optional diagonal scaling matrix.

This class is used to represent a wide range of different factorizations to be passed on to direct solvers and other similar operations. The storage_type represents how the individual factors are stored internally: They may be stored as separate matrices or in a single matrix, and be symmetric or unsymmetric, with the diagonal belonging to both factory, a single factor or being a separate scaling factor (Cholesky vs. LDL^H vs. LU vs. LDU).

Template Parameters

ValueType	the value type used to store the factorization entries	
IndexType	the index type used to represent the sparsity pattern	

43.81.2 Member Function Documentation

43.81.2.1 create_from_combined_lu()

Creates a Factorization from an existing combined representation of an LU factorization.

Parameters

matrix	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular matrix,
	and the last entry to be the transpose of the first entry. If the composition has 3 elements, we expect
	the middle entry to be a diagonal matrix.

Returns

a symmetric Factorization storing the elements from the Composition.

43.81.2.2 create_from_composition()

43.81 gko::experimental::factorization::Factorization<	value Type, Index Type > Class	remplate Reference
Creates a Factorization from an existing composition.		

Parameters

composition	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular
	matrix, and the last entry to be an upper triangular matrix. If the composition has 3 elements, we
	expect the middle entry to be a diagonal matrix.

Returns

a Factorization storing the elements from the Composition.

43.81.2.3 create_from_symm_composition()

Creates a Factorization from an existing symmetric composition.

Parameters

composition	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular
	matrix, and the last entry to be the transpose of the first entry. If the composition has 3
	elements, we expect the middle entry to be a diagonal matrix.

Returns

a symmetric Factorization storing the elements from the Composition.

43.81.2.4 unpack()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<Factorization> gko::experimental::factorization::Factorization
ValueType,
IndexType >::unpack ( ) const
```

Transforms the factorization from a compact representation suitable only for triangular solves to a composition representation that can also be used to access individual factors and multiply with the factorization.

Returns

a new Factorization object containing this factorization represented as storage_type::composition.

The documentation for this class was generated from the following file:

• ginkgo/core/factorization/factorization.hpp

43.82 gko::matrix::Fbcsr< ValueType, IndexType > Class Template Reference

Fixed-block compressed sparse row storage matrix format.

#include <ginkgo/core/matrix/fbcsr.hpp>

Public Member Functions

void convert_to (Csr< ValueType, IndexType > *result) const override

Converts the matrix to CSR format.

void convert_to (SparsityCsr< ValueType, IndexType > *result) const override

Get the block sparsity pattern in CSR-like format.

· void read (const mat data &data) override

Reads a matrix data into Fbcsr format.

void read (const device_mat_data &data) override

Reads a matrix from a device_matrix_data structure.

void read (device mat data &&data) override

Reads a matrix from a device_matrix_data structure.

• void write (mat_data &data) const override

Writes a matrix to a matrix data structure.

std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

std::unique ptr< absolute type > compute absolute () const override

Gets the AbsoluteLinOp.

• void compute_absolute_inplace () override

Compute absolute inplace on each element.

void sort_by_column_index ()

Sorts the values blocks and block-column indices in each row by column index.

bool is_sorted_by_column_index () const

Tests if all row entry pairs (value, col_idx) are sorted by column index.

- value_type * get_values () noexcept
- const value_type * get_const_values () const noexcept
- index_type * get_col_idxs () noexcept
- const index type * get const col idxs () const noexcept
- index type * get row ptrs () noexcept
- const index_type * get_const_row_ptrs () const noexcept
- size_type get_num_stored_elements () const noexcept
- · size type get num stored blocks () const noexcept
- int get_block_size () const noexcept
- index_type get_num_block_rows () const noexcept
- index_type get_num_block_cols () const noexcept
- Fbcsr & operator= (const Fbcsr &)

Copy-assigns an Fbcsr matrix.

• Fbcsr & operator= (Fbcsr &&)

Move-assigns an Fbcsr matrix.

• Fbcsr (const Fbcsr &)

Copy-constructs an Ell matrix.

• Fbcsr (Fbcsr &&)

Move-constructs an Fbcsr matrix.

Static Public Member Functions

static std::unique_ptr< const Fbcsr > create_const (std::shared_ptr< const Executor > exec, const dim
 2 > &size, int blocksize, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&row_ptrs)

Creates a constant (immutable) Fbcsr matrix from a constant array.

43.82.1 Detailed Description

template < typename ValueType = default_precision, typename IndexType = int32 > class gko::matrix::Fbcsr < ValueType, IndexType >

Fixed-block compressed sparse row storage matrix format.

FBCSR is a matrix format meant for matrices having a natural block structure made up of small, dense, disjoint blocks. It is similar to CSR

See also

Csr. However, unlike Csr, each non-zero location stores a small dense block of entries having a constant size. This reduces the number of integers that need to be stored in order to refer to a given non-zero entry, and enables efficient implementation of certain block methods.

The block size is expected to be known in advance and passed to the constructor.

Note

The total number of rows and the number of columns are expected to be divisible by the block size.

The nonzero elements are stored in a 1D array row-wise, and accompanied with a row pointer array which stores the starting index of each block-row. An additional block-column index array is used to identify the block-column of each nonzero block.

The Fbcsr LinOp supports different operations:

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.82.2 Constructor & Destructor Documentation

43.82.2.1 Fbcsr() [1/2]

Copy-constructs an Ell matrix.

Inherits executor and data.

43.82.2.2 Fbcsr() [2/2]

Move-constructs an Fbcsr matrix.

Inherits executor and data. The moved-from object is empty (0x0 with no nonzeros, but valid row pointers).

43.82.3 Member Function Documentation

43.82.3.1 compute absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Fbcsr< ValueType, IndexType >::compute_absolute (
) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove_complex< Fbcsr< ValueType, IndexType >> >.

43.82.3.2 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Fbcsr< ValueType, IndexType >::conj_transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.82.3.3 convert_to() [1/2]

Converts the matrix to CSR format.

Note

Any explicit zeros in the original matrix are retained in the converted result.

43.82.3.4 convert_to() [2/2]

Get the block sparsity pattern in CSR-like format.

Note

The actual non-zero values are never copied; the result always has a value array of size 1 with the value 1.

43.82.3.5 create_const()

Creates a constant (immutable) Fbcsr matrix from a constant array.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
blocksize	the block size of the matrix
values	the value array of the matrix
col_idxs	the block column index array of the matrix
row_ptrs	the block row pointer array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

43.82.3.6 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Fbcsr< ValueType, IndexType >::extract_
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag the vector into which the diagonal will be written

Implements gko::DiagonalExtractable < ValueType >.

43.82.3.7 get_block_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int gko::matrix::Fbcsr< ValueType, IndexType >::get_block_size ( ) const [inline], [noexcept]
```

Returns

The fixed block size for this matrix

43.82.3.8 get col idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_col_idxs () [inline], [noexcept]
```

Returns

The column indexes of the matrix.

References gko::array< ValueType >::get_data().

43.82.3.9 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_col_idxs ( ) const
[inline], [noexcept]
```

Returns

The column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.82.3.10 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_row_ptrs ( ) const
[inline], [noexcept]
```

Returns

The row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.82.3.11 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns

The values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.82.3.12 get_num_block_cols()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_block_cols ( ) const [inline],
[noexcept]
```

Returns

The number of block-columns in the matrix

43.82.3.13 get_num_block_rows()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_block_rows ( ) const [inline],
[noexcept]
```

Returns

The number of block-rows in the matrix

43.82.3.14 get_num_stored_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_blocks ( ) const [inline],
[noexcept]
```

Returns

The number of non-zero blocks explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.82.3.15 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns

The number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.82.3.16 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_row_ptrs () [inline], [noexcept]
```

Returns

The row pointers of the matrix.

References gko::array< ValueType >::get_data().

43.82.3.17 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_values () [inline], [noexcept]
```

Returns

The values of the matrix.

References gko::array< ValueType >::get_data().

43.82.3.18 is_sorted_by_column_index()

```
template<typename ValueType = default_precision, typename IndexType = int32>
bool gko::matrix::Fbcsr< ValueType, IndexType >::is_sorted_by_column_index ( ) const
```

Tests if all row entry pairs (value, col_idx) are sorted by column index.

Returns

True if all row entry pairs (value, col idx) are sorted by column index

43.82.3.19 operator=() [1/2]

Copy-assigns an Fbcsr matrix.

Preserves the executor, copies data and block size from the input.

43.82.3.20 operator=() [2/2]

Move-assigns an Fbcsr matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with no nonzeros, but valid row pointers).

43.82.3.21 read() [1/3]

Reads a matrix from a device matrix data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.82.3.22 read() [2/3]

Reads a matrix_data into Fbcsr format.

Requires the block size to be set beforehand

See also

set_block_size.

Warning

Unlike Csr::read, here explicit non-zeros are NOT dropped.

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.82.3.23 read() [3/3]

Reads a matrix from a device matrix data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.82.3.24 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Fbcsr< ValueType, IndexType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.82.3.25 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/csr.hpp
- · ginkgo/core/matrix/fbcsr.hpp

43.83 gko::solver::Fcg< ValueType > Class Template Reference

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

#include <ginkgo/core/solver/fcg.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
 - Returns a LinOp representing the conjugate transpose of the Transposable object.
- bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

43.83.1 Detailed Description

```
\label{template} \mbox{typename ValueType = default\_precision} > \\ \mbox{class gko::solver::Fcg} < \mbox{ValueType} > \\
```

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

Though this method performs very well for symmetric positive definite matrices, it is in general not suitable for general matrices.

In contrast to the standard CG based on the Polack-Ribiere formula, the flexible CG uses the Fletcher-Reeves formula for creating the orthonormal vectors spanning the Krylov subspace. This increases the computational cost of every Krylov solver iteration but allows for non-constant preconditioners.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of FCG are merged into 2 separate steps.

Template Parameters

ValueType precision of matrix elements

43.83.2 Member Function Documentation

43.83.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Fcg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
99 { return true; }
```

43.83.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Fcg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.83.2.3 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Fcg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/solver/fcg.hpp

43.84 gko::matrix::Fft Class Reference

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

```
#include <ginkgo/core/matrix/fft.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

- void write (matrix_data< std::complex< float >, int32 > &data) const override
- void write (matrix_data< std::complex< float >, int64 > &data) const override
- void write (matrix_data< std::complex< double >, int32 > &data) const override
 - Writes a matrix to a matrix_data structure.

Writes a matrix to a matrix_data structure.

Writes a matrix to a matrix_data structure.

void write (matrix_data< std::complex< double >, int64 > &data) const override

Writes a matrix to a matrix_data structure.

43.84.1 Detailed Description

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

It implements forward and inverse DFT.

For a power-of-two size n with corresponding root of unity $\omega=e^{-2\pi i/n}$ for forward DFT and $\omega=e^{2\pi i/n}$ for inverse DFT it computes

$$x_k = \sum_{j=0}^{n-1} \omega^{jk} b_j$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," Mathematics of Computation, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

43.84.2 Member Function Documentation

43.84.2.1 conj transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.84.2.2 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.84.2.3 write() [1/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< std::complex< double >, int32 >.

43.84.2.4 write() [2/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < double >, int 64 >.$

43.84.2.5 write() [3/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < float >, int 32 >.$

43.84.2.6 write() [4/4]

Writes a matrix to a matrix_data structure.

Parameters

Implements gko::WritableToMatrixData < std::complex < float >, int64 >.

The documentation for this class was generated from the following file:

ginkgo/core/matrix/fft.hpp

43.85 gko::matrix::Fft2 Class Reference

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

#include <ginkgo/core/matrix/fft.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
 Returns a LinOp representing the conjugate transpose of the Transposable object.
- void write (matrix_data < std::complex < float >, int32 > &data) const override
 Writes a matrix to a matrix_data structure.
- void write (matrix_data < std::complex < float >, int64 > &data) const override
 Writes a matrix to a matrix_data structure.
- void write (matrix_data < std::complex < double >, int32 > &data) const override
 Writes a matrix to a matrix_data structure.
- void write (matrix_data < std::complex < double >, int64 > &data) const override
 Writes a matrix to a matrix data structure.

43.85.1 Detailed Description

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

For indexing purposes, the first dimension is the major axis.

It implements complex-to-complex forward and inverse FFT.

For a power-of-two sizes n_1,n_2 with corresponding root of unity $\omega=e^{-2\pi i/(n_1n_2)}$ for forward DFT and $\omega=e^{2\pi i/(n_1n_2)}$ for inverse DFT it computes

$$x_{k_1 n_2 + k_2} = \sum_{i_1 = 0}^{n_1 - 1} \sum_{i_2 = 0}^{n_2 - 1} \omega^{i_1 k_1 + i_2 k_2} b_{i_1 n_2 + i_2}$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," Mathematics of Computation, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

43.85.2 Member Function Documentation

43.85.2.1 conj_transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft2::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.85.2.2 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft2::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.85.2.3 write() [1/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < double >, int 32 >.$

43.85.2.4 write() [2/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < double >, int 64 >.$

43.85.2.5 write() [3/4]

Writes a matrix to a matrix data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< std::complex< float >, int32 >.

43.85.2.6 write() [4/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < float >, int 64 >.$

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/fft.hpp

43.86 gko::matrix::Fft3 Class Reference

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

#include <ginkgo/core/matrix/fft.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

- void write (matrix_data < std::complex < float >, int32 > &data) const override
 Writes a matrix to a matrix data structure.
- void write (matrix_data < std::complex < float >, int64 > &data) const override
 Writes a matrix to a matrix_data structure.
- void write (matrix_data < std::complex < double >, int32 > &data) const override
 Writes a matrix to a matrix_data structure.
- void write (matrix_data < std::complex < double >, int64 > &data) const override
 Writes a matrix to a matrix_data structure.

43.86.1 Detailed Description

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

For indexing purposes, the first dimension is the major axis.

It implements complex-to-complex forward and inverse FFT.

For a power-of-two sizes n_1,n_2,n_3 with corresponding root of unity $\omega=e^{-2\pi i/(n_1n_2n_3)}$ for forward DFT and $\omega=e^{2\pi i/(n_1n_2n_3)}$ for inverse DFT it computes

$$x_{k_1 n_2 n_3 + k_2 n_3 + k_3} = \sum_{i_1 = 0}^{n_1 - 1} \sum_{i_2 = 0}^{n_2 - 1} \sum_{i_3 = 0}^{n_3 - 1} \omega^{i_1 k_1 + i_2 k_2 + i_3 k_3} b_{i_1 n_2 n_3 + i_2 n_3 + i_3}$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," Mathematics of Computation, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

43.86.2 Member Function Documentation

43.86.2.1 conj_transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft3::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.86.2.2 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft3::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.86.2.3 write() [1/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< std::complex< double >, int32 >.

43.86.2.4 write() [2/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< std::complex< double >, int64 >.

43.86.2.5 write() [3/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

 $Implements\ gko::Writable To Matrix Data < std::complex < float >, int 32 >.$

43.86.2.6 write() [4/4]

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< std::complex< float >, int64 >.

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/fft.hpp

43.87 gko::multigrid::FixedCoarsening< ValueType, IndexType > Class Template Reference

FixedCoarsening is a very simple coarse grid generation algorithm.

#include <ginkgo/core/multigrid/fixed_coarsening.hpp>

Public Member Functions

 std::shared_ptr< const LinOp > get_system_matrix () const Returns the system operator (matrix) of the linear system.

43.87.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::multigrid::FixedCoarsening< ValueType, IndexType >

FixedCoarsening is a very simple coarse grid generation algorithm.

It selects the coarse matrix from the fine matrix by with user-specified indices.

The user needs to specify the indices (with global numbering) of the fine matrix, they wish to be in the coarse matrix. The restriction and prolongation matrices will map to and from the coarse space without any interpolation or weighting.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.87.2 Member Function Documentation

43.87.2.1 get_system_matrix()

```
template<typename ValueType = default_precision, typename IndexType = int32> std::shared_ptr<const LinOp> gko::multigrid::FixedCoarsening< ValueType, IndexType >::get_← system_matrix ( ) const [inline]
```

Returns the system operator (matrix) of the linear system.

Returns

the system operator (matrix)

The documentation for this class was generated from the following file:

ginkgo/core/multigrid/fixed_coarsening.hpp

43.88 gko::solver::Gcr< ValueType > Class Template Reference

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

```
#include <ginkgo/core/solver/gcr.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

· bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

• size_type get_krylov_dim () const

Gets the Krylov dimension of the solver.

void set_krylov_dim (size_type other)

Sets the Krylov dimension.

43.88.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::Gcr< ValueType >
```

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of GCR are merged into one step. Modified Gram-Schmidt is used.

Template Parameters

```
ValueType precision of matrix elements
```

43.88.2 Member Function Documentation

43.88.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Gcr< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
95 { return true; }
```

43.88.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gcr< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.88.2.3 get_krylov_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Gcr< ValueType >::get_krylov_dim ( ) const [inline]
```

Gets the Krylov dimension of the solver.

Returns

the Krylov dimension

43.88.2.4 set_krylov_dim()

Sets the Krylov dimension.

Parameters

other the new Krylov dimension

43.88.2.5 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gcr< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

· ginkgo/core/solver/gcr.hpp

43.89 gko::solver::Gmres < ValueType > Class Template Reference

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

#include <ginkgo/core/solver/gmres.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

· bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

size_type get_krylov_dim () const

Gets the Krylov dimension of the solver.

void set_krylov_dim (size_type other)

Sets the Krylov dimension.

43.89.1 Detailed Description

```
\label{template} \begin{tabular}{ll} template < typename \ ValueType = default\_precision > \\ class \ gko::solver::Gmres < ValueType > \\ \end{tabular}
```

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of GMRES are merged into 2 separate steps. Modified Gram-Schmidt is used.

Template Parameters

ValueType	precision of matrix elements
-----------	------------------------------

43.89.2 Member Function Documentation

43.89.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Gmres< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
96 { return true; }
```

43.89.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gmres< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.89.2.3 get_krylov_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Gmres< ValueType >::get_krylov_dim ( ) const [inline]
```

Gets the Krylov dimension of the solver.

Returns

the Krylov dimension

43.89.2.4 set_krylov_dim()

Sets the Krylov dimension.

Parameters

other	the new Krylov dimension
-------	--------------------------

43.89.2.5 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gmres< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/gmres.hpp

43.90 gko::solver::has_with_criteria < SolverType, typename > Struct Template Reference

Helper structure to test if the Factory of SolverType has a function with_criteria.

```
#include <ginkgo/core/solver/solver_traits.hpp>
```

43.90.1 Detailed Description

```
template<typename SolverType, typename = void>
struct gko::solver::has_with_criteria< SolverType, typename >
```

 $\label{thm:continuous} \mbox{Helper structure to test if the Factory of Solver} \mbox{Type has a function $\mbox{with_criteria}$.}$

Contains a constexpr boolean value, which is true if the Factory class of SolverType has a with_criteria, and false otherwise.

Template Parameters

SolverType	Solver to test if its factory has a with_criteria function.
------------	---

The documentation for this struct was generated from the following file:

ginkgo/core/solver/solver_traits.hpp

43.91 gko::solver::has_with_criteria < SolverType, xstd::void_t < decltype(SolverType::build().with_criteria(std::shared_ptr < const stop::CriterionFactory >())) > Struct Template Reference

Helper structure to test if the Factory of SolverType has a function with_criteria.

#include <ginkgo/core/solver/solver_traits.hpp>

43.91.1 Detailed Description

 $\label{template} $$ \textbf{typename SolverType} > $$ \textbf{struct gko::solver::has_with_criteria} < \textbf{SolverType, xstd::void_t} < \textbf{decltype}(\textbf{SolverType::build().with_criteria(std::shared_ptr} < \textbf{const stop::CriterionFactory} > ()))} > $$$

Helper structure to test if the Factory of SolverType has a function with_criteria.

Contains a constexpr boolean value, which is true if the Factory class of SolverType has a with_criteria, and false otherwise.

Template Parameters

SolverType | Solver to test if its factory has a with_criteria function.

The documentation for this struct was generated from the following file:

ginkgo/core/solver/solver_traits.hpp

43.92 gko::hip_stream Class Reference

An RAII wrapper for a custom HIP stream.

#include <ginkgo/core/base/executor.hpp>

Public Member Functions

hip_stream (int device_id=0)

Creates a new custom HIP stream.

• ∼hip_stream ()

Destroys the custom HIP stream, if it wasn't moved-from already.

hip_stream (hip_stream &&)

Move-constructs from an existing stream, which will be emptied.

hip_stream & operator= (hip_stream &&)=delete

Move-assigns from an existing stream, which will be emptied.

CUstream_st * get () const

Returns the native HIP stream handle.

43.92.1 Detailed Description

An RAII wrapper for a custom HIP stream.

The stream will be created on construction and destroyed when the lifetime of the wrapper ends.

43.92.2 Member Function Documentation

43.92.2.1 get()

```
CUstream_st* gko::hip_stream::get ( ) const
```

Returns the native HIP stream handle.

In a moved-from hip_stream, this will return nullptr.

The documentation for this class was generated from the following file:

· ginkgo/core/base/executor.hpp

43.93 gko::HipblasError Class Reference

HipblasError is thrown when a hipBLAS routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• HipblasError (const std::string &file, int line, const std::string &func, int64 error_code)

Initializes a hipBLAS error.

43.93.1 Detailed Description

HipblasError is thrown when a hipBLAS routine throws a non-zero error code.

43.93.2 Constructor & Destructor Documentation

43.93.2.1 HipblasError()

Initializes a hipBLAS error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the hipBLAS routine that failed	
error_code	The resulting hipBLAS error code	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.94 gko::HipError Class Reference

HipError is thrown when a HIP routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

HipError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a HIP error.

43.94.1 Detailed Description

HipError is thrown when a HIP routine throws a non-zero error code.

43.94.2 Constructor & Destructor Documentation

43.94.2.1 HipError()

Initializes a HIP error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the HIP routine that failed	
error_code	The resulting HIP error code	

Generated by Doxygen

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.95 gko::HipExecutor Class Reference

This is the Executor subclass which represents the HIP enhanced device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

std::shared_ptr< Executor > get_master () noexcept override

Returns the master OmpExecutor of this Executor.

• std::shared_ptr< const Executor > get_master () const noexcept override

Returns the master OmpExecutor of this Executor.

· void synchronize () const override

Synchronize the operations launched on the executor with its master.

• int get_device_id () const noexcept

Get the HIP device id of the device associated to this executor.

int get_num_warps_per_sm () const noexcept

Get the number of warps per SM of this executor.

· int get_num_multiprocessor () const noexcept

Get the number of multiprocessor of this executor.

• int get_major_version () const noexcept

Get the major verion of compute capability.

• int get_minor_version () const noexcept

Get the minor verion of compute capability.

• int get_num_warps () const noexcept

Get the number of warps of this executor.

int get_warp_size () const noexcept

Get the warp size of this executor.

hipblasContext * get_hipblas_handle () const

Get the hipblas handle for this executor.

hipsparseContext * get_hipsparse_handle () const

Get the hipsparse handle for this executor.

• int get_closest_numa () const

Get the closest NUMA node.

std::vector< int > get_closest_pus () const

Get the closest PUs.

Static Public Member Functions

• static std::shared_ptr< HipExecutor > create (int device_id, std::shared_ptr< Executor > master, bool device_reset=false, allocation_mode alloc_mode=default_hip_alloc_mode, CUstream_st *stream=nullptr)

Creates a new HipExecutor.

static int get_num_devices ()

Get the number of devices present on the system.

43.95.1 Detailed Description

This is the Executor subclass which represents the HIP enhanced device.

43.95.2 Member Function Documentation

43.95.2.1 create()

```
static std::shared_ptr<HipExecutor> gko::HipExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    bool device_reset = false,
    allocation_mode alloc_mode = default_hip_alloc_mode,
    CUstream_st * stream = nullptr ) [static]
```

Creates a new HipExecutor.

Parameters

device_id	the HIP device id of this device
master	an executor on the host that is used to invoke the device kernels
device_reset	whether to reset the device after the object exits the scope.
alloc_mode	the allocation mode that the executor should operate on. See @allocation_mode for more details

43.95.2.2 get_closest_numa()

```
int gko::HipExecutor::get_closest_numa ( ) const [inline]
```

Get the closest NUMA node.

Returns

the closest NUMA node closest to this device

43.95.2.3 get_closest_pus()

```
\verb|std::vector<int>| gko::HipExecutor::get_closest_pus () const [inline]|
```

Get the closest PUs.

Returns

the array of PUs closest to this device

43.95.2.4 get_hipblas_handle()

```
hipblasContext* gko::HipExecutor::get_hipblas_handle ( ) const [inline]
```

Get the hipblas handle for this executor.

Returns

the hipblas handle (hipblasContext*) for this executor

43.95.2.5 get_hipsparse_handle()

```
hipsparseContext* gko::HipExecutor::get_hipsparse_handle ( ) const [inline]
```

Get the hipsparse handle for this executor.

Returns

the hipsparse handle (hipsparseContext*) for this executor

43.95.2.6 get_master() [1/2]

```
std::shared_ptr<const Executor> gko::HipExecutor::get_master ( ) const [override], [virtual],
[noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.95.2.7 get_master() [2/2]

```
std::shared_ptr<Executor> gko::HipExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

The documentation for this class was generated from the following file:

• ginkgo/core/base/executor.hpp

43.96 gko::HipfftError Class Reference

HipfftError is thrown when a hipFFT routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

HipfftError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a hipFFT error.

43.96.1 Detailed Description

HipfftError is thrown when a hipFFT routine throws a non-zero error code.

43.96.2 Constructor & Destructor Documentation

43.96.2.1 HipfftError()

Initializes a hipFFT error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the hipFFT routine that failed	
error_code	The resulting hipFFT error code	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.97 gko::HiprandError Class Reference

HiprandError is thrown when a hipRAND routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

HiprandError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a hipRAND error.

43.97.1 Detailed Description

HiprandError is thrown when a hipRAND routine throws a non-zero error code.

43.97.2 Constructor & Destructor Documentation

43.97.2.1 HiprandError()

Initializes a hipRAND error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the hipRAND routine that failed	
error_code	The resulting hipRAND error code	

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.98 gko::HipsparseError Class Reference

HipsparseError is thrown when a hipSPARSE routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

HipsparseError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a hipSPARSE error.

43.98.1 Detailed Description

HipsparseError is thrown when a hipSPARSE routine throws a non-zero error code.

43.98.2 Constructor & Destructor Documentation

43.98.2.1 HipsparseError()

Initializes a hipSPARSE error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the hipSPARSE routine that failed	
error_code	The resulting hipSPARSE error code	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.99 gko::HipTimer Class Reference

A timer using events for timing on a HipExecutor.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

• void record (time_point &time) override

Records a time point at the current time.

void wait (time_point &time) override

Waits until all kernels in-process when recording the time point are finished.

• std::chrono::nanoseconds difference_async (const time_point &start, const time_point &stop) override Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

43.99.1 Detailed Description

A timer using events for timing on a HipExecutor.

43.99.2 Member Function Documentation

43.99.2.1 difference_async()

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. timer->wait(stop) needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements gko::Timer.

The documentation for this class was generated from the following file:

• ginkgo/core/base/timer.hpp

43.100 gko::matrix::Hybrid< ValueType, IndexType > Class Template Reference

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

#include <ginkgo/core/matrix/hybrid.hpp>

Classes

· class automatic

automatic is a strategy_type which decides the number of stored elements per row of the ell part automatically.

· class column limit

column_limit is a strategy_type which decides the number of stored elements per row of the ell part by specifying the number of columns.

class imbalance_bounded_limit

imbalance bounded limit is a strategy type which decides the number of stored elements per row of the ell part.

· class imbalance_limit

imbalance_limit is a strategy_type which decides the number of stored elements per row of the ell part according to the percent.

· class minimal_storage_limit

minimal_storage_limit is a strategy_type which decides the number of stored elements per row of the ell part.

class strategy_type

strategy_type is to decide how to set the hybrid config.

Public Member Functions

· void read (const mat data &data) override

Reads a matrix from a matrix_data structure.

· void read (const device mat data &data) override

Reads a matrix from a device matrix data structure.

· void read (device_mat_data &&data) override

Reads a matrix from a device matrix data structure.

· void write (mat data &data) const override

Writes a matrix to a matrix_data structure.

std::unique_ptr< Diagonal
 ValueType > > extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

- $std::unique_ptr < absolute_type > compute_absolute$ () const override

Gets the AbsoluteLinOp.

void compute_absolute_inplace () override

Compute absolute inplace on each element.

value_type * get_ell_values () noexcept

Returns the values of the ell part.

const value_type * get_const_ell_values () const noexcept

Returns the values of the ell part.

index_type * get_ell_col_idxs () noexcept

Returns the column indexes of the ell part.

const index_type * get_const_ell_col_idxs () const noexcept

Returns the column indexes of the ell part.

• size_type get_ell_num_stored_elements_per_row () const noexcept

Returns the number of stored elements per row of ell part.

size_type get_ell_stride () const noexcept

Returns the stride of the ell part.

• size_type get_ell_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the ell part.

value_type & ell_val_at (size_type row, size_type idx) noexcept

Returns the idx-th non-zero element of the row-th row in the ell part.

value_type ell_val_at (size_type row, size_type idx) const noexcept

Returns the idx-th non-zero element of the row-th row in the ell part.

index_type & ell_col_at (size_type row, size_type idx) noexcept

Returns the idx-th column index of the row-th row in the ell part.

• index_type ell_col_at (size_type row, size_type idx) const noexcept

Returns the idx-th column index of the row-th row in the ell part.

const ell_type * get_ell () const noexcept

Returns the matrix of the ell part.

value_type * get_coo_values () noexcept

Returns the values of the coo part.

const value_type * get_const_coo_values () const noexcept

Returns the values of the coo part.

index_type * get_coo_col_idxs () noexcept

Returns the column indexes of the coo part.

const index_type * get_const_coo_col_idxs () const noexcept

Returns the column indexes of the coo part.

index_type * get_coo_row_idxs () noexcept

Returns the row indexes of the coo part.

const index_type * get_const_coo_row_idxs () const noexcept

Returns the row indexes of the coo part.

• size_type get_coo_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the coo part.

const coo_type * get_coo () const noexcept

Returns the matrix of the coo part.

size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

std::shared_ptr< strategy_type > get_strategy () const noexcept

Returns the strategy.

template<typename HybType >

 ${\tt std::shared_ptr} < {\tt typename\ HybType::strategy_type} > {\tt get_strategy\ ()\ const}$

Returns the current strategy allowed in given hybrid format.

Hybrid & operator= (const Hybrid &)

Copy-assigns a Hybrid matrix.

• Hybrid & operator= (Hybrid &&)

Move-assigns a Hybrid matrix.

• Hybrid (const Hybrid &)

Copy-assigns a Hybrid matrix.

• Hybrid (Hybrid &&)

Move-assigns a Hybrid matrix.

43.100.1 Detailed Description

 $\label{template} $$ \ensuremath{\sf template}$$ < typename \ensuremath{\sf ValueType}$ = $$ \ensuremath{\sf dexType}$ = $$ \ensuremath{\sf int32}$ > $$ \ensuremath{\sf class}$ \ensuremath{\sf gko}::matrix::Hybrid< \ensuremath{\sf ValueType}, \ensuremath{\sf IndexType}$ > $$ \ensuremath{\sf lndexType}$ > $$ \ensuremath{\sf lnde$

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

Achieve the excellent performance with a proper partition of ELLPACK and COO.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.100.2 Constructor & Destructor Documentation

43.100.2.1 Hybrid() [1/2]

Copy-assigns a Hybrid matrix.

Inherits the executor, copies the Ell and Coo matrices.

43.100.2.2 Hybrid() [2/2]

Move-assigns a Hybrid matrix.

Inherits the executor, moves the Ell and Coo matrices. The moved-from matrix is empty (0x0 with empty Ell/Coo matrices).

43.100.3 Member Function Documentation

43.100.3.1 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Hybrid< ValueType, IndexType >::compute_absolute
( ) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

 $\label{lem:lemonts} \mbox{Implements gko::} Enable \mbox{AbsoluteComputation} < \mbox{remove_complex} < \mbox{Hybrid} < \mbox{ValueType}, \mbox{IndexType} >>>.$

43.100.3.2 ell_col_at() [1/2]

Returns the idx-th column index of the row-th row in the ell part.

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

```
526 {
527          return ell_->col_at(row, idx);
528 }
```

43.100.3.3 ell_col_at() [2/2]

Returns the idx-th column index of the row-th row in the ell part.

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.100.3.4 ell_val_at() [1/2]

Returns the idx-th non-zero element of the row-th row in the ell part.

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.100.3.5 ell_val_at() [2/2]

Returns the idx-th non-zero element of the row-th row in the ell part.

Parameters

row	the row of the requested element
idx	the idx-th stored element of the row

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.100.3.6 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32> std::unique_ptr<Diagonal<ValueType> > gko::matrix::Hybrid< ValueType, IndexType >::extract← _diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

```
diag the vector into which the diagonal will be written
```

Implements gko::DiagonalExtractable < ValueType >.

43.100.3.7 get_const_coo_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_col_idxs ( )
const [inline], [noexcept]
```

Returns the column indexes of the coo part.

Returns

the column indexes of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.100.3.8 get_const_coo_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_row_idxs ( )
const [inline], [noexcept]
```

Returns the row indexes of the coo part.

Returns

the row indexes of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.100.3.9 get_const_coo_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_values ( ) const
[inline], [noexcept]
```

Returns the values of the coo part.

Returns

the values of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.100.3.10 get_const_ell_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_ell_col_idxs ( )
const [inline], [noexcept]
```

Returns the column indexes of the ell part.

Returns

the column indexes of the ell part

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.100.3.11 get_const_ell_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_ell_values ( ) const
[inline], [noexcept]
```

Returns the values of the ell part.

Returns

the values of the ell part

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.100.3.12 get_coo()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const coo_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo ( ) const [inline],
[noexcept]
```

Returns the matrix of the coo part.

Returns

the matrix of the coo part

43.100.3.13 get_coo_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the coo part.

Returns

the column indexes of the coo part.

43.100.3.14 get_coo_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_coo_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the coo part.

Returns

the number of elements explicitly stored in the coo part

43.100.3.15 get_coo_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_row_idxs () [inline], [noexcept]
```

Returns the row indexes of the coo part.

Returns

the row indexes of the coo part.

43.100.3.16 get_coo_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_values ( ) [inline], [noexcept]
```

Returns the values of the coo part.

Returns

the values of the coo part.

43.100.3.17 get_ell()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const ell_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell ( ) const [inline],
[noexcept]
```

Returns the matrix of the ell part.

Returns

the matrix of the ell part

43.100.3.18 get_ell_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the ell part.

Returns

the column indexes of the ell part

43.100.3.19 get_ell_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the ell part.

Returns

the number of elements explicitly stored in the ell part

43.100.3.20 get ell num stored elements per row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_num_stored_elements_per_row ( )
const [inline], [noexcept]
```

Returns the number of stored elements per row of ell part.

Returns

the number of stored elements per row of ell part

43.100.3.21 get_ell_stride()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_stride ( ) const [inline],
[noexcept]
```

Returns the stride of the ell part.

Returns

the stride of the ell part

43.100.3.22 get_ell_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell_values () [inline], [noexcept]
```

Returns the values of the ell part.

Returns

the values of the ell part

43.100.3.23 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

43.100.3.24 get_strategy() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename HybType >
std::shared_ptr<typename HybType::strategy_type> gko::matrix::Hybrid< ValueType, IndexType
>::get_strategy ( ) const
```

Returns the current strategy allowed in given hybrid format.

Template Parameters

```
HybType hybrid type
```

Returns

the strategy

43.100.3.25 get_strategy() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr< typename HybType::strategy_type > gko::matrix::Hybrid< ValueType, IndexType
>::get_strategy ( ) const [inline], [noexcept]
```

Returns the strategy.

Returns

the strategy

43.100.3.26 operator=() [1/2]

Copy-assigns a Hybrid matrix.

Preserves the executor, copy-assigns the Ell and Coo matrices.

43.100.3.27 operator=() [2/2]

Move-assigns a Hybrid matrix.

Preserves the executor, move-assigns the Ell and Coo matrices. The moved-from matrix is empty (0x0 with empty Ell/Coo matrices).

43.100.3.28 read() [1/3]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.100.3.29 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.100.3.30 read() [3/3]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

 $\label{lem:remark_remark} Reimplemented \ from \ gko:: Readable From Matrix Data < Value Type, \ Index Type >.$

43.100.3.31 write()

Writes a matrix to a matrix_data structure.

Parameters

data	the matrix_data structure
------	---------------------------

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/coo.hpp
- ginkgo/core/matrix/hybrid.hpp

43.101 gko::factorization::lc< ValueType, IndexType > Class Template Reference

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

#include <ginkgo/core/factorization/ic.hpp>

Additional Inherited Members

43.101.1 Detailed Description

template<typename ValueType = gko::default_precision, typename IndexType = gko::int32> class gko::factorization::lc< ValueType, IndexType >

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

More specifically, it consists of a lower triangular factor L and its conjugate transpose L^H with sparsity pattern $S(L + L^H) = S(A)$ fulfilling $LL^H = A$ at every non-zero location of A.

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

ginkgo/core/factorization/ic.hpp

43.102 gko::preconditioner::lc< LSolverType, IndexType > Class Template Reference

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H*x=b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

#include <ginkgo/core/preconditioner/ic.hpp>

Public Member Functions

std::shared_ptr< const l_solver_type > get_l_solver () const

Returns the solver which is used for the provided L matrix.

std::shared_ptr< const lh_solver_type > get_lh_solver () const

Returns the solver which is used for the L^A H matrix.

• std::unique ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

• Ic & operator= (const Ic &other)

Copy-assigns an IC preconditioner.

Ic & operator= (Ic &&other)

Move-assigns an IC preconditioner.

• Ic (const Ic &other)

Copy-constructs an IC preconditioner.

• Ic (Ic &&other)

Move-constructs an IC preconditioner.

43.102.1 Detailed Description

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32> class gko::preconditioner::lc< LSolverType, IndexType >
```

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H*x=b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

It allows to set both the solver for L defaulting to solver::LowerTrs, which is a direct triangular solvers. The solver for $L^{+}H$ is the conjugate-transposed solver for L, ensuring that the preconditioner is symmetric and positive-definite. For this L solver, a factory can be provided (using with_l_solver_factory) to have more control over their behavior. In particular, it is possible to use an iterative method for solving the triangular systems. The default parameters for an iterative triangluar solver are:

- reduction factor = 1e-4
- max iteration = <number of="" rows="" of="" the="" matrix="" given="" to="" the="" solver>=""> Solvers without such criteria can also be used, in which case none are set.

An object of this class can be created with a matrix or a gko::Composition containing two matrices. If created with a matrix, it is factorized before creating the solver. If a gko::Composition (containing two matrices) is used, the first operand will be taken as the L matrix, the second will be considered the L^H matrix, which helps to avoid the otherwise necessary transposition of L inside the solver. Parlc can be directly used, since it orders the factors in the correct way.

Note

When providing a gko::Composition, the first matrix must be the lower matrix (L), and the second matrix must be its conjugate-transpose (L^H). If they are swapped, solving might crash or return the wrong result.

Do not use symmetric solvers (like CG) for the L solver since both matrices (L and $L^{\wedge}H$) are, by design, not symmetric.

This class is not thread safe (even a const object is not) because it uses an internal cache to accelerate multiple (sequential) applies. Using it in parallel can lead to segmentation faults, wrong results and other unwanted behavior.

Template Parameters

LSolverType	type of the solver used for the L matrix. Defaults to solver::LowerTrs	1
IndexType	type of the indices when Parlc is used to generate the L and ${\rm L}^{\wedge}{\rm H}$ factors. Irrelevant otherwise.	1

43.102.2 Constructor & Destructor Documentation

43.102.2.1 lc() [1/2]

Copy-constructs an IC preconditioner.

Inherits the executor, shallow-copies the solvers and parameters.

```
237 : Ic{other.get_executor()} { *this = other; }
```

References gko::PolymorphicObject::get_executor().

43.102.2.2 lc() [2/2]

Move-constructs an IC preconditioner.

Inherits the executor, moves the solvers and parameters. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::PolymorphicObject::get_executor().

43.102.3 Member Function Documentation

43.102.3.1 conj_transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ic< LSolverType, IndexType >::conj_transpose ( )
const [inline], [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

References gko::PolymorphicObject::get_executor(), gko::preconditioner::lc< LSolverType, IndexType >::get_l \cdot \solver(), gko::preconditioner::lc< LSolverType, IndexType >::get_lh_solver(), gko::share(), and gko::transpose().

43.102.3.2 get_l_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::shared_ptr<const l_solver_type> gko::preconditioner::Ic< LSolverType, IndexType >::get
_l_solver ( ) const [inline]
```

Returns the solver which is used for the provided L matrix.

Returns

the solver which is used for the provided L matrix

Referenced by gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose(), and gko::preconditioner::lc< LSolverType, IndexType >::transpose().

43.102.3.3 get_lh_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::shared_ptr<const lh_solver_type> gko::preconditioner::Ic< LSolverType, IndexType > 
::get_lh_solver ( ) const [inline]
```

Returns the solver which is used for the L^H matrix.

Returns

the solver which is used for the L^H matrix

Referenced by gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose(), and gko::preconditioner::lc< LSolverType, IndexType >::transpose().

43.102.3.4 operator=() [1/2]

Copy-assigns an IC preconditioner.

Preserves the executor, shallow-copies the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor.

References gko::clone(), and gko::PolymorphicObject::get_executor().

43.102.3.5 operator=() [2/2]

Move-assigns an IC preconditioner.

Preserves the executor, moves the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::clone(), and gko::PolymorphicObject::get executor().

43.102.3.6 transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ic< LSolverType, IndexType >::transpose ( ) const
[inline], [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

References gko::PolymorphicObject::get_executor(), gko::preconditioner::lc< LSolverType, IndexType >::get_l \leftarrow solver(), gko::preconditioner::lc< LSolverType, IndexType >::get_lh_solver(), gko::share(), and gko::transpose().

The documentation for this class was generated from the following file:

• ginkgo/core/preconditioner/ic.hpp

43.103 gko::matrix::Identity < ValueType > Class Template Reference

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

#include <ginkgo/core/matrix/identity.hpp>

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

43.103.1 Detailed Description

```
template < typename ValueType = default_precision >
class gko::matrix::ldentity < ValueType >
```

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

Thus, objects of the Identity class always represent a square matrix, and don't require any storage for their values. The apply method is implemented as a simple copy (or a linear combination).

Note

This class is useful when composing it with other operators. For example, it can be used instead of a preconditioner in Krylov solvers, if one wants to run a "plain" solver, without using a preconditioner.

Template Parameters

ValueType	precision of matrix elements

43.103.2 Member Function Documentation

43.103.2.1 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Identity< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.103.2.2 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Identity< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/identity.hpp

43.104 gko::matrix::ldentityFactory< ValueType > Class Template Reference

This factory is a utility which can be used to generate Identity operators.

```
#include <ginkgo/core/matrix/identity.hpp>
```

Static Public Member Functions

static std::unique_ptr< IdentityFactory > create (std::shared_ptr< const Executor > exec)
 Creates a new Identity factory.

Additional Inherited Members

43.104.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::matrix::ldentityFactory< ValueType >
```

This factory is a utility which can be used to generate Identity operators.

The factory will generate the <u>Identity</u> matrix with the same dimension as the passed in operator. It will throw an exception if the operator is not square.

Template Parameters

ValueType precision of matrix elements	
--	--

43.104.2 Member Function Documentation

43.104.2.1 create()

Creates a new Identity factory.

Parameters

exec the executor where the Identity operator will be stored

Returns

a unique pointer to the newly created factory

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/identity.hpp

43.105 gko::solver::ldr< ValueType > Class Template Reference

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

```
#include <ginkgo/core/solver/idr.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
 - Returns a LinOp representing the conjugate transpose of the Transposable object.
- · bool apply uses initial guess () const override
 - Return true as iterative solvers use the data in x as an initial guess.
- size_type get_subspace_dim () const

Gets the subspace dimension of the solver.

void set_subspace_dim (const size_type other)

Sets the subspace dimension of the solver.

remove_complex< ValueType > get_kappa () const

Gets the kappa parameter of the solver.

void set_kappa (const remove_complex < ValueType > other)

Sets the kappa parameter of the solver.

· bool get deterministic () const

Gets the deterministic parameter of the solver.

void set_deterministic (const bool other)

Sets the deterministic parameter of the solver.

· bool get_complex_subspace () const

Gets the complex_subspace parameter of the solver.

void set complex subpsace (const bool other)

Sets the complex_subspace parameter of the solver.

43.105.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::ldr< ValueType >
```

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

The implemented version is the one presented in the paper "Algorithm 913: An elegant IDR(s) variant that efficiently exploits biorthogonality properties" by M. B. Van Gijzen and P. Sonneveld.

The method is based on the induced dimension reduction theorem which provides a way to construct subsequent residuals that lie in a sequence of shrinking subspaces. These subspaces are spanned by s vectors which are first generated randomly and then orthonormalized. They are stored in a dense matrix.

Template Parameters

ValueType precision of the elements of the system matrix.

43.105.2 Member Function Documentation

43.105.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Idr< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
103 { return true; }
```

43.105.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Idr< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.105.2.3 get_complex_subspace()

```
template<typename ValueType = default_precision>
bool gko::solver::Idr< ValueType >::get_complex_subspace ( ) const [inline]
```

Gets the complex_subspace parameter of the solver.

Returns

the complex_subspace parameter

43.105.2.4 get_deterministic()

```
template<typename ValueType = default_precision>
bool gko::solver::Idr< ValueType >::get_deterministic ( ) const [inline]
```

Gets the deterministic parameter of the solver.

Returns

the deterministic parameter

43.105.2.5 get_kappa()

```
template<typename ValueType = default_precision>
remove_complex<ValueType> gko::solver::Idr< ValueType >::get_kappa ( ) const [inline]
```

Gets the kappa parameter of the solver.

Returns

the kappa parameter

43.105.2.6 get_subspace_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Idr< ValueType >::get_subspace_dim ( ) const [inline]
```

Gets the subspace dimension of the solver.

Returns

the subspace Dimension

43.105.2.7 set_complex_subpsace()

Sets the complex_subspace parameter of the solver.

Parameters

```
other the new complex_subspace parameter
```

43.105.2.8 set_deterministic()

Sets the deterministic parameter of the solver.

Parameters

```
other the new deterministic parameter
```

43.105.2.9 set_kappa()

Sets the kappa parameter of the solver.

Parameters

other the new kappa	parameter
---------------------	-----------

43.105.2.10 set_subspace_dim()

Sets the subspace dimension of the solver.

Parameters

other	the new subspace Dimension
-------	----------------------------

43.105.2.11 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Idr< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/idr.hpp

43.106 gko::factorization::llu< ValueType, IndexType > Class Template Reference

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

#include <ginkgo/core/factorization/ilu.hpp>

Additional Inherited Members

43.106.1 Detailed Description

template<typename ValueType = gko::default_precision, typename IndexType = gko::int32> class gko::factorization::llu< ValueType, IndexType >

Represents an incomplete LU factorization - ILU(0) - of a sparse matrix.

More specifically, it consists of a lower unitriangular factor L and an upper triangular factor U with sparsity pattern S(L+U) = S(A) fulfilling LU = A at every non-zero location of A.

Template Parameters

ValueType	Type of the values of all matrices used in this class	
IndexType	Type of the indices of all matrices used in this class	

The documentation for this class was generated from the following file:

• ginkgo/core/factorization/ilu.hpp

43.107 gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType > Class Template Reference

The Incomplete LU (ILU) preconditioner solves the equation LUx = b for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

#include <ginkgo/core/preconditioner/ilu.hpp>

Public Member Functions

- std::shared_ptr< const l_solver_type > get_l_solver () const
 - Returns the solver which is used for the provided L matrix.
- $std::shared_ptr < const u_solver_type > get_u_solver$ () const
 - Returns the solver which is used for the provided U matrix.
- std::unique_ptr< LinOp > transpose () const override
 - Returns a LinOp representing the transpose of the Transposable object.
- std::unique_ptr< LinOp > conj_transpose () const override
 - Returns a LinOp representing the conjugate transpose of the Transposable object.
- Ilu & operator= (const Ilu &other)
 - Copy-assigns an ILU preconditioner.
- Ilu & operator= (Ilu &&other)
 - Move-assigns an ILU preconditioner.
- Ilu (const Ilu &other)
 - Copy-constructs an ILU preconditioner.
- Ilu (Ilu &&other)

Move-constructs an ILU preconditioner.

43.107.1 Detailed Description

template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply = false, typename IndexType = int32>

class gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >

The Incomplete LU (ILU) preconditioner solves the equation LUx = b for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

It allows to set both the solver for L and the solver for U independently, while providing the defaults solver::LowerTrs and solver::UpperTrs, which are direct triangular solvers. For these solvers, a factory can be provided (with with—l_solver_factory and with_u_solver_factory) to have more control over their behavior. In particular, it is possible to use an iterative method for solving the triangular systems. The default parameters for an iterative triangluar solver are:

- reduction factor = 1e-4
- max iteration = <number of="" rows="" of="" the="" matrix="" given="" to="" the="" solver>=""> Solvers without such criteria can also be used, in which case none are set.

An object of this class can be created with a matrix or a gko::Composition containing two matrices. If created with a matrix, it is factorized before creating the solver. If a gko::Composition (containing two matrices) is used, the first operand will be taken as the L matrix, the second will be considered the U matrix. Parllu can be directly used, since it orders the factors in the correct way.

Note

When providing a gko::Composition, the first matrix must be the lower matrix (L), and the second matrix must be the upper matrix (U). If they are swapped, solving might crash or return the wrong result.

Do not use symmetric solvers (like CG) for L or U solvers since both matrices (L and U) are, by design, not symmetric.

This class is not thread safe (even a const object is not) because it uses an internal cache to accelerate multiple (sequential) applies. Using it in parallel can lead to segmentation faults, wrong results and other unwanted behavior.

Template Parameters

LSolverType	type of the solver used for the L matrix. Defaults to solver::LowerTrs
USolverType	type of the solver used for the U matrix Defaults to solver::UpperTrs
ReverseApply	default behavior (ReverseApply = false) is first to solve with L (Ly = b) and then with U (Ux = y). When set to true, it will solve first with U, and then with L.
IndexTypeParllu	Type of the indices when Parllu is used to generate both L and U factors. Irrelevant
	otherwise.

43.107.2 Constructor & Destructor Documentation

43.107.2.1 Ilu() [1/2]

Copy-constructs an ILU preconditioner.

Inherits the executor, shallow-copies the solvers and parameters.
255 : Ilu{other.get_executor()} { *this = other; }

References gko::PolymorphicObject::get_executor().

43.107.2.2 Ilu() [2/2]

Move-constructs an ILU preconditioner.

Inherits the executor, moves the solvers and parameters. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::PolymorphicObject::get executor().

43.107.3 Member Function Documentation

43.107.3.1 conj_transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper←
Trs<>, bool ReverseApply = false, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply,
IndexType >::conj_transpose () const [inline], [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

References gko::PolymorphicObject::get_executor(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_l_solver(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_u_solver(), gko::share(), and gko::transpose().

43.107.3.2 get | solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper←
Trs<>, bool ReverseApply = false, typename IndexType = int32>
std::shared_ptr<const l_solver_type> gko::preconditioner::Ilu< LSolverType, USolverType,
ReverseApply, IndexType >::get_l_solver () const [inline]
```

Returns the solver which is used for the provided L matrix.

Returns

the solver which is used for the provided L matrix

Referenced by gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose().

43.107.3.3 get_u_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper←
Trs<>, bool ReverseApply = false, typename IndexType = int32>
std::shared_ptr<const u_solver_type> gko::preconditioner::Ilu< LSolverType, USolverType,
ReverseApply, IndexType >::get_u_solver () const [inline]
```

Returns the solver which is used for the provided U matrix.

Returns

the solver which is used for the provided U matrix

Referenced by gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose().

43.107.3.4 operator=() [1/2]

Copy-assigns an ILU preconditioner.

Preserves the executor, shallow-copies the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor.

References gko::clone(), and gko::PolymorphicObject::get_executor().

43.107.3.5 operator=() [2/2]

Move-assigns an ILU preconditioner.

Preserves the executor, moves the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::clone(), and gko::PolymorphicObject::get_executor().

43.107.3.6 transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper←
Trs<>, bool ReverseApply = false, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply,
IndexType >::transpose () const [inline], [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

References gko::PolymorphicObject::get_executor(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_l_solver(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_u_solver(), gko::share(), and gko::transpose().

The documentation for this class was generated from the following file:

· ginkgo/core/preconditioner/ilu.hpp

43.108 gko::matrix::Hybrid< ValueType, IndexType >::imbalance bounded limit Class Reference

imbalance_bounded_limit is a strategy_type which decides the number of stored elements per row of the ell part.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- imbalance_bounded_limit (double percent=0.8, double ratio=0.0001)
 - Creates a imbalance_bounded_limit strategy.
- size_type compute_ell_num_stored_elements_per_row (array < size_type > *row_nnz) const override
 Computes the number of stored elements per row of the ell part.
- auto get_percentage () const

Get the percent setting.

auto get_ratio () const

Get the ratio setting.

43.108.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit
```

imbalance_bounded_limit is a strategy_type which decides the number of stored elements per row of the ell part.

It uses the imbalance_limit and adds the upper bound of the number of ell's cols by the number of rows.

43.108.2 Member Function Documentation

43.108.2.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32> size_type gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_← num_stored_elements_per_row (

array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

row_nnz	the number of nonzeros of each row
---------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements gko::matrix::Hybrid< ValueType, IndexType >::strategy_type.

References gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements -- __per_row(), and gko::array< ValueType >::get_num_elems().

Referenced by gko::matrix::Hybrid < ValueType, IndexType >::automatic::compute_ell_num_stored_elements_ \leftarrow per_row().

43.108.2.2 get percentage()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_percentage ( )
const [inline]
```

Get the percent setting.

@retrun percent

References gko::matrix::Hybrid < ValueType, IndexType >::imbalance limit::get percentage().

43.108.2.3 get_ratio()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_ratio ( ) const
[inline]
```

Get the ratio setting.

@retrun ratio

The documentation for this class was generated from the following file:

ginkgo/core/matrix/hybrid.hpp

43.109 gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit Class Reference

imbalance_limit is a strategy_type which decides the number of stored elements per row of the ell part according to the percent.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

• imbalance_limit (double percent=0.8)

Creates a imbalance_limit strategy.

- size_type compute_ell_num_stored_elements_per_row (array < size_type > *row_nnz) const override
 Computes the number of stored elements per row of the ell part.
- auto get_percentage () const
 Get the percent setting.

43.109.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit
```

imbalance_limit is a strategy_type which decides the number of stored elements per row of the ell part according to the percent.

It sorts the number of nonzeros of each row and takes the value at the position floor (percent * num_row) as the number of stored elements per row of the ell part. Thus, at least percent rows of all are in the ell part.

43.109.2 Constructor & Destructor Documentation

43.109.2.1 imbalance limit()

Creates a imbalance_limit strategy.

Parameters

percent	the row_nnz[floor(num_rows*percent)] is the number of stored elements per row of the ell part
<i>i</i>	

43.109.3 Member Function Documentation

43.109.3.1 compute_ell_num_stored_elements_per_row()

Computes the number of stored elements per row of the ell part.

Parameters

Returns

the number of stored elements per row of the ell part

Implements gko::matrix::Hybrid< ValueType, IndexType >::strategy_type.

References gko::array< ValueType >::get_data(), and gko::array< ValueType >::get_num_elems().

Referenced by gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_ \hookleftarrow stored_elements_per_row(), and gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::compute \hookleftarrow _ell_num_stored_elements_per_row().

43.109.3.2 get_percentage()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::get_percentage ( ) const
[inline]
```

Get the percent setting.

@retrun percent

Referenced by gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_percentage(), and gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::get_percentage().

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/hybrid.hpp

43.110 gko::stop::ImplicitResidualNorm< ValueType > Class Template Reference

The ImplicitResidualNorm class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

#include <ginkgo/core/stop/residual_norm.hpp>

43.110.1 Detailed Description

template<typename ValueType = default_precision> class gko::stop::ImplicitResidualNorm< ValueType >

The ImplicitResidualNorm class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

- the norm of the right-hand side, implicit_resnorm / norm(right_hand_side) < threshold
- 2. the initial residual, implicit_resnorm / norm(initial_residual) < < threshold.
- 3. one, implicit_resnorm < threshold.

Note

To use this stopping criterion there are some dependencies. The constructor depends on either b or the initial_residual in order to compute their norms. If this is not correctly provided, an exception ::gko ::NotSupported() is thrown.

The documentation for this class was generated from the following file:

• ginkgo/core/stop/residual_norm.hpp

43.111 gko::index_set< IndexType > Class Template Reference

An index set class represents an ordered set of intervals.

#include <ginkgo/core/base/index_set.hpp>

Public Types

using index_type = IndexType

The type of elements stored in the index set.

Public Member Functions

index_set (std::shared_ptr< const Executor > exec) noexcept

Creates an empty index_set tied to the specified Executor.

index_set (std::shared_ptr< const gko::Executor > exec, std::initializer_list< IndexType > init_list, const bool
is sorted=false)

Creates an index set on the specified executor from the initializer list.

• index_set (std::shared_ptr< const gko::Executor > exec, const index_type size, const gko::array< index_type > &indices, const bool is sorted=false)

Creates an index set on the specified executor and the given size.

index_set (std::shared_ptr< const Executor > exec, const index_set &other)

Creates a copy of the input index_set on a different executor.

index_set (const index_set &other)

Creates a copy of the input index_set.

index_set (std::shared_ptr< const Executor > exec, index_set &&other)

Moves the input index_set to a different executor.

index_set (index_set &&other)

Moves the input index_set.

index_set & operator= (const index_set &other)

Copies data from another index_set.

index_set & operator= (index_set &&other)

Moves data from another index_set.

• void clear () noexcept

Deallocates all data used by the index_set.

std::shared_ptr< const Executor > get_executor () const

Returns the executor of the index_set.

• index_type get_size () const

Returns the size of the index set space.

· bool is_contiguous () const

Returns if the index set is contiguous.

• index_type get_num_elems () const

Return the actual number of indices stored in the index set.

index_type get_global_index (index_type local_index) const

Return the global index given a local index.

index_type get_local_index (index_type global_index) const

Return the local index given a global index.

array< index_type > map_local_to_global (const array< index_type > &local_indices, const bool is_

 sorted=false) const

This is an array version of the scalar function above.

This is an array version of the scalar function above.

array< index_type > to_global_indices () const

This function allows the user obtain a decompresed global_indices array from the indices stored in the index set.

array< bool > contains (const array< index_type > &global_indices, const bool is_sorted=false) const

Checks if the individual global indeices exist in the index set.

• bool contains (const index_type global_index) const

Checks if the global index exists in the index set.

index_type get_num_subsets () const

Returns the number of subsets stored in the index set.

const index_type * get_subsets_begin () const

Returns a pointer to the beginning indices of the subsets.

const index_type * get_subsets_end () const

Returns a pointer to the end indices of the subsets.

const index_type * get_superset_indices () const

Returns a pointer to the cumulative indices of the superset of the subsets.

43.111.1 Detailed Description

```
template < typename IndexType = int32 > class gko::index_set < IndexType >
```

An index set class represents an ordered set of intervals.

The index set contains subsets which store the starting and end points of a range, [a,b), storing the first index and one past the last index. As the index set only stores the end-points of ranges, it can be quite efficient in terms of storage.

This class is particularly useful in storing continuous ranges. For example, consider the index set (1, 2, 3, 4, 5, 6, 7, 8, 10, 11, 12, 18, 19, 20, 21, 42). Instead of storing the entire array of indices, one can store intervals ([1,9), [10,13), [18,22), [42,43)), thereby only using half the storage.

We store three arrays, one (subsets_begin) with the starting indices of the subsets in the index set, another (subsets_end) storing one index beyond the end indices of the subsets and the last (superset_cumulative_indices) storing the cumulative number of indices in the subsequent subsets with an initial zero which speeds up the querying. Additionally, the arrays conataining the range boundaries (subsets_begin, subsets_end) are stored in a sorted fashion.

Therefore the storage would look as follows

```
index\_set = (1, 2, 3, 4, 5, 6, 7, 8, 10, 11, 12, 18, 19, 20, 21, 42) subsets\_begin = \{1, 10, 18, 42\} subsets\_end = \{9, 13, 22, 43\} superset\_cumulative\_indices = \{0, 8, 11, 15, 16\}
```

Template Parameters

```
index_type | type of the indices being stored in the index set.
```

43.111.2 Constructor & Destructor Documentation

43.111.2.1 index_set() [1/7]

Creates an empty index set tied to the specified Executor.

Parameters

exec the Executor where the index_set data is allocated

```
98 : exec_(std::move(exec)),
99 index_space_size_{0},
100 num_stored_indices_{0},
101 subsets_begin_{array<index_type>(exec_)},
102 subsets_end_{array<index_type>(exec_)},
103 superset_cumulative_indices_{array<index_type>(exec_)}
104 {}
```

43.111.2.2 index_set() [2/7]

Creates an index set on the specified executor from the initializer list.

Parameters

exec	the Executor where the index set data will be allocated
init_list	the indices that the index set should hold in an initializer_list.
is_sorted	a parameter that specifies if the indices array is sorted or not. true if sorted.

43.111.2.3 index_set() [3/7]

Creates an index set on the specified executor and the given size.

Parameters

exec	the Executor where the index set data will be allocated	
size	the maximum index the index set it allowed to hold. This is the size of the index space.	
indices	the indices that the index set should hold.	
is_sorted	a parameter that specifies if the indices array is sorted or not. true if sorted.	

References gko::array< ValueType >::get_num_elems().

43.111.2.4 index_set() [4/7]

Creates a copy of the input index set on a different executor.

Parameters

exec	the executor where the new index_set will be created
other	the index_set to copy from

43.111.2.5 index_set() [5/7]

Creates a copy of the input index_set.

Parameters

```
other the index_set to copy from
```

43.111.2.6 index_set() [6/7]

Moves the input index_set to a different executor.

Parameters

exec	the executor where the new index_set will be moved to
other	the index_set to move from

43.111.2.7 index_set() [7/7]

template<typename IndexType = int32>

Moves the input index_set.

Parameters

other the index	_set to move from
-----------------	-------------------

43.111.3 Member Function Documentation

43.111.3.1 clear()

```
template<typename IndexType = int32>
void gko::index_set< IndexType >::clear ( ) [inline], [noexcept]
```

Deallocates all data used by the index_set.

The index_set is left in a valid, but empty state, so the same index_set can be used to allocate new memory. Calls to index_set::get_subsets_begin() will return a nullptr.

References gko::array< ValueType >::clear().

43.111.3.2 contains() [1/2]

Checks if the individual global indeices exist in the index set.

Parameters

global_indices	the indices to check.	
is_sorted	a parameter that specifies if the query array is sorted or not. true if sorted.	

Returns

the array that contains element wise whether the corresponding global index in the index set or not.

43.111.3.3 contains() [2/2]

Checks if the global index exists in the index set.

Parameters

global_index the index to check.

Returns

whether the element exists in the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

43.111.3.4 get_executor()

```
template<typename IndexType = int32>
std::shared_ptr<const Executor> gko::index_set< IndexType >::get_executor ( ) const [inline]
```

Returns the executor of the index_set.

Returns

the executor.

43.111.3.5 get_global_index()

Return the global index given a local index.

```
Consider the set idx\_set = (0, 1, 2, 4, 6, 7, 8, 9). This function returns the element at the global index k stored in the index set. For example, idx\_set.get\_global\_index(0) == 0 idx\_set.get\_global\_index(3) == 4 and <math>idx\_set.get\_global\_index(7) == 9
```

Note

This function returns a scalar value and needs a scalar value. For repeated queries, it is more efficient to use the array functions that take and return arrays which allow for more throughput.

Parameters

local index	the local index.

Returns

the global index from the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

43.111.3.6 get_local_index()

Return the local index given a global index.

Consider the set $idx_set = (0, 1, 2, 4, 6, 7, 8, 9)$. This function returns the local index in the index set of the provided index set. For example, $idx_set_get_local_index(0) == 0 idx_set_get_local_index(4) == 3 and <math>idx_set_get_local_index(6) == 4$.

Note

This function returns a scalar value and needs a scalar value. For repeated queries, it is more efficient to use the array functions that take and return arrays which allow for more throughput.

Parameters

global_index	the global index.
--------------	-------------------

Returns

the local index of the element in the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

43.111.3.7 get_num_elems()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_num_elems ( ) const [inline]
```

Return the actual number of indices stored in the index set.

Returns

number of indices stored in the index set

43.111.3.8 get_num_subsets()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_num_subsets ( ) const [inline]
```

Returns the number of subsets stored in the index set.

Returns

the number of stored subsets.

References gko::array< ValueType >::get num elems().

Referenced by gko::index_set< IndexType >::is_contiguous().

43.111.3.9 get_size()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_size ( ) const [inline]
```

Returns the size of the index set space.

Returns

the size of the index set space.

43.111.3.10 get_subsets_begin()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_subsets_begin ( ) const [inline]
```

Returns a pointer to the beginning indices of the subsets.

Returns

a pointer to the beginning indices of the subsets.

References gko::array< ValueType >::get_const_data().

43.111.3.11 get_subsets_end()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_subsets_end ( ) const [inline]
```

Returns a pointer to the end indices of the subsets.

Returns

a pointer to the end indices of the subsets.

References gko::array< ValueType >::get const data().

43.111.3.12 get superset indices()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_superset_indices ( ) const [inline]
```

Returns a pointer to the cumulative indices of the superset of the subsets.

Returns

a pointer to the cumulative indices of the superset of the subsets.

References gko::array< ValueType >::get_const_data().

43.111.3.13 is_contiguous()

```
template<typename IndexType = int32>
bool gko::index_set< IndexType >::is_contiguous ( ) const [inline]
```

Returns if the index set is contiguous.

Returns

if the index set is contiguous.

References gko::index_set< IndexType >::get_num_subsets().

43.111.3.14 map_global_to_local()

This is an array version of the scalar function above.

Parameters

global_indices	the global index array.
is_sorted	a parameter that specifies if the query array is sorted or not. true if sorted.

Returns

the local index array from the index set.

Note

Whenever possible, passing a sorted array is preferred as the queries can be significantly faster.

43.111.3.15 map_local_to_global()

This is an array version of the scalar function above.

Parameters

local_indices	the local index array.
is_sorted	a parameter that specifies if the query array is sorted or not. $\verb true $ if sorted .

Returns

the global index array from the index set.

Note

Whenever possible, passing a sorted array is preferred as the queries can be significantly faster. Passing local indices from [0, size) is equivalent to using the @to_global_indices function.

43.111.3.16 operator=() [1/2]

Copies data from another index_set.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

Returns

this

43.111.3.17 operator=() [2/2]

Moves data from another index_set.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

other	the index	set to move from	
-------	-----------	------------------	--

Returns

this

43.111.3.18 to_global_indices()

```
template<typename IndexType = int32>
array<index_type> gko::index_set< IndexType >::to_global_indices ( ) const
```

This function allows the user obtain a decompresed global_indices array from the indices stored in the index set.

Returns

the decompressed set of indices.

The documentation for this class was generated from the following file:

ginkgo/core/base/index_set.hpp

43.112 gko::solver::lr< ValueType > Class Template Reference

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

```
#include <ginkgo/core/solver/ir.hpp>
```

Public Member Functions

```
    std::unique_ptr< LinOp > transpose () const override
```

Returns a LinOp representing the transpose of the Transposable object.

std::unique ptr< LinOp > conj transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

bool apply_uses_initial_guess () const override

Return true as iterative solvers use the data in x as an initial guess.

std::shared_ptr< const LinOp > get_solver () const

Returns the solver operator used as the inner solver.

void set_solver (std::shared_ptr< const LinOp > new_solver)

Sets the solver operator used as the inner solver.

Ir & operator= (const Ir &)

Copy-assigns an IR solver.

Ir & operator= (Ir &&)

Move-assigns an IR solver.

Ir (const Ir &)

Copy-constructs an IR solver.

• lr (lr &&)

Move-constructs an IR solver.

43.112.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::solver::lr< ValueType >
```

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

Moreover, it can be also considered as preconditioned Richardson iteration with relaxation factor = 1.

For any approximation of the solution solution to the system Ax = b, the residual is defined as: residual = b - A solution. The error in solution, e = x - solution (with x being the exact solution) can be obtained as the solution to the residual equation Ae = residual, since Ae = Ax - A solution = b - A solution = residual. Then, the real solution is computed as $x = relaxation_factor * solution + e$. Instead of accurately solving the residual equation Ae = residual, the solution of the system e can be approximated to obtain the approximation error using a coarse method solver, which is used to update solution, and the entire process is repeated with the updated solution. This yields the iterative refinement method:

```
solution = initial_guess
while not converged:
    residual = b - A solution
    error = solver(A, residual)
    solution = solution + relaxation_factor * error
```

With relaxation_factor equal to 1 (default), the solver is Iterative Refinement, with relaxation_factor equal to a value other than 1, the solver is a Richardson iteration, with possibility for additional preconditioning.

Assuming that solver has accuracy c, i.e., | e - error | <= c | e |, iterative refinement will converge with a convergence rate of c. Indeed, from e - error = x - solution - error = x - solution* (where solution* denotes the value stored in solution after the update) and <math>e = inv(A) residual = inv(A)b - inv(A) A solution = x - solution it follows that | x - solution* | <= c | x - solution |.

Unless otherwise specified via the solver factory parameter, this implementation uses the identity operator (i.e. the solver that approximates the solution of a system Ax = b by setting x := b) as the default inner solver. Such a setting results in a relaxation method known as the Richardson iteration with parameter 1, which is guaranteed to converge for matrices whose spectrum is strictly contained within the unit disc around 1 (i.e., all its eigenvalues lambda have to satisfy the equation '|relaxation_factor * lambda - 1| < 1).

Template Parameters

ValueType	precision of matrix elements
-----------	------------------------------

43.112.2 Constructor & Destructor Documentation

43.112.2.1 Ir() [1/2]

Copy-constructs an IR solver.

Inherits the executor, shallow-copies inner solver, stopping criterion and system matrix.

43.112.2.2 Ir() [2/2]

Move-constructs an IR solver.

Preserves the executor, moves inner solver, stopping criterion and system matrix. The moved-from object is empty (0x0 and nullptr inner solver, stopping criterion and system matrix)

43.112.3 Member Function Documentation

43.112.3.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Ir< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

References gko::solver::provided.

43.112.3.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Ir< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.112.3.3 get_solver()

```
template<typename ValueType = default_precision>
std::shared_ptr<const LinOp> gko::solver::Ir< ValueType >::get_solver ( ) const [inline]
```

Returns the solver operator used as the inner solver.

Returns

the solver operator used as the inner solver

43.112.3.4 operator=() [1/2]

Copy-assigns an IR solver.

Preserves the executor, shallow-copies inner solver, stopping criterion and system matrix. If the executors mismatch, clones inner solver, stopping criterion and system matrix onto this executor.

43.112.3.5 operator=() [2/2]

Move-assigns an IR solver.

Preserves the executor, moves inner solver, stopping criterion and system matrix. If the executors mismatch, clones inner solver, stopping criterion and system matrix onto this executor. The moved-from object is empty (0x0 and nullptr inner solver, stopping criterion and system matrix)

43.112.3.6 set solver()

Sets the solver operator used as the inner solver.

Parameters

```
new_solver the new inner solver
```

43.112.3.7 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Ir< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/ir.hpp

43.113 gko::preconditioner::lsai< lsaiType, ValueType, IndexType > Class Template Reference

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

```
#include <ginkgo/core/preconditioner/isai.hpp>
```

Public Member Functions

std::shared_ptr< const typename std::conditional< IsaiType==isai_type::spd, Comp, Csr >::type > get_approximate_inverse () const

Returns the approximate inverse of the given matrix (either a CSR matrix for IsaiType general, upper or lower or a composition of two CSR matrices for IsaiType spd).

Isai & operator= (const Isai &other)

Copy-assigns an ISAI preconditioner.

Isai & operator= (Isai &&other)

Move-assigns an ISAI preconditioner.

• Isai (const Isai &other)

Copy-constructs an ISAI preconditioner.

Isai (Isai &&other)

Move-constructs an ISAI preconditioner.

std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

43.113.1 Detailed Description

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

Using the preconditioner computes aiA*x, aiU*x, aiL*x or $aiC^T*aiC*x$ (depending on the type of the Isai) for a given vector x (may have multiple right hand sides). aiA, aiU and aiL are the approximate inverses for A, U and L respectively. aiC is an approximation to C, the exact Cholesky factor of B (This is commonly referred to as a Factorized Sparse Approximate Inverse, short FSPAI).

The sparsity pattern used for the approximate inverse of A, L and U is the same as the sparsity pattern of the respective matrix. For B, the sparsity pattern used for the approximate inverse is the same as the sparsity pattern of the lower triangular half of B.

Note that, except for the spd case, for a matrix A generally $ISAI(A)^T = ISAI(A^T)$.

For more details on the algorithm, see the paper Incomplete Sparse Approximate Inverses for Parallel Preconditioning, which is the basis for this work.

Note

GPU implementations can only handle the vector unit width width (warp size for CUDA) as number of elements per row in the sparse matrix. If there are more than width elements per row, the remaining elements will be ignored.

Template Parameters

IsaiType	determines if the ISAI is generated for a general square matrix, a lower triangular matrix, an upper triangular matrix or an spd matrix
ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.113.2 Constructor & Destructor Documentation

43.113.2.1 Isai() [1/2]

Copy-constructs an ISAI preconditioner.

Inherits the executor, shallow-copies the matrix and parameters.

43.113.2.2 Isai() [2/2]

Move-constructs an ISAI preconditioner.

Inherits the executor, moves the matrix and parameters. The moved-from object is empty (0x0 with nullptr matrix and default parameters)

43.113.3 Member Function Documentation

43.113.3.1 conj_transpose()

```
template<isai_type IsaiType, typename ValueType , typename IndexType > std::unique_ptr<LinOp> gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::conj_← transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.113.3.2 get_approximate_inverse()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
std::shared_ptr<const typename std::conditional<IsaiType == isai_type::spd, Comp, Csr>::type>
gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::get_approximate_inverse ( ) const
[inline]
```

Returns the approximate inverse of the given matrix (either a CSR matrix for IsaiType general, upper or lower or a composition of two CSR matrices for IsaiType spd).

Returns

the generated approximate inverse

References gko::as().

43.113.3.3 operator=() [1/2]

Copy-assigns an ISAI preconditioner.

Preserves the executor, shallow-copies the matrix and parameters. Creates a clone of the matrix if it is on the wrong executor.

43.113.3.4 operator=() [2/2]

Move-assigns an ISAI preconditioner.

Preserves the executor, moves the matrix and parameters. Creates a clone of the matrix if it is on the wrong executor. The moved-from object is empty (0x0 with nullptr matrix and default parameters)

43.113.3.5 transpose()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::transpose
( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/preconditioner/isai.hpp

43.114 gko::stop::Iteration Class Reference

The Iteration class is a stopping criterion which stops the iteration process after a preset number of iterations.

```
#include <ginkgo/core/stop/iteration.hpp>
```

43.114.1 Detailed Description

The Iteration class is a stopping criterion which stops the iteration process after a preset number of iterations.

Note

to use this stopping criterion, it is required to update the iteration count for the ::check() method.

The documentation for this class was generated from the following file:

· ginkgo/core/stop/iteration.hpp

43.115 gko::log::iteration_complete_data Struct Reference

Struct representing iteration complete related data.

```
#include <ginkgo/core/log/record.hpp>
```

43.115.1 Detailed Description

Struct representing iteration complete related data.

The documentation for this struct was generated from the following file:

· ginkgo/core/log/record.hpp

43.116 gko::solver::IterativeBase Class Reference

A LinOp implementing this interface stores a stopping criterion factory.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

- std::shared_ptr< const stop::CriterionFactory > get_stop_criterion_factory () const Gets the stopping criterion factory of the solver.
- virtual void set_stop_criterion_factory (std::shared_ptr< const stop::CriterionFactory > new_stop_factory)

 Sets the stopping criterion of the solver.

43.116.1 Detailed Description

A LinOp implementing this interface stores a stopping criterion factory.

43.116.2 Member Function Documentation

43.116.2.1 get_stop_criterion_factory()

```
std::shared\_ptr < const \ stop::CriterionFactory > \ gko::solver::IterativeBase::get\_stop\_criterion \leftarrow \_factory \ ( ) \ const \ [inline]
```

Gets the stopping criterion factory of the solver.

Returns

the stopping criterion factory

Referenced by gko::solver::EnableIterativeBase< Bicg< ValueType > >::operator=().

43.116.2.2 set_stop_criterion_factory()

Sets the stopping criterion of the solver.

Parameters

```
other the new stopping criterion factory
```

Referenced by gko::solver::EnableIterativeBase< Bicg< ValueType > >::set_stop_criterion_factory().

The documentation for this class was generated from the following file:

• ginkgo/core/solver/solver_base.hpp

43.117 gko::preconditioner::Jacobi< ValueType, IndexType > Class Template Reference

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

#include <ginkgo/core/preconditioner/jacobi.hpp>

Public Member Functions

• size_type get_num_blocks () const noexcept

Returns the number of blocks of the operator.

 $\bullet \ \ const \ block_interleaved_storage_scheme < index_type > \& \ get_storage_scheme \ () \ const \ noexcept \\$

Returns the storage scheme used for storing Jacobi blocks.

const value_type * get_blocks () const noexcept

Returns the pointer to the memory used for storing the block data.

const remove complex< value type > * get conditioning () const noexcept

Returns an array of 1-norm condition numbers of the blocks.

size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

void convert_to (matrix::Dense< value_type > *result) const override

Converts the implementer to an object of type result_type.

void move_to (matrix::Dense< value_type > *result) override

Converts the implementer to an object of type result_type by moving data from this object.

· void write (mat data &data) const override

Writes a matrix to a matrix_data structure.

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

• Jacobi & operator= (const Jacobi &other)

Copy-assigns a Jacobi preconditioner.

Jacobi & operator= (Jacobi &&other)

Move-assigns a Jacobi preconditioner.

Jacobi (const Jacobi &other)

Copy-constructs a Jacobi preconditioner.

Jacobi (Jacobi &&other)

Move-assigns a Jacobi preconditioner.

43.117.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::preconditioner::Jacobi< ValueType, IndexType >

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

The Jacobi class implements the inversion of the diagonal blocks using Gauss-Jordan elimination with column pivoting, and stores the inverse explicitly in a customized format.

If the diagonal blocks of the matrix are not explicitly set by the user, the implementation will try to automatically detect the blocks by first finding the natural blocks of the matrix, and then applying the supervariable agglomeration procedure on them. However, if problem-specific knowledge regarding the block diagonal structure is available, it is usually beneficial to explicitly pass the starting rows of the diagonal blocks, as the block detection is merely a heuristic and cannot perfectly detect the diagonal block structure. The current implementation supports blocks of up to 32 rows / columns.

The implementation also includes an improved, adaptive version of the block-Jacobi preconditioner, which can store some of the blocks in lower precision and thus improve the performance of preconditioner application by reducing the amount of memory transfers. This variant can be enabled by setting the Jacobi::Factory's storage optimization parameter. Refer to the documentation of the parameter for more details.

Template Parameters

ValueType	precision of matrix elements
IndexType	integral type used to store pointers to the start of each block

Note

The current implementation supports blocks of up to 32 rows / columns.

When using the adaptive variant, there may be a trade-off in terms of slightly longer preconditioner generation due to extra work required to detect the optimal precision of the blocks.

When the max_block_size is set to 1, specialized kernels are used, both for generation (inverting the diagonals) and application (diagonal scaling) to reduce the overhead involved in the usual (adaptive) block case.

43.117.2 Constructor & Destructor Documentation

43.117.2.1 Jacobi() [1/2]

Copy-constructs a Jacobi preconditioner.

Inherits executor, copies all data and parameters.

43.117.2.2 Jacobi() [2/2]

Move-assigns a Jacobi preconditioner.

Inherits executor, moves all data and parameters. The moved-from object will be empty (0x0 and default parameters).

43.117.3 Member Function Documentation

43.117.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Jacobi< ValueType, IndexType >::conj_transpose (
) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.117.3.2 convert_to()

Converts the implementer to an object of type result type.

Parameters

result	the object used to store the result of the conversion

Implements gko::ConvertibleTo< matrix::Dense< ValueType >>.

43.117.3.3 get_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::preconditioner::Jacobi< ValueType, IndexType >::get_blocks ( ) const
[inline], [noexcept]
```

Returns the pointer to the memory used for storing the block data.

Element (i, j) of block b is stored in position (get_block_pointers() [b] + i) * stride + j of the array.

Returns

the pointer to the memory used for storing the block data

References gko::array< ValueType >::get_const_data().

43.117.3.4 get_conditioning()

```
template<typename ValueType = default_precision, typename IndexType = int32> const remove_complex<value_type>* gko::preconditioner::Jacobi< ValueType, IndexType >::get_← conditioning ( ) const [inline], [noexcept]
```

Returns an array of 1-norm condition numbers of the blocks.

Returns

an array of 1-norm condition numbers of the blocks

Note

This value is valid only if adaptive precision variant is used, and implementations of the standard non-adaptive variant are allowed to omit the calculation of condition numbers.

References gko::array< ValueType >::get_const_data().

43.117.3.5 get_num_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::preconditioner::Jacobi < ValueType, IndexType >::get_num_blocks ( ) const [inline],
[noexcept]
```

Returns the number of blocks of the operator.

Returns

the number of blocks of the operator

43.117.3.6 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored_elements ( )
const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.117.3.7 get_storage_scheme()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const block_interleaved_storage_scheme<index_type>& gko::preconditioner::Jacobi< ValueType,
IndexType >::get_storage_scheme ( ) const [inline], [noexcept]
```

Returns the storage scheme used for storing Jacobi blocks.

Returns

the storage scheme used for storing Jacobi blocks

43.117.3.8 move to()

Converts the implementer to an object of type result_type by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

```
result the object used to emplace the result of the conversion
```

Note

ConvertibleTo::move_to can be implemented by simply calling ConvertibleTo::convert_to. However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implements gko::ConvertibleTo< matrix::Dense< ValueType > >.

43.117.3.9 operator=() [1/2]

Copy-assigns a Jacobi preconditioner.

Preserves executor, copies all data and parameters.

43.117.3.10 operator=() [2/2]

Move-assigns a Jacobi preconditioner.

Preserves executor, moves all data and parameters. The moved-from object will be empty (0x0 and default parameters).

43.117.3.11 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Jacobi< ValueType, IndexType >::transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.117.3.12 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following file:

· ginkgo/core/preconditioner/jacobi.hpp

43.118 gko::KernelNotFound Class Reference

KernelNotFound is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• KernelNotFound (const std::string &file, int line, const std::string &func)

Initializes a KernelNotFound error.

43.118.1 Detailed Description

KernelNotFound is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.

43.118.2 Constructor & Destructor Documentation

43.118.2.1 KernelNotFound()

Initializes a KernelNotFound error.

Parameters

	file	The name of the offending source file	
Ī	line	The source code line number where the error occurred	
func The name of the function where the error occurred		The name of the function where the error occurred	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.119 gko::log::linop_data Struct Reference

Struct representing LinOp related data.

```
#include <ginkgo/core/log/record.hpp>
```

43.119.1 Detailed Description

Struct representing LinOp related data.

The documentation for this struct was generated from the following file:

ginkgo/core/log/record.hpp

43.120 gko::log::linop factory data Struct Reference

Struct representing LinOp factory related data.

#include <ginkgo/core/log/record.hpp>

43.120.1 Detailed Description

Struct representing LinOp factory related data.

The documentation for this struct was generated from the following file:

ginkgo/core/log/record.hpp

43.121 gko::LinOpFactory Class Reference

A LinOpFactory represents a higher order mapping which transforms one linear operator into another.

#include <ginkgo/core/base/lin_op.hpp>

Additional Inherited Members

43.121.1 Detailed Description

A LinOpFactory represents a higher order mapping which transforms one linear operator into another.

In Ginkgo, every linear solver is viewed as a mapping. For example, given an s.p.d linear system Ax=b, the solution $x=A^{-1}b$ can be computed using the CG method. This algorithm can be represented in terms of linear operators and mappings between them as follows:

- A Cg::Factory is a higher order mapping which, given an input operator A, returns a new linear operator A^{-1} stored in "CG format"
- Storing the operator A^{-1} in "CG format" means that the data structure used to store the operator is just a simple pointer to the original matrix A. The application $x=A^{-1}b$ of such an operator can then be implemented by solving the linear system Ax=b using the CG method. This is achieved in code by having a special class for each of those "formats" (e.g. the "Cg" class defines such a format for the CG solver).

Another example of a LinOpFactory is a preconditioner. A preconditioner for a linear operator A is a linear operator M^{-1} , which approximates A^{-1} . In addition, it is stored in a way such that both the data of M^{-1} is cheap to compute from A, and the operation $x=M^{-1}b$ can be computed quickly. These operators are useful to accelerate the convergence of Krylov solvers. Thus, a preconditioner also fits into the LinOpFactory framework:

- The factory maps a linear operator A into a preconditioner M^{-1} which is stored in suitable format (e.g. as a product of two factors in case of ILU preconditioners).
- The resulting linear operator implements the application operation $x = M^{-1}b$ depending on the format the preconditioner is stored in (e.g. as two triangular solves in case of ILU)

43.121.1.1 Example: using CG in Ginkgo

```
{c++}
// Suppose A is a matrix, b a rhs vector, and x an initial guess
// Create a CG which runs for at most 1000 iterations, and stops after
// reducing the residual norm by 6 orders of magnitude
auto cg_factory = solver::Cg<>::build()
    .with_max_iters(1000)
    .with_rel_residual_goal(1e-6)
    .on(cuda);
// create a linear operator which represents the solver
auto cg = cg_factory->generate(A);
// solve the system
cg->apply(b, x);
```

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.122 gko::matrix::Csr< ValueType, IndexType >::load_balance Class Reference

load_balance is a strategy_type which uses the load balance algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

• load_balance ()

Creates a load_balance strategy.

load_balance (std::shared_ptr< const CudaExecutor > exec)

Creates a load_balance strategy with CUDA executor.

load_balance (std::shared_ptr< const HipExecutor > exec)

Creates a load_balance strategy with HIP executor.

load_balance (std::shared_ptr< const DpcppExecutor > exec)

Creates a load_balance strategy with specified parameters.

Creates a load_balance strategy with DPCPP executor.

- load_balance (int64_t nwarps, int warp_size=32, bool cuda_strategy=true, std::string strategy_name="none")
- void process (const array < index_type > &mtx_row_ptrs, array < index_type > *mtx_srow) override
 Computes srow according to row pointers.
- int64_t clac_size (const int64_t nnz) override

Computes the srow size according to the number of nonzeros.

• std::shared_ptr< strategy_type > copy () override

Copy a strategy.

43.122.1 Detailed Description

```
template < typename ValueType = default_precision, typename IndexType = int32 > class gko::matrix::Csr < ValueType, IndexType > ::load_balance
```

load_balance is a strategy_type which uses the load balance algorithm.

43.122.2 Constructor & Destructor Documentation

43.122.2.1 load_balance() [1/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance ( ) [inline]
```

Creates a load_balance strategy.

Warning

this is deprecated! Please rely on the new automatic strategy instantiation or use one of the other constructors.

43.122.2.2 load_balance() [2/5]

Creates a load_balance strategy with CUDA executor.

Parameters

```
exec the CUDA executor
```

43.122.2.3 load_balance() [3/5]

Creates a load_balance strategy with HIP executor.

Parameters

```
exec the HIP executor
```

43.122.2.4 load_balance() [4/5]

template<typename ValueType = default_precision, typename IndexType = int32>

Creates a load_balance strategy with DPCPP executor.

Parameters

```
exec the DPCPP executor
```

Note

TODO: porting - we hardcode the subgroup size is 32

43.122.2.5 load_balance() [5/5]

Creates a load_balance strategy with specified parameters.

Parameters

nwarps	the number of warps in the executor	
warp_size	the warp size of the executor	
cuda_strategy	whether the cuda_strategy needs to be used.	

Note

The warp_size must be the size of full warp. When using this constructor, set_strategy needs to be called with correct parameters which is replaced during the conversion.

43.122.3 Member Function Documentation

43.122.3.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

nnz the number of nonze	ros
-------------------------	-----

Returns

the size of srow

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

References gko::ceildiv(), and gko::min().

43.122.3.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::load_balance::copy (
) [inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.122.3.3 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix	
mtx_srow	the srow of the matrix	

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

 $References\ gko::ceildiv(),\ gko::array<\ Value\ Type>::get_const_data(),\ gko::array<\ Value\ Type>::get_data(),\ gko::array<\ Value\ Type>::get_num_elems().$

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/csr.hpp

43.123 gko::log::Loggable Class Reference

Loggable class is an interface which should be implemented by classes wanting to support logging.

```
#include <ginkgo/core/log/logger.hpp>
```

Public Member Functions

- virtual void add logger (std::shared ptr< const Logger > logger)=0
 - Adds a new logger to the list of subscribed loggers.
- virtual void remove_logger (const Logger *logger)=0

Removes a logger from the list of subscribed loggers.

- virtual const std::vector< std::shared_ptr< const Logger > > & get_loggers () const =0
 - Returns the vector containing all loggers registered at this object.
- virtual void clear_loggers ()=0

Remove all loggers registered at this object.

43.123.1 Detailed Description

Loggable class is an interface which should be implemented by classes wanting to support logging.

For most cases, one can rely on the EnableLogging mixin which provides a default implementation of this interface.

43.123.2 Member Function Documentation

43.123.2.1 add_logger()

Adds a new logger to the list of subscribed loggers.

Parameters

```
logger to add
```

Implemented in gko::Executor.

43.123.2.2 get_loggers()

```
virtual const std::vector<std::shared_ptr<const Logger> >& gko::log::Loggable::get_loggers (
) const [pure virtual]
```

Returns the vector containing all loggers registered at this object.

Returns

the vector containing all registered loggers.

43.123.2.3 remove_logger()

Removes a logger from the list of subscribed loggers.

Parameters

ger the logger to remove

Note

The comparison is done using the logger's object unique identity. Thus, two loggers constructed in the same way are not considered equal.

Implemented in gko::Executor.

The documentation for this class was generated from the following file:

· ginkgo/core/log/logger.hpp

43.124 gko::log::Record::logged_data Struct Reference

Struct storing the actually logged data.

```
#include <ginkgo/core/log/record.hpp>
```

43.124.1 Detailed Description

Struct storing the actually logged data.

The documentation for this struct was generated from the following file:

• ginkgo/core/log/record.hpp

43.125 gko::solver::LowerTrs< ValueType, IndexType > Class Template Reference

LowerTrs is the triangular solver which solves the system L x = b, when L is a lower triangular matrix.

```
#include <ginkgo/core/solver/triangular.hpp>
```

Public Member Functions

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

• LowerTrs (const LowerTrs &)

Copy-assigns a triangular solver.

LowerTrs (LowerTrs &&)

Move-assigns a triangular solver.

LowerTrs & operator= (const LowerTrs &)

Copy-constructs a triangular solver.

LowerTrs & operator= (LowerTrs &&)

Move-constructs a triangular solver.

43.125.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::solver::LowerTrs< ValueType, IndexType >
```

LowerTrs is the triangular solver which solves the system L x = b, when L is a lower triangular matrix.

It works best when passing in a matrix in CSR format. If the matrix is not in CSR, then the generate step converts it into a CSR matrix. The generation fails if the matrix is not convertible to CSR.

Note

As the constructor uses the copy and convert functionality, it is not possible to create a empty solver or a solver with a matrix in any other format other than CSR, if none of the executor modules are being compiled with.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indices

43.125.2 Constructor & Destructor Documentation

43.125.2.1 LowerTrs() [1/2]

Copy-assigns a triangular solver.

Preserves the executor, shallow-copies the system matrix. If the executors mismatch, clones system matrix onto this executor. Solver analysis information will be regenerated.

43.125.2.2 LowerTrs() [2/2]

Move-assigns a triangular solver.

Preserves the executor, moves the system matrix. If the executors mismatch, clones system matrix onto this executor and regenerates solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

43.125.3 Member Function Documentation

43.125.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::LowerTrs< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.125.3.2 operator=() [1/2]

Copy-constructs a triangular solver.

Preserves the executor, shallow-copies the system matrix. Solver analysis information will be regenerated.

43.125.3.3 operator=() [2/2]

Move-constructs a triangular solver.

Preserves the executor, moves the system matrix and solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

43.125.3.4 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::LowerTrs< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/triangular.hpp

43.126 gko::experimental::factorization::Lu< ValueType, IndexType > Class Template Reference

Computes an LU factorization of a sparse matrix.

```
#include <ginkgo/core/factorization/lu.hpp>
```

Public Member Functions

std::unique_ptr< factorization_type > generate (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

static parameters_type build ()
 Creates a new parameter_type to set up the factory.

43.126.1 Detailed Description

```
template<typename ValueType, typename IndexType> class gko::experimental::factorization::Lu< ValueType, IndexType>
```

Computes an LU factorization of a sparse matrix.

This LinOpFactory returns a Factorization storing the L and U factors for the provided system matrix in matrix::Csr format. If no symbolic factorization is provided, it will be computed first.

Template Parameters

ValueType	the type used to store values of the system matrix	
IndexType Generated by Dex	the type used to store sparsity pattern indices of the system matrix	

43.126.2 Member Function Documentation

43.126.2.1 generate()

Note

This function overrides the default LinOpFactory::generate to return a Factorization instead of a generic LinOp, which would need to be cast to Factorization again to access its factors. It is only necessary because smart pointers aren't covariant.

The documentation for this class was generated from the following file:

· ginkgo/core/factorization/lu.hpp

43.127 gko::machine_topology Class Reference

The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.

#include <ginkgo/core/base/machine_topology.hpp>

Public Member Functions

void bind_to_cores (const std::vector< int > &ids, const bool singlify=true) const

Bind the calling process to the CPU cores associated with the ids.

void bind_to_core (const int &id) const

Bind to a single core.

void bind_to_pus (const std::vector< int > &ids, const bool singlify=true) const

Bind the calling process to PUs associated with the ids.

• void bind_to_pu (const int &id) const

Bind to a Processing unit (PU)

const normal_obj_info * get_pu (size_type id) const

Get the object of type PU associated with the id.

const normal_obj_info * get_core (size_type id) const

Get the object of type core associated with the id.

const io_obj_info * get_pci_device (size_type id) const

Get the object of type pci device associated with the id.

const io_obj_info * get_pci_device (const std::string &pci_bus_id) const

Get the object of type pci device associated with the PCI bus id.

• size_type get_num_pus () const

Get the number of PU objects stored in this Topology tree.

• size_type get_num_cores () const

Get the number of core objects stored in this Topology tree.

• size_type get_num_pci_devices () const

Get the number of PCI device objects stored in this Topology tree.

• size_type get_num_numas () const

Get the number of NUMA objects stored in this Topology tree.

Static Public Member Functions

static machine_topology * get_instance ()
 Returns an instance of the machine topology object.

43.127.1 Detailed Description

The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.

Various infomation of the machine are gathered with the help of the Hardware Locality library (hwloc).

This class also provides functionalities to bind objects in the topology to the execution objects. Binding can enhance performance by allowing data to be closer to the executing object.

See the hwloc documentation (https://www.open-mpi.org/projects/hwloc/doc/) for more detailed information on topology detection and binding interfaces.

Note

A global object of machine_topology type is created in a thread safe manner and only destroyed at the end of the program. This means that any subsequent queries will be from the same global object and hence use an extra atomic read.

43.127.2 Member Function Documentation

43.127.2.1 bind_to_core()

Bind to a single core.

Parameters

```
ids The ids of the core to be bound to the calling process.
```

References bind_to_cores(), and get_instance().

43.127.2.2 bind_to_cores()

Bind the calling process to the CPU cores associated with the ids.

Parameters

ids	The ids of cores to be bound.
singlify	The ids of PUs are singlified to prevent possibly expensive migrations by the OS. This means that the
	binding is performed for only one of the ids in the set of ids passed in. See hwloc doc for
	singlify

Referenced by bind_to_core().

43.127.2.3 bind to pu()

Bind to a Processing unit (PU)

Parameters

References bind_to_pus(), and get_instance().

43.127.2.4 bind_to_pus()

Bind the calling process to PUs associated with the ids.

Parameters

ids	The ids of PUs to be bound.
singlify	The ids of PUs are singlified to prevent possibly expensive migrations by the OS. This means that the
	binding is performed for only one of the ids in the set of ids passed in. See hwloc doc for
	singlify

Referenced by bind_to_pu().

43.127.2.5 get_core()

Get the object of type core associated with the id.

Parameters

```
id The id of the core
```

Returns

the core object struct.

43.127.2.6 get_instance()

```
static machine_topology* gko::machine_topology::get_instance ( ) [inline], [static]
```

Returns an instance of the machine_topology object.

Returns

the machine_topology instance

Referenced by bind_to_core(), and bind_to_pu().

43.127.2.7 get_num_cores()

```
size_type gko::machine_topology::get_num_cores ( ) const [inline]
```

Get the number of core objects stored in this Topology tree.

Returns

the number of cores.

43.127.2.8 get_num_numas()

```
size_type gko::machine_topology::get_num_numas ( ) const [inline]
```

Get the number of NUMA objects stored in this Topology tree.

Returns

the number of NUMA objects.

43.127.2.9 get_num_pci_devices()

```
size_type gko::machine_topology::get_num_pci_devices ( ) const [inline]
```

Get the number of PCI device objects stored in this Topology tree.

Returns

the number of PCI devices.

43.127.2.10 get_num_pus()

```
size_type gko::machine_topology::get_num_pus ( ) const [inline]
```

Get the number of PU objects stored in this Topology tree.

Returns

the number of PUs.

43.127.2.11 get_pci_device() [1/2]

Get the object of type pci device associated with the PCI bus id.

Parameters

pci_bus⊷	The PCI bus id of the pci device
_id	

Returns

the PCI object struct.

43.127.2.12 get_pci_device() [2/2]

Get the object of type pci device associated with the id.

Parameters

id The id of the pci device

Returns

the PCI object struct.

43.127.2.13 get_pu()

Get the object of type PU associated with the id.

Parameters

id The id of the PU

Returns

the PU object struct.

The documentation for this class was generated from the following file:

• ginkgo/core/base/machine_topology.hpp

43.128 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference

The Matrix class defines a (MPI-)distributed matrix.

```
#include <ginkgo/core/distributed/matrix.hpp>
```

Public Member Functions

void read_distributed (const device_matrix_data< value_type, global_index_type > &data, ptr_param
 const Partition< local_index_type, global_index_type >> partition)

Reads a square matrix from the device_matrix_data structure and a global partition.

• void read_distributed (const matrix_data< value_type, global_index_type > &data, ptr_param< const Partition< local_index_type, global_index_type >> partition)

Reads a square matrix from the matrix_data structure and a global partition.

void read_distributed (const device_matrix_data< value_type, global_index_type > &data, ptr_param<
 const Partition< local_index_type, global_index_type >> row_partition, ptr_param<
 const Partition< local index_type, global_index_type, global_index_type >> col_partition)

Reads a matrix from the device_matrix_data structure, a global row partition, and a global column partition.

 void read_distributed (const matrix_data< value_type, global_index_type > &data, ptr_param< const Partition< local_index_type, global_index_type >> row_partition, ptr_param< const Partition< local_← index_type, global_index_type >> col_partition)

Reads a matrix from the matrix_data structure, a global row partition, and a global column partition.

std::shared_ptr< const LinOp > get_local_matrix () const

Get read access to the stored local matrix.

std::shared_ptr< const LinOp > get_non_local_matrix () const

Get read access to the stored non-local matrix.

Matrix (const Matrix & other)

Copy constructs a Matrix.

• Matrix (Matrix &&other) noexcept

Move constructs a Matrix.

Matrix & operator= (const Matrix & other)

Copy assigns a Matrix.

Matrix & operator= (Matrix &&other)

Move assigns a Matrix.

43.128.1 Detailed Description

template < typename ValueType = default_precision, typename LocalIndexType = int32, typename GlobalIndexType = int64 > class gko::experimental::distributed::Matrix < ValueType, LocalIndexType, GlobalIndexType >

The Matrix class defines a (MPI-)distributed matrix.

The matrix is stored in a row-wise distributed format. Each process owns a specific set of rows, where the assignment of rows is defined by a row Partition. The following depicts the distribution of global rows according to their assigned part-id (which will usually be the owning process id):

The local rows are further split into two matrices on each process. One matrix, called local, contains only entries from columns that are also owned by the process, while the other one, called non_local, contains entries from columns that are not owned by the process. The non-local matrix is stored in a compressed format, where empty columns are discarded and the remaining columns are renumbered. This splitting is depicted in the following:

```
Non-Local
Part-Id Global
                                               Local
          | .. 1 ! 2 .. ! .. ..
| 3 4 ! .. .. ! .. ..
0
0
                                                     | 3 4 | | ...
                                  --|
          1 .. 5 ! 6 ..
                             ! 7 ..
                                            ----> | 6
                                                         .. |
                .. ! .. 8 ! ..
                                        9 | ----> | 8 .. | | .. .. 9 |
1
          | .. . ! . . 10 ! 11 12 |
| 13 . . ! . . . ! 14 . . |
2
                                                | 11 12 | | .. 10 | | 14 .. | | 13 .. |
```

This uses the same ownership of the columns as for the rows. Additionally, the ownership of the columns may be explicitly defined with an second column partition. If that is not provided, the same row partition will be used for the columns. Using a column partition also allows to create non-square matrices, like the one below:

```
Part-Id Global
                             Local
                                      Non-Local
P_R/P_C
         2 2 0 1
             1 2 .. |
                                        | 1 .. |
0
                               | 2 |
        3 4 .. ..
0
                                    1
                     ١..
        | .. 5 6 ..
| .. .. 8
                      | ----> | .. |
                                         | 6 5 |
2
        | .. .. 10 |
                                | .. .. | | 10 |
```

```
2 | 13 .. .. | | 13 .. | | .. |
```

Here P_R denotes the row partition and P_C denotes the column partition.

The Matrix should be filled using the read_distributed method, e.g.

```
auto part = Partition<...>::build_from_mapping(...);
auto mat = Matrix<...>::create(exec, comm);
mat->read_distributed(matrix_data, part);
```

or if different partitions for the rows and columns are used:

```
auto row_part = Partition<...>::build_from_mapping(...);
auto col_part = Partition<...>::build_from_mapping(...);
auto mat = Matrix<...>::create(exec, comm);
mat->read_distributed(matrix_data, row_part, col_part);
```

This will set the dimensions of the global and local matrices automatically by deducing the sizes from the partitions.

By default the Matrix type uses Csr for both stored matrices. It is possible to explicitly change the datatype for the stored matrices, with the constraint that the new type should implement the LinOp and ReadableFromMatrixData interface. The type can be set by:

```
auto mat = Matrix<ValueType, LocalIndexType[, ...]>::create(
   exec, comm,
   Ell<ValueType, LocalIndexType>::create(exec).get(),
   Coo<ValueType, LocalIndexType>::create(exec).get());
```

Alternatively, the helper function with_matrix_type can be used:

```
auto mat = Matrix<ValueType, LocalIndexType>::create(
  exec, comm,
  with_matrix_type<Ell>(),
  with_matrix_type<Coo>());
```

See also

with_matrix_type

The Matrix LinOp supports the following operations:

Template Parameters

ValueType	The underlying value type.	
LocalIndexType	The index type used by the local matrices.	
GlobalIndexType	The type for global indices.	

43.128.2 Constructor & Destructor Documentation

43.128.2.1 Matrix() [1/2]

Copy constructs a Matrix.

Parameters

other Matrix to copy from.

43.128.2.2 Matrix() [2/2]

Move constructs a Matrix.

Parameters

other Matrix to move from.

43.128.3 Member Function Documentation

43.128.3.1 get_local_matrix()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
std::shared_ptr<const LinOp> gko::experimental::distributed::Matrix< ValueType, LocalIndex
Type, GlobalIndexType >::get_local_matrix ( ) const [inline]
```

Get read access to the stored local matrix.

Returns

Shared pointer to the stored local matrix

43.128.3.2 get_non_local_matrix()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
std::shared_ptr<const LinOp> gko::experimental::distributed::Matrix< ValueType, LocalIndex
Type, GlobalIndexType >::get_non_local_matrix ( ) const [inline]
```

Get read access to the stored non-local matrix.

Returns

Shared pointer to the stored non-local matrix

43.128.3.3 operator=() [1/2]

Copy assigns a Matrix.

Parameters

other | Matrix to copy from, has to have a communicator of the same size as this.

Returns

this.

43.128.3.4 operator=() [2/2]

Move assigns a Matrix.

Parameters

other | Matrix to move from, has to have a communicator of the same size as this.

Returns

this.

43.128.3.5 read distributed() [1/4]

Reads a square matrix from the device_matrix_data structure and a global partition.

The global size of the final matrix is inferred from the size of the partition. Both the number of rows and columns of the device_matrix_data are ignored.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

data	The device_matrix_data structure.	
partition	The global row and column partition.	

43.128.3.6 read_distributed() [2/4]

Reads a matrix from the device_matrix_data structure, a global row partition, and a global column partition.

The global size of the final matrix is inferred from the size of the row partition and the size of the column partition. Both the number of rows and columns of the device_matrix_data are ignored.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

data	The device_matrix_data structure.	
row_partition	The global row partition.	
col_partition	The global col partition.	

43.128.3.7 read_distributed() [3/4]

Reads a square matrix from the matrix_data structure and a global partition.

See also

read distributed

Note

For efficiency it is advised to use the device_matrix_data overload.

43.128.3.8 read_distributed() [4/4]

Reads a matrix from the matrix_data structure, a global row partition, and a global column partition.

See also

read distributed

Note

For efficiency it is advised to use the device_matrix_data overload.

The documentation for this class was generated from the following file:

• ginkgo/core/distributed/matrix.hpp

43.129 gko::matrix_assembly_data< ValueType, IndexType > Class Template Reference

This structure is used as an intermediate type to assemble a sparse matrix.

```
#include <ginkgo/core/base/matrix_assembly_data.hpp>
```

Public Member Functions

- void add_value (index_type row, index_type col, value_type val)
 Sets the matrix value at (row, col).
- void set_value (index_type row, index_type col, value_type val)

Sets the matrix value at (row, col).

value_type get_value (index_type row, index_type col)

Gets the matrix value at (row, col).

bool contains (index_type row, index_type col)

Returns true iff the matrix contains an entry at (row, col).

- dim< 2 > get size () const noexcept
- size_type get_num_stored_elements () const noexcept
- matrix_data< ValueType, IndexType > get_ordered_data () const

43.129.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix_assembly_data< ValueType, IndexType >

This structure is used as an intermediate type to assemble a sparse matrix.

The matrix is stored as a set of nonzero elements, where each element is a triplet of the form (row_index, column
_index, value).

New values can be added by using the matrix_assembly_data::add_value or matrix_assembly_data::set_value

Template Parameters

ValueType	type of matrix values stored in the structure	
IndexType	type of matrix indexes stored in the structure	

43.129.2 Member Function Documentation

43.129.2.1 add_value()

Sets the matrix value at (row, col).

If there is an existing value, it will be set to the sum of the existing and new value, otherwise the value will be inserted.

Parameters

row	the row where the value should be added
col	the column where the value should be added
val	the value to be added to (row, col)

43.129.2.2 contains()

```
template<typename ValueType = default_precision, typename IndexType = int32>
bool gko::matrix_assembly_data< ValueType, IndexType >::contains (
```

```
index_type row,
index_type col ) [inline]
```

Returns true iff the matrix contains an entry at (row, col).

Parameters

row	the row index	
col	the column index	

Returns

true if the value at (row, col) exists, false otherwise

43.129.2.3 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix_assembly_data< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns

the number of non-zeros in the (partially) assembled matrix

43.129.2.4 get_ordered_data()

```
template<typename ValueType = default_precision, typename IndexType = int32>
matrix_data<ValueType, IndexType> gko::matrix_assembly_data< ValueType, IndexType >::get_
ordered_data ( ) const [inline]
```

Returns

a matrix_data instance containing the assembled non-zeros in row-major order to be used by all matrix formats.

Referenced by gko::ReadableFromMatrixData< ValueType, int32 >::read().

43.129.2.5 get_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
dim<2> gko::matrix_assembly_data< ValueType, IndexType >::get_size ( ) const [inline], [noexcept]
```

Returns

the dimensions of the matrix being assembled

43.129.2.6 get_value()

Gets the matrix value at (row, col).

Parameters

row	the row index	
col	the column index	

Returns

the value at (row, col) or 0 if it doesn't exist.

43.129.2.7 set_value()

Sets the matrix value at (row, col).

If there is an existing value, it will be overwritten by the new value.

Parameters

row	the row index	
col	the column index	
val	the value to be written to (row, col)	

The documentation for this class was generated from the following file:

ginkgo/core/base/matrix_assembly_data.hpp

43.130 gko::matrix_data< ValueType, IndexType > Struct Template Reference

This structure is used as an intermediate data type to store a sparse matrix.

#include <ginkgo/core/base/matrix_data.hpp>

Public Member Functions

matrix_data (dim< 2 > size_=dim< 2 >{}, ValueType value=zero< ValueType >())

Initializes a matrix filled with the specified value.

 $\bullet \ \ \text{template}{<} \text{typename RandomDistribution , typename RandomEngine} >$

matrix_data (dim< 2 > size_, RandomDistribution &&dist, RandomEngine &&engine)

Initializes a matrix with random values from the specified distribution.

matrix_data (std::initializer_list< std::initializer_list< ValueType >> values)

List-initializes the structure from a matrix of values.

matrix_data (dim< 2 > size_, std::initializer_list< detail::input_triple< ValueType, IndexType >> nonzeros
 —)

Initializes the structure from a list of nonzeros.

matrix_data (dim< 2 > size_, const matrix_data &block)

Initializes a matrix out of a matrix block via duplication.

template<typename Accessor >

matrix_data (const range< Accessor > &data)

Initializes a matrix from a range.

void ensure_row_major_order ()

Sorts the nonzero vector so the values follow row-major order.

void remove zeros ()

Remove entries with value zero from the matrix data.

void sum_duplicates ()

Sum up all values that refer to the same matrix entry.

Static Public Member Functions

static matrix_data diag (dim< 2 > size_, ValueType value)

Initializes a diagonal matrix.

static matrix_data diag (dim< 2 > size_, std::initializer_list< ValueType > nonzeros_)

Initializes a diagonal matrix using a list of diagonal elements.

static matrix_data diag (dim< 2 > size_, const matrix_data &block)

Initializes a block-diagonal matrix.

template<typename ForwardIterator >

static matrix_data diag (ForwardIterator begin, ForwardIterator end)

Initializes a block-diagonal matrix from a list of diagonal blocks.

static matrix_data diag (std::initializer_list< matrix_data > blocks)

Initializes a block-diagonal matrix from a list of diagonal blocks.

 $\bullet \ \ \text{template}{<} \text{typename RandomDistribution , typename RandomEngine} >$

static matrix_data cond (size_type size, remove_complex< ValueType > condition_number, Random← Distribution &&dist, RandomEngine &&engine, size_type num_reflectors)

Initializes a random dense matrix with a specific condition number.

template<typename RandomDistribution, typename RandomEngine >
 static matrix_data cond (size_type size, remove_complex< ValueType > condition_number, Random
 Distribution &&dist, RandomEngine &&engine)

Initializes a random dense matrix with a specific condition number.

Public Attributes

dim< 2 > size

Size of the matrix.

std::vector< nonzero_type > nonzeros

A vector of tuples storing the non-zeros of the matrix.

43.130.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> struct gko::matrix_data< ValueType, IndexType >

This structure is used as an intermediate data type to store a sparse matrix.

The matrix is stored as a sequence of nonzero elements, where each element is a triple of the form (row_index, column_index, value).

Note

All Ginkgo functions returning such a structure will return the nonzeros sorted in row-major order.

All Ginkgo functions that take this structure as input expect that the nonzeros are sorted in row-major order and that the index pair (row_index, column_index) of each nonzero is unique.

This structure is not optimized for usual access patterns and it can only exist on the CPU. Thus, it should only be used for utility functions which do not have to be optimized for performance.

Template Parameters

ValueType	type of matrix values stored in the structure	
IndexType	type of matrix indexes stored in the structure	

43.130.2 Constructor & Destructor Documentation

43.130.2.1 matrix_data() [1/6]

Initializes a matrix filled with the specified value.

Parameters

size⊷ –	dimensions of the matrix
value	value used to fill the elements of the matrix

```
166
                                                              {}, ValueType value = zero<ValueType>())
167
                  : size{size_}
168
169
                 if (is_zero(value)) {
170
                       return;
172
                  nonzeros.reserve(size[0] * size[1]);
                 for (size_type row = 0; row < size[0]; ++row) {
   for (size_type col = 0; col < size[1]; ++col) {
      nonzeros.emplace_back(row, col, value);
}</pre>
173
174
175
176
```

```
177 }
178 }
```

43.130.2.2 matrix_data() [2/6]

Initializes a matrix with random values from the specified distribution.

Template Parameters

RandomDistribution	random distribution type
RandomEngine	random engine type

Parameters

size⊷	dimensions of the matrix
_	
dist	random distribution of the elements of the matrix
engine	random engine used to generate random values

43.130.2.3 matrix_data() [3/6]

List-initializes the structure from a matrix of values.

Parameters

```
values a 2D braced-init-list of matrix values.
```

43.130.2.4 matrix_data() [4/6]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix_data< ValueType, IndexType >::matrix_data (
```

Initializes the structure from a list of nonzeros.

Parameters

size_	dimensions of the matrix
nonzeros⇔	list of nonzero elements
_	

43.130.2.5 matrix_data() [5/6]

Initializes a matrix out of a matrix block via duplication.

Parameters

size	size of the block-matrix (in blocks)
diag_block	matrix block used to fill the complete matrix

References gko::matrix_data< ValueType, IndexType >::size.

43.130.2.6 matrix_data() [6/6]

Initializes a matrix from a range.

Template Parameters

Accessor	accessor type of the input range
----------	----------------------------------

Parameters

data	range used to initialize the matrix
------	-------------------------------------

References gko::range < Accessor >::length().

43.130.3 Member Function Documentation

43.130.3.1 cond() [1/2]

Initializes a random dense matrix with a specific condition number.

The matrix is generated by applying a series of random Hausholder reflectors to a diagonal matrix with diagonal entries uniformly distributed between sqrt (condition_number) and 1/sqrt (condition_number).

This version of the function applies size - 1 reflectors to each side of the diagonal matrix.

Template Parameters

RandomDistribution	the type of the random distribution
RandomEngine	the type of the random engine

Parameters

size	number of rows and columns of the matrix
condition_number	condition number of the matrix
dist	random distribution used to generate reflectors
engine	random engine used to generate reflectors

Returns

the dense matrix with the specified condition number

References gko::matrix_data < ValueType, IndexType >::cond(), and gko::matrix_data < ValueType, IndexType > ∴ ::size.

43.130.3.2 cond() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename RandomDistribution , typename RandomEngine >
```

Initializes a random dense matrix with a specific condition number.

The matrix is generated by applying a series of random Hausholder reflectors to a diagonal matrix with diagonal entries uniformly distributed between sqrt (condition_number) and 1/sqrt (condition_number).

Template Parameters

RandomDistribution	the type of the random distribution
RandomEngine	the type of the random engine

Parameters

size	number of rows and columns of the matrix
condition_number	condition number of the matrix
dist	random distribution used to generate reflectors
engine	random engine used to generate reflectors
num_reflectors	number of reflectors to apply from each side

Returns

the dense matrix with the specified condition number

References gko::matrix_data< ValueType, IndexType >::size.

Referenced by gko::matrix_data < ValueType, IndexType >::cond().

43.130.3.3 diag() [1/5]

Initializes a block-diagonal matrix.

Parameters

size_	the size of the matrix
diag_block	matrix used to fill diagonal blocks

Returns

the block-diagonal matrix

References gko::matrix_data< ValueType, IndexType >::nonzeros, and gko::matrix_data< ValueType, IndexType >::size.

43.130.3.4 diag() [2/5]

Initializes a diagonal matrix using a list of diagonal elements.

Parameters

size_	dimensions of the matrix
nonzeros⊷	list of diagonal elements

Returns

the diagonal matrix

References gko::matrix_data< ValueType, IndexType >::nonzeros.

43.130.3.5 diag() [3/5]

Initializes a diagonal matrix.

Parameters

size⊷	dimensions of the matrix
value	value used to fill the elements of the matrix

Returns

the diagonal matrix

References gko::is_nonzero(), and gko::matrix_data< ValueType, IndexType >::nonzeros.

Referenced by gko::matrix_data< ValueType, IndexType >::diag().

43.130.3.6 diag() [4/5]

Initializes a block-diagonal matrix from a list of diagonal blocks.

Template Parameters

ForwardIterator	type of list iterator
-----------------	-----------------------

Parameters

begin	the first iterator of the list
end	the last iterator of the list

Returns

the block-diagonal matrix with diagonal blocks set to the blocks between begin (inclusive) and end (exclusive)

References gko::matrix_data< ValueType, IndexType >::nonzeros.

43.130.3.7 diag() [5/5]

Initializes a block-diagonal matrix from a list of diagonal blocks.

Parameters

blocks	a list of blocks to initialize from
--------	-------------------------------------

Returns

the block-diagonal matrix with diagonal blocks set to the blocks passed in blocks

References gko::matrix_data< ValueType, IndexType >::diag().

43.130.3.8 sum_duplicates()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix_data< ValueType, IndexType >::sum_duplicates () [inline]
```

Sum up all values that refer to the same matrix entry.

The result is sorted in row-major order.

References gko::matrix_data< ValueType, IndexType >::ensure_row_major_order(), and gko::matrix_data< ValueType, IndexType >::nonzeros.

43.130.4 Member Data Documentation

43.130.4.1 nonzeros

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::vector<nonzero_type> gko::matrix_data< ValueType, IndexType >::nonzeros
```

A vector of tuples storing the non-zeros of the matrix.

The first two elements of the tuple are the row index and the column index of a matrix element, and its third element is the value at that position.

Referenced by gko::matrix_data < ValueType, IndexType >::diag(), gko::matrix_data < ValueType, IndexType >::ensure_row_major_order(), gko::matrix_data < ValueType, IndexType >::remove_zeros(), and gko::matrix_data < ValueType, IndexType, IndexType >::sum_duplicates().

The documentation for this struct was generated from the following file:

ginkgo/core/base/matrix_data.hpp

43.131 gko::matrix_data_entry< ValueType, IndexType > Struct Template Reference

Type used to store nonzeros.

#include <ginkgo/core/base/matrix_data.hpp>

43.131.1 Detailed Description

```
template<typename ValueType, typename IndexType>
struct gko::matrix_data_entry< ValueType, IndexType>
```

Type used to store nonzeros.

The documentation for this struct was generated from the following file:

• ginkgo/core/base/matrix_data.hpp

43.132 gko::matrix::Csr< ValueType, IndexType >::merge_path Class Reference

merge_path is a strategy_type which uses the merge_path algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- merge_path ()
 - Creates a merge_path strategy.
- void process (const array< index_type > &mtx_row_ptrs, array< index_type > *mtx_srow) override
 Computes srow according to row pointers.
- int64_t clac_size (const int64_t nnz) override

Computes the srow size according to the number of nonzeros.

 std::shared_ptr< strategy_type > copy () override Copy a strategy.

43.132.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Csr< ValueType, IndexType >::merge_path
```

merge_path is a strategy_type which uses the merge_path algorithm.

merge_path is according to Merrill and Garland: Merge-Based Parallel Sparse Matrix-Vector Multiplication

43.132.2 Member Function Documentation

43.132.2.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

nnz	the number of nonzeros
-----	------------------------

Returns

the size of srow

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.132.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::merge_path::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.132.2.3 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix
mtx_srow	the srow of the matrix

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

The documentation for this class was generated from the following file:

ginkgo/core/matrix/csr.hpp

43.133 gko::MetisError Class Reference

MetisError is thrown when METIS routine throws an error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

MetisError (const std::string &file, int line, const std::string &func, const std::string &error)
 Initializes a METIS error.

43.133.1 Detailed Description

MetisError is thrown when METIS routine throws an error code.

43.133.2 Constructor & Destructor Documentation

43.133.2.1 MetisError()

Initializes a METIS error.

Parameters

file	The name of the offending source file	
line	The source code line number where the error occurred	
func	The name of the METIS routine that failed	
error	The resulting METIS error name	

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.134 gko::matrix::Hybrid< ValueType, IndexType >::minimal storage limit Class Reference

minimal_storage_limit is a strategy_type which decides the number of stored elements per row of the ell part.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- minimal_storage_limit ()
 - Creates a minimal_storage_limit strategy.
- size_type compute_ell_num_stored_elements_per_row (array< size_type > *row_nnz) const override
 Computes the number of stored elements per row of the ell part.
- auto get_percentage () const

Get the percent setting.

43.134.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit
```

minimal storage limit is a strategy type which decides the number of stored elements per row of the ell part.

It is determined by the size of ValueType and IndexType, the storage is the minimum among all partition.

43.134.2 Member Function Documentation

43.134.2.1 compute_ell_num_stored_elements_per_row()

Computes the number of stored elements per row of the ell part.

Parameters

rout pp7	the number of nonzeros of each row
TOW_TITE	the number of nonzeros of each row

Returns

the number of stored elements per row of the ell part

Implements gko::matrix::Hybrid < ValueType, IndexType >::strategy_type.

References gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements ← _per_row().

43.134.2.2 get_percentage()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::get_percentage ( )
const [inline]
```

Get the percent setting.

@retrun percent

References gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::get_percentage().

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/hybrid.hpp

43.135 gko::MpiError Class Reference

MpiError is thrown when a MPI routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

MpiError (const std::string &file, int line, const std::string &func, int64 error_code)
 Initializes a MPI error.

43.135.1 Detailed Description

MpiError is thrown when a MPI routine throws a non-zero error code.

43.135.2 Constructor & Destructor Documentation

43.135.2.1 MpiError()

Initializes a MPI error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the MPI routine that failed
error_code	The resulting MPI error code

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.136 gko::solver::Multigrid Class Reference

Multigrid methods have a hierarchy of many levels, whose corase level is a subset of the fine level, of the problem.

#include <ginkgo/core/solver/multigrid.hpp>

Public Member Functions

- bool apply_uses_initial_guess () const override
 - Return true as iterative solvers use the data in x as an initial guess or false if multigrid always set the input as zero.
- std::vector < std::shared_ptr < const gko::multigrid::MultigridLevel >> get_mg_level_list () const
 Gets the list of MultigridLevel operators.
- std::vector< std::shared_ptr< const LinOp >> get_pre_smoother_list () const

Gets the list of pre-smoother operators.

- std::vector< std::shared_ptr< const LinOp >> get_mid_smoother_list () const
 - Gets the list of mid-smoother operators.
- std::vector< std::shared_ptr< const LinOp >> get_post_smoother_list () const

Gets the list of post-smoother operators.

- std::shared_ptr< const LinOp > get_coarsest_solver () const
 - Gets the operator at the coarsest level.
- multigrid::cycle get_cycle () const

Get the cycle of multigrid.

void set_cycle (multigrid::cycle cycle)

Set the cycle of multigrid.

43.136.1 Detailed Description

Multigrid methods have a hierarchy of many levels, whose corase level is a subset of the fine level, of the problem.

The coarse level solves the system on the residual of fine level and fine level will use the coarse solution to correct its own result. Multigrid solves the problem by relatively cheap step in each level and refining the result when prolongating back.

The main step of each level

Presmooth (solve on the fine level)

- · Calculate residual
- Restrict (reduce the problem dimension)
- · Solve residual in next level
- Prolongate (return to the fine level size)
- Postsmooth (correct the answer in fine level)

Ginkgo uses the index from 0 for finest level (original problem size) \sim N for the coarsest level (the coarsest solver), and its level counts is N (N multigrid level generation).

43.136.2 Member Function Documentation

43.136.2.1 apply_uses_initial_guess()

```
bool gko::solver::Multigrid::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess or false if multigrid always set the input as zero.

Returns

bool it is related to parameters variable zero_guess

References gko::solver::provided.

43.136.2.2 get coarsest solver()

```
std::shared_ptr<const LinOp> gko::solver::Multigrid::get_coarsest_solver ( ) const [inline]
```

Gets the operator at the coarsest level.

Returns

the coarsest operator

43.136.2.3 get_cycle()

```
multigrid::cycle gko::solver::Multigrid::get_cycle ( ) const [inline]
```

Get the cycle of multigrid.

Returns

the multigrid::cycle

43.136.2.4 get_mg_level_list()

```
std::vector<std::shared_ptr<const gko::multigrid::MultigridLevel> > gko::solver::Multigrid← ::get_mg_level_list ( ) const [inline]
```

Gets the list of MultigridLevel operators.

Returns

the list of MultigridLevel operators

43.136.2.5 get_mid_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_mid_smoother_list ( )
const [inline]
```

Gets the list of mid-smoother operators.

Returns

the list of mid-smoother operators

43.136.2.6 get_post_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_post_smoother_list ( )
const [inline]
```

Gets the list of post-smoother operators.

Returns

the list of post-smoother operators

43.136.2.7 get_pre_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_pre_smoother_list ( )
const [inline]
```

Gets the list of pre-smoother operators.

Returns

the list of pre-smoother operators

43.136.2.8 set_cycle()

Set the cycle of multigrid.

Parameters

multigrid::cycle the new cycle

The documentation for this class was generated from the following file:

• ginkgo/core/solver/multigrid.hpp

43.137 gko::multigrid::MultigridLevel Class Reference

This class represents two levels in a multigrid hierarchy.

```
#include <ginkgo/core/multigrid/multigrid_level.hpp>
```

Public Member Functions

- virtual std::shared_ptr< const LinOp > get_fine_op () const =0
 Returns the operator on fine level.
- virtual std::shared_ptr< const LinOp > get_restrict_op () const =0
 Returns the restrict operator.
- virtual std::shared_ptr< const LinOp > get_coarse_op () const =0
 Returns the operator on coarse level.
- virtual std::shared_ptr< const LinOp > get_prolong_op () const =0
 Returns the prolong operator.

43.137.1 Detailed Description

This class represents two levels in a multigrid hierarchy.

The MultigridLevel is an interface that allows to get the individual components of multigrid level. Each implementation of a multigrid level should inherit from this interface. Use EnableMultigridLevel<ValueType> to implement this interface with composition by default.

43.137.2 Member Function Documentation

43.137.2.1 get_coarse_op()

virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_coarse_op () const
[pure virtual]

Returns the operator on coarse level.

Returns

the operator on coarse level.

Implemented in gko::multigrid::EnableMultigridLevel< ValueType >.

43.137.2.2 get_fine_op()

virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_fine_op () const
[pure virtual]

Returns the operator on fine level.

Returns

the operator on fine level.

Implemented in gko::multigrid::EnableMultigridLevel < ValueType >.

43.137.2.3 get_prolong_op()

virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_prolong_op () const
[pure virtual]

Returns the prolong operator.

Returns

the prolong operator.

 $Implemented \ in \ gko::multigrid::Enable Multigrid Level < Value Type >.$

43.137.2.4 get_restrict_op()

virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_restrict_op () const
[pure virtual]

Returns the restrict operator.

Returns

the restrict operator.

Implemented in gko::multigrid::EnableMultigridLevel< ValueType >.

The documentation for this class was generated from the following file:

• ginkgo/core/multigrid/multigrid_level.hpp

43.138 gko::log::ProfilerHook::NestedSummaryWriter Class Reference

Recieves the results from ProfilerHook::create_nested_summary().

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

virtual void write_nested (const nested_summary_entry &root, std::chrono::nanoseconds overhead)=0
 Callback to write out the summary results.

43.138.1 Detailed Description

Recieves the results from ProfilerHook::create_nested_summary().

43.138.2 Member Function Documentation

43.138.2.1 write_nested()

Callback to write out the summary results.

Parameters

root	the root range with runtime and count.
overhead	an estimate of the profiler overhead

Implemented in gko::log::ProfilerHook::TableSummaryWriter.

The documentation for this class was generated from the following file:

• ginkgo/core/log/profiler_hook.hpp

43.139 gko::NotCompiled Class Reference

NotCompiled is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

NotCompiled (const std::string &file, int line, const std::string &func, const std::string &module)
 Initializes a NotCompiled error.

43.139.1 Detailed Description

NotCompiled is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.

43.139.2 Constructor & Destructor Documentation

43.139.2.1 NotCompiled()

Initializes a NotCompiled error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func Generated by	The name of the function that has not been compiled
module	The name of the module which contains the function

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.140 gko::NotImplemented Class Reference

NotImplemented is thrown in case an operation has not yet been implemented (but will be implemented in the future).

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

NotImplemented (const std::string &file, int line, const std::string &func)
 Initializes a NotImplemented error.

43.140.1 Detailed Description

NotImplemented is thrown in case an operation has not yet been implemented (but will be implemented in the future).

43.140.2 Constructor & Destructor Documentation

43.140.2.1 NotImplemented()

Initializes a NotImplemented error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the not-yet implemented function

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.141 gko::NotSupported Class Reference

NotSupported is thrown in case it is not possible to perform the requested operation on the given object type.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

NotSupported (const std::string &file, int line, const std::string &func, const std::string &obj_type)
 Initializes a NotSupported error.

43.141.1 Detailed Description

NotSupported is thrown in case it is not possible to perform the requested operation on the given object type.

43.141.2 Constructor & Destructor Documentation

43.141.2.1 NotSupported()

Initializes a NotSupported error.

Parameters

£:1 -	The course of the effective course Cla
file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the function where the error occured
obj_type	The object type on which the requested operation cannot be performed.

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.142 gko::null_deleter< T > Class Template Reference

This is a deleter that does not delete the object.

```
#include <ginkgo/core/base/utils_helper.hpp>
```

Public Member Functions

void operator() (pointer) const noexcept
 Deletes the object.

43.142.1 Detailed Description

```
template < typename T> class gko::null_deleter < T >
```

This is a deleter that does not delete the object.

It is useful where the object has been allocated elsewhere and will be deleted manually.

43.142.2 Member Function Documentation

43.142.2.1 operator()()

Deletes the object.

Parameters

```
ptr pointer to the object being deleted
```

The documentation for this class was generated from the following file:

• ginkgo/core/base/utils_helper.hpp

43.143 gko::nvidia_device Class Reference

nvidia_device handles the number of executor on Nvidia devices and have the corresponding recursive_mutex.

```
#include <ginkgo/core/base/device.hpp>
```

43.143.1 Detailed Description

nvidia_device handles the number of executor on Nvidia devices and have the corresponding recursive_mutex.

The documentation for this class was generated from the following file:

ginkgo/core/base/device.hpp

43.144 gko::OmpExecutor Class Reference

This is the Executor subclass which represents the OpenMP device (typically CPU).

#include <ginkgo/core/base/executor.hpp>

Public Member Functions

- std::shared_ptr< Executor > get_master () noexcept override
 Returns the master OmpExecutor of this Executor.
- std::shared_ptr< const Executor > get_master () const noexcept override

Returns the master OmpExecutor of this Executor.

· void synchronize () const override

Synchronize the operations launched on the executor with its master.

Static Public Member Functions

static std::shared_ptr< OmpExecutor > create ()
 Creates a new OmpExecutor.

43.144.1 Detailed Description

This is the Executor subclass which represents the OpenMP device (typically CPU).

43.144.2 Member Function Documentation

43.144.2.1 get master() [1/2]

std::shared_ptr<const Executor> gko::OmpExecutor::get_master () const [override], [virtual],
[noexcept]

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

43.144.2.2 get_master() [2/2]

```
std::shared_ptr<Executor> gko::OmpExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master OmpExecutor of this Executor.

Returns

the master OmpExecutor of this Executor.

Implements gko::Executor.

The documentation for this class was generated from the following file:

ginkgo/core/base/executor.hpp

43.145 gko::Operation Class Reference

Operations can be used to define functionalities whose implementations differ among devices.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

 virtual const char * get_name () const noexcept
 Returns the operation's name.

43.145.1 Detailed Description

Operations can be used to define functionalities whose implementations differ among devices.

This is done by extending the Operation class and implementing the overloads of the Operation::run() method for all Executor types. When invoking the Executor::run() method with the Operation as input, the library will select the Operation::run() overload corresponding to the dynamic type of the Executor instance.

Consider an overload of operator<< for Executors, which prints some basic device information (e.g. device type and id) of the Executor to a C++ stream:

```
std::ostream& operator (std::ostream &os, const gko::Executor &exec);
```

One possible implementation would be to use RTTI to find the dynamic type of the Executor, However, using the Operation feature of Ginkgo, there is a more elegant approach which utilizes polymorphism. The first step is to define an Operation that will print the desired information for each Executor type.

```
class DeviceInfoPrinter : public gko::Operation {
public:
    explicit DeviceInfoPrinter(std::ostream &os) : os_(os) {}
    void run(const gko::OmpExecutor *)const override { os_ « "OMP"; }
    void run(const gko::CudaExecutor *exec)const override
{ os_ « "CUDA(" « exec->get_device_id() « ")"; }
    void run(const gko::HipExecutor *exec)const override
{ os_ « "HIP(" « exec->get_device_id() « ")"; }
    void run(const gko::DpcppExecutor *exec)const override
{ os_ « "DPC++(" « exec->get_device_id() « ")"; }
    // This is optional, if not overloaded, defaults to OmpExecutor overload
    void run(const gko::ReferenceExecutor *)const override
{ os_ « "Reference CPU"; }
private:
```

```
std::ostream &os_;
}:
```

Using DeviceInfoPrinter, the implementation of operator<< is as simple as calling the run() method of the executor.

```
std::ostream& operator«(std::ostream &os, const gko::Executor &exec)
{
    DeviceInfoPrinter printer(os);
    exec.run(printer);
    return os;
}
```

Now it is possible to write the following code:

which produces the expected output:

```
OMP
CUDA(0)
HIP(0)
DPC++(0)
Reference CPU
```

One might feel that this code is too complicated for such a simple task. Luckily, there is an overload of the Executor::run() method, which is designed to facilitate writing simple operations like this one. The method takes four closures as input: one which is run for OMP, one for CUDA executors, one for HIP executors, and the last one for DPC++ executors. Using this method, there is no need to implement an Operation subclass:

Using this approach, however, it is impossible to distinguish between a OmpExecutor and ReferenceExecutor, as both of them call the OMP closure.

43.145.2 Member Function Documentation

43.145.2.1 get_name()

```
virtual const char* gko::Operation::get_name ( ) const [virtual], [noexcept]
```

Returns the operation's name.

Returns

the operation's name

The documentation for this class was generated from the following file:

ginkgo/core/base/executor.hpp

43.146 gko::log::operation_data Struct Reference

Struct representing Operator related data.

```
#include <ginkgo/core/log/record.hpp>
```

43.146.1 Detailed Description

Struct representing Operator related data.

The documentation for this struct was generated from the following file:

• ginkgo/core/log/record.hpp

43.147 gko::OutOfBoundsError Class Reference

OutOfBoundsError is thrown if a memory access is detected to be out-of-bounds.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

OutOfBoundsError (const std::string &file, int line, size_type index, size_type bound)
 Initializes an OutOfBoundsError.

43.147.1 Detailed Description

OutOfBoundsError is thrown if a memory access is detected to be out-of-bounds.

43.147.2 Constructor & Destructor Documentation

43.147.2.1 OutOfBoundsError()

Initializes an OutOfBoundsError.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
index	The position that was accessed
bound	The first out-of-bound index

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.148 gko::OverflowError Class Reference

OverflowError is thrown when an index calculation for storage requirements overflows.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• OverflowError (const std::string &file, const int line, const std::string &index_type)

43.148.1 Detailed Description

OverflowError is thrown when an index calculation for storage requirements overflows.

This most likely means that the index type is too small.

43.148.2 Constructor & Destructor Documentation

43.148.2.1 OverflowError()

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
index_type	The integer type that overflowed

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.149 gko::log::Papi < ValueType > Class Template Reference

Papi is a Logger which logs every event to the PAPI software.

```
#include <ginkgo/core/log/papi.hpp>
```

Public Member Functions

const std::string get_handle_name () const
 Returns the unique name of this logger, which can be used in the PAPI_read() call.

Static Public Member Functions

 static std::shared_ptr< Papi > create (std::shared_ptr< const gko::Executor >, const Logger::mask_type &enabled_events=Logger::all_events_mask)

```
Creates a Papi Logger.
```

static std::shared_ptr< Papi > create (const Logger::mask_type &enabled_events=Logger::all_events_

 mask)

Creates a Papi Logger.

43.149.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::log::Papi< ValueType >
```

Papi is a Logger which logs every event to the PAPI software.

Thanks to this logger, applications which interface with PAPI can access Ginkgo internal data through PAPI. For an example of usage, see examples/papi_logging/papi_logging.cpp

The logged values for each event are the following:

- · all allocation events: number of bytes per executor
- · all free events: number of calls per executor
- copy_started: number of bytes per executor from (to), in copy_started_from (respectively copy_started_to).
- copy_completed: number of bytes per executor from (to), in copy_completed_from (respectively copy_completed to).
- · all polymorphic objects and operation events: number of calls per executor
- all apply events: number of calls per LinOp (argument "A").
- · all factory events: number of calls per factory
- · criterion check completed event: the residual norm is stored in a record (per criterion)
- iteration_complete event: the number of iteration is counted (per solver)

Template Parameters

ValueType	the type of values stored in the class (e.g. residuals)
-----------	---

43.149.2 Member Function Documentation

43.149.2.1 create() [1/2]

Creates a Papi Logger.

Parameters

enabled_events	the events enabled for this Logger
----------------	------------------------------------

```
226 {
227         return std::shared_ptr<Papi>(new Papi(enabled_events));
228 }
```

43.149.2.2 create() [2/2]

Creates a Papi Logger.

Parameters

```
enabled_events the events enabled for this Logger
```

43.149.2.3 get_handle_name()

```
template<typename ValueType = default_precision>
const std::string gko::log::Papi< ValueType >::get_handle_name ( ) const [inline]
```

Returns the unique name of this logger, which can be used in the PAPI_read() call.

Returns

the unique name of this logger

The documentation for this class was generated from the following file:

ginkgo/core/log/papi.hpp

43.150 gko::factorization::Parlc< ValueType, IndexType > Class Template Reference

ParIC is an incomplete Cholesky factorization which is computed in parallel.

#include <ginkgo/core/factorization/par_ic.hpp>

Additional Inherited Members

43.150.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::factorization::Parlc< ValueType, IndexType >

ParIC is an incomplete Cholesky factorization which is computed in parallel.

L is a lower triangular matrix, which approximates a given matrix A with $A \approx LL^H$. Here, $L + L^H$ has the same sparsity pattern as A, which is also called IC(0).

The ParlC algorithm generates the incomplete factors iteratively, using a fixed-point iteration of the form

$$F(L) = \begin{cases} \sqrt{a_{ii} - \sum_{k=1}^{i-1} |l_{ik}|^2}, & i == j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i < j \end{cases}$$

In general, the entries of L can be iterated in parallel and in asynchronous fashion, the algorithm asymptotically converges to the incomplete factors L and L^H fulfilling $\left(R=A-L\cdot L^H\right)|_{\mathcal{S}}=0|_{\mathcal{S}}$ where \mathcal{S} is the pre-defined sparsity pattern (in case of IC(0) the sparsity pattern of the system matrix A). The number of ParlC sweeps needed for convergence depends on the parallelism level: For sequential execution, a single sweep is sufficient, for fine-grained parallelism, the number of sweeps necessary to get a good approximation of the incomplete factors depends heavily on the problem. On the OpenMP executor, 3 sweeps usually give a decent approximation in our experiments, while GPU executors can take 10 or more iterations.

The ParlC algorithm in Ginkgo follows the design of E. Chow and A. Patel, Fine-grained Parallel Incomplete LU Factorization, SIAM Journal on Scientific Computing, 37, C169-C193 (2015).

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

ginkgo/core/factorization/par_ic.hpp

43.151 gko::factorization::ParIct< ValueType, IndexType > Class Template Reference

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

#include <ginkgo/core/factorization/par_ict.hpp>

Additional Inherited Members

43.151.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::factorization::Parlct< ValueType, IndexType >

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

L is a lower triangular matrix which approximates a given symmetric positive definite matrix A with $A \approx LL^T$. Here, L has a sparsity pattern that is improved iteratively based on its element-wise magnitude. The initial sparsity pattern is chosen based on the lower triangle of A.

One iteration of the ParICT algorithm consists of the following steps:

- 1. Calculating the residual $R = A LL^T$
- 2. Adding new non-zero locations from R to L. The new non-zero locations are initialized based on the corresponding residual value.
- 3. Executing a fixed-point iteration on L according to $F(L) = \begin{cases} \frac{1}{l_{jj}} \left(a_{ij} \sum_{k=1}^{j-1} l_{ik} l_{jk} \right), & i \neq j \\ \sqrt{a_{ij} \sum_{k=1}^{j-1} l_{ik} l_{jk}}, & i = j \end{cases}$
- 4. Removing the smallest entries (by magnitude) from ${\cal L}$
- 5. Executing a fixed-point iteration on the (now sparser) ${\cal L}$

This ParICT algorithm thus improves the sparsity pattern and the approximation of L simultaneously.

The implementation follows the design of H. Anzt et al., ParILUT - A Parallel Threshold ILU for GPUs, 2019 IEEE International Parallel and Distributed Processing Symposium (IPDPS), pp. 231–241.

Template Parameters

ValueType	Type of the values of all matrices used in this class	
IndexType	Type of the indices of all matrices used in this class	

The documentation for this class was generated from the following file:

· ginkgo/core/factorization/par_ict.hpp

43.152 gko::factorization::Parllu< ValueType, IndexType > Class Template Reference

ParILU is an incomplete LU factorization which is computed in parallel.

#include <ginkgo/core/factorization/par_ilu.hpp>

Additional Inherited Members

43.152.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::factorization::Parllu< ValueType, IndexType >

ParILU is an incomplete LU factorization which is computed in parallel.

L is a lower unitriangular, while U is an upper triangular matrix, which approximate a given matrix A with $A \approx LU$. Here, L and U have the same sparsity pattern as A, which is also called ILU(0).

The ParlLU algorithm generates the incomplete factors iteratively, using a fixed-point iteration of the form

$$F(L,U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i > j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i \leq j \end{cases}$$

In general, the entries of L and U can be iterated in parallel and in asynchronous fashion, the algorithm asymptotically converges to the incomplete factors L and U fulfilling $(R=A-L\cdot U)\,|_{\mathcal{S}}=0|_{\mathcal{S}}$ where \mathcal{S} is the pre-defined sparsity pattern (in case of ILU(0) the sparsity pattern of the system matrix A). The number of ParlLU sweeps needed for convergence depends on the parallelism level: For sequential execution, a single sweep is sufficient, for fine-grained parallelism, the number of sweeps necessary to get a good approximation of the incomplete factors depends heavily on the problem. On the OpenMP executor, 3 sweeps usually give a decent approximation in our experiments, while GPU executors can take 10 or more iterations.

The ParlLU algorithm in Ginkgo follows the design of E. Chow and A. Patel, Fine-grained Parallel Incomplete LU Factorization, SIAM Journal on Scientific Computing, 37, C169-C193 (2015).

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

ginkgo/core/factorization/par_ilu.hpp

43.153 gko::factorization::Parllut< ValueType, IndexType > Class Template Reference

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

#include <ginkgo/core/factorization/par_ilut.hpp>

Additional Inherited Members

43.153.1 Detailed Description

template < typename ValueType = default_precision, typename IndexType = int32 > class gko::factorization::Parllut < ValueType, IndexType >

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

L is a lower unitriangular, while U is an upper triangular matrix, which approximate a given matrix A with $A \approx LU$. Here, L and U have a sparsity pattern that is improved iteratively based on their element-wise magnitude. The initial sparsity pattern is chosen based on the ILU(0) factorization of A.

One iteration of the ParILUT algorithm consists of the following steps:

- 1. Calculating the residual R = A LU
- 2. Adding new non-zero locations from R to L and U. The new non-zero locations are initialized based on the corresponding residual value.
- 3. Executing a fixed-point iteration on L and U according to $F(L,U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i>j \\ a_{ij} \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i\leq j \end{cases}$ For a more detailed description of the fixed-point iteration, see Parllu.
- 4. Removing the smallest entries (by magnitude) from ${\cal L}$ and ${\cal U}$
- 5. Executing a fixed-point iteration on the (now sparser) L and U

This ParlLUT algorithm thus improves the sparsity pattern and the approximation of L and U simultaneously.

The implementation follows the design of H. Anzt et al., ParlLUT - A Parallel Threshold ILU for GPUs, 2019 IEEE International Parallel and Distributed Processing Symposium (IPDPS), pp. 231–241.

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

• ginkgo/core/factorization/par_ilut.hpp

43.154 gko::experimental::distributed::Partition < LocalIndexType, GlobalIndexType > Class Template Reference

Represents a partition of a range of indices [0, size) into a disjoint set of parts.

#include <ginkgo/core/distributed/partition.hpp>

Public Member Functions

· size_type get_size () const

Returns the total number of elements represented by this partition.

· size_type get_num_ranges () const noexcept

Returns the number of ranges stored by this partition.

· comm_index_type get_num_parts () const noexcept

Returns the number of parts represented in this partition.

comm_index_type get_num_empty_parts () const noexcept

Returns the number of empty parts within this partition.

const global_index_type * get_range_bounds () const noexcept

Returns the ranges boundary array stored by this partition.

const comm_index_type * get_part_ids () const noexcept

Returns the part IDs of the ranges in this partition.

const local_index_type * get_range_starting_indices () const noexcept

Returns the part-local starting index for each range in this partition.

const local_index_type * get_part_sizes () const noexcept

Returns the part size array.

local_index_type get_part_size (comm_index_type part) const

Returns the size of a part given by its part ID.

bool has_connected_parts ()

Checks if each part has no more than one contiguous range.

• bool has_ordered_parts ()

Checks if the ranges are ordered by their part index.

Static Public Member Functions

static std::unique_ptr< Partition > build_from_mapping (std::shared_ptr< const Executor > exec, const array< comm_index_type > &mapping, comm_index_type num_parts)

Builds a partition from a given mapping global_index -> part_id.

static std::unique_ptr< Partition > build_from_contiguous (std::shared_ptr< const Executor > exec, const array< global_index_type > &ranges)

Builds a partition consisting of contiguous ranges, one for each part.

 static std::unique_ptr< Partition > build_from_global_size_uniform (std::shared_ptr< const Executor > exec, comm_index_type num_parts, global_index_type global_size)

Builds a partition by evenly distributing the global range.

43.154.1 Detailed Description

template<typename LocalIndexType = int32, typename GlobalIndexType = int64> class gko::experimental::distributed::Partition < LocalIndexType, GlobalIndexType >

Represents a partition of a range of indices [0, size) into a disjoint set of parts.

The partition is stored as a set of consecutive ranges [begin, end) with an associated part ID and local index (number of indices in this part before begin). Global indices are stored as 64 bit signed integers (int64), part-local indices use LocalIndexType, Part IDs use 32 bit signed integers (int).

For example, consider the interval [0, 13) that is partitioned into the following ranges:

```
[0,3), [3, 6), [6, 8), [8, 10), [10, 13).
```

These ranges are distributed on three part with:

```
p_0 = [0, 3) + [6, 8) + [10, 13),

p_1 = [3, 6),
p_2 = [8, 10).
```

The part ids can be queried from the get part ids array, and the ranges are represented as offsets, accessed by get range bounds, leading to the offset array:

```
r = [0, 3, 6, 8, 10, 13]
```

so that individual ranges are given by [r[i], r[i + 1]). Since each part may be associated with multiple ranges, it is possible to get the starting index for each range that is local to the owning part, see get_range_starting_indices. These indices can be used to easily iterate over part local data. For example, the above partition has the following starting indices

```
starting_index[0] = 0,
starting_index[1] = 0,
                        // second range of part 1
starting\_index[2] = 3,
starting index[3] = 0,
starting_index[4] = 5, // third range of part 1
```

which you can use to iterate only over the the second range of part 1 (the third global range) with

```
for(int i = 0; i < r[3] - r[2]; ++i){
  data[starting_index[2] + i] = val;</pre>
```

Template Parameters

LocalIndexType	The index type used for part-local indices. To prevent overflows, no single part's size may exceed this index type's maximum value.
GlobalIndexType	The index type used for the global indices. Needs to be at least as large a type as LocalIndexType.

43.154.2 Member Function Documentation

43.154.2.1 build from contiguous()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
static std::unique_ptr<Partition> gko::experimental::distributed::Partition< LocalIndexType,
GlobalIndexType >::build_from_contiguous (
            std::shared_ptr< const Executor > exec,
            const array< global_index_type > & ranges ) [static]
```

Builds a partition consisting of contiguous ranges, one for each part.

Parameters

exec	the Executor on which the partition should be built	
ranges	the boundaries of the ranges representing each part. Part i contains the indices [ranges[i], ranges[i +	
	1]). Has to contain at least one element. The first element has to be 0.	

Returns

a Partition representing the given contiguous partitioning.

43.154.2.2 build_from_global_size_uniform()

Builds a partition by evenly distributing the global range.

Parameters

exec	the Executor on which the partition should be built	
num_parts	the number of parst used in this partition	
global_size	the global size of this partition	

Returns

```
a Partition where each range has either floor(global_size/num_parts) or floor(global_\leftarrow size/num_parts) + 1 indices.
```

43.154.2.3 build_from_mapping()

Builds a partition from a given mapping global_index -> part_id.

Parameters

exec	the Executor on which the partition should be built	
mapping	the mapping from global indices to part IDs.	
num_parts	the number of parts used in the mapping.	

Returns

a Partition representing the given mapping as a set of ranges

43.154.2.4 get_num_empty_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
comm_index_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >
::get_num_empty_parts () const [inline], [noexcept]
```

Returns the number of empty parts within this partition.

Returns

number of empty parts.

43.154.2.5 get_num_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
comm_index_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >
::get_num_parts ( ) const [inline], [noexcept]
```

Returns the number of parts represented in this partition.

Returns

number of parts.

43.154.2.6 get_num_ranges()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
size_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_
num_ranges ( ) const [inline], [noexcept]
```

Returns the number of ranges stored by this partition.

This size refers to the data returned by get_range_bounds().

Returns

number of ranges.

References gko::array< ValueType >::get_num_elems().

43.154.2.7 get_part_ids()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const comm_index_type* gko::experimental::distributed::Partition< LocalIndexType, Global←
IndexType >::get_part_ids ( ) const [inline], [noexcept]
```

Returns the part IDs of the ranges in this partition.

For each range from get_range_bounds(), it stores the part ID in the interval [0, get_num_parts() - 1].

Returns

part ID array.

References gko::array< ValueType >::get_const_data().

43.154.2.8 get_part_size()

Returns the size of a part given by its part ID.

Warning

Triggers a copy from device to host.

Parameters

```
part the part ID.
```

Returns

size of part.

References gko::array< ValueType >::get_const_data(), and gko::PolymorphicObject::get_executor().

43.154.2.9 get_part_sizes()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const local_index_type* gko::experimental::distributed::Partition< LocalIndexType, Global←
IndexType >::get_part_sizes ( ) const [inline], [noexcept]
```

Returns the part size array.

part_sizes[p] stores the total number of indices in part p.

779

Returns

part size array.

References gko::array< ValueType >::get const data().

43.154.2.10 get_range_bounds()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const global_index_type* gko::experimental::distributed::Partition< LocalIndexType, Global←
IndexType >::get_range_bounds ( ) const [inline], [noexcept]
```

Returns the ranges boundary array stored by this partition.

 $range_bounds[i]$ is the beginning (inclusive) and $range_bounds[i+1]$ is the end (exclusive) of the ith range.

Returns

range boundaries array.

References gko::array< ValueType >::get_const_data().

43.154.2.11 get_range_starting_indices()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const local_index_type* gko::experimental::distributed::Partition< LocalIndexType, Global←
IndexType >::get_range_starting_indices ( ) const [inline], [noexcept]
```

Returns the part-local starting index for each range in this partition.

```
Consider the partition on [0, 10) with p_1 = [0-4) + [7-10), p_2 = [4-7).
```

Then range_starting_indices[0] = 0, range_starting_indices[1] = 0, range_ \leftrightarrow starting_indices[2] = 4.

Returns

part-local starting index array.

References gko::array< ValueType >::get const data().

43.154.2.12 get_size()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
size_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_
size () const [inline]
```

Returns the total number of elements represented by this partition.

Returns

number elements.

43.154.2.13 has_connected_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
bool gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::has_\(\cupercolon\)
connected_parts ( )
```

Checks if each part has no more than one contiguous range.

Returns

true if each part has no more than one contiguous range.

43.154.2.14 has_ordered_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
bool gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::has_
ordered_parts ( )
```

Checks if the ranges are ordered by their part index.

Implies that the partition is connected.

Returns

true if the ranges are ordered by their part index.

The documentation for this class was generated from the following files:

- · ginkgo/core/distributed/matrix.hpp
- · ginkgo/core/distributed/partition.hpp

43.155 gko::log::PerformanceHint Class Reference

PerformanceHint is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

```
#include <ginkgo/core/log/performance_hint.hpp>
```

Public Member Functions

· void print_status () const

Writes out the cross-executor writes and allocations that have been stored so far.

Static Public Member Functions

static std::unique_ptr< PerformanceHint > create (std::ostream &os=std::cerr, size_type allocation_size_
 — limit=16, size_type copy_size_limit=16, size_type histogram_max_size=1024)

Creates a PerformanceHint logger.

43.155.1 Detailed Description

PerformanceHint is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

The specific patterns it checks for are:

- · repeated cross-executor copies from or to the same pointer
- · repeated allocation/free pairs of the same size

43.155.2 Member Function Documentation

43.155.2.1 create()

```
static std::unique_ptr<PerformanceHint> gko::log::PerformanceHint::create (
    std::ostream & os = std::cerr,
    size_type allocation_size_limit = 16,
    size_type copy_size_limit = 16,
    size_type histogram_max_size = 1024 ) [inline], [static]
```

Creates a PerformanceHint logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

os	the stream used for this logger
allocation_size_limit	ignore allocations below this limit (bytes)
copy_size_limit	ignore copies below this limit (bytes)
histogram_max_size	how many allocation sizes and/or pointers to keep track of at most?

Returns

an std::unique_ptr to the the constructed object

The documentation for this class was generated from the following file:

· ginkgo/core/log/performance_hint.hpp

43.156 gko::Permutable < IndexType > Class Template Reference

Linear operators which support permutation should implement the Permutable interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual std::unique_ptr< LinOp > permute (const array< IndexType > *permutation_indices) const Returns a LinOp representing the symmetric row and column permutation of the Permutable object.
- virtual std::unique_ptr< LinOp > inverse_permute (const array< IndexType > *permutation_indices) const Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.
- virtual std::unique_ptr< LinOp > row_permute (const array< IndexType > *permutation_indices) const =0

 Returns a LinOp representing the row permutation of the Permutable object.
- virtual std::unique_ptr< LinOp > column_permute (const array< IndexType > *permutation_indices) const
 =0

Returns a LinOp representing the column permutation of the Permutable object.

virtual std::unique_ptr< LinOp > inverse_row_permute (const array< IndexType > *permutation_indices)
 const =0

Returns a LinOp representing the row permutation of the inverse permuted object.

virtual std::unique_ptr< LinOp > inverse_column_permute (const array< IndexType > *permutation_← indices) const =0

Returns a LinOp representing the row permutation of the inverse permuted object.

43.156.1 Detailed Description

```
template<typename IndexType>
class gko::Permutable< IndexType>
```

Linear operators which support permutation should implement the Permutable interface.

It provides functions to permute the rows and columns of a LinOp, independently or symmetrically, and with a regular or inverted permutation.

After a regular row permutation with permutation array perm the row i in the output LinOp contains the row perm[i] from the input LinOp. After an inverse row permutation, the row perm[i] in the output LinOp contains the row i from the input LinOp. Equivalently, after a column permutation, the output stores in column i the column perm[i] from the input, and an inverse column permutation stores in column perm[i] the column i from the input. A symmetric permutation is functionally equivalent to calling as<Permutable>(A->row_ \leftarrow permute (perm)) ->column_permute (perm), but the implementation can provide better performance due to kernel fusion.

43.156.1.1 Example: Permuting a Csr matrix:

```
{c++}
//Permuting an object of LinOp type.
//The object you want to permute.
auto op = matrix::Csr::create(exec);
//Permute the object by first converting it to a Permutable type.
auto perm = op->row_permute(permutation_indices);
```

43.156.2 Member Function Documentation

43.156.2.1 column_permute()

Returns a LinOp representing the column permutation of the Permutable object.

In the resulting LinOp, the column i contains the input column perm[i].

Parameters

```
permutation_indices the array of indices containing the permutation order perm.
```

Returns

a pointer to the new column permuted object

Implemented in gko::matrix::Dense< ValueType >, gko::matrix::Dense< value_type >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType >, and gko::matrix::Csr< ValueType, IndexType >.

43.156.2.2 inverse_column_permute()

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column perm[i] contains the input column i.

Parameters

permutation_indices	the array of indices containing the permutation order perm.
---------------------	---

Returns

a pointer to the new inverse permuted object

 $Implemented \ in \ gko::matrix::Dense<\ ValueType>, \ gko::matrix::Dense<\ ValueType>, \ gko::matrix::Dense<\ ValueType>, \ gko::matrix::Dense<\ ValueType>, \ and \ gko::matrix::Csr<\ ValueType>, \ IndexType>.$

43.156.2.3 inverse permute()

Returns a LinOp representing the symmetric inverse row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (perm[i], perm[j]) contains the input value (i, j).

Parameters

```
permutation_indices the array of indices containing the permutation order.
```

Returns

a pointer to the new permuted object

Reimplemented in gko::matrix::Dense< ValueType >, gko::matrix::Dense< value_type >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType >, and gko::matrix::Csr< ValueType, IndexType >.

43.156.2.4 inverse_row_permute()

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row $\mathtt{perm}\,[\,\mathtt{i}\,]$ contains the input row $\mathtt{i}\,.$

Parameters

permutation indices	the array of indices containing the permutation order perm.

Returns

a pointer to the new inverse permuted object

Implemented in gko::matrix::Dense < ValueType >, gko::matrix::Dense < value_type >, gko::matrix::Dense < ValueType >, gko::matrix::Dense < ValueType >, and gko::matrix::Csr < ValueType, IndexType >.

43.156.2.5 permute()

Returns a LinOp representing the symmetric row and column permutation of the Permutable object.

In the resulting LinOp, the entry at location (i, j) contains the input value (perm[i], perm[j]).

Parameters

Returns

a pointer to the new permuted object

Reimplemented in gko::matrix::Dense < ValueType >, gko::matrix::Dense < value_type >, gko::matrix::Dense < ValueType >, gko::matrix::Dense < ValueType >, and gko::matrix::Csr < ValueType, IndexType >.

43.156.2.6 row_permute()

Returns a LinOp representing the row permutation of the Permutable object.

In the resulting LinOp, the row i contains the input row perm[i].

Parameters

permutation indices	the array of indices containing the permutation order.

Returns

a pointer to the new permuted object

Implemented in gko::matrix::Dense < ValueType >, gko::matrix::Dense < value_type >, gko::matrix::Dense < ValueType >, gko::matrix::Dense < ValueType >, and gko::matrix::Csr < ValueType, IndexType >.

The documentation for this class was generated from the following file:

· ginkgo/core/base/lin op.hpp

43.157 gko::matrix::Permutation < IndexType > Class Template Reference

Permutation is a matrix "format" which stores the row and column permutation arrays which can be used for reordering the rows and columns a matrix.

#include <ginkgo/core/matrix/permutation.hpp>

Public Member Functions

index_type * get_permutation () noexcept

Returns a pointer to the array of permutation.

const index_type * get_const_permutation () const noexcept

Returns a pointer to the array of permutation.

• size_type get_permutation_size () const noexcept

Returns the number of elements explicitly stored in the permutation array.

mask_type get_permute_mask () const

Get the permute masks.

void set_permute_mask (mask_type permute_mask)

Set the permute masks.

Static Public Member Functions

• static std::unique_ptr< const Permutation > create_const (std::shared_ptr< const Executor > exec, size_type size, gko::detail::const_array_view< IndexType > &&perm_idxs, mask_type enabled_const_permute=row permute)

Creates a constant (immutable) Permutation matrix from a constant array.

43.157.1 Detailed Description

```
template<typename IndexType = int32> class gko::matrix::Permutation< IndexType >
```

Permutation is a matrix "format" which stores the row and column permutation arrays which can be used for reordering the rows and columns a matrix.

Template Parameters

IndexType precision of	permutation array indices.
------------------------	----------------------------

Note

This format is used mainly to allow for an abstraction of the permutation/re-ordering and provides the user with an apply method which calls the respective LinOp's permute operation if the respective LinOp implements the Permutable interface. As such it only stores an array of the permutation indices.

43.157.2 Member Function Documentation

43.157.2.1 create const()

Creates a constant (immutable) Permutation matrix from a constant array.

Parameters

exec	the executor to create the matrix on
size	the size of the square matrix
perm_idxs	the permutation index array of the matrix
enabled_permute	the mask describing the type of permutation

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

43.157.2.2 get_const_permutation()

```
template<typename IndexType = int32>
const index_type* gko::matrix::Permutation< IndexType >::get_const_permutation ( ) const [inline],
[noexcept]
```

Returns a pointer to the array of permutation.

Returns

the pointer to the row permutation array.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.157.2.3 get_permutation()

```
template<typename IndexType = int32>
index_type* gko::matrix::Permutation< IndexType >::get_permutation ( ) [inline], [noexcept]
```

Returns a pointer to the array of permutation.

Returns

the pointer to the row permutation array.

References gko::array< ValueType >::get_data().

43.157.2.4 get_permutation_size()

```
template<typename IndexType = int32>
size_type gko::matrix::Permutation< IndexType >::get_permutation_size ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the permutation array.

Returns

the number of elements explicitly stored in the permutation array.

References gko::array< ValueType >::get_num_elems().

43.157.2.5 get_permute_mask()

```
template<typename IndexType = int32>
mask_type gko::matrix::Permutation< IndexType >::get_permute_mask ( ) const [inline]
```

Get the permute masks.

Returns

permute_mask the permute masks

43.157.2.6 set_permute_mask()

Set the permute masks.

Parameters

permute_mask	the permute masks
--------------	-------------------

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/permutation.hpp

43.158 gko::Perturbation < ValueType > Class Template Reference

The Perturbation class can be used to construct a LinOp to represent the operation (identity + scalar * basis * projector).

#include <ginkgo/core/base/perturbation.hpp>

Public Member Functions

- const std::shared_ptr< const LinOp > get_basis () const noexcept
 Returns the basis of the perturbation.
- const std::shared_ptr< const LinOp > get_projector () const noexcept

 Returns the projector of the perturbation.
- const std::shared_ptr< const LinOp > get_scalar () const noexcept
 Returns the scalar of the perturbation.

43.158.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::Perturbation< ValueType >
```

The Perturbation class can be used to construct a LinOp to represent the operation (identity + scalar * basis * projector).

This operator adds a movement along a direction constructed by basis and projector on the LinOp. projector gives the coefficient of basis to decide the direction.

For example, the Householder matrix can be represented with the Perturbation operator as follows. If u is the Householder factor then we can generate the $Householder\ transformation$, $H = (I - 2\ u\ u*)$. In this case, the parameters of Perturbation class are scalar = -2, basis = u, and projector = u*.

Template Parameters

ValueType	precision of input and result vectors
-----------	---------------------------------------

Note

the apply operations of Perturbation class are not thread safe

43.158.2 Member Function Documentation

43.158.2.1 get_basis()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_basis ( ) const [inline],
[noexcept]
```

Returns the basis of the perturbation.

Returns

the basis of the perturbation

```
81 {
82     return basis_;
83 }
```

43.158.2.2 get_projector()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_projector ( ) const
[inline], [noexcept]
```

Returns the projector of the perturbation.

Returns

the projector of the perturbation

43.158.2.3 get_scalar()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_scalar ( ) const [inline],
[noexcept]
```

Returns the scalar of the perturbation.

Returns

the scalar of the perturbation

The documentation for this class was generated from the following file:

ginkgo/core/base/perturbation.hpp

43.159 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference

Parallel graph match (Pgm) is the aggregate method introduced in the paper M.

```
#include <ginkgo/core/multigrid/pgm.hpp>
```

Public Member Functions

- std::shared_ptr< const LinOp > get_system_matrix () const Returns the system operator (matrix) of the linear system.
- IndexType * get_agg () noexcept
 Returns the aggregate group.
- const IndexType * get_const_agg () const noexcept
 Returns the aggregate group.

43.159.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::multigrid::Pgm< ValueType, IndexType >
```

Parallel graph match (Pgm) is the aggregate method introduced in the paper M.

Naumov et al., "AmgX: A Library for GPU Accelerated Algebraic Multigrid and Preconditioned Iterative Methods". Current implementation only contains size = 2 version.

Pgm creates the aggregate group according to the matrix value not the structure. Pgm gives two steps (one-phase handshaking) to group the elements. 1: get the strongest neighbor of each unaggregated element. 2: group the elements whose strongest neighbor is each other. repeating until reaching the given conditions. After that, the un-aggregated elements are assigned to an aggregated group or are left alone.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.159.2 Member Function Documentation

43.159.2.1 get_agg()

```
template<typename ValueType = default_precision, typename IndexType = int32>
IndexType* gko::multigrid::Pgm< ValueType, IndexType >::get_agg ( ) [inline], [noexcept]
```

Returns the aggregate group.

Aggregate group whose size is same as the number of rows. Stores the mapping information from row index to coarse row index. i.e., agg[row_idx] = coarse_row_idx.

Returns

the aggregate group.

```
103 { return agg_.get_data(); }
```

References gko::array< ValueType >::get_data().

43.159.2.2 get_const_agg()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const IndexType* gko::multigrid::Pgm< ValueType, IndexType >::get_const_agg ( ) const [inline],
[noexcept]
```

Returns the aggregate group.

Aggregate group whose size is same as the number of rows. Stores the mapping information from row index to coarse row index. i.e., agg[row_idx] = coarse_row_idx.

Returns

the aggregate group.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.159.2.3 get_system_matrix()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const LinOp> gko::multigrid::Pgm< ValueType, IndexType >::get_system_matrix (
) const [inline]
```

Returns the system operator (matrix) of the linear system.

Returns

the system operator (matrix)

The documentation for this class was generated from the following file:

• ginkgo/core/multigrid/pgm.hpp

43.160 gko::log::polymorphic object data Struct Reference

Struct representing PolymorphicObject related data.

#include <ginkgo/core/log/record.hpp>

43.160.1 Detailed Description

Struct representing PolymorphicObject related data.

The documentation for this struct was generated from the following file:

· ginkgo/core/log/record.hpp

43.161 gko::PolymorphicObject Class Reference

A PolymorphicObject is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

Public Member Functions

- std::unique_ptr< PolymorphicObject > create_default (std::shared_ptr< const Executor > exec) const Creates a new "default" object of the same dynamic type as this object.
- std::unique_ptr< PolymorphicObject > create_default () const

Creates a new "default" object of the same dynamic type as this object.

- std::unique_ptr< PolymorphicObject > clone (std::shared_ptr< const Executor > exec) const
 Creates a clone of the object.
- std::unique_ptr< PolymorphicObject > clone () const

Creates a clone of the object.

• PolymorphicObject * copy_from (const PolymorphicObject *other)

Copies another object into this object.

- template<typename Derived , typename Deleter >

 $std::enable_if_t < std::is_base_of < PolymorphicObject, std::decay_t < Derived > >::value, PolymorphicObject > * copy_from (std::unique_ptr < Derived, Deleter > &&other)$

Moves another object into this object.

• template<typename Derived , typename Deleter >

 $std::enable_if_t < std::is_base_of < PolymorphicObject, std::decay_t < Derived > >::value, PolymorphicObject > * copy_from (const std::unique_ptr < Derived, Deleter > & other)$

Copies another object into this object.

PolymorphicObject * copy_from (const std::shared_ptr< const PolymorphicObject > &other)

Copies another object into this object.

PolymorphicObject * move_from (ptr_param< PolymorphicObject > other)

Moves another object into this object.

PolymorphicObject * clear ()

Transforms the object into its default state.

std::shared_ptr< const Executor > get_executor () const noexcept

Returns the Executor of the object.

43.161.1 Detailed Description

A PolymorphicObject is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.

It defines the basic utilities (copying moving, cloning, clearing the objects) for all such objects. It takes into account that these objects are dynamically allocated, managed by smart pointers, and used polymorphically. Additionally, it assumes their data can be allocated on different executors, and that they can be copied between those executors.

Note

Most of the public methods of this class should not be overridden directly, and are thus not virtual. Instead, there are equivalent protected methods (ending in <method_name>_impl) that should be overriden instead. This allows polymorphic objects to implement default behavior around virtual methods (parameter checking, type casting).

See also

EnablePolymorphicObject if you wish to implement a concrete polymorphic object and have sensible defaults generated automatically. EnableAbstractPolymorphicObject if you wish to implement a new abstract polymorphic object, and have the return types of the methods updated to your type (instead of having them return PolymorphicObject).

43.161.2 Member Function Documentation

43.161.2.1 clear()

```
PolymorphicObject* gko::PolymorphicObject::clear ( ) [inline]
```

Transforms the object into its default state.

Equivalent to this->copy_from(this->create_default()).

See also

clear_impl() when implementing this method

Returns

this

43.161.2.2 clone() [1/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::clone ( ) const [inline]
```

Creates a clone of the object.

This is a shorthand for clone(std::shared_ptr<const Executor>) that uses the executor of this object to construct the new object.

Returns

A clone of the LinOp.

43.161.2.3 clone() [2/2]

Creates a clone of the object.

This is the polymorphic equivalent of the *executor copy constructor* decltype (*this) (exec, this).

Parameters

```
exec the executor where the clone will be created
```

Returns

A clone of the LinOp.

References create_default().

43.161.2.4 copy_from() [1/4]

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

```
copy_from_impl(const PolymorphicObject *)
```

Parameters

other the object to copy	other	the object to copy
--------------------------	-------	--------------------

Returns

this

Referenced by copy_from().

43.161.2.5 copy_from() [2/4]

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

copy_from_impl(const PolymorphicObject *)

Parameters

the object to co	the object to copy
------------------	--------------------

Returns

this

References copy_from().

43.161.2.6 copy_from() [3/4]

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

copy_from_impl(const PolymorphicObject *)

Parameters

other	the object to copy
other	the object to copy

Returns

this

Template Parameters

Derived	the actual pointee type of the parameter, it needs to be derived from PolymorphicObject.
Deleter	the deleter of the unique_ptr parameter

References copy_from().

43.161.2.7 copy_from() [4/4]

Moves another object into this object.

This is the polymorphic equivalent of the move assignment operator.

See also

copy_from_impl(std::unique_ptr<PolymorphicObject>)

Parameters

other the object to mov

Returns

this

Template Parameters

Derived	the actual pointee type of the parameter, it needs to be derived from PolymorphicObject.
Deleter	the deleter of the unique_ptr parameter

43.161.2.8 create_default() [1/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::create_default ( ) const [inline]
```

Creates a new "default" object of the same dynamic type as this object.

This is a shorthand for create_default(std::shared_ptr<const Executor>) that uses the executor of this object to construct the new object.

Returns

a polymorphic object of the same type as this

Referenced by clone().

43.161.2.9 create_default() [2/2]

Creates a new "default" object of the same dynamic type as this object.

This is the polymorphic equivalent of the executor default constructor decltype (*this) (exec);.

Parameters

exec the executor where the object will be created

Returns

a polymorphic object of the same type as this

43.161.2.10 get_executor()

```
std::shared_ptr<const Executor> gko::PolymorphicObject::get_executor ( ) const [inline],
[noexcept]
```

Returns the Executor of the object.

Returns

Executor of the object

Referenced by gko::matrix::Coo< ValueType, IndexType >::apply2(), gko::preconditioner::lc< LSolverType, Index \leftarrow Type >::conj_transpose(), gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType > \leftarrow ::conj_transpose(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_size(), gko::preconditioner::lc< LSolverType, IndexType >::lc(), gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::ilu(), gko::matrix::Csr< ValueType, IndexType >::inv_scale(), gko::preconditioner \leftarrow ::lc< LSolverType, IndexType >::operator=(), gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::operator=(), gko::preconditioner::lc< LSolverType, USolverType, ReverseApply, IndexType >::transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose().

43.161.2.11 move_from()

Moves another object into this object.

This is the polymorphic equivalent of the move assignment operator.

See also

move_from_impl(PolymorphicObject *)

Parameters

other	the object to move from
-------	-------------------------

Returns

this

References gko::ptr_param< T >::get().

The documentation for this class was generated from the following file:

· ginkgo/core/base/polymorphic object.hpp

43.162 gko::precision_reduction Class Reference

This class is used to encode storage precisions of low precision algorithms.

#include <ginkgo/core/base/types.hpp>

Public Types

• using storage_type = uint8

The underlying datatype used to store the encoding.

Public Member Functions

constexpr precision_reduction () noexcept

Creates a default precision_reduction encoding.

- constexpr precision_reduction (storage_type preserving, storage_type nonpreserving) noexcept
 - Creates a precision reduction encoding with the specified number of conversions.
- constexpr operator storage_type () const noexcept

Extracts the raw data of the encoding.

• constexpr storage_type get_preserving () const noexcept

Returns the number of preserving conversions in the encoding.

constexpr storage_type get_nonpreserving () const noexcept

Returns the number of non-preserving conversions in the encoding.

Static Public Member Functions

· constexpr static precision_reduction autodetect () noexcept

Returns a special encoding which instructs the algorithm to automatically detect the best precision.

• constexpr static precision_reduction common (precision_reduction x, precision_reduction y) noexcept Returns the common encoding of input encodings.

43.162.1 Detailed Description

This class is used to encode storage precisions of low precision algorithms.

Some algorithms in Ginkgo can improve their performance by storing parts of the data in lower precision, while doing computation in full precision. This class is used to encode the precisions used to store the data. From the user's perspective, some algorithms can provide a parameter for fine-tuning the storage precision. Commonly, the special value returned by precision_reduction::autodetect() should be used to allow the algorithm to automatically choose an appropriate value, though manually selected values can be used for fine-tuning.

In general, a lower precision floating point value can be obtained by either dropping some of the insignificant bits of the significand (keeping the same number of exponent bits, and thus preserving the range of representable values) or using one of the hardware or software supported conversions between IEEE formats, such as double to float or float to half (reducing both the number of exponent, as well as significand bits, and thus decreasing the range of representable values).

The precision_reduction class encodes the lower precision format relative to the base precision used and the algorithm in question. The encoding is done by specifying the amount of range non-preserving conversions and the amount of range preserving conversions that should be done on the base precision to obtain the lower precision format. For example, starting with a double precision value (11 exp, 52 sig. bits), the encoding specifying 1 non-preserving conversion and 1 preserving conversion would first use a hardware-supported non-preserving conversion to obtain a single precision value (8 exp, 23 sig. bits), followed by a preserving bit truncation to obtain a value with 8 exponent and 7 significand bits. Note that non-preserving conversion are always done first, as preserving conversions usually result in datatypes that are not supported by builtin conversions (thus, it is generally not possible to apply a non-preserving conversion to the result of a preserving conversion).

If the specified conversion is not supported by the algorithm, it will most likely fall back to using full precision for storing the data. Refer to the documentation of specific algorithms using this class for details about such special cases.

43.162.2 Constructor & Destructor Documentation

43.162.2.1 precision_reduction() [1/2]

```
constexpr gko::precision_reduction::precision_reduction ( ) [inline], [constexpr], [noexcept]
```

Creates a default precision_reduction encoding.

This encoding represents the case where no conversions are performed.

Referenced by common().

43.162.2.2 precision_reduction() [2/2]

Creates a precision reduction encoding with the specified number of conversions.

Parameters

preserving	the number of range preserving conversion
nonpreserving	the number of range non-preserving conversions

43.162.3 Member Function Documentation

43.162.3.1 autodetect()

```
constexpr static precision_reduction gko::precision_reduction::autodetect ( ) [inline], [static],
[constexpr], [noexcept]
```

Returns a special encoding which instructs the algorithm to automatically detect the best precision.

Returns

a special encoding instructing the algorithm to automatically detect the best precision.

43.162.3.2 common()

Returns the common encoding of input encodings.

The common encoding is defined as the encoding that does not have more preserving, nor non-preserving conversions than the input encodings.

Parameters

Х	an encoding
у	an encoding

Returns

the common encoding of \boldsymbol{x} and \boldsymbol{y}

References precision reduction().

43.162.3.3 get_nonpreserving()

```
constexpr storage_type gko::precision_reduction::get_nonpreserving ( ) const [inline], [constexpr],
[noexcept]
```

Returns the number of non-preserving conversions in the encoding.

Returns

the number of non-preserving conversions in the encoding.

43.162.3.4 get_preserving()

```
constexpr storage_type gko::precision_reduction::get_preserving ( ) const [inline], [constexpr],
[noexcept]
```

Returns the number of preserving conversions in the encoding.

Returns

the number of preserving conversions in the encoding.

43.162.3.5 operator storage_type()

```
constexpr gko::precision_reduction::operator storage_type ( ) const [inline], [constexpr],
[noexcept]
```

Extracts the raw data of the encoding.

Returns

the raw data of the encoding

The documentation for this class was generated from the following file:

· ginkgo/core/base/types.hpp

43.163 gko::Preconditionable Class Reference

A LinOp implementing this interface can be preconditioned.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual std::shared_ptr< const LinOp > get_preconditioner () const Returns the preconditioner operator used by the Preconditionable.
- virtual void set_preconditioner (std::shared_ptr< const LinOp > new_precond)
 Sets the preconditioner operator used by the Preconditionable.

43.163.1 Detailed Description

A LinOp implementing this interface can be preconditioned.

43.163.2 Member Function Documentation

43.163.2.1 get_preconditioner()

```
virtual std::shared_ptr<const LinOp> gko::Preconditionable::get_preconditioner ( ) const
[inline], [virtual]
```

Returns the preconditioner operator used by the Preconditionable.

Returns

the preconditioner operator used by the Preconditionable

43.163.2.2 set_preconditioner()

Sets the preconditioner operator used by the Preconditionable.

Parameters

new_precond	the new preconditioner operator used by the Preconditionable	
-------------	--	--

The documentation for this class was generated from the following file:

· ginkgo/core/base/lin_op.hpp

43.164 gko::log::ProfilerHook Class Reference

This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Classes

· class NestedSummaryWriter

Recieves the results from ProfilerHook::create_nested_summary().

· class SummaryWriter

Recieves the results from ProfilerHook::create_summary().

class TableSummaryWriter

Writes the results from ProfilerHook::create_summary() and ProfilerHook::create_nested_summary() to a ASCII table in Markdown format.

Public Member Functions

void set_object_name (ptr_param< const PolymorphicObject > obj, std::string name)

Sets the name for an object to be profiled.

void set_synchronization (bool synchronize)

Should the events call executor->synchronize on operations and copy/allocation? This leads to a certain overhead, but makes the execution timeline of kernels synchronous.

profiling_scope_guard user_range (const char *name) const

Creates a scope guard for a user-defined range to be included in the profile.

Static Public Member Functions

static std::shared ptr< ProfilerHook > create tau (bool initialize=true)

Creates a logger annotating Ginkgo events with TAU ranges via PerfStubs.

static std::shared_ptr< ProfilerHook > create_vtune ()

Creates a logger annotating Ginkgo events with VTune ITT ranges.

static std::shared ptr< ProfilerHook > create nvtx (uint32 color argb=color yellow argb)

Creates a logger annotating Ginkgo events with NVTX ranges for CUDA.

static std::shared_ptr< ProfilerHook > create_roctx ()

Creates a logger annotating Ginkgo events with ROCTX ranges for HIP.

static std::shared ptr< ProfilerHook > create for executor (std::shared ptr< const Executor > exec)

Creates a logger annotating Ginkgo events with the most suitable backend for the given executor: NVTX for NSight Systems in CUDA, ROCTX for rocprof in HIP, TAU for everything else.

static std::shared_ptr< ProfilerHook > create_summary (std::shared_ptr< Timer > timer=std::make_←
 shared< CpuTimer > (), std::unique_ptr< SummaryWriter > writer=std::make_unique< TableSummaryWriter
 >(), bool debug check nesting=false)

Creates a logger measuring the runtime of Ginkgo events and printing a summary when it is destroyed.

Creates a logger measuring the runtime of Ginkgo events in a nested fashion and printing a summary when it is destroyed.

• static std::shared_ptr< ProfilerHook > create_custom (hook_function begin, hook_function end)

Creates a logger annotating Ginkgo events with a custom set of functions for range begin and end.

Static Public Attributes

constexpr static uint32 color_yellow_argb = 0xFFFCB05U
 The Ginkgo yellow background color as packed 32 bit ARGB value.

43.164.1 Detailed Description

This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.

It currently supports TAU, VTune, NSightSystems (NVTX) and rocPROF(ROCTX) and custom profiler hooks.

The Logger should be attached to the Executor that is being used to run the application for a full, program-wide annotation, or to individual objects to only highlight events caused directly by them (not operations and memory allocations though)

43.164.2 Member Function Documentation

43.164.2.1 create_nested_summary()

Creates a logger measuring the runtime of Ginkgo events in a nested fashion and printing a summary when it is destroyed.

Parameters

timer	The timer used to record time points.
writer	The NestedSummaryWriter to receive the performance results.
debug_check_nesting	Enable this flag if the output looks like it might contain incorrect nesting. This increases the overhead slightly, but recognizes mismatching push/pop pairs on the range stack.

Note

For this logger to provide reliable GPU timings, either use Timer::create_for_executor or enable synchronization via set_synchronization (true).

43.164.2.2 create_nvtx()

Creates a logger annotating Ginkgo events with NVTX ranges for CUDA.

Parameters

color_argb	The color of the NVTX ranges in the NSight Systems output. It has to be a 32 bit packed ARGB
	value.

43.164.2.3 create_summary()

Creates a logger measuring the runtime of Ginkgo events and printing a summary when it is destroyed.

Parameters

timer	The timer used to record time points.
writer	The SummaryWriter to receive the performance results.
debug_check_nesting	Enable this flag if the output looks like it might contain incorrect nesting. This increases
	the overhead slightly, but recognizes mismatching push/pop pairs on the range stack.

Note

For this logger to provide reliable GPU timings, either use $Timer::create_for_executor$ or enable synchronization via $set_synchronization$ (true).

43.164.2.4 create_tau()

Creates a logger annotating Ginkgo events with TAU ranges via PerfStubs.

Parameters

initialize	Should we call TAU's initialization and finalization functions, or does the application take care of it? The
	initialization will happen immediately, the finalization at program exit.

43.164.2.5 set_object_name()

Sets the name for an object to be profiled.

Every instance of that object in the profile will be replaced by the name instead of its runtime type.

Parameters

obj	the object
name	its name

43.164.2.6 user_range()

Creates a scope guard for a user-defined range to be included in the profile.

Parameters

name the name of the range

Returns

the scope guard. It will begin a range immediately and end it at the end of its scope.

The documentation for this class was generated from the following file:

• ginkgo/core/log/profiler_hook.hpp

43.165 gko::log::profiling scope guard Class Reference

Scope guard that annotates its scope with the provided profiler hooks.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

• profiling scope guard ()

Creates an empty (moved-from) scope guard.

• profiling_scope_guard (const char *name, profile_event_category category, ProfilerHook::hook_function begin, ProfilerHook::hook_function end)

Creates the scope guard.

~profiling_scope_guard ()

Calls the range end function if the scope guard was not moved from.

profiling_scope_guard (profiling_scope_guard &&other)

Move-constructs from another scope guard, other will be left empty.

43.165.1 Detailed Description

Scope guard that annotates its scope with the provided profiler hooks.

43.165.2 Constructor & Destructor Documentation

43.165.2.1 profiling_scope_guard()

Creates the scope guard.

Parameters

name	the name of the profiler range
category	the category of the profiler range
begin	the hook function to begin a range
end	the hook function to end a range

The documentation for this class was generated from the following file:

ginkgo/core/log/profiler_hook.hpp

43.166 gko::ptr param< T > Class Template Reference

This class is used for function parameters in the place of raw pointers.

```
#include <ginkgo/core/base/utils_helper.hpp>
```

Public Member Functions

```
    ptr param (T *ptr)
```

Initializes the ptr_param from a raw pointer.

• template<typename U , std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr> ptr_param (const std::shared_ptr< U > &ptr)

Initializes the ptr_param from a shared_ptr.

template<typename U , typename Deleter , std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr> ptr_param (const std::unique_ptr< U, Deleter > &ptr)

Initializes the ptr_param from a unique_ptr.

• template<typename U , std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr> ptr_param (const ptr_param< U > &ptr)

Initializes the ptr_param from a ptr_param of a derived type.

- T & operator* () const
- T * operator-> () const
- T * get () const
- operator bool () const

43.166.1 Detailed Description

```
template<typename T> class gko::ptr_param< T>
```

This class is used for function parameters in the place of raw pointers.

Pointer parameters should be used for everything that does not involve transfer of ownership. It can be converted to from raw pointers, shared pointers and unique pointers of the specified type or any derived type. This allows functions to be called without having to use gko::lend or calling .get() for every pointer argument. It probably has no use outside of function parameters, as it is immutable.

Template Parameters

```
T the pointed-to type
```

43.166.2 Member Function Documentation

43.166.2.1 get()

```
template<typename T>
T* gko::ptr_param< T >::get ( ) const [inline]
```

Returns

the underlying pointer.

Referenced by gko::matrix::Coo< ValueType, IndexType >::apply2(), gko::as(), gko::Executor::copy_from(), gko \leftarrow ::matrix::Diagonal< ValueType >::inverse_apply(), gko::PolymorphicObject::move_from(), gko::ptr_param< T > \leftarrow ::ptr_param(), and gko::matrix::Diagonal< ValueType >::rapply().

43.166.2.2 operator bool()

```
template<typename T>
gko::ptr_param< T >::operator bool ( ) const [inline], [explicit]
```

Returns

true iff the underlying pointer is non-null.

43.166.2.3 operator*()

```
template<typename T>
T& gko::ptr_param< T >::operator* ( ) const [inline]
```

Returns

a reference to the underlying pointee.

43.166.2.4 operator->()

```
template<typename T>
T* gko::ptr_param< T >::operator-> ( ) const [inline]
```

Returns

the underlying pointer.

The documentation for this class was generated from the following file:

• ginkgo/core/base/utils_helper.hpp

43.167 gko::syn::range< Start, End, Step > Struct Template Reference

range records start, end, step in template

```
#include <ginkgo/core/synthesizer/containers.hpp>
```

43.167.1 Detailed Description

```
template<int Start, int End, int Step = 1>
struct gko::syn::range< Start, End, Step >
```

range records start, end, step in template

Template Parameters

Start	start of range
End	end of range
Step	step of range. default is 1

The documentation for this struct was generated from the following file:

• ginkgo/core/synthesizer/containers.hpp

43.168 gko::range< Accessor > Class Template Reference

A range is a multidimensional view of the memory.

#include <ginkgo/core/base/range.hpp>

Public Types

• using accessor = Accessor

The type of the underlying accessor.

Public Member Functions

∼range ()=default

Use the default destructor.

template<typename... AccessorParams>
 constexpr range (AccessorParams &&... params)

Creates a new range.

template<typename... DimensionTypes>
 constexpr auto operator() (DimensionTypes &&... dimensions) const -> decltype(std::declval< accessor >()(std::forward< DimensionTypes >(dimensions)...))

Returns a value (or a sub-range) with the specified indexes.

- template<typename OtherAccessor >
 const range & operator= (const range< OtherAccessor > &other) const
- const range & operator= (const range & other) const

Assigns another range to this range.

• constexpr size_type length (size_type dimension) const

Returns the length of the specified dimension of the range.

constexpr const accessor * operator-> () const noexcept

Returns a pointer to the accessor.

• constexpr const accessor & get_accessor () const noexcept

`Returns a reference to the accessor.

Static Public Attributes

• static constexpr size_type dimensionality = accessor::dimensionality

The number of dimensions of the range.

43.168.1 Detailed Description

template<typename Accessor> class gko::range< Accessor>

A range is a multidimensional view of the memory.

The range does not store any of its values by itself. Instead, it obtains the values through an accessor (e.g. accessor::row major) which describes how the indexes of the range map to physical locations in memory.

There are several advantages of using ranges instead of plain memory pointers:

- 1. Code using ranges is easier to read and write, as there is no need for index linearizations.
- 2. Code using ranges is safer, as it is impossible to accidentally miscalculate an index or step out of bounds, since range accessors perform bounds checking in debug builds. For performance, this can be disabled in release builds by defining the NDEBUG flag.
- Ranges enable generalized code, as algorithms can be written independent of the memory layout. This does not impede various optimizations based on memory layout, as it is always possible to specialize algorithms for ranges with specific memory layouts.
- 4. Ranges have various pointwise operations predefined, which reduces the amount of loops that need to be written.

43.168.1.1 Range operations

Ranges define a complete set of pointwise unary and binary operators which extend the basic arithmetic operators in C++, as well as a few pointwise operations and mathematical functions useful in ginkgo, and a couple of non-pointwise operations. Compound assignment (+=, *=, etc.) is not yet supported at this moment. Here is a complete list of operations:

- standard unary operations: $+, -, !, \sim$
- standard binary operations: +, * (this is pointwise, not matrix multiplication), /, %, <, >, <=, >=, ==, !=, ||, & &, |, &, ^, <<, >>
- useful unary functions: zero, one, abs, real, imag, conj, squared_norm
- useful binary functions: min, max

All binary pointwise operations also work as expected if one of the operands is a scalar and the other is a range. The scalar operand will have the effect as if it was a range of the same size as the other operand, filled with the specified scalar.

Two "global" functions transpose and mmul are also supported. transpose transposes the first two dimensions of the range (i.e. transpose (r) (i, j, ...) == r(j, i, ...)). mmul performs a (batched) matrix multiply of the ranges - the first two dimensions represent the matrices, while the rest represent the batch. For example, given the ranges r1 and r2 of dimensions (3, 2, 3) and (2, 4, 3), respectively, mmul (r1, r2) will return a range of dimensions (3, 4, 3), obtained by multiplying the 3 frontal slices of the range, and stacking the result back vertically.

43.168.1.2 Compound operations

Multiple range operations can be combined into a single expression. For example, an "axpy" operation can be obtained using y = alpha * x + y, where x an y are ranges, and alpha is a scalar. Range operations are optimized for memory access, and the above code does not allocate additional storage for intermediate ranges alpha * x or aplha * x + y. In fact, the entire computation is done during the assignment, and the results of operations + and * only register the data, and the types of operations that will be computed once the results are needed.

It is possible to store and reuse these intermediate expressions. The following example will overwrite the range \mathbf{x} with it's 4th power:

```
{c++} auto square = x * x; // this is range constructor, not range assignment! x = \text{square}; // overwrites x with x*x (this is range assignment) x = \text{square}; // overwrites new x (x*x) with (x*x)*(x*x) (as is this)
```

43.168.1.3 Caveats

_mmul is not a highly-optimized BLAS-3 version of the matrix multiplication.__ The current design of ranges and accessors prevents that, so if you need a high-perfromance matrix multiplication, you should use one of the libraries that provide that, or implement your own (you can use pointwise range operations to help simplify that). However, range design might get improved in the future to allow efficient implementations of BLAS-3 kernels.

Aliasing the result range in mmul and transpose is not allowed. Constructs like A = transpose(A), A = mmul(A, A), or A = mmul(A, A) + C lead to undefined behavior. However, aliasing input arguments is allowed: C = mmul(A, A), and even C = mmul(A, A) + C is valid code (in the last example, only pointwise operations are aliased). C = mmul(A, A + C) is not valid though.

43.168.1.4 Examples

The range unit tests in core/test/base/range.cpp contain lots of simple 1-line examples of range operations. The accessor unit tests in core/test/base/range.cpp show how to use ranges with concrete accessors, and how to use range slices using spans as arguments to range function call operator. Finally, examples/range contains a complete example where ranges are used to implement a simple version of the right-looking LU factorization.

Template Parameters

```
Accessor underlying accessor of the range
```

43.168.2 Constructor & Destructor Documentation

43.168.2.1 range()

Creates a new range.

Template Parameters

Parameters

43.168.3 Member Function Documentation

43.168.3.1 get_accessor()

```
template<typename Accessor>
constexpr const accessor& gko::range< Accessor >::get_accessor ( ) const [inline], [constexpr],
[noexcept]
```

`Returns a reference to the accessor.

Returns

reference to the accessor

Referenced by gko::range< Accessor >::operator=().

43.168.3.2 length()

Returns the length of the specified dimension of the range.

Parameters

dimension	the dimensions whose length is returned

Returns

the length of the ${\tt dimension}\text{-}{th}$ dimension of the range

Referenced by gko::matrix_data< ValueType, IndexType >::matrix_data().

43.168.3.3 operator()()

Returns a value (or a sub-range) with the specified indexes.

Template Parameters

DimensionTypes	The types of indexes. Supported types depend on the underlying accessor, but are usually	
	either integer types or spans. If at least one index is a span, the returned value will be a	
	sub-range.	

Parameters

dimensions	the indexes of the values.
------------	----------------------------

Returns

```
a value on position (dimensions...).
```

References gko::range < Accessor >::dimensionality.

43.168.3.4 operator->()

```
template<typename Accessor>
constexpr const accessor* gko::range< Accessor >::operator-> ( ) const [inline], [constexpr],
[noexcept]
```

Returns a pointer to the accessor.

Can be used to access data and functions of a specific accessor.

Returns

pointer to the accessor

43.168.3.5 operator=() [1/2]

Assigns another range to this range.

The order of assignment is defined by the accessor of this range, thus the memory access will be optimized for the resulting range, and not for the other range. If the sizes of two ranges do not match, the result is undefined. Sizes of the ranges are checked at runtime in debug builds.

Note

Temporary accessors are allowed to define the implementation of the assignment as deleted, so do not expect r1 * r2 = r2 to work.

Parameters

other	the range to copy the data from
-------	---------------------------------

References gko::range < Accessor >::get_accessor().

43.168.3.6 operator=() [2/2]

This is a version of the function which allows to copy between ranges of different accessors.

Template Parameters

OtherAccessor accessor of the ot	her range
------------------------------------	-----------

The documentation for this class was generated from the following file:

· ginkgo/core/base/range.hpp

43.169 gko::reorder::Rcm< ValueType, IndexType > Class Template Reference

Rcm is a reordering algorithm minimizing the bandwidth of a matrix.

```
#include <ginkgo/core/reorder/rcm.hpp>
```

Public Member Functions

- std::shared_ptr< const PermutationMatrix > get_permutation () const
 Gets the permutation (permutation matrix, output of the algorithm) of the linear operator.
- std::shared_ptr< const PermutationMatrix > get_inverse_permutation () const
 Gets the inverse permutation (permutation matrix, output of the algorithm) of the linear operator.

43.169.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::reorder::Rcm< ValueType, IndexType >
```

Rcm is a reordering algorithm minimizing the bandwidth of a matrix.

Such a reordering typically also significantly reduces fill-in, though usually not as effective as more complex algorithms, specifically AMD and nested dissection schemes. The advantage of this algorithm is its low runtime.

Note

This class is derived from polymorphic object but is not a LinOp as it does not make sense for this class to implement the apply methods. The objective of this class is to generate a reordering/permutation vector (in the form of the Permutation matrix), which can be used to apply to reorder a matrix as required.

There are two "starting strategies" currently available: minimum degree and pseudo-peripheral. These strategies control how a starting vertex for a connected component is choosen, which is then renumbered as first vertex in the component, starting the algorithm from there. In general, the bandwidths obtained by choosing a pseudo-peripheral vertex are slightly smaller than those obtained from choosing a vertex of minimum degree. On the other hand, this strategy is much more expensive, relatively. The algorithm for finding a pseudo-peripheral vertex as described in "Computer Solution of Sparse Linear Systems" (George, Liu, Ng, Oak Ridge National Laboratory, 1994) is implemented here.

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

43.169.2 Member Function Documentation

43.169.2.1 get_inverse_permutation()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const PermutationMatrix> gko::reorder::Rcm< ValueType, IndexType >::get_
inverse_permutation ( ) const [inline]
```

Gets the inverse permutation (permutation matrix, output of the algorithm) of the linear operator.

Returns

the inverse permutation (permutation matrix)

43.169.2.2 get permutation()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const PermutationMatrix> gko::reorder::Rcm< ValueType, IndexType >::get_
permutation ( ) const [inline]
```

Gets the permutation (permutation matrix, output of the algorithm) of the linear operator.

Returns

the permutation (permutation matrix)

The documentation for this class was generated from the following file:

ginkgo/core/reorder/rcm.hpp

43.170 gko::ReadableFromMatrixData< ValueType, IndexType > Class Template Reference

A LinOp implementing this interface can read its data from a matrix_data structure.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual void read (const matrix_data < ValueType, IndexType > &data)=0
 Reads a matrix from a matrix_data structure.
- void read (const matrix_assembly_data < ValueType, IndexType > &data)
 Reads a matrix from a matrix_assembly_data structure.
- virtual void read (const ${\tt device_matrix_data} < {\tt ValueType}, \\ {\tt IndexType} > {\tt \&data})$

Reads a matrix from a device_matrix_data structure.

virtual void read (device_matrix_data< ValueType, IndexType > &&data)

Reads a matrix from a device_matrix_data structure.

43.170.1 Detailed Description

```
template < typename ValueType, typename IndexType > class gko::ReadableFromMatrixData < ValueType, IndexType >
```

A LinOp implementing this interface can read its data from a matrix_data structure.

43.170.2 Member Function Documentation

43.170.2.1 read() [1/4]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Hybrid< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Ell< ValueType, IndexType >, gko::matrix::Coo< ValueType, IndexType > gko::matrix::Sellp< ValueType, IndexType >, and gko::matrix::SparsityCsr< ValueType, IndexType >.

43.170.2.2 read() [2/4]

Reads a matrix from a matrix_assembly_data structure.

Parameters

```
data the matrix_assembly_data structure
```

43.170.2.3 read() [3/4]

Reads a matrix from a matrix data structure.

Parameters

```
data the matrix_data structure
```

Implemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Hybrid< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Ell< ValueType, IndexType >, gko::matrix::Coo< ValueType, IndexType > gko::matrix::Sellp< ValueType, IndexType >, and gko::matrix::SparsityCsr< ValueType, IndexType >.

43.170.2.4 read() [4/4]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Hybrid< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Ell< ValueType, IndexType >, gko::matrix::Coo< ValueType, IndexType > gko::matrix::Sellp< ValueType, IndexType >, and gko::matrix::SparsityCsr< ValueType, IndexType >.

The documentation for this class was generated from the following file:

ginkgo/core/base/lin_op.hpp

43.171 gko::log::Record Class Reference

Record is a Logger which logs every event to an object.

```
#include <ginkgo/core/log/record.hpp>
```

Classes

struct logged data

Struct storing the actually logged data.

Public Member Functions

const logged_data & get () const noexcept
 Returns the logged data.

· logged_data & get () noexcept

Static Public Member Functions

 static std::unique_ptr< Record > create (std::shared_ptr< const Executor > exec, const mask_type &enabled_events=Logger::all_events_mask, size_type max_storage=1)

Creates a Record logger.

static std::unique_ptr< Record > create (const mask_type &enabled_events=Logger::all_events_mask, size_type max_storage=1)

Creates a Record logger.

43.171.1 Detailed Description

Record is a Logger which logs every event to an object.

The object can then be accessed at any time by asking the logger to return it.

Note

Please note that this logger can have significant memory and performance overhead. In particular, when logging events such as the check events, all parameters are cloned. If it is sufficient to clone one parameter, consider implementing a specific logger for this. In addition, it is advised to tune the history size in order to control memory overhead.

43.171.2 Member Function Documentation

43.171.2.1 create() [1/2]

Creates a Record logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

exec	the executor	
enabled_events	the events enabled for this logger. By default all events.	
max_storage	the size of storage (i.e. history) wanted by the user. By default 0 is used, which means unlimited storage. It is advised to control this to reduce memory overhead of this logger.	

Returns

an std::unique_ptr to the the constructed object

43.171.2.2 create() [2/2]

Creates a Record logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

exec	the executor	
enabled_events	the events enabled for this logger. By default all events.	
max_storage	the size of storage (i.e. history) wanted by the user. By default 0 is used, which means unlimited storage. It is advised to control this to reduce memory overhead of this logger.	

Returns

an std::unique_ptr to the the constructed object

43.171.2.3 get() [1/2]

```
const logged_data& gko::log::Record::get ( ) const [inline], [noexcept]
```

Returns the logged data.

Returns

the logged data

43.171.2.4 get() [2/2]

```
logged_data& gko::log::Record::get ( ) [inline], [noexcept]
```

The documentation for this class was generated from the following file:

• ginkgo/core/log/record.hpp

43.172 gko::ReferenceExecutor Class Reference

This is a specialization of the OmpExecutor, which runs the reference implementations of the kernels used for debugging purposes.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

void run (const Operation &op) const override
 Runs the specified Operation using this Executor.

Additional Inherited Members

43.172.1 Detailed Description

This is a specialization of the OmpExecutor, which runs the reference implementations of the kernels used for debugging purposes.

43.172.2 Member Function Documentation

43.172.2.1 run()

Runs the specified Operation using this Executor.

Parameters

```
op the operation to run
```

Implements gko::Executor.

The documentation for this class was generated from the following file:

ginkgo/core/base/executor.hpp

43.173 gko::stop::RelativeResidualNorm< ValueType > Class Template Reference

The RelativeResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

43.173.1 Detailed Description

```
template < typename ValueType = default_precision > class gko::stop::RelativeResidualNorm < ValueType >
```

The RelativeResidualNorm class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

when norm(residual) / norm(right_hand_side) < threshold. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on b in order to compute the norm of the right-hand side. If this is not correctly provided, an exception ::gko::NotSupported() is thrown.

The documentation for this class was generated from the following file:

• ginkgo/core/stop/residual_norm.hpp

43.174 gko::reorder::ReorderingBase< IndexType > Class Template Reference

The ReorderingBase class is a base class for all the reordering algorithms.

```
#include <ginkgo/core/reorder/reordering_base.hpp>
```

Additional Inherited Members

43.174.1 Detailed Description

```
template<typename IndexType = int32> class gko::reorder::ReorderingBase< IndexType >
```

The ReorderingBase class is a base class for all the reordering algorithms.

It contains a factory to instantiate the reorderings. It is up to each specific reordering to decide what to do with the data that is passed to it.

The documentation for this class was generated from the following file:

• ginkgo/core/reorder/reordering_base.hpp

43.175 gko::reorder::ReorderingBaseArgs Struct Reference

This struct is used to pass parameters to the EnableDefaultReorderingBaseFactory::generate() method.

#include <ginkgo/core/reorder/reordering_base.hpp>

43.175.1 Detailed Description

This struct is used to pass parameters to the EnableDefaultReorderingBaseFactory::generate() method.

It is the ComponentsType of ReorderingBaseFactory.

The documentation for this struct was generated from the following file:

ginkgo/core/reorder/reordering base.hpp

43.176 gko::experimental::mpi::request Class Reference

The request class is a light, move-only wrapper around the MPI_Request handle.

#include <ginkgo/core/base/mpi.hpp>

Public Member Functions

• request ()

The default constructor.

• MPI_Request * get ()

Get a pointer to the underlying MPI_Request handle.

• status wait ()

Allows a rank to wait on a particular request handle.

43.176.1 Detailed Description

The request class is a light, move-only wrapper around the MPI_Request handle.

43.176.2 Constructor & Destructor Documentation

43.176.2.1 request()

gko::experimental::mpi::request::request () [inline]

The default constructor.

It creates a null MPI_Request of MPI_REQUEST_NULL type.

43.176.3 Member Function Documentation

43.176.3.1 get()

```
MPI_Request* gko::experimental::mpi::request::get ( ) [inline]
```

Get a pointer to the underlying MPI Request handle.

Returns

a pointer to MPI_Request handle

Referenced by gko::experimental::mpi::communicator::i_all_gather(), gko::experimental::mpi::communicator::i \leftarrow _all_reduce(), gko::experimental::mpi::communicator::i_all_to_all(), gko::experimental::mpi::communicator::i \leftarrow _all_to_all_v(), gko::experimental::mpi::communicator::i_broadcast(), gko::experimental::mpi::communicator::i \leftarrow _recv(), gko::experimental::mpi::communicator::i_gather_v(), gko::experimental::mpi::communicator::i_scan(), gko::experimental::mpi::communicator::i_scan(), gko::experimental::mpi::communicator::i_scan(), gko::experimental::mpi::communicator::i_scanter_v(), gko::experimental::mpi::communicator::i_scanter_v(), gko::experimental::mpi::communicator::i_scanter_v(), gko::experimental::mpi::window< ValueType >::r_accumulate(), gko::experimental::mpi::window< ValueType >::r_get_ \leftarrow accumulate(), and gko::experimental::mpi::window< ValueType >::r_get_ \leftarrow accumulate(), and gko::experimental::mpi::window< ValueType >::r_get_ \leftarrow accumulate(), and gko::experimental::mpi::window< ValueType >::r_get_ \leftarrow accumulate().

43.176.3.2 wait()

```
status gko::experimental::mpi::request::wait ( ) [inline]
```

Allows a rank to wait on a particular request handle.

Parameters

req	The request to wait on.
status	The status variable that can be queried.

References gko::experimental::mpi::status::get().

The documentation for this class was generated from the following file:

ginkgo/core/base/mpi.hpp

43.177 gko::stop::ResidualNorm< ValueType > Class Template Reference

The ResidualNorm class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

#include <ginkgo/core/stop/residual_norm.hpp>

43.177.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::stop::ResidualNorm< ValueType >
```

The ResidualNorm class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

- 1. the norm of the right-hand side, norm(residual) / norm(right_hand_side) < threshold
- 2. the initial residual, norm(residual) / norm(initial_residual) < threshold.
- 3. one, norm(residual) < threshold.

For better performance, the checks are run on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on either b or the initial_residual in order to compute their norms. If this is not correctly provided, an exception ::gko :: NotSupported() is thrown.

The documentation for this class was generated from the following file:

ginkgo/core/stop/residual_norm.hpp

43.178 gko::stop::ResidualNormBase< ValueType > Class Template Reference

The ResidualNormBase class provides a framework for stopping criteria related to the residual norm.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

43.178.1 Detailed Description

```
\label{template} \mbox{template} < \mbox{typename ValueType} > \\ \mbox{class gko::stop::ResidualNormBase} < \mbox{ValueType} > \\ \mbox{}
```

The ResidualNormBase class provides a framework for stopping criteria related to the residual norm.

These criteria differ in the way they initialize starting_tau_, so in the value they compare the residual norm against. The provided check_impl uses the actual residual to check for convergence.

The documentation for this class was generated from the following file:

ginkgo/core/stop/residual_norm.hpp

43.179 gko::stop::ResidualNormReduction < ValueType > Class Template Reference

The ResidualNormReduction class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

#include <ginkgo/core/stop/residual_norm.hpp>

43.179.1 Detailed Description

template<typename ValueType = default_precision> class gko::stop::ResidualNormReduction< ValueType >

The ResidualNormReduction class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

when $norm(residual) / norm(initial_residual) < threshold. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.$

Note

To use this stopping criterion there are some dependencies. The constructor depends on $initial_{\leftarrow}$ residual in order to compute the first relative residual norm. The check method depends on either the residual_norm or the residual being set. When any of those is not correctly provided, an exception ::gko::NotSupported() is thrown.

The documentation for this class was generated from the following file:

ginkgo/core/stop/residual_norm.hpp

43.180 gko::accessor::row_major< ValueType, Dimensionality > Class Template Reference

A row_major accessor is a bridge between a range and the row-major memory layout.

#include <ginkgo/core/base/range_accessors.hpp>

Public Types

using value_type = ValueType
 Type of values returned by the accessor.

using data_type = value_type *

Type of underlying data storage.

Public Member Functions

constexpr value_type & operator() (size_type row, size_type col) const

Returns the data element at position (row, col)

constexpr range < row_major > operator() (const span &rows, const span &cols) const

Returns the sub-range spanning the range (rows, cols)

• constexpr size_type length (size_type dimension) const

Returns the length in dimension dimension.

 $\bullet \ \ \text{template}{<} \text{typename OtherAccessor} >$

void copy_from (const OtherAccessor &other) const

Copies data from another accessor.

Public Attributes

· const data type data

Reference to the underlying data.

const std::array< const size_type, dimensionality > lengths

An array of dimension sizes.

· const size_type stride

Distance between consecutive rows.

Static Public Attributes

static constexpr size_type dimensionality = 2

Number of dimensions of the accessor.

43.180.1 Detailed Description

template<typename ValueType, size_type Dimensionality> class gko::accessor::row_major< ValueType, Dimensionality >

A row_major accessor is a bridge between a range and the row-major memory layout.

You should never try to explicitly create an instance of this accessor. Instead, supply it as a template parameter to a range, and pass the constructor parameters for this class to the range (it will forward it to this class).

Warning

The current implementation is incomplete, and only allows for 2-dimensional ranges.

Template Parameters

ValueType	e type of values this accessor returns	
Dimensionality	number of dimensions of this accessor (has to be 2)	

43.180.2 Member Function Documentation

43.180.2.1 copy_from()

Copies data from another accessor.

Warning

Do not use this function since it is not optimized for a specific executor. It will always be performed sequentially. Please write an optimized version (adjusted to the architecture) by iterating through the values yourself.

Template Parameters

OtherAccessor	type of the other accessor
---------------	----------------------------

Parameters

other	other accessor
-------	----------------

References gko::accessor::row_major< ValueType, Dimensionality >::lengths.

43.180.2.2 length()

Returns the length in dimension dimension.

Parameters

dimension a dimension index

Returns

length in dimension dimension

References gko::accessor::row_major< ValueType, Dimensionality >::lengths.

43.180.2.3 operator()() [1/2]

Returns the sub-range spanning the range (rows, cols)

Parameters

rows	row span
cols	column span

Returns

sub-range spanning the range (rows, cols)

References gko::span::begin, gko::accessor::row_major< ValueType, Dimensionality >::data, gko::span::end, gko::span::is_valid(), gko::accessor::row_major< ValueType, Dimensionality >::lengths, and gko::accessor::row—major< ValueType, Dimensionality >::stride.

43.180.2.4 operator()() [2/2]

Returns the data element at position (row, col)

Parameters

row	row index
col	column index

Returns

data element at (row, col)

References gko::accessor::row_major< ValueType, Dimensionality >::data, gko::accessor::row_major< Value Type, Dimensionality >::lengths, and gko::accessor::row_major< ValueType, Dimensionality >::stride.

The documentation for this class was generated from the following file:

• ginkgo/core/base/range_accessors.hpp

43.181 gko::matrix::RowGatherer< IndexType > Class Template Reference

RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.

#include <ginkgo/core/matrix/row_gatherer.hpp>

Public Member Functions

- index_type * get_row_idxs () noexcept
 - Returns a pointer to the row index array for gathering.
- const index_type * get_const_row_idxs () const noexcept

Returns a pointer to the row index array for gathering.

Static Public Member Functions

static std::unique_ptr< const RowGatherer > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< IndexType > &&row_idxs)

Creates a constant (immutable) RowGatherer matrix from a constant array.

43.181.1 Detailed Description

template<typename IndexType = int32> class gko::matrix::RowGatherer< IndexType >

RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.

Template Parameters

IndexType precision of rowgatherer array indices.

Note

This format is used mainly to allow for an abstraction of the rowgatherer and provides the user with an apply method which calls the respective Dense rowgatherer operation. As such it only stores an array of the rowgatherer indices.

43.181.2 Member Function Documentation

43.181.2.1 create_const()

Creates a constant (immutable) RowGatherer matrix from a constant array.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
row_idxs	the gathered row indices of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

43.181.2.2 get_const_row_idxs()

```
template<typename IndexType = int32>
const index_type* gko::matrix::RowGatherer< IndexType >::get_const_row_idxs ( ) const [inline],
[noexcept]
```

Returns a pointer to the row index array for gathering.

Returns

the pointer to the row index array for gathering.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array < ValueType >::get_const_data().

43.181.2.3 get_row_idxs()

```
template<typename IndexType = int32>
index_type* gko::matrix::RowGatherer< IndexType >::get_row_idxs ( ) [inline], [noexcept]
```

Returns a pointer to the row index array for gathering.

Returns

the pointer to the row index array for gathering.

References gko::array< ValueType >::get_data().

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/row_gatherer.hpp

43.182 gko::ScaledIdentityAddable Class Reference

Adds the operation M < a I + b M for matrix M, identity operator I and scalars a and b, where M is the calling object.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

void add_scaled_identity (ptr_param< const LinOp > const a, ptr_param< const LinOp > const b)
 Scales this and adds another scalar times the identity to it.

43.182.1 Detailed Description

Adds the operation M <- a I + b M for matrix M, identity operator I and scalars a and b, where M is the calling object.

43.182.2 Member Function Documentation

43.182.2.1 add_scaled_identity()

Scales this and adds another scalar times the identity to it.

Parameters

а	Scalar to multiply the identity operator before adding.
b	Scalar to multiply this before adding the scaled identity to it.

References gko::make_temporary_clone().

The documentation for this class was generated from the following file:

· ginkgo/core/base/lin op.hpp

43.183 gko::experimental::reorder::ScaledReordered< ValueType, IndexType > Class Template Reference

Provides an interface to wrap reorderings like Rcm and diagonal scaling like equilibration around a LinOp like e.g.

#include <ginkgo/core/reorder/scaled_reordered.hpp>

Additional Inherited Members

43.183.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::experimental::reorder::ScaledReordered< ValueType, IndexType >

Provides an interface to wrap reorderings like Rcm and diagonal scaling like equilibration around a LinOp like e.g.

a sparse direct solver.

Reorderings can be useful for reducing fill-in in the numerical factorization phase of direct solvers, diagonal scaling can help improve the numerical stability by reducing the condition number of the system matrix.

With a permutation matrix P, a row scaling R and a column scaling C, the inner operator is applied to the system matrix $P*R*A*C*P^{\wedge}T$ instead of A. Instead of A*x = b, the inner operator attempts to solve the equivalent linear system $P*R*A*C*P^{\wedge}T*y = P*R*b$ and retrieves the solution $x = C*P^{\wedge}T*y$. Note: The inner system matrix is computed from a clone of A, so the original system matrix is not changed.

Template Parameters

ValueType	Type of the values of all matrices used in this class
IndexType	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

• ginkgo/core/reorder/scaled_reordered.hpp

43.184 gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference

A Schwarz preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

#include <ginkgo/core/distributed/preconditioner/schwarz.hpp>

Additional Inherited Members

43.184.1 Detailed Description

template < typename ValueType = default_precision, typename LocalIndexType = int32, typename GlobalIndexType = int64 > class gko::experimental::distributed::preconditioner::Schwarz < ValueType, LocalIndexType, GlobalIndexType >

A Schwarz preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

See Iterative Methods for Sparse Linear Systems (Y. Saad) for a general treatment and variations of the method.

Note

Currently overlap and coarse grid correction are not supported (TODO).

Template Parameters

ValueType	precision of matrix elements
IndexType	integral type of the preconditioner

The documentation for this class was generated from the following file:

· ginkgo/core/distributed/preconditioner/schwarz.hpp

43.185 gko::scoped_device_id_guard Class Reference

This move-only class uses RAII to set the device id within a scoped block, if necessary.

#include <qinkqo/core/base/scoped_device_id_quard.hpp>

Public Member Functions

- scoped_device_id_guard (const ReferenceExecutor *exec, int device_id)

 Create a scoped device id from an Reference.
- scoped_device_id_guard (const OmpExecutor *exec, int device_id)

Create a scoped device id from an OmpExecutor.

• scoped_device_id_guard (const CudaExecutor *exec, int device_id)

Create a scoped device id from an CudaExecutor.

scoped_device_id_guard (const HipExecutor *exec, int device_id)

Create a scoped device id from an HipExecutor.

• scoped_device_id_guard (const DpcppExecutor *exec, int device_id)

Create a scoped device id from an DpcppExecutor.

43.185.1 Detailed Description

This move-only class uses RAII to set the device id within a scoped block, if necessary.

The class behaves similar to std::scoped_lock. The scoped guard will make sure that the device code is run on the correct device within one scoped block, when run with multiple devices. Depending on the executor it will record the current device id and set the device id to the one being passed in. After the scope has been exited, the destructor sets the device_id back to the one before entering the scope. The OmpExecutor and DpcppExecutor don't require setting an device id, so in those cases, the class is a no-op.

The device id scope has to be constructed from a executor with concrete type (not plain Executor) and a device id. Only the type of the executor object is relevant, so the pointer will not be accessed, and may even be a nullptr. From the executor type the correct derived class of detail::generic_scoped_device_id_guard is picked. The following illustrates the usage of this class:

```
scoped_device_id_guard g{static_cast<CudaExecutor>(nullptr), 1};
// now the device id is set to 1
}
// now the device id is reverted again
```

43.185.2 Constructor & Destructor Documentation

43.185.2.1 scoped_device_id_guard() [1/5]

Create a scoped device id from an Reference.

The resulting object will be a noop.

Parameters

exec	Not used.
device⊷	Not used.
_id	

43.185.2.2 scoped_device_id_guard() [2/5]

Create a scoped device id from an OmpExecutor.

The resulting object will be a noop.

Parameters

exec	Not used.
device⊷	Not used.
_id	

43.185.2.3 scoped_device_id_guard() [3/5]

Create a scoped device id from an CudaExecutor.

The resulting object will set the cuda device id accordingly.

Parameters

exec	Not used.
device←	The device id to use within the scope.
_id	

43.185.2.4 scoped_device_id_guard() [4/5]

Create a scoped device id from an HipExecutor.

The resulting object will set the hip device id accordingly.

Parameters

exec	Not used.
device⊷	The device id to use within the scope.
_id	

43.185.2.5 scoped_device_id_guard() [5/5]

Create a scoped device id from an DpcppExecutor.

The resulting object will be a noop.

Parameters

exec	Not used.
device⊷	Not used.
_id	

The documentation for this class was generated from the following file:

· ginkgo/core/base/scoped device id guard.hpp

43.186 gko::matrix::Sellp< ValueType, IndexType > Class Template Reference

SELL-P is a matrix format similar to ELL format.

```
#include <ginkgo/core/matrix/sellp.hpp>
```

Public Member Functions

void read (const mat_data &data) override

Reads a matrix from a matrix_data structure.

· void read (const device mat data &data) override

Reads a matrix from a device_matrix_data structure.

void read (device_mat_data &&data) override

Reads a matrix from a device_matrix_data structure.

· void write (mat_data &data) const override

Writes a matrix to a matrix_data structure.

std::unique_ptr< Diagonal< ValueType >> extract_diagonal () const override

Extracts the diagonal entries of the matrix into a vector.

• std::unique_ptr< absolute_type > compute_absolute () const override

Gets the AbsoluteLinOp.

void compute_absolute_inplace () override

Compute absolute inplace on each element.

value_type * get_values () noexcept

Returns the values of the matrix.

const value_type * get_const_values () const noexcept

Returns the values of the matrix.

index_type * get_col_idxs () noexcept

Returns the column indexes of the matrix.

const index_type * get_const_col_idxs () const noexcept

Returns the column indexes of the matrix.

size type * get slice lengths () noexcept

Returns the lengths(columns) of slices.

const size_type * get_const_slice_lengths () const noexcept

Returns the lengths(columns) of slices.

size_type * get_slice_sets () noexcept

Returns the offsets of slices.

const size_type * get_const_slice_sets () const noexcept

Returns the offsets of slices.

• size_type get_slice_size () const noexcept

Returns the size of a slice.

• size type get stride factor () const noexcept

Returns the stride factor(t) of SELL-P.

size_type get_total_cols () const noexcept

Returns the total column number.

size_type get_num_stored_elements () const noexcept

Returns the number of elements explicitly stored in the matrix.

• value_type & val_at (size_type row, size_type slice_set, size_type idx) noexcept

Returns the idx-th non-zero element of the row-th row with slice_set slice set.

value_type val_at (size_type row, size_type slice_set, size_type idx) const noexcept

Returns the idx-th non-zero element of the row-th row with slice_set slice set.

• index_type & col_at (size_type row, size_type slice_set, size_type idx) noexcept

Returns the idx-th column index of the row-th row with slice_set slice set.

index_type col_at (size_type row, size_type slice_set, size_type idx) const noexcept

Returns the idx-th column index of the row-th row with slice_set slice set.

Sellp & operator= (const Sellp &)

Copy-assigns a Sellp matrix.

Sellp & operator= (Sellp &&)

Move-assigns a Sellp matrix.

• Sellp (const Sellp &)

Copy-assigns a Sellp matrix.

Sellp (Sellp &&)

Move-assigns a Sellp matrix.

43.186.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Sellp< ValueType, IndexType >

SELL-P is a matrix format similar to ELL format.

The difference is that SELL-P format divides rows into smaller slices and store each slice with ELL format.

This implementation uses the column index value invalid_index<IndexType>() to mark padding entries that are not part of the sparsity pattern.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indexes

43.186.2 Constructor & Destructor Documentation

43.186.2.1 Sellp() [1/2]

Copy-assigns a Sellp matrix.

Inherits the executor, copies the data and parameters.

43.186.2.2 Sellp() [2/2]

Move-assigns a Sellp matrix.

Inherits the executor, moves the data and parameters. The moved-from object is empty (0x0 with valid slice_sets and unchanged parameters).

43.186.3 Member Function Documentation

43.186.3.1 col_at() [1/2]

Returns the idx-th column index of the row-th row with slice_set slice set.

Parameters

row	the row of the requested element in the slice
_	the slice set of the slice
Generated by Do	rygen idx-th stored element of the row in the slice

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

```
298 {
299     return this
300         ->get_const_col_idxs()[this->linearize_index(row, slice_set, idx)];
301 }
```

43.186.3.2 col_at() [2/2]

Returns the idx-th column index of the row-th row with slice_set slice set.

Parameters

row	the row of the requested element in the slice
slice_set	the slice set of the slice
idx	the idx-th stored element of the row in the slice

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References gko::matrix::Sellp< ValueType, IndexType >::get_col_idxs().

43.186.3.3 compute absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Sellp< ValueType, IndexType >::compute_absolute (
) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove complex< Sellp< ValueType, IndexType >>>.

43.186.3.4 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32> std::unique_ptr<Diagonal<ValueType> > gko::matrix::Sellp< ValueType, IndexType >::extract_← diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

diag the vector into which the diagonal will be written

Implements gko::DiagonalExtractable< ValueType >.

43.186.3.5 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Sellp< ValueType, IndexType >::get_col_idxs () [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References gko::array< ValueType >::get_data().

Referenced by gko::matrix::Sellp< ValueType, IndexType >::col_at().

43.186.3.6 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_col_idxs ( ) const
[inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array < ValueType >::get const data().

43.186.3.7 get_const_slice_lengths()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const size_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_lengths ( ) const
[inline], [noexcept]
```

Returns the lengths(columns) of slices.

Returns

the lengths(columns) of slices.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get const data().

43.186.3.8 get_const_slice_sets()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const size_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_sets ( ) const
[inline], [noexcept]
```

Returns the offsets of slices.

Returns

the offsets of slices.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.186.3.9 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.186.3.10 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Sellp< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.186.3.11 get_slice_lengths()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type* gko::matrix::Sellp< ValueType, IndexType >::get_slice_lengths () [inline], [noexcept]
```

Returns the lengths(columns) of slices.

Returns

the lengths(columns) of slices.

References gko::array< ValueType >::get_data().

43.186.3.12 get_slice_sets()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type* gko::matrix::Sellp< ValueType, IndexType >::get_slice_sets () [inline], [noexcept]
```

Returns the offsets of slices.

Returns

the offsets of slices.

References gko::array< ValueType >::get_data().

43.186.3.13 get_slice_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Sellp< ValueType, IndexType >::get_slice_size ( ) const [inline],
[noexcept]
```

Returns the size of a slice.

Returns

the size of a slice.

43.186.3.14 get_stride_factor()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Sellp< ValueType, IndexType >::get_stride_factor ( ) const [inline],
[noexcept]
```

Returns the stride factor(t) of SELL-P.

Returns

the stride factor(t) of SELL-P.

43.186.3.15 get_total_cols()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Sellp< ValueType, IndexType >::get_total_cols ( ) const [inline],
[noexcept]
```

Returns the total column number.

Returns

the total column number.

References gko::array< ValueType >::get_num_elems().

43.186.3.16 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Sellp< ValueType, IndexType >::get_values () [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References gko::array< ValueType >::get_data().

43.186.3.17 operator=() [1/2]

Copy-assigns a Sellp matrix.

Preserves the executor, copies the data and parameters.

43.186.3.18 operator=() [2/2]

Move-assigns a Sellp matrix.

Preserves the executor, moves the data and parameters. The moved-from object is empty (0x0 with valid slice_sets and unchanged parameters).

43.186.3.19 read() [1/3]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.186.3.20 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.186.3.21 read() [3/3]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

Reimplemented from gko::ReadableFromMatrixData< ValueType, IndexType >.

43.186.3.22 val_at() [1/2]

Returns the idx-th non-zero element of the row-th row with slice_set slice set.

Parameters

row	the row of the requested element in the slice
slice_set	the slice set of the slice
idx	the idx-th stored element of the row in the slice

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References gko::array< ValueType >::get_const_data().

43.186.3.23 val_at() [2/2]

Returns the idx-th non-zero element of the row-th row with slice set slice set.

Parameters

row	the row of the requested element in the slice
slice_set	the slice set of the slice
idx	the idx-th stored element of the row in the slice

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References gko::array< ValueType >::get_data().

43.186.3.24 write()

Writes a matrix to a matrix data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- · ginkgo/core/matrix/csr.hpp
- · ginkgo/core/matrix/sellp.hpp

43.187 gko::span Struct Reference

A span is a lightweight structure used to create sub-ranges from other ranges.

#include <ginkgo/core/base/range.hpp>

Public Member Functions

• constexpr span (size_type point) noexcept

Creates a span representing a point point.

constexpr span (size_type begin, size_type end) noexcept

Creates a span.

constexpr bool is_valid () const

Checks if a span is valid.

• constexpr size_type length () const

Returns the length of a span.

Public Attributes

const size_type begin

Beginning of the span.

· const size_type end

End of the span.

43.187.1 Detailed Description

A span is a lightweight structure used to create sub-ranges from other ranges.

A span s represents a contiguous set of indexes in one dimension of the range, starting on index s.begin (inclusive) and ending at index s.end (exclusive). A span is only valid if its starting index is smaller than its ending index.

Spans can be compared using the == and != operators. Two spans are identical if both their begin and end values are identical.

Spans also have two distinct partial orders defined on them:

```
1. x < y (y > x) if and only if x.end < y.begin
2. x <= y (y >= x) if and only if x.end <= y.begin
```

Note that the orders are in fact partial - there are spans x and y for which none of the following inequalities holds: x < y, x > y, x == y, x <= y, x >= y. An example are spans $span\{0, 2\}$ and $span\{1, 3\}$.

In addition, <= is a distinct order from <, and not just an extension of the strict order to its weak equivalent. Thus, x <= y is not equivalent to $x < y \mid \mid x == y$.

43.187.2 Constructor & Destructor Documentation

43.187.2.1 span() [1/2]

Creates a span representing a point point.

The begin of this span is set to point, and the end to point + 1.

Parameters

43.187.2.2 span() [2/2]

Creates a span.

Parameters

begin	the beginning of the span
end	the end of the span

References begin.

43.187.3 Member Function Documentation

43.187.3.1 is_valid()

```
constexpr bool gko::span::is_valid ( ) const [inline], [constexpr]
```

Checks if a span is valid.

Returns

```
true if and only if this->begin <= this->end
```

References begin, and end.

Referenced by gko::accessor::row_major< ValueType, Dimensionality >::operator()().

43.187.3.2 length()

```
constexpr size_type gko::span::length ( ) const [inline], [constexpr]
```

Returns the length of a span.

Returns

```
this->end - this->begin
```

References begin, and end.

The documentation for this struct was generated from the following file:

• ginkgo/core/base/range.hpp

43.188 gko::matrix::Csr< ValueType, IndexType >::sparselib Class Reference

sparselib is a strategy_type which uses the sparselib csr.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

• sparselib ()

Creates a sparselib strategy.

- void process (const array < index_type > &mtx_row_ptrs, array < index_type > *mtx_srow) override
 Computes srow according to row pointers.
- int64_t clac_size (const int64_t nnz) override

Computes the srow size according to the number of nonzeros.

- std::shared_ptr< $\ensuremath{\mathsf{strategy_type}}\xspace > \ensuremath{\mathsf{copy}}\xspace$ () override

Copy a strategy.

43.188.1 Detailed Description

```
template < typename ValueType = default_precision, typename IndexType = int32 > class gko::matrix::Csr < ValueType, IndexType >::sparselib
```

sparselib is a strategy_type which uses the sparselib csr.

Note

Uses cusparse in cuda and hipsparse in hip.

43.188.2 Member Function Documentation

43.188.2.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

nnz the number of nonzeros

Returns

the size of srow

Implements gko::matrix::Csr< ValueType, IndexType >::strategy type.

43.188.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::sparselib::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

43.188.2.3 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix
mtx_srow	the srow of the matrix

Implements gko::matrix::Csr< ValueType, IndexType >::strategy_type.

The documentation for this class was generated from the following file:

• ginkgo/core/matrix/csr.hpp

43.189 gko::matrix::SparsityCsr< ValueType, IndexType > Class Template Reference

SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

#include <ginkgo/core/matrix/sparsity_csr.hpp>

Public Member Functions

void read (const mat_data &data) override

Reads a matrix from a matrix_data structure.

• void read (const device_mat_data &data) override

Reads a matrix from a device matrix data structure.

void read (device_mat_data &&data) override

Reads a matrix from a device_matrix_data structure.

void write (mat_data &data) const override

Writes a matrix to a matrix_data structure.

std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

std::unique ptr< SparsityCsr > to adjacency matrix () const

Transforms the sparsity matrix to an adjacency matrix.

void sort_by_column_index ()

Sorts each row by column index.

• index_type * get_col_idxs () noexcept

Returns the column indices of the matrix.

const index_type * get_const_col_idxs () const noexcept

Returns the column indices of the matrix.

index_type * get_row_ptrs () noexcept

Returns the row pointers of the matrix.

• const index_type * get_const_row_ptrs () const noexcept

Returns the row pointers of the matrix.

value_type * get_value () noexcept

Returns the value stored in the matrix.

const value_type * get_const_value () const noexcept

Returns the value stored in the matrix.

size_type get_num_nonzeros () const noexcept

Returns the number of elements explicitly stored in the matrix.

SparsityCsr & operator= (const SparsityCsr &)

Copy-assigns a SparsityCsr matrix.

SparsityCsr & operator= (SparsityCsr &&)

Move-assigns a SparsityCsr matrix.

• SparsityCsr (const SparsityCsr &)

Copy-constructs a SparsityCsr matrix.

SparsityCsr (SparsityCsr &&)

Move-constructs a SparsityCsr matrix.

Static Public Member Functions

static std::unique_ptr< const SparsityCsr > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< IndexType > &&col_idxs, gko::detail::const_array_view< IndexType > &&row ptrs, ValueType value=one< ValueType >())

Creates a constant (immutable) SparsityCsr matrix from constant arrays.

43.189.1 Detailed Description

template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::SparsityCsr< ValueType, IndexType >

SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

The values of the nonzero elements are stored as a value array of length 1. All the values in the matrix are equal to this value. By default, this value is set to 1.0. A row pointer array also stores the linearized starting index of each row. An additional column index array is used to identify the column where a nonzero is present.

Template Parameters

ValueType	precision of vectors in apply
IndexType	precision of matrix indexes

43.189.2 Constructor & Destructor Documentation

43.189.2.1 SparsityCsr() [1/2]

Copy-constructs a SparsityCsr matrix.

Inherits executor, strategy and data.

43.189.2.2 SparsityCsr() [2/2]

Move-constructs a SparsityCsr matrix.

Inherits executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor, no nonzeros and valid row pointers).

43.189.3 Member Function Documentation

43.189.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::SparsityCsr< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.189.3.2 create_const()

Creates a constant (immutable) SparsityCsr matrix from constant arrays.

Parameters

exec	the executor to create the matrix on
size	the dimensions of the matrix
values	the value array of the matrix
col_idxs	the column index array of the matrix
row_ptrs	the row pointer array of the matrix
strategy	the strategy the matrix uses for SpMV operations

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

43.189.3.3 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_col_idxs ( ) [inline],
[noexcept]
```

Returns the column indices of the matrix.

Returns

the column indices of the matrix.

References gko::array< ValueType >::get_data().

43.189.3.4 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_col_idxs ( )
const [inline], [noexcept]
```

Returns the column indices of the matrix.

Returns

the column indices of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.189.3.5 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_row_ptrs ( )
const [inline], [noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.189.3.6 get_const_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_value ( ) const
[inline], [noexcept]
```

Returns the value stored in the matrix.

Returns

the value of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

43.189.3.7 get_num_nonzeros()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::SparsityCsr< ValueType, IndexType >::get_num_nonzeros ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_num_elems().

43.189.3.8 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_row_ptrs ( ) [inline],
[noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

References gko::array< ValueType >::get_data().

43.189.3.9 get_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_value ( ) [inline], [noexcept]
```

Returns the value stored in the matrix.

Returns

the value of the matrix.

References gko::array< ValueType >::get_data().

43.189.3.10 operator=() [1/2]

Copy-assigns a SparsityCsr matrix.

Preserves executor, copies everything else.

43.189.3.11 operator=() [2/2]

Move-assigns a SparsityCsr matrix.

Preserves executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor, no nonzeros and valid row pointers).

43.189.3.12 read() [1/3]

Reads a matrix from a device_matrix_data structure.

Parameters

```
data the device matrix data structure.
```

 $\label{lem:lemented_from_gko::ReadableFromMatrixData} Reimplemented from gko::ReadableFromMatrixData < ValueType, IndexType >.$

43.189.3.13 read() [2/3]

Reads a matrix from a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::ReadableFromMatrixData< ValueType, IndexType >.

43.189.3.14 read() [3/3]

Reads a matrix from a device_matrix_data structure.

The structure may be emptied by this function.

Parameters

```
data the device_matrix_data structure.
```

 $\label{lem:lemented$

43.189.3.15 to_adjacency_matrix()

```
template<typename ValueType = default_precision, typename IndexType = int32> std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::to_adjacency← _matrix ( ) const
```

Transforms the sparsity matrix to an adjacency matrix.

As the adjacency matrix has to be square, the input SparsityCsr matrix for this function to work has to be square.

Note

The adjacency matrix in this case is the sparsity pattern but with the diagonal ones removed. This is mainly used for the reordering/partitioning as taken in by graph libraries such as METIS.

43.189.3.16 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::SparsityCsr< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

43.189.3.17 write()

Writes a matrix to a matrix_data structure.

Parameters

```
data the matrix_data structure
```

Implements gko::WritableToMatrixData< ValueType, IndexType >.

The documentation for this class was generated from the following files:

- · ginkgo/core/matrix/csr.hpp
- ginkgo/core/matrix/sparsity_csr.hpp

43.190 gko::experimental::mpi::status Struct Reference

The status struct is a light wrapper around the MPI_Status struct.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

• status ()

The default constructor.

• MPI_Status * get ()

Get a pointer to the underlying MPI_Status object.

• template<typename T >

```
int get_count (const T *data) const
```

Get the count of the number of elements received by the communication call.

43.190.1 Detailed Description

The status struct is a light wrapper around the MPI_Status struct.

43.190.2 Constructor & Destructor Documentation

43.190.2.1 status()

```
gko::experimental::mpi::status::status ( ) [inline]
```

The default constructor.

It creates an empty MPI_Status

43.190.3 Member Function Documentation

43.190.3.1 get()

```
MPI_Status* gko::experimental::mpi::status::get ( ) [inline]
```

Get a pointer to the underlying MPI_Status object.

Returns

a pointer to MPI_Status object

Referenced by gko::experimental::mpi::communicator::recv(), and gko::experimental::mpi::request::wait().

43.190.3.2 get_count()

Get the count of the number of elements received by the communication call.

Template Parameters

 $T \mid$ The datatype of the object that was received.

Parameters

data The data object of type T that was received.

Returns

the count

The documentation for this struct was generated from the following file:

• ginkgo/core/base/mpi.hpp

43.191 gko::stopping_status Class Reference

This class is used to keep track of the stopping status of one vector.

#include <ginkgo/core/stop/stopping_status.hpp>

Public Member Functions

· bool has_stopped () const noexcept

Check if any stopping criteria was fulfilled.

bool has_converged () const noexcept

Check if convergence was reached.

· bool is_finalized () const noexcept

Check if the corresponding vector stores the finalized result.

• uint8 get_id () const noexcept

Get the id of the stopping criterion which caused the stop.

· void reset () noexcept

Clear all flags.

· void stop (uint8 id, bool set finalized=true) noexcept

Call if a stop occured due to a hard limit (and convergence was not reached).

void converge (uint8 id, bool set_finalized=true) noexcept

Call if convergence occured.

· void finalize () noexcept

Set the result to be finalized (it needs to be stopped or converged first).

Friends

• bool operator== (const stopping_status &x, const stopping_status &y) noexcept

• bool operator!= (const stopping_status &x, const stopping_status &y) noexcept Checks if two stopping statuses are different.

43.191.1 Detailed Description

This class is used to keep track of the stopping status of one vector.

Checks if two stopping statuses are equivalent.

43.191.2 Member Function Documentation

43.191.2.1 converge()

Call if convergence occured.

Parameters

id	id of the stopping criteria.
set_finalized	Controls if the current version should count as finalized (set to true) or not (set to false).

References has_stopped().

43.191.2.2 get_id()

```
uint8 gko::stopping_status::get_id ( ) const [inline], [noexcept]
```

Get the id of the stopping criterion which caused the stop.

Returns

Returns the id of the stopping criterion which caused the stop.

Referenced by has_stopped().

43.191.2.3 has_converged()

```
bool gko::stopping_status::has_converged ( ) const [inline], [noexcept]
```

Check if convergence was reached.

Returns

Returns true if convergence was reached.

43.191.2.4 has_stopped()

```
bool gko::stopping_status::has_stopped ( ) const [inline], [noexcept]
```

Check if any stopping criteria was fulfilled.

Returns

Returns true if any stopping criteria was fulfilled.

References get_id().

Referenced by converge(), finalize(), and stop().

43.191.2.5 is_finalized()

```
bool gko::stopping_status::is_finalized ( ) const [inline], [noexcept]
```

Check if the corresponding vector stores the finalized result.

Returns

Returns true if the corresponding vector stores the finalized result.

43.191.2.6 stop()

Call if a stop occured due to a hard limit (and convergence was not reached).

Parameters

id	id of the stopping criteria.
set_finalized	Controls if the current version should count as finalized (set to true) or not (set to false).

References has_stopped().

43.191.3 Friends And Related Function Documentation

43.191.3.1 operator"!=

Checks if two stopping statuses are different.

Parameters

X	a stopping status	
У	a stopping status	

Returns

```
true if and only if ! (x == y)
```

43.191.3.2 operator==

Checks if two stopping statuses are equivalent.

Parameters

Х	a stopping status	
V	a stopping status	

Returns

true if and only if both \boldsymbol{x} and \boldsymbol{y} have the same mask and converged and finalized state

The documentation for this class was generated from the following file:

• ginkgo/core/stop/stopping_status.hpp

43.192 gko::matrix::Hybrid< ValueType, IndexType >::strategy_type Class Reference

strategy_type is to decide how to set the hybrid config.

#include <ginkgo/core/matrix/hybrid.hpp>

Public Member Functions

• strategy_type ()

Creates a strategy_type.

Computes the config of the Hybrid matrix (ell_num_stored_elements_per_row and coo_nnz).

size_type get_ell_num_stored_elements_per_row () const noexcept

Returns the number of stored elements per row of the ell part.

size_type get_coo_nnz () const noexcept

Returns the number of nonzeros of the coo part.

virtual size_type compute_ell_num_stored_elements_per_row (array< size_type > *row_nnz) const =0
 Computes the number of stored elements per row of the ell part.

43.192.1 Detailed Description

```
template < typename ValueType = default_precision, typename IndexType = int32 > class gko::matrix::Hybrid < ValueType, IndexType >::strategy_type
```

strategy type is to decide how to set the hybrid config.

It computes the number of stored elements per row of the ell part and then set the number of residual nonzeros as the number of nonzeros of the coo part.

The practical strategy method should inherit strategy_type and implement its $compute_ell_num_stored_$ \leftarrow $elements_per_row$ function.

43.192.2 Member Function Documentation

43.192.2.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32> virtual size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_ell_← num_stored_elements_per_row (

array< size_type > * row_nnz ) const [pure virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

rou	the number of nonzeros of each row
row nnz	I the number of nonzeros of each row

Returns

the number of stored elements per row of the ell part

Implemented in gko::matrix::Hybrid< ValueType, IndexType >::automatic, gko::matrix::Hybrid< ValueType, IndexType >::minimal_step gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit, gko::matrix::Hybrid< ValueType, IndexType >::imbalance_and gko::matrix::Hybrid< ValueType, IndexType >::column_limit.

Referenced by gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config().

43.192.2.2 compute_hybrid_config()

Computes the config of the Hybrid matrix (ell_num_stored_elements_per_row and coo_nnz).

For now, it copies row_nnz to the reference executor and performs all operations on the reference executor.

Parameters

row_nnz	the number of nonzeros of each row
ell_num_stored_elements_per_row	the output number of stored elements per row of the ell part
coo_nnz	the output number of nonzeros of the coo part

 $References\ gko::matrix::Hybrid<\ ValueType,\ IndexType>::strategy_type::compute_ell_num_stored_elements_{\leftarrow}\ per_row(),\ gko::array<\ ValueType>::get_executor(),\ and\ gko::array<\ ValueType>::get_num_elems().$

43.192.2.3 get_coo_nnz()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::get_coo_nnz ( ) const
[inline], [noexcept]
```

Returns the number of nonzeros of the coo part.

Returns

the number of nonzeros of the coo part

43.192.2.4 get_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::get_ell_num_stored_←
elements_per_row ( ) const [inline], [noexcept]
```

Returns the number of stored elements per row of the ell part.

Returns

the number of stored elements per row of the ell part

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/hybrid.hpp

43.193 gko::matrix::Csr< ValueType, IndexType >::strategy_type Class Reference

strategy_type is to decide how to set the csr algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

• strategy_type (std::string name)

Creates a strategy_type.

std::string get_name ()

Returns the name of strategy.

- virtual void process (const array< index_type > &mtx_row_ptrs, array< index_type > *mtx_srow)=0
 Computes srow according to row pointers.
- virtual int64_t clac_size (const int64_t nnz)=0

Computes the srow size according to the number of nonzeros.

virtual std::shared_ptr< strategy_type > copy ()=0
 Copy a strategy.

43.193.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::matrix::Csr< ValueType, IndexType >::strategy_type
```

strategy_type is to decide how to set the csr algorithm.

The practical strategy method should inherit strategy_type and implement its process, $clac_size$ function and the corresponding device kernel.

43.193.2 Constructor & Destructor Documentation

43.193.2.1 strategy_type()

Creates a strategy_type.

Parameters

name the name of strategy

43.193.3 Member Function Documentation

43.193.3.1 clac_size()

Computes the srow size according to the number of nonzeros.

Parameters

nnz the number of nonzeros

Returns

the size of srow

Implemented in gko::matrix::Csr< ValueType, IndexType >::load_balance, gko::matrix::Csr< ValueType, IndexType >::sparselib, gko::matrix::Csr< ValueType, IndexType >::cusparse, gko::matrix::Csr< ValueType, IndexType >::merge_path, and gko::matrix::Csr< ValueType, IndexType >::classical.

43.193.3.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
virtual std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::strategy_\tipe::copy ( ) [pure virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like automatical do not work when actually shared.

Implemented in gko::matrix::Csr< ValueType, IndexType >::load_balance, gko::matrix::Csr< ValueType, IndexType >::sparselib, gko::matrix::Csr< ValueType, IndexType >::cusparse, gko::matrix::Csr< ValueType, IndexType >::merge_path, and gko::matrix::Csr< ValueType, IndexType >::classical.

43.193.3.3 get_name()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::string gko::matrix::Csr< ValueType, IndexType >::strategy_type::get_name ( ) [inline]
```

Returns the name of strategy.

Returns

the name of strategy

43.193.3.4 process()

Computes srow according to row pointers.

Parameters

mtx_row_ptrs	the row pointers of the matrix
mtx_srow	the srow of the matrix

Implemented in gko::matrix::Csr< ValueType, IndexType >::load_balance, gko::matrix::Csr< ValueType, IndexType >::sparselib, gko::matrix::Csr< ValueType, IndexType >::cusparse, gko::matrix::Csr< ValueType, IndexType >::merge_path, and gko::matrix::Csr< ValueType, IndexType >::classical.

The documentation for this class was generated from the following file:

· ginkgo/core/matrix/csr.hpp

43.194 gko::log::Stream < ValueType > Class Template Reference

Stream is a Logger which logs every event to a stream.

```
#include <ginkgo/core/log/stream.hpp>
```

Static Public Member Functions

- static std::unique_ptr< Stream > create (std::shared_ptr< const Executor > exec, const Logger::mask_type &enabled_events=Logger::all_events_mask, std::ostream &os=std::cout, bool verbose=false)
 - Creates a Stream logger.

Creates a Stream logger.

43.194.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::log::Stream< ValueType >
```

Stream is a Logger which logs every event to a stream.

This can typically be used to log to a file or to the console.

Template Parameters

ValueType	the type of values stored in the class (i.e. ValueType template parameter of the concrete Loggable	1
	this class will log)	

43.194.2 Member Function Documentation

43.194.2.1 create() [1/2]

Creates a Stream logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

exec	the executor
enabled_events	the events enabled for this logger. By default all events.
os	the stream used for this logger
verbose	whether we want detailed information or not. This includes always printing residuals and other information which can give a large output.

Returns

an std::unique_ptr to the the constructed object

```
230 {
231          return std::unique_ptr<Stream>(new Stream(enabled_events, os, verbose));
232    }
```

43.194.2.2 create() [2/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Stream> gko::log::Stream< ValueType >::create (
```

```
std::shared_ptr< const Executor > exec,
const Logger::mask_type & enabled_events = Logger::all_events_mask,
std::ostream & os = std::cout,
bool verbose = false ) [inline], [static]
```

Creates a Stream logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

exec	the executor
enabled_events	the events enabled for this logger. By default all events.
os	the stream used for this logger
verbose	whether we want detailed information or not. This includes always printing residuals and other information which can give a large output.

Returns

an std::unique_ptr to the the constructed object

The documentation for this class was generated from the following file:

• ginkgo/core/log/stream.hpp

43.195 gko::StreamError Class Reference

StreamError is thrown if accessing a stream failed.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

• StreamError (const std::string &file, int line, const std::string &func, const std::string &message)

Initializes a file access error.

43.195.1 Detailed Description

StreamError is thrown if accessing a stream failed.

43.195.2 Constructor & Destructor Documentation

43.195.2.1 StreamError()

Initializes a file access error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The name of the function that tried to access the file
message	The error message

The documentation for this class was generated from the following file:

• ginkgo/core/base/exception.hpp

43.196 gko::log::ProfilerHook::SummaryWriter Class Reference

Recieves the results from ProfilerHook::create_summary().

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

virtual void write (const std::vector< summary_entry > &entries, std::chrono::nanoseconds overhead)=0
 Callback to write out the summary results.

43.196.1 Detailed Description

Recieves the results from ProfilerHook::create_summary().

43.196.2 Member Function Documentation

43.196.2.1 write()

Callback to write out the summary results.

Parameters

entries	the vector of ranges with runtime and count.
overhead	an estimate of the profiler overhead

 $Implemented\ in\ gko:: log:: Profiler Hook:: Table Summary Writer.$

The documentation for this class was generated from the following file:

• ginkgo/core/log/profiler_hook.hpp

43.197 gko::log::ProfilerHook::TableSummaryWriter Class Reference

Writes the results from ProfilerHook::create_summary() and ProfilerHook::create_nested_summary() to a ASCII table in Markdown format.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

- TableSummaryWriter (std::ostream &output=std::cerr, std::string header="Runtime summary")

 Constructs a writer on an output stream.
- void write (const std::vector< summary_entry > &entries, std::chrono::nanoseconds overhead) override Callback to write out the summary results.
- void write_nested (const nested_summary_entry &root, std::chrono::nanoseconds overhead) override
 Callback to write out the summary results.

43.197.1 Detailed Description

Writes the results from ProfilerHook::create_summary() and ProfilerHook::create_nested_summary() to a ASCII table in Markdown format.

43.197.2 Constructor & Destructor Documentation

43.197.2.1 TableSummaryWriter()

Constructs a writer on an output stream.

Parameters

output	the output stream to write the table to.
header	the header to write above the table.

43.197.3 Member Function Documentation

43.197.3.1 write()

Callback to write out the summary results.

Parameters

entries	the vector of ranges with runtime and count.
overhead	an estimate of the profiler overhead

Implements gko::log::ProfilerHook::SummaryWriter.

43.197.3.2 write_nested()

Callback to write out the summary results.

Parameters

root	the root range with runtime and count.
overhead	an estimate of the profiler overhead

Implements gko::log::ProfilerHook::NestedSummaryWriter.

The documentation for this class was generated from the following file:

• ginkgo/core/log/profiler_hook.hpp

43.198 gko::stop::Time Class Reference

The Time class is a stopping criterion which stops the iteration process after a certain amout of time has passed.

```
#include <ginkgo/core/stop/time.hpp>
```

43.198.1 Detailed Description

The Time class is a stopping criterion which stops the iteration process after a certain amout of time has passed.

The documentation for this class was generated from the following file:

· ginkgo/core/stop/time.hpp

43.199 gko::time point Class Reference

An opaque wrapper for a time point generated by a timer.

#include <ginkgo/core/base/timer.hpp>

43.199.1 Detailed Description

An opaque wrapper for a time point generated by a timer.

The documentation for this class was generated from the following file:

· ginkgo/core/base/timer.hpp

43.200 gko::Timer Class Reference

Represents a generic timer that can be used to record time points and measure time differences on host or device streams.

#include <ginkgo/core/base/timer.hpp>

Public Member Functions

• time_point create_time_point ()

Returns a newly created time point.

virtual void record (time_point &time)=0

Records a time point at the current time.

• virtual void wait (time_point &time)=0

Waits until all kernels in-process when recording the time point are finished.

std::chrono::nanoseconds difference (time_point &start, time_point &stop)

Computes the difference between the two time points in nanoseconds.

• virtual std::chrono::nanoseconds difference_async (const time_point &start, const time_point &stop)=0

Computes the difference between the two time points in nanoseconds.

Static Public Member Functions

static std::unique_ptr < Timer > create_for_executor (std::shared_ptr < const Executor > exec)
 Creates the timer type most suitable for recording accurate timings of kernels on the given executor.

43.200.1 Detailed Description

Represents a generic timer that can be used to record time points and measure time differences on host or device streams.

To keep the runtime overhead of timing minimal, time points need to be allocated beforehand using Timer::create_time_point:

```
auto begin = timer->create_time_point();
auto end = timer->create_time_point();
// ...
timer->record(begin);
run_expensive_operation();
timer->record(end);
auto elapsed = timer->difference(begin, end);
```

43.200.2 Member Function Documentation

43.200.2.1 create_for_executor()

Creates the timer type most suitable for recording accurate timings of kernels on the given executor.

Parameters

```
exec the executor to create a Timer for
```

Returns

CpuTimer for ReferenceExecutor and OmpExecutor, CudaTimer for CudaExecutor, HipTimer for HipExecutor or DpcppExecutor.

43.200.2.2 create_time_point()

```
time_point gko::Timer::create_time_point ( )
```

Returns a newly created time point.

Time points may only be used with the timer they were created with.

43.200.2.3 difference()

Computes the difference between the two time points in nanoseconds.

The function synchronizes with stop before computing the difference.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

43.200.2.4 difference_async()

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. timer->wait(stop) needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

start	the first time point (earlier)
end	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implemented in gko::DpcppTimer, gko::HipTimer, gko::CudaTimer, and gko::CpuTimer.

The documentation for this class was generated from the following file:

• ginkgo/core/base/timer.hpp

43.201 gko::Transposable Class Reference

Linear operators which support transposition should implement the Transposable interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

virtual std::unique_ptr< LinOp > transpose () const =0

Returns a LinOp representing the transpose of the Transposable object.

virtual std::unique_ptr< LinOp > conj_transpose () const =0

Returns a LinOp representing the conjugate transpose of the Transposable object.

43.201.1 Detailed Description

Linear operators which support transposition should implement the Transposable interface.

It provides two functionalities, the normal transpose and the conjugate transpose.

The normal transpose returns the transpose of the linear operator without changing any of its elements representing the operation, $B = A^T$.

The conjugate transpose returns the conjugate of each of the elements and additionally transposes the linear operator representing the operation, $B = A^H$.

43.201.1.1 Example: Transposing a Csr matrix:

```
{c++}
//Transposing an object of LinOp type.
//The object you want to transpose.
auto op = matrix::Csr::create(exec);
//Transpose the object by first converting it to a transposable type.
auto trans = op->transpose();
```

43.201.2 Member Function Documentation

43.201.2.1 conj_transpose()

```
virtual std::unique_ptr<LinOp> gko::Transposable::conj_transpose ( ) const [pure virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

```
Implemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< value_type >, gko::matrix::Dense< value_type, IndexType >, gko::matrix::Ft3, gko::solver::UpperTrs< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Fft2, gko::preconditioner::Ic< LSolverType, IndexType >, gko::matrix::SparsityCsr< ValueType, IndexType >, gko::solver::Ir< ValueType >, gko::matrix::Diagonal< ValueType >, gko::solver::LowerTrs< ValueType >, gko::solver::Idr< ValueType >, gko::Combination< ValueType >, gko::matrix::Fft, gko::solver::Bicg< ValueType >, gko::solver::Bicgstab< ValueType >, gko::solver::Fcg< ValueType >, gko::solver::Cg< ValueType >, gko::solver::Direct< ValueType >, IndexType >, gko::solver::Direct< ValueType, IndexType >,
```

43.201.2.2 transpose()

```
virtual std::unique_ptr<LinOp> gko::Transposable::transpose ( ) const [pure virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implemented in gko::matrix::Csr< ValueType, IndexType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType >, gko::matrix::Dense< ValueType, IndexType >, gko::matrix::Dense< ValueType, IndexType >, gko::preconditioner::Jacobi< ValueType, IndexType >, gko::matrix::Fft3, gko::solver::UpperTrs< ValueType, IndexType >, gko::preconditioner::Isai< IsaiType, ValueType, IndexType >, gko::matrix::Fbcsr< ValueType, IndexType >, gko::matrix::Fft2, gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >, gko::preconditioner::Ic< LSolverType, IndexType >, gko::matrix::Diagonal< ValueType >, gko::solver::LowerTrs< ValueType, IndexType >, gko::solver::Idr< ValueType >, gko::Combination< ValueType >, gko::matrix::Fft, gko::solver::Bicg< ValueType >, gko::solver::Bicgstab< ValueType >, gko::solver::Fcg< ValueType >, gko::solver::Cg< ValueType >,

The documentation for this class was generated from the following file:

• ginkgo/core/base/lin_op.hpp

43.202 gko::experimental::mpi::type_impl< T > Struct Template Reference

A struct that is used to determine the MPI_Datatype of a specified type.

```
#include <ginkgo/core/base/mpi.hpp>
```

43.202.1 Detailed Description

```
template < typename T > struct gko::experimental::mpi::type_impl < T >
```

A struct that is used to determine the MPI_Datatype of a specified type.

Template Parameters

```
T | type of which the MPI_Datatype should be inferred.
```

Note

any specialization of this type hast to provide a static function get_type() that returns an MPI Datatype

The documentation for this struct was generated from the following file:

• ginkgo/core/base/mpi.hpp

43.203 gko::syn::type_list< Types > Struct Template Reference

type_list records several types in template

```
#include <qinkgo/core/synthesizer/containers.hpp>
```

43.203.1 Detailed Description

```
template<typename... Types> struct gko::syn::type_list< Types>
```

type_list records several types in template

Template Parameters

```
Types the types in the list
```

The documentation for this struct was generated from the following file:

· ginkgo/core/synthesizer/containers.hpp

43.204 gko::UnsupportedMatrixProperty Class Reference

Exception throws if a matrix does not have a property required by a numerical method.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

UnsupportedMatrixProperty (const std::string &file, const int line, const std::string &msg)
 Initializes the UnsupportedMatrixProperty error.

43.204.1 Detailed Description

Exception throws if a matrix does not have a property required by a numerical method.

Currently, a message is specified at the call-site manually.

43.204.2 Constructor & Destructor Documentation

43.204.2.1 UnsupportedMatrixProperty()

Initializes the UnsupportedMatrixProperty error.

Parameters

file	The name of the offending source file	
line	line The source code line number where the error occurre	
msg	A message describing the property required.	

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.205 gko::stop::Criterion::Updater Class Reference

The Updater class serves for convenient argument passing to the Criterion's check function.

#include <ginkgo/core/stop/criterion.hpp>

Public Member Functions

• Updater (const Updater &)=delete

Prevent copying and moving the object This is to enforce the use of argument passing and calling check at the same time.

bool check (uint8 stopping_id, bool set_finalized, array< stopping_status > *stop_status, bool *one_← changed) const

Calls the parent Criterion object's check method.

43.205.1 Detailed Description

The Updater class serves for convenient argument passing to the Criterion's check function.

The pattern used is a Builder, except Updater builds a function's arguments before calling the function itself, and does not build an object. This allows calling a Criterion's check in the form of: stop_criterion->update() .num_iterations(num_iterations) .ignore_residual_check(ignore_residual_check) .residual_norm(residual_norm) .implicit_sq_residual_norm(implicit_sq_residual_norm) .residual(residual) .solution(solution) .check(converged);

If there is a need for a new form of data to pass to the Criterion, it should be added here.

43.205.2 Member Function Documentation

43.205.2.1 check()

Calls the parent Criterion object's check method.

References gko::stop::Criterion::check().

The documentation for this class was generated from the following file:

ginkgo/core/stop/criterion.hpp

43.206 gko::solver::UpperTrs< ValueType, IndexType > Class Template Reference

UpperTrs is the triangular solver which solves the system U x = b, when U is an upper triangular matrix.

```
#include <ginkgo/core/solver/triangular.hpp>
```

Public Member Functions

• std::unique_ptr< LinOp > transpose () const override

Returns a LinOp representing the transpose of the Transposable object.

• std::unique_ptr< LinOp > conj_transpose () const override

Returns a LinOp representing the conjugate transpose of the Transposable object.

UpperTrs (const UpperTrs &)

Copy-assigns a triangular solver.

UpperTrs (UpperTrs &&)

Move-assigns a triangular solver.

UpperTrs & operator= (const UpperTrs &)

Copy-constructs a triangular solver.

• UpperTrs & operator= (UpperTrs &&)

Move-constructs a triangular solver.

43.206.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32> class gko::solver::UpperTrs< ValueType, IndexType >
```

UpperTrs is the triangular solver which solves the system U x = b, when U is an upper triangular matrix.

It works best when passing in a matrix in CSR format. If the matrix is not in CSR, then the generate step converts it into a CSR matrix. The generation fails if the matrix is not convertible to CSR.

Note

As the constructor uses the copy and convert functionality, it is not possible to create a empty solver or a solver with a matrix in any other format other than CSR, if none of the executor modules are being compiled with.

Template Parameters

ValueType	precision of matrix elements
IndexType	precision of matrix indices

43.206.2 Constructor & Destructor Documentation

43.206.2.1 UpperTrs() [1/2]

Copy-assigns a triangular solver.

Preserves the executor, shallow-copies the system matrix. If the executors mismatch, clones system matrix onto this executor. Solver analysis information will be regenerated.

43.206.2.2 UpperTrs() [2/2]

Move-assigns a triangular solver.

Preserves the executor, moves the system matrix. If the executors mismatch, clones system matrix onto this executor and regenerates solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

43.206.3 Member Function Documentation

43.206.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::UpperTrs< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the Transposable object.

Returns

a pointer to the new conjugate transposed object

Implements gko::Transposable.

43.206.3.2 operator=() [1/2]

Copy-constructs a triangular solver.

Preserves the executor, shallow-copies the system matrix. Solver analysis information will be regenerated.

43.206.3.3 operator=() [2/2]

Move-constructs a triangular solver.

Preserves the executor, moves the system matrix and solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

43.206.3.4 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::UpperTrs< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the Transposable object.

Returns

a pointer to the new transposed object

Implements gko::Transposable.

The documentation for this class was generated from the following file:

• ginkgo/core/solver/triangular.hpp

43.207 gko::UseComposition < ValueType > Class Template Reference

The UseComposition class can be used to store the composition information in LinOp.

```
#include <ginkgo/core/base/composition.hpp>
```

Public Member Functions

- std::shared_ptr< Composition< ValueType >> get_composition () const
 Returns the composition operators.
- std::shared_ptr< const LinOp > get_operator_at (size_type index) const Returns the operator at index-th poistion of composition.

43.207.1 Detailed Description

```
template<typename ValueType = default_precision> class gko::UseComposition< ValueType >
```

The UseComposition class can be used to store the composition information in LinOp.

Template Parameters

ValueType	precision of input and result vectors
-----------	---------------------------------------

43.207.2 Member Function Documentation

43.207.2.1 get_composition()

```
template<typename ValueType = default_precision>
std::shared_ptr<Composition<ValueType> > gko::UseComposition< ValueType >::get_composition (
) const [inline]
```

Returns the composition operators.

Returns

composition

43.207.2.2 get_operator_at()

Returns the operator at index-th poistion of composition.

Returns

index-th operator

Note

when this composition is not set, this function always returns nullptr. However, when this composition is set, it will throw exception when exceeding index.

Exceptions

```
std::out_of_range | if index is out of bound when composition is existed.
```

Referenced by gko::multigrid::EnableMultigridLevel < ValueType >::get_coarse_op(), gko::multigrid::Enable \leftarrow MultigridLevel < ValueType >::get_prolong_op(), and gko::multigrid::EnableMultigridLevel < ValueType >::get_ \leftarrow restrict_op().

The documentation for this class was generated from the following file:

· ginkgo/core/base/composition.hpp

43.208 gko::syn::value_list< T, Values > Struct Template Reference

value list records several values with the same type in template.

#include <ginkgo/core/synthesizer/containers.hpp>

43.208.1 Detailed Description

```
template<typename T, T... Values>
struct gko::syn::value_list< T, Values>
```

value_list records several values with the same type in template.

Template Parameters

Т	the value type of the list
Values	the values in the list

The documentation for this struct was generated from the following file:

· ginkgo/core/synthesizer/containers.hpp

43.209 gko::ValueMismatch Class Reference

ValueMismatch is thrown if two values are not equal.

#include <ginkgo/core/base/exception.hpp>

Public Member Functions

• ValueMismatch (const std::string &file, int line, const std::string &func, size_type val1, size_type val2, const std::string &clarification)

Initializes a value mismatch error.

43.209.1 Detailed Description

ValueMismatch is thrown if two values are not equal.

43.209.2 Constructor & Destructor Documentation

43.209.2.1 ValueMismatch()

Initializes a value mismatch error.

Parameters

file	The name of the offending source file
line	The source code line number where the error occurred
func	The function name where the error occurred
val1	The first value to be compared.
val2	The second value to be compared.
clarification	An additional message further describing the error

The documentation for this class was generated from the following file:

· ginkgo/core/base/exception.hpp

43.210 gko::experimental::distributed::Vector< ValueType > Class Template Reference

Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

```
#include <ginkgo/core/distributed/vector.hpp>
```

Public Member Functions

void read_distributed (const device_matrix_data< ValueType, int64 > &data, ptr_param< const Partition< int64, int64 >> partition)

Reads a vector from the device_matrix_data structure and a global row partition.

void read_distributed (const matrix_data< ValueType, int64 > &data, ptr_param< const Partition< int64, int64 >> partition)

Reads a vector from the matrix_data structure and a global row partition.

- std::unique_ptr< absolute_type > compute_absolute () const override
 Gets the AbsoluteLinOp.
- void compute_absolute_inplace () override

Compute absolute inplace on each element.

std::unique_ptr< complex_type > make_complex () const

Creates a complex copy of the original vectors.

void make_complex (ptr_param < complex_type > result) const

Writes a complex copy of the original vectors to given complex vectors.

std::unique_ptr< real_type > get_real () const

Creates new real vectors and extracts the real part of the original vectors into that.

void get real (ptr param< real type > result) const

Extracts the real part of the original vectors into given real vectors.

std::unique_ptr< real_type > get_imag () const

Creates new real vectors and extracts the imaginary part of the original vectors into that.

void get_imag (ptr_param < real_type > result) const

Extracts the imaginary part of the original vectors into given real vectors.

void fill (ValueType value)

Fill the distributed vectors with a given value.

void scale (ptr_param< const LinOp > alpha)

Scales the vectors with a scalar (aka: BLAS scal).

void inv scale (ptr param< const LinOp > alpha)

Scales the vectors with the inverse of a scalar.

void add_scaled (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)

Adds b scaled by alpha to the vectors (aka: BLAS axpy).

void sub_scaled (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)

Subtracts b scaled by alpha from the vectors (aka: BLAS axpy).

void compute_dot (ptr_param < const LinOp > b, ptr_param < LinOp > result) const

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

void compute_dot (ptr_param < const LinOp > b, ptr_param < LinOp > result, array < char > &tmp) const

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

 $\bullet \ \ void\ compute_conj_dot\ (ptr_param < const\ LinOp > b,\ ptr_param < LinOp > result)\ const$

Computes the column-wise dot product of this (multi-)vector and conj (b) using a global reduction.

void compute_conj_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp) const

Computes the column-wise dot product of this (multi-)vector and conj (b) using a global reduction.

void compute_squared_norm2 (ptr_param< LinOp > result) const

Computes the square of the column-wise Euclidian (L^2) norm of this (multi-)vector using a global reduction.

- void compute_squared_norm2 (ptr_param< LinOp > result, array< char > &tmp) const

Computes the square of the column-wise Euclidian (L^2) norm of this (multi-)vector using a global reduction.

void compute_norm2 (ptr_param< LinOp > result) const

Computes the Euclidian (L^2) norm of this (multi-)vector using a global reduction.

- void compute_norm2 (ptr_param< LinOp > result, array< char > &tmp) const

Computes the Euclidian (L^2) norm of this (multi-)vector using a global reduction.

void compute_norm1 (ptr_param < LinOp > result) const

Computes the column-wise (L^{\wedge} 1) norm of this (multi-)vector.

• void compute_norm1 (ptr_param< LinOp > result, array< char > &tmp) const

Computes the column-wise (L^{\wedge} 1) norm of this (multi-)vector using a global reduction.

value_type & at_local (size_type row, size_type col) noexcept

Returns a single element of the multi-vector.

- value_type at_local (size_type row, size_type col) const noexcept
- ValueType & at_local (size_type idx) noexcept

Returns a single element of the multi-vector.

- ValueType at local (size type idx) const noexcept
- value_type * get_local_values ()

Returns a pointer to the array of local values of the multi-vector.

const value_type * get_const_local_values () const

Returns a pointer to the array of local values of the multi-vector.

const local_vector_type * get_local_vector () const

Direct (read) access to the underlying local local_vector_type vectors.

std::unique_ptr< const real_type > create_real_view () const

Create a real view of the (potentially) complex original multi-vector.

• std::unique ptr< real type > create real view ()

Create a real view of the (potentially) complex original multi-vector.

Static Public Member Functions

static std::unique_ptr< Vector > create_with_config_of (ptr_param< const Vector > other)

Creates a distributed Vector with the same size and stride as another Vector.

static std::unique_ptr< Vector > create_with_type_of (ptr_param< const Vector > other, std::shared_ptr< const Executor > exec)

Creates an empty Vector with the same type as another Vector, but on a different executor.

static std::unique_ptr< Vector > create_with_type_of (ptr_param< const Vector > other, std::shared_ptr< const Executor > exec, const dim< 2 > &global_size, const dim< 2 > &local_size, size_type stride)

Creates an Vector with the same type as another Vector, but on a different executor and with a different size.

 static std::unique_ptr< const Vector > create_const (std::shared_ptr< const Executor > exec, mpi::communicator comm, dim< 2 > global_size, std::unique_ptr< const local_vector_type > local_← vector)

Creates a constant (immutable) distributed Vector from a constant local vector.

 static std::unique_ptr< const Vector > create_const (std::shared_ptr< const Executor > exec, mpi::communicator comm, std::unique_ptr< const local_vector_type > local_vector)

Creates a constant (immutable) distributed Vector from a constant local vector.

43.210.1 Detailed Description

```
template<typename ValueType = double>
class gko::experimental::distributed::Vector< ValueType >
```

Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

The (multi-)vector is distributed by row, which is described by a

See also

Partition. The local vectors are stored using the

```
Dense format. The vector should be filled using the read_distributed method, e.g. auto part = Partition<...>::build_from_mapping(...); auto vector = Vector<...>::create(exec, comm); vector->read_distributed(matrix_data, part);
```

Using this approach the size of the global vectors, as well as the size of the local vectors, will be automatically inferred. It is possible to create a vector with specified global and local sizes and fill the local vectors using the accessor get_local_vector.

Note

Operations between two vectors (axpy, dot product, etc.) are only valid if both vectors where created using the same partition.

Template Parameters

43.210.2 Member Function Documentation

43.210.2.1 add_scaled()

Adds b scaled by alpha to the vectors (aka: BLAS axpy).

Parameters

alpha	If alpha is 1x1 Dense matrix, the all vectors of b are scaled by alpha. If it is a Dense row vector of values, then i-th column vector of b is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of vectors).
b	a (multi-)vector of the same dimension as this

43.210.2.2 at_local() [1/4]

43.210.2.3 at_local() [2/4]

Returns a single element of the multi-vector.

Useful for iterating across all elements of the multi-vector. However, it is less efficient than the two-parameter variant of this method.

Note

the method has to be called on the same Executor the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

43.210.2.4 at_local() [3/4]

43.210.2.5 at_local() [4/4]

Returns a single element of the multi-vector.

Parameters

row	the local row of the requested element
col	the local column of the requested element

Note

the method has to be called on the same Executor the multi-vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

43.210.2.6 compute_absolute()

```
template<typename ValueType = double>
std::unique_ptr<absolute_type> gko::experimental::distributed::Vector< ValueType >::compute
_absolute ( ) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements gko::EnableAbsoluteComputation< remove_complex< Vector< ValueType >> >.

43.210.2.7 compute_conj_dot() [1/2]

Computes the column-wise dot product of this (multi-)vector and conj (b) using a global reduction.

Parameters

b	a (multi-)vector of same dimension as this
result	a Dense row matrix, used to store the dot product (the number of column in result must match the
	number of columns of this)

43.210.2.8 compute_conj_dot() [2/2]

Computes the column-wise dot product of this (multi-)vector and conj(b) using a global reduction.

Parameters

b	a (multi-)vector of same dimension as this
result	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

43.210.2.9 compute_dot() [1/2]

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

b	a (multi-)vector of same dimension as this
result	a Dense row matrix, used to store the dot product (the number of column in result must match the
	number of columns of this)

43.210.2.10 compute_dot() [2/2]

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

Parameters

b	a (multi-)vector of same dimension as this
result	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

43.210.2.11 compute_norm1() [1/2]

Computes the column-wise ($L^{\wedge}1$) norm of this (multi-)vector.

Parameters

resu	ılt	a Dense row matrix, used to store the norm (the number of columns in result must match the number
		of columns of this)

43.210.2.12 compute_norm1() [2/2]

Computes the column-wise $(L^{\wedge}1)$ norm of this (multi-)vector using a global reduction.

result	a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
	the temporary storage to use for partial sums during the reduction computation. It may be resized

43.210.2.13 compute_norm2() [1/2]

Computes the Euclidian (L^2) norm of this (multi-)vector using a global reduction.

Parameters

result	a Dense row matrix, used to store the norm (the number of columns in result must match the number
	of columns of this)

43.210.2.14 compute_norm2() [2/2]

Computes the Euclidian (L^2) norm of this (multi-)vector using a global reduction.

Parameters

a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.
1

43.210.2.15 compute_squared_norm2() [1/2]

Computes the square of the column-wise Euclidian (L^2) norm of this (multi-)vector using a global reduction.

result	a Dense row vector, used to store the norm (the number of columns in the vector must match the
	number of columns of this)

43.210.2.16 compute_squared_norm2() [2/2]

Computes the square of the column-wise Euclidian (L^2) norm of this (multi-)vector using a global reduction.

Parameters

result	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
tmp	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

43.210.2.17 create const() [1/2]

Creates a constant (immutable) distributed Vector from a constant local vector.

Parameters

exec	Executor associated with this vector
comm	Communicator associated with this vector
global_size	The global size of the vector
local_vector	The underlying local vector, of which a view is created

43.210.2.18 create_const() [2/2]

Creates a constant (immutable) distributed Vector from a constant local vector.

The global size will be deduced from the local sizes, which will incur a collective communication.

Parameters

exec	Executor associated with this vector
comm	Communicator associated with this vector
local_vector	The underlying local vector, of which a view is created

43.210.2.19 create_real_view() [1/2]

```
template<typename ValueType = double>
std::unique_ptr<real_type> gko::experimental::distributed::Vector< ValueType >::create_real
_view ( )
```

Create a real view of the (potentially) complex original multi-vector.

If the original vector is real, nothing changes. If the original vector is complex, the result is created by viewing the complex vector with as real with a reinterpret_cast with twice the number of columns and double the stride.

43.210.2.20 create_real_view() [2/2]

```
template<typename ValueType = double>
std::unique_ptr<const real_type> gko::experimental::distributed::Vector< ValueType >::create
_real_view ( ) const
```

Create a real view of the (potentially) complex original multi-vector.

If the original vector is real, nothing changes. If the original vector is complex, the result is created by viewing the complex vector with as real with a reinterpret_cast with twice the number of columns and double the stride.

43.210.2.21 create_with_config_of()

Creates a distributed Vector with the same size and stride as another Vector.

Parameters

other	The other vector whose configuration needs to copied.
-------	---

43.210.2.22 create_with_type_of() [1/2]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create_←
```

Creates an empty Vector with the same type as another Vector, but on a different executor.

Parameters

other	The other multi-vector whose type we target.
exec	The executor of the new multi-vector.

Note

The new multi-vector uses the same communicator as other.

Returns

an empty Vector with the type of other.

43.210.2.23 create_with_type_of() [2/2]

Creates an Vector with the same type as another Vector, but on a different executor and with a different size.

Parameters

other	The other multi-vector whose type we target.
exec	The executor of the new multi-vector.
global_size	The global size of the multi-vector.
local_size	The local size of the multi-vector.
stride	The stride of the new multi-vector.

Returns

a Vector of specified size with the type of other.

43.210.2.24 fill()

Fill the distributed vectors with a given value.

Parameters

```
value the value to be filled
```

43.210.2.25 get_const_local_values()

```
template<typename ValueType = double>
const value_type* gko::experimental::distributed::Vector< ValueType >::get_const_local_values
( ) const
```

Returns a pointer to the array of local values of the multi-vector.

Returns

the pointer to the array of local values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

43.210.2.26 get_local_values()

```
template<typename ValueType = double>
value_type* gko::experimental::distributed::Vector< ValueType >::get_local_values ( )
```

Returns a pointer to the array of local values of the multi-vector.

Returns

the pointer to the array of local values

43.210.2.27 get_local_vector()

```
template<typename ValueType = double>
const local_vector_type* gko::experimental::distributed::Vector< ValueType >::get_local_vector
( ) const
```

Direct (read) access to the underlying local local_vector_type vectors.

Returns

a constant pointer to the underlying local_vector_type vectors

43.210.2.28 inv scale()

Scales the vectors with the inverse of a scalar.

Parameters

alpha

If alpha is 1x1 Dense matrix, the all vectors are scaled by 1 / alpha. If it is a Dense row vector of values, then i-th column vector is scaled with the inverse of the i-th element of alpha (the number of columns of alpha has to match the number of vectors).

43.210.2.29 make_complex() [1/2]

```
template<typename ValueType = double>
std::unique_ptr<complex_type> gko::experimental::distributed::Vector< ValueType >::make_
complex ( ) const
```

Creates a complex copy of the original vectors.

If the original vectors were real, the imaginary part of the result will be zero.

43.210.2.30 make_complex() [2/2]

Writes a complex copy of the original vectors to given complex vectors.

If the original vectors were real, the imaginary part of the result will be zero.

43.210.2.31 read_distributed() [1/2]

Reads a vector from the device_matrix_data structure and a global row partition.

The number of rows of the matrix data is ignored, only its number of columns is relevant. Both the number of local and global rows are inferred from the row partition.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

data	The device_matrix_data structure
partition	The global row partition

43.210.2.32 read_distributed() [2/2]

Reads a vector from the matrix_data structure and a global row partition.

See @read_distributed

Note

For efficiency it is advised to use the device matrix data overload.

43.210.2.33 scale()

Scales the vectors with a scalar (aka: BLAS scal).

Parameters

alpha	If alpha is 1x1 Dense matrx, the all vectors are scaled by alpha. If it is a Dense row vector of values,
	then i-th column vector is scaled with the i-th element of alpha (the number of columns of alpha has to
	match the number of vectors).

43.210.2.34 sub_scaled()

Subtracts b scaled by alpha from the vectors (aka: BLAS axpy).

Parameters

alpha	If alpha is 1x1 Dense matrix, the all vectors of b are scaled by alpha. If it is a Dense row vector of
	values, then i-th column vector of b is scaled with the i-th element of alpha (the number of c
b	a (multi-)vector of the same dimension as this

The documentation for this class was generated from the following files:

- ginkgo/core/distributed/matrix.hpp
- ginkgo/core/distributed/vector.hpp

43.211 gko::version Struct Reference

This structure is used to represent versions of various Ginkgo modules.

```
#include <ginkgo/core/base/version.hpp>
```

Public Attributes

· const uint64 major

The major version number.

· const uint64 minor

The minor version number.

const uint64 patch

The patch version number.

• const char *const tag

Addition tag string that describes the version in more detail.

43.211.1 Detailed Description

This structure is used to represent versions of various Ginkgo modules.

Version structures can be compared using the usual relational operators.

43.211.2 Member Data Documentation

43.211.2.1 tag

```
const char* const gko::version::tag
```

Addition tag string that describes the version in more detail.

It does not participate in comparisons.

Referenced by gko::operator<<().

The documentation for this struct was generated from the following file:

• ginkgo/core/base/version.hpp

43.212 gko::version_info Class Reference

Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:

```
#include <ginkgo/core/base/version.hpp>
```

Static Public Member Functions

static const version_info & get ()
 Returns an instance of version_info.

Public Attributes

· version header_version

Contains version information of the header files.

version core_version

Contains version information of the core library.

version reference_version

Contains version information of the reference module.

version omp_version

Contains version information of the OMP module.

version cuda_version

Contains version information of the CUDA module.

version hip_version

Contains version information of the HIP module.

version dpcpp_version

Contains version information of the DPC++ module.

43.212.1 Detailed Description

Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:

- 1. Versions with different major version number have incompatible interfaces (parts of the earlier interface may not be present anymore, and new interfaces can appear).
- 2. Versions with the same major number X, but different minor numbers Y1 and Y2 numbers keep the same interface as version X.0.0, but additions to the interface in X.0.0 present in X.Y1.0 may not be present in X.Y2.0 and vice versa.
- 3. Versions with the same major an minor version numbers, but different patch numbers have exactly the same interface, but the functionality may be different (something that is not implemented or has a bug in an earlier version may have this implemented or fixed in a later version).

This structure provides versions of different parts of Ginkgo: the headers, the core and the kernel modules (reference, OpenMP, CUDA, HIP, DPCPP). To obtain an instance of version_info filled with information about the current version of Ginkgo, call the version_info::get() static method.

43.212.2 Member Function Documentation

43.212.2.1 get()

```
static const version_info& gko::version_info::get ( ) [inline], [static]
```

Returns an instance of version_info.

Returns

an instance of version info

43.212.3 Member Data Documentation

43.212.3.1 core version

```
version gko::version_info::core_version
```

Contains version information of the core library.

This is the version of the static/shared library called "ginkgo".

43.212.3.2 cuda_version

```
version gko::version_info::cuda_version
```

Contains version information of the CUDA module.

This is the version of the static/shared library called "ginkgo_cuda".

43.212.3.3 dpcpp_version

```
version gko::version_info::dpcpp_version
```

Contains version information of the DPC++ module.

This is the version of the static/shared library called "ginkgo dpcpp".

43.212.3.4 hip_version

```
version gko::version_info::hip_version
```

Contains version information of the HIP module.

This is the version of the static/shared library called "ginkgo_hip".

43.212.3.5 omp_version

```
version gko::version_info::omp_version
```

Contains version information of the OMP module.

This is the version of the static/shared library called "ginkgo_omp".

43.212.3.6 reference_version

```
version gko::version_info::reference_version
```

Contains version information of the reference module.

This is the version of the static/shared library called "ginkgo_reference".

The documentation for this class was generated from the following file:

• ginkgo/core/base/version.hpp

43.213 gko::experimental::mpi::window< ValueType > Class Template Reference

This class wraps the MPI_Window class with RAII functionality.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Types

• enum create_type

The create type for the window object.

· enum lock type

The lock type for passive target synchronization of the windows.

Public Member Functions

• window ()

The default constructor.

• window (window &&other)

The move constructor.

window & operator= (window &&other)

The move assignment operator.

window (std::shared_ptr< const Executor > exec, ValueType *base, int num_elems, const communicator &comm, const int disp_unit=sizeof(ValueType), MPI_Info input_info=MPI_INFO_NULL, create_type c_
 type=create_type::create)

Create a window object with a given data pointer and type.

MPI_Win get_window () const

Get the underlying window object of MPI_Win type.

• void fence (int assert=0) const

The active target synchronization using MPI_Win_fence for the window object.

• void lock (int rank, lock_type lock_t=lock_type::shared, int assert=0) const

Create an epoch using MPI_Win_lock for the window object.

· void unlock (int rank) const

Close the epoch using MPI_Win_unlock for the window object.

· void lock_all (int assert=0) const

Create the epoch on all ranks using MPI_Win_lock_all for the window object.

void unlock_all () const

Close the epoch on all ranks using MPI_Win_unlock_all for the window object.

· void flush (int rank) const

Flush the existing RDMA operations on the target rank for the calling process for the window object.

void flush_local (int rank) const

Flush the existing RDMA operations on the calling rank from the target rank for the window object.

· void flush_all () const

Flush all the existing RDMA operations for the calling process for the window object.

· void flush_all_local () const

Flush all the local existing RDMA operations on the calling rank for the window object.

void sync () const

Synchronize the public and private buffers for the window object.

~window ()

The deleter which calls MPI_Win_free when the window leaves its scope.

 $\bullet \ \ \text{template}{<} \text{typename PutType} >$

void put (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const

Put data into the target window.

• template<typename PutType >

request r_put (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const

Put data into the target window.

template<typename PutType >
 void accumulate (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_
 count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const

Accumulate data into the target window.

template<typename PutType >
 request r_accumulate (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const

(Non-blocking) Accumulate data into the target window.

template<typename GetType >

void get (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, const int target rank, const unsigned int target disp, const int target count) const

Get data from the target window.

 $\bullet \ \ {\it template}{<} {\it typename GetType}>$

request r_get (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const

Get data (with handle) from the target window.

template<typename GetType >

void get_accumulate (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, GetType *result_buffer, const int result_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const

Get Accumulate data from the target window.

template<typename GetType >

request r_get_accumulate (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, GetType *result_buffer, const int result_count, const int target_rank, const unsigned int target ← disp, const int target count, MPI Op operation) const

(Non-blocking) Get Accumulate data (with handle) from the target window.

 $\bullet \ \ \text{template}{<} \text{typename GetType} >$

void fetch_and_op (std::shared_ptr< const Executor > exec, GetType *origin_buffer, GetType *result_buffer, const int target_rank, const unsigned int target_disp, MPI_Op operation) const

Fetch and operate on data from the target window (An optimized version of Get_accumulate).

43.213.1 Detailed Description

```
template<typename ValueType>
class gko::experimental::mpi::window< ValueType >
```

This class wraps the MPI Window class with RAII functionality.

Different create and lock type methods are setup with enums.

MPI_Window is primarily used for one sided communication and this class provides functionalities to fence, lock, unlock and flush the communication buffers.

43.213.2 Constructor & Destructor Documentation

43.213.2.1 window() [1/3]

```
template<typename ValueType >
gko::experimental::mpi::window< ValueType >::window ( ) [inline]
```

The default constructor.

It creates a null window of MPI WIN NULL type.

43.213.2.2 window() [2/3]

The move constructor.

Move the other object and replace it with MPI WIN NULL

Parameters

other	the window object to be moved.
Ottion	The Window object to be moved.

43.213.2.3 window() [3/3]

```
template<typename ValueType >
gko::experimental::mpi::window< ValueType >::window (
    std::shared_ptr< const Executor > exec,
    ValueType * base,
    int num_elems,
    const communicator & comm,
    const int disp_unit = sizeof(ValueType),
    MPI_Info input_info = MPI_INFO_NULL,
    create_type c_type = create_type::create ) [inline]
```

Create a window object with a given data pointer and type.

A collective operation.

exec	The executor, on which the base pointer is located.
base	the base pointer for the window object.
num_elems	the num_elems of type ValueType the window points to.
comm	the communicator whose ranks will have windows created.
disp_unit	the displacement from base for the window object.
input_info	the MPI_Info object used to set certain properties.
c_type	the type of creation method to use to create the window.

References gko::experimental::mpi::communicator::get().

43.213.3 Member Function Documentation

43.213.3.1 accumulate()

```
template<typename ValueType >
template<typename PutType >
void gko::experimental::mpi::window< ValueType >::accumulate (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

Accumulate data into the target window.

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
origin_count	the number of elements to put
target_rank	the rank to put the data to
target_disp	the displacement at the target window
target_count	the request handle for the send call
operation	the reduce operation. See @MPI_Op

References gko::experimental::mpi::window< ValueType >::get_window().

43.213.3.2 fence()

The active target synchronization using MPI_Win_fence for the window object.

This is called on all associated ranks.

_		
ſ	assert	the optimization level. 0 is always valid.

43.213.3.3 fetch and op()

```
template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::fetch_and_op (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    GetType * result_buffer,
    const int target_rank,
    const unsigned int target_disp,
    MPI_Op operation ) const [inline]
```

Fetch and operate on data from the target window (An optimized version of Get_accumulate).

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
target_rank	the rank to get the data from
target_disp	the displacement at the target window
operation	the reduce operation. See @MPI_Op

References gko::experimental::mpi::window< ValueType >::get_window().

43.213.3.4 flush()

Flush the existing RDMA operations on the target rank for the calling process for the window object.

Parameters

```
rank the target rank.
```

43.213.3.5 flush_local()

Flush the existing RDMA operations on the calling rank from the target rank for the window object.

Parameters

```
rank the target rank.
```

43.213.3.6 get()

```
template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::get (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]
```

Get data from the target window.

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
origin_count	the number of elements to get
target_rank	the rank to get the data from
target_disp	the displacement at the target window
target_count	the request handle for the send call

References gko::experimental::mpi::window< ValueType >::get_window().

43.213.3.7 get_accumulate()

```
template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::get_accumulate (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    GetType * result_buffer,
    const int result_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

Get Accumulate data from the target window.

Parameters

exec	The executor, on which the message buffers are located.
origin_buffer	the buffer to send
origin_count	the number of elements to get
result_buffer	the buffer to receive the target data
result_count	the number of elements to get
target_rank	the rank to get the data from
target_disp	the displacement at the target window
target_count	the request handle for the send call
operation	the reduce operation. See @MPI_Op

References gko::experimental::mpi::window< ValueType >::get_window().

43.213.3.8 get window()

```
template<typename ValueType >
MPI_Win gko::experimental::mpi::window< ValueType >::get_window ( ) const [inline]
```

Get the underlying window object of MPI_Win type.

Returns

the underlying window object.

Referenced by gko::experimental::mpi::window< ValueType >::accumulate(), gko::experimental::mpi::window< ValueType >::fetch_and_op(), gko::experimental::mpi::window< ValueType >::get(), gko::experimental:-:mpi::window< ValueType >::get_accumulate(), gko::experimental::mpi::window< ValueType >::put(), gko::experimental::mpi::window< ValueType >::r_accumulate(), gko::experimental::mpi::window< ValueType >::r_cet(), gko::experimental::mpi::window< ValueType >::r_get_accumulate(), and gko::experimental::mpi::window< ValueType >::r_put().

43.213.3.9 lock()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::lock (
    int rank,
    lock_type lock_t = lock_type::shared,
    int assert = 0 ) const [inline]
```

Create an epoch using MPI_Win_lock for the window object.

rank	the target rank.
lock⊷	the type of the lock: shared or exclusive
_t	
assert Generated b	the optimization level. 0 is always valid.

43.213.3.10 lock all()

Create the epoch on all ranks using MPI_Win_lock_all for the window object.

Parameters

```
assert the optimization level. 0 is always valid.
```

43.213.3.11 operator=()

The move assignment operator.

Move the other object and replace it with MPI_WIN_NULL

Parameters

```
other the window object to be moved.
```

43.213.3.12 put()

Put data into the target window.

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send

Parameters

origin_count	the number of elements to put
target_rank	the rank to put the data to
target_disp	the displacement at the target window
target_count	the request handle for the send call

References gko::experimental::mpi::window< ValueType >::get_window().

43.213.3.13 r_accumulate()

```
template<typename ValueType >
template<typename PutType >
request gko::experimental::mpi::window< ValueType >::r_accumulate (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

(Non-blocking) Accumulate data into the target window.

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
origin_count	the number of elements to put
target_rank	the rank to put the data to
target_disp	the displacement at the target window
target_count	the request handle for the send call
operation	the reduce operation. See @MPI_Op

Returns

the request handle for the send call

References gko::experimental::mpi::request::get(), and gko::experimental::mpi::window< ValueType >::get $_\leftarrow$ window().

43.213.3.14 r_get()

```
template<typename ValueType >
template<typename GetType >
```

Get data (with handle) from the target window.

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
origin_count	the number of elements to get
target_rank	the rank to get the data from
target_disp	the displacement at the target window
target_count	the request handle for the send call

Returns

the request handle for the send call

References gko::experimental::mpi::request::get(), and gko::experimental::mpi::window< ValueType >::get_ \leftarrow window().

43.213.3.15 r_get_accumulate()

(Non-blocking) Get Accumulate data (with handle) from the target window.

Parameters

exec	The executor, on which the message buffers are located.
origin_buffer	the buffer to send
origin_count	the number of elements to get
result_buffer	the buffer to receive the target data
result_count	the number of elements to get
target_rank	the rank to get the data from
target_disp	the displacement at the target window
target_count	the request handle for the send call
operation	the reduce operation. See @MPI_Op

Generated by Doxygen

Returns

the request handle for the send call

References gko::experimental::mpi::request::get(), and gko::experimental::mpi::window< ValueType >::get_ \leftarrow window().

43.213.3.16 r_put()

```
template<typename ValueType >
template<typename PutType >
request gko::experimental::mpi::window< ValueType >::r_put (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]
```

Put data into the target window.

Parameters

exec	The executor, on which the message buffer is located.
origin_buffer	the buffer to send
origin_count	the number of elements to put
target_rank	the rank to put the data to
target_disp	the displacement at the target window
target_count	the request handle for the send call

Returns

the request handle for the send call

References gko::experimental::mpi::request::get(), and gko::experimental::mpi::window< ValueType >::get $_\leftarrow$ window().

43.213.3.17 unlock()

Close the epoch using MPI_Win_unlock for the window object.

Parameters

rank the target rank.

The documentation for this class was generated from the following file:

· ginkgo/core/base/mpi.hpp

43.214 gko::solver::workspace_traits< Solver > Struct Template Reference

Traits class providing information on the type and location of workspace vectors inside a solver.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

43.214.1 Detailed Description

```
template < typename Solver > struct gko::solver::workspace_traits < Solver >
```

Traits class providing information on the type and location of workspace vectors inside a solver.

The documentation for this struct was generated from the following file:

ginkgo/core/solver/solver_base.hpp

43.215 gko::WritableToMatrixData< ValueType, IndexType > Class Template Reference

A LinOp implementing this interface can write its data to a matrix_data structure.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

virtual void write (matrix_data < ValueType, IndexType > &data) const =0
 Writes a matrix to a matrix_data structure.

43.215.1 Detailed Description

```
template < typename ValueType, typename IndexType > class gko::WritableToMatrixData < ValueType, IndexType >
```

A LinOp implementing this interface can write its data to a matrix_data structure.

43.215.2 Member Function Documentation

43.215.2.1 write()

Writes a matrix to a matrix_data structure.

Parameters

data the matrix_data structure

Implemented in gko::matrix::Fft3, gko::matrix::Fft4, gko::matrix::Fft4, gko::matrix::Fft5, gko::matrix::Fft5

The documentation for this class was generated from the following file:

ginkgo/core/base/lin_op.hpp

Index

```
ahs
                                                            gko, 324-327
    gko, 324
                                                        as array
                                                            gko::syn, 388
accumulate
                                                       as_const_view
    gko::experimental::mpi::window< ValueType
                                                            gko::array< ValueType >, 403
add logger
                                                        as list
                                                            gko::syn, 387
    gko::Executor, 606
                                                        as_view
    gko::log::Loggable, 719
                                                            gko::array< ValueType >, 403
add scaled
    gko::experimental::distributed::Vector< ValueType
                                                            gko::matrix::Dense< ValueType >, 519, 520
         >, 892
                                                        at local
    gko::matrix::Dense< ValueType >, 518
                                                            gko::experimental::distributed::Vector< ValueType
add_scaled_identity
                                                                  >, 892, 893
    gko::ScaledIdentityAddable, 834
                                                        autodetect
add value
                                                            gko::precision_reduction, 802
    gko::matrix_assembly_data< ValueType, Index-
                                                        automatic
         Type >, 736
                                                            gko, 324
all gather
    gko::experimental::mpi::communicator, 439
                                                        BadDimension
all reduce
                                                            gko::BadDimension, 412
    gko::experimental::mpi::communicator, 439, 440
                                                        bind to core
all to all
                                                            gko::machine topology, 725
    gko::experimental::mpi::communicator, 441
                                                        bind to cores
all to all v
                                                            gko::machine_topology, 725
    gko::experimental::mpi::communicator, 442
                                                        bind to pu
alloc
                                                            gko::machine topology, 726
    gko::Executor, 606
                                                        bind_to_pus
allocation_mode
                                                            gko::machine_topology, 726
    gko, 323
                                                        BlockSizeError
AllocationError
                                                            gko::BlockSizeError< IndexType >, 420
    gko::AllocationError, 394
                                                        broadcast
apply2
                                                            gko::experimental::mpi::communicator, 443
    gko::matrix::Coo < ValueType, IndexType >, 475-
                                                        build from contiguous
         477
                                                            gko::experimental::distributed::Partition< LocalIn-
apply uses initial guess
                                                                 dexType, GlobalIndexType >, 775
    gko::solver::Bicg< ValueType >, 413
                                                        build_from_global_size_uniform
    gko::solver::Bicgstab < ValueType >, 415
                                                            gko::experimental::distributed::Partition< LocalIn-
    gko::solver::Cg< ValueType >, 423
                                                                 dexType, GlobalIndexType >, 776
    gko::solver::Cgs< ValueType >, 425
                                                        build_from_mapping
    gko::solver::Fcg< ValueType >, 627
                                                            gko::experimental::distributed::Partition< LocalIn-
    gko::solver::Gcr< ValueType >, 638
                                                                 dexType, GlobalIndexType >, 776
    gko::solver::Gmres < ValueType >, 641
                                                        build_smoother
    gko::solver::ldr< ValueType >, 673
                                                            gko::solver, 383
    gko::solver::Ir < ValueType >, 699
    gko::solver::Multigrid, 754
                                                        ceildiv
array
                                                            gko, 328
    gko, 323
                                                        check
    gko::array< ValueType >, 399-403
                                                            gko::stop::Criterion, 484
                                                            gko::stop::Criterion::Updater, 883
as
```

922 INDEX

clac_size	gko::AbsoluteComputable, 391
gko::matrix::Csr< ValueType, IndexType >::classical, 428	gko::EnableAbsoluteComputation< AbsoluteLinOp >, 584
gko::matrix::Csr< ValueType, IndexType >::cusparse	e,compute_conj_dot
511	gko::experimental::distributed::Vector< ValueType
gko::matrix::Csr< ValueType, IndexType >::load_bala	ance, >, 893, 894
717	gko::matrix::Dense< ValueType >, 523
gko::matrix::Csr< ValueType, IndexType >::merge_p	• —
748	gko::experimental::distributed::Vector< ValueType
gko::matrix::Csr< ValueType, IndexType >::sparselib	
852	gko::matrix::Dense< ValueType >, 523, 524
gko::matrix::Csr< ValueType, IndexType >::strategy_ 870	
clear	gko::matrix::Hybrid< ValueType, IndexType >::automatic, 411
gko::array< ValueType >, 403	gko::matrix::Hybrid< ValueType, IndexType
gko::index_set< IndexType >, 691	>::column_limit, 430
gko::PolymorphicObject, 795	gko::matrix::Hybrid< ValueType, IndexType
clone	>::imbalance bounded limit, 683
gko, 328, 329	gko::matrix::Hybrid< ValueType, IndexType
gko::PolymorphicObject, 795, 796	>::imbalance_limit, 685
col at	gko::matrix::Hybrid< ValueType, IndexType
gko::matrix::Ell< ValueType, IndexType >, 575,	>::minimal_storage_limit, 751
576	gko::matrix::Hybrid< ValueType, IndexType
gko::matrix::Sellp< ValueType, IndexType >, 841,	>::strategy_type, 867
842	compute_hybrid_config
column_limit	gko::matrix::Hybrid< ValueType, IndexType
gko::matrix::Hybrid< ValueType, IndexType	>::strategy_type, 868
>::column_limit, 430	compute_norm1
column_permute	gko::experimental::distributed::Vector< ValueType
gko::matrix::Csr< ValueType, IndexType >, 489	>, 895
gko::matrix::Dense< ValueType >, 520–522	gko::matrix::Dense< ValueType >, 524
gko::Permutable < IndexType >, 783 Combination	compute_norm2
gko::Combination< ValueType >, 432	gko::experimental::distributed::Vector< ValueType >, 896
combine	gko::matrix::Dense< ValueType >, 526
Stopping criteria, 307	compute_squared_norm2
comm_index_type	gko::experimental::distributed::Vector< ValueType
gko::experimental::distributed, 367	>, 896
common	gko::matrix::Dense< ValueType >, 526, 527
gko::precision_reduction, 802	compute_storage_space
communicator	gko::preconditioner::block_interleaved_storage_scheme<
gko::experimental::mpi::communicator, 438	IndexType >, 417
Composition	concatenate
gko::Composition< ValueType >, 464	gko::syn, 388
compute_absolute	cond
gko::EnableAbsoluteComputation< AbsoluteLinOp >, 584	gko::matrix_data < ValueType, IndexType >, 743 conj
gko::experimental::distributed::Vector< ValueType	gko, 330
>, 893	conj_transpose
gko::matrix::Coo< ValueType, IndexType >, 477	gko::Combination< ValueType >, 432
gko::matrix::Csr< ValueType, IndexType >, 490	gko::Composition < ValueType >, 465
gko::matrix::Dense< ValueType >, 522	gko::experimental::solver::Direct< ValueType, In-
gko::matrix::Diagonal< ValueType >, 555 gko::matrix::Ell< ValueType, IndexType >, 576	<pre>dexType >, 566 gko::matrix::Csr< ValueType, IndexType >, 490</pre>
gko::matrix::Fbcsr< ValueType, IndexType >, 376	gko::matrix::Dense< ValueType, index type >, 430
gko::matrix::Hybrid< ValueType, IndexType >, 655	gko::matrix::Delise< valueType >, 327 gko::matrix::Diagonal< ValueType >, 555
gko::matrix::Tybria < ValueType, IndexType >, 842	gko::matrix::Fbcsr< ValueType, IndexType >, 619
compute_absolute_linop	gko::matrix::Fft, 629
· — · ·	

gko::matrix::Fft2, 632	853
gko::matrix::Fft3, 634	gko::matrix::Csr< ValueType, IndexType >::strategy_type,
gko::matrix::ldentity< ValueType >, 670	870
	py_and_convert_to
>, 855	gko, 330–332
	py_from
>, 667	gko::accessor::row_major< ValueType, Dimen-
gko::preconditioner::llu< LSolverType, USolver-	sionality >, 830
Type, ReverseApply, IndexType >, 680	gko::Executor, 607
gko::preconditioner::lsai< lsaiType, ValueType, In-	gko::PolymorphicObject, 796–798
- · · · · · · - · · · · · · · · · ·	
• •	py_to_host
gko::preconditioner::Jacobi< ValueType, Index-	gko::device_matrix_data< ValueType, IndexType
Type >, 709	>, 549
	py_val_to_host
gko::solver::Bicgstab< ValueType >, 415	gko::Executor, 607
	re_version
gko::solver::Cgs< ValueType >, 425	gko::version_info, 905
3 71 /	eate
gko::solver::Gcr< ValueType >, 639	gko::CudaExecutor, 505
gko::solver::Gmres< ValueType >, 641	gko::DpcppExecutor, 569
gko::solver::ldr< ValueType >, 673	gko::EnableDefaultFactory< ConcreteFactory, Pro-
gko::solver::lr< ValueType >, 700	ductType, ParametersType, PolymorphicBase
gko::solver::LowerTrs< ValueType, IndexType >,	>, 587
722	gko::HipExecutor, 647
gko::solver::UpperTrs< ValueType, IndexType >,	gko::log::Convergence< ValueType >, 469
885	gko::log::Papi< ValueType >, 769
gko::Transposable, 880	gko::log::PerformanceHint, 781
const_view	gko::log::Record, 821, 822
gko::array< ValueType >, 403	gko::log::Stream< ValueType >, 872
contains	gko::matrix::IdentityFactory< ValueType >, 672
	eate_const
gko::matrix_assembly_data< ValueType, Index-	gko::experimental::distributed::Vector< ValueType
Type >, 736	>, 897
contiguous_type	gko::matrix::Coo< ValueType, IndexType >, 477
gko::experimental::mpi::contiguous_type, 467	gko::matrix::Csr< ValueType, IndexType >, 490
converge	gko::matrix::Dense< ValueType >, 528
gko::stopping_status, 864	gko::matrix::Diagonal < ValueType >, 555
	gko::matrix::Ell< ValueType, IndexType >, 576
convert_to	gko::matrix::Fbcsr< ValueType, IndexType >, 576
gko::ConvertibleTo< ResultType >, 472	9
gko::EnablePolymorphicAssignment< Concrete-	gko::matrix::Permutation< IndexType >, 788
Type, ResultType >, 596	gko::matrix::RowGatherer< IndexType >, 833
gko::matrix::Fbcsr< ValueType, IndexType >, 619,	gko::matrix::SparsityCsr< ValueType, IndexType
620	>, 856
	eate_const_view_of
Type >, 709	gko::matrix::Dense< ValueType >, 528
	eate_default
gko, 323	gko::PolymorphicObject, 798, 799
• •	eate_for_executor
gko::Executor, 606	gko::Timer, 878
gko::matrix::Csr< ValueType, IndexType >::classical, cre	
428	gko::experimental::factorization::Factorization<
gko::matrix::Csr< ValueType, IndexType >::cusparse,	ValueType, IndexType >, 614
	eate_from_composition
gko::matrix::Csr< ValueType, IndexType >::load_balanc	
718	ValueType, IndexType >, 614
gko::matrix::Csr< ValueType, IndexType >::merge_path	
749	gko::device_matrix_data< ValueType, IndexType
gko::matrix::Csr< ValueType, IndexType >::sparselib,	>, 549

create_from_symm_composition gko::experimental::factorization::Factorization<	difference gko::Timer, 878 difference_async
create_nested_summary	gko::CpuTimer, 482
gko::log::ProfilerHook, 806	gko::CudaTimer, 508
create_nvtx	gko::DpcppTimer, 572
gko::log::ProfilerHook, 807	gko::HipTimer, 652
create_real_view	gko::Timer, 879
gko::experimental::distributed::Vector< ValueType	dim
>, 898	gko::dim< Dimensionality, DimensionType >, 561
gko::matrix::Dense< ValueType >, 528, 529	DimensionMismatch
create_submatrix	gko::DimensionMismatch, 564
gko::matrix::Csr< ValueType, IndexType >, 491,	DPC++ Executor, 280
492	dpcpp_version
gko::matrix::Dense< ValueType >, 529	gko::version_info, 906
create_summary	EII
gko::log::ProfilerHook, 807	gko::matrix::Ell< ValueType, IndexType >, 575
create_tau	ell_col_at
gko::log::ProfilerHook, 807	gko::matrix::Hybrid< ValueType, IndexType >,
create_time_point	655, 656
gko::Timer, 878	ell val at
create_view_of	gko::matrix::Hybrid< ValueType, IndexType >,
gko::matrix::Dense< ValueType >, 530	656, 657
create_with_config_of	empty out
gko::experimental::distributed::Vector< ValueType	gko::device_matrix_data< ValueType, IndexType
>, 898 gko::matrix::Dense< ValueType >, 530	>, 549
create_with_type_of	EnableDefaultCriterionFactory
gko::experimental::distributed::Vector< ValueType	gko::stop, 386
>, 898, 899	EnableDefaultLinOpFactory
gko::matrix::Dense< ValueType >, 530, 531	Linear Operators, 296
criterion	EnableDefaultReorderingBaseFactory
gko::log, 375	gko::reorder, 380
Csr	EnableIterativeBase
gko::matrix::Csr< ValueType, IndexType >, 489	gko::solver::EnableIterativeBase< DerivedType >,
CublasError	590
gko::CublasError, 502	EnablePreconditionable
CUDA Executor, 279	gko::solver::EnablePreconditionable< DerivedType
cuda version	>, 598
gko::version_info, 905	EnableSolverBase
CudaError	gko::solver::EnableSolverBase< DerivedType, Ma-
gko::CudaError, 504	trixType >, 601
CufftError	environment
gko::CufftError, 509	gko::experimental::mpi::environment, 602
CurandError	Error
gko::CurandError, 510	gko::Error, 603
CusparseError	executor_deleter
gko::CusparseError, 513	gko::executor_deleter< T >, 612
cycle	Executors, 281
gko::solver::multigrid, 384	GKO_REGISTER_HOST_OPERATION, 282
5	GKO_REGISTER_OPERATION, 283
Dense	extract_diagonal
gko::matrix::Dense< ValueType >, 518	gko::DiagonalExtractable $<$ ValueType $>$, 558
device_matrix_data	gko::matrix::Coo< ValueType, IndexType >, 478
gko::device_matrix_data< ValueType, IndexType	gko::matrix::Csr< ValueType, IndexType >, 492
>, 547, 548	gko::matrix::Dense< ValueType >, 532
diag	gko::matrix::Ell< ValueType, IndexType >, 577
gko::matrix_data< ValueType, IndexType >, 744-	gko::matrix::Fbcsr< ValueType, IndexType >, 621
746	gko::matrix::Hybrid< ValueType, IndexType >, 657

$\label{eq:gko:matrix::Sellp} $$ gko::matrix::Sellp< ValueType, IndexType>, 842 $$ extract_diagonal_linop$	gko::experimental::mpi::window< ValueType >, 912
gko::DiagonalExtractable < ValueType >, 558	get_agg
gko::DiagonalLinOpExtractable, 559	gko::multigrid::Pgm< ValueType, IndexType >, 792
Factorizations, 285 factory	get_approximate_inverse gko::preconditioner::lsai< IsaiType, ValueType, In-
gko::log, 375	dexType >, 703
Fbcsr	get_basis gko::Perturbation< ValueType >, 791
gko::matrix::Fbcsr< ValueType, IndexType >, 618,	get_block_offset
619 fence	gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 418
gko::experimental::mpi::window< ValueType >, 910	get_block_size gko::matrix::Fbcsr< ValueType, IndexType >, 621
fetch_and_op	get blocks
gko::experimental::mpi::window< ValueType >, 911	gko::preconditioner::Jacobi< ValueType, Index- Type >, 709
fill	get_closest_numa
gko::array< ValueType >, 404	gko::CudaExecutor, 506
gko::experimental::distributed::Vector< ValueType >, 899	gko::HipExecutor, 647
gko::matrix::Dense< ValueType >, 532	get_closest_pus
flush	gko::CudaExecutor, 506
gko::experimental::mpi::window< ValueType >,	gko::HipExecutor, 647 get_coarse_op
911	gko::multigrid::EnableMultigridLevel< ValueType
flush_local	>, 593
gko::experimental::mpi::window< ValueType >, 911	gko::multigrid::MultigridLevel, 757
free	get_coarsest_solver
gko::Executor, 608	gko::solver::Multigrid, 754
	get_coefficients gko::Combination< ValueType >, 432
gather	get_col_idxs
gko::experimental::mpi::communicator, 444 gather_v	gko::device_matrix_data< ValueType, IndexType
gko::experimental::mpi::communicator, 444	>, 550
generate	gko::matrix::Coo< ValueType, IndexType >, 478
gko::AbstractFactory< AbstractProductType, Com-	gko::matrix::Csr< ValueType, IndexType >, 493
ponentsType >, 393	gko::matrix::Ell< ValueType, IndexType >, 577 gko::matrix::Fbcsr< ValueType, IndexType >, 621
gko::experimental::factorization::Cholesky< Value-	gko::matrix::Fbcsi< value Type, Index Type >, 821 gko::matrix::Sellp< Value Type, Index Type >, 843
Type, IndexType >, 427 gko::experimental::factorization::Lu< ValueType,	gko::matrix::SparsityCsr< ValueType, IndexType >, 856
IndexType >, 724 gko::experimental::reorder::Amd< IndexType >,	get_communicator
395 get	gko::experimental::distributed::DistributedBase, 567
gko::cuda_stream, 503	get_complex_subspace
gko::experimental::mpi::communicator, 445	gko::solver::ldr< ValueType >, 674
gko::experimental::mpi::contiguous_type, 467	get_composition
gko::experimental::mpi::request, 826	gko::UseComposition< ValueType >, 887
gko::experimental::mpi::status, 862	<pre>get_conditioning gko::preconditioner::Jacobi< ValueType, Index-</pre>
gko::experimental::mpi::window< ValueType >, 912	Type >, 709
gko::hip_stream, 644	get_const_agg gko::multigrid::Pam< ValueType_IndexType > 793
gko::log::Record, 822 gko::ptr_param< T >, 810	gko::multigrid::Pgm< ValueType, IndexType >, 793 get_const_col_idxs
gko::version_info, 905	gko::device_matrix_data< ValueType, IndexType
get_accessor	>, 550
gko::range< Accessor >, 815	gko::matrix::Coo< ValueType, IndexType >, 478
get_accumulate	gko::matrix::Csr< ValueType, IndexType >, 493

gko::matrix::Ell< ValueType, IndexType >, 577 gko::matrix::Fbcsr< ValueType, IndexType >, 621	<pre>get_coo_num_stored_elements gko::matrix::Hybrid< ValueType, IndexType >, 660</pre>
gko::matrix::Sellp< ValueType, IndexType >, 843	get_coo_row_idxs
gko::matrix::SparsityCsr< ValueType, IndexType >, 857	gko::matrix::Hybrid< ValueType, IndexType >, 660 get_coo_values
get_const_coo_col_idxs	gko::matrix::Hybrid< ValueType, IndexType >, 660
gko::matrix::Hybrid< ValueType, IndexType >, 657	get_core
get_const_coo_row_idxs	gko::machine_topology, 726
gko::matrix::Hybrid< ValueType, IndexType >, 658	get_count
get_const_coo_values	gko::experimental::mpi::status, 862
gko::matrix::Hybrid< ValueType, IndexType >, 658	get_cublas_handle
get_const_data	gko::CudaExecutor, 506
gko::array< ValueType >, 404	get_cusparse_handle
get_const_ell_col_idxs	gko::CudaExecutor, 506
gko::matrix::Hybrid< ValueType, IndexType >, 658	get_cycle
get_const_ell_values	gko::solver::Multigrid, 754
gko::matrix::Hybrid< ValueType, IndexType >, 659	get_data
get_const_local_values	gko::array< ValueType >, 405
gko::experimental::distributed::Vector< ValueType	get_deterministic
>, 900	gko::solver::ldr< ValueType >, 674
get_const_permutation	get_device_id
gko::matrix::Permutation < IndexType >, 788	gko::DpcppExecutor, 569
get_const_row_idxs	get_device_type
gko::device_matrix_data< ValueType, IndexType	gko::DpcppExecutor, 569
>, 550	get_dynamic_type
gko::matrix::Coo< ValueType, IndexType >, 479 gko::matrix::RowGatherer< IndexType >, 833	gko::name_demangling, 377
get_const_row_ptrs	get_ell gko::matrix::Hybrid< ValueType, IndexType >, 660
gko::matrix::Csr< ValueType, IndexType >, 493	get_ell_col_idxs
gko::matrix::Fbcsr< ValueType, IndexType >, 622	gko::matrix::Hybrid< ValueType, IndexType >, 661
gko::matrix::SparsityCsr< ValueType, IndexType	get_ell_num_stored_elements
>, 857	gko::matrix::Hybrid< ValueType, IndexType >, 661
get_const_slice_lengths	get_ell_num_stored_elements_per_row
gko::matrix::Sellp< ValueType, IndexType >, 843	gko::matrix::Hybrid< ValueType, IndexType >, 661
get_const_slice_sets	gko::matrix::Hybrid< ValueType, IndexType
gko::matrix::Sellp< ValueType, IndexType >, 844	>::strategy_type, 868
get_const_srow	get_ell_stride
gko::matrix::Csr< ValueType, IndexType >, 493	gko::matrix::Hybrid< ValueType, IndexType >, 661
get_const_value	get_ell_values
gko::matrix::SparsityCsr< ValueType, IndexType	gko::matrix::Hybrid< ValueType, IndexType >, 662
>, 857	get_executor
get_const_values	gko::array< ValueType >, 406
gko::device_matrix_data< ValueType, IndexType >, 551	gko::device_matrix_data< ValueType, IndexType >, 551
gko::matrix::Coo< ValueType, IndexType >, 479	gko::index_set< IndexType >, 692
gko::matrix::Csr< ValueType, IndexType >, 494	gko::PolymorphicObject, 799
gko::matrix::Dense< ValueType >, 533	get_fine_op
gko::matrix::Diagonal < ValueType >, 556	gko::multigrid::EnableMultigridLevel< ValueType
gko::matrix::Ell< ValueType, IndexType >, 578	>, 594
gko::matrix::Fbcsr< ValueType, IndexType >, 622	gko::multigrid::MultigridLevel, 757
gko::matrix::Sellp< ValueType, IndexType >, 844	get_global_block_offset
get_coo	gko::preconditioner::block_interleaved_storage_scheme<
gko::matrix::Hybrid< ValueType, IndexType >, 659	IndexType >, 418
get_coo_col_idxs	get_global_index
gko::matrix::Hybrid< ValueType, IndexType >, 659	gko::index_set< IndexType >, 692
get_coo_nnz	get_group_offset
gko::matrix::Hybrid< ValueType, IndexType >::strategy_type, 868	gko::preconditioner::block_interleaved_storage_scheme < IndexType >, 418

get_group_size	get_mid_smoother_list
gko::preconditioner::block_interleaved_storage_sc	
IndexType >, 419	get_name
get_handle_name	gko::matrix::Csr< ValueType, IndexType >::strategy_type,
gko::log::Papi< ValueType >, 769	870
get_hipblas_handle	gko::Operation, 765
gko::HipExecutor, 647	get_non_local_matrix
get_hipsparse_handle	gko::experimental::distributed::Matrix< ValueType,
gko::HipExecutor, 648	LocalIndexType, GlobalIndexType >, 732
get_id	get_nonpreserving
gko::stopping_status, 864	gko::precision_reduction, 803
get_implicit_sq_resnorm	get_num_block_cols
gko::log::Convergence< ValueType >, 470	gko::matrix::Fbcsr< ValueType, IndexType >, 622
get_instance	get_num_block_rows
gko::machine_topology, 727	gko::matrix::Fbcsr< ValueType, IndexType >, 623
get_inverse_permutation	get_num_blocks
gko::reorder::Rcm< ValueType, IndexType >, 818	gko::preconditioner::Jacobi< ValueType, Index-
get_kappa	Type >, 710
gko::solver::ldr< ValueType >, 674	get_num_columns
get_krylov_dim	gko::matrix::Hybrid< ValueType, IndexType
gko::solver::CbGmres< ValueType >, 422	>::column_limit, 430
gko::solver::Gcr< ValueType >, 639	get_num_computing_units
gko::solver::Gmres< ValueType >, 641	gko::DpcppExecutor, 571
get_l_solver	get_num_cores
gko::preconditioner::lc< LSolverType, IndexType	
>, 668	get_num_devices
gko::preconditioner::llu< LSolverType, USolver	
Type, ReverseApply, IndexType >, 680	get_num_elems
get_lh_solver	gko::array< ValueType >, 406
gko::preconditioner::lc< LSolverType, IndexType	
>, 668	>, 551
get_local_index	gko::index_set< IndexType >, 693
gko::index_set< IndexType >, 693	get_num_empty_parts
get_local_matrix	gko::experimental::distributed::Partition< LocalIn-
gko::experimental::distributed::Matrix< ValueType	
LocalIndexType, GlobalIndexType >, 732	get num iterations
get_local_values	gko::log::Convergence< ValueType >, 470
gko::experimental::distributed::Vector< ValueType	
>, 900	gko::matrix::SparsityCsr< ValueType, IndexType
get_local_vector	>, 858
gko::experimental::distributed::Vector< ValueType	•
>, 900	gko::machine_topology, 727
get_loggers	get num parts
gko::log::Loggable, 719	gko::experimental::distributed::Partition< LocalIn-
get_master	dexType, GlobalIndexType >, 777
gko::CudaExecutor, 506, 507	get_num_pci_devices
gko::DpcppExecutor, 570	gko::machine_topology, 727
gko::Executor, 608	get_num_pus
gko::HipExecutor, 648	gko::machine_topology, 728
gko::OmpExecutor, 763	get_num_ranges
get_max_subgroup_size	gko::experimental::distributed::Partition< LocalIn-
gko::DpcppExecutor, 570	dexType, GlobalIndexType >, 777
get_max_workgroup_size	get_num_srow_elements
gko::DpcppExecutor, 570	gko::matrix::Csr< ValueType, IndexType >, 494
get_max_workitem_sizes	get_num_stored_blocks
gko::DpcppExecutor, 571	gko::matrix::Fbcsr< ValueType, IndexType >, 623
get_mg_level_list	get_num_stored_elements
gko::solver::Multigrid, 755	gko::matrix::Coo< ValueType, IndexType >, 479
J == = = = =	5 · · · · · · · · · · · · · · · · · · ·

gko::matrix::Csr $<$ ValueType, IndexType $>$, 494	gko::precision_reduction, 803
gko::matrix::Dense< ValueType >, 533	get_projector
gko::matrix::EII< ValueType, IndexType >, 578	gko::Perturbation < ValueType >, 791
gko::matrix::Fbcsr< ValueType, IndexType >, 623	get_prolong_op
gko::matrix::Hybrid< ValueType, IndexType >, 662	gko::multigrid::EnableMultigridLevel< ValueType
gko::matrix::Sellp< ValueType, IndexType >, 844 gko::matrix_assembly_data< ValueType, Index-	>, 594 gko::multigrid::MultigridLevel, 757
Type >, 737	get_provided_thread_support
gko::preconditioner::Jacobi< ValueType, Index-	gko::experimental::mpi::environment, 602
Type >, 710	get_pu
get_num_stored_elements_per_row	gko::machine_topology, 729
gko::matrix::Ell< ValueType, IndexType >, 578	get_range_bounds
get_num_subsets	gko::experimental::distributed::Partition< LocalIn-
gko::index_set< IndexType >, 694	dexType, GlobalIndexType >, 779
get_operator_at	get_range_starting_indices
gko::UseComposition< ValueType >, 887	gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 779
get_operators gko::Combination< ValueType >, 433	get_ratio
gko::Composition< ValueType >, 465	gko::matrix::Hybrid< ValueType, IndexType
get_ordered_data	>::imbalance_bounded_limit, 683
gko::matrix_assembly_data< ValueType, Index-	get_residual
Type >, 737	gko::log::Convergence< ValueType >, 470
get_parameters	get_residual_norm
gko::EnableDefaultFactory< ConcreteFactory, Pro-	gko::log::Convergence < ValueType >, 471
ductType, ParametersType, PolymorphicBase	get_restrict_op
>, 587	gko::multigrid::EnableMultigridLevel< ValueType
get_part_ids	>, 594
gko::experimental::distributed::Partition< LocalIn-	gko::multigrid::MultigridLevel, 757
<pre>dexType, GlobalIndexType >, 777 get_part_size</pre>	<pre>get_row_idxs gko::device_matrix_data< ValueType, IndexType</pre>
gko::experimental::distributed::Partition< LocalIn-	>, 552
dexType, GlobalIndexType >, 778	gko::matrix::Coo< ValueType, IndexType >, 480
get_part_sizes	gko::matrix::RowGatherer< IndexType >, 833
gko::experimental::distributed::Partition< LocalIn-	get_row_ptrs
dexType, GlobalIndexType >, 778	gko::matrix::Csr< ValueType, IndexType >, 495
get_pci_device	gko::matrix::Fbcsr< ValueType, IndexType >, 623
gko::machine_topology, 728	gko::matrix::SparsityCsr< ValueType, IndexType
get_percentage	>, 858
gko::matrix::Hybrid< ValueType, IndexType	get_scalar
>::imbalance_bounded_limit, 683 gko::matrix::Hybrid< ValueType, IndexType	gko::Perturbation< ValueType >, 791 get significant bit
>::imbalance_limit, 685	gko, 333
gko::matrix::Hybrid< ValueType, IndexType	get_size
>::minimal_storage_limit, 751	gko::device_matrix_data< ValueType, IndexType
get_permutation	>, 552
gko::matrix::Permutation < IndexType >, 789	gko::experimental::distributed::Partition< LocalIn-
gko::reorder::Rcm< ValueType, IndexType >, 818	dexType, GlobalIndexType >, 779
get_permutation_size	gko::index_set< IndexType >, 694
gko::matrix::Permutation < IndexType >, 789	gko::matrix_assembly_data< ValueType, Index-
get_permute_mask	Type >, 737
gko::matrix::Permutation< IndexType >, 789	get_slice_lengths
get_post_smoother_list gko::solver::Multigrid, 755	gko::matrix::Sellp< ValueType, IndexType >, 845 get_slice_sets
get_pre_smoother_list	gko::matrix::Sellp< ValueType, IndexType >, 845
gko::solver::Multigrid, 755	get_slice_size
get_preconditioner	gko::matrix::Sellp< ValueType, IndexType >, 845
gko::Preconditionable, 804	get_solver
get_preserving	gko::solver::lr< ValueType >, 700

<pre>get_srow gko::matrix::Csr< ValueType, IndexType >, 495</pre>	gko::matrix::Fbcsr< ValueType, IndexType >, 624 gko::matrix::Sellp< ValueType, IndexType >, 846
get_static_type	get_walltime
gko::name_demangling, 378	gko::experimental::mpi, 372
get_stop_criterion_factory	get_window
gko::solver::IterativeBase, 706	gko::experimental::mpi::window< ValueType >
get_storage_precision	913
gko::solver::CbGmres< ValueType >, 422	give
get_storage_scheme	gko, 334
gko::preconditioner::Jacobi< ValueType, Index-	gko, 309
Type >, 710	abs, 324
get_strategy	allocation_mode, 323
gko::matrix::Csr< ValueType, IndexType >, 495	array, 323
gko::matrix::Hybrid< ValueType, IndexType >,	as, 324–327
662, 663	automatic, 324
get_stream	ceildiv, 328
gko::CudaExecutor, 507	clone, 328, 329
get_stride	conj, <mark>330</mark>
gko::matrix::Dense< ValueType >, 533	coordinate, 323
gko::matrix::Ell $<$ ValueType, IndexType $>$, 579	copy_and_convert_to, 330-332
gko::preconditioner::block_interleaved_storage_sche	eme< get_significant_bit, 333
IndexType >, 419	get_superior_power, 333
get_stride_factor	give, 334
gko::matrix::Sellp< ValueType, IndexType >, 846	highest_precision, 320
get_subgroup_sizes	imag, 334
gko::DpcppExecutor, 572	is_complex, 335
get_subsets_begin	is_complex_or_scalar, 335
gko::index_set< IndexType >, 694	is_complex_or_scalar_s, 321
get_subsets_end	is_complex_s, 321
gko::index_set< IndexType >, 694	is_finite, 335, 336
get_subspace_dim	is_nan, 336, 337
gko::solver::ldr< ValueType >, 674	is_nonzero, 337
get_superior_power	is_zero, 338
gko, 333	layout_type, 323
get_superset_indices	lend, 338, 339
gko::index_set< IndexType >, 695	log_propagation_mode, 323
get_system_matrix	make_array_view, 339
gko::multigrid::FixedCoarsening< ValueType, In-	make_const_array_view, 340
dexType >, 637	make_const_dense_view, 340
gko::multigrid::Pgm< ValueType, IndexType >, 793	make_dense_view, 342
get_total_cols	make_temporary_clone, 342
gko::matrix::Sellp< ValueType, IndexType >, 846	make_temporary_conversion, 343
get_u_solver	make_temporary_output_clone, 344
gko::preconditioner::llu< LSolverType, USolver-	max, 344
Type, ReverseApply, IndexType >, 680	min, 345
get_value	mixed_precision_dispatch, 345
gko::matrix::SparsityCsr< ValueType, IndexType >, 858	mixed_precision_dispatch_real_complex, 346 nan, 346, 347
gko::matrix_assembly_data< ValueType, Index-	never, 324
Type >, 737	one, 347
get_values	operator!=, 347, 348
gko::device_matrix_data< ValueType, IndexType	operator<<, 349
>, 552	operator==, 351
gko::matrix::Coo< ValueType, IndexType >, 480	pi, 351
gko::matrix::Csr< ValueType, IndexType >, 496	precision_dispatch, 352
gko::matrix::Dense< ValueType >, 533	precision_dispatch_real_complex, 352, 353
gko::matrix::Diagonal< ValueType >, 556	previous_precision, 321
gko::matrix::Ell $<$ ValueType, IndexType $>$, 579	read, 353

read_binary, 354	gko::Combination < ValueType >, 431
read_binary_raw, 354	Combination, 432
read_generic, 355	conj_transpose, 432
read_generic_raw, 356	get_coefficients, 432
read_raw, 356	get_operators, 433
real, 357 reduce_add, 358	operator=, 433 transpose, 433
remove complex, 322	gko::Composition < ValueType >, 463
round_down, 359	Composition, 464
round_up, 359	conj_transpose, 465
safe_divide, 359	get_operators, 465
share, 360	operator=, 465
squared_norm, 361	transpose, 466
to_complex, 322	gko::ConvertibleTo< ResultType >, 471
to_real, 322	convert_to, 472
transpose, 361	move_to, 473
unit_root, 361	gko::CpuTimer, 482
with_matrix_type, 362	difference_async, 482
write, 363	gko::cpx_real_type< T >, 483
write_binary, 363	type, 483
write_binary_raw, 364	gko::CublasError, 501
write_raw, 364	CublasError, 502
zero, 365	gko::cuda_stream, 502
gko::AbsoluteComputable, 391	get, 503
compute_absolute_linop, 391	gko::CudaError, 503
gko::AbstractFactory< AbstractProductType, Compo-	CudaError, 504
nentsType >, 392	gko::CudaExecutor, 504
generate, 393	create, 505
gko::accessor, 366	get_closest_numa, 506
gko::accessor::row_major< ValueType, Dimensionality	get_closest_pus, 506
>, 828	get_cublas_handle, 506
copy_from, 830	get_cusparse_handle, 506
length, 830	get_master, 506, 507
operator(), 831	get_stream, 507
gko::AllocationError, 393	gko::CudaTimer, 507
AllocationError, 394 gko::amd_device, 395	difference_async, 508 gko::CufftError, 509
gko::are_all_integral < Args >, 396	CufftError, 509
gko::array< ValueType >, 397	gko::CurandError, 509
array, 399–403	CurandError, 510
as const view, 403	gko::CusparseError, 512
as view, 403	CusparseError, 513
clear, 403	gko::default_converter< S, R >, 513
const_view, 403	operator(), 514
fill, 404	gko::device_matrix_data< ValueType, IndexType >, 546
get const data, 404	copy_to_host, 549
get_data, 405	create_from_host, 549
get_executor, 406	device_matrix_data, 547, 548
get_num_elems, 406	empty_out, 549
is_owning, 406	get_col_idxs, 550
operator=, 407, 408	get_const_col_idxs, 550
resize_and_reset, 409	get_const_row_idxs, 550
set_executor, 409	get_const_values, 551
view, 410	get_executor, 551
gko::BadDimension, 412	get_num_elems, 551
BadDimension, 412	get_row_idxs, 552
gko::BlockSizeError< IndexType >, 420	get_size, 552
BlockSizeError, 420	get_values, 552

remove_zeros, 552	add_logger, 606
resize_and_reset, 553	alloc, 606
sum_duplicates, 553	copy, 606
gko::device_matrix_data< ValueType, IndexType	copy_from, 607
>::arrays, 410	copy_val_to_host, 607
gko::DiagonalExtractable < ValueType >, 558	free, 608
extract_diagonal, 558	get_master, 608
extract_diagonal_linop, 558	memory_accessible, 609
gko::DiagonalLinOpExtractable, 559	remove_logger, 609
extract_diagonal_linop, 559	run, 609, 610
gko::dim< Dimensionality, DimensionType >, 560	set_log_propagation_mode, 610
dim, 561	should_propagate_log, 611
operator bool, 561	gko::executor_deleter< T >, 611
operator<<, 563	executor_deleter, 612
operator*, 563	operator(), 612
operator==, 563	gko::experimental::distributed, 366
operator[], 562	comm_index_type, 367
gko::DimensionMismatch, 564	make_temporary_conversion, 367, 368
DimensionMismatch, 564	precision_dispatch, 369
gko::DpcppExecutor, 568	precision_dispatch_real_complex, 369, 370
create, 569	gko::experimental::distributed::DistributedBase, 566
get_device_id, 569	get_communicator, 567
get_device_type, 569	operator=, 567
get_master, 570	gko::experimental::distributed::Matrix< ValueType, Lo-
get_max_subgroup_size, 570	callndexType, GloballndexType >, 729
get_max_workgroup_size, 570	get_local_matrix, 732
get_max_workitem_sizes, 571	get_non_local_matrix, 732
get_num_computing_units, 571	Matrix, 731, 732
get_num_devices, 571	operator=, 732, 733
get_subgroup_sizes, 572	read_distributed, 733–735
gko::DpcppTimer, 572	gko::experimental::distributed::Partition< LocalIndex-
difference_async, 572	Type, GlobalIndexType >, 774
gko::enable_parameters_type< ConcreteParameter-	build_from_contiguous, 775
sType, Factory >, 582	build_from_global_size_uniform, 776
on, 583	build_from_mapping, 776
gko::EnableAbsoluteComputation< AbsoluteLinOp >,	get_num_empty_parts, 777
583	get_num_parts, 777
compute_absolute, 584	get_num_ranges, 777
compute_absolute_linop, 584	get_part_ids, 777
gko::EnableAbstractPolymorphicObject< AbstractOb-	get_part_size, 778
ject, PolymorphicBase >, 585	get_part_sizes, 778
gko::EnableCreateMethod< ConcreteType >, 586	get_range_bounds, 779
gko::EnableDefaultFactory< ConcreteFactory, Product-	get_range_starting_indices, 779
Type, ParametersType, PolymorphicBase >,	get_size, 779
586	has_connected_parts, 780
create, 587	has_ordered_parts, 780
get_parameters, 587	gko::experimental::distributed::preconditioner, 371
gko::EnableLinOp< ConcreteLinOp, PolymorphicBase	gko::experimental::distributed::preconditioner::Schwarz<
>, 591	ValueType, LocalIndexType, GlobalIndexType
gko::EnablePolymorphicAssignment< ConcreteType,	>, 836
ResultType >, 595	gko::experimental::distributed::Vector< ValueType >,
convert_to, 596	889
move_to, 596	add_scaled, 892
gko::EnablePolymorphicObject< ConcreteObject, Poly-	at_local, 892, 893
morphicBase >, 597	compute_absolute, 893
gko::Error, 603	compute_conj_dot, 893, 894
Error, 603	compute_dot, 894, 895
gko::Executor, 604	compute_norm1, 895
	· r···= · · · · · · · · · · · · · · · ·

compute_norm2, 896	node_local_rank, 457
compute_squared_norm2, 896	operator!=, 457
create_const, 897	operator==, 457
create real view, 898	rank, 457
create with config of, 898	recv, 458
create_with_type_of, 898, 899	reduce, 458
fill, 899	scan, 459
get_const_local_values, 900	scatter, 460
get_local_values, 900	scatter v, 460
get_local_vector, 900	send, 462
inv_scale, 901	size, 462
make_complex, 901	synchronize, 463
read_distributed, 901, 902	gko::experimental::mpi::contiguous_type, 466
scale, 902	contiguous_type, 467
sub_scaled, 903	get, 467
gko::experimental::EnableDistributedLinOp< Con-	operator=, 468
creteLinOp, PolymorphicBase >, 588	gko::experimental::mpi::environment, 601
gko::experimental::EnableDistributedPolymorphicObject<	
ConcreteObject, PolymorphicBase >, 589	get_provided_thread_support, 602
gko::experimental::factorization::Cholesky< ValueType,	gko::experimental::mpi::request, 825
IndexType >, 426	get, 826
generate, 427	request, 825
gko::experimental::factorization::Factorization< Value-	wait, 826
Type, IndexType >, 613	gko::experimental::mpi::status, 861
create_from_combined_lu, 614	get, 862
create_from_composition, 614	get_count, 862
create_from_symm_composition, 616	status, 862
unpack, 616	gko::experimental::mpi::type_impl< T >, 881
gko::experimental::factorization::Lu< ValueType, Index-	gko::experimental::mpi::window< ValueType >, 906
Type >, 723	accumulate, 910
generate, 724	fence, 910
gko::experimental::mpi, 371	fetch_and_op, 911
get walltime, 372	flush, 911
map_rank_to_device_id, 372	flush_local, 911
wait_all, 373	get, 912
gko::experimental::mpi::communicator, 434	get_accumulate, 912
all_gather, 439	get window, 913
all_reduce, 439, 440	lock, 913
all_to_all, 441	lock all, 914
all_to_all_v, 442	operator=, 914
broadcast, 443	put, 914
communicator, 438	r_accumulate, 915
gather, 444	r_get, 915
gather_v, 444	r_get_accumulate, 916
get, 445	r_put, 917
i_all_gather, 445	unlock, 917
i_all_reduce, 446, 447	window, 908, 909
i_all_to_all, 447, 448	gko::experimental::reorder, 373
i_all_to_all_v, 449	gko::experimental::reorder::Amd< IndexType >, 394
i_broadcast, 450	generate, 395
i_gather, 451	gko::experimental::reorder::ScaledReordered< Value-
i_gather_v, 452	Type, IndexType >, 835
i_recv, 452	gko::experimental::solver::Direct< ValueType, Index-
i_reduce, 453	Type $>$, 565
i_scan, 454	conj_transpose, 566
i_scatter, 454	transpose, 566
i_scatter_v, 455	gko::factorization, 373
i_scatter_v, 455 i_send, 456	gko::factorization::lc< ValueType, IndexType >, 665
1_00114, 100	gromasionzationino < value typo, maex typo >, 000

gko::factorization::Parllut< ValueType, IndexType >, gg 773 gko::hip_stream, 643 get, 644 gko::HipblasError, 644 HipblasError, 644 gko::HipError, 645	create, 469 get_implicit_sq_resnorm, 470 get_num_iterations, 470 get_residual, 470 get_residual_norm, 471 has_converged, 471 og::criterion_data, 485 og::EnableLogging< ConcreteLoggable, Polymor phicBase >, 592 og::executor_data, 611
	og::iteration_complete_data, 705
create, 647 gko::ld	og::linop_data, 713
	og::linop_factory_data, 714
	og::Loggable, 719
	add_logger, 719
· - · · -	get_loggers, 719
	remove_logger, 720
	og::operation_data, 766
	og::Papi< ValueType >, 768
•	create, 769
· · · · · · · · · · · · · · · · · · ·	get_handle_name, 769
	og::PerformanceHint, 781
·	create, 781
	og::polymorphic_object_data, 794 og::ProfilerHook, 805
	create_nested_summary, 806
	create_nvtx, 807
	create_summary, 807
	create_tau, 807
-	set_object_name, 808
	user_range, 808
get_num_elems, 693 gko::ld	og::ProfilerHook::NestedSummaryWriter, 758
get_num_subsets, 694	write_nested, 758
- -	og::ProfilerHook::SummaryWriter, 874
	write, 874
	og::ProfilerHook::TableSummaryWriter, 875
• - • -	TableSummaryWriter, 875
	write, 876
	write_nested, 876
	og::profiling_scope_guard, 809
	profiling_scope_guard, 809
•	og::Record, 821 create, 821, 822
— -	get, 822
	og::Record::logged_data, 720
	og::Stream< ValueType >, 871
	create, 872
-	machine_topology, 724
	bind_to_core, 725
•	bind_to_cores, 725
linop, 375	bind_to_pu, 726
memory, 375	bind_to_pus, 726
object, 375	get_core, 726
	get_instance, 727
. – – • •	get_num_cores, 727
solver, 375	get_num_numas, 727

get_num_pci_devices, 727	copy, 511
get_num_pus, 728	process, 512
get_pci_device, 728	gko::matrix::Csr< ValueType, IndexType >::load_balance
get_pu, 729	715
gko::matrix, 376	clac_size, 717
gko::matrix::Coo< ValueType, IndexType >, 474	copy, 718
apply2, 475–477	load_balance, 716, 717
compute_absolute, 477	process, 718
create_const, 477	gko::matrix::Csr< ValueType, IndexType >::merge_path,
extract_diagonal, 478	748
get_col_idxs, 478	clac_size, 748
get_const_col_idxs, 478	copy, 749
get_const_row_idxs, 479	process, 749
get_const_values, 479	gko::matrix::Csr< ValueType, IndexType >::sparselib,
get_num_stored_elements, 479	852
get_row_idxs, 480	clac_size, 852
get_values, 480	copy, 853
read, 480, 481	process, 853
write, 481	gko::matrix::Csr< ValueType, IndexType >::strategy_type
gko::matrix::Csr< ValueType, IndexType >, 486	869
column permute, 489	clac_size, 870
compute_absolute, 490	copy, 870
conj_transpose, 490	get_name, 870
create_const, 490	process, 871
create_submatrix, 491, 492	strategy_type, 869
Csr, 489	gko::matrix::Dense< ValueType >, 514
extract_diagonal, 492	add_scaled, 518
get_col_idxs, 493	at, 519, 520
get_const_col_idxs, 493	column_permute, 520–522
get_const_row_ptrs, 493	compute_absolute, 522
get_const_srow, 493	compute_conj_dot, 523
get_const_stow, 493	compute_dot, 523, 524
get_const_values, 494 get_num_srow_elements, 494	compute_norm1, 524
	• —
get_num_stored_elements, 494	compute_norm2, 526
get_row_ptrs, 495	compute_squared_norm2, 526, 527
get_srow, 495	conj_transpose, 527
get_strategy, 495	create_const, 528
get_values, 496	create_const_view_of, 528
inv_scale, 496	create_real_view, 528, 529
inverse_column_permute, 496	create_submatrix, 529
inverse_permute, 497	create_view_of, 530
inverse_row_permute, 497	create_with_config_of, 530
operator=, 498	create_with_type_of, 530, 531
permute, 498	Dense, 518
read, 499	extract_diagonal, 532
row_permute, 500	fill, 532
scale, 500	get_const_values, 533
set_strategy, 500	get_num_stored_elements, 533
transpose, 501	get_stride, 533
write, 501	get_values, 533
gko::matrix::Csr< ValueType, IndexType >::classical,	inv_scale, 534
427	inverse_column_permute, 534, 535
clac_size, 428	inverse_permute, 535–537
copy, 428	inverse_row_permute, 537, 538
process, 428	make_complex, 538
gko::matrix::Csr< ValueType, IndexType >::cusparse,	operator=, 539
510	permute, 539, 540
clac_size, 511	row_gather, 541, 542

row_permute, 543, 544	conj_transpose, 632
scale, 544	transpose, 632
sub_scaled, 545	write, 632, 633
transpose, 545	gko::matrix::Fft3, 634
gko::matrix::Diagonal < ValueType >, 554	conj_transpose, 634
compute_absolute, 555	transpose, 635
conj_transpose, 555	write, 635, 636
create_const, 555	gko::matrix::Hybrid< ValueType, IndexType >, 652
get_const_values, 556	compute_absolute, 655
get_values, 556	ell_col_at, 655, 656
inverse_apply, 556	ell_val_at, 656, 657
rapply, 557	extract_diagonal, 657
transpose, 557	get_const_coo_col_idxs, 657
gko::matrix::Ell< ValueType, IndexType >, 573	get_const_coo_row_idxs, 658
col_at, 575, 576	get_const_coo_values, 658
compute_absolute, 576	get_const_ell_col_idxs, 658
create_const, 576	get_const_ell_values, 659
EII, 575	get_coo, 659
extract_diagonal, 577	get_coo_col_idxs, 659
get_col_idxs, 577	get_coo_num_stored_elements, 660
get_const_col_idxs, 577	get_coo_row_idxs, 660
get_const_values, 578	get_coo_values, 660
get_num_stored_elements, 578	get_ell, 660
get num stored elements per row, 578	
•	get_ell_col_idxs, 661
get_stride, 579	get_ell_num_stored_elements, 661
get_values, 579	get_ell_num_stored_elements_per_row, 661
operator=, 579	get_ell_stride, 661
read, 580	get_ell_values, 662
val_at, 581	get_num_stored_elements, 662
write, 582	get_strategy, 662, 663
gko::matrix::Fbcsr< ValueType, IndexType >, 617	Hybrid, 655
compute_absolute, 619	operator=, 663
conj_transpose, 619	read, 663, 664
convert_to, 619, 620	write, 664
create_const, 620	gko::matrix::Hybrid< ValueType, IndexType >::automatic,
extract_diagonal, 621	411
Fbcsr, 618, 619	compute_ell_num_stored_elements_per_row, 411
get_block_size, 621	gko::matrix::Hybrid< ValueType, IndexType >::column_limit,
get_col_idxs, 621	429
get_const_col_idxs, 621	column_limit, 430
get_const_row_ptrs, 622	compute_ell_num_stored_elements_per_row, 430
get_const_values, 622	get_num_columns, 430
get_num_block_cols, 622	gko::matrix::Hybrid < ValueType, IndexType >::imbalance_bounded_limit,
get_num_block_rows, 623	682
get_num_stored_blocks, 623	compute_ell_num_stored_elements_per_row, 683
get_num_stored_elements, 623	get_percentage, 683
get_row_ptrs, 623	get_ratio, 683
get_values, 624	gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit,
is_sorted_by_column_index, 624	684
operator=, 624	compute_ell_num_stored_elements_per_row, 685
read, 625	get_percentage, 685
transpose, 626	imbalance_limit, 684
write, 626	gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit,
gko::matrix::Fft, 628	750
conj_transpose, 629	compute_ell_num_stored_elements_per_row, 751
transpose, 629	get_percentage, 751
write, 629, 630	gko::matrix::Hybrid< ValueType, IndexType >::strategy_type,
gko::matrix::Fft2, 631	866

compute_ell_num_stored_elements_per_row, 867	gko::matrix_assembly_data< ValueType, IndexType >,
compute_hybrid_config, 868	735
get_coo_nnz, 868	add_value, 736
get_ell_num_stored_elements_per_row, 868	contains, 736 get num stored elements, 737
gko::matrix::ldentity< ValueType >, 670	· · ·
conj_transpose, 670	get_ordered_data, 737
transpose, 671	get_size, 737
gko::matrix::ldentityFactory< ValueType >, 671	get_value, 737 set_value, 738
create, 672	gko::matrix_data< ValueType, IndexType >, 738
gko::matrix::Permutation< IndexType >, 787	cond, 743
create_const, 788	diag, 744–746
get_const_permutation, 788	matrix_data, 740–742
get_permutation, 789	nonzeros, 747
get_permutation_size, 789	sum_duplicates, 747
get_permute_mask, 789	gko::matrix_data_entry< ValueType, IndexType >, 747
set_permute_mask, 789	gko::MetisError, 750
gko::matrix::RowGatherer< IndexType >, 832	MetisError, 750
create_const, 833	gko::MpiError, 752
get_const_row_idxs, 833	MpiError, 752
get_row_idxs, 833	gko::multigrid, 377
gko::matrix::Sellp< ValueType, IndexType >, 839	gko::multigrid::EnableMultigridLevel< ValueType >, 593
col_at, 841, 842	get_coarse_op, 593
compute_absolute, 842	get_tine_op, 594
extract_diagonal, 842	get_nne_op, 394 get_prolong_op, 594
get_col_idxs, 843	get_restrict_op, 594
get_const_col_idxs, 843	gko::multigrid::FixedCoarsening< ValueType, IndexType
get_const_slice_lengths, 843	>, 636
get_const_slice_sets, 844	get_system_matrix, 637
get_const_values, 844	gko::multigrid::MultigridLevel, 756
get_num_stored_elements, 844	get_coarse_op, 757
get_slice_lengths, 845	get_soarse_op, 757 get_fine_op, 757
get_slice_sets, 845	get_prolong_op, 757
get_slice_size, 845	get_restrict_op, 757
get_stride_factor, 846	gko::multigrid::Pgm< ValueType, IndexType >, 792
get_total_cols, 846	get_agg, 792
get_values, 846	get_const_agg, 793
operator=, 846, 847	get_system_matrix, 793
read, 847, 848	gko::name_demangling, 377
Sellp, 841	get_dynamic_type, 377
val_at, 848	get_static_type, 378
write, 849	gko::NotCompiled, 759
gko::matrix::SparsityCsr< ValueType, IndexType >, 853	NotCompiled, 759
conj_transpose, 855	gko::NotImplemented, 760
create_const, 856	NotImplemented, 760
get_col_idxs, 856	gko::NotSupported, 761
get_const_col_idxs, 857	NotSupported, 761
get_const_row_ptrs, 857	gko::null_deleter< T >, 761
get_const_value, 857	operator(), 762
get_num_nonzeros, 858	gko::nvidia_device, 762
get_row_ptrs, 858	gko::OmpExecutor, 763
get_value, 858	get_master, 763
operator=, 859	gko::Operation, 764
read, 859, 860	get_name, 765
SparsityCsr, 855	gko::OutOfBoundsError, 766
to_adjacency_matrix, 860	OutOfBoundsError, 766
transpose, 860	gko::OverflowError, 767
write, 861	OverflowError, 767
	Oromone inoi, ror

gko::Permutable < IndexType >, 782	Isai, 703
column_permute, 783	operator=, 704
inverse_column_permute, 783	transpose, 704
inverse_permute, 784 inverse_row_permute, 784	gko::preconditioner::Jacobi< ValueType, IndexType >, 707
	conj_transpose, 709
permute, 786	
row_permute, 786	convert_to, 709
gko::Perturbation < ValueType >, 790	get_blocks, 709 get_conditioning, 709
get_basis, 791	get_conditioning, 709 get_num_blocks, 710
get_projector, 791	get_num_stored_elements, 710
get_scalar, 791 gko::PolymorphicObject, 794	get_storage_scheme, 710
	Jacobi, 708
clear, 795	
clone, 795, 796	move_to, 711 operator=, 711
copy_from, 796–798	·
create_default, 798, 799	transpose, 712
get_executor, 799	write, 712
move_from, 800	gko::ptr_param< T >, 810
gko::precision_reduction, 800	get, 810
autodetect, 802	operator bool, 811
common, 802	operator*, 811
get_nonpreserving, 803	operator->, 811
get_preserving, 803	gko::range < Accessor >, 812
operator storage_type, 803	get_accessor, 815
precision_reduction, 802	length, 815
gko::Preconditionable, 804	operator(), 815
get_preconditioner, 804	operator->, 816
set_preconditioner, 804	operator=, 816, 817
gko::preconditioner, 378	range, 814
ICAL TAND 374	gko::ReadableFromMatrixData< ValueType, IndexType
isai_type, 379	-
gko::preconditioner::block_interleaved_storage_scheme<	>, 819
gko::preconditioner::block_interleaved_storage_scheme < IndexType $>$, 416	>, 819 read, 819, 820
gko::preconditioner::block_interleaved_storage_scheme IndexType >, 416 compute_storage_space, 417	>, 819 read, 819, 820 gko::ReferenceExecutor, 823
gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 416 compute_storage_space, 417 get_block_offset, 418	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823
gko::preconditioner::block_interleaved_storage_scheme <pre>IndexType >, 416 compute_storage_space, 417 get_block_offset, 418 get_global_block_offset, 418</pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379
gko::preconditioner::block_interleaved_storage_scheme <pre> IndexType >, 416 compute_storage_space, 417 get_block_offset, 418 get_global_block_offset, 418 get_group_offset, 418</pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380
gko::preconditioner::block_interleaved_storage_scheme <pre> IndexType >, 416 compute_storage_space, 417 get_block_offset, 418 get_global_block_offset, 418 get_group_offset, 418 get_group_size, 419</pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818
gko::preconditioner::block_interleaved_storage_scheme < IndexType >, 416 compute_storage_space, 417 get_block_offset, 418 get_global_block_offset, 418 get_group_offset, 418 get_group_size, 419 get_stride, 419 group_power, 420 gko::preconditioner::lc< LSolverType, IndexType >, 665	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380
gko::preconditioner::block_interleaved_storage_scheme <	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	>, 819 read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382 gko::solver::ApplyWithInitialGuess, 396
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382 gko::solver::ApplyWithInitialGuess, 396 gko::solver::Bicg< ValueType >, 413
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382 gko::solver::ApplyWithInitialGuess, 396 gko::solver::Bicg< ValueType >, 413 apply_uses_initial_guess, 413
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBase gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382 gko::solver::ApplyWithInitialGuess, 396 gko::solver::Bicg< ValueType >, 413 apply_uses_initial_guess, 413 conj_transpose, 414
gko::preconditioner::block_interleaved_storage_scheme <pre></pre>	read, 819, 820 gko::ReferenceExecutor, 823 run, 823 gko::reorder, 379 EnableDefaultReorderingBaseFactory, 380 gko::reorder::Rcm< ValueType, IndexType >, 817 get_inverse_permutation, 818 get_permutation, 818 gko::reorder::ReorderingBase< IndexType >, 824 gko::reorder::ReorderingBaseArgs, 825 gko::ScaledIdentityAddable, 834 add_scaled_identity, 834 gko::scoped_device_id_guard, 836 scoped_device_id_guard, 837–839 gko::solver, 380 build_smoother, 383 initial_guess_mode, 382 provided, 382 rhs, 382 trisolve_algorithm, 382 zero, 382 gko::solver::ApplyWithInitialGuess, 396 gko::solver::Bicg< ValueType >, 413 apply_uses_initial_guess, 413

apply_uses_initial_guess, 415	get_kappa, 674
conj_transpose, 415	get_subspace_dim, 674
transpose, 416	set_complex_subpsace, 675
gko::solver::CbGmres < ValueType >, 421	set_deterministic, 675
get_krylov_dim, 422	set_kappa, 675
get_storage_precision, 422	set_subspace_dim, 676
set_krylov_dim, 422	transpose, 676
gko::solver::Cg< ValueType >, 423	gko::solver::lr< ValueType >, 698
apply_uses_initial_guess, 423	apply_uses_initial_guess, 699
conj_transpose, 424	conj_transpose, 700
transpose, 424	get_solver, 700
gko::solver::Cgs< ValueType >, 424	Ir, 699
apply_uses_initial_guess, 425	operator=, 700
conj_transpose, 425	set_solver, 701
transpose, 426	transpose, 701
gko::solver::EnableApplyWithInitialGuess< Derived-	gko::solver::IterativeBase, 705
Type >, 585	get_stop_criterion_factory, 706
gko::solver::EnableIterativeBase< DerivedType >, 589	set_stop_criterion_factory, 706
EnableIterativeBase, 590	gko::solver::LowerTrs< ValueType, IndexType >, 720
operator=, 590	conj_transpose, 722
set_stop_criterion_factory, 591	LowerTrs, 721
${\tt gko::solver::EnablePreconditionable} < {\tt DerivedType} \ >,$	operator=, 722
597	transpose, 722
EnablePreconditionable, 598	gko::solver::Multigrid, 753
operator=, 598	apply_uses_initial_guess, 754
set_preconditioner, 599	get_coarsest_solver, 754
gko::solver::EnablePreconditionedIterativeSolver< Val-	get_cycle, 754
ueType, DerivedType $>$, 599	get_mg_level_list, 755
gko::solver::EnableSolverBase< DerivedType, Matrix-	get_mid_smoother_list, 755
Type $>$, 600	get_post_smoother_list, 755
EnableSolverBase, 601	get_pre_smoother_list, 755
operator=, 601	set_cycle, 756
gko::solver::Fcg< ValueType >, 626	gko::solver::multigrid, 384
apply_uses_initial_guess, 627	cycle, 384
conj_transpose, 627	mid_smooth_type, 384
transpose, 628	gko::solver::UpperTrs< ValueType, IndexType >, 884
gko::solver::Gcr< ValueType >, 638	conj_transpose, 885
apply_uses_initial_guess, 638	operator=, 885, 886
conj_transpose, 639	transpose, 886
get_krylov_dim, 639	UpperTrs, 885
set_krylov_dim, 639	gko::solver::workspace_traits< Solver >, 918
transpose, 639	gko::span, 849
gko::solver::Gmres< ValueType >, 640	is_valid, 851
apply_uses_initial_guess, 641	length, 851
conj_transpose, 641	span, 850, 851
get_krylov_dim, 641	gko::stop, 385
set_krylov_dim, 641	EnableDefaultCriterionFactory, 386
transpose, 642	${\tt gko::stop::AbsoluteResidualNorm} < {\tt ValueType} >, {\tt 392}$
gko::solver::has_with_criteria< SolverType, typename	gko::stop::Combined, 434
>, 642	gko::stop::Criterion, 484
${\tt gko::solver::has_with_criteria} < {\tt SolverType, xstd::void_t} <$	check, 484
decltype(SolverType::build().with_criteria(std::sh	ared_uppdate, 485
const stop::CriterionFactory >()))> >, 643	gko::stop::Criterion::Updater, 883
gko::solver::ldr< ValueType >, 672	check, 883
apply_uses_initial_guess, 673	gko::stop::CriterionArgs, 486
conj_transpose, 673	gko::stop::ImplicitResidualNorm< ValueType >, 686
get_complex_subspace, 674	gko::stop::Iteration, 705
get_deterministic, 674	gko::stop::RelativeResidualNorm< ValueType >, 824

gko::stop::ResidualNorm< ValueType >, 826	GKO FACTORY PARAMETER
gko::stop::ResidualNormBase< ValueType >, 827	Linear Operators, 295
gko::stop::ResidualNormReduction< ValueType >, 828	GKO_FACTORY_PARAMETER_SCALAR
gko::stop::Time, 876	Linear Operators, 295
gko::stopping_status, 863	GKO_FACTORY_PARAMETER_VECTOR
converge, 864	Linear Operators, 296
get_id, 864	GKO_REGISTER_HOST_OPERATION
has_converged, 864	Executors, 282
has_stopped, 864	GKO_REGISTER_OPERATION
is_finalized, 865	Executors, 283
operator!=, 865	group_power
operator==, 866	gko::preconditioner::block_interleaved_storage_scheme<
stop, 865	IndexType >, 420
gko::StreamError, 873	
StreamError, 873	has_connected_parts
gko::syn, 387	gko::experimental::distributed::Partition< LocalIn-
as_array, 388	dexType, GlobalIndexType >, 780
as_list, 387	has_converged
concatenate, 388	gko::log::Convergence < ValueType >, 471
gko::syn::range< Start, End, Step >, 811	gko::stopping_status, 864
gko::syn::type_list< Types >, 882	has_ordered_parts
gko::syn::value_list< T, Values >, 888	gko::experimental::distributed::Partition< LocalIn-
gko::time_point, 877	dexType, GlobalIndexType >, 780
gko::Timer, 877	has_stopped
create_for_executor, 878	gko::stopping_status, 864
create_time_point, 878	highest_precision
difference, 878	gko, 320
difference_async, 879	HIP Executor, 286
gko::Transposable, 879	hip_version
conj_transpose, 880	gko::version_info, 906
transpose, 880	HipblasError
gko::UnsupportedMatrixProperty, 882	gko::HipblasError, 644
UnsupportedMatrixProperty, 882	HipError
gko::UseComposition< ValueType >, 886	gko::HipError, 645
get_composition, 887	HipfftError
get_operator_at, 887	gko::HipfftError, 649
gko::ValueMismatch, 888	HiprandError 650
ValueMismatch, 889	gko::HipparaeError, 650
gko::version, 903	HipsparseError gko::HipsparseError, 651
tag, 904	Hybrid
gko::version_info, 904	gko::matrix::Hybrid< ValueType, IndexType >, 655
core_version, 905	gkomatrixriybna value rype, maex rype /, 000
cuda_version, 905	i_all_gather
dpcpp_version, 906	gko::experimental::mpi::communicator, 445
get, 905	i_all_reduce
hip_version, 906	gko::experimental::mpi::communicator, 446, 447
omp_version, 906	i_all_to_all
reference_version, 906	gko::experimental::mpi::communicator, 447, 448
gko::WritableToMatrixData< ValueType, IndexType >,	i_all_to_all_v
918	gko::experimental::mpi::communicator, 449
write, 919	i_broadcast
gko::xstd, 389	gko::experimental::mpi::communicator, 450
GKO_CREATE_FACTORY_PARAMETERS	i_gather
Linear Operators, 293	gko::experimental::mpi::communicator, 451
GKO_ENABLE_BUILD_METHOD	i_gather_v
Linear Operators, 293	gko::experimental::mpi::communicator, 452
GKO_ENABLE_LIN_OP_FACTORY	i_recv
Linear Operators, 294	gko::experimental::mpi::communicator, 452

i_reduce	is_contiguous
gko::experimental::mpi::communicator, 453	gko::index_set< IndexType >, 695
i_scan	is_finalized
gko::experimental::mpi::communicator, 454	gko::stopping_status, 865
i_scatter	is_finite
gko::experimental::mpi::communicator, 454	gko, 335, 336
i_scatter_v	IS_nan
gko::experimental::mpi::communicator, 455	gko, 336, 337
i_send	is_nonzero
gko::experimental::mpi::communicator, 456	gko, 337
lc	is_owning
gko::preconditioner::lc< LSolverType, IndexType	gko::array< ValueType >, 406
>, 667	is_sorted_by_column_index
Ilu	gko::matrix::Fbcsr< ValueType, IndexType >, 624
gko::preconditioner::llu< LSolverType, USolver-	is_valid
Type, ReverseApply, IndexType >, 679	gko::span, 851
imag	is_zero
gko, 334	gko, 338
imbalance_limit	Isai
gko::matrix::Hybrid< ValueType, IndexType	gko::preconditioner::Isai< IsaiType, ValueType, In-
>::imbalance_limit, 684	dexType >, 703
index_set	isai_type
gko::index_set< IndexType >, 688-690	gko::preconditioner, 379
initial_guess_mode	la a a la i
gko::solver, 382	Jacobi
initialize	gko::preconditioner::Jacobi< ValueType, Index-
Linear Operators, 297–299	Type >, 708
internal	Jacobi Preconditioner, 287
gko::log, 375	KernelNotFound
inv_scale	gko::KernelNotFound, 713
gko::experimental::distributed::Vector< ValueType	gkokemeinoti odila, 713
>, 901	layout_type
gko::matrix::Csr< ValueType, IndexType >, 496	gko, 323
gko::matrix::Dense< ValueType >, 534	lend
inverse_apply	gko, 338, 339
gko::matrix::Diagonal< ValueType >, 556	length
inverse_column_permute	gko::accessor::row_major< ValueType, Dimen-
gko::matrix::Csr< ValueType, IndexType >, 496	sionality >, 830
gko::matrix::Dense< ValueType >, 534, 535	gko::range< Accessor >, 815
gko::Permutable < IndexType >, 783	gko::span, 851
inverse_permute	Linear Operators, 288
gko::matrix::Csr< ValueType, IndexType >, 497	EnableDefaultLinOpFactory, 296
gko::matrix::Dense< ValueType >, 535–537	GKO CREATE FACTORY PARAMETERS, 293
gko::Permutable< IndexType >, 784	GKO ENABLE BUILD METHOD, 293
inverse_row_permute	GKO_ENABLE_LIN_OP_FACTORY, 294
gko::matrix::Csr< ValueType, IndexType >, 497	GKO_FACTORY_PARAMETER, 295
gko::matrix::Dense< ValueType >, 537, 538	GKO_FACTORY_PARAMETER_SCALAR, 295
gko::Permutable < IndexType >, 784	GKO_FACTORY_PARAMETER_VECTOR, 296
Ir	initialize, 297–299
gko::solver::Ir< ValueType >, 699	linop
is_complex	gko::log, 375
gko, 335	load_balance
is_complex_or_scalar	gko::matrix::Csr< ValueType, IndexType >::load_balance,
gko, 335	716, 717
is_complex_or_scalar_s	lock
gko, 321	gko::experimental::mpi::window< ValueType >,
is_complex_s	913
gko, 321	lock_all
gno, 521	ioon_aii

gko::experimental::mpi::window< ValueType >, 914	gko::PolymorphicObject, 800 move to
log_propagation_mode	gko::ConvertibleTo< ResultType >, 473
gko, 323	gko::EnablePolymorphicAssignment< Concrete-
Logging, 300	Type, ResultType >, 596
LowerTrs	gko::preconditioner::Jacobi< ValueType, Index-
gko::solver::LowerTrs< ValueType, IndexType >,	Type >, 711
721	MpiError
	gko::MpiError, 752
make_array_view	g
gko, 339	nan
make_complex	gko, 346, 347
gko::experimental::distributed::Vector< ValueType	never
>, 901	gko, 324
gko::matrix::Dense< ValueType >, 538	node_local_rank
make_const_array_view	gko::experimental::mpi::communicator, 457
gko, 340	nonzeros
make_const_dense_view	gko::matrix_data< ValueType, IndexType >, 747
gko, 340	NotCompiled
make_dense_view	gko::NotCompiled, 759
gko, 342	•
make_temporary_clone	NotImplemented
gko, 342	gko::NotImplemented, 760
-	NotSupported
make_temporary_conversion	gko::NotSupported, 761
gko, 343	ahiaat
gko::experimental::distributed, 367, 368	object
make_temporary_output_clone	gko::log, 375
gko, 344	omp_version
map_global_to_local	gko::version_info, 906
gko::index_set< IndexType >, 695	on
map_local_to_global	gko::enable_parameters_type< ConcreteParame-
	tersType, Factory >, 583
gko::index_set< IndexType >, 696	tors type, i dotory >, soo
map_rank_to_device_id	one
map_rank_to_device_id gko::experimental::mpi, 372	one gko, 347
map_rank_to_device_id gko::experimental::mpi, 372 Matrix	one gko, 347 OpenMP Executor, 302
map_rank_to_device_id gko::experimental::mpi, 372 Matrix gko::experimental::distributed::Matrix< ValueType,	one gko, 347 OpenMP Executor, 302 operation
map_rank_to_device_id gko::experimental::mpi, 372 Matrix	one gko, 347 OpenMP Executor, 302
map_rank_to_device_id gko::experimental::mpi, 372 Matrix gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 731, 732 matrix_data	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool
map_rank_to_device_id gko::experimental::mpi, 372 Matrix gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 731, 732	one gko, 347 OpenMP Executor, 302 operation gko::log, 375
map_rank_to_device_id gko::experimental::mpi, 372 Matrix gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 731, 732 matrix_data	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool
map_rank_to_device_id gko::experimental::mpi, 372 Matrix gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 731, 732 matrix_data gko::matrix_data< ValueType, IndexType >, 740-	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!=
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::clim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator*
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::cdim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::dim< Dimensionality, DimensionType >, 563
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::dim< Dimensionality, DimensionType >, 563 gko::ptr_param< T >, 811
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::dim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator()
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator() gko::accessor::row_major< ValueType, Dimension-
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator() gko::accessor::row_major< ValueType, Dimensionality >, 831
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator() gko::accessor::row_major< ValueType, Dimensionality >, 831 gko::default_converter< S, R >, 514
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::clim< Dimensionality, DimensionType >, 561 gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator() gko::accessor::row_major< ValueType, Dimensionality >, 831 gko::default_converter< S, R >, 514 gko::executor_deleter< T >, 612
map_rank_to_device_id	one gko, 347 OpenMP Executor, 302 operation gko::log, 375 operator bool gko::ptr_param< T >, 811 operator storage_type gko::precision_reduction, 803 operator!= gko, 347, 348 gko::experimental::mpi::communicator, 457 gko::stopping_status, 865 operator<< gko, 349 gko::dim< Dimensionality, DimensionType >, 563 operator* gko::ptr_param< T >, 811 operator() gko::accessor::row_major< ValueType, Dimensionality >, 831 gko::default_converter< S, R >, 514

operator->	gko::Permutable < IndexType >, 786
gko::ptr_param< T >, 811	pi
gko::range< Accessor >, 816	gko, 351
operator=	precision_dispatch
gko::array< ValueType >, 407, 408	gko, 352
gko::Combination < ValueType >, 433	gko::experimental::distributed, 369
gko::Composition < ValueType >, 465	precision_dispatch_real_complex
gko::experimental::distributed::DistributedBase,	gko, 352, 353
567	gko::experimental::distributed, 369, 370
gko::experimental::distributed::Matrix< ValueType,	precision_reduction
LocalIndexType, GlobalIndexType >, 732, 733	gko::precision_reduction, 802
gko::experimental::mpi::contiguous_type, 468	Preconditioners, 303
gko::experimental::mpi::window< ValueType >,	previous_precision
914	gko, 321
gko::index_set< IndexType >, 696, 697	process
gko::matrix::Csr< ValueType, IndexType >, 498	gko::matrix::Csr< ValueType, IndexType >::classical,
gko::matrix::Dense< ValueType >, 539	428
gko::matrix::Ell< ValueType, IndexType >, 579	gko::matrix::Csr< ValueType, IndexType >::cusparse,
gko::matrix::Fbcsr< ValueType, IndexType >, 624	512
gko::matrix::Hybrid< ValueType, IndexType >, 663	gko::matrix::Csr< ValueType, IndexType >::load_balance,
gko::matrix::Sellp< ValueType, IndexType >, 846,	718
847	gko::matrix::Csr< ValueType, IndexType >::merge_path,
gko::matrix::SparsityCsr< ValueType, IndexType	749
>, 859	gko::matrix::Csr< ValueType, IndexType >::sparselib,
gko::preconditioner::lc< LSolverType, IndexType	853
>, 668, 669	gko::matrix::Csr< ValueType, IndexType >::strategy_type,
gko::preconditioner::llu< LSolverType, USolver-	871
Type, ReverseApply, IndexType >, 681	profile_event_category
gko::preconditioner::lsai< lsaiType, ValueType, In-	gko::log, 375
dexType >, 704	profiling_scope_guard
gko::preconditioner::Jacobi< ValueType, Index-	gko::log::profiling_scope_guard, 809
Type >, 711	provided
gko::range< Accessor >, 816, 817	gko::solver, 382
${\tt gko::solver::EnableIterativeBase} < {\tt DerivedType} >,$	put
590	gko::experimental::mpi::window< ValueType >,
gko::solver::EnablePreconditionable< DerivedType	914
>, 598	
gko::solver::EnableSolverBase< DerivedType, Ma-	r_accumulate
trixType >, 601	gko::experimentai::mpi::window< value type >,
gko::solver::lr< ValueType >, 700	915
gko::solver::LowerTrs< ValueType, IndexType >,	r_get
722	gko::experimental::mpi::window< ValueType >,
gko::solver::UpperTrs< ValueType, IndexType >,	915
885, 886	r_get_accumulate
operator==	gko::experimental::mpi::window< ValueType >,
gko, 351	916
gko::dim< Dimensionality, DimensionType >, 563	r_put
gko::experimental::mpi::communicator, 457	gko::experimental::mpi::window< ValueType >, 917
gko::stopping_status, 866	
operator[]	range
gko::dim< Dimensionality, DimensionType >, 562	gko::range < Accessor >, 814
OutOfBoundsError	rank
gko::OutOfBoundsError, 766	gko::experimental::mpi::communicator, 457
OverflowError	rapply gko::matrix::Diagonal < ValueType >, 557
gko::OverflowError, 767	read
permute	gko, 353
gko::matrix::Csr< ValueType, IndexType >, 498	gko::matrix::Coo< ValueType, IndexType >, 480,
gko::matrix::Dense< ValueType >, 539, 540	481
J	

gko::matrix::Csr< ValueType, IndexType >, 499 gko::matrix::Ell< ValueType, IndexType >, 580 gko::matrix::Fbcsr< ValueType, IndexType >, 625 gko::matrix::Hybrid< ValueType, IndexType >, 663, 664 gko::matrix::Sellp< ValueType, IndexType >, 847,	gko::matrix::Dense< ValueType >, 541, 542 row_permute gko::matrix::Csr< ValueType, IndexType >, 500 gko::matrix::Dense< ValueType >, 543, 544 gko::Permutable< IndexType >, 786 run
848 gko::matrix::SparsityCsr< ValueType, IndexType	gko::Executor, 609, 610 gko::ReferenceExecutor, 823
>, 859, 860 gko::ReadableFromMatrixData< ValueType, Index- Type >, 819, 820	safe_divide gko, 359
••	scale
read_binary gko, 354	gko::experimental::distributed::Vector< ValueType
read_binary_raw	>, 902
gko, 354	gko::matrix::Csr< ValueType, IndexType >, 500
read_distributed	gko::matrix::Dense< ValueType >, 544
gko::experimental::distributed::Matrix< ValueType,	scan
LocalIndexType, GlobalIndexType >, 733–735	gko::experimental::mpi::communicator, 459 scatter
gko::experimental::distributed::Vector< ValueType >, 901, 902	gko::experimental::mpi::communicator, 460 scatter_v
read_generic	gko::experimental::mpi::communicator, 460
gko, 355	scoped_device_id_guard
read_generic_raw	gko::scoped_device_id_guard, 837-839
gko, 356	Sellp
read_raw gko, 356	gko::matrix::Sellp< ValueType, IndexType >, 841
real	send
gko, 357	gko::experimental::mpi::communicator, 462
recv	set_complex_subpsace
gko::experimental::mpi::communicator, 458	gko::solver::ldr< ValueType >, 675
reduce	set_cycle
gko::experimental::mpi::communicator, 458	gko::solver::Multigrid, 756
reduce_add	set_deterministic
gko, 358	gko::solver::ldr< ValueType >, 675
Reference Executor, 304	set_executor
reference version	gko::array< ValueType >, 409
gko::version_info, 906	set_kappa
remove_complex	gko::solver::ldr< ValueType >, 675
gko, 322	set_krylov_dim gko::solver::CbGmres< ValueType >, 422
remove_logger	gko::solver::Gcr< ValueType >, 639
gko::Executor, 609	gko::solver::Gmres< ValueType >, 641
gko::log::Loggable, 720	set_log_propagation_mode
remove_zeros	gko::Executor, 610
gko::device_matrix_data< ValueType, IndexType	set_object_name
>, 552	gko::log::ProfilerHook, 808
request	set_permute_mask
gko::experimental::mpi::request, 825	gko::matrix::Permutation < IndexType >, 789
resize_and_reset	set_preconditioner
gko::array< ValueType >, 409	gko::Preconditionable, 804
gko::device_matrix_data< ValueType, IndexType	gko::solver::EnablePreconditionable< DerivedType
>, 553	>, 599
rhs	set_solver
gko::solver, 382	gko::solver::lr< ValueType >, 701
round_down	set_stop_criterion_factory
gko, 359	${\tt gko::solver::EnableIterativeBase} < {\tt DerivedType} >,$
round_up	591
gko, 359	gko::solver::IterativeBase, 706
row_gather	set_strategy

gko::matrix::Csr< ValueType, IndexType >, 500	transpose
set_subspace_dim	gko, 361
gko::solver::ldr< ValueType >, 676	gko::Combination < ValueType >, 433
set_value	gko::Composition < ValueType >, 466
gko::matrix_assembly_data< ValueType, Index- Type >, 738	gko::experimental::solver::Direct< ValueType, IndexType >, 566
share	gko::matrix::Csr< ValueType, IndexType >, 501
gko, 360	gko::matrix::Dense< ValueType >, 545
should_propagate_log	gko::matrix::Diagonal < ValueType >, 557
	· · · · · · · · · · · · · · · · · · ·
gko::Executor, 611	gko::matrix::Fbcsr< ValueType, IndexType >, 626
Size	gko::matrix::Fft, 629
gko::experimental::mpi::communicator, 462	gko::matrix::Fft2, 632
solver	gko::matrix::Fft3, 635
gko::log, 375	gko::matrix::ldentity< ValueType >, 671
Solvers, 305	gko::matrix::SparsityCsr< ValueType, IndexType
span	>, 860
gko::span, 850, 851	gko::preconditioner::lc< LSolverType, IndexType
SparsityCsr	>, 669
gko::matrix::SparsityCsr< ValueType, IndexType	gko::preconditioner::llu< LSolverType, USolver-
>, 855	Type, ReverseApply, IndexType >, 681
SpMV employing different Matrix formats, 301	gko::preconditioner::lsai< lsaiType, ValueType, In-
squared_norm	dexType >, 704
gko, 361	gko::preconditioner::Jacobi< ValueType, Index-
status	Type >, 712
gko::experimental::mpi::status, 862	gko::solver::Bicg< ValueType >, 414
stop	gko::solver::Bicgstab< ValueType >, 416
gko::stopping_status, 865	gko::solver::Cg< ValueType >, 424
Stopping criteria, 306	gko::solver::Cgs< ValueType >, 426
combine, 307	gko::solver::Fcg< ValueType >, 628
mode, 307	gko::solver::Gcr< ValueType >, 629
	gko::solver::Gmres< ValueType >, 642
strategy_type	
gko::matrix::Csr< ValueType, IndexType >::strategy_	
869	gko::solver::lr< ValueType >, 701
StreamError	gko::solver::LowerTrs< ValueType, IndexType >,
gko::StreamError, 873	722
sub_scaled	gko::solver::UpperTrs< ValueType, IndexType >,
gko::experimental::distributed::Vector< ValueType	886
>, 903	gko::Transposable, 880
gko::matrix::Dense< ValueType >, 545	trisolve_algorithm
sum_duplicates	gko::solver, 382
gko::device_matrix_data< ValueType, IndexType	type
>, 553	gko::cpx_real_type< T >, 483
gko::matrix_data< ValueType, IndexType >, 747	
synchronize	unit_root
gko::experimental::mpi::communicator, 463	gko, 361
	unlock
TableSummaryWriter	gko::experimental::mpi::window< ValueType >,
gko::log::ProfilerHook::TableSummaryWriter, 875	917
tag	unpack
gko::version, 904	gko::experimental::factorization::Factorization<
to_adjacency_matrix	ValueType, IndexType >, 616
gko::matrix::SparsityCsr< ValueType, IndexType	UnsupportedMatrixProperty
>, 860	gko::UnsupportedMatrixProperty, 882
to_complex	update
gko, 322	gko::stop::Criterion, 485
to_global_indices	UpperTrs
gko::index_set< IndexType >, 697	gko::solver::UpperTrs< ValueType, IndexType >,
to_real	885
gko, 322	user
The state of the s	

```
gko::log, 375
user_range
    gko::log::ProfilerHook, 808
val at
    gko::matrix::Ell< ValueType, IndexType >, 581
     gko::matrix::Sellp< ValueType, IndexType >, 848
ValueMismatch
     gko::ValueMismatch, 889
view
     gko::array< ValueType >, 410
wait
     gko::experimental::mpi::request, 826
wait all
     gko::experimental::mpi, 373
window
    gko::experimental::mpi::window< ValueType >,
         908, 909
with_matrix_type
     gko, 362
write
     ako, 363
    gko::log::ProfilerHook::SummaryWriter, 874
    gko::log::ProfilerHook::TableSummaryWriter, 876
    gko::matrix::Coo < ValueType, IndexType >, 481
    gko::matrix::Csr< ValueType, IndexType >, 501
    gko::matrix::Ell< ValueType, IndexType >, 582
     gko::matrix::Fbcsr< ValueType, IndexType >, 626
     gko::matrix::Fft, 629, 630
     gko::matrix::Fft2, 632, 633
    gko::matrix::Fft3, 635, 636
    gko::matrix::Hybrid< ValueType, IndexType >, 664
    gko::matrix::Sellp< ValueType, IndexType >, 849
    gko::matrix::SparsityCsr< ValueType, IndexType
         >, 861
    gko::preconditioner::Jacobi < ValueType, Index-
         Type >, 712
    gko::WritableToMatrixData < ValueType, IndexType
         >, 919
write_binary
     gko, 363
write_binary_raw
     gko, 364
write nested
     gko::log::ProfilerHook::NestedSummaryWriter, 758
     gko::log::ProfilerHook::TableSummaryWriter, 876
write_raw
     gko, 364
zero
     gko, 365
    gko::solver, 382
```