

Ginkgo

Generated from pipelines/1680922742 branch based on develop. Ginkgo version 1.10.0

Generated by Doxygen 1.8.16

1 Main Page	1
1.0.0.1 Modules	1
2 Installation Instructions	3
2.0.1 Building	3
2.0.2 Building Ginkgo in Windows	5
2.0.3 Building Ginkgo with HIP support	5
2.0.3.1 Changing the paths to search for HIP and other packages	6
2.0.3.2 HIP platform detection of AMD and NVIDIA	6
2.0.4 Third party libraries and packages	6
2.0.5 Installing Ginkgo	7
3 Testing Instructions	9
3.0.1 Running the unit tests	9
3.0.1.1 Using make test	9
3.0.1.2 Using make quick_test	9
3.0.1.3 Using CTest	9
4 Running the benchmarks	11
4.0.1 1: Ginkgo setup and best practice guidelines	11
4.0.2 2: Using ssget to fetch the matrices	12
4.0.3 3: Running benchmarks manually	13
4.0.4 4: Benchmarking overview	14
4.0.5 5: Publishing the results on Github and analyze the results with the GPE (optional)	14
4.0.6 6: Detailed performance analysis and debugging	15
4.0.7 7: Available benchmark options	15
5 Contributing guidelines	19
5.1 Table of Contents	19
5.2 Most important stuff (A TL;DR)	20
5.3 Project structure	20
5.3.1 Extended header files	20
5.3.2 Using library classes	21
5.4 Git related	21
5.4.1 Our git workflow	21
5.4.2 Writing good commit messages	21
5.4.2.1 Attributing credit	21
5.4.3 Creating, Reviewing and Merging Pull Requests	22
5.5 Code style	22
5.5.1 Automatic code formatting	22
5.5.2 Naming scheme	22
5.5.2.1 Filenames	22
5.5.2.2 Macros	23
5.5.2.3 Variables	23

5.5.2.4 Constants	23
5.5.2.5 Functions	23
5.5.2.6 Structures and classes	23
5.5.2.7 Members	23
5.5.2.8 Namespaces	24
5.5.2.9 Template parameters	24
5.5.3 Whitespace	24
5.5.4 Include statement grouping	25
5.5.4.1 Main header	25
5.5.4.2 Some general comments.	26
5.5.5 Other Code Formatting not handled by ClangFormat	26
5.5.5.1 Control flow constructs	26
5.5.5.2 Variable declarations	26
5.5.6 CMake coding style	26
5.5.6.1 Whitespaces	26
5.5.6.2 Use of macros vs functions	26
5.5.6.3 Naming style	26
5.6 Helper scripts	27
5.6.1 Create a new algorithm	27
5.6.2 Converting CUDA code to HIP code	27
5.7 Writing Tests	27
5.7.1 Testing know-how	27
5.7.2 Some general rules.	27
5.7.3 Writing tests for kernels	28
5.8 Documentation style	28
5.8.1 Developer targeted notes	28
5.8.2 Whitespaces	28
5.8.2.1 After named tags such as <code><tt>@param foo</tt></code>	28
5.8.3 Documenting examples	28
5.9 Other programming comments	29
5.9.1 C++ standard stream objects	29
5.9.2 Warnings	29
5.9.3 Avoiding circular dependencies	29
6 Citing Ginkgo	31
6.0.1 The Ginkgo Software	31
6.0.2 On Portability	31
6.0.3 On Software Sustainability	32
6.0.4 On SpMV or solvers performance	32
7 Example programs	33
8 The adaptiveprecision-blockjacobi program	37

9 The batched-solver program	41
10 The cb-gmres program	51
11 The custom-logger program	57
12 The custom-matrix-format program	65
13 The custom-stopping-criterion program	73
14 The distributed-multigrid-preconditioned-solver program	79
15 The distributed-solver program	87
16 The external-lib-interfacing program	93
17 The file-config-solver program	115
18 The ginkgo-overhead program	119
19 The ginkgo-ranges program	121
20 The heat-equation program	123
21 The ilu-preconditioned-solver program	129
22 The inverse-iteration program	133
23 The ir-ilu-preconditioned-solver program	139
24 The iterative-refinement program	145
25 The kokkos-assembly program	149
26 The minimal-cuda-solver program	155
27 The mixed-multigrid-preconditioned-solver program	157
28 The mixed-multigrid-solver program	163
29 The mixed-precision-ir program	169
30 The mixed-spmv program	173
31 The multigrid-preconditioned-solver program	179
32 The multigrid-preconditioned-solver-customized program	183
33 The nine-pt-stencil-solver program	189
34 The papi-logging program	199

35 The par-ilu-convergence program	205
36 The performance-debugging program	211
37 The poisson-solver program	223
38 The preconditioned-solver program	229
39 The preconditioner-export program	233
40 The reordered-preconditioned-solver program	239
41 The schroedinger-splitting program	243
42 The simple-solver program	249
43 The simple-solver-logging program	253
44 The three-pt-stencil-solver program	279
45 Module Documentation	287
45.1 CUDA Executor	287
45.1.1 Detailed Description	287
45.2 DPC++ Executor	288
45.2.1 Detailed Description	288
45.3 Executors	289
45.3.1 Detailed Description	289
45.3.2 Executors in Ginkgo.	290
45.3.3 Macro Definition Documentation	290
45.3.3.1 GKO_REGISTER_HOST_OPERATION	290
45.3.3.2 Example	291
45.3.3.3 GKO_REGISTER_OPERATION	291
45.3.3.4 Example	291
45.4 Factorizations	293
45.4.1 Detailed Description	293
45.5 HIP Executor	294
45.5.1 Detailed Description	294
45.6 Jacobi Preconditioner	295
45.6.1 Detailed Description	295
45.7 Linear Operators	296
45.7.1 Detailed Description	300
45.7.2 Advantages of this approach and usage	300
45.7.3 Linear operator as a concept	300
45.7.4 Macro Definition Documentation	301
45.7.4.1 GKO_CREATE_FACTORY_PARAMETERS	301
45.7.4.2 GKO_ENABLE_BUILD_METHOD	302

45.7.4.3 GKO_ENABLE_LIN_OP_FACTORY	302
45.7.4.4 GKO_FACTORY_PARAMETER	303
45.7.4.5 GKO_FACTORY_PARAMETER_SCALAR	303
45.7.4.6 GKO_FACTORY_PARAMETER_VECTOR	304
45.7.5 Typedef Documentation	304
45.7.5.1 EnableDefaultLinOpFactory	304
45.7.6 Function Documentation	305
45.7.6.1 initialize() [1/4]	305
45.7.6.2 initialize() [2/4]	306
45.7.6.3 initialize() [3/4]	306
45.7.6.4 initialize() [4/4]	307
45.8 Logging	308
45.8.1 Detailed Description	308
45.9 SpMV employing different Matrix formats	309
45.9.1 Detailed Description	310
45.10 OpenMP Executor	311
45.10.1 Detailed Description	311
45.11 Preconditioners	312
45.11.1 Detailed Description	312
45.12 Reference Executor	313
45.12.1 Detailed Description	313
45.13 Solvers	314
45.13.1 Detailed Description	314
45.14 Stopping criteria	315
45.14.1 Detailed Description	316
45.14.2 Enumeration Type Documentation	316
45.14.2.1 mode	316
45.14.3 Function Documentation	316
45.14.3.1 combine()	316
45.15 BatchLinOp	318
45.15.1 Detailed Description	319
45.15.2 Batched Linear operator as a concept	319
45.15.3 Macro Definition Documentation	319
45.15.3.1 GKO_ENABLE_BATCH_LIN_OP_FACTORY	319
45.15.4 Typedef Documentation	320
45.15.4.1 EnableDefaultBatchLinOpFactory	320
46 Namespace Documentation	321
46.1 gko Namespace Reference	321
46.1.1 Detailed Description	333
46.1.2 Typedef Documentation	333
46.1.2.1 highest_precision	333

46.1.2.2 is_complex_or_scalar_s	333
46.1.2.3 is_complex_s	334
46.1.2.4 previous_precision_base	334
46.1.2.5 remove_complex	334
46.1.2.6 to_complex	335
46.1.2.7 to_real	335
46.1.3 Enumeration Type Documentation	335
46.1.3.1 allocation_mode	335
46.1.3.2 layout_type	336
46.1.3.3 log_propagation_mode	336
46.1.4 Function Documentation	336
46.1.4.1 abs()	336
46.1.4.2 as() [1/7]	337
46.1.4.3 as() [2/7]	337
46.1.4.4 as() [3/7]	338
46.1.4.5 as() [4/7]	338
46.1.4.6 as() [5/7]	339
46.1.4.7 as() [6/7]	339
46.1.4.8 as() [7/7]	340
46.1.4.9 ceildiv()	340
46.1.4.10 clone() [1/2]	341
46.1.4.11 clone() [2/2]	341
46.1.4.12 conj()	342
46.1.4.13 copy_and_convert_to() [1/4]	342
46.1.4.14 copy_and_convert_to() [2/4]	343
46.1.4.15 copy_and_convert_to() [3/4]	344
46.1.4.16 copy_and_convert_to() [4/4]	344
46.1.4.17 get_significant_bit()	345
46.1.4.18 get_superior_power()	346
46.1.4.19 give()	346
46.1.4.20 imag()	347
46.1.4.21 is_complex()	347
46.1.4.22 is_complex_or_scalar()	347
46.1.4.23 is_finite() [1/2]	348
46.1.4.24 is_finite() [2/2]	348
46.1.4.25 is_nan() [1/2]	349
46.1.4.26 is_nan() [2/2]	349
46.1.4.27 is_nonzero()	350
46.1.4.28 is_zero()	350
46.1.4.29 lend() [1/2]	351
46.1.4.30 lend() [2/2]	351
46.1.4.31 make_array_view()	352

46.1.4.32 <code>make_const_array_view()</code>	352
46.1.4.33 <code>make_const_dense_view()</code>	353
46.1.4.34 <code>make_dense_view()</code>	353
46.1.4.35 <code>make_temporary_clone()</code>	354
46.1.4.36 <code>make_temporary_conversion()</code>	354
46.1.4.37 <code>make_temporary_output_clone()</code>	355
46.1.4.38 <code>max()</code>	356
46.1.4.39 <code>min()</code>	356
46.1.4.40 <code>mixed_precision_dispatch()</code>	357
46.1.4.41 <code>mixed_precision_dispatch_real_complex()</code>	357
46.1.4.42 <code>nan()</code> [1/2]	358
46.1.4.43 <code>nan()</code> [2/2]	358
46.1.4.44 <code>one()</code> [1/2]	359
46.1.4.45 <code>one()</code> [2/2]	359
46.1.4.46 <code>operator"!="()</code> [1/2]	359
46.1.4.47 <code>operator"!="()</code> [2/2]	360
46.1.4.48 <code>operator<<()</code> [1/2]	360
46.1.4.49 <code>operator<<()</code> [2/2]	361
46.1.4.50 <code>operator==()</code> [1/2]	361
46.1.4.51 <code>operator==()</code> [2/2]	361
46.1.4.52 <code>pi()</code>	362
46.1.4.53 <code>precision_dispatch()</code>	362
46.1.4.54 <code>precision_dispatch_real_complex()</code> [1/3]	363
46.1.4.55 <code>precision_dispatch_real_complex()</code> [2/3]	363
46.1.4.56 <code>precision_dispatch_real_complex()</code> [3/3]	363
46.1.4.57 <code>read()</code>	364
46.1.4.58 <code>read_binary()</code>	364
46.1.4.59 <code>read_binary_raw()</code>	365
46.1.4.60 <code>read_generic()</code>	366
46.1.4.61 <code>read_generic_raw()</code>	366
46.1.4.62 <code>read_raw()</code>	367
46.1.4.63 <code>real()</code>	367
46.1.4.64 <code>reduce_add()</code> [1/2]	368
46.1.4.65 <code>reduce_add()</code> [2/2]	368
46.1.4.66 <code>round_down()</code>	369
46.1.4.67 <code>round_up()</code>	369
46.1.4.68 <code>safe_divide()</code>	370
46.1.4.69 <code>share()</code>	370
46.1.4.70 <code>squared_norm()</code>	371
46.1.4.71 <code>transpose()</code> [1/2]	371
46.1.4.72 <code>transpose()</code> [2/2]	372
46.1.4.73 <code>unit_root()</code>	372

46.1.4.74 <code>with_matrix_type()</code>	373
46.1.4.75 <code>write()</code>	374
46.1.4.76 <code>write_binary()</code>	374
46.1.4.77 <code>write_binary_raw()</code>	375
46.1.4.78 <code>write_raw()</code>	375
46.1.4.79 <code>zero()</code> [1/2]	376
46.1.4.80 <code>zero()</code> [2/2]	376
46.2 <code>gko::accessor</code> Namespace Reference	377
46.2.1 Detailed Description	377
46.3 <code>gko::batch::log</code> Namespace Reference	377
46.3.1 Detailed Description	377
46.4 <code>gko::experimental::distributed</code> Namespace Reference	377
46.4.1 Detailed Description	379
46.4.2 Typedef Documentation	379
46.4.2.1 <code>comm_index_type</code>	379
46.4.3 Enumeration Type Documentation	379
46.4.3.1 <code>assembly_mode</code>	379
46.4.3.2 <code>index_space</code>	380
46.4.4 Function Documentation	380
46.4.4.1 <code>assemble_rows_from_neighbors()</code>	380
46.4.4.2 <code>build_partition_from_local_range()</code>	381
46.4.4.3 <code>build_partition_from_local_size()</code>	381
46.4.4.4 <code>make_temporary_conversion()</code> [1/2]	382
46.4.4.5 <code>make_temporary_conversion()</code> [2/2]	382
46.4.4.6 <code>precision_dispatch()</code>	383
46.4.4.7 <code>precision_dispatch_real_complex()</code> [1/3]	383
46.4.4.8 <code>precision_dispatch_real_complex()</code> [2/3]	384
46.4.4.9 <code>precision_dispatch_real_complex()</code> [3/3]	384
46.5 <code>gko::experimental::distributed::preconditioner</code> Namespace Reference	385
46.5.1 Detailed Description	385
46.6 <code>gko::experimental::mpi</code> Namespace Reference	385
46.6.1 Detailed Description	386
46.6.2 Function Documentation	386
46.6.2.1 <code>get_walltime()</code>	386
46.6.2.2 <code>map_rank_to_device_id()</code>	386
46.6.2.3 <code>wait_all()</code>	387
46.7 <code>gko::experimental::reorder</code> Namespace Reference	387
46.7.1 Detailed Description	387
46.7.2 Enumeration Type Documentation	388
46.7.2.1 <code>mc64_strategy</code>	388
46.8 <code>gko::factorization</code> Namespace Reference	388
46.8.1 Detailed Description	388

46.8.2 Enumeration Type Documentation	388
46.8.2.1 incomplete_algorithm	389
46.9 gko::log Namespace Reference	389
46.9.1 Detailed Description	390
46.9.2 Enumeration Type Documentation	390
46.9.2.1 profile_event_category	390
46.10 gko::matrix Namespace Reference	391
46.10.1 Detailed Description	392
46.10.2 Enumeration Type Documentation	392
46.10.2.1 permute_mode	392
46.11 gko::multigrid Namespace Reference	393
46.11.1 Detailed Description	393
46.12 gko::name_demangling Namespace Reference	393
46.12.1 Detailed Description	394
46.12.2 Function Documentation	394
46.12.2.1 get_dynamic_type()	394
46.12.2.2 get_static_type()	394
46.13 gko::preconditioner Namespace Reference	395
46.13.1 Detailed Description	395
46.13.2 Enumeration Type Documentation	395
46.13.2.1 isai_type	396
46.14 gko::reorder Namespace Reference	396
46.14.1 Detailed Description	396
46.14.2 Typedef Documentation	396
46.14.2.1 EnableDefaultReorderingBaseFactory	396
46.15 gko::solver Namespace Reference	397
46.15.1 Detailed Description	398
46.15.2 Enumeration Type Documentation	399
46.15.2.1 initial_guess_mode	399
46.15.2.2 trisolve_algorithm	399
46.15.3 Function Documentation	399
46.15.3.1 build_smoother() [1/2]	399
46.15.3.2 build_smoother() [2/2]	400
46.16 gko::solver::multigrid Namespace Reference	400
46.16.1 Detailed Description	401
46.16.2 Enumeration Type Documentation	401
46.16.2.1 cycle	401
46.16.2.2 mid_smooth_type	401
46.17 gko::stop Namespace Reference	402
46.17.1 Detailed Description	403
46.17.2 Typedef Documentation	403
46.17.2.1 EnableDefaultCriterionFactory	403

46.18 gko::syn Namespace Reference	403
46.18.1 Detailed Description	404
46.18.2 Typedef Documentation	404
46.18.2.1 as_list	404
46.18.2.2 concatenate	404
46.18.3 Function Documentation	405
46.18.3.1 as_array()	405
46.19 gko::xstd Namespace Reference	405
46.19.1 Detailed Description	406
47 Class Documentation	407
47.1 gko::AbsoluteComputable Class Reference	407
47.1.1 Detailed Description	407
47.1.2 Member Function Documentation	407
47.1.2.1 compute_absolute_linop()	407
47.2 gko::stop::AbsoluteResidualNorm< ValueType > Class Template Reference	408
47.2.1 Detailed Description	408
47.3 gko::AbstractFactory< AbstractProductType, ComponentsType > Class Template Reference	408
47.3.1 Detailed Description	408
47.3.2 Member Function Documentation	409
47.3.2.1 generate()	409
47.4 gko::AllocationError Class Reference	409
47.4.1 Detailed Description	410
47.4.2 Constructor & Destructor Documentation	410
47.4.2.1 AllocationError()	410
47.5 gko::Allocator Class Reference	410
47.5.1 Detailed Description	410
47.6 gko::experimental::reorder::Amd< IndexType > Class Template Reference	411
47.6.1 Detailed Description	411
47.6.2 Member Function Documentation	411
47.6.2.1 generate()	411
47.6.2.2 get_parameters()	412
47.7 gko::amd_device Class Reference	412
47.7.1 Detailed Description	412
47.8 gko::solver::ApplyWithInitialGuess Class Reference	412
47.8.1 Detailed Description	412
47.9 gko::are_all_integral< Args > Struct Template Reference	413
47.9.1 Detailed Description	413
47.10 gko::array< ValueType > Class Template Reference	413
47.10.1 Detailed Description	415
47.10.2 Constructor & Destructor Documentation	415
47.10.2.1 array() [1/11]	415

47.10.2.2 array() [2/11]	416
47.10.2.3 array() [3/11]	416
47.10.2.4 array() [4/11]	416
47.10.2.5 array() [5/11]	417
47.10.2.6 array() [6/11]	417
47.10.2.7 array() [7/11]	418
47.10.2.8 array() [8/11]	418
47.10.2.9 array() [9/11]	419
47.10.2.10 array() [10/11]	419
47.10.2.11 array() [11/11]	419
47.10.3 Member Function Documentation	420
47.10.3.1 as_const_view()	420
47.10.3.2 as_view()	420
47.10.3.3 clear()	420
47.10.3.4 const_view()	420
47.10.3.5 fill()	421
47.10.3.6 get_const_data()	421
47.10.3.7 get_data()	422
47.10.3.8 get_executor()	423
47.10.3.9 get_num_elems()	423
47.10.3.10 get_size()	424
47.10.3.11 is_owning()	424
47.10.3.12 operator=() [1/4]	425
47.10.3.13 operator=() [2/4]	425
47.10.3.14 operator=() [3/4]	426
47.10.3.15 operator=() [4/4]	426
47.10.3.16 resize_and_reset()	427
47.10.3.17 set_executor()	427
47.10.3.18 view()	428
47.11 gko::device_matrix_data< ValueType, IndexType >::arrays Struct Reference	428
47.11.1 Detailed Description	428
47.12 gko::matrix::Hybrid< ValueType, IndexType >::automatic Class Reference	429
47.12.1 Detailed Description	429
47.12.2 Member Function Documentation	429
47.12.2.1 compute_ell_num_stored_elements_per_row()	429
47.13 gko::BadDimension Class Reference	430
47.13.1 Detailed Description	430
47.13.2 Constructor & Destructor Documentation	430
47.13.2.1 BadDimension()	430
47.14 gko::batch_dim< Dimensionality, DimensionType > Struct Template Reference	431
47.14.1 Detailed Description	431
47.14.2 Constructor & Destructor Documentation	431

47.14.2.1 batch_dim()	432
47.14.3 Member Function Documentation	432
47.14.3.1 get_common_size()	432
47.14.3.2 get_num_batch_items()	432
47.14.4 Friends And Related Function Documentation	433
47.14.4.1 operator"!="	433
47.14.4.2 operator=="	433
47.15 gko::batch::log::BatchConvergence< ValueType > Class Template Reference	434
47.15.1 Detailed Description	434
47.15.2 Member Function Documentation	434
47.15.2.1 create()	434
47.15.2.2 get_num_iterations()	435
47.15.2.3 get_residual_norm()	435
47.16 gko::batch::BatchLinOpFactory Class Reference	435
47.16.1 Detailed Description	436
47.16.1.1 Example: using BatchCG in Ginkgo	436
47.17 gko::batch::solver::BatchSolver Class Reference	436
47.17.1 Detailed Description	437
47.17.2 Member Function Documentation	437
47.17.2.1 get_max_iterations()	437
47.17.2.2 get_preconditioner()	437
47.17.2.3 get_system_matrix()	437
47.17.2.4 get_tolerance()	438
47.17.2.5 get_tolerance_type()	438
47.17.2.6 reset_max_iterations()	438
47.17.2.7 reset_tolerance()	438
47.17.2.8 reset_tolerance_type()	439
47.18 gko::solver::Bicg< ValueType > Class Template Reference	439
47.18.1 Detailed Description	439
47.18.2 Member Function Documentation	440
47.18.2.1 apply_uses_initial_guess()	440
47.18.2.2 conj_transpose()	440
47.18.2.3 parse()	440
47.18.2.4 transpose()	441
47.19 gko::batch::solver::Bicgstab< ValueType > Class Template Reference	441
47.19.1 Detailed Description	441
47.20 gko::solver::Bicgstab< ValueType > Class Template Reference	442
47.20.1 Detailed Description	442
47.20.2 Member Function Documentation	443
47.20.2.1 apply_uses_initial_guess()	443
47.20.2.2 conj_transpose()	443
47.20.2.3 parse()	443

47.20.2.4 transpose()	444
47.21 gko::preconditioner::block_interleaved_storage_scheme< IndexType > Struct Template Reference	444
47.21.1 Detailed Description	445
47.21.2 Member Function Documentation	445
47.21.2.1 compute_storage_space()	445
47.21.2.2 get_block_offset()	446
47.21.2.3 get_global_block_offset()	446
47.21.2.4 get_group_offset()	447
47.21.2.5 get_group_size()	447
47.21.2.6 get_stride()	447
47.21.3 Member Data Documentation	448
47.21.3.1 group_power	448
47.22 gko::BlockOperator Class Reference	448
47.22.1 Detailed Description	449
47.22.2 Constructor & Destructor Documentation	449
47.22.2.1 BlockOperator() [1/2]	449
47.22.2.2 BlockOperator() [2/2]	449
47.22.3 Member Function Documentation	449
47.22.3.1 block_at()	449
47.22.3.2 create() [1/2]	450
47.22.3.3 create() [2/2]	450
47.22.3.4 get_block_size()	451
47.22.3.5 operator=() [1/2]	451
47.22.3.6 operator=() [2/2]	451
47.23 gko::BlockSizeError< IndexType > Class Template Reference	451
47.23.1 Detailed Description	451
47.23.2 Constructor & Destructor Documentation	452
47.23.2.1 BlockSizeError()	452
47.24 gko::solver::CbGmres< ValueType > Class Template Reference	452
47.24.1 Detailed Description	453
47.24.2 Member Function Documentation	453
47.24.2.1 get_krylov_dim()	453
47.24.2.2 get_storage_precision()	454
47.24.2.3 parse()	454
47.24.2.4 set_krylov_dim()	454
47.25 gko::batch::solver::Cg< ValueType > Class Template Reference	455
47.25.1 Detailed Description	455
47.26 gko::solver::Cg< ValueType > Class Template Reference	456
47.26.1 Detailed Description	456
47.26.2 Member Function Documentation	456
47.26.2.1 apply_uses_initial_guess()	456
47.26.2.2 conj_transpose()	457

47.26.2.3	parse()	457
47.26.2.4	transpose()	458
47.27	gko::solver::Cgs< ValueType > Class Template Reference	458
47.27.1	Detailed Description	458
47.27.2	Member Function Documentation	459
47.27.2.1	apply_uses_initial_guess()	459
47.27.2.2	conj_transpose()	459
47.27.2.3	parse()	459
47.27.2.4	transpose()	460
47.28	gko::experimental::factorization::Cholesky< ValueType, IndexType > Class Template Reference	460
47.28.1	Detailed Description	461
47.28.2	Member Function Documentation	461
47.28.2.1	generate()	461
47.28.2.2	get_parameters() [1/2]	462
47.28.2.3	get_parameters() [2/2]	462
47.28.2.4	parse()	462
47.29	gko::matrix::Csr< ValueType, IndexType >::classical Class Reference	463
47.29.1	Detailed Description	463
47.29.2	Member Function Documentation	463
47.29.2.1	clac_size()	463
47.29.2.2	copy()	464
47.29.2.3	process()	464
47.30	gko::matrix::Hybrid< ValueType, IndexType >::column_limit Class Reference	465
47.30.1	Detailed Description	465
47.30.2	Constructor & Destructor Documentation	465
47.30.2.1	column_limit()	465
47.30.3	Member Function Documentation	465
47.30.3.1	compute_ell_num_stored_elements_per_row()	466
47.30.3.2	get_num_columns()	466
47.31	gko::Combination< ValueType > Class Template Reference	466
47.31.1	Detailed Description	467
47.31.2	Constructor & Destructor Documentation	467
47.31.2.1	Combination() [1/2]	467
47.31.2.2	Combination() [2/2]	468
47.31.3	Member Function Documentation	468
47.31.3.1	conj_transpose()	468
47.31.3.2	get_coefficients()	468
47.31.3.3	get_operators()	469
47.31.3.4	operator=() [1/2]	469
47.31.3.5	operator=() [2/2]	469
47.31.3.6	transpose()	469
47.32	gko::stop::Combined Class Reference	470

47.32.1 Detailed Description	470
47.33 gko::experimental::mpi::communicator Class Reference	470
47.33.1 Detailed Description	473
47.33.2 Constructor & Destructor Documentation	473
47.33.2.1 communicator() [1/3]	473
47.33.2.2 communicator() [2/3]	474
47.33.2.3 communicator() [3/3]	474
47.33.3 Member Function Documentation	474
47.33.3.1 all_gather()	474
47.33.3.2 all_reduce() [1/2]	476
47.33.3.3 all_reduce() [2/2]	476
47.33.3.4 all_to_all() [1/2]	477
47.33.3.5 all_to_all() [2/2]	478
47.33.3.6 all_to_all_v() [1/2]	478
47.33.3.7 all_to_all_v() [2/2]	479
47.33.3.8 broadcast()	480
47.33.3.9 gather()	480
47.33.3.10 gather_v()	481
47.33.3.11 get()	481
47.33.3.12 i_all_gather()	482
47.33.3.13 i_all_reduce() [1/2]	482
47.33.3.14 i_all_reduce() [2/2]	483
47.33.3.15 i_all_to_all() [1/2]	484
47.33.3.16 i_all_to_all() [2/2]	484
47.33.3.17 i_all_to_all_v() [1/2]	485
47.33.3.18 i_all_to_all_v() [2/2]	486
47.33.3.19 i_broadcast()	487
47.33.3.20 i_gather()	487
47.33.3.21 i_gather_v()	488
47.33.3.22 i_recv()	489
47.33.3.23 i_reduce()	490
47.33.3.24 i_scan()	490
47.33.3.25 i_scatter()	491
47.33.3.26 i_scatter_v()	492
47.33.3.27 i_send()	492
47.33.3.28 node_local_rank()	493
47.33.3.29 operator"!="()	493
47.33.3.30 operator==(())	494
47.33.3.31 rank()	494
47.33.3.32 recv()	494
47.33.3.33 reduce()	495
47.33.3.34 scan()	495

47.33.3.35 scatter()	496
47.33.3.36 scatter_v()	497
47.33.3.37 send()	497
47.33.3.38 size()	498
47.33.3.39 synchronize()	498
47.34 gko::Composition< ValueType > Class Template Reference	498
47.34.1 Detailed Description	499
47.34.2 Constructor & Destructor Documentation	499
47.34.2.1 Composition() [1/2]	500
47.34.2.2 Composition() [2/2]	500
47.34.3 Member Function Documentation	500
47.34.3.1 conj_transpose()	500
47.34.3.2 get_operators()	500
47.34.3.3 operator=() [1/2]	501
47.34.3.4 operator=() [2/2]	501
47.34.3.5 transpose()	501
47.35 gko::experimental::mpi::contiguous_type Class Reference	501
47.35.1 Detailed Description	502
47.35.2 Constructor & Destructor Documentation	502
47.35.2.1 contiguous_type() [1/2]	502
47.35.2.2 contiguous_type() [2/2]	502
47.35.3 Member Function Documentation	503
47.35.3.1 get()	503
47.35.3.2 operator=()	503
47.36 gko::log::Convergence< ValueType > Class Template Reference	503
47.36.1 Detailed Description	504
47.36.2 Member Function Documentation	504
47.36.2.1 create() [1/2]	504
47.36.2.2 create() [2/2]	505
47.36.2.3 get_implicit_sq_resnorm()	505
47.36.2.4 get_num_iterations()	506
47.36.2.5 get_residual()	506
47.36.2.6 get_residual_norm()	506
47.36.2.7 has_converged()	506
47.37 gko::ConvertibleTo< ResultType > Class Template Reference	507
47.37.1 Detailed Description	507
47.37.2 Member Function Documentation	507
47.37.2.1 convert_to()	507
47.37.2.2 move_to()	508
47.38 gko::matrix::Coo< ValueType, IndexType > Class Template Reference	509
47.38.1 Detailed Description	511
47.38.2 Member Function Documentation	511

47.38.2.1 apply2() [1/4]	511
47.38.2.2 apply2() [2/4]	511
47.38.2.3 apply2() [3/4]	512
47.38.2.4 apply2() [4/4]	512
47.38.2.5 compute_absolute()	512
47.38.2.6 create() [1/3]	512
47.38.2.7 create() [2/3]	513
47.38.2.8 create() [3/3]	513
47.38.2.9 create_const()	514
47.38.2.10 extract_diagonal()	514
47.38.2.11 get_col_idxes()	515
47.38.2.12 get_const_col_idxes()	515
47.38.2.13 get_const_row_idxes()	516
47.38.2.14 get_const_values()	516
47.38.2.15 get_num_stored_elements()	516
47.38.2.16 get_row_idxes()	517
47.38.2.17 get_values()	517
47.38.2.18 read() [1/3]	517
47.38.2.19 read() [2/3]	517
47.38.2.20 read() [3/3]	518
47.38.2.21 write()	518
47.39 gko::CpuAllocator Class Reference	518
47.39.1 Detailed Description	519
47.40 gko::CpuAllocatorBase Class Reference	519
47.40.1 Detailed Description	519
47.41 gko::CpuTimer Class Reference	519
47.41.1 Detailed Description	519
47.41.2 Member Function Documentation	520
47.41.2.1 difference_async()	520
47.42 gko::cpx_real_type< T > Struct Template Reference	520
47.42.1 Detailed Description	520
47.42.2 Member Typedef Documentation	521
47.42.2.1 type	521
47.43 gko::stop::Criterion Class Reference	521
47.43.1 Detailed Description	521
47.43.2 Member Function Documentation	522
47.43.2.1 check()	522
47.43.2.2 update()	522
47.44 gko::log::criterion_data Struct Reference	523
47.44.1 Detailed Description	523
47.45 gko::stop::CriterionArgs Struct Reference	523
47.45.1 Detailed Description	523

47.46 gko::batch::matrix::Csr< ValueType, IndexType > Class Template Reference	523
47.46.1 Detailed Description	525
47.46.2 Member Function Documentation	525
47.46.2.1 add_scaled_identity()	525
47.46.2.2 apply() [1/4]	526
47.46.2.3 apply() [2/4]	526
47.46.2.4 apply() [3/4]	526
47.46.2.5 apply() [4/4]	527
47.46.2.6 create() [1/3]	527
47.46.2.7 create() [2/3]	528
47.46.2.8 create() [3/3]	528
47.46.2.9 create_const()	528
47.46.2.10 create_const_view_for_item()	529
47.46.2.11 create_view_for_item()	529
47.46.2.12 get_col_idxxs()	530
47.46.2.13 get_const_col_idxxs()	530
47.46.2.14 get_const_row_ptrs()	531
47.46.2.15 get_const_values()	531
47.46.2.16 get_const_values_for_item()	531
47.46.2.17 get_num_elements_per_item()	532
47.46.2.18 get_num_stored_elements()	532
47.46.2.19 get_row_ptrs()	533
47.46.2.20 get_values()	533
47.46.2.21 get_values_for_item()	533
47.46.2.22 scale()	534
47.47 gko::matrix::Csr< ValueType, IndexType > Class Template Reference	534
47.47.1 Detailed Description	537
47.47.2 Constructor & Destructor Documentation	538
47.47.2.1 Csr() [1/2]	538
47.47.2.2 Csr() [2/2]	538
47.47.3 Member Function Documentation	538
47.47.3.1 column_permute()	538
47.47.3.2 compute_absolute()	539
47.47.3.3 conj_transpose()	539
47.47.3.4 create() [1/4]	539
47.47.3.5 create() [2/4]	540
47.47.3.6 create() [3/4]	540
47.47.3.7 create() [4/4]	541
47.47.3.8 create_const()	541
47.47.3.9 create_submatrix() [1/2]	542
47.47.3.10 create_submatrix() [2/2]	542
47.47.3.11 extract_diagonal()	543

47.47.3.12	get_col_idx()	543
47.47.3.13	get_const_col_idx()	544
47.47.3.14	get_const_row_ptr()	544
47.47.3.15	get_const_srow()	544
47.47.3.16	get_const_values()	545
47.47.3.17	get_num_srow_elements()	545
47.47.3.18	get_num_stored_elements()	545
47.47.3.19	get_row_ptr()	546
47.47.3.20	get_srow()	546
47.47.3.21	get_strategy()	546
47.47.3.22	get_values()	546
47.47.3.23	inv_scale()	546
47.47.3.24	inverse_column_permute()	547
47.47.3.25	inverse_permute()	547
47.47.3.26	inverse_row_permute()	548
47.47.3.27	operator=() [1/2]	548
47.47.3.28	operator=() [2/2]	548
47.47.3.29	permute() [1/3]	549
47.47.3.30	permute() [2/3]	549
47.47.3.31	permute() [3/3]	550
47.47.3.32	permute_reuse() [1/2]	550
47.47.3.33	permute_reuse() [2/2]	551
47.47.3.34	read() [1/3]	551
47.47.3.35	read() [2/3]	551
47.47.3.36	read() [3/3]	552
47.47.3.37	row_permute()	552
47.47.3.38	scale()	553
47.47.3.39	scale_permute() [1/2]	553
47.47.3.40	scale_permute() [2/2]	553
47.47.3.41	set_strategy()	554
47.47.3.42	transpose()	554
47.47.3.43	transpose_reuse()	555
47.47.3.44	write()	555
47.48	gko::CublasError Class Reference	555
47.48.1	Detailed Description	556
47.48.2	Constructor & Destructor Documentation	556
47.48.2.1	CublasError()	556
47.49	gko::cuda_stream Class Reference	556
47.49.1	Detailed Description	557
47.49.2	Constructor & Destructor Documentation	557
47.49.2.1	cuda_stream()	557
47.49.3	Member Function Documentation	557

47.49.3.1 get()	557
47.50 gko::CudaAllocator Class Reference	557
47.50.1 Detailed Description	558
47.51 gko::CudaAllocatorBase Class Reference	558
47.51.1 Detailed Description	558
47.52 gko::CudaError Class Reference	558
47.52.1 Detailed Description	558
47.52.2 Constructor & Destructor Documentation	558
47.52.2.1 CudaError()	558
47.53 gko::CudaExecutor Class Reference	559
47.53.1 Detailed Description	560
47.53.2 Member Function Documentation	560
47.53.2.1 create() [1/2]	560
47.53.2.2 create() [2/2]	561
47.53.2.3 get_blas_handle()	561
47.53.2.4 get_closest_numa()	561
47.53.2.5 get_closest_pus()	562
47.53.2.6 get_cublas_handle()	562
47.53.2.7 get_cusparselib_handle()	562
47.53.2.8 get_description()	562
47.53.2.9 get_master() [1/2]	563
47.53.2.10 get_master() [2/2]	563
47.53.2.11 get_sparselib_handle()	563
47.53.2.12 get_stream()	563
47.53.2.13 run() [1/3]	564
47.53.2.14 run() [2/3]	565
47.53.2.15 run() [3/3]	565
47.54 gko::CudaTimer Class Reference	566
47.54.1 Detailed Description	566
47.54.2 Member Function Documentation	566
47.54.2.1 difference_async()	566
47.55 gko::CufftError Class Reference	567
47.55.1 Detailed Description	567
47.55.2 Constructor & Destructor Documentation	567
47.55.2.1 CufftError()	567
47.56 gko::CurandError Class Reference	568
47.56.1 Detailed Description	568
47.56.2 Constructor & Destructor Documentation	568
47.56.2.1 CurandError()	568
47.57 gko::matrix::Csr< ValueType, IndexType >::cusparselib Class Reference	569
47.57.1 Detailed Description	569
47.57.2 Member Function Documentation	569

47.57.2.1	clac_size()	569
47.57.2.2	copy()	570
47.57.2.3	process()	570
47.58	gko::CusparsError Class Reference	571
47.58.1	Detailed Description	571
47.58.2	Constructor & Destructor Documentation	571
47.58.2.1	CusparsError()	571
47.59	gko::default_converter< S, R > Struct Template Reference	571
47.59.1	Detailed Description	572
47.59.2	Member Function Documentation	572
47.59.2.1	operator()()	572
47.60	gko::deferred_factory_parameter< FactoryType > Class Template Reference	572
47.60.1	Detailed Description	573
47.60.2	Constructor & Destructor Documentation	573
47.60.2.1	deferred_factory_parameter()	573
47.60.3	Member Function Documentation	574
47.60.3.1	on()	574
47.61	gko::batch::matrix::Dense< ValueType > Class Template Reference	574
47.61.1	Detailed Description	575
47.61.2	Member Function Documentation	577
47.61.2.1	add_scaled_identity()	577
47.61.2.2	apply() [1/4]	577
47.61.2.3	apply() [2/4]	578
47.61.2.4	apply() [3/4]	578
47.61.2.5	apply() [4/4]	578
47.61.2.6	at() [1/4]	578
47.61.2.7	at() [2/4]	579
47.61.2.8	at() [3/4]	579
47.61.2.9	at() [4/4]	581
47.61.2.10	create() [1/3]	581
47.61.2.11	create() [2/3]	582
47.61.2.12	create() [3/3]	582
47.61.2.13	create_const()	583
47.61.2.14	create_const_view_for_item()	583
47.61.2.15	create_view_for_item()	584
47.61.2.16	get_const_values()	584
47.61.2.17	get_const_values_for_item()	584
47.61.2.18	get_cumulative_offset()	585
47.61.2.19	get_num_elements_per_item()	585
47.61.2.20	get_num_stored_elements()	586
47.61.2.21	get_values()	586
47.61.2.22	get_values_for_item()	586

47.61.2.23 scale()	587
47.61.2.24 scale_add()	587
47.62 gko::matrix::Dense< ValueType > Class Template Reference	587
47.62.1 Detailed Description	592
47.62.2 Constructor & Destructor Documentation	592
47.62.2.1 Dense() [1/2]	593
47.62.2.2 Dense() [2/2]	593
47.62.3 Member Function Documentation	593
47.62.3.1 add_scaled()	593
47.62.3.2 at() [1/4]	593
47.62.3.3 at() [2/4]	594
47.62.3.4 at() [3/4]	594
47.62.3.5 at() [4/4]	595
47.62.3.6 column_permute() [1/4]	595
47.62.3.7 column_permute() [2/4]	596
47.62.3.8 column_permute() [3/4]	596
47.62.3.9 column_permute() [4/4]	597
47.62.3.10 compute_absolute() [1/2]	597
47.62.3.11 compute_absolute() [2/2]	597
47.62.3.12 compute_conj_dot() [1/2]	597
47.62.3.13 compute_conj_dot() [2/2]	598
47.62.3.14 compute_dot() [1/2]	598
47.62.3.15 compute_dot() [2/2]	598
47.62.3.16 compute_mean() [1/2]	599
47.62.3.17 compute_mean() [2/2]	599
47.62.3.18 compute_norm1() [1/2]	600
47.62.3.19 compute_norm1() [2/2]	600
47.62.3.20 compute_norm2() [1/2]	600
47.62.3.21 compute_norm2() [2/2]	601
47.62.3.22 compute_squared_norm2() [1/2]	601
47.62.3.23 compute_squared_norm2() [2/2]	601
47.62.3.24 conj_transpose() [1/2]	602
47.62.3.25 conj_transpose() [2/2]	602
47.62.3.26 create() [1/3]	602
47.62.3.27 create() [2/3]	603
47.62.3.28 create() [3/3]	603
47.62.3.29 create_const()	604
47.62.3.30 create_const_view_of()	604
47.62.3.31 create_real_view() [1/2]	604
47.62.3.32 create_real_view() [2/2]	605
47.62.3.33 create_submatrix() [1/2]	605
47.62.3.34 create_submatrix() [2/2]	605

47.62.3.35 create_view_of()	606
47.62.3.36 create_with_config_of()	606
47.62.3.37 create_with_type_of() [1/3]	606
47.62.3.38 create_with_type_of() [2/3]	607
47.62.3.39 create_with_type_of() [3/3]	607
47.62.3.40 extract_diagonal() [1/2]	608
47.62.3.41 extract_diagonal() [2/2]	608
47.62.3.42 fill()	608
47.62.3.43 get_const_values()	609
47.62.3.44 get_num_stored_elements()	609
47.62.3.45 get_stride()	609
47.62.3.46 get_values()	610
47.62.3.47 inv_scale()	610
47.62.3.48 inverse_column_permute() [1/4]	610
47.62.3.49 inverse_column_permute() [2/4]	611
47.62.3.50 inverse_column_permute() [3/4]	611
47.62.3.51 inverse_column_permute() [4/4]	611
47.62.3.52 inverse_permute() [1/4]	612
47.62.3.53 inverse_permute() [2/4]	612
47.62.3.54 inverse_permute() [3/4]	612
47.62.3.55 inverse_permute() [4/4]	613
47.62.3.56 inverse_row_permute() [1/4]	613
47.62.3.57 inverse_row_permute() [2/4]	614
47.62.3.58 inverse_row_permute() [3/4]	614
47.62.3.59 inverse_row_permute() [4/4]	614
47.62.3.60 make_complex() [1/2]	615
47.62.3.61 make_complex() [2/2]	615
47.62.3.62 operator=() [1/2]	615
47.62.3.63 operator=() [2/2]	615
47.62.3.64 permute() [1/12]	615
47.62.3.65 permute() [2/12]	616
47.62.3.66 permute() [3/12]	616
47.62.3.67 permute() [4/12]	617
47.62.3.68 permute() [5/12]	617
47.62.3.69 permute() [6/12]	617
47.62.3.70 permute() [7/12]	618
47.62.3.71 permute() [8/12]	618
47.62.3.72 permute() [9/12]	619
47.62.3.73 permute() [10/12]	619
47.62.3.74 permute() [11/12]	619
47.62.3.75 permute() [12/12]	619
47.62.3.76 row_gather() [1/6]	619

47.62.3.77 row_gather() [2/6]	620
47.62.3.78 row_gather() [3/6]	620
47.62.3.79 row_gather() [4/6]	621
47.62.3.80 row_gather() [5/6]	621
47.62.3.81 row_gather() [6/6]	621
47.62.3.82 row_permute() [1/4]	621
47.62.3.83 row_permute() [2/4]	622
47.62.3.84 row_permute() [3/4]	622
47.62.3.85 row_permute() [4/4]	623
47.62.3.86 scale()	623
47.62.3.87 scale_permute() [1/8]	623
47.62.3.88 scale_permute() [2/8]	624
47.62.3.89 scale_permute() [3/8]	624
47.62.3.90 scale_permute() [4/8]	625
47.62.3.91 scale_permute() [5/8]	625
47.62.3.92 scale_permute() [6/8]	625
47.62.3.93 scale_permute() [7/8]	625
47.62.3.94 scale_permute() [8/8]	626
47.62.3.95 sub_scaled()	626
47.62.3.96 transpose() [1/2]	626
47.62.3.97 transpose() [2/2]	626
47.63 gko::device_matrix_data< ValueType, IndexType > Class Template Reference	627
47.63.1 Detailed Description	628
47.63.2 Constructor & Destructor Documentation	629
47.63.2.1 device_matrix_data() [1/4]	629
47.63.2.2 device_matrix_data() [2/4]	629
47.63.2.3 device_matrix_data() [3/4]	629
47.63.2.4 device_matrix_data() [4/4]	630
47.63.3 Member Function Documentation	630
47.63.3.1 copy_to_host()	630
47.63.3.2 create_from_host()	631
47.63.3.3 empty_out()	631
47.63.3.4 get_col_idxs()	631
47.63.3.5 get_const_col_idxs()	632
47.63.3.6 get_const_row_idxs()	632
47.63.3.7 get_const_values()	632
47.63.3.8 get_executor()	633
47.63.3.9 get_num_elems()	633
47.63.3.10 get_num_stored_elements()	633
47.63.3.11 get_row_idxs()	634
47.63.3.12 get_size()	634
47.63.3.13 get_values()	634

47.63.3.14 remove_zeros()	634
47.63.3.15 resize_and_reset() [1/2]	634
47.63.3.16 resize_and_reset() [2/2]	635
47.63.3.17 sum_duplicates()	635
47.64 gko::matrix::Diagonal< ValueType > Class Template Reference	635
47.64.1 Detailed Description	636
47.64.2 Member Function Documentation	637
47.64.2.1 compute_absolute()	637
47.64.2.2 conj_transpose()	637
47.64.2.3 create() [1/3]	637
47.64.2.4 create() [2/3]	638
47.64.2.5 create() [3/3]	638
47.64.2.6 create_const()	639
47.64.2.7 get_const_values()	639
47.64.2.8 get_values()	639
47.64.2.9 inverse_apply()	640
47.64.2.10 rapply()	640
47.64.2.11 transpose()	640
47.65 gko::DiagonalExtractable< ValueType > Class Template Reference	641
47.65.1 Detailed Description	641
47.65.2 Member Function Documentation	641
47.65.2.1 extract_diagonal()	641
47.65.2.2 extract_diagonal_linop()	642
47.66 gko::DiagonalLinOpExtractable Class Reference	642
47.66.1 Detailed Description	642
47.66.2 Member Function Documentation	642
47.66.2.1 extract_diagonal_linop()	643
47.67 gko::dim< Dimensionality, DimensionType > Struct Template Reference	643
47.67.1 Detailed Description	643
47.67.2 Constructor & Destructor Documentation	644
47.67.2.1 dim() [1/2]	644
47.67.2.2 dim() [2/2]	644
47.67.3 Member Function Documentation	644
47.67.3.1 operator bool()	645
47.67.3.2 operator[]() [1/2]	645
47.67.3.3 operator[]() [2/2]	645
47.67.4 Friends And Related Function Documentation	646
47.67.4.1 operator"!="	646
47.67.4.2 operator*	646
47.67.4.3 operator<<	646
47.67.4.4 operator==	647
47.68 gko::DimensionMismatch Class Reference	647

47.68.1 Detailed Description	648
47.68.2 Constructor & Destructor Documentation	648
47.68.2.1 DimensionMismatch()	648
47.69 gko::experimental::solver::Direct< ValueType, IndexType > Class Template Reference	648
47.69.1 Detailed Description	649
47.69.2 Member Function Documentation	649
47.69.2.1 conj_transpose()	649
47.69.2.2 parse()	650
47.69.2.3 transpose()	650
47.70 gko::experimental::distributed::DistributedBase Class Reference	650
47.70.1 Detailed Description	651
47.70.2 Member Function Documentation	651
47.70.2.1 get_communicator()	651
47.70.2.2 operator=() [1/2]	651
47.70.2.3 operator=() [2/2]	652
47.71 gko::DpcppExecutor Class Reference	652
47.71.1 Detailed Description	653
47.71.2 Member Function Documentation	653
47.71.2.1 create()	653
47.71.2.2 get_description()	653
47.71.2.3 get_device_id()	654
47.71.2.4 get_device_type()	654
47.71.2.5 get_master() [1/2]	654
47.71.2.6 get_master() [2/2]	654
47.71.2.7 get_max_subgroup_size()	655
47.71.2.8 get_max_workgroup_size()	655
47.71.2.9 get_max_workitem_sizes()	655
47.71.2.10 get_num_computing_units()	655
47.71.2.11 get_num_devices()	655
47.71.2.12 get_subgroup_sizes()	656
47.71.2.13 run() [1/3]	656
47.71.2.14 run() [2/3]	657
47.71.2.15 run() [3/3]	657
47.72 gko::DpcppTimer Class Reference	658
47.72.1 Detailed Description	658
47.72.2 Member Function Documentation	658
47.72.2.1 difference_async()	658
47.73 gko::batch::matrix::El< ValueType, IndexType > Class Template Reference	659
47.73.1 Detailed Description	660
47.73.2 Member Function Documentation	661
47.73.2.1 add_scaled_identity()	661
47.73.2.2 apply() [1/4]	661

47.73.2.3 apply() [2/4]	662
47.73.2.4 apply() [3/4]	662
47.73.2.5 apply() [4/4]	662
47.73.2.6 create() [1/3]	662
47.73.2.7 create() [2/3]	663
47.73.2.8 create() [3/3]	663
47.73.2.9 create_const()	664
47.73.2.10 create_const_view_for_item()	664
47.73.2.11 create_view_for_item()	665
47.73.2.12 get_col_idxxs()	665
47.73.2.13 get_col_idxxs_for_item()	666
47.73.2.14 get_const_col_idxxs()	666
47.73.2.15 get_const_col_idxxs_for_item()	666
47.73.2.16 get_const_values()	667
47.73.2.17 get_const_values_for_item()	667
47.73.2.18 get_num_elements_per_item()	668
47.73.2.19 get_num_stored_elements()	668
47.73.2.20 get_num_stored_elements_per_row()	669
47.73.2.21 get_values()	669
47.73.2.22 get_values_for_item()	669
47.73.2.23 scale()	670
47.74 gko::matrix::Ell< ValueType, IndexType > Class Template Reference	670
47.74.1 Detailed Description	672
47.74.2 Constructor & Destructor Documentation	672
47.74.2.1 Ell() [1/2]	672
47.74.2.2 Ell() [2/2]	672
47.74.3 Member Function Documentation	672
47.74.3.1 col_at() [1/2]	673
47.74.3.2 col_at() [2/2]	673
47.74.3.3 compute_absolute()	674
47.74.3.4 create() [1/3]	674
47.74.3.5 create() [2/3]	675
47.74.3.6 create() [3/3]	675
47.74.3.7 create_const()	675
47.74.3.8 extract_diagonal()	676
47.74.3.9 get_col_idxxs()	676
47.74.3.10 get_const_col_idxxs()	677
47.74.3.11 get_const_values()	677
47.74.3.12 get_num_stored_elements()	678
47.74.3.13 get_num_stored_elements_per_row()	678
47.74.3.14 get_stride()	678
47.74.3.15 get_values()	678

47.74.3.16 operator=() [1/2]	679
47.74.3.17 operator=() [2/2]	679
47.74.3.18 read() [1/3]	679
47.74.3.19 read() [2/3]	679
47.74.3.20 read() [3/3]	680
47.74.3.21 val_at() [1/2]	680
47.74.3.22 val_at() [2/2]	680
47.74.3.23 write()	681
47.75 gko::enable_parameters_type< ConcreteParametersType, Factory > Class Template Reference	681
47.75.1 Detailed Description	682
47.75.2 Member Function Documentation	682
47.75.2.1 on()	682
47.76 gko::EnableAbsoluteComputation< AbsoluteLinOp > Class Template Reference	682
47.76.1 Detailed Description	683
47.76.2 Member Function Documentation	683
47.76.2.1 compute_absolute()	683
47.76.2.2 compute_absolute_linop()	684
47.77 gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase > Class Template Reference	684
47.77.1 Detailed Description	684
47.78 gko::solver::EnableApplyWithInitialGuess< DerivedType > Class Template Reference	685
47.78.1 Detailed Description	685
47.79 gko::batch::EnableBatchLinOp< ConcreteBatchLinOp, PolymorphicBase > Class Template Reference	685
47.79.1 Detailed Description	686
47.80 gko::batch::solver::EnableBatchSolver< ConcreteSolver, ValueType, PolymorphicBase > Class Template Reference	686
47.80.1 Detailed Description	686
47.81 gko::EnableCreateMethod< ConcreteType > Class Template Reference	687
47.81.1 Detailed Description	687
47.82 gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase > Class Template Reference	687
47.82.1 Detailed Description	688
47.82.2 Member Function Documentation	688
47.82.2.1 create()	688
47.82.2.2 get_parameters()	689
47.83 gko::solver::EnableIterativeBase< DerivedType > Class Template Reference	689
47.83.1 Detailed Description	689
47.83.2 Constructor & Destructor Documentation	690
47.83.2.1 EnableIterativeBase()	690
47.83.3 Member Function Documentation	690
47.83.3.1 operator=()	690
47.83.3.2 set_stop_criterion_factory()	690

47.84	gko::EnableLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference	691
47.84.1	Detailed Description	691
47.85	gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase > Class Template Reference	692
47.85.1	Detailed Description	692
47.86	gko::multigrid::EnableMultigridLevel< ValueType > Class Template Reference	692
47.86.1	Detailed Description	693
47.86.2	Member Function Documentation	693
47.86.2.1	get_coarse_op()	693
47.86.2.2	get_fine_op()	693
47.86.2.3	get_prolong_op()	694
47.86.2.4	get_restrict_op()	694
47.87	gko::EnablePolymorphicAssignment< ConcreteType, ResultType > Class Template Reference	694
47.87.1	Detailed Description	695
47.87.2	Member Function Documentation	696
47.87.2.1	convert_to()	696
47.87.2.2	move_to()	696
47.88	gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase > Class Template Reference	697
47.88.1	Detailed Description	697
47.89	gko::solver::EnablePreconditionable< DerivedType > Class Template Reference	697
47.89.1	Detailed Description	698
47.89.2	Constructor & Destructor Documentation	698
47.89.2.1	EnablePreconditionable()	698
47.89.3	Member Function Documentation	698
47.89.3.1	operator=()	699
47.89.3.2	set_preconditioner()	699
47.90	gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType > Class Template Reference	699
47.90.1	Detailed Description	699
47.91	gko::solver::EnableSolverBase< DerivedType, MatrixType > Class Template Reference	700
47.91.1	Detailed Description	700
47.91.2	Constructor & Destructor Documentation	701
47.91.2.1	EnableSolverBase()	701
47.91.3	Member Function Documentation	701
47.91.3.1	operator=()	701
47.92	gko::experimental::mpi::environment Class Reference	701
47.92.1	Detailed Description	702
47.92.2	Constructor & Destructor Documentation	702
47.92.2.1	environment()	702
47.92.3	Member Function Documentation	702
47.92.3.1	get_provided_thread_support()	702
47.93	gko::Error Class Reference	703
47.93.1	Detailed Description	703

47.93.2 Constructor & Destructor Documentation	703
47.93.2.1 Error()	703
47.94 gko::Executor Class Reference	704
47.94.1 Detailed Description	705
47.94.2 Member Function Documentation	706
47.94.2.1 add_logger()	706
47.94.2.2 alloc()	706
47.94.2.3 copy()	707
47.94.2.4 copy_from()	707
47.94.2.5 copy_val_to_host()	708
47.94.2.6 free()	709
47.94.2.7 get_description()	709
47.94.2.8 get_master() [1/2]	709
47.94.2.9 get_master() [2/2]	709
47.94.2.10 memory_accessible()	710
47.94.2.11 remove_logger()	710
47.94.2.12 run() [1/3]	710
47.94.2.13 run() [2/3]	711
47.94.2.14 run() [3/3]	711
47.94.2.15 set_log_propagation_mode()	712
47.94.2.16 should_propagate_log()	712
47.95 gko::log::executor_data Struct Reference	713
47.95.1 Detailed Description	713
47.96 gko::executor_deleter< T > Class Template Reference	713
47.96.1 Detailed Description	713
47.96.2 Constructor & Destructor Documentation	713
47.96.2.1 executor_deleter()	714
47.96.3 Member Function Documentation	714
47.96.3.1 operator()()	714
47.97 gko::experimental::factorization::Factorization< ValueType, IndexType > Class Template Reference	714
47.97.1 Detailed Description	715
47.97.2 Member Function Documentation	715
47.97.2.1 create_from_combined_lu()	716
47.97.2.2 create_from_composition()	716
47.97.2.3 create_from_symm_composition()	716
47.97.2.4 unpack()	717
47.98 gko::matrix::Fbcsr< ValueType, IndexType > Class Template Reference	717
47.98.1 Detailed Description	719
47.98.2 Constructor & Destructor Documentation	720
47.98.2.1 Fbcsr() [1/2]	720
47.98.2.2 Fbcsr() [2/2]	720
47.98.3 Member Function Documentation	720

47.98.3.1 compute_absolute()	720
47.98.3.2 conj_transpose()	721
47.98.3.3 convert_to() [1/2]	721
47.98.3.4 convert_to() [2/2]	721
47.98.3.5 create() [1/4]	721
47.98.3.6 create() [2/4]	722
47.98.3.7 create() [3/4]	722
47.98.3.8 create() [4/4]	723
47.98.3.9 create_const()	723
47.98.3.10 extract_diagonal()	724
47.98.3.11 get_block_size()	724
47.98.3.12 get_col_idxs()	724
47.98.3.13 get_const_col_idxs()	725
47.98.3.14 get_const_row_ptrs()	725
47.98.3.15 get_const_values()	725
47.98.3.16 get_num_block_cols()	726
47.98.3.17 get_num_block_rows()	726
47.98.3.18 get_num_stored_blocks()	726
47.98.3.19 get_num_stored_elements()	726
47.98.3.20 get_row_ptrs()	727
47.98.3.21 get_values()	727
47.98.3.22 is_sorted_by_column_index()	727
47.98.3.23 operator=() [1/2]	727
47.98.3.24 operator=() [2/2]	728
47.98.3.25 read() [1/3]	728
47.98.3.26 read() [2/3]	728
47.98.3.27 read() [3/3]	728
47.98.3.28 transpose()	729
47.98.3.29 write()	729
47.99 gko::solver::Fcg< ValueType > Class Template Reference	729
47.99.1 Detailed Description	730
47.99.2 Member Function Documentation	730
47.99.2.1 apply_uses_initial_guess()	730
47.99.2.2 conj_transpose()	731
47.99.2.3 parse()	731
47.99.2.4 transpose()	732
47.100 gko::matrix::Fft Class Reference	732
47.100.1 Detailed Description	733
47.100.2 Member Function Documentation	733
47.100.2.1 conj_transpose()	733
47.100.2.2 create() [1/2]	733
47.100.2.3 create() [2/2]	734

47.100.2.4 transpose()	734
47.100.2.5 write() [1/4]	734
47.100.2.6 write() [2/4]	735
47.100.2.7 write() [3/4]	735
47.100.2.8 write() [4/4]	735
47.101 gko::matrix::Fft2 Class Reference	736
47.101.1 Detailed Description	736
47.101.2 Member Function Documentation	737
47.101.2.1 conj_transpose()	737
47.101.2.2 create() [1/3]	737
47.101.2.3 create() [2/3]	737
47.101.2.4 create() [3/3]	738
47.101.2.5 transpose()	738
47.101.2.6 write() [1/4]	738
47.101.2.7 write() [2/4]	739
47.101.2.8 write() [3/4]	739
47.101.2.9 write() [4/4]	739
47.102 gko::matrix::Fft3 Class Reference	740
47.102.1 Detailed Description	740
47.102.2 Member Function Documentation	741
47.102.2.1 conj_transpose()	741
47.102.2.2 create() [1/3]	741
47.102.2.3 create() [2/3]	741
47.102.2.4 create() [3/3]	742
47.102.2.5 transpose()	742
47.102.2.6 write() [1/4]	742
47.102.2.7 write() [2/4]	743
47.102.2.8 write() [3/4]	743
47.102.2.9 write() [4/4]	743
47.103 gko::multigrid::FixedCoarsening< ValueType, IndexType > Class Template Reference	744
47.103.1 Detailed Description	744
47.103.2 Member Function Documentation	744
47.103.2.1 get_system_matrix()	744
47.104 gko::preconditioner::GaussSeidel< ValueType, IndexType > Class Template Reference	745
47.104.1 Detailed Description	745
47.104.2 Member Function Documentation	745
47.104.2.1 generate()	746
47.104.2.2 get_parameters() [1/2]	746
47.104.2.3 get_parameters() [2/2]	746
47.105 gko::solver::Gcr< ValueType > Class Template Reference	746
47.105.1 Detailed Description	747
47.105.2 Member Function Documentation	747

47.105.2.1 apply_uses_initial_guess()	747
47.105.2.2 conj_transpose()	748
47.105.2.3 get_krylov_dim()	748
47.105.2.4 parse()	748
47.105.2.5 set_krylov_dim()	749
47.105.2.6 transpose()	749
47.106 gko::solver::Gmres< ValueType > Class Template Reference	749
47.106.1 Detailed Description	750
47.106.2 Member Function Documentation	750
47.106.2.1 apply_uses_initial_guess()	750
47.106.2.2 conj_transpose()	751
47.106.2.3 get_krylov_dim()	751
47.106.2.4 parse()	751
47.106.2.5 set_krylov_dim()	752
47.106.2.6 transpose()	752
47.107 gko::half Class Reference	752
47.107.1 Detailed Description	753
47.108 gko::solver::has_with_criteria< SolverType, typename > Struct Template Reference	753
47.108.1 Detailed Description	753
47.109 gko::hip_stream Class Reference	753
47.109.1 Detailed Description	754
47.109.2 Constructor & Destructor Documentation	754
47.109.2.1 hip_stream()	754
47.109.3 Member Function Documentation	754
47.109.3.1 get()	754
47.110 gko::HipAllocatorBase Class Reference	755
47.110.1 Detailed Description	755
47.111 gko::HipblasError Class Reference	755
47.111.1 Detailed Description	755
47.111.2 Constructor & Destructor Documentation	755
47.111.2.1 HipblasError()	755
47.112 gko::HipError Class Reference	756
47.112.1 Detailed Description	756
47.112.2 Constructor & Destructor Documentation	756
47.112.2.1 HipError()	756
47.113 gko::HipExecutor Class Reference	757
47.113.1 Detailed Description	758
47.113.2 Member Function Documentation	758
47.113.2.1 create()	758
47.113.2.2 get_blas_handle()	759
47.113.2.3 get_closest_numa()	759
47.113.2.4 get_closest_pus()	759

47.113.2.5 get_description()	759
47.113.2.6 get_hipblas_handle()	760
47.113.2.7 get_hipspase_handle()	760
47.113.2.8 get_master() [1/2]	760
47.113.2.9 get_master() [2/2]	760
47.113.2.10 get_sparselib_handle()	761
47.113.2.11 run() [1/3]	761
47.113.2.12 run() [2/3]	761
47.113.2.13 run() [3/3]	762
47.114 gko::HipfftError Class Reference	762
47.114.1 Detailed Description	763
47.114.2 Constructor & Destructor Documentation	763
47.114.2.1 HipfftError()	763
47.115 gko::HiprandError Class Reference	763
47.115.1 Detailed Description	764
47.115.2 Constructor & Destructor Documentation	764
47.115.2.1 HiprandError()	764
47.116 gko::HipspaseError Class Reference	764
47.116.1 Detailed Description	764
47.116.2 Constructor & Destructor Documentation	765
47.116.2.1 HipspaseError()	765
47.117 gko::HipTimer Class Reference	765
47.117.1 Detailed Description	765
47.117.2 Member Function Documentation	766
47.117.2.1 difference_async()	766
47.118 gko::matrix::Hybrid< ValueType, IndexType > Class Template Reference	766
47.118.1 Detailed Description	768
47.118.2 Constructor & Destructor Documentation	769
47.118.2.1 Hybrid() [1/2]	769
47.118.2.2 Hybrid() [2/2]	769
47.118.3 Member Function Documentation	769
47.118.3.1 compute_absolute()	769
47.118.3.2 create() [1/5]	770
47.118.3.3 create() [2/5]	770
47.118.3.4 create() [3/5]	771
47.118.3.5 create() [4/5]	771
47.118.3.6 create() [5/5]	772
47.118.3.7 ell_col_at() [1/2]	772
47.118.3.8 ell_col_at() [2/2]	773
47.118.3.9 ell_val_at() [1/2]	773
47.118.3.10 ell_val_at() [2/2]	773
47.118.3.11 extract_diagonal()	774

47.118.3.12	get_const_coo_col_idxs()	774
47.118.3.13	get_const_coo_row_idxs()	775
47.118.3.14	get_const_coo_values()	775
47.118.3.15	get_const_ell_col_idxs()	775
47.118.3.16	get_const_ell_values()	776
47.118.3.17	get_coo()	776
47.118.3.18	get_coo_col_idxs()	776
47.118.3.19	get_coo_num_stored_elements()	777
47.118.3.20	get_coo_row_idxs()	777
47.118.3.21	get_coo_values()	777
47.118.3.22	get_ell()	777
47.118.3.23	get_ell_col_idxs()	778
47.118.3.24	get_ell_num_stored_elements()	778
47.118.3.25	get_ell_num_stored_elements_per_row()	778
47.118.3.26	get_ell_stride()	778
47.118.3.27	get_ell_values()	779
47.118.3.28	get_num_stored_elements()	779
47.118.3.29	get_strategy() [1/2]	779
47.118.3.30	get_strategy() [2/2]	779
47.118.3.31	operator=() [1/2]	780
47.118.3.32	operator=() [2/2]	780
47.118.3.33	read() [1/3]	780
47.118.3.34	read() [2/3]	780
47.118.3.35	read() [3/3]	781
47.118.3.36	write()	781
47.119	gko::preconditioner::lc< LSolverType, IndexType > Class Template Reference	782
47.119.1	Detailed Description	782
47.119.2	Constructor & Destructor Documentation	783
47.119.2.1	lc() [1/2]	783
47.119.2.2	lc() [2/2]	783
47.119.3	Member Function Documentation	784
47.119.3.1	conj_transpose()	784
47.119.3.2	get_l_solver()	784
47.119.3.3	get_lh_solver()	784
47.119.3.4	operator=() [1/2]	785
47.119.3.5	operator=() [2/2]	785
47.119.3.6	parse()	785
47.119.3.7	transpose()	786
47.120	gko::factorization::lc< ValueType, IndexType > Class Template Reference	786
47.120.1	Detailed Description	786
47.120.2	Member Function Documentation	787
47.120.2.1	parse()	787

47.121	gko::batch::matrix::Identity< ValueType > Class Template Reference	787
47.121.1	Detailed Description	788
47.121.2	Member Function Documentation	788
47.121.2.1	apply() [1/4]	788
47.121.2.2	apply() [2/4]	789
47.121.2.3	apply() [3/4]	789
47.121.2.4	apply() [4/4]	789
47.121.2.5	create()	790
47.122	gko::matrix::Identity< ValueType > Class Template Reference	790
47.122.1	Detailed Description	791
47.122.2	Member Function Documentation	791
47.122.2.1	conj_transpose()	791
47.122.2.2	create() [1/2]	791
47.122.2.3	create() [2/2]	792
47.122.2.4	transpose()	792
47.123	gko::matrix::IdentityFactory< ValueType > Class Template Reference	792
47.123.1	Detailed Description	793
47.123.2	Member Function Documentation	793
47.123.2.1	create()	793
47.124	gko::solver::ldr< ValueType > Class Template Reference	794
47.124.1	Detailed Description	794
47.124.2	Member Function Documentation	795
47.124.2.1	apply_uses_initial_guess()	795
47.124.2.2	conj_transpose()	795
47.124.2.3	get_complex_subspace()	795
47.124.2.4	get_deterministic()	796
47.124.2.5	get_kappa()	796
47.124.2.6	get_subspace_dim()	796
47.124.2.7	parse()	796
47.124.2.8	set_complex_subspace()	797
47.124.2.9	set_complex_subspace()	797
47.124.2.10	set_deterministic()	797
47.124.2.11	set_kappa()	798
47.124.2.12	set_subspace_dim()	798
47.124.2.13	transpose()	798
47.125	gko::factorization::llu< ValueType, IndexType > Class Template Reference	799
47.125.1	Detailed Description	799
47.125.2	Member Function Documentation	799
47.125.2.1	parse()	799
47.126	gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType > Class Template Reference	800
47.126.1	Detailed Description	801

47.126.2 Constructor & Destructor Documentation	801
47.126.2.1 <code>llu()</code> [1/2]	801
47.126.2.2 <code>llu()</code> [2/2]	802
47.126.3 Member Function Documentation	802
47.126.3.1 <code>conj_transpose()</code>	802
47.126.3.2 <code>get_l_solver()</code>	803
47.126.3.3 <code>get_u_solver()</code>	803
47.126.3.4 <code>operator=()</code> [1/2]	803
47.126.3.5 <code>operator=()</code> [2/2]	804
47.126.3.6 <code>parse()</code>	804
47.126.3.7 <code>transpose()</code>	805
47.127 <code>gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit</code> Class Reference	805
47.127.1 Detailed Description	805
47.127.2 Member Function Documentation	806
47.127.2.1 <code>compute_ell_num_stored_elements_per_row()</code>	806
47.127.2.2 <code>get_percentage()</code>	806
47.127.2.3 <code>get_ratio()</code>	807
47.128 <code>gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit</code> Class Reference	807
47.128.1 Detailed Description	807
47.128.2 Constructor & Destructor Documentation	807
47.128.2.1 <code>imbalance_limit()</code>	807
47.128.3 Member Function Documentation	808
47.128.3.1 <code>compute_ell_num_stored_elements_per_row()</code>	808
47.128.3.2 <code>get_percentage()</code>	808
47.129 <code>gko::stop::ImplicitResidualNorm< ValueType ></code> Class Template Reference	809
47.129.1 Detailed Description	809
47.130 <code>gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType ></code> Struct Template Reference	809
47.130.1 Detailed Description	810
47.130.2 Constructor & Destructor Documentation	811
47.130.2.1 <code>index_map()</code>	811
47.130.3 Member Function Documentation	811
47.130.3.1 <code>get_remote_global_idxs()</code>	811
47.130.3.2 <code>get_remote_local_idxs()</code>	812
47.130.3.3 <code>get_remote_target_ids()</code>	812
47.130.3.4 <code>map_to_local()</code>	812
47.131 <code>gko::index_set< IndexType ></code> Class Template Reference	812
47.131.1 Detailed Description	814
47.131.2 Constructor & Destructor Documentation	814
47.131.2.1 <code>index_set()</code> [1/7]	814
47.131.2.2 <code>index_set()</code> [2/7]	815
47.131.2.3 <code>index_set()</code> [3/7]	815

47.131.2.4 index_set() [4/7]	816
47.131.2.5 index_set() [5/7]	816
47.131.2.6 index_set() [6/7]	816
47.131.2.7 index_set() [7/7]	817
47.131.3 Member Function Documentation	817
47.131.3.1 clear()	817
47.131.3.2 contains() [1/2]	817
47.131.3.3 contains() [2/2]	818
47.131.3.4 get_executor()	818
47.131.3.5 get_global_index()	818
47.131.3.6 get_local_index()	819
47.131.3.7 get_num_elems()	819
47.131.3.8 get_num_subsets()	820
47.131.3.9 get_size()	820
47.131.3.10 get_subsets_begin()	820
47.131.3.11 get_subsets_end()	821
47.131.3.12 get_superset_indices()	821
47.131.3.13 is_contiguous()	821
47.131.3.14 map_global_to_local()	821
47.131.3.15 map_local_to_global()	822
47.131.3.16 operator=() [1/2]	822
47.131.3.17 operator=() [2/2]	823
47.131.3.18 to_global_indices()	823
47.132 gko::InvalidStateError Class Reference	824
47.132.1 Detailed Description	824
47.132.2 Constructor & Destructor Documentation	824
47.132.2.1 InvalidStateError()	824
47.133 gko::solver::lr< ValueType > Class Template Reference	824
47.133.1 Detailed Description	825
47.133.2 Constructor & Destructor Documentation	826
47.133.2.1 lr() [1/2]	826
47.133.2.2 lr() [2/2]	826
47.133.3 Member Function Documentation	826
47.133.3.1 apply_uses_initial_guess()	827
47.133.3.2 conj_transpose()	827
47.133.3.3 get_solver()	827
47.133.3.4 operator=() [1/2]	827
47.133.3.5 operator=() [2/2]	828
47.133.3.6 parse()	828
47.133.3.7 set_solver()	828
47.133.3.8 transpose()	829
47.134 gko::preconditioner::lsai< IsaiType, ValueType, IndexType > Class Template Reference	829

47.134.1 Detailed Description	830
47.134.2 Constructor & Destructor Documentation	830
47.134.2.1 Isai() [1/2]	830
47.134.2.2 Isai() [2/2]	831
47.134.3 Member Function Documentation	831
47.134.3.1 conj_transpose()	831
47.134.3.2 get_approximate_inverse()	831
47.134.3.3 operator=() [1/2]	832
47.134.3.4 operator=() [2/2]	832
47.134.3.5 parse()	832
47.134.3.6 transpose()	832
47.135 gko::stop::iteration Class Reference	833
47.135.1 Detailed Description	833
47.136 gko::log::iteration_complete_data Struct Reference	833
47.136.1 Detailed Description	833
47.137 gko::solver::iterativeBase Class Reference	834
47.137.1 Detailed Description	834
47.137.2 Member Function Documentation	834
47.137.2.1 get_stop_criterion_factory()	834
47.137.2.2 set_stop_criterion_factory()	834
47.138 gko::preconditioner::Jacobi< ValueType, IndexType > Class Template Reference	835
47.138.1 Detailed Description	836
47.138.2 Constructor & Destructor Documentation	836
47.138.2.1 Jacobi() [1/2]	836
47.138.2.2 Jacobi() [2/2]	837
47.138.3 Member Function Documentation	837
47.138.3.1 conj_transpose()	837
47.138.3.2 convert_to()	837
47.138.3.3 get_blocks()	838
47.138.3.4 get_conditioning()	838
47.138.3.5 get_num_blocks()	838
47.138.3.6 get_num_stored_elements()	839
47.138.3.7 get_storage_scheme()	839
47.138.3.8 move_to()	839
47.138.3.9 operator=() [1/2]	840
47.138.3.10 operator=() [2/2]	840
47.138.3.11 parse()	840
47.138.3.12 transpose()	841
47.138.3.13 write()	841
47.139 gko::batch::preconditioner::Jacobi< ValueType, IndexType > Class Template Reference	841
47.139.1 Detailed Description	842
47.139.2 Member Function Documentation	842

47.139.2.1	get_const_block_pointers()	842
47.139.2.2	get_const_blocks()	843
47.139.2.3	get_const_blocks_cumulative_offsets()	843
47.139.2.4	get_const_map_block_to_row()	844
47.139.2.5	get_max_block_size()	844
47.139.2.6	get_num_blocks()	844
47.139.2.7	get_num_stored_elements()	845
47.140	gko::KernelNotFound Class Reference	845
47.140.1	Detailed Description	845
47.140.2	Constructor & Destructor Documentation	845
47.140.2.1	KernelNotFound()	845
47.141	gko::log::linop_data Struct Reference	846
47.141.1	Detailed Description	846
47.142	gko::log::linop_factory_data Struct Reference	846
47.142.1	Detailed Description	846
47.143	gko::LinOpFactory Class Reference	846
47.143.1	Detailed Description	847
47.143.1.1	Example: using CG in Ginkgo	847
47.144	gko::matrix::Csr< ValueType, IndexType >::load_balance Class Reference	847
47.144.1	Detailed Description	848
47.144.2	Constructor & Destructor Documentation	848
47.144.2.1	load_balance() [1/5]	848
47.144.2.2	load_balance() [2/5]	848
47.144.2.3	load_balance() [3/5]	849
47.144.2.4	load_balance() [4/5]	849
47.144.2.5	load_balance() [5/5]	849
47.144.3	Member Function Documentation	850
47.144.3.1	clac_size()	850
47.144.3.2	copy()	850
47.144.3.3	process()	851
47.145	gko::log::Loggable Class Reference	851
47.145.1	Detailed Description	851
47.145.2	Member Function Documentation	851
47.145.2.1	add_logger()	851
47.145.2.2	get_loggers()	852
47.145.2.3	remove_logger()	852
47.146	gko::log::Record::logged_data Struct Reference	852
47.146.1	Detailed Description	853
47.147	gko::solver::LowerTrs< ValueType, IndexType > Class Template Reference	853
47.147.1	Detailed Description	853
47.147.2	Constructor & Destructor Documentation	854
47.147.2.1	LowerTrs() [1/2]	854

47.147.2.2 LowerTrs() [2/2]	854
47.147.3 Member Function Documentation	854
47.147.3.1 conj_transpose()	854
47.147.3.2 operator=() [1/2]	855
47.147.3.3 operator=() [2/2]	855
47.147.3.4 parse()	855
47.147.3.5 transpose()	855
47.148 gko::experimental::factorization::Lu< ValueType, IndexType > Class Template Reference	856
47.148.1 Detailed Description	856
47.148.2 Member Function Documentation	857
47.148.2.1 generate()	857
47.148.2.2 get_parameters() [1/2]	857
47.148.2.3 get_parameters() [2/2]	857
47.148.2.4 parse()	858
47.149 gko::machine_topology Class Reference	858
47.149.1 Detailed Description	859
47.149.2 Member Function Documentation	859
47.149.2.1 bind_to_core()	859
47.149.2.2 bind_to_cores()	860
47.149.2.3 bind_to_pu()	860
47.149.2.4 bind_to_pus()	860
47.149.2.5 get_core()	862
47.149.2.6 get_instance()	862
47.149.2.7 get_num_cores()	862
47.149.2.8 get_num_numas()	863
47.149.2.9 get_num_pci_devices()	863
47.149.2.10 get_num_pus()	863
47.149.2.11 get_pci_device() [1/2]	863
47.149.2.12 get_pci_device() [2/2]	864
47.149.2.13 get_pu()	864
47.150 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference	864
47.150.1 Detailed Description	866
47.150.2 Constructor & Destructor Documentation	867
47.150.2.1 Matrix() [1/2]	867
47.150.2.2 Matrix() [2/2]	868
47.150.3 Member Function Documentation	868
47.150.3.1 col_scale()	868
47.150.3.2 create() [1/7]	868
47.150.3.3 create() [2/7]	869
47.150.3.4 create() [3/7]	870
47.150.3.5 create() [4/7]	870

47.150.3.6 create() [5/7]	871
47.150.3.7 create() [6/7]	872
47.150.3.8 create() [7/7]	873
47.150.3.9 get_local_matrix()	873
47.150.3.10 get_non_local_matrix()	874
47.150.3.11 operator=() [1/2]	874
47.150.3.12 operator=() [2/2]	874
47.150.3.13 read_distributed() [1/4]	875
47.150.3.14 read_distributed() [2/4]	875
47.150.3.15 read_distributed() [3/4]	876
47.150.3.16 read_distributed() [4/4]	876
47.150.3.17 row_scale()	877
47.151 gko::matrix_assembly_data< ValueType, IndexType > Class Template Reference	877
47.151.1 Detailed Description	878
47.151.2 Member Function Documentation	878
47.151.2.1 add_value()	878
47.151.2.2 contains()	878
47.151.2.3 get_num_stored_elements()	879
47.151.2.4 get_ordered_data()	879
47.151.2.5 get_size()	879
47.151.2.6 get_value()	880
47.151.2.7 set_value()	880
47.152 gko::matrix_data< ValueType, IndexType > Struct Template Reference	880
47.152.1 Detailed Description	882
47.152.2 Constructor & Destructor Documentation	882
47.152.2.1 matrix_data() [1/6]	882
47.152.2.2 matrix_data() [2/6]	883
47.152.2.3 matrix_data() [3/6]	883
47.152.2.4 matrix_data() [4/6]	884
47.152.2.5 matrix_data() [5/6]	884
47.152.2.6 matrix_data() [6/6]	884
47.152.3 Member Function Documentation	885
47.152.3.1 cond() [1/2]	885
47.152.3.2 cond() [2/2]	886
47.152.3.3 diag() [1/5]	886
47.152.3.4 diag() [2/5]	887
47.152.3.5 diag() [3/5]	887
47.152.3.6 diag() [4/5]	888
47.152.3.7 diag() [5/5]	888
47.152.3.8 sum_duplicates()	889
47.152.4 Member Data Documentation	889
47.152.4.1 nonzeros	889

47.153 gko::matrix_data_entry< ValueType, IndexType > Struct Template Reference	889
47.153.1 Detailed Description	890
47.154 gko::experimental::reorder::Mc64< ValueType, IndexType > Class Template Reference	890
47.154.1 Detailed Description	890
47.154.2 Member Function Documentation	891
47.154.2.1 generate()	891
47.154.2.2 get_parameters()	891
47.155 gko::matrix::Csr< ValueType, IndexType >::merge_path Class Reference	891
47.155.1 Detailed Description	892
47.155.2 Member Function Documentation	892
47.155.2.1 clac_size()	892
47.155.2.2 copy()	892
47.155.2.3 process()	893
47.156 gko::MetisError Class Reference	893
47.156.1 Detailed Description	893
47.156.2 Constructor & Destructor Documentation	893
47.156.2.1 MetisError()	894
47.157 gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit Class Reference	894
47.157.1 Detailed Description	894
47.157.2 Member Function Documentation	895
47.157.2.1 compute_ell_num_stored_elements_per_row()	895
47.157.2.2 get_percentage()	895
47.158 gko::MpiError Class Reference	895
47.158.1 Detailed Description	896
47.158.2 Constructor & Destructor Documentation	896
47.158.2.1 MpiError()	896
47.159 gko::solver::Multigrid Class Reference	896
47.159.1 Detailed Description	897
47.159.2 Member Function Documentation	897
47.159.2.1 apply_uses_initial_guess()	898
47.159.2.2 get_coarsest_solver()	898
47.159.2.3 get_cycle()	898
47.159.2.4 get_mg_level_list()	898
47.159.2.5 get_mid_smoother_list()	899
47.159.2.6 get_post_smoother_list()	899
47.159.2.7 get_pre_smoother_list()	899
47.159.2.8 parse()	899
47.159.2.9 set_cycle()	900
47.160 gko::multigrid::MultigridLevel Class Reference	900
47.160.1 Detailed Description	901
47.160.2 Member Function Documentation	901
47.160.2.1 get_coarse_op()	901

47.160.2.2 <code>get_fine_op()</code>	901
47.160.2.3 <code>get_prolong_op()</code>	901
47.160.2.4 <code>get_restrict_op()</code>	902
47.161 <code>gko::batch::MultiVector< ValueType ></code> Class Template Reference	902
47.161.1 Detailed Description	903
47.161.2 Member Function Documentation	904
47.161.2.1 <code>add_scaled()</code>	904
47.161.2.2 <code>at()</code> [1/4]	904
47.161.2.3 <code>at()</code> [2/4]	905
47.161.2.4 <code>at()</code> [3/4]	905
47.161.2.5 <code>at()</code> [4/4]	906
47.161.2.6 <code>compute_conj_dot()</code>	906
47.161.2.7 <code>compute_dot()</code>	907
47.161.2.8 <code>compute_norm2()</code>	907
47.161.2.9 <code>create()</code> [1/3]	907
47.161.2.10 <code>create()</code> [2/3]	908
47.161.2.11 <code>create()</code> [3/3]	908
47.161.2.12 <code>create_const()</code>	908
47.161.2.13 <code>create_const_view_for_item()</code>	909
47.161.2.14 <code>create_view_for_item()</code>	909
47.161.2.15 <code>create_with_config_of()</code>	910
47.161.2.16 <code>fill()</code>	910
47.161.2.17 <code>get_common_size()</code>	910
47.161.2.18 <code>get_const_values()</code>	911
47.161.2.19 <code>get_const_values_for_item()</code>	911
47.161.2.20 <code>get_cumulative_offset()</code>	912
47.161.2.21 <code>get_num_batch_items()</code>	912
47.161.2.22 <code>get_num_stored_elements()</code>	913
47.161.2.23 <code>get_size()</code>	913
47.161.2.24 <code>get_values()</code>	913
47.161.2.25 <code>get_values_for_item()</code>	913
47.161.2.26 <code>scale()</code>	914
47.162 <code>gko::log::ProfilerHook::NestedSummaryWriter</code> Class Reference	914
47.162.1 Detailed Description	915
47.162.2 Member Function Documentation	915
47.162.2.1 <code>write_nested()</code>	915
47.163 <code>gko::NotCompiled</code> Class Reference	915
47.163.1 Detailed Description	915
47.163.2 Constructor & Destructor Documentation	916
47.163.2.1 <code>NotCompiled()</code>	916
47.164 <code>gko::NotImplemented</code> Class Reference	916
47.164.1 Detailed Description	916

47.164.2 Constructor & Destructor Documentation	916
47.164.2.1 NotImplemented()	917
47.165 gko::NotSupported Class Reference	917
47.165.1 Detailed Description	917
47.165.2 Constructor & Destructor Documentation	917
47.165.2.1 NotSupported()	917
47.166 gko::null_deleter< T > Class Template Reference	918
47.166.1 Detailed Description	918
47.166.2 Member Function Documentation	918
47.166.2.1 operator>()	918
47.167 gko::nvidia_device Class Reference	919
47.167.1 Detailed Description	919
47.168 gko::OmpExecutor Class Reference	919
47.168.1 Detailed Description	920
47.168.2 Member Function Documentation	920
47.168.2.1 get_description()	920
47.168.2.2 get_master() [1/2]	920
47.168.2.3 get_master() [2/2]	921
47.168.2.4 run() [1/3]	921
47.168.2.5 run() [2/3]	921
47.168.2.6 run() [3/3]	922
47.169 gko::Operation Class Reference	922
47.169.1 Detailed Description	923
47.169.2 Member Function Documentation	924
47.169.2.1 get_name()	924
47.170 gko::log::operation_data Struct Reference	924
47.170.1 Detailed Description	924
47.171 gko::OutOfBoundsError Class Reference	925
47.171.1 Detailed Description	925
47.171.2 Constructor & Destructor Documentation	925
47.171.2.1 OutOfBoundsError()	925
47.172 gko::OverflowError Class Reference	925
47.172.1 Detailed Description	926
47.172.2 Constructor & Destructor Documentation	926
47.172.2.1 OverflowError()	926
47.173 gko::log::Papi< ValueType > Class Template Reference	926
47.173.1 Detailed Description	927
47.173.2 Member Function Documentation	927
47.173.2.1 create() [1/2]	927
47.173.2.2 create() [2/2]	928
47.173.2.3 get_handle()	928
47.173.2.4 get_handle_name()	929

47.174	gko::factorization::Parlc< ValueType, IndexType > Class Template Reference	929
47.174.1	Detailed Description	929
47.174.2	Member Function Documentation	930
47.174.2.1	parse()	930
47.175	gko::factorization::Parlct< ValueType, IndexType > Class Template Reference	930
47.175.1	Detailed Description	931
47.175.2	Member Function Documentation	931
47.175.2.1	parse()	932
47.176	gko::factorization::Parllu< ValueType, IndexType > Class Template Reference	932
47.176.1	Detailed Description	933
47.176.2	Member Function Documentation	933
47.176.2.1	parse()	933
47.177	gko::factorization::Parllut< ValueType, IndexType > Class Template Reference	934
47.177.1	Detailed Description	934
47.177.2	Member Function Documentation	935
47.177.2.1	parse()	935
47.178	gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType > Class Template Reference	935
47.178.1	Detailed Description	936
47.178.2	Member Function Documentation	937
47.178.2.1	build_from_contiguous()	937
47.178.2.2	build_from_global_size_uniform()	938
47.178.2.3	build_from_mapping()	938
47.178.2.4	get_num_empty_parts()	939
47.178.2.5	get_num_parts()	939
47.178.2.6	get_num_ranges()	939
47.178.2.7	get_part_ids()	940
47.178.2.8	get_part_size()	940
47.178.2.9	get_part_sizes()	940
47.178.2.10	get_range_bounds()	941
47.178.2.11	get_range_starting_indices()	941
47.178.2.12	get_size()	942
47.178.2.13	has_connected_parts()	942
47.178.2.14	has_ordered_parts()	942
47.179	gko::log::PerformanceHint Class Reference	943
47.179.1	Detailed Description	943
47.179.2	Member Function Documentation	943
47.179.2.1	create()	943
47.180	gko::Permutable< IndexType > Class Template Reference	944
47.180.1	Detailed Description	945
47.180.1.1	Example: Permuting a Csr matrix:	945
47.180.2	Member Function Documentation	945

47.180.2.1 column_permute()	945
47.180.2.2 inverse_column_permute()	946
47.180.2.3 inverse_permute()	946
47.180.2.4 inverse_row_permute()	947
47.180.2.5 permute()	947
47.180.2.6 row_permute()	948
47.181 gko::matrix::Permutation< IndexType > Class Template Reference	948
47.181.1 Detailed Description	949
47.181.2 Member Function Documentation	949
47.181.2.1 compose()	949
47.181.2.2 compute_inverse()	950
47.181.2.3 create() [1/2]	950
47.181.2.4 create() [2/2]	951
47.181.2.5 create_const() [1/2]	951
47.181.2.6 create_const() [2/2]	951
47.181.2.7 get_const_permutation()	952
47.181.2.8 get_permutation()	952
47.181.2.9 get_permutation_size()	953
47.181.2.10 write()	953
47.182 gko::matrix::Csr< ValueType, IndexType >::permuring_reuse_info Struct Reference	953
47.182.1 Detailed Description	954
47.182.2 Member Function Documentation	954
47.182.2.1 update_values()	954
47.183 gko::Perturbation< ValueType > Class Template Reference	954
47.183.1 Detailed Description	955
47.183.2 Member Function Documentation	955
47.183.2.1 create() [1/3]	955
47.183.2.2 create() [2/3]	956
47.183.2.3 create() [3/3]	956
47.183.2.4 get_basis()	957
47.183.2.5 get_projector()	957
47.183.2.6 get_scalar()	957
47.184 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference	958
47.184.1 Detailed Description	958
47.184.2 Member Function Documentation	958
47.184.2.1 get_agg()	959
47.184.2.2 get_const_agg()	959
47.184.2.3 get_system_matrix()	959
47.184.2.4 parse()	960
47.185 gko::config::pnode Class Reference	960
47.185.1 Detailed Description	961
47.185.2 Constructor & Destructor Documentation	961

47.185.2.1 pnode() [1/7]	962
47.185.2.2 pnode() [2/7]	962
47.185.2.3 pnode() [3/7]	962
47.185.2.4 pnode() [4/7]	963
47.185.2.5 pnode() [5/7]	963
47.185.2.6 pnode() [6/7]	963
47.185.2.7 pnode() [7/7]	963
47.185.3 Member Function Documentation	964
47.185.3.1 get() [1/2]	964
47.185.3.2 get() [2/2]	964
47.185.3.3 get_array()	965
47.185.3.4 get_boolean()	965
47.185.3.5 get_integer()	965
47.185.3.6 get_map()	965
47.185.3.7 get_real()	966
47.185.3.8 get_string()	966
47.185.3.9 get_tag()	966
47.185.3.10 operator bool()	966
47.186 gko::log::polymorphic_object_data Struct Reference	967
47.186.1 Detailed Description	967
47.187 gko::PolymorphicObject Class Reference	967
47.187.1 Detailed Description	968
47.187.2 Member Function Documentation	968
47.187.2.1 clear()	968
47.187.2.2 clone() [1/2]	969
47.187.2.3 clone() [2/2]	969
47.187.2.4 copy_from() [1/4]	969
47.187.2.5 copy_from() [2/4]	970
47.187.2.6 copy_from() [3/4]	970
47.187.2.7 copy_from() [4/4]	971
47.187.2.8 create_default() [1/2]	972
47.187.2.9 create_default() [2/2]	972
47.187.2.10 get_executor()	972
47.187.2.11 move_from()	973
47.188 gko::precision_reduction Class Reference	973
47.188.1 Detailed Description	974
47.188.2 Constructor & Destructor Documentation	974
47.188.2.1 precision_reduction() [1/2]	974
47.188.2.2 precision_reduction() [2/2]	974
47.188.3 Member Function Documentation	975
47.188.3.1 autodetect()	975
47.188.3.2 common()	975

47.188.3.3 <code>get_nonpreserving()</code>	976
47.188.3.4 <code>get_preserving()</code>	976
47.188.3.5 <code>operator storage_type()</code>	976
47.189 <code>gko::Preconditionable</code> Class Reference	976
47.189.1 Detailed Description	977
47.189.2 Member Function Documentation	977
47.189.2.1 <code>get_preconditioner()</code>	977
47.189.2.2 <code>set_preconditioner()</code>	977
47.190 <code>gko::log::ProfilerHook</code> Class Reference	978
47.190.1 Detailed Description	979
47.190.2 Member Function Documentation	979
47.190.2.1 <code>create_nested_summary()</code>	979
47.190.2.2 <code>create_nvtx()</code>	979
47.190.2.3 <code>create_summary()</code>	980
47.190.2.4 <code>create_tau()</code>	980
47.190.2.5 <code>set_object_name()</code>	980
47.190.2.6 <code>user_range()</code>	981
47.191 <code>gko::log::profiling_scope_guard</code> Class Reference	981
47.191.1 Detailed Description	982
47.191.2 Constructor & Destructor Documentation	982
47.191.2.1 <code>profiling_scope_guard()</code>	982
47.192 <code>gko::ptr_param< T ></code> Class Template Reference	982
47.192.1 Detailed Description	983
47.192.2 Member Function Documentation	983
47.192.2.1 <code>get()</code>	983
47.192.2.2 <code>operator bool()</code>	983
47.192.2.3 <code>operator*()</code>	984
47.192.2.4 <code>operator->()</code>	984
47.193 <code>gko::range< Accessor ></code> Class Template Reference	984
47.193.1 Detailed Description	985
47.193.1.1 Range operations	985
47.193.1.2 Compound operations	986
47.193.1.3 Caveats	986
47.193.1.4 Examples	986
47.193.2 Constructor & Destructor Documentation	986
47.193.2.1 <code>range()</code>	986
47.193.3 Member Function Documentation	987
47.193.3.1 <code>get_accessor()</code>	987
47.193.3.2 <code>length()</code>	987
47.193.3.3 <code>operator>()()</code>	988
47.193.3.4 <code>operator->()</code>	988
47.193.3.5 <code>operator=()</code> [1/2]	988

47.193.3.6 operator=() [2/2]	989
47.194 gko::syn::range< Start, End, Step > Struct Template Reference	989
47.194.1 Detailed Description	989
47.195 gko::experimental::reorder::Rcm< IndexType > Class Template Reference	990
47.195.1 Detailed Description	990
47.195.2 Member Function Documentation	991
47.195.2.1 generate()	991
47.195.2.2 get_parameters()	991
47.196 gko::reorder::Rcm< ValueType, IndexType > Class Template Reference	991
47.196.1 Detailed Description	992
47.196.2 Member Function Documentation	992
47.196.2.1 get_inverse_permutation()	992
47.196.2.2 get_permutation()	993
47.197 gko::ReadableFromMatrixData< ValueType, IndexType > Class Template Reference	993
47.197.1 Detailed Description	993
47.197.2 Member Function Documentation	993
47.197.2.1 read() [1/4]	993
47.197.2.2 read() [2/4]	994
47.197.2.3 read() [3/4]	994
47.197.2.4 read() [4/4]	994
47.198 gko::log::Record Class Reference	995
47.198.1 Detailed Description	995
47.198.2 Member Function Documentation	996
47.198.2.1 create() [1/2]	996
47.198.2.2 create() [2/2]	996
47.198.2.3 get() [1/2]	997
47.198.2.4 get() [2/2]	997
47.199 gko::ReferenceExecutor Class Reference	997
47.199.1 Detailed Description	998
47.199.2 Member Function Documentation	998
47.199.2.1 get_description()	998
47.199.2.2 run() [1/4]	998
47.199.2.3 run() [2/4]	999
47.199.2.4 run() [3/4]	999
47.199.2.5 run() [4/4]	999
47.200 gko::config::registry Class Reference	1000
47.200.1 Detailed Description	1000
47.200.2 Constructor & Destructor Documentation	1000
47.200.2.1 registry() [1/2]	1000
47.200.2.2 registry() [2/2]	1001
47.200.3 Member Function Documentation	1001
47.200.3.1 emplace()	1001

47.201	gko::stop::RelativeResidualNorm< ValueType > Class Template Reference	1002
47.201.1	Detailed Description	1002
47.202	gko::reorder::ReorderingBase< IndexType > Class Template Reference	1002
47.202.1	Detailed Description	1003
47.203	gko::reorder::ReorderingBaseArgs Struct Reference	1003
47.203.1	Detailed Description	1003
47.204	gko::experimental::mpi::request Class Reference	1003
47.204.1	Detailed Description	1004
47.204.2	Constructor & Destructor Documentation	1004
47.204.2.1	request()	1004
47.204.3	Member Function Documentation	1004
47.204.3.1	get()	1004
47.204.3.2	wait()	1004
47.205	gko::stop::ResidualNorm< ValueType > Class Template Reference	1005
47.205.1	Detailed Description	1005
47.206	gko::stop::ResidualNormBase< ValueType > Class Template Reference	1006
47.206.1	Detailed Description	1006
47.207	gko::stop::ResidualNormReduction< ValueType > Class Template Reference	1006
47.207.1	Detailed Description	1006
47.208	gko::accessor::row_major< ValueType, Dimensionality > Class Template Reference	1007
47.208.1	Detailed Description	1007
47.208.2	Member Function Documentation	1008
47.208.2.1	copy_from()	1008
47.208.2.2	length()	1008
47.208.2.3	operator>() [1/2]	1009
47.208.2.4	operator>() [2/2]	1009
47.209	gko::matrix::RowGatherer< IndexType > Class Template Reference	1010
47.209.1	Detailed Description	1010
47.209.2	Member Function Documentation	1011
47.209.2.1	create() [1/2]	1011
47.209.2.2	create() [2/2]	1011
47.209.2.3	create_const()	1012
47.209.2.4	get_const_row_idxs()	1012
47.209.2.5	get_row_idxs()	1013
47.210	gko::ScaledIdentityAddable Class Reference	1013
47.210.1	Detailed Description	1013
47.210.2	Member Function Documentation	1013
47.210.2.1	add_scaled_identity()	1013
47.211	gko::matrix::ScaledPermutation< ValueType, IndexType > Class Template Reference	1014
47.211.1	Detailed Description	1015
47.211.2	Member Function Documentation	1015
47.211.2.1	compose()	1015

47.211.2.2 compute_inverse()	1015
47.211.2.3 create() [1/3]	1016
47.211.2.4 create() [2/3]	1016
47.211.2.5 create() [3/3]	1017
47.211.2.6 create_const()	1017
47.211.2.7 get_const_permutation()	1017
47.211.2.8 get_const_scaling_factors()	1018
47.211.2.9 get_permutation()	1018
47.211.2.10 get_scaling_factors()	1019
47.211.2.11 write()	1019
47.212 gko::experimental::reorder::ScaledReordered< ValueType, IndexType > Class Template Reference	1019
47.212.1 Detailed Description	1019
47.213 gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference	1020
47.213.1 Detailed Description	1021
47.213.2 Member Function Documentation	1021
47.213.2.1 apply_uses_initial_guess()	1021
47.213.2.2 parse()	1022
47.214 gko::scoped_device_id_guard Class Reference	1022
47.214.1 Detailed Description	1023
47.214.2 Constructor & Destructor Documentation	1023
47.214.2.1 scoped_device_id_guard() [1/5]	1023
47.214.2.2 scoped_device_id_guard() [2/5]	1023
47.214.2.3 scoped_device_id_guard() [3/5]	1024
47.214.2.4 scoped_device_id_guard() [4/5]	1024
47.214.2.5 scoped_device_id_guard() [5/5]	1024
47.215 gko::segmented_array< T > Struct Template Reference	1025
47.215.1 Detailed Description	1026
47.215.2 Constructor & Destructor Documentation	1026
47.215.2.1 segmented_array() [1/3]	1026
47.215.2.2 segmented_array() [2/3]	1026
47.215.2.3 segmented_array() [3/3]	1027
47.215.3 Member Function Documentation	1027
47.215.3.1 create_from_offsets() [1/2]	1027
47.215.3.2 create_from_offsets() [2/2]	1027
47.215.3.3 create_from_sizes() [1/2]	1028
47.215.3.4 create_from_sizes() [2/2]	1028
47.215.3.5 get_const_flat_data()	1028
47.215.3.6 get_executor()	1028
47.215.3.7 get_flat_data()	1029
47.215.3.8 get_offsets()	1029
47.215.3.9 get_segment_count()	1029

47.215.3.10 get_size()	1029
47.216 gko::matrix::Selp< ValueType, IndexType > Class Template Reference	1030
47.216.1 Detailed Description	1031
47.216.2 Constructor & Destructor Documentation	1031
47.216.2.1 Selp() [1/2]	1032
47.216.2.2 Selp() [2/2]	1032
47.216.3 Member Function Documentation	1032
47.216.3.1 col_at() [1/2]	1032
47.216.3.2 col_at() [2/2]	1033
47.216.3.3 compute_absolute()	1033
47.216.3.4 create() [1/2]	1033
47.216.3.5 create() [2/2]	1034
47.216.3.6 extract_diagonal()	1034
47.216.3.7 get_col_idxs()	1035
47.216.3.8 get_const_col_idxs()	1035
47.216.3.9 get_const_slice_lengths()	1035
47.216.3.10 get_const_slice_sets()	1036
47.216.3.11 get_const_values()	1036
47.216.3.12 get_num_stored_elements()	1036
47.216.3.13 get_slice_lengths()	1037
47.216.3.14 get_slice_sets()	1037
47.216.3.15 get_slice_size()	1037
47.216.3.16 get_stride_factor()	1037
47.216.3.17 get_total_cols()	1038
47.216.3.18 get_values()	1038
47.216.3.19 operator=() [1/2]	1038
47.216.3.20 operator=() [2/2]	1038
47.216.3.21 read() [1/3]	1038
47.216.3.22 read() [2/3]	1039
47.216.3.23 read() [3/3]	1039
47.216.3.24 val_at() [1/2]	1039
47.216.3.25 val_at() [2/2]	1040
47.216.3.26 write()	1040
47.217 gko::log::SolverProgress Class Reference	1041
47.217.1 Detailed Description	1041
47.217.2 Member Function Documentation	1041
47.217.2.1 create_scalar_csv_writer()	1041
47.217.2.2 create_scalar_table_writer()	1042
47.217.2.3 create_vector_storage()	1042
47.218 gko::preconditioner::Sor< ValueType, IndexType > Class Template Reference	1043
47.218.1 Detailed Description	1043
47.218.2 Member Function Documentation	1044

47.218.2.1 generate()	1044
47.218.2.2 get_parameters() [1/2]	1044
47.218.2.3 get_parameters() [2/2]	1044
47.219 gko::span Struct Reference	1045
47.219.1 Detailed Description	1045
47.219.2 Constructor & Destructor Documentation	1045
47.219.2.1 span() [1/2]	1045
47.219.2.2 span() [2/2]	1046
47.219.3 Member Function Documentation	1046
47.219.3.1 is_valid()	1046
47.219.3.2 length()	1046
47.220 gko::matrix::Csr< ValueType, IndexType >::sparselib Class Reference	1047
47.220.1 Detailed Description	1047
47.220.2 Member Function Documentation	1047
47.220.2.1 clac_size()	1047
47.220.2.2 copy()	1048
47.220.2.3 process()	1048
47.221 gko::matrix::SparsityCsr< ValueType, IndexType > Class Template Reference	1048
47.221.1 Detailed Description	1050
47.221.2 Constructor & Destructor Documentation	1050
47.221.2.1 SparsityCsr() [1/2]	1050
47.221.2.2 SparsityCsr() [2/2]	1050
47.221.3 Member Function Documentation	1051
47.221.3.1 conj_transpose()	1051
47.221.3.2 create() [1/4]	1051
47.221.3.3 create() [2/4]	1052
47.221.3.4 create() [3/4]	1052
47.221.3.5 create() [4/4]	1053
47.221.3.6 create_const()	1053
47.221.3.7 get_col_idxxs()	1053
47.221.3.8 get_const_col_idxxs()	1054
47.221.3.9 get_const_row_ptrs()	1054
47.221.3.10 get_const_value()	1055
47.221.3.11 get_num_nonzeros()	1055
47.221.3.12 get_row_ptrs()	1055
47.221.3.13 get_value()	1056
47.221.3.14 operator=() [1/2]	1056
47.221.3.15 operator=() [2/2]	1056
47.221.3.16 read() [1/3]	1056
47.221.3.17 read() [2/3]	1057
47.221.3.18 read() [3/3]	1057
47.221.3.19 to_adjacency_matrix()	1057

47.221.3.20 transpose()	1058
47.221.3.21 write()	1058
47.222 gko::experimental::mpi::status Struct Reference	1058
47.222.1 Detailed Description	1059
47.222.2 Constructor & Destructor Documentation	1059
47.222.2.1 status()	1059
47.222.3 Member Function Documentation	1059
47.222.3.1 get()	1059
47.222.3.2 get_count()	1059
47.223 gko::stopping_status Class Reference	1060
47.223.1 Detailed Description	1060
47.223.2 Member Function Documentation	1061
47.223.2.1 converge()	1061
47.223.2.2 get_id()	1061
47.223.2.3 has_converged()	1061
47.223.2.4 has_stopped()	1062
47.223.2.5 is_finalized()	1062
47.223.2.6 stop()	1062
47.223.3 Friends And Related Function Documentation	1062
47.223.3.1 operator"!="	1063
47.223.3.2 operator=="	1063
47.224 gko::matrix::Csr< ValueType, IndexType >::strategy_type Class Reference	1063
47.224.1 Detailed Description	1064
47.224.2 Constructor & Destructor Documentation	1064
47.224.2.1 strategy_type()	1064
47.224.3 Member Function Documentation	1064
47.224.3.1 clac_size()	1064
47.224.3.2 copy()	1065
47.224.3.3 get_name()	1065
47.224.3.4 process()	1065
47.225 gko::matrix::Hybrid< ValueType, IndexType >::strategy_type Class Reference	1066
47.225.1 Detailed Description	1066
47.225.2 Member Function Documentation	1067
47.225.2.1 compute_ell_num_stored_elements_per_row()	1067
47.225.2.2 compute_hybrid_config()	1067
47.225.2.3 get_coo_nnz()	1068
47.225.2.4 get_ell_num_stored_elements_per_row()	1068
47.226 gko::log::Stream< ValueType > Class Template Reference	1068
47.226.1 Detailed Description	1068
47.226.2 Member Function Documentation	1069
47.226.2.1 create() [1/2]	1069
47.226.2.2 create() [2/2]	1069

47.227	gko::StreamError Class Reference	1070
47.227.1	Detailed Description	1070
47.227.2	Constructor & Destructor Documentation	1070
47.227.2.1	StreamError()	1070
47.228	gko::log::ProfilerHook::SummaryWriter Class Reference	1071
47.228.1	Detailed Description	1071
47.228.2	Member Function Documentation	1071
47.228.2.1	write()	1071
47.229	gko::log::ProfilerHook::TableSummaryWriter Class Reference	1072
47.229.1	Detailed Description	1072
47.229.2	Constructor & Destructor Documentation	1072
47.229.2.1	TableSummaryWriter()	1072
47.229.3	Member Function Documentation	1073
47.229.3.1	write()	1073
47.229.3.2	write_nested()	1073
47.230	gko::stop::Time Class Reference	1073
47.230.1	Detailed Description	1074
47.231	gko::time_point Class Reference	1074
47.231.1	Detailed Description	1074
47.232	gko::Timer Class Reference	1074
47.232.1	Detailed Description	1075
47.232.2	Member Function Documentation	1075
47.232.2.1	create_for_executor()	1075
47.232.2.2	create_time_point()	1075
47.232.2.3	difference()	1075
47.232.2.4	difference_async()	1076
47.233	gko::Transposable Class Reference	1076
47.233.1	Detailed Description	1077
47.233.1.1	Example: Transposing a Csr matrix:	1077
47.233.2	Member Function Documentation	1077
47.233.2.1	conj_transpose()	1077
47.233.2.2	transpose()	1078
47.234	gko::config::type_descriptor Class Reference	1078
47.234.1	Detailed Description	1079
47.234.2	Constructor & Destructor Documentation	1079
47.234.2.1	type_descriptor()	1079
47.235	gko::experimental::mpi::type_impl< T > Struct Template Reference	1079
47.235.1	Detailed Description	1080
47.236	gko::syn::type_list< Types > Struct Template Reference	1080
47.236.1	Detailed Description	1080
47.237	gko::UnsupportedMatrixProperty Class Reference	1080
47.237.1	Detailed Description	1081

47.237.2 Constructor & Destructor Documentation	1081
47.237.2.1 UnsupportedMatrixProperty()	1081
47.238 gko::stop::Criterion::Updater Class Reference	1081
47.238.1 Detailed Description	1082
47.238.2 Member Function Documentation	1082
47.238.2.1 check()	1082
47.239 gko::solver::UpperTrs< ValueType, IndexType > Class Template Reference	1082
47.239.1 Detailed Description	1083
47.239.2 Constructor & Destructor Documentation	1083
47.239.2.1 UpperTrs() [1/2]	1083
47.239.2.2 UpperTrs() [2/2]	1083
47.239.3 Member Function Documentation	1084
47.239.3.1 conj_transpose()	1084
47.239.3.2 operator=() [1/2]	1084
47.239.3.3 operator=() [2/2]	1084
47.239.3.4 parse()	1084
47.239.3.5 transpose()	1085
47.240 gko::UseComposition< ValueType > Class Template Reference	1085
47.240.1 Detailed Description	1085
47.240.2 Member Function Documentation	1086
47.240.2.1 get_composition()	1086
47.240.2.2 get_operator_at()	1086
47.241 gko::syn::value_list< T, Values > Struct Template Reference	1087
47.241.1 Detailed Description	1087
47.242 gko::ValueMismatch Class Reference	1087
47.242.1 Detailed Description	1087
47.242.2 Constructor & Destructor Documentation	1088
47.242.2.1 ValueMismatch()	1088
47.243 gko::experimental::distributed::Vector< ValueType > Class Template Reference	1088
47.243.1 Detailed Description	1090
47.243.2 Member Function Documentation	1091
47.243.2.1 add_scaled()	1091
47.243.2.2 at_local() [1/4]	1091
47.243.2.3 at_local() [2/4]	1092
47.243.2.4 at_local() [3/4]	1092
47.243.2.5 at_local() [4/4]	1092
47.243.2.6 compute_absolute()	1093
47.243.2.7 compute_conj_dot() [1/2]	1093
47.243.2.8 compute_conj_dot() [2/2]	1093
47.243.2.9 compute_dot() [1/2]	1094
47.243.2.10 compute_dot() [2/2]	1094
47.243.2.11 compute_mean() [1/2]	1094

47.243.2.12	compute_mean() [2/2]	1095
47.243.2.13	compute_norm1() [1/2]	1095
47.243.2.14	compute_norm1() [2/2]	1095
47.243.2.15	compute_norm2() [1/2]	1096
47.243.2.16	compute_norm2() [2/2]	1096
47.243.2.17	compute_squared_norm2() [1/2]	1096
47.243.2.18	compute_squared_norm2() [2/2]	1097
47.243.2.19	create() [1/4]	1097
47.243.2.20	create() [2/4]	1098
47.243.2.21	create() [3/4]	1098
47.243.2.22	create() [4/4]	1099
47.243.2.23	create_const() [1/2]	1099
47.243.2.24	create_const() [2/2]	1100
47.243.2.25	create_real_view() [1/2]	1100
47.243.2.26	create_real_view() [2/2]	1100
47.243.2.27	create_with_config_of()	1101
47.243.2.28	create_with_type_of() [1/2]	1101
47.243.2.29	create_with_type_of() [2/2]	1101
47.243.2.30	fill()	1102
47.243.2.31	get_const_local_values()	1102
47.243.2.32	get_local_values()	1103
47.243.2.33	get_local_vector()	1103
47.243.2.34	inv_scale()	1103
47.243.2.35	make_complex() [1/2]	1103
47.243.2.36	make_complex() [2/2]	1104
47.243.2.37	read_distributed() [1/2]	1104
47.243.2.38	read_distributed() [2/2]	1104
47.243.2.39	scale()	1105
47.243.2.40	sub_scaled()	1105
47.244	gko::version Struct Reference	1105
47.244.1	Detailed Description	1106
47.244.2	Member Data Documentation	1106
47.244.2.1	tag	1106
47.245	gko::version_info Class Reference	1106
47.245.1	Detailed Description	1107
47.245.2	Member Function Documentation	1107
47.245.2.1	get()	1107
47.245.3	Member Data Documentation	1107
47.245.3.1	core_version	1107
47.245.3.2	cuda_version	1108
47.245.3.3	dpcpp_version	1108
47.245.3.4	hip_version	1108

47.245.3.5 omp_version	1108
47.245.3.6 reference_version	1108
47.246 gko::experimental::mpi::window< ValueType > Class Template Reference	1108
47.246.1 Detailed Description	1110
47.246.2 Constructor & Destructor Documentation	1110
47.246.2.1 window() [1/3]	1111
47.246.2.2 window() [2/3]	1111
47.246.2.3 window() [3/3]	1111
47.246.3 Member Function Documentation	1112
47.246.3.1 accumulate()	1112
47.246.3.2 fence()	1112
47.246.3.3 fetch_and_op()	1113
47.246.3.4 flush()	1113
47.246.3.5 flush_local()	1113
47.246.3.6 get()	1114
47.246.3.7 get_accumulate()	1114
47.246.3.8 get_window()	1115
47.246.3.9 lock()	1115
47.246.3.10 lock_all()	1116
47.246.3.11 operator=()	1116
47.246.3.12 put()	1116
47.246.3.13 r_accumulate()	1117
47.246.3.14 r_get()	1117
47.246.3.15 r_get_accumulate()	1118
47.246.3.16 r_put()	1119
47.246.3.17 unlock()	1119
47.247 gko::solver::workspace_traits< Solver > Struct Template Reference	1120
47.247.1 Detailed Description	1120
47.248 gko::WritableToMatrixData< ValueType, IndexType > Class Template Reference	1120
47.248.1 Detailed Description	1120
47.248.2 Member Function Documentation	1121
47.248.2.1 write()	1121

Chapter 1

Main Page

This is the main page for the Ginkgo library pdf documentation. The repository is hosted on [github](#). Documentation on aspects such as the build system, can be found at the [Installation Instructions](#) page. The [Example programs](#) can help you get started with using Ginkgo.

1.0.0.1 Modules

The Ginkgo library can be grouped into [modules](#) and these modules form the basic building blocks of Ginkgo. The modules can be summarized as follows:

- [Executors](#) : Where do you want your code to be executed ?
- [Linear Operators](#) : What kind of operation do you want Ginkgo to perform ?
 - [Solvers](#) : Solve a linear system for a given matrix.
 - [Preconditioners](#) : Precondition a system for a solve.
 - [SpMV employing different Matrix formats](#) : Perform a sparse matrix vector multiplication with a particular matrix format.
- [Logging](#) : Monitor your code execution.
- [Stopping criteria](#) : Manage your iteration stopping criteria.

Chapter 2

Installation Instructions

2.0.1 Building

Use the standard CMake build procedure:

```
mkdir build; cd build
cmake [OPTIONS] .. && cmake --build .
```

For Microsoft Visual Studio, use `cmake --build . --config <build_type>` to decide the build type. The possible options are `Debug`, `Release`, `RelWithDebInfo` and `MinSizeRel`.

Replace `[OPTIONS]` with desired `cmake` options for your build. Ginkgo adds the following additional switches to control what is being built:

- `-DGINKGO_DEVEL_TOOLS={ON, OFF}` sets up the build system for development (requires pre-commit, will also download the clang-format pre-commit hook), default is `OFF`. The default behavior installs a pre-commit hook, which disables git commits. If it is set to `ON`, a new pre-commit hook for formatting will be installed (enabling commits again). In both cases the hook may overwrite a user defined pre-commit hook when Ginkgo is used as a submodule.
- `-DGINKGO_MIXED_PRECISION={ON, OFF}` compiles true mixed-precision kernels instead of converting data on the fly, default is `OFF`. Enabling this flag increases the library size, but improves performance of mixed-precision kernels.
- `-DGINKGO_ENABLE_HALF={ON, OFF}` enable half precision support in Ginkgo, default is `ON`. It is `OFF` when the compiler is `MSVC` or `MSYS2` system. If compiling is done with the `CUDA` backend before `CUDA 12.2`, we only support half precision after compute capability 5.3. `CUDA 12.2+` compilers waive the compute capability limitation.
- `-DGINKGO_BUILD_TESTS={ON, OFF}` builds Ginkgo's tests (will download googletest), default is `ON`.
- `-DGINKGO_FAST_TESTS={ON, OFF}` reduces the input sizes for a few slow tests to speed them up, default is `OFF`.
- `-DGINKGO_BUILD_BENCHMARKS={ON, OFF}` builds Ginkgo's benchmarks (will download gflags and nlohmann-json), default is `ON`.
- `-DGINKGO_BUILD_EXAMPLES={ON, OFF}` builds Ginkgo's examples, default is `ON`.
- `-DGINKGO_BUILD_EXTLIB_EXAMPLE={ON, OFF}` builds the interfacing example with `deal.II`, default is `OFF`.
- `-DGINKGO_BUILD_REFERENCE={ON, OFF}` build reference implementations of the kernels, useful for testing, default is `ON`.

- `-DGINKGO_BUILD_OMP={ON, OFF}` builds optimized OpenMP versions of the kernels, default is ON if the selected C++ compiler supports OpenMP, OFF otherwise.
- `-DGINKGO_BUILD_CUDA={ON, OFF}` builds optimized cuda versions of the kernels (requires CUDA), default is ON if a CUDA compiler could be detected, OFF otherwise.
- `-DGINKGO_BUILD_DPCPP={ON, OFF}` is deprecated. Please use `GINKGO_BUILD_SYCL` instead.
- `-DGINKGO_BUILD_SYCL={ON, OFF}` builds optimized SYCL versions of the kernels (requires `CMAKE_CXX_COMPILER` to be set to the `dpcpp` or `icpx` compiler). The default is ON if `CMAKE_CXX_COMPILER` is a SYCL compiler, OFF otherwise. Due to some differences in IEEE 754 floating point numberhandling in the Intel SYCL compilers, Ginkgo tests may fail unless compiled with `-DCMAKE_CXX_FLAGS=-ffp-model=precise`
- `-DGINKGO_BUILD_HIP={ON, OFF}` builds optimized HIP versions of the kernels (requires HIP), default is ON if an installation of HIP could be detected, OFF otherwise.
- `-DCMAKE_HIP_ARCHITECTURES="gpuarch1;gpuarch2"` the AMDGPU targets to be passed to the compiler. If empty, compiler chooses based on the available GPUs.
- `-DGINKGO_BUILD_HWLOC={ON, OFF}` builds Ginkgo with HWLOC. Default is OFF.
- `-DGINKGO_BUILD_DOC={ON, OFF}` creates an HTML version of Ginkgo's documentation from inline comments in the code. The default is OFF.
- `-DGINKGO_DOC_GENERATE_EXAMPLES={ON, OFF}` generates the documentation of examples in Ginkgo. The default is ON.
- `-DGINKGO_DOC_GENERATE_PDF={ON, OFF}` generates a PDF version of Ginkgo's documentation from inline comments in the code. The default is OFF.
- `-DGINKGO_DOC_GENERATE_DEV={ON, OFF}` generates the developer version of Ginkgo's documentation. The default is OFF.
- `-DGINKGO_WITH_CLANG_TIDY={ON, OFF}` makes Ginkgo call `clang-tidy` to find programming issues. The path can be manually controlled with the CMake variable `-DGINKGO_CLANG_TIDY_PATH=<path>`. The default is OFF.
- `-DGINKGO_WITH_IWYU={ON, OFF}` makes Ginkgo call `iwyu` to find include issues. The path can be manually controlled with the CMake variable `-DGINKGO_IWYU_PATH=<path>`. The default is OFF.
- `-DGINKGO_CHECK_CIRCULAR_DEPS={ON, OFF}` enables compile-time checks for circular dependencies between different Ginkgo libraries and self-sufficient headers. Should only be used for development purposes. The default is OFF.
- `-DGINKGO_VERBOSE_LEVEL=integer` sets the verbosity of Ginkgo.
 - 0 disables all output in the main libraries,
 - 1 enables a few important messages related to unexpected behavior (default).
- `GINKGO_INSTALL_RPATH` allows setting any RPATH information when installing the Ginkgo libraries. If this is OFF, the behavior is the same as if all other RPATH flags are set to OFF as well. The default is ON.
- `GINKGO_INSTALL_RPATH_ORIGIN` adds `$ORIGIN` (Linux) or `@loader_path` (MacOS) to the installation RPATH. The default is ON.
- `GINKGO_INSTALL_RPATH_DEPENDENCIES` adds the dependencies to the installation RPATH. The default is OFF.
- `-DCMAKE_INSTALL_PREFIX=path` sets the installation path for `make install`. The default value is usually something like `/usr/local`.
- `-DCMAKE_BUILD_TYPE=type` specifies which configuration will be used for this build of Ginkgo. The default is `RELEASE`. Supported values are CMake's standard build types such as `DEBUG` and `RELEASE` and the Ginkgo specific `COVERAGE`, `ASAN` (AddressSanitizer), `LSAN` (LeakSanitizer), `TSAN` (ThreadSanitizer) and `UBSAN` (undefined behavior sanitizer) types.

- `-DBUILD_SHARED_LIBS={ON, OFF}` builds ginkgo as shared libraries (OFF) or as dynamic libraries (ON), default is ON.
- `-DGINKGO_JACOBI_FULL_OPTIMIZATIONS={ON, OFF}` use all the optimizations for the CUDA Jacobi algorithm. OFF by default. Setting this option to ON may lead to very slow compile time (>20 minutes) for the `jacobi_generate_kernels.cu` file and high memory usage.
- `-DCMAKE_CUDA_HOST_COMPILER=path` instructs the build system to explicitly set CUDA's host compiler to the path given as argument. By default, CUDA uses its toolchain's host compiler. Setting this option may help if you're experiencing linking errors due to ABI incompatibilities. This option is supported since CMake 3.8 but documented starting from 3.10.
- `-DGINKGO_CUDA_ARCHITECTURES=<list>` where <list> is a semicolon (;) separated list of architectures. Supported values are:
 - Auto
 - Kepler, Maxwell, Pascal, Volta, Turing, Ampere
 - CODE, CODE (COMPUTE), (COMPUTE)

Auto will automatically detect the present CUDA-enabled GPU architectures in the system. Kepler, Maxwell, Pascal, Volta and Ampere will add flags for all architectures of that particular NVIDIA GPU generation. COMPUTE and CODE are placeholders that should be replaced with compute and code numbers (e.g. for `compute_70` and `sm_70` COMPUTE and CODE should be replaced with 70. Default is Auto. For a more detailed explanation of this option see the [ARCHITECTURES specification list](#) section in the documentation of the `CudaArchitectureSelector` CMake module.

Additionally, the following CMake options have effect on the build process:

- `-DCMAKE_EXPORT_PACKAGE_REGISTRY={ON, OFF}` if set to ON the build directory will be stored in the current user's CMake package registry.

For example, to build everything (in debug mode), use:

```
cmake .. -BDebug -DCMAKE_BUILD_TYPE=Debug -DGINKGO_DEVEL_TOOLS=ON \
  -DGINKGO_BUILD_TESTS=ON -DGINKGO_BUILD_REFERENCE=ON -DGINKGO_BUILD_OMP=ON \
  -DGINKGO_BUILD_CUDA=ON -DGINKGO_BUILD_HIP=ON
cmake --build Debug
```

NOTE: Ginkgo is known to work with the Unix Makefiles, Ninja, MinGW Makefiles and Visual Studio 16 2019 based generators. Other CMake generators are untested.

2.0.2 Building Ginkgo in Windows

Depending on the configuration settings, some manual work might be required:

- Build Ginkgo with Debug mode: Some Debug build specific issues can appear depending on the machine and environment: When you encounter the error message `ld: error: export ordinal too large`, add the compilation flag `-O1` by adding `-DCMAKE_CXX_FLAGS=-O1` to the CMake invocation.
- Build Ginkgo in MinGW: If encountering the issue `cclplus.exe: out of memory allocating 65536 bytes`, please follow the workaround in [reference](#), or trying to compile ginkgo again might work.

2.0.3 Building Ginkgo with HIP support

Ginkgo provides a [HIP](#) backend. This allows to compile optimized versions of the kernels for either AMD or NVIDIA GPUs. The CMake configuration step will try to auto-detect the presence of HIP either at `/opt/rocm/hip` or at the path specified by `HIP_PATH` as a CMake parameter (`-DHIP_PATH=`) or environment variable (`export HIP_PATH=`), unless `-DGINKGO_BUILD_HIP=ON/OFF` is set explicitly.

2.0.3.1 Changing the paths to search for HIP and other packages

All HIP installation paths can be configured through the use of environment variables or CMake variables. This way of configuring the paths is currently imposed by the HIP tool suite. The variables are the following:

- CMake `-DROCM_PATH=` or environment `export ROCM_PATH=`: sets the ROCM installation path. The default value is `/opt/rocm/`.
- CMake `-DHIP_CLANG_PATH` or environment `export HIP_CLANG_PATH=`: sets the HIP compatible clang binary path. The default value is `${ROCM_PATH}/llvm/bin`.
- CMake `-DHIP_PATH=` or environment `export HIP_PATH=`: sets the HIP installation path. The default value is `${ROCM_PATH}/hip`.
- CMake `-DHIPBLAS_PATH=` or environment `export HIPBLAS_PATH=`: sets the hipBLAS installation path. The default value is `${ROCM_PATH}/hipblas`.
- CMake `-DHIPSPARSE_PATH=` or environment `export HIPSPARSE_PATH=`: sets the hipSPARSE installation path. The default value is `${ROCM_PATH}/hipsparse`.
- CMake `-DHIPFFT_PATH=` or environment `export HIPFFT_PATH=`: sets the hipFFT installation path. The default value is `${ROCM_PATH}/hipfft`.
- CMake `-DROCRAND_PATH=` or environment `export ROCRAND_PATH=`: sets the rocRAND installation path. The default value is `${ROCM_PATH}/rocrand`.
- CMake `-DHIPRAND_PATH=` or environment `export HIPRAND_PATH=`: sets the hipRAND installation path. The default value is `${ROCM_PATH}/hiprand`.
- environment `export CUDA_PATH=`: where `hipcc` can find CUDA if it is not in the default `/usr/local/cuda` path.

2.0.3.2 HIP platform detection of AMD and NVIDIA

Ginkgo relies on CMake to decide which compiler to use for HIP. To choose `nvcc` instead of the default ROCm `clang++`, set the corresponding environment variable:

```
export HIPCXX=nvcc
```

Note that this option is currently not being tested in our CI pipelines.

2.0.4 Third party libraries and packages

Ginkgo relies on third party packages in different cases. These third party packages can be turned off by disabling the relevant options.

- `GINKGO_BUILD_TESTS=ON`: Our tests are implemented with `Google Test`;
- `GINKGO_BUILD_BENCHMARKS=ON`: For argument management we use `gflags` and for JSON parsing we use `nlohmann-json`;
- `GINKGO_BUILD_HWLOC=ON`: `hwloc` to detect and control cores and devices.
- `GINKGO_BUILD_HWLOC=ON` and `GINKGO_BUILD_TESTS=ON`: `libnuma` is required when testing the functions provided through MachineTopology.
- `GINKGO_BUILD_EXAMPLES=ON`: `OpenCV` is required for some examples, they are disabled when OpenCV is not available.

- `GINKGO_BUILD_DOC=ON`: `doxygen` is required to build the documentation and additionally `graphviz` is required to build the class hierarchy graphs.
- `METIS` is required when using the `NestedDissection` reordering functionality. If `METIS` is not found, the functionality is disabled.
- `PAPI` ($\geq 7.1.0$) is required when using the `Papi` logger. If `PAPI` is not found, the functionality is disabled.

Ginkgo attempts to use pre-installed versions of these package if they match version requirements using `find_package`. Otherwise, the configuration step will download the files for each of the packages `GTest`, `gflags`, `nlohmann-json` and `hwloc` and build them internally.

Note that, if the external packages were not installed to the default location, the CMake option `-DCMAKE_PREFIX_PATH=<path-list>` needs to be set to the semicolon (;) separated list of install paths of these external packages. For more Information, see the [CMake documentation for CMAKE_PREFIX_PATH](#) for details.

For convenience, the options `GINKGO_INSTALL_RPATH[_.*]` can be used to bind the installed Ginkgo shared libraries to the path of its dependencies.

2.0.5 Installing Ginkgo

To install Ginkgo into the specified folder, execute the following command in the build folder

```
make install
```

If the installation prefix (see `CMAKE_INSTALL_PREFIX`) is not writable for your user, e.g. when installing Ginkgo system-wide, it might be necessary to prefix the call with `sudo`.

After the installation, CMake can find ginkgo with `find_package(Ginkgo)`. An example can be found in the `test_install`.

Chapter 3

Testing Instructions

3.0.1 Running the unit tests

You need to compile ginkgo with `-DGINKGO_BUILD_TESTS=ON` option to be able to run the tests.

3.0.1.1 Using make test

After configuring Ginkgo, use the following command inside the build folder to run all tests:

```
make test
```

The output should contain several lines of the form:

```
      Start   1:  path/to/test
1/13 Test   #1:  path/to/test .....          Passed    0.01 sec
```

To run only a specific test and see more details results (e.g. if a test failed) run the following from the build folder:

```
./path/to/test
```

where `path/to/test` is the path returned by `make test`.

3.0.1.2 Using make quick_test

After compiling Ginkgo, use the following command inside the build folder to run a small subset of tests that should execute quickly:

```
make quick_test
```

These tests do not use GPU features except for a few device property queries, so they may still fail if Ginkgo was compiled with GPU support, but no such GPU is available. The output is equivalent to `make test`.

3.0.1.3 Using CTest

The tests can also be ran through CTest from the command line, for example when in a configured build directory:

```
ctest -T start -T build -T test -T submit
```

Will start a new test campaign (usually in `Experimental` mode), build Ginkgo with the set configuration, run the tests and submit the results to our CDash dashboard.

Another option is to use Ginkgo's CTest script which is configured to build Ginkgo with default settings, runs the tests and submits the test to our CDash dashboard automatically.

To run the script, use the following command:

```
ctest -S cmake/CTestScript.cmake
```

The default settings are for our own CI system. Feel free to configure the script before launching it through variables or by directly changing its values. A documentation can be found in the script itself.

Chapter 4

Running the benchmarks

In addition to the unit tests designed to verify correctness, Ginkgo also includes an extensive benchmark suite for checking its performance on all Ginkgo supported systems. The purpose of Ginkgo's benchmarking suite is to allow easy and complete reproduction of Ginkgo's performance, and to facilitate performance debugging as well. Most results published in Ginkgo papers are generated thanks to this benchmarking suite and are accessible online under the [ginkgo-data repository](#). These results can also be used for performance comparison in order to ensure that you get similar performance as what is published on this repository.

To compile the benchmarks, the flag `-DGINKGO_BUILD_BENCHMARKS=ON` has to be set during the `cmake` step. In addition, the [ssget command-line utility](#) has to be installed on the system. The purpose of this file is to explain in detail the capacities of this benchmarking suite as well as how to properly setup everything.

There are two ways to benchmark Ginkgo. When compiling the benchmark suite, executables are generated for collecting matrix statistics, running sparse-matrix vector product, solvers (possibly distributed) benchmarks. Another way to run benchmarks is through the convenience script `run_all_benchmarks.sh`, but not all features are exposed through this tool!

Here is a short description of the content of this file:

1. Ginkgo setup and best practice guidelines
2. Installing and using the `ssget` tool to fetch the [SuiteSparse matrices](#).
3. Running benchmarks manually
4. Benchmarking with the script utility
5. How to publish the benchmark results online and use the [Ginkgo Performance Explorer \(GPE\)](#) for performance analysis (optional).
6. Using the benchmark suite for performance debugging thanks to the loggers.
7. Available benchmark customization options with the script utility.

4.0.1 1: Ginkgo setup and best practice guidelines

Before benchmarking Ginkgo, make sure that you follow the general guidelines in order to ensure best performance.

1. The code should be compiled in `Release mode`.

2. Make sure the machine has no competing jobs. On a Linux machine multiple commands can be used, `last` shows the currently opened sessions, `top` or `htop` allows to show the current machine load, and if considering using specific GPUs, `nvidia-smi` or `rocm-smi` can be used to check their load.
3. By default, Ginkgo's benchmarks will always do at least one warm-up run. For better accuracy, every benchmark is also averaged over 10 runs, except for the solver benchmark which are usually fairly long. These parameters can be tuned at the command line to either shorten benchmarking time or improve benchmarking accuracy.

In addition, the following specific options can be considered:

1. When specifically using the adaptive block jacobi preconditioner, enable the `GINKGO_JACOBI_FULL_OPTIMIZATIONS` CMake flag. Be careful that this will use much more memory and time for the compilation due to compiler performance issues with register optimizations, in particular.
2. The current benchmarking setup also allows to benchmark only the overhead by using as either (or for all) preconditioner/spmv/solver, the special `overhead` LinOp. If your purpose is to check Ginkgo's overhead, make sure to try this mode.

4.0.2 2: Using ssget to fetch the matrices

To benchmark `ginkgo`, matrices need to be provided as input in the `Matrix Market` format. A convenient way is to run benchmark with the `SuiteSparse matrix collection`. A helper tool, the `ssget command-line utility` can be used to facilitate downloading and extracting matrices from the `suitesparse` collection. When running the benchmarks with the helper script `run_all_benchmarks.sh` (or calling `make benchmark`), the `ssget` tool is required.

To install `ssget`, access the repository and copy the file `ssget` into a directory present in your `PATH` variable as per the tool's `README.md` instructions. The tool can be installed either in a global system path or a local directory such as `$HOME/.local/bin`. After installing the tool, it is important to review the `ssget` script and configure as needed the variable `ARCHIVE_LOCATION` on line 39. This is where the matrices will be stored into.

The Ginkgo benchmark can be set to run on only a portion of the `SuiteSparse` matrix collection as we will see in the following section. Please note that the entire collection requires roughly 100GB of disk storage in its compressed format, and roughly 25GB of additional disk space for intermediate data (such as uncompressing the archive). Additionally, the benchmark runs usually take a long time (SpMV benchmarks on the complete collection take roughly 24h using the K20 GPU), and will stress the system.

Before proceeding, it can be useful in order to save time to download the matrices as preparation. This can be done by using the `ssget -f -i i` command where `i` is the ID of the matrix to be downloaded. The following loop allows to download the full `SuiteSparse` matrix collection:

```
for i in $(seq 0 $(ssget -n)); do
    ssget -f -i ${i}
done
```

Note that `ssget` can also be used to query properties of the matrix and filter the matrices which are downloaded. For example, the following will download only positive definite matrices with less than 500M non zero elements and 10M columns. Please refer to the `ssget documentation` for more information.

```
for i in $(seq 0 $(ssget -n)); do
    posdef=$(ssget -p posdef -i ${i})
    cols=$(ssget -p cols -i ${i})
    nnz=$(ssget -p nonzeros -i ${i})
    if [ "$posdef" -eq 1 -a "$cols" -lt 10000000 -a "$nnz" -lt 500000000 ]; then
        ssget -f -i ${i}
    fi
done
```

4.0.3 3: Running benchmarks manually

When compiling Ginkgo with the flag `-DGINKGO_BUILD_BENCHMARKS=ON`, a suite of executables will be generated depending on the CMake configuration. These executables are the backbone of the benchmarking suite. Note that all of these executables describe the available options and the required input format when running them with the `--help` option. All executables have multiple variants depending on the precision, by default `double` precision is used for the type of values, but variants with `single` and `complex` (single and double) value types are also available. Here is a non exhaustive list of the available benchmarks:

- `blas/blas`: supports benchmarking many of Ginkgo's BLAS operations: dot products, axpy, copy, etc.
- `conversion/conversion`: conversion between matrix formats.
- `matrix_generator/matrix_generator`: mostly allows generating block diagonal matrices (to benchmark the block-jacobi preconditioner).
- `matrix_statistics/matrix_statistics`: computes size and other matrix statistics (such as variance, load imbalance, ...).
- `preconditioner/preconditioner`: benchmarks most Ginkgo preconditioner.
- `solver/solver`: benchmark most of Ginkgo's solvers in a non distributed setting.
- `sparse_blas/sparse_blas`: benchmarks Sparse BLAS operations, such as SpGEMM, SpGEAM, transpose.
- `spmv/spmv`: benchmarks Ginkgo's matrix formats (Sparse-Matrix Vector product).

Optionally when compiling with MPI support:

- `blas/distributed/multi_vector`: measures BLAS performance on (distributed) multi-vectors.
- `solver/distributed/solver`: distributed solver benchmarks.
- `spmv/distributed/spmv`: distributed matrix Sparse-Matrix Vector (SpMV) product benchmark.

All benchmarks require input data as in a JSON format. The json file has to consist of exactly one array, and within that array the test cases are defined. The exact syntax can change between executables, the `--help` option will explain the necessary JSON input format. For example for the `spmv` benchmark case, and many other benchmarks the following minimal input should be provided:

```
[
  {
    "filename": "path/to/your/matrix",
    "rhs": "path/to/your/rhs"
  },
  { ... }
]
```

The files have to be in matrix market format.

Some benchmarks require some extra fields. For example the solver benchmarks requires the field `"optimal"` ↵
: `{ "spmv": "matrix format (such as csr)" }`. This is automatically populated when running the `spmv` benchmark which finds the optimal (fastest) format among all requested formats.

After writing the necessary data in a JSON file, the benchmark can be called by passing in the input via stdin, i.e.
`./solver < input.json`

The output of our benchmarks is again JSON, and it is printed to stdout, while our status messages are printed to stderr. So, the output can be stored with
`./solver < input.json > output.json`

Note that in most cases, the JSON output by our benchmarks is compatible with other benchmarks, therefore it is possible to first call the `spmv` benchmark, use the resulting output JSON as input to the `solver` benchmark, and finally use the resulting solver JSON output as input to the `preconditioner` benchmark.

4.0.4 4: Benchmarking overview

The benchmark suite is invoked using the `make benchmark` command in the build directory. Under the hood, this command simply calls the script `benchmark/run_all_benchmarks.sh` so it is possible to manually launch this script as well. The behavior of the suite can be modified using environment variables. Assuming the `bash` shell is used, these can either be specified via the `export` command to persist between multiple runs:

```
export VARIABLE="value"
...
make benchmark
```

or specified on the fly, on the same line as the `make benchmark` command:

```
VARIABLE="value" ... make benchmark
```

Since `make` sets any variables passed to it as temporary environment variables, the following shorthand can also be used:

```
make benchmark VARIABLE="value" ...
```

A combination of the above approaches is also possible (e.g. it may be useful to `export` the `SYSTEM_NAME` variable, and specify the others at every benchmark run).

The benchmark suite can take a number of configuration parameters. Benchmarks can be run only for `sparse matrix vector products (spmv)`, for full solvers (with or without preconditioners), or for preconditioners only when supported. The benchmark suite also allows to target a sub-part of the SuiteSparse matrix collection. For details, see the `available benchmark options`. Here are the most important options:

- `BENCHMARK={spmv, solver, preconditioner}` - allows to select the type of benchmark to be ran.
- `EXECUTOR={reference, cuda, hip, omp, dpcpp}` - select the executor and platform the benchmarks should be ran on.
- `SYSTEM_NAME=<name>` - a name which will be used to designate this platform (e.g. V100, RadeonVII, ...).
- `SEGMENTS=<N>` - Split the benchmarked matrix space into `<N>` segments. If specified, `SEGMENT_ID` also has to be set.
- `SEGMENT_ID=<I>` - used in combination with the `SEGMENTS` variable. `<I>` should be an integer between 1 and `<N>`, the number of `SEGMENTS`. If specified, only the `<I>`-th segment of the benchmark suite will be run.
- `BENCHMARK_PRECISION` - defines the precision the benchmarks are run in. Supported values are: "double" (default), "single", "dcomplex" and "scomplex"
- `MATRIX_LIST_FILE=/path/to/matrix_list.file` - allows to list SuiteSparse matrix id or name to benchmark. As an example, a matrix list file containing the following will ensure that benchmarks are ran for only those three matrices:

```
``` 1903 Freescale/circuit5M thermal2 ```
```

#### 4.0.5 5: Publishing the results on Github and analyze the results with the GPE (optional)

The previous experiments generated json files for each matrices, each containing timing, iteration count, achieved precision, ... depending on the type of benchmark run. These files are available in the directory `${ginkgo_build_dir}/benchmark/results/`. These files can be analyzed and processed through any tool (e.g. python). In this section, we describe how to generate the plots by using Ginkgo's `GPE` tool. First, we need to publish the experiments into a Github repository which will be then linked as source input to the GPE. For this, we can simply fork the ginkgo-data repository. To do so, we can go to the github repository and use the forking interface: <https://github.com/ginkgo-project/ginkgo-data/>

Once it's done, we want to clone the repository locally, put all results online and access the GPE for plotting the results. Here are the detailed steps:

```
git clone https://github.com/<username>/ginkgo-data.git $HOME/ginkgo_benchmark/ginkgo-data
send the benchmarked data to the ginkgo-data repository
If needed, remove the old data so that no previous data is left.
rm -r ${HOME}/ginkgo_benchmark/ginkgo-data/data/${SYSTEM_NAME}
rsync -rtv ${ginkgo_build_dir}/benchmark/results/ $HOME/ginkgo_benchmark/ginkgo-data/data/
cd ${HOME}/ginkgo_benchmark/ginkgo-data/data/
The following updates the main '.json' files with the list of data.
Ensure a python 3 installation is available.
./build-list . > list.json
./agregate < list.json > agregate.json
./represent . > represent.json
git config --local user.name "<Name>"
git config --local user.email "<email>"
git commit -am "Ginkgo benchmark ${BENCHMARK} of ${SYSTEM_NAME}..."
git push
```

Note that depending on what data is of interest, you may need to update the scripts `build-list` or `agregate` to change which files you want to agglomerate and summarize (depending on the system name), or which data you want to select (solver results, spmv results, ...).

For the generating the plots in the GPE, here are the steps to go through:

1. Access the GPE: <https://ginkgo-project.github.io/gpe/>
2. Update data root URL, from <https://raw.githubusercontent.com/ginkgo-project/ginkgo-data/master/data> to <https://raw.githubusercontent.com/<username>/ginkgo-data/<branch>/data>
3. Click on the arrow to load the data, select the `Result Summary` entry above.
4. Click on `select an example` to choose a plotting script. Multiple scripts are available by default in different branches. You can use the `jsonata` and `chartjs` languages to develop your own as well.
5. The results should be available in the tab "plot" on the right side. Other tabs allow to access the result of the processed data after invoking the processing script.

## 4.0.6 6: Detailed performance analysis and debugging

Detailed performance analysis can be ran by passing the environment variable `DETAILED=1` to the benchmarking script. This detailed run is available for solvers and allows to log the internal residual after every iteration as well as log the time taken by all operations. These features are also available in the `performance-debugging` example which can be used instead and modified as needed to analyze Ginkgo's performance.

These features are implemented thanks to the loggers located in the file `${ginkgo_src_dir}/benchmark/utlis/loggers`. Ginkgo possesses hooks at all important code location points which can be inspected thanks to the logger. In this fashion, it is easy to use these loggers also for tracking memory allocation sizes and other important library aspects.

## 4.0.7 7: Available benchmark options

There are a set amount of options available for benchmarking. Most important options can be configured through the benchmarking script itself thanks to environment variables. Otherwise, some specific options are not available through the benchmarking scripts but can be directly configured when running the benchmarking program itself. For a list of all options, run for example `${ginkgo_build_dir}/benchmark/solver/solver --help`.

The supported environment variables are described in the following list:

- `BENCHMARK={spmv, solver, preconditioner}` - allows to select the type of benchmark to be ran. Default is `spmv`.

- `spmv` - Runs the sparse matrix-vector product benchmarks on the SuiteSparse collection.
- `solver` - Runs the solver benchmarks on the SuiteSparse collection. The matrix format is determined by running the `spmv` benchmarks first, and using the fastest format determined by that benchmark.
- `preconditioner` - Runs the preconditioner benchmarks on artificially generated block-diagonal matrices.
- `EXECUTOR={reference,cuda,hip,omp,dpcpp}` - select the executor and platform the benchmarks should be ran on. Default is `cuda`.
- `SYSTEM_NAME=<name>` - a name which will be used to designate this platform (e.g. V100, RadeonVII, ...) and not overwrite previous results. Default is `unknown`.
- `SEGMENTS=<N>` - Split the benchmarked matrix space into `<N>` segments. If specified, `SEGMENT_ID` also has to be set. Default is 1.
- `SEGMENT_ID=<I>` - used in combination with the `SEGMENTS` variable. `<I>` should be an integer between 1 and `<N>`, the number of `SEGMENTS`. If specified, only the `<I>`-th segment of the benchmark suite will be run. Default is 1.
- `MATRIX_LIST_FILE=/path/to/matrix_list.file` - allows to list SuiteSparse matrix id or name to benchmark. As an example, a matrix list file containing the following will ensure that benchmarks are ran for only those three matrices: `` 1903 Freescale/circuit5M thermal2 ``\*`DEVICE_ID`- the accelerator device ID to target for the benchmark. The default is 0. \*`DRY_RUN`={true, false}- If set to true, prepares the system for the benchmark runs (downloads the collections, creates the result structure, etc.) and outputs the list of commands that would normally be run, but does not run the benchmarks themselves. Default is false.
- `PRECONDS={jacobi,ic,ilu,paric,parict,parilu,parilut,ic-isai,ilu-isai,paric-isai,pari` the preconditioners to use for either `solver` or `preconditioner` benchmarks. Multiple options can be passed to this variable. Default is `none`.
- `FORMATS={csr,coo,ell,hybrid,sellp,hybridxx,cusparsedxx,hipsparsedxx}` the matrix formats to benchmark for the `spmv` phase of the benchmark. Run `$(ginkgo_build_dir)/benchmark/spmv/spmv --help` for a full list. If needed, multiple options for hybrid with different optimization parameters are available. Depending on the libraries available at build time, vendor library formats (cuSPARSE with `cusparsed` prefix or hipSPARSE with `hipsparsed` prefix) can be used as well. Multiple options can be passed. The default is `csr,coo,ell,hybrid,sellp`.
- `SOLVERS={bicgstab,bicg,cg,cgs,fcg,gmres,cb_gmres_{keep,reducel,reduce2,integer,ired`  
`_trs,upper_trsr}`
  - the solvers which should be benchmarked. Multiple options can be passed. The default is `bicgstab,cg,cgs,fcg,gmres,idr`. Note that `lower/upper_trsr` by default don't use a preconditioner, as they are by default exact direct solvers.
- `SOLVERS_PRECISION=<precision>` - the minimal residual reduction before which the solver should stop. The default is `1e-6`.
- `SOLVERS_MAX_ITERATIONS=<number>` - the maximum number of iterations with which a solver should be ran. The default is 10000.
- `SOLVERS_RHS={1, random, sinus}` - whether to use a vector of all ones, random values or  $b = A * (s / |s|)$  with  $s(idx) = \sin(idx)$  (for complex numbers,  $s(idx) = \sin(2*idx) + i * \sin(2*idx+1)$ ) as the right-hand side in solver benchmarks. Default is 1.
- `SOLVERS_INITIAL_GUESS={rhs, 0, random}` - the initial guess generation of the solvers. `rhs` uses the right-hand side, 0 uses a zero vector and `random` generates a random vector as the initial guess.
- `DETAILED={0, 1}` - selects whether detailed benchmarks should be ran. This generally provides extra, verbose information at the cost of one or more extra benchmark runs. It can be either 0 (off) or 1 (on).
- `GPU_TIMER={true, false}` - If set to true, use the gpu timer, which is valid for cuda/hip executor, to measure the timing. Default is false.

- `SOLVERS_JACOBI_MAX_BS` - sets the maximum block size for the Jacobi preconditioner (if used, otherwise, it does nothing) in the solvers benchmark. The default is '32'.
- `SOLVERS_GMRES_RESTART` - the maximum dimension of the Krylov space to use in GMRES. The default is 100.





## Chapter 5

# Contributing guidelines

We are glad that you are interested in contributing to Ginkgo. Please have a look at our coding guidelines before proposing a pull request.

### 5.1 Table of Contents

Most Important stuff

Project Structure

- Extended header files
- Using library classes

Git related

- Our git Workflow
- Writing good commit messages
- Creating, Reviewing and Merging Pull Requests

Code Style

- Automatic code formatting
- Naming Scheme
- Whitespace
- Include statement grouping
- Other Code Formatting not handled by ClangFormat
- CMake coding style

Helper Scripts

- [Create a new algorithm](#)
- [Converting CUDA code to HIP code](#)

#### Writing Tests

- [Testing know-how](#)
- [Some general rules](#)
- [Writing tests for kernels](#)

#### Documentation style

- [Developer targeted notes](#)
- [Whitespaces](#)
- [Documenting examples](#)

#### Other programming comments

- [C++ standard stream objects](#)
- [Warnings](#)
- [Avoiding circular dependencies](#)

## 5.2 Most important stuff (A TL;DR)

- `GINKGO_DEVEL_TOOLS` needs to be set to `on` to commit. This requires `clang-format` to be installed. See [Automatic code formatting](#) for more details. Once installed, you can run `make format` in your `build/` folder to automatically format your modified files. As `make format` unstages your files post-formatting, you must stage the files again once you have verified that `make format` has done the appropriate formatting, before committing the files.
- See [Our git workflow](#) to get a quick overview of our workflow.
- See [Creating, Reviewing and Merging Pull Requests](#) on how to create a Pull request.

## 5.3 Project structure

Ginkgo is divided into a `core` module with common functionalities independent of the architecture, and several kernel modules (`reference`, `omp`, `cuda`, `hip`, `dpcpp`) which contain low-level computational routines for each supported architecture.

### 5.3.1 Extended header files

Some header files from the core module have to be extended to include special functionality for specific architectures. An example of this is `core/base/math.hpp`, which has a GPU counterpart in `cuda/base/math.↵.hpp`. For such files you should always include the version from the module you are working on, and this file will internally include its `core` counterpart.

### 5.3.2 Using library classes

You can use and call functions of existing classes inside a kernel (that are defined and not just declared in a header file), however, you are not allowed to create new instances of a polymorphic class inside a kernel (or in general inside any kernel module like `cuda/hip/omp/reference`) as this creates circular dependencies between the `core` and the backend library. With this in mind, our CI contains a job which checks if such a circular dependency exists. These checks can be run manually using the `-DGINKGO_CHECK_CIRCULAR_DEPS=ON` option in the CMake configuration.

For example, when creating a new matrix class `AB` by combining existing classes `A` and `B`, the `AB::apply()` function composed of invocations to `A::apply()` and `B::apply()` can only be defined in the `core` module, it is not possible to create instances of `A` and `B` inside the `AB` kernel files. This is to avoid the aforementioned circular dependency issue. An example for such a class is the `Hybrid` matrix format, which uses the `apply()` of the `Ell` and `Coo` matrix formats. Nevertheless, it is possible to call the kernels themselves directly within the same executor. For example, `cuda::dense::add_scaled()` can be called from any other `cuda` kernel.

## 5.4 Git related

Ginkgo uses `git`, the distributed version control system to track code changes and coordinate work among its developers. A general guide to `git` can be found in [its extensive documentation](#).

### 5.4.1 Our git workflow

In Ginkgo, we prioritize keeping a clean history over accurate tracking of commits. `git rebase` is hence our command of choice to make sure that we have a nice and linear history, especially for pulling the latest changes from the `develop` branch. More importantly, rebasing upon `develop` is **required** before the commits of the PR are merged into the `develop` branch.

### 5.4.2 Writing good commit messages

With software sustainability and maintainability in mind, it is important to write commit messages that are short, clear and informative. Ideally, this would be the format to prefer:

```
Summary of the changes in a sentence, max 50 chars.
More detailed comments:
+ Changes that have been added.
- Changes that have been removed.
Related PR: https://github.com/ginkgo-project/ginkgo/pull/<PR-number>
```

You can refer to [this informative guide](#) for more details.

#### 5.4.2.1 Attributing credit

`Git` has a nice feature where it allows you to add a co-author for your commit, if you would like to attribute credits for the changes made in the commit. This can be done by:

```
Commit message.
Co-authored-by: Name <email@domain>
```

In the Ginkgo commit history, this is most common associated with suggested improvements from code reviews.

### 5.4.3 Creating, Reviewing and Merging Pull Requests

- The `develop` branch is the default branch to submit PR's to. From time to time, we merge the `develop` branch to the `main` branch and create tags on the `main` to create new releases of Ginkgo. Therefore, all pull requests must be merged into `develop`.
- Please have a look at the labels and make sure to add the relevant labels.
- You can mark the PR as a WIP if you are still working on it, Ready for Review when it is ready for others to review it.
- Assignees to the PR should be the ones responsible for merging that PR. Currently, it is only possible to assign members within the `ginkgo-project`.
- Each pull request requires at least two approvals before merging.
- PR's created from within the repository will automatically trigger two CI pipelines on pushing to the branch from the which the PR has been created. The Github Actions pipeline tests our framework on Mac OSX and on Windows platforms. Another comprehensive Linux based pipeline is run from a [mirror on gitlab](#) and contains additional checks like static analysis and test coverage.
- Once a PR has been approved and the build has passed, one of the reviewers can mark the PR as `READY TO MERGE`. At this point the creator/assignee of the PR *needs to* verify that the branch is up to date with `develop` and rebase it on `develop` if it is not.

## 5.5 Code style

### 5.5.1 Automatic code formatting

Ginkgo uses `pre-commit` to automatically apply code formatting when committing changes to git. What formatting is applied is managed through `ClangFormat` with a custom `.clang-format` configuration file (mostly based on ClangFormat's *Google* style). **Make sure you have pre-commit set up and running properly** before committing anything that will end up in a pull request against `ginkgo-project/ginkgo` repository. In addition, you should **never** modify the `.clang-format` configuration file shipped with Ginkgo. E.g. if ClangFormat has trouble reading this file on your system, you should install a newer version of ClangFormat, and avoid commenting out parts of the configuration file.

ClangFormat is the primary tool that helps us achieve a uniform look of Ginkgo's codebase, while reducing the learning curve of potential contributors. However, ClangFormat configuration is not expressive enough to incorporate the entire coding style, so there are several additional rules that all contributed code should follow.

*Note:* To learn more about how ClangFormat will format your code, see existing files in Ginkgo, `.clang-format` configuration file shipped with Ginkgo, and ClangFormat's documentation.

### 5.5.2 Naming scheme

#### 5.5.2.1 Filenames

Filenames use `snake_case` and use the following extensions:

- C++ source files: `.cpp`
- C++ header files: `.hpp`

- CUDA source files: `.cu`
- CUDA header files: `.cuh`
- HIP source files: `.hip.cpp`
- HIP header files: `.hip.hpp`
- Common source files used by both CUDA and HIP: `.hpp.inc`
- CMake utility files: `.cmake`
- Shell scripts: `.sh`

*Note:* A C++ source/header file is considered a CUDA file if it contains CUDA code that is not guarded with `#if` guards that disable this code in non-CUDA compilers. I.e. if a file can be compiled by a general C++ compiler, it is not considered a CUDA file.

#### 5.5.2.2 Macros

Macros (both object-like and function-like macros) use `CAPITAL_CASE`. They have to start with `GKO_` to avoid name clashes (even if they are `#undef-ed` in the same file!).

#### 5.5.2.3 Variables

Variables use `snake_case`.

#### 5.5.2.4 Constants

Constants use `snake_case`.

#### 5.5.2.5 Functions

Functions use `snake_case`.

#### 5.5.2.6 Structures and classes

Structures and classes which do not experience polymorphic behavior (i.e. do not contain virtual methods, nor members which experience polymorphic behavior) use `snake_case`.

All other structures and classes use `CamelCase`.

#### 5.5.2.7 Members

All structure / class members use the same naming scheme as they would if they were not members:

- methods use the naming scheme for functions
- data members the naming scheme for variables or constants
- type members for classes / structures

Additionally, non-public data members end with an underscore (`_`).

### 5.5.2.8 Namespaces

Namespaces use `snake_case`.

### 5.5.2.9 Template parameters

- Type template parameters use `CamelCase`, for example `ValueType`.
- Non-type template parameters use `snake_case`, for example `subwarp_size`.

## 5.5.3 Whitespace

Spaces and tabs are handled by ClangFormat, but blank lines are only partially handled (the current configuration doesn't allow for more than 2 blank lines). Thus, contributors should be aware of the following rules for blank lines:

1. Top-level statements and statements directly within namespaces are separated with 2 blank lines. The first / last statement of a namespace is separated by two blank lines from the opening / closing brace of the namespace.

- (a) *exception*: if the first **or** the last statement in the namespace is another namespace, then no blank lines are required *example*: ````c++ namespace foo {`

```
struct x {
};

} // namespace foo

namespace bar {
namespace baz {

void f();

} // namespace baz
} // namespace bar
```
```

2. `_exception_`: in header files whose only purpose is to `_declare_` a bunch of functions (e.g. the `*_kernel.hpp` files) these declarations can be separated by only 1 blank line (note: standard rules apply for all other statements that might be present in that file)`
3. `_exception_`: "related" statement can have 1 blank line between them. "Related" is not a strictly defined adjective in this sense, but is in general one of:
 1. overload of a same function,
 2. function / class template and it's specializations,
 3. macro that modifies the meaning or adds functionality to the previous / following statement.

However, simply calling function `'f'` from function `'g'` does not imply that `'f'` and `'g'` are "related".

1. Statements within structures / classes are separated with 1 blank line. There are no blank lines between the first / last statement in the structure / class.

- (a) *exception*: there is no blank line between an access modifier (`private`, `protected`, `public`) and the following statement. *example*: ````c++ class foo { public: int get_x() const noexcept { return x_; } int &get_x() noexcept { return x_; } private: int x_; }; ````
- 2. Function bodies cannot have multiple consecutive blank lines, and a single blank line can only appear between two logical sections of the function.
- 3. Unit tests should follow the `AAA` pattern, and a single blank line must appear between consecutive "A" sections. No other blank lines are allowed in unit tests.
- 4. Enumeration definitions should have no blank lines between consecutive enumerators.

5.5.4 Include statement grouping

The concrete ordering will be done by `clang-format`. Here are the rules that `clang-format` will follow. In general, all include statements should be present on the top of the file, ordered in the following groups, with *one* blank lines between each group:

1. Main header file (e.g. `core/foo/bar.hpp` included in `core/foo/bar.cpp`, or in the unit test `core/test/foo/bar.cpp`)
2. Standard library headers (e.g. `vector`)
3. Executor specific library headers (e.g. `omp.h`)
4. System third-party library headers (e.g. `papi.h`)
5. Public Ginkgo headers
6. Local headers

Example: A file `core/base/my_file.cpp` might have an include list like this:

```
{c++}
#include "ginkgo/core/base/my_file.hpp"
#include <algorithm>
#include <vector>
#include <tuple>
#include <omp.h>
#include <papi.h>
#include <third_party/blas/cblas.hpp>
#include <third_party/lapack/lapack.hpp>
#include <ginkgo/core/base/executor.hpp>
#include <ginkgo/core/base/lin_op.hpp>
#include <ginkgo/core/base/types.hpp>
#include "core/base/my_file_kernels.hpp"
```

5.5.4.1 Main header

This section presents the handling of the main header attributed to a file. For a given file, the main header is the header that contains the declarations of the functions, classes, etc., which are implemented in this file. In the previous example, this would be `#include "ginkgo/core/base/my_file.hpp"`. The `clang-format` tool figures out the main header. The only intervention from a contributor is to *always* include the main header using `"..."`.

Please note that this only applies to implementation files, so files ending in `.cpp` or `.cu`.

5.5.4.2 Some general comments.

1. Private headers of Ginkgo should not be included within the public Ginkgo header.
2. It is a good idea to keep the headers self-sufficient, See [Google Style guide for reasoning](#). When compiling with `GINKGO_CHECK_CIRCULAR_DEPS` enabled, this property is explicitly checked.
3. The recommendations of the `iwyu` (Include what you use) tool can be used to make sure that the headers are self-sufficient and that the compiled files (`.cu`, `.cpp`, `.hip.cpp`) include only what they use. A [CI pipeline](#) is available that runs with the `iwyu` tool. Please be aware that this tool can be incorrect in some cases.

5.5.5 Other Code Formatting not handled by ClangFormat

5.5.5.1 Control flow constructs

Single line statements should be avoided in all cases. Use of brackets is mandatory for all control flow constructs (e.g. `if`, `for`, `while`, ...).

5.5.5.2 Variable declarations

C++ supports declaring / defining multiple variables using a single *type-specifier*. However, this is often very confusing as references and pointers exhibit strange behavior:

```
{c++}
template <typename T> using pointer = T *;
int *      x, y; // x is a pointer, y is not
pointer<int> x, y; // both x and y are pointers
```

For this reason, **always** declare each variable on a separate line, with its own *type-specifier*.

5.5.6 CMake coding style

5.5.6.1 Whitespaces

All alignment in CMake files should use four spaces.

5.5.6.2 Use of macros vs functions

Macros in CMake do not have a scope. This means that any variable set in this macro will be available to the whole project. In contrast, functions in CMake have local scope and therefore all set variables are local only. In general, wrap all piece of algorithms using temporary variables in a function and use macros to propagate variables to the whole project.

5.5.6.3 Naming style

All Ginkgo specific variables should be prefixed with a `GINKGO_` and all functions by `ginkgo_`.

5.6 Helper scripts

To facilitate easy development within Ginkgo and to encourage coders and scientists who do not want get bogged down by the details of the Ginkgo library, but rather focus on writing the algorithms and the kernels, Ginkgo provides the developers with a few helper scripts.

5.6.1 Create a new algorithm

A `create_new_algorithm.sh` script is available for developers to facilitate easy addition of new algorithms. The options it provides can be queried with `./create_new_algorithm.sh --help`

The main objective of this script is to add files and boiler plate code for the new algorithm using a model and an instance of that model. For example, models can be any one of `factorization`, `matrix`, `preconditioner` or `solver`. For example to create a new solver named `my_solver` similar to `gmres`, you would set the `ModelType` to `solver` and set the `ModelName` to `gmres`. This would duplicate the core algorithm and kernels of the `gmres` algorithm and replace the naming to `my_solver`. Additionally, all the kernels of the new `my_solver` are marked as `GKO_NOT_IMPLEMENTED`. For easy navigation and `.txt` file is created in the folder where the script is run, which lists all the TODO's. These TODO's can also be found in the corresponding files.

5.6.2 Converting CUDA code to HIP code

We provide a `cuda2hip` script that converts `cuda` kernel code into `hip` kernel code. Internally, this script calls the `hipify script` provided by HIP, converting the CUDA syntax to HIP syntax. Additionally, it also automatically replaces the instances of CUDA with HIP as appropriate. Hence, this script can be called on a Ginkgo CUDA file. You can find this script in the `dev_tools/scripts/` folder.

5.7 Writing Tests

Ginkgo uses the `GTest framework` for the unit test framework within Ginkgo. Writing good tests are extremely important to verify the functionality of the new code and to make sure that none of the existing code has been broken.

5.7.1 Testing know-how

- GTest provides a `comprehensive documentation` of the functionality available within Gtest.
- Reduce code duplication with `Testing Fixtures`, `TEST_F`
- Write templated tests using `TYPED_TEST`.

5.7.2 Some general rules.

- Unit tests must follow the `KISS principle`.
- Unit tests must follow the `AAA` pattern, and a single blank line must appear between consecutive "A" sections.

5.7.3 Writing tests for kernels

- Reference kernels, kernels on the `ReferenceExecutor`, are meant to be single threaded reference implementations. Therefore, tests for reference kernels need to be performed with data that can be as small as possible. For example, matrices lesser than 5x5 are acceptable. This allows the reviewers to verify the results for exactness with tools such as MATLAB.
- OpenMP, CUDA, HIP and DPC++ kernels have to be tested against the reference kernels. Hence data for the tests of these kernels can be generated in the test files using helper functions or by using external files to be read through the standard input. In particular for CUDA, HIP and DPC++ the data size should be at least bigger than the architecture's warp size to ensure there is no corner case in the kernels.

5.8 Documentation style

Documentation uses standard Doxygen.

5.8.1 Developer targeted notes

Make use of `@internal` doxygen tag. This can be used for any comment which is not intended for users, but is useful to better understand a piece of code.

5.8.2 Whitespaces

5.8.2.1 After named tags such as `<tt>@param foo</tt>`

The documentation tags which use an additional name should be followed by two spaces in order to better distinguish the text from the doxygen tag. It is also possible to use a line break instead.

5.8.3 Documenting examples

There are two main steps:

1. First, you can just copy over the `doc/` folder (you can copy it from the example most relevant to you) and adapt your example names and such, then you can modify the actual documentation.

- In `tooltip`: A short description of the example.
- In `short-intro`: The name of the example.
- In `results.dox`: Run the example and write the output you get.
- In `kind`: The kind of the example. For different kinds see [the documentation](#). Examples can be of `basic`, `techniques`, `logging`, `stopping_criteria` or `preconditioners`. If your example does not fit any of these categories, feel free to create one.
- In `intro.dox`: You write an explanation of your code with some introduction similar to what you see in an existing example most relevant to you.
- In `builds-on`: You write the examples it builds on.

1. You also need to modify the `examples.hpp.in` file. You add the name of the example in the main section and in the section that you specified in the `doc/kind` file in the example documentation.

5.9 Other programming comments

5.9.1 C++ standard stream objects

These are global objects and are shared inside the same translation unit. Therefore, whenever its state or formatting is changed (e.g. using `std::hex` or floating point formatting) inside library code, make sure to restore the state before returning the control to the user. See this [stackoverflow question](#) for examples on how to do it correctly. This is extremely important for header files.

5.9.2 Warnings

By default, the `-DCMAKE_CXX_FLAGS` should be set to `-Wpedantic` to emit pedantic warnings by default. The CI system additionally also has a step where it compiles for pedantic warnings to be errors.

5.9.3 Avoiding circular dependencies

To facilitate finding circular dependencies issues (see [Using library classes](#) for more details), a CI step `no-circular-deps` was created. For more details on its usage, see [this pipeline](#), where Ginkgo did not abide to this policy and [PR #278](#) which fixed this. Note that doing so is not enough to guarantee with 100% accuracy that no circular dependency is present. For an example of such a case, take a look at [this pipeline](#) where one of the compiler setups detected an incorrect dependency of the `cuda` module (due to `jacobi`) on the `core` module.

Chapter 6

Citing Ginkgo

The main Ginkgo paper describing Ginkgo's purpose, design and interface is available through the following reference:

```
@article{ginkgo-toms-2022,  
  title = {{Ginkgo: A Modern Linear Operator Algebra Framework for High Performance Computing}},  
  volume = {48},  
  copyright = {All rights reserved},  
  issn = {0098-3500},  
  shorttitle = {Ginkgo},  
  url = {https://doi.org/10.1145/3480935},  
  doi = {10.1145/3480935},  
  number = {1},  
  urldate = {2022-02-17},  
  journal = {ACM Transactions on Mathematical Software},  
  author = {Anzt, Hartwig and Cojean, Terry and Flegar, Goran and Göbel, Fritz and Grützmacher, Thomas and  
    Nayak, Pratik and Ribizel, Tobias and Tsai, Yuhsiang Mike and Quintana-Ortí, Enrique S.},  
  month = feb,  
  year = {2022},  
  keywords = {ginkgo, healthy software lifecycle, High performance computing, multi-core and manycore  
    architectures},  
  pages = {2:1--2:33}  
}
```

Multiple topical papers exist on Ginkgo and its algorithms. The following papers can be used to cite specific aspects of the Ginkgo project.

6.0.1 The Ginkgo Software

The Ginkgo software itself was reviewed and has a paper published in the Journal of Open Source Software, which can be cited with the following reference:

```
@article{GinkgoJoss2020,  
  doi = {10.21105/joss.02260},  
  url = {https://doi.org/10.21105/joss.02260},  
  year = {2020},  
  publisher = {The Open Journal},  
  volume = {5},  
  number = {52},  
  pages = {2260},  
  author = {Hartwig Anzt and Terry Cojean and Yen-Chen Chen and Goran Flegar and Fritz Göbel and Thomas  
    Grützmacher and Pratik Nayak and Tobias Ribizel and Yu-Hsiang Tsai},  
  title = {Ginkgo: A high performance numerical linear algebra library},  
  journal = {Journal of Open Source Software}  
}
```

6.0.2 On Portability

```
@misc{tsai2020amdportability,  
  title={Preparing Ginkgo for AMD GPUs -- A Testimonial on Porting CUDA Code to HIP},  
  author={Yuhsiang M. Tsai and Terry Cojean and Tobias Ribizel and Hartwig Anzt},  
  year={2020},  
  eprint={2006.14290},  
  archivePrefix={arXiv},  
  primaryClass={cs.MS}  
}
```

6.0.3 On Software Sustainability

```
@inproceedings{anzt2019pascob,
  author = {Anzt, Hartwig and Chen, Yen-Chen and Cojean, Terry and Dongarra, Jack and Flegar, Goran and Nayak,
    Pratik and Quintana-Ort'\i}, Enrique S. and Tsai, Yuhsiang M. and Wang, Weichung},
  title = {Towards Continuous Benchmarking: An Automated Performance Evaluation Framework for High
    Performance Software},
  year = {2019},
  isbn = {9781450367707},
  publisher = {Association for Computing Machinery},
  address = {New York, NY, USA},
  url = {https://doi.org/10.1145/3324989.3325719},
  doi = {10.1145/3324989.3325719},
  booktitle = {Proceedings of the Platform for Advanced Scientific Computing Conference},
  articleno = {9},
  numpages = {11},
  keywords = {interactive performance visualization, healthy software lifecycle, continuous integration,
    automated performance benchmarking},
  location = {Zurich, Switzerland},
  series = {PASC '19}
}
```

6.0.4 On SpMV or solvers performance

```
@InProceedings{tsai2020amdspmv,
  author="Tsai, Yuhsiang M.
  and Cojean, Terry
  and Anzt, Hartwig",
  editor="Sadayappan, Ponnuswamy
  and Chamberlain, Bradford L.
  and Juckeland, Guido
  and Ltaief, Hatem",
  title="Sparse Linear Algebra on AMD and NVIDIA GPUs -- The Race Is On",
  booktitle="High Performance Computing",
  year="2020",
  publisher="Springer International Publishing",
  address="Cham",
  pages="309--327",
  abstract="Efficiently processing sparse matrices is a central and performance-critical part of many
    scientific simulation codes. Recognizing the adoption of manycore accelerators in HPC, we evaluate in
    this paper the performance of the currently best sparse matrix-vector product (SpMV) implementations
    on high-end GPUs from AMD and NVIDIA. Specifically, we optimize SpMV kernels for the CSR, COO, ELL,
    and HYB format taking the hardware characteristics of the latest GPU technologies into account. We
    compare for 2,800 test matrices the performance of our kernels against AMD's hipSPARSE library and
    NVIDIA's cuSPARSE library, and ultimately assess how the GPU technologies from AMD and NVIDIA compare
    in terms of SpMV performance.",
  isbn="978-3-030-50743-5"
}

@article{anzt2020spmv,
  author = {Anzt, Hartwig and Cojean, Terry and Yen-Chen, Chen and Dongarra, Jack and Flegar, Goran and Nayak,
    Pratik and Tomov, Stanimire and Tsai, Yuhsiang M. and Wang, Weichung},
  title = {Load-Balancing Sparse Matrix Vector Product Kernels on GPUs},
  year = {2020},
  issue_date = {March 2020},
  publisher = {Association for Computing Machinery},
  address = {New York, NY, USA},
  volume = {7},
  number = {1},
  issn = {2329-4949},
  url = {https://doi.org/10.1145/3380930},
  doi = {10.1145/3380930},
  journal = {ACM Trans. Parallel Comput.},
  month = mar,
  articleno = {2},
  numpages = {26},
  keywords = {irregular matrices, GPUs, Sparse Matrix Vector Product (SpMV)}
}

@misc{tsai2020evaluating,
  title={Evaluating the Performance of NVIDIA's A100 Ampere GPU for Sparse Linear Algebra Computations},
  author={Yuhsiang Mike Tsai and Terry Cojean and Hartwig Anzt},
  year={2020},
  eprint={2008.08478},
  archivePrefix={arXiv},
  primaryClass={cs.MS}
}
```

Chapter 7

Example programs

Here you can find example programs that demonstrate the usage of Ginkgo. Some examples are built on one another and some are stand-alone and demonstrate a concept of Ginkgo, which can be used in your own code.

You can browse the available example programs

1. as [a graph](#) that shows how example programs build upon each other.
2. as [a list](#) that provides a short synopsis of each program.
3. or [grouped by topic](#).

By default, all Ginkgo examples are built using CMake.


An example for building the examples and using Ginkgo as an external library without CMake can be found in the script provided for each example, which should be called with the form: `./build.sh PATH_TO_GINKGO_BUILD_DIR`


By default, Ginkgo is compiled with at least `-DGINKGO_BUILD_REFERENCE=ON`. Ginkgo also tries to detect your environment setup (presence of CUDA, ...) to enable the relevant accelerator modules. If you want to target a specific GPU, make sure that Ginkgo is compiled with the accelerator specific module enabled, such as:

1. `-DGINKGO_BUILD_CUDA=ON` option for NVIDIA GPUs.
2. `-DGINKGO_BUILD_HIP=ON` option for AMD or NVIDIA GPUs.
3. `-DGINKGO_BUILD_SYCL=ON` option for Intel GPUs (and possibly any other platform).

Connections between example programs

The following graph shows the connections between example programs and how they build on each other. Click on any of the boxes to go to one of the programs. If you hover your mouse pointer over a box, a brief description of the program should appear.



Legend:**Example programs**

| | |
|---|--|
| The simple-solver program | A minimal CG solver in Ginkgo, which reads a matrix from a file. |
| The minimal-cuda-solver program | A minimal solver on the CUDA executor than can be run on NVIDIA GPU's. |
| The poisson-solver program | Solve an actual physically relevant problem, the poisson problem. The matrix is generated within Ginkgo. |
| The preconditioned-solver program | Using a Jacobi preconditioner to solve a linear system. |
| The ilu-preconditioned-solver program | Using an ILU preconditioner to solve a linear system. |
| The performance-debugging program | Using Loggers to debug the performance within Ginkgo. |
| The three-pt-stencil-solver program | Using a three point stencil to solve the poisson equation with array views. |
| The nine-pt-stencil-solver program | Using a nine point 2D stencil to solve the poisson equation with array views. |
| The external-lib-interfacing program | Using Ginkgo's solver with the external library deal.II. |
| The custom-logger program | Creating a custom logger specifically for comparing the recurrent and the real residual norms. |
| The custom-matrix-format program | Creating a matrix-free stencil solver by using Ginkgo's advanced methods to build your own custom matrix format. |
| The inverse-iteration program | Using Ginkgo to compute eigenvalues of a matrix with the inverse iteration method. |
| The simple-solver-logging program | Using the logging functionality in Ginkgo to get solver and other information to diagnose and debug your code. |
| The papi-logging program | Using the PAPI logging library in Ginkgo to get advanced information about your code and its behaviour. |
| The ginkgo-overhead program | Measuring the overhead of the Ginkgo library. |
| The custom-stopping-criterion program | Creating a custom stopping criterion for the iterative solution process. |
| The ginkgo-ranges program | Using the ranges concept to factorize a matrix with the LU factorization. |

| | |
|---|--|
| The mixed-spmv program | Shows the Ginkgo mixed precision spmv functionality. |
| The mixed-precision-ir program | Manual implementation of a Mixed Precision Iterative Refinement (MPIR) solver. |
| The adaptiveprecision-blockjacobi program | Shows how to use the adaptive precision block-Jacobi preconditioner. |
| The cb-gmres program | Using the Ginkgo CB-GMRES solver (Compressed Basis GMRES). |
| The heat-equation program | Solving a 2D heat equation and showing matrix assembly, vector initialization and solver setup in a more complex setting with output visualization. |
| The iterative-refinement program | Using a low accuracy CG solver as an inner solver to an iterative refinement (IR) method which solves a linear system. |
| The ir-ilu-preconditioned-solver program | Combining iterative refinement with the adaptive precision block-Jacobi preconditioner to approximate triangular systems occurring in ILU preconditioning. |
| The par-ilu-convergence program | Convergence analysis at the examples of parallel incomplete factorization solver. |
| The preconditioner-export program | Explicit generation and storage of preconditioners for given matrices. |
| The multigrid-preconditioned-solver program | Use multigrid as preconditioner to a solver. |
| The mixed-multigrid-solver program | Use multigrid with different precision multigrid_level as a solver. |
| The distributed-solver program | Use a distributed solver to solve a 1D Laplace equation. |

Example programs grouped by topics

| | |
|---|--|
| Solving a simple linear system with choice of executors | The simple-solver program |
| Debug the performance of a solver or preconditioner | The performance-debugging program
The preconditioner-export program |
| Using the CUDA executor | The minimal-cuda-solver program |
| Using preconditioners | The preconditioned-solver program ,
The ilu-preconditioned-solver program ,
The ir-ilu-preconditioned-solver program ,
The adaptiveprecision-blockjacobi program ,
The par-ilu-convergence program ,
The preconditioner-export program
The multigrid-preconditioned-solver program |
| Iterative refinement | The iterative-refinement program ,
The mixed-precision-ir program ,
The ir-ilu-preconditioned-solver program |
| Solving a physically relevant problem | The poisson-solver program ,
The three-pt-stencil-solver program ,
The nine-pt-stencil-solver program ,
The custom-matrix-format program |

| | |
|---|--|
| Reading in a matrix and right hand side from a file | The simple-solver program ,
The minimal-cuda-solver program ,
The preconditioned-solver program ,
The ilu-preconditioned-solver program ,
The inverse-iteration program ,
The simple-solver-logging program ,
The papi-logging program ,
The custom-stopping-criterion program ,
The custom-logger program |
|---|--|

Basic techniques

| | |
|--|---|
| Using Ginkgo with external libraries | The external-lib-interfacing program |
| Customizing Ginkgo | The custom-logger program ,
The custom-stopping-criterion program ,
The custom-matrix-format program |
| Writing your own matrix format | The custom-matrix-format program |
| Using Ginkgo to construct more complex linear algebra routines | The inverse-iteration program |
| Logging within Ginkgo | The simple-solver-logging program ,
The papi-logging program ,
The performance-debugging program
The custom-logger program |
| Constructing your own stopping criterion | The custom-stopping-criterion program |
| Using ranges in Ginkgo | The ginkgo-ranges program |
| Mixed precision | The mixed-spmv program ,
The mixed-precision-ir program ,
The adaptiveprecision-blockjacobi program
The mixed-multigrid-solver program |
| Multigrid | The multigrid-preconditioned-solver program
The mixed-multigrid-solver program |
| Configure a solver from a config file | The file-config-solver program |
| Distributed | The distributed-solver program |

Advanced techniques

Chapter 8

The adaptiveprecision-blockjacobi program

The preconditioned solver example..

This example depends on preconditioned-solver.

This example shows how to use the adaptive precision block-Jacobi preconditioner.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. The preconditioned CG solver is enhanced with a block-Jacobi preconditioner that optimizes the storage format for the distinct inverted diagonal blocks to the numerical requirements. The example features the iteration count and runtime of the CG solver.

The commented program

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS and initial guess as 1

```
gko::size_type size = A->get_size()[0];
```

```

auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);

```

Calculate initial residual by overwriting b

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);

```

copy b again

```

b->copy_from(host_x);

```

Create solver factory

```

const RealValueType reduction_factor = 1e-7;
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(10000u),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor))

```

Add preconditioner, these 2 lines are the only difference from the simple solver example

```

    .with_preconditioner(
        bj::build().with_max_block_size(16u).with_storage_optimization(
            gko::precision_reduction::autodetect()))
    .on(exec);

```

Create solver

```

std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
solver_gen->add_logger(logger);
auto solver = solver_gen->generate(A);

```

Solve system

```

exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);

```

Get residual

```

auto res = gko::as<real_vec>(logger->get_residual_norm());
auto impl_res = gko::as<real_vec>(logger->get_implicit_sq_resnorm());
std::cout << "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
std::cout << "Implicit residual norm squared (r^2):\n";
write(std::cout, impl_res);

```

Print solver statistics

```

std::cout << "CG iteration count: " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG execution time [ms]: "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
}

```

Results

This is the expected output:

```

Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
194.679
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
5.69384e-06
Implicit residual norm squared (r^2):
%%MatrixMarket matrix array real general
1 1
1.27043e-15
CG iteration count: 5
CG execution time [ms]: 0.080041

```

Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 1.;
    }
    auto x = gko::clone(exec, host_x);
    auto b = gko::clone(exec, host_x);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_x);
    const RealValueType reduction_factor = 1e-7;
    auto solver_gen =
        cg::build()
            .with_criteria(gko::stop::Iteration::build().with_max_iters(10000u),
                          gko::stop::ResidualNorm<ValueType>::build()
                              .with_reduction_factor(reduction_factor))
            .with_preconditioner(
                bj::build().with_max_block_size(16u).with_storage_optimization(
                    gko::precision_reduction::autodetect()))
            .on(exec);
    std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
        gko::log::Convergence<ValueType>::create();
    solver_gen->add_logger(logger);
    auto solver = solver_gen->generate(A);
    exec->synchronize();
    std::chrono::nanoseconds time(0);
    auto tic = std::chrono::steady_clock::now();
    solver->apply(b, x);
    auto toc = std::chrono::steady_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
    auto res = gko::as<real_vec>(logger->get_residual_norm());
    auto impl_res = gko::as<real_vec>(logger->get_implicit_sq_resnorm());
    std::cout << "Initial residual norm sqrt(r^T r):\n";
    write(std::cout, initres);
    std::cout << "Final residual norm sqrt(r^T r):\n";
```

```
write(std::cout, res);
std::cout << "Implicit residual norm squared (r^2):\n";
write(std::cout, impl_res);
std::cout << "CG iteration count:          " << logger->get_num_iterations()
<< std::endl;
std::cout << "CG execution time [ms]:  "
<< static_cast<double>(time.count()) / 1000000.0 << std::endl;
}
```

Chapter 9

The batched-solver program

Using and interfacing with a batched solver..

Using batched solvers

This example shows how to use Ginkgo batched solvers with data coming from an application. The "application" in this case is just a function in the example itself; nevertheless, the steps to be taken are shown.

A 'batch' here means a set of small linear systems that can be solved independently, but each system is too small to use an entire computing device. A requirement is that all the systems need to have the same sparsity pattern.

The commented program

```
template <typename InputType>
auto unbatch(const InputType* batch_object)
{
    auto unbatched_mats =
        std::vector<std::unique_ptr<typename InputType::unbatch_type>>{};
    for (size_type b = 0; b < batch_object->get_num_batch_items(); ++b) {
        unbatched_mats.emplace_back(
            batch_object->create_const_view_for_item(b)->clone());
    }
    return unbatched_mats;
}
// namespace detail
```

'Application' structures and functions

Structure to simulate application data related to the linear systems to be solved.

We use raw pointers below to demonstrate how to handle the situation when the application only gives us raw pointers. Ideally, one should use Ginkgo's [gko::array](#) class here. In this example, we assume that the data is in a format that can directly be given to a `batch::matrix::Csr` object.

```
struct ApplSysData {
```

Number of small systems in the batch.

```
size_type nsystems;
```

Number of rows in each system.

```
int nrows;
```

Number of non-zeros in each system matrix.

```
int nnz;
```

Row pointers for one matrix

```
const index_type* row_ptrs;
```

Column indices of non-zeros for one matrix

```
const index_type* col_idxs;
```

Nonzero values for all matrices in the batch, concatenated

```
const value_type* all_values;
```

RHS vectors for all systems in the batch, concatenated

```
const value_type* all_rhs;
};
/ *
 * Generates a batch of tridiagonal systems.
 *
 * @param nrows Number of rows in each system.
 * @param nsystems Number of systems in the batch.
 * @param exec The device executor to use for the solver.
 * @note Normally, the application may not deal with Ginkgo executors, nor do
 * we need it to. Here, we use the executor for backend-independent device
 * memory allocation. The application, for example, might assume Hip (for AMD
 * GPUs) and use 'hipMalloc' directly.
 * /
ApplSysData appl_generate_system(const int nrows, const size_type nsystems,
                                std::shared_ptr<gko::Executor> exec);
```

Deallocate application data.

```
void appl_clean_up(ApplSysData& appl_data, std::shared_ptr<gko::Executor> exec);
int main(int argc, char* argv[])
{
```

Print ginkgo version information

```
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0]
               << " [executor] [num_systems] [num_rows] [print_residuals] "
               << "[num_reps]"
               << std::endl;
    std::exit(-1);
}
```

Where do you want to run your solver ?

The `gko::Executor` class is one of the cornerstones of Ginkgo. Currently, we have support for an `gko::OmpExecutor`, which uses OpenMP multi-threading in most of its kernels, a `gko::ReferenceExecutor`, a single threaded specialization of the OpenMP executor, `gko::CudaExecutor`, `gko::HipExecutor`, `gko::DpcppExecutor` which runs the code on a NVIDIA, AMD and Intel GPUs, respectively.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/ routines within Ginkgo should automatically work and run on the executor with any other changes.

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
};
```



```
    }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
const size_type num_systems = argc >= 3 ? std::atoi(argv[2]) : 2;
const int num_rows = argc >= 4 ? std::atoi(argv[3]) : 32; // per system
```

Whether to print the residuals or not.

```
const bool print_residuals =
    argc >= 5 ? (std::string(argv[4]) == "true") : false;
```

The number of repetitions for the timing.

```
const int num_reps = argc >= 6 ? std::atoi(argv[5]) : 20;
```

Generate data

The "application" generates the batch of linear systems on the device

```
auto appl_sys = appl_generate_system(num_rows, num_systems, exec);
```

Create batch_dim object to describe the dimensions of the batch matrix.

```
auto batch_mat_size =
    gko::batch_dim<2>(num_systems, gko::dim<2>(num_rows, num_rows));
auto batch_vec_size =
    gko::batch_dim<2>(num_systems, gko::dim<2>(num_rows, 1));
```

Use of application-allocated memory

We can either work on the existing memory allocated in the application, or we can copy it for the linear solve. Ginkgo expects the nonzero values for all the small matrices to be allocated contiguously, one matrix after the other.

```
auto vals_view = gko::array<value_type>::const_view(
    exec, num_systems * appl_sys.nnz, appl_sys.all_values);
auto rowptrs_view = gko::array<index_type>::const_view(exec, num_rows + 1,
    appl_sys.row_ptrs);
auto colidxs_view = gko::array<index_type>::const_view(exec, appl_sys.nnz,
    appl_sys.col_idxs);

auto A = gko::share(mtx_type::create_const(
    exec, batch_mat_size, std::move(vals_view), std::move(colidxs_view),
    std::move(rowptrs_view)));
```

RHS and solution vectors

batch_stride object specifies the access stride within the individual matrices (vectors) in the batch. In this case, we specify a stride of 1 as the common value for all the matrices. Create RHS, again reusing application allocation

```
auto b_view = gko::array<value_type>::const_view(
    exec, num_systems * num_rows, appl_sys.all_rhs);
auto b = vec_type::create_const(exec, batch_vec_size, std::move(b_view));
```

Create initial guess as 0 and copy to device

```
auto x = vec_type::create(exec);
auto host_x = vec_type::create(exec->get_master(), batch_vec_size);
for (size_type isys = 0; isys < num_systems; isys++) {
    for (int irow = 0; irow < num_rows; irow++) {
        host_x->at(isys, irow, 0) = gko::zero<value_type>();
    }
}
x->copy_from(host_x.get());
```

Create the batch solver factory

```
const real_type reduction_factor{1e-10};
```

Create a batched solver factory with relevant parameters.

```
auto solver =
    bicgstab::build()
        .with_max_iterations(500)
        .with_tolerance(reduction_factor)
        .with_tolerance_type(gko::batch::stop::tolerance_type::relative)
        .on(exec)
        ->generate(A);
```

Batch logger

Create a logger to obtain the iteration counts and "implicit" residual norms for every system after the solve.

```
std::shared_ptr<const gko::batch::log::BatchConvergence<value_type>>
    logger = gko::batch::log::BatchConvergence<value_type>::create();
```

Generate and solve

add the logger to the solver

```
solver->add_logger(logger);
```

Solve the batch system

```
auto x_clone = gko::clone(x);
```

Warmup

```
for (int i = 0; i < 3; ++i) {
    x_clone->copy_from(x.get());
    solver->apply(b, x_clone);
}
double apply_time = 0.0;
for (int i = 0; i < num_reps; ++i) {
    x_clone->copy_from(x.get());
    exec->synchronize();
    std::chrono::steady_clock::time_point t1 =
        std::chrono::steady_clock::now();
    solver->apply(b, x_clone);
    exec->synchronize();
    std::chrono::steady_clock::time_point t2 =
        std::chrono::steady_clock::now();
    auto time_span =
        std::chrono::duration_cast<std::chrono::duration<double>>(t2 - t1);
    apply_time += time_span.count();
}
x->copy_from(x_clone.get());
```

This is not necessary, but one might want to remove the logger before the next solve using the same solver object.

```
solver->remove_logger(logger.get());
```

Check result

Compute norm of RHS on the device and automatically copy to host

```
auto norm_dim = gko::batch_dim<2>(num_systems, gko::dim<2>(1, 1));
auto host_b_norm = real_vec_type::create(exec->get_master(), norm_dim);
host_b_norm->fill(0.0);
b->compute_norm2(host_b_norm);
```

we need constants on the device

```
auto one = vec_type::create(exec, norm_dim);
one->fill(1.0);
auto neg_one = vec_type::create(exec, norm_dim);
neg_one->fill(-1.0);
```

allocate and compute the residual

```
auto res = vec_type::create(exec, batch_vec_size);
res->copy_from(b);
A->apply(one, x, neg_one, res);
```

allocate and compute residual norm

```
auto host_res_norm = real_vec_type::create(exec->get_master(), norm_dim);
host_res_norm->fill(0.0);
res->compute_norm2(host_res_norm);
auto host_log_resid = gko::make_temporary_clone(
    exec->get_master(), &logger->get_residual_norm());
auto host_log_iters = gko::make_temporary_clone(
    exec->get_master(), &logger->get_num_iterations());
if (print_residuals) {
    std::cout << "Residual norm sqrt(r^T r):\n";
```

"unbatch" converts a batch object into a vector of objects of the corresponding single type, eg. batch::matrix::Dense

```
--> std::vector<Dense>.
auto unb_res = detail::unbatch(host_res_norm.get());
auto unb_bnorm = detail::unbatch(host_b_norm.get());
for (size_type i = 0; i < num_systems; ++i) {
    std::cout << " System no. " << i
              << ": residual norm = " << unb_res[i]->at(0, 0)
              << ", implicit residual norm = "
              << host_log_resid->get_const_data()[i]
              << ", iterations = "
              << host_log_iters->get_const_data()[i] << std::endl;
    const real_type relresnorm =
        unb_res[i]->at(0, 0) / unb_bnorm[i]->at(0, 0);
    if (!(relresnorm <= reduction_factor)) {
        std::cout << "System " << i << " converged only to "
                  << relresnorm << " relative residual." << std::endl;
    }
}
std::cout << "Solver type: "
          << "batch::bicgstab"
          << "\nMatrix size: " << A->get_common_size()
          << "\nNum batch entries: " << A->get_num_batch_items()
          << "\nEntire solve took: " << apply_time / num_reps << " seconds."
          << std::endl;
```

Ginkgo objects are cleaned up automatically; but the "application" still needs to clean up its data in this case.

```
    appl_clean_up(appl_sys, exec);
    return 0;
}
```

Generate the matrix and the vectors. This function emulates the generation of the data from the application.

```
ApplSysData appl_generate_system(const int nrows, const size_type nsystems,
                                std::shared_ptr<gko::Executor> exec)
{
    const int nnz = nrows * 3 - 2;
    std::default_random_engine rgen(15);
    std::normal_distribution<real_type> distb(0.5, 0.1);
    std::vector<real_type> spacings(nsystems * nrows);
    std::generate(spacings.begin(), spacings.end(),
                  [&]() { return distb(rgen); });
    std::vector<value_type> allvalues(nnz * nsystems);
    for (size_type isys = 0; isys < nsystems; isys++) {
        allvalues.at(isys * nnz) = 2.0 / spacings.at(isys * nrows);
        allvalues.at(isys * nnz + 1) = -1.0;
        for (int irow = 0; irow < nrows - 2; irow++) {
            allvalues.at(isys * nnz + 2 + irow * 3) = -1.0;
            allvalues.at(isys * nnz + 2 + irow * 3 + 1) =
                2.0 / spacings.at(isys * nrows + irow + 1);
            allvalues.at(isys * nnz + 2 + irow * 3 + 2) = -1.0;
        }
        allvalues.at(isys * nnz + 2 + (nrows - 2) * 3) = -1.0;
        allvalues.at(isys * nnz + 2 + (nrows - 2) * 3 + 1) =
            2.0 / spacings.at((isys + 1) * nrows - 1);
        assert(isys * nnz + 2 + (nrows - 2) * 3 + 2 == (isys + 1) * nnz);
    }
    std::vector<index_type> rowptrs(nrows + 1);
    rowptrs.at(0) = 0;
    rowptrs.at(1) = 2;
    for (int i = 2; i < nrows; i++) {
        rowptrs.at(i) = rowptrs.at(i - 1) + 3;
    }
    rowptrs.at(nrows) = rowptrs.at(nrows - 1) + 2;
    assert(rowptrs.at(nrows) == nnz);
    std::vector<index_type> colidxs(nnz);
    colidxs.at(0) = 0;
    colidxs.at(1) = 1;
    const int nnz_per_row = 3;
    for (int irow = 1; irow < nrows - 1; irow++) {
        colidxs.at(2 + (irow - 1) * nnz_per_row) = irow - 1;
        colidxs.at(2 + (irow - 1) * nnz_per_row + 1) = irow;
        colidxs.at(2 + (irow - 1) * nnz_per_row + 2) = irow + 1;
    }
    colidxs.at(2 + (nrows - 2) * nnz_per_row) = nrows - 2;
    colidxs.at(2 + (nrows - 2) * nnz_per_row + 1) = nrows - 1;
    assert(2 + (nrows - 2) * nnz_per_row + 1 == nnz - 1);
    std::vector<value_type> allb(nrows * nsystems);
    for (size_type isys = 0; isys < nsystems; isys++) {
        const value_type bval = distb(rgen);
        std::fill(allb.begin() + isys * nrows,
                  allb.begin() + (isys + 1) * nrows, bval);
    }
    index_type* const row_ptrs = exec->alloc<index_type>(nrows + 1);
    exec->copy_from(exec->get_master().get(), static_cast<size_type>(nrows + 1),
                   row_ptrs.data(), row_ptrs);
}
```

```

    index_type* const col_idx = exec->alloc<index_type>(nnz);
    exec->copy_from(exec->get_master().get(), static_cast<size_type>(nnz),
                  col_idx.data(), col_idx);
    value_type* const all_values = exec->alloc<value_type>(nsystems * nnz);
    exec->copy_from(exec->get_master().get(), nsystems * nnz, all_values.data(),
                  all_values);
    value_type* const all_b = exec->alloc<value_type>(nsystems * nrows);
    exec->copy_from(exec->get_master().get(), nsystems * nrows, all_b.data(),
                  all_b);
    return {nsystems, nrows, nnz, row_ptrs, col_idx, all_values, all_b};
}
void appl_clean_up(ApplSysData& appl_data, std::shared_ptr<gko::Executor> exec)
{

```

In general, the application would control non-const pointers; the const casts below would not be needed.

```

    exec->free(const_cast<index_type*>(appl_data.row_ptrs));
    exec->free(const_cast<index_type*>(appl_data.col_idx));
    exec->free(const_cast<value_type*>(appl_data.all_values));
    exec->free(const_cast<value_type*>(appl_data.all_rhs));
}

```

Results

The following is the expected result on the reference executor:

```

Solver type:  batch::bicgstab
Matrix size:  (32, 32)
Num batch entries:  2
Entire solve took:  2.18558e-05 seconds.

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iostream>
#include <map>
#include <random>
#include <string>
#include <ginkgo/ginkgo.hpp>
using value_type = double;
using real_type = gko::remove_complex<value_type>;
using index_type = int;
using size_type = gko::size_type;
using vec_type = gko::batch::MultiVector<value_type>;
using real_vec_type = gko::batch::MultiVector<real_type>;
using mtx_type = gko::batch::matrix::Csr<value_type, index_type>;
using bicgstab = gko::batch::solver::Bicgstab<value_type>;
namespace detail {
template <typename InputType>
auto unbatch(const InputType* batch_object)
{
    auto unbatched_mats =
        std::vector<std::unique_ptr<typename InputType::unbatch_type>>{};
    for (size_type b = 0; b < batch_object->get_num_batch_items(); ++b) {
        unbatched_mats.emplace_back(
            batch_object->create_const_view_for_item(b)->clone());
    }
    return unbatched_mats;
}
} // namespace detail
struct ApplSysData {
    size_type nsystems;
    int nrows;
    int nnz;
    const index_type* row_ptrs;
    const index_type* col_idx;
    const value_type* all_values;
    const value_type* all_rhs;
};
/*
 * Generates a batch of tridiagonal systems.
 *
 * @param nrows  Number of rows in each system.
 * @param nsystems  Number of systems in the batch.

```

```

* @param exec The device executor to use for the solver.
* @note Normally, the application may not deal with Ginkgo executors, nor do
* we need it to. Here, we use the executor for backend-independent device
* memory allocation. The application, for example, might assume Hip (for AMD
* GPUs) and use 'hipMalloc' directly.
*/
ApplSysData appl_generate_system(const int nrows, const size_type nsystems,
                                std::shared_ptr<gko::Executor> exec);
void appl_clean_up(ApplSysData& appl_data, std::shared_ptr<gko::Executor> exec);
int main(int argc, char* argv[])
{
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0]
            << " [executor] [num_systems] [num_rows] [print_residuals] "
            << "[num_reps]"
            << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                 gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                 gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const size_type num_systems = argc >= 3 ? std::atoi(argv[2]) : 2;
    const int num_rows = argc >= 4 ? std::atoi(argv[3]) : 32; // per system
    const bool print_residuals =
        argc >= 5 ? (std::string(argv[4]) == "true") : false;
    const int num_reps = argc >= 6 ? std::atoi(argv[5]) : 20;
    auto appl_sys = appl_generate_system(num_rows, num_systems, exec);
    auto batch_mat_size =
        gko::batch_dim<2>(num_systems, gko::dim<2>(num_rows, num_rows));
    auto batch_vec_size =
        gko::batch_dim<2>(num_systems, gko::dim<2>(num_rows, 1));
    auto vals_view = gko::array<value_type>::const_view(
        exec, num_systems * appl_sys.nnz, appl_sys.all_values);
    auto rowptrs_view = gko::array<index_type>::const_view(exec, num_rows + 1,
        appl_sys.row_ptrs);
    auto colidxs_view = gko::array<index_type>::const_view(exec, appl_sys.nnz,
        appl_sys.col_idxs);
    auto A = gko::share(mtx_type::create_const(
        exec, batch_mat_size, std::move(vals_view), std::move(colidxs_view),
        std::move(rowptrs_view)));
    auto b_view = gko::array<value_type>::const_view(
        exec, num_systems * num_rows, appl_sys.all_rhs);
    auto b = vec_type::create_const(exec, batch_vec_size, std::move(b_view));
    auto x = vec_type::create(exec);
    auto host_x = vec_type::create(exec->get_master(), batch_vec_size);
    for (size_type isys = 0; isys < num_systems; isys++) {
        for (int irow = 0; irow < num_rows; irow++) {
            host_x->at(isys, irow, 0) = gko::zero<value_type>();
        }
    }
    x->copy_from(host_x.get());
    const real_type reduction_factor{1e-10};
    auto solver =
        bicgstab::build()
            .with_max_iterations(500)
            .with_tolerance(reduction_factor)
            .with_tolerance_type(gko::batch::stop::tolerance_type::relative)
            .on(exec)
            ->generate(A);
    std::shared_ptr<const gko::batch::log::BatchConvergence<value_type>
    logger = gko::batch::log::BatchConvergence<value_type>::create();
    solver->add_logger(logger);
    auto x_clone = gko::clone(x);
    for (int i = 0; i < 3; ++i) {
        x_clone->copy_from(x.get());
        solver->apply(b, x_clone);
    }
    double apply_time = 0.0;
    for (int i = 0; i < num_reps; ++i) {

```

```

    x_clone->copy_from(x.get());
    exec->synchronize();
    std::chrono::steady_clock::time_point t1 =
        std::chrono::steady_clock::now();
    solver->apply(b, x_clone);
    exec->synchronize();
    std::chrono::steady_clock::time_point t2 =
        std::chrono::steady_clock::now();
    auto time_span =
        std::chrono::duration_cast<std::chrono::duration<double>>(t2 - t1);
    apply_time += time_span.count();
}
x->copy_from(x_clone.get());
solver->remove_logger(logger.get());
auto norm_dim = gko::batch_dim<2>(num_systems, gko::dim<2>(1, 1));
auto host_b_norm = real_vec_type::create(exec->get_master(), norm_dim);
host_b_norm->fill(0.0);
b->compute_norm2(host_b_norm);
auto one = vec_type::create(exec, norm_dim);
one->fill(1.0);
auto neg_one = vec_type::create(exec, norm_dim);
neg_one->fill(-1.0);
auto res = vec_type::create(exec, batch_vec_size);
res->copy_from(b);
A->apply(one, x, neg_one, res);
auto host_res_norm = real_vec_type::create(exec->get_master(), norm_dim);
host_res_norm->fill(0.0);
res->compute_norm2(host_res_norm);
auto host_log_resid = gko::make_temporary_clone(
    exec->get_master(), &logger->get_residual_norm());
auto host_log_iters = gko::make_temporary_clone(
    exec->get_master(), &logger->get_num_iterations());
if (print_residuals) {
    std::cout << "Residual norm sqrt(r^T r):\n";
    auto unb_res = detail::unbatch(host_res_norm.get());
    auto unb_bnorm = detail::unbatch(host_b_norm.get());
    for (size_type i = 0; i < num_systems; ++i) {
        std::cout << " System no. " << i
            << " : residual norm = " << unb_res[i]->at(0, 0)
            << " , implicit residual norm = "
            << host_log_resid->get_const_data()[i]
            << " , iterations = "
            << host_log_iters->get_const_data()[i] << std::endl;
        const real_type relresnorm =
            unb_res[i]->at(0, 0) / unb_bnorm[i]->at(0, 0);
        if (!relresnorm <= reduction_factor) {
            std::cout << "System " << i << " converged only to "
                << relresnorm << " relative residual." << std::endl;
        }
    }
}
std::cout << "Solver type: "
    << "batch::bicgstab"
    << "\nMatrix size: " << A->get_common_size()
    << "\nNum batch entries: " << A->get_num_batch_items()
    << "\nEntire solve took: " << apply_time / num_reps << " seconds."
    << std::endl;
appl_clean_up(appl_sys, exec);
return 0;
}
ApplSysData appl_generate_system(const int nrows, const size_type nsystems,
    std::shared_ptr<gko::Executor> exec)
{
    const int nnz = nrows * 3 - 2;
    std::default_random_engine rgen(15);
    std::normal_distribution<real_type> distb(0.5, 0.1);
    std::vector<real_type> spacings(nsystems * nrows);
    std::generate(spacings.begin(), spacings.end(),
        [&]() { return distb(rgen); });
    std::vector<value_type> allvalues(nnz * nsystems);
    for (size_type isys = 0; isys < nsystems; isys++) {
        allvalues.at(isys * nnz) = 2.0 / spacings.at(isys * nrows);
        allvalues.at(isys * nnz + 1) = -1.0;
        for (int irow = 0; irow < nrows - 2; irow++) {
            allvalues.at(isys * nnz + 2 + irow * 3) = -1.0;
            allvalues.at(isys * nnz + 2 + irow * 3 + 1) =
                2.0 / spacings.at(isys * nrows + irow + 1);
            allvalues.at(isys * nnz + 2 + irow * 3 + 2) = -1.0;
        }
        allvalues.at(isys * nnz + 2 + (nrows - 2) * 3) = -1.0;
        allvalues.at(isys * nnz + 2 + (nrows - 2) * 3 + 1) =
            2.0 / spacings.at((isys + 1) * nrows - 1);
        assert(isys * nnz + 2 + (nrows - 2) * 3 + 2 == (isys + 1) * nnz);
    }
    std::vector<index_type> rowptrs(nrows + 1);
    rowptrs.at(0) = 0;
    rowptrs.at(1) = 2;

```

```

    for (int i = 2; i < nrows; i++) {
        rowptrs.at(i) = rowptrs.at(i - 1) + 3;
    }
    rowptrs.at(nrows) = rowptrs.at(nrows - 1) + 2;
    assert(rowptrs.at(nrows) == nnz);
    std::vector<index_type> colidxs(nnz);
    colidxs.at(0) = 0;
    colidxs.at(1) = 1;
    const int nnz_per_row = 3;
    for (int irow = 1; irow < nrows - 1; irow++) {
        colidxs.at(2 + (irow - 1) * nnz_per_row) = irow - 1;
        colidxs.at(2 + (irow - 1) * nnz_per_row + 1) = irow;
        colidxs.at(2 + (irow - 1) * nnz_per_row + 2) = irow + 1;
    }
    colidxs.at(2 + (nrows - 2) * nnz_per_row) = nrows - 2;
    colidxs.at(2 + (nrows - 2) * nnz_per_row + 1) = nrows - 1;
    assert(2 + (nrows - 2) * nnz_per_row + 1 == nnz - 1);
    std::vector<value_type> allb(nrows * nsystems);
    for (size_type isys = 0; isys < nsystems; isys++) {
        const value_type bval = distb(rgen);
        std::fill(allb.begin() + isys * nrows,
                  allb.begin() + (isys + 1) * nrows, bval);
    }
    index_type* const row_ptrs = exec->alloc<index_type>(nrows + 1);
    exec->copy_from(exec->get_master().get(), static_cast<size_type>(nrows + 1),
                   rowptrs.data(), row_ptrs);
    index_type* const col_idxxs = exec->alloc<index_type>(nnz);
    exec->copy_from(exec->get_master().get(), static_cast<size_type>(nnz),
                   colidxs.data(), col_idxxs);
    value_type* const all_values = exec->alloc<value_type>(nsystems * nnz);
    exec->copy_from(exec->get_master().get(), nsystems * nnz, all_values.data(),
                   all_values);
    value_type* const all_b = exec->alloc<value_type>(nsystems * nrows);
    exec->copy_from(exec->get_master().get(), nsystems * nrows, allb.data(),
                   all_b);
    return {nsystems, nrows, nnz, row_ptrs, col_idxxs, all_values, all_b};
}

void appl_clean_up(ApplSysData& appl_data, std::shared_ptr<gko::Executor> exec)
{
    exec->free(const_cast<index_type*>(appl_data.row_ptrs));
    exec->free(const_cast<index_type*>(appl_data.col_idxxs));
    exec->free(const_cast<value_type*>(appl_data.all_values));
    exec->free(const_cast<value_type*>(appl_data.all_rhs));
}

```


Chapter 10

The cb-gmres program

The CB-GMRES solver example..

Introduction

About the example

This example showcases the usage of the Ginkgo solver CB-GMRES (Compressed Basis GMRES). A small system is solved with two un-preconditioned CB-GMRES solvers:

1. without compressing the krylov basis; it uses double precision for both the matrix and the krylov basis, and
2. with a compression of the krylov basis; it uses double precision for the matrix and all arithmetic operations, while using single precision for the storage of the krylov basis

Both solves are timed and the residual norm of each solution is computed to show that both solutions are correct.

The commented program

Make sure all previous executor operations have finished before starting the time

```
exec->synchronize();  
auto tic = std::chrono::steady_clock::now();  
solver->apply(b, x_copy);
```

Make sure all computations are done before stopping the time

```
exec->synchronize();  
auto tac = std::chrono::steady_clock::now();  
duration += std::chrono::duration<double>(tac - tic).count();  
}
```

Copy the solution back to x, so the caller has the result

```
x->copy_from(x_copy);  
return duration / static_cast<double>(repeats);  
}  
int main(int argc, char* argv[])  
{
```

Use some shortcuts. In Ginkgo, vectors are seen as a [gko::matrix::Dense](#) with one column/one row. The advantage of this concept is that using multiple vectors is a now a natural extension of adding columns/rows are necessary.

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
```

The `gko::matrix::Csr` class is used here, but any other matrix class such as `gko::matrix::Coo`, `gko::matrix::Hybrid`, `gko::matrix::Ell` or `gko::matrix::Sellp` could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The `gko::solver::CbGmres` is used here, but any other solver class can also be used.

```
using cb_gmres = gko::solver::CbGmres<ValueType>;
```

Print the ginkgo version information.

```
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;
    std::exit(-1);
}
```

Map which generates the appropriate executor

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }
    };
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Note: this matrix is copied from "SOURCE_DIR/matrices" instead of from the local directory. For details, see "examples/cb-gmres/CMakeLists.txt"

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create a uniform right-hand side with a norm2 of 1 on the host (norm2(b) == 1), followed by copying it to the actual executor (to make sure it also works for GPUs)

```
const auto A_size = A->get_size();
auto b_host = vec::create(exec->get_master(), gko::dim<2>{A_size[0], 1});
for (gko::size_type i = 0; i < A_size[0]; ++i) {
    b_host->at(i, 0) =
        ValueType{1} / std::sqrt(static_cast<ValueType>(A_size[0]));
}
auto b_norm = gko::initialize<real_vec>({0.0}, exec);
b_host->compute_norm2(b_norm);
auto b = clone(exec, b_host);
```

As an initial guess, use the right-hand side

```
auto x_keep = clone(b);
auto x_reduce = clone(x_keep);
const RealValueType reduction_factor{1e-6};
```

Generate two solver factories: `_keep` uses the same precision for the krylov basis as the matrix, and `_reduce` uses one precision below it. If `ValueType` is double, then `_reduce` uses float as the krylov basis storage type

```
auto solver_gen_keep =
    cb_gmres::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
                      gko::stop::ResidualNorm<ValueType>::build()
                          .with_baseline(gko::stop::mode::rhs_norm)
                          .with_reduction_factor(reduction_factor))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb_gmres::storage_precision::keep)
        .on(exec);
```

```

auto solver_gen_reduce =
    cb_gmres::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
                       gko::stop::ResidualNorm<ValueType>::build()
                           .with_baseline(gko::stop::mode::rhs_norm)
                           .with_reduction_factor(reduction_factor))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb_gmres::storage_precision::reduced)
        .on(exec);

```

Generate the actual solver from the factory and the matrix.

```

auto solver_keep = solver_gen_keep->generate(A);
auto solver_reduce = solver_gen_reduce->generate(A);

```

Solve both system and measure the time for each.

```

auto time_keep =
    measure_solve_time_in_s(exec, solver_keep.get(), b.get(), x_keep.get());
auto time_reduce = measure_solve_time_in_s(exec, solver_reduce.get(),
                                           b.get(), x_reduce.get());

```

Make sure the output is in scientific notation for easier comparison

```
std::cout << std::scientific;
```

Note: The time might not be significantly different since the matrix is quite small

```

std::cout << "Solve time without compression: " << time_keep << " s\n"
          << "Solve time with compression:      " << time_reduce << " s\n";

```

To measure if your solution has actually converged, the error of the solution is measured. `one`, `neg_one` are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, the (advanced) `apply` method is used.

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res_norm_keep = gko::initialize<real_vec>({0.0}, exec);
auto res_norm_reduce = gko::initialize<real_vec>({0.0}, exec);
auto tmp = gko::clone(b);

```

`tmp = Ax - tmp`

```

A->apply(one, x_keep, neg_one, tmp);
tmp->compute_norm2(res_norm_keep);
std::cout << "\nResidual norm without compression:\n";
write(std::cout, res_norm_keep);
tmp->copy_from(b);
A->apply(one, x_reduce, neg_one, tmp);
tmp->compute_norm2(res_norm_reduce);
std::cout << "\nResidual norm with compression:\n";
write(std::cout, res_norm_reduce);
}

```

Results

The following is the expected result:

```

Solve time without compression: 1.842690e-04 s
Solve time with compression:    1.589936e-04 s
Residual norm without compression:
%%MatrixMarket matrix array real general
1 1
2.430544e-07
Residual norm with compression:
%%MatrixMarket matrix array real general
1 1
3.437257e-07

```

Comments about programming and debugging

The plain program

```
#include <chrono>
#include <cmath>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
double measure_solve_time_in_s(std::shared_ptr<const gko::Executor> exec,
                               gko::LinOp* solver, const gko::LinOp* b,
                               gko::LinOp* x)
{
    constexpr int repeats{5};
    double duration{0};
    auto x_copy = clone(x);
    for (int i = 0; i < repeats; ++i) {
        if (i != 0) {
            x_copy->copy_from(x);
        }
        exec->synchronize();
        auto tic = std::chrono::steady_clock::now();
        solver->apply(b, x_copy);
        exec->synchronize();
        auto tac = std::chrono::steady_clock::now();
        duration += std::chrono::duration<double>(tac - tic).count();
    }
    x->copy_from(x_copy);
    return duration / static_cast<double>(repeats);
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cb_gmres = gko::solver::CbGmres<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    const auto A_size = A->get_size();
    auto b_host = vec::create(exec->get_master(), gko::dim<2>{A_size[0], 1});
    for (gko::size_type i = 0; i < A_size[0]; ++i) {
        b_host->at(i, 0) =
            ValueType{1} / std::sqrt(static_cast<ValueType>(A_size[0]));
    }
    auto b_norm = gko::initialize<real_vec>({0.0}, exec);
    b_host->compute_norm2(b_norm);
    auto b = clone(exec, b_host);
    auto x_keep = clone(b);
    auto x_reduce = clone(x_keep);
    const RealValueType reduction_factor{1e-6};
    auto solver_gen_keep =
        cb_gmres::build()
            .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
                          gko::stop::ResidualNorm<ValueType>::build()
                              .with_baseline(gko::stop::mode::rhs_norm))
```

```

        .with_reduction_factor(reduction_factor))
    .with_krylov_dim(100u)
    .with_storage_precision(
        gko::solver::cb_gmres::storage_precision::keep)
    .on(exec);
auto solver_gen_reduce =
    cb_gmres::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_baseline(gko::stop::mode::rhs_norm)
                .with_reduction_factor(reduction_factor))
        .with_krylov_dim(100u)
        .with_storage_precision(
            gko::solver::cb_gmres::storage_precision::reduce1)
        .on(exec);
auto solver_keep = solver_gen_keep->generate(A);
auto solver_reduce = solver_gen_reduce->generate(A);
auto time_keep =
    measure_solve_time_in_s(exec, solver_keep.get(), b.get(), x_keep.get());
auto time_reduce = measure_solve_time_in_s(exec, solver_reduce.get(),
    b.get(), x_reduce.get());

std::cout << std::scientific;
std::cout << "Solve time without compression: " << time_keep << " s\n"
    << "Solve time with compression: " << time_reduce << " s\n";
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res_norm_keep = gko::initialize<real_vec>({0.0}, exec);
auto res_norm_reduce = gko::initialize<real_vec>({0.0}, exec);
auto tmp = gko::clone(b);
A->apply(one, x_keep, neg_one, tmp);
tmp->compute_norm2(res_norm_keep);
std::cout << "\nResidual norm without compression:\n";
write(std::cout, res_norm_keep);
tmp->copy_from(b);
A->apply(one, x_reduce, neg_one, tmp);
tmp->compute_norm2(res_norm_reduce);
std::cout << "\nResidual norm with compression:\n";
write(std::cout, res_norm_reduce);
}

```


Chapter 11

The custom-logger program

The simple solver with a custom logger example..

This example depends on simple-solver, simple-solver-logging, minimal-cuda-solver.

Introduction

The custom-logger example shows how Ginkgo's API can be leveraged to implement application-specific callbacks for Ginkgo's events. This is the most basic way of extending Ginkgo and a good first step for any application developer who wants to adapt Ginkgo to his specific needs.

Ginkgo's `gko::log::Logger` abstraction provides hooks to the events that happen during the library execution. These hooks concern any low-level event such as memory allocations, deallocations, copies and kernel launches up to high-level events such as linear operator applications and completion of solver iterations.

In this example, a simple logger is implemented to track the solver's recurrent residual norm and compute the true residual norm. At the end of the solver execution, a comparison table is shown on-screen.

About the example

Each example has the following sections:

1. **Introduction:** This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

```
    return mtx->get_executor()->copy_val_to_host(mtx->get_const_values());
}
```

Utility function which computes the norm of a Ginkgo `gko::matrix::Dense` vector.

```
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(
    const gko::matrix::Dense<ValueType>* b)
{
```

Get the executor of the vector

```
auto exec = b->get_executor();
```

Initialize a result scalar containing the value 0.0.

```
auto b_norm =
    gko::initialize<gko::matrix::Dense<gko::remove_complex<ValueType>>>(
        {0.0}, exec);
```

Use the dense `compute_norm2` function to compute the norm.

```
b->compute_norm2(b_norm);
```

Use the other utility function to return the norm contained in `b_norm`

```
    return get_first_element(b_norm.get());
}
```

Custom logger class which intercepts the residual norm scalar and solution vector in order to print a table of real vs recurrent (internal to the solvers) residual norms.

```
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    using RealValueType = gko::remove_complex<ValueType>;
```

Output the logger's data in a table format

```
void write() const
{
```

Print a header for the table

```
std::cout << "Recurrent vs true vs implicit residual norm:"
    << std::endl;
std::cout << '|' << std::setw(10) << "Iteration" << '|' << std::setw(25)
    << "Recurrent Residual Norm" << '|' << std::setw(25)
    << "True Residual Norm" << '|' << std::setw(25)
    << "Implicit Residual Norm" << '|' << std::endl;
```

Print a separation line. Note that for creating 10 characters `std::setw()` should be set to 11.

```
std::cout << '|' << std::setfill('-') << std::setw(11) << '|'
    << std::setw(26) << '|' << std::setw(26) << '|'
    << std::setw(26) << '|' << std::setfill(' ') << std::endl;
```

Print the data one by one in the form

```
std::cout << std::scientific;
for (std::size_t i = 0; i < iterations.size(); i++) {
    std::cout << '|' << std::setw(10) << iterations[i] << '|'
        << std::setw(25) << recurrent_norms[i] << '|'
        << std::setw(25) << real_norms[i] << '|' << std::setw(25)
        << implicit_norms[i] << '|' << std::endl;
}
```

`std::defaultfloat` could be used here but some compilers do not support it properly, e.g. the Intel compiler

```
std::cout.unsetf(std::ios_base::floatfield);
```

Print a separation line

```
std::cout << '|' << std::setfill('-') << std::setw(11) << '|'
    << std::setw(26) << '|' << std::setw(26) << '|'
    << std::setw(26) << '|' << std::setfill(' ') << std::endl;
}
using gko_dense = gko::matrix::Dense<ValueType>;
using gko_real_dense = gko::matrix::Dense<RealValueType>;
```

Customize the logging hook which is called everytime an iteration is completed

```
void on_iteration_complete(const gko::LinOp* solver, const gko::LinOp* b,
    const gko::LinOp* solution,
    const gko::size_type& iteration,
    const gko::LinOp* residual,
```



```

        const gko::LinOp* residual_norm,
        const gko::LinOp* implicit_sq_residual_norm,
        const gko::array<gko::stopping_status>*,
        bool) const override
{

```

If the solver shares a residual norm, log its value

```

if (residual_norm) {
    auto dense_norm = gko::as<gko_real_dense>(residual_norm);

```

Add the norm to the `recurrent_norms` vector

```

recurrent_norms.push_back(get_first_element(dense_norm));

```

Otherwise, use the recurrent residual vector

```

} else {
    auto dense_residual = gko::as<gko_dense>(residual);

```

Compute the residual vector's norm

```

auto norm = compute_norm(dense_residual);

```

Add the computed norm to the `recurrent_norms` vector

```

recurrent_norms.push_back(norm);
}

```

If the solver shares the current solution vector

```

if (solution) {

```

Extract the matrix from the solver

```

auto matrix = gko::as<gko::solver::detail::SolverBaseLinOp>(solver)
    ->get_system_matrix();

```

Store the matrix's executor

```

auto exec = matrix->get_executor();

```

Create a scalar containing the value 1.0

```

auto one = gko::initialize<gko_dense>({1.0}, exec);

```

Create a scalar containing the value -1.0

```

auto neg_one = gko::initialize<gko_dense>({-1.0}, exec);

```

Instantiate a temporary result variable

```

auto res = gko::as<gko_dense>(gko::clone(b));

```

Compute the real residual vector by calling `apply` on the system matrix

```

matrix->apply(one, solution, neg_one, res);

```

Compute the norm of the residual vector and add it to the `real_norms` vector

```

real_norms.push_back(compute_norm(res.get()));
} else {

```

Add to the `real_norms` vector the value -1.0 if it could not be computed

```

real_norms.push_back(-1.0);
}
if (implicit_sq_residual_norm) {
    auto dense_norm =
        gko::as<gko_real_dense>(implicit_sq_residual_norm);

```

Add the norm to the `implicit_norms` vector

```

implicit_norms.push_back(std::sqrt(get_first_element(dense_norm)));
} else {

```

Add to the `implicit_norms` vector the value -1.0 if it could not be computed

```

implicit_norms.push_back(-1.0);
}

```

Add the current iteration number to the `iterations` vector

```

iterations.push_back(iteration);
}

```

Construct the logger

```

    ResidualLogger()
    :   gko::log::Logger(gko::log::Logger::iteration_complete_mask)
    {}
private:

```

Vector which stores all the recurrent residual norms

```
mutable std::vector<RealValueType> recurrent_norms{};
```

Vector which stores all the real residual norms

```
mutable std::vector<RealValueType> real_norms{};
```

Vector which stores all the implicit residual norms

```
mutable std::vector<RealValueType> implicit_norms{};
```

Vector which stores all the iteration numbers

```

    mutable std::vector<std::size_t> iterations{};
};
int main(int argc, char* argv[])
{

```

Use some shortcuts. In Ginkgo, vectors are seen as a `gko::matrix::Dense` with one column/one row. The advantage of this concept is that using multiple vectors is now a natural extension of adding columns/rows are necessary.

```

using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;

```

The `gko::matrix::Csr` class is used here, but any other matrix class such as `gko::matrix::Coo`, `gko::matrix::Hybrid`, `gko::matrix::Ell` or `gko::matrix::Sellp` could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The `gko::solver::Cg` is used here, but any other solver class can also be used.

```
using cg = gko::solver::Cg<ValueType>;
```

Print the ginkgo version information.

```
std::cout << gko::version_info::get() << std::endl;
```

Where do you want to run your solver ?

The `gko::Executor` class is one of the cornerstones of Ginkgo. Currently, we have support for an `gko::OmpExecutor`, which uses OpenMP multi-threading in most of its kernels, a `gko::ReferenceExecutor`, a single threaded specialization of the OpenMP executor and a `gko::CudaExecutor` which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/ routines within Ginkgo should automatically work and run on the executor with any other changes.

```

if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";

```

Figure out where to run the code

```

std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
    [] {
        return gko::CudaExecutor::create(0,
                                         gko::OmpExecutor::create());
    }},
    {"hip",
    [] {
        return gko::HipExecutor::create(0, gko::OmpExecutor::create());
    }},
    {"dpcpp",
    [] {
        return gko::DpcppExecutor::create(0,
                                         gko::OmpExecutor::create());
    }},
    {"reference", [] { return gko::ReferenceExecutor::create(); } }
};

```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Reading your data and transfer to the proper device.

Read the matrix, right hand side and the initial solution using the read function.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. `gko::share` and `gko::give` are the functions that you would need to use.

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor = 1e-7;
```

Creating the solver

Generate the `gko::solver` factory. Ginkgo uses the concept of Factories to build solvers with certain properties. Observe the Fluent interface used here. Here a cg solver is generated with a stopping criteria of maximum iterations of 20 and a residual norm reduction of $1e-15$. You also observe that the stopping criteria(`gko::stop`) are also generated from factories using their build methods. You need to specify the executors which each of the object needs to be built on.

```
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor))
        .on(exec);
```

Instantiate a ResidualLogger logger.

```
auto logger = std::make_shared<ResidualLogger<ValueType>>();
```

Add the previously created logger to the solver factory. The logger will be automatically propagated to all solvers created from this factory.

```
solver_gen->add_logger(logger);
```

Generate the solver from the matrix. The solver factory built in the previous step takes a "matrix"(a `gko::LinOp` to be more general) as an input. In this case we provide it with a full matrix that we previously read, but as the solver only effectively uses the `apply()` method within the provided "matrix" object, you can effectively create a `gko::LinOp` class with your own `apply` implementation to accomplish more tasks. We will see an example of how this can be done in the custom-matrix-format example

```
auto solver = solver_gen->generate(A);
```

Finally, solve the system. The solver, being a `gko::LinOp`, can be applied to a right hand side, `b` to obtain the solution, `x`.

```
solver->apply(b, x);
```

Print the solution to the command line.

```
std::cout << "Solution (x):\n";
write(std::cout, x);
```

Print the table of the residuals obtained from the logger

```
logger->write();
```

To measure if your solution has actually converged, you can measure the error of the solution. `one`, `neg_one` are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, all you need to do is call the `apply` method, which in this case is an `spmv` and equivalent to the LAPACK `z_spmv` routine. Finally, you compute the euclidean 2-norm with the `compute_norm2` function.

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```

Results

The following is the expected result:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Recurrent vs true vs implicit residual norm:
Iteration	Recurrent Residual Norm	True Residual Norm	Implicit Residual Norm
0	4.358899e+00	4.358899e+00	4.358899e+00
1	2.304548e+00	2.304548e+00	2.304548e+00
2	1.467706e+00	1.467706e+00	1.467706e+00
3	9.848751e-01	9.848751e-01	9.848751e-01
4	7.418330e-01	7.418330e-01	7.418330e-01
5	5.136231e-01	5.136231e-01	5.136231e-01
6	3.841650e-01	3.841650e-01	3.841650e-01
7	3.164394e-01	3.164394e-01	3.164394e-01
8	2.277088e-01	2.277088e-01	2.277088e-01
9	1.703121e-01	1.703121e-01	1.703121e-01
10	9.737220e-02	9.737220e-02	9.737220e-02
11	6.168306e-02	6.168306e-02	6.168306e-02
12	4.541231e-02	4.541231e-02	4.541231e-02
13	3.195304e-02	3.195304e-02	3.195304e-02
14	1.616058e-02	1.616058e-02	1.616058e-02
15	6.570152e-03	6.570152e-03	6.570152e-03
16	2.643669e-03	2.643669e-03	2.643669e-03
17	8.588089e-04	8.588089e-04	8.588089e-04
18	2.864613e-04	2.864613e-04	2.864613e-04
19	1.641952e-15	2.107881e-15	1.641952e-15
-----	-----	-----	-----
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
2.10788e-15
```

Comments about programming and debugging

The plain program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <map>
#include <iomanip>
#include <ios>
#include <iostream>
#include <string>
#include <vector>
template <typename ValueType>
ValueType get_first_element(const gko::matrix::Dense<ValueType>* mtx)
{
    return mtx->get_executor()->copy_val_to_host(mtx->get_const_values());
}
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(
    const gko::matrix::Dense<ValueType>* b)
{
    auto exec = b->get_executor();
    auto b_norm =
        gko::initialize<gko::matrix::Dense<gko::remove_complex<ValueType>>>(
            {0.0}, exec);
```

```

        b->compute_norm2(b_norm);
        return get_first_element(b_norm.get());
    }
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    using RealValueType = gko::remove_complex<ValueType>;
    void write() const
    {
        std::cout << "Recurrent vs true vs implicit residual norm:"
        << std::endl;
        std::cout << '|' << std::setw(10) << "Iteration" << '|' << std::setw(25)
        << "Recurrent Residual Norm" << '|' << std::setw(25)
        << "True Residual Norm" << '|' << std::setw(25)
        << "Implicit Residual Norm" << '|' << std::endl;
        std::cout << '|' << std::setfill('-') << std::setw(11) << '|'
        << std::setw(26) << '|' << std::setw(26) << '|'
        << std::setw(26) << '|' << std::setfill(' ') << std::endl;
        std::cout << std::scientific;
        for (std::size_t i = 0; i < iterations.size(); i++) {
            std::cout << '|' << std::setw(10) << iterations[i] << '|'
            << std::setw(25) << recurrent_norms[i] << '|'
            << std::setw(25) << real_norms[i] << '|' << std::setw(25)
            << implicit_norms[i] << '|' << std::endl;
        }
        std::cout.unsetf(std::ios_base::floatfield);
        std::cout << '|' << std::setfill('-') << std::setw(11) << '|'
        << std::setw(26) << '|' << std::setw(26) << '|'
        << std::setw(26) << '|' << std::setfill(' ') << std::endl;
    }
    using gko_dense = gko::matrix::Dense<ValueType>;
    using gko_real_dense = gko::matrix::Dense<RealValueType>;
    void on_iteration_complete(const gko::LinOp* solver, const gko::LinOp* b,
                              const gko::LinOp* solution,
                              const gko::size_type& iteration,
                              const gko::LinOp* residual,
                              const gko::LinOp* residual_norm,
                              const gko::LinOp* implicit_sq_residual_norm,
                              const gko::array<gko::stopping_status>*,
                              bool) const override
    {
        if (residual_norm) {
            auto dense_norm = gko::as<gko_real_dense>(residual_norm);
            recurrent_norms.push_back(get_first_element(dense_norm));
        } else {
            auto dense_residual = gko::as<gko_dense>(residual);
            auto norm = compute_norm(dense_residual);
            recurrent_norms.push_back(norm);
        }
        if (solution) {
            auto matrix = gko::as<gko::solver::detail::SolverBaseLinOp>(solver)
            ->get_system_matrix();
            auto exec = matrix->get_executor();
            auto one = gko::initialize<gko_dense>({1.0}, exec);
            auto neg_one = gko::initialize<gko_dense>({-1.0}, exec);
            auto res = gko::as<gko_dense>(gko::clone(b));
            matrix->apply(one, solution, neg_one, res);
            real_norms.push_back(compute_norm(res.get()));
        } else {
            real_norms.push_back(-1.0);
        }
        if (implicit_sq_residual_norm) {
            auto dense_norm =
                gko::as<gko_real_dense>(implicit_sq_residual_norm);
            implicit_norms.push_back(std::sqrt(get_first_element(dense_norm)));
        } else {
            implicit_norms.push_back(-1.0);
        }
        iterations.push_back(iteration);
    }
    ResidualLogger()
        : gko::log::Logger(gko::log::Logger::iteration_complete_mask)
    {}
private:
    mutable std::vector<RealValueType> recurrent_norms{};
    mutable std::vector<RealValueType> real_norms{};
    mutable std::vector<RealValueType> implicit_norms{};
    mutable std::vector<std::size_t> iterations{};
};
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;

```

```

std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor = 1e-7;
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                       gko::stop::ResidualNorm<ValueType>::build()
                           .with_reduction_factor(reduction_factor))
        .on(exec);
auto logger = std::make_shared<ResidualLogger<ValueType>>();
solver_gen->add_logger(logger);
auto solver = solver_gen->generate(A);
solver->apply(b, x);
std::cout << "Solution (x):\n";
write(std::cout, x);
logger->write();
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}

```

Chapter 12

The custom-matrix-format program

The custom matrix format example..

This example depends on simple-solver, poisson-solver, three-pt-stencil-solver, .

Introduction

This example solves a 1D Poisson equation:

$$\begin{aligned} u &: [0, 1] \rightarrow \mathbb{R} \\ u'' &= f \\ u(0) &= u_0 \\ u(1) &= u_1 \end{aligned}$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

$$\begin{aligned} u(x+h) &= u(x) + u'(x)h + \frac{1}{2}u''(x)h^2 + O(h^3) \\ u(x-h) &= u(x) - u'(x)h + \frac{1}{2}u''(x)h^2 + O(h^3) \\ \hline -u(x-h) + 2u(x) - u(x+h) &= -f(x)h^2 + O(h^3) \end{aligned}$$

For an equidistant grid with K "inner" discretization points $x_1, \dots, x_K$, and step size $h = 1/(K+1)$, the formula produces a system of linear equations

$$\begin{aligned} 2u_1 - u_2 &= -f_1 h^2 + u_0 \\ -u_{k-1} + 2u_k - u_{k+1} &= -f_k h^2, k = 2, \dots, K-1 \\ -u_{K-1} + 2u_K &= -f_K h^2 + u_1 \end{aligned}$$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function f is set to $f(x) = 6x$ (making the solution $u(x) = x^3$), but that can be changed in the `main` function.

The intention of this example is to show how a custom linear operator can be created and integrated into Ginkgo to achieve performance benefits.

About the example

The commented program

```
template <typename ValueType>
class StencilMatrix : public gko::EnableLinOp<StencilMatrix<ValueType>,
        public gko::EnableCreateMethod<StencilMatrix<ValueType> {
public:
```

This constructor will be called by the create method. Here we initialize the coefficients of the stencil.

```
    StencilMatrix(std::shared_ptr<const gko::Executor> exec,
        gko::size_type size = 0, ValueType left = -1.0,
        ValueType center = 2.0, ValueType right = -1.0)
    : gko::EnableLinOp<StencilMatrix>(exec, gko::dim<2>{size}),
      coefficients(exec, {left, center, right})
    {}
protected:
    using vec = gko::matrix::Dense<ValueType>;
    using coef_type = gko::array<ValueType>;
```

Here we implement the application of the linear operator, $x = A * b$. `apply_impl` will be called by the apply method, after the arguments have been moved to the correct executor and the operators checked for conforming sizes.

For simplicity, we assume that there is always only one right hand side and the stride of consecutive elements in the vectors is 1 (both of these are always true in this example).

```
void apply_impl(const gko::LinOp* b, gko::LinOp* x) const override
{
```

we only implement the operator for dense RHS. `gko::as` will throw an exception if its argument is not Dense.

```
    auto dense_b = gko::as<vec>(b);
    auto dense_x = gko::as<vec>(x);
```

we need separate implementations depending on the executor, so we create an operation which maps the call to the correct implementation

```
struct stencil_operation : gko::Operation {
    stencil_operation(const coef_type& coefficients, const vec* b,
        vec* x)
    : coefficients{coefficients}, b{b}, x{x}
    {}
```

OpenMP implementation

```
    void run(std::shared_ptr<const gko::OmpExecutor>) const override
    {
        auto b_values = b->get_const_values();
        auto x_values = x->get_values();
#pragma omp parallel for
        for (std::size_t i = 0; i < x->get_size()[0]; ++i) {
            auto coefs = coefficients.get_const_data();
            auto result = coefs[1] * b_values[i];
            if (i > 0) {
                result += coefs[0] * b_values[i - 1];
            }
            if (i < x->get_size()[0] - 1) {
                result += coefs[2] * b_values[i + 1];
            }
            x_values[i] = result;
        }
    }
```

CUDA implementation

```
void run(std::shared_ptr<const gko::CudaExecutor>) const override
{
    stencil_kernel(x->get_size()[0], coefficients.get_const_data(),
        b->get_const_values(), x->get_values());
}
```

We do not provide an implementation for reference executor. If not provided, Ginkgo will use the implementation for the OpenMP executor when calling it in the reference executor.

```
    const coef_type& coefficients;
    const vec* b;
    vec* x;
};
this->get_executor()->run(
    stencil_operation(coefficients, dense_b, dense_x));
}
```


There is also a version of the apply function which does the operation $x = \alpha * A * b + \beta * x$. This function is commonly used and can often be better optimized than implementing it using $x = A * b$. However, for simplicity, we will implement it exactly like that in this example.

```
void apply_impl(const gko::LinOp* alpha, const gko::LinOp* b,
               const gko::LinOp* beta, gko::LinOp* x) const override
{
    auto dense_b = gko::as<vec>(b);
    auto dense_x = gko::as<vec>(x);
    auto tmp_x = dense_x->clone();
    this->apply_impl(b, tmp_x.get());
    dense_x->scale(beta);
    dense_x->add_scaled(alpha, tmp_x);
}
private:
    coef_type coefficients;
};
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    const auto discretization_points = matrix->get_size()[0];
    auto row_ptrs = matrix->get_row_ptrs();
    auto col_idxxs = matrix->get_col_idxxs();
    auto values = matrix->get_values();
    IndexType pos = 0;
    const ValueType coeffs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (int i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
                values[pos] = coeffs[ofs + 1];
                col_idxxs[pos] = i + ofs;
                ++pos;
            }
        }
        row_ptrs[i + 1] = pos;
    }
}
```

Generates the RHS vector given f and the boundary conditions.

```
template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{
    const auto discretization_points = rhs->get_size()[0];
    auto values = rhs->get_values();
    const ValueType h = 1.0 / (discretization_points + 1);
    for (int i = 0; i < discretization_points; ++i) {
        const ValueType xi = ValueType(i + 1) * h;
        values[i] = -f(xi) * h * h;
    }
    values[0] += u0;
    values[discretization_points - 1] += u1;
}
```

Prints the solution u .

```
template <typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                   const gko::matrix::Dense<ValueType>* u)
{
    std::cout << u0 << '\n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout << u->get_const_values()[i] << '\n';
    }
    std::cout << u1 << std::endl;
}
```

Computes the 1-norm of the error given the computed u and the correct solution function `correct_u`.

```
template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                      const gko::matrix::Dense<ValueType>* u,
                      Closure correct_u)
{
    const auto h = 1.0 / (discretization_points + 1);
    auto error = 0.0;
    for (int i = 0; i < discretization_points; ++i) {
        using std::abs;
        const auto xi = (i + 1) * h;
        error +=
            abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    }
}
```

```

    return error;
}
int main(int argc, char* argv[])
{

```

Some shortcuts

```

using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;

```

Figure out where to run the code

```

if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100u;
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};

```

executor where Ginkgo will perform the computation

```

const auto exec = exec_map.at(executor_string)(); // throws if not valid

```

executor used by the application

```

const auto app_exec = exec->get_master();

```

problem:

```

auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType{6} * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);

```

initialize vectors

```

auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
}
const RealValueType reduction_factor{1e-7};

```

Generate solver and solve the system

```

cg::build()
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(discretization_points),
        gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
            reduction_factor))
    .on(exec)

```

notice how our custom StencilMatrix can be used in the same way as any built-in type

```

->generate(StencilMatrix<ValueType>::create(exec, discretization_points,
                                           -1, 2, -1))

->apply(rhs, u);
std::cout << "\nSolve complete."
          << "\nThe average relative error is "
          << calculate_error(discretization_points, u.get(), correct_u) /
              discretization_points
          << std::endl;
}

```

Results

Comments about programming and debugging

The plain program

```
#include <iostream>
#include <map>
#include <string>
#include <omp.h>
#include <ginkgo/ginkgo.hpp>
template <typename ValueType>
void stencil_kernel(std::size_t size, const ValueType* coefs,
                  const ValueType* b, ValueType* x);
template <typename ValueType>
class StencilMatrix : public gko::EnableLinOp<StencilMatrix<ValueType>,
                  public gko::EnableCreateMethod<StencilMatrix<ValueType> {
public:
    StencilMatrix(std::shared_ptr<const gko::Executor> exec,
                  gko::size_type size = 0, ValueType left = -1.0,
                  ValueType center = 2.0, ValueType right = -1.0)
        : gko::EnableLinOp<StencilMatrix>(exec, gko::dim<2>{size}),
          coefficients(exec, {left, center, right})
    {}
protected:
    using vec = gko::matrix::Dense<ValueType>;
    using coef_type = gko::array<ValueType>;
    void apply_impl(const gko::LinOp* b, gko::LinOp* x) const override
    {
        auto dense_b = gko::as<vec>(b);
        auto dense_x = gko::as<vec>(x);
        struct stencil_operation : gko::Operation {
            stencil_operation(const coef_type& coefficients, const vec* b,
                             vec* x)
                : coefficients{coefficients}, b{b}, x{x}
            {}
            void run(std::shared_ptr<const gko::OmpExecutor>) const override
            {
                auto b_values = b->get_const_values();
                auto x_values = x->get_values();
                #pragma omp parallel for
                for (std::size_t i = 0; i < x->get_size()[0]; ++i) {
                    auto coefs = coefficients.get_const_data();
                    auto result = coefs[1] * b_values[i];
                    if (i > 0) {
                        result += coefs[0] * b_values[i - 1];
                    }
                    if (i < x->get_size()[0] - 1) {
                        result += coefs[2] * b_values[i + 1];
                    }
                    x_values[i] = result;
                }
            }
        }
        void run(std::shared_ptr<const gko::CudaExecutor>) const override
    {
        stencil_kernel(x->get_size()[0], coefficients.get_const_data(),
                      b->get_const_values(), x->get_values());
    }
    const coef_type& coefficients;
    const vec* b;
    vec* x;
};
this->get_executor()->run(
    stencil_operation(coefficients, dense_b, dense_x));
}
void apply_impl(const gko::LinOp* alpha, const gko::LinOp* b,
               const gko::LinOp* beta, gko::LinOp* x) const override
{
    auto dense_b = gko::as<vec>(b);
    auto dense_x = gko::as<vec>(x);
    auto tmp_x = dense_x->clone();
    this->apply_impl(b, tmp_x.get());
    dense_x->scale(beta);
    dense_x->add_scaled(alpha, tmp_x);
}
private:
    coef_type coefficients;
};
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    const auto discretization_points = matrix->get_size()[0];
```

```

    auto row_ptrs = matrix->get_row_ptrs();
    auto col_idxes = matrix->get_col_idxes();
    auto values = matrix->get_values();
    IndexType pos = 0;
    const ValueType coeffs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (int i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
                values[pos] = coeffs[ofs + 1];
                col_idxes[pos] = i + ofs;
                ++pos;
            }
        }
        row_ptrs[i + 1] = pos;
    }
}

template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{
    const auto discretization_points = rhs->get_size()[0];
    auto values = rhs->get_values();
    const ValueType h = 1.0 / (discretization_points + 1);
    for (int i = 0; i < discretization_points; ++i) {
        const ValueType xi = ValueType(i + 1) * h;
        values[i] = -f(xi) * h * h;
    }
    values[0] += u0;
    values[discretization_points - 1] += u1;
}

template <typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                   const gko::matrix::Dense<ValueType>* u)
{
    std::cout << u0 << '\n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout << u->get_const_values()[i] << '\n';
    }
    std::cout << u1 << std::endl;
}

template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                      const gko::matrix::Dense<ValueType>* u,
                      Closure correct_u)
{
    const auto h = 1.0 / (discretization_points + 1);
    auto error = 0.0;
    for (int i = 0; i < discretization_points; ++i) {
        using std::abs;
        const auto xi = (i + 1) * h;
        error +=
            abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    }
    return error;
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const unsigned int discretization_points =
        argc >= 3 ? std::atoi(argv[2]) : 100u;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }}
    };

```

```

    }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }});
const auto exec = exec_map.at(executor_string()); // throws if not valid
const auto app_exec = exec->get_master();
auto correct_u = [] (ValueType x) { return x * x * x; };
auto f = [] (ValueType x) { return ValueType{6} * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
}
const RealValueType reduction_factor{1e-7};
cg::build()
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(discretization_points),
        gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
            reduction_factor))
    .on(exec)
    ->generate(StencilMatrix<ValueType>::create(exec, discretization_points,
        -1, 2, -1))
    ->apply(rhs, u);
std::cout << "\nSolve complete."
    << "\nThe average relative error is "
    << calculate_error(discretization_points, u.get(), correct_u) /
        discretization_points
    << std::endl;
}

```


Chapter 13

The custom-stopping-criterion program

The custom stopping criterion creation example..

This example depends on simple-solver, minimal-cuda-solver.

Introduction

About the example

The commented program

```
};
GKO_ENABLE_CRITERION_FACTORY(ByInteraction, parameters, Factory);
GKO_ENABLE_BUILD_METHOD(Factory);
protected:
    bool check_impl(gko::uint8 stoppingId, bool setFinalized,
                    gko::array<gko::stopping_status>* stop_status,
                    bool* one_changed, const Criterion::Updater&) override
{
    bool result = *(parameters_.stop_iteration_process);
    if (result) {
        this->set_all_statuses(stoppingId, setFinalized, stop_status);
        *one_changed = true;
    }
    return result;
}
explicit ByInteraction(std::shared_ptr<const gko::Executor> exec)
    : EnablePolymorphicObject<ByInteraction, Criterion>(std::move(exec))
{}
explicit ByInteraction(const Factory* factory,
                       const gko::stop::CriterionArgs& args)
    : EnablePolymorphicObject<ByInteraction, Criterion>(
        factory->get_executor(),
        parameters_{factory->get_parameters()})
{}
};
void run_solver(volatile bool* stop_iteration_process,
                std::shared_ptr<gko::Executor> exec)
{
```

Some shortcuts

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using bicg = gko::solver::Bicgstab<ValueType>;
```

Read Data

```

auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);

```

Create solver factory and solve system

```

auto solver =
    bicg::build()
        .with_criteria(ByInteraction::build().with_stop_iteration_process(
            stop_iteration_process))
        .on(exec)
        ->generate(A);
solver->add_logger(gko::log::Stream<ValueType>::create(
    gko::log::Logger::iteration_complete_mask, std::cout, true));
solver->apply(b, x);
std::cout << "Solver stopped" << std::endl;

```

Print solution

```

std::cout << "Solution (x): \n";
write(std::cout, x);

```

Calculate residual

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r): \n";
write(std::cout, res);
}
int main(int argc, char* argv[])
{

```

Print version information

```

std::cout << gko::version_info::get() << std::endl;

```

Figure out where to run the code

```

if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}

```

Figure out where to run the code

```

const auto executor_string = argc >= 2 ? argv[1] : "reference";

```

Figure out where to run the code

```

std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}};

```

executor where Ginkgo will perform the computation

```

const auto exec = exec_map.at(executor_string)(); // throws if not valid

```

Declare a user controlled boolean for the iteration process

```

volatile bool stop_iteration_process{};

```

Create a new a thread to launch the solver

```

std::thread t(run_solver, &stop_iteration_process, exec);

```

Look for an input command "stop" in the console, which sets the boolean to true

```

std::cout << "Type 'stop' to stop the iteration process" << std::endl;
std::string command;
while (std::cin >> command) {

```



```

        if (command == "stop") {
            break;
        } else {
            std::cout << "Unknown command" << std::endl;
        }
    }
    std::cout << "User input command 'stop' - The solver will stop!"
              << std::endl;
    stop_iteration_process = true;
    t.join();
}

```

Results

This is the expected output:

```

.
.
.
.
.
.
[LOG] >> iteration 22516 completed with solver LinOp[gko::solver::Bicgstab<double>,0x7fe6a4003710] with
      residual LinOp[gko::matrix::Dense<double>,0x7fe6a40050b0], solution
      LinOp[gko::matrix::Dense<double>,0x7fe6a40048e0] and residual_norm LinOp[gko::LinOp const*,0]
LinOp[gko::matrix::Dense<double>,0x7fe6a40050b0] [
  5.17803e-164
 -7.6865e-165
 -2.06149e-164
 -4.84737e-165
 -3.36597e-164
  2.22353e-164
  1.47594e-165
 -1.78592e-165
 -6.17274e-166
 -3.02681e-166
  7.82009e-166
  8.57102e-165
 -1.28879e-164
 -2.62076e-165
  2.55329e-165
 -5.95988e-166
 -5.79273e-166
 -5.20172e-166
 -6.79458e-166
]
// Typing 'stop' stops the solver.
User input command 'stop' - The solver will stop
LinOp[gko::matrix::Dense<double>,0x7fe6a40048e0] [
  0.252218
  0.108645
  0.0662811
  0.0630433
  0.0384088
  0.0396536
  0.0402648
  0.0338935
  0.0193098
  0.0234653
  0.0211499
  0.0196413
  0.0199151
  0.0181674
  0.0162722
  0.0150714
  0.0107016
  0.0121141
  0.0123025
]
Solver stopped
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098

```

```

0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
6.50306e-16

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <thread>
#include <ginkgo/ginkgo.hpp>
class ByInteraction
: public gko::EnablePolymorphicObject<ByInteraction, gko::stop::Criterion> {
    friend class gko::EnablePolymorphicObject<ByInteraction,
                                                gko::stop::Criterion>;
    using Criterion = gko::stop::Criterion;
public:
    GKO_CREATE_FACTORY_PARAMETERS(parameters, Factory)
    {
        std::add_pointer<volatile bool>::type GKO_FACTORY_PARAMETER_SCALAR(
            stop_iteration_process, nullptr);
    };
    GKO_ENABLE_CRITERION_FACTORY(ByInteraction, parameters, Factory);
    GKO_ENABLE_BUILD_METHOD(Factory);
protected:
    bool check_impl(gko::uint8 stoppingId, bool setFinalized,
                    gko::array<gko::stopping_status>* stop_status,
                    bool* one_changed, const Criterion::Updater&) override
    {
        bool result = *(parameters_.stop_iteration_process);
        if (result) {
            this->set_all_statuses(stoppingId, setFinalized, stop_status);
            *one_changed = true;
        }
        return result;
    }
    explicit ByInteraction(std::shared_ptr<const gko::Executor> exec)
        : EnablePolymorphicObject<ByInteraction, Criterion>(std::move(exec))
    {}
    explicit ByInteraction(const Factory* factory,
                           const gko::stop::CriterionArgs& args)
        : EnablePolymorphicObject<ByInteraction, Criterion>(
            factory->get_executor()),
          parameters_{factory->get_parameters()}
    {}
};

void run_solver(volatile bool* stop_iteration_process,
                std::shared_ptr<gko::Executor> exec)
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using bicg = gko::solver::Bicgstab<ValueType>;
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    auto solver =
        bicg::build()
            .with_criteria(ByInteraction::build().with_stop_iteration_process(
                stop_iteration_process))
            .on(exec)
            ->generate(A);
    solver->add_logger(gko::log::Stream<ValueType>::create(
        gko::log::Logger::iteration_complete_mask, std::cout, true));
}

```

```

solver->apply(b, x);
std::cout << "Solver stopped" << std::endl;
std::cout << "Solution (x): \n";
write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r): \n";
write(std::cout, res);
}
int main(int argc, char* argv[])
{
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    volatile bool stop_iteration_process{};
    std::thread t(run_solver, &stop_iteration_process, exec);
    std::cout << "Type 'stop' to stop the iteration process" << std::endl;
    std::string command;
    while (std::cin >> command) {
        if (command == "stop") {
            break;
        } else {
            std::cout << "Unknown command" << std::endl;
        }
    }
    std::cout << "User input command 'stop' - The solver will stop!"
              << std::endl;
    stop_iteration_process = true;
    t.join();
}

```


The distributed-multigrid-preconditioned-solver program

This example depends on distributed-solver.

This distributed multigrid preconditioned solver example should help you understand customizing Ginkgo multigrid in a distributed setting. The example will solve a simple 1D Laplace equation where the system can be distributed row-wise to multiple processes. Note. Because the stencil for the discretized Laplacian is configured with equal weight, the coarsening method does not perform well on this kind of problem. To run the solver with multiple processes, use `mpirun -n NUM_PROCS ./distributed-solver [executor] [num_grid←_points] [num_iterations]`.

The commented program

[illegible]

We can use here the same solver type as you would use in a non-distributed program. Please note that not all solvers support distributed systems at the moment.

```
using solver = gko::solver::Cg<ValueType>;
```

We use the Schwarz preconditioner to extend non-distributed preconditioners, like our Jacobi, to the distributed case. The Schwarz preconditioner wraps another preconditioner, and applies it only to the local part of a distributed matrix. This will be used as our distributed multigrid smoother.

```
using schwarz = gko::experimental::distributed::preconditioner::Schwarz<
    ValueType, LocalIndexType, GlobalIndexType>;
using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
```

Multigrid and Pgm can accept the distributed matrix, so we still use the same type as the non-distributed case.

```
using mg = gko::solver::Multigrid;
using pgm = gko::multigrid::Pgm<ValueType, LocalIndexType>;
```

Create an MPI communicator get the rank of the calling process.

```
const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);
const auto rank = comm.rank();
```

User Input Handling

User input settings:

- The executor, defaults to reference.
- The number of grid points, defaults to 100.
- The number of iterations, defaults to 1000.

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    if (rank == 0) {
        std::cerr << "Usage: " << argv[0]
            << " [executor] [num_grid_points] [num_iterations] "
            << std::endl;
    }
    std::exit(-1);
}
ValueType t_init = gko::experimental::mpi::get_walltime();
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const auto grid_dim =
    static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
const auto num_iters =
    static_cast<gko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
const std::map<std::string,
    std::function<std::shared_ptr<gko::Executor>(MPI_Comm)>>
    executor_factory_mpi{
    {"reference",
        [] (MPI_Comm) { return gko::ReferenceExecutor::create(); }},
    {"omp", [] (MPI_Comm) { return gko::OmpExecutor::create(); }},
    {"cuda",
        [] (MPI_Comm comm) {
            int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::CudaExecutor::get_num_devices());
            return gko::CudaExecutor::create(
                device_id, gko::ReferenceExecutor::create());
        }},
    {"hip",
        [] (MPI_Comm comm) {
            int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::HipExecutor::get_num_devices());
            return gko::HipExecutor::create(
                device_id, gko::ReferenceExecutor::create());
        }},
    {"dpcpp", [] (MPI_Comm comm) {
        int device_id = 0;
        if (gko::DpcppExecutor::get_num_devices("gpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("gpu"));
        } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("cpu"));
        } else {
            throw std::runtime_error("No suitable DPC++ devices");
        }
        return gko::DpcppExecutor::create(
            device_id, gko::ReferenceExecutor::create());
    }}};
auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
```

Creating the Distributed Matrix and Vectors

As a first step, we create a partition of the rows. The partition consists of ranges of consecutive rows which are assigned a part-id. These part-ids will be used for the distributed data structures to determine which rows will be stored locally. In this example each rank has (nearly) the same number of rows, so we can use the following specialized constructor. See `gko::distributed::Partition` for other modes of creating a partition.

```
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
```

Assemble the matrix using a 3-pt stencil and fill the right-hand-side with a sine value. The distributed matrix supports only constructing an empty matrix of zero size and filling in the values with `gko::experimental::distributed::Matrix::read_distributed`. Only the data that belongs to the rows by this rank will be assembled.

```
gko::matrix_data<ValueType, GlobalIndexType> A_data;
gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {
    if (i > 0) {
        A_data.nonzeros.emplace_back(i, i - 1, -1);
    }
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid_dim - 1) {
        A_data.nonzeros.emplace_back(i, i + 1, -1);
    }
    b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
    x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
}
```

Take timings.

```
comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();
```

Read the matrix data, currently this is only supported on CPU executors. This will also set up the communication pattern needed for the distributed matrix-vector multiplication.

```
auto A_host = gko::share(dist_mtx::create(exec->get_master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);
```

After reading, the matrix and vector can be moved to the chosen executor, since the distributed matrix supports SpMV also on devices.

```
auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);
```

Take timings.

```
comm.synchronize();
ValueType t_read_setup_end = gko::experimental::mpi::get_walltime();
```

Solve the Distributed System

Generate the solver

Setup the multigrid factory with customized smoother and coarse solver Because BlockJacobi does not support distributed matrix, we need wrap it in Schwarz.

```
auto schwarz_bj_factory =
    gko::share(schwarz::build().with_local_solver(bj::build()).on(exec));
auto smoother_factory = gko::share(gko::solver::build_smoother(
    schwarz_bj_factory, 2u, static_cast<ValueType>(0.9)));
```

Cg supports distributed matrix, so we can use it as we did in non-distributed case

```
auto coarsest_factory = gko::share(
    solver::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
        .on(exec));
```

The multigrid preconditioner uses the Schwarz Jacobi as smoother and Cg as coarse solver

```
auto mg_factory = gko::share(
    mg::build()
        .with_mg_level(pgm::build().with_deterministic(true))
        .with_pre_smoother(smoother_factory)
        .with_coarsest_solver(coarsest_factory)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec));
```

Setup the stopping criterion and logger

```
const gko::remove_complex<ValueType> reduction_factor{1e-8};
std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
auto Ainv = solver::build()
    .with_preconditioner(mg_factory)
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(num_iters),
        gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
    .on(exec)
    ->generate(A);
```

Add logger to the generated solver to log the iteration count and residual norm

```
Ainv->add_logger(logger);
```

Take timings.

```
comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();
```

Apply the distributed solver, this is the same as in the non-distributed case.

```
Ainv->apply(b, x);
```

Take timings.

```
comm.synchronize();
ValueType t_end = gko::experimental::mpi::get_walltime();
```

Get the residual.

```
auto res_norm = gko::clone(exec->get_master(),
    gko::as<vec>(logger->get_residual_norm()));
```

Printing Results

Print the achieved residual norm and timings on rank 0.

```
if (comm.rank() == 0) {
```

clang-format off

```
std::cout << "\nNum rows in matrix: " << num_rows
    << "\nNum ranks: " << comm.size()
    << "\nFinal Res norm: " << res_norm->at(0, 0)
    << "\nIteration count: " << logger->get_num_iterations()
    << "\nInit time: " << t_init_end - t_init
    << "\nRead time: " << t_read_setup_end - t_init
    << "\nSolver generate time: " << t_solver_generate_end - t_read_setup_end
    << "\nSolver apply time: " << t_end - t_solver_generate_end
    << "\nTotal time: " << t_end - t_init
    << std::endl;
```

clang-format on

```
    }
}
```


Results

This is the expected output for `mpirun -n 4 ./distributed-multigrid-preconditioned-solver`

```
:
Num rows in matrix: 100
Num ranks: 4
Final Res norm: 1.61045e-08
Iteration count: 18
Init time: 0.000117699
Read time: 0.000522518
Solver generate time: 0.000430548
Solver apply time: 0.00183804
Total time: 0.00279111
```

The timings may vary depending on the machine.

The plain program

```
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
{
    const gko::experimental::mpi::environment env(argc, argv);
    using GlobalIndexType = gko::int64;
    using LocalIndexType = gko::int32;
    using ValueType = double;
    using dist_vec = gko::experimental::distributed::Vector<ValueType>;
    using dist_mtx =
        gko::experimental::distributed::Matrix<ValueType, LocalIndexType,
            GlobalIndexType>;

    using vec = gko::matrix::Dense<ValueType>;
    using part_type =
        gko::experimental::distributed::Partition<LocalIndexType,
            GlobalIndexType>;

    using solver = gko::solver::Cg<ValueType>;
    using schwarz = gko::experimental::distributed::preconditioner::Schwarz<
        ValueType, LocalIndexType, GlobalIndexType>;
    using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
    using mg = gko::solver::Multigrid;
    using pgm = gko::multigrid::Pgm<ValueType, LocalIndexType>;
    const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);
    const auto rank = comm.rank();
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        if (rank == 0) {
            std::cerr << "Usage: " << argv[0]
                << " [executor] [num_grid_points] [num_iterations] "
                << std::endl;
        }
        std::exit(-1);
    }
    ValueType t_init = gko::experimental::mpi::get_walltime();
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const auto grid_dim =
        static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
    const auto num_iters =
        static_cast<gko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
    const std::map<std::string,
        std::function<std::shared_ptr<gko::Executor>(MPI_Comm)>>
        executor_factory_mpi{
        {"reference",
            [] (MPI_Comm) { return gko::ReferenceExecutor::create(); }},
        {"omp", [] (MPI_Comm) { return gko::OmpExecutor::create(); }},
        {"cuda",
            [] (MPI_Comm comm) {
                int device_id = gko::experimental::mpi::map_rank_to_device_id(
                    comm, gko::CudaExecutor::get_num_devices());
                return gko::CudaExecutor::create(
                    device_id, gko::ReferenceExecutor::create());
            }},
        {"hip",
            [] (MPI_Comm comm) {
                int device_id = gko::experimental::mpi::map_rank_to_device_id(
                    comm, gko::HipExecutor::get_num_devices());
                return gko::HipExecutor::create(
                    device_id, gko::ReferenceExecutor::create());
            }},
        {"dpcpp", [] (MPI_Comm comm) {
```

```

    int device_id = 0;
    if (gko::DpcppExecutor::get_num_devices("gpu")) {
        device_id = gko::experimental::mpi::map_rank_to_device_id(
            comm, gko::DpcppExecutor::get_num_devices("gpu"));
    } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
        device_id = gko::experimental::mpi::map_rank_to_device_id(
            comm, gko::DpcppExecutor::get_num_devices("cpu"));
    } else {
        throw std::runtime_error("No suitable DPC++ devices");
    }
    return gko::DpcppExecutor::create(
        device_id, gko::ReferenceExecutor::create());
    }
};

auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
gko::matrix_data<ValueType, GlobalIndexType> A_data;
gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {
    if (i > 0) {
        A_data.nonzeros.emplace_back(i, i - 1, -1);
    }
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid_dim - 1) {
        A_data.nonzeros.emplace_back(i, i + 1, -1);
    }
    b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
    x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
}
comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();
auto A_host = gko::share(dist_mtx::create(exec->get_master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);
auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);
comm.synchronize();
ValueType t_read_setup_end = gko::experimental::mpi::get_walltime();
auto schwarz_bj_factory =
    gko::share(schwarz::build().with_local_solver(bj::build()).on(exec));
auto smoother_factory = gko::share(gko::solver::build_smoother(
    schwarz_bj_factory, 2u, static_cast<ValueType>(0.9)));
auto coarsest_factory = gko::share(
    solver::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
        .on(exec));
auto mg_factory = gko::share(
    mg::build()
        .with_mg_level(pgm::build().with_deterministic(true))
        .with_pre_smoother(smoother_factory)
        .with_coarsest_solver(coarsest_factory)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec));
const gko::remove_complex<ValueType> reduction_factor{1e-8};
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
auto Ainv = solver::build()
    .with_preconditioner(mg_factory)
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(num_iters),
        gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
    .on(exec)
    ->generate(A);
Ainv->add_logger(logger);
comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();
Ainv->apply(b, x);
comm.synchronize();
ValueType t_end = gko::experimental::mpi::get_walltime();
auto res_norm = gko::clone(exec->get_master(),
    gko::as<vec>(logger->get_residual_norm()));

```

```
if (comm.rank() == 0) {
    std::cout << "\nNum rows in matrix: " << num_rows
               << "\nNum ranks: " << comm.size()
               << "\nFinal Res norm: " << res_norm->at(0, 0)
               << "\nIteration count: " << logger->get_num_iterations()
               << "\nInit time: " << t_init_end - t_init
               << "\nRead time: " << t_read_setup_end - t_init
               << "\nSolver generate time: " << t_solver_generate_end - t_read_setup_end
               << "\nSolver apply time: " << t_end - t_solver_generate_end
               << "\nTotal time: " << t_end - t_init
               << std::endl;
}
```


Chapter 15

The distributed-solver program

The distributed solver example..

This example depends on simple-solver, three-pt-stencil-solver.

Introduction

This distributed solver example should help you understand the basics of using Ginkgo in a distributed setting. The example will solve a simple 1D Laplace equation where the system can be distributed row-wise to multiple processes. To run the solver with multiple processes, use `mpirun -n NUM_PROCS ./distributed-solver [executor] [num_grid_points] [num_iterations]`.

If you are using GPU devices, please make sure that you run this example with at most as many processes as you have GPU devices available.

The commented program

As vector type we use the following, which implements a subset of `gko::matrix::Dense`.

```
using dist_vec = gko::experimental::distributed::Vector<ValueType>;
```

As matrix type we simply use the following type, which can read distributed data and be applied to a distributed vector.

```
using dist_mtx =  
    gko::experimental::distributed::Matrix<ValueType, LocalIndexType,  
        GlobalIndexType>;
```

We still need a localized vector type to be used as scalars in the advanced apply operations.

```
using vec = gko::matrix::Dense<ValueType>;
```

The partition type describes how the rows of the matrices are distributed.

```
using part_type =  
    gko::experimental::distributed::Partition<LocalIndexType,  
        GlobalIndexType>;
```

We can use here the same solver type as you would use in a non-distributed program. Please note that not all solvers support distributed systems at the moment.

```
using solver = gko::solver::Cg<ValueType>;  
using schwarz = gko::experimental::distributed::preconditioner::Schwarz<  
    ValueType, LocalIndexType, GlobalIndexType>;  
using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
```

Create an MPI communicator get the rank of the calling process.

```
const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);  
const auto rank = comm.rank();
```

User Input Handling

User input settings:

- The executor, defaults to reference.
- The number of grid points, defaults to 100.
- The number of iterations, defaults to 1000.

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    if (rank == 0) {
        std::cerr << "Usage: " << argv[0]
            << " [executor] [num_grid_points] [num_iterations] "
            << std::endl;
    }
    std::exit(-1);
}
ValueType t_init = gko::experimental::mpi::get_walltime();
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const auto grid_dim =
    static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
const auto num_iters =
    static_cast<gko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
const std::map<std::string,
    std::function<std::shared_ptr<gko::Executor>(MPI_Comm)>>
    executor_factory_mpi{
    {"reference",
        [] (MPI_Comm) { return gko::ReferenceExecutor::create(); }},
    {"omp", [] (MPI_Comm) { return gko::OmpExecutor::create(); }},
    {"cuda",
        [] (MPI_Comm comm) {
            int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::CudaExecutor::get_num_devices());
            return gko::CudaExecutor::create(
                device_id, gko::ReferenceExecutor::create());
        }},
    {"hip",
        [] (MPI_Comm comm) {
            int device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::HipExecutor::get_num_devices());
            return gko::HipExecutor::create(
                device_id, gko::ReferenceExecutor::create());
        }},
    {"dpcpp", [] (MPI_Comm comm) {
        int device_id = 0;
        if (gko::DpcppExecutor::get_num_devices("gpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("gpu"));
        } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("cpu"));
        } else {
            throw std::runtime_error("No suitable DPC++ devices");
        }
        return gko::DpcppExecutor::create(
            device_id, gko::ReferenceExecutor::create());
    }
    }
};
auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
```

Creating the Distributed Matrix and Vectors

As a first step, we create a partition of the rows. The partition consists of ranges of consecutive rows which are assigned a part-id. These part-ids will be used for the distributed data structures to determine which rows will be stored locally. In this example each rank has (nearly) the same number of rows, so we can use the following specialized constructor. See `gko::distributed::Partition` for other modes of creating a partition.

```
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
```

Assemble the matrix using a 3-pt stencil and fill the right-hand-side with a sine value. The distributed matrix supports only constructing an empty matrix of zero size and filling in the values with `gko::experimental::distributed::Matrix::read_distributed`. Only the data that belongs to the rows by this rank will be assembled.

```
gko::matrix_data<ValueType, GlobalIndexType> A_data;
```

```

gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {
    if (i > 0) {
        A_data.nonzeros.emplace_back(i, i - 1, -1);
    }
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid_dim - 1) {
        A_data.nonzeros.emplace_back(i, i + 1, -1);
    }
    b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
    x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
}

```

Take timings.

```

comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();

```

Read the matrix data, currently this is only supported on CPU executors. This will also set up the communication pattern needed for the distributed matrix-vector multiplication.

```

auto A_host = gko::share(dist_mtx::create(exec->get_master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);

```

After reading, the matrix and vector can be moved to the chosen executor, since the distributed matrix supports SpMV also on devices.

```

auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);

```

Take timings.

```

comm.synchronize();
ValueType t_read_setup_end = gko::experimental::mpi::get_walltime();

```

Solve the Distributed System

Generate the solver, this is the same as in the non-distributed case. with a local block diagonal preconditioner.

Setup the local block diagonal solver factory.

```

auto local_solver = gko::share(bj::build().on(exec));

```

Setup the stopping criterion and logger

```

const gko::remove_complex<ValueType> reduction_factor{1e-8};
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
auto Ainv = solver::build()
    .with_preconditioner(
        schwarz::build().with_local_solver(local_solver))
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(num_iters),
        gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
    .on(exec)
    ->generate(A);

```

Add logger to the generated solver to log the iteration count and residual norm

```

Ainv->add_logger(logger);

```

Take timings.

```

comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();

```

Apply the distributed solver, this is the same as in the non-distributed case.

```

Ainv->apply(b, x);

```

Take timings.

```

comm.synchronize();
ValueType t_end = gko::experimental::mpi::get_walltime();

```

Get the residual.

```

auto res_norm = gko::clone(exec->get_master(),
    gko::as<vec>(logger->get_residual_norm()));

```

Printing Results

Print the achieved residual norm and timings on rank 0.

```
if (comm.rank() == 0) {

clang-format off
std::cout << "\nNum rows in matrix: " << num_rows
  << "\nNum ranks: " << comm.size()
  << "\nFinal Res norm: " << res_norm->at(0, 0)
  << "\nIteration count: " << logger->get_num_iterations()
  << "\nInit time: " << t_init_end - t_init
  << "\nRead time: " << t_read_setup_end - t_init
  << "\nSolver generate time: " << t_solver_generate_end - t_read_setup_end
  << "\nSolver apply time: " << t_end - t_solver_generate_end
  << "\nTotal time: " << t_end - t_init
  << std::endl;

clang-format on
}
}
```

Results

This is the expected output for `mpirun -n 4 ./distributed-solver`:

```
Num rows in matrix: 100
Num ranks: 4
Final Res norm: 5.58392e-12
Iteration count: 7
Init time: 0.0663887
Read time: 0.0729806
Solver generate time: 7.6348e-05
Solver apply time: 0.0680783
Total time: 0.141351
```

The timings may vary depending on the machine.

The plain program

```
#include <ginkgo/ginkgo.hpp>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
{
    const gko::experimental::mpi::environment env(argc, argv);
    using GlobalIndexType = gko::int64;
    using LocalIndexType = gko::int32;
    using ValueType = double;
    using dist_vec = gko::experimental::distributed::Vector<ValueType>;
    using dist_mtx =
        gko::experimental::distributed::Matrix<ValueType, LocalIndexType,
                                                GlobalIndexType>;

    using vec = gko::matrix::Dense<ValueType>;
    using part_type =
        gko::experimental::distributed::Partition<LocalIndexType,
                                                GlobalIndexType>;

    using solver = gko::solver::Cg<ValueType>;
    using schwarz = gko::experimental::distributed::preconditioner::Schwarz<
        ValueType, LocalIndexType, GlobalIndexType>;
    using bj = gko::preconditioner::Jacobi<ValueType, LocalIndexType>;
    const auto comm = gko::experimental::mpi::communicator(MPI_COMM_WORLD);
    const auto rank = comm.rank();
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        if (rank == 0) {
            std::cerr << "Usage: " << argv[0]
                << " " << "[executor] [num_grid_points] [num_iterations] "
                << std::endl;
        }
        std::exit(-1);
    }
    ValueType t_init = gko::experimental::mpi::get_walltime();
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
```



```

const auto grid_dim =
    static_cast<gko::size_type>(argc >= 3 ? std::atoi(argv[2]) : 100);
const auto num_iters =
    static_cast<gko::size_type>(argc >= 4 ? std::atoi(argv[3]) : 1000);
const std::map<std::string,
    std::function<std::shared_ptr<gko::Executor>(MPI_Comm)>>
    executor_factory_mpi{
    {"reference",
    [] (MPI_Comm) { return gko::ReferenceExecutor::create(); }},
    {"omp", [] (MPI_Comm) { return gko::OmpExecutor::create(); }},
    {"cuda",
    [] (MPI_Comm comm) {
        int device_id = gko::experimental::mpi::map_rank_to_device_id(
            comm, gko::CudaExecutor::get_num_devices());
        return gko::CudaExecutor::create(
            device_id, gko::ReferenceExecutor::create());
    }},
    {"hip",
    [] (MPI_Comm comm) {
        int device_id = gko::experimental::mpi::map_rank_to_device_id(
            comm, gko::HipExecutor::get_num_devices());
        return gko::HipExecutor::create(
            device_id, gko::ReferenceExecutor::create());
    }},
    {"dpcpp", [] (MPI_Comm comm) {
        int device_id = 0;
        if (gko::DpcppExecutor::get_num_devices("gpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("gpu"));
        } else if (gko::DpcppExecutor::get_num_devices("cpu")) {
            device_id = gko::experimental::mpi::map_rank_to_device_id(
                comm, gko::DpcppExecutor::get_num_devices("cpu"));
        } else {
            throw std::runtime_error("No suitable DPC++ devices");
        }
        return gko::DpcppExecutor::create(
            device_id, gko::ReferenceExecutor::create());
    }}};

auto exec = executor_factory_mpi.at(executor_string)(MPI_COMM_WORLD);
const auto num_rows = grid_dim;
auto partition = gko::share(part_type::build_from_global_size_uniform(
    exec->get_master(), comm.size(),
    static_cast<GlobalIndexType>(num_rows)));
gko::matrix_data<ValueType, GlobalIndexType> A_data;
gko::matrix_data<ValueType, GlobalIndexType> b_data;
gko::matrix_data<ValueType, GlobalIndexType> x_data;
A_data.size = {num_rows, num_rows};
b_data.size = {num_rows, 1};
x_data.size = {num_rows, 1};
const auto range_start = partition->get_range_bounds()[rank];
const auto range_end = partition->get_range_bounds()[rank + 1];
for (int i = range_start; i < range_end; i++) {
    if (i > 0) {
        A_data.nonzeros.emplace_back(i, i - 1, -1);
    }
    A_data.nonzeros.emplace_back(i, i, 2);
    if (i < grid_dim - 1) {
        A_data.nonzeros.emplace_back(i, i + 1, -1);
    }
    b_data.nonzeros.emplace_back(i, 0, std::sin(i * 0.01));
    x_data.nonzeros.emplace_back(i, 0, gko::zero<ValueType>());
}
comm.synchronize();
ValueType t_init_end = gko::experimental::mpi::get_walltime();
auto A_host = gko::share(dist_mtx::create(exec->get_master(), comm));
auto x_host = dist_vec::create(exec->get_master(), comm);
auto b_host = dist_vec::create(exec->get_master(), comm);
A_host->read_distributed(A_data, partition);
b_host->read_distributed(b_data, partition);
x_host->read_distributed(x_data, partition);
auto A = gko::share(dist_mtx::create(exec, comm));
auto x = dist_vec::create(exec, comm);
auto b = dist_vec::create(exec, comm);
A->copy_from(A_host);
b->copy_from(b_host);
x->copy_from(x_host);
comm.synchronize();
ValueType t_read_setup_end = gko::experimental::mpi::get_walltime();
auto local_solver = gko::share(bj::build().on(exec));
const gko::remove_complex<ValueType> reduction_factor{1e-8};
std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
auto Ain = solver::build()
    .with_preconditioner(
        schwarz::build().with_local_solver(local_solver))
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(num_iters),

```

```

        gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
        .on(exec)
        ->generate(A);
Ainv->add_logger(logger);
comm.synchronize();
ValueType t_solver_generate_end = gko::experimental::mpi::get_walltime();
Ainv->apply(b, x);
comm.synchronize();
ValueType t_end = gko::experimental::mpi::get_walltime();
auto res_norm = gko::clone(exec->get_master(),
    gko::as<vec>(logger->get_residual_norm()));
if (comm.rank() == 0) {
    std::cout << "\nNum rows in matrix: " << num_rows
        << "\nNum ranks: " << comm.size()
        << "\nFinal Res norm: " << res_norm->at(0, 0)
        << "\nIteration count: " << logger->get_num_iterations()
        << "\nInit time: " << t_init_end - t_init
        << "\nRead time: " << t_read_setup_end - t_init
        << "\nSolver generate time: " << t_solver_generate_end - t_read_setup_end
        << "\nSolver apply time: " << t_end - t_solver_generate_end
        << "\nTotal time: " << t_end - t_init
        << std::endl;
}
}

```

Chapter 16

The external-lib-interfacing program

The external library(deal.II) interfacing example..

Introduction

About the example

The commented program

```
#include <deal.II/base/function.h>
#include <deal.II/base/logstream.h>
#include <deal.II/base/quadrature_lib.h>
#include <deal.II/dofs/dof_accessor.h>
#include <deal.II/dofs/dof_handler.h>
#include <deal.II/dofs/dof_tools.h>
#include <deal.II/fe/fe_q.h>
#include <deal.II/fe/fe_values.h>
#include <deal.II/grid/grid_generator.h>
#include <deal.II/grid/grid_out.h>
#include <deal.II/grid/grid_refinement.h>
#include <deal.II/grid/tria.h>
#include <deal.II/grid/tria_accessor.h>
#include <deal.II/grid/tria_iterator.h>
#include <deal.II/lac/constraint_matrix.h>
#include <deal.II/lac/dynamic_sparsity_pattern.h>
#include <deal.II/lac/full_matrix.h>
#include <deal.II/lac/precondition.h>
#include <deal.II/lac/solver_bicgstab.h>
#include <deal.II/lac/sparse_matrix.h>
#include <deal.II/lac/vector.h>
#include <deal.II/numerics/data_out.h>
#include <deal.II/numerics/matrix_tools.h>
#include <deal.II/numerics/vector_tools.h>
```

The following two files provide classes and information for multithreaded programs. In the first one, the classes and functions are declared which we need to do assembly in parallel (i.e. the `WorkStream` namespace). The second file has a class `MultithreadInfo` which can be used to query the number of processors in your system, which is often useful when deciding how many threads to start in parallel.

```
#include <deal.II/base/multithread_info.h>
#include <deal.II/base/work_stream.h>
```

The next new include file declares a base class `TensorFunction` not unlike the `Function` class, but with the difference that the return value is tensor-valued rather than scalar or vector-valued.

```
#include <deal.II/base/tensor_function.h>
#include <deal.II/numerics/error_estimator.h>
```

Ginkgo's header file

```
#include <ginkgo/ginkgo.hpp>
```

This is C++, as we want to write some output to disk:

```
#include <fstream>
#include <iostream>
```

The last step is as in previous programs:

```
namespace Step9 {
using namespace dealii;
```

AdvectionProblem class declaration

Following we declare the main class of this program. It is very much like the main classes of previous examples, so we again only comment on the differences.

```
template <int dim>
class AdvectionProblem {
public:
    AdvectionProblem();
    AdvectionProblem();
    void run();
private:
    void setup_system();
```

The next set of functions will be used to assemble the matrix. However, unlike in the previous examples, the `assemble_system()` function will not do the work itself, but rather will delegate the actual assembly to helper functions `assemble_local_system()` and `copy_local_to_global()`. The rationale is that matrix assembly can be parallelized quite well, as the computation of the local contributions on each cell is entirely independent of other cells, and we only have to synchronize when we add the contribution of a cell to the global matrix.

The strategy for parallelization we choose here is one of the possibilities mentioned in detail in the threads module in the documentation. Specifically, we will use the WorkStream approach discussed there. Since there is so much documentation in this module, we will not repeat the rationale for the design choices here (for example, if you read through the module mentioned above, you will understand what the purpose of the `AssemblyScratchData` and `AssemblyCopyData` structures is). Rather, we will only discuss the specific implementation.

If you read the page mentioned above, you will find that in order to parallelize assembly, we need two data structures – one that corresponds to data that we need during local integration ("scratch data", i.e., things we only need as temporary storage), and one that carries information from the local integration to the function that then adds the local contributions to the corresponding elements of the global matrix. The former of these typically contains the `FEValues` and `FEFaceValues` objects, whereas the latter has the local matrix, local right hand side, and information about which degrees of freedom live on the cell for which we are assembling a local contribution. With this information, the following should be relatively self-explanatory:

```
struct AssemblyScratchData {
    AssemblyScratchData(const FiniteElement<dim>& fe);
    AssemblyScratchData(const AssemblyScratchData& scratch_data);
    FEValues<dim> fe_values;
    FEFaceValues<dim> fe_face_values;
};
struct AssemblyCopyData {
    FullMatrix<double> cell_matrix;
    Vector<double> cell_rhs;
    std::vector<types::global_dof_index> local_dof_indices;
};
void assemble_system();
void local_assemble_system(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    AssemblyScratchData& scratch, AssemblyCopyData& copy_data);
void copy_local_to_global(const AssemblyCopyData& copy_data);
```

The following functions again are as in previous examples, as are the subsequent variables.

```
void solve();
void refine_grid();
void output_results(const unsigned int cycle) const;
Triangulation<dim> triangulation;
DoFHandler<dim> dof_handler;
FE_Q<dim> fe;
ConstraintMatrix hanging_node_constraints;
SparsityPattern sparsity_pattern;
SparseMatrix<double> system_matrix;
Vector<double> solution;
Vector<double> system_rhs;
};
```

Equation data declaration

Next we declare a class that describes the advection field. This, of course, is a vector field with as many components as there are space dimensions. One could now use a class derived from the `Function` base class, as we have done for boundary values and coefficients in previous examples, but there is another possibility in the library, namely

a base class that describes tensor valued functions. In contrast to the usual `Function` objects, we provide the compiler with knowledge on the size of the objects of the return type. This enables the compiler to generate efficient code, which is not so simple for usual vector-valued functions where memory has to be allocated on the heap (thus, the `Function::vector_value` function has to be given the address of an object into which the result is to be written, in order to avoid copying and memory allocation and deallocation on the heap). In addition to the known size, it is possible not only to return vectors, but also tensors of higher rank; however, this is not very often requested by applications, to be honest...

The interface of the `TensorFunction` class is relatively close to that of the `Function` class, so there is probably no need to comment in detail the following declaration:

```
template <int dim>
class AdvectionField : public TensorFunction<1, dim> {
public:
    AdvectionField() : TensorFunction<1, dim>() {}
    virtual Tensor<1, dim> value(const Point<dim>& p) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                           std::vector<Tensor<1, dim>& values) const;
```

In previous examples, we have used assertions that throw exceptions in several places. However, we have never seen how such exceptions are declared. This can be done as follows:

```
DeclException2(ExcDimensionMismatch, unsigned int, unsigned int,
               << "The vector has size " << arg1 << " but should have "
               << arg2 << " elements.");
```

The syntax may look a little strange, but is reasonable. The format is basically as follows: use the name of one of the macros `DeclExceptionN`, where `N` denotes the number of additional parameters which the exception object shall take. In this case, as we want to throw the exception when the sizes of two vectors differ, we need two arguments, so we use `DeclException2`. The first parameter then describes the name of the exception, while the following declare the data types of the parameters. The last argument is a sequence of output directives that will be piped into the `std::cerr` object, thus the strange format with the leading `<<` operator and the like. Note that we can access the parameters which are passed to the exception upon construction (i.e. within the `Assert` call) by using the names `arg1` through `argN`, where `N` is the number of arguments as defined by the use of the respective macro `DeclExceptionN`.

To learn how the preprocessor expands this macro into actual code, please refer to the documentation of the exception classes in the base library. Suffice it to say that by this macro call, the respective exception class is declared, which also has error output functions already implemented.

```
};
```

The following two functions implement the interface described above. The first simply implements the function as described in the introduction, while the second uses the same trick to avoid calling a virtual function as has already been introduced in the previous example program. Note the check for the right sizes of the arguments in the second function, which should always be present in such functions; it is our experience that many if not most programming errors result from incorrectly initialized arrays, incompatible parameters to functions and the like; using assertion as in this case can eliminate many of these problems.

```
template <int dim>
Tensor<1, dim> AdvectionField<dim>::value(const Point<dim>& p) const
{
    Point<dim> value;
    value[0] = 2;
    for (unsigned int i = 1; i < dim; ++i)
        value[i] = 1 + 0.8 * std::sin(8 * numbers::PI * p[0]);
    return value;
}

template <int dim>
void AdvectionField<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<Tensor<1, dim>& values) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = AdvectionField<dim>::value(points[i]);
}
```

Besides the advection field, we need two functions describing the source terms (right hand side) and the boundary values. First for the right hand side, which follows the same pattern as in previous examples. As described in the introduction, the source is a constant function in the vicinity of a source point, which we denote by the constant static variable `center_point`. We set the values of this center using the same template tricks as we have shown

in the step-7 example program. The rest is simple and has been shown previously, including the way to avoid virtual function calls in the `value_list` function.

```
template <int dim>
class RightHandSide : public Function<dim> {
public:
    RightHandSide() : Function<dim>() {}
    virtual double value(const Point<dim>& p,
                        const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                          std::vector<double>& values,
                          const unsigned int component = 0) const;
private:
    static const Point<dim> center_point;
};
template <>
const Point<1> RightHandSide<1>::center_point = Point<1>(-0.75);
template <>
const Point<2> RightHandSide<2>::center_point = Point<2>(-0.75, -0.75);
template <>
const Point<3> RightHandSide<3>::center_point = Point<3>(-0.75, -0.75, -0.75);
```

The only new thing here is that we check for the value of the `component` parameter. As this is a scalar function, it is obvious that it only makes sense if the desired component has the index zero, so we assert that this is indeed the case. `ExcIndexRange` is a global predefined exception (probably the one most often used, we therefore made it global instead of local to some class), that takes three parameters: the index that is outside the allowed range, the first element of the valid range and the one past the last (i.e. again the half-open interval so often used in the C++ standard library):

```
template <int dim>
double RightHandSide<dim>::value(const Point<dim>& p,
                                const unsigned int component) const
{
    (void)component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
    const double diameter = 0.1;
    return ((p - center_point).norm_square() < diameter * diameter
            ? .1 / std::pow(diameter, dim)
            : 0);
}
template <int dim>
void RightHandSide<dim>::value_list(const std::vector<Point<dim>& points,
                                   std::vector<double>& values,
                                   const unsigned int component) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = RightHandSide<dim>::value(points[i], component);
}
```

Finally for the boundary values, which is just another class derived from the `Function` base class:

```
template <int dim>
class BoundaryValues : public Function<dim> {
public:
    BoundaryValues() : Function<dim>() {}
    virtual double value(const Point<dim>& p,
                        const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                          std::vector<double>& values,
                          const unsigned int component = 0) const;
};
template <int dim>
double BoundaryValues<dim>::value(const Point<dim>& p,
                                const unsigned int component) const
{
    (void)component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
    const double sine_term =
        std::sin(16 * numbers::PI * std::sqrt(p.norm_square()));
    const double weight = std::exp(-5 * p.norm_square()) / std::exp(-5.);
    return sine_term * weight;
}
template <int dim>
void BoundaryValues<dim>::value_list(const std::vector<Point<dim>& points,
                                   std::vector<double>& values,
                                   const unsigned int component) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = BoundaryValues<dim>::value(points[i], component);
}
```

GradientEstimation class declaration

Now, finally, here comes the class that will compute the difference approximation of the gradient on each cell and weighs that with a power of the mesh size, as described in the introduction. This class is a simple version of the `DerivativeApproximation` class in the library, that uses similar techniques to obtain finite difference approximations of the gradient of a finite element field, or of higher derivatives.

The class has one public static function `estimate` that is called to compute a vector of error indicators, and a few private functions that do the actual work on all active cells. As in other parts of the library, we follow an informal convention to use vectors of floats for error indicators rather than the common vectors of doubles, as the additional accuracy is not necessary for estimated values.

In addition to these two functions, the class declares two exceptions which are raised when a cell has no neighbors in each of the space directions (in which case the matrix described in the introduction would be singular and can't be inverted), while the other one is used in the more common case of invalid parameters to a function, namely a vector of wrong size.

Two other comments: first, the class has no non-static member functions or variables, so this is not really a class, but rather serves the purpose of a `namespace` in C++. The reason that we chose a class over a namespace is that this way we can declare functions that are private. This can be done with namespaces as well, if one declares some functions in header files in the namespace and implements these and other functions in the implementation file. The functions not declared in the header file are still in the namespace but are not callable from outside. However, as we have only one file here, it is not possible to hide functions in the present case.

The second comment is that the dimension template parameter is attached to the function rather than to the class itself. This way, you don't have to specify the template parameter yourself as in most other cases, but the compiler can figure its value out itself from the dimension of the DoF handler object that one passes as first argument.

Before jumping into the fray with the implementation, let us also comment on the parallelization strategy. We have already introduced the necessary framework for using the `WorkStream` concept in the declaration of the main class of this program above. We will use it again here. In the current context, this means that we have to define (i) classes for scratch and copy objects, (ii) a function that does the local computation on one cell, and (iii) a function that copies the local result into a global object. Given this general framework, we will, however, deviate from it a bit. In particular, `WorkStream` was generally invented for cases where each local computation on a cell *adds* to a global object – for example, when assembling linear systems where we add local contributions into a global matrix and right hand side. `WorkStream` is designed to handle the potential conflict of multiple threads trying to do this addition at the same time, and consequently has to provide for some way to ensure that only thread gets to do this at a time. Here, however, the situation is slightly different: we compute contributions from every cell individually, but then all we need to do is put them into an element of an output vector that is unique to each cell. Consequently, there is no risk that the write operations from two cells might conflict, and the elaborate machinery of `WorkStream` to avoid conflicting writes is not necessary. Consequently, what we will do is this: We still need a scratch object that holds, for example, the `FEValues` object. However, we only create a fake, empty copy data structure. Likewise, we do need the function that computes local contributions, but since it can already put the result into its final location, we do not need a copy-local-to-global function and will instead give the `WorkStream::run()` function an empty function object – the equivalent to a NULL function pointer.

```
class GradientEstimation {
public:
    template <int dim>
    static void estimate(const DoFHandler<dim>& dof,
                       const Vector<double>& solution,
                       Vector<float>& error_per_cell);
    DeclException2(ExcInvalidVectorLength, int, int,
                  « "Vector has length " « arg1 « ", but should have "
                  « arg2);
    DeclException0(ExcInsufficientDirections);
private:
    template <int dim>
    struct EstimateScratchData {
        EstimateScratchData(const FiniteElement<dim>& fe,
                           const Vector<double>& solution,
                           Vector<float>& error_per_cell);
        EstimateScratchData(const EstimateScratchData& data);
        FEValues<dim> fe_midpoint_value;
        const Vector<double>& solution;
        Vector<float>& error_per_cell;
    };
};
```

```

};
struct EstimateCopyData {};
template <int dim>
static void estimate_cell(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    EstimateScratchData<dim>& scratch_data,
    const EstimateCopyData& copy_data);
};

```

AdvectionProblem class implementation

Now for the implementation of the main class. Constructor, destructor and the function `setup_system` follow the same pattern that was used previously, so we need not comment on these three function:

```

template <int dim>
AdvectionProblem<dim>::AdvectionProblem() : dof_handler(triangulation), fe(1)
{}
template <int dim>
AdvectionProblem<dim>::AdvectionProblem()
{
    dof_handler.clear();
}
template <int dim>
void AdvectionProblem<dim>::setup_system()
{
    dof_handler.distribute_dofs(fe);
    hanging_node_constraints.clear();
    DoFTools::make_hanging_node_constraints(dof_handler,
                                           hanging_node_constraints);

    hanging_node_constraints.close();
    DynamicSparsityPattern dsp(dof_handler.n_dofs(), dof_handler.n_dofs());
    DoFTools::make_sparsity_pattern(dof_handler, dsp, hanging_node_constraints,
                                   /* *keep_constrained_dofs = */ true);
    sparsity_pattern.copy_from(dsp);
    system_matrix.reinit(sparsity_pattern);
    solution.reinit(dof_handler.n_dofs());
    system_rhs.reinit(dof_handler.n_dofs());
}

```

In the following function, the matrix and right hand side are assembled. As stated in the documentation of the main class above, it does not do this itself, but rather delegates to the function following next, utilizing the `WorkStream` concept discussed in threads .

If you have looked through the threads module, you will have seen that assembling in parallel does not take an incredible amount of extra code as long as you diligently describe what the scratch and copy data objects are, and if you define suitable functions for the local assembly and the copy operation from local contributions to global objects. This done, the following will do all the heavy lifting to get these operations done on multiple threads on as many cores as you have in your system:

```

template <int dim>
void AdvectionProblem<dim>::assemble_system()
{
    WorkStream::run(dof_handler.begin_active(), dof_handler.end(), *this,
                   &AdvectionProblem::local_assemble_system,
                   &AdvectionProblem::copy_local_to_global,
                   AssemblyScratchData(fe), AssemblyCopyData());
}

```

After the matrix has been assembled in parallel, we still have to eliminate hanging node constraints. This is something that can't be done on each of the threads separately, so we have to do it now. Note also, that unlike in previous examples, there are no boundary conditions to be applied to the system of equations. This, of course, is due to the fact that we have included them into the weak formulation of the problem.

```

    hanging_node_constraints.condense(system_matrix);
    hanging_node_constraints.condense(system_rhs);
}

```

As already mentioned above, we need to have scratch objects for the parallel computation of local contributions. These objects contain `FEValues` and `FEFaceValues` objects, and so we will need to have constructors and copy constructors that allow us to create them. In initializing them, note first that we use bilinear elements, so Gauss formulae with two points in each space direction are sufficient. For the cell terms we need the values and gradients of the shape functions, the quadrature points in order to determine the source density and the advection field at a given point, and the weights of the quadrature points times the determinant of the Jacobian at these points. In contrast, for the boundary integrals, we don't need the gradients, but rather the normal vectors to the cells. This determines which update flags we will have to pass to the constructors of the members of the class:


```

template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const FiniteElement<dim>& fe)
:   fe_values(fe, QGauss<dim>(2),
        update_values | update_gradients | update_quadrature_points |
        update_JxW_values),
    fe_face_values(fe, QGauss<dim - 1>(2),
        update_values | update_quadrature_points |
        update_JxW_values | update_normal_vectors)
{}

template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const AssemblyScratchData& scratch_data)
:   fe_values(scratch_data.fe_values.get_fe(),
        scratch_data.fe_values.get_quadrature(),
        update_values | update_gradients | update_quadrature_points |
        update_JxW_values),
    fe_face_values(scratch_data.fe_face_values.get_fe(),
        scratch_data.fe_face_values.get_quadrature(),
        update_values | update_quadrature_points |
        update_JxW_values | update_normal_vectors)
{}

```

Now, this is the function that does the actual work. It is not very different from the `assemble_system` functions of previous example programs, so we will again only comment on the differences. The mathematical stuff follows closely what we have said in the introduction.

There are a number of points worth mentioning here, though. The first one is that we have moved the `FEValues` and `FEFaceValues` objects into the `ScratchData` object. We have done so because the alternative would have been to simply create one every time we get into this function – i.e., on every cell. It now turns out that the `FEValues` classes were written with the explicit goal of moving everything that remains the same from cell to cell into the construction of the object, and only do as little work as possible in `FEValues::reinit()` whenever we move to a new cell. What this means is that it would be very expensive to create a new object of this kind in this function as we would have to do it for every cell – exactly the thing we wanted to avoid with the `FEValues` class. Instead, what we do is create it only once (or a small number of times) in the scratch objects and then re-use it as often as we can.

This begs the question of whether there are other objects we create in this function whose creation is expensive compared to its use. Indeed, at the top of the function, we declare all sorts of objects. The `AdvectionField`, `RightHandSide` and `BoundaryValues` do not cost much to create, so there is no harm here. However, allocating memory in creating the `rhs_values` and similar variables below typically costs a significant amount of time, compared to just accessing the (temporary) values we store in them. Consequently, these would be candidates for moving into the `AssemblyScratchData` class. We will leave this as an exercise.

```

template <int dim>
void AdvectionProblem<dim>::local_assemble_system(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    AssemblyScratchData& scratch_data, AssemblyCopyData& copy_data)
{

```

First of all, we will need some objects that describe boundary values, right hand side function and the advection field. As we will only perform actions on these objects that do not change them, we declare them as constant, which can enable the compiler in some cases to perform additional optimizations.

```

const AdvectionField<dim> advection_field;
const RightHandSide<dim> right_hand_side;
const BoundaryValues<dim> boundary_values;

```

Then we define some abbreviations to avoid unnecessarily long lines:

```

const unsigned int dofs_per_cell = fe.dofs_per_cell;
const unsigned int n_q_points =
    scratch_data.fe_values.get_quadrature().size();
const unsigned int n_face_q_points =
    scratch_data.fe_face_values.get_quadrature().size();

```

We declare cell matrix and cell right hand side...

```

copy_data.cell_matrix.reinit(dofs_per_cell, dofs_per_cell);
copy_data.cell_rhs.reinit(dofs_per_cell);

```

... an array to hold the global indices of the degrees of freedom of the cell on which we are presently working...

```

copy_data.local_dof_indices.resize(dofs_per_cell);

```

... and array in which the values of right hand side, advection direction, and boundary values will be stored, for cell and face integrals respectively:

```
std::vector<double> rhs_values(n_q_points);
std::vector<Tensor<1, dim>> advection_directions(n_q_points);
std::vector<double> face_boundary_values(n_face_q_points);
std::vector<Tensor<1, dim>> face_advection_directions(n_face_q_points);
```

... then initialize the FEValues object...

```
scratch_data.fe_values.reinit(cell);
```

... obtain the values of right hand side and advection directions at the quadrature points...

```
advection_field.value_list(scratch_data.fe_values.get_quadrature_points(),
                           advection_directions);
right_hand_side.value_list(scratch_data.fe_values.get_quadrature_points(),
                           rhs_values);
```

... set the value of the streamline diffusion parameter as described in the introduction...

```
const double delta = 0.1 * cell->diameter();
```

... and assemble the local contributions to the system matrix and right hand side as also discussed above:

```
for (unsigned int q_point = 0; q_point < n_q_points; ++q_point)
  for (unsigned int i = 0; i < dofs_per_cell; ++i) {
    for (unsigned int j = 0; j < dofs_per_cell; ++j)
      copy_data.cell_matrix(i, j) +=
        ((advection_directions[q_point] *
          scratch_data.fe_values.shape_grad(j, q_point) *
          (scratch_data.fe_values.shape_value(i, q_point) +
            delta *
              (advection_directions[q_point] *
                scratch_data.fe_values.shape_grad(i, q_point)))) *
          scratch_data.fe_values.JxW(q_point));
    copy_data.cell_rhs(i) +=
      ((scratch_data.fe_values.shape_value(i, q_point) +
        delta * (advection_directions[q_point] *
          scratch_data.fe_values.shape_grad(i, q_point))) *
        rhs_values[q_point] * scratch_data.fe_values.JxW(q_point));
  }
```

Besides the cell terms which we have built up now, the bilinear form of the present problem also contains terms on the boundary of the domain. Therefore, we have to check whether any of the faces of this cell are on the boundary of the domain, and if so assemble the contributions of this face as well. Of course, the bilinear form only contains contributions from the inflow part of the boundary, but to find out whether a certain part of a face of the present cell is part of the inflow boundary, we have to have information on the exact location of the quadrature points and on the direction of flow at this point; we obtain this information using the FEFaceValues object and only decide within the main loop whether a quadrature point is on the inflow boundary.

```
for (unsigned int face = 0; face < GeometryInfo<dim>::faces_per_cell;
      ++face)
  if (cell->face(face)->at_boundary()) {
```

Ok, this face of the present cell is on the boundary of the domain. Just as for the usual FEValues object which we have used in previous examples and also above, we have to reinitialize the FEFaceValues object for the present face:

```
scratch_data.fe_face_values.reinit(cell, face);
```

For the quadrature points at hand, we ask for the values of the inflow function and for the direction of flow:

```
boundary_values.value_list(
  scratch_data.fe_face_values.get_quadrature_points(),
  face_boundary_values);
advection_field.value_list(
  scratch_data.fe_face_values.get_quadrature_points(),
  face_advection_directions);
```

Now loop over all quadrature points and see whether it is on the inflow or outflow part of the boundary. This is determined by a test whether the advection direction points inwards or outwards of the domain (note that the normal vector points outwards of the cell, and since the cell is at the boundary, the normal vector points outward of the domain, so if the advection direction points into the domain, its scalar product with the normal vector must be negative):

```
for (unsigned int q_point = 0; q_point < n_face_q_points; ++q_point)
  if (scratch_data.fe_face_values.normal_vector(q_point) *
      face_advection_directions[q_point] <
      0)
```

If the is part of the inflow boundary, then compute the contributions of this face to the global matrix and right hand side, using the values obtained from the FEFaceValues object and the formulae discussed in the introduction:

```

    for (unsigned int i = 0; i < dofs_per_cell; ++i) {
        for (unsigned int j = 0; j < dofs_per_cell; ++j)
            copy_data.cell_matrix(i, j) -=
                (face_advection_directions[q_point] *
                 scratch_data.fe_face_values.normal_vector(
                     q_point) *
                 scratch_data.fe_face_values.shape_value(
                     i, q_point) *
                 scratch_data.fe_face_values.shape_value(
                     j, q_point) *
                 scratch_data.fe_face_values.JxW(q_point));
        copy_data.cell_rhs(i) -=
            (face_advection_directions[q_point] *
             scratch_data.fe_face_values.normal_vector(
                 q_point) *
             face_boundary_values[q_point] *
             scratch_data.fe_face_values.shape_value(i,
                                                         q_point) *
             scratch_data.fe_face_values.JxW(q_point));
    }
}

```

Now go on by transferring the local contributions to the system of equations into the global objects. The first step was to obtain the global indices of the degrees of freedom on this cell.

```

cell->get_dof_indices(copy_data.local_dof_indices);
}

```

The second function we needed to write was the one that copies the local contributions the previous function has computed and put into the copy data object, into the global matrix and right hand side vector objects. This is essentially what we always had as the last block of code when assembling something on every cell. The following should therefore be pretty obvious:

```

template <int dim>
void AdvectionProblem<dim>::copy_local_to_global(
    const AssemblyCopyData& copy_data)
{
    for (unsigned int i = 0; i < copy_data.local_dof_indices.size(); ++i) {
        for (unsigned int j = 0; j < copy_data.local_dof_indices.size(); ++j)
            system_matrix.add(copy_data.local_dof_indices[i],
                              copy_data.local_dof_indices[j],
                              copy_data.cell_matrix(i, j));
        system_rhs(copy_data.local_dof_indices[i]) += copy_data.cell_rhs(i);
    }
}

```

Following is the function that solves the linear system of equations. As the system is no more symmetric positive definite as in all the previous examples, we can't use the Conjugate Gradients method anymore. Rather, we use a solver that is tailored to nonsymmetric systems like the one at hand, the BiCGStab method. As preconditioner, we use the Block Jacobi method.

```

template <int dim>
void AdvectionProblem<dim>::solve()
{

```

Assert that the system be symmetric.

```

Assert(system_matrix.m() == system_matrix.n(), ExcNotQuadratic());
auto num_rows = system_matrix.m();

```

Make a copy of the rhs to use with Ginkgo.

```

std::vector<double> rhs(num_rows);
std::copy(system_rhs.begin(), system_rhs.begin() + num_rows, rhs.begin());

```

Ginkgo setup Some shortcuts: A vector is a Dense matrix with co-dimension 1. The matrix is setup in CSR. But various formats can be used. Look at Ginkgo's documentation.

```

using vec = gko::matrix::Dense<>;
using mtx = gko::matrix::Csr<>;
using bicgstab = gko::solver::Bicgstab<>;
using bj = gko::preconditioner::Jacobi<>;
using val_array = gko::array<double>;

```

Where the code is to be executed. Can be changed to `omp` or `cuda` to run on multiple threads or on gpu's

```

std::shared_ptr<gko::Executor> exec = gko::ReferenceExecutor::create();

```

Setup Ginkgo's data structures

```

auto b = vec::create(exec, gko::dim<2>(num_rows, 1),
                    val_array::view(exec, num_rows, rhs.data()), 1);
auto x = vec::create(exec, gko::dim<2>(num_rows, 1));
auto A = mtx::create(exec, gko::dim<2>(num_rows),

```

```

        system_matrix.n_nonzero_elements());
mtx::value_type* values = A->get_values();
mtx::index_type* row_ptr = A->get_row_ptrs();
mtx::index_type* col_idx = A->get_col_idxs();

```

Convert to standard CSR format As deal.ii does not expose its system matrix pointers, we construct them individually.

```

row_ptr[0] = 0;
for (auto row = 1; row <= num_rows; ++row) {
    row_ptr[row] = row_ptr[row - 1] + system_matrix.get_row_length(row - 1);
}
std::vector<mtx::index_type> ptrs(num_rows + 1);
std::copy(A->get_row_ptrs(), A->get_row_ptrs() + num_rows + 1,
          ptrs.begin());
for (auto row = 0; row < system_matrix.m(); ++row) {
    for (auto p = system_matrix.begin(row); p != system_matrix.end(row);
         ++p) {

```

write entry into the first free one for this row

```

col_idx[ptrs[row]] = p->column();
values[ptrs[row]] = p->value();

```

then move pointer ahead

```

        ++ptrs[row];
    }
}

```

Ginkgo solve The stopping criteria is set at maximum iterations of 1000 and a reduction factor of 1e-12. For other options, refer to Ginkgo's documentation.

```

auto solver_gen =
    bicgstab::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1000),
            gko::stop::ResidualNorm<>::build().with_reduction_factor(1e-12))
        .with_preconditioner(bj::build())
        .on(exec);
auto solver = solver_gen->generate(gko::give(A));

```

Solve system

```

solver->apply(b, x);

```

Copy the solution vector back to deal.ii's data structures.

```

std::copy(x->get_values(), x->get_values() + num_rows, solution.begin());
/ *****
* deal.ii internal solver. Here for reference.
SolverControl          solver_control (1000, 1e-12);
SolverBicgstab<>       bicgstab (solver_control);
PreconditionJacobi<>   preconditioner;
preconditioner.initialize(system_matrix, 1.0);
bicgstab.solve (system_matrix, solution, system_rhs,
                preconditioner);
***** /

```

Distribute the possible hanging nodes constraints

```

    hanging_node_constraints.distribute(solution);
}

```

The following function refines the grid according to the quantity described in the introduction. The respective computations are made in the class GradientEstimation. The only difference to previous examples is that we refine a little more aggressively (0.5 instead of 0.3 of the number of cells).

```

template <int dim>
void AdvectionProblem<dim>::refine_grid()
{
    Vector<float> estimated_error_per_cell(triangulation.n_active_cells());
    GradientEstimation::estimate(dof_handler, solution,
                                estimated_error_per_cell);
    GridRefinement::refine_and_coarsen_fixed_number(
        triangulation, estimated_error_per_cell, 0.5, 0.03);
    triangulation.execute_coarsening_and_refinement();
}

```

Writing output to disk is done in the same way as in the previous examples. Indeed, the function is identical to the one in step-6.

```

template <int dim>
void AdvectionProblem<dim>::output_results(const unsigned int cycle)const
{

```

```

    {
        GridOut grid_out;
        std::ofstream output("grid-" + std::to_string(cycle) + ".eps");
        grid_out.write_eps(triangulation, output);
    }
    {
        DataOut<dim> data_out;
        data_out.attach_dof_handler(dof_handler);
        data_out.add_data_vector(solution, "solution");
        data_out.build_patches();
        std::ofstream output("solution-" + std::to_string(cycle) + ".vtk");
        data_out.write_vtk(output);
    }
}

```

... as is the main loop (setup – solve – refine)

```

template <int dim>
void AdvectionProblem<dim>::run()
{
    for (unsigned int cycle = 0; cycle < 6; ++cycle) {
        std::cout << "Cycle " << cycle << ':' << std::endl;
        if (cycle == 0) {
            GridGenerator::hyper_cube(triangulation, -1, 1);
            triangulation.refine_global(4);
        } else {
            refine_grid();
        }
        std::cout << "    Number of active cells: " <<
            << triangulation.n_active_cells() << std::endl;
        setup_system();
        std::cout << "    Number of degrees of freedom: " << dof_handler.n_dofs()
            << std::endl;
        assemble_system();
        solve();
        output_results(cycle);
    }
}

```

GradientEstimation class implementation

Now for the implementation of the GradientEstimation class. Let us start by defining constructors for the EstimateScratchData class used by the estimate_cell() function:

```

template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const FiniteElement<dim>& fe, const Vector<double>& solution,
    Vector<float>& error_per_cell)
: fe_midpoint_value(fe, QMidpoint<dim>()),
  update_values | update_quadrature_points(),
  solution(solution),
  error_per_cell(error_per_cell)
{}

template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const EstimateScratchData& scratch_data)
: fe_midpoint_value(scratch_data.fe_midpoint_value.get_fe(),
    scratch_data.fe_midpoint_value.get_quadrature(),
    update_values | update_quadrature_points(),
    solution(scratch_data.solution),
    error_per_cell(scratch_data.error_per_cell)
{}

```

Next for the implementation of the GradientEstimation class. The first function does not much except for delegating work to the other function, but there is a bit of setup at the top.

Before starting with the work, we check that the vector into which the results are written has the right size. Programming mistakes in which one forgets to size arguments correctly at the calling site are quite common. Because the resulting damage from not catching such errors is often subtle (e.g., corruption of data somewhere in memory, or non-reproducible results), it is well worth the effort to check for such things.

```

template <int dim>
void GradientEstimation::estimate(const DoFHandler<dim>& dof_handler,
    const Vector<double>& solution,
    Vector<float>& error_per_cell)
{
    Assert(error_per_cell.size() ==
        dof_handler.get_triangulation().n_active_cells(),
        ExcInvalidVectorLength(
            error_per_cell.size(),

```

```

        dof_handler.get_triangulation().n_active_cells());
    WorkStream::run(dof_handler.begin_active(), dof_handler.end(),
        &GradientEstimation::template estimate_cell<dim>,
        std::function<void(const EstimateCopyData&)>(),
        EstimateScratchData<dim>(dof_handler.get_fe(), solution,
            error_per_cell),
        EstimateCopyData());
}

```

Following now the function that actually computes the finite difference approximation to the gradient. The general outline of the function is to first compute the list of active neighbors of the present cell and then compute the quantities described in the introduction for each of the neighbors. The reason for this order is that it is not a one-liner to find a given neighbor with locally refined meshes. In principle, an optimized implementation would find neighbors and the quantities depending on them in one step, rather than first building a list of neighbors and in a second step their contributions but we will gladly leave this as an exercise. As discussed before, the worker function passed to `WorkStream::run` works on "scratch" objects that keep all temporary objects. This way, we do not need to create and initialize objects that are expensive to initialize within the function that does the work, every time it is called for a given cell. Such an argument is passed as the second argument. The third argument would be a "copy-data" object (see threads for more information) but we do not actually use any of these here. Because `WorkStream::run()` insists on passing three arguments, we declare this function with three arguments, but simply ignore the last one.

(This is unsatisfactory from an esthetic perspective. It can be avoided, at the cost of some other trickery. If you allow, let us here show how. First, assume that we had declared this function to only take two arguments by omitting the unused last one. Now, `WorkStream::run` still wants to call this function with three arguments, so we need to find a way to "forget" the third argument in the call. Simply passing `WorkStream::run` the pointer to the function as we do above will not do this – the compiler will complain that a function declared to have two arguments is called with three arguments. However, we can do this by passing the following as the third argument when calling `WorkStream::run()` above:

```

std::function<void (const typename DoFHandler<dim>::active_cell_iterator
&,
                    EstimateScratchData<dim> &,
                    EstimateCopyData &)>
    (std::bind (&GradientEstimation::template estimate_cell<dim>,
        std::placeholders::_1,
        std::placeholders::_2))

```

This creates a function object taking three arguments, but when it calls the underlying function object, it simply only uses the first and second argument – we simply "forget" to use the third argument :-). In the end, this isn't completely obvious either, and so we didn't implement it, but hey – it can be done!

Now for the details:

```

template <int dim>
void GradientEstimation::estimate_cell(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    EstimateScratchData<dim>& scratch_data, const EstimateCopyData&)
{

```

We need space for the tensor $\mathbf{Y}$, which is the sum of outer products of the $\mathbf{y}$ -vectors.

```
Tensor<2, dim> Y;
```

Then we allocate a vector to hold iterators to all active neighbors of a cell. We reserve the maximal number of active neighbors in order to avoid later reallocations. Note how this maximal number of active neighbors is computed here.

```

std::vector<typename DoFHandler<dim>::active_cell_iterator>
    active_neighbors;
active_neighbors.reserve(GeometryInfo<dim>::faces_per_cell *
    GeometryInfo<dim>::max_children_per_face);

```

First initialize the `FEValues` object, as well as the $\mathbf{Y}$ tensor:

```
scratch_data.fe_midpoint_value.reinit(cell);
```

Then allocate the vector that will be the sum over the $\mathbf{y}$ -vectors times the approximate directional derivative:

```
Tensor<1, dim> projected_gradient;
```

Now before going on first compute a list of all active neighbors of the present cell. We do so by first looping over all faces and see whether the neighbor there is active, which would be the case if it is on the same level as the present cell or one level coarser (note that a neighbor can only be once coarser than the present cell, as we only allow a

maximal difference of one refinement over a face in deal.II). Alternatively, the neighbor could be on the same level and be further refined; then we have to find which of its children are next to the present cell and select these (note that if a child of a neighbor of an active cell that is next to this active cell, needs necessarily be active itself, due to the one-refinement rule cited above).

Things are slightly different in one space dimension, as there the one-refinement rule does not exist: neighboring active cells may differ in as many refinement levels as they like. In this case, the computation becomes a little more difficult, but we will explain this below.

Before starting the loop over all neighbors of the present cell, we have to clear the array storing the iterators to the active neighbors, of course.

```
active_neighbors.clear();
for (unsigned int face_no = 0; face_no < GeometryInfo<dim>::faces_per_cell;
    ++face_no)
    if (!cell->at_boundary(face_no)) {
```

First define an abbreviation for the iterator to the face and the neighbor

```
const typename DoFHandler<dim>::face_iterator face =
    cell->face(face_no);
const typename DoFHandler<dim>::cell_iterator neighbor =
    cell->neighbor(face_no);
```

Then check whether the neighbor is active. If it is, then it is on the same level or one level coarser (if we are not in 1D), and we are interested in it in any case.

```
if (neighbor->active())
    active_neighbors.push_back(neighbor);
else {
```

If the neighbor is not active, then check its children.

```
if (dim == 1) {
```

To find the child of the neighbor which bounds to the present cell, successively go to its right child if we are left of the present cell ($n==0$), or go to the left child if we are on the right ($n==1$), until we find an active cell.

```
typename DoFHandler<dim>::cell_iterator neighbor_child =
    neighbor;
while (neighbor_child->has_children())
    neighbor_child =
        neighbor_child->child(face_no == 0 ? 1 : 0);
```

As this used some non-trivial geometrical intuition, we might want to check whether we did it right, i.e. check whether the neighbor of the cell we found is indeed the cell we are presently working on. Checks like this are often useful and have frequently uncovered errors both in algorithms like the line above (where it is simple to involuntarily exchange $n==1$ for $n==0$ or the like) and in the library (the assumptions underlying the algorithm above could either be wrong, wrongly documented, or are violated due to an error in the library). One could in principle remove such checks after the program works for some time, but it might be a good thing to leave it in anyway to check for changes in the library or in the algorithm above.

Note that if this check fails, then this is certainly an error that is irrecoverable and probably qualifies as an internal error. We therefore use a predefined exception class to throw here.

```
Assert(
    neighbor_child->neighbor(face_no == 0 ? 1 : 0) == cell,
    ExcInternalError());
```

If the check succeeded, we push the active neighbor we just found to the stack we keep:

```
    active_neighbors.push_back(neighbor_child);
} else
```

If we are not in 1d, we collect all neighbor children 'behind' the subfaces of the current face

```
    for (unsigned int subface_no = 0;
        subface_no < face->n_children(); ++subface_no)
        active_neighbors.push_back(
            cell->neighbor_child_on_subface(face_no,
                                           subface_no));
    }
```

OK, now that we have all the neighbors, let's start the computation on each of them. First we do some preliminaries: find out about the center of the present cell and the solution at this point. The latter is obtained as a vector of

function values at the quadrature points, of which there are only one, of course. Likewise, the position of the center is the position of the first (and only) quadrature point in real space.

```
const Point<dim> this_center =
    scratch_data.fe_midpoint_value.quadrature_point(0);
std::vector<double> this_midpoint_value(1);
scratch_data.fe_midpoint_value.get_function_values(scratch_data.solution,
                                                    this_midpoint_value);
```

Now loop over all active neighbors and collect the data we need. Allocate a vector just like `this_midpoint_value` which we will use to store the value of the solution in the midpoint of the neighbor cell. We allocate it here already, since that way we don't have to allocate memory repeatedly in each iteration of this inner loop (memory allocation is a rather expensive operation):

```
std::vector<double> neighbor_midpoint_value(1);
typename std::vector<typename DoFHandler<dim>::active_cell_iterator>::
    const_iterator neighbor_ptr = active_neighbors.begin();
for (; neighbor_ptr != active_neighbors.end(); ++neighbor_ptr) {
```

First define an abbreviation for the iterator to the active neighbor cell:

```
const typename DoFHandler<dim>::active_cell_iterator neighbor =
    *neighbor_ptr;
```

Then get the center of the neighbor cell and the value of the finite element function thereon. Note that for this information we have to reinitialize the `FEValues` object for the neighbor cell.

```
scratch_data.fe_midpoint_value.reinit(neighbor);
const Point<dim> neighbor_center =
    scratch_data.fe_midpoint_value.quadrature_point(0);
scratch_data.fe_midpoint_value.get_function_values(
    scratch_data.solution, neighbor_midpoint_value);
```

Compute the vector $\mathbf{y}$ connecting the centers of the two cells. Note that as opposed to the introduction, we denote by $\mathbf{y}$ the normalized difference vector, as this is the quantity used everywhere in the computations.

```
Tensor<1, dim> y = neighbor_center - this_center;
const double distance = y.norm();
y /= distance;
```

Then add up the contribution of this cell to the $\mathbf{Y}$ matrix...

```
for (unsigned int i = 0; i < dim; ++i)
    for (unsigned int j = 0; j < dim; ++j) Y[i][j] += y[i] * y[j];
```

... and update the sum of difference quotients:

```
projected_gradient +=
    (neighbor_midpoint_value[0] - this_midpoint_value[0]) / distance *
    y;
}
```

If now, after collecting all the information from the neighbors, we can determine an approximation of the gradient for the present cell, then we need to have passed over vectors $\mathbf{y}$ which span the whole space, otherwise we would not have all components of the gradient. This is indicated by the invertibility of the matrix.

If the matrix should not be invertible, this means that the present cell had an insufficient number of active neighbors. In contrast to all previous cases, where we raised exceptions, this is, however, not a programming error: it is a runtime error that can happen in optimized mode even if it ran well in debug mode, so it is reasonable to try to catch this error also in optimized mode. For this case, there is the `AssertThrow` macro: it checks the condition like the `Assert` macro, but not only in debug mode; it then outputs an error message, but instead of terminating the program as in the case of the `Assert` macro, the exception is thrown using the `throw` command of C++. This way, one has the possibility to catch this error and take reasonable counter actions. One such measure would be to refine the grid globally, as the case of insufficient directions can not occur if every cell of the initial grid has been refined at least once.

```
AssertThrow(determinant(Y) != 0, ExcInsufficientDirections());
```

If, on the other hand the matrix is invertible, then invert it, multiply the other quantity with it and compute the estimated error using this quantity and the right powers of the mesh width:

```
const Tensor<2, dim> Y_inverse = invert(Y);
Tensor<1, dim> gradient = Y_inverse * projected_gradient;
```

The last part of this function is the one where we write into the element of the output vector what we have just computed. The address of this vector has been stored in the scratch data object, and all we have to do is know how to get at the correct element inside this vector – but we can ask the cell we're on the how-manyth active cell it is for this:

```
scratch_data.error_per_cell(cell->active_cell_index()) =
    (std::pow(cell->diameter(), 1 + 1.0 * dim / 2) *
     std::sqrt(gradient.norm_square()));
}
// namespace Step9
```


Main function

The main function is similar to the previous examples. The main difference is that we use `MultithreadInfo` to set the maximum number of threads (see [Parallel computing with multiple processors accessing shared memory](#) documentation module for more explanation). The number of threads used is the minimum of the environment variable `DEAL_II_NUM_THREADS` and the parameter of `set_thread_limit`. If no value is given to `set_thread_limit`, the default value from the Intel Threading Building Blocks (TBB) library is used. If the call to `set_thread_limit` is omitted, the number of threads will be chosen by TBB independently of `DEAL_II_NUM_THREADS`.

```
int main()
{
    try {
        dealii::MultithreadInfo::set_thread_limit();
        Step9::AdvectionProblem<2> advection_problem_2d;
        advection_problem_2d.run();
    } catch (std::exception& exc) {
        std::cerr << std::endl
                  << std::endl
                  << "-----"
                  << std::endl;
        std::cerr << "Exception on processing: " << std::endl
                  << exc.what() << std::endl
                  << "Aborting!" << std::endl
                  << "-----"
                  << std::endl;
        return 1;
    } catch (...) {
        std::cerr << std::endl
                  << std::endl
                  << "-----"
                  << std::endl;
        std::cerr << "Unknown exception!" << std::endl
                  << "Aborting!" << std::endl
                  << "-----"
                  << std::endl;
        return 1;
    }
    return 0;
}
```

Results

Comments about programming and debugging

The plain program

```
/* -----
 *
 * Copyright (C) 2000 - 2018 by the deal.II authors
 *
 * This file is part of the deal.II library.
 *
 * The deal.II library is free software; you can use it, redistribute
 * it, and/or modify it under the terms of the GNU Lesser General
 * Public License as published by the Free Software Foundation; either
 * version 2.1 of the License, or (at your option) any later version.
 * The full text of the license can be found in the file LICENSE at
 * the top level of the deal.II distribution.
 *
 * -----
 *
 * Author:  Wolfgang Bangerth, University of Heidelberg, 2000
 */
/* -----
 *
 * This file has been taken verbatim from the deal.ii (version 9.0)
 * examples directory and modified.
 *
 * This example aims to demonstrate the ease with which Ginkgo can
 * be interfaced with other libraries. The only modification/ addition
 * has been to the AdvectionProblem::solve () function.
 *
 */
```

```

*/
#include <deal.II/base/function.h>
#include <deal.II/base/logstream.h>
#include <deal.II/base/quadrature_lib.h>
#include <deal.II/dofs/dof_accessor.h>
#include <deal.II/dofs/dof_handler.h>
#include <deal.II/dofs/dof_tools.h>
#include <deal.II/fe/fe_q.h>
#include <deal.II/fe/fe_values.h>
#include <deal.II/grid/grid_generator.h>
#include <deal.II/grid/grid_out.h>
#include <deal.II/grid/grid_refinement.h>
#include <deal.II/grid/tria.h>
#include <deal.II/grid/tria_accessor.h>
#include <deal.II/grid/tria_iterator.h>
#include <deal.II/lac/constraint_matrix.h>
#include <deal.II/lac/dynamic_sparsity_pattern.h>
#include <deal.II/lac/full_matrix.h>
#include <deal.II/lac/precondition.h>
#include <deal.II/lac/solver_bicgstab.h>
#include <deal.II/lac/sparse_matrix.h>
#include <deal.II/lac/vector.h>
#include <deal.II/numerics/data_out.h>
#include <deal.II/numerics/matrix_tools.h>
#include <deal.II/numerics/vector_tools.h>
#include <deal.II/base/multithread_info.h>
#include <deal.II/base/work_stream.h>
#include <deal.II/base/tensor_function.h>
#include <deal.II/numerics/error_estimator.h>
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
namespace Step9 {
using namespace dealii;
template <int dim>
class AdvectionProblem {
public:
    AdvectionProblem();
    AdvectionProblem();
    void run();
private:
    void setup_system();
    struct AssemblyScratchData {
        AssemblyScratchData(const FiniteElement<dim>& fe);
        AssemblyScratchData(const AssemblyScratchData& scratch_data);
        FEValues<dim> fe_values;
        FEFaceValues<dim> fe_face_values;
    };
    struct AssemblyCopyData {
        FullMatrix<double> cell_matrix;
        Vector<double> cell_rhs;
        std::vector<types::global_dof_index> local_dof_indices;
    };
    void assemble_system();
    void local_assemble_system(
        const typename DoFHandler<dim>::active_cell_iterator& cell,
        AssemblyScratchData& scratch, AssemblyCopyData& copy_data);
    void copy_local_to_global(const AssemblyCopyData& copy_data);
    void solve();
    void refine_grid();
    void output_results(const unsigned int cycle) const;
    Triangulation<dim> triangulation;
    DoFHandler<dim> dof_handler;
    FE_Q<dim> fe;
    ConstraintMatrix hanging_node_constraints;
    SparsityPattern sparsity_pattern;
    SparseMatrix<double> system_matrix;
    Vector<double> solution;
    Vector<double> system_rhs;
};
template <int dim>
class AdvectionField : public TensorFunction<1, dim> {
public:
    AdvectionField() : TensorFunction<1, dim>() {}
    virtual Tensor<1, dim> value(const Point<dim>& p) const;
    virtual void value_list(const std::vector<Point<dim>& points,
        std::vector<Tensor<1, dim>& values) const;
    DeclException2(ExcDimensionMismatch, unsigned int, unsigned int,
        « "The vector has size " « arg1 « " but should have "
        « arg2 « " elements.");
};
template <int dim>
Tensor<1, dim> AdvectionField<dim>::value(const Point<dim>& p) const
{
    Point<dim> value;
    value[0] = 2;
    for (unsigned int i = 1; i < dim; ++i)

```

```

        value[i] = 1 + 0.8 * std::sin(8 * numbers::PI * p[0]);
    }
    return value;
}
template <int dim>
void AdvectionField<dim>::value_list(const std::vector<Point<dim>& points,
                                     std::vector<Tensor<1, dim>& values) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = AdvectionField<dim>::value(points[i]);
}
template <int dim>
class RightHandSide : public Function<dim> {
public:
    RightHandSide() : Function<dim>() {}
    virtual double value(const Point<dim>& p,
                        const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                           std::vector<double>& values,
                           const unsigned int component = 0) const;
private:
    static const Point<dim> center_point;
};
template <>
const Point<1> RightHandSide<1>::center_point = Point<1>(-0.75);
template <>
const Point<2> RightHandSide<2>::center_point = Point<2>(-0.75, -0.75);
template <>
const Point<3> RightHandSide<3>::center_point = Point<3>(-0.75, -0.75, -0.75);
template <int dim>
double RightHandSide<dim>::value(const Point<dim>& p,
                                const unsigned int component) const
{
    (void)component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
    const double diameter = 0.1;
    return ((p - center_point).norm_square() < diameter * diameter
           ? .1 / std::pow(diameter, dim)
           : 0);
}
template <int dim>
void RightHandSide<dim>::value_list(const std::vector<Point<dim>& points,
                                    std::vector<double>& values,
                                    const unsigned int component) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = RightHandSide<dim>::value(points[i], component);
}
template <int dim>
class BoundaryValues : public Function<dim> {
public:
    BoundaryValues() : Function<dim>() {}
    virtual double value(const Point<dim>& p,
                        const unsigned int component = 0) const;
    virtual void value_list(const std::vector<Point<dim>& points,
                           std::vector<double>& values,
                           const unsigned int component = 0) const;
};
template <int dim>
double BoundaryValues<dim>::value(const Point<dim>& p,
                                const unsigned int component) const
{
    (void)component;
    Assert(component == 0, ExcIndexRange(component, 0, 1));
    const double sine_term =
        std::sin(16 * numbers::PI * std::sqrt(p.norm_square()));
    const double weight = std::exp(-5 * p.norm_square()) / std::exp(-5.);
    return sine_term * weight;
}
template <int dim>
void BoundaryValues<dim>::value_list(const std::vector<Point<dim>& points,
                                    std::vector<double>& values,
                                    const unsigned int component) const
{
    Assert(values.size() == points.size(),
           ExcDimensionMismatch(values.size(), points.size()));
    for (unsigned int i = 0; i < points.size(); ++i)
        values[i] = BoundaryValues<dim>::value(points[i], component);
}
class GradientEstimation {
public:
    template <int dim>
    static void estimate(const DoFHandler<dim>& dof,
                       const Vector<double>& solution,

```

```

        Vector<float>& error_per_cell);
DeclException2(ExcInvalidVectorLength, int, int,
    « "Vector has length " « arg1 « ", but should have "
    « arg2);
DeclException0(ExcInsufficientDirections);
private:
template <int dim>
struct EstimateScratchData {
    EstimateScratchData(const FiniteElement<dim>& fe,
        const Vector<double>& solution,
        Vector<float>& error_per_cell);
    EstimateScratchData(const EstimateScratchData& data);
    FEValues<dim> fe_midpoint_value;
    const Vector<double>& solution;
    Vector<float>& error_per_cell;
};
struct EstimateCopyData {};
template <int dim>
static void estimate_cell(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    EstimateScratchData<dim>& scratch_data,
    const EstimateCopyData& copy_data);
};
template <int dim>
AdvectionProblem<dim>::AdvectionProblem() : dof_handler(triangulation), fe(1)
{}
template <int dim>
AdvectionProblem<dim>::AdvectionProblem()
{
    dof_handler.clear();
}
template <int dim>
void AdvectionProblem<dim>::setup_system()
{
    dof_handler.distribute_dofs(fe);
    hanging_node_constraints.clear();
    DoFTools::make_hanging_node_constraints(dof_handler,
        hanging_node_constraints);
    hanging_node_constraints.close();
    DynamicSparsityPattern dsp(dof_handler.n_dofs(), dof_handler.n_dofs());
    DoFTools::make_sparsity_pattern(dof_handler, dsp, hanging_node_constraints,
        /*keep_constrained_dofs = */ true);
    sparsity_pattern.copy_from(dsp);
    system_matrix.reinit(sparsity_pattern);
    solution.reinit(dof_handler.n_dofs());
    system_rhs.reinit(dof_handler.n_dofs());
}
template <int dim>
void AdvectionProblem<dim>::assemble_system()
{
    WorkStream::run(dof_handler.begin_active(), dof_handler.end(), *this,
        &AdvectionProblem::local_assemble_system,
        &AdvectionProblem::copy_local_to_global,
        AssemblyScratchData(fe), AssemblyCopyData());
    hanging_node_constraints.condense(system_matrix);
    hanging_node_constraints.condense(system_rhs);
}
template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const FiniteElement<dim>& fe)
: fe_values(fe, QGauss<dim>(2),
    update_values | update_gradients | update_quadrature_points |
    update_JxW_values),
    fe_face_values(fe, QGauss<dim - 1>(2),
    update_values | update_quadrature_points |
    update_JxW_values | update_normal_vectors)
{}
template <int dim>
AdvectionProblem<dim>::AssemblyScratchData::AssemblyScratchData(
    const AssemblyScratchData& scratch_data)
: fe_values(scratch_data.fe_values.get_fe(),
    scratch_data.fe_values.get_quadrature(),
    update_values | update_gradients | update_quadrature_points |
    update_JxW_values),
    fe_face_values(scratch_data.fe_face_values.get_fe(),
    scratch_data.fe_face_values.get_quadrature(),
    update_values | update_quadrature_points |
    update_JxW_values | update_normal_vectors)
{}
template <int dim>
void AdvectionProblem<dim>::local_assemble_system(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    AssemblyScratchData& scratch_data, AssemblyCopyData& copy_data)
{
    const AdvectionField<dim> advection_field;
    const RightHandSide<dim> right_hand_side;
    const BoundaryValues<dim> boundary_values;

```

```

const unsigned int dofs_per_cell = fe.dofs_per_cell;
const unsigned int n_q_points =
    scratch_data.fe_values.get_quadrature().size();
const unsigned int n_face_q_points =
    scratch_data.fe_face_values.get_quadrature().size();
copy_data.cell_matrix.reinit(dofs_per_cell, dofs_per_cell);
copy_data.cell_rhs.reinit(dofs_per_cell);
copy_data.local_dof_indices.resize(dofs_per_cell);
std::vector<double> rhs_values(n_q_points);
std::vector<Tensor<1, dim>> advection_directions(n_q_points);
std::vector<double> face_boundary_values(n_face_q_points);
std::vector<Tensor<1, dim>> face_advection_directions(n_face_q_points);
scratch_data.fe_values.reinit(cell);
advection_field.value_list(scratch_data.fe_values.get_quadrature_points(),
    advection_directions);
right_hand_side.value_list(scratch_data.fe_values.get_quadrature_points(),
    rhs_values);
const double delta = 0.1 * cell->diameter();
for (unsigned int q_point = 0; q_point < n_q_points; ++q_point)
    for (unsigned int i = 0; i < dofs_per_cell; ++i) {
        for (unsigned int j = 0; j < dofs_per_cell; ++j)
            copy_data.cell_matrix(i, j) +=
                ((advection_directions[q_point] *
                    scratch_data.fe_values.shape_grad(j, q_point) *
                    (scratch_data.fe_values.shape_value(i, q_point) +
                        delta *
                            (advection_directions[q_point] *
                                scratch_data.fe_values.shape_grad(i, q_point)))) *
                    scratch_data.fe_values.JxW(q_point))) *
                copy_data.cell_rhs(i) +=
                    ((scratch_data.fe_values.shape_value(i, q_point) +
                        delta * (advection_directions[q_point] *
                            scratch_data.fe_values.shape_grad(i, q_point))) *
                        rhs_values[q_point] * scratch_data.fe_values.JxW(q_point));
    }
for (unsigned int face = 0; face < GeometryInfo<dim>::faces_per_cell;
    ++face)
    if (cell->face(face)->at_boundary()) {
        scratch_data.fe_face_values.reinit(cell, face);
        boundary_values.value_list(
            scratch_data.fe_face_values.get_quadrature_points(),
            face_boundary_values);
        advection_field.value_list(
            scratch_data.fe_face_values.get_quadrature_points(),
            face_advection_directions);
        for (unsigned int q_point = 0; q_point < n_face_q_points; ++q_point)
            if (scratch_data.fe_face_values.normal_vector(q_point) *
                face_advection_directions[q_point] <
                    0)
                for (unsigned int i = 0; i < dofs_per_cell; ++i) {
                    for (unsigned int j = 0; j < dofs_per_cell; ++j)
                        copy_data.cell_matrix(i, j) -=
                            (face_advection_directions[q_point] *
                                scratch_data.fe_face_values.normal_vector(
                                    q_point) *
                                    scratch_data.fe_face_values.shape_value(
                                        i, q_point) *
                                    scratch_data.fe_face_values.shape_value(
                                        j, q_point) *
                                    scratch_data.fe_face_values.JxW(q_point));
                        copy_data.cell_rhs(i) -=
                            (face_advection_directions[q_point] *
                                scratch_data.fe_face_values.normal_vector(
                                    q_point) *
                                    face_boundary_values[q_point] *
                                    scratch_data.fe_face_values.shape_value(i,
                                        q_point) *
                                    scratch_data.fe_face_values.JxW(q_point));
                    }
                }
        cell->get_dof_indices(copy_data.local_dof_indices);
    }
}
template <int dim>
void AdvectionProblem<dim>::copy_local_to_global(
    const AssemblyCopyData& copy_data)
{
    for (unsigned int i = 0; i < copy_data.local_dof_indices.size(); ++i) {
        for (unsigned int j = 0; j < copy_data.local_dof_indices.size(); ++j)
            system_matrix.add(copy_data.local_dof_indices[i],
                copy_data.local_dof_indices[j],
                copy_data.cell_matrix(i, j));
        system_rhs(copy_data.local_dof_indices[i]) += copy_data.cell_rhs(i);
    }
}
template <int dim>
void AdvectionProblem<dim>::solve()
{

```

```

Assert(system_matrix.m() == system_matrix.n(), ExcNotQuadratic());
auto num_rows = system_matrix.m();
std::vector<double> rhs(num_rows);
std::copy(system_rhs.begin(), system_rhs.begin() + num_rows, rhs.begin());
using vec = gko::matrix::Dense<>;
using mtx = gko::matrix::Csr<>;
using bicgstab = gko::solver::Bicgstab<>;
using bj = gko::preconditioner::Jacobi<>;
using val_array = gko::array<double>;
std::shared_ptr<gko::Executor> exec = gko::ReferenceExecutor::create();
auto b = vec::create(exec, gko::dim<2>(num_rows, 1),
    val_array::view(exec, num_rows, rhs.data(), 1);
auto x = vec::create(exec, gko::dim<2>(num_rows, 1));
auto A = mtx::create(exec, gko::dim<2>(num_rows),
    system_matrix.n_nonzero_elements());
mtx::value_type* values = A->get_values();
mtx::index_type* row_ptr = A->get_row_ptrs();
mtx::index_type* col_idx = A->get_col_idxs();
row_ptr[0] = 0;
for (auto row = 1; row <= num_rows; ++row) {
    row_ptr[row] = row_ptr[row - 1] + system_matrix.get_row_length(row - 1);
}
std::vector<mtx::index_type> ptrs(num_rows + 1);
std::copy(A->get_row_ptrs(), A->get_row_ptrs() + num_rows + 1,
    ptrs.begin());
for (auto row = 0; row < system_matrix.m(); ++row) {
    for (auto p = system_matrix.begin(row); p != system_matrix.end(row);
        ++p) {
        col_idx[ptrs[row]] = p->column();
        values[ptrs[row]] = p->value();
        ++ptrs[row];
    }
}
auto solver_gen =
    bicgstab::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(1000),
            gko::stop::ResidualNorm<>::build().with_reduction_factor(1e-12))
        .with_preconditioner(bj::build())
        .on(exec);
auto solver = solver_gen->generate(gko::give(A));
solver->apply(b, x);
std::copy(x->get_values(), x->get_values() + num_rows, solution.begin());
/*****
 * deal.ii internal solver. Here for reference.
 SolverControl          solver_control (1000, 1e-12);
 SolverBicgstab<>       bicgstab (solver_control);

PreconditionJacobi<> preconditioner;
preconditioner.initialize(system_matrix, 1.0);

bicgstab.solve (system_matrix, solution, system_rhs,
preconditioner);
*****/
hanging_node_constraints.distribute(solution);
}
template <int dim>
void AdvectionProblem<dim>::refine_grid()
{
    Vector<float> estimated_error_per_cell(triangulation.n_active_cells());
    GradientEstimation::estimate(dof_handler, solution,
        estimated_error_per_cell);
    GridRefinement::refine_and_coarsen_fixed_number(
        triangulation, estimated_error_per_cell, 0.5, 0.03);
    triangulation.execute_coarsening_and_refinement();
}
template <int dim>
void AdvectionProblem<dim>::output_results(const unsigned int cycle) const
{
    {
        GridOut grid_out;
        std::ofstream output("grid-" + std::to_string(cycle) + ".eps");
        grid_out.write_eps(triangulation, output);
    }
    {
        DataOut<dim> data_out;
        data_out.attach_dof_handler(dof_handler);
        data_out.add_data_vector(solution, "solution");
        data_out.build_patches();
        std::ofstream output("solution-" + std::to_string(cycle) + ".vtk");
        data_out.write_vtk(output);
    }
}
template <int dim>
void AdvectionProblem<dim>::run()
{
    for (unsigned int cycle = 0; cycle < 6; ++cycle) {

```

```

std::cout << "Cycle " << cycle << ':' << std::endl;
if (cycle == 0) {
    GridGenerator::hyper_cube(triangulation, -1, 1);
    triangulation.refine_global(4);
} else {
    refine_grid();
}
std::cout << "    Number of active cells:          "
    << triangulation.n_active_cells() << std::endl;
setup_system();
std::cout << "    Number of degrees of freedom: " << dof_handler.n_dofs()
    << std::endl;
assemble_system();
solve();
output_results(cycle);
}
}

template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const FiniteElement<dim>& fe, const Vector<double>& solution,
    Vector<float>& error_per_cell)
    : fe_midpoint_value(fe, QMidpoint<dim>()),
      update_values | update_quadrature_points(),
      solution(solution),
      error_per_cell(error_per_cell)
{}

template <int dim>
GradientEstimation::EstimateScratchData<dim>::EstimateScratchData(
    const EstimateScratchData& scratch_data)
    : fe_midpoint_value(scratch_data.fe_midpoint_value.get_fe(),
        scratch_data.fe_midpoint_value.get_quadrature(),
        update_values | update_quadrature_points()),
      solution(scratch_data.solution),
      error_per_cell(scratch_data.error_per_cell)
{}

template <int dim>
void GradientEstimation::estimate(const DoFHandler<dim>& dof_handler,
    const Vector<double>& solution,
    Vector<float>& error_per_cell)
{
    Assert(error_per_cell.size() ==
        dof_handler.get_triangulation().n_active_cells(),
        ExcInvalidVectorLength(
            error_per_cell.size(),
            dof_handler.get_triangulation().n_active_cells()));
    WorkStream::run(dof_handler.begin_active(), dof_handler.end(),
        &GradientEstimation::template estimate_cell<dim>,
        std::function<void(const EstimateCopyData&)>(),
        EstimateScratchData<dim>(dof_handler.get_fe(), solution,
            error_per_cell),
        EstimateCopyData());
}

template <int dim>
void GradientEstimation::estimate_cell(
    const typename DoFHandler<dim>::active_cell_iterator& cell,
    EstimateScratchData<dim>& scratch_data, const EstimateCopyData&)
{
    Tensor<2, dim> Y;
    std::vector<typename DoFHandler<dim>::active_cell_iterator>
        active_neighbors;
    active_neighbors.reserve(GeometryInfo<dim>::faces_per_cell *
        GeometryInfo<dim>::max_children_per_face);
    scratch_data.fe_midpoint_value.reinit(cell);
    Tensor<1, dim> projected_gradient;
    active_neighbors.clear();
    for (unsigned int face_no = 0; face_no < GeometryInfo<dim>::faces_per_cell;
        ++face_no)
        if (!cell->at_boundary(face_no)) {
            const typename DoFHandler<dim>::face_iterator face =
                cell->face(face_no);
            const typename DoFHandler<dim>::cell_iterator neighbor =
                cell->neighbor(face_no);
            if (neighbor->active())
                active_neighbors.push_back(neighbor);
            else {
                if (dim == 1) {
                    typename DoFHandler<dim>::cell_iterator neighbor_child =
                        neighbor;
                    while (neighbor_child->has_children())
                        neighbor_child =
                            neighbor_child->child(face_no == 0 ? 1 : 0);
                    Assert(
                        neighbor_child->neighbor(face_no == 0 ? 1 : 0) == cell,
                        ExcInternalError());
                    active_neighbors.push_back(neighbor_child);
                } else
                    for (unsigned int subface_no = 0;

```

```

        subface_no < face->n_children(); ++subface_no)
        active_neighbors.push_back(
            cell->neighbor_child_on_subface(face_no,
                                           subface_no));
    }
}
const Point<dim> this_center =
    scratch_data.fe_midpoint_value.quadrature_point(0);
std::vector<double> this_midpoint_value(1);
scratch_data.fe_midpoint_value.get_function_values(scratch_data.solution,
                                                  this_midpoint_value);

std::vector<double> neighbor_midpoint_value(1);
typename std::vector<typename DoFHandler<dim>::active_cell_iterator>::
    const_iterator neighbor_ptr = active_neighbors.begin();
for (; neighbor_ptr != active_neighbors.end(); ++neighbor_ptr) {
    const typename DoFHandler<dim>::active_cell_iterator neighbor =
        *neighbor_ptr;
    scratch_data.fe_midpoint_value.reinit(neighbor);
    const Point<dim> neighbor_center =
        scratch_data.fe_midpoint_value.quadrature_point(0);
    scratch_data.fe_midpoint_value.get_function_values(
        scratch_data.solution, neighbor_midpoint_value);
    Tensor<1, dim> y = neighbor_center - this_center;
    const double distance = y.norm();
    y /= distance;
    for (unsigned int i = 0; i < dim; ++i)
        for (unsigned int j = 0; j < dim; ++j) Y[i][j] += y[i] * y[j];
    projected_gradient +=
        (neighbor_midpoint_value[0] - this_midpoint_value[0]) / distance *
        y;
}
AssertThrow(determinant(Y) != 0, ExcInsufficientDirections());
const Tensor<2, dim> Y_inverse = invert(Y);
Tensor<1, dim> gradient = Y_inverse * projected_gradient;
scratch_data.error_per_cell(cell->active_cell_index()) =
    (std::pow(cell->diameter(), 1 + 1.0 * dim / 2) *
     std::sqrt(gradient.norm_square()));
}
} // namespace Step9
int main()
{
    try {
        dealii::MultithreadInfo::set_thread_limit();
        Step9::AdvectionProblem<2> advection_problem_2d;
        advection_problem_2d.run();
    } catch (std::exception& exc) {
        std::cerr << std::endl
                  << std::endl
                  << "-----"
                  << std::endl;
        std::cerr << "Exception on processing: " << std::endl
                  << exc.what() << std::endl
                  << "Aborting!" << std::endl
                  << "-----"
                  << std::endl;
        return 1;
    } catch (...) {
        std::cerr << std::endl
                  << std::endl
                  << "-----"
                  << std::endl;
        std::cerr << "Unknown exception!" << std::endl
                  << "Aborting!" << std::endl
                  << "-----"
                  << std::endl;
        return 1;
    }
    return 0;
}

```


Chapter 17

The file-config-solver program

The file config solver example..

This example depends on simple-solver.

This example shows how to use file to configure solver.

In this example, we first read in a matrix from a file. Then we read the config file to configure the solver. The example application reports the generation and run time of the solver.

The commented program

```
const std::string matrix_path = argc >= 4 ? argv[3] : "data/A.mtx";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(
             0, gko::ReferenceExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream(matrix_path), exec));
```

Create RHS as 1 and initial guess as 0

```
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
}
```

```

auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);

```

Calculate initial residual by overwriting b

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);

```

Copy b again

```

b->copy_from(host_b);

```

Read the json config file to configure the ginkgo solver. The following files, which are mapped to corresponding examples, are available cg.json: simple-solver blockjacobi-cg.json: preconditioned-solver ir.json: iterative-refinement parilu.json: ilu-preconditioned-solver (by using factoization parameter directly) pgm-multigrid-cg.json: multigrid-preconditioned-solver (set min_coarse_rows additionally due to this small example matrix) mixed-pgm-multigrid-cg.json: mixed-multigrid-preconditioned-solver (assuming there are always more than one level)

```

auto config = gko::ext::config::parse_json_file(configfile);

```

Create the registry, which allows passing the existing data into config This example does not use existing data.

```

auto reg = gko::config::registry();

```

Create the default type descriptor, which gives the default common type (value/index) for solver generation. If the solver does not specify value type, the solver will use these types.

```

auto td = gko::config::make_type_descriptor<ValueType, IndexType>();

```

generate the linopfactory on the given executors

```

auto solver_gen = gko::config::parse(config, reg, td).on(exec);

```

Create solver

```

const auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
const auto gen_toc = std::chrono::steady_clock::now();
const auto gen_time = std::chrono::duration_cast<std::chrono::milliseconds>(
    gen_toc - gen_tic);

```

Add logger

```

std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);

```

Solve system

```

exec->synchronize();
const auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
const auto toc = std::chrono::steady_clock::now();
const auto time =
    std::chrono::duration_cast<std::chrono::milliseconds>(toc - tic);

```

Print out the solver config

```

std::cout << "Config file: " << configfile << std::endl;
std::ifstream f(configfile);
std::cout << f.rdbuf() << std::endl;

```

Calculate residual

```

auto res = gko::as<vec>(logger->get_residual_norm());
std::cout << "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);

```

Print solver statistics

```

std::cout << "Solver iteration count: " << logger->get_num_iterations()
    << std::endl;
std::cout << "Solver generation time [ms]: "
    << static_cast<double>(gen_time.count()) << std::endl;
std::cout << "Solver execution time [ms]: "
    << static_cast<double>(time.count()) << std::endl;
std::cout << "Solver execution time per iteration[ms]: "
    << static_cast<double>(time.count()) /
        logger->get_num_iterations()
    << std::endl;
}

```

Results

This is the expected output:

```
Config file:  config/cg.json
{
  "type": "solver::Cg",
  "criteria": [
    {
      "type": "Iteration",
      "max_iters": 20
    },
    {
      "type": "ResidualNorm",
      "reduction_factor": 1e-7
    }
  ]
}
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
25.9808
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
58.79
Solver iteration count:          20
Solver generation time [ms]:    0.065244
Solver execution time [ms]:     0.793764
Solver execution time per iteration[ms]: 0.0396882
```

Comments about programming and debugging

The plain program

```
#include <chrono>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
#include <ginkgo/extensions/config/json_config.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    std::cout << argv[0] << " executor configfile matrix" << std::endl;
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const auto configfile = argc >= 3 ? argv[2] : "config/cg.json";
    const std::string matrix_path = argc >= 4 ? argv[3] : "data/A.mtx";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream(matrix_path), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    }
```

```

    }
    auto x = vec::create(exec);
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_b);
    auto config = gko::ext::config::parse_json_file(configfile);
    auto reg = gko::config::registry();
    auto td = gko::config::make_type_descriptor<ValueType, IndexType>();
    auto solver_gen = gko::config::parse(config, reg, td).on(exec);
    const auto gen_tic = std::chrono::steady_clock::now();
    auto solver = solver_gen->generate(A);
    exec->synchronize();
    const auto gen_toc = std::chrono::steady_clock::now();
    const auto gen_time = std::chrono::duration_cast<std::chrono::milliseconds>(
        gen_toc - gen_tic);
    std::shared_ptr<const gko::log::Convergence<ValueType> logger =
        gko::log::Convergence<ValueType>::create();
    solver->add_logger(logger);
    exec->synchronize();
    const auto tic = std::chrono::steady_clock::now();
    solver->apply(b, x);
    exec->synchronize();
    const auto toc = std::chrono::steady_clock::now();
    const auto time =
        std::chrono::duration_cast<std::chrono::milliseconds>(toc - tic);
    std::cout << "Config file: " << configfile << std::endl;
    std::ifstream f(configfile);
    std::cout << f.rdbuf() << std::endl;
    auto res = gko::as<vec>(logger->get_residual_norm());
    std::cout << "Initial residual norm sqrt(r^T r):  \n";
    write(std::cout, initres);
    std::cout << "Final residual norm sqrt(r^T r):  \n";
    write(std::cout, res);
    std::cout << "Solver iteration count:          " << logger->get_num_iterations()
        << std::endl;
    std::cout << "Solver generation time [ms]:  "
        << static_cast<double>(gen_time.count()) << std::endl;
    std::cout << "Solver execution time [ms]:  "
        << static_cast<double>(time.count()) << std::endl;
    std::cout << "Solver execution time per iteration[ms]:  "
        << static_cast<double>(time.count()) /
            logger->get_num_iterations()
        << std::endl;
}

```

Chapter 18

The ginkgo-overhead program

The ginkgo overhead measurement example..

Introduction

About the example

The commented program

```
num_iters = std::atol(argv[1]);
if (num_iters == 0) {
    print_usage_and_exit(argv[0]);
}
std::cout << gko::version_info::get() << std::endl;
auto exec = gko::ReferenceExecutor::create();
auto cg_factory =
    cg::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(num_iters))
        .on(exec);
auto A = gko::initialize<mtx>({1.0}, exec);
auto b = gko::initialize<vec>({std::nan("")}, exec);
auto x = gko::initialize<vec>({0.0}, exec);
auto tic = std::chrono::steady_clock::now();
auto solver = cg_factory->generate(gko::give(A));
solver->apply(x, b);
exec->synchronize();
auto tac = std::chrono::steady_clock::now();
auto time = std::chrono::duration_cast<std::chrono::nanoseconds>(tac - tic);
std::cout << "Running " << num_iters
    << " iterations of the CG solver took a total of "
    << static_cast<double>(time.count()) /
        static_cast<double>(std::nano::den)
    << " seconds." << std::endl
    << "\tAverage library overhead:          "
    << static_cast<double>(time.count()) /
        static_cast<double>(num_iters)
    << " [nanoseconds / iteration]" << std::endl;
}
```

Results

This is the expected output:

```
Running 1000000 iterations of the CG solver took a total of 1.60337 seconds.
Average library overhead:          1603.37 [nanoseconds / iteration]
```

Comments about programming and debugging

The plain program

```
#include <chrono>
#include <cmath>
#include <iostream>
#include <ginkgo/ginkgo.hpp>
[[noreturn]] void print_usage_and_exit(const char* name)
{
    std::cerr << "Usage: " << name << " [NUM_ITERS]" << std::endl;
    std::exit(-1);
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    long unsigned num_iters = 1000000;
    if (argc > 2) {
        print_usage_and_exit(argv[0]);
    }
    if (argc == 2) {
        num_iters = std::atol(argv[1]);
        if (num_iters == 0) {
            print_usage_and_exit(argv[0]);
        }
    }
    std::cout << gko::version_info::get() << std::endl;
    auto exec = gko::ReferenceExecutor::create();
    auto cg_factory =
        cg::build()
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(num_iters))
            .on(exec);
    auto A = gko::initialize<mtx>({1.0}, exec);
    auto b = gko::initialize<vec>({std::nan("")}, exec);
    auto x = gko::initialize<vec>({0.0}, exec);
    auto tic = std::chrono::steady_clock::now();
    auto solver = cg_factory->generate(gko::give(A));
    solver->apply(x, b);
    exec->synchronize();
    auto tac = std::chrono::steady_clock::now();
    auto time = std::chrono::duration_cast<std::chrono::nanoseconds>(tac - tic);
    std::cout << "Running " << num_iters
        << " iterations of the CG solver took a total of "
        << static_cast<double>(time.count()) /
            static_cast<double>(std::nano::den)
        << " seconds." << std::endl
        << "\tAverage library overhead: "
        << static_cast<double>(time.count()) /
            static_cast<double>(num_iters)
        << " [nanoseconds / iteration]" << std::endl;
}
```

Chapter 19

The ginkgo-ranges program

The ranges and accessor example..

Introduction

About the example

The commented program

a utility function for printing the factorization on screen

```
template <typename Accessor>
void print_lu(const gko::range<Accessor>& A)
{
    std::cout << std::setprecision(2) << std::fixed;
    std::cout << "L = [";
    for (int i = 0; i < A.length(0); ++i) {
        std::cout << "\n ";
        for (int j = 0; j < A.length(1); ++j) {
            std::cout << (i > j ? A(i, j) : (i == j) * 1.) << " ";
        }
        std::cout << "\n]\n\nU = [";
        for (int i = 0; i < A.length(0); ++i) {
            std::cout << "\n ";
            for (int j = 0; j < A.length(1); ++j) {
                std::cout << (i <= j ? A(i, j) : 0.) << " ";
            }
        }
        std::cout << "\n]" << std::endl;
    }
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
```

Print version information

```
std::cout << gko::version_info::get() << std::endl;
```

Create some test data, add some padding just to demonstrate how to use it with ranges. clang-format off

```
ValueType data[] = {
    2., 4., 5., -1.0,
    4., 11., 12., -1.0,
    6., 24., 24., -1.0
};
```

clang-format on

Create a 3-by-3 range, with a 2D row-major accessor using data as the underlying storage. Set the stride (a.k.a. "LDA") to 4.

```
auto A =
    gko::range<gko::accessor::row_major<ValueType, 2>>(data, 3u, 3u, 4u);
```

use the LU factorization routine defined above to factorize the matrix

```
factorize(A);
```

print the factorization on screen

```
    print_lu(A);
}
```

Results

This is the expected output:

```
L = [
  1.00 0.00 0.00
  2.00 1.00 0.00
  3.00 4.00 1.00
]
U = [
  2.00 4.00 5.00
  0.00 3.00 2.00
  0.00 0.00 1.00
]
```

Comments about programming and debugging

The plain program

```
#include <iomanip>
#include <iostream>
#include <ginkgo/ginkgo.hpp>
template <typename Accessor>
void factorize(const gko::range<Accessor>& A)
{
    using gko::span;
    assert(A.length(0) == A.length(1));
    for (gko::size_type i = 0; i < A.length(0) - 1; ++i) {
        const auto trail = span{i + 1, A.length(0)};
        A(trail, i) = A(trail, i) / A(i, i);
        A(trail, trail) = A(trail, trail) - mmul(A(trail, i), A(i, trail));
    }
}

template <typename Accessor>
void print_lu(const gko::range<Accessor>& A)
{
    std::cout << std::setprecision(2) << std::fixed;
    std::cout << "L = [";
    for (int i = 0; i < A.length(0); ++i) {
        std::cout << "\n ";
        for (int j = 0; j < A.length(1); ++j) {
            std::cout << (i > j ? A(i, j) : (i == j) * 1.) << " ";
        }
        std::cout << "\n]\nU = [";
        for (int i = 0; i < A.length(0); ++i) {
            std::cout << "\n ";
            for (int j = 0; j < A.length(1); ++j) {
                std::cout << (i <= j ? A(i, j) : 0.) << " ";
            }
        }
        std::cout << "\n]" << std::endl;
    }
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    std::cout << gko::version_info::get() << std::endl;
    ValueType data[] = {
        2., 4., 5., -1.0,
        4., 11., 12., -1.0,
        6., 24., 24., -1.0
    };
    auto A =
        gko::range<gko::accessor::row_major<ValueType, 2>>(data, 3u, 3u, 4u);
    factorize(A);
    print_lu(A);
}
```


Chapter 20

The heat-equation program

The heat equation example..

This example depends on simple-solver, three-pt-stencil-solver.

Introduction

This example solves a 2D heat conduction equation

$$u : [0, d]^2 \rightarrow \mathbb{R}$$
$$\partial_t u = \delta u + f$$

with Dirichlet boundary conditions and given initial condition and constant-in-time source function f .

The partial differential equation (PDE) is solved with a finite difference spatial discretization on an equidistant grid: For n grid points, and grid distance $h = 1/n$ we write

$$u_{i,j}' = \alpha \frac{u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1} - 4u_{i,j}}{h^2} + f_{i,j}$$

We then build an implicit Euler integrator by discretizing with time step τ

$$\frac{u_{i,j}^{k+1} - u_{i,j}^k}{\tau} = \alpha \frac{u_{i-1,j}^{k+1} + u_{i+1,j}^{k+1} - u_{i,j-1}^{k+1} - u_{i,j+1}^{k+1} + 4u_{i,j}^{k+1}}{h^2} + f_{i,j}$$

and solve the resulting linear system for u^{k+1} using Ginkgo's CG solver preconditioned with an incomplete Cholesky factorization for each time step, occasionally writing the resulting grid values into a video file using OpenCV and a custom color mapping.

The intention of this example is to provide a mini-app showing matrix assembly, vector initialization, solver setup and the use of Ginkgo in a more complex setting.

About the example

The commented program

The intention of [this](#) example is to provide a mini-app showing matrix assembly, vector initialization, [solver](#) setup and the use of Ginkgo in a more complex setting.

```
*****<DESCRIPTION>***** /
#include <chrono>
#include <fstream>
#include <iostream>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
#include <ginkgo/ginkgo.hpp>
```

This function implements a simple Ginkgo-themed clamped color mapping for values in the range [0,5].

```
void set_val(unsigned char* data, double value)
{
```

RGB values for the 6 colors used for values 0, 1, ..., 5 We will interpolate linearly between these values.

```
double col_r[] = {255, 221, 129, 201, 249, 255};
double col_g[] = {255, 220, 130, 161, 158, 204};
double col_b[] = {255, 220, 133, 93, 24, 8};
value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
auto d = std::max(0.0, std::min(1.0, value - i));
```

OpenCV uses BGR instead of RGB by default, revert indices

```
data[2] = static_cast<unsigned char>(col_r[i + 1] * d + col_r[i] * (1 - d));
data[1] = static_cast<unsigned char>(col_g[i + 1] * d + col_g[i] * (1 - d));
data[0] = static_cast<unsigned char>(col_b[i + 1] * d + col_b[i] * (1 - d));
}
```

Initialize video output with given dimension and FPS (frames per seconds)

```
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
{
    cv::Size videosize{n, n};
    auto output =
        std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', '1');
    output.first.open("heat.mp4", fourcc, fps, videosize);
    return output;
}
```

Write the current frame to video output using the above color mapping

```
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                    const double* data)
{
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
        for (int j = 0; j < n; j++) {
            set_val(&row[3 * j], data[i * n + j]);
        }
        output.first.write(output.second);
    }
}

int main(int argc, char* argv[])
{
    using mtx = gko::matrix::Csr<>;
    using vec = gko::matrix::Dense<>;
```

Problem parameters: simulation length

```
auto t0 = 5.0;
```

diffusion factor

```
auto diffusion = 0.0005;
```

scaling factor for heat source

```
auto source_scale = 2.5;
```

Simulation parameters: inner grid points per discretization direction

```
auto n = 256;
```

number of simulation steps per second

```
auto steps_per_sec = 500;
```

number of video frames per second

```
auto fps = 25;
```

number of grid points

```
auto n2 = n * n;
```

grid point distance (ignoring boundary points)

```
auto h = 1.0 / (n + 1);
auto h2 = h * h;
```

time step size for the simulation

```
auto tau = 1.0 / steps_per_sec;
```

create a CUDA executor with an associated OpenMP host executor

```
auto exec = gko::CudaExecutor::create(0, gko::OmpExecutor::create());
```

load heat source and initial state vectors

```
std::ifstream initial_stream("data/gko_logo_2d.mtx");
std::ifstream source_stream("data/gko_text_2d.mtx");
auto source = gko::read<vec>(source_stream, exec);
auto in_vector = gko::read<vec>(initial_stream, exec);
```

create output vector with initial guess for

```
auto out_vector = in_vector->clone();
```

create scalar for source update

```
auto tau_source_scalar = gko::initialize<vec>({source_scale * tau}, exec);
```

create stencil matrix as shared_ptr for solver

```
auto stencil_matrix = gko::share(mtx::create(exec));
```

assemble matrix

```
gko::matrix_data<> mtx_data{gko::dim<2>(n2, n2)};
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        auto c = i * n + j;
        auto c_val = diffusion * tau * 4.0 / h2 + 1.0;
        auto off_val = -diffusion * tau / h2;
```

for each grid point: insert 5 stencil points with Dirichlet boundary conditions, i.e. with zero boundary value

```
        if (i > 0) {
            mtx_data.nonzeros.emplace_back(c, c - n, off_val);
        }
        if (j > 0) {
            mtx_data.nonzeros.emplace_back(c, c - 1, off_val);
        }
        mtx_data.nonzeros.emplace_back(c, c, c_val);
        if (j < n - 1) {
            mtx_data.nonzeros.emplace_back(c, c + 1, off_val);
        }
        if (i < n - 1) {
            mtx_data.nonzeros.emplace_back(c, c + n, off_val);
        }
    }
}
stencil_matrix->read(mtx_data);
```

prepare video output

```
auto output = build_output(n, fps);
```

build CG solver on stencil with incomplete Cholesky preconditioner stopping at 1e-10 relative accuracy

```
auto solver =
    gko::solver::Cg<>::build()
        .with_preconditioner(gko::preconditioner::Ic<>::build())
        .with_criteria(gko::stop::ResidualNorm<>::build()
            .with_baseline(gko::stop::mode::rhs_norm)
            .with_reduction_factor(1e-10))
        .on(exec)
        ->generate(stencil_matrix);
```

time stamp of the last output frame (initialized to a sentinel value)

```
double last_t = -t0;
```

execute implicit Euler method: for each timestep, solve stencil system

```

for (double t = 0; t < t0; t += tau) {

if enough time has passed, output the next video frame
if (t - last_t > 1.0 / fps) {
    last_t = t;
    std::cout << t << std::endl;
    output_timestep(
        output, n,
        gko::make_temporary_clone(exec->get_master(), in_vector)
        ->get_const_values());
}

add heat source contribution
in_vector->add_scaled(tau_source_scalar, source);

execute Euler step
solver->apply(in_vector, out_vector);

swap input and output
    std::swap(in_vector, out_vector);
}
}

```

Results

The program will generate a video file named heat.mp4 and output the timestamp of each generated frame.

Comments about programming and debugging

The plain program

```

/*****<DESCRIPTION>*****/
This example solves a 2D heat conduction equation

u : [0, d]^2 \rightarrow R
\partial_t u = \Delta u + f

with Dirichlet boundary conditions and given initial condition and
constant-in-time source function f.

The partial differential equation (PDE) is solved with a finite difference
spatial discretization on an equidistant grid: For 'n' grid points,
and grid distance @f$h = 1/n@f$ we write

u_{i,j}' = \alpha (u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1}
- 4 u_{i,j}) / h^2
+ f_{i,j}

We then build an implicit Euler integrator by discretizing with time step @f$\tau@f$

(u_{i,j}^{k+1} - u_{i,j}^k) / \tau =
\alpha (u_{i-1,j}^{k+1} - u_{i+1,j}^{k+1}
+ u_{i,j-1}^{k+1} - u_{i,j+1}^{k+1} - 4 u_{i,j}^{k+1}) / h^2
+ f_{i,j}

and solve the resulting linear system for @f$u_{\cdot}^{k+1}@f$ using Ginkgo's CG
solver preconditioned with an incomplete Cholesky factorization for each time
step, occasionally writing the resulting grid values into a video file using
OpenCV and a custom color mapping.

The intention of this example is to provide a mini-app showing matrix assembly,
vector initialization, solver setup and the use of Ginkgo in a more complex
setting.
*****/
#include <chrono>
#include <fstream>
#include <iostream>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
#include <ginkgo/ginkgo.hpp>
void set_val(unsigned char* data, double value)
{

```

```

double col_r[] = {255, 221, 129, 201, 249, 255};
double col_g[] = {255, 220, 130, 161, 158, 204};
double col_b[] = {255, 220, 133, 93, 24, 8};
value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
auto d = std::max(0.0, std::min(1.0, value - i));
data[2] = static_cast<unsigned char>(col_r[i + 1] * d + col_r[i] * (1 - d));
data[1] = static_cast<unsigned char>(col_g[i + 1] * d + col_g[i] * (1 - d));
data[0] = static_cast<unsigned char>(col_b[i + 1] * d + col_b[i] * (1 - d));
}
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
{
    cv::Size videosize{n, n};
    auto output =
        std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', 'l');
    output.first.open("heat.mp4", fourcc, fps, videosize);
    return output;
}
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                    const double* data)
{
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
        for (int j = 0; j < n; j++) {
            set_val(&row[3 * j], data[i * n + j]);
        }
    }
    output.first.write(output.second);
}
int main(int argc, char* argv[])
{
    using mtx = gko::matrix::Csr<>;
    using vec = gko::matrix::Dense<>;
    auto t0 = 5.0;
    auto diffusion = 0.0005;
    auto source_scale = 2.5;
    auto n = 256;
    auto steps_per_sec = 500;
    auto fps = 25;
    auto n2 = n * n;
    auto h = 1.0 / (n + 1);
    auto h2 = h * h;
    auto tau = 1.0 / steps_per_sec;
    auto exec = gko::CudaExecutor::create(0, gko::OmpExecutor::create());
    std::ifstream initial_stream("data/gko_logo_2d.mtx");
    std::ifstream source_stream("data/gko_text_2d.mtx");
    auto source = gko::read<vec>(source_stream, exec);
    auto in_vector = gko::read<vec>(initial_stream, exec);
    auto out_vector = in_vector->clone();
    auto tau_source_scalar = gko::initialize<vec>({source_scale * tau}, exec);
    auto stencil_matrix = gko::share(mtx::create(exec));
    gko::matrix_data<> mtx_data{gko::dim<2>(n2, n2)};
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            auto c = i * n + j;
            auto c_val = diffusion * tau * 4.0 / h2 + 1.0;
            auto off_val = -diffusion * tau / h2;
            if (i > 0) {
                mtx_data.nonzeros.emplace_back(c, c - n, off_val);
            }
            if (j > 0) {
                mtx_data.nonzeros.emplace_back(c, c - 1, off_val);
            }
            mtx_data.nonzeros.emplace_back(c, c, c_val);
            if (j < n - 1) {
                mtx_data.nonzeros.emplace_back(c, c + 1, off_val);
            }
            if (i < n - 1) {
                mtx_data.nonzeros.emplace_back(c, c + n, off_val);
            }
        }
    }
    stencil_matrix->read(mtx_data);
    auto output = build_output(n, fps);
    auto solver =
        gko::solver::Cg<>::build()
        .with_preconditioner(gko::preconditioner::Ic<>::build())
        .with_criteria(gko::stop::ResidualNorm<>::build())
        .with_baseline(gko::stop::mode::rhs_norm)
        .with_reduction_factor(1e-10)
        .on(exec)
        ->generate(stencil_matrix);
    double last_t = -t0;
    for (double t = 0; t < t0; t += tau) {
        if (t - last_t > 1.0 / fps) {
            last_t = t;

```

```
        std::cout << t << std::endl;
        output_timestep(
            output, n,
            gko::make_temporary_clone(exec->get_master(), in_vector)
                ->get_const_values());
    }
    in_vector->add_scaled(tau_source_scalar, source);
    solver->apply(in_vector, out_vector);
    std::swap(in_vector, out_vector);
}
}
```

Chapter 21

The ilu-preconditioned-solver program

The ILU-preconditioned solver example..

This example depends on simple-solver.

Introduction

About the example

This example shows how to use incomplete factors generated via the ParILU algorithm to generate an incomplete factorization (ILU) preconditioner, how to specify the sparse triangular solves in the ILU preconditioner application, and how to generate an ILU-preconditioned solver and apply it to a specific problem.

The commented program

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
```

Generate incomplete factors using ParILU

```
auto par_ilu_fact =
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
```

Generate concrete factorization for input matrix

```
auto par_ilu = gko::share(par_ilu_fact->generate(A));
```

Generate an ILU preconditioner factory by setting lower and upper triangular solver - in this case the exact triangular solves

```
auto ilu_pre_factory =
    gko::preconditioner::Ilu<gko::solver::LowerTrs<ValueType, IndexType>,
        gko::solver::UpperTrs<ValueType, IndexType>,
        false>::build()
        .on(exec);
```

Use incomplete factors to generate ILU preconditioner

```
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
```

Use preconditioner inside GMRES solver factory Generating a solver factory tied to a specific preconditioner makes sense if there are several very similar systems to solve, and the same solver+preconditioner combination is expected to be effective.

```
const RealValueType reduction_factor{1e-7};
auto ilu_gmres_factory =
    gmres::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor))
        .with_generated_preconditioner(ilu_preconditioner)
        .on(exec);
```

Generate preconditioned solver for a specific target system

```
auto ilu_gmres = ilu_gmres_factory->generate(A);
```

Solve system

```
ilu_gmres->apply(b, x);
```

Print solution

```
std::cout << "Solution (x):\n";
write(std::cout, x);
```

Calculate residual

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```

Results

This is the expected output:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.46249e-08
```


Comments about programming and debugging

The plain program

```
#include <cstdlib>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using gmres = gko::solver::Gmres<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    auto par_ilu_fact =
        gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
    auto par_ilu = gko::share(par_ilu_fact->generate(A));
    auto ilu_pre_factory =
        gko::preconditioner::Ilu<gko::solver::LowerTrs<ValueType, IndexType>,
                                gko::solver::UpperTrs<ValueType, IndexType>,
                                false>::build()
        .on(exec);
    auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
    const RealValueType reduction_factor{1e-7};
    auto ilu_gmres_factory =
        gmres::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1000u),
                      gko::stop::ResidualNorm<ValueType>::build()
                      .with_reduction_factor(reduction_factor))
        .with_generated_preconditioner(ilu_preconditioner)
        .on(exec);
    auto ilu_gmres = ilu_gmres_factory->generate(A);
    ilu_gmres->apply(b, x);
    std::cout << "Solution (x):\n";
    write(std::cout, x);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(res);
    std::cout << "Residual norm sqrt(r^T r):\n";
    write(std::cout, res);
}
```


Chapter 22

The inverse-iteration program

The inverse iteration example..

This example depends on simple-solver, .

Introduction

This example shows how components available in Ginkgo can be used to implement higher-level numerical methods. The method used here will be the shifted inverse iteration method for eigenvalue computation which find the eigenvalue and eigenvector of A closest to z , for some scalar z . The method requires repeatedly solving the shifted linear system $(A - zI)x = b$, as well as performing matrix-vector products with the matrix A . Here is the complete pseudocode of the method:

```
x_0 = initial guess
for i = 0 .. max_iterations:
    solve (A - zI) y_i = x_i for y_i+1
    x_(i+1) = y_i / || y_i || # compute next eigenvector approximation
    g_(i+1) = x_(i+1)^* A x_(i+1) # approximate eigenvalue (Rayleigh quotient)
    if ||A x_(i+1) - g_(i+1)x_(i+1)|| < tol * g_(i+1): # check convergence
        break
```

About the example

The commented program

Figure out where to run the code

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
```

```

        gko::OmpExecutor::create());
    }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}};

```

executor where Ginkgo will perform the computation

```

const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto this_exec = exec->get_master();

```

linear system solver parameters

```

auto system_max_iterations = 100u;
auto system_residual_goal = real_precision{1e-16};

```

eigensolver parameters

```

auto max_iterations = 20u;
auto residual_goal = real_precision{1e-8};
auto z = precision{20.0, 2.0};

```

Read data

```

auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));

```

Generate shifted matrix $A - zI$

- we avoid duplicating memory by not storing both A and $A - zI$, but compute $A - zI$ on the fly by using Ginkgo's utilities for creating linear combinations of operators

```

auto one = share(gko::initialize<vec>({precision{1.0}}, exec));
auto neg_one = share(gko::initialize<vec>({-precision{1.0}}, exec));
auto neg_z = gko::initialize<vec>({-z}, exec);
auto system_matrix = share(gko::Combination<precision>::create(
    one, A, gko::initialize<vec>({-z}, exec),
    gko::matrix::Identity<precision>::create(exec, A->get_size()[0])));

```

Generate solver operator $(A - zI)^{-1}$

```

auto solver =
    solver_type::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(
            system_max_iterations),
            gko::stop::ResidualNorm<precision>::build()
                .with_reduction_factor(system_residual_goal))
        .on(exec)
        ->generate(system_matrix);

```

inverse iterations

start with guess $[1, 1, \dots, 1]$

```

auto x = [&] {
    auto work = vec::create(this_exec, gko::dim<2>{A->get_size()[0], 1});
    const auto n = work->get_size()[0];
    for (int i = 0; i < n; ++i) {
        work->get_values()[i] = precision{1.0} / sqrt(n);
    }
    return clone(exec, work);
}();
auto y = clone(x);
auto tmp = clone(x);
auto norm = gko::initialize<real_vec>({1.0}, exec);
auto inv_norm = clone(this_exec, one);
auto g = clone(one);
for (auto i = 0u; i < max_iterations; ++i) {
    std::cout << "{ ";

```

$(A - zI)y = x$

```

solver->apply(x, y);
system_matrix->apply(one, y, neg_one, x);
x->compute_norm2(norm);
std::cout << "\"system_residual\": "
    << clone(this_exec, norm)->get_values()[0] << ", ";
x->copy_from(y);

```

$x = y / \|y\|$

```

x->compute_norm2(norm);
inv_norm->get_values()[0] =
    real_precision{1.0} / clone(this_exec, norm)->get_values()[0];
x->scale(clone(exec, inv_norm));

```

```

g = x^A * A x
A->apply(x, tmp);
x->compute_dot(tmp, g);
auto g_val = clone(this_exec, g)->get_values()[0];
std::cout << "\"eigenvalue\": " << g_val << ", ";

||Ax - gx|| < tol * g
    auto v = gko::initialize<vec>({-g_val}, exec);
    tmp->add_scaled(v, x);
    tmp->compute_norm2(norm);
    auto res_val = clone(exec->get_master(), norm)->get_values()[0];
    std::cout << "\"residual\": " << res_val / g_val << ", " << std::endl;
    if (abs(res_val) < residual_goal * abs(g_val)) {
        break;
    }
}
}

```

Results

This is the expected output:

```

{ "system_residual": +1.61736920e-14, "eigenvalue": (+2.03741410e+01,-1.17744356e-16), "residual":
  (+2.92231055e-01,+1.68883476e-18) },
{ "system_residual": +4.98014795e-15, "eigenvalue": (+1.94878474e+01,+1.25948378e-15), "residual":
  (+7.94370276e-02,-5.13395071e-18) },
{ "system_residual": +3.39296916e-15, "eigenvalue": (+1.93282121e+01,-1.19329332e-15), "residual":
  (+4.11149623e-02,+2.53837290e-18) },
{ "system_residual": +3.35953656e-15, "eigenvalue": (+1.92638912e+01,+3.28657016e-16), "residual":
  (+2.34717040e-02,-4.00445585e-19) },
{ "system_residual": +2.91474009e-15, "eigenvalue": (+1.92409166e+01,+3.65597737e-16), "residual":
  (+1.34709547e-02,-2.55962367e-19) },
{ "system_residual": +3.09863953e-15, "eigenvalue": (+1.92331106e+01,-1.07919176e-15), "residual":
  (+7.72060707e-03,+4.33212063e-19) },
{ "system_residual": +2.31198069e-15, "eigenvalue": (+1.92305014e+01,-2.89755360e-16), "residual":
  (+4.42106625e-03,+6.66143651e-20) },
{ "system_residual": +3.02771202e-15, "eigenvalue": (+1.92296339e+01,+8.04259901e-16), "residual":
  (+2.53081312e-03,-1.05848687e-19) },
{ "system_residual": +2.02954523e-15, "eigenvalue": (+1.92293461e+01,+7.81834016e-16), "residual":
  (+1.44862114e-03,-5.88985854e-20) },
{ "system_residual": +2.31762332e-15, "eigenvalue": (+1.92292506e+01,-1.11718775e-16), "residual":
  (+8.29183451e-04,+4.81741912e-21) },
{ "system_residual": +8.12541038e-15, "eigenvalue": (+1.92292190e+01,-6.55606254e-16), "residual":
  (+4.74636702e-04,+1.61823936e-20) },
{ "system_residual": +2.77259926e-15, "eigenvalue": (+1.92292085e+01,+4.30588140e-16), "residual":
  (+2.71701077e-04,-6.08403935e-21) },
{ "system_residual": +8.87888675e-14, "eigenvalue": (+1.92292051e+01,+9.67936313e-18), "residual":
  (+1.55539937e-04,-7.82937998e-23) },
{ "system_residual": +2.85077117e-15, "eigenvalue": (+1.92292039e+01,-4.52923128e-16), "residual":
  (+8.90457139e-05,+2.09737561e-21) },
{ "system_residual": +6.46865302e-14, "eigenvalue": (+1.92292035e+01,+1.58710681e-17), "residual":
  (+5.09805252e-05,-4.20774259e-23) },
{ "system_residual": +4.18913713e-15, "eigenvalue": (+1.92292034e+01,+1.06839590e-15), "residual":
  (+2.91887365e-05,-1.62175862e-21) },
{ "system_residual": +1.06421578e-11, "eigenvalue": (+1.92292034e+01,+3.26089685e-17), "residual":
  (+1.67126561e-05,-2.83413965e-23) },
{ "system_residual": +2.97434420e-13, "eigenvalue": (+1.92292034e+01,-7.85427712e-16), "residual":
  (+9.56961199e-06,+3.90876227e-22) },
{ "system_residual": +1.63230281e-11, "eigenvalue": (+1.92292033e+01,+3.69307000e-16), "residual":
  (+5.47975753e-06,-1.05241636e-22) },
{ "system_residual": +6.14939758e-14, "eigenvalue": (+1.92292033e+01,+1.36057865e-15), "residual":
  (+3.13794996e-06,-2.22028320e-22) },

```

Comments about programming and debugging

The plain program

```

#include <cmath>
#include <complex>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>

```

```

int main(int argc, char* argv[])
{
    using precision = std::complex<double>;
    using real_precision = gko::remove_complex<precision>;
    using vec = gko::matrix::Dense<precision>;
    using real_vec = gko::matrix::Dense<real_precision>;
    using mtx = gko::matrix::Csr<precision>;
    using solver_type = gko::solver::Bicgstab<precision>;
    using std::abs;
    using std::sqrt;
    std::cout << gko::version_info::get() << std::endl;
    std::cout << std::scientific << std::setprecision(8) << std::showpos;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto this_exec = exec->get_master();
    auto system_max_iterations = 100u;
    auto system_residual_goal = real_precision{1e-16};
    auto max_iterations = 20u;
    auto residual_goal = real_precision{1e-8};
    auto z = precision{20.0, 2.0};
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto one = share(gko::initialize<vec>({precision{1.0}}, exec));
    auto neg_one = share(gko::initialize<vec>({-precision{1.0}}, exec));
    auto neg_z = gko::initialize<vec>({-z}, exec);
    auto system_matrix = share(gko::Combination<precision>::create(
        one, A, gko::initialize<vec>({-z}, exec),
        gko::matrix::Identity<precision>::create(exec, A->get_size()[0])));
    auto solver =
        solver_type::build()
            .with_criteria(gko::stop::Iteration::build().with_max_iters(
                system_max_iterations),
                gko::stop::ResidualNorm<precision>::build()
                    .with_reduction_factor(system_residual_goal))
            .on(exec)
            ->generate(system_matrix);
    auto x = [&] {
        auto work = vec::create(this_exec, gko::dim<2>(A->get_size()[0], 1));
        const auto n = work->get_size()[0];
        for (int i = 0; i < n; ++i) {
            work->get_values()[i] = precision{1.0} / sqrt(n);
        }
        return clone(exec, work);
    }();
    auto y = clone(x);
    auto tmp = clone(x);
    auto norm = gko::initialize<real_vec>({1.0}, exec);
    auto inv_norm = clone(this_exec, one);
    auto g = clone(one);
    for (auto i = 0u; i < max_iterations; ++i) {
        std::cout << "{ ";
        solver->apply(x, y);
        system_matrix->apply(one, y, neg_one, x);
        x->compute_norm2(norm);
        std::cout << "\nsystem_residual\": "
                  << clone(this_exec, norm)->get_values()[0] << ", ";
        x->copy_from(y);
        x->compute_norm2(norm);
        inv_norm->get_values()[0] =
            real_precision{1.0} / clone(this_exec, norm)->get_values()[0];
        x->scale(clone(exec, inv_norm));
        A->apply(x, tmp);
        x->compute_dot(tmp, g);
        auto g_val = clone(this_exec, g)->get_values()[0];
        std::cout << "\neigenvalue\": " << g_val << ", ";
        auto v = gko::initialize<vec>({-g_val}, exec);
    }
}

```

```
tmp->add_scaled(v, x);
tmp->compute_norm2(norm);
auto res_val = clone(exec->get_master(), norm)->get_values()[0];
std::cout << "\"residual\": " << res_val / g_val << " }," << std::endl;
if (abs(res_val) < residual_goal * abs(g_val)) {
    break;
}
}
```


Chapter 23

The ir-ilu-preconditioned-solver program

The IR-ILU preconditioned solver example..

This example depends on ilu-preconditioned-solver, iterative-refinement.

Introduction

About the example

This example shows how to combine iterative refinement with the adaptive precision block-Jacobi preconditioner in order to approximately solve the triangular systems occurring in ILU preconditioning. Using an adaptive precision block-Jacobi preconditioner matrix as inner solver for the iterative refinement method is equivalent to doing adaptive precision block-Jacobi relaxation in the triangular solves. This example roughly approximates the triangular solves with five adaptive precision block-Jacobi sweeps with a maximum block size of 16.

This example is motivated by "Multiprecision block-Jacobi for Iterative Triangular Solves" (Göbel, Anzt, Cojean, Flegar, Quintana-Ortí, Euro-Par 2020). The theory and a detailed analysis can be found there.

The commented program

```
std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int sweeps = argc == 3 ? std::atoi(argv[2]) : 5u;
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}};

executor where Ginkgo will perform the computation
```

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS and initial guess as 1

```
gko::size_type num_rows = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(num_rows, 1));
for (gko::size_type i = 0; i < num_rows; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto clone_x = gko::clone(exec, x);
```

Generate incomplete factors using ParILU

```
auto par_ilu_fact =
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
```

Generate concrete factorization for input matrix

```
auto par_ilu = gko::share(par_ilu_fact->generate(A));
```

Generate an iterative refinement factory to be used as a triangular solver in the preconditioner application. The generated method is equivalent to doing five block-Jacobi sweeps with a maximum block size of 16.

```
auto bj_factory = gko::share(
    bj::build()
        .with_max_block_size(16u)
        .with_storage_optimization(gko::precision_reduction::autodetect())
        .on(exec));
auto trisolve_factory =
    ir::build()
        .with_solver(bj_factory)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(sweeps))
        .on(exec);
```

Generate an ILU preconditioner factory by setting lower and upper triangular solver - in this case the previously defined iterative refinement method.

```
auto ilu_pre_factory = gko::preconditioner::Ilu<ir, ir>::build()
    .with_l_solver(gko::clone(trisolve_factory))
    .with_u_solver(gko::clone(trisolve_factory))
    .on(exec);
```

Use incomplete factors to generate ILU preconditioner

```
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
```

Create stopping criteria for Gmres

```
const RealValueType reduction_factor{1e-12};
auto iter_stop = gko::share(
    gko::stop::Iteration::build().with_max_iters(1000u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_reduction_factor(reduction_factor)
    .on(exec));
```

Use preconditioner inside GMRES solver factory Generating a solver factory tied to a specific preconditioner makes sense if there are several very similar systems to solve, and the same solver+preconditioner combination is expected to be effective.

```
auto ilu_gmres_factory =
    gmres::build()
        .with_criteria(iter_stop, tol_stop)
        .with_generated_preconditioner(ilu_preconditioner)
        .on(exec);
```

Generate preconditioned solver for a specific target system

```
auto ilu_gmres = ilu_gmres_factory->generate(A);
```

Add logger

```
std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
ilu_gmres->add_logger(logger);
```

Warmup run

```
ilu_gmres->apply(b, x);
```

Solve system 100 times and take the average time.

```
std::chrono::nanoseconds time(0);
for (int i = 0; i < 100; i++) {
    x->copy_from(clone_x);
    auto tic = std::chrono::high_resolution_clock::now();
    ilu_gmres->apply(b, x);
    auto toc = std::chrono::high_resolution_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
}
std::cout << "Using " << sweeps << " block-Jacobi sweeps.\n";
```

Print solution

```
std::cout << "Solution (x):\n";
write(std::cout, x);
```

Get residual

```
auto res = gko::as<vec>(logger->get_residual_norm());
std::cout << "GMRES iteration count: " << logger->get_num_iterations()
    << "\n";
std::cout << "GMRES execution time [ms]: "
    << static_cast<double>(time.count()) / 1000000000.0 << "\n";
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```

Results

This is the expected output:

```
Using 5 block-Jacobi sweeps.
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
GMRES iteration count:      8
GMRES execution time [ms]: 0.377673
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.65303e-12
```

Comments about programming and debugging

The plain program

```
#include <cstdlib>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
```

```

using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using gmres = gko::solver::Gmres<ValueType>;
using ir = gko::solver::Ir<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const unsigned int sweeps = argc == 3 ? std::atoi(argv[2]) : 5u;
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type num_rows = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(num_rows, 1));
for (gko::size_type i = 0; i < num_rows; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto clone_x = gko::clone(exec, x);
auto par_ilu_fact =
    gko::factorization::ParIlu<ValueType, IndexType>::build().on(exec);
auto par_ilu = gko::share(par_ilu_fact->generate(A));
auto bj_factory = gko::share(
    bj::build()
        .with_max_block_size(16u)
        .with_storage_optimization(gko::precision_reduction::autodetect())
        .on(exec));
auto trisolve_factory =
    ir::build()
        .with_solver(bj_factory)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(sweeps))
        .on(exec);
auto ilu_pre_factory = gko::preconditioner::Ilu<ir, ir>::build()
    .with_l_solver(gko::clone(trisolve_factory))
    .with_u_solver(gko::clone(trisolve_factory))
    .on(exec);
auto ilu_preconditioner = gko::share(ilu_pre_factory->generate(par_ilu));
const RealValueType reduction_factor{1e-12};
auto iter_stop = gko::share(
    gko::stop::Iteration::build().with_max_iters(1000u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_reduction_factor(reduction_factor)
    .on(exec));

auto ilu_gmres_factory =
    gmres::build()
        .with_criteria(iter_stop, tol_stop)
        .with_generated_preconditioner(ilu_preconditioner)
        .on(exec);
auto ilu_gmres = ilu_gmres_factory->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
ilu_gmres->add_logger(logger);
ilu_gmres->apply(b, x);
std::chrono::nanoseconds time(0);
for (int i = 0; i < 100; i++) {
    x->copy_from(clone_x);
    auto tic = std::chrono::high_resolution_clock::now();
    ilu_gmres->apply(b, x);
    auto toc = std::chrono::high_resolution_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
}
std::cout << "Using " << sweeps << " block-Jacobi sweeps.\n";
std::cout << "Solution (x):\n";
write(std::cout, x);
auto res = gko::as<vec>(logger->get_residual_norm());
std::cout << "GMRES iteration count: " << logger->get_num_iterations()

```

```
        « "\n";
std::cout << "GMRES execution time [ms]: "
        << static_cast<double>(time.count()) / 1000000000.0 << "\n";
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```


Chapter 24

The iterative-refinement program

The iterative refinement solver example..

This example depends on simple-solver.

This example shows how to use the iterative refinement solver.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. An inaccurate CG solver is used as the inner solver to an iterative refinement (IR) method which solves a linear system. The example features the iteration count and runtime of the IR solver.

The commented program

```
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); } }
};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS and initial guess as 1

```
gko::size_type size = A->get_size()[0];
auto host_x = gko::matrix::Dense<ValueType>::create(exec->get_master(),
                                                    gko::dim<2>(size, 1));

for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
```

```
auto b = gko::clone(exec, host_x);
```

Calculate initial residual by overwriting b

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
```

copy b again

```
b->copy_from(host_x);
```

Create solver factory

```
gko::size_type max_iters = 10000u;
RealValueType outer_reduction_factor{1e-12};
RealValueType inner_reduction_factor{1e-2};
auto solver_gen =
    ir::build()
        .with_solver(cg::build().with_criteria(
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(inner_reduction_factor)))
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(max_iters),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(outer_reduction_factor))
        .on(exec);
```

Create solver

```
auto solver = solver_gen->generate(A);
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
```

Solve system

```
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
```

Get residual

```
auto res = gko::as<real_vec>(logger->get_residual_norm());
std::cout << "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
```

Print solver statistics

```
std::cout << "IR iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "IR execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
}
```

Results

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
194.679
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.23821e-11
IR iteration count:          24
IR execution time [ms]:  0.794962
```


Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using ir = gko::solver::Ir<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = gko::matrix::Dense<ValueType>::create(exec->get_master(),
                                                         gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 1.;
    }
    auto x = gko::clone(exec, host_x);
    auto b = gko::clone(exec, host_x);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_x);
    gko::size_type max_iters = 10000u;
    RealValueType outer_reduction_factor{1e-12};
    RealValueType inner_reduction_factor{1e-2};
    auto solver_gen =
        ir::build()
            .with_solver(cg::build().with_criteria(
                gko::stop::ResidualNorm<ValueType>::build()
                    .with_reduction_factor(inner_reduction_factor)))
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(max_iters),
                gko::stop::ResidualNorm<ValueType>::build()
                    .with_reduction_factor(outer_reduction_factor))
            .on(exec);
    auto solver = solver_gen->generate(A);
    std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
        gko::log::Convergence<ValueType>::create();
    solver->add_logger(logger);
    exec->synchronize();
    std::chrono::nanoseconds time(0);
    auto tic = std::chrono::steady_clock::now();
    solver->apply(b, x);
    auto toc = std::chrono::steady_clock::now();
    time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
    auto res = gko::as<real_vec>(logger->get_residual_norm());
```

```
std::cout << "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):\n";
write(std::cout, res);
std::cout << "IR iteration count:          " << logger->get_num_iterations()
          << std::endl;
std::cout << "IR execution time [ms]:  "
          << static_cast<double>(time.count()) / 1000000.0 << std::endl;
}
```

Chapter 25

The kokkos-assembly program

The Kokkos assembly example..

This example depends on simple-solver, poisson-solver, three-pt-stencil-solver, .

Introduction

This example solves a 1D Poisson equation:

$$u : [0, 1] \rightarrow \mathbb{R} \text{ } u'' = f u(0) = u_0 u(1) = u_1$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter).

The resulting CSR matrix is assembled using Kokkos kernels. This example show how to use Ginkgo data with Kokkos kernels.

Notes

If this example is built as part of Ginkgo, it is advised to configure Ginkgo with `-DGINKGO_WITH_CCACHE=OFF` to prevent incompatibilities with Kokkos' compiler wrapper for `nvcc`.

The commented program

```
*
* @tparam ValueType_c The value type of the matrix elements, might have
*                     other cv qualifiers than ValueType
* @tparam IndexType_c The index type of the matrix elements, might have
*                     other cv qualifiers than IndexType
* /
template <typename ValueType_c, typename IndexType_c>
struct type {
    using index_array = typename index_mapper::template type<IndexType_c>;
    using value_array = typename value_mapper::template type<ValueType_c>;
    / **
    * Constructor based on size and raw pointers
    *
    * @param size The number of stored elements
    * @param row_idx Pointer to the row indices
    * @param col_idx Pointer to the column indices
```

```

    * @param values Pointer to the values
    *
    * @return An object which has each gko::array of the
    *         device_matrix_data mapped to a Kokkos view
    * /
    static type map(size_type size, IndexType_c* row_idxs,
                    IndexType_c* col_idxs, ValueType_c* values)
    {
        return {index_mapper::map(row_idxs, size),
                index_mapper::map(col_idxs, size),
                value_mapper::map(values, size)};
    }
    index_array row_idxs;
    index_array col_idxs;
    value_array values;
};
static type<ValueType, IndexType> map(
    device_matrix_data<ValueType, IndexType>& md)
{
    assert_compatibility<MemorySpace>(md);
    return type<ValueType, IndexType>::map(
        md.get_num_stored_elements(), md.get_row_idxs(), md.get_col_idxs(),
        md.get_values());
}
static type<const ValueType, const IndexType> map(
    const device_matrix_data<ValueType, IndexType>& md)
{
    assert_compatibility<MemorySpace>(md);
    return type<const ValueType, const IndexType>::map(
        md.get_num_stored_elements(), md.get_const_row_idxs(),
        md.get_const_col_idxs(), md.get_const_values());
}
};
} // namespace gko::ext::kokkos::detail

```

Creates a stencil matrix in CSR format for the given number of discretization points.

```

template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    auto exec = matrix->get_executor();
    const auto discretization_points = matrix->get_size()[0];

```

Over-allocate storage for the matrix elements. Each row has 3 entries, except for the first and last one. To handle each row uniformly, we allocate space for 3x the number of rows.

```

gko::device_matrix_data<ValueType, IndexType> md(exec, matrix->get_size(),
                                                discretization_points * 3);

```

Create Kokkos views on Ginkgo data.

```

auto k_md = gko::ext::kokkos::map_data(md);

```

Create the matrix entries. This also creates zero entries for the first and second row to handle all rows uniformly.

```

Kokkos::parallel_for(
    "generate_stencil_matrix", md.get_num_stored_elements(),
    KOKKOS_LAMBDA(int i) {
        const ValueType coefs[] = {-1, 2, -1};
        auto ofs = static_cast<IndexType>((i % 3) - 1);
        auto row = static_cast<IndexType>(i / 3);
        auto col = row + ofs;
    });

```

To prevent branching, a mask is used to set the entry to zero, if the column is out-of-bounds

```

    auto mask =
        static_cast<IndexType>(0 <= col && col < discretization_points);
    k_md.row_idxs[i] = mask * row;
    k_md.col_idxs[i] = mask * col;
    k_md.values[i] = mask * coefs[ofs + 1];
});

```

Add up duplicate (zero) entries.

```

md.sum_duplicates();

```

Build Csr matrix.

```

matrix->read(std::move(md));
}

```

Generates the RHS vector given f and the boundary conditions.

```

template <typename Closure, typename ValueType>
void generate_rhs(Closure&& f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{

```

```

const auto discretization_points = rhs->get_size()[0];
auto k_rhs = gko::ext::kokkos::map_data(rhs);
Kokkos::parallel_for(
    "generate_rhs", discretization_points, KOKKOS_LAMBDA(int i) {
        const ValueType h = 1.0 / (discretization_points + 1);
        const ValueType xi = ValueType(i + 1) * h;
        k_rhs(i, 0) = -f(xi) * h * h;
        if (i == 0) {
            k_rhs(i, 0) += u0;
        }
        if (i == discretization_points - 1) {
            k_rhs(i, 0) += u1;
        }
    });
}

```

Computes the 1-norm of the error given the computed `u` and the correct solution function `correct_u`.

```

template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
    const gko::matrix::Dense<ValueType>* u,
    Closure&& correct_u)
{
    auto k_u = gko::ext::kokkos::map_data(u);
    auto error = 0.0;
    Kokkos::parallel_reduce(
        "calculate_error", discretization_points,
        KOKKOS_LAMBDA(int i, double& lsum) {
            const auto h = 1.0 / (discretization_points + 1);
            const auto xi = (i + 1) * h;
            lsum += Kokkos::abs((k_u(i, 0) - correct_u(xi)) /
                Kokkos::abs(correct_u(xi)));
        },
        error);
    return error;
}
int main(int argc, char* argv[])
{

```

Some shortcuts

```

using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType>;

```

Print help message. For details on the kokkos-options see <https://kokkos.github.io/kokkos-core-wiki/ProgrammingGuide/Initialization.html#initialization-by-command-line-arguments>

```

if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0]
        << " [discretization_points] [kokkos-options]" << std::endl;
    Kokkos::ScopeGuard kokkos(argc, argv); // print Kokkos help
    std::exit(1);
}
Kokkos::ScopeGuard kokkos(argc, argv);
const auto discretization_points =
    static_cast<gko::size_type>(argc >= 2 ? std::atoi(argv[1]) : 100u);

```

chooses the executor that corresponds to the Kokkos DefaultExecutionSpace

```

auto exec = gko::ext::kokkos::create_default_executor();

```

problem:

```

auto correct_u = [] KOKKOS_FUNCTION(ValueType x) { return x * x * x; };
auto f = [] KOKKOS_FUNCTION(ValueType x) { return ValueType{6} * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);

```

initialize vectors

```

auto rhs = vec::create(exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(exec, gko::dim<2>(discretization_points, 1));
u->fill(0.0);

```

initialize the stencil matrix

```

auto A = share(mtx::create(
    exec, gko::dim<2>(discretization_points, discretization_points)));
generate_stencil_matrix(A.get());
const RealValueType reduction_factor{1e-7};

```

Generate solver and solve the system

```
cg::build()
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(discretization_points),
        gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
            reduction_factor)
    ).with_preconditioner(bj::build())
    .on(exec)
    ->generate(A)
    ->apply(rhs, u);
std::cout << "\nSolve complete."
            << "\nThe average relative error is "
            << calculate_error(discretization_points, u.get(), correct_u) /
                discretization_points
            << std::endl;
}
```

Results

Example output:

```
> ./kokkos-assembly
Solve complete.
The average relative error is 1.05488e-11
```

The actual error depends on the used hardware.

The plain program

```
#include <iostream>
#include <string>
#include <Kokkos_Core.hpp>
#include <ginkgo/ginkgo.hpp>
#include <ginkgo/extensions/kokkos.hpp>
namespace gko::ext::kokkos::detail {
template <typename ValueType, typename IndexType, typename MemorySpace>
struct mapper<device_matrix_data<ValueType, IndexType>, MemorySpace> {
    using index_mapper = mapper<array<IndexType>, MemorySpace>;
    using value_mapper = mapper<array<ValueType>, MemorySpace>;
    template <typename ValueType_c, typename IndexType_c>
    struct type {
        using index_array = typename index_mapper::template type<IndexType_c>;
        using value_array = typename value_mapper::template type<ValueType_c>;
        static type map(size_type size, IndexType_c* row_idxs,
            IndexType_c* col_idxs, ValueType_c* values)
        {
            return {index_mapper::map(row_idxs, size),
                index_mapper::map(col_idxs, size),
                value_mapper::map(values, size)};
        }
        index_array row_idxs;
        index_array col_idxs;
        value_array values;
    };
    static type<ValueType, IndexType> map(
        device_matrix_data<ValueType, IndexType>& md)
    {
        assert_compatibility<MemorySpace>(md);
        return type<ValueType, IndexType>::map(
            md.get_num_stored_elements(), md.get_row_idxs(), md.get_col_idxs(),
            md.get_values());
    }
    static type<const ValueType, const IndexType> map(
        const device_matrix_data<ValueType, IndexType>& md)
    {
        assert_compatibility<MemorySpace>(md);
        return type<const ValueType, const IndexType>::map(
            md.get_num_stored_elements(), md.get_const_row_idxs(),
            md.get_const_col_idxs(), md.get_const_values());
    }
};
} // namespace gko::ext::kokkos::detail
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    auto exec = matrix->get_executor();
    const auto discretization_points = matrix->get_size()[0];
```

```

gko::device_matrix_data<ValueType, IndexType> md(exec, matrix->get_size(),
                                                discretization_points * 3);

auto k_md = gko::ext::kokkos::map_data(md);
Kokkos::parallel_for(
    "generate_stencil_matrix", md.get_num_stored_elements(),
    KOKKOS_LAMBDA(int i) {
        const ValueType coefs[] = {-1, 2, -1};
        auto ofs = static_cast<IndexType>((i % 3) - 1);
        auto row = static_cast<IndexType>(i / 3);
        auto col = row + ofs;
        auto mask =
            static_cast<IndexType>(0 <= col && col < discretization_points);
        k_md.row_idxs[i] = mask * row;
        k_md.col_idxs[i] = mask * col;
        k_md.values[i] = mask * coefs[ofs + 1];
    });
md.sum_duplicates();
matrix->read(std::move(md));
}

template <typename Closure, typename ValueType>
void generate_rhs(Closure&& f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{
    const auto discretization_points = rhs->get_size()[0];
    auto k_rhs = gko::ext::kokkos::map_data(rhs);
    Kokkos::parallel_for(
        "generate_rhs", discretization_points, KOKKOS_LAMBDA(int i) {
            const ValueType h = 1.0 / (discretization_points + 1);
            const ValueType xi = ValueType(i + 1) * h;
            k_rhs(i, 0) = -f(xi) * h * h;
            if (i == 0) {
                k_rhs(i, 0) += u0;
            }
            if (i == discretization_points - 1) {
                k_rhs(i, 0) += u1;
            }
        });
}

template <typename Closure, typename ValueType>
double calculate_error(int discretization_points,
                      const gko::matrix::Dense<ValueType>* u,
                      Closure&& correct_u)
{
    auto k_u = gko::ext::kokkos::map_data(u);
    auto error = 0.0;
    Kokkos::parallel_reduce(
        "calculate_error", discretization_points,
        KOKKOS_LAMBDA(int i, double& lsum) {
            const auto h = 1.0 / (discretization_points + 1);
            const auto xi = (i + 1) * h;
            lsum += Kokkos::abs((k_u(i, 0) - correct_u(xi)) /
                               Kokkos::abs(correct_u(xi)));
        },
        error);
    return error;
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType>;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0]
            << " [discretization_points] [kokkos-options]" << std::endl;
        Kokkos::ScopeGuard kokkos(argc, argv); // print Kokkos help
        std::exit(1);
    }
    Kokkos::ScopeGuard kokkos(argc, argv);
    const auto discretization_points =
        static_cast<gko::size_type>(argc >= 2 ? std::atoi(argv[1]) : 100u);
    auto exec = gko::ext::kokkos::create_default_executor();
    auto correct_u = [] KOKKOS_FUNCTION(ValueType x) { return x * x * x; };
    auto f = [] KOKKOS_FUNCTION(ValueType x) { return ValueType{6} * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    auto rhs = vec::create(exec, gko::dim<2>(discretization_points, 1));
    generate_rhs(f, u0, u1, rhs.get());
    auto u = vec::create(exec, gko::dim<2>(discretization_points, 1));
    u->fill(0.0);
    auto A = share(mtx::create(
        exec, gko::dim<2>(discretization_points, discretization_points)));
    generate_stencil_matrix(A.get());
    const RealValueType reduction_factor{1e-7};

```

```
cg::build()
  .with_criteria(
    gko::stop::Iteration::build().with_max_iters(discretization_points),
    gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
      reduction_factor))
  .with_preconditioner(bj::build())
  .on(exec)
  ->generate(A)
  ->apply(rhs, u);
std::cout << "\nSolve complete."
  << "\nThe average relative error is "
  << calculate_error(discretization_points, u.get(), correct_u) /
    discretization_points
  << std::endl;
}
```


Chapter 26

The minimal-cuda-solver program

The minimal CUDA solver example..

This example depends on simple-solver.

Introduction

This is a minimal example that solves a system with Ginkgo. The matrix, right hand side and initial guess are read from standard input, and the result is written to standard output. The system matrix is stored in CSR format, and the system solved using the CG method, preconditioned with the block-Jacobi preconditioner. All computations are done on the GPU.

The easiest way to use the example data from the `data/` folder is to concatenate the matrix, the right hand side and the initial solution (in that exact order), and pipe the result to the `minimal_solver_cuda` executable:

```
cat data/A.mtx data/b.mtx data/x0.mtx | ./minimal-cuda-solver
```

About the example

The commented program

Results

The following is the expected result when using the data contained in the folder `data` as input:

```
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
```

Comments about programming and debugging

The plain program

```
#include <iostream>
#include <ginkgo/ginkgo.hpp>
int main()
{
    auto gpu = gko::CudaExecutor::create(0, gko::OmpExecutor::create());
    auto A = gko::read<gko::matrix::Csr<>>(std::cin, gpu);
    auto b = gko::read<gko::matrix::Dense<>>(std::cin, gpu);
    auto x = gko::read<gko::matrix::Dense<>>(std::cin, gpu);
    auto solver =
        gko::solver::Cg<>::build()
            .with_preconditioner(gko::preconditioner::Jacobi<>::build())
            .with_criteria(
                gko::stop::Iteration::build().with_max_iters(20u),
                gko::stop::ResidualNorm<>::build().with_reduction_factor(1e-15))
            .on(gpu);
    solver->generate(give(A))->apply(b, x);
    write(std::cout, x);
}
```

Chapter 27

The mixed-multigrid-preconditioned-solver program

The preconditioned solver example..

This example depends on multigrid-preconditioned-solver, mixed-multigrid-solver.

This example shows how to use the mixed-precision multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a mixed-precision multigrid preconditioner. The example features the generating time and runtime of the CG solver.

The commented program

Print version information

```
std::cout << gko::version_info::get() << std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(
             0, gko::ReferenceExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout << "Using mixed precision? " << use_mixed << std::endl;
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS as 1 and initial guess as 0

```
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
}
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
```

Calculate initial residual by overwriting b

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
```

copy b again

```
b->copy_from(host_b);
```

Prepare the stopping criteria

```
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::absolute)
    .with_reduction_factor(tolerance)
    .on(exec));
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
```

Create smoother factory (ir with bj)

```
auto inner_solver_gen =
    gko::share(bj::build().with_max_block_size(1u).on(exec));
auto inner_solver_gen_f =
    gko::share(bj_f::build().with_max_block_size(1u).on(exec));
auto smoother_gen = gko::share(
    ir::build()
    .with_solver(inner_solver_gen)
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));
auto smoother_gen_f = gko::share(
    ir_f::build()
    .with_solver(inner_solver_gen_f)
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));
```

Create MultigridLevel factory

```
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen_f =
    gko::share(pgm_f::build().with_deterministic(true).on(exec));
```

Create CoarsestSolver factory

```
auto coarsest_gen = gko::share(
    ir::build()
    .with_solver(inner_solver_gen)
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));
auto coarsest_gen_f = gko::share(
    ir_f::build()
    .with_solver(inner_solver_gen_f)
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));
```

Create multigrid factory

```
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
    multigrid_gen =
```

```

mg::build()
    .with_max_levels(10u)
    .with_min_coarse_rows(2u)
    .with_pre_smoother(smoother_gen, smoother_gen_f)
    .with_post_uses_pre(true)
    .with_mg_level(mg_level_gen, mg_level_gen_f)
    .with_level_selector([](const gko::size_type level,
                           const gko::LinOp*) -> gko::size_type {

        return level >= 1 ? 1 : 0;

    })
    .with_coarsest_solver(coarsest_gen_f)
    .with_default_initial_guess(
        gko::solver::initial_guess_mode::zero)
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec);
} else {
    multigrid_gen =
        mg::build()
            .with_max_levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen)
            .with_coarsest_solver(coarsest_gen)
            .with_default_initial_guess(
                gko::solver::initial_guess_mode::zero)
            .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
            .on(exec);
}

```

The first (index 0) level will use the first `mg_level_gen`, `smoother_gen` which are the factories with `ValueType`. The rest of levels (≥ 1) will use the second (index 1) `mg_level_gen2` and `smoother_gen2` which use the `MixedType`. The rest of levels will use different type than the normal multigrid.

Create solver factory

```

auto solver_gen = cg::build()
    .with_criteria(iter_stop, tol_stop)
    .with_preconditioner(multigrid_gen)
    .on(exec);

```

Create solver

```

std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);

```

Solve system

```

exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);

```

Calculate residual

```

auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);

```

Print solver statistics

```

std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```

Results

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.3589
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.69858e-09
CG iteration count:          39
CG generation time [ms]:    2.04293
CG execution time [ms]:     22.3874
CG execution time per iteration[ms]: 0.574036
```

Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using MixedType = float;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using fcg = gko::solver::Fcg<ValueType>;
    using cg = gko::solver::Cg<ValueType>;
    using ir = gko::solver::Ir<ValueType>;
    using mg = gko::solver::Multigrid;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    using cg_f = gko::solver::Cg<MixedType>;
    using ir_f = gko::solver::Ir<MixedType>;
    using bj_f = gko::preconditioner::Jacobi<MixedType, IndexType>;
    using pgm_f = gko::multigrid::Pgm<MixedType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
    const bool use_mixed = mixed_int != 0; // nonzero uses mixed
    std::cout << "Using mixed precision? " << use_mixed << std::endl;
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    }
    auto x = vec::create(exec);
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
```

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_b);
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::absolute)
    .with_reduction_factor(tolerance)
    .on(exec));

std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
auto inner_solver_gen =
    gko::share(bj::build().with_max_block_size(1u).on(exec));
auto inner_solver_gen_f =
    gko::share(bj_f::build().with_max_block_size(1u).on(exec));
auto smoother_gen = gko::share(
    ir::build()
    .with_solver(inner_solver_gen)
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));
auto smoother_gen_f = gko::share(
    ir_f::build()
    .with_solver(inner_solver_gen_f)
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen_f =
    gko::share(pgm_f::build().with_deterministic(true).on(exec));
auto coarsest_gen = gko::share(
    ir::build()
    .with_solver(inner_solver_gen)
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));
auto coarsest_gen_f = gko::share(
    ir_f::build()
    .with_solver(inner_solver_gen_f)
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));

std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
    multigrid_gen =
        mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(2u)
        .with_pre_smoother(smoother_gen, smoother_gen_f)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen, mg_level_gen_f)
        .with_level_selector([](const gko::size_type level,
            const gko::LinOp* -> gko::size_type {
                return level >= 1 ? 1 : 0;
            })
        .with_coarsest_solver(coarsest_gen_f)
        .with_default_initial_guess(
            gko::solver::initial_guess_mode::zero)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec);
} else {
    multigrid_gen =
        mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(2u)
        .with_pre_smoother(smoother_gen)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen)
        .with_coarsest_solver(coarsest_gen)
        .with_default_initial_guess(
            gko::solver::initial_guess_mode::zero)
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec);
}
auto solver_gen = cg::build()
    .with_criteria(iter_stop, tol_stop)
    .with_preconditioner(multigrid_gen)
    .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();

```

```

auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);
std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```


Chapter 28

The mixed-multigrid-solver program

The mixed multigrid solver example..

This example depends on simple-solver.

This example shows how to use the mixed-precision multigrid solver.

In this example, we first read in a matrix from a file, then generate a right-hand side and an initial guess. The multigrid solver can mix different precision of MultigridLevel. The example features the generating time and runtime of the multigrid solver.

The commented program

```
std::cout << gko::version_info::get() << std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(
             0, gko::ReferenceExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
const bool use_mixed = mixed_int != 0; // nonzero uses mixed
std::cout << "Using mixed precision? " << use_mixed << std::endl;
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS as 1 and initial guess as 0

```
gko::size_type size = A->get_size()[0];
```

```

auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
}
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);

```

Calculate initial residual by overwriting b

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);

```

copy b again

```

b->copy_from(host_b);

```

Prepare the stopping criteria

```

const gko::remove_complex<ValueType> tolerance = 1e-12;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::absolute)
    .with_reduction_factor(tolerance)
    .on(exec));

```

Create smoother factory (ir with bj)

```

auto smoother_gen = gko::share(
    ir::build()
    .with_solver(bj::build().with_max_block_size(1u))
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));
auto smoother_gen2 = gko::share(
    ir2::build()
    .with_solver(bj2::build().with_max_block_size(1u))
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec));

```

Create RestrictProlong factory

```

auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen2 =
    gko::share(pgm2::build().with_deterministic(true).on(exec));

```

Create CoarsesSolver factory

```

auto coarsest_solver_gen = gko::share(
    ir::build()
    .with_solver(bj::build().with_max_block_size(1u))
    .with_relaxation_factor(static_cast<ValueType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));
auto coarsest_solver_gen2 = gko::share(
    ir2::build()
    .with_solver(bj2::build().with_max_block_size(1u))
    .with_relaxation_factor(static_cast<MixedType>(0.9))
    .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
    .on(exec));

```

Create multigrid factory

```

std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
    multigrid_gen =
        mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(2u)
        .with_pre_smoother(smoother_gen, smoother_gen2)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen, mg_level_gen2)
        .with_level_selector([](const gko::size_type level,
            const gko::LinOp*) -> gko::size_type {

```

The first (index 0) level will use the first `mg_level_gen`, `smoother_gen` which are the factories with `ValueType`. The rest of levels (≥ 1) will use the second (index 1) `mg_level_gen2` and `smoother_gen2` which use the `MixedType`. The rest of levels will use different type than the normal multigrid.

```

        return level >= 1 ? 1 : 0;
    })
    .with_coarsest_solver(coarsest_solver_gen2)
    .with_criteria(iter_stop, tol_stop)
    .on(exec);
} else {
    multigrid_gen = mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(2u)
        .with_pre_smoother(smoother_gen)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen)
        .with_coarsest_solver(coarsest_solver_gen)
        .with_criteria(iter_stop, tol_stop)
        .on(exec);
}
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();

auto solver = solver_gen->generate(A);
auto solver = multigrid_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);

```

Add logger

```

std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);

```

Solve system

```

exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);

```

Calculate residual explicitly, because the residual is not available inside of the multigrid solver

```

auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);

```

Print solver statistics

```

std::cout << "Multigrid iteration count:          "
    << logger->get_num_iterations() << std::endl;
std::cout << "Multigrid generation time [ms]: "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "Multigrid execution time [ms]: "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "Multigrid execution time per iteration[ms]: "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```

Results

This is the expected output:

```

Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.3589
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
6.31088e-14
Multigrid iteration count:          9
Multigrid generation time [ms]:  3.35361
Multigrid execution time [ms]:  10.048
Multigrid execution time per iteration[ms]:  1.11644

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using MixedType = float;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using fcg = gko::solver::Fcg<ValueType>;
    using cg = gko::solver::Cg<MixedType>;
    using ir = gko::solver::Ir<ValueType>;
    using ir2 = gko::solver::Ir<MixedType>;
    using mg = gko::solver::Multigrid;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using bj2 = gko::preconditioner::Jacobi<MixedType, IndexType>;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    using pgm2 = gko::multigrid::Pgm<MixedType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const int mixed_int = argc >= 3 ? std::atoi(argv[2]) : 1;
    const bool use_mixed = mixed_int != 0; // nonzero uses mixed
    std::cout << "Using mixed precision? " << use_mixed << std::endl;
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    }
    auto x = vec::create(exec);
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_b);
    const gko::remove_complex<ValueType> tolerance = 1e-12;
    auto iter_stop =
        gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
    auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
        .with_baseline(gko::stop::mode::absolute)
        .with_reduction_factor(tolerance)
        .on(exec));
    auto smoother_gen = gko::share(
        ir::build()
        .with_solver(bj::build().with_max_block_size(1u))
        .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec));
    auto smoother_gen2 = gko::share(
        ir2::build()
        .with_solver(bj2::build().with_max_block_size(1u))

```

```

        .with_relaxation_factor(static_cast<MixedType>(0.9))
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec));
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto mg_level_gen2 =
    gko::share(pgm2::build().with_deterministic(true).on(exec));
auto coarsest_solver_gen = gko::share(
    ir::build()
        .with_solver(bj::build().with_max_block_size(1u))
        .with_relaxation_factor(static_cast<ValueType>(0.9))
        .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
        .on(exec));
auto coarsest_solver_gen2 = gko::share(
    ir2::build()
        .with_solver(bj2::build().with_max_block_size(1u))
        .with_relaxation_factor(static_cast<MixedType>(0.9))
        .with_criteria(gko::stop::Iteration::build().with_max_iters(4u))
        .on(exec));
std::shared_ptr<gko::LinOpFactory> multigrid_gen;
if (use_mixed) {
    multigrid_gen =
        mg::build()
            .with_max_levels(10u)
            .with_min_coarse_rows(2u)
            .with_pre_smoother(smoother_gen, smoother_gen2)
            .with_post_uses_pre(true)
            .with_mg_level(mg_level_gen, mg_level_gen2)
            .with_level_selector([](const gko::size_type level,
                                   const gko::LinOp* ) -> gko::size_type {
                return level >= 1 ? 1 : 0;
            })
            .with_coarsest_solver(coarsest_solver_gen2)
            .with_criteria(iter_stop, tol_stop)
            .on(exec);
} else {
    multigrid_gen = mg::build()
        .with_max_levels(10u)
        .with_min_coarse_rows(2u)
        .with_pre_smoother(smoother_gen)
        .with_post_uses_pre(true)
        .with_mg_level(mg_level_gen)
        .with_coarsest_solver(coarsest_solver_gen)
        .with_criteria(iter_stop, tol_stop)
        .on(exec);
}
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = multigrid_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);
std::cout << "Multigrid iteration count:          "
    << logger->get_num_iterations() << std::endl;
std::cout << "Multigrid generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "Multigrid execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "Multigrid execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```


Chapter 29

The mixed-precision-ir program

The Mixed Precision Iterative Refinement (MPIR) solver example..

This example depends on iterative-refinement.

This example manually implements a Mixed Precision Iterative Refinement (MPIR) solver.

In this example, we first read in a matrix from file, then generate a right-hand side and an initial guess. An inaccurate CG solver in single precision is used as the inner solver to an iterative refinement (IR) in double precision method which solves a linear system.

The commented program

Print version information

```
std::cout << gko::version_info::get() << std::endl;
```

Figure out where to run the code

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }
    };
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS and initial guess as 1

```
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
```

Calculate initial residual by overwriting b

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres_vec = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres_vec);
```

Build lower-precision system matrix and residual

```
auto solver_A = solver_mtx::create(exec);
auto inner_residual = solver_vec::create(exec);
auto outer_residual = vec::create(exec);
A->convert_to(solver_A);
b->convert_to(outer_residual);
```

restore b

```
b->copy_from(host_x);
```

Create inner solver

```
auto inner_solver =
    cg::build()
        .with_criteria(
            gko::stop::ResidualNorm<SolverType>::build()
                .with_reduction_factor(inner_reduction_factor),
            gko::stop::Iteration::build().with_max_iters(max_inner_iters))
        .on(exec)
        ->generate(give(solver_A));
```

Solve system

```
exec->synchronize();
std::chrono::nanoseconds time(0);
auto res_vec = gko::initialize<real_vec>({0.0}, exec);
auto initres = exec->copy_val_to_host(initres_vec->get_const_values());
auto inner_solution = solver_vec::create(exec);
auto outer_delta = vec::create(exec);
auto tic = std::chrono::steady_clock::now();
int iter = -1;
while (true) {
    ++iter;
```

convert residual to inner precision

```
outer_residual->convert_to(inner_residual);
outer_residual->compute_norm2(res_vec);
auto res = exec->copy_val_to_host(res_vec->get_const_values());
```

break if we exceed the number of iterations or have converged

```
if (iter > max_outer_iters || res / initres < outer_reduction_factor) {
    break;
}
```

Use the inner solver to solve $A * \text{inner_solution} = \text{inner_residual}$ with residual as initial guess.

```
inner_solution->copy_from(inner_residual);
inner_solver->apply(inner_residual, inner_solution);
```

convert inner solution to outer precision

```
inner_solution->convert_to(outer_delta);
```

 $x = x + \text{inner_solution}$

```
x->add_scaled(one, outer_delta);
```

residual = $b - A * x$

```
outer_residual->copy_from(b);
A->apply(neg_one, x, one, outer_residual);
}
auto toc = std::chrono::steady_clock::now();
```



```
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
```

Calculate residual

```
A->apply(one, x, neg_one, b);
b->compute_norm2(res_vec);
std::cout << "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres_vec);
std::cout << "Final residual norm sqrt(r^T r):\n";
write(std::cout, res_vec);
```

Print solver statistics

```
std::cout << "MPIR iteration count:          " << iter << std::endl;
std::cout << "MPIR execution time [ms]:  "
            << static_cast<double>(time.count()) / 1000000.0 << std::endl;
}
```

Results

This is the expected output:

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
194.679
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.22728e-10
MPIR iteration count:          25
MPIR execution time [ms]:  0.846559
```

Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using SolverType = float;
    using RealSolverType = gko::remove_complex<SolverType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using solver_vec = gko::matrix::Dense<SolverType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using solver_mtx = gko::matrix::Csr<SolverType, IndexType>;
    using cg = gko::solver::Cg<SolverType>;
    gko::size_type max_outer_iters = 100u;
    gko::size_type max_inner_iters = 100u;
    RealValueType outer_reduction_factor{1e-12};
    RealSolverType inner_reduction_factor{1e-2};
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage:  " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
        exec_map{
            {"omp", [] { return gko::OmpExecutor::create(); }},
            {"cuda",
             [] {
                 return gko::CudaExecutor::create(0,
                                                    gko::OmpExecutor::create());
             }},
            {"hip",
             [] {
                 return gko::HipExecutor::create(0, gko::OmpExecutor::create());
             }}
        }
```

```

    }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(
             0, gko::ReferenceExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }});
const auto exec = exec_map.at(executor_string()); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 1.;
}
auto x = gko::clone(exec, host_x);
auto b = gko::clone(exec, host_x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres_vec = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres_vec);
auto solver_A = solver_mtx::create(exec);
auto inner_residual = solver_vec::create(exec);
auto outer_residual = vec::create(exec);
A->convert_to(solver_A);
b->convert_to(outer_residual);
b->copy_from(host_x);
auto inner_solver =
    cg::build()
        .with_criteria(
            gko::stop::ResidualNorm<SolverType>::build()
                .with_reduction_factor(inner_reduction_factor),
            gko::stop::Iteration::build().with_max_iters(max_inner_iters))
        .on(exec)
        ->generate(give(solver_A));
exec->synchronize();
std::chrono::nanoseconds time(0);
auto res_vec = gko::initialize<real_vec>({0.0}, exec);
auto initres = exec->copy_val_to_host(initres_vec->get_const_values());
auto inner_solution = solver_vec::create(exec);
auto outer_delta = vec::create(exec);
auto tic = std::chrono::steady_clock::now();
int iter = -1;
while (true) {
    ++iter;
    outer_residual->convert_to(inner_residual);
    outer_residual->compute_norm2(res_vec);
    auto res = exec->copy_val_to_host(res_vec->get_const_values());
    if (iter > max_outer_iters || res / initres < outer_reduction_factor) {
        break;
    }
    inner_solution->copy_from(inner_residual);
    inner_solver->apply(inner_residual, inner_solution);
    inner_solution->convert_to(outer_delta);
    x->add_scaled(one, outer_delta);
    outer_residual->copy_from(b);
    A->apply(neg_one, x, one, outer_residual);
}
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
A->apply(one, x, neg_one, b);
b->compute_norm2(res_vec);
std::cout << "Initial residual norm sqrt(r^T r):\n";
write(std::cout, initres_vec);
std::cout << "Final residual norm sqrt(r^T r):\n";
write(std::cout, res_vec);
std::cout << "MPIR iteration count: " << iter << std::endl;
std::cout << "MPIR execution time [ms]: "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
}

```

Chapter 30

The mixed-spmv program

The mixed spmv example..

Introduction

This mixed spmv example should give the usage of Ginkgo mixed precision. This example is meant for you to understand how Ginkgo works with different precision of data. We encourage you to play with the code, change the parameters and see what is best suited for your purposes.

About the example

Each example has the following sections:

1. **Introduction:** This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

```
* @param value_dist  distribution of array values
* @param engine      a random engine
*
* @return ValueType
* /
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<!gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
{
    return value_dist(gen);
}
/ **
* Specialization for complex types.
```

```

*
* @copydoc get_rand_value
* /
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
{
    return ValueType(value_dist(gen), value_dist(gen));
}
/ **
* timing the apply operation A->apply(b, x). It will runs 2 warmup and get
* average time among 10 times.
*
* @return seconds
* /
double timing(std::shared_ptr<const gko::Executor> exec,
              std::shared_ptr<const gko::LinOp> A,
              std::shared_ptr<const gko::LinOp> b,
              std::shared_ptr<gko::LinOp> x)
{
    int warmup = 2;
    int rep = 10;
    for (int i = 0; i < warmup; i++) {
        A->apply(b, x);
    }
    double total_sec = 0;
    for (int i = 0; i < rep; i++) {

```

always clone the x in each apply

```
auto xx = x->clone();
```

synchronize to make sure data is already on device

```
exec->synchronize();
auto start = std::chrono::steady_clock::now();
A->apply(b, xx);
```

synchronize to make sure the operation is done

```
exec->synchronize();
auto stop = std::chrono::steady_clock::now();
```

get the duration in seconds

```
std::chrono::duration<double> duration_time = stop - start;
total_sec += duration_time.count();
if (i + 1 == rep) {
```

copy the result back to x

```

        x->copy_from(xx);
    }
    return total_sec / rep;
}
// namespace
int main(int argc, char* argv[])
{

```

Use some shortcuts. In Ginkgo, vectors are seen as a `gko::matrix::Dense` with one column/one row. The advantage of this concept is that using multiple vectors is now a natural extension of adding columns/rows are necessary.

```

using HighPrecision = double;
using RealValueType = gko::remove_complex<HighPrecision>;
using LowPrecision = float;
using IndexType = int;
using hp_vec = gko::matrix::Dense<HighPrecision>;
using lp_vec = gko::matrix::Dense<LowPrecision>;
using real_vec = gko::matrix::Dense<RealValueType>;

```

The `gko::matrix::Ell` class is used here, but any other matrix class such as `gko::matrix::Coo`, `gko::matrix::Hybrid`, `gko::matrix::Csr` or `gko::matrix::Sellp` could also be used. Note. the behavior will depends GINKGO_MIXED_PRECISION flags and the actual implementation from different matrices.

```

using hp_mtx = gko::matrix::Ell<HighPrecision, IndexType>;
using lp_mtx = gko::matrix::Ell<LowPrecision, IndexType>;

```

Print the ginkgo version information.

```

std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;
    std::exit(-1);
}

```

Where do you want to run your operation?

The `gko::Executor` class is one of the cornerstones of Ginkgo. Currently, we have support for an `gko::OmpExecutor`, which uses OpenMP multi-threading in most of its kernels, a `gko::ReferenceExecutor`, a single threaded specialization of the OpenMP executor and a `gko::CudaExecutor` which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/ routines within Ginkgo should automatically work and run on the executor with any other changes.

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }
    };
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Preparing your data and transfer to the proper device.

Read the matrix using the read function and set the right hand side randomly.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. `gko::share` and `gko::give` are the functions that you would need to use.

read the matrix into HighPrecision and LowPrecision.

```
auto hp_A = share(gko::read<hp_mtx>(std::ifstream("data/A.mtx"), exec));
auto lp_A = share(gko::read<lp_mtx>(std::ifstream("data/A.mtx"), exec));
```

Set the shortcut for each dimension

```
auto A_dim = hp_A->get_size();
auto b_dim = gko::dim<2>{A_dim[1], 1};
auto x_dim = gko::dim<2>{A_dim[0], b_dim[1]};
auto host_b = hp_vec::create(exec->get_master(), b_dim);
```

fill the b vector with some random data

```
std::default_random_engine rand_engine(32);
auto dist = std::uniform_real_distribution<RealValueType>(0.0, 1.0);
for (int i = 0; i < host_b->get_size()[0]; i++) {
    host_b->at(i, 0) = get_rand_value<HighPrecision>(dist, rand_engine);
}
```

copy the data from host to device

```
auto hp_b = share(gko::clone(exec, host_b));
auto lp_b = share(lp_vec::create(exec));
lp_b->copy_from(hp_b);
```

create several result x vector in different precision

```
auto hp_x = share(hp_vec::create(exec, x_dim));
auto lp_x = share(lp_vec::create(exec, x_dim));
auto hplp_x = share(hp_x->clone());
auto lplp_x = share(hp_x->clone());
auto lphp_x = share(hp_x->clone());
```

Measure the time of apply

We measure the time among different combination of apply operation.

```

Hp * Hp -> Hp
auto hp_sec = timing(exec, hp_A, hp_b, hp_x);

Lp * Lp -> Lp
auto lp_sec = timing(exec, lp_A, lp_b, lp_x);

Hp * Lp -> Hp
auto hlp_sec = timing(exec, hp_A, lp_b, hlp_x);

Lp * Lp -> Hp
auto lplp_sec = timing(exec, lp_A, lp_b, lplp_x);

Lp * Hp -> Hp
auto lphp_sec = timing(exec, lp_A, hp_b, lphp_x);

```

To measure error of result. `neg_one` is an object that represent the number -1.0 which allows for a uniform interface when computing on any device. To compute the residual, all you need to do is call the `add_scaled` method, which in this case is an axpy and equivalent to the LAPACK axpy routine. Finally, you compute the euclidean 2-norm with the `compute_norm2` function.

```

auto neg_one = gko::initialize<hp_vec>({-1.0}, exec);
auto hp_x_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto hlp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lphp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff = hp_x->clone();
auto hlp_diff = hp_x->clone();
auto lplp_diff = hp_x->clone();
auto lphp_diff = hp_x->clone();
hp_x->compute_norm2(hp_x_norm);
lp_diff->add_scaled(neg_one, lp_x);
lp_diff->compute_norm2(lp_diff_norm);
hlp_diff->add_scaled(neg_one, hlp_x);
hlp_diff->compute_norm2(hlp_diff_norm);
lplp_diff->add_scaled(neg_one, lplp_x);
lplp_diff->compute_norm2(lplp_diff_norm);
lphp_diff->add_scaled(neg_one, lphp_x);
lphp_diff->compute_norm2(lphp_diff_norm);
exec->synchronize();
std::cout.precision(10);
std::cout << std::scientific;
std::cout << "High Precision time(s): " << hp_sec << std::endl;
std::cout << "High Precision result norm: " << hp_x_norm->at(0)
<< std::endl;
std::cout << "Low Precision time(s): " << lp_sec << std::endl;
std::cout << "Low Precision relative error: "
<< lp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Hp * Lp -> Hp time(s): " << hlp_sec << std::endl;
std::cout << "Hp * Lp -> Hp relative error: "
<< hlp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Lp * Lp -> Hp time(s): " << lplp_sec << std::endl;
std::cout << "Lp * Lp -> Hp relative error: "
<< lplp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Lp * Hp -> Hp time(s): " << lphp_sec << std::endl;
std::cout << "Lp * Hp -> Hp relative error: "
<< lphp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
}

```

Results

The following is the expected result (omp):

```

High Precision time(s): 2.0568800000e-05
High Precision result norm: 1.7725534898e+05
Low Precision time(s): 2.0955600000e-05
Low Precision relative error: 9.1052887738e-08
Hp * Lp -> Hp time(s): 2.1186100000e-05
Hp * Lp -> Hp relative error: 3.7799774251e-08
Lp * Lp -> Hp time(s): 2.0312300000e-05
Lp * Lp -> Hp relative error: 5.7910008031e-08
Lp * Hp -> Hp time(s): 2.0312300000e-05
Lp * Hp -> Hp relative error: 3.7173133506e-08

```

Comments about programming and debugging

The plain program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <chrono>
#include <random>
namespace {
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<!gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
{
    return value_dist(gen);
}
template <typename ValueType, typename ValueDistribution, typename Engine>
typename std::enable_if<gko::is_complex_s<ValueType>::value, ValueType>::type
get_rand_value(ValueDistribution&& value_dist, Engine&& gen)
{
    return ValueType(value_dist(gen), value_dist(gen));
}
double timing(std::shared_ptr<const gko::Executor> exec,
              std::shared_ptr<const gko::LinOp> A,
              std::shared_ptr<const gko::LinOp> b,
              std::shared_ptr<gko::LinOp> x)
{
    int warmup = 2;
    int rep = 10;
    for (int i = 0; i < warmup; i++) {
        A->apply(b, x);
    }
    double total_sec = 0;
    for (int i = 0; i < rep; i++) {
        auto xx = x->clone();
        exec->synchronize();
        auto start = std::chrono::steady_clock::now();
        A->apply(b, xx);
        exec->synchronize();
        auto stop = std::chrono::steady_clock::now();
        std::chrono::duration<double> duration_time = stop - start;
        total_sec += duration_time.count();
        if (i + 1 == rep) {
            x->copy_from(xx);
        }
    }
    return total_sec / rep;
}
// namespace
int main(int argc, char* argv[])
{
    using HighPrecision = double;
    using RealValueType = gko::remove_complex<HighPrecision>;
    using LowPrecision = float;
    using IndexType = int;
    using hp_vec = gko::matrix::Dense<HighPrecision>;
    using lp_vec = gko::matrix::Dense<LowPrecision>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using hp_mtx = gko::matrix::Ell<HighPrecision, IndexType>;
    using lp_mtx = gko::matrix::Ell<LowPrecision, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }}
    }
```

```

    }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }};
const auto exec = exec_map.at(executor_string()); // throws if not valid
auto hp_A = share(gko::read<hp_mtx>(std::ifstream("data/A.mtx"), exec));
auto lp_A = share(gko::read<lp_mtx>(std::ifstream("data/A.mtx"), exec));
auto A_dim = hp_A->get_size();
auto b_dim = gko::dim<2>{A_dim[1], 1};
auto x_dim = gko::dim<2>{A_dim[0], b_dim[1]};
auto host_b = hp_vec::create(exec->get_master(), b_dim);
std::default_random_engine rand_engine(32);
auto dist = std::uniform_real_distribution<RealValueType>(0.0, 1.0);
for (int i = 0; i < host_b->get_size()[0]; i++) {
    host_b->at(i, 0) = get_rand_value<HighPrecision>(dist, rand_engine);
}
auto hp_b = share(gko::clone(exec, host_b));
auto lp_b = share(lp_vec::create(exec));
lp_b->copy_from(hp_b);
auto hp_x = share(hp_vec::create(exec, x_dim));
auto lp_x = share(lp_vec::create(exec, x_dim));
auto hplp_x = share(hp_x->clone());
auto lplp_x = share(lp_x->clone());
auto lphp_x = share(hp_x->clone());
auto hp_sec = timing(exec, hp_A, hp_b, hp_x);
auto lp_sec = timing(exec, lp_A, lp_b, lp_x);
auto hplp_sec = timing(exec, hp_A, lp_b, hplp_x);
auto lplp_sec = timing(exec, lp_A, lp_b, lplp_x);
auto lphp_sec = timing(exec, lp_A, hp_b, lphp_x);
auto neg_one = gko::initialize<hp_vec>({-1.0}, exec);
auto hp_x_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto hplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lplp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lphp_diff_norm = gko::initialize<real_vec>({0.0}, exec->get_master());
auto lp_diff = hp_x->clone();
auto hplp_diff = hp_x->clone();
auto lplp_diff = lp_x->clone();
auto lphp_diff = hp_x->clone();
hp_x->compute_norm2(hp_x_norm);
lp_diff->add_scaled(neg_one, lp_x);
lp_diff->compute_norm2(lp_diff_norm);
hplp_diff->add_scaled(neg_one, hplp_x);
hplp_diff->compute_norm2(hplp_diff_norm);
lplp_diff->add_scaled(neg_one, lplp_x);
lplp_diff->compute_norm2(lplp_diff_norm);
lphp_diff->add_scaled(neg_one, lphp_x);
lphp_diff->compute_norm2(lphp_diff_norm);
exec->synchronize();
std::cout.precision(10);
std::cout << std::scientific;
std::cout << "High Precision time(s): " << hp_sec << std::endl;
std::cout << "High Precision result norm: " << hp_x_norm->at(0)
    << std::endl;
std::cout << "Low Precision time(s): " << lp_sec << std::endl;
std::cout << "Low Precision relative error: "
    << lp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Hp * Lp -> Hp time(s): " << hplp_sec << std::endl;
std::cout << "Hp * Lp -> Hp relative error: "
    << hplp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Lp * Lp -> Hp time(s): " << lplp_sec << std::endl;
std::cout << "Lp * Lp -> Hp relative error: "
    << lplp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
std::cout << "Lp * Hp -> Hp time(s): " << lphp_sec << std::endl;
std::cout << "Lp * Hp -> Hp relative error: "
    << lphp_diff_norm->at(0) / hp_x_norm->at(0) << "\n";
}

```


Chapter 31

The multigrid-preconditioned-solver program

The preconditioned solver example..

This example depends on preconditioned-solver.

This example shows how to use the multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a multigrid preconditioner. The example features the generating time and runtime of the CG solver.

The commented program

```
{ "cuda",  
  [] {  
    return gko::CudaExecutor::create(0,  
                                      gko::OmpExecutor::create());  
  }},  
{ "hip",  
  [] {  
    return gko::HipExecutor::create(0, gko::OmpExecutor::create());  
  }},  
{ "dpcpp",  
  [] {  
    return gko::DpcppExecutor::create(  
      0, gko::ReferenceExecutor::create());  
  }},  
{ "reference", [] { return gko::ReferenceExecutor::create(); } };
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS as 1 and initial guess as 0

```
gko::size_type size = A->get_size()[0];  
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));  
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));  
for (auto i = 0; i < size; i++) {  
  host_x->at(i, 0) = 0.;  
  host_b->at(i, 0) = 1.;  
}  
auto x = vec::create(exec);  
auto b = vec::create(exec);  
x->copy_from(host_x);  
b->copy_from(host_b);
```

Calculate initial residual by overwriting b

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);

```

copy b again

```
b->copy_from(host_b);
```

Create multigrid factory

```

std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid_gen =
    mg::build()
        .with_mg_level(pgm::build().with_deterministic(true))
        .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
        .on(exec);
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(100u),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_baseline(gko::stop::mode::absolute)
                .with_reduction_factor(tolerance))
        .with_preconditioner(multigrid_gen)
        .on(exec);

```

Create solver

```

std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);

```

Add logger

```

std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);

```

Solve system

```

exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);

```

Calculate residual

```

auto res = gko::as<vec>(logger->get_residual_norm());
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);

```

Print solver statistics

```

std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```

Results

This is the expected output:

```

Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1

```

```

4.3589
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1.69858e-09
CG iteration count:          39
CG generation time [ms]:    2.04293
CG execution time [ms]:     22.3874
CG execution time per iteration[ms]: 0.574036

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using mg = gko::solver::Multigrid;
    using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string()); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    gko::size_type size = A->get_size()[0];
    auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
    for (auto i = 0; i < size; i++) {
        host_x->at(i, 0) = 0.;
        host_b->at(i, 0) = 1.;
    }
    auto x = vec::create(exec);
    auto b = vec::create(exec);
    x->copy_from(host_x);
    b->copy_from(host_b);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto initres = gko::initialize<vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(initres);
    b->copy_from(host_b);
    std::shared_ptr<gko::LinOpFactory> multigrid_gen;
    multigrid_gen =
        mg::build()
            .with_mg_level(pgm::build().with_deterministic(true))
            .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
            .on(exec);
    const gko::remove_complex<ValueType> tolerance = 1e-8;
    auto solver_gen =
        cg::build()
            .with_criteria(gko::stop::Iteration::build().with_max_iters(100u),
                          gko::stop::ResidualNorm<ValueType>::build()
                              .with_baseline(gko::stop::mode::absolute)
                              .with_reduction_factor(tolerance))

```

```

        .with_preconditioner(multigrid_gen)
        .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(logger);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::as<vec>(logger->get_residual_norm());
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);
std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```

Chapter 32

The multigrid-preconditioned-solver-customized program

The customized multigrid preconditioned solver example..

This example depends on multigrid-preconditioned-solver.

This example shows how to customize the multigrid preconditioner.

In this example, we first read in a matrix from a file. The preconditioned CG solver is enhanced with a multigrid preconditioner. Several non-default options are used to create this preconditioner. The example features the generating time and runtime of the CG solver.

The commented program

```
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(
             0, gko::ReferenceExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }};
}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
```

Create RHS as 1 and initial guess as 0

```
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
}
```

```

}
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);

```

Calculate initial residual by overwriting b

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);

```

copy b again

```

b->copy_from(host_b);

```

Prepare the stopping criteria

```

const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::absolute)
    .with_reduction_factor(tolerance)
    .on(exec));

auto exact_tol_stop =
    gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::rhs_norm)
    .with_reduction_factor(1e-14)
    .on(exec));
std::shared_ptr<const gko::log::Convergence<ValueType> logger =
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);

```

Now we customize some settings of the multigrid preconditioner. First we choose a smoother. Since the input matrix is spd, we use iterative refinement with two iterations and an lc solver.

```

auto ic_gen = gko::share(
    ic::build()
    .with_factorization(gko::factorization::Ic<ValueType, int>::build())
    .on(exec));
auto smoother_gen = gko::share(
    gko::solver::build_smoother(ic_gen, 2u, static_cast<ValueType>(0.9)));

```

Use Pgm as the MultigridLevel factory.

```

auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));

```

Next we select a CG solver for the coarsest level. Again, since the input matrix is known to be spd, and the Pgm restriction preserves this characteristic, we can safely choose the CG. We reuse the lc factory here to generate an lc preconditioner. It is important to solve until machine precision here to get a good convergence rate.

```

auto coarsest_gen = gko::share(cg::build()
    .with_preconditioner(ic_gen)
    .with_criteria(iter_stop, exact_tol_stop)
    .on(exec));

```

Here we put the customized options together and create the multigrid factory.

```

std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid_gen =
    mg::build()
    .with_max_levels(10u)
    .with_min_coarse_rows(32u)
    .with_pre_smoother(smoother_gen)
    .with_post_uses_pre(true)
    .with_mg_level(mg_level_gen)
    .with_coarsest_solver(coarsest_gen)
    .with_default_initial_guess(gko::solver::initial_guess_mode::zero)
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec);

```

Create solver factory

```

auto solver_gen = cg::build()
    .with_criteria(iter_stop, tol_stop)
    .with_preconditioner(multigrid_gen)
    .on(exec);

```

Create solver

```

std::chrono::nanoseconds gen_time(0);

```

```

auto gen_tic = std::chrono::steady_clock::now();
auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);

```

Solve system

```

exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);

```

Calculate residual

```

auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r): \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r): \n";
write(std::cout, res);

```

Print solver statistics

```

std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]: "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]: "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]: "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```

Results

This is the expected output:

```

Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
25.9808
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
5.81328e-09
CG iteration count:          12
CG generation time [ms]:    1.41642
CG execution time [ms]:     6.59244
CG execution time per iteration[ms]: 0.54937

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using ir = gko::solver::Ir<ValueType>;
    using mg = gko::solver::Multigrid;

```

```

using ic = gko::preconditioner::Ic<gko::solver::LowerTrs<ValueType>>;
using pgm = gko::multigrid::Pgm<ValueType, IndexType>;
std::cout << gko::version_info::get() << std::endl;
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(
                 0, gko::ReferenceExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
gko::size_type size = A->get_size()[0];
auto host_x = vec::create(exec->get_master(), gko::dim<2>(size, 1));
auto host_b = vec::create(exec->get_master(), gko::dim<2>(size, 1));
for (auto i = 0; i < size; i++) {
    host_x->at(i, 0) = 0.;
    host_b->at(i, 0) = 1.;
}
auto x = vec::create(exec);
auto b = vec::create(exec);
x->copy_from(host_x);
b->copy_from(host_b);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto initres = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(initres);
b->copy_from(host_b);
const gko::remove_complex<ValueType> tolerance = 1e-8;
auto iter_stop =
    gko::share(gko::stop::Iteration::build().with_max_iters(100u).on(exec));
auto tol_stop = gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::absolute)
    .with_reduction_factor(tolerance)
    .on(exec));

auto exact_tol_stop =
    gko::share(gko::stop::ResidualNorm<ValueType>::build()
    .with_baseline(gko::stop::mode::rhs_norm)
    .with_reduction_factor(1e-14)
    .on(exec));

std::shared_ptr<const gko::log::Convergence<ValueType>> logger =
    gko::log::Convergence<ValueType>::create();
iter_stop->add_logger(logger);
tol_stop->add_logger(logger);
auto ic_gen = gko::share(
    ic::build()
    .with_factorization(gko::factorization::Ic<ValueType, int>::build())
    .on(exec));
auto smoother_gen = gko::share(
    gko::solver::build_smoother(ic_gen, 2u, static_cast<ValueType>(0.9)));
auto mg_level_gen =
    gko::share(pgm::build().with_deterministic(true).on(exec));
auto coarsest_gen = gko::share(cg::build()
    .with_preconditioner(ic_gen)
    .with_criteria(iter_stop, exact_tol_stop)
    .on(exec));

std::shared_ptr<gko::LinOpFactory> multigrid_gen;
multigrid_gen =
    mg::build()
    .with_max_levels(10u)
    .with_min_coarse_rows(32u)
    .with_pre_smoother(smoother_gen)
    .with_post_uses_pre(true)
    .with_mg_level(mg_level_gen)
    .with_coarsest_solver(coarsest_gen)
    .with_default_initial_guess(gko::solver::initial_guess_mode::zero)
    .with_criteria(gko::stop::Iteration::build().with_max_iters(1u))
    .on(exec);
auto solver_gen = cg::build()
    .with_criteria(iter_stop, tol_stop)
    .with_preconditioner(multigrid_gen)
    .on(exec);
std::chrono::nanoseconds gen_time(0);
auto gen_tic = std::chrono::steady_clock::now();

```

```

auto solver = solver_gen->generate(A);
exec->synchronize();
auto gen_toc = std::chrono::steady_clock::now();
gen_time +=
    std::chrono::duration_cast<std::chrono::nanoseconds>(gen_toc - gen_tic);
exec->synchronize();
std::chrono::nanoseconds time(0);
auto tic = std::chrono::steady_clock::now();
solver->apply(b, x);
exec->synchronize();
auto toc = std::chrono::steady_clock::now();
time += std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic);
auto res = gko::initialize<vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Initial residual norm sqrt(r^T r):  \n";
write(std::cout, initres);
std::cout << "Final residual norm sqrt(r^T r):  \n";
write(std::cout, res);
std::cout << "CG iteration count:          " << logger->get_num_iterations()
    << std::endl;
std::cout << "CG generation time [ms]:  "
    << static_cast<double>(gen_time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time [ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 << std::endl;
std::cout << "CG execution time per iteration[ms]:  "
    << static_cast<double>(time.count()) / 1000000.0 /
        logger->get_num_iterations()
    << std::endl;
}

```


Chapter 33

The nine-pt-stencil-solver program

The 9-point stencil example..

This example depends on simple-solver, three-pt-stencil-solver, poisson-solver, .

Introduction

This example solves a 2D Poisson equation:

$$\begin{aligned} &[\Omega = (0,1)^2 \setminus \Omega_b = [0,1]^2 \setminus \text{boundary} \setminus \partial\Omega = \Omega_b \setminus \Omega \\ &u : \Omega_b \rightarrow \mathbb{R} \setminus u'' = f \text{ in } \Omega \setminus u = u_D \text{ in } \partial\Omega \setminus] \end{aligned}$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization may be done by any order Taylor polynomial. For an equidistant grid with K "inner" discretization points $((x_1, y_1), \dots, (x_k, y_1), (x_1, y_2), \dots, (x_k, y_k, z_1))$ step size $(h = 1 / (K + 1))$ and a stencil $(\text{in } \mathbb{R}^{3 \times 3})$, the formula produces a system of linear equations

$$(\sum_{a,b=-1}^1 \text{stencil}(a,b) * u_{\{(i+a,j+b)\}} = -f_k h^2), \text{ on any inner node with a neighborhood of inner nodes}$$

On any node, where neighbor is on the border, the neighbor is replaced with a $(-\text{stencil}(a,b) * u_{\{(i+a,j+b)\}})$ and added to the right hand side vector. For example a node with a neighborhood of only edge nodes may look like this

$$[\sum_{a,b=-1}^1 (1,0) \text{stencil}(a,b) * u_{\{(i+a,j+b)\}} = -f_k h^2 - \sum_{a=-1}^1 \text{stencil}(a,1) * u_{\{(i+a,j+1)\}}]$$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function f is set to $(f(x,y) = 6x + 6y)$ (making the solution $(u(x,y) = x^3$

- $y^3)$), but that can be changed in the `main` function. Also the stencil values for the core, the faces, the edge and the corners can be changed when passing additional parameters.

The intention of this is to show how generation of stencil values and the right hand side vector changes when increasing the dimension.

About the example

The commented program

preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x,y) = 6x + 6y' (making the solution 'u(x,y) = x^3 + y^3'), but that can be changed in the 'main' function. Also the stencil values for the core, the faces, the edge and the corners can be changed when passing additional parameters. The intention of this is to show how generation of stencil values and the right hand side vector changes when increasing the dimension.

```
*****<DESCRIPTION>***** /
#include <array>
#include <chrono>
#include <iostream>
#include <map>
#include <string>
#include <vector>
#include <ginkgo/ginkgo.hpp>
```

Stencil values. Ordering can be seen in the main function Can also be changed by passing additional parameter when executing

```
constexpr double default_alpha = 10.0 / 3.0;
constexpr double default_beta = -2.0 / 3.0;
constexpr double default_gamma = -1.0 / 6.0;
/* * Possible alternative default values are
 * default_alpha = 8.0;
 * default_beta = -1.0;
 * default_gamma = -1.0;
 * /
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType dp, IndexType* row_ptrs,
                           IndexType* col_idxxs, ValueType* values,
                           ValueType* coefs)
{
    IndexType pos = 0;
    const size_t dp_2 = dp * dp;
    row_ptrs[0] = pos;
    for (IndexType k = 0; k < dp; ++k) {
        for (IndexType i = 0; i < dp; ++i) {
            const size_t index = i + k * dp;
            for (IndexType j = -1; j <= 1; ++j) {
                for (IndexType l = -1; l <= 1; ++l) {
                    const IndexType offset = 1 + 1 + 3 * (j + 1);
                    if ((k + j) >= 0 && (k + j) < dp && (i + l) >= 0 &&
                        (i + l) < dp) {
                        values[pos] = coefs[offset];
                        col_idxxs[pos] = index + 1 + dp * j;
                        ++pos;
                    }
                }
            }
            row_ptrs[index + 1] = pos;
        }
    }
}
```

Generates the RHS vector given f and the boundary conditions.

```
template <typename Closure, typename ClosureT, typename ValueType,
         typename IndexType>
void generate_rhs(IndexType dp, Closure f, ClosureT u, ValueType* rhs,
                 ValueType* coefs)
{
    const size_t dp_2 = dp * dp;
    const ValueType h = 1.0 / (dp + 1.0);
    for (IndexType i = 0; i < dp; ++i) {
        const auto yi = ValueType(i + 1) * h;
        for (IndexType j = 0; j < dp; ++j) {
            const auto xi = ValueType(j + 1) * h;
            const auto index = i * dp + j;
            rhs[index] = -f(xi, yi) * h * h;
        }
    }
}
```

Iterating over the edges to add boundary values and adding the overlapping 3x1 to the rhs

```
for (size_t i = 0; i < dp; ++i) {
    const auto xi = ValueType(i + 1) * h;
```

```

    const auto index_top = i;
    const auto index_bot = i + dp * (dp - 1);
    rhs[index_top] -= u(xi - h, 0.0) * coefs[0];
    rhs[index_top] -= u(xi, 0.0) * coefs[1];
    rhs[index_top] -= u(xi + h, 0.0) * coefs[2];
    rhs[index_bot] -= u(xi - h, 1.0) * coefs[6];
    rhs[index_bot] -= u(xi, 1.0) * coefs[7];
    rhs[index_bot] -= u(xi + h, 1.0) * coefs[8];
}
for (size_t i = 0; i < dp; ++i) {
    const auto yi = ValueType(i + 1) * h;
    const auto index_left = i * dp;
    const auto index_right = i * dp + (dp - 1);
    rhs[index_left] -= u(0.0, yi - h) * coefs[0];
    rhs[index_left] -= u(0.0, yi) * coefs[3];
    rhs[index_left] -= u(0.0, yi + h) * coefs[6];
    rhs[index_right] -= u(1.0, yi - h) * coefs[2];
    rhs[index_right] -= u(1.0, yi) * coefs[5];
    rhs[index_right] -= u(1.0, yi + h) * coefs[8];
}

```

remove the double corner values

```

    rhs[0] += u(0.0, 0.0) * coefs[0];
    rhs[(dp - 1)] += u(1.0, 0.0) * coefs[2];
    rhs[(dp - 1) * dp] += u(0.0, 1.0) * coefs[6];
    rhs[dp * dp - 1] += u(1.0, 1.0) * coefs[8];
}

```

Prints the solution u.

```

template <typename ValueType, typename IndexType>
void print_solution(IndexType dp, const ValueType* u)
{
    for (IndexType i = 0; i < dp; ++i) {
        for (IndexType j = 0; j < dp; ++j) {
            std::cout << u[i * dp + j] << ' ';
        }
        std::cout << '\n';
    }
    std::cout << std::endl;
}

```

Computes the 1-norm of the error given the computed u and the correct solution function correct_u.

```

template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType dp, const ValueType* u,
                                                Closure correct_u)
{
    const ValueType h = 1.0 / (dp + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (IndexType j = 0; j < dp; ++j) {
        const auto xi = ValueType(j + 1) * h;
        for (IndexType i = 0; i < dp; ++i) {
            using std::abs;
            const auto yi = ValueType(i + 1) * h;
            error +=
                abs(u[i * dp + j] - correct_u(xi, yi)) / abs(correct_u(xi, yi));
        }
    }
    return error;
}

template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                  unsigned int discretization_points, IndexType* row_ptrs,
                  IndexType* col_idxs, ValueType* values, ValueType* rhs,
                  ValueType* u, gko::remove_complex<ValueType> reduction_factor)
{
}

```

Some shortcuts

```

using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto& dp = discretization_points;
const gko::size_type dp_2 = dp * dp;

```

Figure out where to run the code

```

std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
    [] {

```

```

        return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
    },
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); } }];

```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

executor where the application initialized the data

```
const auto app_exec = exec->get_master();
```

Tell Ginkgo to use the data in our application

Matrix: we have to set the executor of the matrix to the one where we want SpMV's to run (in this case `exec`). When creating array views, we have to specify the executor where the data is (in this case `app_exec`).

If the two do not match, Ginkgo will automatically create a copy of the data on `exec` (however, it will not copy the data back once it is done

- here this is not important since we are not modifying the matrix).

```

auto matrix = mtx::create(
    exec, gko::dim<2>(dp_2),
    val_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), values),
    idx_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), col_idxs),
    idx_array::view(app_exec, dp_2 + 1, row_ptrs));

```

RHS: similar to matrix

```

auto b = vec::create(exec, gko::dim<2>(dp_2, 1),
                    val_array::view(app_exec, dp_2, rhs), 1);

```

Solution: we have to be careful here - if the executors are different, once we compute the solution the array will not be automatically copied back to the original memory locations. Fortunately, whenever `apply` is called on a linear operator (e.g. matrix, solver) the arguments automatically get copied to the executor where the operator is, and copied back once the operation is completed. Thus, in this case, we can just define the solution on `app_exec`, and it will be automatically transferred to/from `exec` if needed.

```

auto x = vec::create(app_exec, gko::dim<2>(dp_2, 1),
                    val_array::view(app_exec, dp_2, u), 1);

```

Generate solver

```

auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(dp_2),
                      gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor))
        .with_preconditioner(bj::build())
        .on(exec);
auto solver = solver_gen->generate(gko::give(matrix));

```

Solve system

```

    solver->apply(b, x);
}
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;

```

Print version information

```

std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && std::string(argv[1]) == "--help") {
    std::cerr
        << "Usage: " << argv[0]
        << " [executor] [DISCRETIZATION_POINTS] [alpha] [beta] [gamma]"
        << std::endl;

```

```

    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const IndexType discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
const ValueType alpha_c = argc >= 4 ? std::atof(argv[3]) : default_alpha;
const ValueType beta_c = argc >= 5 ? std::atof(argv[4]) : default_beta;
const ValueType gamma_c = argc >= 6 ? std::atof(argv[5]) : default_gamma;

```

clang-format off

```

std::array<ValueType, 9> coefs{
    gamma_c, beta_c, gamma_c,
    beta_c, alpha_c, beta_c,
    gamma_c, beta_c, gamma_c};

```

clang-format on

```

const auto dp = discretization_points;
const size_t dp_2 = dp * dp;

```

problem:

```

auto correct_u = [] (ValueType x, ValueType y) {
    return x * x * x + y * y * y;
};
auto f = [] (ValueType x, ValueType y) {
    return ValueType(6) * x + ValueType(6) * y;
};

```

matrix

```

std::vector<IndexType> row_ptrs(dp_2 + 1);
std::vector<IndexType> col_idxs((3 * dp - 2) * (3 * dp - 2));
std::vector<ValueType> values((3 * dp - 2) * (3 * dp - 2));

```

right hand side

```

std::vector<ValueType> rhs(dp_2);

```

solution

```

std::vector<ValueType> u(dp_2, 0.0);
generate_stencil_matrix(dp, row_ptrs.data(), col_idxs.data(), values.data(),
    coefs.data());

```

looking for solution $u = x^3$: $f = 6x$, $u(0) = 0$, $u(1) = 1$

```

generate_rhs(dp, f, correct_u, rhs.data(), coefs.data());
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
auto start_time = std::chrono::steady_clock::now();
solve_system(executor_string, dp, row_ptrs.data(), col_idxs.data(),
    values.data(), rhs.data(), u.data(), reduction_factor);
auto stop_time = std::chrono::steady_clock::now();
auto runtime_duration =
    static_cast<double>(
        std::chrono::duration_cast<std::chrono::nanoseconds>(stop_time -
            start_time)
        .count()) *
    1e-6;

```

Uncomment to print the solution print_solution(dp, u.data());

```

std::cout << "The average relative error is "
    << calculate_error(dp, u.data(), correct_u) /
        static_cast<gko::remove_complex<ValueType>>(dp_2)
    << std::endl;
std::cout << "The runtime is " << std::to_string(runtime_duration) << " ms"
    << std::endl;
}

```

Results

The expected output should be

```

The average relative error is 6.35715e-06
The runtime is 167.320520 ms

```

Comments about programming and debugging

The plain program

```

/*****<DESCRIPTION>*****/
This example solves a 2D Poisson equation:

```

```

\Omega = (0,1)^2
\Omega_b = [0,1]^2 (with boundary)
\partial\Omega = \Omega_b \backslash \Omega
u : \Omega_b \rightarrow \mathbb{R}
u'' = f in \Omega
u = u_D on \partial\Omega

```

using a finite difference method on an equidistant grid with 'K' discretization points ('K' can be controlled with a command line parameter). The discretization may be done by any order Taylor polynomial. For an equidistant grid with K "inner" discretization points (x1,y1), ..., (xk,y1), (x1,y2), ..., (xk,yk) step size $h = 1 / (K + 1)$ and a stencil $\setminus R^{\{3 \times 3\}}$, the formula produces a system of linear equations

```

\sum_{a,b=-1}^1 stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2, on any inner node with
a neighborhood of inner nodes

```

On any node, where neighbor is on the border, the neighbor is replaced with a '-stencil(a,b) * u_{i+a,j+b}' and added to the right hand side vector. For example a node with a neighborhood of only edge nodes may look like this

```

\sum_{a,b=-1}^1 stencil(a,b) * u_{(i+a,j+b)} = -f_k h^2 - \sum_{a=-1}^1
stencil(a,1) * u_{(i+a,j+1)}

```

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x,y) = 6x + 6y' (making the solution 'u(x,y) = x^3 + y^3'), but that can be changed in the 'main' function. Also the stencil values for the core, the faces, the edge and the corners can be changed when passing additional parameters.

The intention of this is to show how generation of stencil values and the right hand side vector changes when increasing the dimension.

```

*****/
#include <array>
#include <chrono>
#include <iostream>
#include <map>
#include <string>
#include <vector>
#include <ginkgo/ginkgo.hpp>
constexpr double default_alpha = 10.0 / 3.0;
constexpr double default_beta = -2.0 / 3.0;
constexpr double default_gamma = -1.0 / 6.0;
/* Possible alternative default values are
* default_alpha = 8.0;
* default_beta = -1.0;
* default_gamma = -1.0;
*/
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType dp, IndexType* row_ptrs,
                           IndexType* col_idxs, ValueType* values,
                           ValueType* coeffs)
{
    IndexType pos = 0;
    const size_t dp_2 = dp * dp;
    row_ptrs[0] = pos;
    for (IndexType k = 0; k < dp; ++k) {
        for (IndexType i = 0; i < dp; ++i) {
            const size_t index = i + k * dp;
            for (IndexType j = -1; j <= 1; ++j) {
                for (IndexType l = -1; l <= 1; ++l) {
                    const IndexType offset = 1 + 1 + 3 * (j + 1);
                    if ((k + j) >= 0 && (k + j) < dp && (i + l) >= 0 &&
                        (i + l) < dp) {
                        values[pos] = coeffs[offset];
                        col_idxs[pos] = index + 1 + dp * j;
                        ++pos;
                    }
                }
            }
            row_ptrs[index + 1] = pos;
        }
    }
}
template <typename Closure, typename ClosureT, typename ValueType,

```



```

        typename IndexType>
void generate_rhs(IndexType dp, Closure f, ClosureT u, ValueType* rhs,
                 ValueType* coefs)
{
    const size_t dp_2 = dp * dp;
    const ValueType h = 1.0 / (dp + 1.0);
    for (IndexType i = 0; i < dp; ++i) {
        const auto yi = ValueType(i + 1) * h;
        for (IndexType j = 0; j < dp; ++j) {
            const auto xi = ValueType(j + 1) * h;
            const auto index = i * dp + j;
            rhs[index] = -f(xi, yi) * h * h;
        }
    }
    for (size_t i = 0; i < dp; ++i) {
        const auto xi = ValueType(i + 1) * h;
        const auto index_top = i;
        const auto index_bot = i + dp * (dp - 1);
        rhs[index_top] -= u(xi - h, 0.0) * coefs[0];
        rhs[index_top] -= u(xi, 0.0) * coefs[1];
        rhs[index_top] -= u(xi + h, 0.0) * coefs[2];
        rhs[index_bot] -= u(xi - h, 1.0) * coefs[6];
        rhs[index_bot] -= u(xi, 1.0) * coefs[7];
        rhs[index_bot] -= u(xi + h, 1.0) * coefs[8];
    }
    for (size_t i = 0; i < dp; ++i) {
        const auto yi = ValueType(i + 1) * h;
        const auto index_left = i * dp;
        const auto index_right = i * dp + (dp - 1);
        rhs[index_left] -= u(0.0, yi - h) * coefs[0];
        rhs[index_left] -= u(0.0, yi) * coefs[3];
        rhs[index_left] -= u(0.0, yi + h) * coefs[6];
        rhs[index_right] -= u(1.0, yi - h) * coefs[2];
        rhs[index_right] -= u(1.0, yi) * coefs[5];
        rhs[index_right] -= u(1.0, yi + h) * coefs[8];
    }
    rhs[0] += u(0.0, 0.0) * coefs[0];
    rhs[(dp - 1)] += u(1.0, 0.0) * coefs[2];
    rhs[(dp - 1) * dp] += u(0.0, 1.0) * coefs[6];
    rhs[dp * dp - 1] += u(1.0, 1.0) * coefs[8];
}

template <typename ValueType, typename IndexType>
void print_solution(IndexType dp, const ValueType* u)
{
    for (IndexType i = 0; i < dp; ++i) {
        for (IndexType j = 0; j < dp; ++j) {
            std::cout << u[i * dp + j] << ' ';
        }
        std::cout << '\n';
    }
    std::cout << std::endl;
}

template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType dp, const ValueType* u,
                                              Closure correct_u)
{
    const ValueType h = 1.0 / (dp + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (IndexType j = 0; j < dp; ++j) {
        const auto xi = ValueType(j + 1) * h;
        for (IndexType i = 0; i < dp; ++i) {
            using std::abs;
            const auto yi = ValueType(i + 1) * h;
            error +=
                abs(u[i * dp + j] - correct_u(xi, yi)) / abs(correct_u(xi, yi));
        }
    }
    return error;
}

template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                 unsigned int discretization_points, IndexType* row_ptrs,
                 IndexType* col_idxs, ValueType* values, ValueType* rhs,
                 ValueType* u, gko::remove_complex<ValueType> reduction_factor)
{
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using val_array = gko::array<ValueType>;
    using idx_array = gko::array<IndexType>;
    const auto& dp = discretization_points;
    const gko::size_type dp_2 = dp * dp;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>
            exec_map{
                {"omp", [] { return gko::OmpExecutor::create(); }},
                {"cuda",

```

```

    [] {
        return gko::CudaExecutor::create(0,
                                          gko::OmpExecutor::create());
    },
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                             gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }});
const auto exec = exec_map.at(executor_string()); // throws if not valid
const auto app_exec = exec->get_master();
auto matrix = mtx::create(
    exec, gko::dim<2>(dp_2),
    val_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), values),
    idx_array::view(app_exec, (3 * dp - 2) * (3 * dp - 2), col_idxxs),
    idx_array::view(app_exec, dp_2 + 1, row_ptrs));
auto b = vec::create(exec, gko::dim<2>(dp_2, 1),
    val_array::view(app_exec, dp_2, rhs), 1);
auto x = vec::create(app_exec, gko::dim<2>(dp_2, 1),
    val_array::view(app_exec, dp_2, u), 1);
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(dp_2),
                      gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor))
        .with_preconditioner(bj::build())
        .on(exec);
auto solver = solver_gen->generate(gko::give(matrix));
solver->apply(b, x);
}
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && std::string(argv[1]) == "--help") {
        std::cerr
            << "Usage: " << argv[0]
            << " [executor] [DISCRETIZATION_POINTS] [alpha] [beta] [gamma]"
            << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const IndexType discretization_points =
        argc >= 3 ? std::atoi(argv[2]) : 100;
    const ValueType alpha_c = argc >= 4 ? std::atof(argv[3]) : default_alpha;
    const ValueType beta_c = argc >= 5 ? std::atof(argv[4]) : default_beta;
    const ValueType gamma_c = argc >= 6 ? std::atof(argv[5]) : default_gamma;
    std::array<ValueType, 9> coeffs{
        gamma_c, beta_c, gamma_c,
        beta_c, alpha_c, beta_c,
        gamma_c, beta_c, gamma_c};
    const auto dp = discretization_points;
    const size_t dp_2 = dp * dp;
    auto correct_u = [] (ValueType x, ValueType y) {
        return x * x * x + y * y * y;
    };
    auto f = [] (ValueType x, ValueType y) {
        return ValueType(6) * x + ValueType(6) * y;
    };
    std::vector<IndexType> row_ptrs(dp_2 + 1);
    std::vector<IndexType> col_idxxs((3 * dp - 2) * (3 * dp - 2));
    std::vector<ValueType> values((3 * dp - 2) * (3 * dp - 2));
    std::vector<ValueType> rhs(dp_2);
    std::vector<ValueType> u(dp_2, 0.0);
    generate_stencil_matrix(dp, row_ptrs.data(), col_idxxs.data(), values.data(),
        coeffs.data());
    generate_rhs(dp, f, correct_u, rhs.data(), coeffs.data());
    const gko::remove_complex<ValueType> reduction_factor = 1e-7;
    auto start_time = std::chrono::steady_clock::now();
    solve_system(executor_string, dp, row_ptrs.data(), col_idxxs.data(),
        values.data(), rhs.data(), u.data(), reduction_factor);
    auto stop_time = std::chrono::steady_clock::now();
    auto runtime_duration =
        static_cast<double>(
            std::chrono::duration_cast<std::chrono::nanoseconds>(stop_time -
                                                                    start_time)
                .count()) *
        1e-6;
    std::cout << "The average relative error is "
        << calculate_error(dp, u.data(), correct_u) /
            static_cast<gko::remove_complex<ValueType>>(dp_2)

```

```
        « std::endl;
std::cout « "The runtime is " « std::to_string(runtime_duration) « " ms"
        « std::endl;
}
```


Chapter 34

The papi-logging program

The papi logging example..

This example depends on simple-solver-logging.

Introduction

About the example

The commented program

```
}
template <typename T>
std::string to_string(T* ptr)
{
    std::ostringstream os;
    os << reinterpret_cast<gko::uintptr>(ptr);
    return os.str();
}
// namespace
int init_papi_counters(std::string solver_name, std::string A_name)
{
```

Initialize PAPI, add events and start it up

```
    int eventset = PAPI_NULL;
    int ret_val = PAPI_library_init(PAPI_VER_CURRENT);
    if (ret_val != PAPI_VER_CURRENT) {
        std::cerr << "Error at PAPI_library_init()" << std::endl;
        std::exit(-1);
    }
    ret_val = PAPI_create_eventset(&eventset);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_create_eventset()" << std::endl;
        std::exit(-1);
    }
    std::string simple_apply_string("sde::ginkgo0::linop_apply_completed::");
    std::string advanced_apply_string(
        "sde::ginkgo0::linop_advanced_apply_completed::");
    papi_add_event(simple_apply_string + solver_name, eventset);
    papi_add_event(simple_apply_string + A_name, eventset);
    papi_add_event(advanced_apply_string + A_name, eventset);
    ret_val = PAPI_start(eventset);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_start()" << std::endl;
        std::exit(-1);
    }
    return eventset;
}
void print_papi_counters(int eventset)
{
```

Stop PAPI and read the `linop_apply_completed` event for all of them

```
long long int values[3];
int ret_val = PAPI_stop(eventset, values);
if (PAPI_OK != ret_val) {
    std::cerr << "Error at PAPI_stop()" << std::endl;
    std::exit(-1);
}
PAPI_shutdown();
```

Print all values returned from PAPI

```
std::cout << "PAPI SDE counters:" << std::endl;
std::cout << "solver did " << values[0] << " applies." << std::endl;
std::cout << "A did " << values[1] << " simple applies." << std::endl;
std::cout << "A did " << values[2] << " advanced applies." << std::endl;
}
int main(int argc, char* argv[])
{
```

Some shortcuts

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
```

Print version information

```
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
```

Generate solver

```
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor))
        .on(exec);
auto solver = solver_gen->generate(A);
```

In this example, we split as much as possible the Ginkgo solver/logger and the PAPI interface. Note that the PAPI ginkgo namespaces are of the form `sde::ginkgo<x>` where `<x>` starts from 0 and is incremented with every new PAPI logger.

```
int eventset =
    init_papi_counters(to_string(solver.get()), to_string(A.get()));
```

Create a PAPI logger and add it to relevant LinOps

```
auto logger = gko::log::Papi<ValueType>::create(
    gko::log::Logger::linop_apply_completed_mask |
    gko::log::Logger::linop_advanced_apply_completed_mask);
solver->add_logger(logger);
A->add_logger(logger);
```

Solve system

```
solver->apply(b, x);
```

Stop PAPI event gathering and print the counters

```
print_papi_counters(eventset);
```

Print solution

```
std::cout << "Solution (x):  \n";
write(std::cout, x);
```

Calculate residual

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):  \n";
write(std::cout, res);
}
```

Results

The following is the expected result:

```
PAPI SDE counters:
solver did 1 applies.
A did 20 simple applies.
A did 1 advanced applies.
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
8.87107e-16
```

Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <thread>
#include <papi.h>
#include <ginkgo/ginkgo.hpp>
```

```

namespace {
void papi_add_event(const std::string& event_name, int& eventset)
{
    int code;
    int ret_val = PAPI_event_name_to_code(event_name.c_str(), &code);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_name_to_code()" << std::endl;
        std::exit(-1);
    }
    ret_val = PAPI_add_event(eventset, code);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_name_to_code()" << std::endl;
        std::exit(-1);
    }
}

template <typename T>
std::string to_string(T* ptr)
{
    std::ostringstream os;
    os << reinterpret_cast<gko::uintptr>(ptr);
    return os.str();
}

// namespace
int init_papi_counters(std::string solver_name, std::string A_name)
{
    int eventset = PAPI_NULL;
    int ret_val = PAPI_library_init(PAPI_VER_CURRENT);
    if (ret_val != PAPI_VER_CURRENT) {
        std::cerr << "Error at PAPI_library_init()" << std::endl;
        std::exit(-1);
    }
    ret_val = PAPI_create_eventset(&eventset);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_create_eventset()" << std::endl;
        std::exit(-1);
    }
    std::string simple_apply_string("sde::ginkgo0::linop_apply_completed:");
    std::string advanced_apply_string(
        "sde::ginkgo0::linop_advanced_apply_completed:");
    papi_add_event(simple_apply_string + solver_name, eventset);
    papi_add_event(simple_apply_string + A_name, eventset);
    papi_add_event(advanced_apply_string + A_name, eventset);
    ret_val = PAPI_start(eventset);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_start()" << std::endl;
        std::exit(-1);
    }
    return eventset;
}

void print_papi_counters(int eventset)
{
    long long int values[3];
    int ret_val = PAPI_stop(eventset, values);
    if (PAPI_OK != ret_val) {
        std::cerr << "Error at PAPI_stop()" << std::endl;
        std::exit(-1);
    }
    PAPI_shutdown();
    std::cout << "PAPI SDE counters:" << std::endl;
    std::cout << "solver did " << values[0] << " applies." << std::endl;
    std::cout << "A did " << values[1] << " simple applies." << std::endl;
    std::cout << "A did " << values[2] << " advanced applies." << std::endl;
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>()>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",

```



```

    [] {
        return gko::HipExecutor::create(0, gko::OmpExecutor::create());
    },
    {"dpcpp",
    [] {
        return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
    },
    {"reference", [] { return gko::ReferenceExecutor::create(); }});
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                          .with_reduction_factor(reduction_factor))
        .on(exec);
auto solver = solver_gen->generate(A);
int eventset =
    init_papi_counters(to_string(solver.get()), to_string(A.get()));
auto logger = gko::log::Papi<ValueType>::create(
    gko::log::Logger::linop_apply_completed_mask |
    gko::log::Logger::linop_advanced_apply_completed_mask);
solver->add_logger(logger);
A->add_logger(logger);
solver->apply(b, x);
print_papi_counters(eventset);
std::cout << "Solution (x):  \n";
write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):  \n";
write(std::cout, res);
}

```


Chapter 35

The par-ilu-convergence program

The ParILU convergence example..

This example depends on simple-solver.

Introduction

About the example

This example can be used to inspect the convergence behavior of parallel incomplete factorizations.

The commented program

```
auto try_generate(Function fun) -> decltype(fun())
{
    decltype(fun()) result;
    try {
        result = fun();
    } catch (const gko::Error& err) {
        std::cerr << "Error: " << err.what() << '\n';
        std::exit(-1);
    }
    return result;
}

template <typename ValueType, typename IndexType>
double compute_ilu_residual_norm(
    const gko::matrix::Csr<ValueType, IndexType>* residual,
    const gko::matrix::Csr<ValueType, IndexType>* mtx)
{
    gko::matrix_data<ValueType, IndexType> residual_data;
    gko::matrix_data<ValueType, IndexType> mtx_data;
    residual->write(residual_data);
    mtx->write(mtx_data);
    residual_data.sort_row_major();
    mtx_data.sort_row_major();
    auto it = mtx_data.nonzeros.begin();
    double residual_norm{};
    for (auto entry : residual_data.nonzeros) {
        auto ref_row = it->row;
        auto ref_col = it->column;
        if (entry.row == ref_row && entry.column == ref_col) {
            residual_norm += gko::squared_norm(entry.value);
            ++it;
        }
    }
    return std::sqrt(residual_norm);
}

int main(int argc, char* argv[])
{

```

```
using ValueType = double;
using IndexType = int;
```

print usage message

```
if (argc < 2 || executors.find(argv[1]) == executors.end()) {
    std::cerr << "Usage:  executable"
                << "  <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
                << "[<parilu|parilut|paric|parict> [<max-iterations>] "
                << "[<num-repetitions>] [<fill-in-limit>]\n";
    return -1;
}
```

generate executor based on first argument

```
auto exec = try_generate([&] { return executors.at(argv[1])(); });
```

set matrix and preconditioner name with default values

```
std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
std::string precondition = argc < 4 ? "parilu" : argv[3];
int max_iterations = argc < 5 ? 10 : std::stoi(argv[4]);
int num_repetitions = argc < 6 ? 10 : std::stoi(argv[5]);
double limit = argc < 7 ? 2 : std::stod(argv[6]);
```

load matrix file into Csr format

```
auto mtx = gko::share(try_generate([&] {
    std::ifstream mtx_stream{matrix};
    if (!mtx_stream) {
        throw GKO_STREAM_ERROR("Unable to open matrix file");
    }
    std::cerr << "Reading " << matrix << std::endl;
    return gko::read<gko::matrix::Csr<ValueType, IndexType>(mtx_stream,
                                                             exec);
}));
auto factory_generator =
    [&](gko::size_type iteration) -> std::shared_ptr<gko::LinOpFactory> {
    if (precond == "parilu") {
        return gko::factorization::ParIlut<ValueType, IndexType>::build()
            .with_iterations(iteration)
            .on(exec);
    } else if (precond == "paric") {
        return gko::factorization::ParIc<ValueType, IndexType>::build()
            .with_iterations(iteration)
            .on(exec);
    } else if (precond == "parilut") {
        return gko::factorization::ParIlut<ValueType, IndexType>::build()
            .with_fill_in_limit(limit)
            .with_iterations(iteration)
            .on(exec);
    } else if (precond == "parict") {
        return gko::factorization::ParIct<ValueType, IndexType>::build()
            .with_fill_in_limit(limit)
            .with_iterations(iteration)
            .on(exec);
    } else {
        GKO_NOT_IMPLEMENTED;
    }
};
auto one = gko::initialize<gko::matrix::Dense<ValueType>({1.0}, exec);
auto minus_one =
    gko::initialize<gko::matrix::Dense<ValueType>({-1.0}, exec);
for (int it = 1; it <= max_iterations; ++it) {
    auto factory = factory_generator(it);
    std::cout << it << " ";
    std::vector<long> times;
    std::vector<double> residuals;
    for (int rep = 0; rep < num_repetitions; ++rep) {
        auto tic = std::chrono::high_resolution_clock::now();
        auto result =
            gko::as<gko::Composition<ValueType>(factory->generate(mtx));
        exec->synchronize();
        auto toc = std::chrono::high_resolution_clock::now();
        auto residual = gko::clone(exec, mtx);
        result->get_operators()[0]->apply(one, result->get_operators()[1],
                                         minus_one, residual);

        times.push_back(
            std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic)
                .count());
        residuals.push_back(
            compute_ilu_residual_norm(residual.get(), mtx.get()));
    }
    for (auto el : times) {
        std::cout << el << " ";
    }
    for (auto el : residuals) {
```



```

auto it = mtx_data.nonzeros.begin();
double residual_norm{};
for (auto entry : residual_data.nonzeros) {
    auto ref_row = it->row;
    auto ref_col = it->column;
    if (entry.row == ref_row && entry.column == ref_col) {
        residual_norm += gko::squared_norm(entry.value);
        ++it;
    }
}
return std::sqrt(residual_norm);
}
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    if (argc < 2 || executors.find(argv[1]) == executors.end()) {
        std::cerr << "Usage: executable"
            << " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
            << "[<parilu|parilut|paric|parict> [<max-iterations>] "
            << "[<num-repetitions>] [<fill-in-limit>]\n";
        return -1;
    }
    auto exec = try_generate([&] { return executors.at(argv[1])(); });
    std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
    std::string precondition = argc < 4 ? "parilu" : argv[3];
    int max_iterations = argc < 5 ? 10 : std::stoi(argv[4]);
    int num_repetitions = argc < 6 ? 10 : std::stoi(argv[5]);
    double limit = argc < 7 ? 2 : std::stod(argv[6]);
    auto mtx = gko::share(try_generate([&] {
        std::ifstream mtx_stream(matrix);
        if (!mtx_stream) {
            throw GKO_STREAM_ERROR("Unable to open matrix file");
        }
        std::cerr << "Reading " << matrix << std::endl;
        return gko::readgko::matrix::Csr<ValueType, IndexType>(mtx_stream,
            exec);
    }));
    auto factory_generator =
    [&](gko::size_type iteration) -> std::shared_ptr<gko::LinOpFactory> {
        if (precondition == "parilu") {
            return gko::factorization::ParIl<ValueType, IndexType>::build()
                .with_iterations(iteration)
                .on(exec);
        } else if (precondition == "paric") {
            return gko::factorization::ParIc<ValueType, IndexType>::build()
                .with_iterations(iteration)
                .on(exec);
        } else if (precondition == "parilut") {
            return gko::factorization::ParIlut<ValueType, IndexType>::build()
                .with_fill_in_limit(limit)
                .with_iterations(iteration)
                .on(exec);
        } else if (precondition == "parict") {
            return gko::factorization::ParIct<ValueType, IndexType>::build()
                .with_fill_in_limit(limit)
                .with_iterations(iteration)
                .on(exec);
        } else {
            GKO_NOT_IMPLEMENTED;
        }
    };
    auto one = gko::initialize<gko::matrix::Dense<ValueType>>({1.0}, exec);
    auto minus_one =
    gko::initialize<gko::matrix::Dense<ValueType>>({-1.0}, exec);
    for (int it = 1; it <= max_iterations; ++it) {
        auto factory = factory_generator(it);
        std::cout << it << ' ';
        std::vector<long> times;
        std::vector<double> residuals;
        for (int rep = 0; rep < num_repetitions; ++rep) {
            auto tic = std::chrono::high_resolution_clock::now();
            auto result =
            gko::as<gko::Composition<ValueType>>(factory->generate(mtx));
            exec->synchronize();
            auto toc = std::chrono::high_resolution_clock::now();
            auto residual = gko::clone(exec, mtx);
            result->get_operators()[0]->apply(one, result->get_operators()[1],
                minus_one, residual);
            times.push_back(
                std::chrono::duration_cast<std::chrono::nanoseconds>(toc - tic)
                    .count());
            residuals.push_back(
                compute_ilu_residual_norm(residual.get(), mtx.get()));
        }
        for (auto el : times) {
            std::cout << el << ' ';

```

```
    }  
    for (auto el : residuals) {  
        std::cout << el << ' ';  
    }  
    std::cout << '\n';  
}  
}
```


Chapter 36

The performance-debugging program

The simple solver with performance debugging example..

This example depends on simple-solver-logging, minimal-cuda-solver.

Introduction

About the example

This example runs a solver on a test problem and shows how to use loggers to debug performance and convergence rate.

The commented program

creates a zero vector

```
template <typename ValueType>
std::unique_ptr<vec<ValueType>> create_vector(
    std::shared_ptr<const gko::Executor> exec, gko::size_type size,
    ValueType value)
{
    auto res = vec<ValueType>::create(exec);
    res->read(gko::matrix_data<ValueType>(gko::dim<2>{size, 1}, value));
    return res;
}
```

utilities for computing norms and residuals

```
template <typename ValueType>
ValueType get_first_element(const vec<ValueType>* norm)
{
    return norm->get_executor()->copy_val_to_host(norm->get_const_values());
}

template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(const vec<ValueType>* b)
{
    auto exec = b->get_executor();
    auto b_norm = gko::initialize<real_vec<ValueType>>({0.0}, exec);
    b->compute_norm2(b_norm);
    return get_first_element(b_norm.get());
}

template <typename ValueType>
gko::remove_complex<ValueType> compute_residual_norm(
    const gko::LinOp* system_matrix, const vec<ValueType>* b,
    const vec<ValueType>* x)
{
    auto exec = system_matrix->get_executor();
    auto one = gko::initialize<vec<ValueType>>({1.0}, exec);
}
```

```

    auto neg_one = gko::initialize<vec<ValueType>>({-1.0}, exec);
    auto res = gko::clone(b);
    system_matrix->apply(one, x, neg_one, res);
    return compute_norm(res.get());
}
} // namespace utils
namespace loggers {

```

A logger that accumulates the time of all operations. For each operation type (allocations, free, copy, internal operations i.e. kernels), the timing is taken before and after. This can create significant overhead since to ensure proper timings, calls to `synchronize` are required.

```

struct OperationLogger : gko::log::Logger {
    void on_allocation_started(const gko::Executor* exec,
                             const gko::size_type&)const override
    {
        this->start_operation(exec, "allocate");
    }
    void on_allocation_completed(const gko::Executor* exec,
                                const gko::size_type&,
                                const gko::uintptr&)const override
    {
        this->end_operation(exec, "allocate");
    }
    void on_free_started(const gko::Executor* exec,
                        const gko::uintptr&)const override
    {
        this->start_operation(exec, "free");
    }
    void on_free_completed(const gko::Executor* exec,
                           const gko::uintptr&)const override
    {
        this->end_operation(exec, "free");
    }
    void on_copy_started(const gko::Executor* from, const gko::Executor* to,
                        const gko::uintptr&, const gko::uintptr&,
                        const gko::size_type&)const override
    {
        from->synchronize();
        this->start_operation(to, "copy");
    }
    void on_copy_completed(const gko::Executor* from, const gko::Executor* to,
                           const gko::uintptr&, const gko::uintptr&,
                           const gko::size_type&)const override
    {
        from->synchronize();
        this->end_operation(to, "copy");
    }
    void on_operation_launched(const gko::Executor* exec,
                              const gko::Operation* op)const override
    {
        this->start_operation(exec, op->get_name());
    }
    void on_operation_completed(const gko::Executor* exec,
                               const gko::Operation* op)const override
    {
        this->end_operation(exec, op->get_name());
    }
    void write_data(std::ostream& ostream)
    {
        for (const auto& entry : total) {
            ostream << "\t" << entry.first.c_str() << ": "
                << std::chrono::duration_cast<std::chrono::nanoseconds>(
                    entry.second)
                    .count()
                << std::endl;
        }
    }
private:

```

Helper which synchronizes and starts the time before every operation.

```

void start_operation(const gko::Executor* exec,
                   const std::string& name)const
{
    nested.emplace_back(0);
    exec->synchronize();
    start[name] = std::chrono::steady_clock::now();
}

```

Helper to compute the end time and store the operation's time at its end. Also time nested operations.

```

void end_operation(const gko::Executor* exec, const std::string& name)const
{
    exec->synchronize();
    const auto end = std::chrono::steady_clock::now();
}

```

```
const auto diff = end - start[name];
```

make sure timings for nested operations are not counted twice

```
total[name] += diff - nested.back();
nested.pop_back();
if (nested.size() > 0) {
    nested.back() += diff;
}
mutable std::map<std::string, std::chrono::steady_clock::time_point> start;
mutable std::map<std::string, std::chrono::steady_clock::duration> total;
```

the position *i* of this vector holds the total time spend on child operations on nesting level *i*

```
mutable std::vector<std::chrono::steady_clock::duration> nested;
};
```

This logger tracks the persistently allocated data

```
struct StorageLogger : gko::log::Logger {
```

Store amount of bytes allocated on every allocation

```
void on_allocation_completed(const gko::Executor*,
                             const gko::size_type& num_bytes,
                             const gko::uintptr& location) const override
{
    storage[location] = num_bytes;
}
```

Reset the amount of bytes on every free

```
void on_free_completed(const gko::Executor*,
                       const gko::uintptr& location) const override
{
    storage[location] = 0;
}
```

Write the data after summing the total from all allocations

```
void write_data(std::ostream& ostream)
{
    gko::size_type total{};
    for (const auto& e : storage) {
        total += e.second;
    }
    ostream << "Storage: " << total << std::endl;
}
private:
mutable std::unordered_map<gko::uintptr, gko::size_type> storage;
};
```

Logs true and recurrent residuals of the solver

```
template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
```

Depending on the available information, store the norm or compute it from the residual. If the true residual norm could not be computed, store the value -1.0 .

```
void on_iteration_complete(const gko::LinOp*, const gko::size_type&,
                           const gko::LinOp* residual,
                           const gko::LinOp* solution,
                           const gko::LinOp* residual_norm) const override
{
    if (residual_norm) {
        rec_res_norms.push_back(utils::get_first_element(
            gko::as<real_vec<ValueType>>(residual_norm)));
    } else {
        rec_res_norms.push_back(
            utils::compute_norm(gko::as<vec<ValueType>>(residual)));
    }
    if (solution) {
        true_res_norms.push_back(utils::compute_residual_norm(
            matrix, b, gko::as<vec<ValueType>>(solution)));
    } else {
        true_res_norms.push_back(-1.0);
    }
}
ResidualLogger(const gko::LinOp* matrix, const vec<ValueType>* b)
: gko::log::Logger(gko::log::Logger::iteration_complete_mask),
  matrix{matrix},
  b{b}
{}
void write_data(std::ostream& ostream)
{
}
```

```

    ostream << "Recurrent Residual Norms: " << std::endl;
    ostream << "[" << std::endl;
    for (const auto& entry : rec_res_norms) {
        ostream << "\t" << entry << std::endl;
    }
    ostream << "];" << std::endl;
    ostream << "True Residual Norms: " << std::endl;
    ostream << "[" << std::endl;
    for (const auto& entry : true_res_norms) {
        ostream << "\t" << entry << std::endl;
    }
    ostream << "];" << std::endl;
}
private:
    const gko::LinOp* matrix;
    const vec<ValueType>* b;
    mutable std::vector<gko::remove_complex<ValueType>> rec_res_norms;
    mutable std::vector<gko::remove_complex<ValueType>> true_res_norms;
};
} // namespace loggers
namespace {

```

Print usage help

```

void print_usage(const char* filename)
{
    std::cerr << "Usage: " << filename << " [executor] [matrix file]"
               << std::endl;
    std::cerr << "matrix file should be a file in matrix market format. "
               << "The file data/A.mtx is provided as an example."
               << std::endl;
    std::exit(-1);
}
template <typename ValueType>
void print_vector(const gko::matrix::Dense<ValueType>* vec)
{
    auto elements_to_print = std::min(gko::size_type(10), vec->get_size()[0]);
    std::cout << "[" << std::endl;
    for (int i = 0; i < elements_to_print; ++i) {
        std::cout << "\t" << vec->at(i) << std::endl;
    }
    std::cout << "];" << std::endl;
}
} // namespace
int main(int argc, char* argv[])
{

```

Parametrize the benchmark here Pick a value type

```

using ValueType = double;
using IndexType = int;

```

Pick a matrix format

```

using mtx = gko::matrix::Csr<ValueType, IndexType>;

```

Pick a solver

```

using solver = gko::solver::Cg<ValueType>;

```

Pick a preconditioner type

```

using preconditioner = gko::matrix::IdentityFactory<ValueType>;

```

Pick a residual norm reduction value

```

const gko::remove_complex<ValueType> reduction_factor = 1e-12;

```

Pick an output file name

```

const auto of_name = "log.txt";

```

Simple shortcut

```

using vec = gko::matrix::Dense<ValueType>;

```

Print version information

```

std::cout << gko::version_info::get() << std::endl;

```

Figure out where to run the code

```

if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}

```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read the input matrix file directory

```
std::string input_mtx = "data/A.mtx";
if (argc == 3) {
    input_mtx = std::string(argv[2]);
}
```

Read data: A is read from disk Create a StorageLogger to track the size of A

```
auto storage_logger = std::make_shared<loggers::StorageLogger>();
```

Add the logger to the executor

```
exec->add_logger(storage_logger);
```

Read the matrix A from file

```
auto A = gko::share(gko::read<mtx>(std::ifstream(input_mtx), exec));
```

Remove the storage logger

```
exec->remove_logger(storage_logger);
```

Pick a maximum iteration count

```
const auto max_iters = A->get_size()[0];
```

Generate b and x vectors

```
auto b = utils::create_vector<ValueType>(exec, A->get_size()[0], 1.0);
auto x = utils::create_vector<ValueType>(exec, A->get_size()[0], 0.0);
```

Declare the solver factory. The preconditioner's arguments should be adapted if needed.

```
auto solver_factory =
    solver::build()
        .with_criteria(
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor),
            gko::stop::Iteration::build().with_max_iters(max_iters)
        )
        .with_preconditioner(preconditioner::create(exec))
        .on(exec);
```

Declare the output file for all our loggers

```
std::ofstream output_file(of_name);
```

Do a warmup run

```
{
```

Clone x to not overwrite the original one

```
auto x_clone = gko::clone(x);
```

Generate and call apply on a solver

```
    solver_factory->generate(A)->apply(b, x_clone);
    exec->synchronize();
}
```

Do a timed run

```
{
```

Clone x to not overwrite the original one

```
auto x_clone = gko::clone(x);
```

Synchronize ensures no operation are ongoing

```
exec->synchronize();
```

Time before generate

```
auto g_tic = std::chrono::steady_clock::now();
```

Generate a solver

```
auto generated_solver = solver_factory->generate(A);
exec->synchronize();
```

Time after generate

```
auto g_tac = std::chrono::steady_clock::now();
```

Compute the generation time

```
auto generate_time =
    std::chrono::duration_cast<std::chrono::nanoseconds>(g_tac - g_tic);
```

Write the generate time to the output file

```
output_file << "Generate time (ns): " << generate_time.count()
            << std::endl;
```

Similarly time the apply

```
exec->synchronize();
auto a_tic = std::chrono::steady_clock::now();
generated_solver->apply(b, x_clone);
exec->synchronize();
auto a_tac = std::chrono::steady_clock::now();
auto apply_time =
    std::chrono::duration_cast<std::chrono::nanoseconds>(a_tac - a_tic);
output_file << "Apply time (ns): " << apply_time.count() << std::endl;
```

Compute the residual norm

```
auto residual =
    utils::compute_residual_norm(A.get(), b.get(), x_clone.get());
output_file << "Residual_norm: " << residual << std::endl;
}
```

Log the internal operations using the OperationLogger without timing

```
{
```

Create an OperationLogger to analyze the generate step

```
auto gen_logger = std::make_shared<loggers::OperationLogger>();
```

Add the generate logger to the executor

```
exec->add_logger(gen_logger);
```

Generate a solver

```
auto generated_solver = solver_factory->generate(A);
```

Remove the generate logger from the executor

```
exec->remove_logger(gen_logger);
```

Write the data to the output file

```
output_file << "Generate operations times (ns):" << std::endl;
gen_logger->write_data(output_file);
```

Create an OperationLogger to analyze the apply step

```
auto apply_logger = std::make_shared<loggers::OperationLogger>();
exec->add_logger(apply_logger);
```

Create a ResidualLogger to log the recurrent residual

```
auto res_logger = std::make_shared<loggers::ResidualLogger<ValueType>>(
    A.get(), b.get());
generated_solver->add_logger(res_logger);
```

Solve the system

```
generated_solver->apply(b, x);
exec->remove_logger(apply_logger);
```

Write the data to the output file

```
output_file << "Apply operations times (ns):" << std::endl;
apply_logger->write_data(output_file);
res_logger->write_data(output_file);
}
```

Print solution

```
std::cout << "Solution, first ten entries: \n";
print_vector(x.get());
```

Print output file location

```
std::cout << "The performance and residual data can be found in " << of_name
<< std::endl;
}
```

Results

This is the expected standard output:

```
Solution, first ten entries:
[
  0.252218
  0.108645
  0.0662811
  0.0630433
  0.0384088
  0.0396536
  0.0402648
  0.0338935
  0.0193098
  0.0234653
];
The performance and residual data can be found in log.txt
```

Here is a sample output in the file log.txt:

```
Generate time (ns): 861
Apply time (ns): 108144
Residual_norm: 2.10788e-15
Generate operations times (ns):
Apply operations times (ns):
  allocate: 14991
  cg::initialize#8: 872
  cg::step_1#5: 7683
  cg::step_2#7: 7756
  copy: 7751
  csr::advanced_spmv#5: 21819
  csr::spmv#3: 20429
  dense::compute_dot#3: 18043
  dense::compute_norm2#2: 16726
  free: 8857
  residual_norm::residual_norm#9: 3614
Recurrent Residual Norms:
[
  4.3589
  2.30455
  1.46771
  0.984875
  0.741833
  0.513623
  0.384165
  0.316439
  0.227709
  0.170312
  0.0973722
  0.0616831
  0.0454123
  0.031953
  0.0161606
  0.00657015
  0.00264367
  0.000858809
  0.000286461
  1.64195e-15
];
```

```

True Residual Norms:
[
    4.3589
    2.30455
    1.46771
    0.984875
    0.741833
    0.513623
    0.384165
    0.316439
    0.227709
    0.170312
    0.0973722
    0.0616831
    0.0454123
    0.031953
    0.0161606
    0.00657015
    0.00264367
    0.000858809
    0.000286461
    2.10788e-15
];

```

Comments about programming and debugging

The plain program

```

#include <algorithm>
#include <array>
#include <chrono>
#include <cstdlib>
#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>
#include <ostream>
#include <sstream>
#include <string>
#include <unordered_map>
#include <utility>
#include <vector>
#include <ginkgo/ginkgo.hpp>
template <typename ValueType>
using vec = gko::matrix::Dense<ValueType>;
template <typename ValueType>
using real_vec = gko::matrix::Dense<gko::remove_complex<ValueType>>;
namespace utils {
template <typename ValueType>
std::unique_ptr<vec<ValueType>> create_vector(
    std::shared_ptr<const gko::Executor> exec, gko::size_type size,
    ValueType value)
{
    auto res = vec<ValueType>::create(exec);
    res->read(gko::matrix_data<ValueType>(gko::dim<2>{size, 1}, value));
    return res;
}
template <typename ValueType>
ValueType get_first_element(const vec<ValueType>* norm)
{
    return norm->get_executor()->copy_val_to_host(norm->get_const_values());
}
template <typename ValueType>
gko::remove_complex<ValueType> compute_norm(const vec<ValueType>* b)
{
    auto exec = b->get_executor();
    auto b_norm = gko::initialize<real_vec<ValueType>>({0.0}, exec);
    b->compute_norm2(b_norm);
    return get_first_element(b_norm.get());
}
template <typename ValueType>
gko::remove_complex<ValueType> compute_residual_norm(
    const gko::LinOp* system_matrix, const vec<ValueType>* b,
    const vec<ValueType>* x)
{
    auto exec = system_matrix->get_executor();
    auto one = gko::initialize<vec<ValueType>>({1.0}, exec);
    auto neg_one = gko::initialize<vec<ValueType>>({-1.0}, exec);
    auto res = gko::clone(b);
    system_matrix->apply(one, x, neg_one, res);
    return compute_norm(res.get());
}

```



```

}
} // namespace utils
namespace loggers {
struct OperationLogger : gko::log::Logger {
    void on_allocation_started(const gko::Executor* exec,
                             const gko::size_type&)const override
    {
        this->start_operation(exec, "allocate");
    }
    void on_allocation_completed(const gko::Executor* exec,
                                const gko::size_type&,
                                const gko::uintptr&)const override
    {
        this->end_operation(exec, "allocate");
    }
    void on_free_started(const gko::Executor* exec,
                        const gko::uintptr&)const override
    {
        this->start_operation(exec, "free");
    }
    void on_free_completed(const gko::Executor* exec,
                           const gko::uintptr&)const override
    {
        this->end_operation(exec, "free");
    }
    void on_copy_started(const gko::Executor* from, const gko::Executor* to,
                        const gko::uintptr&, const gko::uintptr&,
                        const gko::size_type&)const override
    {
        from->synchronize();
        this->start_operation(to, "copy");
    }
    void on_copy_completed(const gko::Executor* from, const gko::Executor* to,
                           const gko::uintptr&, const gko::uintptr&,
                           const gko::size_type&)const override
    {
        from->synchronize();
        this->end_operation(to, "copy");
    }
    void on_operation_launched(const gko::Executor* exec,
                               const gko::Operation* op)const override
    {
        this->start_operation(exec, op->get_name());
    }
    void on_operation_completed(const gko::Executor* exec,
                               const gko::Operation* op)const override
    {
        this->end_operation(exec, op->get_name());
    }
    void write_data(std::ostream& ostream)
    {
        for (const auto& entry : total) {
            ostream << "\t" << entry.first.c_str() << ": "
                << std::chrono::duration_cast<std::chrono::nanoseconds>(
                    entry.second)
                    .count()
                << std::endl;
        }
    }
private:
    void start_operation(const gko::Executor* exec,
                        const std::string& name)const
    {
        nested.emplace_back(0);
        exec->synchronize();
        start[name] = std::chrono::steady_clock::now();
    }
    void end_operation(const gko::Executor* exec, const std::string& name)const
    {
        exec->synchronize();
        const auto end = std::chrono::steady_clock::now();
        const auto diff = end - start[name];
        total[name] += diff - nested.back();
        nested.pop_back();
        if (nested.size() > 0) {
            nested.back() += diff;
        }
    }
    mutable std::map<std::string, std::chrono::steady_clock::time_point> start;
    mutable std::map<std::string, std::chrono::steady_clock::duration> total;
    mutable std::vector<std::chrono::steady_clock::duration> nested;
};
struct StorageLogger : gko::log::Logger {
    void on_allocation_completed(const gko::Executor*,
                                const gko::size_type& num_bytes,
                                const gko::uintptr& location)const override
    {

```

```

        storage[location] = num_bytes;
    }
    void on_free_completed(const gko::Executor*,
                          const gko::uintptr& location) const override
    {
        storage[location] = 0;
    }
    void write_data(std::ostream& ostream)
    {
        gko::size_type total{};
        for (const auto& e : storage) {
            total += e.second;
        }
        ostream << "Storage: " << total << std::endl;
    }
private:
    mutable std::unordered_map<gko::uintptr, gko::size_type> storage;
};

template <typename ValueType>
struct ResidualLogger : gko::log::Logger {
    void on_iteration_complete(const gko::LinOp*, const gko::size_type&,
                             const gko::LinOp* residual,
                             const gko::LinOp* solution,
                             const gko::LinOp* residual_norm) const override
    {
        if (residual_norm) {
            rec_res_norms.push_back(utils::get_first_element(
                gko::as<real_vec<ValueType>>(residual_norm)));
        } else {
            rec_res_norms.push_back(
                utils::compute_norm(gko::as<vec<ValueType>>(residual)));
        }
        if (solution) {
            true_res_norms.push_back(utils::compute_residual_norm(
                matrix, b, gko::as<vec<ValueType>>(solution)));
        } else {
            true_res_norms.push_back(-1.0);
        }
    }
};

ResidualLogger(const gko::LinOp* matrix, const vec<ValueType>* b)
    : gko::log::Logger(gko::log::Logger::iteration_complete_mask),
      matrix{matrix},
      b{b}
{}

void write_data(std::ostream& ostream)
{
    ostream << "Recurrent Residual Norms: " << std::endl;
    ostream << "[" << std::endl;
    for (const auto& entry : rec_res_norms) {
        ostream << "\t" << entry << std::endl;
    }
    ostream << "]" << std::endl;
    ostream << "True Residual Norms: " << std::endl;
    ostream << "[" << std::endl;
    for (const auto& entry : true_res_norms) {
        ostream << "\t" << entry << std::endl;
    }
    ostream << "]" << std::endl;
}

private:
    const gko::LinOp* matrix;
    const vec<ValueType>* b;
    mutable std::vector<gko::remove_complex<ValueType>> rec_res_norms;
    mutable std::vector<gko::remove_complex<ValueType>> true_res_norms;
};

// namespace loggers
namespace {
void print_usage(const char* filename)
{
    std::cerr << "Usage: " << filename << " [executor] [matrix file]"
              << std::endl;
    std::cerr << "matrix file should be a file in matrix market format. "
              << "The file data/A.mtx is provided as an example."
              << std::endl;
    std::exit(-1);
}

template <typename ValueType>
void print_vector(const gko::matrix::Dense<ValueType>* vec)
{
    auto elements_to_print = std::min(gko::size_type(10), vec->get_size()[0]);
    std::cout << "[" << std::endl;
    for (int i = 0; i < elements_to_print; ++i) {
        std::cout << "\t" << vec->at(i) << std::endl;
    }
    std::cout << "]" << std::endl;
}

// namespace

```

```

int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using solver = gko::solver::Cg<ValueType>;
    using preconditioner = gko::matrix::IdentityFactory<ValueType>;
    const gko::remove_complex<ValueType> reduction_factor = 1e-12;
    const auto of_name = "log.txt";
    using vec = gko::matrix::Dense<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    std::string input_mtx = "data/A.mtx";
    if (argc == 3) {
        input_mtx = std::string(argv[2]);
    }
    auto storage_logger = std::make_shared<loggers::StorageLogger>();
    exec->add_logger(storage_logger);
    auto A = gko::share(gko::read<mtx>(std::ifstream(input_mtx), exec));
    exec->remove_logger(storage_logger);
    const auto max_iters = A->get_size()[0];
    auto b = utils::create_vector<ValueType>(exec, A->get_size()[0], 1.0);
    auto x = utils::create_vector<ValueType>(exec, A->get_size()[0], 0.0);
    auto solver_factory =
        solver::build()
            .with_criteria(
                gko::stop::ResidualNorm<ValueType>::build()
                    .with_reduction_factor(reduction_factor),
                gko::stop::Iteration::build().with_max_iters(max_iters))
            .with_preconditioner(preconditioner::create(exec))
            .on(exec);
    std::ofstream output_file(of_name);
    {
        auto x_clone = gko::clone(x);
        solver_factory->generate(A)->apply(b, x_clone);
        exec->synchronize();
    }
    {
        auto x_clone = gko::clone(x);
        exec->synchronize();
        auto g_tic = std::chrono::steady_clock::now();
        auto generated_solver = solver_factory->generate(A);
        exec->synchronize();
        auto g_tac = std::chrono::steady_clock::now();
        auto generate_time =
            std::chrono::duration_cast<std::chrono::nanoseconds>(g_tac - g_tic);
        output_file << "Generate time (ns): " << generate_time.count()
            << std::endl;
        exec->synchronize();
        auto a_tic = std::chrono::steady_clock::now();
        generated_solver->apply(b, x_clone);
        exec->synchronize();
        auto a_tac = std::chrono::steady_clock::now();
        auto apply_time =
            std::chrono::duration_cast<std::chrono::nanoseconds>(a_tac - a_tic);
        output_file << "Apply time (ns): " << apply_time.count() << std::endl;
        auto residual =
            utils::compute_residual_norm(A.get(), b.get(), x_clone.get());
        output_file << "Residual_norm: " << residual << std::endl;
    }
    {
        auto gen_logger = std::make_shared<loggers::OperationLogger>();
        exec->add_logger(gen_logger);
        auto generated_solver = solver_factory->generate(A);
    }
}

```

```
exec->remove_logger(gen_logger);
output_file << "Generate operations times (ns):" << std::endl;
gen_logger->write_data(output_file);
auto apply_logger = std::make_shared<loggers::OperationLogger>();
exec->add_logger(apply_logger);
auto res_logger = std::make_shared<loggers::ResidualLogger<ValueType>>(
    A.get(), b.get());
generated_solver->add_logger(res_logger);
generated_solver->apply(b, x);
exec->remove_logger(apply_logger);
output_file << "Apply operations times (ns):" << std::endl;
apply_logger->write_data(output_file);
res_logger->write_data(output_file);
}
std::cout << "Solution, first ten entries: \n";
print_vector(x.get());
std::cout << "The performance and residual data can be found in " << of_name
           << std::endl;
}
```

Chapter 37

The poisson-solver program

The poisson solver example..

This example depends on simple-solver.

Introduction

This example solves a 1D Poisson equation:

$$\begin{aligned} u &: [0, 1] \rightarrow R \\ u'' &= f \\ u(0) &= u_0 \\ u(1) &= u_1 \end{aligned}$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

$$\begin{aligned} u(x+h) &= u(x) + u'(x)h + \frac{1}{2}u''(x)h^2 + O(h^3) \\ u(x-h) &= u(x) - u'(x)h + \frac{1}{2}u''(x)h^2 + O(h^3) \\ \hline -u(x-h) + 2u(x) - u(x+h) &= -f(x)h^2 + O(h^3) \end{aligned}$$

For an equidistant grid with K "inner" discretization points $x_1, \dots, x_K$, and step size $h = 1/(K+1)$, the formula produces a system of linear equations

$$\begin{aligned} 2u_1 - u_2 &= -f_1 h^2 + u_0 \\ -u_{k-1} + 2u_k - u_{k+1} &= -f_k h^2, k = 2, \dots, K-1 \\ -u_{K-1} + 2u_K &= -f_K h^2 + u_1 \end{aligned}$$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function `f` is set to `f(x) = 6x` (making the solution `u(x) = x3`), but that can be changed in the `main` function.

The intention of the example is to show how Ginkgo can be used to build an application solving a real-world problem, which includes a solution of a large, sparse linear system as a component.

About the example

The commented program

```
        row_ptrs[i + 1] = pos;
    }
}
```

Generates the RHS vector given f and the boundary conditions.

```
template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{
    const auto discretization_points = rhs->get_size()[0];
    auto values = rhs->get_values();
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
    for (gko::size_type i = 0; i < discretization_points; ++i) {
        const auto xi = static_cast<ValueType>(i + 1) * h;
        values[i] = -f(xi) * h * h;
    }
    values[0] += u0;
    values[discretization_points - 1] += u1;
}
```

Prints the solution u .

```
template <typename Closure, typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                   const gko::matrix::Dense<ValueType>* u)
{
    std::cout << u0 << '\n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout << u->get_const_values()[i] << '\n';
    }
    std::cout << u1 << std::endl;
}
```

Computes the 1-norm of the error given the computed u and the correct solution function `correct_u`.

```
template <typename Closure, typename ValueType>
gko::remove_complex<ValueType> calculate_error(
    int discretization_points, const gko::matrix::Dense<ValueType>* u,
    Closure correct_u)
{
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (int i = 0; i < discretization_points; ++i) {
        using std::abs;
        const auto xi = static_cast<ValueType>(i + 1) * h;
        error +=
            abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    }
    return error;
}

int main(int argc, char* argv[])
{

```

Some shortcuts

```
using ValueType = double;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
```

Print version information

```
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0]
               << " [executor] [DISCRETIZATION_POINTS]" << std::endl;
    std::exit(-1);
}
```

Get number of discretization points

```
const unsigned int discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;
```

Get the executor string

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
```

Figure out where to run the code

```
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string()); // throws if not valid
```

executor used by the application

```
const auto app_exec = exec->get_master();
```

Set up the problem: define the exact solution, the right hand side and the Dirichlet boundary condition.

```
auto correct_u = [] (ValueType x) { return x * x * x; };
auto f = [] (ValueType x) { return ValueType(6) * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);
```

initialize matrix and vectors

```
auto matrix = mtx::create(app_exec, gko::dim<2>(discretization_points),
                          3 * discretization_points - 2);
generate_stencil_matrix(matrix.get());
auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
generate_rhs(f, u0, u1, rhs.get());
auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
for (int i = 0; i < u->get_size()[0]; ++i) {
    u->get_values()[i] = 0.0;
}
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
```

Generate solver and solve the system

```
cg::build()
    .with_criteria(
        gko::stop::Iteration::build().with_max_iters(discretization_points),
        gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
            reduction_factor))
    .with_preconditioner(bj::build())
    .on(exec)
    ->generate(clone(exec, matrix)) // copy the matrix to the executor
    ->apply(rhs, u);
```

Uncomment to print the solution print_solution<ValueType>(u0, u1, u.get());

```
std::cout << "Solve complete.\n";
std::cout << "The average relative error is "
            << calculate_error(discretization_points, u.get(), correct_u) /
            << static_cast<gko::remove_complex<ValueType>>(
                discretization_points)
            << std::endl;
}
```

Results

This is the expected output:

```
Solve complete.
The average relative error is 2.52236e-11
```

Comments about programming and debugging

The plain program

```
#include <iostream>
#include <map>
#include <string>
#include <vector>
#include <ginkgo/ginkgo.hpp>
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(gko::matrix::Csr<ValueType, IndexType>* matrix)
{
    const auto discretization_points = matrix->get_size()[0];
    auto row_ptrs = matrix->get_row_ptrs();
    auto col_idxs = matrix->get_col_idxs();
    auto values = matrix->get_values();
    int pos = 0;
    const ValueType coeffs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (int i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
                values[pos] = coeffs[ofs + 1];
                col_idxs[pos] = i + ofs;
                ++pos;
            }
        }
        row_ptrs[i + 1] = pos;
    }
}

template <typename Closure, typename ValueType>
void generate_rhs(Closure f, ValueType u0, ValueType u1,
                 gko::matrix::Dense<ValueType>* rhs)
{
    const auto discretization_points = rhs->get_size()[0];
    auto values = rhs->get_values();
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
    for (gko::size_type i = 0; i < discretization_points; ++i) {
        const auto xi = static_cast<ValueType>(i + 1) * h;
        values[i] = -f(xi) * h * h;
    }
    values[0] += u0;
    values[discretization_points - 1] += u1;
}

template <typename Closure, typename ValueType>
void print_solution(ValueType u0, ValueType u1,
                   const gko::matrix::Dense<ValueType>* u)
{
    std::cout << u0 << '\n';
    for (int i = 0; i < u->get_size()[0]; ++i) {
        std::cout << u->get_const_values()[i] << '\n';
    }
    std::cout << u1 << std::endl;
}

template <typename Closure, typename ValueType>
gko::remove_complex<ValueType> calculate_error(
    int discretization_points, const gko::matrix::Dense<ValueType>* u,
    Closure correct_u)
{
    const ValueType h = 1.0 / static_cast<ValueType>(discretization_points + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (int i = 0; i < discretization_points; ++i) {
        using std::abs;
        const auto xi = static_cast<ValueType>(i + 1) * h;
        error +=
            abs(u->get_const_values()[i] - correct_u(xi)) / abs(correct_u(xi));
    }
    return error;
}

int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0]
                  << " [executor] [DISCRETIZATION_POINTS]" << std::endl;
        std::exit(-1);
    }
    const unsigned int discretization_points =
```



```

    argc >= 3 ? std::atoi(argv[2]) : 100;
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
    auto correct_u = [] (ValueType x) { return x * x * x; };
    auto f = [] (ValueType x) { return ValueType(6) * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);
    auto matrix = mtx::create(app_exec, gko::dim<2>(discretization_points),
                              3 * discretization_points - 2);
    generate_stencil_matrix(matrix.get());
    auto rhs = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
    generate_rhs(f, u0, u1, rhs.get());
    auto u = vec::create(app_exec, gko::dim<2>(discretization_points, 1));
    for (int i = 0; i < u->get_size()[0]; ++i) {
        u->get_values()[i] = 0.0;
    }
    const gko::remove_complex<ValueType> reduction_factor = 1e-7;
    cg::build()
        .with_criteria(
            gko::stop::Iteration::build().with_max_iters(discretization_points),
            gko::stop::ResidualNorm<ValueType>::build().with_reduction_factor(
                reduction_factor))
        .with_preconditioner(bj::build())
        .on(exec)
        ->generate(clone(exec, matrix)) // copy the matrix to the executor
        ->apply(rhs, u);
    std::cout << "Solve complete.\n";
    std::cout << "The average relative error is "
        << calculate_error(discretization_points, u.get(), correct_u) /
            static_cast<gko::remove_complex<ValueType>>(
                discretization_points)
        << std::endl;
}

```


Chapter 38

The preconditioned-solver program

The preconditioned solver example..

This example depends on simple-solver.

Introduction

About the example

The commented program

```
}
```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); } }
};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
```

Create solver factory

```
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                       gko::stop::ResidualNorm<ValueType>::build())
```

```
.with_reduction_factor(reduction_factor))
```

Add preconditioner, these 2 lines are the only difference from the simple solver example

```
.with_preconditioner(bj::build().with_max_block_size(8u))
.on(exec);
```

Create solver

```
auto solver = solver_gen->generate(A);
```

Solve system

```
solver->apply(b, x);
```

Print solution

```
std::cout << "Solution (x):\n";
write(std::cout, x);
```

Calculate residual

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```

Results

This is the expected output:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.82005e-08
```

Comments about programming and debugging

The plain program

```
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
```

```

using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                       gko::stop::ResidualNorm<ValueType>::build()
                           .with_reduction_factor(reduction_factor))
        .with_preconditioner(bj::build().with_max_block_size(8u))
        .on(exec);
auto solver = solver_gen->generate(A);
solver->apply(b, x);
std::cout << "Solution (x):\n";
write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}

```


Chapter 39

The preconditioner-export program

The preconditioner export example..

This example depends on simple-solver.

Introduction

About the example

This example shows how to explicitly generate and store preconditioners for a given matrix. It can also be used to inspect and debug the preconditioner generation.

The commented program

```
void output(gko::ptr_param<const gko::WritableToMatrixData<double, int>> mtx,
            std::string name)
{
    std::ofstream stream(name);
    std::cerr << "Writing " << name << std::endl;
    gko::write(stream, mtx, gko::layout_type::coordinate);
}
template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
{
    decltype(fun()) result;
    try {
        result = fun();
    } catch (const gko::Error& err) {
        std::cerr << "Error: " << err.what() << '\n';
        std::exit(-1);
    }
    return result;
}
int main(int argc, char* argv[])
{
    print usage message
    if (argc < 2 || executors.find(argv[1]) == executors.end()) {
        std::cerr << "Usage: executable"
            << " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
            << "[<jacobi|ilu|parilu|parilut|ilu-isai|parilu-isai|parilut-"
            << "isai> [<preconditioner args>]\n";
        std::cerr << "Jacobi parameters:  [<max-block-size>] [<accuracy>] "
            << "[<storage-optimization:auto|0|1|2>]\n";
        std::cerr << "ParILU parameters:  [<iteration-count>]\n";
        std::cerr
            << "ParILUT parameters:  [<iteration-count>] [<fill-in-limit>]\n";
    }
```

```

std::cerr << "ILU-ISAI parameters:  [<sparsity-power>]\n";
std::cerr << "ParILU-ISAI parameters:  [<iteration-count>] "
            "<sparsity-power>]\n";
std::cerr << "ParILUT-ISAI parameters:  [<iteration-count>] "
            "<fill-in-limit>] [<sparsity-power>]\n";
return -1;
}

```

generate executor based on first argument

```
auto exec = try_generate([&] { return executors.at(argv[1])(); });
```

set matrix and preconditioner name with default values

```
std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
std::string precondition = argc < 4 ? "jacobi" : argv[3];
```

load matrix file into Csr format

```

auto mtx = gko::share(try_generate([&] {
    std::ifstream mtx_stream(matrix);
    if (!mtx_stream) {
        throw GKO_STREAM_ERROR("Unable to open matrix file");
    }
    std::cerr << "Reading " << matrix << std::endl;
    return gko::read<gko::matrix::Csr>(mtx_stream, exec);
}));

```

concatenate remaining arguments for filename

```

std::string output_suffix;
for (auto i = 4; i < argc; ++i) {
    output_suffix = output_suffix + "-" + argv[i];
}

```

handle different preconditioners

```
if (precond == "jacobi") {
```

jacobi: max_block_size, accuracy, storage_optimization

```

auto factory_parameter = gko::preconditioner::Jacobi<>::build();
if (argc >= 5) {
    factory_parameter.with_max_block_size(
        static_cast<gko::uint32>(std::stoi(argv[4])));
}
if (argc >= 6) {
    factory_parameter.with_accuracy(std::stod(argv[5]));
}
if (argc >= 7) {
    factory_parameter.with_storage_optimization(
        std::string{argv[6]} == "auto"
        ? gko::precision_reduction::autodetect()
        : gko::precision_reduction(0, std::stoi(argv[6])));
}
auto factory = factory_parameter.on(exec);
auto jacobi = try_generate([&] { return factory->generate(mtx); });
output(jacobi, matrix + ".jacobi" + output_suffix);
} else if (precond == "ilu") {

```

ilu: no parameters

```

auto ilu = gko::as<gko::Composition>(try_generate([&] {
    return gko::factorization::Ilu<>::build().on(exec)->generate(mtx);
}));
output(gko::as<gko::matrix::Csr>(ilu->get_operators()[0]),
    matrix + ".ilu-l");
output(gko::as<gko::matrix::Csr>(ilu->get_operators()[1]),
    matrix + ".ilu-u");
} else if (precond == "parilu") {

```

parilu: iterations

```

auto factory_parameter = gko::factorization::ParIlu<>::build();
if (argc >= 5) {
    factory_parameter.with_iterations(
        static_cast<gko::size_type>(std::stoi(argv[4])));
}
auto factory = factory_parameter.on(exec);
auto ilu = gko::as<gko::Composition>(
    try_generate([&] { return factory->generate(mtx); }));
output(gko::as<gko::matrix::Csr>(ilu->get_operators()[0]),
    matrix + ".parilu" + output_suffix + "-l");
output(gko::as<gko::matrix::Csr>(ilu->get_operators()[1]),
    matrix + ".parilu" + output_suffix + "-u");
} else if (precond == "parilut") {

```


parilut: iterations, fill-in limit

```

auto factory_parameter = gko::factorization::ParIlut<>::build();
if (argc >= 5) {
    factory_parameter.with_iterations(
        static_cast<gko::size_type>(std::stoi(argv[4])));
}
if (argc >= 6) {
    factory_parameter.with_fill_in_limit(std::stod(argv[5]));
}
auto factory = factory_parameter.on(exec);
auto ilut = gko::as<gko::Composition<>>(
    try_generate([&] { return factory->generate(mtx); }));
output(gko::as<gko::matrix::Csr<>>(ilut->get_operators()[0]),
    matrix + ".parilut" + output_suffix + "-l");
output(gko::as<gko::matrix::Csr<>>(ilut->get_operators()[1]),
    matrix + ".parilut" + output_suffix + "-u");
} else if (precond == "ilu-isai") {

```

ilu-isai: sparsity power

```

auto fact_factory =
    gko::share(gko::factorization::Ilu<>::build().on(exec));
int sparsity_power = 1;
if (argc >= 5) {
    sparsity_power = std::stoi(argv[4]);
}
auto factory =
    gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
        gko::preconditioner::UpperIsai<>>::build()
        .with_factorization(fact_factory)
        .with_l_solver(gko::preconditioner::LowerIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .with_u_solver(gko::preconditioner::UpperIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .on(exec);
auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
output(ilu_isai->get_l_solver()->get_approximate_inverse(),
    matrix + ".ilu-isai" + output_suffix + "-l");
output(ilu_isai->get_u_solver()->get_approximate_inverse(),
    matrix + ".ilu-isai" + output_suffix + "-u");
} else if (precond == "parilu-isai") {

```

parilu-isai: iterations, sparsity power

```

auto fact_parameter = gko::factorization::ParIlu<>::build();
int sparsity_power = 1;
if (argc >= 5) {
    fact_parameter.with_iterations(
        static_cast<gko::size_type>(std::stoi(argv[4])));
}
if (argc >= 6) {
    sparsity_power = std::stoi(argv[5]);
}
auto fact_factory = gko::share(fact_parameter.on(exec));
auto factory =
    gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
        gko::preconditioner::UpperIsai<>>::build()
        .with_factorization(fact_factory)
        .with_l_solver(gko::preconditioner::LowerIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .with_u_solver(gko::preconditioner::UpperIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .on(exec);
auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
output(ilu_isai->get_l_solver()->get_approximate_inverse(),
    matrix + ".parilu-isai" + output_suffix + "-l");
output(ilu_isai->get_u_solver()->get_approximate_inverse(),
    matrix + ".parilu-isai" + output_suffix + "-u");
} else if (precond == "parilut-isai") {

```

parilut-isai: iterations, fill-in limit, sparsity power

```

auto fact_parameter = gko::factorization::ParIlut<>::build();
int sparsity_power = 1;
if (argc >= 5) {
    fact_parameter.with_iterations(
        static_cast<gko::size_type>(std::stoi(argv[4])));
}
if (argc >= 6) {
    fact_parameter.with_fill_in_limit(std::stod(argv[5]));
}
if (argc >= 7) {
    sparsity_power = std::stoi(argv[6]);
}
auto fact_factory = gko::share(fact_parameter.on(exec));
auto factory =
    gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,

```

```

        gko::preconditioner::UpperIsai<>>::build()
        .with_factorization(fact_factory)
        .with_l_solver(gko::preconditioner::LowerIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .with_u_solver(gko::preconditioner::UpperIsai<>::build()
            .with_sparsity_power(sparsity_power))
        .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
        matrix + ".parilut-isai" + output_suffix + "-l");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
        matrix + ".parilut-isai" + output_suffix + "-u");
    }
}

```

Results

This is the expected output:

```

Usage: ./preconditioner-export <reference|omp|cuda|hip|dpcpp> [<matrix-file>]
        [<jacobi|ilu|parilu|parilut|ilu-isai|parilu-isai|parilut-isai>] [<preconditioner args>]
Jacobi parameters:  [<max-block-size>] [<accuracy>] [<storage-optimization:auto|0|1|2>]
ParILU parameters:  [<iteration-count>]
ParILUT parameters: [<iteration-count>] [<fill-in-limit>]
ILU-ISAI parameters: [<sparsity-power>]
ParILU-ISAI parameters: [<iteration-count>] [<sparsity-power>]
ParILUT-ISAI parameters: [<iteration-count>] [<fill-in-limit>] [<sparsity-power>]

```

When specifying an executor:

```

Reading data/A.mtx
Writing data/A.mtx.jacobi

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <functional>
#include <iostream>
#include <map>
#include <memory>
#include <string>
#include <ginkgo/ginkgo.hpp>
const std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    executors{{"reference", [] { return gko::ReferenceExecutor::create(); }},
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
            [] {
                return gko::CudaExecutor::create(
                    0, gko::ReferenceExecutor::create());
            }},
        {"hip",
            [] {
                return gko::HipExecutor::create(
                    0, gko::ReferenceExecutor::create());
            }},
        {"dpcpp", [] {
            return gko::DpcppExecutor::create(
                0, gko::ReferenceExecutor::create());
        }}};

void output(gko::ptr_param<const gko::WritableToMatrixData<double, int>> mtx,
    std::string name)
{
    std::ofstream stream{name};
    std::cerr << "Writing " << name << std::endl;
    gko::write(stream, mtx, gko::layout_type::coordinate);
}

template <typename Function>
auto try_generate(Function fun) -> decltype(fun())
{
    decltype(fun()) result;
    try {
        result = fun();
    } catch (const gko::Error& err) {
        std::cerr << "Error: " << err.what() << '\n';
    }
}

```

```

        std::exit(-1);
    }
    return result;
}
int main(int argc, char* argv[])
{
    if (argc < 2 || executors.find(argv[1]) == executors.end()) {
        std::cerr << "Usage: executable"
            << " <reference|omp|cuda|hip|dpcpp> [<matrix-file>] "
            << "[<jacobi|ilu|parilu|parilut|ilu-isai|parilu-isai|parilut-|"
            << "isai] [<preconditioner args>]\n";
        std::cerr << "Jacobi parameters:  [<max-block-size>] [<accuracy>] "
            << "[<storage-optimization:auto|0|1|2>]\n";
        std::cerr << "ParILU parameters:  [<iteration-count>]\n";
        std::cerr
            << "ParILUT parameters:  [<iteration-count>] [<fill-in-limit>]\n";
        std::cerr << "ILU-ISAI parameters:  [<sparsity-power>]\n";
        std::cerr << "ParILU-ISAI parameters:  [<iteration-count>] "
            << "[<sparsity-power>]\n";
        std::cerr << "ParILUT-ISAI parameters:  [<iteration-count>] "
            << "[<fill-in-limit>] [<sparsity-power>]\n";
        return -1;
    }
    auto exec = try_generate([&] { return executors.at(argv[1])(); });
    std::string matrix = argc < 3 ? "data/A.mtx" : argv[2];
    std::string precondition = argc < 4 ? "jacobi" : argv[3];
    auto mtx = gko::share(try_generate([&] {
        std::ifstream mtx_stream(matrix);
        if (!mtx_stream) {
            throw GKO_STREAM_ERROR("Unable to open matrix file");
        }
        std::cerr << "Reading " << matrix << std::endl;
        return gko::read<gko::matrix::Csr>(mtx_stream, exec);
    }));
    std::string output_suffix;
    for (auto i = 4; i < argc; ++i) {
        output_suffix = output_suffix + "-" + argv[i];
    }
    if (precondition == "jacobi") {
        auto factory_parameter = gko::preconditioner::Jacobi<>::build();
        if (argc >= 5) {
            factory_parameter.with_max_block_size(
                static_cast<gko::uint32>(std::stoi(argv[4])));
        }
        if (argc >= 6) {
            factory_parameter.with_accuracy(std::stod(argv[5]));
        }
        if (argc >= 7) {
            factory_parameter.with_storage_optimization(
                std::string(argv[6]) == "auto"
                ? gko::precision_reduction::autodetect()
                : gko::precision_reduction(0, std::stoi(argv[6])));
        }
        auto factory = factory_parameter.on(exec);
        auto jacobi = try_generate([&] { return factory->generate(mtx); });
        output(jacobi, matrix + ".jacobi" + output_suffix);
    } else if (precondition == "ilu") {
        auto ilu = gko::as<gko::Composition>(try_generate([&] {
            return gko::factorization::Ilut<>::build().on(exec)->generate(mtx);
        }));
        output(gko::as<gko::matrix::Csr>(ilu->get_operators()[0]),
            matrix + ".ilu-l");
        output(gko::as<gko::matrix::Csr>(ilu->get_operators()[1]),
            matrix + ".ilu-u");
    } else if (precondition == "parilu") {
        auto factory_parameter = gko::factorization::ParIlut<>::build();
        if (argc >= 5) {
            factory_parameter.with_iterations(
                static_cast<gko::size_type>(std::stoi(argv[4])));
        }
        auto factory = factory_parameter.on(exec);
        auto ilu = gko::as<gko::Composition>(
            try_generate([&] { return factory->generate(mtx); }));
        output(gko::as<gko::matrix::Csr>(ilu->get_operators()[0]),
            matrix + ".parilu" + output_suffix + "-l");
        output(gko::as<gko::matrix::Csr>(ilu->get_operators()[1]),
            matrix + ".parilu" + output_suffix + "-u");
    } else if (precondition == "parilut") {
        auto factory_parameter = gko::factorization::ParIlut<>::build();
        if (argc >= 5) {
            factory_parameter.with_iterations(
                static_cast<gko::size_type>(std::stoi(argv[4])));
        }
        if (argc >= 6) {
            factory_parameter.with_fill_in_limit(std::stod(argv[5]));
        }
        auto factory = factory_parameter.on(exec);
    }

```

```

auto ilut = gko::as<gko::Composition>({
    try_generate([&] { return factory->generate(mtx); });
    output(gko::as<gko::matrix::Csr<>>(ilut->get_operators()[0]),
        matrix + ".parilut" + output_suffix + "-l");
    output(gko::as<gko::matrix::Csr<>>(ilut->get_operators()[1]),
        matrix + ".parilut" + output_suffix + "-u");
}) else if (precond == "ilu-isai") {
    auto fact_factory =
        gko::share(gko::factorization::Ilu<>::build().on(exec));
    int sparsity_power = 1;
    if (argc >= 5) {
        sparsity_power = std::stoi(argv[4]);
    }
    auto factory =
        gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
            gko::preconditioner::UpperIsai<>>::build()
            .with_factorization(fact_factory)
            .with_l_solver(gko::preconditioner::LowerIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .with_u_solver(gko::preconditioner::UpperIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
        matrix + ".ilu-isai" + output_suffix + "-l");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
        matrix + ".ilu-isai" + output_suffix + "-u");
} else if (precond == "parilu-isai") {
    auto fact_parameter = gko::factorization::ParIlu<>::build();
    int sparsity_power = 1;
    if (argc >= 5) {
        fact_parameter.with_iterations(
            static_cast<gko::size_type>(std::stoi(argv[4])));
    }
    if (argc >= 6) {
        sparsity_power = std::stoi(argv[5]);
    }
    auto fact_factory = gko::share(fact_parameter.on(exec));
    auto factory =
        gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
            gko::preconditioner::UpperIsai<>>::build()
            .with_factorization(fact_factory)
            .with_l_solver(gko::preconditioner::LowerIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .with_u_solver(gko::preconditioner::UpperIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
        matrix + ".parilu-isai" + output_suffix + "-l");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
        matrix + ".parilu-isai" + output_suffix + "-u");
} else if (precond == "parilut-isai") {
    auto fact_parameter = gko::factorization::ParIlut<>::build();
    int sparsity_power = 1;
    if (argc >= 5) {
        fact_parameter.with_iterations(
            static_cast<gko::size_type>(std::stoi(argv[4])));
    }
    if (argc >= 6) {
        fact_parameter.with_fill_in_limit(std::stod(argv[5]));
    }
    if (argc >= 7) {
        sparsity_power = std::stoi(argv[6]);
    }
    auto fact_factory = gko::share(fact_parameter.on(exec));
    auto factory =
        gko::preconditioner::Ilu<gko::preconditioner::LowerIsai<>,
            gko::preconditioner::UpperIsai<>>::build()
            .with_factorization(fact_factory)
            .with_l_solver(gko::preconditioner::LowerIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .with_u_solver(gko::preconditioner::UpperIsai<>::build()
                .with_sparsity_power(sparsity_power))
            .on(exec);
    auto ilu_isai = try_generate([&] { return factory->generate(mtx); });
    output(ilu_isai->get_l_solver()->get_approximate_inverse(),
        matrix + ".parilut-isai" + output_suffix + "-l");
    output(ilu_isai->get_u_solver()->get_approximate_inverse(),
        matrix + ".parilut-isai" + output_suffix + "-u");
}
}

```

Chapter 40

The reordered-preconditioned-solver program

The reordered preconditioned solver example..

This example depends on preconditioned-solver.

Introduction

About the example

This uses an RCM reordering on an input matrix to construct a ParILU preconditioned solver and solve a linear system.

The commented program

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>()>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
```

```

auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
auto reordering =
    gko::experimental::reorder::Rcm<IndexType>::build().on(exec)->generate(
        A);

```

Permute matrix and vectors

```

auto A_reordered = share(A->permute(reordering));
auto b_reordered = b->permute(reordering, gko::matrix::permute_mode::rows);

```

this reordering is not necessary, but it maps the initial guess to the unreordered case

```

auto x_reordered = x->permute(reordering, gko::matrix::permute_mode::rows);
const RealValueType reduction_factor{1e-7};

```

Create solver factory

```

auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
            gko::stop::ResidualNorm<ValueType>::build()
                .with_reduction_factor(reduction_factor))
        .with_preconditioner(ilu::build())
        .on(exec);

```

Create solver

```

auto solver = solver_gen->generate(A_reordered);

```

Solve system

```

solver->apply(b_reordered, x_reordered);

```

Revert permutation to get the unpermuted solution

```

x_reordered->permute(reordering, x,
    gko::matrix::permute_mode::inverse_rows);

```

Print solution

```

std::cout << "Solution (x):\n";
write(std::cout, x);

```

Calculate residual

```

auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}

```

Results

This is the expected output:

```

Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
4.82005e-08

```

The plain program

```
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <ginkgo/ginkgo.hpp>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using ilu =
        gko::preconditioner::ilu<gko::solver::LowerTrs<ValueType, IndexType>,
                                gko::solver::UpperTrs<ValueType, IndexType>,
                                false, IndexType>;

    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    auto reordering =
        gko::experimental::reorder::Rcm<IndexType>::build().on(exec)->generate(
        A);
    auto A_reordered = share(A->permute(reordering));
    auto b_reordered = b->permute(reordering, gko::matrix::permute_mode::rows);
    auto x_reordered = x->permute(reordering, gko::matrix::permute_mode::rows);
    const RealValueType reduction_factor{1e-7};
    auto solver_gen =
        cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                      .with_reduction_factor(reduction_factor))
        .with_preconditioner(ilu::build())
        .on(exec);
    auto solver = solver_gen->generate(A_reordered);
    solver->apply(b_reordered, x_reordered);
    x_reordered->permute(reordering, x,
                       gko::matrix::permute_mode::inverse_rows);
    std::cout << "Solution (x):\n";
    write(std::cout, x);
    auto one = gko::initialize<vec>({1.0}, exec);
    auto neg_one = gko::initialize<vec>({-1.0}, exec);
    auto res = gko::initialize<real_vec>({0.0}, exec);
    A->apply(one, x, neg_one, b);
    b->compute_norm2(res);
    std::cout << "Residual norm sqrt(r^T r):\n";
    write(std::cout, res);
}
```


Chapter 41

The schroedinger-splitting program

The Schroedinger equation example..

This example depends on heat-equation.

Introduction

This example shows how to use the FFT and iFFT implementations in Ginkgo to solve the non-linear Schrödinger equation with a splitting method.

The non-linear Schrödinger equation (NLS) is given by

$$i\partial_t\theta = -\delta\theta + |\theta|^2\theta$$

Here θ is the wave function of a single particle in two dimensions. Its magnitude $|\theta|^2$ describes the probability distribution of the particle's position.

This equation can be split in to its linear (1) and non-linear (2) part

$$(1) \quad i\partial_t\theta = -\delta\theta$$

$$(2) \quad i\partial_t\theta = |\theta|^2\theta$$

For both of these equations, we can compute exact solutions, assuming periodic boundary conditions and using the Fourier series expansion for (1) and using the fact that $|\theta|^2$ is constant in (2):

$$(\hat{1}) \quad \partial_t\hat{\theta}_k = -i|k|^2\hat{\theta}_k$$

$$(2') \quad \partial_t|\theta|^2 = i|\theta|^2(\theta - \bar{\theta}) = 0$$

The exact solutions are then given by

$$(\hat{1}) \quad \hat{\theta}(t) = e^{-i|k|^2t}\hat{\theta}(0)$$

$$(2') \quad \theta(t) = e^{-i|\theta|^2t}\theta(0)$$

These partial solutions can be used to approximate a solution to the full NLS by alternating between small time steps for (1) and (2).

For nicer visual results, we add another constant potential term $V(x)\theta$ to the non-linear part, which turns it into the Gross–Pitaevskii equation.

About the example

The commented program

```
\f}
The exact solutions are then given by
\{align*}{
    (\hat{1}) \quad \hat{\theta}(t) \quad \&= \quad e^{-i |k|^2 t} \quad \hat{\theta}(0) \\
    (2') \quad \theta(t) \quad \&= \quad e^{-i |\theta|^2 t} \quad \theta(0)
}\f}
```

These partial solutions can be used to approximate a solution to the full NLS by alternating between small time steps for (1) and (2).

For nicer visual results, we add another constant potential term $V(x)$ θ to the non-linear part, which turns it into the Gross-Pitaevskii equation.

```
*****<DESCRIPTION>***** /
```

```
#include <algorithm>
#include <chrono>
#include <fstream>
#include <iostream>
#include <utility>

#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>

#include <ginkgo/ginkgo.hpp>
```

This function implements a simple Ginkgo-themed clamped color mapping for values in the range [0,5].

```
void set_val(unsigned char* data, double value)
{
```

RGB values for the 6 colors used for values 0, 1, ..., 5 We will interpolate linearly between these values.

```
double col_r[] = {255, 221, 129, 201, 249, 255};
double col_g[] = {255, 220, 130, 161, 158, 204};
double col_b[] = {255, 220, 133, 93, 24, 8};
value = std::max(0.0, value);
auto i = std::max(0, std::min(4, int(value)));
auto d = std::max(0.0, std::min(1.0, value - i));
```

OpenCV uses BGR instead of RGB by default, revert indices

```
data[2] = static_cast<unsigned char>(col_r[i + 1] * d + col_r[i] * (1 - d));
data[1] = static_cast<unsigned char>(col_g[i + 1] * d + col_g[i] * (1 - d));
data[0] = static_cast<unsigned char>(col_b[i + 1] * d + col_b[i] * (1 - d));
}
```

Initialize video output with given dimension and FPS (frames per seconds)

```
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
{
    cv::Size videosize{n, n};
    auto output =
        std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', 'l');
    output.first.open("nls.mp4", fourcc, fps, videosize);
    return output;
}
```

Write the current frame to video output using the above color mapping

```
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                    const std::complex<double>* data)
{
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
        for (int j = 0; j < n; j++) {
            set_val(&row[3 * j], abs(data[i * n + j]));
        }
        output.first.write(output.second);
    }
}

int main(int argc, char* argv[])
{
    using vec = gko::matrix::Dense<std::complex<double>>;
    using real_vec = gko::matrix::Dense<double>;
    using fft2 = gko::matrix::Fft2;
```

Problem parameters: simulation length

```
const auto t0 = 15.0;
```

scaling factor for non-linearity

```
const auto nonlinear_scale = 1.0;
```

scaling factor for potential

```
const auto potential_scale = 3.0;
```

Simulation parameters: time scaling factor

```
const auto time_scale = 0.25;
```

number of grid points in each dimension

```
const auto n = 256;
```

number of simulation steps per second

```
const auto steps_per_sec = 1000;
```

number of video frames per second

```
const auto fps = 25;
```

number of grid points

```
const auto n2 = n * n;
```

phase difference between neighboring grid points

```
const auto h = 2.0 * gko::pi<double>() / n;
const auto h2 = h * h;
```

time step size for the simulation

```
const auto tau = 1.0 / steps_per_sec;
const auto idx = [&](int i, int j) { return i * n + j; };
```

create an OpenMP executor

```
auto exec = gko::OmpExecutor::create();
```

load initial state vector

```
std::ifstream initial_stream("data/gko_logo_2d.mtx");
std::ifstream potential_stream("data/gko_text_2d.mtx");
auto amplitude = gko::read<vec>(initial_stream, exec);
auto potential = gko::read<real_vec>(potential_stream, exec);
```

create vector for frequency space representation

```
auto frequency = vec::create(exec, amplitude->get_size());
```

create Fourier matrix

```
auto fft = fft2::create(exec, n, n);
auto ifft = fft->conj_transpose();
```

prepare video output

```
auto output = build_output(n, fps);
```

time stamp of the last output frame (sentinel value)

```
double last_t = -t0;
```

execute splitting method: time step in linear part, then non-linear part

```
for (double t = 0; t < t0; t += tau) {
```

if enough time has passed, output the next frame

```
if (t - last_t > 1.0 / fps) {
    last_t = t;
    std::cout << t << std::endl;
    output_timestep(output, n, amplitude->get_const_values());
}
```

time step in linear part

```
fft->apply(amplitude, frequency);
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        frequency->at(idx(i, j)) *=
            std::polar(1.0, -h2 * (i * i + j * j) * tau * time_scale);
    }
}
```

scale by FFT*iFFT normalization factor

```
        frequency->at(idx(i, j)) *= 1.0 / n2;
    }
}
ifft->apply(frequency, amplitude);
```

time step in non-linear part

```
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            amplitude->at(idx(i, j)) *= std::polar(
                1.0, -(nonlinear_scale *
                    gko::squared_norm(amplitude->at(idx(i, j))) +
                    potential_scale * potential->at(idx(i, j))) *
                    tau * time_scale);
        }
    }
}
```

Results

The program will generate a video file named nls.mp4 and output the timestamp of each generated frame.

Comments about programming and debugging

The plain program

```
/******<DESCRIPTION>*****
This example shows how to use the FFT and iFFT implementations in Ginkgo
to solve the non-linear Schrödinger equation with a splitting method.
```

The non-linear Schrödinger equation (NLS) is given by

```
@f$
i \partial_t \theta = -\delta \theta + |\theta|^2 \theta
@f$
```

Here θ is the wave function of a single particle in two dimensions. Its magnitude $|\theta|^2$ describes the probability distribution of the particle's position.

This equation can be split in to its linear (1) and non-linear (2) part

```
\f{align*}{
(1) \quad i \partial_t \theta = -\delta \theta \\
(2) \quad i \partial_t \theta = |\theta|^2 \theta
}
```

For both of these equations, we can compute exact solutions, assuming periodic boundary conditions and using the Fourier series expansion for (1) and using the fact that $|\theta|^2$ is constant in (2):

```
\f{align*}{
(\hat{1}) \quad \partial_t \hat{\theta}_k = -i |k|^2 \hat{\theta}_k \\
(2') \quad \partial_t |\theta|^2 = i |\theta|^2 (\theta - \bar{\theta}) = 0
}
```

The exact solutions are then given by

```
\f{align*}{
(\hat{1}) \quad \hat{\theta}(t) = e^{-i |k|^2 t} \hat{\theta}(0) \\
(2') \quad \theta(t) = e^{-i |\theta|^2 t} \theta(0)
}
```

These partial solutions can be used to approximate a solution to the full NLS by alternating between small time steps for (1) and (2).

For nicer visual results, we add another constant potential term $V(x)$ to the non-linear part, which turns it into the Gross-Pitaevskii equation.

```
*****<DESCRIPTION>*****/
#include <algorithm>
#include <chrono>
#include <fstream>
```

```

#include <iostream>
#include <utility>
#include <opencv2/core.hpp>
#include <opencv2/videoio.hpp>
#include <ginkgo/ginkgo.hpp>
void set_val(unsigned char* data, double value)
{
    double col_r[] = {255, 221, 129, 201, 249, 255};
    double col_g[] = {255, 220, 130, 161, 158, 204};
    double col_b[] = {255, 220, 133, 93, 24, 8};
    value = std::max(0.0, value);
    auto i = std::max(0, std::min(4, int(value)));
    auto d = std::max(0.0, std::min(1.0, value - i));
    data[2] = static_cast<unsigned char>(col_r[i + 1] * d + col_r[i] * (1 - d));
    data[1] = static_cast<unsigned char>(col_g[i + 1] * d + col_g[i] * (1 - d));
    data[0] = static_cast<unsigned char>(col_b[i + 1] * d + col_b[i] * (1 - d));
}
std::pair<cv::VideoWriter, cv::Mat> build_output(int n, double fps)
{
    cv::Size videosize{n, n};
    auto output =
        std::make_pair(cv::VideoWriter{}, cv::Mat{videosize, CV_8UC3});
    auto fourcc = cv::VideoWriter::fourcc('a', 'v', 'c', 'l');
    output.first.open("nls.mp4", fourcc, fps, videosize);
    return output;
}
void output_timestep(std::pair<cv::VideoWriter, cv::Mat>& output, int n,
                    const std::complex<double>* data)
{
    for (int i = 0; i < n; i++) {
        auto row = output.second.ptr(i);
        for (int j = 0; j < n; j++) {
            set_val(&row[3 * j], abs(data[i * n + j]));
        }
    }
    output.first.write(output.second);
}
int main(int argc, char* argv[])
{
    using vec = gko::matrix::Dense<std::complex<double>>;
    using real_vec = gko::matrix::Dense<double>;
    using fft2 = gko::matrix::Fft2;
    const auto t0 = 15.0;
    const auto nonlinear_scale = 1.0;
    const auto potential_scale = 3.0;
    const auto time_scale = 0.25;
    const auto n = 256;
    const auto steps_per_sec = 1000;
    const auto fps = 25;
    const auto n2 = n * n;
    const auto h = 2.0 * gko::pi<double>() / n;
    const auto h2 = h * h;
    const auto tau = 1.0 / steps_per_sec;
    const auto idx = [&](int i, int j) { return i * n + j; };
    auto exec = gko::OmpExecutor::create();
    std::ifstream initial_stream("data/gko_logo_2d.mtx");
    std::ifstream potential_stream("data/gko_text_2d.mtx");
    auto amplitude = gko::read<vec>(initial_stream, exec);
    auto potential = gko::read<real_vec>(potential_stream, exec);
    auto frequency = vec::create(exec, amplitude->get_size());
    auto fft = fft2::create(exec, n, n);
    auto ifft = fft->conj_transpose();
    auto output = build_output(n, fps);
    double last_t = -t0;
    for (double t = 0; t < t0; t += tau) {
        if (t - last_t > 1.0 / fps) {
            last_t = t;
            std::cout << t << std::endl;
            output_timestep(output, n, amplitude->get_const_values());
        }
        fft->apply(amplitude, frequency);
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                frequency->at(idx(i, j)) *=
                    std::polar(1.0, -h2 * (i * i + j * j) * tau * time_scale);
                frequency->at(idx(i, j)) *= 1.0 / n2;
            }
        }
        ifft->apply(frequency, amplitude);
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                amplitude->at(idx(i, j)) *= std::polar(
                    1.0, -(nonlinear_scale *
                        gko::squared_norm(amplitude->at(idx(i, j))) +
                        potential_scale * potential->at(idx(i, j))) *
                        tau * time_scale);
            }
        }
    }
}

```

```
}  
}  
}
```

Chapter 42

The simple-solver program

The simple solver example..

Introduction

This simple solver example should help you get started with Ginkgo. This example is meant for you to understand how Ginkgo works and how you can solve a simple linear system with Ginkgo. We encourage you to play with the code, change the parameters and see what is best suited for your purposes.

About the example

Each example has the following sections:

1. **Introduction:** This gives an overview of the example and mentions any interesting aspects in the example that might help the reader.
2. **The commented program:** This section is intended for you to understand the details of the example so that you can play with it and understand Ginkgo and its features better.
3. **Results:** This section shows the results of the code when run. Though the results may not be completely the same, you can expect the behaviour to be similar.
4. **The plain program:** This is the complete code without any comments to have an complete overview of the code.

The commented program

`gko::matrix::Sellp` could also be used.

```
using mtx = gko::matrix::Csr<ValueType, IndexType>;
```

The `gko::solver::Cg` is used here, but any other solver class can also be used.

```
using cg = gko::solver::Cg<ValueType>;
```

Print the ginkgo version information.

```
std::cout << gko::version_info::get() << std::endl;
```

Print help on how to execute this example.

```
if (argc == 2 && (std::string(argv[1]) == "--help")) {  
    std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;  
    std::exit(-1);  
}
```

Where do you want to run your solver ?

The `gko::Executor` class is one of the cornerstones of Ginkgo. Currently, we have support for an `gko::OmpExecutor`, which uses OpenMP multi-threading in most of its kernels, a `gko::ReferenceExecutor`, a single threaded specialization of the OpenMP executor and a `gko::CudaExecutor` which runs the code on a NVIDIA GPU if available.

Note

With the help of C++, you see that you only ever need to change the executor and all the other functions/ routines within Ginkgo should automatically work and run on the executor with any other changes.

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}}
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Reading your data and transfer to the proper device.

Read the matrix, right hand side and the initial solution using the read function.

Note

Ginkgo uses C++ smart pointers to automatically manage memory. To this end, we use our own object ownership transfer functions that under the hood call the required smart pointer functions to manage object ownership. `gko::share` and `gko::give` are the functions that you would need to use.

```
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
```

Creating the solver

Generate the `gko::solver` factory. Ginkgo uses the concept of Factories to build solvers with certain properties. Observe the Fluent interface used here. Here a cg solver is generated with a stopping criteria of maximum iterations of 20 and a residual norm reduction of 1e-7. You also observe that the stopping criteria(`gko::stop`) are also generated from factories using their build methods. You need to specify the executors which each of the object needs to be built on.

```
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                          .with_reduction_factor(reduction_factor))
        .on(exec);
```

Generate the solver from the matrix. The solver factory built in the previous step takes a "matrix"(a `gko::LinOp` to be more general) as an input. In this case we provide it with a full matrix that we previously read, but as the solver only

effectively uses the `apply()` method within the provided "matrix" object, you can effectively create a `gko::LinOp` class with your own `apply` implementation to accomplish more tasks. We will see an example of how this can be done in the custom-matrix-format example

```
auto solver = solver_gen->generate(A);
```

Finally, solve the system. The solver, being a `gko::LinOp`, can be applied to a right hand side, `b` to obtain the solution, `x`.

```
solver->apply(b, x);
```

Print the solution to the command line.

```
std::cout << "Solution (x):\n";
write(std::cout, x);
```

To measure if your solution has actually converged, you can measure the error of the solution. `one`, `neg_one` are objects that represent the numbers which allow for a uniform interface when computing on any device. To compute the residual, all you need to do is call the `apply` method, which in this case is an `spmv` and equivalent to the LAPACK `z_spmv` routine. Finally, you compute the euclidean 2-norm with the `compute_norm2` function.

```
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}
```

Results

The following is the expected result:

```
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
2.10788e-15
```

Comments about programming and debugging

The plain program

```
#include <ginkgo/ginkgo.hpp>
#include <fstream>
#include <iostream>
#include <map>
#include <string>
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
```

```

using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor] " << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
const auto exec = exec_map.at(executor_string)(); // throws if not valid
auto A = gko::share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
const RealValueType reduction_factor{1e-7};
auto solver_gen =
    cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(20u),
                      gko::stop::ResidualNorm<ValueType>::build()
                        .with_reduction_factor(reduction_factor))
        .on(exec);
auto solver = solver_gen->generate(A);
solver->apply(b, x);
std::cout << "Solution (x):\n";
write(std::cout, x);
auto one = gko::initialize<vec>({1.0}, exec);
auto neg_one = gko::initialize<vec>({-1.0}, exec);
auto res = gko::initialize<real_vec>({0.0}, exec);
A->apply(one, x, neg_one, b);
b->compute_norm2(res);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, res);
}

```

Chapter 43

The simple-solver-logging program

The simple solver with logging example..

This example depends on simple-solver, minimal-cuda-solver.

Introduction

About the example

The commented program

```
{
```

Some shortcuts

```
using ValueType = double;
using RealValueType = gko::remove_complex<ValueType>;
using IndexType = int;
using vec = gko::matrix::Dense<ValueType>;
using real_vec = gko::matrix::Dense<RealValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
```

Print version information

```
std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && (std::string(argv[1]) == "--help")) {
    std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
    std::exit(-1);
}
```

Figure out where to run the code

```
const auto executor_string = argc >= 2 ? argv[1] : "reference";
std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}};
```

executor where Ginkgo will perform the computation

```
const auto exec = exec_map.at(executor_string)(); // throws if not valid
```

Read data

```
auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
```

Let's declare a logger which prints to `std::cout` instead of printing to a file. We log all events except for all linop factory and polymorphic object events. Events masks are group of events which are provided for convenience.

```
std::shared_ptr<gko::log::Stream<ValueType>> stream_logger =
    gko::log::Stream<ValueType>::create(
        gko::log::Logger::all_events_mask ^
        gko::log::Logger::linop_factory_events_mask ^
        gko::log::Logger::polymorphic_object_events_mask,
        std::cout);
```

Add stream_logger to the executor

```
exec->add_logger(stream_logger);
```

Add stream_logger only to the ResidualNorm criterion Factory Note that the logger will get automatically propagated to every criterion generated from this factory.

```
const RealValueType reduction_factor{1e-7};
using ResidualCriterionFactory =
    gko::stop::ResidualNorm<ValueType>::Factory;
std::shared_ptr<ResidualCriterionFactory> residual_criterion =
    ResidualCriterionFactory::create()
        .with_reduction_factor(reduction_factor)
        .on(exec);
residual_criterion->add_logger(stream_logger);
```

Generate solver

```
auto solver_gen =
    cg::build()
        .with_criteria(residual_criterion,
            gko::stop::Iteration::build().with_max_iters(20u))
        .on(exec);
auto solver = solver_gen->generate(A);
```

First we add facilities to only print to a file. It's possible to select events, using masks, e.g. only iterations mask:

`gko::log::Logger::iteration_complete_mask`. See the documentation of Logger class for more information.

```
std::ofstream filestream("my_file.txt");
solver->add_logger(gko::log::Stream<ValueType>::create(
    gko::log::Logger::all_events_mask, filestream));
solver->add_logger(stream_logger);
```

This adds a simple logger that only reports convergence state at the end of the solver. Specifically it captures the last residual norm, the final number of iterations, and the converged or not converged status.

```
std::shared_ptr<gko::log::Convergence<ValueType>> convergence_logger =
    gko::log::Convergence<ValueType>::create();
solver->add_logger(convergence_logger);
```

Add another logger which puts all the data in an object, we can later retrieve this object in our code. Here we only have want Executor and criterion check completed events.

```
std::shared_ptr<gko::log::Record> record_logger =
    gko::log::Record::create(gko::log::Logger::executor_events_mask |
        gko::log::Logger::iteration_complete_mask);
exec->add_logger(record_logger);
solver->add_logger(record_logger);
```

Solve system

```
solver->apply(b, x);
```

Print the residual of the last criterion check event (where convergence happened)

```
auto residual =
    record_logger->get().iteration_completed.back()->residual.get();
auto residual_d = gko::as<vec>(residual);
print_vector("Residual", residual_d);
```

Print solution

```
std::cout << "Solution (x):\n";
write(std::cout, x);
std::cout << "Residual norm sqrt(r^T r):\n";
write(std::cout, gko::as<vec>(convergence_logger->get_residual_norm()));
std::cout << "Number of iterations "
    << convergence_logger->get_num_iterations() << std::endl;
std::cout << "Convergence status " << std::boolalpha
    << convergence_logger->has_converged() << std::endl;
}
```

Results

This is the expected output:

```
[LOG] >> apply started on A LinOp[gko::solver::Cg<double>,0x2142d60] with b
LinOp[gko::matrix::Dense<double>,0x2142140] and x LinOp[gko::matrix::Dense<double>,0x2143450]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142280] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482f0] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21486b0] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148010] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[1]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90] with
Bytes[1]
[LOG] >> Operation[gko::solver::cg::initialize_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::initialize_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*,
gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>
const*, gko::matrix::Dense<double>*>,0x7ffd93d14aa0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*,
gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>
const*, gko::matrix::Dense<double>*>,0x7ffd93d14aa0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[2]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0] with
Bytes[2]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148e50] with
Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0] with
Bytes[8]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14a20] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14a20] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 0 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 0 with ID
1 and finalized set to 1
```

```

[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>>*,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>>*,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>*,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>*,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 0 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149550] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149550] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149550] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149730] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149730] with
Bytes[152]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>**,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>**,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>**,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>**,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>*,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 1 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 1 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>*,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>*,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>*,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]

```

```

[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 1 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149980] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149980] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b80] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b80] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149550]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149550]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*&,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*&,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*&,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*&,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*&,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 2 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 2 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*&,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*&,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]

```



```

[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 2 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149290] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149290] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149290] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149690] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149690] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b80]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149980]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>* &,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>* &,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 3 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 3 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const* &,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>* &,
gko::array<bool>* &, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const* &,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>* &,
gko::array<bool>* &, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 3 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with

```



```

Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149ae0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149ae0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149690]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149290]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149290]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 4 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 4 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 4 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
Bytes[152]

```

```

[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149310] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149310] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149ae0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 5 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 5 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 5 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]

```

```

[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cc0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cc0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149310]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 6 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 6 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 6 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21494f0] with

```

```

Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21494f0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cc0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 7 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 7 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 7 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149730] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149730] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21497d0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21497d0] with
Bytes[152]

```

```

[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21494f0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 8 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 8 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 8 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149200] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21492a0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21492a0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21497d0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149730]

```



```

[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 9 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 9 with ID
1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 9 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21496c0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21496c0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21492a0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149200]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]

```

```

[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 10 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 10 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 10 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149450] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149760] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149760] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21496c0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]

```

```

[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 11 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 11 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 11 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149860] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149860] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149860] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149900] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149900] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149760]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149450]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,

```



```

gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,&
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,&
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 12 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 12 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 12 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21499a0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21499a0] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21493d0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21493d0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21493d0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149900]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149860]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149860]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on

```

```

Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,&
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*>,&
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 13 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 13 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*>, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*>, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*>, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*>, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 13 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149490] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149490] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149580] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149580] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21493d0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21493d0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499a0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]

```

```

[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 14 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 14 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 14 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b50] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149b50] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149b50] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21499c0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x21499c0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149580]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149490]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on

```

```

Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 15 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 15 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*&, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*&,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*&, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*&,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 15 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149a70] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149a70] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21499c0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b50]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b50]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*&, gko::matrix::Dense<double>*&,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]

```

```

gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 16 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 16 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 16 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149970] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149970] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b10] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149b10] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149a70]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with

```



```

Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 17 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 17 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 17 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149780] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149780] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149890] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149890] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149b10]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149970]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*>,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]

```

```

[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 18 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 18 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 18 with
ID 1 and finalized set to 1. It changed one RHS 0, stopped the iteration process 0
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149620] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cf0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149cf0] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149780]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_1_operation<gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::spmv_operation<gko::matrix::Csr<double, int> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14b80] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::solver::cg::step_2_operation<gko::matrix::Dense<double>*&,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::matrix::Dense<double>*, gko::matrix::Dense<double>*,
gko::array<gko::stopping_status>*>,0x7ffd93d14ef0] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x21482f0] with
Bytes[152]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_dot_operation<gko::matrix::Dense<double> const*,

```

```

gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14c50] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> iteration 19 completed with solver LinOp[gko::solver::Cg<double>,0x2142d60] with residual
LinOp[gko::matrix::Dense<double>,0x2147b30], solution LinOp[gko::matrix::Dense<double>,0x2143450] and
residual_norm LinOp[gko::LinOp const*,0]
[LOG] >> check started for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 19 with
ID 1 and finalized set to 1
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*,
gko::matrix::Dense<double>*>,0x7ffd93d14ad0] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*&>,0x7ffd93d14b90] started on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::stop::residual_norm::residual_norm_operation<gko::matrix::Dense<double> const*&,
gko::matrix::Dense<double>*, double&, unsigned char&, bool&, gko::array<gko::stopping_status>*&,
gko::array<bool>*&, bool*, bool*&>,0x7ffd93d14b90] completed on
Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> check completed for stop::Criterion[gko::stop::ResidualNorm<double>,0x2148db0] at iteration 19 with
ID 1 and finalized set to 1. It changed one RHS 1, stopped the iteration process 1
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149890] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x21480a0] to Location[0x2149890] with
Bytes[152]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[152]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149340] with
Bytes[152]
[LOG] >> copy started from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
Bytes[152]
[LOG] >> copy completed from Executor[gko::ReferenceExecutor,0x21400d0] to
Executor[gko::ReferenceExecutor,0x21400d0] from Location[0x2143e90] to Location[0x2149340] with
Bytes[152]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149cf0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149620]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148ee0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147ce0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148e50]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148e50]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2147c90]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482b0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a40]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148a60]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148010]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2148010]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21486b0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21486b0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21484d0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482f0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21482f0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x21480a0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143410]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142280]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142280]
[LOG] >> apply completed on A LinOp[gko::solver::Cg<double>,0x2142d60] with b
LinOp[gko::matrix::Dense<double>,0x2142140] and x LinOp[gko::matrix::Dense<double>,0x2143450]
Last memory copied was of size 98 FROM executor 0x21400d0 pointer 2143e90 TO executor 0x21400d0 pointer
2149340
Residual = [
8.1654e-19
-1.51449e-17
2.23854e-17
-1.0842e-19
6.09864e-20
-1.92446e-18
1.97867e-18
-4.58075e-18
-1.55854e-18
-2.64274e-17
4.20128e-17

```



```

-8.71427e-18
-2.62919e-18
-5.49947e-17
5.51893e-17
-1.57022e-16
-4.2034e-17
-8.71951e-16
1.37837e-15
];
Solution (x):
%%MatrixMarket matrix array real general
19 1
0.252218
0.108645
0.0662811
0.0630433
0.0384088
0.0396536
0.0402648
0.0338935
0.0193098
0.0234653
0.0211499
0.0196413
0.0199151
0.0181674
0.0162722
0.0150714
0.0107016
0.0121141
0.0123025
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0] with Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870] with Bytes[8]
[LOG] >> allocation started on Executor[gko::ReferenceExecutor,0x21400d0] with Bytes[8]
[LOG] >> allocation completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500] with Bytes[8]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*, gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14e50] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::csr::advanced_spmv_operation<gko::matrix::Dense<double> const*, gko::matrix::Csr<double, int> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14e50] completed on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14f70] started on Executor[gko::ReferenceExecutor,0x21400d0]
[LOG] >> Operation[gko::matrix::dense::compute_norm2_operation<gko::matrix::Dense<double> const*, gko::matrix::Dense<double>*>,0x7ffd93d14f70] completed on Executor[gko::ReferenceExecutor,0x21400d0]
Residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
2.10788e-15
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149500]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149870]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2149bb0]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143e90]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143e90]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143590]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143590]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142b10]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2142b10]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143c30]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143c30]
[LOG] >> free started on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143790]
[LOG] >> free completed on Executor[gko::ReferenceExecutor,0x21400d0] at Location[0x2143790]

```

Comments about programming and debugging

The plain program

```

#include <fstream>
#include <iomanip>
#include <iostream>
#include <map>

```

```

#include <string>
#include <ginkgo/ginkgo.hpp>
namespace {
template <typename ValueType>
void print_vector(const std::string& name,
                  const gko::matrix::Dense<ValueType>* vec)
{
    std::cout << name << " = [" << std::endl;
    for (int i = 0; i < vec->get_size()[0]; ++i) {
        std::cout << "      " << vec->at(i, 0) << std::endl;
    }
    std::cout << "];" << std::endl;
}
} // namespace
int main(int argc, char* argv[])
{
    using ValueType = double;
    using RealValueType = gko::remove_complex<ValueType>;
    using IndexType = int;
    using vec = gko::matrix::Dense<ValueType>;
    using real_vec = gko::matrix::Dense<RealValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && (std::string(argv[1]) == "--help")) {
        std::cerr << "Usage: " << argv[0] << " [executor]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>()>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                                gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); }}};
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    auto A = share(gko::read<mtx>(std::ifstream("data/A.mtx"), exec));
    auto b = gko::read<vec>(std::ifstream("data/b.mtx"), exec);
    auto x = gko::read<vec>(std::ifstream("data/x0.mtx"), exec);
    std::shared_ptr<gko::log::Stream<ValueType> stream_logger =
        gko::log::Stream<ValueType>::create(
            gko::log::Logger::all_events_mask ^
            gko::log::Logger::linop_factory_events_mask ^
            gko::log::Logger::polymorphic_object_events_mask,
            std::cout);
    exec->add_logger(stream_logger);
    const RealValueType reduction_factor{1e-7};
    using ResidualCriterionFactory =
        gko::stop::ResidualNorm<ValueType>::Factory;
    std::shared_ptr<ResidualCriterionFactory> residual_criterion =
        ResidualCriterionFactory::create()
            .with_reduction_factor(reduction_factor)
            .on(exec);
    residual_criterion->add_logger(stream_logger);
    auto solver_gen =
        cg::build()
            .with_criteria(residual_criterion,
                          gko::stop::Iteration::build().with_max_iters(20u))
            .on(exec);
    auto solver = solver_gen->generate(A);
    std::ofstream filestream("my_file.txt");
    solver->add_logger(gko::log::Stream<ValueType>::create(
        gko::log::Logger::all_events_mask, filestream));
    solver->add_logger(stream_logger);
    std::shared_ptr<gko::log::Convergence<ValueType> convergence_logger =
        gko::log::Convergence<ValueType>::create();
    solver->add_logger(convergence_logger);
    std::shared_ptr<gko::log::Record> record_logger =
        gko::log::Record::create(gko::log::Logger::executor_events_mask |
                                gko::log::Logger::iteration_complete_mask);
    exec->add_logger(record_logger);
    solver->add_logger(record_logger);
    solver->apply(b, x);
    auto residual =
        record_logger->get().iteration_completed.back()->residual.get();

```

```
    auto residual_d = gko::as<vec>(residual);
    print_vector("Residual", residual_d);
    std::cout << "Solution (x):\n";
    write(std::cout, x);
    std::cout << "Residual norm sqrt(r^T r):\n";
    write(std::cout, gko::as<vec>(convergence_logger->get_residual_norm()));
    std::cout << "Number of iterations "
                << convergence_logger->get_num_iterations() << std::endl;
    std::cout << "Convergence status " << std::boolalpha
                << convergence_logger->has_converged() << std::endl;
}
```


Chapter 44

The three-pt-stencil-solver program

The 3-point stencil example..

This example depends on simple-solver, poisson-solver.

Introduction

This example solves a 1D Poisson equation:

$$\begin{aligned}u &: [0, 1] \rightarrow R \\ u'' &= f \\ u(0) &= u_0 \\ u(1) &= u_1\end{aligned}$$

using a finite difference method on an equidistant grid with K discretization points (K can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

$$\begin{aligned}u(x+h) &= u(x) + u'(x)h + 1/2u''(x)h^2 + O(h^3) \\ u(x-h) &= u(x) - u'(x)h + 1/2u''(x)h^2 + O(h^3) \\ \hline -u(x-h) + 2u(x) - u(x+h) &= -f(x)h^2 + O(h^3)\end{aligned}$$

For an equidistant grid with K "inner" discretization points $x_1, \dots, x_K$, and step size $h = 1/(K+1)$, the formula produces a system of linear equations

$$\begin{aligned}2u_1 - u_2 &= -f_1h^2 + u_0 \\ -u_{k-1} + 2u_k - u_{k+1} &= -f_kh^2, k = 2, \dots, K-1 \\ -u_{K-1} + 2u_K &= -f_Kh^2 + u_1\end{aligned}$$

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function ' f ' is set to ' $f(x) = 6x$ ' (making the solution ' $u(x) = x^3$ '), but that can be changed in the `main` function.

The intention of the example is to show how Ginkgo can be integrated into existing software - the `generate_stencil_matrix`, `generate_rhs`, `print_solution`, `compute_error` and `main` function do not reference Ginkgo at all (i.e. they could have been there before the application developer decided to use Ginkgo, and the only part where Ginkgo is introduced is inside the `solve_system` function).

About the example

The commented program

The function `f` is set to `f(x) = 6x` (making the solution `u(x) = x3`), but that can be changed in the `'main'` function. The intention of the example is to show how Ginkgo can be integrated into existing software – the `'generate_stencil_matrix'`, `'generate_rhs'`, `'print_solution'`, `'compute_error'` and `'main'` function do not reference Ginkgo at all (i.e. they could have been there before the application developer decided to use Ginkgo, and the only part where Ginkgo is introduced is inside the `'solve_system'` function).

```
*****<DESCRIPTION>***** /
#include <iostream>
#include <map>
#include <string>
#include <vector>
#include <ginkgo/ginkgo.hpp>
```

Creates a stencil matrix in CSR format for the given number of discretization points.

```
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType discretization_points,
                           IndexType* row_ptrs, IndexType* col_idxs,
                           ValueType* values)
{
    IndexType pos = 0;
    const ValueType coefs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (IndexType i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
                values[pos] = coefs[ofs + 1];
                col_idxs[pos] = i + ofs;
                ++pos;
            }
        }
        row_ptrs[i + 1] = pos;
    }
}
```

Generates the RHS vector given `f` and the boundary conditions.

```
template <typename Closure, typename ValueType, typename IndexType>
void generate_rhs(IndexType discretization_points, Closure f, ValueType u0,
                 ValueType u1, ValueType* rhs)
{
    const ValueType h = 1.0 / (discretization_points + 1);
    for (IndexType i = 0; i < discretization_points; ++i) {
        const ValueType xi = ValueType(i + 1) * h;
        rhs[i] = -f(xi) * h * h;
    }
    rhs[0] += u0;
    rhs[discretization_points - 1] += u1;
}
```

Prints the solution `u`.

```
template <typename ValueType, typename IndexType>
void print_solution(IndexType discretization_points, ValueType u0, ValueType u1,
                   const ValueType* u)
{
    std::cout << u0 << '\n';
    for (IndexType i = 0; i < discretization_points; ++i) {
        std::cout << u[i] << '\n';
    }
    std::cout << u1 << std::endl;
}
```

Computes the 1-norm of the error given the computed `u` and the correct solution function `correct_u`.

```
template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType discretization_points,
                                                const ValueType* u,
                                                Closure correct_u)
{
    const ValueType h = 1.0 / (discretization_points + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (IndexType i = 0; i < discretization_points; ++i) {
        using std::abs;
        const ValueType xi = ValueType(i + 1) * h;
        error += abs(u[i] - correct_u(xi)) / abs(correct_u(xi));
    }
    return error;
}
```

```

}
template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                  IndexType discretization_points, IndexType* row_ptrs,
                  IndexType* col_idxxs, ValueType* values, ValueType* rhs,
                  ValueType* u, gko::remove_complex<ValueType> reduction_factor)
{

```

Some shortcuts

```

using vec = gko::matrix::Dense<ValueType>;
using mtx = gko::matrix::Csr<ValueType, IndexType>;
using cg = gko::solver::Cg<ValueType>;
using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
using val_array = gko::array<ValueType>;
using idx_array = gko::array<IndexType>;
const auto& dp = discretization_points;

```

Figure out where to run the code

```

std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
exec_map{
    {"omp", [] { return gko::OmpExecutor::create(); }},
    {"cuda",
     [] {
         return gko::CudaExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"hip",
     [] {
         return gko::HipExecutor::create(0, gko::OmpExecutor::create());
     }},
    {"dpcpp",
     [] {
         return gko::DpcppExecutor::create(0,
                                           gko::OmpExecutor::create());
     }},
    {"reference", [] { return gko::ReferenceExecutor::create(); }}};

```

executor where Ginkgo will perform the computation

```

const auto exec = exec_map.at(executor_string()); // throws if not valid

```

executor where the application initialized the data

```

const auto app_exec = exec->get_master();

```

Tell Ginkgo to use the data in our application

Matrix: we have to set the executor of the matrix to the one where we want SpMV's to run (in this case `exec`). When creating array views, we have to specify the executor where the data is (in this case `app_exec`).

If the two do not match, Ginkgo will automatically create a copy of the data on `exec` (however, it will not copy the data back once it is done

- here this is not important since we are not modifying the matrix).

```

auto matrix = mtx::create(exec, gko::dim<2>(dp),
                          val_array::view(app_exec, 3 * dp - 2, values),
                          idx_array::view(app_exec, 3 * dp - 2, col_idxxs),
                          idx_array::view(app_exec, dp + 1, row_ptrs));

```

RHS: similar to matrix

```

auto b = vec::create(exec, gko::dim<2>(dp, 1),
                    val_array::view(app_exec, dp, rhs), 1);

```

Solution: we have to be careful here - if the executors are different, once we compute the solution the array will not be automatically copied back to the original memory locations. Fortunately, whenever `apply` is called on a linear operator (e.g. `matrix`, `solver`) the arguments automatically get copied to the executor where the operator is, and copied back once the operation is completed. Thus, in this case, we can just define the solution on `app_exec`, and it will be automatically transferred to/from `exec` if needed.

```

auto x = vec::create(app_exec, gko::dim<2>(dp, 1),
                    val_array::view(app_exec, dp, u), 1);

```

Generate solver

```

auto solver_gen =
    cg::build()

```

```

        .with_criteria(gko::stop::Iteration::build().with_max_iters(
            gko::size_type(dp)),
            gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
        .with_preconditioner(bj::build())
        .on(exec);
auto solver = solver_gen->generate(gko::give(matrix));

```

Solve system

```

    solver->apply(b, x);
}
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;

```

Print version information

```

std::cout << gko::version_info::get() << std::endl;
if (argc == 2 && std::string(argv[1]) == "--help") {
    std::cerr << "Usage: " << argv[0]
        << " [executor] [DISCRETIZATION_POINTS]" << std::endl;
    std::exit(-1);
}
const auto executor_string = argc >= 2 ? argv[1] : "reference";
const IndexType discretization_points =
    argc >= 3 ? std::atoi(argv[2]) : 100;

```

problem:

```

auto correct_u = [](ValueType x) { return x * x * x; };
auto f = [](ValueType x) { return ValueType(6) * x; };
auto u0 = correct_u(0);
auto u1 = correct_u(1);

```

matrix

```

std::vector<IndexType> row_ptrs(discretization_points + 1);
std::vector<IndexType> col_idxs(3 * discretization_points - 2);
std::vector<ValueType> values(3 * discretization_points - 2);

```

right hand side

```

std::vector<ValueType> rhs(discretization_points);

```

solution

```

std::vector<ValueType> u(discretization_points, 0.0);
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
generate_stencil_matrix(discretization_points, row_ptrs.data(),
    col_idxs.data(), values.data());

```

looking for solution $u = x^3$: $f = 6x$, $u(0) = 0$, $u(1) = 1$

```

generate_rhs(discretization_points, f, u0, u1, rhs.data());
solve_system(executor_string, discretization_points, row_ptrs.data(),
    col_idxs.data(), values.data(), rhs.data(), u.data(),
    reduction_factor);

```

Uncomment to print the solution `print_solution<ValueType, IndexType>(discretization_points, 0, 1, u.data());`

```

    std::cout << "The average relative error is "
        << calculate_error(discretization_points, u.data(), correct_u) /
            discretization_points
        << std::endl;
}

```

Results

This is the expected output:

The average relative error is 2.52236e-11

Comments about programming and debugging

The plain program

```

/*****<DESCRIPTION>*****/
This example solves a 1D Poisson equation:

```

```

u : [0, 1] -> R
u'' = f
u(0) = u0
u(1) = u1

```

using a finite difference method on an equidistant grid with 'K' discretization points ('K' can be controlled with a command line parameter). The discretization is done via the second order Taylor polynomial:

```

u(x + h) = u(x) - u'(x)h + 1/2 u''(x)h^2 + O(h^3)
u(x - h) = u(x) + u'(x)h + 1/2 u''(x)h^2 + O(h^3)  / +
-----
-u(x - h) + 2u(x) + -u(x + h) = -f(x)h^2 + O(h^3)

```

For an equidistant grid with K "inner" discretization points $x_1, \dots, x_K$, and step size $h = 1 / (K + 1)$, the formula produces a system of linear equations

```

2u_1 - u_2      = -f_1 h^2 + u0
-u_(k-1) + 2u_k - u_(k+1) = -f_k h^2,      k = 2, ..., K - 1
-u_(K-1) + 2u_K      = -f_K h^2 + u1

```

which is then solved using Ginkgo's implementation of the CG method preconditioned with block-Jacobi. It is also possible to specify on which executor Ginkgo will solve the system via the command line. The function 'f' is set to 'f(x) = 6x' (making the solution 'u(x) = x^3'), but that can be changed in the 'main' function.

The intention of the example is to show how Ginkgo can be integrated into existing software - the 'generate_stencil_matrix', 'generate_rhs', 'print_solution', 'compute_error' and 'main' function do not reference Ginkgo at all (i.e. they could have been there before the application developer decided to use Ginkgo, and the only part where Ginkgo is introduced is inside the 'solve_system' function.

```

*****/
#include <iostream>
#include <map>
#include <string>
#include <vector>
#include <ginkgo/ginkgo.hpp>
template <typename ValueType, typename IndexType>
void generate_stencil_matrix(IndexType discretization_points,
                           IndexType* row_ptrs, IndexType* col_idxs,
                           ValueType* values)
{
    IndexType pos = 0;
    const ValueType coefs[] = {-1, 2, -1};
    row_ptrs[0] = pos;
    for (IndexType i = 0; i < discretization_points; ++i) {
        for (auto ofs : {-1, 0, 1}) {
            if (0 <= i + ofs && i + ofs < discretization_points) {
                values[pos] = coefs[ofs + 1];
                col_idxs[pos] = i + ofs;
                ++pos;
            }
        }
        row_ptrs[i + 1] = pos;
    }
}

template <typename Closure, typename ValueType, typename IndexType>
void generate_rhs(IndexType discretization_points, Closure f, ValueType u0,
                 ValueType u1, ValueType* rhs)
{
    const ValueType h = 1.0 / (discretization_points + 1);
    for (IndexType i = 0; i < discretization_points; ++i) {
        const ValueType xi = ValueType(i + 1) * h;
        rhs[i] = -f(xi) * h * h;
    }
    rhs[0] += u0;
    rhs[discretization_points - 1] += u1;
}

template <typename ValueType, typename IndexType>
void print_solution(IndexType discretization_points, ValueType u0, ValueType u1,
                   const ValueType* u)
{
    std::cout << u0 << '\n';
    for (IndexType i = 0; i < discretization_points; ++i) {

```

```

        std::cout << u[i] << '\n';
    }
    std::cout << u1 << std::endl;
}
template <typename Closure, typename ValueType, typename IndexType>
gko::remove_complex<ValueType> calculate_error(IndexType discretization_points,
                                              const ValueType* u,
                                              Closure correct_u)
{
    const ValueType h = 1.0 / (discretization_points + 1);
    gko::remove_complex<ValueType> error = 0.0;
    for (IndexType i = 0; i < discretization_points; ++i) {
        using std::abs;
        const ValueType xi = ValueType(i + 1) * h;
        error += abs(u[i] - correct_u(xi)) / abs(correct_u(xi));
    }
    return error;
}
template <typename ValueType, typename IndexType>
void solve_system(const std::string& executor_string,
                 IndexType discretization_points, IndexType* row_ptrs,
                 IndexType* col_idxs, ValueType* values, ValueType* rhs,
                 ValueType* u, gko::remove_complex<ValueType> reduction_factor)
{
    using vec = gko::matrix::Dense<ValueType>;
    using mtx = gko::matrix::Csr<ValueType, IndexType>;
    using cg = gko::solver::Cg<ValueType>;
    using bj = gko::preconditioner::Jacobi<ValueType, IndexType>;
    using val_array = gko::array<ValueType>;
    using idx_array = gko::array<IndexType>;
    const auto& dp = discretization_points;
    std::map<std::string, std::function<std::shared_ptr<gko::Executor>()>>
    exec_map{
        {"omp", [] { return gko::OmpExecutor::create(); }},
        {"cuda",
         [] {
             return gko::CudaExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"hip",
         [] {
             return gko::HipExecutor::create(0, gko::OmpExecutor::create());
         }},
        {"dpcpp",
         [] {
             return gko::DpcppExecutor::create(0,
                                              gko::OmpExecutor::create());
         }},
        {"reference", [] { return gko::ReferenceExecutor::create(); } }},
    const auto exec = exec_map.at(executor_string)(); // throws if not valid
    const auto app_exec = exec->get_master();
    auto matrix = mtx::create(exec, gko::dim<2>(dp),
                             val_array::view(app_exec, 3 * dp - 2, values),
                             idx_array::view(app_exec, 3 * dp - 2, col_idxs),
                             idx_array::view(app_exec, dp + 1, row_ptrs));
    auto b = vec::create(exec, gko::dim<2>(dp, 1),
                        val_array::view(app_exec, dp, rhs), 1);
    auto x = vec::create(app_exec, gko::dim<2>(dp, 1),
                        val_array::view(app_exec, dp, u), 1);
    auto solver_gen =
        cg::build()
        .with_criteria(gko::stop::Iteration::build().with_max_iters(
            gko::size_type(dp)),
            gko::stop::ResidualNorm<ValueType>::build()
            .with_reduction_factor(reduction_factor))
        .with_preconditioner(bj::build())
        .on(exec);
    auto solver = solver_gen->generate(gko::give(matrix));
    solver->apply(b, x);
}
int main(int argc, char* argv[])
{
    using ValueType = double;
    using IndexType = int;
    std::cout << gko::version_info::get() << std::endl;
    if (argc == 2 && std::string(argv[1]) == "--help") {
        std::cerr << "Usage: " << argv[0]
                  << " [executor] [DISCRETIZATION_POINTS]" << std::endl;
        std::exit(-1);
    }
    const auto executor_string = argc >= 2 ? argv[1] : "reference";
    const IndexType discretization_points =
        argc >= 3 ? std::atoi(argv[2]) : 100;
    auto correct_u = [] (ValueType x) { return x * x * x; };
    auto f = [] (ValueType x) { return ValueType(6) * x; };
    auto u0 = correct_u(0);
    auto u1 = correct_u(1);

```

```
std::vector<IndexType> row_ptrs(discretization_points + 1);
std::vector<IndexType> col_idxes(3 * discretization_points - 2);
std::vector<ValueType> values(3 * discretization_points - 2);
std::vector<ValueType> rhs(discretization_points);
std::vector<ValueType> u(discretization_points, 0.0);
const gko::remove_complex<ValueType> reduction_factor = 1e-7;
generate_stencil_matrix(discretization_points, row_ptrs.data(),
                       col_idxes.data(), values.data());
generate_rhs(discretization_points, f, u0, u1, rhs.data());
solve_system(executor_string, discretization_points, row_ptrs.data(),
            col_idxes.data(), values.data(), rhs.data(), u.data(),
            reduction_factor);
std::cout << "The average relative error is "
            << calculate_error(discretization_points, u.data(), correct_u) /
              discretization_points
            << std::endl;
}
```


Chapter 45

Module Documentation

45.1 CUDA Executor

A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.

Classes

- class [gko::CudaExecutor](#)

This is the [Executor](#) subclass which represents the CUDA device.

45.1.1 Detailed Description

A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.

45.2 DPC++ Executor

A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.

Classes

- class [gko::DpcppExecutor](#)

This is the [Executor](#) subclass which represents a DPC++ enhanced device.

45.2.1 Detailed Description

A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.

45.3 Executors

A module dedicated to the implementation and usage of the executors in Ginkgo.

Modules

- [CUDA Executor](#)
A module dedicated to the implementation and usage of the CUDA executor in Ginkgo.
- [DPC++ Executor](#)
A module dedicated to the implementation and usage of the DPC++ executor in Ginkgo.
- [HIP Executor](#)
A module dedicated to the implementation and usage of the HIP executor in Ginkgo.
- [OpenMP Executor](#)
A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.
- [Reference Executor](#)
A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

Classes

- class [gko::Operation](#)
Operations can be used to define functionalities whose implementations differ among devices.
- class [gko::Executor](#)
The first step in using the Ginkgo library consists of creating an executor.
- class [gko::executor_deleter< T >](#)
This is a deleter that uses an executor's `free` method to deallocate the data.
- class [gko::OmpExecutor](#)
This is the [Executor](#) subclass which represents the OpenMP device (typically CPU).
- class [gko::ReferenceExecutor](#)
This is a specialization of the [OmpExecutor](#), which runs the reference implementations of the kernels used for debugging purposes.
- class [gko::CudaExecutor](#)
This is the [Executor](#) subclass which represents the CUDA device.
- class [gko::HipExecutor](#)
This is the [Executor](#) subclass which represents the HIP enhanced device.
- class [gko::DpcppExecutor](#)
This is the [Executor](#) subclass which represents a DPC++ enhanced device.

Macros

- `#define GKO\_REGISTER\_OPERATION(_name, _kernel)`
Binds a set of device-specific kernels to an Operation.
- `#define GKO\_REGISTER\_HOST\_OPERATION(_name, _kernel)`
Binds a host-side kernel (independent of executor type) to an Operation.

45.3.1 Detailed Description

A module dedicated to the implementation and usage of the executors in Ginkgo.

Below, we provide a brief introduction to executors in Ginkgo, how they have been implemented, how to best make use of them and how to add new executors.

45.3.2 Executors in Ginkgo.

The first step in using the Ginkgo library consists of creating an executor. Executors are used to specify the location for the data of linear algebra objects, and to determine where the operations will be executed. Ginkgo currently supports three different executor types:

- [OpenMP Executor](#) specifies that the data should be stored and the associated operations executed on an OpenMP-supporting device (e.g. host CPU);
- [CUDA Executor](#) specifies that the data should be stored and the operations executed on the NVIDIA GPU accelerator;
- [HIP Executor](#) uses the HIP library to compile code for either NVIDIA or AMD GPU accelerator;
- [DPC++ Executor](#) uses the DPC++ compiler for any DPC++ supported hardware (e.g. Intel CPUs, GPU, FPGAs, ...);
- [Reference Executor](#) executes a non-optimized reference implementation, which can be used to debug the library.

45.3.3 Macro Definition Documentation

45.3.3.1 GKO_REGISTER_HOST_OPERATION

```
#define GKO_REGISTER_HOST_OPERATION(
    _name,
    _kernel )
```

Value:

```
template <typename... Args>
auto make_##_name(Args&&... args)
{
    return ::gko::detail::make_register_operation(
        #_kernel,
        [&args...](auto) { _kernel(std::forward<Args>(args)...); });
}
static_assert(true,
    "This assert is used to counter the false positive extra " \
    "semi-colon warnings")
```

Binds a host-side kernel (independent of executor type) to an Operation.

It also defines a helper function which creates the associated operation. Any input arguments passed to the helper function are forwarded to the kernel when the operation is executed. The kernel name is searched for in the namespace where this macro is called. Host operations are used to make computations that are not part of the device kernels visible to profiling loggers and benchmarks.

Parameters

| | |
|----------------------|---|
| <code>_name</code> | operation name |
| <code>_kernel</code> | kernel which will be bound to the operation |

45.3.3.2 Example

```
{c++}
void host_kernel(int) {
    // do some expensive computations
}
// Bind the kernels to the operation
GKO_REGISTER_HOST_OPERATION(my_op, host_kernel);
int main() {
    // create executor
    auto ref = ReferenceExecutor::create();
    // create the operation
    auto op = make_my_op(5); // x = 5
    ref->run(op); // run host kernel
}
```

45.3.3.3 GKO_REGISTER_OPERATION

```
#define GKO_REGISTER_OPERATION(
    _name,
    _kernel )
```

Binds a set of device-specific kernels to an Operation.

It also defines a helper function which creates the associated operation. Any input arguments passed to the helper function are forwarded to the kernel when the operation is executed.

The kernels used to bind the operation are searched in `kernels::DEV_TYPE` namespace, where `DEV_TYPE` is replaced by `omp`, `cuda`, `hip`, `dpcpp` and `reference`.

Parameters

| | |
|----------------------|---|
| <code>_name</code> | operation name |
| <code>_kernel</code> | kernel which will be bound to the operation |

45.3.3.4 Example

```
{c++}
// define the omp, cuda, hip and reference kernels which will be bound to the
// operation
namespace kernels {
namespace omp {
void my_kernel(int x) {
    // omp code
}
}
namespace cuda {
void my_kernel(int x) {
    // cuda code
}
}
namespace hip {
void my_kernel(int x) {
    // hip code
}
}
namespace dpcpp {
void my_kernel(int x) {
    // dpcpp code
}
}
namespace reference {
void my_kernel(int x) {
    // reference code
}
}
```

```
// Bind the kernels to the operation
GKO_REGISTER_OPERATION(my_op, my_kernel);
int main() {
    // create executors
    auto omp = OmpExecutor::create();
    auto cuda = CudaExecutor::create(0, omp);
    auto hip = HipExecutor::create(0, omp);
    auto dpcpp = DpcppExecutor::create(0, omp);
    auto ref = ReferenceExecutor::create();
    // create the operation
    auto op = make_my_op(5); // x = 5
    omp->run(op); // run omp kernel
    cuda->run(op); // run cuda kernel
    hip->run(op); // run hip kernel
    dpcpp->run(op); // run DPC++ kernel
    ref->run(op); // run reference kernel
}
```

45.4 Factorizations

A module dedicated to the implementation and usage of the Factorizations in Ginkgo.

Namespaces

- [gko::factorization](#)

The Factorization namespace.

Classes

- class [gko::factorization::lc< ValueType, IndexType >](#)
Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.
- class [gko::factorization::llu< ValueType, IndexType >](#)
Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.
- class [gko::factorization::Parlc< ValueType, IndexType >](#)
ParlC is an incomplete Cholesky factorization which is computed in parallel.
- class [gko::factorization::Parlct< ValueType, IndexType >](#)
ParlCT is an incomplete threshold-based Cholesky factorization which is computed in parallel.
- class [gko::factorization::Parllu< ValueType, IndexType >](#)
ParllU is an incomplete LU factorization which is computed in parallel.
- class [gko::factorization::Parllut< ValueType, IndexType >](#)
ParllUT is an incomplete threshold-based LU factorization which is computed in parallel.

45.4.1 Detailed Description

A module dedicated to the implementation and usage of the Factorizations in Ginkgo.

45.5 HIP Executor

A module dedicated to the implementation and usage of the HIP executor in Ginkgo.

Classes

- class [gko::HipExecutor](#)

This is the [Executor](#) subclass which represents the HIP enhanced device.

45.5.1 Detailed Description

A module dedicated to the implementation and usage of the HIP executor in Ginkgo.

45.6 Jacobi Preconditioner

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

Classes

- class [gko::batch::preconditioner::Jacobi< ValueType, IndexType >](#)
A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks (stored in a dense row major fashion) of the source operator.
- struct [gko::preconditioner::block_interleaved_storage_scheme< IndexType >](#)
Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.
- class [gko::preconditioner::Jacobi< ValueType, IndexType >](#)
A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

45.6.1 Detailed Description

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

45.7 Linear Operators

A module dedicated to the implementation and usage of the Linear operators in Ginkgo.

Modules

- [Factorizations](#)
A module dedicated to the implementation and usage of the Factorizations in Ginkgo.
- [SpMV employing different Matrix formats](#)
A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.
- [Preconditioners](#)
A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.
- [Solvers](#)
A module dedicated to the implementation and usage of the Solvers in Ginkgo.

Classes

- class [gko::Combination< ValueType >](#)
*The [Combination](#) class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$*
- class [gko::Composition< ValueType >](#)
*The [Composition](#) class can be used to compose linear operators $op1, op2, \dots, opn$ and obtain the operator $op1 * op2 * \dots$*
- class [gko::LinOpFactory](#)
A [LinOpFactory](#) represents a higher order mapping which transforms one linear operator into another.
- class [gko::ReadableFromMatrixData< ValueType, IndexType >](#)
A [LinOp](#) implementing this interface can read its data from a [matrix_data](#) structure.
- class [gko::WritableToMatrixData< ValueType, IndexType >](#)
A [LinOp](#) implementing this interface can write its data to a [matrix_data](#) structure.
- class [gko::Preconditionable](#)
A [LinOp](#) implementing this interface can be preconditioned.
- class [gko::DiagonalLinOpExtractable](#)
The diagonal of a [LinOp](#) can be extracted.
- class [gko::DiagonalExtractable< ValueType >](#)
The diagonal of a [LinOp](#) implementing this interface can be extracted.
- class [gko::EnableAbsoluteComputation< AbsoluteLinOp >](#)
The [EnableAbsoluteComputation](#) mixin provides the default implementations of `compute_absolute_linop` and the `absolute` interface.
- class [gko::EnableLinOp< ConcreteLinOp, PolymorphicBase >](#)
The [EnableLinOp](#) mixin can be used to provide sensible default implementations of the majority of the [LinOp](#) and [PolymorphicObject](#) interface.
- class [gko::Perturbation< ValueType >](#)
*The [Perturbation](#) class can be used to construct a [LinOp](#) to represent the operation $(identity + scalar * basis * projector)$.*
- class [gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >](#)
A [Schwarz](#) preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.
- class [gko::experimental::distributed::Vector< ValueType >](#)
[Vector](#) is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.
- class [gko::factorization::lc< ValueType, IndexType >](#)

- Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.*

 - class `gko::factorization::llu< ValueType, IndexType >`
- Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.*

 - class `gko::factorization::Parlc< ValueType, IndexType >`
- Parlc is an incomplete Cholesky factorization which is computed in parallel.*

 - class `gko::factorization::Parlct< ValueType, IndexType >`
- Parlct is an incomplete threshold-based Cholesky factorization which is computed in parallel.*

 - class `gko::factorization::Parllu< ValueType, IndexType >`
- Parllu is an incomplete LU factorization which is computed in parallel.*

 - class `gko::factorization::Parllut< ValueType, IndexType >`
- Parllut is an incomplete threshold-based LU factorization which is computed in parallel.*

 - class `gko::matrix::Coo< ValueType, IndexType >`
- COO stores a matrix in the coordinate matrix format.*

 - class `gko::matrix::Csr< ValueType, IndexType >`
- CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).*

 - class `gko::matrix::Dense< ValueType >`
- Dense is a matrix format which explicitly stores all values of the matrix.*

 - class `gko::matrix::Diagonal< ValueType >`
- This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).*

 - class `gko::matrix::Ell< ValueType, IndexType >`
- ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.*

 - class `gko::matrix::Fbcsr< ValueType, IndexType >`
- Fixed-block compressed sparse row storage matrix format.*

 - class `gko::matrix::Fft`
- This LinOp implements a 1D Fourier matrix using the FFT algorithm.*

 - class `gko::matrix::Fft2`
- This LinOp implements a 2D Fourier matrix using the FFT algorithm.*

 - class `gko::matrix::Fft3`
- This LinOp implements a 3D Fourier matrix using the FFT algorithm.*

 - class `gko::matrix::Hybrid< ValueType, IndexType >`
- HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.*

 - class `gko::matrix::Identity< ValueType >`
- This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).*

 - class `gko::matrix::IdentityFactory< ValueType >`
- This factory is a utility which can be used to generate `Identity` operators.*

 - class `gko::matrix::Permutation< IndexType >`
- Permutation is a matrix format that represents a permutation matrix, i.e.*

 - class `gko::matrix::RowGatherer< IndexType >`
- RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.*

 - class `gko::matrix::ScaledPermutation< ValueType, IndexType >`
- ScaledPermutation is a matrix combining a permutation with scaling factors.*

 - class `gko::matrix::Sellp< ValueType, IndexType >`
- SELL-P is a matrix format similar to ELL format.*

 - class `gko::matrix::SparsityCsr< ValueType, IndexType >`
- SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).*

 - class `gko::multigrid::FixedCoarsening< ValueType, IndexType >`

- FixedCoarsening* is a very simple coarse grid generation algorithm.
- class `gko::multigrid::Pgm< ValueType, IndexType >`
Parallel graph match (Pgm) is the aggregate method introduced in the paper M.
 - class `gko::preconditioner::Ic< LSolverType, IndexType >`
*The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).*
 - class `gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >`
The Incomplete LU (ILU) preconditioner solves the equation $LUx = b$ for a given lower triangular matrix L , an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).
 - class `gko::preconditioner::Isai< IsaiType, ValueType, IndexType >`
The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A , lower triangular matrix L , upper triangular matrix U or symmetric positive (spd) matrix B .
 - class `gko::preconditioner::Jacobi< ValueType, IndexType >`
A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.
 - class `gko::solver::Bicg< ValueType >`
BICG or the Biconjugate gradient method is a Krylov subspace solver.
 - class `gko::solver::Bicgstab< ValueType >`
BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.
 - class `gko::solver::CbGmres< ValueType >`
CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
 - class `gko::solver::Cg< ValueType >`
CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
 - class `gko::solver::Cgs< ValueType >`
CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.
 - class `gko::solver::Fcg< ValueType >`
FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
 - class `gko::solver::Gcr< ValueType >`
GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.
 - class `gko::solver::Gmres< ValueType >`
GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
 - class `gko::solver::Idr< ValueType >`
IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.
 - class `gko::solver::Ir< ValueType >`
Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.
 - class `gko::solver::Multigrid`
Multigrid methods have a hierarchy of many levels, whose coarse level is a subset of the fine level, of the problem.
 - class `gko::solver::EnableSolverBase< DerivedType, MatrixType >`
A LinOp deriving from this CRTP class stores a system matrix.
 - class `gko::solver::IterativeBase`
A LinOp implementing this interface stores a stopping criterion factory.
 - class `gko::solver::EnableIterativeBase< DerivedType >`
A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.
 - class `gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType >`
A LinOp implementing this interface stores a system matrix and stopping criterion factory.
 - class `gko::solver::LowerTrs< ValueType, IndexType >`

LowerTrs is the triangular solver which solves the system $Lx = b$, when L is a lower triangular matrix.

- class `gko::solver::UpperTrs< ValueType, IndexType >`

UpperTrs is the triangular solver which solves the system $Ux = b$, when U is an upper triangular matrix.

Macros

- #define `GKO_CREATE_FACTORY_PARAMETERS(_parameters_name, _factory_name)`
This Macro will generate a new type containing the parameters for the factory `_factory_name`.
- #define `GKO_ENABLE_BUILD_METHOD(_factory_name)`
Defines a build method for the factory, simplifying its construction by removing the repetitive typing of factory's name.
- #define `GKO_FACTORY_PARAMETER(_name, ...)`
Creates a factory parameter in the factory parameters structure.
- #define `GKO_FACTORY_PARAMETER_SCALAR(_name, _default) GKO_FACTORY_PARAMETER(_name, _default)`
Creates a scalar factory parameter in the factory parameters structure.
- #define `GKO_FACTORY_PARAMETER_VECTOR(_name, ...) GKO_FACTORY_PARAMETER(_name, _VA_ARGS_)`
Creates a vector factory parameter in the factory parameters structure.
- #define `GKO_ENABLE_LIN_OP_FACTORY(_lin_op, _parameters_name, _factory_name)`
This macro will generate a default implementation of a LinOpFactory for the LinOp subclass it is defined in.

Typedefs

- template<typename ConcreteFactory, typename ConcreteLinOp, typename ParametersType, typename PolymorphicBase = LinOpFactory>
using `gko::EnableDefaultLinOpFactory` = `EnableDefaultFactory< ConcreteFactory, ConcreteLinOp, ParametersType, PolymorphicBase >`
This is an alias for the `EnableDefaultFactory` mixin, which correctly sets the template parameters to enable a subclass of `LinOpFactory`.

Functions

- template<typename Matrix, typename... TArgs>
std::unique_ptr< Matrix > `gko::initialize` (size_type stride, std::initializer_list< typename Matrix::value_type > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)
Creates and initializes a column-vector.
- template<typename Matrix, typename... TArgs>
std::unique_ptr< Matrix > `gko::initialize` (std::initializer_list< typename Matrix::value_type > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)
Creates and initializes a column-vector.
- template<typename Matrix, typename... TArgs>
std::unique_ptr< Matrix > `gko::initialize` (size_type stride, std::initializer_list< std::initializer_list< typename Matrix::value_type > > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)
Creates and initializes a matrix.
- template<typename Matrix, typename... TArgs>
std::unique_ptr< Matrix > `gko::initialize` (std::initializer_list< std::initializer_list< typename Matrix::value_type > > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)
Creates and initializes a matrix.

45.7.1 Detailed Description

A module dedicated to the implementation and usage of the Linear operators in Ginkgo.

Below we elaborate on one of the most important concepts of Ginkgo, the linear operator. The linear operator (LinOp) is a base class for all linear algebra objects in Ginkgo. The main benefit of having a single base class for the entire collection of linear algebra objects (as opposed to having separate hierarchies for matrices, solvers and preconditioners) is the generality it provides.

45.7.2 Advantages of this approach and usage

A common interface often allows for writing more generic code. If a user's routine requires only operations provided by the LinOp interface, the same code can be used for any kind of linear operators, independent of whether these are matrices, solvers or preconditioners. This feature is also extensively used in Ginkgo itself. For example, a preconditioner used inside a Krylov solver is a LinOp. This allows the user to supply a wide variety of preconditioners: either the ones which were designed to be used in this scenario (like ILU or block-Jacobi), a user-supplied matrix which is known to be a good preconditioner for the specific problem, or even another solver (e.g., if constructing a flexible GMRES solver).

For example, a matrix free implementation would require the user to provide an apply implementation and instead of passing the generated matrix to the solver, they would have to provide their apply implementation for all the executors needed and no other code needs to be changed. See [The custom-matrix-format program](#) example for more details.

45.7.3 Linear operator as a concept

The linear operator (LinOp) is a base class for all linear algebra objects in Ginkgo. The main benefit of having a single base class for the entire collection of linear algebra objects (as opposed to having separate hierarchies for matrices, solvers and preconditioners) is the generality it provides.

First, since all subclasses provide a common interface, the library users are exposed to a smaller set of routines. For example, a matrix-vector product, a preconditioner application, or even a system solve are just different terms given to the operation of applying a certain linear operator to a vector. As such, Ginkgo uses the same routine name, `LinOp::apply()` for each of these operations, where the actual operation performed depends on the type of linear operator involved in the operation.

Second, a common interface often allows for writing more generic code. If a user's routine requires only operations provided by the LinOp interface, the same code can be used for any kind of linear operators, independent of whether these are matrices, solvers or preconditioners. This feature is also extensively used in Ginkgo itself. For example, a preconditioner used inside a Krylov solver is a LinOp. This allows the user to supply a wide variety of preconditioners: either the ones which were designed to be used in this scenario (like ILU or block-Jacobi), a user-supplied matrix which is known to be a good preconditioner for the specific problem, or even another solver (e.g., if constructing a flexible GMRES solver).

A key observation for providing a unified interface for matrices, solvers, and preconditioners is that the most common operation performed on all of them can be expressed as an application of a linear operator to a vector:

- the sparse matrix-vector product with a matrix A is a linear operator application $y = Ax$;
- the application of a preconditioner is a linear operator application $y = M^{-1}x$, where M is an approximation of the original system matrix A (thus a preconditioner represents an "approximate inverse" operator M^{-1}).
- the system solve $Ax = b$ can be viewed as linear operator application $x = A^{-1}b$ (it goes without saying that the implementation of linear system solves does not follow this conceptual idea), so a linear system solver can be viewed as a representation of the operator A^{-1} .

Finally, direct manipulation of LinOp objects is rarely required in simple scenarios. As an illustrative example, one could construct a fixed-point iteration routine $x_{k+1} = Lx_k + b$ as follows:

```
std::unique_ptr<matrix::Dense<>> calculate_fixed_point(
    int iters, const LinOp *L, const matrix::Dense<> *x0
    const matrix::Dense<> *b)
{
    auto x = gko::clone(x0);
    auto tmp = gko::clone(x0);
    auto one = Dense<>::create(L->get_executor(), {1.0,});
    for (int i = 0; i < iters; ++i) {
        L->apply(tmp, x);
        x->add_scaled(one, b);
        tmp->copy_from(x);
    }
    return x;
}
```

Here, if L is a matrix, `LinOp::apply()` refers to the matrix vector product, and `L->apply(a, b)` computes $b = L \cdot a$. `x->add_scaled(one, b)` is the axpy vector update $x := x + b$.

The interesting part of this example is the `apply()` routine at line 4 of the function body. Since this routine is part of the `LinOp` base class, the fixed-point iteration routine can calculate a fixed point not only for matrices, but for any type of linear operator.

Linear Operators

45.7.4 Macro Definition Documentation

45.7.4.1 GKO_CREATE_FACTORY_PARAMETERS

```
#define GKO_CREATE_FACTORY_PARAMETERS(
    _parameters_name,
    _factory_name )
```

Value:

```
public:
    \
    class _factory_name;
    struct _parameters_name##_type
        : public ::gko::enable_parameters_type<_parameters_name##_type, \
        _factory_name>
```

This Macro will generate a new type containing the parameters for the factory `_factory_name`.

For more details, see [GKO_ENABLE_LIN_OP_FACTORY\(\)](#). It is required to use this macro **before** calling the macro [GKO_ENABLE_LIN_OP_FACTORY\(\)](#). It is also required to use the same names for all parameters between both macros.

Parameters

| | |
|-------------------------------|--|
| <code>_parameters_name</code> | name of the parameters member in the class |
| <code>_factory_name</code> | name of the generated factory type |

45.7.4.2 GKO_ENABLE_BUILD_METHOD

```
#define GKO_ENABLE_BUILD_METHOD(
    _factory_name )
```

Value:

```
static auto build()->decltype(_factory_name::create())
{
    return _factory_name::create();
}
static_assert(true,
    "This assert is used to counter the false positive extra "
    "semi-colon warnings")
```

Defines a build method for the factory, simplifying its construction by removing the repetitive typing of factory's name.

Parameters

| | |
|----------------------------|--|
| <code>_factory_name</code> | the factory for which to define the method |
|----------------------------|--|

45.7.4.3 GKO_ENABLE_LIN_OP_FACTORY

```
#define GKO_ENABLE_LIN_OP_FACTORY(
    _lin_op,
    _parameters_name,
    _factory_name )
```

This macro will generate a default implementation of a LinOpFactory for the LinOp subclass it is defined in.

It is required to first call the macro [GKO_CREATE_FACTORY_PARAMETERS\(\)](#) before this one in order to instantiate the parameters type first.

The list of parameters for the factory should be defined in a code block after the macro definition, and should contain a list of `GKO_FACTORY_PARAMETER_*` declarations. The class should provide a constructor with signature `↔ _lin_op(const _factory_name *, std::shared_ptr<const LinOp>)` which the factory will use a callback to construct the object.

A minimal example of a linear operator is the following:

```
```c++ struct MyLinOp : public EnableLinOp<MyLinOp> { GKO_ENABLE_LIN_OP_FACTORY(MyLinOp, my_parameters, Factory)
// a factory parameter named "my_value", of type int and default // value of 5 int GKO_FACTORY_PARAMETER_SCALAR(my_value,
// a factory parameter named my_pair of type std::pair<int, int> // and default value {5, 5} std::pair<int,
int> GKO_FACTORY_PARAMETER_VECTOR(my_pair, 5, 5); }; // constructor needed by EnableLinOp explicit
MyLinOp(std::shared_ptr<const Executor> exec) { : EnableLinOp<MyLinOp>(exec) {} // constructor needed by
the factory explicit MyLinOp(const Factory *factory, std::shared_ptr<const LinOp> matrix) : EnableLinOp<↔
MyLinOp>(factory->get_executor()), matrix->get_size()), // store factory's parameters locally my_parameters_↔
{factory->get_parameters()}, { int value = my_parameters._my_value; // do something with value }
MyLinOp can then be created as follows:
```c++
auto exec = gko::ReferenceExecutor::create();
// create a factory with default 'my_value' parameter
auto fact = MyLinOp::build().on(exec);
// create a operator using the factory:
auto my_op = fact->generate(gko::matrix::Identity::create(exec, 2));
std::cout << my_op->get_my_parameters().my_value; // prints 5
// create a factory with custom 'my_value' parameter
auto fact = MyLinOp::build().with_my_value(0).on(exec);
// create a operator using the factory:
auto my_op = fact->generate(gko::matrix::Identity::create(exec, 2));
std::cout << my_op->get_my_parameters().my_value; // prints 0
```

Note

It is possible to combine both the `#GKO_CREATE_FACTORY_PARAMETER_*`() macros with this one in a unique macro for class **templates** (not with regular classes). Splitting this into two distinct macros allows to use them in all contexts. See <https://stackoverflow.com/q/50202718/9385966> for more details.

Parameters

<code>_lin_op</code>	concrete operator for which the factory is to be created [CRTP parameter]
<code>_parameters_name</code>	name of the parameters member in the class (its type is <code><_parameters_name>_type</code> , the protected member's name is <code><_parameters_name>_</code> , and the public getter's name is <code>get_<_parameters_name>()</code>)
<code>_factory_name</code>	name of the generated factory type

45.7.4.4 GKO_FACTORY_PARAMETER

```
#define GKO_FACTORY_PARAMETER(  
    _name,  
    ... )
```

Value:

```
_name{__VA_ARGS__};  
  
template <typename... Args>  
auto with_##_name(Args&&... _value)  
->std::decay_t<decltype(*(this->self()))>&  
{  
    using type = decltype(this->_name);  
    this->_name = type{std::forward<Args>(_value)...};  
    return *(this->self());  
}  
static_assert(true,  
    "This assert is used to counter the false positive extra "  
    "semi-colon warnings")
```

Creates a factory parameter in the factory parameters structure.

Parameters

<code>_name</code>	name of the parameter
<code>VA_ARGS</code>	default value of the parameter

See also

[GKO_ENABLE_LIN_OP_FACTORY](#) for more details, and usage example

45.7.4.5 GKO_FACTORY_PARAMETER_SCALAR

```
#define GKO_FACTORY_PARAMETER_SCALAR(  
    _name,  
    _default ) GKO_FACTORY_PARAMETER(_name, _default)
```

Creates a scalar factory parameter in the factory parameters structure.

Scalar in this context means that the constructor for this type only takes a single parameter.

Parameters

<code>_name</code>	name of the parameter
<code>_default</code>	default value of the parameter

See also

[GKO_ENABLE_LIN_OP_FACTORY](#) for more details, and usage example

45.7.4.6 GKO_FACTORY_PARAMETER_VECTOR

```
#define GKO_FACTORY_PARAMETER_VECTOR(
    _name,
    ... ) GKO_FACTORY_PARAMETER(_name, __VA_ARGS__)
```

Creates a vector factory parameter in the factory parameters structure.

Vector in this context means that the constructor for this type takes multiple parameters.

Parameters

<code>_name</code>	name of the parameter
<code>_default</code>	default value of the parameter

See also

[GKO_ENABLE_LIN_OP_FACTORY](#) for more details, and usage example

45.7.5 Typedef Documentation

45.7.5.1 EnableDefaultLinOpFactory

```
template<typename ConcreteFactory , typename ConcreteLinOp , typename ParametersType , typename
PolymorphicBase = LinOpFactory>
using gko::EnableDefaultLinOpFactory = typedef EnableDefaultFactory<ConcreteFactory, Concrete←
LinOp, ParametersType, PolymorphicBase>
```

This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of [LinOpFactory](#).

Template Parameters

<i>ConcreteFactory</i>	the concrete factory which is being implemented [CRTP parameter]
<i>ConcreteLinOp</i>	the concrete LinOp type which this factory produces, needs to have a constructor which takes a const ConcreteFactory *, and an std::shared_ptr<const LinOp> as parameters.
<i>ParametersType</i>	a subclass of enable_parameters_type template which defines all of the parameters of the factory
<i>PolymorphicBase</i>	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of LinOpFactory

45.7.6 Function Documentation

45.7.6.1 initialize() [1/4]

```

template<typename Matrix , typename... TArgs>
std::unique_ptr<Matrix> gko::initialize (
    size_type stride,
    std::initializer_list< std::initializer_list< typename Matrix::value_type >>
    vals,
    std::shared_ptr< const Executor > exec,
    TArgs &&... create_args )

```

Creates and initializes a matrix.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type.

Template Parameters

<i>Matrix</i>	matrix type to initialize (Dense has to implement the ConvertibleTo<Matrix> interface)
<i>TArgs</i>	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

<i>stride</i>	row stride for the temporary Dense matrix
<i>vals</i>	values used to initialize the matrix
<i>exec</i>	Executor associated to the matrix
<i>create_args</i>	additional arguments passed to Matrix::create, not including the Executor , which is passed as the first argument

```

1638 {
1639     using dense = matrix::Dense<typename Matrix::value_type>;
1640     size_type num_rows = vals.size();
1641     size_type num_cols = num_rows > 0 ? begin(vals)->size() : 1;
1642     auto tmp =
1643         dense::create(exec->get_master(), dim<2>{num_rows, num_cols}, stride);
1644     size_type ridx = 0;
1645     for (const auto& row : vals) {
1646         size_type cidx = 0;
1647         for (const auto& elem : row) {
1648             tmp->at(ridx, cidx) = elem;
1649             ++cidx;

```

```

1650         }
1651         ++ridx;
1652     }
1653     auto mtx = Matrix::create(exec, std::forward<TArgs>(create_args)...);
1654     tmp->move_to(mtx);
1655     return mtx;
1656 }

```

References `gko::matrix::Dense< ValueType >::at()`.

45.7.6.2 initialize() [2/4]

```

template<typename Matrix , typename... TArgs>
std::unique_ptr<Matrix> gko::initialize (
    size_type stride,
    std::initializer_list< typename Matrix::value_type > vals,
    std::shared_ptr< const Executor > exec,
    TArgs &&... create_args )

```

Creates and initializes a column-vector.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type.

Template Parameters

<i>Matrix</i>	matrix type to initialize (Dense has to implement the <code>ConvertibleTo<Matrix></code> interface)
<i>TArgs</i>	argument types for <code>Matrix::create</code> method (not including the implied Executor as the first argument)

Parameters

<i>stride</i>	row stride for the temporary Dense matrix
<i>vals</i>	values used to initialize the vector
<i>exec</i>	Executor associated to the vector
<i>create_args</i>	additional arguments passed to <code>Matrix::create</code> , not including the Executor , which is passed as the first argument

References `gko::matrix::Dense< ValueType >::at()`.

45.7.6.3 initialize() [3/4]

```

template<typename Matrix , typename... TArgs>
std::unique_ptr<Matrix> gko::initialize (
    std::initializer_list< std::initializer_list< typename Matrix::value_type >>
    vals,
    std::shared_ptr< const Executor > exec,
    TArgs &&... create_args )

```

Creates and initializes a matrix.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type. The stride of the intermediate Dense matrix is set to the number of columns of the initializer list.

Template Parameters

<i>Matrix</i>	matrix type to initialize (Dense has to implement the ConvertibleTo<Matrix> interface)
<i>TArgs</i>	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

<i>vals</i>	values used to initialize the matrix
<i>exec</i>	Executor associated to the matrix
<i>create_args</i>	additional arguments passed to Matrix::create, not including the Executor , which is passed as the first argument

45.7.6.4 initialize() [4/4]

```
template<typename Matrix , typename... TArgs>
std::unique_ptr<Matrix> gko::initialize (
    std::initializer_list< typename Matrix::value_type > vals,
    std::shared_ptr< const Executor > exec,
    TArgs &&... create_args )
```

Creates and initializes a column-vector.

This function first creates a temporary Dense matrix, fills it with passed in values, and then converts the matrix to the requested type. The stride of the intermediate Dense matrix is set to 1.

Template Parameters

<i>Matrix</i>	matrix type to initialize (Dense has to implement the ConvertibleTo<Matrix> interface)
<i>TArgs</i>	argument types for Matrix::create method (not including the implied Executor as the first argument)

Parameters

<i>vals</i>	values used to initialize the vector
<i>exec</i>	Executor associated to the vector
<i>create_args</i>	additional arguments passed to Matrix::create, not including the Executor , which is passed as the first argument

45.8 Logging

A module dedicated to the implementation and usage of the Logging in Ginkgo.

Namespaces

- [gko::batch::log](#)
The logger namespace .
- [gko::log](#)
The logger namespace .

Classes

- class [gko::batch::log::BatchConvergence< ValueType >](#)
Logs the final residuals and iteration counts for a batch solver.
- class [gko::log::Convergence< ValueType >](#)
[Convergence](#) is a Logger which logs data strictly from the `criterion_check_completed` event.
- class [gko::log::Papi< ValueType >](#)
[Papi](#) is a Logger which logs every event to the PAPI software.
- class [gko::log::PerformanceHint](#)
[PerformanceHint](#) is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.
- class [gko::log::Stream< ValueType >](#)
[Stream](#) is a Logger which logs every event to a stream.

45.8.1 Detailed Description

A module dedicated to the implementation and usage of the Logging in Ginkgo.

The Logger class represents a simple Logger object. It comprises all masks and events internally. Every new logging event addition should be done here. The Logger class also provides a default implementation for most events which do nothing, therefore it is not an obligation to change all classes which derive from Logger, although it is good practice. The logger class is built using event masks to control which events should be logged, and which should not.

45.9 SpMV employing different Matrix formats

A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.

Classes

- class `gko::experimental::distributed::Vector< ValueType >`
Vector is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.
- class `gko::batch::matrix::Csr< ValueType, IndexType >`
Csr is a general sparse matrix format that stores the column indices for each nonzero entry and a cumulative sum of the number of nonzeros in each row.
- class `gko::batch::matrix::Dense< ValueType >`
Dense is a batch matrix format which explicitly stores all values of the matrix in each of the batches.
- class `gko::batch::matrix::Ell< ValueType, IndexType >`
Ell is a sparse matrix format that stores the same number of nonzeros in each row, enabling coalesced accesses.
- class `gko::batch::matrix::Identity< ValueType >`
The batch *Identity* matrix, which represents a batch of *Identity* matrices.
- class `gko::matrix::Coo< ValueType, IndexType >`
COO stores a matrix in the coordinate matrix format.
- class `gko::matrix::Csr< ValueType, IndexType >`
CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).
- class `gko::matrix::Dense< ValueType >`
Dense is a matrix format which explicitly stores all values of the matrix.
- class `gko::matrix::Diagonal< ValueType >`
This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).
- class `gko::matrix::Ell< ValueType, IndexType >`
ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.
- class `gko::matrix::Fbcsr< ValueType, IndexType >`
Fixed-block compressed sparse row storage matrix format.
- class `gko::matrix::Fft`
This LinOp implements a 1D Fourier matrix using the FFT algorithm.
- class `gko::matrix::Fft2`
This LinOp implements a 2D Fourier matrix using the FFT algorithm.
- class `gko::matrix::Fft3`
This LinOp implements a 3D Fourier matrix using the FFT algorithm.
- class `gko::matrix::Hybrid< ValueType, IndexType >`
HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.
- class `gko::matrix::Identity< ValueType >`
This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).
- class `gko::matrix::IdentityFactory< ValueType >`
This factory is a utility which can be used to generate *Identity* operators.
- class `gko::matrix::Permutation< IndexType >`
Permutation is a matrix format that represents a permutation matrix, i.e.
- class `gko::matrix::ScaledPermutation< ValueType, IndexType >`
ScaledPermutation is a matrix combining a permutation with scaling factors.
- class `gko::matrix::Sellp< ValueType, IndexType >`
SELL-P is a matrix format similar to ELL format.
- class `gko::matrix::SparsityCsr< ValueType, IndexType >`
SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

45.9.1 Detailed Description

A module dedicated to the implementation and usage of the various Matrix Formats in Ginkgo.

45.10 OpenMP Executor

A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.

Classes

- class [gko::OmpExecutor](#)

This is the [Executor](#) subclass which represents the OpenMP device (typically CPU).

45.10.1 Detailed Description

A module dedicated to the implementation and usage of the OpenMP executor in Ginkgo.

45.11 Preconditioners

A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.

Modules

- [Jacobi Preconditioner](#)

A module dedicated to the implementation and usage of the Jacobi Preconditioner in Ginkgo.

Namespaces

- [gko::experimental::distributed::preconditioner](#)

The Preconditioner namespace.

- [gko::preconditioner](#)

The Preconditioner namespace.

Classes

- class [gko::Preconditionable](#)

A LinOp implementing this interface can be preconditioned.

- class [gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >](#)

A [Schwarz](#) preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

- class [gko::batch::preconditioner::Jacobi< ValueType, IndexType >](#)

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks (stored in a dense row major fashion) of the source operator.

- class [gko::preconditioner::GaussSeidel< ValueType, IndexType >](#)

This class generates the Gauss-Seidel preconditioner.

- class [gko::preconditioner::lc< LSolverType, IndexType >](#)

*The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).*

- class [gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >](#)

The Incomplete LU (ILU) preconditioner solves the equation $LUx = b$ for a given lower triangular matrix L , an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

- class [gko::preconditioner::lsai< IsaiType, ValueType, IndexType >](#)

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A , lower triangular matrix L , upper triangular matrix U or symmetric positive (spd) matrix B .

- class [gko::preconditioner::Jacobi< ValueType, IndexType >](#)

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

- class [gko::preconditioner::Sor< ValueType, IndexType >](#)

This class generates the (S)SOR preconditioner.

45.11.1 Detailed Description

A module dedicated to the implementation and usage of the Preconditioners in Ginkgo.

45.12 Reference Executor

A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

Classes

- class [gko::ReferenceExecutor](#)

This is a specialization of the [OmpExecutor](#), which runs the reference implementations of the kernels used for debugging purposes.

45.12.1 Detailed Description

A module dedicated to the implementation and usage of the Reference executor in Ginkgo.

45.13 Solvers

A module dedicated to the implementation and usage of the Solvers in Ginkgo.

Namespaces

- [gko::solver](#)
The ginkgo Solve namespace.

Classes

- class [gko::batch::solver::Bicgstab< ValueType >](#)
BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.
- class [gko::batch::solver::Cg< ValueType >](#)
Cg or the Conjugate Gradient is a Krylov subspace solver.
- class [gko::batch::solver::BatchSolver](#)
The [BatchSolver](#) is a base class for all batched solvers and provides the common getters and setter for these batched solver classes.
- class [gko::solver::Bicg< ValueType >](#)
BICG or the Biconjugate gradient method is a Krylov subspace solver.
- class [gko::solver::Bicgstab< ValueType >](#)
BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.
- class [gko::solver::CbGmres< ValueType >](#)
CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
- class [gko::solver::Cg< ValueType >](#)
CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
- class [gko::solver::Cgs< ValueType >](#)
CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.
- class [gko::solver::Fcg< ValueType >](#)
FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
- class [gko::solver::Gcr< ValueType >](#)
GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.
- class [gko::solver::Gmres< ValueType >](#)
GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
- class [gko::solver::Idr< ValueType >](#)
IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.
- class [gko::solver::Ir< ValueType >](#)
Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.
- class [gko::solver::Multigrid](#)
Multigrid methods have a hierarchy of many levels, whose coarse level is a subset of the fine level, of the problem.
- class [gko::solver::LowerTrs< ValueType, IndexType >](#)
LowerTrs is the triangular solver which solves the system $Lx = b$, when L is a lower triangular matrix.
- class [gko::solver::UpperTrs< ValueType, IndexType >](#)
UpperTrs is the triangular solver which solves the system $Ux = b$, when U is an upper triangular matrix.

45.13.1 Detailed Description

A module dedicated to the implementation and usage of the Solvers in Ginkgo.

45.14 Stopping criteria

A module dedicated to the implementation and usage of the Stopping Criteria in Ginkgo.

Namespaces

- [gko::stop](#)

The Stopping criterion namespace.

Classes

- class [gko::stop::Combined](#)

The [Combined](#) class is used to combine multiple criterions together through an OR operation.

- class [gko::stop::Iteration](#)

The [Iteration](#) class is a stopping criterion which stops the iteration process after a preset number of iterations.

- class [gko::stop::ResidualNormBase< ValueType >](#)

The [ResidualNormBase](#) class provides a framework for stopping criteria related to the residual norm.

- class [gko::stop::ResidualNorm< ValueType >](#)

The [ResidualNorm](#) class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

- class [gko::stop::ImplicitResidualNorm< ValueType >](#)

The [ImplicitResidualNorm](#) class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

- class [gko::stop::ResidualNormReduction< ValueType >](#)

The [ResidualNormReduction](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

- class [gko::stop::RelativeResidualNorm< ValueType >](#)

The [RelativeResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

- class [gko::stop::AbsoluteResidualNorm< ValueType >](#)

The [AbsoluteResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

- class [gko::stopping_status](#)

This class is used to keep track of the stopping status of one vector.

- class [gko::stop::Time](#)

The [Time](#) class is a stopping criterion which stops the iteration process after a certain amount of time has passed.

Enumerations

- enum [gko::stop::mode](#)

The mode for the residual norm criterion.

Functions

- `template<typename FactoryContainer >
std::shared_ptr< const CriterionFactory > gko::stop::combine (FactoryContainer &&factories)`

Combines multiple criterion factories into a single combined criterion factory.

45.14.1 Detailed Description

A module dedicated to the implementation and usage of the Stopping Criteria in Ginkgo.

45.14.2 Enumeration Type Documentation

45.14.2.1 mode

```
enum gko::stop::mode [strong]
```

The mode for the residual norm criterion.

- absolute: Check for tolerance against residual norm. $\|r\| \leq \tau$
- initial_resnorm: Check for tolerance relative to the initial residual norm. $\|r\| \leq \tau \times \|r_0\|$
- rhs_norm: Check for tolerance relative to the rhs norm. $\|r\| \leq \tau \times \|b\|$

```
38 { absolute, initial_resnorm, rhs_norm };
```

45.14.3 Function Documentation

45.14.3.1 combine()

```
template<typename FactoryContainer >
std::shared_ptr<const CriterionFactory> gko::stop::combine (
    FactoryContainer && factories )
```

Combines multiple criterion factories into a single combined criterion factory.

This function treats a singleton container as a special case and avoids creating an additional object and just returns the input factory.

Template Parameters

<i>FactoryContainer</i>	a random access container type
-------------------------	--------------------------------

Parameters

<i>factories</i>	a list of factories to combined
------------------	---------------------------------

Returns

a combined criterion factory if the input contains multiple factories or the input factory if the input contains only one factory

```
110 {  
111     switch (factories.size()) {  
112     case 0:  
113         GKO_NOT_SUPPORTED(nullptr);  
114     case 1:  
115         if (factories[0] == nullptr) {  
116             GKO_NOT_SUPPORTED(nullptr);  
117         }  
118         return factories[0];  
119     default:  
120         if (factories[0] == nullptr) {  
121             // first factory must be valid to capture executor  
122             GKO_NOT_SUPPORTED(nullptr);  
123         } else {  
124             auto exec = factories[0]->get_executor();  
125             return Combined::build()  
126                 .with_criteria(std::forward<FactoryContainer>(factories))  
127                 .on(exec);  
128         }  
129     }  
130 }
```

45.15 BatchLinOp

Classes

- class `gko::batch::BatchLinOpFactory`
A *BatchLinOpFactory* represents a higher order mapping which transforms one batch linear operator into another.
- class `gko::batch::EnableBatchLinOp< ConcreteBatchLinOp, PolymorphicBase >`
The *EnableBatchLinOp* mixin can be used to provide sensible default implementations of the majority of the *BatchLinOp* and *PolymorphicObject* interface.
- class `gko::batch::matrix::Csr< ValueType, IndexType >`
Csr is a general sparse matrix format that stores the column indices for each nonzero entry and a cumulative sum of the number of nonzeros in each row.
- class `gko::batch::matrix::Dense< ValueType >`
Dense is a batch matrix format which explicitly stores all values of the matrix in each of the batches.
- class `gko::batch::matrix::Ell< ValueType, IndexType >`
Ell is a sparse matrix format that stores the same number of nonzeros in each row, enabling coalesced accesses.
- class `gko::batch::matrix::Identity< ValueType >`
The batch *Identity* matrix, which represents a batch of *Identity* matrices.
- class `gko::batch::preconditioner::Jacobi< ValueType, IndexType >`
A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks (stored in a dense row major fashion) of the source operator.
- class `gko::batch::solver::Bicgstab< ValueType >`
BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.
- class `gko::batch::solver::Cg< ValueType >`
Cg or the Conjugate Gradient is a Krylov subspace solver.

Macros

- `#define GKO_ENABLE_BATCH_LIN_OP_FACTORY(_batch_lin_op, _parameters_name, _factory_name)`
This macro will generate a default implementation of a *BatchLinOpFactory* for the *BatchLinOp* subclass it is defined in.

Typedefs

- `template<typename ConcreteFactory, typename ConcreteBatchLinOp, typename ParametersType, typename PolymorphicBase = BatchLinOpFactory>`
using `gko::batch::EnableDefaultBatchLinOpFactory = EnableDefaultFactory< ConcreteFactory, ConcreteBatchLinOp, ParametersType, PolymorphicBase >`
This is an alias for the *EnableDefaultFactory* mixin, which correctly sets the template parameters to enable a subclass of *BatchLinOpFactory*.

45.15.1 Detailed Description

45.15.2 Batched Linear operator as a concept

A batch linear operator (BatchLinOp) forms the base class for all batched linear algebra objects. In general, it follows the same structure as the LinOp class, but has some crucial differences which make it not strictly representable through or with the LinOp class.

A batched operator is defined as a set of independent linear operators which have no communication/information exchange between them. Therefore, any collective operations between the batches is not possible and not implemented. This allows for each batch to be computed and operated on in an embarrassingly parallel fashion.

A key difference between the LinOp and the BatchLinOp class is that the apply between BatchLinOps is no longer supported. The user can apply a BatchLinOp to a [batch::MultiVector](#) but not to any general BatchLinOp.

Therefore, the BatchLinOp serves only as a base class providing necessary core functionality from Polymorphic object and store the dimensions of the batched object.

Note

Apply to [batch::MultiVector](#) objects are handled by the concrete LinOp and may be moved to the base BatchLinOp class in the future.

[BatchLinOp](#)

45.15.3 Macro Definition Documentation

45.15.3.1 GKO_ENABLE_BATCH_LIN_OP_FACTORY

```
#define GKO_ENABLE_BATCH_LIN_OP_FACTORY(
    _batch_lin_op,
    _parameters_name,
    _factory_name )
```

This macro will generate a default implementation of a BatchLinOpFactory for the BatchLinOp subclass it is defined in.

It is required to first call the macro [GKO_CREATE_FACTORY_PARAMETERS\(\)](#) before this one in order to instantiate the parameters type first.

The list of parameters for the factory should be defined in a code block after the macro definition, and should contain a list of GKO_FACTORY_PARAMETER_* declarations. The class should provide a constructor with signature `_batch_lin_op(const _factory_name *, std::shared_ptr<const BatchLinOp>)` which the factory will use a callback to construct the object.

A minimal example of a batch linear operator is the following:

```
```c++ struct MyBatchLinOp : public EnableBatchLinOp<MyBatchLinOp> { GKO_ENABLE_BATCH_LIN_OP_FACTORY(MyBatchLinOp,
{ // a factory parameter named "my_value", of type int and default // value of 5 int GKO_FACTORY_PARAMETER_SCALAR(my_value,
// a factory parameter named my_pair of type std::pair<int, int> // and default value {5, 5} std::pair<int,
int> GKO_FACTORY_PARAMETER_VECTOR(my_pair, 5, 5); }); // constructor needed by EnableBatchLinOp
```

```
explicit MyBatchLinOp(std::shared_ptr<const Executor> exec) { : EnableBatchLinOp<MyBatchLinOp>(exec) {}
// constructor needed by the factory explicit MyBatchLinOp(const Factory *factory, std::shared_ptr<const BatchLinOp> matrix) : EnableBatchLinOp<MyBatchLinOp>(factory->get_executor()), matrix->get_size(), // store
factory's parameters locally my_parameters_{factory->get_parameters()} { int value = my_parameters_.my_value;
// do something with value }
MyBatchLinOp can then be created as follows:
```c++
auto exec = gko::ReferenceExecutor::create();
// create a factory with default 'my_value' parameter
auto fact = MyBatchLinOp::build().on(exec);
// create a operator using the factory:
auto my_op = fact->generate(gko::batch::matrix::Identity::create(exec, 2));
std::cout << my_op->get_my_parameters().my_value; // prints 5
// create a factory with custom 'my_value' parameter
auto fact = MyBatchLinOp::build().with_my_value(0).on(exec);
// create a operator using the factory:
auto my_op = fact->generate(gko::batch::matrix::Identity::create(exec, 2));
std::cout << my_op->get_my_parameters().my_value; // prints 0
```

Note

It is possible to combine both the `#GKO_CREATE_FACTORY_PARAMETER_*`() macros with this one in a unique macro for class **templates** (not with regular classes). Splitting this into two distinct macros allows to use them in all contexts. See <https://stackoverflow.com/q/50202718/9385966> for more details.

Parameters

<code>_batch_lin_op</code>	concrete operator for which the factory is to be created [CRTP parameter]
<code>_parameters_name</code>	name of the parameters member in the class (its type is <code><_parameters_name>_type</code> , the protected member's name is <code><_parameters_name>_</code> , and the public getter's name is <code>get_<_parameters_name>()</code>)
<code>_factory_name</code>	name of the generated factory type

45.15.4 Typedef Documentation

45.15.4.1 EnableDefaultBatchLinOpFactory

```
template<typename ConcreteFactory , typename ConcreteBatchLinOp , typename ParametersType ,
typename PolymorphicBase = BatchLinOpFactory>
using gko::batch::EnableDefaultBatchLinOpFactory = typedef EnableDefaultFactory<ConcreteFactory, ConcreteBatchLinOp, ParametersType, PolymorphicBase>
```

This is an alias for the `EnableDefaultFactory` mixin, which correctly sets the template parameters to enable a subclass of `BatchLinOpFactory`.

Template Parameters

<i>ConcreteFactory</i>	the concrete factory which is being implemented [CRTP parameter]
<i>ConcreteBatchLinOp</i>	the concrete BatchLinOp type which this factory produces, needs to have a constructor which takes a const ConcreteFactory *, and an std::shared_ptr<const BatchLinOp> as parameters.
<i>ParametersType</i>	a subclass of <code>enable_parameters_type</code> template which defines all of the parameters of the factory
<i>PolymorphicBase</i>	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of <code>BatchLinOpFactory</code>

Chapter 46

Namespace Documentation

46.1 gko Namespace Reference

The Ginkgo namespace.

Namespaces

- [accessor](#)
The accessor namespace.
- [factorization](#)
The Factorization namespace.
- [log](#)
The logger namespace .
- [matrix](#)
The matrix namespace.
- [multigrid](#)
The multigrid components namespace.
- [name_demangling](#)
The name demangling namespace.
- [preconditioner](#)
The Preconditioner namespace.
- [reorder](#)
The Reorder namespace.
- [solver](#)
The ginkgo Solve namespace.
- [stop](#)
The Stopping criterion namespace.
- [syn](#)
The Synthesizer namespace.
- [xstd](#)
The namespace for functionalities after C++14 standard.

Classes

- class [AbsoluteComputable](#)
The [AbsoluteComputable](#) is an interface that allows to get the component wise absolute of a [LinOp](#).
- class [AbstractFactory](#)
The [AbstractFactory](#) is a generic interface template that enables easy implementation of the abstract factory design pattern.
- class [AllocationError](#)
[AllocationError](#) is thrown if a memory allocation fails.
- class [Allocator](#)
Provides generic allocation and deallocation functionality to be used by an [Executor](#).
- class [amd_device](#)
[amd_device](#) handles the number of executor on Amd devices and have the corresponding recursive_mutex.
- struct [are_all_integral](#)
Evaluates if all template arguments Args fulfill std::is_integral.
- class [array](#)
An array is a container which encapsulates fixed-sized arrays, stored on the [Executor](#) tied to the array.
- class [BadDimension](#)
[BadDimension](#) is thrown if an operation is being applied to a [LinOp](#) with bad dimensions.
- struct [batch_dim](#)
A type representing the dimensions of a multidimensional batch object.
- class [BlockOperator](#)
A [BlockOperator](#) represents a linear operator that is partitioned into multiple blocks.
- class [BlockSizeError](#)
[Error](#) that denotes issues between block sizes and matrix dimensions.
- class [Combination](#)
The [Combination](#) class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$
- class [Composition](#)
The [Composition](#) class can be used to compose linear operators $op1, op2, \dots, opn$ and obtain the operator $op1 * op2 * \dots$
- class [ConvertibleTo](#)
[ConvertibleTo](#) interface is used to mark that the implementer can be converted to the object of [ResultType](#).
- class [CpuAllocator](#)
[Allocator](#) using new/delete.
- class [CpuAllocatorBase](#)
Implement this interface to provide an allocator for [OmpExecutor](#) or [ReferenceExecutor](#).
- class [CpuTimer](#)
A timer using std::chrono::steady_clock for timing.
- struct [cpx_real_type](#)
Access the underlying real type of a complex number.
- class [CublasError](#)
[CublasError](#) is thrown when a cuBLAS routine throws a non-zero error code.
- class [cuda_stream](#)
An RAII wrapper for a custom CUDA stream.
- class [CudaAllocator](#)
[Allocator](#) using cudaMalloc.
- class [CudaAllocatorBase](#)
Implement this interface to provide an allocator for [CudaExecutor](#).
- class [CudaError](#)
[CudaError](#) is thrown when a CUDA routine throws a non-zero error code.
- class [CudaExecutor](#)

- This is the [Executor](#) subclass which represents the CUDA device.
- class [CudaTimer](#)

A timer using events for timing on a [CudaExecutor](#).
 - class [CufftError](#)

[CufftError](#) is thrown when a cuFFT routine throws a non-zero error code.
 - class [CurandError](#)

[CurandError](#) is thrown when a cuRAND routine throws a non-zero error code.
 - class [CusparsError](#)

[CusparsError](#) is thrown when a cuSPARSE routine throws a non-zero error code.
 - struct [default_converter](#)

Used to convert objects of type S to objects of type R using `static_cast`.
 - class [deferred_factory_parameter](#)

Represents a factory parameter of factory type that can either be initialized by a pre-existing factory or by passing in a `factory_parameters` object whose `.on(exec)` will be called to instantiate a factory.
 - class [device_matrix_data](#)

This type is a device-side equivalent to [matrix_data](#).
 - class [DiagonalExtractable](#)

The diagonal of a `LinOp` implementing this interface can be extracted.
 - class [DiagonalLinOpExtractable](#)

The diagonal of a `LinOp` can be extracted.
 - struct [dim](#)

A type representing the dimensions of a multidimensional object.
 - class [DimensionMismatch](#)

[DimensionMismatch](#) is thrown if an operation is being applied to `LinOps` of incompatible size.
 - class [DpcppExecutor](#)

This is the [Executor](#) subclass which represents a DPC++ enhanced device.
 - class [DpcppTimer](#)

A timer using kernels for timing on a [DpcppExecutor](#) in profiling mode.
 - class [enable_parameters_type](#)

The [enable_parameters_type](#) mixin is used to create a base implementation of the factory parameters structure.
 - class [EnableAbsoluteComputation](#)

The [EnableAbsoluteComputation](#) mixin provides the default implementations of `compute_absolute_linop` and the `absolute` interface.
 - class [EnableAbstractPolymorphicObject](#)

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new abstract object.
 - class [EnableCreateMethod](#)

This mixin implements a static `create()` method on `ConcreteType` that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an `std::unique_ptr` to such an object.
 - class [EnableDefaultFactory](#)

This mixin provides a default implementation of a concrete factory.
 - class [EnableLinOp](#)

The [EnableLinOp](#) mixin can be used to provide sensible default implementations of the majority of the `LinOp` and `PolymorphicObject` interface.
 - class [EnablePolymorphicAssignment](#)

This mixin is used to enable a default [PolymorphicObject::copy_from\(\)](#) implementation for objects that have implemented conversions between them.
 - class [EnablePolymorphicObject](#)

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new concrete polymorphic object.
 - class [Error](#)

The [Error](#) class is used to report exceptional behaviour in library functions.

- class [Executor](#)
The first step in using the Ginkgo library consists of creating an executor.
- class [executor_deleter](#)
This is a deleter that uses an executor's `free` method to deallocate the data.
- class [half](#)
A class providing basic support for half precision floating point types.
- class [hip_stream](#)
An RAII wrapper for a custom HIP stream.
- class [HipAllocatorBase](#)
Implement this interface to provide an allocator for [HipExecutor](#).
- class [HipblasError](#)
[HipblasError](#) is thrown when a hipBLAS routine throws a non-zero error code.
- class [HipError](#)
[HipError](#) is thrown when a HIP routine throws a non-zero error code.
- class [HipExecutor](#)
This is the [Executor](#) subclass which represents the HIP enhanced device.
- class [HipfftError](#)
[HipfftError](#) is thrown when a hipFFT routine throws a non-zero error code.
- class [HiprandError](#)
[HiprandError](#) is thrown when a hipRAND routine throws a non-zero error code.
- class [HipsparseError](#)
[HipsparseError](#) is thrown when a hipSPARSE routine throws a non-zero error code.
- class [HipTimer](#)
A timer using events for timing on a [HipExecutor](#).
- class [index_set](#)
An index set class represents an ordered set of intervals.
- class [InvalidStateError](#)
Exception thrown if an object is in an invalid state.
- class [KernelNotFound](#)
[KernelNotFound](#) is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.
- class [LinOpFactory](#)
A [LinOpFactory](#) represents a higher order mapping which transforms one linear operator into another.
- class [machine_topology](#)
The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.
- class [matrix_assembly_data](#)
This structure is used as an intermediate type to assemble a sparse matrix.
- struct [matrix_data](#)
This structure is used as an intermediate data type to store a sparse matrix.
- struct [matrix_data_entry](#)
Type used to store nonzeros.
- class [MetisError](#)
[MetisError](#) is thrown when METIS routine throws an error code.
- class [MpiError](#)
[MpiError](#) is thrown when a MPI routine throws a non-zero error code.
- class [NotCompiled](#)
[NotCompiled](#) is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.
- class [NotImplemented](#)
[NotImplemented](#) is thrown in case an operation has not yet been implemented (but will be implemented in the future).
- class [NotSupported](#)

- NotSupported* is thrown in case it is not possible to perform the requested operation on the given object type.
- class `null_deleter`

This is a deleter that does not delete the object.
 - class `nvidia_device`

nvidia_device handles the number of executor on Nvidia devices and have the corresponding *recursive_mutex*.
 - class `OmpExecutor`

This is the *Executor* subclass which represents the OpenMP device (typically CPU).
 - class `Operation`

Operations can be used to define functionalities whose implementations differ among devices.
 - class `OutOfBoundsError`

OutOfBoundsError is thrown if a memory access is detected to be out-of-bounds.
 - class `OverflowError`

OverflowError is thrown when an index calculation for storage requirements overflows.
 - class `Permutable`

Linear operators which support permutation should implement the *Permutable* interface.
 - class `Perturbation`

The *Perturbation* class can be used to construct a *LinOp* to represent the operation $(identity + scalar * basis * projector)$.
 - class `PolymorphicObject`

A *PolymorphicObject* is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.
 - class `precision_reduction`

This class is used to encode storage precisions of low precision algorithms.
 - class `Preconditionable`

A *LinOp* implementing this interface can be preconditioned.
 - class `ptr_param`

This class is used for function parameters in the place of raw pointers.
 - class `range`

A range is a multidimensional view of the memory.
 - class `ReadableFromMatrixData`

A *LinOp* implementing this interface can read its data from a *matrix_data* structure.
 - class `ReferenceExecutor`

This is a specialization of the *OmpExecutor*, which runs the reference implementations of the kernels used for debugging purposes.
 - class `ScaledIdentityAddable`

Adds the operation $M \leftarrow a I + b M$ for matrix *M*, identity operator *I* and scalars *a* and *b*, where *M* is the calling object.
 - class `scoped_device_id_guard`

This move-only class uses RAII to set the device id within a scoped block, if necessary.
 - struct `segmented_array`

A minimal interface for a segmented array.
 - struct `span`

A span is a lightweight structure used to create sub-ranges from other ranges.
 - class `stopping_status`

This class is used to keep track of the stopping status of one vector.
 - class `StreamError`

StreamError is thrown if accessing a stream failed.
 - class `time_point`

An opaque wrapper for a time point generated by a timer.
 - class `Timer`

Represents a generic timer that can be used to record time points and measure time differences on host or device streams.
 - class `Transposable`

- Linear operators which support transposition should implement the [Transposable](#) interface.
- class [UnsupportedMatrixProperty](#)
Exception throws if a matrix does not have a property required by a numerical method.
- class [UseComposition](#)
The [UseComposition](#) class can be used to store the composition information in [LinOp](#).
- class [ValueMismatch](#)
[ValueMismatch](#) is thrown if two values are not equal.
- struct [version](#)
This structure is used to represent versions of various Ginkgo modules.
- class [version_info](#)
Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:
- class [WritableToMatrixData](#)
A [LinOp](#) implementing this interface can write its data to a [matrix_data](#) structure.

Typedefs

- template<typename ConcreteFactory , typename ConcreteLinOp , typename ParametersType , typename PolymorphicBase = [LinOp](#)<Factory>
using [EnableDefaultLinOpFactory](#) = [EnableDefaultFactory](#)< ConcreteFactory, ConcreteLinOp, ParametersType, PolymorphicBase >
This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of [LinOpFactory](#).
- template<typename T >
using [is_complex_s](#) = detail::is_complex_impl< T >
Allows to check if T is a complex value during compile time by accessing the `value` attribute of this struct.
- template<typename T >
using [is_complex_or_scalar_s](#) = detail::is_complex_or_scalar_impl< T >
Allows to check if T is a complex or scalar value during compile time by accessing the `value` attribute of this struct.
- template<typename T >
using [remove_complex](#) = typename detail::remove_complex_s< T >::type
Obtain the type which removed the complex of complex/scalar type or the template parameter of class by accessing the `type` attribute of this struct.
- template<typename T >
using [to_complex](#) = typename detail::to_complex_s< T >::type
Obtain the type which adds the complex of complex/scalar type or the template parameter of class by accessing the `type` attribute of this struct.
- template<typename T >
using [to_real](#) = [remove_complex](#)< T >
`to_real` is alias of `remove_complex`
- template<typename T >
using [next_precision_base](#) = typename detail::next_precision_base_impl< T >::type
Obtains the next type in the singly-linked precision list.
- template<typename T >
using [previous_precision_base](#) = [next_precision_base](#)< T >
Obtains the previous type in the singly-linked precision list.
- template<typename T >
using [next_precision](#) = [next_precision_base](#)< T >
Obtains the next type in the singly-linked precision list with half.
- template<typename T >
using [reduce_precision](#) = typename detail::reduce_precision_impl< T >::type
Obtains the next type in the hierarchy with lower precision than T.

- `template<typename T >`
`using increase\_precision = typename detail::increase_precision_impl< T >::type`
Obtains the next type in the hierarchy with higher precision than T.
- `template<typename... Ts>`
`using highest\_precision = typename detail::highest_precision_variadic< Ts... >::type`
Obtains the smallest arithmetic type that is able to store elements of all template parameter types exactly.
- `template<typename T , size_type Limit = sizeof(uint16) * byte_size>`
`using truncate\_type = std::conditional_t< detail::type_size_impl< T >::value >= 2 * Limit, typename detail::truncate_type_impl< T >::type, T >`
*Truncates the type by half (by dropping bits), but ensures that it is at least *Limit* bits wide.*
- `using size\_type = std::size_t`
Integral type used for allocation quantities.
- `using int8 = std::int8_t`
8-bit signed integral type.
- `using int16 = std::int16_t`
16-bit signed integral type.
- `using int32 = std::int32_t`
32-bit signed integral type.
- `using int64 = std::int64_t`
64-bit signed integral type.
- `using uint8 = std::uint8_t`
8-bit unsigned integral type.
- `using uint16 = std::uint16_t`
16-bit unsigned integral type.
- `using uint32 = std::uint32_t`
32-bit unsigned integral type.
- `using uint64 = std::uint64_t`
64-bit unsigned integral type.
- `using uintptr = std::uintptr_t`
Unsigned integer type capable of holding a pointer to void.
- `using float16 = half`
Half precision floating point type.
- `using float32 = float`
Single precision floating point type.
- `using float64 = double`
Double precision floating point type.
- `using full\_precision = double`
The most precise floating-point type.
- `using default\_precision = double`
Precision used if no precision is explicitly specified.

Enumerations

- `enum log\_propagation\_mode { log\_propagation\_mode::never, log\_propagation\_mode::automatic }`
How Logger events are propagated to their Executor.
- `enum allocation\_mode`
Specify the mode of allocation for CUDA/HIP GPUs.
- `enum layout\_type { layout\_type::array, layout\_type::coordinate }`
Specifies the layout type when writing data in matrix market format.

Functions

- `template<typename ValueType >`
`ValueType reduce_add (const array< ValueType > &input_arr, const ValueType init_val=0)`
Reduce (sum) the values in the array.
- `template<typename ValueType >`
`void reduce_add (const array< ValueType > &input_arr, array< ValueType > &result)`
Reduce (sum) the values in the array.
- `template<typename ValueType >`
`array< ValueType > make_array_view (std::shared_ptr< const Executor > exec, size_type size, ValueType *data)`
Helper function to create an array view deducing the value type.
- `template<typename ValueType >`
`detail::const_array_view< ValueType > make_const_array_view (std::shared_ptr< const Executor > exec, size_type size, const ValueType *data)`
Helper function to create a const array view deducing the value type.
- `template<typename DimensionType >`
`batch_dim< 2, DimensionType > transpose (const batch_dim< 2, DimensionType > &input)`
Returns a batch_dim object with its dimensions swapped for batched operators.
- `template<typename DimensionType >`
`constexpr dim< 2, DimensionType > transpose (const dim< 2, DimensionType > &dimensions) noexcept`
Returns a dim<2> object with its dimensions swapped.
- `template<typename T >`
`constexpr bool is_complex ()`
Checks if T is a complex type.
- `template<typename T >`
`constexpr bool is_complex_or_scalar ()`
Checks if T is a complex/scalar type.
- `template<typename T >`
`constexpr reduce_precision< T > round_down (T val)`
Reduces the precision of the input parameter.
- `template<typename T >`
`constexpr increase_precision< T > round_up (T val)`
Increases the precision of the input parameter.
- `constexpr int64 ceildiv (int64 num, int64 den)`
Performs integer division with rounding up.
- `template<typename T >`
`constexpr T zero ()`
Returns the additive identity for T.
- `template<typename T >`
`constexpr T zero (const T &)`
Returns the additive identity for T.
- `template<typename T >`
`constexpr T one ()`
Returns the multiplicative identity for T.
- `template<typename T >`
`constexpr T one (const T &)`
Returns the multiplicative identity for T.
- `template<typename T >`
`constexpr bool is_zero (T value)`
Returns true if and only if the given value is zero.
- `template<typename T >`
`constexpr bool is_nonzero (T value)`

Returns true if and only if the given value is not zero.

- `template<typename T >`
`constexpr T max (const T &x, const T &y)`

Returns the larger of the arguments.

- `template<typename T >`
`constexpr T min (const T &x, const T &y)`

Returns the smaller of the arguments.

- `template<typename T >`
`constexpr auto real (const T &x)`

Returns the real part of the object.

- `template<typename T >`
`constexpr auto imag (const T &x)`

Returns the imaginary part of the object.

- `template<typename T >`
`constexpr auto conj (const T &x)`

Returns the conjugate of an object.

- `template<typename T >`
`constexpr auto squared\_norm (const T &x) -> decltype(real(conj(x) * x))`

Returns the squared norm of the object.

- `template<typename T >`
`constexpr std::enable_if_t<!is_complex_s< T >::value, T > abs (const T &x)`

Returns the absolute value of the object.

- `template<typename T >`
`constexpr T pi ()`

Returns the value of pi.

- `template<typename T >`
`constexpr std::complex< remove\_complex< T > > unit\_root (int64 n, int64 k=1)`

*Returns the value of $\exp(2 * \pi * i * k / n)$, i.e.*

- `template<typename T >`
`constexpr uint32 get\_significant\_bit (const T &n, uint32 hint=0u) noexcept`

Returns the position of the most significant bit of the number.

- `template<typename T >`
`constexpr T get\_superior\_power (const T &base, const T &limit, const T &hint=T{1}) noexcept`

*Returns the smallest power of *base* not smaller than *limit*.*

- `template<typename T >`
`std::enable_if_t<!is_complex_s< T >::value, bool > is\_finite (const T &value)`

Checks if a floating point number is finite, meaning it is neither +/- infinity nor NaN.

- `template<typename T >`
`std::enable_if_t<is_complex_s< T >::value, bool > is\_finite (const T &value)`

Checks if all components of a complex value are finite, meaning they are neither +/- infinity nor NaN.

- `template<typename T >`
`T safe\_divide (T a, T b)`

Computes the quotient of the given parameters, guarding against division by zero.

- `template<typename T >`
`std::enable_if_t<!is_complex_s< T >::value, bool > is\_nan (const T &value)`

Checks if a floating point number is NaN.

- `template<typename T >`
`std::enable_if_t<is_complex_s< T >::value, bool > is\_nan (const T &value)`

Checks if any component of a complex value is NaN.

- `template<typename T >`
`constexpr std::enable_if_t<!is_complex_s< T >::value, T > nan ()`

Returns a quiet NaN of the given type.

- `template<typename T >`
`constexpr std::enable_if_t< is\_complex\_s< T >::value, T > nan ()`
Returns a complex with both components quiet NaN.
- `template<typename ValueType = default_precision, typename IndexType = int32>`
`matrix\_data< ValueType, IndexType > read\_raw (std::istream &is)`
Reads a matrix stored in matrix market format from an input stream.
- `template<typename ValueType = default_precision, typename IndexType = int32>`
`matrix\_data< ValueType, IndexType > read\_binary\_raw (std::istream &is)`
Reads a matrix stored in Ginkgo's binary matrix format from an input stream.
- `template<typename ValueType = default_precision, typename IndexType = int32>`
`matrix\_data< ValueType, IndexType > read\_generic\_raw (std::istream &is)`
Reads a matrix stored in either binary or matrix market format from an input stream.
- `template<typename ValueType , typename IndexType >`
`void write\_raw (std::ostream &os, const matrix\_data< ValueType, IndexType > &data, layout\_type layout=layout\_type::coordinate)`
Writes a [matrix_data](#) structure to a stream in matrix market format.
- `template<typename ValueType , typename IndexType >`
`void write\_binary\_raw (std::ostream &os, const matrix\_data< ValueType, IndexType > &data)`
Writes a [matrix_data](#) structure to a stream in binary format.
- `template<typename MatrixType , typename StreamType , typename... MatrixArgs>`
`std::unique_ptr< MatrixType > read (StreamType &&is, MatrixArgs &&... args)`
Reads a matrix stored in matrix market format from an input stream.
- `template<typename MatrixType , typename StreamType , typename... MatrixArgs>`
`std::unique_ptr< MatrixType > read\_binary (StreamType &&is, MatrixArgs &&... args)`
Reads a matrix stored in binary format from an input stream.
- `template<typename MatrixType , typename StreamType , typename... MatrixArgs>`
`std::unique_ptr< MatrixType > read\_generic (StreamType &&is, MatrixArgs &&... args)`
Reads a matrix stored either in binary or matrix market format from an input stream.
- `template<typename MatrixPtrType , typename StreamType >`
`void write (StreamType &&os, MatrixPtrType &&matrix, layout\_type layout=detail::mtx\_io\_traits< std::remove_cv_t< detail::pointee< MatrixPtrType > >::default_layout)`
Writes a matrix into an output stream in matrix market format.
- `template<typename MatrixPtrType , typename StreamType >`
`void write\_binary (StreamType &&os, MatrixPtrType &&matrix)`
Writes a matrix into an output stream in binary format.
- `template<typename R , typename T >`
`std::unique_ptr< R, std::function< void(R *)> > copy\_and\_convert\_to (std::shared_ptr< const Executor > exec, T *obj)`
Converts the object to R and places it on [Executor](#) exec.
- `template<typename R , typename T >`
`std::unique_ptr< const R, std::function< void(const R *)> > copy\_and\_convert\_to (std::shared_ptr< const Executor > exec, const T *obj)`
Converts the object to R and places it on [Executor](#) exec.
- `template<typename R , typename T >`
`std::shared_ptr< R > copy\_and\_convert\_to (std::shared_ptr< const Executor > exec, std::shared_ptr< T > obj)`
Converts the object to R and places it on [Executor](#) exec.
- `template<typename R , typename T >`
`std::shared_ptr< const R > copy\_and\_convert\_to (std::shared_ptr< const Executor > exec, std::shared_ptr< const T > obj)`
- `template<typename ValueType , typename Ptr >`
`detail::temporary\_conversion< std::conditional_t< std::is_const< detail::pointee< Ptr > >::value, const matrix::Dense< ValueType >, matrix::Dense< ValueType > > > make\_temporary\_conversion (Ptr &&matrix)`

Convert the given LinOp from matrix::Dense<...> to matrix::Dense<ValueType>.

- template<typename ValueType , typename Function , typename... Args>
void [precision_dispatch](#) (Function fn, Args *... linops)

Calls the given function with each given argument LinOp temporarily converted into matrix::Dense<ValueType> as parameters.
- template<typename ValueType , typename Function >
void [precision_dispatch_real_complex](#) (Function fn, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType> as parameters.*
- template<typename ValueType , typename Function >
void [precision_dispatch_real_complex](#) (Function fn, const LinOp *alpha, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType> as parameters.*
- template<typename ValueType , typename Function >
void [precision_dispatch_real_complex](#) (Function fn, const LinOp *alpha, const LinOp *in, const LinOp *beta, LinOp *out)

Calls the given function with the given LinOps temporarily converted to matrix::Dense<ValueType> as parameters.*
- template<typename ValueType , typename Function >
void [mixed_precision_dispatch](#) (Function fn, const LinOp *in, LinOp *out)

Calls the given function with each given argument LinOp converted into matrix::Dense<ValueType> as parameters.
- template<typename ValueType , typename Function , std::enable_if_t< is_complex< ValueType >()> * = nullptr>
void [mixed_precision_dispatch_real_complex](#) (Function fn, const LinOp *in, LinOp *out)

Calls the given function with the given LinOps cast to their dynamic type matrix::Dense<ValueType> as parameters.*
- template<typename Ptr >
detail::temporary_clone< detail::pointee< Ptr > > [make_temporary_clone](#) (std::shared_ptr< const [Executor](#) > exec, Ptr &&ptr)

Creates a temporary_clone.
- template<typename Ptr >
detail::temporary_clone< detail::pointee< Ptr > > [make_temporary_output_clone](#) (std::shared_ptr< const [Executor](#) > exec, Ptr &&ptr)

Creates a uninitialized temporary_clone that will be copied back to the input afterwards.
- constexpr bool [operator==](#) ([precision_reduction](#) x, [precision_reduction](#) y) noexcept

Checks if two [precision_reduction](#) encodings are equal.
- constexpr bool [operator!=](#) ([precision_reduction](#) x, [precision_reduction](#) y) noexcept

Checks if two [precision_reduction](#) encodings are different.
- template<typename IndexType >
constexpr IndexType [invalid_index](#) ()

Value for an invalid signed index type.
- template<typename Pointer >
detail::cloned_type< Pointer > [clone](#) (const Pointer &p)

Creates a unique clone of the object pointed to by p.
- template<typename Pointer >
detail::cloned_type< Pointer > [clone](#) (std::shared_ptr< const [Executor](#) > exec, const Pointer &p)

Creates a unique clone of the object pointed to by p on [Executor](#) exec.
- template<typename OwningPointer >
detail::shared_type< OwningPointer > [share](#) (OwningPointer &&p)

Marks the object pointed to by p as shared.
- template<typename OwningPointer >
std::remove_reference< OwningPointer >::type && [give](#) (OwningPointer &&p)

Marks that the object pointed to by p can be given to the callee.
- template<typename Pointer >
std::enable_if< detail::have_ownership_s< Pointer >::value, detail::pointee< Pointer > * >::type [lend](#) (const Pointer &p)

Returns a non-owning (plain) pointer to the object pointed to by p.

- `template<typename Pointer >
std::enable_if<!detail::have_ownership_s< Pointer >::value, detail::pointee< Pointer > * >::type lend (const
Pointer &p)`
Returns a non-owning (plain) pointer to the object pointed to by p.
- `template<typename T , typename U >
std::decay_t< T > * as (U *obj)`
Performs polymorphic type conversion.
- `template<typename T , typename U >
const std::decay_t< T > * as (const U *obj)`
Performs polymorphic type conversion.
- `template<typename T , typename U >
std::decay_t< T > * as (ptr_param< U > obj)`
Performs polymorphic type conversion on a ptr_param.
- `template<typename T , typename U >
const std::decay_t< T > * as (ptr_param< const U > obj)`
Performs polymorphic type conversion.
- `template<typename T , typename U >
std::unique_ptr< std::decay_t< T > > as (std::unique_ptr< U > &&obj)`
Performs polymorphic type conversion of a unique_ptr.
- `template<typename T , typename U >
std::shared_ptr< std::decay_t< T > > as (std::shared_ptr< U > obj)`
Performs polymorphic type conversion of a shared_ptr.
- `template<typename T , typename U >
std::shared_ptr< const std::decay_t< T > > as (std::shared_ptr< const U > obj)`
Performs polymorphic type conversion of a shared_ptr.
- `std::ostream & operator<< (std::ostream &os, const version &ver)`
Prints version information to a stream.
- `std::ostream & operator<< (std::ostream &os, const version_info &ver_info)`
Prints library version information in human-readable format to a stream.
- `template<template< typename, typename > class MatrixType, typename... Args>
auto with_matrix_type (Args &&... create_args)`
This function returns a type that delays a call to MatrixType::create.
- `template<typename VecPtr >
std::unique_ptr< matrix::Dense< typename detail::pointee< VecPtr >::value_type > > make_dense_view
(VecPtr &&vector)`
Creates a view of a given Dense vector.
- `template<typename VecPtr >
std::unique_ptr< const matrix::Dense< typename detail::pointee< VecPtr >::value_type > > make_const_dense_view
(VecPtr &&vector)`
Creates a view of a given Dense vector.
- `template<typename Matrix , typename... TArgs>
std::unique_ptr< Matrix > initialize (size_type stride, std::initializer_list< typename Matrix::value_type > vals,
std::shared_ptr< const Executor > exec, TArgs &&... create_args)`
Creates and initializes a column-vector.
- `template<typename Matrix , typename... TArgs>
std::unique_ptr< Matrix > initialize (std::initializer_list< typename Matrix::value_type > vals, std::shared_ptr<
const Executor > exec, TArgs &&... create_args)`
Creates and initializes a column-vector.
- `template<typename Matrix , typename... TArgs>
std::unique_ptr< Matrix > initialize (size_type stride, std::initializer_list< std::initializer_list< typename
Matrix::value_type > > vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)`
Creates and initializes a matrix.

- `template<typename Matrix, typename... TArgs>
std::unique_ptr< Matrix > initialize (std::initializer_list< std::initializer_list< typename Matrix::value_type >>
vals, std::shared_ptr< const Executor > exec, TArgs &&... create_args)`
Creates and initializes a matrix.
- `bool operator== (const stopping_status &x, const stopping_status &y) noexcept`
Checks if two stopping statuses are equivalent.
- `bool operator!= (const stopping_status &x, const stopping_status &y) noexcept`
Checks if two stopping statuses are different.

Variables

- `constexpr size_type byte_size = CHAR_BIT`
Number of bits in a byte.

46.1.1 Detailed Description

The Ginkgo namespace.

46.1.2 Typedef Documentation

46.1.2.1 highest_precision

```
template<typename... Ts>
using gko::highest_precision = typedef typename detail::highest_precision_variadic<Ts...>↵
::type
```

Obtains the smallest arithmetic type that is able to store elements of all template parameter types exactly.

All template type parameters need to be either real or complex types, mixing them is not possible.

Formally, it computes a right-fold over the type list, with the highest precision of a pair of real arithmetic types T1, T2 computed as `decltype(T1{} + T2{})`, or `std::complex<highest_precision<remove_↵complex<T1>, remove_complex<T2>>>` for complex types.

46.1.2.2 is_complex_or_scalar_s

```
template<typename T >
using gko::is_complex_or_scalar_s = typedef detail::is_complex_or_scalar_impl<T>
```

Allows to check if T is a complex or scalar value during compile time by accessing the `value` attribute of this struct.

If `value` is `true`, T is a complex/scalar type, if it is `false`, T is not a complex/scalar type.

Template Parameters

<i>T</i>	type to check
----------	---------------

46.1.2.3 is_complex_s

```
template<typename T >
using gko::is_complex_s = typedef detail::is_complex_impl<T>
```

Allows to check if T is a complex value during compile time by accessing the `value` attribute of this struct.

If `value` is `true`, T is a complex type, if it is `false`, T is not a complex type.

Template Parameters

<i>T</i>	type to check
----------	---------------

46.1.2.4 previous_precision_base

```
template<typename T >
using gko::previous_precision_base = typedef next_precision_base<T>
```

Obtains the previous type in the singly-linked precision list.

Note

Currently our lists contains only two elements, so this is the same as `next_precision`.

46.1.2.5 remove_complex

```
template<typename T >
using gko::remove_complex = typedef typename detail::remove_complex_s<T>::type
```

Obtain the type which removed the complex of complex/scalar type or the template parameter of class by accessing the `type` attribute of this struct.

Template Parameters

<i>T</i>	type to remove complex
----------	------------------------

Note

`remove_complex<class>` can not be used in friend class declaration.

46.1.2.6 to_complex

```
template<typename T >
using gko::to_complex = typedef typename detail::to_complex_s<T>::type
```

Obtain the type which adds the complex of complex/scalar type or the template parameter of class by accessing the `type` attribute of this struct.

Template Parameters

<i>T</i>	type to complex_type
----------	----------------------

Note

`to_complex<class>` can not be used in friend class declaration. the followings are the error message from different combination. `friend to_complex<Csr>;` error: can not recognize it is class correctly. `friend class to_complex<Csr>;` error: using alias template specialization friend class `to_complex_s<Csr<ValueType, IndexType>>::type;` error: can not recognize it is class correctly.

46.1.2.7 to_real

```
template<typename T >
using gko::to_real = typedef remove_complex<T>
```

`to_real` is alias of `remove_complex`

Template Parameters

<i>T</i>	type to real
----------	--------------

46.1.3 Enumeration Type Documentation

46.1.3.1 allocation_mode

```
enum gko::allocation_mode [strong]
```

Specify the mode of allocation for CUDA/HIP GPUs.

`device` allocates memory on the device and Unified Memory model is not used.

`unified_global` allocates memory on the device, but is accessible by the host through the Unified memory model.

`unified_host` allocates memory on the host and it is not available on devices which do not have concurrent accesses switched on, but this access can be explicitly switched on, when necessary.

```
62 { device, unified_global, unified_host };
```

46.1.3.2 layout_type

```
enum gko::layout_type [strong]
```

Specifies the layout type when writing data in matrix market format.

Enumerator

array	The matrix should be written as dense matrix in column-major order.
coordinate	The matrix should be written as a sparse matrix in coordinate format.

```
92                                     {
96     array,
100     coordinate
101 };
```

46.1.3.3 log_propagation_mode

```
enum gko::log_propagation_mode [strong]
```

How Logger events are propagated to their [Executor](#).

Enumerator

never	Events only get reported at loggers attached to the triggering object. (Except for allocation/free, copy and Operations, since they happen at the Executor).
automatic	Events get reported to loggers attached to the triggering object and propagating loggers (Logger::needs_propagation() return true) attached to its Executor .

46.1.4 Function Documentation

46.1.4.1 abs()

```
template<typename T >
constexpr std::enable_if_t<!is_complex_s<T>::value, T> gko::abs (
    const T & x ) [inline], [constexpr]
```

Returns the absolute value of the object.

Template Parameters

<i>T</i>	the type of the object
----------	------------------------

Parameters

<i>x</i>	the object
----------	------------

Returns

`x >= zero<T>() ? x : -x;`

```
933 {
934     return x >= zero<T>() ? x : -x;
935 }
```

Referenced by `is_finite()`.

46.1.4.2 `as()` [1/7]

```
template<typename T , typename U >
const std::decay_t<T>* gko::as (
    const U * obj ) [inline]
```

Performs polymorphic type conversion.

This is the constant version of the function.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the object which should be converted
------------	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws [NotSupported](#).

```
334 {
335     if (auto p = dynamic_cast<const std::decay_t<T>*>(obj)) {
336         return p;
337     } else {
338         throw NotSupported(__FILE__, __LINE__,
339             std::string{"gko::as<" +
340                 name_demangling::get_type_name(typeid(T)) + ">",
341                 name_demangling::get_type_name(typeid(*obj))});
342     }
343 }
```

46.1.4.3 `as()` [2/7]

```
template<typename T , typename U >
const std::decay_t<T>* gko::as (
    ptr_param< const U > obj ) [inline]
```

Performs polymorphic type conversion.

This is the constant version of the function.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the object which should be converted
------------	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws [NotSupported](#).

References `gko::ptr_param< T >::get()`.

46.1.4.4 as() [3/7]

```
template<typename T , typename U >
std::decay_t<T>* gko::as (
    ptr_param< U > obj ) [inline]
```

Performs polymorphic type conversion on a [ptr_param](#).

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the object which should be converted
------------	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws [NotSupported](#).

References `gko::ptr_param< T >::get()`.

46.1.4.5 as() [4/7]

```
template<typename T , typename U >
std::shared_ptr<const std::decay_t<T> > gko::as (
    std::shared_ptr< const U > obj ) [inline]
```

Performs polymorphic type conversion of a `shared_ptr`.

This is the constant version of the function.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the shared_ptr to the object which should be converted.
------------	---

Returns

If successful, returns a shared_ptr to the subtype, otherwise throws [NotSupported](#). This pointer shares ownership with the input pointer.

46.1.4.6 as() [5/7]

```
template<typename T , typename U >
std::shared_ptr<std::decay_t<T> > gko::as (
    std::shared_ptr< U > obj ) [inline]
```

Performs polymorphic type conversion of a shared_ptr.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the shared_ptr to the object which should be converted.
------------	---

Returns

If successful, returns a shared_ptr to the subtype, otherwise throws [NotSupported](#). This pointer shares ownership with the input pointer.

46.1.4.7 as() [6/7]

```
template<typename T , typename U >
std::unique_ptr<std::decay_t<T> > gko::as (
    std::unique_ptr< U > && obj ) [inline]
```

Performs polymorphic type conversion of a unique_ptr.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the unique_ptr to the object which should be converted. If successful, it will be reset to a nullptr.
------------	---

Returns

If successful, returns a unique_ptr to the subtype, otherwise throws [NotSupported](#).

46.1.4.8 as() [7/7]

```
template<typename T , typename U >
std::decay_t<T>* gko::as (
    U * obj ) [inline]
```

Performs polymorphic type conversion.

Template Parameters

<i>T</i>	requested result type
<i>U</i>	static type of the passed object

Parameters

<i>obj</i>	the object which should be converted
------------	--------------------------------------

Returns

If successful, returns a pointer to the subtype, otherwise throws [NotSupported](#).

Referenced by `gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::get_approximate_inverse()`.

46.1.4.9 ceildiv()

```
constexpr int64 gko::ceildiv (
    int64 num,
    int64 den ) [inline], [constexpr]
```

Performs integer division with rounding up.

Parameters

<i>num</i>	numerator
<i>den</i>	denominator

Returns

returns the ceiled quotient.

Referenced by `gko::matrix::Csr< ValueType, IndexType >::load_balance::clac_size()`, `gko::preconditioner::block↵_interleaved_storage_scheme< index_type >::compute_storage_space()`, and `gko::matrix::Csr< ValueType, IndexType >::load_balance::process()`.

46.1.4.10 clone() [1/2]

```
template<typename Pointer >
detail::cloned_type<Pointer> gko::clone (
    const Pointer & p ) [inline]
```

Creates a unique clone of the object pointed to by `p`.

The pointee (i.e. `*p`) needs to have a clone method that returns a `std::unique_ptr` in order for this method to work.

Template Parameters

<i>Pointer</i>	type of pointer to the object (plain or smart pointer)
----------------	--

Parameters

<i>p</i>	a pointer to the object
----------	-------------------------

Note

The difference between this function and directly calling `LinOp::clone()` is that this one preserves the static type of the object.

Referenced by `gko::preconditioner::lc< LSolverType, IndexType >::operator=()`, `gko::preconditioner::llu< L↵SolverType, USolverType, ReverseApply, IndexType >::operator=()`, `gko::solver::EnablePreconditionable< Bicg< ValueType > >::set_preconditioner()`, and `gko::solver::EnableIterativeBase< Bicg< ValueType > >::set_stop_↵criterion_factory()`.

46.1.4.11 clone() [2/2]

```
template<typename Pointer >
detail::cloned_type<Pointer> gko::clone (
```

```
std::shared_ptr< const Executor > exec,
const Pointer & p ) [inline]
```

Creates a unique clone of the object pointed to by `p` on `Executor` `exec`.

The pointee (i.e. `*p`) needs to have a clone method that takes an executor and returns a `std::unique_ptr` in order for this method to work.

Template Parameters

<i>Pointer</i>	type of pointer to the object (plain or smart pointer)
----------------	--

Parameters

<i>exec</i>	the executor where the cloned object should be stored
<i>p</i>	a pointer to the object

Note

The difference between this function and directly calling `LinOp::clone()` is that this one preserves the static type of the object.

46.1.4.12 conj()

```
template<typename T >
constexpr auto gko::conj (
    const T & x ) [inline], [constexpr]
```

Returns the conjugate of an object.

Parameters

<i>x</i>	the number to conjugate
----------	-------------------------

Returns

conjugate of the object (by default, the object itself)

Referenced by `squared_norm()`.

46.1.4.13 copy_and_convert_to() [1/4]

```
template<typename R , typename T >
std::unique_ptr<const R, std::function<void(const R*)> > gko::copy_and_convert_to (
```

```
std::shared_ptr< const Executor > exec,
const T * obj )
```

Converts the object to R and places it on [Executor](#) exec.

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

<i>R</i>	the type to which the object should be converted
<i>T</i>	the type of the input object

Parameters

<i>exec</i>	the executor where the result should be placed
<i>obj</i>	the object that should be converted

Returns

a unique pointer (with dynamically bound deleter) to the converted object

Note

This is a version of the function which adds the const qualifier to the result if the input had the same qualifier.

```
587 {
588     return detail::copy_and_convert_to_impl<const R>(std::move(exec), obj);
589 }
```

46.1.4.14 copy_and_convert_to() [2/4]

```
template<typename R , typename T >
std::shared_ptr<const R> gko::copy_and_convert_to (
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< const T > obj )
```

This is the version that takes in the `std::shared_ptr` and returns a `std::shared_ptr`

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

<i>R</i>	the type to which the object should be converted
<i>T</i>	the type of the input object

Parameters

<i>exec</i>	the executor where the result should be placed
-------------	--

Parameters

<i>obj</i>	the object that should be converted
------------	-------------------------------------

Returns

a shared pointer to the converted object

Note

This is a version of the function which adds the const qualifier to the result if the input had the same qualifier.

46.1.4.15 copy_and_convert_to() [3/4]

```
template<typename R , typename T >
std::shared_ptr<R> gko::copy_and_convert_to (
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< T > obj )
```

Converts the object to R and places it on [Executor](#) exec.

This is the version that takes in the std::shared_ptr and returns a std::shared_ptr

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

<i>R</i>	the type to which the object should be converted
<i>T</i>	the type of the input object

Parameters

<i>exec</i>	the executor where the result should be placed
<i>obj</i>	the object that should be converted

Returns

a shared pointer to the converted object

46.1.4.16 copy_and_convert_to() [4/4]

```
template<typename R , typename T >
std::unique_ptr<R, std::function<void(R*)> > gko::copy_and_convert_to (
```

```
std::shared_ptr< const Executor > exec,
T * obj )
```

Converts the object to R and places it on [Executor](#) `exec`.

If the object is already of the requested type and on the requested executor, the copy and conversion is avoided and a reference to the original object is returned instead.

Template Parameters

<i>R</i>	the type to which the object should be converted
<i>T</i>	the type of the input object

Parameters

<i>exec</i>	the executor where the result should be placed
<i>obj</i>	the object that should be converted

Returns

a unique pointer (with dynamically bound deleter) to the converted object

46.1.4.17 get_significant_bit()

```
template<typename T >
constexpr uint32 gko::get_significant_bit (
    const T & n,
    uint32 hint = 0u ) [constexpr], [noexcept]
```

Returns the position of the most significant bit of the number.

This is the same as the rounded down base-2 logarithm of the number.

Template Parameters

<i>T</i>	a numeric type supporting bit shift and comparison
----------	--

Parameters

<i>n</i>	a number
<i>hint</i>	a lower bound for the position of the significant bit

Returns

maximum of `hint` and the significant bit position of `n`

46.1.4.18 get_superior_power()

```
template<typename T >
constexpr T gko::get_superior_power (
    const T & base,
    const T & limit,
    const T & hint = T{1} ) [constexpr, [noexcept]
```

Returns the smallest power of `base` not smaller than `limit`.

Template Parameters

<i>T</i>	a numeric type supporting multiplication and comparison
----------	---

Parameters

<i>base</i>	the base of the power to be returned
<i>limit</i>	the lower limit on the size of the power returned
<i>hint</i>	a lower bound on the result, has to be a power of <code>base</code>

Returns

the smallest power of `base` not smaller than `limit`

46.1.4.19 give()

```
template<typename OwningPointer >
std::remove_reference<OwningPointer>::type&& gko::give (
    OwningPointer && p ) [inline]
```

Marks that the object pointed to by `p` can be given to the callee.

Effectively calls `std::move(p)`.

Template Parameters

<i>OwningPointer</i>	type of pointer with ownership to the object (has to be a smart pointer)
----------------------	--

Parameters

<i>p</i>	a pointer to the object
----------	-------------------------

Note

The original pointer `p` becomes invalid after this call.

46.1.4.20 `imag()`

```
template<typename T >
constexpr auto gko::imag (
    const T & x ) [inline], [constexpr]
```

Returns the imaginary part of the object.

Template Parameters

<i>T</i>	type of the object
----------	--------------------

Parameters

<i>x</i>	the object
----------	------------

Returns

imaginary part of the object (by default, [zero<T>\(\)](#))

46.1.4.21 `is_complex()`

```
template<typename T >
constexpr bool gko::is_complex ( ) [inline], [constexpr]
```

Checks if T is a complex type.

Template Parameters

<i>T</i>	type to check
----------	---------------

Returns

`true` if T is a complex type, `false` otherwise

46.1.4.22 `is_complex_or_scalar()`

```
template<typename T >
constexpr bool gko::is_complex_or_scalar ( ) [inline], [constexpr]
```

Checks if T is a complex/scalar type.

Template Parameters

<i>T</i>	type to check
----------	---------------

Returns

`true` if `T` is a complex/scalar type, `false` otherwise

46.1.4.23 `is_finite()` [1/2]

```
template<typename T >
std::enable_if_t<!is_complex_s<T>::value, bool> gko::is_finite (
    const T & value ) [inline]
```

Checks if a floating point number is finite, meaning it is neither +/- infinity nor NaN.

Template Parameters

<i>T</i>	type of the value to check
----------	----------------------------

Parameters

<i>value</i>	value to check
--------------	----------------

Returns

`true` if the value is finite, meaning it are neither +/- infinity nor NaN.

References `abs()`.

Referenced by `is_finite()`.

46.1.4.24 `is_finite()` [2/2]

```
template<typename T >
std::enable_if_t<is_complex_s<T>::value, bool> gko::is_finite (
    const T & value ) [inline]
```

Checks if all components of a complex value are finite, meaning they are neither +/- infinity nor NaN.

Template Parameters

<i>T</i>	complex type of the value to check
----------	------------------------------------

Parameters

<i>value</i>	complex value to check
--------------	------------------------

Returns

`true` if both components of the given value are finite, meaning they are neither +/- infinity nor NaN.

References `is_finite()`.

46.1.4.25 `is_nan()` [1/2]

```
template<typename T >
std::enable_if_t<!is_complex_s<T>::value, bool> gko::is_nan (
    const T & value ) [inline]
```

Checks if a floating point number is NaN.

Template Parameters

<i>T</i>	type of the value to check
----------	----------------------------

Parameters

<i>value</i>	value to check
--------------	----------------

Returns

`true` if the value is NaN.

Referenced by `is_nan()`.

46.1.4.26 `is_nan()` [2/2]

```
template<typename T >
std::enable_if_t<is_complex_s<T>::value, bool> gko::is_nan (
    const T & value ) [inline]
```

Checks if any component of a complex value is NaN.

Template Parameters

<i>T</i>	complex type of the value to check
----------	------------------------------------

Parameters

<i>value</i>	complex value to check
--------------	------------------------

Returns

`true` if any component of the given value is NaN.

References `is_nan()`.

46.1.4.27 is_nonzero()

```
template<typename T >
constexpr bool gko::is_nonzero (
    T value ) [inline], [constexpr]
```

Returns true if and only if the given value is not zero.

Template Parameters

<i>T</i>	the type of the value
----------	-----------------------

Parameters

<i>value</i>	the given value
--------------	-----------------

Returns

true iff the given value is not zero, i.e. `value != zero<T>()`

Referenced by `gko::matrix_data< ValueType, IndexType >::diag()`.

46.1.4.28 is_zero()

```
template<typename T >
constexpr bool gko::is_zero (
    T value ) [inline], [constexpr]
```

Returns true if and only if the given value is zero.

Template Parameters

<i>T</i>	the type of the value
----------	-----------------------

Parameters

<i>value</i>	the given value
--------------	-----------------

Returns

true iff the given value is zero, i.e. `value == zero<T>()`

Referenced by `gko::matrix_data<ValueType, IndexType>::remove_zeros()`.

46.1.4.29 `lend()` [1/2]

```
template<typename Pointer >
std::enable_if<detail::have_ownership_s<Pointer>::value, detail::pointee<Pointer>*>::type
gko::lend (
    const Pointer & p ) [inline]
```

Returns a non-owning (plain) pointer to the object pointed to by `p`.

Template Parameters

<i>Pointer</i>	type of pointer to the object (plain or smart pointer)
----------------	--

Parameters

<i>p</i>	a pointer to the object
----------	-------------------------

Note

This is the overload for owning (smart) pointers, that behaves the same as calling `.get()` on the smart pointer.

46.1.4.30 `lend()` [2/2]

```
template<typename Pointer >
std::enable_if<!detail::have_ownership_s<Pointer>::value, detail::pointee<Pointer>*>::type
gko::lend (
    const Pointer & p ) [inline]
```

Returns a non-owning (plain) pointer to the object pointed to by `p`.

Template Parameters

<i>Pointer</i>	type of pointer to the object (plain or smart pointer)
----------------	--

Parameters

<i>p</i>	a pointer to the object
----------	-------------------------

Note

This is the overload for non-owning (plain) pointers, that just returns *p*.

46.1.4.31 `make_array_view()`

```
template<typename ValueType >
array<ValueType> gko::make_array_view (
    std::shared_ptr< const Executor > exec,
    size_type size,
    ValueType * data )
```

Helper function to create an array view deducing the value type.

Parameters

<i>exec</i>	the executor on which the array resides
<i>size</i>	the number of elements for the array
<i>data</i>	the pointer to the array we create a view on.

Template Parameters

<i>ValueType</i>	the type of the array elements
------------------	--------------------------------

Returns

```
array<ValueType>::view(exec, size, data)

789 {
790     return array<ValueType>::view(exec, size, data);
791 }
```

References `make_array_view()`.

Referenced by `make_array_view()`.

46.1.4.32 `make_const_array_view()`

```
template<typename ValueType >
detail::const_array_view<ValueType> gko::make_const_array_view (
    std::shared_ptr< const Executor > exec,
    size_type size,
    const ValueType * data )
```

Helper function to create a const array view deducing the value type.

Parameters

<i>exec</i>	the executor on which the array resides
<i>size</i>	the number of elements for the array
<i>data</i>	the pointer to the array we create a view on.

Template Parameters

<i>ValueType</i>	the type of the array elements
------------------	--------------------------------

Returns

`array<ValueType>::const_view(exec, size, data)`

References `make_const_array_view()`.

Referenced by `make_const_array_view()`.

46.1.4.33 make_const_dense_view()

```
template<typename VecPtr >
std::unique_ptr< const matrix::Dense<typename detail::pointee<VecPtr>::value_type> > gko↵
::make_const_dense_view (
    VecPtr && vector )
```

Creates a view of a given Dense vector.

Template Parameters

<i>VecPtr</i>	a (smart or raw) pointer to the vector.
---------------	---

Parameters

<i>vector</i>	the vector on which to create the view
---------------	--

References `gko::matrix::Dense< ValueType >::create_const_view_of()`.

46.1.4.34 make_dense_view()

```
template<typename VecPtr >
std::unique_ptr<matrix::Dense<typename detail::pointee<VecPtr>::value_type> > gko::make_↵
dense_view (
    VecPtr && vector )
```

Creates a view of a given Dense vector.

Template Parameters

<i>VecPtr</i>	a (smart or raw) pointer to the vector.
---------------	---

Parameters

<i>vector</i>	the vector on which to create the view
---------------	--

References `gko::matrix::Dense< ValueType >::create_view_of()`.

46.1.4.35 `make_temporary_clone()`

```
template<typename Ptr >
detail::temporary_clone<detail::pointee<Ptr> > gko::make_temporary_clone (
    std::shared_ptr< const Executor > exec,
    Ptr && ptr )
```

Creates a temporary_clone.

This is a helper function which avoids the need to explicitly specify the type of the object, as would be the case if using the constructor of temporary_clone.

Parameters

<i>exec</i>	the executor where the clone will be created
<i>ptr</i>	a pointer to the object of which the clone will be created

Template Parameters

<i>Ptr</i>	the (raw or smart) pointer type to be temporarily cloned
------------	--

```
210 {
211     using T = detail::pointee<Ptr>;
212     return detail::temporary_clone<T>(std::move(exec), std::forward<Ptr>(ptr));
213 }
```

Referenced by `gko::ScaledIdentityAddable::add_scaled_identity()`, `gko::matrix::Csr< ValueType, IndexType >::inv_scale()`, and `gko::matrix::Csr< ValueType, IndexType >::scale()`.

46.1.4.36 `make_temporary_conversion()`

```
template<typename ValueType , typename Ptr >
detail::temporary_conversion<std::conditional_t< std::is_const<detail::pointee<Ptr> >::value, const matrix::Dense<ValueType>, matrix::Dense<ValueType> > > gko::make_temporary_
_conversion (
    Ptr && matrix )
```

Convert the given LinOp from `matrix::Dense<...>` to `matrix::Dense<ValueType>`.

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by `next_precision_base<...>` starting from the `ValueType` template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

<i>matrix</i>	the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type <code>matrix::Dense<ValueType></code> , otherwise it will be converted to this type if possible.
---------------	---

Returns

a `detail::temporary_conversion` pointing to the (potentially converted) object.

Exceptions

<i>NotSupported</i>	if the input matrix cannot be converted to <code>matrix::Dense<ValueType></code>
---------------------	--

Template Parameters

<i>ValueType</i>	the value type into whose associated <code>matrix::Dense</code> type to convert the input LinOp.
------------------	--

```

48 {
49     using Pointee = detail::pointee<Ptr>;
50     using Dense = matrix::Dense<ValueType>;
51     using NextDense = matrix::Dense<next_precision<ValueType>>;
52     using NextNextDense =
53         matrix::Dense<next_precision<next_precision<ValueType>>>;
54     using MaybeConstDense =
55         std::conditional_t<std::is_const<Pointee>::value, const Dense, Dense>;
56     auto result = detail::temporary_conversion<
57         MaybeConstDense>::template create<NextDense, NextNextDense>(matrix);
58     if (!result) {
59         GKO_NOT_SUPPORTED(matrix);
60     }
61     return result;
62 }
```

46.1.4.37 make_temporary_output_clone()

```

template<typename Ptr >
detail::temporary_clone<detail::pointee<Ptr> > gko::make_temporary_output_clone (
    std::shared_ptr< const Executor > exec,
    Ptr && ptr )
```

Creates a uninitialized `temporary_clone` that will be copied back to the input afterwards.

It can be used for output parameters to avoid an unnecessary copy in `make_temporary_clone`.

This is a helper function which avoids the need to explicitly specify the type of the object, as would be the case if using the constructor of `temporary_clone`.

Parameters

<i>exec</i>	the executor where the uninitialized clone will be created
<i>ptr</i>	a pointer to the object of which the clone will be created

Generated by Doxygen

Template Parameters

<i>Ptr</i>	the (raw or smart) pointer type to be temporarily cloned
------------	--

46.1.4.38 max()

```
template<typename T >
constexpr T gko::max (
    const T & x,
    const T & y ) [inline], [constexpr]
```

Returns the larger of the arguments.

Template Parameters

<i>T</i>	type of the arguments
----------	-----------------------

Parameters

<i>x</i>	first argument
<i>y</i>	second argument

Returns

$x \geq y ? x : y$

46.1.4.39 min()

```
template<typename T >
constexpr T gko::min (
    const T & x,
    const T & y ) [inline], [constexpr]
```

Returns the smaller of the arguments.

Template Parameters

<i>T</i>	type of the arguments
----------	-----------------------

Parameters

<i>x</i>	first argument
<i>y</i>	second argument

Returns

$$x \leq y ? x : y$$

Referenced by `gko::matrix::Csr< ValueType, IndexType >::load_balance::clac_size()`.

46.1.4.40 mixed_precision_dispatch()

```
template<typename ValueType , typename Function >
void gko::mixed_precision_dispatch (
    Function fn,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with each given argument `LinOp` converted into `matrix::Dense<ValueType>` as parameters.

If `GINKGO_MIXED_PRECISION` is defined, this means that the function will be called with its dynamic type as a static type, so the (templated/generic) function will be instantiated with all pairs of `Dense<ValueType>` and `Dense<next_precision_base<ValueType>>` parameter types, and the appropriate overload will be called based on the dynamic type of the parameter.

If `GINKGO_MIXED_PRECISION` is not defined, it will behave exactly like `precision_dispatch`.

Parameters

<i>fn</i>	the given function. It will be called with one const and one non-const <code>matrix::Dense<...></code> parameter based on the dynamic type of the inputs (<code>GINKGO_MIXED_PRECISION</code>) or of type <code>matrix::Dense<ValueType></code> (no <code>GINKGO_MIXED_PRECISION</code>).
<i>in</i>	The first parameter to be cast (<code>GINKGO_MIXED_PRECISION</code>) or converted (no <code>GINKGO_MIXED_PRECISION</code>) and used to call <code>fn</code> .
<i>out</i>	The second parameter to be cast (<code>GINKGO_MIXED_PRECISION</code>) or converted (no <code>GINKGO_MIXED_PRECISION</code>) and used to call <code>fn</code> .

Template Parameters

<i>ValueType</i>	the value type to use for the parameters of <code>fn</code> (no <code>GINKGO_MIXED_PRECISION</code>). With <code>GINKGO_MIXED_PRECISION</code> enabled, it only matters whether this type is complex or real.
<i>Function</i>	the function pointer, lambda or other functor type to call with the converted arguments.

46.1.4.41 mixed_precision_dispatch_real_complex()

```
template<typename ValueType , typename Function , std::enable_if_t< is_complex< ValueType >()> * = nullptr>
void gko::mixed_precision_dispatch_real_complex (
    Function fn,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with the given LinOps cast to their dynamic type `matrix::Dense<ValueType>*` as parameters.

If `ValueType` is real and both `in` and `out` are complex, uses `matrix::Dense::get_real_view()` to convert them into real matrices after precision conversion.

See also

[mixed_precision_dispatch\(\)](#)

46.1.4.42 `nan()` [1/2]

```
template<typename T >
constexpr std::enable_if_t<!is_complex_s<T>::value, T> gko::nan ( ) [inline], [constexpr]
```

Returns a quiet NaN of the given type.

Template Parameters

<i>T</i>	the type of the object
----------	------------------------

Returns

NaN.

Referenced by `nan()`.

46.1.4.43 `nan()` [2/2]

```
template<typename T >
constexpr std::enable_if_t<is_complex_s<T>::value, T> gko::nan ( ) [inline], [constexpr]
```

Returns a complex with both components quiet NaN.

Template Parameters

<i>T</i>	the type of the object
----------	------------------------

Returns

`complex{NaN, NaN}`.

References `nan()`.

46.1.4.44 one() [1/2]

```
template<typename T >
constexpr T gko::one ( ) [inline], [constexpr]
```

Returns the multiplicative identity for T.

Returns

the multiplicative identity for T

Referenced by `unit_root()`.

46.1.4.45 one() [2/2]

```
template<typename T >
constexpr T gko::one (
    const T & ) [inline], [constexpr]
```

Returns the multiplicative identity for T.

Returns

the multiplicative identity for T

Note

This version takes an unused reference argument to avoid complicated calls like `one<decltype(x)>()`. Instead, it allows `one(x)`.

46.1.4.46 operator"!="() [1/2]

```
bool gko::operator!= (
    const stopping_status & x,
    const stopping_status & y ) [inline], [noexcept]
```

Checks if two stopping statuses are different.

Parameters

<i>x</i>	a stopping status
<i>y</i>	a stopping status

Returns

true if and only if ! (x == y)

```
151 {
152     return x.data_ != y.data_;
153 }
```

46.1.4.47 operator"!=()" [2/2]

```
constexpr bool gko::operator!= (
    precision_reduction x,
    precision_reduction y ) [constexpr], [noexcept]
```

Checks if two [precision_reduction](#) encodings are different.

Parameters

<i>x</i>	an encoding
<i>y</i>	an encoding

Returns

true if and only if *x* and *y* are different encodings.

```
378 {
379     using st = precision_reduction::storage_type;
380     return static_cast<st>(x) != static_cast<st>(y);
381 }
```

46.1.4.48 operator<<() [1/2]

```
std::ostream& gko::operator<< (
    std::ostream & os,
    const version & ver ) [inline]
```

Prints version information to a stream.

Parameters

<i>os</i>	output stream
<i>ver</i>	version structure

Returns

OS

```
102 {
103     os << ver.major << "." << ver.minor << "." << ver.patch;
104     if (ver.tag) {
105         os << " (" << ver.tag << ")";
106     }
107     return os;
```

```
108 }
```

References `gko::version::major`, `gko::version::minor`, `gko::version::patch`, and `gko::version::tag`.

46.1.4.49 `operator<<()` [2/2]

```
std::ostream& gko::operator<< (
    std::ostream & os,
    const version_info & ver_info )
```

Prints library version information in human-readable format to a stream.

Parameters

<code>os</code>	output stream
<code>ver_info</code>	version information

Returns

OS

46.1.4.50 `operator==()` [1/2]

```
bool gko::operator==(
    const stopping_status & x,
    const stopping_status & y ) [inline], [noexcept]
```

Checks if two stopping statuses are equivalent.

Parameters

<code>x</code>	a stopping status
<code>y</code>	a stopping status

Returns

true if and only if both `x` and `y` have the same mask and converged and finalized state

46.1.4.51 `operator==()` [2/2]

```
constexpr bool gko::operator==(
    precision_reduction x,
    precision_reduction y ) [constexpr], [noexcept]
```

Checks if two `precision_reduction` encodings are equal.

Parameters

<i>x</i>	an encoding
<i>y</i>	an encoding

Returns

true if and only if *x* and *y* are the same encodings

46.1.4.52 pi()

```
template<typename T >
constexpr T gko::pi ( ) [inline], [constexpr]
```

Returns the value of pi.

Template Parameters

<i>T</i>	the value type to return
----------	--------------------------

46.1.4.53 precision_dispatch()

```
template<typename ValueType , typename Function , typename... Args>
void gko::precision_dispatch (
    Function fn,
    Args *... linops )
```

Calls the given function with each given argument LinOp temporarily converted into `matrix::Dense<ValueType>` as parameters.

Parameters

<i>fn</i>	the given function. It will be passed one (potentially const) <code>matrix::Dense<ValueType>*</code> parameter per parameter in the parameter pack <i>linops</i> .
<i>linops</i>	the given arguments to be converted and passed on to <i>fn</i> .

Template Parameters

<i>ValueType</i>	the value type to use for the parameters of <i>fn</i> .
<i>Function</i>	the function pointer, lambda or other functor type to call with the converted arguments.
<i>Args</i>	the argument type list.

46.1.4.54 precision_dispatch_real_complex() [1/3]

```
template<typename ValueType , typename Function >
void gko::precision_dispatch_real_complex (
    Function fn,
    const LinOp * alpha,
    const LinOp * in,
    const LinOp * beta,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `matrix::Dense<ValueType>*` as parameters.

If `ValueType` is real and both `in` and `out` are complex, uses `matrix::Dense::get_real_view()` to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.1.4.55 precision_dispatch_real_complex() [2/3]

```
template<typename ValueType , typename Function >
void gko::precision_dispatch_real_complex (
    Function fn,
    const LinOp * alpha,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `matrix::Dense<ValueType>*` as parameters.

If `ValueType` is real and both `in` and `out` are complex, uses [matrix::Dense::create_real_view\(\)](#) to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.1.4.56 precision_dispatch_real_complex() [3/3]

```
template<typename ValueType , typename Function >
void gko::precision_dispatch_real_complex (
    Function fn,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `matrix::Dense<ValueType>*` as parameters.

If `ValueType` is real and both input vectors are complex, uses [matrix::Dense::create_real_view\(\)](#) to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.1.4.57 read()

```
template<typename MatrixType , typename StreamType , typename... MatrixArgs>
std::unique_ptr<MatrixType> gko::read (
    StreamType && is,
    MatrixArgs &&... args ) [inline]
```

Reads a matrix stored in matrix market format from an input stream.

Template Parameters

<i>MatrixType</i>	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from disk.
<i>StreamType</i>	type of stream used to write the data to
<i>MatrixArgs</i>	additional argument types passed to MatrixType constructor

Parameters

<i>is</i>	input stream from which to read the data
<i>args</i>	additional arguments passed to MatrixType constructor

Returns

A MatrixType LinOp filled with data from filename

References read_raw().

Referenced by gko::ReadableFromMatrixData< ValueType, int32 >::read().

46.1.4.58 read_binary()

```
template<typename MatrixType , typename StreamType , typename... MatrixArgs>
std::unique_ptr<MatrixType> gko::read_binary (
    StreamType && is,
    MatrixArgs &&... args ) [inline]
```

Reads a matrix stored in binary format from an input stream.

Template Parameters

<i>MatrixType</i>	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from disk.
<i>StreamType</i>	type of stream used to write the data to
<i>MatrixArgs</i>	additional argument types passed to MatrixType constructor

Parameters

<i>is</i>	input stream from which to read the data
<i>args</i>	additional arguments passed to MatrixType constructor

Returns

A MatrixType LinOp filled with data from filename

References [read_binary_raw\(\)](#).

46.1.4.59 read_binary_raw()

```
template<typename ValueType = default_precision, typename IndexType = int32>
matrix_data<ValueType, IndexType> gko::read_binary_raw (
    std::istream & is )
```

Reads a matrix stored in Ginkgo's binary matrix format from an input stream.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

The binary format has the following structure (in system endianness):

1. A 32 byte header consisting of 4 uint64_t values: magic = GINKGO__: The highest two bytes stand for value and index type. value type: S (float), D (double), C (complex<float>), Z(complex<double>) index type: I (int32), L (int64) num_rows: Number of rows num_cols: Number of columns num_entries: Number of (row, column, value) tuples to follow
2. Following are num_entries blocks of size sizeof(IndexType) * 2 + sizeof(ValueType). Each consists of a row index stored as IndexType, followed by a column index stored as IndexType and a value stored as ValueType.

Template Parameters

<i>ValueType</i>	type of matrix values
<i>IndexType</i>	type of matrix indexes

Parameters

<i>is</i>	input stream from which to read the data
-----------	--

Returns

A [matrix_data](#) structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using [gko::read_binary](#) instead.

Referenced by [read_binary\(\)](#).

46.1.4.60 read_generic()

```
template<typename MatrixType , typename StreamType , typename... MatrixArgs>
std::unique_ptr<MatrixType> gko::read_generic (
    StreamType && is,
    MatrixArgs &&... args ) [inline]
```

Reads a matrix stored either in binary or matrix market format from an input stream.

Template Parameters

<i>MatrixType</i>	a ReadableFromMatrixData LinOp type used to store the matrix once it's been read from disk.
<i>StreamType</i>	type of stream used to write the data to
<i>MatrixArgs</i>	additional argument types passed to MatrixType constructor

Parameters

<i>is</i>	input stream from which to read the data
<i>args</i>	additional arguments passed to MatrixType constructor

Returns

A MatrixType LinOp filled with data from filename

References `read_generic_raw()`.

46.1.4.61 read_generic_raw()

```
template<typename ValueType = default_precision, typename IndexType = int32>
matrix\_data<ValueType, IndexType> gko::read_generic_raw (
    std::istream & is )
```

Reads a matrix stored in either binary or matrix market format from an input stream.

Template Parameters

<i>ValueType</i>	type of matrix values
<i>IndexType</i>	type of matrix indexes

Parameters

<i>is</i>	input stream from which to read the data
-----------	--

Returns

A [matrix_data](#) structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using [gko::read_generic](#) instead.

Referenced by `read_generic()`.

46.1.4.62 read_raw()

```
template<typename ValueType = default_precision, typename IndexType = int32>
matrix_data<ValueType, IndexType> gko::read_raw (
    std::istream & is )
```

Reads a matrix stored in matrix market format from an input stream.

Template Parameters

<i>ValueType</i>	type of matrix values
<i>IndexType</i>	type of matrix indexes

Parameters

<i>is</i>	input stream from which to read the data
-----------	--

Returns

A [matrix_data](#) structure containing the matrix. The nonzero elements are sorted in lexicographic order of their (row, column) indexes.

Note

This is an advanced routine that will return the raw matrix data structure. Consider using [gko::read](#) instead.

Referenced by `read()`.

46.1.4.63 real()

```
template<typename T >
constexpr auto gko::real (
    const T & x ) [inline], [constexpr]
```

Returns the real part of the object.

Template Parameters

<i>T</i>	type of the object
----------	--------------------

Parameters

<i>x</i>	the object
----------	------------

Returns

real part of the object (by default, the object itself)

Referenced by `squared_norm()`.

46.1.4.64 `reduce_add()` [1/2]

```
template<typename ValueType >
void gko::reduce_add (
    const array< ValueType > & input_arr,
    array< ValueType > & result )
```

Reduce (sum) the values in the array.

Template Parameters

<i>The</i>	type of the input data
------------	------------------------

Parameters

<i>in</i>	<i>input_arr</i>	the input array to be reduced
<i>in, out</i>	<i>result</i>	the reduced value. The result is written into the first entry and the value in the first entry is used as the initial value for the reduce.

46.1.4.65 `reduce_add()` [2/2]

```
template<typename ValueType >
ValueType gko::reduce_add (
    const array< ValueType > & input_arr,
    const ValueType init_val = 0 )
```

Reduce (sum) the values in the array.

Template Parameters

<i>The</i>	type of the input data
------------	------------------------

Parameters

<i>in</i>	<i>input_arr</i>	the input array to be reduced
<i>in</i>	<i>init_val</i>	the initial value

Returns

the reduced value

46.1.4.66 round_down()

```
template<typename T >
constexpr reduce\_precision<T> gko::round_down (
    T val ) [inline], [constexpr]
```

Reduces the precision of the input parameter.

Template Parameters

<i>T</i>	the original precision
----------	------------------------

Parameters

<i>val</i>	the value to round down
------------	-------------------------

Returns

the rounded down value

46.1.4.67 round_up()

```
template<typename T >
constexpr increase\_precision<T> gko::round_up (
    T val ) [inline], [constexpr]
```

Increases the precision of the input parameter.

Template Parameters

<i>T</i>	the original precision
----------	------------------------

Parameters

<i>val</i>	the value to round up
------------	-----------------------

Returns

the rounded up value

46.1.4.68 safe_divide()

```
template<typename T >
T gko::safe_divide (
    T a,
    T b ) [inline]
```

Computes the quotient of the given parameters, guarding against division by zero.

Template Parameters

<i>T</i>	value type of the parameters
----------	------------------------------

Parameters

<i>a</i>	the dividend
<i>b</i>	the divisor

Returns

the value of a / b if b is non-zero, zero otherwise.

46.1.4.69 share()

```
template<typename OwingPointer >
detail::shared_type<OwingPointer> gko::share (
    OwingPointer && p ) [inline]
```

Marks the object pointed to by p as shared.

Effectively converts a pointer with ownership to `std::shared_ptr`.

Template Parameters

<i>OwingPointer</i>	type of pointer with ownership to the object (has to be a smart pointer)
---------------------	--

Parameters

<i>p</i>	a pointer to the object. It must be a temporary or explicitly marked movable (rvalue reference).
----------	--

Note

The original pointer `p` becomes invalid after this call.

Referenced by `gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose()`, `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose()`, `gko::preconditioner::lc< LSolverType, IndexType >::transpose()`, and `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose()`.

46.1.4.70 squared_norm()

```
template<typename T >
constexpr auto gko::squared_norm (
    const T & x ) -> decltype(real(conj(x) * x))    [inline], [constexpr]
```

Returns the squared norm of the object.

Template Parameters

<i>T</i>	type of the object.
----------	---------------------

Returns

The squared norm of the object.

References `conj()`, and `real()`.

46.1.4.71 transpose() [1/2]

```
template<typename DimensionType >
batch_dim<2, DimensionType> gko::transpose (
    const batch_dim< 2, DimensionType > & input ) [inline]
```

Returns a `batch_dim` object with its dimensions swapped for batched operators.

Template Parameters

<i>DimensionType</i>	datatype used to represent each dimension
----------------------	---

Parameters

<i>dimensions</i>	original object
-------------------	-----------------

Returns

a `batch_dim` object with dimensions swapped

```

121 {
122     return batch_dim<2, DimensionType>(input.get_num_batch_items(),
123                                     transpose(input.get_common_size()));
124 }
```

References `gko::batch_dim< Dimensionality, DimensionType >::get_common_size()`, and `gko::batch_dim< Dimensionality, DimensionType >::get_num_batch_items()`.

Referenced by `gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose()`, `gko::preconditioner::ilu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose()`, `gko::preconditioner::lc< LSolverType, IndexType >::transpose()`, and `gko::preconditioner::ilu< LSolverType, USolverType, ReverseApply, IndexType >::transpose()`.

46.1.4.72 transpose() [2/2]

```

template<typename DimensionType >
constexpr dim<2, DimensionType> gko::transpose (
    const dim< 2, DimensionType > & dimensions ) [inline], [constexpr], [noexcept]
```

Returns a `dim<2>` object with its dimensions swapped.

Template Parameters

<i>DimensionType</i>	datatype used to represent each dimension
----------------------	---

Parameters

<i>dimensions</i>	original object
-------------------	-----------------

Returns

a `dim<2>` object with its dimensions swapped

```

253 {
254     return {dimensions[1], dimensions[0]};
255 }
```

46.1.4.73 unit_root()

```

template<typename T >
constexpr std::complex<remove_complex<T> > gko::unit_root (
    int64 n,
    int64 k = 1 ) [inline], [constexpr]
```

Returns the value of $\exp(2 * \pi * i * k / n)$, i.e.

an n th root of unity.

Parameters

n	the denominator of the argument
k	the numerator of the argument. Defaults to 1.

Template Parameters

T	the corresponding real value type.
-----	------------------------------------

References one().

46.1.4.74 with_matrix_type()

```
template<template< typename, typename > class MatrixType, typename... Args>
auto gko::with_matrix_type (
    Args &&... create_args )
```

This function returns a type that delays a call to MatrixType::create.

It can be used to set the used value and index type, as well as the executor at a later stage.

For example, the following code creates first a temporary object, which is then used later to construct an operator of the previously defined base type:

```
auto type = gko::with_matrix_type<gko::matrix::Csr>();
...
std::unique_ptr<LinOp> concrete_op
if(flag1){
    concrete_op = type.template create<double, int>(exec);
} else {
    concrete_op = type.template create<float, int>(exec);
}
```

Note

This is mainly a helper function to specify the local matrix type for a [gko::experimental::distributed::Matrix](#) more easily.

Template Parameters

<i>MatrixType</i>	A template type that accepts two types, the first one will be set to the value type, the second one to the index type.
<i>Args</i>	Types of the arguments passed to MatrixType::create.

Parameters

<i>create_args</i>	arguments that will be forwarded to MatrixType::create
--------------------	--

Returns

A type with a function `create<value_type, index_type>(executor)`.

```

126 {
127     return detail::MatrixTypeBuilderFromValueAndIndex<MatrixType, Args...>{
128         std::forward_as_tuple(create_args...)};
129 }

```

46.1.4.75 write()

```

template<typename MatrixPtrType , typename StreamType >
void gko::write (
    StreamType && os,
    MatrixPtrType && matrix,
    layout_type layout = detail::mtx_io_traits< std::remove_cv_t<detail::pointee<MatrixPtrType>>><
::default_layout ) [inline]

```

Writes a matrix into an output stream in matrix market format.

Template Parameters

<i>MatrixPtrType</i>	a (smart or raw) pointer to a WritableToMatrixData object providing data to be written.
<i>StreamType</i>	type of stream used to write the data to

Parameters

<i>os</i>	output stream where the data is to be written
<i>matrix</i>	the matrix to write
<i>layout</i>	the layout used in the output

References `write_raw()`.

46.1.4.76 write_binary()

```

template<typename MatrixPtrType , typename StreamType >
void gko::write_binary (
    StreamType && os,
    MatrixPtrType && matrix ) [inline]

```

Writes a matrix into an output stream in binary format.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

Template Parameters

<i>MatrixPtrType</i>	a (smart or raw) pointer to a WritableToMatrixData object providing data to be written.
<i>StreamType</i>	type of stream used to write the data to

Parameters

<i>os</i>	output stream where the data is to be written
<i>matrix</i>	the matrix to write

References `write_binary_raw()`.

46.1.4.77 `write_binary_raw()`

```
template<typename ValueType , typename IndexType >
void gko::write_binary_raw (
    std::ostream & os,
    const matrix\_data< ValueType, IndexType > & data )
```

Writes a [matrix_data](#) structure to a stream in binary format.

Note that this format depends on the processor's endianness, so files from a big endian processor can't be read from a little endian processor and vice-versa.

Template Parameters

<i>ValueType</i>	type of matrix values
<i>IndexType</i>	type of matrix indexes

Parameters

<i>os</i>	output stream where the data is to be written
<i>data</i>	the matrix data to write

Note

This is an advanced routine that writes the raw matrix data structure. If you are trying to write an existing matrix, consider using [gko::write_binary](#) instead.

Referenced by `write_binary()`.

46.1.4.78 `write_raw()`

```
template<typename ValueType , typename IndexType >
void gko::write_raw (
    std::ostream & os,
    const matrix\_data< ValueType, IndexType > & data,
    layout\_type layout = layout\_type::coordinate )
```

Writes a [matrix_data](#) structure to a stream in matrix market format.

Template Parameters

<i>ValueType</i>	type of matrix values
<i>IndexType</i>	type of matrix indexes

Parameters

<i>os</i>	output stream where the data is to be written
<i>data</i>	the matrix data to write
<i>layout</i>	the layout used in the output

Note

This is an advanced routine that writes the raw matrix data structure. If you are trying to write an existing matrix, consider using [gko::write](#) instead.

Referenced by `write()`.

46.1.4.79 zero() [1/2]

```
template<typename T >
constexpr T gko::zero ( ) [inline], [constexpr]
```

Returns the additive identity for T.

Returns

additive identity for T

Referenced by `gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::strategy_type()`.

46.1.4.80 zero() [2/2]

```
template<typename T >
constexpr T gko::zero (
    const T & ) [inline], [constexpr]
```

Returns the additive identity for T.

Returns

additive identity for T

Note

This version takes an unused reference argument to avoid complicated calls like `zero<decltype(x)>()`. Instead, it allows `zero(x)`.

46.2 gko::accessor Namespace Reference

The accessor namespace.

Classes

- class [row_major](#)

A [row_major](#) accessor is a bridge between a range and the row-major memory layout.

46.2.1 Detailed Description

The accessor namespace.

46.3 gko::batch::log Namespace Reference

The logger namespace .

Classes

- class [BatchConvergence](#)

Logs the final residuals and iteration counts for a batch solver.

46.3.1 Detailed Description

The logger namespace .

[Logging](#)

46.4 gko::experimental::distributed Namespace Reference

The distributed namespace.

Namespaces

- [preconditioner](#)

The Preconditioner namespace.

Classes

- class [DistributedBase](#)
A base class for distributed objects.
- struct [index_map](#)
This class defines mappings between global and local indices.
- class [Matrix](#)
The [Matrix](#) class defines a (MPI-)distributed matrix.
- class [Partition](#)
Represents a partition of a range of indices [0, size) into a disjoint set of parts.
- class [Vector](#)
[Vector](#) is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

Typedefs

- using [comm_index_type](#) = int
Index type for enumerating processes in a distributed application.

Enumerations

- enum [index_space](#) { [index_space::local](#), [index_space::non_local](#), [index_space::combined](#) }
Index space classification for the locally stored indices.
- enum [assembly_mode](#)
[assembly_mode](#) defines how the [read_distributed](#) function of the distributed matrix treats non-local indices in the (device_)matrix_data:

Functions

- template<typename ValueType >
gko::detail::temporary_conversion< [Vector](#)< ValueType > > [make_temporary_conversion](#) (LinOp *matrix)
Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector<Value↔Type>.
- template<typename ValueType >
gko::detail::temporary_conversion< const [Vector](#)< ValueType > > [make_temporary_conversion](#) (const LinOp *matrix)
Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector<Value↔Type>.
- template<typename ValueType , typename Function , typename... Args>
void [precision_dispatch](#) (Function fn, Args *... linops)
Calls the given function with each given argument LinOp temporarily converted into experimental::distributed::↔Vector<ValueType> as parameters.
- template<typename ValueType , typename Function >
void [precision_dispatch_real_complex](#) (Function fn, const LinOp *in, LinOp *out)
Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector<Value↔Type> as parameters.*
- template<typename ValueType , typename Function >
void [precision_dispatch_real_complex](#) (Function fn, const LinOp *alpha, const LinOp *in, LinOp *out)
Calls the given function with the given LinOps temporarily converted to experimental::distributed::Vector<Value↔Type> as parameters.*

- `template<typename ValueType , typename Function >`
`void precision\_dispatch\_real\_complex (Function fn, const LinOp *alpha, const LinOp *in, const LinOp *beta, LinOp *out)`
Calls the given function with the given LinOps temporarily converted to `experimental::distributed::Vector<ValueType>*` as parameters.
- `template<typename ValueType , typename LocalIndexType , typename GlobalIndexType >`
`device\_matrix\_data< ValueType, GlobalIndexType > assemble\_rows\_from\_neighbors (mpi::communicator comm, const device\_matrix\_data< ValueType, GlobalIndexType > &input, ptr\_param< const Partition< LocalIndexType, GlobalIndexType >> partition)`
Assembles [device_matrix_data](#) entries owned by this MPI rank from other ranks and communicates entries located on this MPI rank owned by other ranks to their respective owners.
- `template<typename LocalIndexType , typename GlobalIndexType >`
`std::unique_ptr< Partition< LocalIndexType, GlobalIndexType > > build\_partition\_from\_local\_range (std::shared_ptr< const Executor > exec, mpi::communicator comm, span local_range)`
Builds a partition from a local range.
- `template<typename LocalIndexType , typename GlobalIndexType >`
`std::unique_ptr< Partition< LocalIndexType, GlobalIndexType > > build\_partition\_from\_local\_size (std::shared_ptr< const Executor > exec, mpi::communicator comm, size\_type local_size)`
Builds a partition from a local size.

46.4.1 Detailed Description

The distributed namespace.

46.4.2 Typedef Documentation

46.4.2.1 `comm_index_type`

```
using gko::experimental::distributed::comm\_index\_type = typedef int
```

Index type for enumerating processes in a distributed application.

Conforms to the MPI C interface of e.g. MPI rank or size

46.4.3 Enumeration Type Documentation

46.4.3.1 `assembly_mode`

```
enum gko::experimental::distributed::assembly\_mode [strong]
```

`assembly_mode` defines how the `read_distributed` function of the distributed matrix treats non-local indices in the (device_)[matrix_data](#):

- `communicate` communicates the overlap between ranks and adds up all local contributions. Indices smaller than 0 or larger than the global size of the matrix are ignored.
- `local_only` does not communicate any overlap but ignores all non-local indices.

46.4.3.2 index_space

```
enum gko::experimental::distributed::index_space [strong]
```

Index space classification for the locally stored indices.

The definitions of the enum values is clarified in [index_map](#).

Enumerator

local	indices that are locally owned
non_local	indices that are owned by other processes
combined	both local and non_local indices

```
23         {
24     local,
25     non_local,
26     combined
27 };
```

46.4.4 Function Documentation

46.4.4.1 assemble_rows_from_neighbors()

```
template<typename ValueType , typename LocalIndexType , typename GlobalIndexType >
device_matrix_data<ValueType, GlobalIndexType> gko::experimental::distributed::assemble_rows←
_from_neighbors (
    mpi::communicator comm,
    const device_matrix_data< ValueType, GlobalIndexType > & input,
    ptr_param< const Partition< LocalIndexType, GlobalIndexType >> partition )
```

Assembles [device_matrix_data](#) entries owned by this MPI rank from other ranks and communicates entries located on this MPI rank owned by other ranks to their respective owners.

This can be useful e.g. in a finite element code where each rank assembles a local contribution to a global system matrix and the global matrix has to be assembled by summing up the local contributions on rank boundaries. The partition used is only relevant for row ownership.

Parameters

<i>comm</i>	the communicator used to assemble the global matrix.
<i>input</i>	the device_matrix_data structure.
<i>partition</i>	the partition used to determine row ownership.

Returns

the globally assembled [device_matrix_data](#) structure for this MPI rank.

46.4.4.2 build_partition_from_local_range()

```
template<typename LocalIndexType , typename GlobalIndexType >
std::unique_ptr<Partition<LocalIndexType, GlobalIndexType> > gko::experimental::distributed←
::build_partition_from_local_range (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    span local_range )
```

Builds a partition from a local range.

Parameters

<i>exec</i>	the Executor on which the partition should be built.
<i>comm</i>	the communicator used to determine the global partition.
<i>local_range</i>	the start and end indices of the local range.

Warning

This throws, if the resulting partition would contain gaps. That means that for a partition of size n every local range $r_i = [s_i, e_i)$ either $s_i \neq 0$ and another local range $r_j = [s_j, e_j = s_i)$ exists, or $e_i \neq n$ and another local range $r_j = [s_j = e_i, e_j)$ exists.

Returns

a [Partition](#) where each range has the individual `local_start` and `local_ends`.

46.4.4.3 build_partition_from_local_size()

```
template<typename LocalIndexType , typename GlobalIndexType >
std::unique_ptr<Partition<LocalIndexType, GlobalIndexType> > gko::experimental::distributed←
::build_partition_from_local_size (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    size_type local_size )
```

Builds a partition from a local size.

Parameters

<i>exec</i>	the Executor on which the partition should be built.
<i>comm</i>	the communicator used to determine the global partition.
<i>local_range</i>	the number of the locally owned indices

Returns

a [Partition](#) where each range has the specified local size. More specifically, if this is called on process i with local size s_i , then the range i has size s_i , and range $r_i = [start, start + s_i)$, where $start = \sum_{j=0}^{i-1} s_j$.

46.4.4.4 make_temporary_conversion() [1/2]

```
template<typename ValueType >
gko::detail::temporary_conversion<const Vector<ValueType> > gko::experimental::distributed::
::make_temporary_conversion (
    const LinOp * matrix )
```

Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector<ValueType>.

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by next_precision_base<...> starting from the ValueType template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

<i>matrix</i>	the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type experimental::distributed::Vector<ValueType>, otherwise it will be converted to this type if possible.
---------------	---

Returns

a detail::temporary_conversion pointing to the (potentially converted) object.

Exceptions

<i>NotSupported</i>	if the input matrix cannot be converted to experimental::distributed::Vector<ValueType>
---------------------	---

Template Parameters

<i>ValueType</i>	the value type into whose associated Vector type to convert the input LinOp.
------------------	--

46.4.4.5 make_temporary_conversion() [2/2]

```
template<typename ValueType >
gko::detail::temporary_conversion<Vector<ValueType> > gko::experimental::distributed::make_
temporary_conversion (
    LinOp * matrix )
```

Convert the given LinOp from experimental::distributed::Vector<...> to experimental::distributed::Vector<ValueType>.

The conversion tries to convert the input LinOp to all Dense types with value type recursively reachable by next_precision_base<...> starting from the ValueType template parameter. This means that all real-to-real and complex-to-complex conversions for default precisions are being considered. If the input matrix is non-const, the contents of the modified converted object will be converted back to the input matrix when the returned object is destroyed. This may lead to a loss of precision!

Parameters

<i>matrix</i>	the input matrix which is supposed to be converted. It is wrapped unchanged if it is already of type <code>experimental::distributed::Vector<ValueType></code> , otherwise it will be converted to this type if possible.
---------------	---

Returns

a `detail::temporary_conversion` pointing to the (potentially converted) object.

Exceptions

<i>NotSupported</i>	if the input matrix cannot be converted to <code>experimental::distributed::Vector<ValueType></code>
---------------------	--

Template Parameters

<i>ValueType</i>	the value type into whose associated Vector type to convert the input LinOp.
------------------	--

46.4.4.6 precision_dispatch()

```
template<typename ValueType , typename Function , typename... Args>
void gko::experimental::distributed::precision_dispatch (
    Function fn,
    Args *... linops )
```

Calls the given function with each given argument LinOp temporarily converted into `experimental::distributed::Vector<ValueType>` as parameters.

Parameters

<i>fn</i>	the given function. It will be passed one (potentially const) <code>experimental::distributed::Vector<ValueType>*</code> parameter per parameter in the parameter pack <code>linops</code> .
<i>linops</i>	the given arguments to be converted and passed on to <code>fn</code> .

Template Parameters

<i>ValueType</i>	the value type to use for the parameters of <code>fn</code> .
<i>Function</i>	the function pointer, lambda or other functor type to call with the converted arguments.
<i>Args</i>	the argument type list.

46.4.4.7 precision_dispatch_real_complex() [1/3]

```
template<typename ValueType , typename Function >
void gko::experimental::distributed::precision_dispatch_real_complex (
```

```
Function fn,
const LinOp * alpha,
const LinOp * in,
const LinOp * beta,
LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `experimental::distributed::Vector<ValueType>*` as parameters.

If `ValueType` is real and both input vectors are complex, uses `experimental::distributed::Vector::get_real_view()` to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.4.4.8 precision_dispatch_real_complex() [2/3]

```
template<typename ValueType , typename Function >
void gko::experimental::distributed::precision_dispatch_real_complex (
    Function fn,
    const LinOp * alpha,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `experimental::distributed::Vector<ValueType>*` as parameters.

If `ValueType` is real and both input vectors are complex, uses `experimental::distributed::Vector::get_real_view()` to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.4.4.9 precision_dispatch_real_complex() [3/3]

```
template<typename ValueType , typename Function >
void gko::experimental::distributed::precision_dispatch_real_complex (
    Function fn,
    const LinOp * in,
    LinOp * out )
```

Calls the given function with the given LinOps temporarily converted to `experimental::distributed::Vector<ValueType>*` as parameters.

If `ValueType` is real and both input vectors are complex, uses `experimental::distributed::Vector::get_real_view()` to convert them into real matrices after precision conversion.

See also

[precision_dispatch\(\)](#)

46.5 gko::experimental::distributed::preconditioner Namespace Reference

The Preconditioner namespace.

Classes

- class [Schwarz](#)

A [Schwarz](#) preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

46.5.1 Detailed Description

The Preconditioner namespace.

46.6 gko::experimental::mpi Namespace Reference

The mpi namespace, contains wrapper for many MPI functions.

Classes

- class [communicator](#)
A thin wrapper of MPI_Comm that supports most MPI calls.
- class [contiguous_type](#)
A move-only wrapper for a contiguous MPI_Datatype.
- class [environment](#)
Class that sets up and finalizes the MPI environment.
- class [request](#)
The request class is a light, move-only wrapper around the MPI_Request handle.
- struct [status](#)
The status struct is a light wrapper around the MPI_Status struct.
- struct [type_impl](#)
A struct that is used to determine the MPI_Datatype of a specified type.
- class [window](#)
This class wraps the MPI_Window class with RAII functionality.

Enumerations

- enum [thread_type](#)
This enum specifies the threading type to be used when creating an MPI environment.

Functions

- constexpr bool `is_gpu_aware` ()
Return if GPU aware functionality is available.
- int `map_rank_to_device_id` (MPI_Comm comm, int num_devices)
Maps each MPI rank to a single device id in a round robin manner.
- std::vector< `status` > `wait_all` (std::vector< `request` > &req)
Allows a rank to wait on multiple request handles.
- bool `requires_host_buffer` (const std::shared_ptr< const `Executor` > &exec, const `communicator` &comm)
Checks if the combination of `Executor` and communicator requires passing MPI buffers from the host memory.
- double `get_walltime` ()
Get the rank in the communicator of the calling process.

46.6.1 Detailed Description

The mpi namespace, contains wrapper for many MPI functions.

46.6.2 Function Documentation

46.6.2.1 `get_walltime()`

```
double gko::experimental::mpi::get_walltime ( ) [inline]
```

Get the rank in the communicator of the calling process.

Parameters

<code>comm</code>	the communicator
-------------------	------------------

```
1502 { return MPI_Wtime(); }
```

46.6.2.2 `map_rank_to_device_id()`

```
int gko::experimental::mpi::map_rank_to_device_id (
    MPI_Comm comm,
    int num_devices )
```

Maps each MPI rank to a single device id in a round robin manner.

Parameters

<code>comm</code>	used to determine the node-local rank, if no suitable environment variable is available.
<code>num_devices</code>	the number of devices per node.

Returns

device id that this rank should use.

46.6.2.3 wait_all()

```
std::vector<status> gko::experimental::mpi::wait_all (
    std::vector< request > & req ) [inline]
```

Allows a rank to wait on multiple request handles.

Parameters

<i>req</i>	The vector of request handles to be waited on.
------------	--

Returns

status The vector of status objects that can be queried.

46.7 gko::experimental::reorder Namespace Reference

The Reorder namespace.

Classes

- class [Amd](#)
Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.
- class [Mc64](#)
MC64 is an algorithm for permuting large entries to the diagonal of a sparse matrix.
- class [Rcm](#)
Rcm (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.
- class [ScaledReordered](#)
Provides an interface to wrap reorderings like [Rcm](#) and diagonal scaling like equilibration around a LinOp like e.g.

Enumerations

- enum [mc64_strategy](#)
Strategy defining the goal of the MC64 reordering.

46.7.1 Detailed Description

The Reorder namespace.

46.7.2 Enumeration Type Documentation

46.7.2.1 mc64_strategy

```
enum gko::experimental::reorder::mc64_strategy [strong]
```

Strategy defining the goal of the MC64 reordering.

max_diagonal_product aims at maximizing the product of absolute diagonal entries. max_diag_sum aims at maximizing the sum of absolute values for the diagonal entries.

```
44 { max_diagonal_product, max_diagonal_sum };
```

46.8 gko::factorization Namespace Reference

The Factorization namespace.

Classes

- class [lc](#)
Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.
- class [llu](#)
Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.
- class [Parlc](#)
ParlC is an incomplete Cholesky factorization which is computed in parallel.
- class [Parlct](#)
ParlCT is an incomplete threshold-based Cholesky factorization which is computed in parallel.
- class [Parllu](#)
ParlLU is an incomplete LU factorization which is computed in parallel.
- class [Parllut](#)
ParlLUT is an incomplete threshold-based LU factorization which is computed in parallel.

Enumerations

- enum [incomplete_algorithm](#)
An enum class for algorithm selection in the incomplete factorization.

46.8.1 Detailed Description

The Factorization namespace.

46.8.2 Enumeration Type Documentation

46.8.2.1 incomplete_algorithm

```
enum gko::factorization::incomplete_algorithm [strong]
```

An enum class for algorithm selection in the incomplete factorization.

`sparselib` is only available for CUDA, HIP, and reference. `syncfree` is Ginkgo's implementation by using the Lu/Cholesky factorization components with given sparsity.

```
19 { sparselib, syncfree };
```

46.9 gko::log Namespace Reference

The logger namespace .

Classes

- class [Convergence](#)
Convergence is a *Logger* which logs data strictly from the *criterion_check_completed* event.
- struct [criterion_data](#)
Struct representing *Criterion* related data.
- class [EnableLogging](#)
EnableLogging is a *mixin* which should be inherited by any class which wants to enable logging.
- struct [executor_data](#)
Struct representing *Executor* related data.
- struct [iteration_complete_data](#)
Struct representing iteration complete related data.
- struct [linop_data](#)
Struct representing *LinOp* related data.
- struct [linop_factory_data](#)
Struct representing *LinOp* factory related data.
- class [Loggable](#)
Loggable class is an interface which should be implemented by classes wanting to support logging.
- struct [operation_data](#)
Struct representing *Operator* related data.
- class [Papi](#)
Papi is a *Logger* which logs every event to the PAPI software.
- class [PerformanceHint](#)
PerformanceHint is a *Logger* which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.
- struct [polymorphic_object_data](#)
Struct representing *PolymorphicObject* related data.
- class [ProfilerHook](#)
This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.
- class [profiling_scope_guard](#)
Scope guard that annotates its scope with the provided profiler hooks.
- class [Record](#)
Record is a *Logger* which logs every event to an object.
- class [SolverProgress](#)
This Logger outputs the value of all scalar values (and potentially vectors) stored internally by the solver after each iteration.
- class [Stream](#)
Stream is a *Logger* which logs every event to a stream.

Enumerations

- enum `profile_event_category` {
`profile_event_category::memory`, `profile_event_category::operation`, `profile_event_category::object`,
`profile_event_category::linop`,
`profile_event_category::factory`, `profile_event_category::solver`, `profile_event_category::criterion`, `profile_event_category::user`,
`profile_event_category::internal` }

Categorization of logger events.

46.9.1 Detailed Description

The logger namespace .

The Logging namespace.

[Logging](#)

46.9.2 Enumeration Type Documentation

46.9.2.1 `profile_event_category`

enum `gko::log::profile_event_category` [strong]

Categorization of logger events.

Enumerator

<code>memory</code>	Memory allocation.
<code>operation</code>	Kernel execution and data movement.
<code>object</code>	PolymorphicObject events.
<code>linop</code>	LinOp events.
<code>factory</code>	LinOpFactory events.
<code>solver</code>	Solver events.
<code>criterion</code>	Stopping criterion events.
<code>user</code>	User-defined events.
<code>internal</code>	For development use.

```

22                                     {
23
24     memory,
25     operation,
26     object,
27     linop,
28     factory,
29     solver,
30     criterion,
31     user,
32     internal,
33 };
34
35
36
37
38
39
40
41

```

46.10 gko::matrix Namespace Reference

The matrix namespace.

Classes

- class [Coo](#)
COO stores a matrix in the coordinate matrix format.
- class [Csr](#)
CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).
- class [Dense](#)
Dense is a matrix format which explicitly stores all values of the matrix.
- class [Diagonal](#)
This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).
- class [Ell](#)
ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.
- class [Fbcsr](#)
Fixed-block compressed sparse row storage matrix format.
- class [Fft](#)
This LinOp implements a 1D Fourier matrix using the FFT algorithm.
- class [Fft2](#)
This LinOp implements a 2D Fourier matrix using the FFT algorithm.
- class [Fft3](#)
This LinOp implements a 3D Fourier matrix using the FFT algorithm.
- class [Hybrid](#)
HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.
- class [Identity](#)
This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).
- class [IdentityFactory](#)
This factory is a utility which can be used to generate [Identity](#) operators.
- class [Permutation](#)
Permutation is a matrix format that represents a permutation matrix, i.e.
- class [RowGatherer](#)
RowGatherer is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.
- class [ScaledPermutation](#)
ScaledPermutation is a matrix combining a permutation with scaling factors.
- class [Sellp](#)
SELL-P is a matrix format similar to ELL format.
- class [SparsityCsr](#)
SparsityCsr is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

Enumerations

- enum `permute_mode` : unsigned {
`permute_mode::none` = 0b000u, `permute_mode::rows` = 0b001u, `permute_mode::columns` = 0b010u,
`permute_mode::symmetric` = 0b011u,
`permute_mode::inverse` = 0b100u, `permute_mode::inverse_rows` = 0b101u, `permute_mode::inverse_columns`
= 0b110u, `permute_mode::inverse_symmetric` = 0b111u }

Specifies how a permutation will be applied to a matrix.

Functions

- `permute_mode operator|` (`permute_mode` a, `permute_mode` b)
Combines two permutation modes.
- `permute_mode operator&` (`permute_mode` a, `permute_mode` b)
Computes the intersection of two permutation modes.
- `permute_mode operator^` (`permute_mode` a, `permute_mode` b)
Computes the symmetric difference of two permutation modes.
- `std::ostream & operator<<` (`std::ostream &stream`, `permute_mode` mode)
Prints a permutation mode.

46.10.1 Detailed Description

The matrix namespace.

46.10.2 Enumeration Type Documentation

46.10.2.1 permute_mode

```
enum gko::matrix::permute_mode : unsigned [strong]
```

Specifies how a permutation will be applied to a matrix.

For the effect of the different permutation modes, see the following table.

mode	entry mapping	matrix representation
none	$A'(i, j) = A(i, j)$	$A' = A$
rows	$A'(i, j) = A(p[i], j)$	$A' = PA$
columns	$A'(i, j) = A(i, p[j])$	$A' = AP^T$
inverse_rows	$A'(p[i], j) = A(i, j)$	$A' = P^{-1}A$
inverse_columns	$A'(i, p[j]) = A(i, j)$	$A' = AP^{-T}$
symmetric	$A'(i, j) = A(p[i], p[j])$	$A' = PAP^T$
inverse_symmetric	$A'(p[i], p[j]) = A(i, j)$	$A' = P^{-1}AP^{-T}$

Enumerator

none	Neither rows nor columns will be permuted.
rows	The rows will be permuted.
columns	The columns will be permuted.
symmetric	The rows and columns will be permuted. This is equivalent to permute_mode::rows permute_mode::columns .
inverse	The permutation will be inverted before being applied.
inverse_rows	The rows will be permuted using the inverse permutation. This is equivalent to permute_mode::rows permute_mode::inverse .
inverse_columns	The columns will be permuted using the inverse permutation. This is equivalent to permute_mode::columns permute_mode::inverse .
inverse_symmetric	The rows and columns will be permuted using the inverse permutation. This is equivalent to permute_mode::symmetric permute_mode::inverse .

```

42                                     : unsigned {
44     none = 0b000u,
46     rows = 0b001u,
48     columns = 0b010u,
53     symmetric = 0b011u,
55     inverse = 0b100u,
60     inverse_rows = 0b101u,
65     inverse_columns = 0b110u,
70     inverse_symmetric = 0b111u
71 };

```

46.11 gko::multigrid Namespace Reference

The multigrid components namespace.

Classes

- class [EnableMultigridLevel](#)
The [EnableMultigridLevel](#) gives the default implementation of [MultigridLevel](#) with composition and provides `set_↔multigrid_level` function.
- class [FixedCoarsening](#)
[FixedCoarsening](#) is a very simple coarse grid generation algorithm.
- class [MultigridLevel](#)
This class represents two levels in a multigrid hierarchy.
- class [Pgm](#)
Parallel graph match ([Pgm](#)) is the aggregate method introduced in the paper [M](#).

46.11.1 Detailed Description

The multigrid components namespace.

46.12 gko::name_demangling Namespace Reference

The name demangling namespace.

Functions

- `template<typename T >`
`std::string get\_static\_type (const T &)`

*This function uses name demangling facilities to get the name of the static type (*T*) of the object passed in arguments.*

- `template<typename T >`
`std::string get\_dynamic\_type (const T &t)`

This function uses name demangling facilities to get the name of the dynamic type of the object passed in arguments.

46.12.1 Detailed Description

The name demangling namespace.

46.12.2 Function Documentation

46.12.2.1 `get_dynamic_type()`

```
template<typename T >
std::string gko::name_demangling::get_dynamic_type (
    const T & t )
```

This function uses name demangling facilities to get the name of the dynamic type of the object passed in arguments.

Template Parameters

<i>T</i>	the type of the object to demangle
----------	------------------------------------

Parameters

<i>t</i>	the object we get the dynamic type of
----------	---------------------------------------

```
73 {
74     return get_type_name(typeid(t));
75 }
```

46.12.2.2 `get_static_type()`

```
template<typename T >
std::string gko::name_demangling::get_static_type (
    const T & )
```

This function uses name demangling facilities to get the name of the static type (*T*) of the object passed in arguments.

Template Parameters

<i>T</i>	the type of the object to demangle
----------	------------------------------------

Parameters

<i>unused</i>	
---------------	--

46.13 gko::preconditioner Namespace Reference

The Preconditioner namespace.

Classes

- struct [block_interleaved_storage_scheme](#)
Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.
- class [GaussSeidel](#)
This class generates the Gauss-Seidel preconditioner.
- class [lc](#)
*The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).*
- class [llu](#)
The Incomplete LU (ILU) preconditioner solves the equation $LUx = b$ for a given lower triangular matrix L , an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).
- class [lsai](#)
The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A , lower triangular matrix L , upper triangular matrix U or symmetric positive (spd) matrix B .
- class [Jacobi](#)
A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.
- class [Sor](#)
This class generates the (S)SOR preconditioner.

Enumerations

- enum [isai_type](#)
This enum lists the types of the ISAI preconditioner.

46.13.1 Detailed Description

The Preconditioner namespace.

46.13.2 Enumeration Type Documentation

46.13.2.1 isai_type

```
enum gko::preconditioner::isai_type [strong]
```

This enum lists the types of the ISAI preconditioner.

ISAI can either be generated for a general square matrix, a lower triangular matrix, an upper triangular matrix or an spd matrix.

```
36 { lower, upper, general, spd };
```

46.14 gko::reorder Namespace Reference

The Reorder namespace.

Classes

- class [Rcm](#)
Rcm (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.
- class [ReorderingBase](#)
The [ReorderingBase](#) class is a base class for all the reordering algorithms.
- struct [ReorderingBaseArgs](#)
This struct is used to pass parameters to the `EnableDefaultReorderingBaseFactory::generate()` method.

Typedefs

- `template<typename IndexType = int32>`
`using ReorderingBaseFactory = AbstractFactory< ReorderingBase< IndexType >, ReorderingBaseArgs >`
Declares an Abstract Factory specialized for [ReorderingBases](#).
- `template<typename ConcreteFactory, typename ConcreteReorderingBase, typename ParametersType, typename IndexType = int32, typename PolymorphicBase = ReorderingBaseFactory<IndexType>>`
`using EnableDefaultReorderingBaseFactory = EnableDefaultFactory< ConcreteFactory, ConcreteReorderingBase, ParametersType, PolymorphicBase >`
This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of [ReorderingBaseFactory](#).

46.14.1 Detailed Description

The Reorder namespace.

The reordering namespace.

46.14.2 Typedef Documentation

46.14.2.1 EnableDefaultReorderingBaseFactory

```
template<typename ConcreteFactory, typename ConcreteReorderingBase, typename ParametersType,
, typename IndexType = int32, typename PolymorphicBase = ReorderingBaseFactory<IndexType>>
using gko::reorder::EnableDefaultReorderingBaseFactory = typedef EnableDefaultFactory<ConcreteFactory, ConcreteReorderingBase, ParametersType, PolymorphicBase>
```

This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of [ReorderingBaseFactory](#).

Template Parameters

<i>ConcreteFactory</i>	the concrete factory which is being implemented [CRTP parameter]
<i>ConcreteReorderingBase</i>	the concrete ReorderingBase type which this factory produces, needs to have a constructor which takes a const <code>ConcreteFactory *</code> , and a const <code>ReorderingBaseArgs *</code> as parameters.
<i>ParametersType</i>	a subclass of enable_parameters_type template which defines all of the parameters of the factory
<i>PolymorphicBase</i>	parent of <code>ConcreteFactory</code> in the polymorphic hierarchy, has to be a subclass of <code>ReorderingBaseFactory</code>

46.15 gko::solver Namespace Reference

The ginkgo Solve namespace.

Namespaces

- [multigrid](#)

The solver multigrid namespace.

Classes

- class [ApplyWithInitialGuess](#)
[ApplyWithInitialGuess](#) provides a way to give the input guess for apply function.
- class [Bicg](#)
BICG or the Biconjugate gradient method is a Krylov subspace solver.
- class [Bicgstab](#)
BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.
- class [CbGmres](#)
CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
- class [Cg](#)
CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
- class [Cgs](#)
CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.
- class [EnableApplyWithInitialGuess](#)
[EnableApplyWithInitialGuess](#) providing default operation for [ApplyWithInitialGuess](#) with correct validation and log.
- class [EnableIterativeBase](#)
A `LinOp` deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.
- class [EnablePreconditionable](#)
Mixin providing default operation for [Preconditionable](#) with correct value semantics.
- class [EnablePreconditionedIterativeSolver](#)
A `LinOp` implementing this interface stores a system matrix and stopping criterion factory.
- class [EnableSolverBase](#)
A `LinOp` deriving from this CRTP class stores a system matrix.

- class [Fcg](#)
FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.
- class [Gcr](#)
GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.
- class [Gmres](#)
GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.
- struct [has_with_criteria](#)
Helper structure to test if the Factory of SolverType has a function `with_criteria`.
- class [ldr](#)
IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.
- class [lr](#)
Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.
- class [IterativeBase](#)
A LinOp implementing this interface stores a stopping criterion factory.
- class [LowerTrs](#)
LowerTrs is the triangular solver which solves the system $Lx = b$, when L is a lower triangular matrix.
- class [Multigrid](#)
Multigrid methods have a hierarchy of many levels, whose coarse level is a subset of the fine level, of the problem.
- class [UpperTrs](#)
UpperTrs is the triangular solver which solves the system $Ux = b$, when U is an upper triangular matrix.
- struct [workspace_traits](#)
Traits class providing information on the type and location of workspace vectors inside a solver.

Enumerations

- enum [initial_guess_mode](#) { [initial_guess_mode::zero](#), [initial_guess_mode::rhs](#), [initial_guess_mode::provided](#) }
Give a initial guess mode about the input of the apply method.
- enum [trisolve_algorithm](#)
A helper for algorithm selection in the triangular solvers.

Functions

- template<typename ValueType >
auto [build_smoother](#) (std::shared_ptr< const [LinOpFactory](#) > factory, [size_type](#) iteration=1, ValueType relaxation_factor=0.9)
build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relaxation_factor).
- template<typename ValueType >
auto [build_smoother](#) (std::shared_ptr< const LinOp > solver, [size_type](#) iteration=1, ValueType relaxation_factor=0.9)
build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relaxation_factor).

46.15.1 Detailed Description

The ginkgo Solve namespace.

The ginkgo Solver namespace.

46.15.2 Enumeration Type Documentation

46.15.2.1 initial_guess_mode

```
enum gko::solver::initial_guess_mode [strong]
```

Give a initial guess mode about the input of the apply method.

Enumerator

zero	the input is zero
rhs	the input is right hand side
provided	the input is provided

```
33                                     {
37     zero,
41     rhs,
45     provided
46 };
```

46.15.2.2 trisolve_algorithm

```
enum gko::solver::trisolve_algorithm [strong]
```

A helper for algorithm selection in the triangular solvers.

It currently only matters for the Cuda executor as there, we have a choice between the Ginkgo syncfree and cuS_{PARSE} implementations.

```
40 { sparselib, syncfree };
```

46.15.3 Function Documentation

46.15.3.1 build_smoother() [1/2]

```
template<typename ValueType >
auto gko::solver::build_smoother (
    std::shared_ptr< const LinOp > solver,
    size_type iteration = 1,
    ValueType relaxation_factor = 0.9 )
```

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relaxation_factor).

Parameters

<i>solver</i>	the shared pointer of solver
<i>iteration</i>	the maximum number of iteration, which default is 1
<i>relaxation_factor</i>	the relaxation factor for Richardson

Returns

the pointer of Ir(Richardson)

Note

this is the overload function for LinOp.

```

328 {
329     auto exec = solver->get_executor();
330     return Ir<ValueType>::build()
331         .with_generated_solver(solver)
332         .with_relaxation_factor(relaxation_factor)
333         .with_criteria(gko::stop::Iteration::build().with_max_iters(iteration))
334         .on(exec);
335 }
```

46.15.3.2 build_smoother() [2/2]

```

template<typename ValueType >
auto gko::solver::build_smoother (
    std::shared_ptr< const LinOpFactory > factory,
    size_type iteration = 1,
    ValueType relaxation_factor = 0.9 )
```

build_smoother gives a shortcut to build a smoother by IR(Richardson) with limited stop criterion(iterations and relaxation_factor).

Parameters

<i>factory</i>	the shared pointer of factory
<i>iteration</i>	the maximum number of iteration, which default is 1
<i>relaxation_factor</i>	the relaxation factor for Richardson

Returns

the pointer of Ir(Richardson)

46.16 gko::solver::multigrid Namespace Reference

The solver multigrid namespace.

Enumerations

- enum [cycle](#)
cycle defines which kind of multigrid cycle can be used.
- enum [mid_smooth_type](#)
mid_smooth_type gives the options to handle the middle smoother behavior between the two cycles in the same level.

46.16.1 Detailed Description

The solver multigrid namespace.

46.16.2 Enumeration Type Documentation

46.16.2.1 cycle

```
enum gko::solver::multigrid::cycle [strong]
```

cycle defines which kind of multigrid cycle can be used.

It contains V, W, and F cycle.

- V, W cycle uses the algorithm according to Briggs, Henson, and McCormick: A multigrid tutorial 2nd Edition.
- F cycle uses the algorithm according to Trottenberg, Oosterlee, and Schuller: [Multigrid](#) 1st Edition. F cycle first uses the recursive call but second uses the V-cycle call such that F-cycle is between V and W cycle.

46.16.2.2 mid_smooth_type

```
enum gko::solver::multigrid::mid_smooth_type [strong]
```

mid_smooth_type gives the options to handle the middle smoother behavior between the two cycles in the same level.

It only affects the behavior when there's no operation between the post smoother of previous cycle and the pre smoother of next cycle. Thus, it only affects W cycle and F cycle.

- both: gives the same behavior as the original algorithm, which use posts smoother from previous cycle and pre smoother from next cycle.
- post_smoother: only uses the post smoother of previous cycle in the mid smoother
- pre_smoother: only uses the pre smoother of next cycle in the mid smoother
- standalone: uses the defined smoother in the mid smoother

46.17 gko::stop Namespace Reference

The Stopping criterion namespace.

Classes

- class [AbsoluteResidualNorm](#)
The [AbsoluteResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.
- class [Combined](#)
The [Combined](#) class is used to combine multiple criterions together through an OR operation.
- class [Criterion](#)
The [Criterion](#) class is a base class for all stopping criteria.
- struct [CriterionArgs](#)
This struct is used to pass parameters to the `EnableDefaultCriterionFactory::generate()` method.
- class [ImplicitResidualNorm](#)
The [ImplicitResidualNorm](#) class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.
- class [Iteration](#)
The [Iteration](#) class is a stopping criterion which stops the iteration process after a preset number of iterations.
- class [RelativeResidualNorm](#)
The [RelativeResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.
- class [ResidualNorm](#)
The [ResidualNorm](#) class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.
- class [ResidualNormBase](#)
The [ResidualNormBase](#) class provides a framework for stopping criteria related to the residual norm.
- class [ResidualNormReduction](#)
The [ResidualNormReduction](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.
- class [Time](#)
The [Time](#) class is a stopping criterion which stops the iteration process after a certain amount of time has passed.

Typedefs

- using [CriterionFactory](#) = [AbstractFactory](#)< [Criterion](#), [CriterionArgs](#) >
Declares an Abstract Factory specialized for [Criteria](#)s.
- template<typename ConcreteFactory , typename ConcreteCriterion , typename ParametersType , typename PolymorphicBase = [CriterionFactory](#)>
 using [EnableDefaultCriterionFactory](#) = [EnableDefaultFactory](#)< ConcreteFactory, ConcreteCriterion, ParametersType, PolymorphicBase >
This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of [CriterionFactory](#).

Enumerations

- enum [mode](#)
The mode for the residual norm criterion.

Functions

- `template<typename FactoryContainer >
std::shared_ptr< const CriterionFactory > combine (FactoryContainer &&factories)`
Combines multiple criterion factories into a single combined criterion factory.

46.17.1 Detailed Description

The Stopping criterion namespace.

[Stopping criteria](#)

46.17.2 Typedef Documentation

46.17.2.1 EnableDefaultCriterionFactory

```
template<typename ConcreteFactory , typename ConcreteCriterion , typename ParametersType ,
typename PolymorphicBase = CriterionFactory>
using gko::stop::EnableDefaultCriterionFactory = typedef EnableDefaultFactory<ConcreteFactory,
ConcreteCriterion, ParametersType, PolymorphicBase>
```

This is an alias for the [EnableDefaultFactory](#) mixin, which correctly sets the template parameters to enable a subclass of CriterionFactory.

Template Parameters

<i>ConcreteFactory</i>	the concrete factory which is being implemented [CRTP parameter]
<i>ConcreteCriterion</i>	the concrete Criterion type which this factory produces, needs to have a constructor which takes a const ConcreteFactory *, and a const CriterionArgs * as parameters.
<i>ParametersType</i>	a subclass of enable_parameters_type template which defines all of the parameters of the factory
<i>PolymorphicBase</i>	parent of ConcreteFactory in the polymorphic hierarchy, has to be a subclass of CriterionFactory

46.18 gko::syn Namespace Reference

The Synthesizer namespace.

Classes

- struct [range](#)
range records start, end, step in template
- struct [type_list](#)

- *type_list* records several types in template
- struct *value_list*
value_list records several values with the same type in template.

Typedefs

- template<typename List1 , typename List2 >
using *concatenate* = typename detail::concatenate_impl< List1, List2 >::type
concatenate combines two *value_list* into one *value_list*.
- template<typename T >
using *as_list* = typename detail::as_list_impl< T >::type
as_list<T> gives the alias type of *as_list_impl*<T>::type.

Functions

- template<typename T , T... Value>
constexpr *std::array*< T, sizeof...(Value)> *as_array* (*value_list*< T, Value... > vl)
as_array<T> returns the array from *value_list*.

46.18.1 Detailed Description

The Synthesizer namespace.

46.18.2 Typedef Documentation

46.18.2.1 as_list

```
template<typename T >
using gko::syn::as_list = typedef typename detail::as_list_impl<T>::type
```

as_list<T> gives the alias type of *as_list_impl*<T>::type.

It gives a list (itself) if input is already a list, or generates list type from range input.

Template Parameters

<i>T</i>	list or range
----------	---------------

46.18.2.2 concatenate

```
template<typename List1 , typename List2 >
using gko::syn::concatenate = typedef typename detail::concatenate_impl<List1, List2>::type
```

concatenate combines two [value_list](#) into one [value_list](#).

Template Parameters

<i>List1</i>	the first list
<i>List2</i>	the second list

46.18.3 Function Documentation

46.18.3.1 as_array()

```
template<typename T , T... Value>
constexpr std::array<T, sizeof...(Value)> gko::syn::as_array (
    value\_list< T, Value... > vl ) [constexpr]
```

[as_array](#)<T> returns the array from [value_list](#).

It will be helpful if using for in runtime on the array.

Template Parameters

<i>T</i>	the type of value_list
<i>Value</i>	the values of value_list

Parameters

value_list	the input value_list
----------------------------	--------------------------------------

Returns

[std::array](#) the [std::array](#) contains the values of [value_list](#)

```
176 {
177     return std::array<T, sizeof...(Value)>{Value...};
178 }
```

References [gko::array](#).

46.19 gko::xstd Namespace Reference

The namespace for functionalities after C++14 standard.

Typedefs

- template<typename... Ts>
using [void_t](#) = typename detail::make_void< Ts... >::type

Use the custom implementation, since the `std::void_t` used in `is_matrix_type_builder` seems to trigger a compiler bug in GCC 7.5.

46.19.1 Detailed Description

The namespace for functionalities after C++14 standard.

Chapter 47

Class Documentation

47.1 gko::AbsoluteComputable Class Reference

The [AbsoluteComputable](#) is an interface that allows to get the component wise absolute of a LinOp.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual std::unique_ptr< LinOp > [compute_absolute_linop](#) () const =0
Gets the absolute LinOp.
- virtual void [compute_absolute_inplace](#) ()=0
Compute absolute inplace on each element.

47.1.1 Detailed Description

The [AbsoluteComputable](#) is an interface that allows to get the component wise absolute of a LinOp.

Use `EnableAbsoluteComputation<AbsoluteLinOp>` to implement this interface.

47.1.2 Member Function Documentation

47.1.2.1 compute_absolute_linop()

```
virtual std::unique_ptr<LinOp> gko::AbsoluteComputable::compute_absolute_linop ( ) const  
[pure virtual]
```

Gets the absolute LinOp.

Returns

a pointer to the new absolute LinOp

Implemented in [gko::EnableAbsoluteComputation< AbsoluteLinOp >](#), [gko::EnableAbsoluteComputation< remove_complex< Hybrid>](#), [gko::EnableAbsoluteComputation< remove_complex< Sellp< ValueType, IndexType > > >](#), [gko::EnableAbsoluteComputation< remove_complex< Coo< ValueType, IndexType > > >](#), [gko::EnableAbsoluteComputation< remove_complex< Vector< ValueType > > >](#), [gko::EnableAbsoluteComputation< remove_complex< Csr< ValueType, IndexType > > >](#), and [gko::EnableAbsoluteComputation< remove_complex< SparseMatrix< ValueType, IndexType > > >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/lin_op.hpp

47.2 `gko::stop::AbsoluteResidualNorm< ValueType >` Class Template Reference

The `AbsoluteResidualNorm` class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.2.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::stop::AbsoluteResidualNorm< ValueType >
```

The `AbsoluteResidualNorm` class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold, i.e.

when `norm(residual) < threshold`. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on `b` in order to get the number of right-hand sides. If this is not correctly provided, an exception `gko::NotSupported()` is thrown.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/residual_norm.hpp`

47.3 `gko::AbstractFactory< AbstractProductType, ComponentsType >` Class Template Reference

The `AbstractFactory` is a generic interface template that enables easy implementation of the abstract factory design pattern.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

- `template<typename... Args>`
`std::unique_ptr< abstract_product_type > generate (Args &&... args) const`
Creates a new product from the given components.

47.3.1 Detailed Description

```
template<typename AbstractProductType, typename ComponentsType>
class gko::AbstractFactory< AbstractProductType, ComponentsType >
```

The `AbstractFactory` is a generic interface template that enables easy implementation of the abstract factory design pattern.

The interface provides the `AbstractFactory::generate()` method that can produce products of type `AbstractProductType` using an object of `ComponentsType` (which can be constructed on the fly from parameters to its constructors). The `generate()` method is not declared as virtual, as this allows subclasses to hide the method with a variant that preserves the compile-time type of the objects. Instead, implementers should override the `generate_impl()` method, which is declared virtual.

Implementers of concrete factories should consider using the `EnableDefaultFactory` mixin to obtain default implementations of utility methods of `PolymorphicObject` and `AbstractFactory`.

Template Parameters

<i>AbstractProductType</i>	the type of products the factory produces
<i>ComponentsType</i>	the type of components the factory needs to produce the product

47.3.2 Member Function Documentation

47.3.2.1 generate()

```
template<typename AbstractProductType, typename ComponentsType>
template<typename... Args>
std::unique_ptr<abstract_product_type> gko::AbstractFactory< AbstractProductType, ComponentsType >::generate (
    Args &&... args ) const [inline]
```

Creates a new product from the given components.

The method will create an `ComponentsType` object from the arguments of this method, and pass it to the `generate_impl()` function which will create a new `AbstractProductType`.

Template Parameters

<i>Args</i>	types of arguments passed to the constructor of <code>ComponentsType</code>
-------------	---

Parameters

<i>args</i>	arguments passed to the constructor of <code>ComponentsType</code>
-------------	--

Returns

an instance of `AbstractProductType`

```
68     {
69         auto product =
70             this->generate_impl(components_type{std::forward<Args>(args)...});
71         for (auto logger : this->loggers_) {
72             product->add_logger(logger);
73         }
74         return product;
75     }
```

The documentation for this class was generated from the following file:

- `ginkgo/core/base/abstract_factory.hpp`

47.4 gko::AllocationError Class Reference

`AllocationError` is thrown if a memory allocation fails.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [AllocationError](#) (const std::string &file, int line, const std::string &device, [size_type](#) bytes)
Initializes an allocation error.

47.4.1 Detailed Description

[AllocationError](#) is thrown if a memory allocation fails.

47.4.2 Constructor & Destructor Documentation

47.4.2.1 AllocationError()

```
gko::AllocationError::AllocationError (
    const std::string & file,
    int line,
    const std::string & device,
    size\_type bytes ) [inline]
```

Initializes an allocation error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>device</i>	The device on which the error occurred
<i>bytes</i>	The size of the memory block whose allocation failed.

```
547         : Error(file, line,
548               device + ": failed to allocate memory block of " +
549               std::to_string(bytes) + "B")
550     {}
```

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.5 gko::Allocator Class Reference

Provides generic allocation and deallocation functionality to be used by an [Executor](#).

```
#include <ginkgo/core/base/memory.hpp>
```

47.5.1 Detailed Description

Provides generic allocation and deallocation functionality to be used by an [Executor](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/memory.hpp

47.6 gko::experimental::reorder::Amd< IndexType > Class Template Reference

Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.

```
#include <ginkgo/core/reorder/amd.hpp>
```

Public Member Functions

- const parameters_type & [get_parameters](#) ()
Returns the parameters used to construct the factory.
- std::unique_ptr< [permutation_type](#) > [generate](#) (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

- static parameters_type [build](#) ()
Creates a new parameter_type to set up the factory.

47.6.1 Detailed Description

```
template<typename IndexType = int32>
class gko::experimental::reorder::Amd< IndexType >
```

Computes a Approximate Minimum Degree (AMD) reordering of an input matrix.

Template Parameters

<i>IndexType</i>	the type used to store sparsity pattern indices of the system matrix
------------------	--

47.6.2 Member Function Documentation

47.6.2.1 generate()

```
template<typename IndexType = int32>
std::unique_ptr<permutation_type> gko::experimental::reorder::Amd< IndexType >::generate (
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a Permutation instead of a generic LinOp, which would need to be cast to Permutation again to access its indices. It is only necessary because smart pointers aren't covariant.

47.6.2.2 `get_parameters()`

```
template<typename IndexType = int32>
const parameters_type& gko::experimental::reorder::Amd< IndexType >::get_parameters ( ) [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
67 { return parameters_; }
```

The documentation for this class was generated from the following file:

- `ginkgo/core/reorder/amd.hpp`

47.7 `gko::amd_device` Class Reference

`amd_device` handles the number of executor on Amd devices and have the corresponding recursive_mutex.

```
#include <ginkgo/core/base/device.hpp>
```

47.7.1 Detailed Description

`amd_device` handles the number of executor on Amd devices and have the corresponding recursive_mutex.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/device.hpp`

47.8 `gko::solver::ApplyWithInitialGuess` Class Reference

`ApplyWithInitialGuess` provides a way to give the input guess for apply function.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

47.8.1 Detailed Description

`ApplyWithInitialGuess` provides a way to give the input guess for apply function.

All functionalities have the protected access specifier. It should only be used internally.

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/solver_base.hpp`

47.9 gko::are_all_integral< Args > Struct Template Reference

Evaluates if all template arguments Args fulfill std::is_integral.

```
#include <ginkgo/core/base/types.hpp>
```

47.9.1 Detailed Description

```
template<typename... Args>
struct gko::are_all_integral< Args >
```

Evaluates if all template arguments Args fulfill std::is_integral.

If that is the case, this class inherits from std::true_type, otherwise, it inherits from std::false_type. If no values are passed in, std::true_type is inherited from.

Template Parameters

Args...	Arguments to test for std::is_integral
---------	--

The documentation for this struct was generated from the following file:

- ginkgo/core/base/types.hpp

47.10 gko::array< ValueType > Class Template Reference

An array is a container which encapsulates fixed-sized arrays, stored on the [Executor](#) tied to the array.

```
#include <ginkgo/core/base/array.hpp>
```

Public Types

- using [value_type](#) = ValueType
The type of elements stored in the array.
- using [default_deleter](#) = [executor_deleter](#)< [value_type](#)[]>
The default deleter type used by array.
- using [view_deleter](#) = [null_deleter](#)< [value_type](#)[]>
The deleter type used for views.

Public Member Functions

- `array ()` noexcept
Creates an empty array not tied to any executor.
- `array (std::shared_ptr< const Executor > exec)` noexcept
Creates an empty array tied to the specified Executor.
- `array (std::shared_ptr< const Executor > exec, size_type size)`
Creates an array on the specified Executor.
- `template<typename DeleterType >`
`array (std::shared_ptr< const Executor > exec, size_type size, value_type *data, DeleterType deleter)`
Creates an array from existing memory.
- `array (std::shared_ptr< const Executor > exec, size_type size, value_type *data)`
Creates an array from existing memory.
- `template<typename RandomAccessIterator >`
`array (std::shared_ptr< const Executor > exec, RandomAccessIterator begin, RandomAccessIterator end)`
Creates an array on the specified Executor and initializes it with values.
- `template<typename T >`
`array (std::shared_ptr< const Executor > exec, std::initializer_list< T > init_list)`
Creates an array on the specified Executor and initializes it with values.
- `array (std::shared_ptr< const Executor > exec, const array &other)`
Creates a copy of another array on a different executor.
- `array (const array &other)`
Creates a copy of another array.
- `array (std::shared_ptr< const Executor > exec, array &&other)`
Moves another array to a different executor.
- `array (array &&other)`
Moves another array.
- `array< ValueType > as_view ()`
Returns a non-owning view of the memory owned by this array.
- `detail::const_array_view< ValueType > as_const_view () const`
Returns a non-owning constant view of the memory owned by this array.
- `array & operator= (const array &other)`
Copies data from another array or view.
- `array & operator= (array &&other)`
Moves data from another array or view.
- `template<typename OtherValueType >`
`std::enable_if_t<!std::is_same< ValueType, OtherValueType >::value, array > & operator= (const array< OtherValueType > &other)`
Copies and converts data from another array with another data type.
- `array & operator= (const detail::const_array_view< ValueType > &other)`
Copies data from a const_array_view.
- `void clear ()` noexcept
Deallocates all data used by the array.
- `void resize_and_reset (size_type size)`
Resizes the array so it is able to hold the specified number of elements.
- `void fill (const value_type value)`
Fill the array with the given value.
- `size_type get_size () const` noexcept
Returns the number of elements in the array.
- `size_type get_num_elems () const` noexcept
Returns the number of elements in the array.
- `value_type * get_data ()` noexcept

- Returns a pointer to the block of memory used to store the elements of the array.*
- const [value_type](#) * [get_const_data](#) () const noexcept
Returns a constant pointer to the block of memory used to store the elements of the array.
- std::shared_ptr< const [Executor](#) > [get_executor](#) () const noexcept
Returns the [Executor](#) associated with the array.
- void [set_executor](#) (std::shared_ptr< const [Executor](#) > exec)
Changes the [Executor](#) of the array, moving the allocated data to the new [Executor](#).
- bool [is_owning](#) ()
Tells whether this array owns its data or not.

Static Public Member Functions

- static [array view](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size, [value_type](#) *data)
Creates an array from existing memory.
- static detail::const_array_view< ValueType > [const_view](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size, const [value_type](#) *data)
Creates a constant (immutable) array from existing memory.

47.10.1 Detailed Description

```
template<typename ValueType>
class gko::array< ValueType >
```

An array is a container which encapsulates fixed-sized arrays, stored on the [Executor](#) tied to the array.

The array stores and transfers its data as **raw** memory, which means that the constructors of its elements are not called when constructing, copying or moving the array. Thus, the array class is most suitable for storing POD types.

Template Parameters

<i>ValueType</i>	the type of elements stored in the array
------------------	--

47.10.2 Constructor & Destructor Documentation

47.10.2.1 array() [1/11]

```
template<typename ValueType>
gko::array< ValueType >::array ( ) [inline], [noexcept]
```

Creates an empty array not tied to any executor.

An array without an assigned executor can only be empty. Attempts to change its size (e.g. via the `resize_and_reset` method) will result in an exception. If such an array is used as the right hand side of an assignment or move assignment expression, the data of the target array will be cleared, but its executor will not be modified.

The executor can later be set by using the `set_executor` method. If an array with no assigned executor is assigned or moved to, it will inherit the executor of the source array.

47.10.2.2 array() [2/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec ) [inline], [explicit], [noexcept]
```

Creates an empty array tied to the specified [Executor](#).

Parameters

<i>exec</i>	the Executor where the array data is allocated
-------------	--

47.10.2.3 array() [3/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    size_type size ) [inline]
```

Creates an array on the specified [Executor](#).

Parameters

<i>exec</i>	the Executor where the array data will be allocated
<i>size</i>	the amount of memory (expressed as the number of <code>value_type</code> elements) allocated on the Executor

47.10.2.4 array() [4/11]

```
template<typename ValueType>
template<typename DeleterType >
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    size_type size,
    value_type * data,
    DeleterType deleter ) [inline]
```

Creates an array from existing memory.

The memory will be managed by the array, and deallocated using the specified deleter (e.g. use `std::default_delete` for data allocated with `new`).

Template Parameters

<i>DeleterType</i>	type of the deleter
--------------------	---------------------

Parameters

<i>exec</i>	executor where <code>data</code> is located
<i>size</i>	number of elements in <code>data</code>
<i>data</i>	chunk of memory used to create the array
<i>deleter</i>	the deleter used to free the memory

See also

[array::view\(\)](#) to create an [array](#) that does not deallocate memory

[array\(std::shared_ptr<cont Executor>, size_type, value_type*\)](#) to deallocate the memory using [Executor::free\(\)](#) method

47.10.2.5 array() [5/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    size_type size,
    value_type * data ) [inline]
```

Creates an array from existing memory.

The memory will be managed by the array, and deallocated using the [Executor::free](#) method.

Parameters

<i>exec</i>	executor where <code>data</code> is located
<i>size</i>	number of elements in <code>data</code>
<i>data</i>	chunk of memory used to create the array

47.10.2.6 array() [6/11]

```
template<typename ValueType>
template<typename RandomAccessIterator >
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    RandomAccessIterator begin,
    RandomAccessIterator end ) [inline]
```

Creates an array on the specified [Executor](#) and initializes it with values.

Template Parameters

<i>RandomAccessIterator</i>	type of the iterators
-----------------------------	-----------------------

Parameters

<i>exec</i>	the Executor where the array data will be allocated
<i>begin</i>	start of range of values
<i>end</i>	end of range of values

47.10.2.7 array() [7/11]

```
template<typename ValueType>
template<typename T >
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    std::initializer_list< T > init_list ) [inline]
```

Creates an array on the specified [Executor](#) and initializes it with values.

Template Parameters

<i>T</i>	type of values used to initialize the array (T has to be implicitly convertible to <i>value_type</i>)
----------	--

Parameters

<i>exec</i>	the Executor where the array data will be allocated
<i>init_list</i>	list of values used to initialize the array

47.10.2.8 array() [8/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    const array< ValueType > & other ) [inline]
```

Creates a copy of another array on a different executor.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

<i>exec</i>	the executor where the new array will be created
<i>other</i>	the array to copy from

47.10.2.9 array() [9/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    const array< ValueType > & other ) [inline]
```

Creates a copy of another array.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

<i>other</i>	the array to copy from
--------------	------------------------

47.10.2.10 array() [10/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    std::shared_ptr< const Executor > exec,
    array< ValueType > && other ) [inline]
```

Moves another array to a different executor.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

<i>exec</i>	the executor where the new array will be moved
<i>other</i>	the array to move

47.10.2.11 array() [11/11]

```
template<typename ValueType>
gko::array< ValueType >::array (
    array< ValueType > && other ) [inline]
```

Moves another array.

This does not invoke the constructors of the elements, instead they are copied as POD types.

Parameters

<i>other</i>	the array to move
--------------	-------------------

47.10.3 Member Function Documentation

47.10.3.1 `as_const_view()`

```
template<typename ValueType>
detail::const_array_view<ValueType> gko::array< ValueType >::as_const_view ( ) const [inline]
```

Returns a non-owning constant view of the memory owned by this array.

It can only be used until this array gets deleted, cleared or resized.

47.10.3.2 `as_view()`

```
template<typename ValueType>
array<ValueType> gko::array< ValueType >::as_view ( ) [inline]
```

Returns a non-owning view of the memory owned by this array.

It can only be used until this array gets deleted, cleared or resized.

47.10.3.3 `clear()`

```
template<typename ValueType>
void gko::array< ValueType >::clear ( ) [inline], [noexcept]
```

Deallocates all data used by the array.

The array is left in a valid, but empty state, so the same array can be used to allocate new memory. Calls to `array::get_data()` will return a `nullptr`.

Referenced by `gko::index_set< IndexType >::clear()`, and `gko::array< index_type >::operator=()`.

47.10.3.4 `const_view()`

```
template<typename ValueType>
static detail::const_array_view<ValueType> gko::array< ValueType >::const_view (
    std::shared_ptr< const Executor > exec,
    size_type size,
    const value_type * data ) [inline], [static]
```

Creates a constant (immutable) array from existing memory.

The array does not take ownership of the memory, and will not deallocate it once it goes out of scope. This array type cannot use the function `resize_and_reset` since it does not own the data it should resize.

Parameters

<i>exec</i>	executor where <code>data</code> is located
<i>size</i>	number of elements in <code>data</code>
<i>data</i>	chunk of memory used to create the array

Returns

an array constructed from `data`

Referenced by `gko::array< index_type >::as_const_view()`.

47.10.3.5 fill()

```
template<typename ValueType>
void gko::array< ValueType >::fill (
    const value_type value )
```

Fill the array with the given value.

Parameters

<i>value</i>	the value to be filled
--------------	------------------------

47.10.3.6 get_const_data()

```
template<typename ValueType>
const value_type* gko::array< ValueType >::get_const_data ( ) const [inline], [noexcept]
```

Returns a constant pointer to the block of memory used to store the elements of the array.

Returns

a constant pointer to the block of memory used to store the elements of the array

Referenced by `gko::array< index_type >::as_const_view()`, `gko::batch::matrix::Dense< ValueType >::at()`, `gko::batch::MultiVector< ValueType >::at()`, `gko::matrix::Dense< value_type >::at()`, `gko::preconditioner::Jacobi< ValueType, IndexType >::get_blocks()`, `gko::preconditioner::Jacobi< ValueType, IndexType >::get_conditioning()`, `gko::multigrid::Pgm< ValueType, IndexType >::get_const_agg()`, `gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_block_pointers()`, `gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_blocks()`, `gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_blocks_cumulative_offsets()`, `gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_col_idxs()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_const_col_idxs()`, `gko::batch::matrix::Ell< ValueType, IndexType >::get_const_col_idxs()`, `gko::matrix::Sellp< ValueType, IndexType >::get_const_col_idxs()`, `gko::matrix::Coo< ValueType, IndexType >::get_const_col_idxs()`, `gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs()`, `gko::device::matrix_data< ValueType, IndexType >::get_const_col_idxs()`, `gko::matrix::Fbcsr< ValueType, IndexType >::get_const_col_idxs()`, `gko::matrix::Csr< ValueType, IndexType >::get_const_col_idxs()`, `gko::batch::matrix::Ell<`

ValueType, IndexType >::get_const_col_idxs_for_item(), gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_map_block_to_row(), gko::matrix::ScaledPermutation< ValueType, IndexType >::get_const_permutation(), gko::matrix::Permutation< IndexType >::get_const_permutation(), gko::matrix::RowGatherer< IndexType >::get_const_row_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_const_row_idxs(), gko::device_matrix_data< ValueType, IndexType >::get_const_row_idxs(), gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_row_ptrs(), gko::batch::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs(), gko::matrix::Fbcsr< ValueType, IndexType >::get_const_row_ptrs(), gko::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs(), gko::matrix::ScaledPermutation< ValueType, IndexType >::get_const_scaling_factors(), gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_lengths(), gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_sets(), gko::matrix::Csr< ValueType, IndexType >::get_const_srow(), gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_value(), gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values(), gko::matrix::Diagonal< ValueType >::get_const_values(), gko::batch::matrix::Ell< ValueType, IndexType >::get_const_values(), gko::batch::matrix::Dense< ValueType >::get_const_values(), gko::matrix::Sellp< ValueType, IndexType >::get_const_values(), gko::matrix::Coo< ValueType, IndexType >::get_const_values(), gko::matrix::Ell< ValueType, IndexType >::get_const_values(), gko::batch::MultiVector< ValueType >::get_const_values(), gko::device_matrix_data< ValueType, IndexType >::get_const_values(), gko::matrix::Fbcsr< ValueType, IndexType >::get_const_values(), gko::matrix::Dense< value_type >::get_const_values(), gko::matrix::Csr< ValueType, IndexType >::get_const_values(), gko::batch::MultiVector< ValueType >::get_const_values_for_item(), gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values_for_item(), gko::batch::matrix::Dense< ValueType >::get_const_values_for_item(), gko::batch::matrix::Ell< ValueType, IndexType >::get_const_values_for_item(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_ids(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_sizes(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_range_bounds(), gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_range_starting_indices(), gko::index_set< IndexType >::get_subsets_begin(), gko::index_set< IndexType >::get_subsets_end(), gko::index_set< IndexType >::get_superset_indices(), gko::matrix::Csr< ValueType, IndexType >::classical::process(), gko::matrix::Csr< ValueType, IndexType >::load_balance::process(), gko::matrix::Ell< ValueType, IndexType >::val_at(), and gko::matrix::Sellp< ValueType, IndexType >::val_at().

47.10.3.7 get_data()

```
template<typename ValueType>
value_type* gko::array< ValueType >::get_data ( ) [inline], [noexcept]
```

Returns a pointer to the block of memory used to store the elements of the array.

Returns

a pointer to the block of memory used to store the elements of the array

Referenced by gko::array< index_type >::as_view(), gko::batch::matrix::Dense< ValueType >::at(), gko::batch::MultiVector< ValueType >::at(), gko::matrix::Dense< value_type >::at(), gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row(), gko::multigrid::Pgm< ValueType, IndexType >::get_agg(), gko::matrix::SparsityCsr< ValueType, IndexType >::get_col_idxs(), gko::batch::matrix::Csr< ValueType, IndexType >::get_col_idxs(), gko::batch::matrix::Ell< ValueType, IndexType >::get_col_idxs(), gko::matrix::Sellp< ValueType, IndexType >::get_col_idxs(), gko::matrix::Coo< ValueType, IndexType >::get_col_idxs(), gko::matrix::Ell< ValueType, IndexType >::get_col_idxs(), gko::device_matrix_data< ValueType, IndexType >::get_col_idxs(), gko::matrix::Fbcsr< ValueType, IndexType >::get_col_idxs(), gko::matrix::Csr< ValueType, IndexType >::get_col_idxs(), gko::batch::matrix::Ell< ValueType, IndexType >::get_col_idxs_for_item(), gko::matrix::ScaledPermutation< ValueType, IndexType >::get_permutation(), gko::matrix::Permutation< IndexType >::get_permutation(), gko::matrix::RowGatherer< IndexType >::get_row_idxs(), gko::matrix::Coo< ValueType,

`IndexType >::get_row_idxs()`, `gko::device_matrix_data< ValueType, IndexType >::get_row_idxs()`, `gko::matrix< SparsityCsr< ValueType, IndexType >::get_row_ptrs()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_row_ptrs()`, `gko::matrix::Fbcsr< ValueType, IndexType >::get_row_ptrs()`, `gko::matrix::Csr< ValueType, IndexType >::get_row_ptrs()`, `gko::matrix::ScaledPermutation< ValueType, IndexType >::get_scaling_factors()`, `gko::matrix::Sellp< ValueType, IndexType >::get_slice_lengths()`, `gko::matrix::Sellp< ValueType, IndexType >::get_slice_sets()`, `gko::matrix::Csr< ValueType, IndexType >::get_srow()`, `gko::matrix::SparsityCsr< ValueType, IndexType >::get_value()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_values()`, `gko::batch::matrix::Ell< ValueType, IndexType >::get_values()`, `gko::matrix::Diagonal< ValueType >::get_values()`, `gko::batch::matrix::Dense< ValueType >::get_values()`, `gko::matrix::Sellp< ValueType, IndexType >::get_values()`, `gko::matrix::Coo< ValueType, IndexType >::get_values()`, `gko::matrix::Ell< ValueType, IndexType >::get_values()`, `gko::batch::MultiVector< ValueType >::get_values()`, `gko::device_matrix_data< ValueType, IndexType >::get_values()`, `gko::matrix::Fbcsr< ValueType, IndexType >::get_values()`, `gko::matrix::Dense< ValueType >::get_values()`, `gko::matrix::Csr< ValueType, IndexType >::get_values()`, `gko::batch::MultiVector< ValueType >::get_values_for_item()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_values_for_item()`, `gko::batch::matrix::Dense< ValueType >::get_values_for_item()`, `gko::batch::matrix::Ell< ValueType, IndexType >::get_values_for_item()`, `gko::array< index_type >::operator=()`, `gko::matrix::Csr< ValueType, IndexType >::load_balance::process()`, `gko::matrix::Ell< ValueType, IndexType >::val_at()`, and `gko::matrix::Sellp< ValueType, IndexType >::val_at()`.

47.10.3.8 get_executor()

```
template<typename ValueType>
std::shared_ptr<const Executor> gko::array< ValueType >::get_executor ( ) const [inline],
[noexcept]
```

Returns the [Executor](#) associated with the array.

Returns

the [Executor](#) associated with the array

Referenced by `gko::array< index_type >::as_const_view()`, `gko::array< index_type >::as_view()`, `gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config()`, `gko::device_matrix_data< ValueType, IndexType >::get_executor()`, `gko::matrix::Csr< ValueType, IndexType >::classical::process()`, and `gko::matrix::Csr< ValueType, IndexType >::load_balance::process()`.

47.10.3.9 get_num_elems()

```
template<typename ValueType>
size_type gko::array< ValueType >::get_num_elems ( ) const [inline], [noexcept]
```

Returns the number of elements in the array.

Returns

the number of elements in the array

47.10.3.10 get_size()

```
template<typename ValueType>
size_type gko::array< ValueType >::get_size ( ) const [inline], [noexcept]
```

Returns the number of elements in the array.

Returns

the number of elements in the array

Referenced by `gko::array< index_type >::as_const_view()`, `gko::array< index_type >::as_view()`, `gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row()`, `gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_stored_elements_per_row()`, `gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values_for_item()`, `gko::matrix::SparsityCsr< ValueType, IndexType >::get_num_nonzeros()`, `gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_num_ranges()`, `gko::matrix::Csr< ValueType, IndexType >::get_num_srow_elements()`, `gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_blocks()`, `gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored_elements()`, `gko::batch::matrix::Ell< ValueType, IndexType >::get_num_stored_elements()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_num_stored_elements()`, `gko::device_matrix_data< ValueType, IndexType >::get_num_stored_elements()`, `gko::matrix::Ell< ValueType, IndexType >::get_num_stored_elements()`, `gko::matrix::Coo< ValueType, IndexType >::get_num_stored_elements()`, `gko::batch::MultiVector< ValueType >::get_num_stored_elements()`, `gko::batch::matrix::Dense< ValueType >::get_num_stored_elements()`, `gko::matrix::Sellp< ValueType, IndexType >::get_num_stored_elements()`, `gko::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored_elements()`, `gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_elements()`, `gko::matrix::Dense< value_type >::get_num_stored_elements()`, `gko::matrix::Csr< ValueType, IndexType >::get_num_stored_elements()`, `gko::index_set< IndexType >::get_num_subsets()`, `gko::matrix::Sellp< ValueType, IndexType >::get_total_cols()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_values_for_item()`, `gko::index_set< IndexType >::index_set()`, `gko::array< index_type >::operator=()`, `gko::matrix::Csr< ValueType, IndexType >::classical::process()`, and `gko::matrix::Csr< ValueType, IndexType >::load_balance::process()`.

47.10.3.11 is_owning()

```
template<typename ValueType>
bool gko::array< ValueType >::is_owning ( ) [inline]
```

Tells whether this array owns its data or not.

Views do not own their data and this has multiple implications. They cannot be resized since the data is not owned by the array which stores a view. It is also unclear whether custom deleter types are owning types as they could be a user-created view-type, therefore only proper array which use the `default_deleter` are considered owning types.

Returns

whether this array can be resized or not.

Referenced by `gko::array< index_type >::operator=()`.

47.10.3.12 operator=() [1/4]

```
template<typename ValueType>
array& gko::array< ValueType >::operator= (
    array< ValueType > && other ) [inline]
```

Moves data from another array or view.

Only the pointer and deleter type change, a copy only happens when targeting another executor's data. This means that in the following situation:

```
gko::array<int> a; // an existing array or view
gko::array<int> b; // an existing array or view
b = std::move(a);
```

Depending on whether *a* and *b* are array or view, this happens:

- *a* and *b* are views, *b* becomes the only valid view of *a*;
- *a* and *b* are arrays, *b* becomes the only valid array of *a*;
- *a* is a view and *b* is an array, *b* frees its data and becomes the only valid view of *a* ();
- *a* is an array and *b* is a view, *b* becomes the only valid array of *a*.

In all the previous cases, *a* becomes empty (e.g., a `nullptr`).

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of *other*.

Parameters

<i>other</i>	the array to move data from
--------------	-----------------------------

Returns

this

47.10.3.13 operator=() [2/4]

```
template<typename ValueType>
array& gko::array< ValueType >::operator= (
    const array< ValueType > & other ) [inline]
```

Copies data from another array or view.

In the case of an array target, the array is resized to match the source's size. In the case of a view target, if the dimensions are not compatible a `gko::OutOfBoundsError` is thrown.

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of *other*.

Parameters

<i>other</i>	the array to copy from
--------------	------------------------

Returns

this

47.10.3.14 operator=() [3/4]

```
template<typename ValueType>
template<typename OtherValueType >
std::enable_if_t<!std::is_same<ValueType, OtherValueType>::value, array>& gko::array< ValueType >::operator= (
    const array< OtherValueType > & other ) [inline]
```

Copies and converts data from another array with another data type.

In the case of an array target, the array is resized to match the source's size. In the case of a view target, if the dimensions are not compatible a [gko::OutOfBoundsError](#) is thrown.

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

<i>other</i>	the array to copy from
--------------	------------------------

Template Parameters

<i>OtherValueType</i>	the value type of <i>other</i>
-----------------------	--------------------------------

Returns

this

47.10.3.15 operator=() [4/4]

```
template<typename ValueType>
array& gko::array< ValueType >::operator= (
    const detail::const_array_view< ValueType > & other ) [inline]
```

Copies data from a `const_array_view`.

In the case of an array target, the array is resized to match the source's size. In the case of a view target, if the dimensions are not compatible a [gko::OutOfBoundsError](#) is thrown.

This does not invoke the constructors of the elements, instead they are copied as POD types.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

<i>other</i>	the <code>const_array_view</code> to copy from
--------------	--

Returns

this

47.10.3.16 `resize_and_reset()`

```
template<typename ValueType>
void gko::array< ValueType >::resize_and_reset (
    size_type size ) [inline]
```

Resizes the array so it is able to hold the specified number of elements.

For a view and other non-owning array types, this throws an exception since these types cannot be resized.

All data stored in the array will be lost.

If the array is not assigned an executor, an exception will be thrown.

Parameters

<i>size</i>	the amount of memory (expressed as the number of <code>value_type</code> elements) allocated on the Executor
-------------	--

Referenced by `gko::array< index_type >::operator=()`.

47.10.3.17 `set_executor()`

```
template<typename ValueType>
void gko::array< ValueType >::set_executor (
    std::shared_ptr< const Executor > exec ) [inline]
```

Changes the [Executor](#) of the array, moving the allocated data to the new [Executor](#).

Parameters

<i>exec</i>	the Executor where the data will be moved to
-------------	--

47.10.3.18 view()

```
template<typename ValueType>
static array gko::array< ValueType >::view (
    std::shared_ptr< const Executor > exec,
    size_type size,
    value_type * data ) [inline], [static]
```

Creates an array from existing memory.

The array does not take ownership of the memory, and will not deallocate it once it goes out of scope. This array type cannot use the function `resize_and_reset` since it does not own the data it should resize.

Parameters

<i>exec</i>	executor where <i>data</i> is located
<i>size</i>	number of elements in <i>data</i>
<i>data</i>	chunk of memory used to create the array

Returns

an array constructed from *data*

Referenced by `gko::array< index_type >::as_view()`.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/array.hpp`

47.11 gko::device_matrix_data< ValueType, IndexType >::arrays Struct Reference

Stores the internal arrays of a `device_matrix_data` object.

```
#include <ginkgo/core/base/device_matrix_data.hpp>
```

47.11.1 Detailed Description

```
template<typename ValueType, typename IndexType>
struct gko::device_matrix_data< ValueType, IndexType >::arrays
```

Stores the internal arrays of a `device_matrix_data` object.

The documentation for this struct was generated from the following file:

- `ginkgo/core/base/device_matrix_data.hpp`

47.12 gko::matrix::Hybrid< ValueType, IndexType >::automatic Class Reference

automatic is a [strategy_type](#) which decides the number of stored elements per row of the ell part automatically.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- [automatic](#) ()
Creates an automatic strategy.
- [size_type compute_ell_num_stored_elements_per_row](#) (array< [size_type](#) > *row_nnz) const override
Computes the number of stored elements per row of the ell part.

47.12.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::automatic
```

automatic is a [strategy_type](#) which decides the number of stored elements per row of the ell part automatically.

47.12.2 Member Function Documentation

47.12.2.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::automatic::compute_ell_num_stored_↵
elements_per_row (
    array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<i>row_nnz</i>	the number of nonzeros of each row
----------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type](#).

References [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_stored_↵_elements_per_row\(\)](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/hybrid.hpp`

47.13 gko::BadDimension Class Reference

`BadDimension` is thrown if an operation is being applied to a `LinOp` with bad dimensions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- `BadDimension` (const std::string &file, int line, const std::string &func, const std::string &op_name, [size_type](#) op_num_rows, [size_type](#) op_num_cols, const std::string &clarification)
Initializes a bad dimension error.

47.13.1 Detailed Description

`BadDimension` is thrown if an operation is being applied to a `LinOp` with bad dimensions.

47.13.2 Constructor & Destructor Documentation

47.13.2.1 BadDimension()

```
gko::BadDimension::BadDimension (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & op_name,
    size\_type op_num_rows,
    size\_type op_num_cols,
    const std::string & clarification ) [inline]
```

Initializes a bad dimension error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The function name where the error occurred
<i>op_name</i>	The name of the operator
<i>op_num_rows</i>	The row dimension of the operator
<i>op_num_cols</i>	The column dimension of the operator
<i>clarification</i>	An additional message further describing the error

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.14 gko::batch_dim< Dimensionality, DimensionType > Struct Template Reference

A type representing the dimensions of a multidimensional batch object.

```
#include <ginkgo/core/base/batch_dim.hpp>
```

Public Member Functions

- [size_type get_num_batch_items](#) () const
Get the number of batch items stored.
- [dim](#)< dimensionality, dimension_type > [get_common_size](#) () const
Get the common size of the batch items.
- [batch_dim](#) ()
The default constructor.
- [batch_dim](#) (const [size_type](#) num_batch_items, const [dim](#)< dimensionality, dimension_type > &common_size)
Creates a [batch_dim](#) object which stores a uniform size for all batch entries.

Friends

- bool [operator==](#) (const [batch_dim](#) &x, const [batch_dim](#) &y)
Checks if two [batch_dim](#) objects are equal.
- bool [operator!=](#) (const [batch_dim](#)< Dimensionality, DimensionType > &x, const [batch_dim](#)< Dimensionality, DimensionType > &y)
Checks if two [batch_dim](#) objects are different.

47.14.1 Detailed Description

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
struct gko::batch_dim< Dimensionality, DimensionType >
```

A type representing the dimensions of a multidimensional batch object.

Template Parameters

<i>Dimensionality</i>	number of dimensions of the object
<i>DimensionType</i>	datatype used to represent each dimension

47.14.2 Constructor & Destructor Documentation

47.14.2.1 batch_dim()

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
gko::batch_dim< Dimensionality, DimensionType >::batch_dim (
    const size_type num_batch_items,
    const dim< dimensionality, dimension_type > & common_size ) [inline], [explicit]
```

Creates a `batch_dim` object which stores a uniform size for all batch entries.

Parameters

<code>num_batch_items</code>	the number of batch items to be stored
<code>common_size</code>	the common size of all the batch items stored

Note

Use this constructor when uniform batches need to be stored.

47.14.3 Member Function Documentation

47.14.3.1 get_common_size()

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
dim<dimensionality, dimension_type> gko::batch_dim< Dimensionality, DimensionType >::get_↵
common_size ( ) const [inline]
```

Get the common size of the batch items.

Returns

`common_size`

Referenced by `gko::batch::MultiVector< ValueType >::get_common_size()`, and `gko::transpose()`.

47.14.3.2 get_num_batch_items()

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
size_type gko::batch_dim< Dimensionality, DimensionType >::get_num_batch_items ( ) const
[inline]
```

Get the number of batch items stored.

Returns

`num_batch_items`

Referenced by `gko::batch::MultiVector< ValueType >::get_num_batch_items()`, and `gko::transpose()`.

47.14.4 Friends And Related Function Documentation

47.14.4.1 operator"!="

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
bool operator!= (
    const batch_dim< Dimensionality, DimensionType > & x,
    const batch_dim< Dimensionality, DimensionType > & y ) [friend]
```

Checks if two [batch_dim](#) objects are different.

Template Parameters

<i>Dimensionality</i>	number of dimensions of the dim objects
<i>DimensionType</i>	datatype used to represent each dimension

Parameters

<i>x</i>	first object
<i>y</i>	second object

Returns

`!(x == y)`

47.14.4.2 operator==

```
template<size_type Dimensionality = 2, typename DimensionType = size_type>
bool operator== (
    const batch_dim< Dimensionality, DimensionType > & x,
    const batch_dim< Dimensionality, DimensionType > & y ) [friend]
```

Checks if two [batch_dim](#) objects are equal.

Parameters

<i>x</i>	first object
<i>y</i>	second object

Returns

true if and only if all dimensions of both objects are equal.

The documentation for this struct was generated from the following file:

- `ginkgo/core/base/batch_dim.hpp`

47.15 gko::batch::log::BatchConvergence< ValueType > Class Template Reference

Logs the final residuals and iteration counts for a batch solver.

```
#include <ginkgo/core/log/batch_logger.hpp>
```

Public Member Functions

- const [array](#)< index_type > & [get_num_iterations](#) () const noexcept
- const [array](#)< real_type > & [get_residual_norm](#) () const noexcept

Static Public Member Functions

- static std::unique_ptr< [BatchConvergence](#) > [create](#) (const mask_type &enabled_events=gko::log::Logger↵
::batch_solver_completed_mask)
Creates a convergence logger.

47.15.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::log::BatchConvergence< ValueType >
```

Logs the final residuals and iteration counts for a batch solver.

The purpose of this logger is to give simple access to standard data generated by the solver once it has converged.

Note

The final logged residuals are the implicit residuals that have been computed within the solver process. Depending on the solver algorithm, this may be significantly different from the true residual ($\|b - Ax\|$).

47.15.2 Member Function Documentation

47.15.2.1 create()

```
template<typename ValueType = default_precision>
static std::unique_ptr<BatchConvergence> gko::batch::log::BatchConvergence< ValueType >↵
::create (
    const mask_type & enabled_events = gko::log::Logger::batch\_solver\_completed\_mask
) [inline], [static]
```

Creates a convergence logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object. TODO: See if the objects can be pre-allocated beforehand instead of being copied in the on_<> event

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.

Returns

an `std::unique_ptr` to the the constructed object

```

95     {
96         return std::unique_ptr<BatchConvergence>(
97             new BatchConvergence(enabled_events));
98     }
```

47.15.2.2 get_num_iterations()

```

template<typename ValueType = default_precision>
const array<index_type>& gko::batch::log::BatchConvergence< ValueType >::get_num_iterations (
) const [inline], [noexcept]
```

Returns

The number of iterations for entire batch

47.15.2.3 get_residual_norm()

```

template<typename ValueType = default_precision>
const array<real_type>& gko::batch::log::BatchConvergence< ValueType >::get_residual_norm (
) const [inline], [noexcept]
```

Returns

The residual norms for the entire batch.

The documentation for this class was generated from the following file:

- `ginkgo/core/log/batch_logger.hpp`

47.16 gko::batch::BatchLinOpFactory Class Reference

A [BatchLinOpFactory](#) represents a higher order mapping which transforms one batch linear operator into another.

```
#include <ginkgo/core/base/batch_lin_op.hpp>
```

Additional Inherited Members

47.16.1 Detailed Description

A [BatchLinOpFactory](#) represents a higher order mapping which transforms one batch linear operator into another.

In a similar fashion to LinOps, BatchLinOps are also "generated" from the [BatchLinOpFactory](#). A function of this class is to provide a generate method, which internally calls the generate_impl(), which the concrete BatchLinOps have to implement.

47.16.1.1 Example: using BatchCG in Ginkgo

```
{c++}
// Suppose A is a batch matrix, batch_b, a batch rhs vector, and batch_x, an
// initial guess
// Create a BatchCG which runs for at most 1000 iterations, and stops after
// reducing the residual norm by 6 orders of magnitude
auto batch_cg_factory = solver::BatchCg<>::build()
    .with_max_iters(1000)
    .with_rel_residual_goal(1e-6)
    .on(cuda);
// create a batch linear operator which represents the solver
auto batch_cg = batch_cg_factory->generate(A);
// solve the system
batch_cg->apply(batch_b, batch_x);
```

The documentation for this class was generated from the following file:

- ginkgo/core/base/batch_lin_op.hpp

47.17 gko::batch::solver::BatchSolver Class Reference

The [BatchSolver](#) is a base class for all batched solvers and provides the common getters and setter for these batched solver classes.

```
#include <ginkgo/core/solver/batch_solver_base.hpp>
```

Public Member Functions

- std::shared_ptr< const BatchLinOp > [get_system_matrix](#) () const
Returns the system operator (matrix) of the linear system.
- std::shared_ptr< const BatchLinOp > [get_preconditioner](#) () const
Returns the generated preconditioner.
- double [get_tolerance](#) () const
Get the residual tolerance used by the solver.
- void [reset_tolerance](#) (double res_tol)
Update the residual tolerance to be used by the solver.
- int [get_max_iterations](#) () const
Get the maximum number of iterations set on the solver.
- void [reset_max_iterations](#) (int max_iterations)
Set the maximum number of iterations for the solver to use, independent of the factory that created it.
- ::gko::batch::stop::tolerance_type [get_tolerance_type](#) () const
Get the tolerance type.
- void [reset_tolerance_type](#) (::gko::batch::stop::tolerance_type tol_type)
Set the type of tolerance check to use inside the solver.

47.17.1 Detailed Description

The [BatchSolver](#) is a base class for all batched solvers and provides the common getters and setter for these batched solver classes.

47.17.2 Member Function Documentation

47.17.2.1 get_max_iterations()

```
int gko::batch::solver::BatchSolver::get_max_iterations ( ) const [inline]
```

Get the maximum number of iterations set on the solver.

Returns

Maximum number of iterations.

```
77 { return this->max_iterations; }
```

47.17.2.2 get_preconditioner()

```
std::shared_ptr<const BatchLinOp> gko::batch::solver::BatchSolver::get_preconditioner ( )  
const [inline]
```

Returns the generated preconditioner.

Returns

the generated preconditioner.

47.17.2.3 get_system_matrix()

```
std::shared_ptr<const BatchLinOp> gko::batch::solver::BatchSolver::get_system_matrix ( ) const  
[inline]
```

Returns the system operator (matrix) of the linear system.

Returns

the system operator (matrix)

47.17.2.4 `get_tolerance()`

```
double gko::batch::solver::BatchSolver::get_tolerance ( ) const [inline]
```

Get the residual tolerance used by the solver.

Returns

The residual tolerance.

47.17.2.5 `get_tolerance_type()`

```
::gko::batch::stop::tolerance_type gko::batch::solver::BatchSolver::get_tolerance_type ( )  
const [inline]
```

Get the tolerance type.

Returns

The tolerance type.

47.17.2.6 `reset_max_iterations()`

```
void gko::batch::solver::BatchSolver::reset_max_iterations (   
    int max_iterations ) [inline]
```

Set the maximum number of iterations for the solver to use, independent of the factory that created it.

Parameters

<i>max_iterations</i>	The maximum number of iterations for the solver.
-----------------------	--

47.17.2.7 `reset_tolerance()`

```
void gko::batch::solver::BatchSolver::reset_tolerance (   
    double res_tol ) [inline]
```

Update the residual tolerance to be used by the solver.

Parameters

<i>res_tol</i>	The residual tolerance to be used for subsequent invocations of the solver.
----------------	---

47.17.2.8 reset_tolerance_type()

```
void gko::batch::solver::BatchSolver::reset_tolerance_type (
    ::gko::batch::stop::tolerance_type tol_type ) [inline]
```

Set the type of tolerance check to use inside the solver.

Parameters

<code>tol_type</code>	The tolerance type.
-----------------------	---------------------

The documentation for this class was generated from the following file:

- ginkgo/core/solver/batch_solver_base.hpp

47.18 gko::solver::Bicg< ValueType > Class Template Reference

BICG or the Biconjugate gradient method is a Krylov subspace solver.

```
#include <ginkgo/core/solver/bicg.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.

Static Public Member Functions

- static `parameters_type parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType >())
Create the parameters from the property_tree.

47.18.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Bicg< ValueType >
```

BICG or the Biconjugate gradient method is a Krylov subspace solver.

Being a generic solver, it is capable of solving general matrices, including non-s.p.d matrices. Though, the memory and the computational requirement of the BiCG solver are higher than of its s.p.d solver counterpart, it has the capability to solve generic systems.

BiCG is based on the bi-Lanczos tridiagonalization method and in exact arithmetic should terminate in at most N iterations ($2N$ MV's, with A and A^H). It forms the basis of many of the cheaper methods such as BiCGSTAB and CGS.

Reference: R.Fletcher, Conjugate gradient methods for indefinite systems, doi: 10.1007/BFb0080116

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.18.2 Member Function Documentation

47.18.2.1 `apply_uses_initial_guess()`

```
template<typename ValueType = default_precision>
bool gko::solver::Bicg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
73 { return true; }
```

47.18.2.2 `conj_transpose()`

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.18.2.3 `parse()`

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Bicg< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.18.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/bicg.hpp

47.19 gko::batch::solver::Bicgstab< ValueType > Class Template Reference

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

```
#include <ginkgo/core/solver/batch_bicgstab.hpp>
```

Additional Inherited Members

47.19.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::solver::Bicgstab< ValueType >
```

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

Being a generic solver, it is capable of solving general matrices, including non-s.p.d matrices.

This solver solves a batch of linear systems using the [Bicgstab](#) algorithm. Each linear system in the batch can converge independently.

Unless otherwise specified via the `preconditioner` factory parameter, this implementation does not use any preconditioner by default. The type of tolerance (absolute or relative) and the maximum number of iterations to be used in the stopping criterion can be set via the factory parameters.

Note

The tolerance check is against the internal residual computed within the solver process. This implicit (internal) residual, can diverge from the true residual ($\|b - Ax\|$). A posteriori checks (by computing the true residual, $\|b - Ax\|$) are recommended to ensure that the solution has converged to the desired tolerance.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/solver/batch_bicgstab.hpp

47.20 gko::solver::Bicgstab< ValueType > Class Template Reference

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

```
#include <ginkgo/core/solver/bicgstab.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- bool [apply_uses_initial_guess](#) () const override
Return true as iterative solvers use the data in x as an initial guess.

Static Public Member Functions

- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType >())
Create the parameters from the property_tree.

47.20.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Bicgstab< ValueType >
```

BiCGSTAB or the Bi-Conjugate Gradient-Stabilized is a Krylov subspace solver.

Being a generic solver, it is capable of solving general matrices, including non-s.p.d matrices. Though, the memory and the computational requirement of the BiCGSTAB solver are higher than of its s.p.d solver counterpart, it has the capability to solve generic systems. It was developed by stabilizing the BiCG method.

Template Parameters

<i>ValueType</i>	precision of the elements of the system matrix.
------------------	---

47.20.2 Member Function Documentation

47.20.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Bicgstab< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
72 { return true; }
```

47.20.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicgstab< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.20.2.3 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Bicgstab< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.20.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Bicgstab< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/bicgstab.hpp

47.21 gko::preconditioner::block_interleaved_storage_scheme< IndexType > Struct Template Reference

Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.

```
#include <ginkgo/core/preconditioner/jacobi.hpp>
```

Public Member Functions

- IndexType [get_group_size](#) () const noexcept
Returns the number of elements in the group.
- [size_type compute_storage_space](#) ([size_type](#) num_blocks) const noexcept
Computes the storage space required for the requested number of blocks.
- IndexType [get_group_offset](#) (IndexType block_id) const noexcept
Returns the offset of the group belonging to the block with the given ID.
- IndexType [get_block_offset](#) (IndexType block_id) const noexcept
Returns the offset of the block with the given ID within its group.
- IndexType [get_global_block_offset](#) (IndexType block_id) const noexcept
Returns the offset of the block with the given ID.
- IndexType [get_stride](#) () const noexcept
Returns the stride between columns of the block.

Public Attributes

- IndexType [block_offset](#)
The offset between consecutive blocks within the group.
- IndexType [group_offset](#)
The offset between two block groups.
- uint32 [group_power](#)
Then base 2 power of the group.

47.21.1 Detailed Description

```
template<typename IndexType>
struct gko::preconditioner::block_interleaved_storage_scheme< IndexType >
```

Defines the parameters of the interleaved block storage scheme used by block-Jacobi blocks.

Template Parameters

<i>IndexType</i>	type used for storing indices of the matrix
------------------	---

47.21.2 Member Function Documentation

47.21.2.1 compute_storage_space()

```
template<typename IndexType>
size_type gko::preconditioner::block_interleaved_storage_scheme< IndexType >::compute_storage↵
_space (
    size_type num_blocks ) const [inline], [noexcept]
```

Computes the storage space required for the requested number of blocks.

Parameters

<i>num_blocks</i>	the total number of blocks that needs to be stored
-------------------	--

Returns

the total memory (as the number of elements) that need to be allocated for the scheme

Note

To simplify using the method in situations where the number of blocks is not known, for a special input `size↵_type{} - 1` the method returns 0 to avoid overallocation of memory.

```

88     {
89         return (num_blocks + 1 == size_type{0})
90             ? size_type{0}
91             : ceildiv(num_blocks, this->get_group_size()) * group_offset;
92     }

```

47.21.2.2 get_block_offset()

```

template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_block_offset
(
    IndexType block_id ) const [inline], [noexcept]

```

Returns the offset of the block with the given ID within its group.

Parameters

<i>block</i> ↔ <i>_id</i>	the ID of the block
------------------------------	---------------------

Returns

the offset of the block with ID `block_id` within its group

Referenced by `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_global_block_offset()`.

47.21.2.3 get_global_block_offset()

```

template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_global_↔
block_offset (
    IndexType block_id ) const [inline], [noexcept]

```

Returns the offset of the block with the given ID.

Parameters

<i>block</i> ↔ <i>_id</i>	the ID of the block
------------------------------	---------------------

Returns

the offset of the block with ID `block_id`

47.21.2.4 get_group_offset()

```
template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_group_offset
(
    IndexType block_id ) const [inline], [noexcept]
```

Returns the offset of the group belonging to the block with the given ID.

Parameters

<i>block_id</i>	the ID of the block
-----------------	---------------------

Returns

the offset of the group belonging to block with ID `block_id`

Referenced by `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_global_block_offset()`.

47.21.2.5 get_group_size()

```
template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_group_size (
) const [inline], [noexcept]
```

Returns the number of elements in the group.

Returns

the number of elements in the group

Referenced by `gko::preconditioner::block_interleaved_storage_scheme< index_type >::compute_storage_space()`, and `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_block_offset()`.

47.21.2.6 get_stride()

```
template<typename IndexType>
IndexType gko::preconditioner::block_interleaved_storage_scheme< IndexType >::get_stride ( )
const [inline], [noexcept]
```

Returns the stride between columns of the block.

Returns

stride between columns of the block

47.21.3 Member Data Documentation

47.21.3.1 group_power

```
template<typename IndexType>
uint32 gko::preconditioner::block_interleaved_storage_scheme< IndexType >::group_power
```

Then base 2 power of the group.

I.e. the group contains $1 \ll \text{group_power}$ elements.

Referenced by `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_group_offset()`, `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_group_size()`, and `gko::preconditioner::block_interleaved_storage_scheme< index_type >::get_stride()`.

The documentation for this struct was generated from the following file:

- `ginkgo/core/preconditioner/jacobi.hpp`

47.22 gko::BlockOperator Class Reference

A [BlockOperator](#) represents a linear operator that is partitioned into multiple blocks.

```
#include <ginkgo/core/base/block_operator.hpp>
```

Public Member Functions

- `dim< 2 > get_block_size ()` const
Get the block dimension of this, i.e.
- `const LinOp * block_at (size_type i, size_type j)` const
Const access to a specific block.
- `BlockOperator (const BlockOperator &other)`
Copy constructs a BlockOperator.
- `BlockOperator (BlockOperator &&other) noexcept`
Move constructs a BlockOperator.
- `BlockOperator & operator= (const BlockOperator &other)`
Copy assigns a BlockOperator.
- `BlockOperator & operator= (BlockOperator &&other)`
Move assigns a BlockOperator.

Static Public Member Functions

- `static std::unique_ptr< BlockOperator > create (std::shared_ptr< const Executor > exec)`
Create empty BlockOperator.
- `static std::unique_ptr< BlockOperator > create (std::shared_ptr< const Executor > exec, std::vector< std::vector< std::shared_ptr< const LinOp >>> blocks)`
Create BlockOperator from the given blocks.

47.22.1 Detailed Description

A [BlockOperator](#) represents a linear operator that is partitioned into multiple blocks.

For example, a [BlockOperator](#) can be used to define the operator:

```
| A B |
| C D |
```

where A, B, C, D itself are matrices of compatible size. This can be created with:

```
{c++}
std::shared_ptr<LinOp> A = ...;
std::shared_ptr<LinOp> B = ...;
std::shared_ptr<LinOp> C = ...;
std::shared_ptr<LinOp> D = ...;
auto bop = BlockOperator::create(exec, {{A, B}, {C, D}});
```

The requirements on the individual blocks passed to the create method are:

- In each block-row, all blocks have the same number of rows
- In each block-column, all blocks have the same number of columns
- Each block-row must have the same number of blocks It is possible to set blocks to zero, by passing in a nullptr. But every block-row and block-column must contain at least one non-nullptr block.

The constructor will store all passed in blocks on the same executor as the [BlockOperator](#), which will requires copying any block that is associated with a different executor.

47.22.2 Constructor & Destructor Documentation

47.22.2.1 BlockOperator() [1/2]

```
gko::BlockOperator::BlockOperator (
    const BlockOperator & other )
```

Copy constructs a [BlockOperator](#).

The executor of other is used for this. The blocks of other are deep-copied into this, using clone.

47.22.2.2 BlockOperator() [2/2]

```
gko::BlockOperator::BlockOperator (
    BlockOperator && other ) [noexcept]
```

Move constructs a [BlockOperator](#).

The executor of other is used for this. All remaining data of other is moved into this. After this operation, other will be empty.

47.22.3 Member Function Documentation

47.22.3.1 block_at()

```
const LinOp* gko::BlockOperator::block_at (
    size_type i,
    size_type j ) const [inline]
```

Const access to a specific block.

Parameters

<i>i</i>	block row.
<i>j</i>	block column.

Returns

the block stored at (i, j).

```

97     {
98         GKO_ENSURE_IN_DIMENSION_BOUNDS(i, j, block_size_);
99         return blocks_[i * block_size_[1] + j].get();
100    }
```

47.22.3.2 create() [1/2]

```

static std::unique_ptr<BlockOperator> gko::BlockOperator::create (
    std::shared_ptr< const Executor > exec ) [static]
```

Create empty [BlockOperator](#).

Parameters

<i>exec</i>	the executor of this.
-------------	-----------------------

Returns

empty [BlockOperator](#).

47.22.3.3 create() [2/2]

```

static std::unique_ptr<BlockOperator> gko::BlockOperator::create (
    std::shared_ptr< const Executor > exec,
    std::vector< std::vector< std::shared_ptr< const LinOp >>> blocks ) [static]
```

Create [BlockOperator](#) from the given blocks.

Parameters

<i>exec</i>	the executor of this.
<i>blocks</i>	the blocks of this operator. The blocks will be used in a row-major form.

Returns

[BlockOperator](#) with the given blocks.

47.22.3.4 get_block_size()

```
dim<2> gko::BlockOperator::get_block_size ( ) const [inline]
```

Get the block dimension of this, i.e.

the number of blocks per row and column.

Returns

The block dimension of this.

47.22.3.5 operator=() [1/2]

```
BlockOperator& gko::BlockOperator::operator= (
    BlockOperator && other )
```

Move assigns a [BlockOperator](#).

The executor of this is not modified. All data of other (except its executor) is moved into this. If the executor of this and other differ, the blocks will be copied to the executor of this. After this operation, other will be empty.

47.22.3.6 operator=() [2/2]

```
BlockOperator& gko::BlockOperator::operator= (
    const BlockOperator & other )
```

Copy assigns a [BlockOperator](#).

The executor of this is not modified. The blocks of other are deep-copied into this, using clone.

The documentation for this class was generated from the following file:

- ginkgo/core/base/block_operator.hpp

47.23 gko::BlockSizeError< IndexType > Class Template Reference

[Error](#) that denotes issues between block sizes and matrix dimensions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [BlockSizeError](#) (const std::string &file, const int line, const int block_size, const IndexType size)

47.23.1 Detailed Description

```
template<typename IndexType>
class gko::BlockSizeError< IndexType >
```

[Error](#) that denotes issues between block sizes and matrix dimensions.

Template Parameters

<i>IndexType</i>	Type of index used by the linear algebra object that is incompatible with the required block size.
------------------	--

47.23.2 Constructor & Destructor Documentation

47.23.2.1 BlockSizeError()

```
template<typename IndexType >
gko::BlockSizeError< IndexType >::BlockSizeError (
    const std::string & file,
    const int line,
    const int block_size,
    const IndexType size ) [inline]
```

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>block_size</i>	Size of small dense blocks in a matrix
<i>size</i>	The size that is not exactly divided by the block size

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.24 gko::solver::CbGmres< ValueType > Class Template Reference

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

```
#include <ginkgo/core/solver/cb_gmres.hpp>
```

Public Member Functions

- `size_type get_krylov_dim () const`
Returns the Krylov dimension.
- `void set_krylov_dim (size_type other)`
Sets the Krylov dimension.
- `cb_gmres::storage_precision get_storage_precision () const`
Returns the storage precision used internally.

Static Public Member Functions

- static parameters_type `parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType >())

Create the parameters from the property_tree.

47.24.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::CbGmres< ValueType >
```

CB-GMRES or the compressed basis generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CB-GMRES are merged into 2 separate steps. Classical Gram-Schmidt with reorthogonalization is used.

The Krylov basis can be stored in reduced precision (compressed) to reduce memory accesses, while all computations (including Krylov basis operations) are performed in the same arithmetic precision `ValueType`. By default, the Krylov basis are stored in one precision lower than `ValueType`.

Template Parameters

<code>ValueType</code>	the arithmetic precision and the precision of matrix elements
------------------------	---

47.24.2 Member Function Documentation

47.24.2.1 `get_krylov_dim()`

```
template<typename ValueType = default_precision>
size_type gko::solver::CbGmres< ValueType >::get_krylov_dim ( ) const [inline]
```

Returns the Krylov dimension.

Returns

the Krylov dimension

```
110 { return parameters_.krylov_dim; }
```

47.24.2.2 `get_storage_precision()`

```
template<typename ValueType = default_precision>
cb_gmres::storage_precision gko::solver::CbGmres< ValueType >::get_storage_precision ( ) const
[inline]
```

Returns the storage precision used internally.

Returns

the storage precision used internally

47.24.2.3 `parse()`

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::CbGmres< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.24.2.4 `set_krylov_dim()`

```
template<typename ValueType = default_precision>
void gko::solver::CbGmres< ValueType >::set_krylov_dim (
    size_type other ) [inline]
```

Sets the Krylov dimension.

Parameters

<i>other</i>	the new Krylov dimension
--------------	--------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/solver/cb_gmres.hpp

47.25 gko::batch::solver::Cg< ValueType > Class Template Reference

[Cg](#) or the Conjugate Gradient is a Krylov subspace solver.

```
#include <ginkgo/core/solver/batch_cg.hpp>
```

Additional Inherited Members

47.25.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::solver::Cg< ValueType >
```

[Cg](#) or the Conjugate Gradient is a Krylov subspace solver.

It is a short recurrence solver that is generally used to solve linear systems with SPD matrices.

This solver solves a batch of linear systems using the [Cg](#) algorithm. Each linear system in the batch can converge independently.

Unless otherwise specified via the `preconditioner` factory parameter, this implementation does not use any preconditioner by default. The type of tolerance (absolute or relative) and the maximum number of iterations to be used in the stopping criterion can be set via the factory parameters.

Note

The tolerance check is against the internal residual computed within the solver process. This implicit (internal) residual can diverge from the true residual ($\|b - Ax\|$). A posteriori checks (by computing the true residual, $\|b - Ax\|$) are recommended to ensure that the solution has converged to the desired tolerance.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/solver/batch_cg.hpp

47.26 gko::solver::Cg< ValueType > Class Template Reference

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

```
#include <ginkgo/core/solver/cg.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.26.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Cg< ValueType >
```

CG or the conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

Though this method performs very well for symmetric positive definite matrices, it is in general not suitable for general matrices.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CG are merged into 2 separate steps.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.26.2 Member Function Documentation

47.26.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Cg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
67 { return true; }
```

47.26.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.26.2.3 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Cg< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.26.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/cg.hpp`

47.27 gko::solver::Cgs< ValueType > Class Template Reference

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

```
#include <ginkgo/core/solver/cgs.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ( ) const` override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ( ) const` override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ( ) const` override
Return true as iterative solvers use the data in x as an initial guess.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.27.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Cgs< ValueType >
```

CGS or the conjugate gradient square method is an iterative type Krylov subspace method which is suitable for general systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of CGS are merged into 3 separate steps.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.27.2 Member Function Documentation

47.27.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Cgs< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
64 { return true; }
```

47.27.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cgs< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.27.2.3 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Cgs< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.27.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Cgs< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/cgs.hpp

47.28 gko::experimental::factorization::Cholesky< ValueType, IndexType > Class Template Reference

Computes a [Cholesky](#) factorization of a symmetric, positive-definite sparse matrix.

```
#include <ginkgo/core/factorization/cholesky.hpp>
```

Public Member Functions

- const parameters_type & [get_parameters](#) ()
Returns the parameters used to construct the factory.
- const parameters_type & [get_parameters](#) () const
Returns the parameters used to construct the factory.
- std::unique_ptr< [factorization_type](#) > [generate](#) (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

- static parameters_type [build](#) ()
Creates a new parameter_type to set up the factory.
- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType, IndexType >())
Create the parameters from the property_tree.

47.28.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::experimental::factorization::Cholesky< ValueType, IndexType >
```

Computes a [Cholesky](#) factorization of a symmetric, positive-definite sparse matrix.

This [LinOpFactory](#) returns a [Factorization](#) storing the L and L^H factors for the provided system matrix in [matrix::Csr](#) format. If no symbolic factorization is provided, it will be computed first. It expects all fill-in entries to be present in the symbolic factorization. If the symbolic factorization is missing some entries, please refer to [lc](#).

Template Parameters

<i>ValueType</i>	the type used to store values of the system matrix
<i>IndexType</i>	the type used to store sparsity pattern indices of the system matrix

47.28.2 Member Function Documentation

47.28.2.1 generate()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<factorization\_type> gko::experimental::factorization::Cholesky< ValueType,
IndexType >::generate (
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default [LinOpFactory::generate](#) to return a [Factorization](#) instead of a generic [LinOp](#), which would need to be cast to [Factorization](#) again to access its factors. It is only necessary because smart pointers aren't covariant.

47.28.2.2 get_parameters() [1/2]

```
template<typename ValueType , typename IndexType >
const parameters_type& gko::experimental::factorization::Cholesky< ValueType, IndexType >↵
::get_parameters ( ) [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
79 { return parameters_; }
```

47.28.2.3 get_parameters() [2/2]

```
template<typename ValueType , typename IndexType >
const parameters_type& gko::experimental::factorization::Cholesky< ValueType, IndexType >↵
::get_parameters ( ) const [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

47.28.2.4 parse()

```
template<typename ValueType , typename IndexType >
static parameters_type gko::experimental::factorization::Cholesky< ValueType, IndexType >↵
::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/cholesky.hpp

47.29 gko::matrix::Csr< ValueType, IndexType >::classical Class Reference

classical is a [strategy_type](#) which uses the same number of threads on each row.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [classical](#) ()
Creates a classical strategy.
- void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow) override
Computes srow according to row pointers.
- int64_t [clac_size](#) (const int64_t nnz) override
Computes the srow size according to the number of nonzeros.
- std::shared_ptr< [strategy_type](#) > [copy](#) () override
Copy a strategy.

47.29.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::classical
```

classical is a [strategy_type](#) which uses the same number of threads on each row.

Classical strategy uses multithreads to calculate on parts of rows and then do a reduction of these threads results. The number of threads per row depends on the max number of stored elements per row.

47.29.2 Member Function Documentation

47.29.2.1 clac_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int64_t gko::matrix::Csr< ValueType, IndexType >::classical::clac_size (
    const int64_t nnz ) [inline], [override], [virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.29.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::classical::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.29.2.3 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::classical::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [inline], [override], [virtual]
```

Computes srow according to row pointers.

Parameters

<i>mtx_row_ptrs</i>	the row pointers of the matrix
<i>mtx_srow</i>	the srow of the matrix

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

References [gko::array< ValueType >::get_const_data\(\)](#), [gko::array< ValueType >::get_executor\(\)](#), and [gko::array< ValueType >::get_size\(\)](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/csr.hpp`

47.30 gko::matrix::Hybrid< ValueType, IndexType >::column_limit Class Reference

`column_limit` is a [strategy_type](#) which decides the number of stored elements per row of the ell part by specifying the number of columns.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- `column_limit` ([size_type](#) num_column=0)
Creates a `column_limit` strategy.
- `size_type compute_ell_num_stored_elements_per_row` ([array](#)< [size_type](#) > *row_nnz) const override
Computes the number of stored elements per row of the ell part.
- auto `get_num_columns` () const
Get the number of columns limit.

47.30.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::column_limit
```

`column_limit` is a [strategy_type](#) which decides the number of stored elements per row of the ell part by specifying the number of columns.

47.30.2 Constructor & Destructor Documentation

47.30.2.1 column_limit()

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Hybrid< ValueType, IndexType >::column_limit::column_limit (
    size\_type num_column = 0 ) [inline], [explicit]
```

Creates a `column_limit` strategy.

Parameters

<code>num_column</code>	the specified number of columns of the ell part
-------------------------	---

47.30.3 Member Function Documentation

47.30.3.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::column_limit::compute_ell_num_stored_↵
elements_per_row (
    array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<code>row_nnz</code>	the number of nonzeros of each row
----------------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type](#).

47.30.3.2 get_num_columns()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::column_limit::get_num_columns ( ) const
[inline]
```

Get the number of columns limit.

Returns

the number of columns limit

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/hybrid.hpp`

47.31 gko::Combination< ValueType > Class Template Reference

The [Combination](#) class can be used to construct a linear combination of multiple linear operators $c1 * op1 + c2 * op2 + \dots$

```
#include <ginkgo/core/base/combination.hpp>
```


Public Member Functions

- `const std::vector< std::shared_ptr< const LinOp > > & get_coefficients ()` const noexcept
Returns a list of coefficients of the combination.
- `const std::vector< std::shared_ptr< const LinOp > > & get_operators ()` const noexcept
Returns a list of operators of the combination.
- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `Combination & operator= (const Combination &)`
Copy-assigns a [Combination](#).
- `Combination & operator= (Combination &&)`
Move-assigns a [Combination](#).
- `Combination (const Combination &)`
Copy-constructs a [Combination](#).
- `Combination (Combination &&)`
Move-constructs a [Combination](#).

47.31.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::Combination< ValueType >
```

The [Combination](#) class can be used to construct a linear combination of multiple linear operators $c_1 * op_1 + c_2 * op_2 + \dots$

- $ck * opk$.

[Combination](#) ensures that all LinOps passed to its constructor use the same executor, and if not, copies the operators to the executor of the first operator.

Template Parameters

<i>ValueType</i>	precision of input and result vectors
------------------	---------------------------------------

47.31.2 Constructor & Destructor Documentation

47.31.2.1 Combination() [1/2]

```
template<typename ValueType = default_precision>
gko::Combination< ValueType >::Combination (
    const Combination< ValueType > & )
```

Copy-constructs a [Combination](#).

This inherits the executor of the input [Combination](#) and all of its operators with shared ownership.

47.31.2.2 Combination() [2/2]

```
template<typename ValueType = default_precision>
gko::Combination< ValueType >::Combination (
    Combination< ValueType > && )
```

Move-constructs a [Combination](#).

This inherits the executor of the input [Combination](#) and all of its operators. The moved-from object is empty (0x0 LinOp without operators) afterwards.

47.31.3 Member Function Documentation

47.31.3.1 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Combination< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.31.3.2 get_coefficients()

```
template<typename ValueType = default_precision>
const std::vector<std::shared_ptr<const LinOp> >& gko::Combination< ValueType >::get_↵
coefficients ( ) const [inline], [noexcept]
```

Returns a list of coefficients of the combination.

Returns

a list of coefficients

```
48     {
49         return coefficients_;
50     }
```

47.31.3.3 get_operators()

```
template<typename ValueType = default_precision>
const std::vector<std::shared_ptr<const LinOp> >& gko::Combination< ValueType >::get_←
operators ( ) const [inline], [noexcept]
```

Returns a list of operators of the combination.

Returns

a list of operators

47.31.3.4 operator=() [1/2]

```
template<typename ValueType = default_precision>
Combination& gko::Combination< ValueType >::operator= (
    Combination< ValueType > && )
```

Move-assigns a [Combination](#).

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor, otherwise they share ownership. The moved-from object is empty (0x0 LinOp without operators) afterwards.

47.31.3.5 operator=() [2/2]

```
template<typename ValueType = default_precision>
Combination& gko::Combination< ValueType >::operator= (
    const Combination< ValueType > & )
```

Copy-assigns a [Combination](#).

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor.

47.31.3.6 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Combination< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/combination.hpp

47.32 gko::stop::Combined Class Reference

The [Combined](#) class is used to combine multiple criterions together through an OR operation.

```
#include <ginkgo/core/stop/combined.hpp>
```

47.32.1 Detailed Description

The [Combined](#) class is used to combine multiple criterions together through an OR operation.

The typical use case is to stop the iteration process if any of the criteria is fulfilled, e.g. a number of iterations, the relative residual norm has reached a threshold, etc.

The documentation for this class was generated from the following file:

- ginkgo/core/stop/combined.hpp

47.33 gko::experimental::mpi::communicator Class Reference

A thin wrapper of MPI_Comm that supports most MPI calls.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

- [communicator](#) (const MPI_Comm &comm, bool force_host_buffer=false)
Non-owning constructor for an existing communicator of type MPI_Comm.
- [communicator](#) (const MPI_Comm &comm, int color, int key)
Create a communicator object from an existing MPI_Comm object using color and key.
- [communicator](#) (const [communicator](#) &comm, int color, int key)
Create a communicator object from an existing MPI_Comm object using color and key.
- const MPI_Comm & [get](#) () const
Return the underlying MPI_Comm object.
- int [size](#) () const
Return the size of the communicator (number of ranks).
- int [rank](#) () const
Return the rank of the calling process in the communicator.
- int [node_local_rank](#) () const
Return the node local rank of the calling process in the communicator.
- bool [operator==](#) (const [communicator](#) &rhs) const
Compare two communicator objects for equality.
- bool [operator!=](#) (const [communicator](#) &rhs) const
Compare two communicator objects for non-equality.
- void [synchronize](#) () const
This function is used to synchronize the ranks in the communicator.
- template<typename SendType >
void [send](#) (std::shared_ptr< const [Executor](#) > exec, const SendType *send_buffer, const int send_count, const int destination_rank, const int send_tag) const

Send (Blocking) data from calling process to destination rank.

- template<typename SendType >
request i_send (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, const int destination_rank, const int send_tag) const

Send (Non-blocking, Immediate return) data from calling process to destination rank.

- template<typename RecvType >
status recv (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count, const int source_rank, const int recv_tag) const

Receive data from source rank.

- template<typename RecvType >
request i_recv (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count, const int source_rank, const int recv_tag) const

Receive (Non-blocking, Immediate return) data from source rank.

- template<typename BroadcastType >
void broadcast (std::shared_ptr< const Executor > exec, BroadcastType *buffer, int count, int root_rank) const

Broadcast data from calling process to all ranks in the communicator.

- template<typename BroadcastType >
request i_broadcast (std::shared_ptr< const Executor > exec, BroadcastType *buffer, int count, int root_rank) const

(Non-blocking) Broadcast data from calling process to all ranks in the communicator

- template<typename ReduceType >
void reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv_buffer, int count, MPI_Op operation, int root_rank) const

Reduce data into root from all calling processes on the same communicator.

- template<typename ReduceType >
request i_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv_buffer, int count, MPI_Op operation, int root_rank) const

(Non-blocking) Reduce data into root from all calling processes on the same communicator.

- template<typename ReduceType >
void all_reduce (std::shared_ptr< const Executor > exec, ReduceType *recv_buffer, int count, MPI_Op operation) const

(In-place) Reduce data from all calling processes from all calling processes on same communicator.

- template<typename ReduceType >
request i_all_reduce (std::shared_ptr< const Executor > exec, ReduceType *recv_buffer, int count, MPI_Op operation) const

(In-place, non-blocking) Reduce data from all calling processes from all calling processes on same communicator.

- template<typename ReduceType >
void all_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv_buffer, int count, MPI_Op operation) const

Reduce data from all calling processes from all calling processes on same communicator.

- template<typename ReduceType >
request i_all_reduce (std::shared_ptr< const Executor > exec, const ReduceType *send_buffer, ReduceType *recv_buffer, int count, MPI_Op operation) const

Reduce data from all calling processes from all calling processes on same communicator.

- template<typename SendType, typename RecvType >
void gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count, int root_rank) const

Gather data onto the root rank from all ranks in the communicator.

- template<typename SendType, typename RecvType >
request i_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count, int root_rank) const

(Non-blocking) Gather data onto the root rank from all ranks in the communicator.

- `template<typename SendType , typename RecvType >`
`void gather\_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int *recv_counts, const int *displacements, int root_rank) const`
Gather data onto the root rank from all ranks in the communicator with offsets.
- `template<typename SendType , typename RecvType >`
`request\_i\_gather\_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int *recv_counts, const int *displacements, int root_rank) const`
(Non-blocking) Gather data onto the root rank from all ranks in the communicator with offsets.
- `template<typename SendType , typename RecvType >`
`void all\_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count) const`
Gather data onto all ranks from all ranks in the communicator.
- `template<typename SendType , typename RecvType >`
`request\_i\_all\_gather (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count) const`
(Non-blocking) Gather data onto all ranks from all ranks in the communicator.
- `template<typename SendType , typename RecvType >`
`void scatter (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count, int root_rank) const`
Scatter data from root rank to all ranks in the communicator.
- `template<typename SendType , typename RecvType >`
`request\_i\_scatter (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count, int root_rank) const`
(Non-blocking) Scatter data from root rank to all ranks in the communicator.
- `template<typename SendType , typename RecvType >`
`void scatter\_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int *send_counts, const int *displacements, RecvType *recv_buffer, const int recv_count, int root_rank) const`
Scatter data from root rank to all ranks in the communicator with offsets.
- `template<typename SendType , typename RecvType >`
`request\_i\_scatter\_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int *send_counts, const int *displacements, RecvType *recv_buffer, const int recv_count, int root_rank) const`
(Non-blocking) Scatter data from root rank to all ranks in the communicator with offsets.
- `template<typename RecvType >`
`void all\_to\_all (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count) const`
(In-place) Communicate data from all ranks to all other ranks in place (MPI_Alltoall).
- `template<typename RecvType >`
`request\_i\_all\_to\_all (std::shared_ptr< const Executor > exec, RecvType *recv_buffer, const int recv_count) const`
(In-place, Non-blocking) Communicate data from all ranks to all other ranks in place (MPI_lalltoall).
- `template<typename SendType , typename RecvType >`
`void all\_to\_all (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count) const`
Communicate data from all ranks to all other ranks (MPI_Alltoall).
- `template<typename SendType , typename RecvType >`
`request\_i\_all\_to\_all (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int send_count, RecvType *recv_buffer, const int recv_count) const`
(Non-blocking) Communicate data from all ranks to all other ranks (MPI_lalltoall).
- `template<typename SendType , typename RecvType >`
`void all\_to\_all\_v (std::shared_ptr< const Executor > exec, const SendType *send_buffer, const int *send_counts, const int *send_offsets, RecvType *recv_buffer, const int *recv_counts, const int *recv_offsets) const`
Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).
- `void all\_to\_all\_v (std::shared_ptr< const Executor > exec, const void *send_buffer, const int *send_counts, const int *send_offsets, MPI_Datatype send_type, void *recv_buffer, const int *recv_counts, const int *recv_offsets, MPI_Datatype recv_type) const`

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

- `request i_all_to_all_v` (std::shared_ptr< const [Executor](#) > exec, const void *send_buffer, const int *send_counts, const int *send_offsets, MPI_Datatype send_type, void *recv_buffer, const int *recv_counts, const int *recv_offsets, MPI_Datatype recv_type) const

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

- template<typename SendType , typename RecvType >
`request i_all_to_all_v` (std::shared_ptr< const [Executor](#) > exec, const SendType *send_buffer, const int *send_counts, const int *send_offsets, RecvType *recv_buffer, const int *recv_counts, const int *recv_offsets) const

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

- template<typename ScanType >
void `scan` (std::shared_ptr< const [Executor](#) > exec, const ScanType *send_buffer, ScanType *recv_buffer, int count, MPI_Op operation) const

Does a scan operation with the given operator.

- template<typename ScanType >
`request i_scan` (std::shared_ptr< const [Executor](#) > exec, const ScanType *send_buffer, ScanType *recv_buffer, int count, MPI_Op operation) const

Does a scan operation with the given operator.

47.33.1 Detailed Description

A thin wrapper of MPI_Comm that supports most MPI calls.

A wrapper class that takes in the given MPI communicator. If a bare MPI_Comm is provided, the wrapper takes no ownership of the MPI_Comm. Thus the MPI_Comm must remain valid throughout the lifetime of the communicator. If the communicator was created through splitting, the wrapper takes ownership of the MPI_Comm. In this case, as the class or object goes out of scope, the underlying MPI_Comm is freed.

Note

All MPI calls that work on a buffer take in an [Executor](#) as an additional argument. This argument specifies the memory space the buffer lives in.

47.33.2 Constructor & Destructor Documentation

47.33.2.1 communicator() [1/3]

```
gko::experimental::mpi::communicator::communicator (
    const MPI_Comm & comm,
    bool force_host_buffer = false ) [inline]
```

Non-owning constructor for an existing communicator of type MPI_Comm.

The MPI_Comm object will not be deleted after the communicator object has been freed and an explicit MPI_Comm_free needs to be called on the original MPI_Comm object.

Parameters

<code>comm</code>	The input MPI_Comm object.
<code>force_host_buffer</code>	If set to true, always communicates through host memory

47.33.2.2 communicator() [2/3]

```
gko::experimental::mpi::communicator::communicator (
    const MPI_Comm & comm,
    int color,
    int key ) [inline]
```

Create a communicator object from an existing MPI_Comm object using color and key.

Parameters

<i>comm</i>	The input MPI_Comm object.
<i>color</i>	The color to split the original comm object
<i>key</i>	The key to split the comm object

47.33.2.3 communicator() [3/3]

```
gko::experimental::mpi::communicator::communicator (
    const communicator & comm,
    int color,
    int key ) [inline]
```

Create a communicator object from an existing MPI_Comm object using color and key.

Parameters

<i>comm</i>	The input communicator object.
<i>color</i>	The color to split the original comm object
<i>key</i>	The key to split the comm object

References [get\(\)](#).

47.33.3 Member Function Documentation

47.33.3.1 all_gather()

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::all_gather (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
```



```
RecvType * recv_buffer,  
const int recv_count ) const [inline]
```

Gather data onto all ranks from all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References [get\(\)](#).

47.33.3.2 all_reduce() [1/2]

```
template<typename ReduceType >
void gko::experimental::mpi::communicator::all_reduce (
    std::shared_ptr< const Executor > exec,
    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Reduce data from all calling processes from all calling processes on same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the data to reduce
<i>recv_buffer</i>	the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the reduce operation. See @MPI_Op

Template Parameters

<i>ReduceType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-------------------	---

References [get\(\)](#).

47.33.3.3 all_reduce() [2/2]

```
template<typename ReduceType >
```

```
void gko::experimental::mpi::communicator::all_reduce (
    std::shared_ptr< const Executor > exec,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

(In-place) Reduce data from all calling processes from all calling processes on same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>recv_buffer</i>	the data to reduce and the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the MPI_Op type reduce operation.

Template Parameters

<i>ReduceType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-------------------	---

References [get\(\)](#).

47.33.3.4 all_to_all() [1/2]

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::all_to_all (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count ) const [inline]
```

Communicate data from all ranks to all other ranks (MPI_Alltoall).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to receive
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References `get()`.

47.33.3.5 `all_to_all()` [2/2]

```
template<typename RecvType >
void gko::experimental::mpi::communicator::all_to_all (
    std::shared_ptr< const Executor > exec,
    RecvType * recv_buffer,
    const int recv_count ) const [inline]
```

(In-place) Communicate data from all ranks to all other ranks in place (MPI_Alltoall).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>buffer</i>	the buffer to send and the buffer receive
<i>recv_count</i>	the number of elements to receive
<i>comm</i>	the communicator

Template Parameters

<i>RecvType</i>	the type of the data to send. Has to be a type which has a specialization of <code>type_impl</code> that defines its <code>MPI_Datatype</code> .
-----------------	--

Note

This overload uses `MPI_IN_PLACE` and the source and destination buffers are the same.

References `get()`.

47.33.3.6 `all_to_all_v()` [1/2]

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::all_to_all_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int * send_counts,
    const int * send_offsets,
    RecvType * recv_buffer,
    const int * recv_counts,
    const int * recv_offsets ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>send_offsets</i>	the offsets for the send buffer
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>recv_offsets</i>	the offsets for the recv buffer
<i>comm</i>	the communicator

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

47.33.3.7 all_to_all_v() [2/2]

```
void gko::experimental::mpi::communicator::all_to_all_v (
    std::shared_ptr< const Executor > exec,
    const void * send_buffer,
    const int * send_counts,
    const int * send_offsets,
    MPI_Datatype send_type,
    void * recv_buffer,
    const int * recv_counts,
    const int * recv_offsets,
    MPI_Datatype recv_type ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_Alltoallv).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>send_offsets</i>	the offsets for the send buffer
<i>send_type</i>	the MPI_Datatype for the send buffer
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>recv_offsets</i>	the offsets for the recv buffer
<i>recv_type</i>	the MPI_Datatype for the recv buffer
<i>comm</i>	the communicator

References [get\(\)](#).

47.33.3.8 broadcast()

```
template<typename BroadcastType >
void gko::experimental::mpi::communicator::broadcast (
    std::shared_ptr< const Executor > exec,
    BroadcastType * buffer,
    int count,
    int root_rank ) const [inline]
```

Broadcast data from calling process to all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>buffer</i>	the buffer to broadcast
<i>count</i>	the number of elements to broadcast
<i>root_rank</i>	the rank to broadcast from

Template Parameters

<i>BroadcastType</i>	the type of the data to broadcast. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
----------------------	--

References [get\(\)](#).

47.33.3.9 gather()

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::gather (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

Gather data onto the root rank from all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>root_rank</i>	the rank to gather into

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References [get\(\)](#).

47.33.3.10 gather_v()

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::gather_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int * recv_counts,
    const int * displacements,
    int root_rank ) const [inline]
```

Gather data onto the root rank from all ranks in the communicator with offsets.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>displacements</i>	the offsets for the buffer
<i>root_rank</i>	the rank to gather into

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References [get\(\)](#).

47.33.3.11 get()

```
const MPI_Comm& gko::experimental::mpi::communicator::get ( ) const [inline]
```

Return the underlying MPI_Comm object.

Returns

the MPI_Comm object

Referenced by `all_gather()`, `all_reduce()`, `all_to_all()`, `all_to_all_v()`, `broadcast()`, `communicator()`, `gather()`, `gather_v()`, `i_all_gather()`, `i_all_reduce()`, `i_all_to_all()`, `i_all_to_all_v()`, `i_broadcast()`, `i_gather()`, `i_gather_v()`, `i_recv()`, `i_reduce()`, `i_scan()`, `i_scatter()`, `i_scatter_v()`, `i_send()`, `recv()`, `reduce()`, `scan()`, `scatter()`, `scatter_v()`, `send()`, `synchronize()`, and `gko::experimental::mpi::window< ValueType >::window()`.

47.33.3.12 i_all_gather()

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_all_gather (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count ) const [inline]
```

(Non-blocking) Gather data onto all ranks from all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of <code>type_impl</code> that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.13 i_all_reduce() [1/2]

```
template<typename ReduceType >
request gko::experimental::mpi::communicator::i_all_reduce (
    std::shared_ptr< const Executor > exec,
```



```

    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]

```

Reduce data from all calling processes from all calling processes on same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the data to reduce
<i>recv_buffer</i>	the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the reduce operation. See @MPI_Op

Template Parameters

<i>ReduceType</i>	the type of the data to reduce. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-------------------	---

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.14 i_all_reduce() [2/2]

```

template<typename ReduceType >
request gko::experimental::mpi::communicator::i_all_reduce (
    std::shared_ptr< const Executor > exec,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]

```

(In-place, non-blocking) Reduce data from all calling processes from all calling processes on same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>recv_buffer</i>	the data to reduce and the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the reduce operation. See @MPI_Op

Template Parameters

<i>ReduceType</i>	the type of the data to reduce. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-------------------	---

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.15 i_all_to_all() [1/2]

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_all_to_all (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count ) const [inline]
```

(Non-blocking) Communicate data from all ranks to all other ranks (MPI_lalltoall).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to receive
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of <code>type_impl</code> that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.16 i_all_to_all() [2/2]

```
template<typename RecvType >
request gko::experimental::mpi::communicator::i_all_to_all (
    std::shared_ptr< const Executor > exec,
    RecvType * recv_buffer,
    const int recv_count ) const [inline]
```

(In-place, Non-blocking) Communicate data from all ranks to all other ranks in place (MPI_lalltoall).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>buffer</i>	the buffer to send and the buffer receive
<i>recv_count</i>	the number of elements to receive
<i>comm</i>	the communicator

Template Parameters

<i>RecvType</i>	the type of the data to receive. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	--

Returns

the request handle for the call

Note

This overload uses MPI_IN_PLACE and the source and destination buffers are the same.

References [gko::experimental::mpi::request::get\(\)](#), and [get\(\)](#).

47.33.3.17 i_all_to_all_v() [1/2]

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_all_to_all_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int * send_counts,
    const int * send_offsets,
    RecvType * recv_buffer,
    const int * recv_counts,
    const int * recv_offsets ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>send_offsets</i>	the offsets for the send buffer
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>recv_offsets</i>	the offsets for the recv buffer

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References [i_all_to_all_v\(\)](#).

47.33.3.18 i_all_to_all_v() [2/2]

```
request gko::experimental::mpi::communicator::i_all_to_all_v (
    std::shared_ptr< const Executor > exec,
    const void * send_buffer,
    const int * send_counts,
    const int * send_offsets,
    MPI_Datatype send_type,
    void * recv_buffer,
    const int * recv_counts,
    const int * recv_offsets,
    MPI_Datatype recv_type ) const [inline]
```

Communicate data from all ranks to all other ranks with offsets (MPI_lalltoallv).

See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>send_offsets</i>	the offsets for the send buffer
<i>send_type</i>	the MPI_Datatype for the send buffer
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>recv_offsets</i>	the offsets for the recv buffer
<i>recv_type</i>	the MPI_Datatype for the recv buffer

Returns

the request handle for the call

Note

This overload allows specifying the MPI_Datatype for both the send and received data.

References gko::experimental::mpi::request::get(), and get().

Referenced by i_all_to_all_v().

47.33.3.19 i_broadcast()

```
template<typename BroadcastType >
request gko::experimental::mpi::communicator::i_broadcast (
    std::shared_ptr< const Executor > exec,
    BroadcastType * buffer,
    int count,
    int root_rank ) const [inline]
```

(Non-blocking) Broadcast data from calling process to all ranks in the communicator

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>buffer</i>	the buffer to broadcast
<i>count</i>	the number of elements to broadcast
<i>root_rank</i>	the rank to broadcast from

Template Parameters

<i>BroadcastType</i>	the type of the data to broadcast. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
----------------------	--

Returns

the request handle for the call

References gko::experimental::mpi::request::get(), and get().

47.33.3.20 i_gather()

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_gather (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

(Non-blocking) Gather data onto the root rank from all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>root_rank</i>	the rank to gather into

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.21 i_gather_v()

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_gather_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int * recv_counts,
    const int * displacements,
    int root_rank ) const [inline]
```

(Non-blocking) Gather data onto the root rank from all ranks in the communicator with offsets.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>displacements</i>	the offsets for the buffer
<i>root_rank</i>	the rank to gather into

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.22 i_recv()

```
template<typename RecvType >
request gko::experimental::mpi::communicator::i_recv (
    std::shared_ptr< const Executor > exec,
    RecvType * recv_buffer,
    const int recv_count,
    const int source_rank,
    const int recv_tag ) const [inline]
```

Receive (Non-blocking, Immediate return) data from source rank.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>recv_buffer</i>	the buffer to send
<i>recv_count</i>	the number of elements to receive
<i>source_rank</i>	the rank to receive the data from
<i>recv_tag</i>	the tag for the recv call

Template Parameters

<i>RecvType</i>	the type of the data to receive. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	--

Returns

the request handle for the recv call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.23 i_reduce()

```
template<typename ReduceType >
request gko::experimental::mpi::communicator::i_reduce (
    std::shared_ptr< const Executor > exec,
    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation,
    int root_rank ) const [inline]
```

(Non-blocking) Reduce data into root from all calling processes on the same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>send_buffer</i>	the buffer to reduce
<i>recv_buffer</i>	the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the MPI_Op type reduce operation.

Template Parameters

<i>ReduceType</i>	the type of the data to reduce. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-------------------	---

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.24 i_scan()

```
template<typename ScanType >
request gko::experimental::mpi::communicator::i_scan (
    std::shared_ptr< const Executor > exec,
    const ScanType * send_buffer,
    ScanType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Does a scan operation with the given operator.

(MPI_Iscan). See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to scan from
<i>recv_buffer</i>	the result buffer
<i>recv_count</i>	the number of elements to scan
<i>operation</i>	the operation type to be used for the scan. See @MPI_Op

Template Parameters

<i>ScanType</i>	the type of the data to scan. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	---

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.25 i_scatter()

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_scatter (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

(Non-blocking) Scatter data from root rank to all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.26 i_scatter_v()

```
template<typename SendType , typename RecvType >
request gko::experimental::mpi::communicator::i_scatter_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int * send_counts,
    const int * displacements,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

(Non-blocking) Scatter data from root rank to all ranks in the communicator with offsets.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>displacements</i>	the offsets for the buffer
<i>comm</i>	the communicator

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

Returns

the request handle for the call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.27 i_send()

```
template<typename SendType >
request gko::experimental::mpi::communicator::i_send (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    const int destination_rank,
    const int send_tag ) const [inline]
```

Send (Non-blocking, Immediate return) data from calling process to destination rank.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>destination_rank</i>	the rank to send the data to
<i>send_tag</i>	the tag for the send call

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	---

Returns

the request handle for the send call

References `gko::experimental::mpi::request::get()`, and `get()`.

47.33.3.28 node_local_rank()

```
int gko::experimental::mpi::communicator::node_local_rank ( ) const [inline]
```

Return the node local rank of the calling process in the communicator.

Returns

the node local rank

47.33.3.29 operator"!="()

```
bool gko::experimental::mpi::communicator::operator!= (
    const communicator & rhs ) const [inline]
```

Compare two communicator objects for non-equality.

Returns

if the two comm objects are not equal

47.33.3.30 operator==()

```
bool gko::experimental::mpi::communicator::operator== (
    const communicator & rhs ) const [inline]
```

Compare two communicator objects for equality.

Returns

if the two comm objects are equal

47.33.3.31 rank()

```
int gko::experimental::mpi::communicator::rank ( ) const [inline]
```

Return the rank of the calling process in the communicator.

Returns

the rank

47.33.3.32 recv()

```
template<typename RecvType >
status gko::experimental::mpi::communicator::recv (
    std::shared_ptr< const Executor > exec,
    RecvType * recv_buffer,
    const int recv_count,
    const int source_rank,
    const int recv_tag ) const [inline]
```

Receive data from source rank.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>recv_buffer</i>	the buffer to receive
<i>recv_count</i>	the number of elements to receive
<i>source_rank</i>	the rank to receive the data from
<i>recv_tag</i>	the tag for the recv call

Template Parameters

<i>RecvType</i>	the type of the data to receive. Has to be a type which has a specialization of <code>type_impl</code> that defines its <code>MPI_Datatype</code> .
-----------------	---

Returns

the status of completion of this call

References `gko::experimental::mpi::status::get()`, and `get()`.

47.33.3.33 reduce()

```
template<typename ReduceType >
void gko::experimental::mpi::communicator::reduce (
    std::shared_ptr< const Executor > exec,
    const ReduceType * send_buffer,
    ReduceType * recv_buffer,
    int count,
    MPI_Op operation,
    int root_rank ) const [inline]
```

Reduce data into root from all calling processes on the same communicator.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>send_buffer</i>	the buffer to reduce
<i>recv_buffer</i>	the reduced result
<i>count</i>	the number of elements to reduce
<i>operation</i>	the MPI_Op type reduce operation.

Template Parameters

<i>ReduceType</i>	the type of the data to reduce. Has to be a type which has a specialization of <code>type_impl</code> that defines its MPI_Datatype.
-------------------	--

References `get()`.

47.33.3.34 scan()

```
template<typename ScanType >
void gko::experimental::mpi::communicator::scan (
    std::shared_ptr< const Executor > exec,
    const ScanType * send_buffer,
    ScanType * recv_buffer,
    int count,
    MPI_Op operation ) const [inline]
```

Does a scan operation with the given operator.

(MPI_Scan). See MPI documentation for more details.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to scan from
<i>recv_buffer</i>	the result buffer
<i>recv_count</i>	the number of elements to scan
<i>operation</i>	the operation type to be used for the scan. See @MPI_Op

Template Parameters

<i>ScanType</i>	the type of the data to scan. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	---

References [get\(\)](#).

47.33.3.35 scatter()

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::scatter (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

Scatter data from root rank to all ranks in the communicator.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References [get\(\)](#).

47.33.3.36 scatter_v()

```
template<typename SendType , typename RecvType >
void gko::experimental::mpi::communicator::scatter_v (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int * send_counts,
    const int * displacements,
    RecvType * recv_buffer,
    const int recv_count,
    int root_rank ) const [inline]
```

Scatter data from root rank to all ranks in the communicator with offsets.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>send_buffer</i>	the buffer to gather from
<i>send_count</i>	the number of elements to send
<i>recv_buffer</i>	the buffer to gather into
<i>recv_count</i>	the number of elements to receive
<i>displacements</i>	the offsets for the buffer
<i>comm</i>	the communicator

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
<i>RecvType</i>	the type of the data to receive. The same restrictions as for SendType apply.

References [get\(\)](#).

47.33.3.37 send()

```
template<typename SendType >
void gko::experimental::mpi::communicator::send (
    std::shared_ptr< const Executor > exec,
    const SendType * send_buffer,
    const int send_count,
    const int destination_rank,
    const int send_tag ) const [inline]
```

Send (Blocking) data from calling process to destination rank.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>send_buffer</i>	the buffer to send
<i>send_count</i>	the number of elements to send
<i>destination_rank</i>	the rank to send the data to
<i>send_tag</i>	the tag for the send call

Template Parameters

<i>SendType</i>	the type of the data to send. Has to be a type which has a specialization of type_impl that defines its MPI_Datatype.
-----------------	---

References [get\(\)](#).

47.33.3.38 size()

```
int gko::experimental::mpi::communicator::size ( ) const [inline]
```

Return the size of the communicator (number of ranks).

Returns

the size

47.33.3.39 synchronize()

```
void gko::experimental::mpi::communicator::synchronize ( ) const [inline]
```

This function is used to synchronize the ranks in the communicator.

Calls [MPI_Barrier](#)

References [get\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/mpi.hpp](#)

47.34 gko::Composition< ValueType > Class Template Reference

The [Composition](#) class can be used to compose linear operators `op1`, `op2`, ..., `opn` and obtain the operator `op1 * op2 * ...`.

```
#include <ginkgo/core/base/composition.hpp>
```


Public Member Functions

- `const std::vector< std::shared_ptr< const LinOp > > & get_operators ()` const noexcept
Returns a list of operators of the composition.
- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `Composition & operator= (const Composition &)`
Copy-assigns a [Composition](#).
- `Composition & operator= (Composition &&)`
Move-assigns a [Composition](#).
- `Composition (const Composition &)`
Copy-constructs a [Composition](#).
- `Composition (Composition &&)`
Move-constructs a [Composition](#).

47.34.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::Composition< ValueType >
```

The [Composition](#) class can be used to compose linear operators `op1`, `op2`, ..., `opn` and obtain the operator `op1 * op2 * ...`

- `opn`.

All LinOps of the [Composition](#) must operate on Dense inputs. For an operator `op_k` that require an initial guess for their `apply`, [Composition](#) provides either

- the output of the previous `op_{k+1} -> apply` if `op_k` has square dimension
- zero if `op_k` is rectangular as an initial guess.

[Composition](#) ensures that all LinOps passed to its constructor use the same executor, and if not, copies the operators to the executor of the first operator.

Template Parameters

<i>ValueType</i>	precision of input and result vectors
------------------	---------------------------------------

47.34.2 Constructor & Destructor Documentation

47.34.2.1 Composition() [1/2]

```
template<typename ValueType = default_precision>
gko::Composition< ValueType >::Composition (
    const Composition< ValueType > & )
```

Copy-constructs a [Composition](#).

This inherits the executor of the input [Composition](#) and all of its operators with shared ownership.

47.34.2.2 Composition() [2/2]

```
template<typename ValueType = default_precision>
gko::Composition< ValueType >::Composition (
    Composition< ValueType > && )
```

Move-constructs a [Composition](#).

This inherits the executor of the input [Composition](#) and all of its operators. The moved-from object is empty (0x0 LinOp without operators) afterwards.

47.34.3 Member Function Documentation**47.34.3.1 conj_transpose()**

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Composition< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.34.3.2 get_operators()

```
template<typename ValueType = default_precision>
const std::vector<std::shared_ptr<const LinOp> >& gko::Composition< ValueType >::get_↵
operators ( ) const [inline], [noexcept]
```

Returns a list of operators of the composition.

Returns

a list of operators

```
56     {
57         return operators_;
58     }
```

47.34.3.3 operator=() [1/2]

```
template<typename ValueType = default_precision>
Composition& gko::Composition< ValueType >::operator= (
    Composition< ValueType > && )
```

Move-assigns a [Composition](#).

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor, otherwise they share ownership. The moved-from object is empty (0x0 LinOp without operators) afterwards.

47.34.3.4 operator=() [2/2]

```
template<typename ValueType = default_precision>
Composition& gko::Composition< ValueType >::operator= (
    const Composition< ValueType > & )
```

Copy-assigns a [Composition](#).

The executor is not modified, and the wrapped LinOps are only being cloned if they are on a different executor.

47.34.3.5 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::Composition< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/composition.hpp

47.35 gko::experimental::mpi::contiguous_type Class Reference

A move-only wrapper for a contiguous MPI_Datatype.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

- [contiguous_type](#) (int count, MPI_Datatype old_type)
Constructs a wrapper for a contiguous MPI_Datatype.
- [contiguous_type](#) ()
Constructs empty wrapper with MPI_DATATYPE_NULL.
- [contiguous_type](#) (const [contiguous_type](#) &)=delete
Disallow copying of wrapper type.
- [contiguous_type](#) & [operator=](#) (const [contiguous_type](#) &)=delete
Disallow copying of wrapper type.
- [contiguous_type](#) ([contiguous_type](#) &&other) noexcept
Move constructor, leaves other with MPI_DATATYPE_NULL.
- [contiguous_type](#) & [operator=](#) ([contiguous_type](#) &&other) noexcept
Move assignment, leaves other with MPI_DATATYPE_NULL.
- [~contiguous_type](#) ()
Destructs object by freeing wrapped MPI_Datatype.
- MPI_Datatype [get](#) () const
Access the underlying MPI_Datatype.

47.35.1 Detailed Description

A move-only wrapper for a contiguous MPI_Datatype.

The underlying MPI_Datatype is automatically created and committed when an object of this type is constructed, and freed when it is destructed.

47.35.2 Constructor & Destructor Documentation

47.35.2.1 contiguous_type() [1/2]

```
gko::experimental::mpi::contiguous_type::contiguous_type (
    int count,
    MPI_Datatype old_type ) [inline]
```

Constructs a wrapper for a contiguous MPI_Datatype.

Parameters

<i>count</i>	the number of old_type elements the new datatype contains.
<i>old_type</i>	the MPI_Datatype that is contained.

47.35.2.2 contiguous_type() [2/2]

```
gko::experimental::mpi::contiguous_type::contiguous_type (
```

```
contiguous_type && other ) [inline], [noexcept]
```

Move constructor, leaves other with MPI_DATATYPE_NULL.

Parameters

<i>other</i>	to be moved from object.
--------------	--------------------------

47.35.3 Member Function Documentation

47.35.3.1 get()

```
MPI_Datatype gko::experimental::mpi::contiguous_type::get ( ) const [inline]
```

Access the underlying MPI_Datatype.

Returns

the underlying MPI_Datatype.

47.35.3.2 operator=()

```
contiguous_type& gko::experimental::mpi::contiguous_type::operator= (
    contiguous_type && other ) [inline], [noexcept]
```

Move assignment, leaves other with MPI_DATATYPE_NULL.

Parameters

<i>other</i>	to be moved from object.
--------------	--------------------------

Returns

this object.

The documentation for this class was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.36 gko::log::Convergence< ValueType > Class Template Reference

[Convergence](#) is a Logger which logs data strictly from the `criterion_check_completed` event.

```
#include <ginkgo/core/log/convergence.hpp>
```

Public Member Functions

- bool `has_converged` () const noexcept
Returns true if the solver has converged.
- void `reset_convergence_status` ()
Resets the convergence status to false.
- const `size_type` & `get_num_iterations` () const noexcept
Returns the number of iterations.
- const `LinOp` * `get_residual` () const noexcept
Returns the residual.
- const `LinOp` * `get_residual_norm` () const noexcept
Returns the residual norm.
- const `LinOp` * `get_implicit_sq_resnorm` () const noexcept
Returns the implicit squared residual norm.

Static Public Member Functions

- static `std::unique_ptr< Convergence >` `create` (`std::shared_ptr< const Executor >`, const `mask_type` & `enabled_events=Logger::criterion_events_mask|Logger::iteration_complete_mask`)
Creates a convergence logger.
- static `std::unique_ptr< Convergence >` `create` (const `mask_type` & `enabled_events=Logger::criterion_↵events_mask|Logger::iteration_complete_mask`)
Creates a convergence logger.

47.36.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::log::Convergence< ValueType >
```

`Convergence` is a `Logger` which logs data strictly from the `criterion_check_completed` event.

The purpose of this logger is to give a simple access to standard data generated by the solver once it has stopped with minimal overhead.

This logger also computes the residual norm from the residual when the residual norm was not available. This can add some slight overhead.

47.36.2 Member Function Documentation

47.36.2.1 `create()` [1/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Convergence> gko::log::Convergence< ValueType >::create (
    const mask_type & enabled_events = Logger::criterion_events_mask | Logger::iteration↵
    _complete_mask ) [inline], [static]
```

Creates a convergence logger.

This dynamically allocates the memory, constructs the object and returns an `std::unique_ptr` to this object.

Parameters

<i>enabled_events</i>	the events enabled for this logger. By default all events.
-----------------------	--

Returns

an std::unique_ptr to the the constructed object

```

93     {
94         return std::unique_ptr<Convergence>(new Convergence(enabled_events));
95     }

```

47.36.2.2 create() [2/2]

```

template<typename ValueType = default_precision>
static std::unique_ptr<Convergence> gko::log::Convergence< ValueType >::create (
    std::shared_ptr< const Executor > ,
    const mask_type & enabled_events = Logger::criterion_events_mask | Logger::iteration↵
_complete_mask ) [inline], [static]

```

Creates a convergence logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.

Returns

an std::unique_ptr to the the constructed object

47.36.2.3 get_implicit_sq_resnorm()

```

template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_implicit_sq_resnorm ( ) const [inline],
[noexcept]

```

Returns the implicit squared residual norm.

Returns

the implicit squared residual norm

47.36.2.4 get_num_iterations()

```
template<typename ValueType = default_precision>
const size_type& gko::log::Convergence< ValueType >::get_num_iterations ( ) const [inline], [noexcept]
```

Returns the number of iterations.

Returns

the number of iterations

47.36.2.5 get_residual()

```
template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_residual ( ) const [inline], [noexcept]
```

Returns the residual.

Returns

the residual

47.36.2.6 get_residual_norm()

```
template<typename ValueType = default_precision>
const LinOp* gko::log::Convergence< ValueType >::get_residual_norm ( ) const [inline], [noexcept]
```

Returns the residual norm.

Returns

the residual norm

47.36.2.7 has_converged()

```
template<typename ValueType = default_precision>
bool gko::log::Convergence< ValueType >::has_converged ( ) const [inline], [noexcept]
```

Returns true if the solver has converged.

Returns

the bool flag for convergence status

The documentation for this class was generated from the following file:

- ginkgo/core/log/convergence.hpp

47.37 gko::ConvertibleTo< ResultType > Class Template Reference

[ConvertibleTo](#) interface is used to mark that the implementer can be converted to the object of ResultType.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

Public Member Functions

- virtual void [convert_to](#) (result_type *result) const =0
Converts the implementer to an object of type result_type.
- virtual void [move_to](#) (result_type *result)=0
Converts the implementer to an object of type result_type by moving data from this object.

47.37.1 Detailed Description

```
template<typename ResultType>
class gko::ConvertibleTo< ResultType >
```

[ConvertibleTo](#) interface is used to mark that the implementer can be converted to the object of ResultType.

This interface is used to enable conversions between polymorphic objects. To mark that an object of type U can be converted to an object of type V , U should implement `ConvertibleTo<V>`. Then, the implementation of `PolymorphicObject::copy_from` automatically generated by `EnablePolymorphicObject` mixin will use RTTI to figure out that U implements the interface and convert it using the `convert_to` / `move_to` methods of the interface.

As an example, the following function:

```
{c++}
void my_function(const U *u, V *v) {
    v->copy_from(u);
}
```

will convert object u to object v by checking that u can be dynamically casted to `ConvertibleTo<V>`, and calling `ConvertibleTo<V>::convert_to(V*)` to do the actual conversion.

In case u is passed as a `unique_ptr`, call to `convert_to` will be replaced by a call to `move_to` and trigger move semantics.

Template Parameters

<i>ResultType</i>	the type to which the implementer can be converted to, has to be a subclass of PolymorphicObject
-------------------	--

47.37.2 Member Function Documentation

47.37.2.1 convert_to()

```
template<typename ResultType>
virtual void gko::ConvertibleTo< ResultType >::convert\_to (
    result_type * result ) const [pure virtual]
```

Converts the implementer to an object of type `result_type`.

Parameters

<i>result</i>	the object used to store the result of the conversion
---------------	---

Implemented in `gko::EnablePolymorphicAssignment< ConcreteType, ResultType >`, `gko::EnablePolymorphicAssignment< Factoriza`
`gko::EnablePolymorphicAssignment< Lu< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Isai< IsaiType, Value`
`gko::EnablePolymorphicAssignment< SparsityCsr< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Vector< Valu`
`gko::EnablePolymorphicAssignment< Mc64< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Diagonal< ValueTy`
`gko::EnablePolymorphicAssignment< Pgm< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Dense< ValueType`
`gko::EnablePolymorphicAssignment< UpperTrs< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Amd< IndexTy`
`gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndexType, GlobalIndexType > >`, `gko::EnablePolymorphicAssignme`
`gko::EnablePolymorphicAssignment< Identity< ValueType > >`, `gko::EnablePolymorphicAssignment< ConcreteLinOp >`,
`gko::EnablePolymorphicAssignment< Rcm< IndexType > >`, `gko::EnablePolymorphicAssignment< Fft3 >`,
`gko::EnablePolymorphicAssignment< Partition< LocalIndexType, GlobalIndexType > >`, `gko::EnablePolymorphicAssignment< Com`
`gko::EnablePolymorphicAssignment< ConcreteBatchLinOp >`, `gko::EnablePolymorphicAssignment< RowGatherer< IndexType > >`
`gko::EnablePolymorphicAssignment< Cholesky< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Fft2 >`,
`gko::EnablePolymorphicAssignment< FixedCoarsening< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Fbcsr<`
`gko::EnablePolymorphicAssignment< LowerTrs< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Combination<`
`gko::EnablePolymorphicAssignment< Bicgstab< ValueType > >`, `gko::EnablePolymorphicAssignment< MultiVector< ValueType >`
`gko::EnablePolymorphicAssignment< Multigrid >`, `gko::EnablePolymorphicAssignment< Gmres< ValueType > >`,
`gko::EnablePolymorphicAssignment< GaussSeidel< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< CbGmres<`
`gko::EnablePolymorphicAssignment< Csr< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< ConcreteSolver >`,
`gko::EnablePolymorphicAssignment< Ir< ValueType > >`, `gko::EnablePolymorphicAssignment< Coo< ValueType, IndexType > >`,
`gko::EnablePolymorphicAssignment< Fcg< ValueType > >`, `gko::EnablePolymorphicAssignment< Rcm< ValueType, IndexType >`
`gko::EnablePolymorphicAssignment< Sor< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Fft >`,
`gko::EnablePolymorphicAssignment< BlockOperator >`, `gko::EnablePolymorphicAssignment< Direct< ValueType, IndexType > >`,
`gko::EnablePolymorphicAssignment< Idr< ValueType > >`, `gko::EnablePolymorphicAssignment< ConcreteFactory >`,
`gko::EnablePolymorphicAssignment< Cgs< ValueType > >`, `gko::EnablePolymorphicAssignment< Ell< ValueType, IndexType > >`
`gko::EnablePolymorphicAssignment< Ilu< LSolverType, USolverType, ReverseApply, IndexType > >`, `gko::EnablePolymorphicAssign`
`gko::EnablePolymorphicAssignment< Permutation< IndexType > >`, `gko::EnablePolymorphicAssignment< ScaledPermutation< Va`
`gko::EnablePolymorphicAssignment< Jacobi< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Gcr< ValueType >`
`gko::EnablePolymorphicAssignment< Schwarz< ValueType, LocalIndexType, GlobalIndexType > >`, `gko::EnablePolymorphicAssign`
`gko::EnablePolymorphicAssignment< ScaledReordered< ValueType, IndexType > >`, `gko::EnablePolymorphicAssignment< Sellp<`
`gko::EnablePolymorphicAssignment< Bicg< ValueType > >`, `gko::EnablePolymorphicAssignment< Perturbation< ValueType > >`,
and `gko::preconditioner::Jacobi< ValueType, IndexType >`.

47.37.2.2 move_to()

```
template<typename ResultType>
virtual void gko::ConvertibleTo< ResultType >::move_to (
    result_type * result ) [pure virtual]
```

Converts the implementer to an object of type `result_type` by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

<i>result</i>	the object used to emplace the result of the conversion
---------------	---

Note

[ConvertibleTo::move_to](#) can be implemented by simply calling [ConvertibleTo::convert_to](#). However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implemented in [gko::EnablePolymorphicAssignment< ConcreteType, ResultType >](#), [gko::EnablePolymorphicAssignment< Factoriza](#)
[gko::EnablePolymorphicAssignment< Lu< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Isai< IsaiType, Value](#)
[gko::EnablePolymorphicAssignment< SparsityCsr< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Vector< Valu](#)
[gko::EnablePolymorphicAssignment< Mc64< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Diagonal< ValueTy](#)
[gko::EnablePolymorphicAssignment< Pgm< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Dense< ValueType](#)
[gko::EnablePolymorphicAssignment< UpperTrs< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Amd< IndexTy](#)
[gko::EnablePolymorphicAssignment< Matrix< ValueType, LocalIndexType, GlobalIndexType > >](#), [gko::EnablePolymorphicAssignme](#)
[gko::EnablePolymorphicAssignment< Identity< ValueType > >](#), [gko::EnablePolymorphicAssignment< ConcreteLinOp >](#),
[gko::EnablePolymorphicAssignment< Rcm< IndexType > >](#), [gko::EnablePolymorphicAssignment< Fft3 >](#),
[gko::EnablePolymorphicAssignment< Partition< LocalIndexType, GlobalIndexType > >](#), [gko::EnablePolymorphicAssignment< Com](#)
[gko::EnablePolymorphicAssignment< ConcreteBatchLinOp >](#), [gko::EnablePolymorphicAssignment< RowGatherer< IndexType > >](#)
[gko::EnablePolymorphicAssignment< Cholesky< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Fft2 >](#),
[gko::EnablePolymorphicAssignment< FixedCoarsening< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Fbcsr<](#)
[gko::EnablePolymorphicAssignment< LowerTrs< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Combination<](#)
[gko::EnablePolymorphicAssignment< Bicgstab< ValueType > >](#), [gko::EnablePolymorphicAssignment< MultiVector< ValueType >](#)
[gko::EnablePolymorphicAssignment< Multigrid >](#), [gko::EnablePolymorphicAssignment< Gmres< ValueType > >](#),
[gko::EnablePolymorphicAssignment< GaussSeidel< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< CbGmres<](#)
[gko::EnablePolymorphicAssignment< Csr< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< ConcreteSolver >](#),
[gko::EnablePolymorphicAssignment< Ir< ValueType > >](#), [gko::EnablePolymorphicAssignment< Coo< ValueType, IndexType > >](#),
[gko::EnablePolymorphicAssignment< Fcg< ValueType > >](#), [gko::EnablePolymorphicAssignment< Rcm< ValueType, IndexType >](#)
[gko::EnablePolymorphicAssignment< Sor< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Fft >](#),
[gko::EnablePolymorphicAssignment< BlockOperator >](#), [gko::EnablePolymorphicAssignment< Direct< ValueType, IndexType > >](#),
[gko::EnablePolymorphicAssignment< Idr< ValueType > >](#), [gko::EnablePolymorphicAssignment< ConcreteFactory >](#),
[gko::EnablePolymorphicAssignment< Cgs< ValueType > >](#), [gko::EnablePolymorphicAssignment< Ell< ValueType, IndexType > >](#)
[gko::EnablePolymorphicAssignment< Ilu< LSolverType, USolverType, ReverseApply, IndexType > >](#), [gko::EnablePolymorphicAssign](#)
[gko::EnablePolymorphicAssignment< Permutation< IndexType > >](#), [gko::EnablePolymorphicAssignment< ScaledPermutation< Va](#)
[gko::EnablePolymorphicAssignment< Jacobi< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Gcr< ValueType >](#)
[gko::EnablePolymorphicAssignment< Schwarz< ValueType, LocalIndexType, GlobalIndexType > >](#), [gko::EnablePolymorphicAssign](#)
[gko::EnablePolymorphicAssignment< ScaledReordered< ValueType, IndexType > >](#), [gko::EnablePolymorphicAssignment< Sellp<](#)
[gko::EnablePolymorphicAssignment< Bicg< ValueType > >](#), [gko::EnablePolymorphicAssignment< Perturbation< ValueType > >](#),
and [gko::preconditioner::Jacobi< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/polymorphic_object.hpp](#)

47.38 gko::matrix::Coo< ValueType, IndexType > Class Template Reference

COO stores a matrix in the coordinate matrix format.

```
#include <ginkgo/core/matrix/coo.hpp>
```

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override

- Reads a matrix from a [device_matrix_data](#) structure.*
- void [read](#) ([device_mat_data](#) &&data) override
- Reads a matrix from a [device_matrix_data](#) structure.*
- void [write](#) ([mat_data](#) &data) const override
- Writes a matrix to a [matrix_data](#) structure.*
- std::unique_ptr< [Diagonal](#)< ValueType > > [extract_diagonal](#) () const override
- Extracts the diagonal entries of the matrix into a vector.*
- std::unique_ptr< absolute_type > [compute_absolute](#) () const override
- Gets the AbsoluteLinOp.*
- void [compute_absolute_inplace](#) () override
- Compute absolute inplace on each element.*
- value_type * [get_values](#) () noexcept
- Returns the values of the matrix.*
- const value_type * [get_const_values](#) () const noexcept
- Returns the values of the matrix.*
- index_type * [get_col_idxxs](#) () noexcept
- Returns the column indexes of the matrix.*
- const index_type * [get_const_col_idxxs](#) () const noexcept
- Returns the column indexes of the matrix.*
- index_type * [get_row_idxxs](#) () noexcept
- Returns the row indexes of the matrix.*
- const index_type * [get_const_row_idxxs](#) () const noexcept
- [size_type](#) [get_num_stored_elements](#) () const noexcept
- Returns the number of elements explicitly stored in the matrix.*
- LinOp * [apply2](#) ([ptr_param](#)< const LinOp > b, [ptr_param](#)< LinOp > x)
- Applies [Coo](#) matrix axpy to a vector (or a sequence of vectors).*
- const LinOp * [apply2](#) ([ptr_param](#)< const LinOp > b, [ptr_param](#)< LinOp > x) const
- LinOp * [apply2](#) ([ptr_param](#)< const LinOp > alpha, [ptr_param](#)< const LinOp > b, [ptr_param](#)< LinOp > x)
- Performs the operation $x = \text{alpha} * \text{Coo} * b + x$.*
- const LinOp * [apply2](#) ([ptr_param](#)< const LinOp > alpha, [ptr_param](#)< const LinOp > b, [ptr_param](#)< LinOp > x) const

Static Public Member Functions

- static std::unique_ptr< [Coo](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size=[dim](#)< 2 > {}, [size_type](#) num_nonzeros={})
- Creates an uninitialized COO matrix of the specified size.*
- static std::unique_ptr< [Coo](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, [array](#)< value_type > values, [array](#)< index_type > col_idxxs, [array](#)< index_type > row_idxxs)
- Creates a COO matrix from already allocated (and initialized) row index, column index and value arrays.*
- template<typename InputValueType, typename InputColumnIndexType, typename InputRowIndexType >
static std::unique_ptr< [Coo](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, std::initializer_list< InputValueType > values, std::initializer_list< InputColumnIndexType > col_idxxs, std::initializer_list< InputRowIndexType > row_idxxs)
- create(std::shared_ptr<const Executor> ,*
- static std::unique_ptr< const [Coo](#) > [create_const](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&col_idxxs, gko::detail::const_array_view< IndexType > &&row_idxxs)
- Creates a constant (immutable) [Coo](#) matrix from a set of constant arrays.*

47.38.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Coo< ValueType, IndexType >
```

COO stores a matrix in the coordinate matrix format.

The nonzero elements are stored in an array row-wise (but not necessarily sorted by column index within a row). Two extra arrays contain the row and column indexes of each nonzero element of the matrix.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.38.2 Member Function Documentation

47.38.2.1 `apply2()` [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
LinOp* gko::matrix::Coo< ValueType, IndexType >::apply2 (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x )
```

Performs the operation $x = \text{alpha} * \text{Coo} * b + x$.

Parameters

<i>alpha</i>	scaling of the result of <code>Coo * b</code>
<i>b</i>	vector(s) on which the operator is applied
<i>x</i>	output vector(s)

Returns

this

47.38.2.2 `apply2()` [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const LinOp* gko::matrix::Coo< ValueType, IndexType >::apply2 (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x ) const
```

47.38.2.3 apply2() [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
LinOp* gko::matrix::Coo< ValueType, IndexType >::apply2 (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x )
```

Applies [Coo](#) matrix axpy to a vector (or a sequence of vectors).

Performs the operation $x = \text{Coo} * b + x$

Parameters

<i>b</i>	the input vector(s) on which the operator is applied
<i>x</i>	the output vector(s) where the result is stored

Returns

this

47.38.2.4 apply2() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const LinOp* gko::matrix::Coo< ValueType, IndexType >::apply2 (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x ) const
```

47.38.2.5 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Coo< ValueType, IndexType >::compute_absolute ( )
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Coo< ValueType, IndexType > > >](#).

47.38.2.6 create() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Coo> gko::matrix::Coo< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< value_type > values,
    array< index_type > col_idxs,
    array< index_type > row_idxs ) [static]
```

Creates a COO matrix from already allocated (and initialized) row index, column index and value arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values
<i>col_idx</i> s	array of column indexes
<i>row_idx</i> s	array of row pointers

Note

If one of `row_idx`s, `col_idx`s or `values` is not an rvalue, not an array of `IndexType`, `IndexType` and `ValueType`, respectively, or is on the wrong executor, an internal copy of that array will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

47.38.2.7 create() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType , typename InputColumnType , typename InputRowIndexType
>
static std::unique_ptr<Coo> gko::matrix::Coo< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::initializer_list< InputValueType > values,
    std::initializer_list< InputColumnType > col_idx,
    std::initializer_list< InputRowIndexType > row_idx ) [inline], [static]

create(std::shared_ptr<const Executor>,

create(std::shared_ptr<const Executor>, const dim<2>&, array<value_type>, array<index_type>, array<index_type>)
286     {
287         return create(exec, size, array<value_type>(exec, std::move(values)),
288                       array<index_type>(exec, std::move(col_idx)),
289                       array<index_type>(exec, std::move(row_idx)));
290     }
```

References `gko::matrix::Coo< ValueType, IndexType >::create()`.

47.38.2.8 create() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Coo> gko::matrix::Coo< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = dim< 2 > {},
    size_type num_nonzeros = {} ) [static]
```

Creates an uninitialized COO matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_nonzeros</i>	number of nonzeros

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::matrix::Coo< ValueType, IndexType >::create()`.

47.38.2.9 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Coo> gko::matrix::Coo< ValueType, IndexType >::create_const (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    gko::detail::const_array_view< ValueType > && values,
    gko::detail::const_array_view< IndexType > && col_idx,
    gko::detail::const_array_view< IndexType > && row_idx ) [static]
```

Creates a constant (immutable) [Coo](#) matrix from a set of constant arrays.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>col_idx</i>	the column index array of the matrix
<i>row_ptr</i>	the row index array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

47.38.2.10 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Coo< ValueType, IndexType >::extract_↵
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.38.2.11 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Coo< ValueType, IndexType >::get_col_idxs ( ) [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References [gko::array< ValueType >::get_data\(\)](#).

47.38.2.12 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Coo< ValueType, IndexType >::get_const_col_idxs ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References [gko::array< ValueType >::get_const_data\(\)](#).

47.38.2.13 get_const_row_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Coo< ValueType, IndexType >::get_const_row_idxes ( ) const [inline],
[noexcept]
```

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.38.2.14 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Coo< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.38.2.15 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Coo< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.38.2.16 get_row_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Coo< ValueType, IndexType >::get_row_idxes ( ) [inline], [noexcept]
```

Returns the row indexes of the matrix.

Returns

the row indexes of the matrix.

References `gko::array< ValueType >::get_data()`.

47.38.2.17 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Coo< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References `gko::array< ValueType >::get_data()`.

47.38.2.18 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Coo< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a `device_matrix_data` structure.

Parameters

<i>data</i>	the <code>device_matrix_data</code> structure.
-------------	--

Reimplemented from `gko::ReadableFromMatrixData< ValueType, IndexType >`.

47.38.2.19 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Coo< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.38.2.20 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Coo< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.38.2.21 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Coo< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/coo.hpp

47.39 gko::CpuAllocator Class Reference

[Allocator](#) using new/delete.

```
#include <ginkgo/core/base/memory.hpp>
```

47.39.1 Detailed Description

[Allocator](#) using new/delete.

The documentation for this class was generated from the following file:

- ginkgo/core/base/memory.hpp

47.40 gko::CpuAllocatorBase Class Reference

Implement this interface to provide an allocator for [OmpExecutor](#) or [ReferenceExecutor](#).

```
#include <ginkgo/core/base/memory.hpp>
```

47.40.1 Detailed Description

Implement this interface to provide an allocator for [OmpExecutor](#) or [ReferenceExecutor](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/memory.hpp

47.41 gko::CpuTimer Class Reference

A timer using `std::chrono::steady_clock` for timing.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void [record](#) ([time_point](#) &time) override
Records a time point at the current time.
- void [wait](#) ([time_point](#) &time) override
Waits until all kernels in-process when recording the time point are finished.
- `std::chrono::nanoseconds` [difference_async](#) (const [time_point](#) &start, const [time_point](#) &stop) override
Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

47.41.1 Detailed Description

A timer using `std::chrono::steady_clock` for timing.

47.41.2 Member Function Documentation

47.41.2.1 `difference_async()`

```
std::chrono::nanoseconds gko::CpuTimer::difference_async (
    const time_point & start,
    const time_point & stop ) [override], [virtual]
```

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. `timer->wait(stop)` needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements [gko::Timer](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/timer.hpp`

47.42 `gko::cpx_real_type< T >` Struct Template Reference

Access the underlying real type of a complex number.

```
#include <ginkgo/core/base/math.hpp>
```

Public Types

- using `type` = T
The type.

47.42.1 Detailed Description

```
template<typename T>
struct gko::cpx_real_type< T >
```

Access the underlying real type of a complex number.

Template Parameters

<i>T</i>	the type being checked.
----------	-------------------------

47.42.2 Member Typedef Documentation

47.42.2.1 type

```
template<typename T >
using gko::cpx_real_type< T >::type = T
```

The type.

When the type is not complex, return the type itself.

The documentation for this struct was generated from the following file:

- ginkgo/core/base/math.hpp

47.43 gko::stop::Criterion Class Reference

The [Criterion](#) class is a base class for all stopping criteria.

```
#include <ginkgo/core/stop/criterion.hpp>
```

Classes

- class [Updater](#)
The [Updater](#) class serves for convenient argument passing to the [Criterion](#)'s check function.

Public Member Functions

- [Updater](#) update ()
Returns the updater object.
- bool [check](#) (uint8 stopping_id, bool set_finalized, array< [stopping_status](#) > *stop_status, bool *one_↵ changed, const [Updater](#) &updater)
This checks whether convergence was reached for a certain criterion.

47.43.1 Detailed Description

The [Criterion](#) class is a base class for all stopping criteria.

It contains a factory to instantiate criteria. It is up to each specific stopping criterion to decide what to do with the data that is passed to it.

Note that depending on the criterion, convergence may not have happened after stopping.

47.43.2 Member Function Documentation

47.43.2.1 check()

```
bool gko::stop::Criterion::check (
    uint8 stopping_id,
    bool set_finalized,
    array< stopping_status > * stop_status,
    bool * one_changed,
    const Updater & updater ) [inline]
```

This checks whether convergence was reached for a certain criterion.

The actual implantation of the criterion goes here.

Parameters

<i>stopping_id</i>	id of the stopping criterion
<i>set_finalized</i>	Controls if the current version should count as finalized or not
<i>stop_status</i>	status of the stopping criterion
<i>one_changed</i>	indicates if the status of a vector has changed
<i>updater</i>	the Updater object containing all the information

Returns

whether convergence was completely reached

```
140 {
141     this->template log<log::Logger::criterion_check_started>(
142         this, updater.num_iterations_, updater.residual_,
143         updater.residual_norm_, updater.solution_, stopping_id,
144         set_finalized);
145     auto all_converged = this->check_impl(
146         stopping_id, set_finalized, stop_status, one_changed, updater);
147     this->template log<log::Logger::criterion_check_completed>(
148         this, updater.num_iterations_, updater.residual_,
149         updater.residual_norm_, updater.implicit_sq_residual_norm_,
150         updater.solution_, stopping_id, set_finalized, stop_status,
151         *one_changed, all_converged);
152     return all_converged;
153 }
```

Referenced by `gko::stop::Criterion::Updater::check()`.

47.43.2.2 update()

```
Updater gko::stop::Criterion::update ( ) [inline]
```

Returns the updater object.

Returns

the updater object

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/criterion.hpp`

47.44 gko::log::criterion_data Struct Reference

Struct representing Criterion related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.44.1 Detailed Description

Struct representing Criterion related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.45 gko::stop::CriterionArgs Struct Reference

This struct is used to pass parameters to the EnableDefaultCriterionFactoryCriterionFactory::generate() method.

```
#include <ginkgo/core/stop/criterion.hpp>
```

47.45.1 Detailed Description

This struct is used to pass parameters to the EnableDefaultCriterionFactoryCriterionFactory::generate() method.

It is the ComponentsType of CriterionFactory.

Note

Dependly on the use case, some of these parameters can be `nullptr` as only some stopping criterion require them to be set. An example is the [ResidualNormReduction](#) which really requires the `initial←_residual` to be set.

The documentation for this struct was generated from the following file:

- ginkgo/core/stop/criterion.hpp

47.46 gko::batch::matrix::Csr< ValueType, IndexType > Class Template Reference

[Csr](#) is a general sparse matrix format that stores the column indices for each nonzero entry and a cumulative sum of the number of nonzeros in each row.

```
#include <ginkgo/core/matrix/batch_csr.hpp>
```

Public Member Functions

- `std::unique_ptr< unbatch_type > create_view_for_item (size_type item_id)`
Creates a mutable view (of `matrix::Csr` type) of one item of the `batch::matrix::Csr<value_type>` object.
- `std::unique_ptr< const unbatch_type > create_const_view_for_item (size_type item_id) const`
Creates a mutable view (of `matrix::Csr` type) of one item of the `batch::matrix::Csr<value_type>` object.
- `value_type * get_values () noexcept`
Returns a pointer to the array of values of the matrix.
- `const value_type * get_const_values () const noexcept`
Returns a pointer to the array of values of the matrix.
- `index_type * get_col_idxs () noexcept`
Returns a pointer to the array of column indices of the matrix.
- `const index_type * get_const_col_idxs () const noexcept`
Returns a pointer to the array of column indices of the matrix.
- `index_type * get_row_ptrs () noexcept`
Returns a pointer to the array of row pointers of the matrix.
- `const index_type * get_const_row_ptrs () const noexcept`
Returns a pointer to the array of row pointers of the matrix.
- `size_type get_num_stored_elements () const noexcept`
Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.
- `size_type get_num_elements_per_item () const noexcept`
Returns the number of stored elements in each batch item.
- `value_type * get_values_for_item (size_type batch_id) noexcept`
Returns a pointer to the array of values of the matrix for a specific batch item.
- `const value_type * get_const_values_for_item (size_type batch_id) const noexcept`
Returns a pointer to the array of values of the matrix for a specific batch item.
- `Csr * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector.
- `Csr * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector with a linear combination of the given input vector.
- `const Csr * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x) const`
- `const Csr * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x) const`
- `void scale (const array< value_type > &row_scale, const array< value_type > &col_scale)`
Performs in-place row and column scaling for this matrix.
- `void add_scaled_identity (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> beta)`
*Performs the operation $this = \alpha * I + \beta * this$.*

Static Public Member Functions

- `static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size=batch_dim< 2 >{}, size_type num_nonzeros_per_item={})`
Creates an uninitialized `Csr` matrix of the specified size.
- `static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, array< value_type > values, array< index_type > col_idxs, array< index_type > row_ptrs)`
Creates a `Csr` matrix from an already allocated (and initialized) array.

- `template<typename InputValueType, typename ColIndexType, typename RowPtrType >`
`static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size,`
`std::initializer_list< InputValueType > values, std::initializer_list< ColIndexType > col_idxes, std::initializer_↵`
`list< RowPtrType > row_ptrs)`
`create(std::shared_ptr<const Executor>,`
- `static std::unique_ptr< const Csr > create_const (std::shared_ptr< const Executor > exec, const`
`batch_dim< 2 > &sizes, gko::detail::const_array_view< value_type > &&values, gko::detail::const_↵`
`array_view< index_type > &&col_idxes, gko::detail::const_array_view< index_type > &&row_ptrs)`
Creates a constant (immutable) batch csr matrix from a constant array.

47.46.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::batch::matrix::Csr< ValueType, IndexType >
```

Csr is a general sparse matrix format that stores the column indices for each nonzero entry and a cumulative sum of the number of nonzeros in each row.

It is one of the most popular sparse matrix formats due to its versatility and ability to store a wide range of sparsity patterns in an efficient fashion.

Note

It is assumed that the sparsity pattern of all the items in the batch is the same and therefore only a single copy of the sparsity pattern is stored.

Currently only IndexType of int32 is supported.

Template Parameters

<i>ValueType</i>	value precision of matrix elements
<i>IndexType</i>	index precision of matrix elements

47.46.2 Member Function Documentation

47.46.2.1 add_scaled_identity()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::batch::matrix::Csr< ValueType, IndexType >::add_scaled_identity (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> beta )
```

Performs the operation $\text{this} = \alpha * I + \beta * \text{this}$.

Parameters

<i>alpha</i>	the scalar for identity
<i>beta</i>	the scalar to multiply this matrix

Note

Performs the operation in-place for this batch matrix

This operation fails in case this matrix does not have all its diagonal entries.

47.46.2.2 apply() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Csr* gko::batch::matrix::Csr< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector with a linear combination of the given input vector.

Represents the matrix vector multiplication, $x = \alpha * A$

- $b + \beta * x$, where x and b are both multi-vectors.

Parameters

<i>alpha</i>	the scalar to scale the matrix-vector product with
<i>b</i>	the multi-vector to be applied to
<i>beta</i>	the scalar to scale the x vector with
<i>x</i>	the output multi-vector

47.46.2.3 apply() [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const Csr* gko::batch::matrix::Csr< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x ) const
```

47.46.2.4 apply() [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Csr* gko::batch::matrix::Csr< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector.

Represents the matrix vector multiplication, $x = A * b$, where x and b are both multi-vectors.

Parameters

<i>b</i>	the multi-vector to be applied to
<i>x</i>	the output multi-vector

47.46.2.5 apply() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const Csr* gko::batch::matrix::Csr< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x ) const
```

47.46.2.6 create() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Csr> gko::batch::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    array< value_type > values,
    array< index_type > col_idxes,
    array< index_type > row_ptrs ) [static]
```

Creates a [Csr](#) matrix from an already allocated (and initialized) array.

The column indices array needs to be the same for all batch items.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values
<i>col_idxes</i>	the <i>col_idxes</i> array of a single batch item of the matrix.
<i>row_ptrs</i>	the <i>row_ptrs</i> array of a single batch item of the matrix.

Note

If *values* is not an rvalue, not an array of *ValueType*, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.46.2.7 create() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType , typename ColIndexType , typename RowPtrType >
static std::unique_ptr<Csr> gko::batch::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    std::initializer_list< InputValueType > values,
    std::initializer_list< ColIndexType > col_idxxs,
    std::initializer_list< RowPtrType > row_ptrs ) [inline], [static]
```

create(std::shared_ptr<const Executor>,

```
create(std::shared_ptr<const Executor>, const batch_dim<2>&, array<value_type>, array<index_type>,
array<index_type>)
```

```
278     {
279         return create(exec, size, array<value_type>{exec, std::move(values)},
280                     array<index_type>{exec, std::move(col_idxxs)},
281                     array<index_type>{exec, std::move(row_ptrs)});
282     }
```

References gko::batch::matrix::Csr< ValueType, IndexType >::create().

47.46.2.8 create() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Csr> gko::batch::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size = batch_dim< 2 >{},
    size_type num_nonzeros_per_item = {} ) [static]
```

Creates an uninitialized Csr matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_nonzeros_per_item</i>	number of nonzeros in each item of the batch matrix

Referenced by gko::batch::matrix::Csr< ValueType, IndexType >::create().

47.46.2.9 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Csr> gko::batch::matrix::Csr< ValueType, IndexType >::create_←
const (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & sizes,
    gko::detail::const_array_view< value_type > && values,
```

```
gko::detail::const_array_view< index_type > && col_idxes,
gko::detail::const_array_view< index_type > && row_ptrs ) [static]
```

Creates a constant (immutable) batch csr matrix from a constant array.

Only a single sparsity pattern (column indices and row pointers) is stored and hence the user needs to ensure that each batch item has the same sparsity pattern.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>col_idxes</i>	the col_idxes array of a single batch item of the matrix.
<i>row_ptrs</i>	the row_ptrs array of a single batch item of the matrix.

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

A smart pointer to the newly created matrix.

47.46.2.10 create_const_view_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<const unbatch_type> gko::batch::matrix::Csr< ValueType, IndexType >::create↵
const_view_for_item (
    size_type item_id ) const
```

Creates a mutable view (of [matrix::Csr](#) type) of one item of the [batch::matrix::Csr<value_type>](#) object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<i>item↵ _id</i>	The index of the batch item
----------------------	-----------------------------

Returns

a [batch::matrix::Csr](#) object with the data from the batch item at the given index.

47.46.2.11 create_view_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<unbatch_type> gko::batch::matrix::Csr< ValueType, IndexType >::create_view↵
for_item (
    size_type item_id )
```

Creates a mutable view (of `matrix::Csr` type) of one item of the `batch::matrix::Csr<value_type>` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<code>item↔ _id</code>	The index of the batch item
----------------------------	-----------------------------

Returns

a `batch::matrix::Csr` object with the data from the batch item at the given index.

47.46.2.12 `get_col_idxs()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_col_idxs ( ) [inline], [noexcept]
```

Returns a pointer to the array of column indices of the matrix.

Returns

the pointer to the array of column indices

References `gko::array< ValueType >::get_data()`.

47.46.2.13 `get_const_col_idxs()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_const_col_idxs ( )
const [inline], [noexcept]
```

Returns a pointer to the array of column indices of the matrix.

Returns

the pointer to the array of column indices

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.46.2.14 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs ( )
const [inline], [noexcept]
```

Returns a pointer to the array of row pointers of the matrix.

Returns

the pointer to the array of row pointers

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.46.2.15 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values ( ) const
[inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.46.2.16 get_const_values_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values_for_item (
    size_type batch_id ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch</i> ↔ <i>_id</i>	the id of the batch item.
------------------------------	---------------------------

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`, `gko::batch::matrix::Csr< ValueType, IndexType >::get_↔
num_elements_per_item()`, and `gko::array< ValueType >::get_size()`.

47.46.2.17 get_num_elements_per_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::matrix::Csr< ValueType, IndexType >::get_num_elements_per_item ( ) const
[inline], [noexcept]
```

Returns the number of stored elements in each batch item.

Returns

the number of stored elements per batch item.

References `gko::batch::matrix::Csr< ValueType, IndexType >::get_num_stored_elements()`.

Referenced by `gko::batch::matrix::Csr< ValueType, IndexType >::get_const_values_for_item()`, and `gko::batch↔
::matrix::Csr< ValueType, IndexType >::get_values_for_item()`.

47.46.2.18 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::matrix::Csr< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.

Returns

the number of elements explicitly stored in the vector, cumulative across all the batch items

References `gko::array< ValueType >::get_size()`.

Referenced by `gko::batch::matrix::Csr< ValueType, IndexType >::get_num_elements_per_item()`.

47.46.2.19 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_row_ptrs ( ) [inline], [noexcept]
```

Returns a pointer to the array of row pointers of the matrix.

Returns

the pointer to the array of row pointers

References gko::array< ValueType >::get_data().

47.46.2.20 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data().

47.46.2.21 get_values_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::batch::matrix::Csr< ValueType, IndexType >::get_values_for_item (
    size_type batch_id ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch_id</i>	the id of the batch item.
-----------------	---------------------------

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data(), gko::batch::matrix::Csr< ValueType, IndexType >::get_num_elements_per_item(), and gko::array< ValueType >::get_size().

47.46.2.22 scale()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::batch::matrix::Csr< ValueType, IndexType >::scale (
    const array< value_type > & row_scale,
    const array< value_type > & col_scale )
```

Performs in-place row and column scaling for this matrix.

Parameters

<i>row_scale</i>	the row scalars
<i>col_scale</i>	the column scalars

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/batch_csr.hpp

47.47 gko::matrix::Csr< ValueType, IndexType > Class Template Reference

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

```
#include <ginkgo/core/matrix/csr.hpp>
```

Classes

- class [classical](#)
classical is a [strategy_type](#) which uses the same number of threads on each row.
- class [cusparse](#)
cusparse is a [strategy_type](#) which uses the sparselib csr.
- class [load_balance](#)
load_balance is a [strategy_type](#) which uses the load balance algorithm.
- class [merge_path](#)
merge_path is a [strategy_type](#) which uses the [merge_path](#) algorithm.
- struct [permuting_reuse_info](#)
A struct describing a transformation of the matrix that reorders the values of the matrix into the transformed matrix.
- class [sparselib](#)
sparselib is a [strategy_type](#) which uses the sparselib csr.
- class [strategy_type](#)
strategy_type is to decide how to set the csr algorithm.

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- std::pair< std::unique_ptr< [Csr](#) >, [permuting_reuse_info](#) > [transpose_reuse](#) () const
Computes the necessary data to update a transposed matrix from its original matrix.
- std::unique_ptr< [Csr](#) > [permute](#) ([ptr_param](#)< const [Permutation](#)< index_type >> permutation, [permute_mode](#) mode=[permute_mode::symmetric](#)) const
Creates a permuted copy A' of this matrix A with the given permutation P .
- std::unique_ptr< [Csr](#) > [permute](#) ([ptr_param](#)< const [Permutation](#)< index_type >> row_permutation, [ptr_param](#)< const [Permutation](#)< index_type >> column_permutation, bool invert=false) const
Creates a non-symmetrically permuted copy A' of this matrix A with the given row and column permutations P and Q .
- std::pair< std::unique_ptr< [Csr](#) >, [permuting_reuse_info](#) > [permute_reuse](#) ([ptr_param](#)< const [Permutation](#)< index_type >> permutation, [permute_mode](#) mode=[permute_mode::symmetric](#)) const
Computes the operations necessary to propagate changed values from a matrix A to a permuted matrix.
- std::pair< std::unique_ptr< [Csr](#) >, [permuting_reuse_info](#) > [permute_reuse](#) ([ptr_param](#)< const [Permutation](#)< index_type >> row_permutation, [ptr_param](#)< const [Permutation](#)< index_type >> column_permutation, bool invert=false) const
Computes the operations necessary to propagate changed values from a matrix A to a permuted matrix.
- std::unique_ptr< [Csr](#) > [scale_permute](#) ([ptr_param](#)< const [ScaledPermutation](#)< value_type, index_type >> permutation, [permute_mode](#)=[permute_mode::symmetric](#)) const
Creates a scaled and permuted copy of this matrix.
- std::unique_ptr< [Csr](#) > [scale_permute](#) ([ptr_param](#)< const [ScaledPermutation](#)< value_type, index_type >> row_permutation, [ptr_param](#)< const [ScaledPermutation](#)< value_type, index_type >> column_permutation, bool invert=false) const
Creates a scaled and permuted copy of this matrix.
- std::unique_ptr< LinOp > [permute](#) (const [array](#)< IndexType > *permutation_indices) const override
Returns a LinOp representing the symmetric row and column permutation of the [Permutable](#) object.
- std::unique_ptr< LinOp > [inverse_permute](#) (const [array](#)< IndexType > *inverse_permutation_indices) const override
Returns a LinOp representing the symmetric inverse row and column permutation of the [Permutable](#) object.
- std::unique_ptr< LinOp > [row_permute](#) (const [array](#)< IndexType > *permutation_indices) const override
Returns a LinOp representing the row permutation of the [Permutable](#) object.
- std::unique_ptr< LinOp > [column_permute](#) (const [array](#)< IndexType > *permutation_indices) const override
Returns a LinOp representing the column permutation of the [Permutable](#) object.
- std::unique_ptr< LinOp > [inverse_row_permute](#) (const [array](#)< IndexType > *inverse_permutation_indices) const override
Returns a LinOp representing the row permutation of the inverse permuted object.
- std::unique_ptr< LinOp > [inverse_column_permute](#) (const [array](#)< IndexType > *inverse_permutation_indices) const override
Returns a LinOp representing the row permutation of the inverse permuted object.
- std::unique_ptr< [Diagonal](#)< ValueType >> [extract_diagonal](#) () const override

- Extracts the diagonal entries of the matrix into a vector.*

 - `std::unique_ptr< absolute_type > compute_absolute ()` const override

Gets the AbsoluteLinOp.
- `void compute_absolute_inplace ()` override

Compute absolute inplace on each element.
- `void sort_by_column_index ()`

Sorts all (value, col_idx) pairs in each row by column index.
- `value_type * get_values ()` noexcept

Returns the values of the matrix.
- `const value_type * get_const_values ()` const noexcept

Returns the values of the matrix.
- `std::unique_ptr< Dense< ValueType > > create_value_view ()`

Creates a Dense view of the value array of this matrix as a column vector of dimensions nnz x 1.
- `std::unique_ptr< const Dense< ValueType > > create_const_value_view ()` const

Creates a const Dense view of the value array of this matrix as a column vector of dimensions nnz x 1.
- `index_type * get_col_idxs ()` noexcept

Returns the column indexes of the matrix.
- `const index_type * get_const_col_idxs ()` const noexcept

Returns the column indexes of the matrix.
- `index_type * get_row_ptrs ()` noexcept

Returns the row pointers of the matrix.
- `const index_type * get_const_row_ptrs ()` const noexcept

Returns the row pointers of the matrix.
- `index_type * get_srow ()` noexcept

Returns the starting rows.
- `const index_type * get_const_srow ()` const noexcept

Returns the starting rows.
- `size_type get_num_srow_elements ()` const noexcept

Returns the number of the srow stored elements (involved warps)
- `size_type get_num_stored_elements ()` const noexcept

Returns the number of elements explicitly stored in the matrix.
- `std::shared_ptr< strategy_type > get_strategy ()` const noexcept

Returns the strategy.
- `void set_strategy (std::shared_ptr< strategy_type > strategy)`

Set the strategy.
- `void scale (ptr_param< const LinOp > alpha)`

Scales the matrix with a scalar.
- `void inv_scale (ptr_param< const LinOp > alpha)`

Scales the matrix with the inverse of a scalar.
- `std::unique_ptr< Csr< ValueType, IndexType > > create_submatrix (const index_set< IndexType > &row_index_set, const index_set< IndexType > &column_index_set)` const

Creates a submatrix from this Csr matrix given row and column index_set objects.
- `std::unique_ptr< Csr< ValueType, IndexType > > create_submatrix (const span &row_span, const span &column_span)` const

Creates a submatrix from this Csr matrix given row and column spans.
- `Csr & operator= (const Csr &)`

Copy-assigns a Csr matrix.
- `Csr & operator= (Csr &&)`

Move-assigns a Csr matrix.
- `Csr (const Csr &)`

Copy-constructs a Csr matrix.
- `Csr (Csr &&)`

Move-constructs a Csr matrix.

Static Public Member Functions

- static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, std::shared_ptr< strategy_type > strategy)
Creates an uninitialized CSR matrix of the specified size.
- static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size={}, size_type num_nonzeros={}, std::shared_ptr< strategy_type > strategy=nullptr)
Creates an uninitialized CSR matrix of the specified size.
- static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, array< value_type > values, array< index_type > col_idx, array< index_type > row_ptrs, std::shared_ptr< strategy_type > strategy=nullptr)
Creates a CSR matrix from already allocated (and initialized) row pointer, column index and value arrays.
- template<typename InputValueType, typename InputColumnIndexType, typename InputRowPtrType >
static std::unique_ptr< Csr > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, std::initializer_list< InputValueType > values, std::initializer_list< InputColumnIndexType > col_idx, std::initializer_list< InputRowPtrType > row_ptrs)
create(std::shared_ptr<const Executor> ,
- static std::unique_ptr< const Csr > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&col_idx, gko::detail::const_array_view< IndexType > &&row_ptrs, std::shared_ptr< strategy_type > strategy=nullptr)
Creates a constant (immutable) Csr matrix from a set of constant arrays.

47.47.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >
```

CSR is a matrix format which stores only the nonzero coefficients by compressing each row of the matrix (compressed sparse row format).

The nonzero elements are stored in a 1D array row-wise, and accompanied with a row pointer array which stores the starting index of each row. An additional column index array is used to identify the column of each nonzero element.

The Csr LinOp supports different operations:

```
matrix::Csr *A, *B, *C;           // matrices
matrix::Dense *b, *x;             // vectors tall-and-skinny matrices
matrix::Dense *alpha, *beta;      // scalars of dimension 1x1
matrix::Identity *I;              // identity matrix
// Applying to Dense matrices computes an SpMV/SpMM product
A->apply(b, x)                     // x = A*b
A->apply(alpha, b, beta, x)        // x = alpha*A*b + beta*x
// Applying to Csr matrices computes a SpGEMM product of two sparse matrices
A->apply(B, C)                     // C = A*B
A->apply(alpha, B, beta, C)        // C = alpha*A*B + beta*C
// Applying to an Identity matrix computes a SpGEAM sparse matrix addition
A->apply(alpha, I, beta, B)        // B = alpha*A + beta*B
```

Both the SpGEMM and SpGEAM operation require the input matrices to be sorted by column index, otherwise the algorithms will produce incorrect results.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.47.2 Constructor & Destructor Documentation

47.47.2.1 Csr() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::Csr (
    const Csr< ValueType, IndexType > & )
```

Copy-constructs a [Csr](#) matrix.

Inherits executor, strategy and data.

47.47.2.2 Csr() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::Csr (
    Csr< ValueType, IndexType > && )
```

Move-constructs a [Csr](#) matrix.

Inherits executor and strategy, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor and strategy, no nonzeros and valid row pointers).

47.47.3 Member Function Documentation

47.47.3.1 column_permute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::column_permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the column permutation of the [Permutable](#) object.

In the resulting LinOp, the column `i` contains the input column `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^T .

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order <code>perm</code> .
----------------------------------	---

Returns

a pointer to the new column permuted object

Implements [gko::Permutable< IndexType >](#).

47.47.3.2 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Csr< ValueType, IndexType >::compute_absolute ( )
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Csr< ValueType, IndexType > > >](#).

47.47.3.3 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::conj_transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.47.3.4 create() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< value_type > values,
    array< index_type > col_idxs,
    array< index_type > row_ptrs,
    std::shared_ptr< strategy_type > strategy = nullptr ) [static]
```

Creates a CSR matrix from already allocated (and initialized) row pointer, column index and value arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values
<i>col_idx</i> s	array of column indexes
<i>row_ptr</i> s	array of row pointers
<i>strategy</i>	the strategy the matrix uses for SpMV operations

Note

If one of `row_ptr`s, `col_idx`s or `values` is not an rvalue, not an array of `IndexType`, `IndexType` and `ValueType`, respectively, or is on the wrong executor, an internal copy of that array will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.47.3.5 `create()` [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType , typename InputColumnIndexType , typename InputRowPtrType >
static std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::initializer_list< InputValueType > values,
    std::initializer_list< InputColumnIndexType > col_idx,
    std::initializer_list< InputRowPtrType > row_ptr ) [inline], [static]
```

`create(std::shared_ptr<const Executor>,`

`create(std::shared_ptr<const Executor>, const dim<2>&, array<value_type>, array<index_type>, array<index_type>)`

```
1171 {
1172     return create(exec, size, array<value_type>(exec, std::move(values)),
1173                 array<index_type>(exec, std::move(col_idx)),
1174                 array<index_type>(exec, std::move(row_ptr)));
1175 }
```

References `gko::matrix::Csr< ValueType, IndexType >::create()`.

47.47.3.6 `create()` [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = {},
    size_type num_nonzeros = {},
    std::shared_ptr< strategy_type > strategy = nullptr ) [static]
```

Creates an uninitialized CSR matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_nonzeros</i>	number of nonzeros
<i>strategy</i>	the strategy of CSR, or the default strategy if set to nullptr

Returns

A smart pointer to the newly created matrix.

47.47.3.7 create() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< strategy_type > strategy ) [static]
```

Creates an uninitialized CSR matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>strategy</i>	the strategy of CSR

Returns

A smart pointer to the newly created matrix.

Referenced by gko::matrix::Csr< ValueType, IndexType >::create().

47.47.3.8 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Csr> gko::matrix::Csr< ValueType, IndexType >::create_const (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    gko::detail::const_array_view< ValueType > && values,
    gko::detail::const_array_view< IndexType > && col_idx,
    gko::detail::const_array_view< IndexType > && row_ptr,
    std::shared_ptr< strategy_type > strategy = nullptr ) [static]
```

Creates a constant (immutable) [Csr](#) matrix from a set of constant arrays.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>col_idx</i>	the column index array of the matrix
<i>row_ptr</i>	the row pointer array of the matrix
<i>strategy</i>	the strategy the matrix uses for SpMV operations

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

A smart pointer to the newly created matrix.

47.47.3.9 create_submatrix() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr<ValueType, IndexType> > gko::matrix::Csr< ValueType, IndexType >::create←
_submatrix (
    const index_set< IndexType > & row_index_set,
    const index_set< IndexType > & column_index_set ) const
```

Creates a submatrix from this [Csr](#) matrix given row and column [index_set](#) objects.

Parameters

<i>row_index_set</i>	the row index set containing the set of rows to be in the submatrix.
<i>column_index_set</i>	the col index set containing the set of columns to be in the submatrix.

Returns

A new CSR matrix with the elements that belong to the row and columns of this matrix as specified by the index sets.

Note

This is not a view but creates a new, separate CSR matrix.

47.47.3.10 create_submatrix() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr<ValueType, IndexType> > gko::matrix::Csr< ValueType, IndexType >::create←
_submatrix (
    const span & row_span,
    const span & column_span ) const
```

Creates a submatrix from this [Csr](#) matrix given row and column spans.

Parameters

<i>row_span</i>	the row span containing the contiguous set of rows to be in the submatrix.
<i>column_span</i>	the column span containing the contiguous set of columns to be in the submatrix.

Returns

A new CSR matrix with the elements that belong to the row and columns of this matrix as specified by the index sets.

Note

This is not a view but creates a new, separate CSR matrix.

47.47.3.11 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Csr< ValueType, IndexType >::extract_↵
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.47.3.12 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_col_idxs ( ) [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References [gko::array< ValueType >::get_data\(\)](#).

47.47.3.13 get_const_col_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_col_idxes ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.47.3.14 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_row_ptrs ( ) const [inline],
[noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.47.3.15 get_const_srow()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Csr< ValueType, IndexType >::get_const_srow ( ) const [inline],
[noexcept]
```

Returns the starting rows.

Returns

the starting rows.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.47.3.16 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Csr< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.47.3.17 get_num_srow_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Csr< ValueType, IndexType >::get_num_srow_elements ( ) const [inline],
[noexcept]
```

Returns the number of the srow stored elements (involved warps)

Returns

the number of the srow stored elements (involved warps)

References gko::array< ValueType >::get_size().

47.47.3.18 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Csr< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_size().

47.47.3.19 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_row_ptrs ( ) [inline], [noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

References `gko::array< ValueType >::get_data()`.

47.47.3.20 get_srow()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Csr< ValueType, IndexType >::get_srow ( ) [inline], [noexcept]
```

Returns the starting rows.

Returns

the starting rows.

References `gko::array< ValueType >::get_data()`.

47.47.3.21 get_strategy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::get_strategy ( )
const [inline], [noexcept]
```

Returns the strategy.

Returns

the strategy

47.47.3.22 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Csr< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References `gko::array< ValueType >::get_data()`.

47.47.3.23 inv_scale()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::inv_scale (
    ptr_param< const LinOp > alpha ) [inline]
```

Scales the matrix with the inverse of a scalar.

Parameters

<i>alpha</i>	The entire matrix is scaled by 1 / alpha. alpha has to be a 1x1 Dense matrix.
--------------	---

References `gko::PolymorphicObject::get_executor()`, and `gko::make_temporary_clone()`.

47.47.3.24 inverse_column_permute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::inverse_column_permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column `perm[i]` contains the input column `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation AP^{-T} .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new inverse permuted object

Implements [gko::Permutable< IndexType >](#).

47.47.3.25 inverse_permute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::inverse_permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the symmetric inverse row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location `(perm[i], perm[j])` contains the input value `(i, j)`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation $P^{-1}AP^{-T}$.

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< IndexType >](#).

47.47.3.26 inverse_row_permute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::inverse_row_permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row `perm[i]` contains the input row `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}A$.

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new inverse permuted object

Implements [gko::Permutable< IndexType >](#).

47.47.3.27 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Csr& gko::matrix::Csr< ValueType, IndexType >::operator= (
    const Csr< ValueType, IndexType > & )
```

Copy-assigns a [Csr](#) matrix.

Preserves executor, copies everything else.

47.47.3.28 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Csr& gko::matrix::Csr< ValueType, IndexType >::operator= (
    Csr< ValueType, IndexType > && )
```

Move-assigns a [Csr](#) matrix.

Preserves executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor and strategy, no nonzeros and valid row pointers).

47.47.3.29 permute() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the symmetric row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location (i, j) contains the input value $(\text{perm}[i], \text{perm}[j])$.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PAP^T .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< IndexType >](#).

47.47.3.30 permute() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::permute (
    ptr_param< const Permutation< index_type >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

Creates a permuted copy A' of this matrix A with the given permutation P .

By default, this computes a symmetric permutation ([permute_mode::symmetric](#)). For the effect of the different permutation modes, see [permute_mode](#)

Parameters

<i>permutation</i>	The input permutation.
<i>mode</i>	The permutation mode. If permute_mode::inverse is set, we use the inverse permutation P^{-1} instead of P . If permute_mode::rows is set, the rows will be permuted. If permute_mode::columns is set, the columns will be permuted.

Returns

The permuted matrix.

47.47.3.31 permute() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::permute (
    ptr_param< const Permutation< index_type >> row_permutation,
    ptr_param< const Permutation< index_type >> column_permutation,
    bool invert = false ) const
```

Creates a non-symmetrically permuted copy A' of this matrix A with the given row and column permutations P and Q .

The operation will compute $A'(i, j) = A(p[i], q[j])$, or $A' = PAQ^T$ if `invert` is `false`, and $A'(p[i], q[j]) = A(i, j)$, or $A' = P^{-1}AQ^{-T}$ if `invert` is `true`.

Parameters

<i>row_permutation</i>	The permutation P to apply to the rows
<i>column_permutation</i>	The permutation Q to apply to the columns
<i>invert</i>	If set to <code>false</code> , uses the input permutations, otherwise uses their inverses P^{-1}, Q^{-1}

Returns

The permuted matrix.

47.47.3.32 permute_reuse() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::pair<std::unique_ptr<Csr>, permuting_reuse_info> gko::matrix::Csr< ValueType, IndexType >::permute_reuse (
    ptr_param< const Permutation< index_type >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

Computes the operations necessary to propagate changed values from a matrix A to a permuted matrix.

The semantics of this function match those of `permute(ptr_param<const Permutation<index_type>>, permute_mode)`. Updating values works as follows:

```
auto [permuted, reuse] = matrix->permute_reuse(permutation, mode);
change_values(matrix);
reuse->update_values(matrix, permuted);
```

Parameters

<i>permutation</i>	The input permutation.
<i>mode</i>	The permutation mode. If <code>permute_mode::inverse</code> is set, we use the inverse permutation P^{-1} instead of P . If <code>permute_mode::rows</code> is set, the rows will be permuted. If <code>permute_mode::columns</code> is set, the columns will be permuted.

Returns

an `std::pair` consisting of the permuted matrix and the reuse info that can be used to update values in the permuted matrix.

47.47.3.33 permute_reuse() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::pair<std::unique_ptr<Csr>, permuting_reuse_info> gko::matrix::Csr< ValueType, IndexType
>::permute_reuse (
    ptr_param< const Permutation< index_type >> row_permutation,
    ptr_param< const Permutation< index_type >> column_permutation,
    bool invert = false ) const
```

Computes the operations necessary to propagate changed values from a matrix *A* to a permuted matrix.

The semantics of this function match those of `permute(ptr_param<const Permutation<index_type>>, ptr_param<const Permutation<index_type>>, bool)`. Updating values works as follows:

```
auto [permuted, reuse] = matrix->permute_reuse(row_perm, col_perm, inv);
change_values(matrix);
reuse->update_values(matrix, permuted);
```

Parameters

<i>row_permutation</i>	The permutation P to apply to the rows
<i>column_permutation</i>	The permutation Q to apply to the columns
<i>invert</i>	If set to <code>false</code> , uses the input permutations, otherwise uses their inverses P^{-1}, Q^{-1}

Returns

an `std::pair` consisting of the permuted matrix and the reuse info that can be used to update values in the permuted matrix.

47.47.3.34 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a `device_matrix_data` structure.

Parameters

<i>data</i>	the <code>device_matrix_data</code> structure.
-------------	--

Reimplemented from `gko::ReadableFromMatrixData< ValueType, IndexType >`.

47.47.3.35 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
```

```
void gko::matrix::Csr< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.47.3.36 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.47.3.37 row_permute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::row_permute (
    const array< IndexType > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the row permutation of the [Permutable](#) object.

In the resulting LinOp, the row *i* contains the input row `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation PA .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Implements [gko::Permutable< IndexType >](#).

47.47.3.38 scale()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::scale (
    ptr_param< const LinOp > alpha ) [inline]
```

Scales the matrix with a scalar.

Parameters

<i>alpha</i>	The entire matrix is scaled by alpha. alpha has to be a 1x1 Dense matrix.
--------------	---

References [gko::PolymorphicObject::get_executor\(\)](#), and [gko::make_temporary_clone\(\)](#).

47.47.3.39 scale_permute() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, index_type >> permutation,
    permute_mode = permute_mode::symmetric ) const
```

Creates a scaled and permuted copy of this matrix.

For an explanation of the permutation modes, see [permute\(ptr_param<const Permutation<index_type>>, permute_mode\)](#)

Parameters

<i>permutation</i>	The scaled permutation.
<i>mode</i>	The permutation mode.

Returns

The permuted matrix.

47.47.3.40 scale_permute() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Csr> gko::matrix::Csr< ValueType, IndexType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, index_type >> row_permutation,
```

```
ptr_param< const ScaledPermutation< value_type, index_type >> column_permutation,
bool invert = false ) const
```

Creates a scaled and permuted copy of this matrix.

For an explanation of the parameters, see [permute\(ptr_param<const Permutation<index_type>>, ptr_param<const Permutation<index_type>>, permute_mode\)](#)

Parameters

<i>row_permutation</i>	The scaled row permutation.
<i>column_permutation</i>	The scaled column permutation.
<i>invert</i>	If set to <code>false</code> , uses the input permutations, otherwise uses their inverses P^{-1}, Q^{-1}

Returns

The permuted matrix.

47.47.3.41 set_strategy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::set_strategy (
    std::shared_ptr< strategy_type > strategy ) [inline]
```

Set the strategy.

Parameters

<i>strategy</i>	the csr strategy
-----------------	------------------

47.47.3.42 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Csr< ValueType, IndexType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.47.3.43 transpose_reuse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::pair<std::unique_ptr<Csr>, permuting_reuse_info> gko::matrix::Csr< ValueType, IndexType
>::transpose_reuse ( ) const
```

Computes the necessary data to update a transposed matrix from its original matrix.

```
auto [transposed, reuse] = matrix->transpose_reuse();
change_values(matrix);
reuse->update_values(matrix, transposed);
```

Returns

the reuse info struct that can be used to update values in the transposed matrix.

47.47.3.44 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following files:

- ginkgo/core/distributed/matrix.hpp
- ginkgo/core/matrix/csr.hpp

47.48 gko::CublasError Class Reference

[CublasError](#) is thrown when a cuBLAS routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [CublasError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a cuBLAS error.

47.48.1 Detailed Description

[CublasError](#) is thrown when a cuBLAS routine throws a non-zero error code.

47.48.2 Constructor & Destructor Documentation

47.48.2.1 CublasError()

```
gko::CublasError::CublasError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a cuBLAS error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the cuBLAS routine that failed
<i>error_code</i>	The resulting cuBLAS error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.49 gko::cuda_stream Class Reference

An RAII wrapper for a custom CUDA stream.

```
#include <ginkgo/core/base/stream.hpp>
```

Public Member Functions

- [cuda_stream](#) ()
Creates an empty stream wrapper, representing the default stream.
- [cuda_stream](#) (int device_id)
Creates a new custom CUDA stream on the given device.
- [~cuda_stream](#) ()
Destroys the custom CUDA stream, if it isn't empty.
- [cuda_stream](#) ([cuda_stream](#) &&)
Move-constructs from an existing stream, which will be emptied.
- [cuda_stream](#) & [operator=](#) ([cuda_stream](#) &&)=delete
Move-assigns from an existing stream, which will be emptied.
- CUstream_st * [get](#) () const
Returns the native CUDA stream handle.

47.49.1 Detailed Description

An RAI wrapper for a custom CUDA stream.

The stream will be created on construction and destroyed when the lifetime of the wrapper ends.

47.49.2 Constructor & Destructor Documentation

47.49.2.1 cuda_stream()

```
gko::cuda_stream::cuda_stream (
    int device_id )
```

Creates a new custom CUDA stream on the given device.

Parameters

<i>device_id</i>	the device ID to create the stream on.
------------------	--

47.49.3 Member Function Documentation

47.49.3.1 get()

```
CUstream_st* gko::cuda_stream::get ( ) const
```

Returns the native CUDA stream handle.

In an empty [cuda_stream](#), this will return nullptr.

The documentation for this class was generated from the following file:

- ginkgo/core/base/stream.hpp

47.50 gko::CudaAllocator Class Reference

[Allocator](#) using cudaMalloc.

```
#include <ginkgo/core/base/memory.hpp>
```

47.50.1 Detailed Description

[Allocator](#) using `cudaMalloc`.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/memory.hpp`

47.51 `gko::CudaAllocatorBase` Class Reference

Implement this interface to provide an allocator for [CudaExecutor](#).

```
#include <ginkgo/core/base/memory.hpp>
```

47.51.1 Detailed Description

Implement this interface to provide an allocator for [CudaExecutor](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/memory.hpp`

47.52 `gko::CudaError` Class Reference

[CudaError](#) is thrown when a CUDA routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [CudaError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a CUDA error.

47.52.1 Detailed Description

[CudaError](#) is thrown when a CUDA routine throws a non-zero error code.

47.52.2 Constructor & Destructor Documentation

47.52.2.1 `CudaError()`

```
gko::CudaError::CudaError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a CUDA error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the CUDA routine that failed
<i>error_code</i>	The resulting CUDA error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.53 gko::CudaExecutor Class Reference

This is the [Executor](#) subclass which represents the CUDA device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- `std::shared_ptr< Executor > get\_master ()` noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `std::shared_ptr< const Executor > get\_master ()` const noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `void synchronize ()` const override
Synchronize the operations launched on the executor with its master.
- `std::string get\_description ()` const override
- `int get\_device\_id ()` const noexcept
Get the CUDA device id of the device associated to this executor.
- `int get\_num\_warps\_per\_sm ()` const noexcept
Get the number of warps per SM of this executor.
- `int get\_num\_multiprocessor ()` const noexcept
Get the number of multiprocessor of this executor.
- `int get\_num\_warps ()` const noexcept
Get the number of warps of this executor.
- `int get\_warp\_size ()` const noexcept
Get the warp size of this executor.
- `int get\_major\_version ()` const noexcept
Get the major version of compute capability.
- `int get\_minor\_version ()` const noexcept
Get the minor version of compute capability.
- `cublasContext * get\_cublas\_handle ()` const
Get the cublas handle for this executor.
- `cublasContext * get\_blas\_handle ()` const
Get the cublas handle for this executor.
- `cusparseContext * get\_cusparses\_handle ()` const
Get the cusparse handle for this executor.
- `cusparseContext * get\_sparselib\_handle ()` const
Get the cusparse handle for this executor.

- `std::vector< int > get\_closest\_pus () const`
Get the closest PUs.
- `int get\_closest\_numa () const`
Get the closest NUMA node.
- `CUstream_st * get\_stream () const`
Returns the CUDA stream used by this executor.
- `virtual void run (const Operation &op) const=0`
Runs the specified [Operation](#) using this [Executor](#).
- `template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void run (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const
ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.
- `template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename Closure↵
Dpcpp >
void run (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↵
Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.

Static Public Member Functions

- `static std::shared_ptr< CudaExecutor > create (int device_id, std::shared_ptr< Executor > master, bool
device_reset, allocation_mode alloc_mode=default_cuda_alloc_mode, CUstream_st *stream=nullptr)`
Creates a new [CudaExecutor](#).
- `static std::shared_ptr< CudaExecutor > create (int device_id, std::shared_ptr< Executor > mas-
ter, std::shared_ptr< CudaAllocatorBase > alloc=std::make_shared< CudaAllocator >(), CUstream_st
*stream=nullptr)`
Creates a new [CudaExecutor](#) with a custom allocator and device stream.
- `static int get\_num\_devices ()`
Get the number of devices present on the system.

47.53.1 Detailed Description

This is the [Executor](#) subclass which represents the CUDA device.

47.53.2 Member Function Documentation

47.53.2.1 `create()` [1/2]

```
static std::shared_ptr<CudaExecutor> gko::CudaExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    bool device_reset,
    allocation\_mode alloc_mode = default_cuda_alloc_mode,
    CUstream_st * stream = nullptr ) [static]
```

Creates a new [CudaExecutor](#).

Parameters

<i>device_id</i>	the CUDA device id of this device
<i>master</i>	an executor on the host that is used to invoke the device kernels
<i>device_reset</i>	this option no longer has any effect.
<i>alloc_mode</i>	the allocation mode that the executor should operate on. See @allocation_mode for more details
<i>stream</i>	the stream to execute operations on.

47.53.2.2 create() [2/2]

```
static std::shared_ptr<CudaExecutor> gko::CudaExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    std::shared_ptr< CudaAllocatorBase > alloc = std::make_shared< CudaAllocator > (),
    CUstream_st * stream = nullptr ) [static]
```

Creates a new [CudaExecutor](#) with a custom allocator and device stream.

Parameters

<i>device_id</i>	the CUDA device id of this device
<i>master</i>	an executor on the host that is used to invoke the device kernels.
<i>alloc</i>	the allocator to use for device memory allocations.
<i>stream</i>	the stream to execute operations on.

47.53.2.3 get_blas_handle()

```
cublasContext* gko::CudaExecutor::get_blas_handle ( ) const [inline]
```

Get the cublas handle for this executor.

Returns

the cublas handle (cublasContext*) for this executor

47.53.2.4 get_closest_numa()

```
int gko::CudaExecutor::get_closest_numa ( ) const [inline]
```

Get the closest NUMA node.

Returns

the closest NUMA node closest to this device

47.53.2.5 `get_closest_pus()`

```
std::vector<int> gko::CudaExecutor::get_closest_pus ( ) const [inline]
```

Get the closest PUs.

Returns

the array of PUs closest to this device

47.53.2.6 `get_cublas_handle()`

```
cublasContext* gko::CudaExecutor::get_cublas_handle ( ) const [inline]
```

Get the cublas handle for this executor.

Returns

the cublas handle (cublasContext*) for this executor

47.53.2.7 `get_cuspars_handle()`

```
cusparsContext* gko::CudaExecutor::get_cuspars_handle ( ) const [inline]
```

Get the cuspars handle for this executor.

Returns

the cuspars handle (cusparsContext*) for this executor

47.53.2.8 `get_description()`

```
std::string gko::CudaExecutor::get_description ( ) const [override], [virtual]
```

Returns

a textual representation of the executor and its device.

Implements [gko::Executor](#).

47.53.2.9 get_master() [1/2]

```
std::shared_ptr<const Executor> gko::CudaExecutor::get_master ( ) const [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.53.2.10 get_master() [2/2]

```
std::shared_ptr<Executor> gko::CudaExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.53.2.11 get_sparselib_handle()

```
cusparseContext* gko::CudaExecutor::get_sparselib_handle ( ) const [inline]
```

Get the cusparse handle for this executor.

Returns

the cusparse handle (cusparseContext*) for this executor

47.53.2.12 get_stream()

```
CUstream_st* gko::CudaExecutor::get_stream ( ) const [inline]
```

Returns the CUDA stream used by this executor.

Can be nullptr for the default stream.

Returns

the stream used to execute kernels and memory operations.

47.53.2.13 run() [1/3]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

47.53.2.14 run() [2/3]

```
virtual void gko::Executor::run
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

47.53.2.15 run() [3/3]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureReference ,
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of op_ref
<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.54 gko::CudaTimer Class Reference

A timer using events for timing on a [CudaExecutor](#).

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void [record](#) ([time_point](#) &time) override
Records a time point at the current time.
- void [wait](#) ([time_point](#) &time) override
Waits until all kernels in-process when recording the time point are finished.
- std::chrono::nanoseconds [difference_async](#) (const [time_point](#) &start, const [time_point](#) &stop) override
Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

47.54.1 Detailed Description

A timer using events for timing on a [CudaExecutor](#).

Note

When using a [CudaExecutor](#) with a custom stream, make sure that the stream's lifetime is longer than the lifetime of this timer.

47.54.2 Member Function Documentation

47.54.2.1 difference_async()

```
std::chrono::nanoseconds gko::CudaTimer::difference_async (
    const time\_point & start,
    const time\_point & stop ) [override], [virtual]
```

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. `timer->wait(stop)` needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements [gko::Timer](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/timer.hpp`

47.55 gko::CufftError Class Reference

[CufftError](#) is thrown when a cuFFT routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [CufftError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a cuFFT error.

47.55.1 Detailed Description

[CufftError](#) is thrown when a cuFFT routine throws a non-zero error code.

47.55.2 Constructor & Destructor Documentation

47.55.2.1 CufftError()

```
gko::CufftError::CufftError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a cuFFT error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the cuFFT routine that failed
<i>error_code</i>	The resulting cuFFT error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.56 gko::CurandError Class Reference

[CurandError](#) is thrown when a cuRAND routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [CurandError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a cuRAND error.

47.56.1 Detailed Description

[CurandError](#) is thrown when a cuRAND routine throws a non-zero error code.

47.56.2 Constructor & Destructor Documentation

47.56.2.1 CurandError()

```
gko::CurandError::CurandError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a cuRAND error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the cuRAND routine that failed
<i>error_code</i>	The resulting cuRAND error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.57 gko::matrix::Csr< ValueType, IndexType >::cuspars Class Reference

cuspars is a [strategy_type](#) which uses the sparselib csr.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [cuspars](#) ()
Creates a cuspars strategy.
- void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow) override
Computes srow according to row pointers.
- int64_t [clac_size](#) (const int64_t nnz) override
Computes the srow size according to the number of nonzeros.
- std::shared_ptr< [strategy_type](#) > [copy](#) () override
Copy a strategy.

47.57.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::cuspars
```

cuspars is a [strategy_type](#) which uses the sparselib csr.

Note

cuspars is also known to the hip executor which converts between cuda and hip.

47.57.2 Member Function Documentation

47.57.2.1 clac_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int64_t gko::matrix::Csr< ValueType, IndexType >::cuspars::clac_size (
    const int64_t nnz ) [inline], [override], [virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.57.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::cusparse::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.57.2.3 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::cusparse::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [inline], [override], [virtual]
```

Computes srow according to row pointers.

Parameters

<i>mtx_row_ptrs</i>	the row pointers of the matrix
<i>mtx_srow</i>	the srow of the matrix

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/csr.hpp`

47.58 gko::CusparsedError Class Reference

[CusparsedError](#) is thrown when a cuSPARSE routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [CusparsedError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a cuSPARSE error.

47.58.1 Detailed Description

[CusparsedError](#) is thrown when a cuSPARSE routine throws a non-zero error code.

47.58.2 Constructor & Destructor Documentation

47.58.2.1 CusparsedError()

```
gko::CusparsedError::CusparsedError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a cuSPARSE error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the cuSPARSE routine that failed
<i>error_code</i>	The resulting cuSPARSE error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.59 gko::default_converter< S, R > Struct Template Reference

Used to convert objects of type *S* to objects of type *R* using `static_cast`.

```
#include <ginkgo/core/base/math.hpp>
```

Public Member Functions

- `R operator() (S val)`
Converts the object to result type.

47.59.1 Detailed Description

```
template<typename S, typename R>
struct gko::default_converter< S, R >
```

Used to convert objects of type `S` to objects of type `R` using `static_cast`.

Template Parameters

<code>S</code>	source type
<code>R</code>	result type

47.59.2 Member Function Documentation

47.59.2.1 operator>()

```
template<typename S , typename R >
R gko::default_converter< S, R >::operator() (
    S val ) [inline]
```

Converts the object to result type.

Parameters

<code>val</code>	the object to convert
------------------	-----------------------

Returns

the converted object

The documentation for this struct was generated from the following file:

- `ginkgo/core/base/math.hpp`

47.60 gko::deferred_factory_parameter< FactoryType > Class Template Reference

Represents a factory parameter of factory type that can either initialized by a pre-existing factory or by passing in a `factory_parameters` object whose `.on(exec)` will be called to instantiate a factory.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

- [deferred_factory_parameter](#) ()=default
Creates an empty deferred factory parameter.
- [deferred_factory_parameter](#) (std::nullptr_t)
Creates a deferred factory parameter returning a nullptr.
- `template<typename ConcreteFactoryType , std::enable_if_t< detail::is_pointer_convertible< ConcreteFactoryType, FactoryType >::value > * = nullptr>`
[deferred_factory_parameter](#) (std::shared_ptr< ConcreteFactoryType > factory)
Creates a deferred factory parameter from a preexisting factory with shared ownership.
- `template<typename ConcreteFactoryType , typename Deleter , std::enable_if_t< detail::is_pointer_convertible< ConcreteFactoryType, FactoryType >::value > * = nullptr>`
[deferred_factory_parameter](#) (std::unique_ptr< ConcreteFactoryType, Deleter > factory)
Creates a deferred factory parameter by taking ownership of a preexisting factory with unique ownership.
- `template<typename ParametersType , typename U = decltype(std::declval<ParametersType>()).on( std::shared_ptr<const Executor>{})), std::enable_if_t< detail::is_pointer_convertible< typename U::element_type, FactoryType >::value > * = nullptr>`
[deferred_factory_parameter](#) (ParametersType parameters)
Creates a deferred factory parameter object from a factory_parameters-like object.
- `std::shared_ptr< FactoryType > on` (std::shared_ptr< const [Executor](#) > exec) const
Instantiates the deferred parameter into an actual factory.
- `bool is_empty` () const
Returns true iff the parameter is empty.

47.60.1 Detailed Description

```
template<typename FactoryType>
class gko::deferred_factory_parameter< FactoryType >
```

Represents a factory parameter of factory type that can either initialized by a pre-existing factory or by passing in a `factory_parameters` object whose `.on(exec)` will be called to instantiate a factory.

Template Parameters

<i>FactoryType</i>	the type of factory that can be instantiated from this object.
--------------------	--

47.60.2 Constructor & Destructor Documentation

47.60.2.1 deferred_factory_parameter()

```
template<typename FactoryType>
template<typename ParametersType , typename U = decltype(std::declval<ParametersType>()).on(
std::shared_ptr<const Executor>{})), std::enable_if_t< detail::is_pointer_convertible<
typename U::element_type, FactoryType >::value > * = nullptr>
gko::deferred_factory_parameter< FactoryType >::deferred_factory_parameter (
    ParametersType parameters ) [inline]
```

Creates a deferred factory parameter object from a `factory_parameters`-like object.

To instantiate the actual factory, the parameter's `.on(exec)` function will be called.

47.60.3 Member Function Documentation

47.60.3.1 on()

```
template<typename FactoryType>
std::shared_ptr<FactoryType> gko::deferred_factory_parameter< FactoryType >::on (
    std::shared_ptr< const Executor > exec ) const [inline]
```

Instantiates the deferred parameter into an actual factory.

This will throw if the deferred factory parameter is empty.

The documentation for this class was generated from the following file:

- ginkgo/core/base/abstract_factory.hpp

47.61 gko::batch::matrix::Dense< ValueType > Class Template Reference

[Dense](#) is a batch matrix format which explicitly stores all values of the matrix in each of the batches.

```
#include <ginkgo/core/matrix/batch_dense.hpp>
```

Public Member Functions

- `std::unique_ptr< unbatch_type > create_view_for_item (size_type item_id)`
Creates a mutable view (of [gko::matrix::Dense](#) type) of one item of the [batch::matrix::Dense<value_type>](#) object.
- `std::unique_ptr< const unbatch_type > create_const_view_for_item (size_type item_id) const`
Creates a mutable view (of [gko::matrix::Dense](#) type) of one item of the [batch::matrix::Dense<value_type>](#) object.
- `size_type get_cumulative_offset (size_type batch_id) const`
Get the cumulative storage size offset.
- `value_type * get_values () noexcept`
Returns a pointer to the array of values of the multi-vector.
- `const value_type * get_const_values () const noexcept`
Returns a pointer to the array of values of the multi-vector.
- `value_type & at (size_type batch_id, size_type row, size_type col)`
Returns a single element for a particular batch item.
- `value_type at (size_type batch_id, size_type row, size_type col) const`
Returns a single element for a particular batch item.
- `ValueType & at (size_type batch_id, size_type idx) noexcept`
Returns a single element for a particular batch item.
- `ValueType at (size_type batch_id, size_type idx) const noexcept`
Returns a single element for a particular batch item.
- `value_type * get_values_for_item (size_type batch_id) noexcept`
Returns a pointer to the array of values of the matrix for a specific batch item.
- `const value_type * get_const_values_for_item (size_type batch_id) const noexcept`

- Returns a pointer to the array of values of the matrix for a specific batch item.*
- `size_type get_num_stored_elements ()` const noexcept
Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.
- `size_type get_num_elements_per_item ()` const noexcept
Returns the number of stored elements in each batch item.
- `Dense * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector.
- `Dense * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector with a linear combination of the given input vector.
- `const Dense * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x)` const
Apply the matrix to a multi-vector with a linear combination of the given input vector.
- `const Dense * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x)` const
- `void scale (const array< value_type > &row_scale, const array< value_type > &col_scale)`
Performs in-place row and column scaling for this matrix.
- `void scale_add (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const batch::matrix::Dense< value_type >> b)`
*Performs the operation $this = alpha * this + b$.*
- `void add_scaled_identity (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> beta)`
*Performs the operation $this = alpha * I + beta * this$.*

Static Public Member Functions

- `static std::unique_ptr< Dense > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size=batch_dim< 2 >{})`
Creates an uninitialized Dense matrix of the specified size.
- `static std::unique_ptr< Dense > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, array< value_type > values)`
Creates a Dense matrix from an already allocated (and initialized) array.
- `template<typename InputValueType >
static std::unique_ptr< Dense > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, std::initializer_list< InputValueType > values)`
create(std::shared_ptr<const Executor>,
- `static std::unique_ptr< const Dense > create_const (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &sizes, gko::detail::const_array_view< ValueType > &&values)`
Creates a constant (immutable) batch dense matrix from a constant array.

47.61.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::matrix::Dense< ValueType >
```

`Dense` is a batch matrix format which explicitly stores all values of the matrix in each of the batches.

The values in each of the batches are stored in row-major format (values belonging to the same row appear consecutive in the memory and the values of each batch item are also stored consecutively in memory).

Note

Though the storage layout is the same as the multi-vector object, the class semantics and the operations it aims to provide are different. Hence it is recommended to create multi-vector objects if the user means to view the data as a set of vectors.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.61.2 Member Function Documentation

47.61.2.1 add_scaled_identity()

```
template<typename ValueType = default_precision>
void gko::batch::matrix::Dense< ValueType >::add_scaled_identity (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> beta )
```

Performs the operation $\text{this} = \alpha * I + \beta * \text{this}$.

Parameters

<i>alpha</i>	the scalar for identity
<i>beta</i>	the scalar to multiply this matrix

Note

Performs the operation in-place for this batch matrix

47.61.2.2 apply() [1/4]

```
template<typename ValueType = default_precision>
Dense* gko::batch::matrix::Dense< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector with a linear combination of the given input vector.

Represents the matrix vector multiplication, $x = \alpha * A$

- $b + \beta * x$, where x and b are both multi-vectors.

Parameters

<i>alpha</i>	the scalar to scale the matrix-vector product with
<i>b</i>	the multi-vector to be applied to
<i>beta</i>	the scalar to scale the x vector with
<i>x</i>	the output multi-vector

47.61.2.3 apply() [2/4]

```
template<typename ValueType = default_precision>
const Dense* gko::batch::matrix::Dense< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x ) const
```

47.61.2.4 apply() [3/4]

```
template<typename ValueType = default_precision>
Dense* gko::batch::matrix::Dense< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector.

Represents the matrix vector multiplication, $x = A * b$, where x and b are both multi-vectors.

Parameters

b	the multi-vector to be applied to
x	the output multi-vector

47.61.2.5 apply() [4/4]

```
template<typename ValueType = default_precision>
const Dense* gko::batch::matrix::Dense< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x ) const
```

47.61.2.6 at() [1/4]

```
template<typename ValueType = default_precision>
ValueType gko::batch::matrix::Dense< ValueType >::at (
    size_type batch_id,
    size_type idx ) const [inline], [noexcept]
```

Returns a single element for a particular batch item.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU [Dense](#) object from the OMP may result in incorrect behaviour)

```

188     {
189         GKO_ASSERT(batch_id < this->get_num_batch_items());
190         return values_.get_const_data()[linearize_index(batch_id, idx)];
191     }

```

References `gko::array< ValueType >::get_const_data()`.

47.61.2.7 at() [2/4]

```

template<typename ValueType = default_precision>
ValueType& gko::batch::matrix::Dense< ValueType >::at (
    size_type batch_id,
    size_type idx ) [inline], [noexcept]

```

Returns a single element for a particular batch item.

Useful for iterating across all elements of the matrix. However, it is less efficient than the two-parameter variant of this method.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>idx</i>	a linear index of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU [Dense](#) object from the OMP may result in incorrect behaviour)

References `gko::array< ValueType >::get_data()`.

47.61.2.8 at() [3/4]

```

template<typename ValueType = default_precision>
value_type& gko::batch::matrix::Dense< ValueType >::at (

```

```
size_type batch_id,  
size_type row,  
size_type col ) [inline]
```

Returns a single element for a particular batch item.

Parameters

<i>batch</i> ↔ <i>_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU [Dense](#) object from the OMP may result in incorrect behaviour)

References `gko::array< ValueType >::get_data()`.

47.61.2.9 at() [4/4]

```
template<typename ValueType = default_precision>
value_type gko::batch::matrix::Dense< ValueType >::at (
    size_type batch_id,
    size_type row,
    size_type col ) const [inline]
```

Returns a single element for a particular batch item.

Parameters

<i>batch</i> ↔ <i>_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU [Dense](#) object from the OMP may result in incorrect behaviour)

References `gko::array< ValueType >::get_const_data()`.

47.61.2.10 create() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::batch::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    array< value_type > values ) [static]
```

Creates a [Dense](#) matrix from an already allocated (and initialized) array.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	sizes of the batch matrices in a batch_dim object
<i>values</i>	array of matrix values

Note

If *values* is not an rvalue, not an array of `ValueType`, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.61.2.11 create() [2/3]

```
template<typename ValueType = default_precision>
template<typename InputValueType >
static std::unique_ptr<Dense> gko::batch::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch\_dim< 2 > & size,
    std::initializer_list< InputValueType > values ) [inline], [static]
```

`create(std::shared_ptr<const Executor>,`

`create(std::shared_ptr<const Executor>, const batch\_dim<2>&, array<value_type>)`

References `gko::batch::matrix::Dense< ValueType >::create()`.

47.61.2.12 create() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::batch::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch\_dim< 2 > & size = batch\_dim< 2 >{} ) [static]
```

Creates an uninitialized [Dense](#) matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix

Returns

A smart pointer to the newly created matrix.

Referenced by gko::batch::matrix::Dense< ValueType >::create().

47.61.2.13 create_const()

```
template<typename ValueType = default_precision>
static std::unique_ptr<const Dense> gko::batch::matrix::Dense< ValueType >::create_const (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & sizes,
    gko::detail::const_array_view< ValueType > && values ) [static]
```

Creates a constant (immutable) batch dense matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

A smart pointer to the newly created matrix.

47.61.2.14 create_const_view_for_item()

```
template<typename ValueType = default_precision>
std::unique_ptr<const unbatch_type> gko::batch::matrix::Dense< ValueType >::create_const_↵
view_for_item (
    size_type item_id ) const
```

Creates a mutable view (of gko::matrix::Dense type) of one item of the batch::matrix::Dense<value_type> object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<i>item_↵ _id</i>	The index of the batch item
-----------------------	-----------------------------

Returns

a `gko::matrix::Dense` object with the data from the batch item at the given index.

47.61.2.15 create_view_for_item()

```
template<typename ValueType = default_precision>
std::unique_ptr<unbatch_type> gko::batch::matrix::Dense< ValueType >::create_view_for_item (
    size_type item_id )
```

Creates a mutable view (of `gko::matrix::Dense` type) of one item of the `batch::matrix::Dense<value_type>` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<i>item</i> ↔ <i>_id</i>	The index of the batch item
-----------------------------	-----------------------------

Returns

a `gko::matrix::Dense` object with the data from the batch item at the given index.

47.61.2.16 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::batch::matrix::Dense< ValueType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns a pointer to the array of values of the multi-vector.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.61.2.17 get_const_values_for_item()

```
template<typename ValueType = default_precision>
const value_type* gko::batch::matrix::Dense< ValueType >::get_const_values_for_item (
    size_type batch_id ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch</i> ↔ _id	the id of the batch item.
-----------------------	---------------------------

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data(), and gko::batch::matrix::Dense< ValueType >::get_↔cumulative_offset().

47.61.2.18 get_cumulative_offset()

```
template<typename ValueType = default_precision>
size_type gko::batch::matrix::Dense< ValueType >::get_cumulative_offset (
    size_type batch_id ) const [inline]
```

Get the cumulative storage size offset.

Parameters

<i>batch</i> ↔ _id	the batch id
-----------------------	--------------

Returns

the cumulative offset

Referenced by gko::batch::matrix::Dense< ValueType >::get_const_values_for_item(), and gko::batch::matrix::↔Dense< ValueType >::get_values_for_item().

47.61.2.19 get_num_elements_per_item()

```
template<typename ValueType = default_precision>
size_type gko::batch::matrix::Dense< ValueType >::get_num_elements_per_item ( ) const [inline],
[noexcept]
```

Returns the number of stored elements in each batch item.

Returns

the number of stored elements per batch item.

References gko::batch::matrix::Dense< ValueType >::get_num_stored_elements().

47.61.2.20 get_num_stored_elements()

```
template<typename ValueType = default_precision>
size_type gko::batch::matrix::Dense< ValueType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.

Returns

the number of elements explicitly stored in the vector, cumulative across all the batch items

References `gko::array< ValueType >::get_size()`.

Referenced by `gko::batch::matrix::Dense< ValueType >::get_num_elements_per_item()`.

47.61.2.21 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::batch::matrix::Dense< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the multi-vector.

Returns

the pointer to the array of values

References `gko::array< ValueType >::get_data()`.

47.61.2.22 get_values_for_item()

```
template<typename ValueType = default_precision>
value_type* gko::batch::matrix::Dense< ValueType >::get_values_for_item (
    size_type batch_id ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch_id</i>	the id of the batch item.
-----------------	---------------------------

Returns

the pointer to the array of values

References `gko::batch::matrix::Dense< ValueType >::get_cumulative_offset()`, and `gko::array< ValueType >::get_data()`.

47.61.2.23 `scale()`

```
template<typename ValueType = default_precision>
void gko::batch::matrix::Dense< ValueType >::scale (
    const array< value_type > & row_scale,
    const array< value_type > & col_scale )
```

Performs in-place row and column scaling for this matrix.

Parameters

<i>row_scale</i>	the row scalars
<i>col_scale</i>	the column scalars

47.61.2.24 `scale_add()`

```
template<typename ValueType = default_precision>
void gko::batch::matrix::Dense< ValueType >::scale_add (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const batch::matrix::Dense< value_type >> b )
```

Performs the operation `this = alpha*this + b`.

Parameters

<i>alpha</i>	the scalar to multiply this matrix
<i>b</i>	the matrix to add

Note

Performs the operation in-place for this batch matrix

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/batch_dense.hpp`

47.62 `gko::matrix::Dense< ValueType >` Class Template Reference

`Dense` is a matrix format which explicitly stores all values of the matrix.

```
#include <ginkgo/core/matrix/dense.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `void transpose (ptr_param< Dense > output)` const
Writes the transposed matrix into the given output matrix.
- `void conj_transpose (ptr_param< Dense > output)` const
Writes the conjugate-transposed matrix into the given output matrix.
- `void fill (const ValueType value)`
Fill the dense matrix with a given value.
- `std::unique_ptr< Dense > permute (ptr_param< const Permutation< int32 >> permutation, permute_mode mode=permute_mode::symmetric)` const
Creates a permuted copy A' of this matrix A with the given permutation P .
- `std::unique_ptr< Dense > permute (ptr_param< const Permutation< int64 >> permutation, permute_mode mode=permute_mode::symmetric)` const
- `void permute (ptr_param< const Permutation< int32 >> permutation, ptr_param< Dense > output, permute_mode mode)` const
Overload of `permute(ptr_param<const Permutation<int32>>, permute_mode)` that writes the permuted copy into an existing [Dense](#) matrix.
- `void permute (ptr_param< const Permutation< int64 >> permutation, ptr_param< Dense > output, permute_mode mode)` const
- `std::unique_ptr< Dense > permute (ptr_param< const Permutation< int32 >> row_permutation, ptr_param< const Permutation< int32 >> column_permutation, bool invert=false)` const
Creates a non-symmetrically permuted copy A' of this matrix A with the given row and column permutations P and Q .
- `std::unique_ptr< Dense > permute (ptr_param< const Permutation< int64 >> row_permutation, ptr_param< const Permutation< int64 >> column_permutation, bool invert=false)` const
- `void permute (ptr_param< const Permutation< int32 >> row_permutation, ptr_param< const Permutation< int32 >> column_permutation, ptr_param< Dense > output, bool invert=false)` const
Overload of `permute(ptr_param<const Permutation<int32>>, ptr_param<const Permutation<int32>>, permute_mode)` that writes the permuted copy into an existing [Dense](#) matrix.
- `void permute (ptr_param< const Permutation< int64 >> row_permutation, ptr_param< const Permutation< int64 >> column_permutation, ptr_param< Dense > output, bool invert=false)` const
- `std::unique_ptr< Dense > scale_permute (ptr_param< const ScaledPermutation< value_type, int32 >> permutation, permute_mode mode=permute_mode::symmetric)` const
Creates a scaled and permuted copy of this matrix.
- `std::unique_ptr< Dense > scale_permute (ptr_param< const ScaledPermutation< value_type, int64 >> permutation, permute_mode mode=permute_mode::symmetric)` const
- `void scale_permute (ptr_param< const ScaledPermutation< value_type, int32 >> permutation, ptr_param< Dense > output, permute_mode mode)` const
Overload of `scale_permute(ptr_param<const ScaledPermutation<value_type, int32>>, permute_mode)` that writes the permuted copy into an existing [Dense](#) matrix.
- `void scale_permute (ptr_param< const ScaledPermutation< value_type, int64 >> permutation, ptr_param< Dense > output, permute_mode mode)` const
- `std::unique_ptr< Dense > scale_permute (ptr_param< const ScaledPermutation< value_type, int32 >> row_permutation, ptr_param< const ScaledPermutation< value_type, int32 >> column_permutation, bool invert=false)` const
Creates a scaled and permuted copy of this matrix.
- `std::unique_ptr< Dense > scale_permute (ptr_param< const ScaledPermutation< value_type, int64 >> row_permutation, ptr_param< const ScaledPermutation< value_type, int64 >> column_permutation, bool invert=false)` const

- void `scale_permute` (`ptr_param`< const `ScaledPermutation`< `value_type`, `int32` >> `row_permutation`, `ptr_param`< const `ScaledPermutation`< `value_type`, `int32` >> `column_permutation`, `ptr_param`< `Dense` > `output`, `bool invert=false`) const
Overload of `scale_permute(ptr_param<const ScaledPermutation<value_type, int32>>, ptr_param<const ScaledPermutation<value_type, int32>>, bool)` that writes the permuted copy into an existing `Dense` matrix.
- void `scale_permute` (`ptr_param`< const `ScaledPermutation`< `value_type`, `int64` >> `row_permutation`, `ptr_param`< const `ScaledPermutation`< `value_type`, `int64` >> `column_permutation`, `ptr_param`< `Dense` > `output`, `bool invert=false`) const
- std::unique_ptr< `LinOp` > `permute` (const `array`< `int32` > *`permutation_indices`) const override
Returns a `LinOp` representing the symmetric row and column permutation of the `Permutable` object.
- std::unique_ptr< `LinOp` > `permute` (const `array`< `int64` > *`permutation_indices`) const override
Returns a `LinOp` representing the symmetric row and column permutation of the `Permutable` object.
- void `permute` (const `array`< `int32` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
Writes the symmetrically permuted matrix into the given output matrix.
- void `permute` (const `array`< `int64` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
- std::unique_ptr< `LinOp` > `inverse_permute` (const `array`< `int32` > *`permutation_indices`) const override
Returns a `LinOp` representing the symmetric inverse row and column permutation of the `Permutable` object.
- std::unique_ptr< `LinOp` > `inverse_permute` (const `array`< `int64` > *`permutation_indices`) const override
Returns a `LinOp` representing the symmetric inverse row and column permutation of the `Permutable` object.
- void `inverse_permute` (const `array`< `int32` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
Writes the inverse symmetrically permuted matrix into the given output matrix.
- void `inverse_permute` (const `array`< `int64` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
- std::unique_ptr< `LinOp` > `row_permute` (const `array`< `int32` > *`permutation_indices`) const override
Returns a `LinOp` representing the row permutation of the `Permutable` object.
- std::unique_ptr< `LinOp` > `row_permute` (const `array`< `int64` > *`permutation_indices`) const override
Returns a `LinOp` representing the row permutation of the `Permutable` object.
- void `row_permute` (const `array`< `int32` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
Writes the row-permuted matrix into the given output matrix.
- void `row_permute` (const `array`< `int64` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
- std::unique_ptr< `Dense` > `row_gather` (const `array`< `int32` > *`gather_indices`) const
Create a `Dense` matrix consisting of the given rows from this matrix.
- std::unique_ptr< `Dense` > `row_gather` (const `array`< `int64` > *`gather_indices`) const
Create a `Dense` matrix consisting of the given rows from this matrix.
- void `row_gather` (const `array`< `int32` > *`gather_indices`, `ptr_param`< `LinOp` > `row_collection`) const
Copies the given rows from this matrix into `row_collection`
- void `row_gather` (const `array`< `int64` > *`gather_indices`, `ptr_param`< `LinOp` > `row_collection`) const
- void `row_gather` (`ptr_param`< const `LinOp` > `alpha`, const `array`< `int32` > *`gather_indices`, `ptr_param`< const `LinOp` > `beta`, `ptr_param`< `LinOp` > `row_collection`) const
Copies the given rows from this matrix into `row_collection` with scaling.
- void `row_gather` (`ptr_param`< const `LinOp` > `alpha`, const `array`< `int64` > *`gather_indices`, `ptr_param`< const `LinOp` > `beta`, `ptr_param`< `LinOp` > `row_collection`) const
- std::unique_ptr< `LinOp` > `column_permute` (const `array`< `int32` > *`permutation_indices`) const override
Returns a `LinOp` representing the column permutation of the `Permutable` object.
- std::unique_ptr< `LinOp` > `column_permute` (const `array`< `int64` > *`permutation_indices`) const override
Returns a `LinOp` representing the column permutation of the `Permutable` object.
- void `column_permute` (const `array`< `int32` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
Writes the column-permuted matrix into the given output matrix.
- void `column_permute` (const `array`< `int64` > *`permutation_indices`, `ptr_param`< `Dense` > `output`) const
- std::unique_ptr< `LinOp` > `inverse_row_permute` (const `array`< `int32` > *`permutation_indices`) const override
Returns a `LinOp` representing the row permutation of the inverse permuted object.
- std::unique_ptr< `LinOp` > `inverse_row_permute` (const `array`< `int64` > *`permutation_indices`) const override
Returns a `LinOp` representing the row permutation of the inverse permuted object.

- void [inverse_row_permute](#) (const array< int32 > *permutation_indices, ptr_param< Dense > output) const
Writes the inverse row-permuted matrix into the given output matrix.
- void [inverse_row_permute](#) (const array< int64 > *permutation_indices, ptr_param< Dense > output) const
- std::unique_ptr< LinOp > [inverse_column_permute](#) (const array< int32 > *permutation_indices) const override
Returns a LinOp representing the row permutation of the inverse permuted object.
- std::unique_ptr< LinOp > [inverse_column_permute](#) (const array< int64 > *permutation_indices) const override
Returns a LinOp representing the row permutation of the inverse permuted object.
- void [inverse_column_permute](#) (const array< int32 > *permutation_indices, ptr_param< Dense > output) const
Writes the inverse column-permuted matrix into the given output matrix.
- void [inverse_column_permute](#) (const array< int64 > *permutation_indices, ptr_param< Dense > output) const
- std::unique_ptr< Diagonal< ValueType > > [extract_diagonal](#) () const override
Extracts the diagonal entries of the matrix into a vector.
- void [extract_diagonal](#) (ptr_param< Diagonal< ValueType > > output) const
Writes the diagonal of this matrix into an existing diagonal matrix.
- std::unique_ptr< absolute_type > [compute_absolute](#) () const override
Gets the AbsoluteLinOp.
- void [compute_absolute](#) (ptr_param< absolute_type > output) const
Writes the absolute values of this matrix into an existing matrix.
- void [compute_absolute_inplace](#) () override
Compute absolute inplace on each element.
- std::unique_ptr< complex_type > [make_complex](#) () const
Creates a complex copy of the original matrix.
- void [make_complex](#) (ptr_param< complex_type > result) const
Writes a complex copy of the original matrix to a given complex matrix.
- std::unique_ptr< real_type > [get_real](#) () const
Creates a new real matrix and extracts the real part of the original matrix into that.
- void [get_real](#) (ptr_param< real_type > result) const
Extracts the real part of the original matrix into a given real matrix.
- std::unique_ptr< real_type > [get_imag](#) () const
Creates a new real matrix and extracts the imaginary part of the original matrix into that.
- void [get_imag](#) (ptr_param< real_type > result) const
Extracts the imaginary part of the original matrix into a given real matrix.
- value_type * [get_values](#) () noexcept
Returns a pointer to the array of values of the matrix.
- const value_type * [get_const_values](#) () const noexcept
Returns a pointer to the array of values of the matrix.
- size_type [get_stride](#) () const noexcept
Returns the stride of the matrix.
- size_type [get_num_stored_elements](#) () const noexcept
Returns the number of elements explicitly stored in the matrix.
- value_type & [at](#) (size_type row, size_type col) noexcept
Returns a single element of the matrix.
- value_type [at](#) (size_type row, size_type col) const noexcept
Returns a single element of the matrix.
- ValueType & [at](#) (size_type idx) noexcept
Returns a single element of the matrix.
- ValueType [at](#) (size_type idx) const noexcept

- Returns a single element of the matrix.*
- void [scale](#) (ptr_param< const LinOp > alpha)
 - Scales the matrix with a scalar (aka: BLAS scal).*
- void [inv_scale](#) (ptr_param< const LinOp > alpha)
 - Scales the matrix with the inverse of a scalar.*
- void [add_scaled](#) (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)
 - Adds b scaled by alpha to the matrix (aka: BLAS axpy).*
- void [sub_scaled](#) (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)
 - Subtracts b scaled by alpha from the matrix (aka: BLAS axpy).*
- void [compute_dot](#) (ptr_param< const LinOp > b, ptr_param< LinOp > result) const
 - Computes the column-wise dot product of this matrix and b.*
- void [compute_dot](#) (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the column-wise dot product of this matrix and b.*
- void [compute_conj_dot](#) (ptr_param< const LinOp > b, ptr_param< LinOp > result) const
 - Computes the column-wise dot product of conj(this matrix) and b.*
- void [compute_conj_dot](#) (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the column-wise dot product of conj(this matrix) and b.*
- void [compute_norm2](#) (ptr_param< LinOp > result) const
 - Computes the column-wise Euclidean (L^2) norm of this matrix.*
- void [compute_norm2](#) (ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the column-wise Euclidean (L^2) norm of this matrix.*
- void [compute_norm1](#) (ptr_param< LinOp > result) const
 - Computes the column-wise (L^1) norm of this matrix.*
- void [compute_norm1](#) (ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the column-wise (L^1) norm of this matrix.*
- void [compute_squared_norm2](#) (ptr_param< LinOp > result) const
 - Computes the square of the column-wise Euclidean (L^2) norm of this matrix.*
- void [compute_squared_norm2](#) (ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the square of the column-wise Euclidean (L^2) norm of this matrix.*
- void [compute_mean](#) (ptr_param< LinOp > result) const
 - Computes the column-wise arithmetic mean of this matrix.*
- void [compute_mean](#) (ptr_param< LinOp > result, array< char > &tmp) const
 - Computes the column-wise arithmetic mean of this matrix.*
- std::unique_ptr< [Dense](#) > [create_submatrix](#) (const span &rows, const span &columns, const size_type stride)
 - Create a submatrix from the original matrix.*
- std::unique_ptr< [Dense](#) > [create_submatrix](#) (const span &rows, const span &columns)
 - Create a submatrix from the original matrix.*
- std::unique_ptr< real_type > [create_real_view](#) ()
 - Create a real view of the (potentially) complex original matrix.*
- std::unique_ptr< const real_type > [create_real_view](#) () const
 - Create a real view of the (potentially) complex original matrix.*
- [Dense](#) & operator= (const [Dense](#) &)
 - Copy-assigns a Dense matrix.*
- [Dense](#) & operator= ([Dense](#) &&)
 - Move-assigns a Dense matrix.*
- [Dense](#) (const [Dense](#) &)
 - Copy-constructs a Dense matrix.*
- [Dense](#) ([Dense](#) &&)
 - Move-constructs a Dense matrix.*

Static Public Member Functions

- static std::unique_ptr< [Dense](#) > [create_with_config_of](#) (ptr_param< const [Dense](#) > other)
Creates a [Dense](#) matrix with the same size and stride as another [Dense](#) matrix.
- static std::unique_ptr< [Dense](#) > [create_with_type_of](#) (ptr_param< const [Dense](#) > other, std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size=dim< 2 >{})
Creates a [Dense](#) matrix with the same type as another [Dense](#) matrix but on a different executor and with a different size.
- static std::unique_ptr< [Dense](#) > [create_with_type_of](#) (ptr_param< const [Dense](#) > other, std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size, size_type stride)
Creates a [Dense](#) matrix with the same type as another [Dense](#) matrix but on a different executor and with a different size.
- static std::unique_ptr< [Dense](#) > [create_with_type_of](#) (ptr_param< const [Dense](#) > other, std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size, const dim< 2 > &local_size, size_type stride)
Creates a [Dense](#) matrix with the same type as another [Dense](#) matrix but on a different executor and with a different size.
- static std::unique_ptr< [Dense](#) > [create_view_of](#) (ptr_param< [Dense](#) > other)
Creates a [Dense](#) matrix, where the underlying array is a view of another [Dense](#) matrix' array.
- static std::unique_ptr< const [Dense](#) > [create_const_view_of](#) (ptr_param< const [Dense](#) > other)
Creates a immutable [Dense](#) matrix, where the underlying array is a view of another [Dense](#) matrix' array.
- static std::unique_ptr< [Dense](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size={}, size_type stride=0)
Creates an uninitialized [Dense](#) matrix of the specified size.
- static std::unique_ptr< [Dense](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size, array< value_type > values, size_type stride)
Creates a [Dense](#) matrix from an already allocated (and initialized) array.
- template<typename InputValueType >
static std::unique_ptr< [Dense](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size, std::initializer_list< InputValueType > values, size_type stride)
create(std::shared_ptr<const Executor> ,
- static std::unique_ptr< const [Dense](#) > [create_const](#) (std::shared_ptr< const [Executor](#) > exec, const dim< 2 > &size, gko::detail::const_array_view< ValueType > &&values, size_type stride)
Creates a constant (immutable) [Dense](#) matrix from a constant array.

47.62.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::matrix::Dense< ValueType >
```

[Dense](#) is a matrix format which explicitly stores all values of the matrix.

The values are stored in row-major format (values belonging to the same row appear consecutive in the memory). Optionally, rows can be padded for better memory access.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

Note

While this format is not very useful for storing sparse matrices, it is often suitable to store vectors, and sets of vectors.

47.62.2 Constructor & Destructor Documentation

47.62.2.1 Dense() [1/2]

```
template<typename ValueType = default_precision>
gko::matrix::Dense< ValueType >::Dense (
    const Dense< ValueType > & )
```

Copy-constructs a [Dense](#) matrix.

Inherits executor and dimensions, but copies data without padding.

47.62.2.2 Dense() [2/2]

```
template<typename ValueType = default_precision>
gko::matrix::Dense< ValueType >::Dense (
    Dense< ValueType > && )
```

Move-constructs a [Dense](#) matrix.

Inherits executor, dimensions and data with padding. The moved-from object is empty (0x0 with empty Array).

47.62.3 Member Function Documentation**47.62.3.1 add_scaled()**

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::add_scaled (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b )
```

Adds *b* scaled by *alpha* to the matrix (aka: BLAS axpy).

Parameters

<i>alpha</i>	If <i>alpha</i> is 1x1 Dense matrix, the entire matrix is scaled by <i>alpha</i> . If it is a Dense row vector of values, then <i>i</i> -th column of the matrix is scaled with the <i>i</i> -th element of <i>alpha</i> (the number of columns of <i>alpha</i> has to match the number of columns of the matrix).
<i>b</i>	a matrix of the same dimension as this

47.62.3.2 at() [1/4]

```
template<typename ValueType = default_precision>
ValueType gko::matrix::Dense< ValueType >::at (
    size_type idx ) const [inline], [noexcept]
```

Returns a single element of the matrix.

Useful for iterating across all elements of the matrix. However, it is less efficient than the two-parameter variant of this method.

Parameters

<i>idx</i>	a linear index of the requested element (ignoring the stride)
------------	---

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.62.3.3 `at()` [2/4]

```
template<typename ValueType = default_precision>
ValueType& gko::matrix::Dense< ValueType >::at (
    size_type idx ) [inline], [noexcept]
```

Returns a single element of the matrix.

Useful for iterating across all elements of the matrix. However, it is less efficient than the two-parameter variant of this method.

Parameters

<i>idx</i>	a linear index of the requested element (ignoring the stride)
------------	---

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.62.3.4 `at()` [3/4]

```
template<typename ValueType = default_precision>
value_type gko::matrix::Dense< ValueType >::at (
    size_type row,
    size_type col ) const [inline], [noexcept]
```

Returns a single element of the matrix.

Parameters

<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.62.3.5 at() [4/4]

```
template<typename ValueType = default_precision>
value_type& gko::matrix::Dense< ValueType >::at (
    size_type row,
    size_type col ) [inline], [noexcept]
```

Returns a single element of the matrix.

Parameters

<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

Referenced by `gko::initialize()`.

47.62.3.6 column_permute() [1/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::column_permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the column permutation of the [Permutable](#) object.

In the resulting `LinOp`, the column `i` contains the input column `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^T .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new column permuted object

Implements [gko::Permutable< int32 >](#).

47.62.3.7 column_permute() [2/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::column_permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the column-permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size() [1]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::column_permute(const array<int32>*)`

47.62.3.8 column_permute() [3/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::column_permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the column permutation of the [Permutable](#) object.

In the resulting `LinOp`, the column `i` contains the input column `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^T .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new column permuted object

Implements [gko::Permutable< int64 >](#).

47.62.3.9 column_permute() [4/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::column_permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.10 compute_absolute() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<absolute_type> gko::matrix::Dense< ValueType >::compute_absolute ( ) const
[override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Dense< ValueType > > >](#).

47.62.3.11 compute_absolute() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_absolute (
    ptr_param< absolute_type > output ) const
```

Writes the absolute values of this matrix into an existing matrix.

Parameters

<i>output</i>	The output matrix. Its size must match the size of this matrix.
---------------	---

See also

[Dense::compute_absolute\(\)](#)

47.62.3.12 compute_conj_dot() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_conj_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result ) const
```

Computes the column-wise dot product of `conj(this matrix)` and `b`.

Parameters

<i>b</i>	a Dense matrix of same dimension as this
<i>result</i>	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)

47.62.3.13 compute_conj_dot() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_conj_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise dot product of `conj(this matrix)` and `b`.

Parameters

<i>b</i>	a Dense matrix of same dimension as this
<i>result</i>	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.14 compute_dot() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result ) const
```

Computes the column-wise dot product of this matrix and `b`.

Parameters

<i>b</i>	a Dense matrix of same dimension as this
<i>result</i>	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)

47.62.3.15 compute_dot() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_dot (
```

```
ptr_param< const LinOp > b,
ptr_param< LinOp > result,
array< char > & tmp ) const
```

Computes the column-wise dot product of this matrix and *b*.

Parameters

<i>b</i>	a Dense matrix of same dimension as this
<i>result</i>	a Dense row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.16 compute_mean() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_mean (
    ptr_param< LinOp > result ) const
```

Computes the column-wise arithmetic mean of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the mean (the number of columns in the vector must match the number of columns of this)
---------------	---

47.62.3.17 compute_mean() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_mean (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise arithmetic mean of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the mean (the number of columns in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.18 compute_norm1() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_norm1 (
    ptr_param< LinOp > result ) const
```

Computes the column-wise (L^1) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
---------------	---

47.62.3.19 compute_norm1() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_norm1 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise (L^1) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.20 compute_norm2() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_norm2 (
    ptr_param< LinOp > result ) const
```

Computes the column-wise Euclidean (L^2) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
---------------	---

47.62.3.21 compute_norm2() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_norm2 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise Euclidean (L^2) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.22 compute_squared_norm2() [1/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_squared_norm2 (
    ptr_param< LinOp > result ) const
```

Computes the square of the column-wise Euclidean (L^2) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
---------------	---

47.62.3.23 compute_squared_norm2() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::compute_squared_norm2 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the square of the column-wise Euclidean (L^2) norm of this matrix.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.62.3.24 conj_transpose() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.62.3.25 conj_transpose() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::conj_transpose (
    ptr_param< Dense< ValueType > > output ) const
```

Writes the conjugate-transposed matrix into the given output matrix.

Parameters

<i>output</i>	The output matrix. It must have the dimensions <code>gko::transpose(this->get_size())</code>
---------------	---

47.62.3.26 create() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< value_type > values,
    size_type stride ) [static]
```

Creates a [Dense](#) matrix from an already allocated (and initialized) array.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values
<i>stride</i>	stride of the rows (i.e. offset between the first elements of two consecutive rows, expressed as the number of matrix elements)

Note

If `values` is not an rvalue, not an array of `ValueType`, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.62.3.27 create() [2/3]

```
template<typename ValueType = default_precision>
template<typename InputValueType >
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::initializer_list< InputValueType > values,
    size_type stride ) [inline], [static]
```

`create(std::shared_ptr<const Executor>,`

`create(std::shared_ptr<const Executor>, const dim<2>&, array<value_type>, size_type)`

47.62.3.28 create() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = {},
    size_type stride = 0 ) [static]
```

Creates an uninitialized `Dense` matrix of the specified size.

Parameters

<i>exec</i>	<code>Executor</code> associated to the matrix
<i>size</i>	size of the matrix
<i>stride</i>	stride of the rows (i.e. offset between the first elements of two consecutive rows, expressed as the number of matrix elements). If it is set to 0, <code>size[1]</code> will be used instead.

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::matrix::Dense< value_type >::create()`.

47.62.3.29 create_const()

```
template<typename ValueType = default_precision>
static std::unique_ptr<const Dense> gko::matrix::Dense< ValueType >::create_const (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    gko::detail::const_array_view< ValueType > && values,
    size_type stride ) [static]
```

Creates a constant (immutable) [Dense](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>stride</i>	the row-stride of the matrix

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

47.62.3.30 create_const_view_of()

```
template<typename ValueType = default_precision>
static std::unique_ptr<const Dense> gko::matrix::Dense< ValueType >::create_const_view_of (
    ptr_param< const Dense< ValueType > > > other ) [inline], [static]
```

Creates an immutable [Dense](#) matrix, where the underlying array is a view of another [Dense](#) matrix' array.

Parameters

<i>other</i>	The other matrix on which to create the view
--------------	--

Returns

A immutable [Dense](#) matrix that is a view of other

Referenced by `gko::make_const_dense_view()`.

47.62.3.31 create_real_view() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<real_type> gko::matrix::Dense< ValueType >::create_real_view ( )
```

Create a real view of the (potentially) complex original matrix.

If the original matrix is real, nothing changes. If the original matrix is complex, the result is created by viewing the complex matrix with as real with a `reinterpret_cast` with twice the number of columns and double the stride.

47.62.3.32 create_real_view() [2/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<const real_type> gko::matrix::Dense< ValueType >::create_real_view ( ) const
```

Create a real view of the (potentially) complex original matrix.

If the original matrix is real, nothing changes. If the original matrix is complex, the result is created by viewing the complex matrix with as real with a reinterpret_cast with twice the number of columns and double the stride.

47.62.3.33 create_submatrix() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_submatrix (
    const span & rows,
    const span & columns ) [inline]
```

Create a submatrix from the original matrix.

Parameters

<i>rows</i>	row span
<i>columns</i>	column span

47.62.3.34 create_submatrix() [2/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_submatrix (
    const span & rows,
    const span & columns,
    const size_type stride ) [inline]
```

Create a submatrix from the original matrix.

Warning: defining stride for this create_submatrix method might cause wrong memory access. Better use the create_submatrix(rows, columns) method instead.

Parameters

<i>rows</i>	row span
<i>columns</i>	column span
<i>stride</i>	stride of the new submatrix.

Referenced by gko::matrix::Dense< value_type >::create_submatrix().

47.62.3.35 create_view_of()

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_view_of (
    ptr_param< Dense< ValueType > > other ) [inline], [static]
```

Creates a [Dense](#) matrix, where the underlying array is a view of another [Dense](#) matrix' array.

Parameters

<i>other</i>	The other matrix on which to create the view
--------------	--

Returns

A [Dense](#) matrix that is a view of other

Referenced by `gko::make_dense_view()`.

47.62.3.36 create_with_config_of()

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_with_config_of (
    ptr_param< const Dense< ValueType > > other ) [inline], [static]
```

Creates a [Dense](#) matrix with the same size and stride as another [Dense](#) matrix.

Parameters

<i>other</i>	The other matrix whose configuration needs to copied.
--------------	---

47.62.3.37 create_with_type_of() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_with_type_of (
    ptr_param< const Dense< ValueType > > other,
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    const dim< 2 > & local_size,
    size_type stride ) [inline], [static]
```

Parameters

<i>local_size</i>	Unused
<i>stride</i>	The stride of the new matrix.

Note

This is an overload to stay consistent with [gko::experimental::distributed::Vector](#)

47.62.3.38 create_with_type_of() [2/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_with_type_of (
    ptr_param< const Dense< ValueType > > other,
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size_type stride ) [inline], [static]
```

Parameters

<i>stride</i>	The stride of the new matrix.
---------------	-------------------------------

Note

This is an overload which allows full parameter specification.

47.62.3.39 create_with_type_of() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::create_with_type_of (
    ptr_param< const Dense< ValueType > > other,
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = dim<2>{} ) [inline], [static]
```

Creates a [Dense](#) matrix with the same type as another [Dense](#) matrix but on a different executor and with a different size.

Parameters

<i>other</i>	The other matrix whose type we target.
<i>exec</i>	The executor of the new matrix.
<i>size</i>	The size of the new matrix.
<i>stride</i>	The stride of the new matrix.

Returns

a [Dense](#) matrix with the type of other.

47.62.3.40 extract_diagonal() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Dense< ValueType >::extract_diagonal ( )
const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.62.3.41 extract_diagonal() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::extract_diagonal (
    ptr_param< Diagonal< ValueType >> output ) const
```

Writes the diagonal of this matrix into an existing diagonal matrix.

Parameters

<i>output</i>	The output matrix. Its size must match the size of this matrix's diagonal.
---------------	--

See also

[Dense::extract_diagonal\(\)](#)

47.62.3.42 fill()

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::fill (
    const ValueType value )
```

Fill the dense matrix with a given value.

Parameters

<i>value</i>	the value to be filled
--------------	------------------------

47.62.3.43 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::matrix::Dense< ValueType >::get_const_values ( ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.62.3.44 get_num_stored_elements()

```
template<typename ValueType = default_precision>
size_type gko::matrix::Dense< ValueType >::get_num_stored_elements ( ) const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

47.62.3.45 get_stride()

```
template<typename ValueType = default_precision>
size_type gko::matrix::Dense< ValueType >::get_stride ( ) const [inline], [noexcept]
```

Returns the stride of the matrix.

Returns

the stride of the matrix.

Referenced by gko::matrix::Dense< value_type >::create_submatrix().

47.62.3.46 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::matrix::Dense< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

47.62.3.47 inv_scale()

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inv_scale (
    ptr_param< const LinOp > alpha )
```

Scales the matrix with the inverse of a scalar.

Parameters

<i>alpha</i>	If alpha is 1x1 Dense matrix, the entire matrix is scaled by 1 / alpha. If it is a Dense row vector of values, then i-th column of the matrix is scaled with the inverse of the i-th element of alpha (the number of columns of alpha has to match the number of columns of the matrix).
--------------	--

47.62.3.48 inverse_column_permute() [1/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_column_permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column `perm[i]` contains the input column `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation AP^{-T} .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new inverse permuted object

Implements [gko::Permutable< int32 >](#).

47.62.3.49 inverse_column_permute() [2/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_column_permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the inverse column-permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size() [1]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::inverse_column_permute(const array<int32>*)`

47.62.3.50 inverse_column_permute() [3/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_column_permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the row permutation of the inverse permuted object.

In the resulting `LinOp`, the column `perm[i]` contains the input column `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^{-T} .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new inverse permuted object

Implements `gko::Permutable< int64 >`.

47.62.3.51 inverse_column_permute() [4/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_column_permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.52 inverse_permute() [1/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the symmetric inverse row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location $(\text{perm}[i], \text{perm}[j])$ contains the input value (i, j) .

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}AP^{-T}$.

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< int32 >](#).

47.62.3.53 inverse_permute() [2/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the inverse symmetrically permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size()[0]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::inverse_permute(const array<int32>*)`

47.62.3.54 inverse_permute() [3/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the symmetric inverse row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location $(\text{perm}[i], \text{perm}[j])$ contains the input value (i, j) .

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}AP^{-T}$.

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order.
----------------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< int64 >](#).

47.62.3.55 inverse_permute() [4/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.56 inverse_row_permute() [1/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_row_permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row `perm[i]` contains the input row `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}A$.

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order <code>perm</code> .
----------------------------------	---

Returns

a pointer to the new inverse permuted object

Implements [gko::Permutable< int32 >](#).

47.62.3.57 inverse_row_permute() [2/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_row_permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the inverse row-permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size() [0]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::inverse_row_permute(const array<int32>*)`

47.62.3.58 inverse_row_permute() [3/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::inverse_row_permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the row permutation of the inverse permuted object.

In the resulting `LinOp`, the row `perm[i]` contains the input row `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation $P^{-1}A$.

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order <code>perm</code> .
----------------------------	---

Returns

a pointer to the new inverse permuted object

Implements `gko::Permutable< int64 >`.

47.62.3.59 inverse_row_permute() [4/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::inverse_row_permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.60 make_complex() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<complex_type> gko::matrix::Dense< ValueType >::make_complex ( ) const
```

Creates a complex copy of the original matrix.

If the original matrix was real, the imaginary part of the result will be zero.

47.62.3.61 make_complex() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::make_complex (
    ptr_param< complex_type > result ) const
```

Writes a complex copy of the original matrix to a given complex matrix.

If the original matrix was real, the imaginary part of the result will be zero.

47.62.3.62 operator=() [1/2]

```
template<typename ValueType = default_precision>
Dense& gko::matrix::Dense< ValueType >::operator= (
    const Dense< ValueType > & )
```

Copy-assigns a [Dense](#) matrix.

Preserves the executor, reallocates the matrix with minimal stride if the dimensions don't match, then copies the data over, ignoring padding.

47.62.3.63 operator=() [2/2]

```
template<typename ValueType = default_precision>
Dense& gko::matrix::Dense< ValueType >::operator= (
    Dense< ValueType > && )
```

Move-assigns a [Dense](#) matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with empty Array).

47.62.3.64 permute() [1/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a LinOp representing the symmetric row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location (i, j) contains the input value $(\text{perm}[i], \text{perm}[j])$.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PAP^T .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< int32 >](#).

47.62.3.65 permute() [2/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the symmetrically permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size() [0]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::permute(const array<int32>*)`

47.62.3.66 permute() [3/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the symmetric row and column permutation of the [Permutable](#) object.

In the resulting `LinOp`, the entry at location (i, j) contains the input value $(\text{perm}[i], \text{perm}[j])$.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PAP^T .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Reimplemented from [gko::Permutable< int64 >](#).

47.62.3.67 permute() [4/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.68 permute() [5/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int32 >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

Creates a permuted copy A' of this matrix A with the given permutation P .

By default, this computes a symmetric permutation ([permute_mode::symmetric](#)). For the effect of the different permutation modes, see [permute_mode](#).

Parameters

<i>permutation</i>	The input permutation.
<i>mode</i>	The permutation mode. If permute_mode::inverse is set, we use the inverse permutation P^{-1} instead of P . If permute_mode::rows is set, the rows will be permuted. If permute_mode::columns is set, the columns will be permuted.

Returns

The permuted matrix.

47.62.3.69 permute() [6/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int32 >> permutation,
    ptr_param< Dense< ValueType > > output,
    permute_mode mode ) const
```

Overload of `permute(ptr_param<const Permutation<int32>>, permute_mode)` that writes the permuted copy into an existing [Dense](#) matrix.

Parameters

<i>output</i>	the output matrix.
---------------	--------------------

47.62.3.70 permute() [7/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int32 >> row_permutation,
    ptr_param< const Permutation< int32 >> column_permutation,
    bool invert = false ) const
```

Creates a non-symmetrically permuted copy A' of this matrix A with the given row and column permutations P and Q .

The operation will compute $A'(i, j) = A(p[i], q[j])$, or $A' = PAQ^T$ if `invert` is `false`, and $A'(p[i], q[j]) = A(i, j)$, or $A' = P^{-1}AQ^{-T}$ if `invert` is `true`.

Parameters

<i>row_permutation</i>	The permutation P to apply to the rows
<i>column_permutation</i>	The permutation Q to apply to the columns
<i>invert</i>	If set to <code>false</code> , uses the input permutations, otherwise uses their inverses P^{-1}, Q^{-1}

Returns

The permuted matrix.

47.62.3.71 permute() [8/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int32 >> row_permutation,
    ptr_param< const Permutation< int32 >> column_permutation,
    ptr_param< Dense< ValueType > > output,
    bool invert = false ) const
```

Overload of `permute(ptr_param<const Permutation<int32>>, ptr_param<const Permutation<int32>>, permute_mode)` that writes the permuted copy into an existing `Dense` matrix.

Parameters

<i>output</i>	the output matrix.
---------------	--------------------

47.62.3.72 permute() [9/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int64 >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

47.62.3.73 permute() [10/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int64 >> permutation,
    ptr_param< Dense< ValueType > > output,
    permute_mode mode ) const
```

47.62.3.74 permute() [11/12]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int64 >> row_permutation,
    ptr_param< const Permutation< int64 >> column_permutation,
    bool invert = false ) const
```

47.62.3.75 permute() [12/12]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::permute (
    ptr_param< const Permutation< int64 >> row_permutation,
    ptr_param< const Permutation< int64 >> column_permutation,
    ptr_param< Dense< ValueType > > output,
    bool invert = false ) const
```

47.62.3.76 row_gather() [1/6]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::row_gather (
    const array< int32 > * gather_indices ) const
```

Create a [Dense](#) matrix consisting of the given rows from this matrix.

Parameters

<i>gather_indices</i>	pointer to an array containing row indices from this matrix. It may contain duplicates.
-----------------------	---

Returns

Dense matrix on the same executor with the same number of columns and `gather_indices->get_size()` rows containing the gathered rows from this matrix: `output(i, j) = input(gather_indices(i), j)`

47.62.3.77 row_gather() [2/6]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_gather (
    const array< int32 > * gather_indices,
    ptr_param< LinOp > row_collection ) const
```

Copies the given rows from this matrix into `row_collection`

Parameters

<i>gather_indices</i>	pointer to an array containing row indices from this matrix. It may contain duplicates.
<i>row_collection</i>	pointer to a LinOp that will store the gathered rows: <code>row_collection(i, j) = input(gather_indices(i), j)</code> It must have the same number of columns as this matrix and <code>gather_indices->get_size()</code> rows.

47.62.3.78 row_gather() [3/6]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::row_gather (
    const array< int64 > * gather_indices ) const
```

Create a **Dense** matrix consisting of the given rows from this matrix.

Parameters

<i>gather_indices</i>	pointer to an array containing row indices from this matrix. It may contain duplicates.
-----------------------	---

Returns

Dense matrix on the same executor with the same number of columns and `gather_indices->get_size()` rows containing the gathered rows from this matrix: `output(i, j) = input(gather_indices(i), j)`

47.62.3.79 row_gather() [4/6]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_gather (
    const array< int64 > * gather_indices,
    ptr_param< LinOp > row_collection ) const
```

47.62.3.80 row_gather() [5/6]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_gather (
    ptr_param< const LinOp > alpha,
    const array< int32 > * gather_indices,
    ptr_param< const LinOp > beta,
    ptr_param< LinOp > row_collection ) const
```

Copies the given rows from this matrix into `row_collection` with scaling.

Parameters

<i>alpha</i>	scaling the result of row gathering
<i>gather_indices</i>	pointer to an array containing row indices from this matrix. It may contain duplicates.
<i>beta</i>	scaling the input <code>row_collection</code>
<i>row_collection</i>	pointer to a <code>LinOp</code> that will store the gathered rows: <code>row_collection(i, j) = input(gather_indices(i), j)</code> It must have the same number of columns as this matrix and <code>gather_indices->get_size()</code> rows.

47.62.3.81 row_gather() [6/6]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_gather (
    ptr_param< const LinOp > alpha,
    const array< int64 > * gather_indices,
    ptr_param< const LinOp > beta,
    ptr_param< LinOp > row_collection ) const
```

47.62.3.82 row_permute() [1/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::row_permute (
    const array< int32 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the row permutation of the [Permutable](#) object.

In the resulting `LinOp`, the row `i` contains the input row `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(i)}$, this represents the operation PA .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Implements [gko::Permutable< int32 >](#).

47.62.3.83 row_permute() [2/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_permute (
    const array< int32 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

Writes the row-permuted matrix into the given output matrix.

Parameters

<i>permutation_indices</i>	The array containing permutation indices. It must have <code>this->get_size() [0]</code> elements.
<i>output</i>	The output matrix. It must have the dimensions <code>this->get_size()</code>

See also

`Dense::row_permute(const array<int32>*)`

47.62.3.84 row_permute() [3/4]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::row_permute (
    const array< int64 > * permutation_indices ) const [override], [virtual]
```

Returns a `LinOp` representing the row permutation of the [Permutable](#) object.

In the resulting `LinOp`, the row `i` contains the input row `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PA .

Parameters

<i>permutation_indices</i>	the array of indices containing the permutation order.
----------------------------	--

Returns

a pointer to the new permuted object

Implements [gko::Permutable< int64 >](#).

47.62.3.85 row_permute() [4/4]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::row_permute (
    const array< int64 > * permutation_indices,
    ptr_param< Dense< ValueType > > output ) const
```

47.62.3.86 scale()

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::scale (
    ptr_param< const LinOp > alpha )
```

Scales the matrix with a scalar (aka: BLAS scal).

Parameters

<i>alpha</i>	If alpha is 1x1 Dense matrix, the entire matrix is scaled by alpha. If it is a Dense row vector of values, then i-th column of the matrix is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of columns of the matrix).
--------------	---

47.62.3.87 scale_permute() [1/8]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int32 >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

Creates a scaled and permuted copy of this matrix.

For an explanation of the permutation modes, see `permute(ptr_param<const Permutation<index_type>>, permute_mode)`

Parameters

<i>permutation</i>	The scaled permutation.
<i>mode</i>	The permutation mode.

Returns

The permuted matrix.

47.62.3.88 scale_permute() [2/8]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int32 >> permutation,
    ptr_param< Dense< ValueType > > output,
    permute_mode mode ) const
```

Overload of `scale_permute(ptr_param<const ScaledPermutation<value_type, int32>>, permute_mode)` that writes the permuted copy into an existing `Dense` matrix.

Parameters

<i>output</i>	the output matrix.
---------------	--------------------

47.62.3.89 scale_permute() [3/8]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int32 >> row_permutation,
    ptr_param< const ScaledPermutation< value_type, int32 >> column_permutation,
    bool invert = false ) const
```

Creates a scaled and permuted copy of this matrix.

For an explanation of the parameters, see `permute(ptr_param<const Permutation<index_type>>, ptr_param<const Permutation<index_type>>, permute_mode)`

Parameters

<i>row_permutation</i>	The scaled row permutation.
<i>column_permutation</i>	The scaled column permutation.
<i>invert</i>	If set to <code>false</code> , uses the input permutations, otherwise uses their inverses P^{-1}, Q^{-1}

Returns

The permuted matrix.

47.62.3.90 scale_permute() [4/8]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int32 >> row_permutation,
    ptr_param< const ScaledPermutation< value_type, int32 >> column_permutation,
    ptr_param< Dense< ValueType > > output,
    bool invert = false ) const
```

Overload of `scale_permute(ptr_param<const ScaledPermutation<value_type, int32>>, ptr_param<const ScaledPermutation<value_type, int32>>, bool)` that writes the permuted copy into an existing `Dense` matrix.

Parameters

<i>output</i>	the output matrix.
---------------	--------------------

47.62.3.91 scale_permute() [5/8]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int64 >> permutation,
    permute_mode mode = permute_mode::symmetric ) const
```

47.62.3.92 scale_permute() [6/8]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int64 >> permutation,
    ptr_param< Dense< ValueType > > output,
    permute_mode mode ) const
```

47.62.3.93 scale_permute() [7/8]

```
template<typename ValueType = default_precision>
std::unique_ptr<Dense> gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int64 >> row_permutation,
    ptr_param< const ScaledPermutation< value_type, int64 >> column_permutation,
    bool invert = false ) const
```

47.62.3.94 scale_permute() [8/8]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::scale_permute (
    ptr_param< const ScaledPermutation< value_type, int64 >> row_permutation,
    ptr_param< const ScaledPermutation< value_type, int64 >> column_permutation,
    ptr_param< Dense< ValueType > > output,
    bool invert = false ) const
```

47.62.3.95 sub_scaled()

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::sub_scaled (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b )
```

Subtracts *b* scaled by *alpha* from the matrix (aka: BLAS axpy).

Parameters

<i>alpha</i>	If <i>alpha</i> is 1x1 Dense matrix, <i>b</i> is scaled by <i>alpha</i> . If it is a Dense row vector of values, then <i>i</i> -th column of <i>b</i> is scaled with the <i>i</i> -th element of <i>alpha</i> (the number of columns of <i>alpha</i> has to match the number of columns of the matrix).
<i>b</i>	a matrix of the same dimension as this

47.62.3.96 transpose() [1/2]

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Dense< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a *LinOp* representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.62.3.97 transpose() [2/2]

```
template<typename ValueType = default_precision>
void gko::matrix::Dense< ValueType >::transpose (
    ptr_param< Dense< ValueType > > output ) const
```

Writes the transposed matrix into the given output matrix.

Parameters

<i>output</i>	The output matrix. It must have the dimensions <code>gko::transpose(this->get_size())</code>
---------------	---

The documentation for this class was generated from the following files:

- ginkgo/core/base/dense_cache.hpp
- ginkgo/core/matrix/dense.hpp

47.63 gko::device_matrix_data< ValueType, IndexType > Class Template Reference

This type is a device-side equivalent to [matrix_data](#).

```
#include <ginkgo/core/base/device_matrix_data.hpp>
```

Classes

- struct [arrays](#)
Stores the internal arrays of a [device_matrix_data](#) object.

Public Member Functions

- [device_matrix_data](#) (std::shared_ptr< const [Executor](#) > exec, [dim](#)< 2 > size={}, [size_type](#) num_entries=0)
Initializes a new [device_matrix_data](#) object.
- [device_matrix_data](#) (std::shared_ptr< const [Executor](#) > exec, const [device_matrix_data](#) &data)
Initializes a [device_matrix_data](#) object by copying an existing object on another executor.
- [device_matrix_data](#) (std::shared_ptr< const [Executor](#) > exec, [dim](#)< 2 > size, [array](#)< index_type > row_idx, [array](#)< index_type > col_idx, [array](#)< value_type > values)
Initializes a new [device_matrix_data](#) object from existing data.
- template<typename InputValueType, typename RowIndexType, typename ColIndexType >
[device_matrix_data](#) (std::shared_ptr< const [Executor](#) > exec, [dim](#)< 2 > size, std::initializer_list< RowIndexType > row_idx, std::initializer_list< ColIndexType > col_idx, std::initializer_list< InputValueType > values)
Initializes a new [device_matrix_data](#) object from existing data.
- [host_type copy_to_host](#) () const
Copies the [device_matrix_data](#) entries to the host to return a regular [matrix_data](#) object with the same dimensions and entries.
- void [fill_zero](#) ()
Fills the matrix entries with zeros.
- void [sort_row_major](#) ()
Sorts the matrix entries in row-major order. This means that they will be sorted by row index first, and then by column index inside each row.
- void [remove_zeros](#) ()
Removes all zero entries from the storage.
- void [sum_duplicates](#) ()
Sums up all duplicate entries pointing to the same non-zero location.
- std::shared_ptr< const [Executor](#) > [get_executor](#) () const
Returns the executor used to store the [device_matrix_data](#) entries.

- `dim< 2 > get_size ()` const
Returns the dimensions of the matrix.
- `size_type get_num_elems ()` const
Returns the number of stored elements of the matrix.
- `size_type get_num_stored_elements ()` const
Returns the number of stored elements of the matrix.
- `index_type * get_row_idxes ()`
Returns a pointer to the row index array.
- `const index_type * get_const_row_idxes ()` const
Returns a pointer to the constant row index array.
- `index_type * get_col_idxes ()`
Returns a pointer to the column index array.
- `const index_type * get_const_col_idxes ()` const
Returns a pointer to the constant column index array.
- `value_type * get_values ()`
Returns a pointer to the value array.
- `const value_type * get_const_values ()` const
Returns a pointer to the constant value array.
- `void resize_and_reset (size_type new_num_entries)`
Resizes the internal storage to the given number of stored matrix entries.
- `void resize_and_reset (dim< 2 > new_size, size_type new_num_entries)`
Resizes the matrix and internal storage to the given dimensions.
- `arrays_empty_out ()`
Moves out the internal arrays of the [device_matrix_data](#) object and resets it to an empty 0x0 matrix.

Static Public Member Functions

- static `device_matrix_data create_from_host` (std::shared_ptr< const [Executor](#) > exec, const `host_type` &data)
Creates a [device_matrix_data](#) object from the given host data on the given executor.

47.63.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::device_matrix_data< ValueType, IndexType >
```

This type is a device-side equivalent to [matrix_data](#).

It stores the data necessary to initialize any matrix format in Ginkgo in individual value, column and row index arrays together with associated matrix dimensions. [matrix_data](#) uses array-of-Structs storage (AoS), while [device_matrix_data](#) uses Struct-of-Arrays (SoA).

Note

To be used with a Ginkgo matrix type, the entry array must be sorted in row-major order, i.e. by row index, then by column index within rows. This can be achieved by calling the `sort_row_major` function.

The data must not contain any duplicate (row, column) pairs.

Template Parameters

<i>ValueType</i>	the type used to store matrix values
<i>IndexType</i>	the type used to store matrix row and column indices

47.63.2 Constructor & Destructor Documentation

47.63.2.1 device_matrix_data() [1/4]

```
template<typename ValueType, typename IndexType>
gko::device_matrix_data< ValueType, IndexType >::device_matrix_data (
    std::shared_ptr< const Executor > exec,
    dim< 2 > size = {},
    size_type num_entries = 0 ) [explicit]
```

Initializes a new `device_matrix_data` object.

It uses the given executor to allocate storage for the given number of entries and matrix dimensions.

Parameters

<i>exec</i>	the executor to be used to store the matrix entries
<i>size</i>	the matrix dimensions
<i>num_entries</i>	the number of entries to be stored

47.63.2.2 device_matrix_data() [2/4]

```
template<typename ValueType, typename IndexType>
gko::device_matrix_data< ValueType, IndexType >::device_matrix_data (
    std::shared_ptr< const Executor > exec,
    const device_matrix_data< ValueType, IndexType > & data )
```

Initializes a `device_matrix_data` object by copying an existing object on another executor.

Parameters

<i>exec</i>	the executor to be used to store the matrix entries
<i>data</i>	the <code>device_matrix_data</code> object to copy, potentially stored on another executor.

47.63.2.3 device_matrix_data() [3/4]

```
template<typename ValueType, typename IndexType>
```

```

gko::device_matrix_data< ValueType, IndexType >::device_matrix_data (
    std::shared_ptr< const Executor > exec,
    dim< 2 > size,
    array< index_type > row_idx,
    array< index_type > col_idx,
    array< value_type > values )

```

Initializes a new `device_matrix_data` object from existing data.

Parameters

<i>size</i>	the matrix dimensions
<i>values</i>	the array containing the matrix values
<i>col_idx</i>	the array containing the matrix column indices
<i>row_idx</i>	the array containing the matrix row indices

47.63.2.4 device_matrix_data() [4/4]

```

template<typename ValueType, typename IndexType>
template<typename InputValueType, typename RowIndexType, typename ColIndexType >
gko::device_matrix_data< ValueType, IndexType >::device_matrix_data (
    std::shared_ptr< const Executor > exec,
    dim< 2 > size,
    std::initializer_list< RowIndexType > row_idx,
    std::initializer_list< ColIndexType > col_idx,
    std::initializer_list< InputValueType > values ) [inline]

```

Parameters

<i>size</i>	the matrix dimensions
<i>values</i>	the array containing the matrix values
<i>col_idx</i>	the array containing the matrix column indices
<i>row_idx</i>	the array containing the matrix row indices

```

91      : device_matrix_data(exec, size,
92                          array<index_type>(exec, std::move(row_idx)),
93                          array<index_type>(exec, std::move(col_idx)),
94                          array<value_type>(exec, std::move(values)))
95      {}

```

47.63.3 Member Function Documentation

47.63.3.1 copy_to_host()

```

template<typename ValueType, typename IndexType>
host_type gko::device_matrix_data< ValueType, IndexType >::copy_to_host ( ) const

```

Copies the `device_matrix_data` entries to the host to return a regular `matrix_data` object with the same dimensions and entries.

Returns

a [matrix_data](#) object with the same dimensions and entries.

Referenced by `gko::ReadableFromMatrixData< ValueType, int32 >::read()`.

47.63.3.2 create_from_host()

```
template<typename ValueType, typename IndexType>
static device_matrix_data gko::device_matrix_data< ValueType, IndexType >::create_from_host (
    std::shared_ptr< const Executor > exec,
    const host_type & data ) [static]
```

Creates a [device_matrix_data](#) object from the given host data on the given executor.

Parameters

<i>exec</i>	the executor to create the device_matrix_data on.
<i>data</i>	the data to be wrapped or copied into a device_matrix_data .

Returns

a [device_matrix_data](#) object with the same size and entries as `data` copied to the device executor.

47.63.3.3 empty_out()

```
template<typename ValueType, typename IndexType>
arrays gko::device_matrix_data< ValueType, IndexType >::empty_out ( )
```

Moves out the internal arrays of the [device_matrix_data](#) object and resets it to an empty 0x0 matrix.

Returns

a struct containing the internal arrays.

47.63.3.4 get_col_idxs()

```
template<typename ValueType, typename IndexType>
index_type* gko::device_matrix_data< ValueType, IndexType >::get_col_idxs ( ) [inline]
```

Returns a pointer to the column index array.

Returns

a pointer to the column index array

References `gko::array< ValueType >::get_data()`.

47.63.3.5 get_const_col_idxs()

```
template<typename ValueType, typename IndexType>
const index_type* gko::device_matrix_data< ValueType, IndexType >::get_const_col_idxs ( )
const [inline]
```

Returns a pointer to the constant column index array.

Returns

a pointer to the constant column index array

References gko::array< ValueType >::get_const_data().

47.63.3.6 get_const_row_idxs()

```
template<typename ValueType, typename IndexType>
const index_type* gko::device_matrix_data< ValueType, IndexType >::get_const_row_idxs ( )
const [inline]
```

Returns a pointer to the constant row index array.

Returns

a pointer to the constant row index array

References gko::array< ValueType >::get_const_data().

47.63.3.7 get_const_values()

```
template<typename ValueType, typename IndexType>
const value_type* gko::device_matrix_data< ValueType, IndexType >::get_const_values ( ) const
[inline]
```

Returns a pointer to the constant value array.

Returns

a pointer to the constant value array

References gko::array< ValueType >::get_const_data().

47.63.3.8 get_executor()

```
template<typename ValueType, typename IndexType>
std::shared_ptr<const Executor> gko::device_matrix_data< ValueType, IndexType >::get_executor
( ) const [inline]
```

Returns the executor used to store the `device_matrix_data` entries.

Returns

the executor used to store the `device_matrix_data` entries.

References `gko::array< ValueType >::get_executor()`.

47.63.3.9 get_num_elems()

```
template<typename ValueType, typename IndexType>
size_type gko::device_matrix_data< ValueType, IndexType >::get_num_elems ( ) const [inline]
```

Returns the number of stored elements of the matrix.

Returns

the number of stored elements of the matrix.

References `gko::device_matrix_data< ValueType, IndexType >::get_num_stored_elements()`.

47.63.3.10 get_num_stored_elements()

```
template<typename ValueType, typename IndexType>
size_type gko::device_matrix_data< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline]
```

Returns the number of stored elements of the matrix.

Returns

the number of stored elements of the matrix.

References `gko::array< ValueType >::get_size()`.

Referenced by `gko::device_matrix_data< ValueType, IndexType >::get_num_elems()`.

47.63.3.11 get_row_idx()

```
template<typename ValueType, typename IndexType>
index_type* gko::device_matrix_data< ValueType, IndexType >::get_row_idx ( ) [inline]
```

Returns a pointer to the row index array.

Returns

a pointer to the row index array

References `gko::array< ValueType >::get_data()`.

47.63.3.12 get_size()

```
template<typename ValueType, typename IndexType>
dim<2> gko::device_matrix_data< ValueType, IndexType >::get_size ( ) const [inline]
```

Returns the dimensions of the matrix.

Returns

the dimensions of the matrix.

47.63.3.13 get_values()

```
template<typename ValueType, typename IndexType>
value_type* gko::device_matrix_data< ValueType, IndexType >::get_values ( ) [inline]
```

Returns a pointer to the value array.

Returns

a pointer to the value array

References `gko::array< ValueType >::get_data()`.

47.63.3.14 remove_zeros()

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::remove_zeros ( )
```

Removes all zero entries from the storage.

This does not modify the storage if there are no zero entries, and keeps the relative order of nonzero entries otherwise.

47.63.3.15 resize_and_reset() [1/2]

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::resize_and_reset (
    dim< 2 > new_size,
    size_type new_num_entries )
```

Resizes the matrix and internal storage to the given dimensions.

The resulting storage should be assumed uninitialized.

Parameters

<i>new_size</i>	the new matrix dimensions.
<i>new_num_entries</i>	the new number of stored matrix entries.

47.63.3.16 `resize_and_reset()` [2/2]

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::resize_and_reset (
    size_type new_num_entries )
```

Resizes the internal storage to the given number of stored matrix entries.

The resulting storage should be assumed uninitialized.

Parameters

<i>new_num_entries</i>	the new number of stored matrix entries.
------------------------	--

47.63.3.17 `sum_duplicates()`

```
template<typename ValueType, typename IndexType>
void gko::device_matrix_data< ValueType, IndexType >::sum_duplicates ( )
```

Sums up all duplicate entries pointing to the same non-zero location.

The output will be sorted in row-major order, and it will only reallocate if duplicates exist.

The documentation for this class was generated from the following file:

- ginkgo/core/base/device_matrix_data.hpp

47.64 `gko::matrix::Diagonal< ValueType >` Class Template Reference

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

```
#include <ginkgo/core/matrix/diagonal.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `std::unique_ptr< absolute_type > compute_absolute ()` const override
Gets the AbsoluteLinOp.
- `void compute_absolute_inplace ()` override
Compute absolute inplace on each element.
- `value_type * get_values ()` noexcept
Returns a pointer to the array of values of the matrix.
- `const value_type * get_const_values ()` const noexcept
Returns a pointer to the array of values of the matrix.
- `void rapply (ptr_param< const LinOp > b, ptr_param< LinOp > x)` const
Applies the diagonal matrix from the right side to a matrix b, which means scales the columns of b with the according diagonal entries.
- `void inverse_apply (ptr_param< const LinOp > b, ptr_param< LinOp > x)` const
Applies the inverse of the diagonal matrix to a matrix b, which means scales the columns of b with the inverse of the according diagonal entries.

Static Public Member Functions

- `static std::unique_ptr< Diagonal > create (std::shared_ptr< const Executor > exec, size_type size=0)`
Creates an [Diagonal](#) matrix of the specified size.
- `static std::unique_ptr< Diagonal > create (std::shared_ptr< const Executor > exec, const size_type size, array< value_type > values)`
Creates a [Diagonal](#) matrix from an already allocated (and initialized) array.
- `template<typename InputValueType >`
`static std::unique_ptr< Diagonal > create (std::shared_ptr< const Executor > exec, const size_type size, std::initializer_list< InputValueType > values)`
create(std::shared_ptr<const
- `static std::unique_ptr< const Diagonal > create_const (std::shared_ptr< const Executor > exec, size_type size, gko::detail::const_array_view< ValueType > &&values)`
Creates a constant (immutable) [Diagonal](#) matrix from a constant array.

47.64.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::matrix::Diagonal< ValueType >
```

This class is a utility which efficiently implements the diagonal matrix (a linear operator which scales a vector row wise).

Objects of the [Diagonal](#) class always represent a square matrix, and require one array to store their values.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes of a CSR matrix the diagonal is applied or converted to.

47.64.2 Member Function Documentation

47.64.2.1 compute_absolute()

```
template<typename ValueType = default_precision>
std::unique_ptr<absolute_type> gko::matrix::Diagonal< ValueType >::compute_absolute ( ) const
[override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Diagonal< ValueType > > >](#).

47.64.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Diagonal< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.64.2.3 create() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Diagonal> gko::matrix::Diagonal< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const size_type size,
    array< value_type > values ) [static]
```

Creates a [Diagonal](#) matrix from an already allocated (and initialized) array.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values

Note

If `values` is not an rvalue, not an array of `ValueType`, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.64.2.4 create() [2/3]

```
template<typename ValueType = default_precision>
template<typename InputValueType >
static std::unique_ptr<Diagonal> gko::matrix::Diagonal< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const size_type size,
    std::initializer_list< InputValueType > values ) [inline], [static]
```

`create(std::shared_ptr<const`

```
create(std::shared_ptr<const executor>="", const size_type, array<value_type>)
222     {
223         return create(exec, size, array<value_type>{exec, std::move(values)});
224     }
```

References `gko::matrix::Diagonal< ValueType >::create()`.

47.64.2.5 create() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Diagonal> gko::matrix::Diagonal< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    size_type size = 0 ) [static]
```

Creates an `Diagonal` matrix of the specified size.

Parameters

<i>exec</i>	<code>Executor</code> associated to the matrix
<i>size</i>	size of the matrix

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::matrix::Diagonal< ValueType >::create()`.

47.64.2.6 create_const()

```
template<typename ValueType = default_precision>
static std::unique_ptr<const Diagonal> gko::matrix::Diagonal< ValueType >::create_const (
    std::shared_ptr< const Executor > exec,
    size_type size,
    gko::detail::const_array_view< ValueType > && values ) [static]
```

Creates a constant (immutable) [Diagonal](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the size of the square matrix
<i>values</i>	the value array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

47.64.2.7 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::matrix::Diagonal< ValueType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.64.2.8 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::matrix::Diagonal< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

References `gko::array< ValueType >::get_data()`.

47.64.2.9 inverse_apply()

```
template<typename ValueType = default_precision>
void gko::matrix::Diagonal< ValueType >::inverse_apply (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x ) const [inline]
```

Applies the inverse of the diagonal matrix to a matrix b, which means scales the columns of b with the inverse of the according diagonal entries.

Parameters

<i>b</i>	the input vector(s) on which the inverse of the diagonal matrix is applied
<i>x</i>	the output vector(s) where the result is stored

References `gko::ptr_param< T >::get()`.

47.64.2.10 rapply()

```
template<typename ValueType = default_precision>
void gko::matrix::Diagonal< ValueType >::rapply (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > x ) const [inline]
```

Applies the diagonal matrix from the right side to a matrix b, which means scales the columns of b with the according diagonal entries.

Parameters

<i>b</i>	the input vector(s) on which the diagonal matrix is applied
<i>x</i>	the output vector(s) where the result is stored

References `gko::ptr_param< T >::get()`.

47.64.2.11 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Diagonal< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a `LinOp` representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following files:

- `ginkgo/core/base/lin_op.hpp`
- `ginkgo/core/matrix/diagonal.hpp`

47.65 gko::DiagonalExtractable< ValueType > Class Template Reference

The diagonal of a LinOp implementing this interface can be extracted.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > extract\_diagonal\_linop ()` const override
Extracts the diagonal entries of the matrix into a vector.
- `virtual std::unique_ptr< matrix::Diagonal< ValueType > > extract\_diagonal ()` const =0
Extracts the diagonal entries of the matrix into a vector.

47.65.1 Detailed Description

```
template<typename ValueType>
class gko::DiagonalExtractable< ValueType >
```

The diagonal of a LinOp implementing this interface can be extracted.

`extract_diagonal` extracts the elements whose col and row index are the same and stores the result in a min(nrows, ncols) x 1 dense matrix.

47.65.2 Member Function Documentation

47.65.2.1 `extract_diagonal()`

```
template<typename ValueType >
virtual std::unique_ptr<matrix::Diagonal<ValueType> > gko::DiagonalExtractable< ValueType
>::extract_diagonal ( ) const [pure virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Ell< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), and [gko::matrix::Sellp< ValueType, IndexType >](#).

47.65.2.2 `extract_diagonal_linop()`

```
template<typename ValueType >
std::unique_ptr<LinOp> gko::DiagonalExtractable< ValueType >::extract_diagonal_linop ( )
const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Returns

linop the linop of diagonal format

Implements [gko::DiagonalLinOpExtractable](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/lin_op.hpp`

47.66 `gko::DiagonalLinOpExtractable` Class Reference

The diagonal of a LinOp can be extracted.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual `std::unique_ptr< LinOp > extract\_diagonal\_linop () const =0`
Extracts the diagonal entries of the matrix into a vector.

47.66.1 Detailed Description

The diagonal of a LinOp can be extracted.

It will be implemented by `DiagonalExtractable<ValueType>`, so the class does not need to implement it. `extract_diagonal_linop` returns a linop which extracts the elements whose col and row index are the same and stores the result in a `min(nrows, ncols) x 1` dense matrix.

47.66.2 Member Function Documentation

47.66.2.1 extract_diagonal_linop()

```
virtual std::unique_ptr<LinOp> gko::DiagonalLinOpExtractable::extract_diagonal_linop ( ) const
[pure virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Returns

linop the linop of diagonal format

Implemented in [gko::DiagonalExtractable< ValueType >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/lin_op.hpp

47.67 gko::dim< Dimensionality, DimensionType > Struct Template Reference

A type representing the dimensions of a multidimensional object.

```
#include <ginkgo/core/base/dim.hpp>
```

Public Member Functions

- constexpr [dim](#) ()
Creates a dimension object with all dimensions set to zero.
- constexpr [dim](#) (const dimension_type &size)
Creates a dimension object with all dimensions set to the same value.
- template<typename... Rest, std::enable_if_t< sizeof...(Rest)==Dimensionality - 1 > * = nullptr>
constexpr [dim](#) (const dimension_type &first, const Rest &... rest)
Creates a dimension object with the specified dimensions.
- constexpr const dimension_type & [operator\[\]](#) (const [size_type](#) &dimension) const noexcept
Returns the requested dimension.
- dimension_type & [operator\[\]](#) (const [size_type](#) &dimension) noexcept
- constexpr [operator bool](#) () const
Checks if all dimensions evaluate to true.

Friends

- constexpr friend bool [operator==](#) (const [dim](#) &x, const [dim](#) &y)
Checks if two dim objects are equal.
- constexpr friend bool [operator!=](#) (const [dim](#) &x, const [dim](#) &y)
Checks if two dim objects are not equal.
- constexpr friend [dim operator*](#) (const [dim](#) &x, const [dim](#) &y)
Multiplies two dim objects.
- std::ostream & [operator<<](#) (std::ostream &os, const [dim](#) &x)
A stream operator overload for dim.

47.67.1 Detailed Description

```
template<size_type Dimensionality, typename DimensionType = size_type>
struct gko::dim< Dimensionality, DimensionType >
```

A type representing the dimensions of a multidimensional object.

Template Parameters

<i>Dimensionality</i>	number of dimensions of the object
<i>DimensionType</i>	datatype used to represent each dimension

47.67.2 Constructor & Destructor Documentation

47.67.2.1 `dim()` [1/2]

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr gko::dim< Dimensionality, DimensionType >::dim (
    const dimension_type & size ) [inline], [explicit], [constexpr]
```

Creates a dimension object with all dimensions set to the same value.

Parameters

<i>size</i>	the size of each dimension
-------------	----------------------------

47.67.2.2 `dim()` [2/2]

```
template<size_type Dimensionality, typename DimensionType = size_type>
template<typename... Rest, std::enable_if_t< sizeof...(Rest)==Dimensionality - 1 > * = nullptr>
constexpr gko::dim< Dimensionality, DimensionType >::dim (
    const dimension_type & first,
    const Rest &... rest ) [inline], [constexpr]
```

Creates a dimension object with the specified dimensions.

If the number of dimensions given is less than the dimensionality of the object, the remaining dimensions are set to the same value as the last value given.

For example, in the context of matrices `dim<2>{2, 3}` creates the dimensions for a 2-by-3 matrix.

Parameters

<i>first</i>	first dimension
<i>rest</i>	other dimensions

47.67.3 Member Function Documentation

47.67.3.1 operator bool()

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr gko::dim< Dimensionality, DimensionType >::operator bool ( ) const [inline], [explicit],
[constexpr]
```

Checks if all dimensions evaluate to true.

For standard arithmetic types, this is equivalent to all dimensions being different than zero.

Returns

true if and only if all dimensions evaluate to true

Note

This operator is explicit to avoid implicit dim-to-int casts. It will still be used in contextual conversions (if, &&, ||, !)

47.67.3.2 operator[]() [1/2]

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr const dimension_type& gko::dim< Dimensionality, DimensionType >::operator[] (
    const size_type & dimension ) const [inline], [constexpr], [noexcept]
```

Returns the requested dimension.

For example, if d is a [dim<2>](#) object representing matrix dimensions, d[0] returns the number of rows, and d[1] returns the number of columns.

Parameters

<i>dimension</i>	the requested dimension
------------------	-------------------------

Returns

the dimension-th dimension

47.67.3.3 operator[]() [2/2]

```
template<size_type Dimensionality, typename DimensionType = size_type>
dimension_type& gko::dim< Dimensionality, DimensionType >::operator[] (
    const size_type & dimension ) [inline], [noexcept]
```

47.67.4 Friends And Related Function Documentation

47.67.4.1 operator"!="

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr friend bool operator!= (
    const dim< Dimensionality, DimensionType > & x,
    const dim< Dimensionality, DimensionType > & y ) [friend]
```

Checks if two dim objects are not equal.

Parameters

<i>x</i>	first object
<i>y</i>	second object

Returns

false if and only if all dimensions of both objects are equal.

47.67.4.2 operator*

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr friend dim operator* (
    const dim< Dimensionality, DimensionType > & x,
    const dim< Dimensionality, DimensionType > & y ) [friend]
```

Multiplies two dim objects.

Parameters

<i>x</i>	first object
<i>y</i>	second object

Returns

a dim object representing the size of the tensor product $x * y$

47.67.4.3 operator<<

```
template<size_type Dimensionality, typename DimensionType = size_type>
std::ostream& operator<< (
```

```
std::ostream & os,
const dim< Dimensionality, DimensionType > & x ) [friend]
```

A stream operator overload for dim.

Parameters

<code>os</code>	stream object
<code>x</code>	dim object

Returns

a stream object appended with the dim output

47.67.4.4 operator==

```
template<size_type Dimensionality, typename DimensionType = size_type>
constexpr friend bool operator== (
    const dim< Dimensionality, DimensionType > & x,
    const dim< Dimensionality, DimensionType > & y ) [friend]
```

Checks if two dim objects are equal.

Parameters

<code>x</code>	first object
<code>y</code>	second object

Returns

true if and only if all dimensions of both objects are equal.

The documentation for this struct was generated from the following file:

- ginkgo/core/base/dim.hpp

47.68 gko::DimensionMismatch Class Reference

[DimensionMismatch](#) is thrown if an operation is being applied to LinOps of incompatible size.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [DimensionMismatch](#) (const std::string &file, int line, const std::string &func, const std::string &first_name, [size_type](#) first_rows, [size_type](#) first_cols, const std::string &second_name, [size_type](#) second_rows, [size_type](#) second_cols, const std::string &clarification)
Initializes a dimension mismatch error.

47.68.1 Detailed Description

[DimensionMismatch](#) is thrown if an operation is being applied to LinOps of incompatible size.

47.68.2 Constructor & Destructor Documentation

47.68.2.1 DimensionMismatch()

```
gko::DimensionMismatch::DimensionMismatch (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & first_name,
    size_type first_rows,
    size_type first_cols,
    const std::string & second_name,
    size_type second_rows,
    size_type second_cols,
    const std::string & clarification ) [inline]
```

Initializes a dimension mismatch error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The function name where the error occurred
<i>first_name</i>	The name of the first operator
<i>first_rows</i>	The output dimension of the first operator
<i>first_cols</i>	The input dimension of the first operator
<i>second_name</i>	The name of the second operator
<i>second_rows</i>	The output dimension of the second operator
<i>second_cols</i>	The input dimension of the second operator
<i>clarification</i>	An additional message describing the error further

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.69 gko::experimental::solver::Direct< ValueType, IndexType > Class Template Reference

A direct solver based on a factorization into lower and upper triangular factors (with an optional diagonal scaling).

```
#include <ginkgo/core/solver/direct.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `Direct (const Direct &)`
Creates a copy of the solver.
- `Direct (Direct &&)`
Moves from the given solver, leaving it empty.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType, IndexType >())`
Create the parameters from the property_tree.

47.69.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::experimental::solver::Direct< ValueType, IndexType >
```

A direct solver based on a factorization into lower and upper triangular factors (with an optional diagonal scaling).

The solver is built from the Factorization returned by the provided [LinOpFactory](#).

Template Parameters

<i>ValueType</i>	the type used to store values of the system matrix
<i>IndexType</i>	the type used to store sparsity pattern indices of the system matrix

47.69.2 Member Function Documentation

47.69.2.1 conj_transpose()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::experimental::solver::Direct< ValueType, IndexType >::conj_↵
transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.69.2.2 parse()

```
template<typename ValueType , typename IndexType >
static parameters_type gko::experimental::solver::Direct< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type\_descriptor & td_for_child = config::make\_type\_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

47.69.2.3 transpose()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::experimental::solver::Direct< ValueType, IndexType >::transpose (
) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/solver/direct.hpp](#)

47.70 gko::experimental::distributed::DistributedBase Class Reference

A base class for distributed objects.

```
#include <ginkgo/core/distributed/base.hpp>
```


Public Member Functions

- [DistributedBase & operator=](#) (const [DistributedBase](#) &)
Copy assignment that doesn't change the used [mpi::communicator](#).
- [DistributedBase & operator=](#) ([DistributedBase](#) &&) noexcept
Move assignment that doesn't change the used [mpi::communicator](#).
- [mpi::communicator get_communicator](#) () const
Access the used [mpi::communicator](#).

47.70.1 Detailed Description

A base class for distributed objects.

This class stores and gives access to the used [mpi::communicator](#) object.

Note

The communicator is not changed on assignment.

47.70.2 Member Function Documentation

47.70.2.1 [get_communicator\(\)](#)

```
mpi::communicator gko::experimental::distributed::DistributedBase::get_communicator ( ) const
[inline]
```

Access the used [mpi::communicator](#).

Returns

used [mpi::communicator](#)

```
56 { return comm_; }
```

47.70.2.2 [operator=\(\)](#) [1/2]

```
DistributedBase& gko::experimental::distributed::DistributedBase::operator= (
    const DistributedBase & ) [inline]
```

Copy assignment that doesn't change the used [mpi::communicator](#).

Returns

unmodified *this

47.70.2.3 operator=() [2/2]

```
DistributedBase& gko::experimental::distributed::DistributedBase::operator= (
    DistributedBase && ) [inline], [noexcept]
```

Move assignment that doesn't change the used `mpi::communicator`.

Returns

unmodified *this

The documentation for this class was generated from the following file:

- ginkgo/core/distributed/base.hpp

47.71 gko::DpcppExecutor Class Reference

This is the `Executor` subclass which represents a DPC++ enhanced device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- `std::shared_ptr< Executor > get_master ()` noexcept override
Returns the master `OmpExecutor` of this `Executor`.
- `std::shared_ptr< const Executor > get_master ()` const noexcept override
Returns the master `OmpExecutor` of this `Executor`.
- `void synchronize ()` const override
Synchronize the operations launched on the executor with its master.
- `std::string get_description ()` const override
- `int get_device_id ()` const noexcept
Get the DPCPP device id of the device associated to this executor.
- `const std::vector< int > & get_subgroup_sizes ()` const noexcept
Get the available subgroup sizes for this device.
- `int get_num_computing_units ()` const noexcept
Get the number of Computing Units of this executor.
- `int get_num_subgroups ()` const noexcept
Get the number of subgroups of this executor.
- `const std::vector< int > & get_max_workitem_sizes ()` const noexcept
Get the maximum work item sizes.
- `int get_max_workgroup_size ()` const noexcept
Get the maximum workgroup size.
- `int get_max_subgroup_size ()` const noexcept
Get the maximum subgroup size.
- `std::string get_device_type ()` const noexcept
Get a string representing the device type.
- `virtual void run (const Operation &op)` const=0
Runs the specified `Operation` using this `Executor`.
- `template<typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename ClosureDpcpp >`
`void run (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp)` const
Runs one of the passed in functors, depending on the `Executor` type.
- `template<typename ClosureReference, typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename Closure↵`
`Dpcpp >`
`void run (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↵`
`Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp)` const
Runs one of the passed in functors, depending on the `Executor` type.

Static Public Member Functions

- static std::shared_ptr< [DpcppExecutor](#) > [create](#) (int device_id, std::shared_ptr< [Executor](#) > master, std::string device_type="all", dpcpp_queue_property property=dpcpp_queue_property::in_order)
Creates a new [DpcppExecutor](#).
- static int [get_num_devices](#) (std::string device_type)
Get the number of devices present on the system.

47.71.1 Detailed Description

This is the [Executor](#) subclass which represents a DPC++ enhanced device.

47.71.2 Member Function Documentation

47.71.2.1 [create\(\)](#)

```
static std::shared_ptr<DpcppExecutor> gko::DpcppExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    std::string device_type = "all",
    dpcpp_queue_property property = dpcpp_queue_property::in_order ) [static]
```

Creates a new [DpcppExecutor](#).

Parameters

<i>device_id</i>	the DPCPP device id of this device
<i>master</i>	an executor on the host that is used to invoke the device kernels
<i>device_type</i>	a string representing the type of device to consider (accelerator, cpu, gpu or all).

47.71.2.2 [get_description\(\)](#)

```
std::string gko::DpcppExecutor::get_description ( ) const [override], [virtual]
```

Returns

a textual representation of the executor and its device.

Implements [gko::Executor](#).

47.71.2.3 `get_device_id()`

```
int gko::DpcppExecutor::get_device_id ( ) const [inline], [noexcept]
```

Get the DPCPP device id of the device associated to this executor.

Returns

the DPCPP device id of the device associated to this executor

47.71.2.4 `get_device_type()`

```
std::string gko::DpcppExecutor::get_device_type ( ) const [inline], [noexcept]
```

Get a string representing the device type.

Returns

a string representing the device type

47.71.2.5 `get_master()` [1/2]

```
std::shared_ptr<const Executor> gko::DpcppExecutor::get_master ( ) const [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.71.2.6 `get_master()` [2/2]

```
std::shared_ptr<Executor> gko::DpcppExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.71.2.7 get_max_subgroup_size()

```
int gko::DpcppExecutor::get_max_subgroup_size ( ) const [inline], [noexcept]
```

Get the maximum subgroup size.

Returns

the maximum subgroup size

47.71.2.8 get_max_workgroup_size()

```
int gko::DpcppExecutor::get_max_workgroup_size ( ) const [inline], [noexcept]
```

Get the maximum workgroup size.

Returns

the maximum workgroup size

47.71.2.9 get_max_workitem_sizes()

```
const std::vector<int>& gko::DpcppExecutor::get_max_workitem_sizes ( ) const [inline], [noexcept]
```

Get the maximum work item sizes.

Returns

the maximum work item sizes

47.71.2.10 get_num_computing_units()

```
int gko::DpcppExecutor::get_num_computing_units ( ) const [inline], [noexcept]
```

Get the number of Computing Units of this executor.

Returns

the number of Computing Units of this executor

47.71.2.11 get_num_devices()

```
static int gko::DpcppExecutor::get_num_devices (
    std::string device_type ) [static]
```

Get the number of devices present on the system.

Parameters

<i>device_type</i>	a string representing the device type
--------------------	---------------------------------------

Returns

the number of devices present on the system

47.71.2.12 get_subgroup_sizes()

```
const std::vector<int>& gko::DpcppExecutor::get_subgroup_sizes ( ) const [inline], [noexcept]
```

Get the available subgroup sizes for this device.

Returns

the available subgroup sizes for this device

47.71.2.13 run() [1/3]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

47.71.2.14 run() [2/3]

```
virtual void gko::Executor::run
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

47.71.2.15 run() [3/3]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureReference ,
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of op_ref
<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.72 gko::DpcppTimer Class Reference

A timer using kernels for timing on a [DpcppExecutor](#) in profiling mode.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void [record](#) ([time_point](#) &time) override
Records a time point at the current time.
- void [wait](#) ([time_point](#) &time) override
Waits until all kernels in-process when recording the time point are finished.
- std::chrono::nanoseconds [difference_async](#) (const [time_point](#) &start, const [time_point](#) &stop) override
Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

47.72.1 Detailed Description

A timer using kernels for timing on a [DpcppExecutor](#) in profiling mode.

47.72.2 Member Function Documentation

47.72.2.1 difference_async()

```
std::chrono::nanoseconds gko::DpcppTimer::difference_async (
    const time\_point & start,
    const time\_point & stop ) [override], [virtual]
```

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. `timer->wait(stop)` needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements [gko::Timer](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/timer.hpp

47.73 gko::batch::matrix::Ell< ValueType, IndexType > Class Template Reference

[Ell](#) is a sparse matrix format that stores the same number of nonzeros in each row, enabling coalesced accesses.

```
#include <ginkgo/core/matrix/batch_ell.hpp>
```

Public Member Functions

- `std::unique_ptr< unbatch_type > create_view_for_item (size_type item_id)`
Creates a mutable view (of [matrix::Ell](#) type) of one item of the [batch::matrix::Ell<value_type>](#) object.
- `std::unique_ptr< const unbatch_type > create_const_view_for_item (size_type item_id) const`
Creates a mutable view (of [matrix::Ell](#) type) of one item of the [batch::matrix::Ell<value_type>](#) object.
- `value_type * get_values () noexcept`
Returns a pointer to the array of values of the matrix.
- `const value_type * get_const_values () const noexcept`
Returns a pointer to the array of values of the matrix.
- `index_type * get_col_idxs () noexcept`
Returns a pointer to the array of column indices of the matrix.
- `const index_type * get_const_col_idxs () const noexcept`
Returns a pointer to the array of column indices of the matrix.
- `index_type get_num_stored_elements_per_row () const noexcept`
Returns the number of elements per row explicitly stored.
- `size_type get_num_stored_elements () const noexcept`
Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.
- `size_type get_num_elements_per_item () const noexcept`
Returns the number of stored elements in each batch item.
- `index_type * get_col_idxs_for_item (size_type batch_id) noexcept`
Returns a pointer to the array of col_idxs of the matrix.
- `const index_type * get_const_col_idxs_for_item (size_type batch_id) const noexcept`
Returns a pointer to the array of col_idxs of the matrix.
- `value_type * get_values_for_item (size_type batch_id) noexcept`
Returns a pointer to the array of values of the matrix for a specific batch item.
- `const value_type * get_const_values_for_item (size_type batch_id) const noexcept`
Returns a pointer to the array of values of the matrix for a specific batch item.
- `Ell * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector.
- `Ell * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector with a linear combination of the given input vector.
- `const Ell * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x) const`

- `const Ell * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x) const`
- `void scale (const array< value_type > &row_scale, const array< value_type > &col_scale)`
Performs in-place row and column scaling for this matrix.
- `void add_scaled_identity (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> beta)`
*Performs the operation $this = \alpha * I + \beta * this$.*

Static Public Member Functions

- `static std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size=batch_dim< 2 >{}, const IndexType num_elems_per_row=0)`
Creates an uninitialized Ell matrix of the specified size.
- `static std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, const IndexType num_elems_per_row, array< value_type > values, array< index_type > col_idxs)`
Creates a Ell matrix from an already allocated (and initialized) array.
- `template<typename InputValueType, typename ColIndexType >`
`static std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, const IndexType num_elems_per_row, std::initializer_list< InputValueType > values, std::initializer_list< ColIndexType > col_idxs)`
create(std::shared_ptr<const Executor>,
- `static std::unique_ptr< const Ell > create_const (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &sizes, const index_type num_elems_per_row, gko::detail::const_array_view< value_type > &&values, gko::detail::const_array_view< index_type > &&col_idxs)`
Creates a constant (immutable) batch ell matrix from a constant array.

47.73.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::batch::matrix::Ell< ValueType, IndexType >
```

Ell is a sparse matrix format that stores the same number of nonzeros in each row, enabling coalesced accesses.

It is suitable for sparsity patterns that have a similar number of nonzeros in every row. The values are stored in a column-major fashion similar to the monolithic `gko::matrix::Ell` class.

Similar to the monolithic `gko::matrix::Ell` class, `invalid_index<IndexType>` is used as the column index for padded zero entries.

Note

It is also assumed that the sparsity pattern of all the items in the batch is the same and therefore only a single copy of the sparsity pattern is stored.

Currently only `IndexType` of `int32` is supported.

Template Parameters

<i>ValueType</i>	value precision of matrix elements
<i>IndexType</i>	index precision of matrix elements

47.73.2 Member Function Documentation

47.73.2.1 add_scaled_identity()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::batch::matrix::Ell< ValueType, IndexType >::add_scaled_identity (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> beta )
```

Performs the operation $\text{this} = \alpha * \text{I} + \beta * \text{this}$.

Parameters

<i>alpha</i>	the scalar for identity
<i>beta</i>	the scalar to multiply this matrix

Note

Performs the operation in-place for this batch matrix.

This operation fails in case this matrix does not have all its diagonal entries.

47.73.2.2 apply() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Ell* gko::batch::matrix::Ell< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector with a linear combination of the given input vector.

Represents the matrix vector multiplication, $x = \alpha * A$

- $b + \beta * x$, where x and b are both multi-vectors.

Parameters

<i>alpha</i>	the scalar to scale the matrix-vector product with
<i>b</i>	the multi-vector to be applied to
<i>beta</i>	the scalar to scale the x vector with
<i>x</i>	the output multi-vector

47.73.2.3 apply() [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const Ell* gko::batch::matrix::Ell< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x ) const
```

47.73.2.4 apply() [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Ell* gko::batch::matrix::Ell< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector.

Represents the matrix vector multiplication, $x = A * b$, where x and b are both multi-vectors.

Parameters

<i>b</i>	the multi-vector to be applied to
<i>x</i>	the output multi-vector

47.73.2.5 apply() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const Ell* gko::batch::matrix::Ell< ValueType, IndexType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x ) const
```

47.73.2.6 create() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Ell> gko::batch::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    const IndexType num_elems_per_row,
    array< value_type > values,
    array< index_type > col_idxs ) [static]
```

Creates a [Ell](#) matrix from an already allocated (and initialized) array.

The column indices array needs to be the same for all batch items.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_elems_per_row</i>	the number of elements to be stored in each row
<i>values</i>	array of matrix values
<i>col_idx</i> s	the <i>col_idx</i> s array of a single batch item of the matrix.

Note

If *values* is not an rvalue, not an array of *ValueType*, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.73.2.7 create() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType , typename ColIndexType >
static std::unique_ptr<Ell> gko::batch::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    const IndexType num_elems_per_row,
    std::initializer_list< InputValueType > values,
    std::initializer_list< ColIndexType > col_idxes ) [inline], [static]

create(std::shared_ptr<const Executor>,

create(std::shared_ptr<const Executor>, const batch_dim<2>&, const IndexType, array<value_type>,
array<index_type>)
292     {
293         return create(exec, size, num_elems_per_row,
294             array<value_type>{exec, std::move(values)},
295             array<index_type>{exec, std::move(col_idxes)});
296     }
```

References `gko::batch::matrix::Ell< ValueType, IndexType >::create()`.

47.73.2.8 create() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Ell> gko::batch::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size = batch_dim< 2 >{},
    const IndexType num_elems_per_row = 0 ) [static]
```

Creates an uninitialized [Ell](#) matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_elems_per_row</i>	the number of elements to be stored in each row

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::batch::matrix::Ell< ValueType, IndexType >::create()`.

47.73.2.9 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Ell> gko::batch::matrix::Ell< ValueType, IndexType >::create_↵
const (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & sizes,
    const index_type num_elems_per_row,
    gko::detail::const_array_view< value_type > && values,
    gko::detail::const_array_view< index_type > && col_idxs ) [static]
```

Creates a constant (immutable) batch ell matrix from a constant array.

The column indices array needs to be the same for all batch items.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>num_elems_per_row</i>	the number of elements to be stored in each row
<i>values</i>	the value array of the matrix
<i>col_idxs</i>	the col_idxs array of a single batch item of the matrix.

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

A smart pointer to the newly created matrix.

47.73.2.10 create_const_view_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<const unbatch_type> gko::batch::matrix::Ell< ValueType, IndexType >::create_↵
_const_view_for_item (
    size_type item_id ) const
```

Creates a mutable view (of [matrix::Ell](#) type) of one item of the `batch::matrix::Ell<value_type>` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<code>item↔ _id</code>	The index of the batch item
----------------------------	-----------------------------

Returns

a [batch::matrix::Ell](#) object with the data from the batch item at the given index.

47.73.2.11 create_view_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<unbatch_type> gko::batch::matrix::Ell< ValueType, IndexType >::create_view_↔
for_item (
    size_type item_id )
```

Creates a mutable view (of [matrix::Ell](#) type) of one item of the `batch::matrix::Ell<value_type>` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<code>item↔ _id</code>	The index of the batch item
----------------------------	-----------------------------

Returns

a [batch::matrix::Ell](#) object with the data from the batch item at the given index.

47.73.2.12 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_col_idxs ( ) [inline], [noexcept]
```

Returns a pointer to the array of column indices of the matrix.

Returns

the pointer to the array of column indices

References `gko::array< ValueType >::get_data()`.

47.73.2.13 get_col_idxes_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_col_idxes_for_item (
    size_type batch_id ) [inline], [noexcept]
```

Returns a pointer to the array of col_idxes of the matrix.

This is shared across all batch items.

Parameters

<i>batch_id</i>	the id of the batch item.
-----------------	---------------------------

Returns

the pointer to the array of col_idxes

References gko::array< ValueType >::get_data().

47.73.2.14 get_const_col_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_const_col_idxes ( )
const [inline], [noexcept]
```

Returns a pointer to the array of column indices of the matrix.

Returns

the pointer to the array of column indices

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.73.2.15 get_const_col_idxes_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_const_col_idxes_for_item
(
    size_type batch_id ) const [inline], [noexcept]
```

Returns a pointer to the array of col_idxes of the matrix.

This is shared across all batch items.

Parameters

<i>batch</i> ↔ <i>_id</i>	the id of the batch item.
------------------------------	---------------------------

Returns

the pointer to the array of col_idxes

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.73.2.16 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_const_values ( ) const
[inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.73.2.17 get_const_values_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_const_values_for_item (
    size_type batch_id ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch</i> ↔ <i>_id</i>	the id of the batch item.
------------------------------	---------------------------

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`, and `gko::batch::matrix::Ell< ValueType, IndexType >↔::get_num_elements_per_item()`.

47.73.2.18 get_num_elements_per_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::matrix::Ell< ValueType, IndexType >::get_num_elements_per_item ( ) const
[inline], [noexcept]
```

Returns the number of stored elements in each batch item.

Returns

the number of stored elements per batch item.

References `gko::batch::matrix::Ell< ValueType, IndexType >::get_num_stored_elements()`.

Referenced by `gko::batch::matrix::Ell< ValueType, IndexType >::get_const_values_for_item()`, and `gko::batch↔::matrix::Ell< ValueType, IndexType >::get_values_for_item()`.

47.73.2.19 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::matrix::Ell< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.

Returns

the number of elements explicitly stored in the vector, cumulative across all the batch items

References `gko::array< ValueType >::get_size()`.

Referenced by `gko::batch::matrix::Ell< ValueType, IndexType >::get_num_elements_per_item()`.

47.73.2.20 get_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::batch::matrix::Ell< ValueType, IndexType >::get_num_stored_elements_per_row (
) const [inline], [noexcept]
```

Returns the number of elements per row explicitly stored.

Returns

the number of elements stored in each row of the ELL matrix. Same for each batch item

47.73.2.21 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix.

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data().

47.73.2.22 get_values_for_item()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::batch::matrix::Ell< ValueType, IndexType >::get_values_for_item (
    size_type batch_id ) [inline], [noexcept]
```

Returns a pointer to the array of values of the matrix for a specific batch item.

Parameters

<i>batch_id</i>	the id of the batch item.
-----------------	---------------------------

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data(), and gko::batch::matrix::Ell< ValueType, IndexType >::get_num_elements_per_item().

47.73.2.23 scale()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::batch::matrix::Ell< ValueType, IndexType >::scale (
    const array< value_type > & row_scale,
    const array< value_type > & col_scale )
```

Performs in-place row and column scaling for this matrix.

Parameters

<i>row_scale</i>	the row scalars
<i>col_scale</i>	the column scalars

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/batch_ell.hpp

47.74 gko::matrix::Ell< ValueType, IndexType > Class Template Reference

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

```
#include <ginkgo/core/matrix/ell.hpp>
```

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< [Diagonal](#)< ValueType > > [extract_diagonal](#) () const override
Extracts the diagonal entries of the matrix into a vector.
- std::unique_ptr< absolute_type > [compute_absolute](#) () const override
Gets the AbsoluteLinOp.
- void [compute_absolute_inplace](#) () override
Compute absolute inplace on each element.
- value_type * [get_values](#) () noexcept
Returns the values of the matrix.
- const value_type * [get_const_values](#) () const noexcept
Returns the values of the matrix.
- index_type * [get_col_idxs](#) () noexcept
Returns the column indexes of the matrix.

- `const index_type * get_const_col_idxxs ()` const noexcept
Returns the column indexes of the matrix.
- `size_type get_num_stored_elements_per_row ()` const noexcept
Returns the number of stored elements per row.
- `size_type get_stride ()` const noexcept
Returns the stride of the matrix.
- `size_type get_num_stored_elements ()` const noexcept
Returns the number of elements explicitly stored in the matrix.
- `value_type & val_at (size_type row, size_type idx)` noexcept
*Returns the *idx*-th non-zero element of the *row*-th row .*
- `value_type val_at (size_type row, size_type idx)` const noexcept
*Returns the *idx*-th non-zero element of the *row*-th row .*
- `index_type & col_at (size_type row, size_type idx)` noexcept
*Returns the *idx*-th column index of the *row*-th row .*
- `index_type col_at (size_type row, size_type idx)` const noexcept
*Returns the *idx*-th column index of the *row*-th row .*
- `Ell & operator= (const Ell &)`
*Copy-assigns an *Ell* matrix.*
- `Ell & operator= (Ell &&)`
*Move-assigns an *Ell* matrix.*
- `Ell (const Ell &)`
*Copy-constructs an *Ell* matrix.*
- `Ell (Ell &&)`
*Move-constructs an *Ell* matrix.*

Static Public Member Functions

- static `std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size={}, size_type num_stored_elements_per_row=0, size_type stride=0)`
*Creates an uninitialized *Ell* matrix of the specified size.*
- static `std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, array< value_type > values, array< index_type > col_idxxs, size_type num_stored_elements_per_row, size_type stride)`
*Creates an *ELL* matrix from already allocated (and initialized) column index and value arrays.*
- template<typename InputValueType, typename InputColumnIndexType >
static `std::unique_ptr< Ell > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, std::initializer_list< InputValueType > values, std::initializer_list< InputColumnIndexType > col_idxxs, size_type num_stored_elements_per_row, size_type stride)`
create(std::shared_ptr<const Executor>,
- static `std::unique_ptr< const Ell > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&col_idxxs, size_type num_stored_elements_per_row, size_type stride)`
*Creates a constant (immutable) *Ell* matrix from a set of constant arrays.*

47.74.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Ell< ValueType, IndexType >
```

ELL is a matrix format where stride with explicit zeros is used such that all rows have the same number of stored elements.

The number of elements stored in each row is the largest number of nonzero elements in any of the rows (obtainable through [get_num_stored_elements_per_row\(\)](#) method). This removes the need of a row pointer like in the CSR format, and allows for SIMD processing of the distinct rows. For efficient processing, the nonzero elements and the corresponding column indices are stored in column-major fashion. The columns are padded to the length by user-defined stride parameter whose default value is the number of rows of the matrix.

This implementation uses the column index value [invalid_index<IndexType>\(\)](#) to mark padding entries that are not part of the sparsity pattern.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.74.2 Constructor & Destructor Documentation

47.74.2.1 Ell() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Ell< ValueType, IndexType >::Ell (
    const Ell< ValueType, IndexType > & )
```

Copy-constructs an [Ell](#) matrix.

Inherits executor and dimensions, but copies data without padding.

47.74.2.2 Ell() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Ell< ValueType, IndexType >::Ell (
    Ell< ValueType, IndexType > && )
```

Move-constructs an [Ell](#) matrix.

Inherits executor, dimensions and data with padding. The moved-from object is empty (0x0 with empty Array).

47.74.3 Member Function Documentation

47.74.3.1 col_at() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Ell< ValueType, IndexType >::col_at (
    size_type row,
    size_type idx ) const [inline], [noexcept]
```

Returns the `idx`-th column index of the `row`-th row .

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <code>idx</code> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

```
239 {
240     return this->get_const_col_idxs()[this->linearize_index(row, idx)];
241 }
```

References `gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs()`.

47.74.3.2 col_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type& gko::matrix::Ell< ValueType, IndexType >::col_at (
    size_type row,
    size_type idx ) [inline], [noexcept]
```

Returns the `idx`-th column index of the `row`-th row .

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <code>idx</code> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References `gko::matrix::Ell< ValueType, IndexType >::get_col_idxs()`.

47.74.3.3 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Ell< ValueType, IndexType >::compute_absolute ( )
const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Ell< ValueType, IndexType > > >](#).

47.74.3.4 create() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Ell> gko::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< value_type > values,
    array< index_type > col_idxs,
    size_type num_stored_elements_per_row,
    size_type stride ) [static]
```

Creates an ELL matrix from already allocated (and initialized) column index and value arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>values</i>	array of matrix values
<i>col_idxs</i>	array of column indexes
<i>num_stored_elements_per_row</i>	the number of stored elements per row
<i>stride</i>	stride of the rows

Note

If one of `col_idxs` or `values` is not an rvalue, not an array of `IndexType` and `ValueType`, respectively, or is on the wrong executor, an internal copy of that array will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.74.3.5 create() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType, typename InputColumnIndexType >
static std::unique_ptr<Ell> gko::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::initializer_list< InputValueType > values,
    std::initializer_list< InputColumnIndexType > col_idxes,
    size_type num_stored_elements_per_row,
    size_type stride ) [inline], [static]
```

create(std::shared_ptr<const Executor>,

create(std::shared_ptr<const Executor>, const dim<2>&, array<value_type>, array<index_type>, size_type, size_type)

References gko::matrix::Ell< ValueType, IndexType >::create().

47.74.3.6 create() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Ell> gko::matrix::Ell< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = {},
    size_type num_stored_elements_per_row = 0,
    size_type stride = 0 ) [static]
```

Creates an uninitialized [Ell](#) matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_stored_elements_per_row</i>	the number of stored elements per row
<i>stride</i>	stride of the columns. If it is set to 0, size[0] will be used instead.

Returns

A smart pointer to the newly created matrix.

Referenced by gko::matrix::Ell< ValueType, IndexType >::create().

47.74.3.7 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Ell> gko::matrix::Ell< ValueType, IndexType >::create_const (
```

```

std::shared_ptr< const Executor > exec,
const dim< 2 > & size,
gko::detail::const_array_view< ValueType > && values,
gko::detail::const_array_view< IndexType > && col_idx,
size_type num_stored_elements_per_row,
size_type stride ) [static]

```

Creates a constant (immutable) [Ell](#) matrix from a set of constant arrays.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>col_idx</i>	the column index array of the matrix
<i>num_stored_elements_per_row</i>	the number of stored nonzeros per row
<i>stride</i>	the column-stride of the value and column index array

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

47.74.3.8 [extract_diagonal\(\)](#)

```

template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Ell< ValueType, IndexType >::extract_↵
diagonal ( ) const [override], [virtual]

```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.74.3.9 [get_col_idxes\(\)](#)

```

template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Ell< ValueType, IndexType >::get_col_idxes ( ) [inline], [noexcept]

```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References gko::array< ValueType >::get_data().

Referenced by gko::matrix::Ell< ValueType, IndexType >::col_at().

47.74.3.10 get_const_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Ell< ValueType, IndexType >::get_const_col_idxs ( ) const [inline],
[noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

Referenced by gko::matrix::Ell< ValueType, IndexType >::col_at().

47.74.3.11 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Ell< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.74.3.12 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_num_stored_elements ( ) const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.74.3.13 get_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_num_stored_elements_per_row ( ) const [inline], [noexcept]
```

Returns the number of stored elements per row.

Returns

the number of stored elements per row.

47.74.3.14 get_stride()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Ell< ValueType, IndexType >::get_stride ( ) const [inline], [noexcept]
```

Returns the stride of the matrix.

Returns

the stride of the matrix.

47.74.3.15 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Ell< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References `gko::array< ValueType >::get_data()`.

47.74.3.16 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Ell& gko::matrix::Ell< ValueType, IndexType >::operator= (
    const Ell< ValueType, IndexType > & )
```

Copy-assigns an [Ell](#) matrix.

Preserves the executor, reallocates the matrix with minimal stride if the dimensions don't match, then copies the data over, ignoring padding.

47.74.3.17 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Ell& gko::matrix::Ell< ValueType, IndexType >::operator= (
    Ell< ValueType, IndexType > && )
```

Move-assigns an [Ell](#) matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with empty Array).

47.74.3.18 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Ell< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.74.3.19 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Ell< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.74.3.20 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Ell< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.74.3.21 val_at() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type gko::matrix::Ell< ValueType, IndexType >::val_at (
    size_type row,
    size_type idx ) const [inline], [noexcept]
```

Returns the *idx*-th non-zero element of the *row*-th row .

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <i>idx</i> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References [gko::array< ValueType >::get_const_data\(\)](#).

47.74.3.22 val_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type& gko::matrix::Ell< ValueType, IndexType >::val_at (
    size_type row,
    size_type idx ) [inline], [noexcept]
```

Returns the *idx*-th non-zero element of the *row*-th row .

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the idx-th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

References `gko::array< ValueType >::get_data()`.

47.74.3.23 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Ell< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/csr.hpp
- ginkgo/core/matrix/ell.hpp

47.75 gko::enable_parameters_type< ConcreteParametersType, Factory > Class Template Reference

The [enable_parameters_type](#) mixin is used to create a base implementation of the factory parameters structure.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

- `template<typename... Args>`
`ConcreteParametersType & with_loggers (Args &&... _value)`
Provides the loggers to be added to the factory and its generated objects in a fluent interface.
- `std::unique_ptr< Factory > on (std::shared_ptr< const Executor > exec) const`
Creates a new factory on the specified executor.

47.75.1 Detailed Description

```
template<typename ConcreteParametersType, typename Factory>
class gko::enable_parameters_type< ConcreteParametersType, Factory >
```

The [enable_parameters_type](#) mixin is used to create a base implementation of the factory parameters structure.

It provides only the [on\(\)](#) method which can be used to instantiate the factory give the parameters stored in the structure.

Template Parameters

<i>ConcreteParametersType</i>	the concrete parameters type which is being implemented [CRTP parameter]
<i>Factory</i>	the concrete factory for which these parameters are being used

47.75.2 Member Function Documentation

47.75.2.1 on()

```
template<typename ConcreteParametersType, typename Factory>
std::unique_ptr<Factory> gko::enable\_parameters\_type< ConcreteParametersType, Factory >::on (
    std::shared_ptr< const Executor > exec ) const [inline]
```

Creates a new factory on the specified executor.

Parameters

<i>exec</i>	the executor where the factory will be created
-------------	--

Returns

a new factory instance

The documentation for this class was generated from the following file:

- `ginkgo/core/base/abstract_factory.hpp`

47.76 [gko::EnableAbsoluteComputation](#)< [AbsoluteLinOp](#) > Class Template Reference

The [EnableAbsoluteComputation](#) mixin provides the default implementations of `compute_absolute_linop` and the absolute interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```


Public Member Functions

- `std::unique_ptr< LinOp > compute_absolute_linop ()` const override
Gets the absolute LinOp.
- `virtual std::unique_ptr< absolute_type > compute_absolute ()` const =0
Gets the AbsoluteLinOp.

47.76.1 Detailed Description

```
template<typename AbsoluteLinOp>
class gko::EnableAbsoluteComputation< AbsoluteLinOp >
```

The [EnableAbsoluteComputation](#) mixin provides the default implementations of `compute_absolute_linop` and the absolute interface.

`compute_absolute` gets a new `AbsoluteLinOp`. `compute_absolute_inplace` applies absolute inplace, so it still keeps the `value_type` of the class.

Template Parameters

<i>AbsoluteLinOp</i>	the absolute LinOp which is being returned [CRTP parameter]
----------------------	---

47.76.2 Member Function Documentation

47.76.2.1 compute_absolute()

```
template<typename AbsoluteLinOp>
virtual std::unique_ptr<absolute_type> gko::EnableAbsoluteComputation< AbsoluteLinOp >↵
::compute_absolute ( ) const [pure virtual]
```

Gets the `AbsoluteLinOp`.

Returns

a pointer to the new absolute object

Implemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::experimental::distributed::Vector< gko::matrix::Eil< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), [gko::matrix::Sellp< ValueType, IndexType >](#), and [gko::matrix::Diagonal< ValueType >](#).

Referenced by [gko::EnableAbsoluteComputation< remove_complex< Eil< ValueType, IndexType > > >↵](#) `::compute_absolute_linop()`.

47.76.2.2 compute_absolute_linop()

```
template<typename AbsoluteLinOp>
std::unique_ptr<LinOp> gko::EnableAbsoluteComputation< AbsoluteLinOp >::compute_absolute_↵
linop ( ) const [inline], [override], [virtual]
```

Gets the absolute LinOp.

Returns

a pointer to the new absolute LinOp

Implements [gko::AbsoluteComputable](#).

```
801 {
802     return this->compute_absolute();
803 }
```

The documentation for this class was generated from the following file:

- [ginkgo/core/base/lin_op.hpp](#)

47.77 gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase > Class Template Reference

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new abstract object.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

47.77.1 Detailed Description

```
template<typename AbstractObject, typename PolymorphicBase = PolymorphicObject>
class gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase >
```

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new abstract object.

It uses method hiding to update the parameter and return types from [PolymorphicObject](#) to `AbstractObject` wherever it makes sense. As opposed to [EnablePolymorphicObject](#), it does not implement [PolymorphicObject](#)'s virtual methods.

Template Parameters

<i>AbstractObject</i>	the abstract class which is being implemented [CRTP parameter]
<i>PolymorphicBase</i>	parent of AbstractObject in the polymorphic hierarchy, has to be a subclass of polymorphic object

See also

[EnablePolymorphicObject](#) for creating a concrete subclass of [PolymorphicObject](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/polymorphic_object.hpp

47.78 gko::solver::EnableApplyWithInitialGuess< DerivedType > Class Template Reference

[EnableApplyWithInitialGuess](#) providing default operation for [ApplyWithInitialGuess](#) with correct validation and log.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

47.78.1 Detailed Description

```
template<typename DerivedType>
class gko::solver::EnableApplyWithInitialGuess< DerivedType >
```

[EnableApplyWithInitialGuess](#) providing default operation for [ApplyWithInitialGuess](#) with correct validation and log.

It ensures that vectors of `apply_with_initial_guess` will always have the same executor as the object this mixin is used in, creating a clone on the correct executor if necessary.

Template Parameters

<i>DerivedType</i>	The type that this Mixin is used in. It must provide <code>get_size()</code> and <code>get_executor()</code> functions that return correctly initialized values and the logger functionality.
--------------------	---

The documentation for this class was generated from the following file:

- ginkgo/core/solver/solver_base.hpp

47.79 gko::batch::EnableBatchLinOp< ConcreteBatchLinOp, PolymorphicBase > Class Template Reference

The [EnableBatchLinOp](#) mixin can be used to provide sensible default implementations of the majority of the [BatchLinOp](#) and [PolymorphicObject](#) interface.

```
#include <ginkgo/core/base/batch_lin_op.hpp>
```

Additional Inherited Members

47.79.1 Detailed Description

```
template<typename ConcreteBatchLinOp, typename PolymorphicBase = BatchLinOp>
class gko::batch::EnableBatchLinOp< ConcreteBatchLinOp, PolymorphicBase >
```

The [EnableBatchLinOp](#) mixin can be used to provide sensible default implementations of the majority of the [BatchLinOp](#) and [PolymorphicObject](#) interface.

The goal of the mixin is to facilitate the development of new [BatchLinOp](#), by enabling the implementers to focus on the important parts of their operator, while the library takes care of generating the trivial utility functions. The mixin will provide default implementations for the entire [PolymorphicObject](#) interface, including a default implementation of `copy_from` between objects of the new [BatchLinOp](#) type.

Implementers of new [BatchLinOps](#) are required to specify only the following aspects:

1. Creation of the [BatchLinOp](#): This can be facilitated via either [EnableCreateMethod](#) mixin (used mostly for matrix formats), or `GKO_ENABLE_BATCH_LIN_OP_FACTORY` macro (used for operators created from other operators, like preconditioners and solvers).

Template Parameters

<i>ConcreteBatchLinOp</i>	the concrete BatchLinOp which is being implemented [CRTP parameter]
<i>PolymorphicBase</i>	parent of ConcreteBatchLinOp in the polymorphic hierarchy, has to be a subclass of BatchLinOp

The documentation for this class was generated from the following file:

- `ginkgo/core/base/batch_lin_op.hpp`

47.80 gko::batch::solver::EnableBatchSolver< ConcreteSolver, ValueType, PolymorphicBase > Class Template Reference

This mixin provides apply and common iterative solver functionality to all the batched solvers.

```
#include <ginkgo/core/solver/batch_solver_base.hpp>
```

Additional Inherited Members

47.80.1 Detailed Description

```
template<typename ConcreteSolver, typename ValueType, typename PolymorphicBase = BatchLinOp>
class gko::batch::solver::EnableBatchSolver< ConcreteSolver, ValueType, PolymorphicBase >
```

This mixin provides apply and common iterative solver functionality to all the batched solvers.

Template Parameters

<i>ConcreteSolver</i>	The concrete solver class.
<i>ValueType</i>	The value type of the multivectors.
<i>PolymorphicBase</i>	The base class; must be a subclass of <code>BatchLinOp</code> .

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/batch_solver_base.hpp`

47.81 `gko::EnableCreateMethod`< `ConcreteType` > Class Template Reference

This mixin implements a static `create()` method on `ConcreteType` that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an `std::unique_ptr` to such an object.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

47.81.1 Detailed Description

```
template<typename ConcreteType>
class gko::EnableCreateMethod< ConcreteType >
```

This mixin implements a static `create()` method on `ConcreteType` that dynamically allocates the memory, uses the passed-in arguments to construct the object, and returns an `std::unique_ptr` to such an object.

Template Parameters

<i>ConcreteObject</i>	the concrete type for which <code>create()</code> is being implemented [CRTP parameter]
-----------------------	---

The documentation for this class was generated from the following file:

- `ginkgo/core/base/polymorphic_object.hpp`

47.82 `gko::EnableDefaultFactory`< `ConcreteFactory`, `ProductType`, `ParametersType`, `PolymorphicBase` > Class Template Reference

This mixin provides a default implementation of a concrete factory.

```
#include <ginkgo/core/base/abstract_factory.hpp>
```

Public Member Functions

- `const parameters_type & get_parameters()` `const noexcept`
Returns the parameters of the factory.

Static Public Member Functions

- static parameters_type [create](#) ()

Creates a new ParametersType object which can be used to instantiate a new ConcreteFactory.

47.82.1 Detailed Description

```
template<typename ConcreteFactory, typename ProductType, typename ParametersType, typename PolymorphicBase>
class gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >
```

This mixin provides a default implementation of a concrete factory.

It implements all the methods of [AbstractFactory](#) and [PolymorphicObject](#). Its implementation of the `generate_impl()` method delegates the creation of the product by calling the `ProductType::ProductType(const ConcreteFactory *, const components_type &)` constructor. The factory also supports parameters by using the `ParametersType` structure, which is defined by the user.

For a simple example, see `IntFactory` in `core/test/base/abstract_factory.cpp`.

Template Parameters

<i>ConcreteFactory</i>	the concrete factory which is being implemented [CRTP parameter]
<i>ProductType</i>	the concrete type of products which this factory produces, has to be a subclass of <code>PolymorphicBase::abstract_product_type</code>
<i>ParametersType</i>	a type representing the parameters of the factory, has to inherit from the enable_parameters_type mixin
<i>PolymorphicBase</i>	parent of <code>ConcreteFactory</code> in the polymorphic hierarchy, has to be a subclass of AbstractFactory

47.82.2 Member Function Documentation

47.82.2.1 `create()`

```
template<typename ConcreteFactory , typename ProductType , typename ParametersType , typename
PolymorphicBase >
static parameters_type gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >::create ( ) [inline], [static]
```

Creates a new ParametersType object which can be used to instantiate a new ConcreteFactory.

This method does not construct the factory directly, but returns a new `parameters_type` object, which can be used to set the parameters of the factory. Once the parameters have been set, the `parameters_type::on()` method can be used to obtain an instance of the factory with those parameters.

Returns

a default `parameters_type` object

47.82.2.2 get_parameters()

```
template<typename ConcreteFactory , typename ProductType , typename ParametersType , typename
PolymorphicBase >
const parameters_type& gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >::get_parameters ( ) const [inline], [noexcept]
```

Returns the parameters of the factory.

Returns

the parameters of the factory

The documentation for this class was generated from the following file:

- ginkgo/core/base/abstract_factory.hpp

47.83 gko::solver::EnableIterativeBase< DerivedType > Class Template Reference

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

- [EnableIterativeBase](#) & [operator=](#) (const [EnableIterativeBase](#) &other)
Creates a shallow copy of the provided stopping criterion, clones it onto this executor if executors don't match.
- [EnableIterativeBase](#) & [operator=](#) ([EnableIterativeBase](#) &&other)
Moves the provided stopping criterion, clones it onto this executor if executors don't match.
- [EnableIterativeBase](#) (const [EnableIterativeBase](#) &other)
Creates a shallow copy of the provided stopping criterion.
- [EnableIterativeBase](#) ([EnableIterativeBase](#) &&other)
Moves the provided stopping criterion.
- void [set_stop_criterion_factory](#) (std::shared_ptr< const [stop::CriterionFactory](#) > new_stop_factory) override
Sets the stopping criterion of the solver.

47.83.1 Detailed Description

```
template<typename DerivedType>
class gko::solver::EnableIterativeBase< DerivedType >
```

A LinOp deriving from this CRTP class stores a stopping criterion factory and allows applying with a guess.

Template Parameters

<i>DerivedType</i>	the CRTP type that derives from this
--------------------	--------------------------------------

47.83.2 Constructor & Destructor Documentation

47.83.2.1 EnableIterativeBase()

```
template<typename DerivedType>
gko::solver::EnableIterativeBase< DerivedType >::EnableIterativeBase (
    EnableIterativeBase< DerivedType > && other ) [inline]
```

Moves the provided stopping criterion.

The moved-from object has a nullptr stopping criterion.

47.83.3 Member Function Documentation

47.83.3.1 operator=()

```
template<typename DerivedType>
EnableIterativeBase& gko::solver::EnableIterativeBase< DerivedType >::operator= (
    EnableIterativeBase< DerivedType > && other ) [inline]
```

Moves the provided stopping criterion, clones it onto this executor if executors don't match.

The moved-from object has a nullptr stopping criterion.

47.83.3.2 set_stop_criterion_factory()

```
template<typename DerivedType>
void gko::solver::EnableIterativeBase< DerivedType >::set_stop_criterion_factory (
    std::shared_ptr< const stop::CriterionFactory > new_stop_factory ) [inline],
[override], [virtual]
```

Sets the stopping criterion of the solver.

Parameters

<i>other</i>	the new stopping criterion factory
--------------	------------------------------------

Reimplemented from [gko::solver::IterativeBase](#).

Referenced by [gko::solver::EnableIterativeBase< Bicg< ValueType > >::operator=\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/solver/solver_base.hpp](#)

47.84 gko::EnableLinOp< ConcreteLinOp, PolymorphicBase > Class Template Reference

The [EnableLinOp](#) mixin can be used to provide sensible default implementations of the majority of the LinOp and [PolymorphicObject](#) interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Additional Inherited Members

47.84.1 Detailed Description

```
template<typename ConcreteLinOp, typename PolymorphicBase = LinOp>
class gko::EnableLinOp< ConcreteLinOp, PolymorphicBase >
```

The [EnableLinOp](#) mixin can be used to provide sensible default implementations of the majority of the LinOp and [PolymorphicObject](#) interface.

The goal of the mixin is to facilitate the development of new LinOp, by enabling the implementers to focus on the important parts of their operator, while the library takes care of generating the trivial utility functions. The mixin will provide default implementations for the entire [PolymorphicObject](#) interface, including a default implementation of `copy_from` between objects of the new LinOp type. It will also hide the default `LinOp::apply()` methods with versions that preserve the static type of the object.

Implementers of new LinOps are required to specify only the following aspects:

1. Creation of the LinOp: This can be facilitated via either [EnableCreateMethod](#) mixin (used mostly for matrix formats), or `GKO_ENABLE_LIN_OP_FACTORY` macro (used for operators created from other operators, like preconditioners and solvers).
2. Application of the LinOp: Implementers have to override the two overloads of the `LinOp::apply_impl()` virtual methods.

Note

This mixin can't be used with concrete types that derive from [experimental::distributed::DistributedBase](#). In that case use `experimental::EnableDistributedLinOp` instead.

Template Parameters

<i>ConcreteLinOp</i>	the concrete LinOp which is being implemented [CRTP parameter]
<i>PolymorphicBase</i>	parent of ConcreteLinOp in the polymorphic hierarchy, has to be a subclass of LinOp

The documentation for this class was generated from the following file:

- `ginkgo/core/base/lin_op.hpp`

47.85 gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase > Class Template Reference

[EnableLogging](#) is a mixin which should be inherited by any class which wants to enable logging.

```
#include <ginkgo/core/log/logger.hpp>
```

47.85.1 Detailed Description

```
template<typename ConcreteLoggable, typename PolymorphicBase = Loggable>
class gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase >
```

[EnableLogging](#) is a mixin which should be inherited by any class which wants to enable logging.

All the received events are passed to the loggers this class contains.

Template Parameters

<i>ConcreteLoggable</i>	the object being logged [CRTP parameter]
<i>PolymorphicBase</i>	the polymorphic base of this class. By default it is Loggable . Change it if you want to use a new superclass of Loggable as polymorphic base of this class.

The documentation for this class was generated from the following file:

- ginkgo/core/log/logger.hpp

47.86 gko::multigrid::EnableMultigridLevel< ValueType > Class Template Reference

The [EnableMultigridLevel](#) gives the default implementation of [MultigridLevel](#) with composition and provides `set↔_multigrid_level` function.

```
#include <ginkgo/core/multigrid/multigrid_level.hpp>
```

Public Member Functions

- `std::shared_ptr< const LinOp > get\_fine\_op () const` override
Returns the operator on fine level.
- `std::shared_ptr< const LinOp > get\_restrict\_op () const` override
Returns the restrict operator.
- `std::shared_ptr< const LinOp > get\_coarse\_op () const` override
Returns the operator on coarse level.
- `std::shared_ptr< const LinOp > get\_prolong\_op () const` override
Returns the prolong operator.

47.86.1 Detailed Description

```
template<typename ValueType>
class gko::multigrid::EnableMultigridLevel< ValueType >
```

The [EnableMultigridLevel](#) gives the default implementation of [MultigridLevel](#) with composition and provides `set↔_multigrid_level` function.

A class inherit from [EnableMultigridLevel](#) should use the `this->get_compositions()->apply(...)` as its own `apply`, which represents `op(b) = prolong(coarse(restrict(b)))`.

47.86.2 Member Function Documentation

47.86.2.1 get_coarse_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_coarse_op
( ) const [inline], [override], [virtual]
```

Returns the operator on coarse level.

Returns

the operator on coarse level.

Implements [gko::multigrid::MultigridLevel](#).

```
97 {
98     return this->get_operator_at(1);
99 }
```

References [gko::UseComposition< ValueType >::get_operator_at\(\)](#).

47.86.2.2 get_fine_op()

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_fine_op (
) const [inline], [override], [virtual]
```

Returns the operator on fine level.

Returns

the operator on fine level.

Implements [gko::multigrid::MultigridLevel](#).

47.86.2.3 `get_prolong_op()`

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_prolong↵
_op ( ) const [inline], [override], [virtual]
```

Returns the prolong operator.

Returns

the prolong operator.

Implements [gko::multigrid::MultigridLevel](#).

References [gko::UseComposition< ValueType >::get_operator_at\(\)](#).

47.86.2.4 `get_restrict_op()`

```
template<typename ValueType >
std::shared_ptr<const LinOp> gko::multigrid::EnableMultigridLevel< ValueType >::get_restrict↵
_op ( ) const [inline], [override], [virtual]
```

Returns the restrict operator.

Returns

the restrict operator.

Implements [gko::multigrid::MultigridLevel](#).

References [gko::UseComposition< ValueType >::get_operator_at\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/multigrid/multigrid_level.hpp](#)

47.87 `gko::EnablePolymorphicAssignment< ConcreteType, ResultType >` Class Template Reference

This mixin is used to enable a default [PolymorphicObject::copy_from\(\)](#) implementation for objects that have implemented conversions between them.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

Public Member Functions

- void [convert_to](#) (result_type *result) const override
Converts the implementer to an object of type result_type.
- void [move_to](#) (result_type *result) override
Converts the implementer to an object of type result_type by moving data from this object.

47.87.1 Detailed Description

```
template<typename ConcreteType, typename ResultType = ConcreteType>  
class gko::EnablePolymorphicAssignment< ConcreteType, ResultType >
```

This mixin is used to enable a default [PolymorphicObject::copy_from\(\)](#) implementation for objects that have implemented conversions between them.

The requirement is that there is either a conversion constructor from `ConcreteType` in `ResultType`, or a conversion operator to `ResultType` in `ConcreteType`.

Template Parameters

<i>ConcreteType</i>	the concrete type from which the copy_from is being enabled [CRTP parameter]
<i>ResultType</i>	the type to which copy_from is being enabled

47.87.2 Member Function Documentation

47.87.2.1 convert_to()

```
template<typename ConcreteType, typename ResultType = ConcreteType>
void gko::EnablePolymorphicAssignment< ConcreteType, ResultType >::convert_to (
    result_type * result ) const [inline], [override], [virtual]
```

Converts the implementer to an object of type result_type.

Parameters

<i>result</i>	the object used to store the result of the conversion
---------------	---

Implements [gko::ConvertibleTo< ResultType >](#).

47.87.2.2 move_to()

```
template<typename ConcreteType, typename ResultType = ConcreteType>
void gko::EnablePolymorphicAssignment< ConcreteType, ResultType >::move_to (
    result_type * result ) [inline], [override], [virtual]
```

Converts the implementer to an object of type result_type by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

<i>result</i>	the object used to emplace the result of the conversion
---------------	---

Note

[ConvertibleTo::move_to](#) can be implemented by simply calling [ConvertibleTo::convert_to](#). However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implements [gko::ConvertibleTo< ResultType >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/polymorphic_object.hpp

47.88 gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase > Class Template Reference

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new concrete polymorphic object.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

47.88.1 Detailed Description

```
template<typename ConcreteObject, typename PolymorphicBase = PolymorphicObject>
class gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase >
```

This mixin inherits from (a subclass of) [PolymorphicObject](#) and provides a base implementation of a new concrete polymorphic object.

The mixin changes parameter and return types of appropriate public methods of [PolymorphicObject](#) in the same way [EnableAbstractPolymorphicObject](#) does. In addition, it also provides default implementations of [PolymorphicObject](#)'s virtual methods by using the *executor default constructor* and the assignment operator of [ConcreteObject](#). Consequently, the following is a minimal example of [PolymorphicObject](#):

```
{c++}
struct MyObject : EnablePolymorphicObject<MyObject> {
    MyObject (std::shared_ptr<const Executor> exec)
        : EnablePolymorphicObject<MyObject> (std::move(exec))
    {}
};
```

In a way, this mixin can be viewed as an extension of default constructor/destructor/assignment operators.

Note

This mixin does not enable copying the polymorphic object to the object of the same type (i.e. it does not implement the [ConvertibleTo<ConcreteObject>](#) interface). To enable a default implementation of this interface see the [EnablePolymorphicAssignment](#) mixin.

Template Parameters

<i>ConcreteObject</i>	the concrete type which is being implemented [CRTP parameter]
<i>PolymorphicBase</i>	parent of <i>ConcreteObject</i> in the polymorphic hierarchy, has to be a subclass of polymorphic object

The documentation for this class was generated from the following file:

- `ginkgo/core/base/polymorphic_object.hpp`

47.89 gko::solver::EnablePreconditionable< DerivedType > Class Template Reference

Mixin providing default operation for [Preconditionable](#) with correct value semantics.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

- void [set_preconditioner](#) (std::shared_ptr< const LinOp > new_precond) override
Sets the preconditioner operator used by the [Preconditionable](#).
- [EnablePreconditionable](#) & operator= (const [EnablePreconditionable](#) &other)
Creates a shallow copy of the provided preconditioner, clones it onto this executor if executors don't match.
- [EnablePreconditionable](#) & operator= ([EnablePreconditionable](#) &&other)
Moves the provided preconditioner, clones it onto this executor if executors don't match.
- [EnablePreconditionable](#) (const [EnablePreconditionable](#) &other)
Creates a shallow copy of the provided preconditioner.
- [EnablePreconditionable](#) ([EnablePreconditionable](#) &&other)
Moves the provided preconditioner.

47.89.1 Detailed Description

```
template<typename DerivedType>
class gko::solver::EnablePreconditionable< DerivedType >
```

Mixin providing default operation for [Preconditionable](#) with correct value semantics.

It ensures that the preconditioner stored in this class will always have the same executor as the object this mixin is used in, creating a clone on the correct executor if necessary.

Template Parameters

<i>DerivedType</i>	The type that this Mixin is used in. It must provide <code>get_size()</code> and <code>get_executor()</code> functions that return correctly initialized values when the EnablePreconditionable constructor is called, i.e. the constructor must be provided by a base class added before EnablePreconditionable , since the member initialization order also applying to multiple inheritance.
--------------------	---

47.89.2 Constructor & Destructor Documentation

47.89.2.1 [EnablePreconditionable\(\)](#)

```
template<typename DerivedType>
gko::solver::EnablePreconditionable< DerivedType >::EnablePreconditionable (
    EnablePreconditionable< DerivedType > && other ) [inline]
```

Moves the provided preconditioner.

The moved-from object has a nullptr preconditioner.

47.89.3 Member Function Documentation

47.89.3.1 operator=()

```
template<typename DerivedType>
EnablePreconditionable& gko::solver::EnablePreconditionable< DerivedType >::operator= (
    EnablePreconditionable< DerivedType > && other ) [inline]
```

Moves the provided preconditioner, clones it onto this executor if executors don't match.

The moved-from object has a nullptr preconditioner.

47.89.3.2 set_preconditioner()

```
template<typename DerivedType>
void gko::solver::EnablePreconditionable< DerivedType >::set_preconditioner (
    std::shared_ptr< const LinOp > new_precond ) [inline], [override], [virtual]
```

Sets the preconditioner operator used by the [Preconditionable](#).

Parameters

<i>new_precond</i>	the new preconditioner operator used by the Preconditionable
--------------------	--

Reimplemented from [gko::Preconditionable](#).

Referenced by `gko::solver::EnablePreconditionable< Bicg< ValueType > >::operator=()`.

The documentation for this class was generated from the following file:

- ginkgo/core/solver/solver_base.hpp

47.90 gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType > Class Template Reference

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Additional Inherited Members

47.90.1 Detailed Description

```
template<typename ValueType, typename DerivedType>
class gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType >
```

A LinOp implementing this interface stores a system matrix and stopping criterion factory.

Template Parameters

<i>ValueType</i>	the value type that iterative solver uses for its vectors
<i>DerivedType</i>	the CRTP type that derives from this

The documentation for this class was generated from the following file:

- ginkgo/core/solver/solver_base.hpp

47.91 gko::solver::EnableSolverBase< DerivedType, MatrixType > Class Template Reference

A LinOp deriving from this CRTP class stores a system matrix.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

- [EnableSolverBase](#) & [operator=](#) (const [EnableSolverBase](#) &other)
Creates a shallow copy of the provided system matrix, clones it onto this executor if executors don't match.
- [EnableSolverBase](#) & [operator=](#) ([EnableSolverBase](#) &&other)
Moves the provided system matrix, clones it onto this executor if executors don't match.
- [EnableSolverBase](#) (const [EnableSolverBase](#) &other)
Creates a shallow copy of the provided system matrix.
- [EnableSolverBase](#) ([EnableSolverBase](#) &&other)
Moves the provided system matrix.
- `std::vector< int > get\_workspace\_scalars ()` const override
Returns the IDs of all scalars (workspace vectors with system dimension-independent size, usually 1 x num_rhs).
- `std::vector< int > get\_workspace\_vectors ()` const override
Returns the IDs of all vectors (workspace vectors with system dimension-dependent size, usually system_matrix_size x num_rhs).

47.91.1 Detailed Description

```
template<typename DerivedType, typename MatrixType = LinOp>
class gko::solver::EnableSolverBase< DerivedType, MatrixType >
```

A LinOp deriving from this CRTP class stores a system matrix.

Template Parameters

<i>DerivedType</i>	the CRTP type that derives from this
<i>MatrixType</i>	the concrete matrix type to be stored as system_matrix

47.91.2 Constructor & Destructor Documentation

47.91.2.1 EnableSolverBase()

```
template<typename DerivedType, typename MatrixType = LinOp>
gko::solver::EnableSolverBase< DerivedType, MatrixType >::EnableSolverBase (
    EnableSolverBase< DerivedType, MatrixType > && other ) [inline]
```

Moves the provided system matrix.

The moved-from object has a nullptr system matrix.

47.91.3 Member Function Documentation

47.91.3.1 operator=()

```
template<typename DerivedType, typename MatrixType = LinOp>
EnableSolverBase& gko::solver::EnableSolverBase< DerivedType, MatrixType >::operator= (
    EnableSolverBase< DerivedType, MatrixType > && other ) [inline]
```

Moves the provided system matrix, clones it onto this executor if executors don't match.

The moved-from object has a nullptr system matrix.

The documentation for this class was generated from the following file:

- ginkgo/core/solver/solver_base.hpp

47.92 gko::experimental::mpi::environment Class Reference

Class that sets up and finalizes the MPI environment.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

- int [get_provided_thread_support](#) () const
Return the provided thread support.
- [environment](#) (int &argc, char **&argv, const [thread_type](#) thread_t=thread_type::serialized)
Call MPI_Init_thread and initialize the MPI environment.
- [~environment](#) ()
Call MPI_Finalize at the end of the scope of this class.

47.92.1 Detailed Description

Class that sets up and finalizes the MPI environment.

This class is a simple RAII wrapper to MPI_Init and MPI_Finalize.

MPI_Init must have been called before calling any MPI functions.

Note

If MPI_Init has already been called, then this class should not be used.

47.92.2 Constructor & Destructor Documentation

47.92.2.1 environment()

```
gko::experimental::mpi::environment::environment (
    int & argc,
    char **& argv,
    const thread_type thread_t = thread_type::serialized ) [inline]
```

Call MPI_Init_thread and initialize the MPI environment.

Parameters

<i>argc</i>	the number of arguments to the main function.
<i>argv</i>	the arguments provided to the main function.
<i>thread_t</i>	the type of threading for initialization. See @thread_type

47.92.3 Member Function Documentation

47.92.3.1 get_provided_thread_support()

```
int gko::experimental::mpi::environment::get_provided_thread_support ( ) const [inline]
```

Return the provided thread support.

Returns

the provided thread support

The documentation for this class was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.93 gko::Error Class Reference

The [Error](#) class is used to report exceptional behaviour in library functions.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [Error](#) (const std::string &file, int line, const std::string &what)
Initializes an error.
- virtual const char * [what](#) () const noexcept override
Returns a human-readable string with a more detailed description of the error.

47.93.1 Detailed Description

The [Error](#) class is used to report exceptional behaviour in library functions.

Ginkgo uses C++ exception mechanism to this end, and the [Error](#) class represents a base class for all types of errors. The exact list of errors which could occur during the execution of a certain library routine is provided in the documentation of that routine, along with a short description of the situation when that error can occur. During runtime, these errors can be detected by using standard C++ try-catch blocks, and a human-readable error description can be obtained by calling the [Error::what\(\)](#) method.

As an example, trying to compute a matrix-vector product with arguments of incompatible size will result in a [DimensionMismatch](#) error, which is demonstrated in the following program.

```
#include <ginkgo.h>
#include <iostream>
using namespace gko;
int main()
{
    auto omp = create<OmpExecutor>();
    auto A = randn_fill<matrix::Csr<float>>(5, 5, 0f, 1f, omp);
    auto x = fill<matrix::Dense<float>>(6, 1, 1f, omp);
    try {
        auto y = apply(A, x);
    } catch(Error e) {
        // an error occurred, write the message to screen and exit
        std::cout << e.what() << std::endl;
        return -1;
    }
    return 0;
}
```

47.93.2 Constructor & Destructor Documentation

47.93.2.1 Error()

```
gko::Error::Error (
    const std::string & file,
    int line,
    const std::string & what ) [inline]
```

Initializes an error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>what</i>	The error message

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.94 gko::Executor Class Reference

The first step in using the Ginkgo library consists of creating an executor.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- virtual void [run](#) (const [Operation](#) &op) const =0
Runs the specified [Operation](#) using this [Executor](#).
- template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void [run](#) (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const
Runs one of the passed in functors, depending on the [Executor](#) type.
- template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename Closure↔
Dpcpp >
void [run](#) (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↔
Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const
Runs one of the passed in functors, depending on the [Executor](#) type.
- template<typename T >
T * [alloc](#) (size_type num_elems) const
Allocates memory in this [Executor](#).
- void [free](#) (void *ptr) const noexcept
Frees memory previously allocated with [Executor::alloc\(\)](#).
- template<typename T >
void [copy_from](#) (ptr_param< const [Executor](#) > src_exec, size_type num_elems, const T *src_ptr, T *dest↔
_ptr) const
Copies data from another [Executor](#).
- template<typename T >
void [copy](#) (size_type num_elems, const T *src_ptr, T *dest_ptr) const
Copies data within this [Executor](#).
- template<typename T >
T [copy_val_to_host](#) (const T *ptr) const
Retrieves a single element at the given location from executor memory.
- virtual std::shared_ptr< [Executor](#) > [get_master](#) () noexcept=0
Returns the master [OmpExecutor](#) of this [Executor](#).
- virtual std::shared_ptr< const [Executor](#) > [get_master](#) () const noexcept=0
Returns the master [OmpExecutor](#) of this [Executor](#).
- virtual void [synchronize](#) () const =0

Synchronize the operations launched on the executor with its master.

- void `add_logger` (std::shared_ptr< const log::Logger > logger) override
- void `remove_logger` (const log::Logger *logger) override
- void `set_log_propagation_mode` (log_propagation_mode mode)

Sets the logger event propagation mode for the executor.

- bool `should_propagate_log` () const

Returns true iff events occurring at an object created on this executor should be logged at propagating loggers attached to this executor, and there is at least one such propagating logger.

- bool `memory_accessible` (const std::shared_ptr< const Executor > &other) const

Verifies whether the executors share the same memory.

- virtual std::string `get_description` () const =0

47.94.1 Detailed Description

The first step in using the Ginkgo library consists of creating an executor.

Executors are used to specify the location for the data of linear algebra objects, and to determine where the operations will be executed. Ginkgo currently supports five different executor types:

- `OmpExecutor` specifies that the data should be stored and the associated operations executed on an OpenMP-supporting device (e.g. host CPU);
- `CudaExecutor` specifies that the data should be stored and the operations executed on the NVIDIA GPU accelerator;
- `HipExecutor` specifies that the data should be stored and the operations executed on either an NVIDIA or AMD GPU accelerator;
- `DpcppExecutor` specifies that the data should be stored and the operations executed on an hardware supporting DPC++;
- `ReferenceExecutor` executes a non-optimized reference implementation, which can be used to debug the library.

The following code snippet demonstrates the simplest possible use of the Ginkgo library:

```
auto omp = gko::create<gko::OmpExecutor>();
auto A = gko::read_from_mtx<gko::matrix::Csr<float>>("A.mtx", omp);
```

First, we create a OMP executor, which will be used in the next line to specify where we want the data for the matrix A to be stored. The second line will read a matrix from the matrix market file 'A.mtx', and store the data on the CPU in CSR format (`gko::matrix::Csr` is a Ginkgo matrix class which stores its data in CSR format). At this point, matrix A is bound to the CPU, and any routines called on it will be performed on the CPU. This approach is usually desired in sparse linear algebra, as the cost of individual operations is several orders of magnitude lower than the cost of copying the matrix to the GPU.

If matrix A is going to be reused multiple times, it could be beneficial to copy it over to the accelerator, and perform the operations there, as demonstrated by the next code snippet:

```
auto cuda = gko::create<gko::CudaExecutor>(0, omp);
auto dA = gko::copy_to<gko::matrix::Csr<float>>(A.get(), cuda);
```

The first line of the snippet creates a new CUDA executor. Since there may be multiple NVIDIA GPUs present on the system, the first parameter instructs the library to use the first device (i.e. the one with device ID zero, as in `cudaSetDevice()` routine from the CUDA runtime API). In addition, since GPUs are not stand-alone processors, it is required to pass a "master" `OmpExecutor` which will be used to schedule the requested CUDA kernels on the accelerator.

The second command creates a copy of the matrix A on the GPU. Notice the use of the `get()` method. As Ginkgo aims to provide automatic memory management of its objects, the result of calling `gko::read_from_mtx()` is a smart pointer (`std::unique_ptr`) to the created object. On the other hand, as the library will not hold a reference to A once the copy is completed, the input parameter for `gko::copy_to()` is a plain pointer. Thus, the `get()` method is used to convert from a `std::unique_ptr` to a plain pointer, as expected by `gko::copy_to()`.

As a side note, the `gko::copy_to` routine is far more powerful than just copying data between different devices. It can also be used to convert data between different formats. For example, if the above code used [gko::matrix::Ell](#) as the template parameter, dA would be stored on the GPU, in ELLPACK format.

Finally, if all the processing of the matrix is supposed to be done on the GPU, and a CPU copy of the matrix is not required, we could have read the matrix to the GPU directly:

```
auto omp = gko::create<gko::OmpExecutor>();
auto cuda = gko::create<gko::CudaExecutor>(0, omp);
auto dA = gko::read_from_mtx<gko::matrix::Csr<float>>("A.mtx", cuda);
```

Notice that even though reading the matrix directly from a file to the accelerator is not supported, the library is designed to abstract away the intermediate step of reading the matrix to the CPU memory. This is a general design approach taken by the library: in case an operation is not supported by the device, the data will be copied to the CPU, the operation performed there, and finally the results copied back to the device. This approach makes using the library more concise, as explicit copies are not required by the user. Nevertheless, this feature should be taken into account when considering performance implications of using such operations.

47.94.2 Member Function Documentation

47.94.2.1 add_logger()

```
void gko::Executor::add_logger (
    std::shared_ptr< const log::Logger > logger ) [inline], [override], [virtual]
```

Note

This specialization keeps track of whether any propagating loggers were attached to the executor.

See also

`Logger::needs_propagation()`

Implements [gko::log::Loggable](#).

47.94.2.2 alloc()

```
template<typename T >
T* gko::Executor::alloc (
    size_type num_elems ) const [inline]
```

Allocates memory in this [Executor](#).

Template Parameters

<i>T</i>	datatype to allocate
----------	----------------------

Parameters

<i>num_elems</i>	number of elements of type T to allocate
------------------	--

Exceptions

<i>AllocationError</i>	if the allocation failed
--	--------------------------

Returns

pointer to allocated memory

47.94.2.3 copy()

```
template<typename T >
void gko::Executor::copy (
    size_type num_elems,
    const T * src_ptr,
    T * dest_ptr ) const [inline]
```

Copies data within this [Executor](#).

Template Parameters

<i>T</i>	datatype to copy
----------	------------------

Parameters

<i>num_elems</i>	number of elements of type T to copy
<i>src_ptr</i>	pointer to a block of memory containing the data to be copied
<i>dest_ptr</i>	pointer to an allocated block of memory where the data will be copied to

References [copy_from\(\)](#).

47.94.2.4 copy_from()

```
template<typename T >
void gko::Executor::copy_from (
    ptr_param< const Executor > src_exec,
```

```

size_type num_elems,
const T * src_ptr,
T * dest_ptr ) const [inline]

```

Copies data from another [Executor](#).

Template Parameters

<i>T</i>	datatype to copy
----------	------------------

Parameters

<i>src_exec</i>	Executor from which the memory will be copied
<i>num_elems</i>	number of elements of type T to copy
<i>src_ptr</i>	pointer to a block of memory containing the data to be copied
<i>dest_ptr</i>	pointer to an allocated block of memory where the data will be copied to

References `gko::ptr_param< T >::get()`.

Referenced by `copy()`.

47.94.2.5 `copy_val_to_host()`

```

template<typename T >
T gko::Executor::copy_val_to_host (
    const T * ptr ) const [inline]

```

Retrieves a single element at the given location from executor memory.

Template Parameters

<i>T</i>	datatype to copy
----------	------------------

Parameters

<i>ptr</i>	the pointer to the element to be copied
------------	---

Returns

the value stored at `ptr`

References `get_master()`.

47.94.2.6 free()

```
void gko::Executor::free (
    void * ptr ) const [inline], [noexcept]
```

Frees memory previously allocated with [Executor::alloc\(\)](#).

If `ptr` is a `nullptr`, the function has no effect.

Parameters

<i>ptr</i>	pointer to the allocated memory block
------------	---------------------------------------

47.94.2.7 get_description()

```
virtual std::string gko::Executor::get_description ( ) const [pure virtual]
```

Returns

a textual representation of the executor and its device.

Implemented in [gko::DpcppExecutor](#), [gko::HipExecutor](#), [gko::CudaExecutor](#), [gko::ReferenceExecutor](#), and [gko::OmpExecutor](#).

47.94.2.8 get_master() [1/2]

```
virtual std::shared_ptr<const Executor> gko::Executor::get_master ( ) const [pure virtual],
[noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implemented in [gko::DpcppExecutor](#), [gko::HipExecutor](#), [gko::CudaExecutor](#), and [gko::OmpExecutor](#).

47.94.2.9 get_master() [2/2]

```
virtual std::shared_ptr<Executor> gko::Executor::get_master ( ) [pure virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implemented in [gko::DpcppExecutor](#), [gko::HipExecutor](#), [gko::CudaExecutor](#), and [gko::OmpExecutor](#).

Referenced by `copy_val_to_host()`.

47.94.2.10 memory_accessible()

```
bool gko::Executor::memory_accessible (
    const std::shared_ptr< const Executor > & other ) const [inline]
```

Verifies whether the executors share the same memory.

Parameters

<i>other</i>	the other Executor to compare against
--------------	---

Returns

whether the executors this and other share the same memory.

47.94.2.11 remove_logger()

```
void gko::Executor::remove_logger (
    const log::Logger * logger ) [inline], [override], [virtual]
```

Note

This specialization keeps track of whether any propagating loggers were attached to the executor.

See also

[Logger::needs_propagation\(\)](#)

Implements [gko::log::Loggable](#).

47.94.2.12 run() [1/3]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename Closure←
Dpcpp >
void gko::Executor::run (
    const ClosureOmp & op_omp,
    const ClosureCuda & op_cuda,
    const ClosureHip & op_hip,
    const ClosureDpcpp & op_dpcpp ) const [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

References [run\(\)](#).

47.94.2.13 [run\(\)](#) [2/3]

```
virtual void gko::Executor::run (
    const Operation & op ) const [pure virtual]
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

Implemented in [gko::ReferenceExecutor](#).

Referenced by [run\(\)](#).

47.94.2.14 [run\(\)](#) [3/3]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    std::string name,
    const ClosureReference & op_ref,
    const ClosureOmp & op_omp,
    const ClosureCuda & op_cuda,
    const ClosureHip & op_hip,
    const ClosureDpcpp & op_dpcpp ) const [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of <i>op_ref</i>
<i>ClosureOmp</i>	type of <i>op_omp</i>
<i>ClosureCuda</i>	type of <i>op_cuda</i>
<i>ClosureHip</i>	type of <i>op_hip</i>
<i>ClosureDpcpp</i>	type of <i>op_dpcpp</i>

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

References [run\(\)](#).

47.94.2.15 `set_log_propagation_mode()`

```
void gko::Executor::set_log_propagation_mode (
    log_propagation_mode mode ) [inline]
```

Sets the logger event propagation mode for the executor.

This controls whether events that happen at objects created on this executor will also be logged at propagating loggers attached to the executor.

See also

[Logger::needs_propagation\(\)](#)

47.94.2.16 `should_propagate_log()`

```
bool gko::Executor::should_propagate_log ( ) const [inline]
```

Returns true iff events occurring at an object created on this executor should be logged at propagating loggers attached to this executor, and there is at least one such propagating logger.

See also

[Logger::needs_propagation\(\)](#)

[Executor::set_log_propagation_mode\(log_propagation_mode\)](#)

References [gko::automatic](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/executor.hpp](#)

47.95 gko::log::executor_data Struct Reference

Struct representing [Executor](#) related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.95.1 Detailed Description

Struct representing [Executor](#) related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.96 gko::executor_deleter< T > Class Template Reference

This is a deleter that uses an executor's `free` method to deallocate the data.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- [executor_deleter](#) (std::shared_ptr< const [Executor](#) > exec)
Creates a new deleter.
- void [operator\(\)](#) (pointer ptr) const
Deletes the object.

47.96.1 Detailed Description

```
template<typename T>
class gko::executor_deleter< T >
```

This is a deleter that uses an executor's `free` method to deallocate the data.

Template Parameters

<i>T</i>	the type of object being deleted
----------	----------------------------------

47.96.2 Constructor & Destructor Documentation

47.96.2.1 executor_deleter()

```
template<typename T >
gko::executor_deleter< T >::executor_deleter (
    std::shared_ptr< const Executor > exec ) [inline], [explicit]
```

Creates a new deleter.

Parameters

<i>exec</i>	the executor used to free the data
-------------	------------------------------------

47.96.3 Member Function Documentation

47.96.3.1 operator>()()

```
template<typename T >
void gko::executor_deleter< T >::operator() (
    pointer ptr ) const [inline]
```

Deletes the object.

Parameters

<i>ptr</i>	pointer to the object being deleted
------------	-------------------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.97 gko::experimental::factorization::Factorization< ValueType, IndexType > Class Template Reference

Represents a generic factorization consisting of two triangular factors (upper and lower) and an optional diagonal scaling matrix.

```
#include <ginkgo/core/factorization/factorization.hpp>
```

Public Member Functions

- `std::unique_ptr< Factorization > unpack () const`
Transforms the factorization from a compact representation suitable only for triangular solves to a composition representation that can also be used to access individual factors and multiply with the factorization.
- `storage_type get_storage_type () const`

- Returns the storage type used by this factorization.*
- `std::shared_ptr< const matrix\_type > get\_lower\_factor () const`
Returns the lower triangular factor of the factorization, if available, nullptr otherwise.
- `std::shared_ptr< const diag\_type > get\_diagonal () const`
Returns the diagonal scaling matrix of the factorization, if available, nullptr otherwise.
- `std::shared_ptr< const matrix\_type > get\_upper\_factor () const`
Returns the upper triangular factor of the factorization, if available, nullptr otherwise.
- `std::shared_ptr< const matrix\_type > get\_combined () const`
Returns the matrix storing a compact representation of the factorization, if available, nullptr otherwise.
- `Factorization (const Factorization &)`
Creates a deep copy of the factorization.
- `Factorization (Factorization &&)`
Moves from the given factorization, leaving it empty.

Static Public Member Functions

- `static std::unique_ptr< Factorization > create\_from\_composition (std::unique_ptr< composition\_type > composition)`
Creates a [Factorization](#) from an existing composition.
- `static std::unique_ptr< Factorization > create\_from\_symm\_composition (std::unique_ptr< composition\_type > composition)`
Creates a [Factorization](#) from an existing symmetric composition.
- `static std::unique_ptr< Factorization > create\_from\_combined\_lu (std::unique_ptr< matrix\_type > matrix)`
Creates a [Factorization](#) from an existing combined representation of an LU factorization.

47.97.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::experimental::factorization::Factorization< ValueType, IndexType >
```

Represents a generic factorization consisting of two triangular factors (upper and lower) and an optional diagonal scaling matrix.

This class is used to represent a wide range of different factorizations to be passed on to direct solvers and other similar operations. The `storage_type` represents how the individual factors are stored internally: They may be stored as separate matrices or in a single matrix, and be symmetric or unsymmetric, with the diagonal belonging to both factors, a single factor or being a separate scaling factor ([Cholesky](#) vs. LDL^H vs. LU vs. LDU).

Template Parameters

<i>ValueType</i>	the value type used to store the factorization entries
<i>IndexType</i>	the index type used to represent the sparsity pattern

47.97.2 Member Function Documentation

47.97.2.1 create_from_combined_lu()

```
template<typename ValueType , typename IndexType >
static std::unique_ptr<Factorization> gko::experimental::factorization::Factorization< Value←
Type, IndexType >::create_from_combined_lu (
    std::unique_ptr< matrix_type > matrix ) [static]
```

Creates a [Factorization](#) from an existing combined representation of an LU factorization.

Parameters

<i>matrix</i>	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular matrix, and the last entry to be the transpose of the first entry. If the composition has 3 elements, we expect the middle entry to be a diagonal matrix.
---------------	---

Returns

a symmetric [Factorization](#) storing the elements from the [Composition](#).

47.97.2.2 create_from_composition()

```
template<typename ValueType , typename IndexType >
static std::unique_ptr<Factorization> gko::experimental::factorization::Factorization< Value←
Type, IndexType >::create_from_composition (
    std::unique_ptr< composition_type > composition ) [static]
```

Creates a [Factorization](#) from an existing composition.

Parameters

<i>composition</i>	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular matrix, and the last entry to be an upper triangular matrix. If the composition has 3 elements, we expect the middle entry to be a diagonal matrix.
--------------------	---

Returns

a [Factorization](#) storing the elements from the [Composition](#).

47.97.2.3 create_from_symm_composition()

```
template<typename ValueType , typename IndexType >
static std::unique_ptr<Factorization> gko::experimental::factorization::Factorization< Value←
Type, IndexType >::create_from_symm_composition (
    std::unique_ptr< composition_type > composition ) [static]
```

Creates a [Factorization](#) from an existing symmetric composition.

Parameters

<i>composition</i>	the composition consisting of 2 or 3 elements. We expect the first entry to be a lower triangular matrix, and the last entry to be the transpose of the first entry. If the composition has 3 elements, we expect the middle entry to be a diagonal matrix.
--------------------	---

Returns

a symmetric [Factorization](#) storing the elements from the [Composition](#).

47.97.2.4 unpack()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<Factorization> gko::experimental::factorization::Factorization< ValueType,
IndexType >::unpack ( ) const
```

Transforms the factorization from a compact representation suitable only for triangular solves to a composition representation that can also be used to access individual factors and multiply with the factorization.

Returns

a new [Factorization](#) object containing this factorization represented as `storage_type::composition`.

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/factorization.hpp

47.98 gko::matrix::Fbcsr< ValueType, IndexType > Class Template Reference

Fixed-block compressed sparse row storage matrix format.

```
#include <ginkgo/core/matrix/fbcsr.hpp>
```

Public Member Functions

- void [convert_to](#) ([Csr](#)< [ValueType](#), [IndexType](#) > *result) const override
Converts the matrix to CSR format.
- void [convert_to](#) ([SparsityCsr](#)< [ValueType](#), [IndexType](#) > *result) const override
Get the block sparsity pattern in CSR-like format.
- void [read](#) (const [mat_data](#) &data) override
Reads a [matrix_data](#) into [Fbcsr](#) format.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< [LinOp](#) > [transpose](#) () const override
Returns a [LinOp](#) representing the transpose of the [Transposable](#) object.
- std::unique_ptr< [LinOp](#) > [conj_transpose](#) () const override
Returns a [LinOp](#) representing the conjugate transpose of the [Transposable](#) object.
- std::unique_ptr< [Diagonal](#)< [ValueType](#) > > [extract_diagonal](#) () const override
Extracts the diagonal entries of the matrix into a vector.
- std::unique_ptr< [absolute_type](#) > [compute_absolute](#) () const override
Gets the [AbsoluteLinOp](#).
- void [compute_absolute_inplace](#) () override
Compute absolute inplace on each element.
- void [sort_by_column_index](#) ()
Sorts the values blocks and block-column indices in each row by column index.
- bool [is_sorted_by_column_index](#) () const
Tests if all row entry pairs (value, col_idx) are sorted by column index.
- [value_type](#) * [get_values](#) () noexcept
- const [value_type](#) * [get_const_values](#) () const noexcept
- [index_type](#) * [get_col_idxs](#) () noexcept
- const [index_type](#) * [get_const_col_idxs](#) () const noexcept
- [index_type](#) * [get_row_ptrs](#) () noexcept
- const [index_type](#) * [get_const_row_ptrs](#) () const noexcept
- [size_type](#) [get_num_stored_elements](#) () const noexcept
- [size_type](#) [get_num_stored_blocks](#) () const noexcept
- int [get_block_size](#) () const noexcept
- [index_type](#) [get_num_block_rows](#) () const noexcept
- [index_type](#) [get_num_block_cols](#) () const noexcept
- [Fbcsr](#) & [operator=](#) (const [Fbcsr](#) &)
Copy-assigns an [Fbcsr](#) matrix.
- [Fbcsr](#) & [operator=](#) ([Fbcsr](#) &&)
Move-assigns an [Fbcsr](#) matrix.
- [Fbcsr](#) (const [Fbcsr](#) &)
Copy-constructs an [Ell](#) matrix.
- [Fbcsr](#) ([Fbcsr](#) &&)
Move-constructs an [Fbcsr](#) matrix.

Static Public Member Functions

- static std::unique_ptr< [Fbcsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, int block_size=1)
Creates an uninitialized FBCSR matrix with the given block size.
- static std::unique_ptr< [Fbcsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, [size_type](#) num_nonzeros, int block_size)
Creates an uninitialized FBCSR matrix of the specified size.
- static std::unique_ptr< [Fbcsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, int block_size, [array](#)< value_type > values, [array](#)< index_type > col_idxs, [array](#)< index_type > row_ptrs)
Creates a FBCSR matrix from already allocated (and initialized) row pointer, column index and value arrays.
- template<typename InputValueType, typename InputColumnIndexType, typename InputRowPtrType >
static std::unique_ptr< [Fbcsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, int block_size, std::initializer_list< InputValueType > values, std::initializer_list< InputColumnIndexType > col_idxs, std::initializer_list< InputRowPtrType > row_ptrs)
create(std::shared_ptr<const Executor>,
- static std::unique_ptr< const [Fbcsr](#) > [create_const](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, int blocksize, gko::detail::const_array_view< ValueType > &&values, gko::detail::const_array_view< IndexType > &&col_idxs, gko::detail::const_array_view< IndexType > &&row_ptrs)
Creates a constant (immutable) [Fbcsr](#) matrix from a constant array.

47.98.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Fbcsr< ValueType, IndexType >
```

Fixed-block compressed sparse row storage matrix format.

FBCSR is a matrix format meant for matrices having a natural block structure made up of small, dense, disjoint blocks. It is similar to CSR

See also

[Csr](#). However, unlike [Csr](#), each non-zero location stores a small dense block of entries having a constant size. This reduces the number of integers that need to be stored in order to refer to a given non-zero entry, and enables efficient implementation of certain block methods.

The block size is expected to be known in advance and passed to the constructor.

Note

The total number of rows and the number of columns are expected to be divisible by the block size.

The nonzero elements are stored in a 1D array row-wise, and accompanied with a row pointer array which stores the starting index of each block-row. An additional block-column index array is used to identify the block-column of each nonzero block.

The [Fbcsr](#) LinOp supports different operations:

```
matrix::Fbcsr *A, *B, *C;           // matrices
matrix::Dense *b, *x;              // vectors tall-and-skinny matrices
matrix::Dense *alpha, *beta;       // scalars of dimension 1x1
// Applying to Dense matrices computes an SpMV/SpMM product
A->apply(b, x)                      // x = A*b
A->apply(alpha, b, beta, x)         // x = alpha*A*b + beta*x
```

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.98.2 Constructor & Destructor Documentation**47.98.2.1 Fbcsr() [1/2]**

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Fbcsr< ValueType, IndexType >::Fbcsr (
    const Fbcsr< ValueType, IndexType > & )
```

Copy-constructs an [ElI](#) matrix.

Inherits executor and data.

47.98.2.2 Fbcsr() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Fbcsr< ValueType, IndexType >::Fbcsr (
    Fbcsr< ValueType, IndexType > && )
```

Move-constructs an [Fbcsr](#) matrix.

Inherits executor and data. The moved-from object is empty (0x0 with no nonzeros, but valid row pointers).

47.98.3 Member Function Documentation**47.98.3.1 compute_absolute()**

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Fbcsr< ValueType, IndexType >::compute_absolute (
    ) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Fbcsr< ValueType, IndexType > > >](#).

47.98.3.2 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Fbcsr< ValueType, IndexType >::conj_transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.98.3.3 convert_to() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::convert_to (
    Csr< ValueType, IndexType > * result ) const [override]
```

Converts the matrix to CSR format.

Note

Any explicit zeros in the original matrix are retained in the converted result.

47.98.3.4 convert_to() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::convert_to (
    SparsityCsr< ValueType, IndexType > * result ) const [override]
```

Get the block sparsity pattern in CSR-like format.

Note

The actual non-zero values are never copied; the result always has a value array of size 1 with the value 1.

47.98.3.5 create() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Fbcsr> gko::matrix::Fbcsr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    int block_size,
    array< value_type > values,
    array< index_type > col_idxs,
    array< index_type > row_ptrs ) [static]
```

Creates a FBCSR matrix from already allocated (and initialized) row pointer, column index and value arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>block_size</i>	Size of the small square dense nonzero blocks
<i>values</i>	the value array of the matrix, stored in column-major order for each block
<i>col_idx</i> s	the block column index array of the matrix
<i>row_ptr</i> s	the block row pointer array of the matrix

Note

If one of *row_ptr*s, *col_idx*s or *values* is not an rvalue, not an array of `IndexType`, `IndexType` and `ValueType`, respectively, or is on the wrong executor, an internal copy of that array will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.98.3.6 `create()` [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename InputValueType, typename InputColumnIndexType, typename InputRowPtrType >
static std::unique_ptr<Fbcsr> gko::matrix::Fbcsr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    int block_size,
    std::initializer_list< InputValueType > values,
    std::initializer_list< InputColumnIndexType > col_idx,
    std::initializer_list< InputRowPtrType > row_ptr ) [inline], [static]
```

`create(std::shared_ptr<const Executor>,`

`create(std::shared_ptr<const Executor>, const dim<2>&, int, array<value_type>, array<index_type>,`

`array<index_type>)`

```
393     {
394         return create(exec, size, block_size,
395                       array<value_type>{exec, std::move(values)},
396                       array<index_type>{exec, std::move(col_idx)},
397                       array<index_type>{exec, std::move(row_ptr)});
398     }
```

References `gko::matrix::Fbcsr< ValueType, IndexType >::create()`.

47.98.3.7 `create()` [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Fbcsr> gko::matrix::Fbcsr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size\_type num_nonzeros,
    int block_size ) [static]
```

Creates an uninitialized FBCSR matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_nonzeros</i>	number of stored nonzeros. It needs to be a multiple of <code>block_size * block_size</code> .
<i>block_size</i>	size of the small dense square blocks

Returns

A smart pointer to the newly created matrix.

47.98.3.8 create() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Fbcsr> gko::matrix::Fbcsr< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    int block_size = 1 ) [static]
```

Creates an uninitialized FBCSR matrix with the given block size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>block_size</i>	The desired size of the dense square nonzero blocks; defaults to 1.

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::matrix::Fbcsr< ValueType, IndexType >::create()`.

47.98.3.9 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const Fbcsr> gko::matrix::Fbcsr< ValueType, IndexType >::create_const
(
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    int blocksize,
    gko::detail::const_array_view< ValueType > && values,
    gko::detail::const_array_view< IndexType > && col_idxs,
    gko::detail::const_array_view< IndexType > && row_ptrs ) [static]
```

Creates a constant (immutable) [Fbcsr](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>blocksize</i>	the block size of the matrix
<i>values</i>	the value array of the matrix, stored in column-major order for each block
<i>col_idx</i> s	the block column index array of the matrix
<i>row_ptr</i> s	the block row pointer array of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

47.98.3.10 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Fbcsr< ValueType, IndexType >::extract_↵
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.98.3.11 get_block_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int gko::matrix::Fbcsr< ValueType, IndexType >::get_block_size ( ) const [inline], [noexcept]
```

Returns

The fixed block size for this matrix

47.98.3.12 get_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_col_idx
```

s () [inline], [noexcept]

Returns

The column indexes of the matrix.

References [gko::array< ValueType >::get_data\(\)](#).

47.98.3.13 get_const_col_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_col_idxes ( ) const
[inline], [noexcept]
```

Returns

The column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.98.3.14 get_const_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_row_ptrs ( ) const
[inline], [noexcept]
```

Returns

The row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.98.3.15 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns

The values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.98.3.16 get_num_block_cols()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_block_cols ( ) const [inline],
[noexcept]
```

Returns

The number of block-columns in the matrix

47.98.3.17 get_num_block_rows()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_block_rows ( ) const [inline],
[noexcept]
```

Returns

The number of block-rows in the matrix

47.98.3.18 get_num_stored_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_blocks ( ) const [inline],
[noexcept]
```

Returns

The number of non-zero blocks explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.98.3.19 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Fbcsr< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns

The number of elements explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.98.3.20 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_row_ptrs ( ) [inline], [noexcept]
```

Returns

The row pointers of the matrix.

References `gko::array< ValueType >::get_data()`.

47.98.3.21 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Fbcsr< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns

The values of the matrix.

References `gko::array< ValueType >::get_data()`.

47.98.3.22 is_sorted_by_column_index()

```
template<typename ValueType = default_precision, typename IndexType = int32>
bool gko::matrix::Fbcsr< ValueType, IndexType >::is_sorted_by_column_index ( ) const
```

Tests if all row entry pairs (value, col_idx) are sorted by column index.

Returns

True if all row entry pairs (value, col_idx) are sorted by column index

47.98.3.23 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Fbcsr& gko::matrix::Fbcsr< ValueType, IndexType >::operator= (
    const Fbcsr< ValueType, IndexType > & )
```

Copy-assigns an `Fbcsr` matrix.

Preserves the executor, copies data and block size from the input.

47.98.3.24 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Fbcsr& gko::matrix::Fbcsr< ValueType, IndexType >::operator= (
    Fbcsr< ValueType, IndexType > && )
```

Move-assigns an [Fbcsr](#) matrix.

Preserves the executor, moves the data over preserving size and stride. Leaves the moved-from object in an empty state (0x0 with no nonzeros, but valid row pointers).

47.98.3.25 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.98.3.26 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a [matrix_data](#) into [Fbcsr](#) format.

Requires the block size to be set beforehand

See also

[set_block_size](#).

Warning

Unlike [Csr::read](#), here explicit non-zeros are NOT dropped.

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.98.3.27 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.98.3.28 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::Fbcsr< ValueType, IndexType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.98.3.29 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Fbcsr< ValueType, IndexType >::write (
    mat\_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following files:

- [ginkgo/core/matrix/csr.hpp](#)
- [ginkgo/core/matrix/fbcsr.hpp](#)

47.99 gko::solver::Fcg< ValueType > Class Template Reference

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

```
#include <ginkgo/core/solver/fcg.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.99.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Fcg< ValueType >
```

FCG or the flexible conjugate gradient method is an iterative type Krylov subspace method which is suitable for symmetric positive definite methods.

Though this method performs very well for symmetric positive definite matrices, it is in general not suitable for general matrices.

In contrast to the standard CG based on the Polack-Ribiere formula, the flexible CG uses the Fletcher-Reeves formula for creating the orthonormal vectors spanning the Krylov subspace. This increases the computational cost of every Krylov solver iteration but allows for non-constant preconditioners.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of FCG are merged into 2 separate steps.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.99.2 Member Function Documentation

47.99.2.1 `apply_uses_initial_guess()`

```
template<typename ValueType = default_precision>
bool gko::solver::Fcg< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
72 { return true; }
```

47.99.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Fcg< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.99.2.3 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Fcg< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.99.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Fcg< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/fcg.hpp

47.100 gko::matrix::Fft Class Reference

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

```
#include <ginkgo/core/matrix/fft.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- void [write](#) ([matrix_data](#)< std::complex< float >, [int32](#) > &data) const override
Writes a matrix to a [matrix_data](#) structure.
- void [write](#) ([matrix_data](#)< std::complex< float >, [int64](#) > &data) const override
Writes a matrix to a [matrix_data](#) structure.
- void [write](#) ([matrix_data](#)< std::complex< double >, [int32](#) > &data) const override
Writes a matrix to a [matrix_data](#) structure.
- void [write](#) ([matrix_data](#)< std::complex< double >, [int64](#) > &data) const override
Writes a matrix to a [matrix_data](#) structure.

Static Public Member Functions

- static std::unique_ptr< [Fft](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec)
Creates an empty Fourier matrix.
- static std::unique_ptr< [Fft](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size=0, bool [inverse](#)=false)
Creates an Fourier matrix with the given dimensions.

47.100.1 Detailed Description

This LinOp implements a 1D Fourier matrix using the FFT algorithm.

It implements forward and inverse DFT.

For a power-of-two size n with corresponding root of unity $\omega = e^{-2\pi i/n}$ for forward DFT and $\omega = e^{2\pi i/n}$ for inverse DFT it computes

$$x_k = \sum_{j=0}^{n-1} \omega^{jk} b_j$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," Mathematics of Computation, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

47.100.2 Member Function Documentation

47.100.2.1 conj_transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.100.2.2 create() [1/2]

```
static std::unique_ptr<Fft> gko::matrix::Fft::create (
    std::shared_ptr< const Executor > exec ) [static]
```

Creates an empty Fourier matrix.

Parameters

<i>exec</i>	Executor associated to the matrix
-------------	---

Returns

A smart pointer to the newly created matrix.

47.100.2.3 create() [2/2]

```
static std::unique_ptr<Fft> gko::matrix::Fft::create (
    std::shared_ptr< const Executor > exec,
    size_type size = 0,
    bool inverse = false ) [static]
```

Creates an Fourier matrix with the given dimensions.

Parameters

<i>size</i>	size of the matrix
<i>inverse</i>	true to compute an inverse DFT instead of a normal DFT

Returns

A smart pointer to the newly created matrix.

47.100.2.4 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.100.2.5 write() [1/4]

```
void gko::matrix::Fft::write (
    matrix_data< std::complex< double >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int32 >](#).

47.100.2.6 write() [2/4]

```
void gko::matrix::Fft::write (
    matrix\_data< std::complex< double >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int64 >](#).

47.100.2.7 write() [3/4]

```
void gko::matrix::Fft::write (
    matrix\_data< std::complex< float >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int32 >](#).

47.100.2.8 write() [4/4]

```
void gko::matrix::Fft::write (
    matrix\_data< std::complex< float >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int64 >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/fft.hpp](#)

47.101 gko::matrix::Fft2 Class Reference

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

```
#include <ginkgo/core/matrix/fft.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `void write (matrix_data< std::complex< float >, int32 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< float >, int64 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< double >, int32 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< double >, int64 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.

Static Public Member Functions

- `static std::unique_ptr< Fft2 > create (std::shared_ptr< const Executor > exec)`
Creates an empty Fourier matrix.
- `static std::unique_ptr< Fft2 > create (std::shared_ptr< const Executor > exec, size_type size)`
Creates an Fourier matrix with the given dimensions.
- `static std::unique_ptr< Fft2 > create (std::shared_ptr< const Executor > exec, size_type size1, size_type size2, bool inverse=false)`
Creates an Fourier matrix with the given dimensions.

47.101.1 Detailed Description

This LinOp implements a 2D Fourier matrix using the FFT algorithm.

For indexing purposes, the first dimension is the major axis.

It implements complex-to-complex forward and inverse FFT.

For a power-of-two sizes n_1, n_2 with corresponding root of unity $\omega = e^{-2\pi i/(n_1 n_2)}$ for forward DFT and $\omega = e^{2\pi i/(n_1 n_2)}$ for inverse DFT it computes

$$x_{k_1 n_2 + k_2} = \sum_{i_1=0}^{n_1-1} \sum_{i_2=0}^{n_2-1} \omega^{i_1 k_1 + i_2 k_2} b_{i_1 n_2 + i_2}$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," *Mathematics of Computation*, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

47.101.2 Member Function Documentation

47.101.2.1 conj_transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft2::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.101.2.2 create() [1/3]

```
static std::unique_ptr<Fft2> gko::matrix::Fft2::create (
    std::shared_ptr< const Executor > exec ) [static]
```

Creates an empty Fourier matrix.

Parameters

<i>exec</i>	Executor associated to the matrix
-------------	---

Returns

A smart pointer to the newly created matrix.

47.101.2.3 create() [2/3]

```
static std::unique_ptr<Fft2> gko::matrix::Fft2::create (
    std::shared_ptr< const Executor > exec,
    size_type size ) [static]
```

Creates an Fourier matrix with the given dimensions.

Parameters

<i>size</i>	size of both FFT dimensions
-------------	-----------------------------

Returns

A smart pointer to the newly created matrix.

47.101.2.4 create() [3/3]

```
static std::unique_ptr<Fft2> gko::matrix::Fft2::create (
    std::shared_ptr< const Executor > exec,
    size_type size1,
    size_type size2,
    bool inverse = false ) [static]
```

Creates an Fourier matrix with the given dimensions.

Parameters

<i>size1</i>	size of the first FFT dimension
<i>size2</i>	size of the second FFT dimension
<i>inverse</i>	true to compute an inverse DFT instead of a normal DFT

Returns

A smart pointer to the newly created matrix.

47.101.2.5 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft2::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.101.2.6 write() [1/4]

```
void gko::matrix::Fft2::write (
    matrix_data< std::complex< double >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int32 >](#).

47.101.2.7 write() [2/4]

```
void gko::matrix::Fft2::write (
    matrix\_data< std::complex< double >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int64 >](#).

47.101.2.8 write() [3/4]

```
void gko::matrix::Fft2::write (
    matrix\_data< std::complex< float >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int32 >](#).

47.101.2.9 write() [4/4]

```
void gko::matrix::Fft2::write (
    matrix\_data< std::complex< float >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int64 >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/fft.hpp](#)

47.102 gko::matrix::Fft3 Class Reference

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

```
#include <ginkgo/core/matrix/fft.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `void write (matrix_data< std::complex< float >, int32 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< float >, int64 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< double >, int32 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.
- `void write (matrix_data< std::complex< double >, int64 > &data)` const override
Writes a matrix to a [matrix_data](#) structure.

Static Public Member Functions

- `static std::unique_ptr< Fft3 > create (std::shared_ptr< const Executor > exec)`
Creates an empty Fourier matrix.
- `static std::unique_ptr< Fft3 > create (std::shared_ptr< const Executor > exec, size_type size)`
Creates an Fourier matrix with the given dimensions.
- `static std::unique_ptr< Fft3 > create (std::shared_ptr< const Executor > exec, size_type size1, size_type size2, size_type size3, bool inverse=false)`
Creates an Fourier matrix with the given dimensions.

47.102.1 Detailed Description

This LinOp implements a 3D Fourier matrix using the FFT algorithm.

For indexing purposes, the first dimension is the major axis.

It implements complex-to-complex forward and inverse FFT.

For a power-of-two sizes n_1, n_2, n_3 with corresponding root of unity $\omega = e^{-2\pi i/(n_1 n_2 n_3)}$ for forward DFT and $\omega = e^{2\pi i/(n_1 n_2 n_3)}$ for inverse DFT it computes

$$x_{k_1 n_2 n_3 + k_2 n_3 + k_3} = \sum_{i_1=0}^{n_1-1} \sum_{i_2=0}^{n_2-1} \sum_{i_3=0}^{n_3-1} \omega^{i_1 k_1 + i_2 k_2 + i_3 k_3} b_{i_1 n_2 n_3 + i_2 n_3 + i_3}$$

without normalization factors.

The Reference and OpenMP implementations support only power-of-two input sizes, as they use the Radix-2 algorithm by J. W. Cooley and J. W. Tukey, "An Algorithm for the Machine Calculation of Complex Fourier Series," *Mathematics of Computation*, vol. 19, no. 90, pp. 297–301, 1965, doi: 10.2307/2003354. The CUDA and HIP implementations use cuSPARSE/hipSPARSE with full support for non-power-of-two input sizes and special optimizations for products of small prime powers.

47.102.2 Member Function Documentation

47.102.2.1 conj_transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft3::conj_transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.102.2.2 create() [1/3]

```
static std::unique_ptr<Fft3> gko::matrix::Fft3::create (
    std::shared_ptr< const Executor > exec ) [static]
```

Creates an empty Fourier matrix.

Parameters

<i>exec</i>	Executor associated to the matrix
-------------	---

Returns

A smart pointer to the newly created matrix.

47.102.2.3 create() [2/3]

```
static std::unique_ptr<Fft3> gko::matrix::Fft3::create (
    std::shared_ptr< const Executor > exec,
    size_type size ) [static]
```

Creates an Fourier matrix with the given dimensions.

Parameters

<i>size</i>	size of all FFT dimensions
-------------	----------------------------

Returns

A smart pointer to the newly created matrix.

47.102.2.4 create() [3/3]

```
static std::unique_ptr<Fft3> gko::matrix::Fft3::create (
    std::shared_ptr< const Executor > exec,
    size_type size1,
    size_type size2,
    size_type size3,
    bool inverse = false ) [static]
```

Creates an Fourier matrix with the given dimensions.

Parameters

<i>size1</i>	size of the first FFT dimension
<i>size2</i>	size of the second FFT dimension
<i>size3</i>	size of the third FFT dimension
<i>inverse</i>	true to compute an inverse DFT instead of a normal DFT

Returns

A smart pointer to the newly created matrix.

47.102.2.5 transpose()

```
std::unique_ptr<LinOp> gko::matrix::Fft3::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.102.2.6 write() [1/4]

```
void gko::matrix::Fft3::write (
    matrix_data< std::complex< double >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int32 >](#).

47.102.2.7 write() [2/4]

```
void gko::matrix::Fft3::write (
    matrix\_data< std::complex< double >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< double >, int64 >](#).

47.102.2.8 write() [3/4]

```
void gko::matrix::Fft3::write (
    matrix\_data< std::complex< float >, int32 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int32 >](#).

47.102.2.9 write() [4/4]

```
void gko::matrix::Fft3::write (
    matrix\_data< std::complex< float >, int64 > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< std::complex< float >, int64 >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/fft.hpp](#)

47.103 [gko::multigrid::FixedCoarsening< ValueType, IndexType >](#) Class Template Reference

[FixedCoarsening](#) is a very simple coarse grid generation algorithm.

```
#include <ginkgo/core/multigrid/fixed_coarsening.hpp>
```

Public Member Functions

- `std::shared_ptr< const LinOp > get\_system\_matrix () const`
Returns the system operator (matrix) of the linear system.

47.103.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::multigrid::FixedCoarsening< ValueType, IndexType >
```

[FixedCoarsening](#) is a very simple coarse grid generation algorithm.

It selects the coarse matrix from the fine matrix by with user-specified indices.

The user needs to specify the indices (with global numbering) of the fine matrix, they wish to be in the coarse matrix. The restriction and prolongation matrices will map to and from the coarse space without any interpolation or weighting.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.103.2 Member Function Documentation

47.103.2.1 [get_system_matrix\(\)](#)

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const LinOp> gko::multigrid::FixedCoarsening< ValueType, IndexType >::get\_↵
system\_matrix ( ) const [inline]
```

Returns the system operator (matrix) of the linear system.

Returns

the system operator (matrix)

```
58     {  
59         return system_matrix_;  
60     }
```

The documentation for this class was generated from the following file:

- ginkgo/core/multigrid/fixed_coarsening.hpp

47.104 gko::preconditioner::GaussSeidel< ValueType, IndexType > Class Template Reference

This class generates the Gauss-Seidel preconditioner.

```
#include <ginkgo/core/preconditioner/gauss_seidel.hpp>
```

Public Member Functions

- const parameters_type & [get_parameters](#) ()
Returns the parameters used to construct the factory.
- const parameters_type & [get_parameters](#) () const
Returns the parameters used to construct the factory.
- std::unique_ptr< [composition_type](#) > [generate](#) (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

- static parameters_type [build](#) ()
Creates a new parameter_type to set up the factory.

47.104.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>  
class gko::preconditioner::GaussSeidel< ValueType, IndexType >
```

This class generates the Gauss-Seidel preconditioner.

This is the special case of the relaxation factor $\omega = 1$ of the (S)SOR preconditioner.

See also

[Sor](#)

47.104.2 Member Function Documentation

47.104.2.1 generate()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<composition_type> gko::preconditioner::GaussSeidel< ValueType, IndexType >↵
::generate (
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a Factorization instead of a generic LinOp, which would need to be cast to Factorization again to access its factors. It is only necessary because smart pointers aren't covariant.

47.104.2.2 get_parameters() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const parameters_type& gko::preconditioner::GaussSeidel< ValueType, IndexType >::get_parameters
( ) [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
70 { return parameters_; }
```

47.104.2.3 get_parameters() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const parameters_type& gko::preconditioner::GaussSeidel< ValueType, IndexType >::get_parameters
( ) const [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

The documentation for this class was generated from the following file:

- ginkgo/core/preconditioner/gauss_seidel.hpp

47.105 gko::solver::Gcr< ValueType > Class Template Reference

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

```
#include <ginkgo/core/solver/gcr.hpp>
```


Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.
- `size_type get_krylov_dim ()` const
Gets the Krylov dimension of the solver.
- `void set_krylov_dim (size_type other)`
Sets the Krylov dimension.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.105.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Gcr< ValueType >
```

GCR or the generalized conjugate residual method is an iterative type Krylov subspace method similar to GMRES which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of GCR are merged into one step. Modified Gram-Schmidt is used.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.105.2 Member Function Documentation

47.105.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Gcr< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
68 { return true; }
```

47.105.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gcr< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.105.2.3 get_krylov_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Gcr< ValueType >::get_krylov_dim ( ) const [inline]
```

Gets the Krylov dimension of the solver.

Returns

the Krylov dimension

47.105.2.4 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Gcr< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.105.2.5 set_krylov_dim()

```
template<typename ValueType = default_precision>
void gko::solver::Gcr< ValueType >::set_krylov_dim (
    size_type other ) [inline]
```

Sets the Krylov dimension.

Parameters

<i>other</i>	the new Krylov dimension
--------------	--------------------------

47.105.2.6 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gcr< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/gcr.hpp

47.106 gko::solver::Gmres< ValueType > Class Template Reference

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

```
#include <ginkgo/core/solver/gmres.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.
- `size_type get_krylov_dim ()` const
Gets the Krylov dimension of the solver.
- `void set_krylov_dim (size_type other)`
Sets the Krylov dimension.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.106.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Gmres< ValueType >
```

GMRES or the generalized minimal residual method is an iterative type Krylov subspace method which is suitable for nonsymmetric linear systems.

The implementation in Ginkgo makes use of the merged kernel to make the best use of data locality. The inner operations in one iteration of GMRES are merged into 2 separate steps. Modified Gram-Schmidt is used.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.106.2 Member Function Documentation

47.106.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::Gmres< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
92 { return true; }
```

47.106.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gmres< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.106.2.3 get_krylov_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Gmres< ValueType >::get_krylov_dim ( ) const [inline]
```

Gets the Krylov dimension of the solver.

Returns

the Krylov dimension

47.106.2.4 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Gmres< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.106.2.5 set_krylov_dim()

```
template<typename ValueType = default_precision>
void gko::solver::Gmres< ValueType >::set_krylov_dim (
    size_type other ) [inline]
```

Sets the Krylov dimension.

Parameters

<i>other</i>	the new Krylov dimension
--------------	--------------------------

47.106.2.6 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Gmres< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/gmres.hpp

47.107 gko::half Class Reference

A class providing basic support for half precision floating point types.

```
#include <ginkgo/core/base/half.hpp>
```

47.107.1 Detailed Description

A class providing basic support for half precision floating point types.

For now the only features are reduced storage compared to single precision and conversions from and to single precision floating point type.

The documentation for this class was generated from the following file:

- ginkgo/core/base/half.hpp

47.108 gko::solver::has_with_criteria< SolverType, typename > Struct Template Reference

Helper structure to test if the Factory of SolverType has a function `with_criteria`.

```
#include <ginkgo/core/solver/solver_traits.hpp>
```

47.108.1 Detailed Description

```
template<typename SolverType, typename = void>
struct gko::solver::has_with_criteria< SolverType, typename >
```

Helper structure to test if the Factory of SolverType has a function `with_criteria`.

Contains a constexpr boolean `value`, which is true if the Factory class of SolverType has a `with_criteria`, and false otherwise.

Template Parameters

<i>SolverType</i>	Solver to test if its factory has a <code>with_criteria</code> function.
-------------------	--

The documentation for this struct was generated from the following file:

- ginkgo/core/solver/solver_traits.hpp

47.109 gko::hip_stream Class Reference

An RAII wrapper for a custom HIP stream.

```
#include <ginkgo/core/base/stream.hpp>
```

Public Member Functions

- [hip_stream](#) ()
Creates an empty stream wrapper, representing the default stream.
- [hip_stream](#) (int device_id)
Creates a new custom HIP stream on the given device.
- [~hip_stream](#) ()
Destroys the custom HIP stream, if it isn't empty.
- [hip_stream](#) (hip_stream &&)
Move-constructs from an existing stream, which will be emptied.
- [hip_stream](#) & operator= (hip_stream &&)=delete
Move-assigns from an existing stream, which will be emptied.
- CUstream_st * [get](#) () const
Returns the native HIP stream handle.

47.109.1 Detailed Description

An RAII wrapper for a custom HIP stream.

The stream will be created on construction and destroyed when the lifetime of the wrapper ends.

47.109.2 Constructor & Destructor Documentation

47.109.2.1 hip_stream()

```
gko::hip_stream::hip_stream (
    int device_id )
```

Creates a new custom HIP stream on the given device.

Parameters

<i>device_id</i>	the device ID to create the stream on.
------------------	--

47.109.3 Member Function Documentation

47.109.3.1 get()

```
CUstream_st* gko::hip_stream::get ( ) const
```


Returns the native HIP stream handle.

In an empty [hip_stream](#), this will return nullptr.

The documentation for this class was generated from the following file:

- ginkgo/core/base/stream.hpp

47.110 gko::HipAllocatorBase Class Reference

Implement this interface to provide an allocator for [HipExecutor](#).

```
#include <ginkgo/core/base/memory.hpp>
```

47.110.1 Detailed Description

Implement this interface to provide an allocator for [HipExecutor](#).

The documentation for this class was generated from the following file:

- ginkgo/core/base/memory.hpp

47.111 gko::HipblasError Class Reference

[HipblasError](#) is thrown when a hipBLAS routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [HipblasError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a hipBLAS error.

47.111.1 Detailed Description

[HipblasError](#) is thrown when a hipBLAS routine throws a non-zero error code.

47.111.2 Constructor & Destructor Documentation

47.111.2.1 HipblasError()

```
gko::HipblasError::HipblasError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a hipBLAS error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the hipBLAS routine that failed
<i>error_code</i>	The resulting hipBLAS error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.112 gko::HipError Class Reference

[HipError](#) is thrown when a HIP routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [HipError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a HIP error.

47.112.1 Detailed Description

[HipError](#) is thrown when a HIP routine throws a non-zero error code.

47.112.2 Constructor & Destructor Documentation

47.112.2.1 HipError()

```
gko::HipError::HipError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a HIP error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the HIP routine that failed
<i>error_code</i>	The resulting HIP error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.113 gko::HipExecutor Class Reference

This is the [Executor](#) subclass which represents the HIP enhanced device.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- `std::shared_ptr< Executor > get_master ()` noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `std::shared_ptr< const Executor > get_master ()` const noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `void synchronize ()` const override
Synchronize the operations launched on the executor with its master.
- `std::string get_description ()` const override
- `int get_device_id ()` const noexcept
Get the HIP device id of the device associated to this executor.
- `int get_num_warps_per_sm ()` const noexcept
Get the number of warps per SM of this executor.
- `int get_num_multiprocessor ()` const noexcept
Get the number of multiprocessor of this executor.
- `int get_major_version ()` const noexcept
Get the major version of compute capability.
- `int get_minor_version ()` const noexcept
Get the minor version of compute capability.
- `int get_num_warps ()` const noexcept
Get the number of warps of this executor.
- `int get_warp_size ()` const noexcept
Get the warp size of this executor.
- `hipblasContext * get_hipblas_handle ()` const
Get the hipblas handle for this executor.
- `hipblasContext * get_blas_handle ()` const
Get the hipblas handle for this executor.
- `hipsparseContext * get_hipsparse_handle ()` const
Get the hipsparse handle for this executor.
- `hipsparseContext * get_sparselib_handle ()` const
Get the hipsparse handle for this executor.
- `int get_closest_numa ()` const
Get the closest NUMA node.
- `std::vector< int > get_closest_pus ()` const
Get the closest PUs.
- `virtual void run (const Operation &op)` const=0
Runs the specified [Operation](#) using this [Executor](#).

- `template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >`
`void run (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.
- `template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename Closure↵`
`Dpcpp >`
`void run (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↵`
`Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.

Static Public Member Functions

- `static std::shared_ptr< HipExecutor > create (int device_id, std::shared_ptr< Executor > master, bool device_reset, allocation\_mode alloc_mode=default_hip_alloc_mode, CUstream_st *stream=nullptr)`
Creates a new [HipExecutor](#).
- `static int get_num_devices ()`
Get the number of devices present on the system.

47.113.1 Detailed Description

This is the [Executor](#) subclass which represents the HIP enhanced device.

47.113.2 Member Function Documentation

47.113.2.1 create()

```
static std::shared_ptr<HipExecutor> gko::HipExecutor::create (
    int device_id,
    std::shared_ptr< Executor > master,
    bool device_reset,
    allocation\_mode alloc_mode = default_hip_alloc_mode,
    CUstream_st * stream = nullptr ) [static]
```

Creates a new [HipExecutor](#).

Parameters

<i>device_id</i>	the HIP device id of this device
<i>master</i>	an executor on the host that is used to invoke the device kernels
<i>device_reset</i>	whether to reset the device after the object exits the scope.
<i>alloc_mode</i>	the allocation mode that the executor should operate on. See @allocation_mode for more details

47.113.2.2 get_blas_handle()

```
hipblasContext* gko::HipExecutor::get_blas_handle ( ) const [inline]
```

Get the hipblas handle for this executor.

Returns

the hipblas handle (hipblasContext*) for this executor

Referenced by `get_hipblas_handle()`.

47.113.2.3 get_closest_numa()

```
int gko::HipExecutor::get_closest_numa ( ) const [inline]
```

Get the closest NUMA node.

Returns

the closest NUMA node closest to this device

47.113.2.4 get_closest_pus()

```
std::vector<int> gko::HipExecutor::get_closest_pus ( ) const [inline]
```

Get the closest PUs.

Returns

the array of PUs closest to this device

47.113.2.5 get_description()

```
std::string gko::HipExecutor::get_description ( ) const [override], [virtual]
```

Returns

a textual representation of the executor and its device.

Implements [gko::Executor](#).

47.113.2.6 `get_hipblas_handle()`

```
hipblasContext* gko::HipExecutor::get_hipblas_handle ( ) const [inline]
```

Get the hipblas handle for this executor.

Returns

the hipblas handle (`hipblasContext*`) for this executor

References `get_blas_handle()`.

47.113.2.7 `get_hipsparse_handle()`

```
hipsparseContext* gko::HipExecutor::get_hipsparse_handle ( ) const [inline]
```

Get the hipsparse handle for this executor.

Returns

the hipsparse handle (`hipsparseContext*`) for this executor

References `get_sparselib_handle()`.

47.113.2.8 `get_master()` [1/2]

```
std::shared_ptr<const Executor> gko::HipExecutor::get_master ( ) const [override], [virtual],  
[noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.113.2.9 `get_master()` [2/2]

```
std::shared_ptr<Executor> gko::HipExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.113.2.10 get_sparselib_handle()

```
hipsparseContext* gko::HipExecutor::get_sparselib_handle ( ) const [inline]
```

Get the hipsparse handle for this executor.

Returns

the hipsparse handle (hipsparseContext*) for this executor

Referenced by `get_hipsparse_handle()`.

47.113.2.11 run() [1/3]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

47.113.2.12 run() [2/3]

```
virtual void gko::Executor::run
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

47.113.2.13 run() [3/3]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureReference ,
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of op_ref
<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.114 gko::HipfftError Class Reference

[HipfftError](#) is thrown when a hipFFT routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```


Public Member Functions

- [HipfftError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a hipFFT error.

47.114.1 Detailed Description

[HipfftError](#) is thrown when a hipFFT routine throws a non-zero error code.

47.114.2 Constructor & Destructor Documentation

47.114.2.1 HipfftError()

```
gko::HipfftError::HipfftError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a hipFFT error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the hipFFT routine that failed
<i>error_code</i>	The resulting hipFFT error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.115 gko::HiprandError Class Reference

[HiprandError](#) is thrown when a hipRAND routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [HiprandError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a hipRAND error.

47.115.1 Detailed Description

[HiprandError](#) is thrown when a hipRAND routine throws a non-zero error code.

47.115.2 Constructor & Destructor Documentation

47.115.2.1 HiprandError()

```
gko::HiprandError::HiprandError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a hipRAND error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the hipRAND routine that failed
<i>error_code</i>	The resulting hipRAND error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.116 gko::HipsparseError Class Reference

[HipsparseError](#) is thrown when a hipSPARSE routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [HipsparseError](#) (const std::string &file, int line, const std::string &func, int64 error_code)
Initializes a hipSPARSE error.

47.116.1 Detailed Description

[HipsparseError](#) is thrown when a hipSPARSE routine throws a non-zero error code.

47.116.2 Constructor & Destructor Documentation

47.116.2.1 HipsparseError()

```
gko::HipsparseError::HipsparseError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a hipSPARSE error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the hipSPARSE routine that failed
<i>error_code</i>	The resulting hipSPARSE error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.117 gko::HipTimer Class Reference

A timer using events for timing on a [HipExecutor](#).

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- void [record](#) ([time_point](#) &time) override
Records a time point at the current time.
- void [wait](#) ([time_point](#) &time) override
Waits until all kernels in-process when recording the time point are finished.
- std::chrono::nanoseconds [difference_async](#) (const [time_point](#) &start, const [time_point](#) &stop) override
Computes the difference between the two time points in nanoseconds.

Additional Inherited Members

47.117.1 Detailed Description

A timer using events for timing on a [HipExecutor](#).

Note

When using a [HipExecutor](#) with a custom stream, make sure that the stream's lifetime is longer than the lifetime of this timer.

47.117.2 Member Function Documentation

47.117.2.1 `difference_async()`

```
std::chrono::nanoseconds gko::HipTimer::difference_async (
    const time\_point & start,
    const time\_point & stop ) [override], [virtual]
```

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. `timer->wait(stop)` needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implements [gko::Timer](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/timer.hpp`

47.118 `gko::matrix::Hybrid< ValueType, IndexType >` Class Template Reference

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Classes

- class [automatic](#)
automatic is a [strategy_type](#) which decides the number of stored elements per row of the ell part automatically.
- class [column_limit](#)
column_limit is a [strategy_type](#) which decides the number of stored elements per row of the ell part by specifying the number of columns.
- class [imbalance_bounded_limit](#)
imbalance_bounded_limit is a [strategy_type](#) which decides the number of stored elements per row of the ell part.
- class [imbalance_limit](#)
imbalance_limit is a [strategy_type](#) which decides the number of stored elements per row of the ell part according to the percent.
- class [minimal_storage_limit](#)
minimal_storage_limit is a [strategy_type](#) which decides the number of stored elements per row of the ell part.
- class [strategy_type](#)
strategy_type is to decide how to set the hybrid config.

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< [Diagonal](#)< ValueType > > [extract_diagonal](#) () const override
Extracts the diagonal entries of the matrix into a vector.
- std::unique_ptr< [absolute_type](#) > [compute_absolute](#) () const override
Gets the AbsoluteLinOp.
- void [compute_absolute_inplace](#) () override
Compute absolute inplace on each element.
- value_type * [get_ell_values](#) () noexcept
Returns the values of the ell part.
- const value_type * [get_const_ell_values](#) () const noexcept
Returns the values of the ell part.
- index_type * [get_ell_col_idxs](#) () noexcept
Returns the column indexes of the ell part.
- const index_type * [get_const_ell_col_idxs](#) () const noexcept
Returns the column indexes of the ell part.
- size_type [get_ell_num_stored_elements_per_row](#) () const noexcept
Returns the number of stored elements per row of ell part.
- size_type [get_ell_stride](#) () const noexcept
Returns the stride of the ell part.
- size_type [get_ell_num_stored_elements](#) () const noexcept
Returns the number of elements explicitly stored in the ell part.
- value_type & [ell_val_at](#) (size_type row, size_type idx) noexcept
*Returns the *idx*-th non-zero element of the *row*-th row in the ell part.*
- value_type [ell_val_at](#) (size_type row, size_type idx) const noexcept
*Returns the *idx*-th non-zero element of the *row*-th row in the ell part.*
- index_type & [ell_col_at](#) (size_type row, size_type idx) noexcept
*Returns the *idx*-th column index of the *row*-th row in the ell part.*
- index_type [ell_col_at](#) (size_type row, size_type idx) const noexcept
*Returns the *idx*-th column index of the *row*-th row in the ell part.*
- const ell_type * [get_ell](#) () const noexcept
Returns the matrix of the ell part.
- value_type * [get_coo_values](#) () noexcept
Returns the values of the coo part.
- const value_type * [get_const_coo_values](#) () const noexcept
Returns the values of the coo part.
- index_type * [get_coo_col_idxs](#) () noexcept
Returns the column indexes of the coo part.
- const index_type * [get_const_coo_col_idxs](#) () const noexcept
Returns the column indexes of the coo part.
- index_type * [get_coo_row_idxs](#) () noexcept
Returns the row indexes of the coo part.
- const index_type * [get_const_coo_row_idxs](#) () const noexcept

- Returns the row indexes of the coo part.*
- `size_type get_coo_num_stored_elements ()` const noexcept
Returns the number of elements explicitly stored in the coo part.
- `const coo_type * get_coo ()` const noexcept
Returns the matrix of the coo part.
- `size_type get_num_stored_elements ()` const noexcept
Returns the number of elements explicitly stored in the matrix.
- `std::shared_ptr< strategy_type > get_strategy ()` const noexcept
Returns the strategy.
- `template<typename HybType > std::shared_ptr< typename HybType::strategy_type > get_strategy ()` const
Returns the current strategy allowed in given hybrid format.
- `Hybrid & operator= (const Hybrid &)`
Copy-assigns a Hybrid matrix.
- `Hybrid & operator= (Hybrid &&)`
Move-assigns a Hybrid matrix.
- `Hybrid (const Hybrid &)`
Copy-assigns a Hybrid matrix.
- `Hybrid (Hybrid &&)`
Move-assigns a Hybrid matrix.

Static Public Member Functions

- `static std::unique_ptr< Hybrid > create (std::shared_ptr< const Executor > exec, std::shared_ptr< strategy_type > strategy=std::make_shared< automatic >())`
Creates an uninitialized Hybrid matrix of specified method.
- `static std::unique_ptr< Hybrid > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, std::shared_ptr< strategy_type > strategy=std::make_shared< automatic >())`
Creates an uninitialized Hybrid matrix of the specified size and method.
- `static std::unique_ptr< Hybrid > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, size_type num_stored_elements_per_row, std::shared_ptr< strategy_type > strategy=std::make_shared< automatic >())`
Creates an uninitialized Hybrid matrix of the specified size and method.
- `static std::unique_ptr< Hybrid > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, size_type num_stored_elements_per_row, size_type stride, std::shared_ptr< strategy_type > strategy)`
Creates an uninitialized Hybrid matrix of the specified size and method.
- `static std::unique_ptr< Hybrid > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, size_type num_stored_elements_per_row, size_type stride, size_type num_nonzeros={}, std::shared_ptr< strategy_type > strategy=std::make_shared< automatic >())`
Creates an uninitialized Hybrid matrix of the specified size and method.

47.118.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >
```

HYBRID is a matrix format which splits the matrix into ELLPACK and COO format.

Achieve the excellent performance with a proper partition of ELLPACK and COO.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.118.2 Constructor & Destructor Documentation**47.118.2.1 Hybrid() [1/2]**

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Hybrid< ValueType, IndexType >::Hybrid (
    const Hybrid< ValueType, IndexType > & )
```

Copy-assigns a [Hybrid](#) matrix.

Inherits the executor, copies the [EII](#) and [Coo](#) matrices.

47.118.2.2 Hybrid() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Hybrid< ValueType, IndexType >::Hybrid (
    Hybrid< ValueType, IndexType > && )
```

Move-assigns a [Hybrid](#) matrix.

Inherits the executor, moves the [EII](#) and [Coo](#) matrices. The moved-from matrix is empty (0x0 with empty [EII](#)/[Coo](#) matrices).

47.118.3 Member Function Documentation**47.118.3.1 compute_absolute()**

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Hybrid< ValueType, IndexType >::compute_absolute
( ) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Hybrid< ValueType, IndexType > > >](#).

47.118.3.2 create() [1/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Hybrid> gko::matrix::Hybrid< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size_type num_stored_elements_per_row,
    size_type stride,
    size_type num_nonzeros = {},
    std::shared_ptr< strategy_type > strategy = std::make_shared< automatic >() )
[static]
```

Creates an uninitialized [Hybrid](#) matrix of the specified size and method.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_stored_elements_per_row</i>	the number of stored elements per row
<i>stride</i>	stride of the rows
<i>num_nonzeros</i>	number of nonzeros
<i>strategy</i>	strategy of deciding the Hybrid config

Returns

A smart pointer to the newly created matrix.

47.118.3.3 create() [2/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Hybrid> gko::matrix::Hybrid< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size_type num_stored_elements_per_row,
    size_type stride,
    std::shared_ptr< strategy_type > strategy ) [static]
```

Creates an uninitialized [Hybrid](#) matrix of the specified size and method.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_stored_elements_per_row</i>	the number of stored elements per row
<i>stride</i>	stride of the rows
<i>strategy</i>	strategy of deciding the Hybrid config

Returns

A smart pointer to the newly created matrix.

47.118.3.4 create() [3/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Hybrid> gko::matrix::Hybrid< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size_type num_stored_elements_per_row,
    std::shared_ptr< strategy_type > strategy = std::make_shared< automatic >() )
[static]
```

Creates an uninitialized [Hybrid](#) matrix of the specified size and method.

(ell_stride is set to the number of rows of the matrix.)

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>num_stored_elements_per_row</i>	the number of stroed elements per row
<i>strategy</i>	strategy of deciding the Hybrid config

Returns

A smart pointer to the newly created matrix.

47.118.3.5 create() [4/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Hybrid> gko::matrix::Hybrid< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::shared_ptr< strategy_type > strategy = std::make_shared< automatic >() )
[static]
```

Creates an uninitialized [Hybrid](#) matrix of the specified size and method.

(ell_num_stored_elements_per_row is set to the number of cols of the matrix. ell_stride is set to the number of rows of the matrix.)

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>strategy</i>	strategy of deciding the Hybrid config

Returns

A smart pointer to the newly created matrix.

47.118.3.6 create() [5/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Hybrid> gko::matrix::Hybrid< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< strategy_type > strategy = std::make_shared< automatic >() )
[static]
```

Creates an uninitialized [Hybrid](#) matrix of specified method.

(ell_num_stored_elements_per_row is set to the number of cols of the matrix. ell_stride is set to the number of rows of the matrix.)

Parameters

<i>exec</i>	Executor associated to the matrix
<i>strategy</i>	strategy of deciding the Hybrid config

Returns

A smart pointer to the newly created matrix.

47.118.3.7 ell_col_at() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Hybrid< ValueType, IndexType >::ell_col_at (
    size_type row,
    size_type idx ) const [inline], [noexcept]
```

Returns the *idx*-th column index of the *row*-th row in the ell part.

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <i>idx</i> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

```
514 {
515     return ell->col_at(row, idx);
516 }
```

47.118.3.8 ell_col_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type& gko::matrix::Hybrid< ValueType, IndexType >::ell_col_at (
    size_type row,
    size_type idx ) [inline], [noexcept]
```

Returns the `idx`-th column index of the `row`-th row in the ell part.

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <code>idx</code> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.118.3.9 ell_val_at() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type gko::matrix::Hybrid< ValueType, IndexType >::ell_val_at (
    size_type row,
    size_type idx ) const [inline], [noexcept]
```

Returns the `idx`-th non-zero element of the `row`-th row in the ell part.

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the <code>idx</code> -th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.118.3.10 ell_val_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type& gko::matrix::Hybrid< ValueType, IndexType >::ell_val_at (
    size_type row,
    size_type idx ) [inline], [noexcept]
```

Returns the `idx`-th non-zero element of the `row`-th row in the ell part.

Parameters

<i>row</i>	the row of the requested element
<i>idx</i>	the idx-th stored element of the row

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.118.3.11 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Hybrid< ValueType, IndexType >::extract↵
_diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.118.3.12 get_const_coo_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_col_idxs ( )
const [inline], [noexcept]
```

Returns the column indexes of the coo part.

Returns

the column indexes of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.118.3.13 get_const_coo_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_row_idxs ( )
const [inline], [noexcept]
```

Returns the row indexes of the coo part.

Returns

the row indexes of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.118.3.14 get_const_coo_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_coo_values ( ) const
[inline], [noexcept]
```

Returns the values of the coo part.

Returns

the values of the coo part.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.118.3.15 get_const_ell_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_ell_col_idxs ( )
const [inline], [noexcept]
```

Returns the column indexes of the ell part.

Returns

the column indexes of the ell part

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.118.3.16 get_const_ell_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_const_ell_values ( ) const
[inline], [noexcept]
```

Returns the values of the ell part.

Returns

the values of the ell part

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.118.3.17 get_coo()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const coo_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo ( ) const [inline],
[noexcept]
```

Returns the matrix of the coo part.

Returns

the matrix of the coo part

47.118.3.18 get_coo_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_col_idxs ( ) [inline], [noexcept]
```

Returns the column indexes of the coo part.

Returns

the column indexes of the coo part.

47.118.3.19 get_coo_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_coo_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the coo part.

Returns

the number of elements explicitly stored in the coo part

47.118.3.20 get_coo_row_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_row_idxs ( ) [inline], [noexcept]
```

Returns the row indexes of the coo part.

Returns

the row indexes of the coo part.

47.118.3.21 get_coo_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_coo_values ( ) [inline], [noexcept]
```

Returns the values of the coo part.

Returns

the values of the coo part.

47.118.3.22 get_ell()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const ell_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell ( ) const [inline],
[noexcept]
```

Returns the matrix of the ell part.

Returns

the matrix of the ell part

47.118.3.23 get_ell_col_idxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell_col_idxs ( ) [inline], [noexcept]
```

Returns the column indexes of the ell part.

Returns

the column indexes of the ell part

47.118.3.24 get_ell_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the ell part.

Returns

the number of elements explicitly stored in the ell part

47.118.3.25 get_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_num_stored_elements_per_row ( )
const [inline], [noexcept]
```

Returns the number of stored elements per row of ell part.

Returns

the number of stored elements per row of ell part

47.118.3.26 get_ell_stride()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_ell_stride ( ) const [inline],
[noexcept]
```

Returns the stride of the ell part.

Returns

the stride of the ell part

47.118.3.27 get_ell_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Hybrid< ValueType, IndexType >::get_ell_values ( ) [inline], [noexcept]
```

Returns the values of the ell part.

Returns

the values of the ell part

47.118.3.28 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

47.118.3.29 get_strategy() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename HybType >
std::shared_ptr<typename HybType::strategy_type> gko::matrix::Hybrid< ValueType, IndexType
>::get_strategy ( ) const
```

Returns the current strategy allowed in given hybrid format.

Template Parameters

<i>HybType</i>	hybrid type
----------------	-------------

Returns

the strategy

47.118.3.30 get_strategy() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr< typename HybType::strategy_type > gko::matrix::Hybrid< ValueType, IndexType
>::get_strategy ( ) const [inline], [noexcept]
```

Returns the strategy.

Returns

the strategy

47.118.3.31 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Hybrid& gko::matrix::Hybrid< ValueType, IndexType >::operator= (
    const Hybrid< ValueType, IndexType > & )
```

Copy-assigns a [Hybrid](#) matrix.

Preserves the executor, copy-assigns the [EII](#) and [Coo](#) matrices.

47.118.3.32 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Hybrid& gko::matrix::Hybrid< ValueType, IndexType >::operator= (
    Hybrid< ValueType, IndexType > && )
```

Move-assigns a [Hybrid](#) matrix.

Preserves the executor, move-assigns the [EII](#) and [Coo](#) matrices. The moved-from matrix is empty (0x0 with empty EII/Coo matrices).

47.118.3.33 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Hybrid< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.118.3.34 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Hybrid< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.118.3.35 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Hybrid< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.118.3.36 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Hybrid< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/coo.hpp
- ginkgo/core/matrix/hybrid.hpp

47.119 `gko::preconditioner::lc< LSolverType, IndexType >` Class Template Reference

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

```
#include <ginkgo/core/preconditioner/ic.hpp>
```

Public Member Functions

- `std::shared_ptr< const l_solver_type > get_l_solver () const`
Returns the solver which is used for the provided L matrix.
- `std::shared_ptr< const lh_solver_type > get_lh_solver () const`
Returns the solver which is used for the L^H matrix.
- `std::unique_ptr< LinOp > transpose () const override`
Returns a `LinOp` representing the transpose of the `Transposable` object.
- `std::unique_ptr< LinOp > conj_transpose () const override`
Returns a `LinOp` representing the conjugate transpose of the `Transposable` object.
- `lc & operator= (const lc &other)`
Copy-assigns an IC preconditioner.
- `lc & operator= (lc &&other)`
Move-assigns an IC preconditioner.
- `lc (const lc &other)`
Copy-constructs an IC preconditioner.
- `lc (lc &&other)`
Move-constructs an IC preconditioner.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< value_type, index_type >())`
Create the parameters from the property_tree.

47.119.1 Detailed Description

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
class gko::preconditioner::lc< LSolverType, IndexType >
```

The Incomplete Cholesky (IC) preconditioner solves the equation $LL^H * x = b$ for a given lower triangular matrix L and the right hand side b (can contain multiple right hand sides).

It allows to set both the solver for L defaulting to `solver::LowerTrs`, which is a direct triangular solvers. The solver for L^H is the conjugate-transposed solver for L , ensuring that the preconditioner is symmetric and positive-definite. For this L solver, a factory can be provided (using `with_l_solver`) to have more control over their behavior. In particular, it is possible to use an iterative method for solving the triangular systems. The default parameters for an iterative triangular solver are:

- reduction factor = $1e-4$

- max iteration = <number of="" rows="" of="" the="" matrix="" given="" to="" the="" solver>=""> Solvers without such criteria can also be used, in which case none are set.

An object of this class can be created with a matrix or a [gko::Composition](#) containing two matrices. If created with a matrix, it is factorized before creating the solver. If a [gko::Composition](#) (containing two matrices) is used, the first operand will be taken as the L matrix, the second will be considered the L^H matrix, which helps to avoid the otherwise necessary transposition of L inside the solver. Parlc can be directly used, since it orders the factors in the correct way.

Note

When providing a [gko::Composition](#), the first matrix must be the lower matrix (L), and the second matrix must be its conjugate-transpose (L^H). If they are swapped, solving might crash or return the wrong result.

Do not use symmetric solvers (like CG) for the L solver since both matrices (L and L^H) are, by design, not symmetric.

This class is not thread safe (even a const object is not) because it uses an internal cache to accelerate multiple (sequential) applies. Using it in parallel can lead to segmentation faults, wrong results and other unwanted behavior.

Template Parameters

<i>LSolverType</i>	type of the solver used for the L matrix. Defaults to solver::LowerTrs
<i>IndexType</i>	type of the indices when Parlc is used to generate the L and L^H factors. Irrelevant otherwise.

47.119.2 Constructor & Destructor Documentation

47.119.2.1 lc() [1/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
gko::preconditioner::lc< LSolverType, IndexType >::lc (
    const lc< LSolverType, IndexType > & other ) [inline]
```

Copy-constructs an IC preconditioner.

Inherits the executor, shallow-copies the solvers and parameters.

```
318 :   lc{other.get_executor()} { *this = other; }
```

References [gko::PolymorphicObject::get_executor\(\)](#).

47.119.2.2 lc() [2/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
gko::preconditioner::lc< LSolverType, IndexType >::lc (
    lc< LSolverType, IndexType > && other ) [inline]
```

Move-constructs an IC preconditioner.

Inherits the executor, moves the solvers and parameters. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References [gko::PolymorphicObject::get_executor\(\)](#).

47.119.3 Member Function Documentation

47.119.3.1 conj_transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ic< LSolverType, IndexType >::conj_transpose ( )
const [inline], [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

References [gko::PolymorphicObject::get_executor\(\)](#), [gko::preconditioner::Ic< LSolverType, IndexType >::get_lhs_solver\(\)](#), [gko::preconditioner::Ic< LSolverType, IndexType >::get_rhs_solver\(\)](#), [gko::share\(\)](#), and [gko::transpose\(\)](#).

47.119.3.2 get_lhs_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::shared_ptr<const lhs_solver_type> gko::preconditioner::Ic< LSolverType, IndexType >::get_lhs_solver ( ) const [inline]
```

Returns the solver which is used for the provided L matrix.

Returns

the solver which is used for the provided L matrix

Referenced by [gko::preconditioner::Ic< LSolverType, IndexType >::conj_transpose\(\)](#), and [gko::preconditioner::Ic< LSolverType, IndexType >::transpose\(\)](#).

47.119.3.3 get_rhs_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::shared_ptr<const rhs_solver_type> gko::preconditioner::Ic< LSolverType, IndexType >::get_rhs_solver ( ) const [inline]
```

Returns the solver which is used for the L^H matrix.

Returns

the solver which is used for the L^H matrix

Referenced by [gko::preconditioner::Ic< LSolverType, IndexType >::conj_transpose\(\)](#), and [gko::preconditioner::Ic< LSolverType, IndexType >::transpose\(\)](#).

47.119.3.4 operator=() [1/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
Ic& gko::preconditioner::Ic< LSolverType, IndexType >::operator= (
    const Ic< LSolverType, IndexType > & other ) [inline]
```

Copy-assigns an IC preconditioner.

Preserves the executor, shallow-copies the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor.

References gko::clone(), and gko::PolymorphicObject::get_executor().

47.119.3.5 operator=() [2/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
Ic& gko::preconditioner::Ic< LSolverType, IndexType >::operator= (
    Ic< LSolverType, IndexType > && other ) [inline]
```

Move-assigns an IC preconditioner.

Preserves the executor, moves the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::clone(), and gko::PolymorphicObject::get_executor().

47.119.3.6 parse()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
static parameters_type gko::preconditioner::Ic< LSolverType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor<value↔
_type, index_type>() ) [inline], [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

Note

only support the following for `I_solver`: `Ir`, `Gmres`, `LowerTrs`, and `Lowerlsai`

47.119.3.7 transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ic< LSolverType, IndexType >::transpose ( ) const
[inline], [override], [virtual]
```

Returns a `LinOp` representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

References [gko::PolymorphicObject::get_executor\(\)](#), [gko::preconditioner::Ic< LSolverType, IndexType >::get_I_solver\(\)](#), [gko::preconditioner::Ic< LSolverType, IndexType >::get_lh_solver\(\)](#), [gko::share\(\)](#), and [gko::transpose\(\)](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/preconditioner/ic.hpp`

47.120 gko::factorization::Ic< ValueType, IndexType > Class Template Reference

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

```
#include <ginkgo/core/factorization/ic.hpp>
```

Static Public Member Functions

- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType, IndexType >())
Create the parameters from the property_tree.

Additional Inherited Members**47.120.1 Detailed Description**

```
template<typename ValueType = gko::default_precision, typename IndexType = gko::int32>
class gko::factorization::Ic< ValueType, IndexType >
```

Represents an incomplete Cholesky factorization (IC(0)) of a sparse matrix.

More specifically, it consists of a lower triangular factor L and its conjugate transpose L^H with sparsity pattern $S(L + L^H) = S(A)$ fulfilling $LL^H = A$ at every non-zero location of A .

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.120.2 Member Function Documentation

47.120.2.1 parse()

```
template<typename ValueType = gko::default_precision, typename IndexType = gko::int32>
static parameters_type gko::factorization::Ic< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/ic.hpp

47.121 gko::batch::matrix::Identity< ValueType > Class Template Reference

The batch Identity matrix, which represents a batch of Identity matrices.

```
#include <ginkgo/core/matrix/batch_identity.hpp>
```

Public Member Functions

- `Identity * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector.
- `Identity * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x)`
Apply the matrix to a multi-vector with a linear combination of the given input vector.
- `const Identity * apply (ptr_param< const MultiVector< value_type >> b, ptr_param< MultiVector< value_type >> x) const`
- `const Identity * apply (ptr_param< const MultiVector< value_type >> alpha, ptr_param< const MultiVector< value_type >> b, ptr_param< const MultiVector< value_type >> beta, ptr_param< MultiVector< value_type >> x) const`

Static Public Member Functions

- `static std::unique_ptr< Identity > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size=batch_dim< 2 >{})`
Creates an Identity matrix of the specified size.

47.121.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::matrix::Identity< ValueType >
```

The batch `Identity` matrix, which represents a batch of `Identity` matrices.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.121.2 Member Function Documentation

47.121.2.1 `apply()` [1/4]

```
template<typename ValueType = default_precision>
Identity* gko::batch::matrix::Identity< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector with a linear combination of the given input vector.

Represents the matrix vector multiplication, $x = \alpha * I$

- $b + \beta * x$, where x and b are both multi-vectors.

Parameters

<i>alpha</i>	the scalar to scale the matrix-vector product with
<i>b</i>	the multi-vector to be applied to
<i>beta</i>	the scalar to scale the x vector with
<i>x</i>	the output multi-vector

47.121.2.2 apply() [2/4]

```
template<typename ValueType = default_precision>
const Identity* gko::batch::matrix::Identity< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> alpha,
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< const MultiVector< value_type >> beta,
    ptr_param< MultiVector< value_type >> x ) const
```

47.121.2.3 apply() [3/4]

```
template<typename ValueType = default_precision>
Identity* gko::batch::matrix::Identity< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x )
```

Apply the matrix to a multi-vector.

Represents the matrix vector multiplication, $x = I * b$, where x and b are both multi-vectors.

Parameters

<i>b</i>	the multi-vector to be applied to
<i>x</i>	the output multi-vector

47.121.2.4 apply() [4/4]

```
template<typename ValueType = default_precision>
const Identity* gko::batch::matrix::Identity< ValueType >::apply (
    ptr_param< const MultiVector< value_type >> b,
    ptr_param< MultiVector< value_type >> x ) const
```

47.121.2.5 create()

```
template<typename ValueType = default_precision>
static std::unique_ptr<Identity> gko::batch::matrix::Identity< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size = batch_dim< 2 >{} ) [static]
```

Creates an [Identity](#) matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the batch matrices in a batch_dim object

Returns

A smart pointer to the newly created matrix.

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/batch_identity.hpp

47.122 gko::matrix::Identity< ValueType > Class Template Reference

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

```
#include <ginkgo/core/matrix/identity.hpp>
```

Public Member Functions

- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Static Public Member Functions

- static std::unique_ptr< [Identity](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [dim](#)< 2 > size)
Creates an [Identity](#) matrix of the specified size.
- static std::unique_ptr< [Identity](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size=0)
Creates an [Identity](#) matrix of the specified size.

47.122.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::matrix::Identity< ValueType >
```

This class is a utility which efficiently implements the identity matrix (a linear operator which maps each vector to itself).

Thus, objects of the [Identity](#) class always represent a square matrix, and don't require any storage for their values. The apply method is implemented as a simple copy (or a linear combination).

Note

This class is useful when composing it with other operators. For example, it can be used instead of a preconditioner in Krylov solvers, if one wants to run a "plain" solver, without using a preconditioner.

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.122.2 Member Function Documentation

47.122.2.1 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Identity< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.122.2.2 create() [1/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Identity> gko::matrix::Identity< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    dim< 2 > size ) [static]
```

Creates an [Identity](#) matrix of the specified size.

Parameters

<i>size</i>	size of the matrix (must be square)
-------------	-------------------------------------

47.122.2.3 create() [2/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Identity> gko::matrix::Identity< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    size_type size = 0 ) [static]
```

Creates an [Identity](#) matrix of the specified size.

Parameters

<i>size</i>	size of the matrix
-------------	--------------------

47.122.2.4 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::matrix::Identity< ValueType >::transpose ( ) const [override],
[virtual]
```

Returns a `LinOp` representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/identity.hpp`

47.123 gko::matrix::IdentityFactory< ValueType > Class Template Reference

This factory is a utility which can be used to generate [Identity](#) operators.

```
#include <ginkgo/core/matrix/identity.hpp>
```

Static Public Member Functions

- static std::unique_ptr< IdentityFactory > create (std::shared_ptr< const Executor > exec)
Creates a new Identity factory.

Additional Inherited Members

47.123.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::matrix::IdentityFactory< ValueType >
```

This factory is a utility which can be used to generate Identity operators.

The factory will generate the Identity matrix with the same dimension as the passed in operator. It will throw an exception if the operator is not square.

Template Parameters

ValueType	precision of matrix elements
-----------	------------------------------

47.123.2 Member Function Documentation

47.123.2.1 create()

```
template<typename ValueType = default_precision>
static std::unique_ptr<IdentityFactory> gko::matrix::IdentityFactory< ValueType >::create (
    std::shared_ptr< const Executor > exec ) [inline], [static]
```

Creates a new Identity factory.

Parameters

exec	the executor where the Identity operator will be stored
------	---

Returns

a unique pointer to the newly created factory

```
105 {
106     return std::unique_ptr<IdentityFactory>(
107         new IdentityFactory(std::move(exec));
108 }
```

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/identity.hpp

47.124 gko::solver::Idr< ValueType > Class Template Reference

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

```
#include <ginkgo/core/solver/idr.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.
- `size_type get_subspace_dim ()` const
Gets the subspace dimension of the solver.
- `void set_subspace_dim (const size_type other)`
Sets the subspace dimension of the solver.
- `remove_complex< ValueType > get_kappa ()` const
Gets the kappa parameter of the solver.
- `void set_kappa (const remove_complex< ValueType > other)`
Sets the kappa parameter of the solver.
- `bool get_deterministic ()` const
Gets the deterministic parameter of the solver.
- `void set_deterministic (const bool other)`
Sets the deterministic parameter of the solver.
- `bool get_complex_subspace ()` const
Gets the complex_subspace parameter of the solver.
- `void set_complex_subspace (const bool other)`
Sets the complex_subspace parameter of the solver.
- `void set_complex_subspace (const bool other)`
Sets the complex_subspace parameter of the solver.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.124.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::Idr< ValueType >
```

IDR(s) is an efficient method for solving large nonsymmetric systems of linear equations.

The implemented version is the one presented in the paper "Algorithm 913: An elegant IDR(s) variant that efficiently exploits biorthogonality properties" by M. B. Van Gijzen and P. Sonneveld.

The method is based on the induced dimension reduction theorem which provides a way to construct subsequent residuals that lie in a sequence of shrinking subspaces. These subspaces are spanned by s vectors which are first generated randomly and then orthonormalized. They are stored in a dense matrix.

Template Parameters

<i>ValueType</i>	precision of the elements of the system matrix.
------------------	---

47.124.2 Member Function Documentation

47.124.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::ldr< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

```
76 { return true; }
```

47.124.2.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::ldr< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.124.2.3 get_complex_subspace()

```
template<typename ValueType = default_precision>
bool gko::solver::ldr< ValueType >::get_complex_subspace ( ) const [inline]
```

Gets the complex_subspace parameter of the solver.

Returns

the complex_subspace parameter

47.124.2.4 get_deterministic()

```
template<typename ValueType = default_precision>
bool gko::solver::Idr< ValueType >::get_deterministic ( ) const [inline]
```

Gets the deterministic parameter of the solver.

Returns

the deterministic parameter

47.124.2.5 get_kappa()

```
template<typename ValueType = default_precision>
remove_complex<ValueType> gko::solver::Idr< ValueType >::get_kappa ( ) const [inline]
```

Gets the kappa parameter of the solver.

Returns

the kappa parameter

47.124.2.6 get_subspace_dim()

```
template<typename ValueType = default_precision>
size_type gko::solver::Idr< ValueType >::get_subspace_dim ( ) const [inline]
```

Gets the subspace dimension of the solver.

Returns

the subspace Dimension

47.124.2.7 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Idr< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.124.2.8 set_complex_subspace()

```
template<typename ValueType = default_precision>
void gko::solver::ldr< ValueType >::set_complex_subspace (
    const bool other ) [inline]
```

Sets the `complex_subspace` parameter of the solver.

Parameters

<i>other</i>	the new <code>complex_subspace</code> parameter
--------------	---

References `gko::solver::ldr< ValueType >::set_complex_subspace()`.

47.124.2.9 set_complex_subspace()

```
template<typename ValueType = default_precision>
void gko::solver::ldr< ValueType >::set_complex_subspace (
    const bool other ) [inline]
```

Sets the `complex_subspace` parameter of the solver.

Parameters

<i>other</i>	the new <code>complex_subspace</code> parameter
--------------	---

Referenced by `gko::solver::ldr< ValueType >::set_complex_subspace()`.

47.124.2.10 set_deterministic()

```
template<typename ValueType = default_precision>
void gko::solver::ldr< ValueType >::set_deterministic (
    const bool other ) [inline]
```

Sets the deterministic parameter of the solver.

Parameters

<i>other</i>	the new deterministic parameter
--------------	---------------------------------

47.124.2.11 set_kappa()

```
template<typename ValueType = default_precision>
void gko::solver::Idr< ValueType >::set_kappa (
    const remove_complex< ValueType > other ) [inline]
```

Sets the kappa parameter of the solver.

Parameters

<i>other</i>	the new kappa parameter
--------------	-------------------------

47.124.2.12 set_subspace_dim()

```
template<typename ValueType = default_precision>
void gko::solver::Idr< ValueType >::set_subspace_dim (
    const size_type other ) [inline]
```

Sets the subspace dimension of the solver.

Parameters

<i>other</i>	the new subspace Dimension
--------------	----------------------------

47.124.2.13 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Idr< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/idr.hpp`

47.125 gko::factorization::ilu< ValueType, IndexType > Class Template Reference

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

```
#include <ginkgo/core/factorization/ilu.hpp>
```

Static Public Member Functions

- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

Additional Inherited Members

47.125.1 Detailed Description

```
template<typename ValueType = gko::default_precision, typename IndexType = gko::int32>
class gko::factorization::ilu< ValueType, IndexType >
```

Represents an incomplete LU factorization – ILU(0) – of a sparse matrix.

More specifically, it consists of a lower unitriangular factor L and an upper triangular factor U with sparsity pattern $S(L + U) = S(A)$ fulfilling $LU = A$ at every non-zero location of A .

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.125.2 Member Function Documentation

47.125.2.1 parse()

```
template<typename ValueType = gko::default_precision, typename IndexType = gko::int32>
static parameters_type gko::factorization::ilu< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type\_descriptor & td_for_child = config::make\_type\_descriptor< ValueType, IndexType >
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/ilu.hpp

47.126 gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType > Class Template Reference

The Incomplete LU (ILU) preconditioner solves the equation $LUx = b$ for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

```
#include <ginkgo/core/preconditioner/ilu.hpp>
```

Public Member Functions

- `std::shared_ptr< const l_solver_type > get\_l\_solver () const`
Returns the solver which is used for the provided L matrix.
- `std::shared_ptr< const u_solver_type > get\_u\_solver () const`
Returns the solver which is used for the provided U matrix.
- `std::unique_ptr< LinOp > transpose () const override`
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj\_transpose () const override`
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `llu & operator= (const llu &other)`
Copy-assigns an ILU preconditioner.
- `llu & operator= (llu &&other)`
Move-assigns an ILU preconditioner.
- `llu (const llu &other)`
Copy-constructs an ILU preconditioner.
- `llu (llu &&other)`
Move-constructs an ILU preconditioner.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type\_descriptor &td_for_child=config::make_type_descriptor< value_type, index_type >())`
Create the parameters from the property_tree.

47.126.1 Detailed Description

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply =
false, typename IndexType = int32>
class gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >
```

The Incomplete LU (ILU) preconditioner solves the equation $LUx = b$ for a given lower triangular matrix L, an upper triangular matrix U and the right hand side b (can contain multiple right hand sides).

It allows to set both the solver for L and the solver for U independently, while providing the defaults [solver::LowerTrs](#) and [solver::UpperTrs](#), which are direct triangular solvers. For these solvers, a factory can be provided (with `with_l_solver` and `with_u_solver`) to have more control over their behavior. In particular, it is possible to use an iterative method for solving the triangular systems. The default parameters for an iterative triangular solver are:

- reduction factor = 1e-4
- max iteration = <number of="" rows="" of="" the="" matrix="" given="" to="" the="" solver>=""> Solvers without such criteria can also be used, in which case none are set.

An object of this class can be created with a matrix or a [gko::Composition](#) containing two matrices. If created with a matrix, it is factorized before creating the solver. If a [gko::Composition](#) (containing two matrices) is used, the first operand will be taken as the L matrix, the second will be considered the U matrix. Parllu can be directly used, since it orders the factors in the correct way.

Note

When providing a [gko::Composition](#), the first matrix must be the lower matrix (*L*), and the second matrix must be the upper matrix (*U*). If they are swapped, solving might crash or return the wrong result.

Do not use symmetric solvers (like CG) for L or U solvers since both matrices (L and U) are, by design, not symmetric.

This class is not thread safe (even a const object is not) because it uses an internal cache to accelerate multiple (sequential) applies. Using it in parallel can lead to segmentation faults, wrong results and other unwanted behavior.

Template Parameters

<i>LSolverType</i>	type of the solver used for the L matrix. Defaults to solver::LowerTrs
<i>USolverType</i>	type of the solver used for the U matrix Defaults to solver::UpperTrs
<i>ReverseApply</i>	default behavior (ReverseApply = false) is first to solve with L (Ly = b) and then with U (Ux = y). When set to true, it will solve first with U, and then with L.
<i>IndexTypeParllu</i>	Type of the indices when Parllu is used to generate both L and U factors. Irrelevant otherwise.

47.126.2 Constructor & Destructor Documentation

47.126.2.1 llu() [1/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply = false, typename IndexType = int32>
```

```
gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::Ilu (
    const Ilu< LSolverType, USolverType, ReverseApply, IndexType > & other ) [inline]
```

Copy-constructs an ILU preconditioner.

Inherits the executor, shallow-copies the solvers and parameters.

```
365 : Ilu{other.get_executor()} { *this = other; }
```

References gko::PolymorphicObject::get_executor().

47.126.2.2 Ilu() [2/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper↵
Trs<>, bool ReverseApply = false, typename IndexType = int32>
gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::Ilu (
    Ilu< LSolverType, USolverType, ReverseApply, IndexType > && other ) [inline]
```

Move-constructs an ILU preconditioner.

Inherits the executor, moves the solvers and parameters. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::PolymorphicObject::get_executor().

47.126.3 Member Function Documentation

47.126.3.1 conj_transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper↵
Trs<>, bool ReverseApply = false, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply,
IndexType >::conj_transpose ( ) const [inline], [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

References gko::PolymorphicObject::get_executor(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_l_solver(), gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_u_solver(), gko::share(), and gko::transpose().

47.126.3.2 get_l_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply = false, typename IndexType = int32>
std::shared_ptr<const l_solver_type> gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::get_l_solver ( ) const [inline]
```

Returns the solver which is used for the provided L matrix.

Returns

the solver which is used for the provided L matrix

Referenced by gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose().

47.126.3.3 get_u_solver()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply = false, typename IndexType = int32>
std::shared_ptr<const u_solver_type> gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::get_u_solver ( ) const [inline]
```

Returns the solver which is used for the provided U matrix.

Returns

the solver which is used for the provided U matrix

Referenced by gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose(), and gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose().

47.126.3.4 operator=() [1/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTrs<>, bool ReverseApply = false, typename IndexType = int32>
llu& gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::operator= (
    const llu< LSolverType, USolverType, ReverseApply, IndexType > & other ) [inline]
```

Copy-assigns an ILU preconditioner.

Preserves the executor, shallow-copies the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor.

References gko::clone(), and gko::PolymorphicObject::get_executor().

47.126.3.5 operator=() [2/2]

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper↵
Trs<>, bool ReverseApply = false, typename IndexType = int32>
Ilu& gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >::operator=
(
    Ilu< LSolverType, USolverType, ReverseApply, IndexType > && other ) [inline]
```

Move-assigns an ILU preconditioner.

Preserves the executor, moves the solvers and parameters. Creates a clone of the solvers if they are on the wrong executor. The moved-from object is empty (0x0 with nullptr solvers and default parameters)

References gko::clone(), and gko::PolymorphicObject::get_executor().

47.126.3.6 parse()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::Upper↵
Trs<>, bool ReverseApply = false, typename IndexType = int32>
static parameters_type gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, Index↵
Type >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor<value↵
_type, index_type>() ) [inline], [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

Note

only support the following pairs for <l_solver, u_solver>: <lr, lr>, <Gmres, Gmres>, <LowerTrs, Upper↵ Trs>, and <Lowerlsai, Upperlsai>

47.126.3.7 transpose()

```
template<typename LSolverType = solver::LowerTrs<>, typename USolverType = solver::UpperTriTrs<>, bool ReverseApply = false, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::ilu< LSolverType, USolverType, ReverseApply, IndexType >::transpose ( ) const [inline], [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

References [gko::PolymorphicObject::get_executor\(\)](#), [gko::preconditioner::ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_l_solver\(\)](#), [gko::preconditioner::ilu< LSolverType, USolverType, ReverseApply, IndexType >::get_u_solver\(\)](#), [gko::share\(\)](#), and [gko::transpose\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/preconditioner/ilu.hpp](#)

47.127 gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit Class Reference

[imbalance_bounded_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- [imbalance_bounded_limit](#) (double percent=0.8, double ratio=0.0001)
Creates a [imbalance_bounded_limit](#) strategy.
- [size_type compute_ell_num_stored_elements_per_row](#) (array< [size_type](#) > *row_nnz) const override
Computes the number of stored elements per row of the ell part.
- auto [get_percentage](#) () const
Get the percent setting.
- auto [get_ratio](#) () const
Get the ratio setting.

47.127.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit
```

[imbalance_bounded_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part.

It uses the [imbalance_limit](#) and adds the upper bound of the number of ell's cols by the number of rows.

47.127.2 Member Function Documentation

47.127.2.1 `compute_ell_num_stored_elements_per_row()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_stored_elements_per_row (
    array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<code>row_nnz</code>	the number of nonzeros of each row
----------------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type](#).

References [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row\(\)](#), and [gko::array< ValueType >::get_size\(\)](#).

Referenced by [gko::matrix::Hybrid< ValueType, IndexType >::automatic::compute_ell_num_stored_elements_per_row\(\)](#).

47.127.2.2 `get_percentage()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_percentage ( )
const [inline]
```

Get the percent setting.

Returns

percent

References [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::get_percentage\(\)](#).

47.127.2.3 get_ratio()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_ratio ( ) const
[inline]
```

Get the ratio setting.

Returns

ratio

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/hybrid.hpp

47.128 gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit Class Reference

[imbalance_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part according to the percent.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- [imbalance_limit](#) (double percent=0.8)
Creates a [imbalance_limit](#) strategy.
- [size_type compute_ell_num_stored_elements_per_row](#) (array< [size_type](#) > *row_nnz) const override
Computes the number of stored elements per row of the ell part.
- auto [get_percentage](#) () const
Get the percent setting.

47.128.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit
```

[imbalance_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part according to the percent.

It sorts the number of nonzeros of each row and takes the value at the position `floor(percent * num_row)` as the number of stored elements per row of the ell part. Thus, at least `percent` rows of all are in the ell part.

47.128.2 Constructor & Destructor Documentation

47.128.2.1 imbalance_limit()

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::imbalance_limit (
    double percent = 0.8 ) [inline], [explicit]
```

Creates a [imbalance_limit](#) strategy.

Parameters

<i>percent</i>	the row_nnz[floor(num_rows*percent)] is the number of stored elements per row of the ell part
----------------	---

47.128.3 Member Function Documentation

47.128.3.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_↵
stored_elements_per_row (
    array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<i>row_nnz</i>	the number of nonzeros of each row
----------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type](#).

References [gko::array< ValueType >::get_data\(\)](#), and [gko::array< ValueType >::get_size\(\)](#).

Referenced by [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::compute_ell_num_↵](#)
[stored_elements_per_row\(\)](#), and [gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::compute_↵](#)
[_ell_num_stored_elements_per_row\(\)](#).

47.128.3.2 get_percentage()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::get_percentage ( ) const
[inline]
```

Get the percent setting.

Returns

percent

Referenced by [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit::get_percentage\(\)](#), and
[gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::get_percentage\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/hybrid.hpp](#)

47.129 gko::stop::ImplicitResidualNorm< ValueType > Class Template Reference

The [ImplicitResidualNorm](#) class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.129.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::stop::ImplicitResidualNorm< ValueType >
```

The [ImplicitResidualNorm](#) class is a stopping criterion which stops the iteration process when the implicit residual norm is below a certain threshold relative to.

1. the norm of the right-hand side, $\text{implicit_resnorm} \leq \text{threshold} * \text{norm}(\text{right_hand_side})$
2. the initial residual, $\text{implicit_resnorm} \leq \text{threshold} * \text{norm}(\text{initial_residual})$.
3. one, $\text{implicit_resnorm} \leq \text{threshold}$.

Note

To use this stopping criterion there are some dependencies. The constructor depends on either `b` or the `initial_residual` in order to compute their norms. If this is not correctly provided, an exception `gko::NotSupported()` is thrown.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/residual_norm.hpp`

47.130 gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType > Struct Template Reference

This class defines mappings between global and local indices.

```
#include <ginkgo/core/distributed/index_map.hpp>
```

Public Member Functions

- `array< LocalIndexType > map_to_local` (const `array< GlobalIndexType >` &global_ids, `index_space` `index↔_space_v`) const
Maps global indices to local indices.
- `size_type get_global_size` () const
get size of `index_space::local`
- `size_type get_local_size` () const
get size of `index_space::local`
- `size_type get_non_local_size` () const
get size of `index_space::non_local`
- `index_map` (std::shared_ptr< const `Executor` > exec, std::shared_ptr< const `partition_type` > partition, `comm_index_type` rank, const `array< GlobalIndexType >` &recv_connections)
Creates a new index map.
- `index_map` (std::shared_ptr< const `Executor` > exec)
Creates an empty index map.
- const `segmented_array< GlobalIndexType >` & `get_remote_global_idx`s () const
get the index set R_k for this rank.
- const `segmented_array< LocalIndexType >` & `get_remote_local_idx`s () const
get the index set R_k , but mapped to their respective local index space.
- const `array< comm_index_type >` & `get_remote_target_ids` () const
get the rank ids which contain indices accessed by this rank.
- std::shared_ptr< const `Executor` > `get_executor` () const
get the associated executor.

47.130.1 Detailed Description

```
template<typename LocalIndexType, typename GlobalIndexType = int64>
struct gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >
```

This class defines mappings between global and local indices.

Given an index space $I = [0, \dots, N)$ that is partitioned into P disjoint subsets $I_k, k = 1, \dots, P$, this class defines for each subset an extended global index set $\hat{I}_k \supset I_k$. The extended index set contains the global indices owned by part k , as well as remote indices $R_k = \hat{I}_k \setminus I_k$, which are also accessed by part k , but owned by parts $l \neq k$. At the core, this class provides mappings from the global index space I into different local index spaces. The combined local index space (`index_space::combined`) is then defined as $[0, \dots, |\hat{I}_k|)$. Additionally, the combined index space can be separated into locally owned (`index_space::local`) and non-locally owned (`index_space::non_local`). The locally owned indices are defined as $[0, \dots, |I_k|)$, and the non-locally owned as $[0, \dots, |R_k|)$. With these index sets, the following mappings are defined:

- $c_k : \hat{I}_k \mapsto [0, \dots, |\hat{I}_k|)$ which maps global indices into the combined/full local index space (denoted as `index_space::combined`),
- $l_k : I_k \mapsto [0, \dots, |I_k|)$ which maps global indices into the locally owned index space (denoted as `index_space::local`),
- $r_k : R_k \mapsto [0, \dots, |R_k|)$ which maps global indices into the non-locally owned index space (denoted as `index_space::non_local`).

The required map can be selected by passing the appropriate type of an `index_space`.

The index map for I_k has no knowledge about any other index maps for $I_l, l \neq k$. In particular, any global index passed to the `map_to_local` map that is not part of the specified index space, will be mapped to an `invalid_index`.

Template Parameters

<i>LocalIndexType</i>	type for local indices
<i>GlobalIndexType</i>	type for global indices

47.130.2 Constructor & Destructor Documentation

47.130.2.1 index_map()

```
template<typename LocalIndexType , typename GlobalIndexType = int64>
gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >::index_map (
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< const partition_type > partition,
    comm_index_type rank,
    const array< GlobalIndexType > & recv_connections )
```

Creates a new index map.

The passed in recv_connections may contain duplicates, which will be filtered out.

Parameters

<i>exec</i>	the executor
<i>partition</i>	the partition of the global index set
<i>rank</i>	the id of the global index space subset
<i>recv_connections</i>	the global indices that are not owned by this rank, but accessed by it

47.130.3 Member Function Documentation

47.130.3.1 get_remote_global_idxxs()

```
template<typename LocalIndexType , typename GlobalIndexType = int64>
const segmented_array<GlobalIndexType>& gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >::get_remote_global_idxxs ( ) const
```

get the index set R_k for this rank.

The indices are ordered by their owning rank and global index.

47.130.3.2 get_remote_local_idxxs()

```
template<typename LocalIndexType , typename GlobalIndexType = int64>
const segmented_array<LocalIndexType>& gko::experimental::distributed::index_map< Local↔
IndexType, GlobalIndexType >::get_remote_local_idxxs ( ) const
```

get the index set R_k , but mapped to their respective local index space.

The indices are grouped by their owning rank and sorted according to their global index within each group.

The set $R_k = \hat{I}_k \setminus I_k$ can also be written as the union of the intersection of $\hat{I}_k$ with other disjoint sets $I_l, l \neq k$, i.e. $R_k = \bigcup_{j \neq k} \hat{I}_k \cap I_j = \bigcup_{j \neq k} R_{k,j}$. The set $R_{k,j}$ can then be mapped by l_j to get the local indices wrt. part j . The indices here are mapped by l_j .

47.130.3.3 get_remote_target_ids()

```
template<typename LocalIndexType , typename GlobalIndexType = int64>
const array<comm_index_type>& gko::experimental::distributed::index_map< LocalIndexType,
GlobalIndexType >::get_remote_target_ids ( ) const
```

get the rank ids which contain indices accessed by this rank.

The order matches the order of the sets in get_remote_global_idxxs and get_remote_local_idxxs.

47.130.3.4 map_to_local()

```
template<typename LocalIndexType , typename GlobalIndexType = int64>
array<LocalIndexType> gko::experimental::distributed::index_map< LocalIndexType, Global↔
IndexType >::map_to_local (
    const array< GlobalIndexType > & global_ids,
    index_space index_space_v ) const
```

Maps global indices to local indices.

Parameters

<i>global_ids</i>	the global indices to map
<i>index_space</i> ↔ _v	the index space in which the returned local indices are defined

Returns

the mapped local indices. Any global index that is not in the specified index space is mapped to invalid_index.

The documentation for this struct was generated from the following file:

- ginkgo/core/distributed/index_map.hpp

47.131 gko::index_set< IndexType > Class Template Reference

An index set class represents an ordered set of intervals.

```
#include <ginkgo/core/base/index_set.hpp>
```

Public Types

- using `index_type` = `IndexType`
The type of elements stored in the index set.

Public Member Functions

- `index_set` (`std::shared_ptr< const Executor > exec`) noexcept
Creates an empty `index_set` tied to the specified `Executor`.
- `index_set` (`std::shared_ptr< const gko::Executor > exec`, `std::initializer_list< IndexType > init_list`, `const bool is_sorted=false`)
Creates an index set on the specified executor from the initializer list.
- `index_set` (`std::shared_ptr< const gko::Executor > exec`, `const index_type size`, `const gko::array< index_type > &indices`, `const bool is_sorted=false`)
Creates an index set on the specified executor and the given size.
- `index_set` (`std::shared_ptr< const Executor > exec`, `const index_set &other`)
Creates a copy of the input `index_set` on a different executor.
- `index_set` (`const index_set &other`)
Creates a copy of the input `index_set`.
- `index_set` (`std::shared_ptr< const Executor > exec`, `index_set &&other`)
Moves the input `index_set` to a different executor.
- `index_set` (`index_set &&other`)
Moves the input `index_set`.
- `index_set & operator=` (`const index_set &other`)
Copies data from another `index_set`.
- `index_set & operator=` (`index_set &&other`)
Moves data from another `index_set`.
- `void clear` () noexcept
Deallocates all data used by the `index_set`.
- `std::shared_ptr< const Executor > get_executor` () const
Returns the executor of the `index_set`.
- `index_type get_size` () const
Returns the size of the index set space.
- `bool is_contiguous` () const
Returns if the index set is contiguous.
- `index_type get_num_elems` () const
Return the actual number of indices stored in the index set.
- `index_type get_global_index` (`index_type local_index`) const
Return the global index given a local index.
- `index_type get_local_index` (`index_type global_index`) const
Return the local index given a global index.
- `array< index_type > map_local_to_global` (`const array< index_type > &local_indices`, `const bool is_↔ sorted=false`) const
This is an array version of the scalar function above.
- `array< index_type > map_global_to_local` (`const array< index_type > &global_indices`, `const bool is_↔ sorted=false`) const
This is an array version of the scalar function above.
- `array< index_type > to_global_indices` () const
This function allows the user obtain a decompressed `global_indices` array from the indices stored in the index set.
- `array< bool > contains` (`const array< index_type > &global_indices`, `const bool is_sorted=false`) const
Checks if the individual global indeices exist in the index set.

- `bool contains (const index_type global_index) const`
Checks if the global index exists in the index set.
- `index_type get_num_subsets () const`
Returns the number of subsets stored in the index set.
- `const index_type * get_subsets_begin () const`
Returns a pointer to the beginning indices of the subsets.
- `const index_type * get_subsets_end () const`
Returns a pointer to the end indices of the subsets.
- `const index_type * get_superset_indices () const`
Returns a pointer to the cumulative indices of the superset of the subsets.

47.131.1 Detailed Description

```
template<typename IndexType = int32>
class gko::index_set< IndexType >
```

An index set class represents an ordered set of intervals.

The index set contains subsets which store the starting and end points of a range, [a,b), storing the first index and one past the last index. As the index set only stores the end-points of ranges, it can be quite efficient in terms of storage.

This class is particularly useful in storing continuous ranges. For example, consider the index set (1, 2, 3, 4, 5, 6, 7, 8, 10, 11, 12, 18, 19, 20, 21, 42). Instead of storing the entire array of indices, one can store intervals ([1,9), [10,13), [18,22), [42,43)), thereby only using half the storage.

We store three arrays, one (subsets_begin) with the starting indices of the subsets in the index set, another (subsets_end) storing one index beyond the end indices of the subsets and the last (superset_cumulative_indices) storing the cumulative number of indices in the subsequent subsets with an initial zero which speeds up the querying. Additionally, the arrays containing the range boundaries (subsets_begin, subsets_end) are stored in a sorted fashion.

Therefore the storage would look as follows

```
index_set = (1, 2, 3, 4, 5, 6, 7, 8, 10, 11, 12, 18, 19, 20, 21, 42) subsets_begin = {1, 10, 18, 42}
subsets_end = {9, 13, 22, 43} superset_cumulative_indices = {0, 8, 11, 15, 16}
```

Template Parameters

<code>index_type</code>	type of the indices being stored in the index set.
-------------------------	--

47.131.2 Constructor & Destructor Documentation

47.131.2.1 index_set() [1/7]

```
template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
```

```
std::shared_ptr< const Executor > exec ) [inline], [explicit], [noexcept]
```

Creates an empty `index_set` tied to the specified `Executor`.

Parameters

<code>exec</code>	the <code>Executor</code> where the <code>index_set</code> data is allocated
-------------------	--

```

69         :   exec_(std::move(exec)),
70             index_space_size_{0},
71             num_stored_indices_{0},
72             subsets_begin_{array<index_type>(exec_)},
73             subsets_end_{array<index_type>(exec_)},
74             superset_cumulative_indices_{array<index_type>(exec_)}
75     {}

```

47.131.2.2 index_set() [2/7]

```

template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    std::shared_ptr< const gko::Executor > exec,
    std::initializer_list< IndexType > init_list,
    const bool is_sorted = false ) [inline], [explicit]

```

Creates an index set on the specified executor from the initializer list.

Parameters

<code>exec</code>	the <code>Executor</code> where the index set data will be allocated
<code>init_list</code>	the indices that the index set should hold in an <code>initializer_list</code> .
<code>is_sorted</code>	a parameter that specifies if the indices array is sorted or not. <code>true</code> if sorted.

47.131.2.3 index_set() [3/7]

```

template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    std::shared_ptr< const gko::Executor > exec,
    const index_type size,
    const gko::array< index_type > & indices,
    const bool is_sorted = false ) [inline], [explicit]

```

Creates an index set on the specified executor and the given size.

Parameters

<code>exec</code>	the <code>Executor</code> where the index set data will be allocated
<code>size</code>	the maximum index the index set it allowed to hold. This is the size of the index space.
<code>indices</code>	the indices that the index set should hold.
<code>is_sorted</code>	a parameter that specifies if the indices array is sorted or not. <code>true</code> if sorted.

References `gko::array< ValueType >::get_size()`.

47.131.2.4 `index_set()` [4/7]

```
template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    std::shared_ptr< const Executor > exec,
    const index_set< IndexType > & other ) [inline]
```

Creates a copy of the input `index_set` on a different executor.

Parameters

<i>exec</i>	the executor where the new <code>index_set</code> will be created
<i>other</i>	the <code>index_set</code> to copy from

47.131.2.5 `index_set()` [5/7]

```
template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    const index_set< IndexType > & other ) [inline]
```

Creates a copy of the input `index_set`.

Parameters

<i>other</i>	the <code>index_set</code> to copy from
--------------	---

47.131.2.6 `index_set()` [6/7]

```
template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    std::shared_ptr< const Executor > exec,
    index_set< IndexType > && other ) [inline]
```

Moves the input `index_set` to a different executor.

Parameters

<i>exec</i>	the executor where the new <code>index_set</code> will be moved to
<i>other</i>	the <code>index_set</code> to move from

47.131.2.7 index_set() [7/7]

```
template<typename IndexType = int32>
gko::index_set< IndexType >::index_set (
    index_set< IndexType > && other ) [inline]
```

Moves the input [index_set](#).

Parameters

<i>other</i>	the index_set to move from
--------------	--

47.131.3 Member Function Documentation

47.131.3.1 clear()

```
template<typename IndexType = int32>
void gko::index_set< IndexType >::clear ( ) [inline], [noexcept]
```

Deallocates all data used by the [index_set](#).

The [index_set](#) is left in a valid, but empty state, so the same [index_set](#) can be used to allocate new memory. Calls to [index_set::get_subsets_begin\(\)](#) will return a `nullptr`.

References `gko::array< ValueType >::clear()`.

47.131.3.2 contains() [1/2]

```
template<typename IndexType = int32>
array<bool> gko::index_set< IndexType >::contains (
    const array< index_type > & global_indices,
    const bool is_sorted = false ) const
```

Checks if the individual global indeices exist in the index set.

Parameters

<i>global_indices</i>	the indices to check.
<i>is_sorted</i>	a parameter that specifies if the query array is sorted or not. <code>true</code> if sorted.

Returns

the array that contains element wise whether the corresponding global index in the index set or not.

47.131.3.3 contains() [2/2]

```
template<typename IndexType = int32>
bool gko::index_set< IndexType >::contains (
    const index_type global_index ) const
```

Checks if the global index exists in the index set.

Parameters

<i>global_index</i>	the index to check.
---------------------	---------------------

Returns

whether the element exists in the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

47.131.3.4 get_executor()

```
template<typename IndexType = int32>
std::shared_ptr<const Executor> gko::index_set< IndexType >::get_executor ( ) const [inline]
```

Returns the executor of the [index_set](#).

Returns

the executor.

47.131.3.5 get_global_index()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_global_index (
    index_type local_index ) const
```

Return the global index given a local index.

Consider the set `idx_set = (0, 1, 2, 4, 6, 7, 8, 9)`. This function returns the element at the global index `k` stored in the index set. For example, `idx_set.get_global_index(0) == 0` `idx_set.get_global_index(3) == 4` and `idx_set.get_global_index(7) == 9`

Note

This function returns a scalar value and needs a scalar value. For repeated queries, it is more efficient to use the array functions that take and return arrays which allow for more throughput.

Parameters

<i>local_index</i>	the local index.
--------------------	------------------

Returns

the global index from the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

47.131.3.6 get_local_index()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_local_index (
    index_type global_index ) const
```

Return the local index given a global index.

Consider the set `idx_set = (0, 1, 2, 4, 6, 7, 8, 9)`. This function returns the local index in the index set of the provided index set. For example, `idx_set.get_local_index(0) == 0` `idx_set.get_local_index(4) == 3` and `idx_set.get_local_index(6) == 4`.

Note

This function returns a scalar value and needs a scalar value. For repeated queries, it is more efficient to use the array functions that take and return arrays which allow for more throughput.

Parameters

<i>global_index</i>	the global index.
---------------------	-------------------

Returns

the local index of the element in the index set.

Warning

This single entry query can have significant kernel launch overheads and should be avoided if possible.

47.131.3.7 get_num_elems()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_num_elems ( ) const [inline]
```

Return the actual number of indices stored in the index set.

Returns

number of indices stored in the index set

47.131.3.8 get_num_subsets()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_num_subsets ( ) const [inline]
```

Returns the number of subsets stored in the index set.

Returns

the number of stored subsets.

References `gko::array< ValueType >::get_size()`.

Referenced by `gko::index_set< IndexType >::is_contiguous()`.

47.131.3.9 get_size()

```
template<typename IndexType = int32>
index_type gko::index_set< IndexType >::get_size ( ) const [inline]
```

Returns the size of the index set space.

Returns

the size of the index set space.

47.131.3.10 get_subsets_begin()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_subsets_begin ( ) const [inline]
```

Returns a pointer to the beginning indices of the subsets.

Returns

a pointer to the beginning indices of the subsets.

References `gko::array< ValueType >::get_const_data()`.

47.131.3.11 get_subsets_end()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_subsets_end ( ) const [inline]
```

Returns a pointer to the end indices of the subsets.

Returns

a pointer to the end indices of the subsets.

References gko::array< ValueType >::get_const_data().

47.131.3.12 get_superset_indices()

```
template<typename IndexType = int32>
const index_type* gko::index_set< IndexType >::get_superset_indices ( ) const [inline]
```

Returns a pointer to the cumulative indices of the superset of the subsets.

Returns

a pointer to the cumulative indices of the superset of the subsets.

References gko::array< ValueType >::get_const_data().

47.131.3.13 is_contiguous()

```
template<typename IndexType = int32>
bool gko::index_set< IndexType >::is_contiguous ( ) const [inline]
```

Returns if the index set is contiguous.

Returns

if the index set is contiguous.

References gko::index_set< IndexType >::get_num_subsets().

47.131.3.14 map_global_to_local()

```
template<typename IndexType = int32>
array<index_type> gko::index_set< IndexType >::map_global_to_local (
    const array< index_type > & global_indices,
    const bool is_sorted = false ) const
```

This is an array version of the scalar function above.

Parameters

<i>global_indices</i>	the global index array.
<i>is_sorted</i>	a parameter that specifies if the query array is sorted or not. <code>true</code> if sorted.

Returns

the local index array from the index set.

Note

Whenever possible, passing a sorted array is preferred as the queries can be significantly faster.

47.131.3.15 map_local_to_global()

```
template<typename IndexType = int32>
array<index_type> gko::index_set< IndexType >::map_local_to_global (
    const array< index_type > & local_indices,
    const bool is_sorted = false ) const
```

This is an array version of the scalar function above.

Parameters

<i>local_indices</i>	the local index array.
<i>is_sorted</i>	a parameter that specifies if the query array is sorted or not. <code>true</code> if sorted .

Returns

the global index array from the index set.

Note

Whenever possible, passing a sorted array is preferred as the queries can be significantly faster.

Passing local indices from `[0, size)` is equivalent to using the `@to_global_indices` function.

47.131.3.16 operator=() [1/2]

```
template<typename IndexType = int32>
index_set& gko::index_set< IndexType >::operator= (
    const index_set< IndexType > & other ) [inline]
```

Copies data from another `index_set`.

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of `other`.

Parameters

<i>other</i>	the index_set to copy from
--------------	--

Returns

this

47.131.3.17 operator=() [2/2]

```
template<typename IndexType = int32>
index_set& gko::index_set< IndexType >::operator= (
    index_set< IndexType > && other ) [inline]
```

Moves data from another [index_set](#).

The executor of this is preserved. In case this does not have an assigned executor, it will inherit the executor of other.

Parameters

<i>other</i>	the index_set to move from
--------------	--

Returns

this

47.131.3.18 to_global_indices()

```
template<typename IndexType = int32>
array<index_type> gko::index_set< IndexType >::to_global_indices ( ) const
```

This function allows the user obtain a decompressed `global_indices` array from the indices stored in the index set.

Returns

the decompressed set of indices.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/index_set.hpp`

47.132 gko::InvalidStateError Class Reference

Exception thrown if an object is in an invalid state.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [InvalidStateError](#) (const std::string &file, int line, const std::string &func, const std::string &clarification)
Initializes an invalid state error.

47.132.1 Detailed Description

Exception thrown if an object is in an invalid state.

47.132.2 Constructor & Destructor Documentation

47.132.2.1 InvalidStateError()

```
gko::InvalidStateError::InvalidStateError (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & clarification ) [inline]
```

Initializes an invalid state error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The function name where the error occurred
<i>clarification</i>	A message describing the invalid state

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.133 gko::solver::lr< ValueType > Class Template Reference

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

```
#include <ginkgo/core/solver/ir.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `bool apply_uses_initial_guess ()` const override
Return true as iterative solvers use the data in x as an initial guess.
- `std::shared_ptr< const LinOp > get_solver ()` const
Returns the solver operator used as the inner solver.
- `void set_solver (std::shared_ptr< const LinOp > new_solver)`
Sets the solver operator used as the inner solver.
- `lr & operator= (const lr &)`
Copy-assigns an IR solver.
- `lr & operator= (lr &&)`
Move-assigns an IR solver.
- `lr (const lr &)`
Copy-constructs an IR solver.
- `lr (lr &&)`
Move-constructs an IR solver.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType >())`
Create the parameters from the property_tree.

47.133.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::solver::lr< ValueType >
```

Iterative refinement (IR) is an iterative method that uses another coarse method to approximate the error of the current solution via the current residual.

Moreover, it can be also considered as preconditioned Richardson iteration with relaxation factor = 1.

For any approximation of the solution `solution` to the system $Ax = b$, the residual is defined as: `residual = b - A solution`. The error in `solution`, $e = x - \text{solution}$ (with x being the exact solution) can be obtained as the solution to the residual equation $Ae = \text{residual}$, since $Ae = Ax - A \text{ solution} = b - A \text{ solution} = \text{residual}$. Then, the real solution is computed as $x = \text{relaxation_factor} * \text{solution} + e$. Instead of accurately solving the residual equation $Ae = \text{residual}$, the solution of the system e can be approximated to obtain the approximation `error` using a coarse method `solver`, which is used to update `solution`, and the entire process is repeated with the updated `solution`. This yields the iterative refinement method:

```
solution = initial_guess
while not converged:
    residual = b - A solution
    error = solver(A, residual)
    solution = solution + relaxation_factor * error
```

With `relaxation_factor` equal to 1 (default), the solver is Iterative Refinement, with `relaxation_factor` equal to a value other than 1, the solver is a Richardson iteration, with possibility for additional preconditioning.

Assuming that `solver` has accuracy `c`, i.e., $|e - \text{error}| \leq c |e|$, iterative refinement will converge with a convergence rate of `c`. Indeed, from $e - \text{error} = x - \text{solution} - \text{error} = x - \text{solution*}$ (where `solution*` denotes the value stored in `solution` after the update) and $e = \text{inv}(A) \text{ residual} = \text{inv}(A)b - \text{inv}(A) A \text{ solution} = x - \text{solution}$ it follows that $|x - \text{solution*}| \leq c |x - \text{solution}|$.

Unless otherwise specified via the `solver` factory parameter, this implementation uses the identity operator (i.e. the solver that approximates the solution of a system $Ax = b$ by setting $x := b$) as the default inner solver. Such a setting results in a relaxation method known as the Richardson iteration with parameter 1, which is guaranteed to converge for matrices whose spectrum is strictly contained within the unit disc around 1 (i.e., all its eigenvalues `lambda` have to satisfy the equation $|\text{relaxation_factor} * \text{lambda} - 1| < 1$).

Template Parameters

<i>ValueType</i>	precision of matrix elements
------------------	------------------------------

47.133.2 Constructor & Destructor Documentation

47.133.2.1 `Ir()` [1/2]

```
template<typename ValueType = default_precision>
gko::solver::Ir< ValueType >::Ir (
    const Ir< ValueType > & )
```

Copy-constructs an IR solver.

Inherits the executor, shallow-copies inner solver, stopping criterion and system matrix.

47.133.2.2 `Ir()` [2/2]

```
template<typename ValueType = default_precision>
gko::solver::Ir< ValueType >::Ir (
    Ir< ValueType > && )
```

Move-constructs an IR solver.

Preserves the executor, moves inner solver, stopping criterion and system matrix. The moved-from object is empty (0x0 and nullptr inner solver, stopping criterion and system matrix)

47.133.3 Member Function Documentation

47.133.3.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision>
bool gko::solver::lr< ValueType >::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess.

Returns

true as iterative solvers use the data in x as an initial guess.

References gko::solver::provided.

47.133.3.2 conj_transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::lr< ValueType >::conj_transpose ( ) const [override],
[virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.133.3.3 get_solver()

```
template<typename ValueType = default_precision>
std::shared_ptr<const LinOp> gko::solver::lr< ValueType >::get_solver ( ) const [inline]
```

Returns the solver operator used as the inner solver.

Returns

the solver operator used as the inner solver

47.133.3.4 operator=() [1/2]

```
template<typename ValueType = default_precision>
lr& gko::solver::lr< ValueType >::operator= (
    const lr< ValueType > & )
```

Copy-assigns an IR solver.

Preserves the executor, shallow-copies inner solver, stopping criterion and system matrix. If the executors mismatch, clones inner solver, stopping criterion and system matrix onto this executor.

47.133.3.5 operator=() [2/2]

```
template<typename ValueType = default_precision>
Ir& gko::solver::Ir< ValueType >::operator= (
    Ir< ValueType > && )
```

Move-assigns an IR solver.

Preserves the executor, moves inner solver, stopping criterion and system matrix. If the executors mismatch, clones inner solver, stopping criterion and system matrix onto this executor. The moved-from object is empty (0x0 and nullptr inner solver, stopping criterion and system matrix)

47.133.3.6 parse()

```
template<typename ValueType = default_precision>
static parameters_type gko::solver::Ir< ValueType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value type of this class.

Returns

parameters

47.133.3.7 set_solver()

```
template<typename ValueType = default_precision>
void gko::solver::Ir< ValueType >::set_solver (
    std::shared_ptr< const LinOp > new_solver )
```

Sets the solver operator used as the inner solver.

Parameters

<i>new_solver</i>	the new inner solver
-------------------	----------------------

47.133.3.8 transpose()

```
template<typename ValueType = default_precision>
std::unique_ptr<LinOp> gko::solver::Ir< ValueType >::transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/ir.hpp

47.134 gko::preconditioner::lsai< IsaiType, ValueType, IndexType > Class Template Reference

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

```
#include <ginkgo/core/preconditioner/isai.hpp>
```

Public Member Functions

- `std::shared_ptr< const typename std::conditional< IsaiType==isai_type::spd, Comp, Csr >::type > get_approximate_inverse () const`
Returns the approximate inverse of the given matrix (either a CSR matrix for IsaiType general, upper or lower or a composition of two CSR matrices for IsaiType spd).
- `lsai & operator= (const lsai &other)`
Copy-assigns an ISAI preconditioner.
- `lsai & operator= (lsai &&other)`
Move-assigns an ISAI preconditioner.
- `lsai (const lsai &other)`
Copy-constructs an ISAI preconditioner.
- `lsai (lsai &&other)`
Move-constructs an ISAI preconditioner.
- `std::unique_ptr< LinOp > transpose () const override`
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose () const override`
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Static Public Member Functions

- static parameters_type `parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

47.134.1 Detailed Description

```
template<isai_type IsaiType, typename ValueType, typename IndexType>
class gko::preconditioner::Isai< IsaiType, ValueType, IndexType >
```

The Incomplete Sparse Approximate Inverse (ISAI) Preconditioner generates an approximate inverse matrix for a given square matrix A, lower triangular matrix L, upper triangular matrix U or symmetric positive (spd) matrix B.

Using the preconditioner computes $aiA * x$, $aiU * x$, $aiL * x$ or $aiC^T * aiC * x$ (depending on the type of the `Isai`) for a given vector x (may have multiple right hand sides). `aiA`, `aiU` and `aiL` are the approximate inverses for A, U and L respectively. `aiC` is an approximation to C, the exact Cholesky factor of B (This is commonly referred to as a Factorized Sparse Approximate Inverse, short FSPAI).

The sparsity pattern used for the approximate inverse of A, L and U is the same as the sparsity pattern of the respective matrix. For B, the sparsity pattern used for the approximate inverse is the same as the sparsity pattern of the lower triangular half of B.

Note that, except for the spd case, for a matrix A generally $ISAI(A)^T \neq ISAI(A^T)$.

For more details on the algorithm, see the paper [Incomplete Sparse Approximate Inverses for Parallel Preconditioning](#), which is the basis for this work.

Note

GPU implementations can only handle the vector unit width `width` (warp size for CUDA) as number of elements per row in the sparse matrix. If there are more than `width` elements per row, the remaining elements will be ignored.

Template Parameters

<code>IsaiType</code>	determines if the ISAI is generated for a general square matrix, a lower triangular matrix, an upper triangular matrix or an spd matrix
<code>ValueType</code>	precision of matrix elements
<code>IndexType</code>	precision of matrix indexes

47.134.2 Constructor & Destructor Documentation

47.134.2.1 `Isai()` [1/2]

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::Isai (
    const Isai< IsaiType, ValueType, IndexType > & other )
```

Copy-constructs an ISAI preconditioner.

Inherits the executor, shallow-copies the matrix and parameters.

47.134.2.2 Isai() [2/2]

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::Isai (
    Isai< IsaiType, ValueType, IndexType > && other )
```

Move-constructs an ISAI preconditioner.

Inherits the executor, moves the matrix and parameters. The moved-from object is empty (0x0 with nullptr matrix and default parameters)

47.134.3 Member Function Documentation

47.134.3.1 conj_transpose()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
std::unique_ptr<LinOp> gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::conj_↔
transpose ( ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.134.3.2 get_approximate_inverse()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
std::shared_ptr<const typename std::conditional<IsaiType == isai_type::spd, Comp, Csr>::type>
gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::get_approximate_inverse ( ) const
[inline]
```

Returns the approximate inverse of the given matrix (either a CSR matrix for IsaiType general, upper or lower or a composition of two CSR matrices for IsaiType spd).

Returns

the generated approximate inverse

References [gko::as\(\)](#).

47.134.3.3 operator=() [1/2]

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
Isai& gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::operator= (
    const Isai< IsaiType, ValueType, IndexType > & other )
```

Copy-assigns an ISAI preconditioner.

Preserves the executor, shallow-copies the matrix and parameters. Creates a clone of the matrix if it is on the wrong executor.

47.134.3.4 operator=() [2/2]

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
Isai& gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::operator= (
    Isai< IsaiType, ValueType, IndexType > && other )
```

Move-assigns an ISAI preconditioner.

Preserves the executor, moves the matrix and parameters. Creates a clone of the matrix if it is on the wrong executor. The moved-from object is empty (0x0 with nullptr matrix and default parameters)

47.134.3.5 parse()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
static parameters_type gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

47.134.3.6 transpose()

```
template<isai_type IsaiType, typename ValueType , typename IndexType >
```

```
std::unique_ptr<LinOp> gko::preconditioner::Isai< IsaiType, ValueType, IndexType >::transpose
( ) const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/preconditioner/isai.hpp

47.135 gko::stop::Iteration Class Reference

The [Iteration](#) class is a stopping criterion which stops the iteration process after a preset number of iterations.

```
#include <ginkgo/core/stop/iteration.hpp>
```

47.135.1 Detailed Description

The [Iteration](#) class is a stopping criterion which stops the iteration process after a preset number of iterations.

Note

to use this stopping criterion, it is required to update the iteration count for the `::check()` method.

The documentation for this class was generated from the following file:

- ginkgo/core/stop/iteration.hpp

47.136 gko::log::iteration_complete_data Struct Reference

Struct representing iteration complete related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.136.1 Detailed Description

Struct representing iteration complete related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.137 gko::solver::IterativeBase Class Reference

A LinOp implementing this interface stores a stopping criterion factory.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

Public Member Functions

- `std::shared_ptr< const stop::CriterionFactory > get\_stop\_criterion\_factory () const`
Gets the stopping criterion factory of the solver.
- `virtual void set\_stop\_criterion\_factory (std::shared_ptr< const stop::CriterionFactory > new_stop_factory)`
Sets the stopping criterion of the solver.

47.137.1 Detailed Description

A LinOp implementing this interface stores a stopping criterion factory.

47.137.2 Member Function Documentation

47.137.2.1 `get_stop_criterion_factory()`

```
std::shared_ptr<const stop::CriterionFactory> gko::solver::IterativeBase::get_stop_criterion↔  
_factory ( ) const [inline]
```

Gets the stopping criterion factory of the solver.

Returns

the stopping criterion factory

Referenced by `gko::solver::EnableIterativeBase< Bicg< ValueType > >::operator=()`.

47.137.2.2 `set_stop_criterion_factory()`

```
virtual void gko::solver::IterativeBase::set_stop_criterion_factory (  
    std::shared_ptr< const stop::CriterionFactory > new_stop_factory ) [inline],  
[virtual]
```

Sets the stopping criterion of the solver.

Parameters

<i>other</i>	the new stopping criterion factory
--------------	------------------------------------

Reimplemented in [gko::solver::EnableIterativeBase< DerivedType >](#), [gko::solver::EnableIterativeBase< Bicgstab< ValueType > >](#), [gko::solver::EnableIterativeBase< Multigrid >](#), [gko::solver::EnableIterativeBase< Gmres< ValueType > >](#), [gko::solver::EnableIterativeBase< CbGmres< ValueType > >](#), [gko::solver::EnableIterativeBase< Ir< ValueType > >](#), [gko::solver::EnableIterativeBase< Fcg< ValueType > >](#), [gko::solver::EnableIterativeBase< Idr< ValueType > >](#), [gko::solver::EnableIterativeBase< Cgs< ValueType > >](#), [gko::solver::EnableIterativeBase< Gcr< ValueType > >](#), [gko::solver::EnableIterativeBase< Cg< ValueType > >](#), and [gko::solver::EnableIterativeBase< Bicg< ValueType > >](#).

Referenced by [gko::solver::EnableIterativeBase< Bicg< ValueType > >::set_stop_criterion_factory\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/solver/solver_base.hpp](#)

47.138 gko::preconditioner::Jacobi< ValueType, IndexType > Class Template Reference

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

```
#include <ginkgo/core/preconditioner/jacobi.hpp>
```

Public Member Functions

- [size_type get_num_blocks](#) () const noexcept
Returns the number of blocks of the operator.
- const [block_interleaved_storage_scheme](#)< index_type > & [get_storage_scheme](#) () const noexcept
Returns the storage scheme used for storing [Jacobi](#) blocks.
- const value_type * [get_blocks](#) () const noexcept
Returns the pointer to the memory used for storing the block data.
- const [remove_complex](#)< value_type > * [get_conditioning](#) () const noexcept
Returns an array of 1-norm condition numbers of the blocks.
- [size_type get_num_stored_elements](#) () const noexcept
Returns the number of elements explicitly stored in the matrix.
- void [convert_to](#) ([matrix::Dense](#)< value_type > *result) const override
Converts the implementer to an object of type result_type.
- void [move_to](#) ([matrix::Dense](#)< value_type > *result) override
Converts the implementer to an object of type result_type by moving data from this object.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- [Jacobi](#) & operator= (const [Jacobi](#) &other)
Copy-assigns a [Jacobi](#) preconditioner.
- [Jacobi](#) & operator= ([Jacobi](#) &&other)
Move-assigns a [Jacobi](#) preconditioner.
- [Jacobi](#) (const [Jacobi](#) &other)
Copy-constructs a [Jacobi](#) preconditioner.
- [Jacobi](#) ([Jacobi](#) &&other)
Move-assigns a [Jacobi](#) preconditioner.

Static Public Member Functions

- static parameters_type `parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

47.138.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::preconditioner::Jacobi< ValueType, IndexType >
```

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks of the source operator.

The `Jacobi` class implements the inversion of the diagonal blocks using Gauss-Jordan elimination with column pivoting, and stores the inverse explicitly in a customized format.

If the diagonal blocks of the matrix are not explicitly set by the user, the implementation will try to automatically detect the blocks by first finding the natural blocks of the matrix, and then applying the supervariable agglomeration procedure on them. However, if problem-specific knowledge regarding the block diagonal structure is available, it is usually beneficial to explicitly pass the starting rows of the diagonal blocks, as the block detection is merely a heuristic and cannot perfectly detect the diagonal block structure. The current implementation supports blocks of up to 32 rows / columns.

The implementation also includes an improved, adaptive version of the block-Jacobi preconditioner, which can store some of the blocks in lower precision and thus improve the performance of preconditioner application by reducing the amount of memory transfers. This variant can be enabled by setting the `Jacobi::Factory`'s `storage_optimization` parameter. Refer to the documentation of the parameter for more details.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	integral type used to store pointers to the start of each block

Note

The current implementation supports blocks of up to 32 rows / columns.

When using the adaptive variant, there may be a trade-off in terms of slightly longer preconditioner generation due to extra work required to detect the optimal precision of the blocks.

When the `max_block_size` is set to 1, specialized kernels are used, both for generation (inverting the diagonals) and application (diagonal scaling) to reduce the overhead involved in the usual (adaptive) block case.

47.138.2 Constructor & Destructor Documentation

47.138.2.1 `Jacobi()` [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::preconditioner::Jacobi< ValueType, IndexType >::Jacobi (
    const Jacobi< ValueType, IndexType > & other )
```

Copy-constructs a [Jacobi](#) preconditioner.

Inherits executor, copies all data and parameters.

47.138.2.2 Jacobi() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::preconditioner::Jacobi< ValueType, IndexType >::Jacobi (
    Jacobi< ValueType, IndexType > && other )
```

Move-assigns a [Jacobi](#) preconditioner.

Inherits executor, moves all data and parameters. The moved-from object will be empty (0x0 and default parameters).

47.138.3 Member Function Documentation

47.138.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Jacobi< ValueType, IndexType >::conj_transpose (
    ) const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.138.3.2 convert_to()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::preconditioner::Jacobi< ValueType, IndexType >::convert_to (
    matrix::Dense< value_type > * result ) const [override], [virtual]
```

Converts the implementer to an object of type result_type.

Parameters

<i>result</i>	the object used to store the result of the conversion
---------------	---

Implements [gko::ConvertibleTo< matrix::Dense< ValueType > >](#).

47.138.3.3 get_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::preconditioner::Jacobi< ValueType, IndexType >::get_blocks ( ) const
[inline], [noexcept]
```

Returns the pointer to the memory used for storing the block data.

Element (i, j) of block b is stored in position (get_block_pointers()[b] + i) * stride + j of the array.

Returns

the pointer to the memory used for storing the block data

References gko::array< ValueType >::get_const_data().

47.138.3.4 get_conditioning()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const remove_complex<value_type>* gko::preconditioner::Jacobi< ValueType, IndexType >::get_↵
conditioning ( ) const [inline], [noexcept]
```

Returns an array of 1-norm condition numbers of the blocks.

Returns

an array of 1-norm condition numbers of the blocks

Note

This value is valid only if adaptive precision variant is used, and implementations of the standard non-adaptive variant are allowed to omit the calculation of condition numbers.

References gko::array< ValueType >::get_const_data().

47.138.3.5 get_num_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::preconditioner::Jacobi< ValueType, IndexType >::get_num_blocks ( ) const [inline],
[noexcept]
```

Returns the number of blocks of the operator.

Returns

the number of blocks of the operator

47.138.3.6 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored_elements ( )
const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.138.3.7 get_storage_scheme()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const block_interleaved_storage_scheme<index_type>& gko::preconditioner::Jacobi< ValueType,
IndexType >::get_storage_scheme ( ) const [inline], [noexcept]
```

Returns the storage scheme used for storing [Jacobi](#) blocks.

Returns

the storage scheme used for storing [Jacobi](#) blocks

47.138.3.8 move_to()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::preconditioner::Jacobi< ValueType, IndexType >::move_to (
    matrix::Dense< value_type > * result ) [override], [virtual]
```

Converts the implementer to an object of type `result_type` by moving data from this object.

This method is used when the implementer is a temporary object, and move semantics can be used.

Parameters

<i>result</i>	the object used to replace the result of the conversion
---------------	---

Note

[ConvertibleTo::move_to](#) can be implemented by simply calling [ConvertibleTo::convert_to](#). However, this operation can often be optimized by exploiting the fact that implementer's data can be moved to the result.

Implements [gko::ConvertibleTo< matrix::Dense< ValueType > >](#).

47.138.3.9 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Jacobi& gko::preconditioner::Jacobi< ValueType, IndexType >::operator= (
    const Jacobi< ValueType, IndexType > & other )
```

Copy-assigns a [Jacobi](#) preconditioner.

Preserves executor, copies all data and parameters.

47.138.3.10 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Jacobi& gko::preconditioner::Jacobi< ValueType, IndexType >::operator= (
    Jacobi< ValueType, IndexType > && other )
```

Move-assigns a [Jacobi](#) preconditioner.

Preserves executor, moves all data and parameters. The moved-from object will be empty (0x0 and default parameters).

47.138.3.11 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::preconditioner::Jacobi< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

Note

[Jacobi](#) does not support block_pointers and storage_optimization array.

47.138.3.12 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::preconditioner::Jacobi< ValueType, IndexType >::transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.138.3.13 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::preconditioner::Jacobi< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/preconditioner/jacobi.hpp

47.139 gko::batch::preconditioner::Jacobi< ValueType, IndexType > Class Template Reference

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks (stored in a dense row major fashion) of the source operator.

```
#include <ginkgo/core/preconditioner/batch_jacobi.hpp>
```

Public Member Functions

- const index_type * [get_const_block_pointers](#) () const noexcept
Returns the block pointers.
- const index_type * [get_const_map_block_to_row](#) () const noexcept
Returns the mapping between the blocks and the row id.

- `const index_type * get\_const\_blocks\_cumulative\_offsets ()` `const noexcept`
Returns the cumulative blocks storage array.
- `uint32 get\_max\_block\_size ()` `const noexcept`
Returns the max block size.
- `size_type get\_num\_blocks ()` `const noexcept`
Returns the number of blocks in an individual batch entry.
- `const value_type * get\_const\_blocks ()` `const noexcept`
Returns the pointer to the memory used for storing the block data.
- `size_type get\_num\_stored\_elements ()` `const noexcept`
Returns the number of elements explicitly stored in the dense blocks.

47.139.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::batch::preconditioner::Jacobi< ValueType, IndexType >
```

A block-Jacobi preconditioner is a block-diagonal linear operator, obtained by inverting the diagonal blocks (stored in a dense row major fashion) of the source operator.

With the batched preconditioners, it is required that all items in the batch have the same sparsity pattern. The detection of the blocks and the block pointers require that the sparsity pattern of all the items be the same. Other cases is undefined behaviour. The input batch matrix must be in `batch::Csr` matrix format or must be convertible to `batch::Csr` matrix format. The block detection algorithm and the conversion to dense blocks kernels require this assumption.

Note

In a fashion similar to the non-batched [Jacobi](#) preconditioner, the maximum possible size of the diagonal blocks is equal to the maximum warp size on the device (32 for NVIDIA GPUs, 64 for AMD GPUs).

Template Parameters

<i>ValueType</i>	value precision of matrix elements
<i>IndexType</i>	index precision of matrix elements

47.139.2 Member Function Documentation

47.139.2.1 `get_const_block_pointers()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_↵
block_pointers ( ) const [inline], [noexcept]
```

Returns the block pointers.

Note

Returns nullptr in case of a scalar jacobi preconditioner (max_block_size = 1).

Returns

the block pointers

```
68     {
69         return block_pointers_.get_const_data();
70     }
```

References gko::array< ValueType >::get_const_data().

47.139.2.2 get_const_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_blocks
( ) const [inline], [noexcept]
```

Returns the pointer to the memory used for storing the block data.

Element (i, j) of the block, which belongs to the batch entry with index = "batch_id" and has local id = "block_id" within its batch entry is stored at the address = `get_const_blocks()` + `detail::get_global_block_offset(batch_id, num_blocks, block_id, cumulative_blocks_storage) + i * detail::get_stride(block_id, block_pointers) + j`

Note

Returns nullptr in case of a scalar jacobi preconditioner (max_block_size = 1). The blocks array is empty in case of scalar jacobi preconditioner as the preconditioner is generated inside the batched solver kernel.

Returns

the pointer to the memory used for storing the block data

References gko::array< ValueType >::get_const_data().

47.139.2.3 get_const_blocks_cumulative_offsets()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_↵
blocks_cumulative_offsets ( ) const [inline], [noexcept]
```

Returns the cumulative blocks storage array.

Note

Returns nullptr in case of a scalar jacobi preconditioner (max_block_size = 1).

References gko::array< ValueType >::get_const_data().

47.139.2.4 get_const_map_block_to_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_const_map_block_to_row ( ) const [inline], [noexcept]
```

Returns the mapping between the blocks and the row id.

Note

Returns nullptr in case of a scalar jacobi preconditioner (max_block_size = 1).

References gko::array< ValueType >::get_const_data().

47.139.2.5 get_max_block_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
uint32 gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_max_block_size ( )
const [inline], [noexcept]
```

Returns the max block size.

Returns

the max block size

47.139.2.6 get_num_blocks()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_num_blocks ( ) const
[inline], [noexcept]
```

Returns the number of blocks in an individual batch entry.

Returns

the number of blocks in an individual batch entry.

47.139.2.7 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::batch::preconditioner::Jacobi< ValueType, IndexType >::get_num_stored_elements
( ) const [inline], [noexcept]
```

Returns the number of elements explicitly stored in the dense blocks.

Note

Returns 0 in case of scalar jacobi preconditioner as the preconditioner is generated inside the batched solver kernels, hence, blocks array storage is not required.

Returns

the number of elements explicitly stored in the dense blocks.

References `gko::array< ValueType >::get_size()`.

The documentation for this class was generated from the following file:

- `ginkgo/core/preconditioner/batch_jacobi.hpp`

47.140 gko::KernelNotFound Class Reference

[KernelNotFound](#) is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [KernelNotFound](#) (const std::string &file, int line, const std::string &func)
Initializes a [KernelNotFound](#) error.

47.140.1 Detailed Description

[KernelNotFound](#) is thrown if Ginkgo cannot find a kernel which satisfies the criteria imposed by the input arguments.

47.140.2 Constructor & Destructor Documentation

47.140.2.1 KernelNotFound()

```
gko::KernelNotFound::KernelNotFound (
    const std::string & file,
    int line,
    const std::string & func ) [inline]
```

Initializes a [KernelNotFound](#) error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the function where the error occurred

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.141 gko::log::linop_data Struct Reference

Struct representing LinOp related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.141.1 Detailed Description

Struct representing LinOp related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.142 gko::log::linop_factory_data Struct Reference

Struct representing LinOp factory related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.142.1 Detailed Description

Struct representing LinOp factory related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.143 gko::LinOpFactory Class Reference

A [LinOpFactory](#) represents a higher order mapping which transforms one linear operator into another.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Additional Inherited Members

47.143.1 Detailed Description

A [LinOpFactory](#) represents a higher order mapping which transforms one linear operator into another.

In Ginkgo, every linear solver is viewed as a mapping. For example, given an s.p.d linear system $Ax = b$, the solution $x = A^{-1}b$ can be computed using the CG method. This algorithm can be represented in terms of linear operators and mappings between them as follows:

- A `Cg::Factory` is a higher order mapping which, given an input operator A , returns a new linear operator A^{-1} stored in "CG format"
- Storing the operator A^{-1} in "CG format" means that the data structure used to store the operator is just a simple pointer to the original matrix A . The application $x = A^{-1}b$ of such an operator can then be implemented by solving the linear system $Ax = b$ using the CG method. This is achieved in code by having a special class for each of those "formats" (e.g. the "Cg" class defines such a format for the CG solver).

Another example of a [LinOpFactory](#) is a preconditioner. A preconditioner for a linear operator A is a linear operator M^{-1} , which approximates A^{-1} . In addition, it is stored in a way such that both the data of M^{-1} is cheap to compute from A , and the operation $x = M^{-1}b$ can be computed quickly. These operators are useful to accelerate the convergence of Krylov solvers. Thus, a preconditioner also fits into the [LinOpFactory](#) framework:

- The factory maps a linear operator A into a preconditioner M^{-1} which is stored in suitable format (e.g. as a product of two factors in case of ILU preconditioners).
- The resulting linear operator implements the application operation $x = M^{-1}b$ depending on the format the preconditioner is stored in (e.g. as two triangular solves in case of ILU)

47.143.1.1 Example: using CG in Ginkgo

```
{c++}
// Suppose A is a matrix, b a rhs vector, and x an initial guess
// Create a CG which runs for at most 1000 iterations, and stops after
// reducing the residual norm by 6 orders of magnitude
auto cg_factory = solver::Cg<>::build()
    .with_max_iters(1000)
    .with_rel_residual_goal(1e-6)
    .on(cuda);
// create a linear operator which represents the solver
auto cg = cg_factory->generate(A);
// solve the system
cg->apply(b, x);
```

The documentation for this class was generated from the following file:

- `ginkgo/core/base/lin_op.hpp`

47.144 gko::matrix::Csr< ValueType, IndexType >::load_balance Class Reference

[load_balance](#) is a [strategy_type](#) which uses the load balance algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [load_balance](#) ()
Creates a [load_balance](#) strategy.
- [load_balance](#) (std::shared_ptr< const [CudaExecutor](#) > exec)
Creates a [load_balance](#) strategy with CUDA executor.
- [load_balance](#) (std::shared_ptr< const [HipExecutor](#) > exec)
Creates a [load_balance](#) strategy with HIP executor.
- [load_balance](#) (std::shared_ptr< const [DpcppExecutor](#) > exec)
Creates a [load_balance](#) strategy with DPCPP executor.
- [load_balance](#) (int64_t nwarps, int warp_size=32, bool cuda_strategy=true, std::string strategy_name="none")
Creates a [load_balance](#) strategy with specified parameters.
- void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow) override
Computes srow according to row pointers.
- int64_t [clac_size](#) (const int64_t nnz) override
Computes the srow size according to the number of nonzeros.
- std::shared_ptr< [strategy_type](#) > [copy](#) () override
Copy a strategy.

47.144.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::load_balance
```

[load_balance](#) is a [strategy_type](#) which uses the load balance algorithm.

47.144.2 Constructor & Destructor Documentation

47.144.2.1 [load_balance](#)() [1/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance ( ) [inline]
```

Creates a [load_balance](#) strategy.

Warning

this is deprecated! Please rely on the new automatic strategy instantiation or use one of the other constructors.

47.144.2.2 [load_balance](#)() [2/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance (
    std::shared_ptr< const CudaExecutor > exec ) [inline]
```

Creates a [load_balance](#) strategy with CUDA executor.

Parameters

<i>exec</i>	the CUDA executor
-------------	-------------------

47.144.2.3 load_balance() [3/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance (
    std::shared_ptr< const HipExecutor > exec ) [inline]
```

Creates a [load_balance](#) strategy with HIP executor.

Parameters

<i>exec</i>	the HIP executor
-------------	------------------

47.144.2.4 load_balance() [4/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance (
    std::shared_ptr< const DpcppExecutor > exec ) [inline]
```

Creates a [load_balance](#) strategy with DPCPP executor.

Parameters

<i>exec</i>	the DPCPP executor
-------------	--------------------

Note

TODO: porting - we hardcode the subgroup size is 32

47.144.2.5 load_balance() [5/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::load_balance::load_balance (
    int64_t nwarps,
    int warp_size = 32,
    bool cuda_strategy = true,
    std::string strategy_name = "none" ) [inline]
```

Creates a [load_balance](#) strategy with specified parameters.

Parameters

<i>nwarps</i>	the number of warps in the executor
<i>warp_size</i>	the warp size of the executor
<i>cuda_strategy</i>	whether the <code>cuda_strategy</code> needs to be used.

Note

The `warp_size` must be the size of full warp. When using this constructor, `set_strategy` needs to be called with correct parameters which is replaced during the conversion.

47.144.3 Member Function Documentation

47.144.3.1 `clac_size()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
int64_t gko::matrix::Csr< ValueType, IndexType >::load_balance::clac_size (
    const int64_t nnz ) [inline], [override], [virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

References [gko::ceildiv\(\)](#), and [gko::min\(\)](#).

47.144.3.2 `copy()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::load_balance::copy (
) [inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.144.3.3 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::load_balance::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [inline], [override], [virtual]
```

Computes srow according to row pointers.

Parameters

<i>mtx_row_ptrs</i>	the row pointers of the matrix
<i>mtx_srow</i>	the srow of the matrix

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

References [gko::ceildiv\(\)](#), [gko::array< ValueType >::get_const_data\(\)](#), [gko::array< ValueType >::get_data\(\)](#), [gko::array< ValueType >::get_executor\(\)](#), and [gko::array< ValueType >::get_size\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/csr.hpp](#)

47.145 gko::log::Loggable Class Reference

[Loggable](#) class is an interface which should be implemented by classes wanting to support logging.

```
#include <ginkgo/core/log/logger.hpp>
```

Public Member Functions

- virtual void [add_logger](#) (std::shared_ptr< const Logger > logger)=0
Adds a new logger to the list of subscribed loggers.
- virtual void [remove_logger](#) (const Logger *logger)=0
Removes a logger from the list of subscribed loggers.
- virtual const std::vector< std::shared_ptr< const Logger > > & [get_loggers](#) () const =0
Returns the vector containing all loggers registered at this object.
- virtual void [clear_loggers](#) ()=0
Remove all loggers registered at this object.

47.145.1 Detailed Description

[Loggable](#) class is an interface which should be implemented by classes wanting to support logging.

For most cases, one can rely on the [EnableLogging](#) mixin which provides a default implementation of this interface.

47.145.2 Member Function Documentation

47.145.2.1 add_logger()

```
virtual void gko::log::Loggable::add_logger (
    std::shared_ptr< const Logger > logger ) [pure virtual]
```

Adds a new logger to the list of subscribed loggers.

Parameters

<i>logger</i>	the logger to add
---------------	-------------------

Implemented in [gko::Executor](#).

47.145.2.2 get_loggers()

```
virtual const std::vector<std::shared_ptr<const Logger> >& gko::log::Loggable::get_loggers (
) const [pure virtual]
```

Returns the vector containing all loggers registered at this object.

Returns

the vector containing all registered loggers.

47.145.2.3 remove_logger()

```
virtual void gko::log::Loggable::remove_logger (
    const Logger * logger ) [pure virtual]
```

Removes a logger from the list of subscribed loggers.

Parameters

<i>logger</i>	the logger to remove
---------------	----------------------

Note

The comparison is done using the logger's object unique identity. Thus, two loggers constructed in the same way are not considered equal.

Implemented in [gko::Executor](#).

The documentation for this class was generated from the following file:

- ginkgo/core/log/logger.hpp

47.146 gko::log::Record::logged_data Struct Reference

Struct storing the actually logged data.

```
#include <ginkgo/core/log/record.hpp>
```

47.146.1 Detailed Description

Struct storing the actually logged data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.147 gko::solver::LowerTrs< ValueType, IndexType > Class Template Reference

[LowerTrs](#) is the triangular solver which solves the system $Lx = b$, when L is a lower triangular matrix.

```
#include <ginkgo/core/solver/triangular.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- `LowerTrs (const LowerTrs &)`
Copy-assigns a triangular solver.
- `LowerTrs (LowerTrs &&)`
Move-assigns a triangular solver.
- `LowerTrs & operator= (const LowerTrs &)`
Copy-constructs a triangular solver.
- `LowerTrs & operator= (LowerTrs &&)`
Move-constructs a triangular solver.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType, IndexType >())`
Create the parameters from the property_tree.

47.147.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::solver::LowerTrs< ValueType, IndexType >
```

[LowerTrs](#) is the triangular solver which solves the system $Lx = b$, when L is a lower triangular matrix.

It works best when passing in a matrix in CSR format. If the matrix is not in CSR, then the generate step converts it into a CSR matrix. The generation fails if the matrix is not convertible to CSR.

Note

As the constructor uses the copy and convert functionality, it is not possible to create a empty solver or a solver with a matrix in any other format other than CSR, if none of the executor modules are being compiled with.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indices

47.147.2 Constructor & Destructor Documentation**47.147.2.1 LowerTrs() [1/2]**

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::solver::LowerTrs< ValueType, IndexType >::LowerTrs (
    const LowerTrs< ValueType, IndexType > & )
```

Copy-assigns a triangular solver.

Preserves the executor, shallow-copies the system matrix. If the executors mismatch, clones system matrix onto this executor. Solver analysis information will be regenerated.

47.147.2.2 LowerTrs() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::solver::LowerTrs< ValueType, IndexType >::LowerTrs (
    LowerTrs< ValueType, IndexType > && )
```

Move-assigns a triangular solver.

Preserves the executor, moves the system matrix. If the executors mismatch, clones system matrix onto this executor and regenerates solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

47.147.3 Member Function Documentation**47.147.3.1 conj_transpose()**

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::LowerTrs< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.147.3.2 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
LowerTrs& gko::solver::LowerTrs< ValueType, IndexType >::operator= (
    const LowerTrs< ValueType, IndexType > & )
```

Copy-constructs a triangular solver.

Preserves the executor, shallow-copies the system matrix. Solver analysis information will be regenerated.

47.147.3.3 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
LowerTrs& gko::solver::LowerTrs< ValueType, IndexType >::operator= (
    LowerTrs< ValueType, IndexType > && )
```

Move-constructs a triangular solver.

Preserves the executor, moves the system matrix and solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

47.147.3.4 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::solver::LowerTrs< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

47.147.3.5 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::LowerTrs< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/triangular.hpp`

47.148 `gko::experimental::factorization::Lu< ValueType, IndexType >` Class Template Reference

Computes an LU factorization of a sparse matrix.

```
#include <ginkgo/core/factorization/lu.hpp>
```

Public Member Functions

- `const parameters_type & get_parameters ()`
Returns the parameters used to construct the factory.
- `const parameters_type & get_parameters () const`
Returns the parameters used to construct the factory.
- `std::unique_ptr< factorization_type > generate (std::shared_ptr< const LinOp > system_matrix) const`

Static Public Member Functions

- `static parameters_type build ()`
Creates a new parameter_type to set up the factory.
- `static parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType, IndexType >())`
Create the parameters from the property_tree.

47.148.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::experimental::factorization::Lu< ValueType, IndexType >
```

Computes an LU factorization of a sparse matrix.

This [LinOpFactory](#) returns a [Factorization](#) storing the L and U factors for the provided system matrix in [matrix::Csr](#) format. If no symbolic factorization is provided, it will be computed first. It expects all fill-in entries to be present in the symbolic factorization. If the symbolic factorization is missing some entries, please refer to `llu`.

Template Parameters

<i>ValueType</i>	the type used to store values of the system matrix
<i>IndexType</i>	the type used to store sparsity pattern indices of the system matrix

47.148.2 Member Function Documentation

47.148.2.1 generate()

```
template<typename ValueType , typename IndexType >
std::unique_ptr<factorization_type> gko::experimental::factorization::Lu< ValueType, IndexType >::generate (
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a [Factorization](#) instead of a generic LinOp, which would need to be cast to [Factorization](#) again to access its factors. It is only necessary because smart pointers aren't covariant.

47.148.2.2 get_parameters() [1/2]

```
template<typename ValueType , typename IndexType >
const parameters_type& gko::experimental::factorization::Lu< ValueType, IndexType >::get_parameters ( ) [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
109 { return parameters_; }
```

47.148.2.3 get_parameters() [2/2]

```
template<typename ValueType , typename IndexType >
const parameters_type& gko::experimental::factorization::Lu< ValueType, IndexType >::get_parameters ( ) const [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

47.148.2.4 parse()

```
template<typename ValueType , typename IndexType >
static parameters_type gko::experimental::factorization::Lu< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/lu.hpp

47.149 gko::machine_topology Class Reference

The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.

```
#include <ginkgo/core/base/machine_topology.hpp>
```

Public Member Functions

- void [bind_to_cores](#) (const std::vector< int > &ids, const bool singlify=true) const
Bind the calling process to the CPU cores associated with the ids.
- void [bind_to_core](#) (const int &id) const
Bind to a single core.
- void [bind_to_pus](#) (const std::vector< int > &ids, const bool singlify=true) const
Bind the calling process to PUs associated with the ids.
- void [bind_to_pu](#) (const int &id) const
Bind to a Processing unit (PU)
- const normal_obj_info * [get_pu](#) (size_type id) const
Get the object of type PU associated with the id.
- const normal_obj_info * [get_core](#) (size_type id) const
Get the object of type core associated with the id.

- `const io_obj_info * get_pci_device (size_type id) const`
Get the object of type pci device associated with the id.
- `const io_obj_info * get_pci_device (const std::string &pci_bus_id) const`
Get the object of type pci device associated with the PCI bus id.
- `size_type get_num_pus () const`
Get the number of PU objects stored in this Topology tree.
- `size_type get_num_cores () const`
Get the number of core objects stored in this Topology tree.
- `size_type get_num_pci_devices () const`
Get the number of PCI device objects stored in this Topology tree.
- `size_type get_num_numas () const`
Get the number of NUMA objects stored in this Topology tree.

Static Public Member Functions

- static `machine_topology * get_instance ()`
Returns an instance of the [machine_topology](#) object.

47.149.1 Detailed Description

The machine topology class represents the hierarchical topology of a machine, including NUMA nodes, cores and PCI Devices.

Various information of the machine are gathered with the help of the Hardware Locality library (hwloc).

This class also provides functionalities to bind objects in the topology to the execution objects. Binding can enhance performance by allowing data to be closer to the executing object.

See the hwloc documentation (<https://www.open-mpi.org/projects/hwloc/doc/>) for more detailed information on topology detection and binding interfaces.

Note

A global object of `machine_topology` type is created in a thread safe manner and only destroyed at the end of the program. This means that any subsequent queries will be from the same global object and hence use an extra atomic read.

47.149.2 Member Function Documentation

47.149.2.1 bind_to_core()

```
void gko::machine_topology::bind_to_core (
    const int & id ) const [inline]
```

Bind to a single core.

Parameters

<i>ids</i>	The ids of the core to be bound to the calling process.
------------	---

```

212     {
213         machine_topology::get_instance()->bind_to_cores(std::vector<int>{id});
214     }

```

References `bind_to_cores()`, and `get_instance()`.

47.149.2.2 bind_to_cores()

```

void gko::machine_topology::bind_to_cores (
    const std::vector< int > & ids,
    const bool singlify = true ) const [inline]

```

Bind the calling process to the CPU cores associated with the ids.

Parameters

<i>ids</i>	The ids of cores to be bound.
<i>singlify</i>	The ids of PUs are singlified to prevent possibly expensive migrations by the OS. This means that the binding is performed for only one of the ids in the set of ids passed in. See hwloc doc for <i>singlify</i>

Referenced by `bind_to_core()`.

47.149.2.3 bind_to_pu()

```

void gko::machine_topology::bind_to_pu (
    const int & id ) const [inline]

```

Bind to a Processing unit (PU)

Parameters

<i>ids</i>	The ids of PUs to be bound to the calling process.
------------	--

References `bind_to_pus()`, and `get_instance()`.

47.149.2.4 bind_to_pus()

```

void gko::machine_topology::bind_to_pus (
    const std::vector< int > & ids,
    const bool singlify = true ) const [inline]

```

Bind the calling process to PUs associated with the ids.

Parameters

<i>ids</i>	The ids of PUs to be bound.
<i>singlify</i>	The ids of PUs are singlified to prevent possibly expensive migrations by the OS. This means that the binding is performed for only one of the ids in the set of ids passed in. See hwloc doc for singlify

Referenced by `bind_to_pu()`.

47.149.2.5 get_core()

```
const normal_obj_info* gko::machine_topology::get_core (
    size_type id ) const [inline]
```

Get the object of type core associated with the id.

Parameters

<i>id</i>	The id of the core
-----------	--------------------

Returns

the core object struct.

47.149.2.6 get_instance()

```
static machine_topology* gko::machine_topology::get_instance ( ) [inline], [static]
```

Returns an instance of the [machine_topology](#) object.

Returns

the [machine_topology](#) instance

Referenced by `bind_to_core()`, and `bind_to_pu()`.

47.149.2.7 get_num_cores()

```
size_type gko::machine_topology::get_num_cores ( ) const [inline]
```

Get the number of core objects stored in this Topology tree.

Returns

the number of cores.

47.149.2.8 get_num_numas()

```
size_type gko::machine_topology::get_num_numas ( ) const [inline]
```

Get the number of NUMA objects stored in this Topology tree.

Returns

the number of NUMA objects.

47.149.2.9 get_num_pci_devices()

```
size_type gko::machine_topology::get_num_pci_devices ( ) const [inline]
```

Get the number of PCI device objects stored in this Topology tree.

Returns

the number of PCI devices.

47.149.2.10 get_num_pus()

```
size_type gko::machine_topology::get_num_pus ( ) const [inline]
```

Get the number of PU objects stored in this Topology tree.

Returns

the number of PUs.

47.149.2.11 get_pci_device() [1/2]

```
const io_obj_info* gko::machine_topology::get_pci_device (
    const std::string & pci_bus_id ) const
```

Get the object of type pci device associated with the PCI bus id.

Parameters

<i>pci_bus_id</i>	The PCI bus id of the pci device
-------------------	----------------------------------

Returns

the PCI object struct.

47.149.2.12 get_pci_device() [2/2]

```
const io_obj_info* gko::machine_topology::get_pci_device (
    size_type id ) const [inline]
```

Get the object of type pci device associated with the id.

Parameters

<i>id</i>	The id of the pci device
-----------	--------------------------

Returns

the PCI object struct.

47.149.2.13 get_pu()

```
const normal_obj_info* gko::machine_topology::get_pu (
    size_type id ) const [inline]
```

Get the object of type PU associated with the id.

Parameters

<i>id</i>	The id of the PU
-----------	------------------

Returns

the PU object struct.

The documentation for this class was generated from the following file:

- ginkgo/core/base/machine_topology.hpp

47.150 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference

The [Matrix](#) class defines a (MPI-)distributed matrix.

```
#include <ginkgo/core/distributed/matrix.hpp>
```

Public Member Functions

- void [read_distributed](#) (const [device_matrix_data](#)< value_type, global_index_type > &data, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> partition, [assembly_mode](#) assembly_type=assembly_mode::local_only)
Reads a square matrix from the [device_matrix_data](#) structure and a global partition.
- void [read_distributed](#) (const [matrix_data](#)< value_type, global_index_type > &data, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> partition, [assembly_mode](#) assembly_type=assembly_mode::local_only)
Reads a square matrix from the [matrix_data](#) structure and a global partition.
- void [read_distributed](#) (const [device_matrix_data](#)< value_type, global_index_type > &data, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> row_partition, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> col_partition, [assembly_mode](#) assembly_type=assembly_mode::local_only)
Reads a matrix from the [device_matrix_data](#) structure, a global row partition, and a global column partition.
- void [read_distributed](#) (const [matrix_data](#)< value_type, global_index_type > &data, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> row_partition, std::shared_ptr< const [Partition](#)< local_index_type, global_index_type >> col_partition, [assembly_mode](#) assembly_type=assembly_mode::local_only)
Reads a matrix from the [matrix_data](#) structure, a global row partition, and a global column partition.
- std::shared_ptr< const LinOp > [get_local_matrix](#) () const
Get read access to the stored local matrix.
- std::shared_ptr< const LinOp > [get_non_local_matrix](#) () const
Get read access to the stored non-local matrix.
- [Matrix](#) (const [Matrix](#) &other)
Copy constructs a [Matrix](#).
- [Matrix](#) ([Matrix](#) &&other) noexcept
Move constructs a [Matrix](#).
- [Matrix](#) & operator= (const [Matrix](#) &other)
Copy assigns a [Matrix](#).
- [Matrix](#) & operator= ([Matrix](#) &&other)
Move assigns a [Matrix](#).
- void [col_scale](#) (ptr_param< const [global_vector_type](#) > scaling_factors)
Scales the columns of the matrix by the respective entries of the vector.
- void [row_scale](#) (ptr_param< const [global_vector_type](#) > scaling_factors)
Scales the rows of the matrix by the respective entries of the vector.

Static Public Member Functions

- static std::unique_ptr< [Matrix](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [mpi::communicator](#) comm)
Creates an empty distributed matrix.
- template<typename MatrixType , typename = std::enable_if_t<gko::detail::is_matrix_type_builder< MatrixType, ValueType, LocalIndexType>::value>>
static std::unique_ptr< [Matrix](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [mpi::communicator](#) comm, MatrixType matrix_template)
Creates an empty distributed matrix with specified type for local matrices.
- template<typename LocalMatrixType , typename NonLocalMatrixType , typename = std::enable_if_t< gko::detail::is_matrix_type_builder< LocalMatrixType, ValueType, LocalIndexType>::value && gko::detail::is_matrix_type_builder< NonLocalMatrixType, ValueType, LocalIndexType>::value>>
static std::unique_ptr< [Matrix](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [mpi::communicator](#) comm, LocalMatrixType local_matrix_template, NonLocalMatrixType non_local_matrix_template)
Creates an empty distributed matrix with specified types for the local matrix and the non-local matrix.

- static `std::unique_ptr< Matrix > create` (`std::shared_ptr< const Executor > exec`, `mpi::communicator comm`, `ptr_param< const LinOp > matrix_template`)
Creates an empty distributed matrix with specified type for local matrices.
- static `std::unique_ptr< Matrix > create` (`std::shared_ptr< const Executor > exec`, `mpi::communicator comm`, `ptr_param< const LinOp > local_matrix_template`, `ptr_param< const LinOp > non_local_matrix_template`)
Creates an empty distributed matrix with specified types for the local matrix and the non-local matrix.
- static `std::unique_ptr< Matrix > create` (`std::shared_ptr< const Executor > exec`, `mpi::communicator comm`, `dim< 2 > size`, `std::shared_ptr< LinOp > local_linop`)
Creates a local-only distributed matrix with existent `LinOp`.
- static `std::unique_ptr< Matrix > create` (`std::shared_ptr< const Executor > exec`, `mpi::communicator comm`, `dim< 2 > size`, `std::shared_ptr< LinOp > local_linop`, `std::shared_ptr< LinOp > non_local_linop`, `std::vector< comm_index_type > recv_sizes`, `std::vector< comm_index_type > recv_offsets`, `array< local_index_type > recv_gather_idxs`)
Creates distributed matrix with existent local and non-local `LinOp` and the corresponding mapping to collect the non-local data from the other ranks.

47.150.1 Detailed Description

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename GlobalIndexType = int64>
class gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >
```

The `Matrix` class defines a (MPI-)distributed matrix.

The matrix is stored in a row-wise distributed format. Each process owns a specific set of rows, where the assignment of rows is defined by a row `Partition`. The following depicts the distribution of global rows according to their assigned part-id (which will usually be the owning process id):

Part-Id	Global Rows	Part-Id	Local Rows
0	.. 1 2	0	.. 1 2
1	3 4		13 14 ..
2	.. 5 6 .. 7 ..	1	3 4
2	 8 .. 9		 10 11 12
1	 10 11 12	2	.. 5 6 .. 7 ..
0	13 14 ..		 8 .. 9

The local rows are further split into two matrices on each process. One matrix, called `local`, contains only entries from columns that are also owned by the process, while the other one, called `non_local`, contains entries from columns that are not owned by the process. The non-local matrix is stored in a compressed format, where empty columns are discarded and the remaining columns are renumbered. This splitting is depicted in the following:

Part-Id	Global	Local	Non-Local
0	.. 1 ! 2 .. !		.. 1 2
0	3 4 ! !		3 4 ..
1	.. 5 ! 6 .. ! 7 ..	6 ..	5 7 ..
1	 ! .. 8 ! .. 9	8 ..	 9
2	 ! .. 10 ! 11 12	11 12	.. 10
2	13 .. ! ! 14 ..	14 ..	13 ..

This uses the same ownership of the columns as for the rows. Additionally, the ownership of the columns may be explicitly defined with an second column partition. If that is not provided, the same row partition will be used for the columns. Using a column partition also allows to create non-square matrices, like the one below:

Part-Id	Global	Local	Non-Local
P_R/P_C	2 2 0 1		
0	.. 1 2 ..	2	1 ..
0	3 4	..	3 4
1	.. 5 6 ..	6 ..	5 7 ..
1	 8	8	
2	 10		10
2	13	13 ..	..

Here `P_R` denotes the row partition and `P_C` denotes the column partition.

The [Matrix](#) should be filled using the `read_distributed` method, e.g.

```
auto part = Partition<...>::build_from_mapping(...);
auto mat = Matrix<...>::create(exec, comm);
mat->read_distributed(matrix_data, part);
```

or if different partitions for the rows and columns are used:

```
auto row_part = Partition<...>::build_from_mapping(...);
auto col_part = Partition<...>::build_from_mapping(...);
auto mat = Matrix<...>::create(exec, comm);
mat->read_distributed(matrix_data, row_part, col_part);
```

This will set the dimensions of the global and local matrices automatically by deducing the sizes from the partitions.

By default the [Matrix](#) type uses Csr for both stored matrices. It is possible to explicitly change the datatype for the stored matrices, with the constraint that the new type should implement the `LinOp` and [ReadableFromMatrixData](#) interface. The type can be set by:

```
auto mat = Matrix<ValueType, LocalIndexType[, ...]>::create(
    exec, comm,
    Ell<ValueType, LocalIndexType>::create(exec).get(),
    Coo<ValueType, LocalIndexType>::create(exec).get());
```

Alternatively, the helper function `with_matrix_type` can be used:

```
auto mat = Matrix<ValueType, LocalIndexType>::create(
    exec, comm,
    with_matrix_type<Ell>(),
    with_matrix_type<Coo>());
```

See also

[with_matrix_type](#)

The [Matrix](#) `LinOp` supports the following operations:

```
experimental::distributed::Matrix *A;           // distributed matrix
experimental::distributed::Vector *b, *x;       // distributed multi-vectors
matrix::Dense *alpha, *beta; // scalars of dimension 1x1
// Applying to distributed multi-vectors computes an SpMV/SpMM product
A->apply(b, x) // x = A*b
A->apply(alpha, b, beta, x) // x = alpha*A*b + beta*x
```

Template Parameters

<i>ValueType</i>	The underlying value type.
<i>LocalIndexType</i>	The index type used by the local matrices.
<i>GlobalIndexType</i>	The type for global indices.

47.150.2 Constructor & Destructor Documentation

47.150.2.1 Matrix() [1/2]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::Matrix (
    const Matrix< ValueType, LocalIndexType, GlobalIndexType > & other )
```

Copy constructs a [Matrix](#).

Parameters

<i>other</i>	Matrix to copy from.
--------------	--------------------------------------

47.150.2.2 Matrix() [2/2]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::Matrix (
    Matrix< ValueType, LocalIndexType, GlobalIndexType > && other ) [noexcept]
```

Move constructs a [Matrix](#).

Parameters

<i>other</i>	Matrix to move from.
--------------	--------------------------------------

47.150.3 Member Function Documentation**47.150.3.1 col_scale()**

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::col_scale (
    ptr_param< const global_vector_type > scaling_factors )
```

Scales the columns of the matrix by the respective entries of the vector.

The vector's row partition has to be the same as the matrix's column partition. The scaling is done in-place.

Parameters

<i>scaling_factors</i>	The vector containing the scaling factors.
------------------------	--

47.150.3.2 create() [1/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, Local↵
IndexType, GlobalIndexType >::create (
```

```
std::shared_ptr< const Executor > exec,  
mpi::communicator comm ) [static]
```

Creates an empty distributed matrix.

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix. The default is the MPI_COMM_WORLD.

Returns

A smart pointer to the newly created matrix.

Referenced by gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::create().

47.150.3.3 create() [2/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename  
GlobalIndexType = int64>  
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::create (  
    std::shared_ptr< const Executor > exec,  
    mpi::communicator comm,  
    dim< 2 > size,  
    std::shared_ptr< LinOp > local_linop ) [static]
```

Creates a local-only distributed matrix with existent LinOp.

Note

It use the input to build up the distributed matrix

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>size</i>	the global size
<i>local_linop</i>	the local linop

Returns

A smart pointer to the newly created matrix.

47.150.3.4 create() [3/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, Local↵
IndexType, GlobalIndexType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    dim< 2 > size,
    std::shared_ptr< LinOp > local_linop,
    std::shared_ptr< LinOp > non_local_linop,
    std::vector< comm_index_type > recv_sizes,
    std::vector< comm_index_type > recv_offsets,
    array< local_index_type > recv_gather_idxs ) [static]
```

Creates distributed matrix with existent local and non-local LinOp and the corresponding mapping to collect the non-local data from the other ranks.

Note

It use the input to build up the distributed matrix

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>size</i>	the global size
<i>local_linop</i>	the local linop
<i>non_local_linop</i>	the non-local linop
<i>recv_sizes</i>	the size of non-local receiver
<i>recv_offset</i>	the offset of non-local receiver
<i>recv_gather_idxs</i>	the gathering index of non-local receiver

Returns

A smart pointer to the newly created matrix.

47.150.3.5 create() [4/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
template<typename LocalMatrixType , typename NonLocalMatrixType , typename = std::enable_if↵
t< gko::detail::is_matrix_type_builder< LocalMatrixType, ValueType, LocalIndexType>::value &&
gko::detail::is_matrix_type_builder< NonLocalMatrixType, ValueType, LocalIndexType>::value>>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, Local↵
IndexType, GlobalIndexType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    LocalMatrixType local_matrix_template,
    NonLocalMatrixType non_local_matrix_template ) [inline], [static]
```

Creates an empty distributed matrix with specified types for the local matrix and the non-local matrix.

Note

This is mainly a convenience wrapper for `Matrix(std::shared_ptr<const Executor>, mpi::communicator, const LinOp*, const LinOp*)`

Template Parameters

<i>LocalMatrixType</i>	A type that has a <code>create<ValueType, IndexType>(exec)</code> function to create a smart pointer of a type derived from <code>LinOp</code> and ReadableFromMatrixData .
------------------------	---

See also

[with_matrix_type](#)

Template Parameters

<i>NonLocalMatrixType</i>	A (possible different) type with the same constraints as <code>LocalMatrixType</code> .
---------------------------	---

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>local_matrix_template</i>	the local matrix will be constructed with the same type as <code>create</code> returns. It should be the return value of <code>make_matrix_template</code> .
<i>non_local_matrix_template</i>	the non-local matrix will be constructed with the same type as <code>create</code> returns. It should be the return value of <code>make_matrix_template</code> .

Returns

A smart pointer to the newly created matrix.

References `gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::create()`.

47.150.3.6 create() [5/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
template<typename MatrixType, typename = std::enable_if_t<gko::detail::is_matrix_type_↵
builder< MatrixType, ValueType, LocalIndexType>::value>>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, Local↵
IndexType, GlobalIndexType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    MatrixType matrix_template ) [inline], [static]
```

Creates an empty distributed matrix with specified type for local matrices.

Note

This is mainly a convenience wrapper for `Matrix(std::shared_ptr<const Executor>, mpi::communicator, const LinOp*)`

Template Parameters

<i>MatrixType</i>	A type that has a <code>create<ValueType, IndexType>(exec)</code> function to create a smart pointer of a type derived from <code>LinOp</code> and ReadableFromMatrixData .
-------------------	---

See also

[with_matrix_type](#)

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>matrix_template</i>	the local matrices will be constructed with the same type as <code>create</code> returns. It should be the return value of <code>make_matrix_template</code> .

Returns

A smart pointer to the newly created matrix.

References `gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::create()`.

47.150.3.7 create() [6/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, Local←
IndexType, GlobalIndexType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    ptr_param< const LinOp > local_matrix_template,
    ptr_param< const LinOp > non_local_matrix_template ) [static]
```

Creates an empty distributed matrix with specified types for the local matrix and the non-local matrix.

Note

It internally clones the passed in `local_matrix_template` and `non_local_matrix_template`. Therefore, those `LinOps` should be empty.

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>local_matrix_template</i>	the local matrix will be constructed with the same runtime type.
<i>non_local_matrix_template</i>	the non-local matrix will be constructed with the same runtime type.

Returns

A smart pointer to the newly created matrix.

47.150.3.8 create() [7/7]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
static std::unique_ptr<Matrix> gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    ptr_param< const LinOp > matrix_template ) [static]
```

Creates an empty distributed matrix with specified type for local matrices.

Note

It internally clones the passed in matrix_template. Therefore, the LinOp should be empty.

Parameters

<i>exec</i>	Executor associated with this matrix.
<i>comm</i>	Communicator associated with this matrix.
<i>matrix_template</i>	the local matrices will be constructed with the same runtime type.

Returns

A smart pointer to the newly created matrix.

47.150.3.9 get_local_matrix()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
std::shared_ptr<const LinOp> gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::get_local_matrix ( ) const [inline]
```

Get read access to the stored local matrix.

Returns

Shared pointer to the stored local matrix

47.150.3.10 get_non_local_matrix()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
std::shared_ptr<const LinOp> gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::get_non_local_matrix ( ) const [inline]
```

Get read access to the stored non-local matrix.

Returns

Shared pointer to the stored non-local matrix

47.150.3.11 operator=() [1/2]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
Matrix& gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::operator= (
    const Matrix< ValueType, LocalIndexType, GlobalIndexType > & other )
```

Copy assigns a [Matrix](#).

Parameters

<i>other</i>	Matrix to copy from, has to have a communicator of the same size as this.
--------------	---

Returns

this.

47.150.3.12 operator=() [2/2]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
Matrix& gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >::operator= (
    Matrix< ValueType, LocalIndexType, GlobalIndexType > && other )
```

Move assigns a [Matrix](#).

Parameters

<i>other</i>	Matrix to move from, has to have a communicator of the same size as this.
--------------	---

Returns

this.

47.150.3.13 read_distributed() [1/4]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::read_distributed (
    const device_matrix_data< value_type, global_index_type > & data,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> partition,
    assembly_mode assembly_type = assembly_mode::local_only )
```

Reads a square matrix from the [device_matrix_data](#) structure and a global partition.

The global size of the final matrix is inferred from the size of the partition. Both the number of rows and columns of the [device_matrix_data](#) are ignored.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

<i>data</i>	The device_matrix_data structure.
<i>partition</i>	The global row and column partition.
<i>assembly_type</i>	The mode of assembly.

Returns

the [index_map](#) induced by the partitions and the matrix structure

47.150.3.14 read_distributed() [2/4]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::read_distributed (
    const device_matrix_data< value_type, global_index_type > & data,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> row↵
partition,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> col↵
partition,
    assembly_mode assembly_type = assembly_mode::local_only )
```

Reads a matrix from the [device_matrix_data](#) structure, a global row partition, and a global column partition.

The global size of the final matrix is inferred from the size of the row partition and the size of the column partition. Both the number of rows and columns of the [device_matrix_data](#) are ignored.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

<i>data</i>	The device_matrix_data structure.
<i>row_partition</i>	The global row partition.
<i>col_partition</i>	The global col partition.
<i>assembly_type</i>	The mode of assembly.

Returns

the [index_map](#) induced by the partitions and the matrix structure

47.150.3.15 read_distributed() [3/4]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::read_distributed (
    const matrix\_data< value_type, global_index_type > & data,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> partition,
    assembly\_mode assembly_type = assembly\_mode::local\_only )
```

Reads a square matrix from the [matrix_data](#) structure and a global partition.

See also

[read_distributed](#)

Note

For efficiency it is advised to use the [device_matrix_data](#) overload.

47.150.3.16 read_distributed() [4/4]

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::read_distributed (
    const matrix\_data< value_type, global_index_type > & data,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> row_↵
partition,
    std::shared_ptr< const Partition< local_index_type, global_index_type >> col_↵
partition,
    assembly\_mode assembly_type = assembly\_mode::local\_only )
```

Reads a matrix from the [matrix_data](#) structure, a global row partition, and a global column partition.

See also

[read_distributed](#)

Note

For efficiency it is advised to use the [device_matrix_data](#) overload.

47.150.3.17 row_scale()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
void gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >↵
::row_scale (
    ptr_param< const global_vector_type > scaling_factors )
```

Scales the rows of the matrix by the respective entries of the vector.

The vector and the matrix have to have the same row partition. The scaling is done in-place.

Parameters

<i>scaling_factors</i>	The vector containing the scaling factors.
------------------------	--

The documentation for this class was generated from the following file:

- ginkgo/core/distributed/matrix.hpp

47.151 gko::matrix_assembly_data< ValueType, IndexType > Class Template Reference

This structure is used as an intermediate type to assemble a sparse matrix.

```
#include <ginkgo/core/base/matrix_assembly_data.hpp>
```

Public Member Functions

- void [add_value](#) (index_type row, index_type col, value_type val)
Sets the matrix value at (row, col).
- void [set_value](#) (index_type row, index_type col, value_type val)
Sets the matrix value at (row, col).
- value_type [get_value](#) (index_type row, index_type col)
Gets the matrix value at (row, col).
- bool [contains](#) (index_type row, index_type col)
Returns true iff the matrix contains an entry at (row, col).
- dim< 2 > [get_size](#) () const noexcept
- size_type [get_num_stored_elements](#) () const noexcept
- matrix_data< ValueType, IndexType > [get_ordered_data](#) () const

47.151.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix_assembly_data< ValueType, IndexType >
```

This structure is used as an intermediate type to assemble a sparse matrix.

The matrix is stored as a set of nonzero elements, where each element is a triplet of the form (row_index, column_index, value).

New values can be added by using the [matrix_assembly_data::add_value](#) or [matrix_assembly_data::set_value](#)

Template Parameters

<i>ValueType</i>	type of matrix values stored in the structure
<i>IndexType</i>	type of matrix indexes stored in the structure

47.151.2 Member Function Documentation

47.151.2.1 add_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix_assembly_data< ValueType, IndexType >::add_value (
    index_type row,
    index_type col,
    value_type val ) [inline]
```

Sets the matrix value at (row, col).

If there is an existing value, it will be set to the sum of the existing and new value, otherwise the value will be inserted.

Parameters

<i>row</i>	the row where the value should be added
<i>col</i>	the column where the value should be added
<i>val</i>	the value to be added to (row, col)

```
79      {
80          auto ind = std::make_pair(row, col);
81          nonzeros_[ind] += val;
82      }
```

47.151.2.2 contains()

```
template<typename ValueType = default_precision, typename IndexType = int32>
bool gko::matrix_assembly_data< ValueType, IndexType >::contains (
```

```
index_type row,
index_type col ) [inline]
```

Returns true iff the matrix contains an entry at (row, col).

Parameters

<i>row</i>	the row index
<i>col</i>	the column index

Returns

true if the value at (row, col) exists, false otherwise

47.151.2.3 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix_assembly_data< ValueType, IndexType >::get_num_stored_elements ( ) const
[inline], [noexcept]
```

Returns

the number of non-zeros in the (partially) assembled matrix

47.151.2.4 get_ordered_data()

```
template<typename ValueType = default_precision, typename IndexType = int32>
matrix_data<ValueType, IndexType> gko::matrix_assembly_data< ValueType, IndexType >::get_↵
ordered_data ( ) const [inline]
```

Returns

a [matrix_data](#) instance containing the assembled non-zeros in row-major order to be used by all matrix formats.

Referenced by `gko::ReadableFromMatrixData< ValueType, int32 >::read()`.

47.151.2.5 get_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
dim<2> gko::matrix_assembly_data< ValueType, IndexType >::get_size ( ) const [inline], [noexcept]
```

Returns

the dimensions of the matrix being assembled

47.151.2.6 get_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type gko::matrix_assembly_data< ValueType, IndexType >::get_value (
    index_type row,
    index_type col ) [inline]
```

Gets the matrix value at (row, col).

Parameters

<i>row</i>	the row index
<i>col</i>	the column index

Returns

the value at (row, col) or 0 if it doesn't exist.

47.151.2.7 set_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix_assembly_data< ValueType, IndexType >::set_value (
    index_type row,
    index_type col,
    value_type val ) [inline]
```

Sets the matrix value at (row, col).

If there is an existing value, it will be overwritten by the new value.

Parameters

<i>row</i>	the row index
<i>col</i>	the column index
<i>val</i>	the value to be written to (row, col)

The documentation for this class was generated from the following file:

- ginkgo/core/base/matrix_assembly_data.hpp

47.152 gko::matrix_data< ValueType, IndexType > Struct Template Reference

This structure is used as an intermediate data type to store a sparse matrix.

```
#include <ginkgo/core/base/matrix_data.hpp>
```

Public Member Functions

- `matrix_data (dim< 2 > size___dim< 2 > {}, ValueType value=zero< ValueType >())`
Initializes a matrix filled with the specified value.
- `template<typename RandomDistribution, typename RandomEngine >
matrix_data (dim< 2 > size___, RandomDistribution &&dist, RandomEngine &&engine)`
Initializes a matrix with random values from the specified distribution.
- `matrix_data (std::initializer_list< std::initializer_list< ValueType >> values)`
List-initializes the structure from a matrix of values.
- `matrix_data (dim< 2 > size___, std::initializer_list< detail::input_triple< ValueType, IndexType >> nonzeros↵
_)`
Initializes the structure from a list of nonzeros.
- `matrix_data (dim< 2 > size___, const matrix_data &block)`
Initializes a matrix out of a matrix block via duplication.
- `template<typename Accessor >
matrix_data (const range< Accessor > &data)`
Initializes a matrix from a range.
- `void sort_row_major ()`
Sorts the nonzero vector so the values follow row-major order.
- `void ensure_row_major_order ()`
Sorts the nonzero vector so the values follow row-major order.
- `void remove_zeros ()`
Remove entries with value zero from the matrix data.
- `void sum_duplicates ()`
Sum up all values that refer to the same matrix entry.

Static Public Member Functions

- `static matrix_data diag (dim< 2 > size___, ValueType value)`
Initializes a diagonal matrix.
- `static matrix_data diag (dim< 2 > size___, std::initializer_list< ValueType > nonzeros_)`
Initializes a diagonal matrix using a list of diagonal elements.
- `static matrix_data diag (dim< 2 > size___, const matrix_data &block)`
Initializes a block-diagonal matrix.
- `template<typename ForwardIterator >
static matrix_data diag (ForwardIterator begin, ForwardIterator end)`
Initializes a block-diagonal matrix from a list of diagonal blocks.
- `static matrix_data diag (std::initializer_list< matrix_data > blocks)`
Initializes a block-diagonal matrix from a list of diagonal blocks.
- `template<typename RandomDistribution, typename RandomEngine >
static matrix_data cond (size_type size, remove_complex< ValueType > condition_number, Random↵
Distribution &&dist, RandomEngine &&engine, size_type num_reflectors)`
Initializes a random dense matrix with a specific condition number.
- `template<typename RandomDistribution, typename RandomEngine >
static matrix_data cond (size_type size, remove_complex< ValueType > condition_number, Random↵
Distribution &&dist, RandomEngine &&engine)`
Initializes a random dense matrix with a specific condition number.

Public Attributes

- `dim< 2 > size`
Size of the matrix.
- `std::vector< nonzero_type > nonzeros`
A vector of tuples storing the non-zeros of the matrix.

47.152.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
struct gko::matrix_data< ValueType, IndexType >
```

This structure is used as an intermediate data type to store a sparse matrix.

The matrix is stored as a sequence of nonzero elements, where each element is a triple of the form (row_index, column_index, value).

Note

All Ginkgo functions returning such a structure will return the nonzeros sorted in row-major order.

All Ginkgo functions that take this structure as input expect that the nonzeros are sorted in row-major order and that the index pair (row_index, column_index) of each nonzero is unique.

This structure is not optimized for usual access patterns and it can only exist on the CPU. Thus, it should only be used for utility functions which do not have to be optimized for performance.

Template Parameters

<i>ValueType</i>	type of matrix values stored in the structure
<i>IndexType</i>	type of matrix indexes stored in the structure

47.152.2 Constructor & Destructor Documentation

47.152.2.1 matrix_data() [1/6]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix_data< ValueType, IndexType >::matrix_data (
    dim< 2 > size_ = dim<2>{},
    ValueType value = zero<ValueType>() ) [inline]
```

Initializes a matrix filled with the specified value.

Parameters

<i>size_</i>	dimensions of the matrix
<i>value</i>	value used to fill the elements of the matrix


```

137                                     {}, ValueType value = zero<ValueType>())
138     : size{size_}
139     {
140         if (is_zero(value)) {
141             return;
142         }
143         nonzeros.reserve(size[0] * size[1]);
144         for (size_type row = 0; row < size[0]; ++row) {
145             for (size_type col = 0; col < size[1]; ++col) {
146                 nonzeros.emplace_back(row, col, value);
147             }
148         }
149     }

```

47.152.2.2 matrix_data() [2/6]

```

template<typename ValueType = default_precision, typename IndexType = int32>
template<typename RandomDistribution, typename RandomEngine>
gko::matrix_data< ValueType, IndexType >::matrix_data (
    dim< 2 > size_,
    RandomDistribution && dist,
    RandomEngine && engine ) [inline]

```

Initializes a matrix with random values from the specified distribution.

Template Parameters

<i>RandomDistribution</i>	random distribution type
<i>RandomEngine</i>	random engine type

Parameters

<i>size</i> _↔	dimensions of the matrix
<i>—</i>	
<i>dist</i>	random distribution of the elements of the matrix
<i>engine</i>	random engine used to generate random values

47.152.2.3 matrix_data() [3/6]

```

template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix_data< ValueType, IndexType >::matrix_data (
    std::initializer_list< std::initializer_list< ValueType >> values ) [inline]

```

List-initializes the structure from a matrix of values.

Parameters

<i>values</i>	a 2D braced-init-list of matrix values.
---------------	---

47.152.2.4 matrix_data() [4/6]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix_data< ValueType, IndexType >::matrix_data (
    dim< 2 > size_,
    std::initializer_list< detail::input_triple< ValueType, IndexType >> nonzeros_ )
[inline]
```

Initializes the structure from a list of nonzeros.

Parameters

<i>size_</i>	dimensions of the matrix
<i>nonzeros_</i>	list of nonzero elements

47.152.2.5 matrix_data() [5/6]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix_data< ValueType, IndexType >::matrix_data (
    dim< 2 > size_,
    const matrix_data< ValueType, IndexType > & block ) [inline]
```

Initializes a matrix out of a matrix block via duplication.

Parameters

<i>size</i>	size of the block-matrix (in blocks)
<i>diag_block</i>	matrix block used to fill the complete matrix

References `gko::matrix_data< ValueType, IndexType >::size`.

47.152.2.6 matrix_data() [6/6]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename Accessor >
gko::matrix_data< ValueType, IndexType >::matrix_data (
    const range< Accessor > & data ) [inline]
```

Initializes a matrix from a range.

Template Parameters

<i>Accessor</i>	accessor type of the input range
-----------------	----------------------------------

Parameters

<i>data</i>	range used to initialize the matrix
-------------	-------------------------------------

References gko::range< Accessor >::length().

47.152.3 Member Function Documentation

47.152.3.1 cond() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename RandomDistribution, typename RandomEngine>
static matrix_data gko::matrix_data< ValueType, IndexType >::cond (
    size_type size,
    remove_complex< ValueType > condition_number,
    RandomDistribution && dist,
    RandomEngine && engine ) [inline], [static]
```

Initializes a random dense matrix with a specific condition number.

The matrix is generated by applying a series of random Householder reflectors to a diagonal matrix with diagonal entries uniformly distributed between `sqrt(condition_number)` and `1/sqrt(condition_number)`.

This version of the function applies `size - 1` reflectors to each side of the diagonal matrix.

Template Parameters

<i>RandomDistribution</i>	the type of the random distribution
<i>RandomEngine</i>	the type of the random engine

Parameters

<i>size</i>	number of rows and columns of the matrix
<i>condition_number</i>	condition number of the matrix
<i>dist</i>	random distribution used to generate reflectors
<i>engine</i>	random engine used to generate reflectors

Returns

the dense matrix with the specified condition number

References gko::matrix_data< ValueType, IndexType >::cond(), and gko::matrix_data< ValueType, IndexType >↵::size.

47.152.3.2 cond() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename RandomDistribution, typename RandomEngine>
static matrix_data gko::matrix_data< ValueType, IndexType >::cond (
    size_type size,
    remove_complex< ValueType > condition_number,
    RandomDistribution && dist,
    RandomEngine && engine,
    size_type num_reflectors ) [inline], [static]
```

Initializes a random dense matrix with a specific condition number.

The matrix is generated by applying a series of random Hausholder reflectors to a diagonal matrix with diagonal entries uniformly distributed between `sqrt(condition_number)` and `1/sqrt(condition_number)`.

Template Parameters

<i>RandomDistribution</i>	the type of the random distribution
<i>RandomEngine</i>	the type of the random engine

Parameters

<i>size</i>	number of rows and columns of the matrix
<i>condition_number</i>	condition number of the matrix
<i>dist</i>	random distribution used to generate reflectors
<i>engine</i>	random engine used to generate reflectors
<i>num_reflectors</i>	number of reflectors to apply from each side

Returns

the dense matrix with the specified condition number

References `gko::matrix_data< ValueType, IndexType >::size`.

Referenced by `gko::matrix_data< ValueType, IndexType >::cond()`.

47.152.3.3 diag() [1/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static matrix_data gko::matrix_data< ValueType, IndexType >::diag (
    dim< 2 > size_,
    const matrix_data< ValueType, IndexType > & block ) [inline], [static]
```

Initializes a block-diagonal matrix.

Parameters

<i>size_</i>	the size of the matrix
<i>diag_block</i>	matrix used to fill diagonal blocks

Returns

the block-diagonal matrix

References gko::matrix_data< ValueType, IndexType >::nonzeros, and gko::matrix_data< ValueType, IndexType >::size.

47.152.3.4 diag() [2/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static matrix_data gko::matrix_data< ValueType, IndexType >::diag (
    dim< 2 > size_,
    std::initializer_list< ValueType > nonzeros_ ) [inline], [static]
```

Initializes a diagonal matrix using a list of diagonal elements.

Parameters

<i>size_</i>	dimensions of the matrix
<i>nonzeros_↔</i>	list of diagonal elements
—	

Returns

the diagonal matrix

References gko::matrix_data< ValueType, IndexType >::nonzeros.

47.152.3.5 diag() [3/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static matrix_data gko::matrix_data< ValueType, IndexType >::diag (
    dim< 2 > size_,
    ValueType value ) [inline], [static]
```

Initializes a diagonal matrix.

Parameters

<i>size_↔</i>	dimensions of the matrix
—	
<i>value</i>	value used to fill the elements of the matrix

Returns

the diagonal matrix

References `gko::is_nonzero()`, and `gko::matrix_data< ValueType, IndexType >::nonzeros`.

Referenced by `gko::matrix_data< ValueType, IndexType >::diag()`.

47.152.3.6 `diag()` [4/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename ForwardIterator >
static matrix_data gko::matrix_data< ValueType, IndexType >::diag (
    ForwardIterator begin,
    ForwardIterator end ) [inline], [static]
```

Initializes a block-diagonal matrix from a list of diagonal blocks.

Template Parameters

<i>ForwardIterator</i>	type of list iterator
------------------------	-----------------------

Parameters

<i>begin</i>	the first iterator of the list
<i>end</i>	the last iterator of the list

Returns

the block-diagonal matrix with diagonal blocks set to the blocks between *begin* (inclusive) and *end* (exclusive)

References `gko::matrix_data< ValueType, IndexType >::nonzeros`.

47.152.3.7 `diag()` [5/5]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static matrix_data gko::matrix_data< ValueType, IndexType >::diag (
    std::initializer_list< matrix_data< ValueType, IndexType > > blocks ) [inline],
[static]
```

Initializes a block-diagonal matrix from a list of diagonal blocks.

Parameters

<i>blocks</i>	a list of blocks to initialize from
---------------	-------------------------------------

Returns

the block-diagonal matrix with diagonal blocks set to the blocks passed in *blocks*

References gko::matrix_data< ValueType, IndexType >::diag().

47.152.3.8 sum_duplicates()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix_data< ValueType, IndexType >::sum_duplicates ( ) [inline]
```

Sum up all values that refer to the same matrix entry.

The result is sorted in row-major order.

References gko::matrix_data< ValueType, IndexType >::nonzeros, and gko::matrix_data< ValueType, IndexType >::sort_row_major().

47.152.4 Member Data Documentation

47.152.4.1 nonzeros

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::vector<nonzero_type> gko::matrix_data< ValueType, IndexType >::nonzeros
```

A vector of tuples storing the non-zeros of the matrix.

The first two elements of the tuple are the row index and the column index of a matrix element, and its third element is the value at that position.

Referenced by gko::matrix_data< ValueType, IndexType >::diag(), gko::matrix_data< ValueType, IndexType >::remove_zeros(), gko::matrix_data< ValueType, IndexType >::sort_row_major(), and gko::matrix_data< ValueType, IndexType >::sum_duplicates().

The documentation for this struct was generated from the following file:

- ginkgo/core/base/matrix_data.hpp

47.153 gko::matrix_data_entry< ValueType, IndexType > Struct Template Reference

Type used to store nonzeros.

```
#include <ginkgo/core/base/matrix_data.hpp>
```

47.153.1 Detailed Description

```
template<typename ValueType, typename IndexType>
struct gko::matrix_data_entry< ValueType, IndexType >
```

Type used to store nonzeros.

The documentation for this struct was generated from the following file:

- ginkgo/core/base/matrix_data.hpp

47.154 gko::experimental::reorder::Mc64< ValueType, IndexType > Class Template Reference

MC64 is an algorithm for permuting large entries to the diagonal of a sparse matrix.

```
#include <ginkgo/core/reorder/mc64.hpp>
```

Public Member Functions

- const parameters_type & [get_parameters](#) () const
Returns the parameters used to construct the factory.
- std::unique_ptr< [result_type](#) > [generate](#) (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

- static parameters_type [build](#) ()
Creates a new parameter_type to set up the factory.

47.154.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::experimental::reorder::Mc64< ValueType, IndexType >
```

MC64 is an algorithm for permuting large entries to the diagonal of a sparse matrix.

This approach can increase numerical stability of e.g. an LU factorization without pivoting. Under the assumption of working on a nonsingular square matrix, the algorithm computes a minimum weight perfect matching on a weighted edge bipartite graph of the matrix. It is described in detail in "On Algorithms for Permuting Large Entries to the Diagonal of a Sparse Matrix" (Duff, Koster, 2001, DOI: 10.1137/S0895479899358443). There are two strategies for choosing the weights supported:

- Maximizing the product of the absolute values on the diagonal. For this strategy, the weights are computed as $c(i, j) = \log_2(a_i) - \log_2(|a(i, j)|)$ if $a(i, j) \neq 0$ and $c(i, j) = \infty$ otherwise. Here, a_i is the maximum absolute value in row i of the matrix A . In this case, the implementation computes a row permutation P and row and column scaling coefficients L and R such that the matrix $P*L*A*R$ has values with unity absolute value on the diagonal and smaller or equal entries everywhere else.
- Maximizing the sum of the absolute values on the diagonal. For this strategy, the weights are computed as $c(i, j) = a_i - |a(i, j)|$ if $a(i, j) \neq 0$ and $c(i, j) = \infty$ otherwise. In this case, no scaling coefficients are computed.

This class creates a [Combination](#) of two ScaledPermutations representing the row and column permutation and scaling factors computed by this algorithm.

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.154.2 Member Function Documentation

47.154.2.1 generate()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<result_type> gko::experimental::reorder::Mc64< ValueType, IndexType >::generate
(
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a Permutation instead of a generic LinOp, which would need to be cast to ScaledPermutation again to access its indices. It is only necessary because smart pointers aren't covariant.

47.154.2.2 get_parameters()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const parameters_type& gko::experimental::reorder::Mc64< ValueType, IndexType >::get_parameters
( ) const [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

The documentation for this class was generated from the following file:

- ginkgo/core/reorder/mc64.hpp

47.155 gko::matrix::Csr< ValueType, IndexType >::merge_path Class Reference

[merge_path](#) is a [strategy_type](#) which uses the [merge_path](#) algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [merge_path](#) ()
Creates a [merge_path](#) strategy.
- void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow) override
Computes srow according to row pointers.
- int64_t [clac_size](#) (const int64_t nnz) override
Computes the srow size according to the number of nonzeros.
- std::shared_ptr< [strategy_type](#) > [copy](#) () override
Copy a strategy.

47.155.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::merge_path
```

[merge_path](#) is a [strategy_type](#) which uses the [merge_path](#) algorithm.

[merge_path](#) is according to Merrill and Garland: Merge-Based Parallel Sparse Matrix-Vector Multiplication

47.155.2 Member Function Documentation

47.155.2.1 clac_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int64_t gko::matrix::Csr< ValueType, IndexType >::merge_path::clac_size (
    const int64_t nnz ) [inline], [override], [virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.155.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy\_type> gko::matrix::Csr< ValueType, IndexType >::merge_path::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.155.2.3 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::merge_path::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [inline], [override], [virtual]
```

Computes srow according to row pointers.

Parameters

<code>mtx_row_ptrs</code>	the row pointers of the matrix
<code>mtx_srow</code>	the srow of the matrix

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/csr.hpp`

47.156 gko::MetisError Class Reference

[MetisError](#) is thrown when METIS routine throws an error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [MetisError](#) (const std::string &file, int line, const std::string &func, const std::string &error)
Initializes a METIS error.

47.156.1 Detailed Description

[MetisError](#) is thrown when METIS routine throws an error code.

47.156.2 Constructor & Destructor Documentation

47.156.2.1 MetisError()

```
gko::MetisError::MetisError (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & error ) [inline]
```

Initializes a METIS error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the METIS routine that failed
<i>error</i>	The resulting METIS error name

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.157 gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit Class Reference

[minimal_storage_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- [minimal_storage_limit](#) ()
Creates a [minimal_storage_limit](#) strategy.
- [size_type compute_ell_num_stored_elements_per_row](#) (array< [size_type](#) > *row_nnz) const override
Computes the number of stored elements per row of the ell part.
- auto [get_percentage](#) () const
Get the percent setting.

47.157.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit
```

[minimal_storage_limit](#) is a [strategy_type](#) which decides the number of stored elements per row of the ell part.

It is determined by the size of ValueType and IndexType, the storage is the minimum among all partition.

47.157.2 Member Function Documentation

47.157.2.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::compute_ell_num_stored_elements_per_row (
    array< size_type > * row_nnz ) const [inline], [override], [virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<code>row_nnz</code>	the number of nonzeros of each row
----------------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implements [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type](#).

References [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::compute_ell_num_stored_elements_per_row\(\)](#).

47.157.2.2 get_percentage()

```
template<typename ValueType = default_precision, typename IndexType = int32>
auto gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit::get_percentage ( )
const [inline]
```

Get the percent setting.

Returns

percent

References [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit::get_percentage\(\)](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/hybrid.hpp](#)

47.158 gko::MpiError Class Reference

[MpiError](#) is thrown when a MPI routine throws a non-zero error code.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [MpiError](#) (const std::string &file, int line, const std::string &func, [int64](#) error_code)
Initializes a MPI error.

47.158.1 Detailed Description

[MpiError](#) is thrown when a MPI routine throws a non-zero error code.

47.158.2 Constructor & Destructor Documentation

47.158.2.1 MpiError()

```
gko::MpiError::MpiError (
    const std::string & file,
    int line,
    const std::string & func,
    int64 error_code ) [inline]
```

Initializes a MPI error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the MPI routine that failed
<i>error_code</i>	The resulting MPI error code

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.159 gko::solver::Multigrid Class Reference

[Multigrid](#) methods have a hierarchy of many levels, whose coarse level is a subset of the fine level, of the problem.

```
#include <ginkgo/core/solver/multigrid.hpp>
```

Public Member Functions

- bool [apply_uses_initial_guess](#) () const override
Return true as iterative solvers use the data in x as an initial guess or false if multigrid always set the input as zero.
- std::vector< std::shared_ptr< const [gko::multigrid::MultigridLevel](#) > > [get_mg_level_list](#) () const

Gets the list of MultigridLevel operators.

- `std::vector< std::shared_ptr< const LinOp > > get_pre_smoother_list () const`

Gets the list of pre-smoother operators.

- `std::vector< std::shared_ptr< const LinOp > > get_mid_smoother_list () const`

Gets the list of mid-smoother operators.

- `std::vector< std::shared_ptr< const LinOp > > get_post_smoother_list () const`

Gets the list of post-smoother operators.

- `std::shared_ptr< const LinOp > get_coarsest_solver () const`

Gets the operator at the coarsest level.

- `multigrid::cycle get_cycle () const`

Get the cycle of multigrid.

- `void set_cycle (multigrid::cycle cycle)`

Set the cycle of multigrid.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor<>())`

Create the parameters from the property_tree.

47.159.1 Detailed Description

[Multigrid](#) methods have a hierarchy of many levels, whose coarse level is a subset of the fine level, of the problem.

The coarse level solves the system on the residual of fine level and fine level will use the coarse solution to correct its own result. [Multigrid](#) solves the problem by relatively cheap step in each level and refining the result when prolongating back.

The main step of each level

- Presmooth (solve on the fine level)
- Calculate residual
- Restrict (reduce the problem dimension)
- Solve residual in next level
- Prolongate (return to the fine level size)
- Postsmooth (correct the answer in fine level)

Ginkgo uses the index from 0 for finest level (original problem size) $\sim N$ for the coarsest level (the coarsest solver), and its level counts is N (N multigrid level generation).

47.159.2 Member Function Documentation

47.159.2.1 apply_uses_initial_guess()

```
bool gko::solver::Multigrid::apply_uses_initial_guess ( ) const [inline], [override]
```

Return true as iterative solvers use the data in x as an initial guess or false if multigrid always set the input as zero.

Returns

bool it is related to parameters variable `zero_guess`

```
123     {
124         return this->get_default_initial_guess() ==
125             initial_guess_mode::provided;
126     }
```

References `gko::solver::provided`.

47.159.2.2 get_coarsest_solver()

```
std::shared_ptr<const LinOp> gko::solver::Multigrid::get_coarsest_solver ( ) const [inline]
```

Gets the operator at the coarsest level.

Returns

the coarsest operator

47.159.2.3 get_cycle()

```
multigrid::cycle gko::solver::Multigrid::get_cycle ( ) const [inline]
```

Get the cycle of multigrid.

Returns

the `multigrid::cycle`

47.159.2.4 get_mg_level_list()

```
std::vector<std::shared_ptr<const gko::multigrid::MultigridLevel> > gko::solver::Multigrid↵
::get_mg_level_list ( ) const [inline]
```

Gets the list of MultigridLevel operators.

Returns

the list of MultigridLevel operators

47.159.2.5 get_mid_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_mid_smoother_list ( )
const [inline]
```

Gets the list of mid-smoother operators.

Returns

the list of mid-smoother operators

47.159.2.6 get_post_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_post_smoother_list ( )
const [inline]
```

Gets the list of post-smoother operators.

Returns

the list of post-smoother operators

47.159.2.7 get_pre_smoother_list()

```
std::vector<std::shared_ptr<const LinOp> > gko::solver::Multigrid::get_pre_smoother_list ( )
const [inline]
```

Gets the list of pre-smoother operators.

Returns

the list of pre-smoother operators

47.159.2.8 parse()

```
static parameters_type gko::solver::Multigrid::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor<>() )
[static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs

Returns

parameters

Note

[Multigrid](#) does not support the selector option in file config

47.159.2.9 set_cycle()

```
void gko::solver::Multigrid::set_cycle (
    multigrid::cycle cycle ) [inline]
```

Set the cycle of multigrid.

Parameters

multigrid::cycle	the new cycle
----------------------------------	---------------

The documentation for this class was generated from the following file:

- `ginkgo/core/solver/multigrid.hpp`

47.160 gko::multigrid::MultigridLevel Class Reference

This class represents two levels in a multigrid hierarchy.

```
#include <ginkgo/core/multigrid/multigrid_level.hpp>
```

Public Member Functions

- virtual `std::shared_ptr< const LinOp >` [get_fine_op](#) () const =0
Returns the operator on fine level.
- virtual `std::shared_ptr< const LinOp >` [get_restrict_op](#) () const =0
Returns the restrict operator.
- virtual `std::shared_ptr< const LinOp >` [get_coarse_op](#) () const =0
Returns the operator on coarse level.
- virtual `std::shared_ptr< const LinOp >` [get_prolong_op](#) () const =0
Returns the prolong operator.

47.160.1 Detailed Description

This class represents two levels in a multigrid hierarchy.

The [MultigridLevel](#) is an interface that allows to get the individual components of multigrid level. Each implementation of a multigrid level should inherit from this interface. Use [EnableMultigridLevel<ValueType>](#) to implement this interface with composition by default.

47.160.2 Member Function Documentation

47.160.2.1 get_coarse_op()

```
virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_coarse_op ( ) const  
[pure virtual]
```

Returns the operator on coarse level.

Returns

the operator on coarse level.

Implemented in [gko::multigrid::EnableMultigridLevel< ValueType >](#).

47.160.2.2 get_fine_op()

```
virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_fine_op ( ) const  
[pure virtual]
```

Returns the operator on fine level.

Returns

the operator on fine level.

Implemented in [gko::multigrid::EnableMultigridLevel< ValueType >](#).

47.160.2.3 get_prolong_op()

```
virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_prolong_op ( ) const  
[pure virtual]
```

Returns the prolong operator.

Returns

the prolong operator.

Implemented in [gko::multigrid::EnableMultigridLevel< ValueType >](#).

47.160.2.4 `get_restrict_op()`

```
virtual std::shared_ptr<const LinOp> gko::multigrid::MultigridLevel::get_restrict_op ( ) const
[pure virtual]
```

Returns the restrict operator.

Returns

the restrict operator.

Implemented in [gko::multigrid::EnableMultigridLevel< ValueType >](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/multigrid/multigrid_level.hpp`

47.161 `gko::batch::MultiVector< ValueType >` Class Template Reference

[MultiVector](#) stores multiple vectors in a batched fashion and is useful for batched operations.

```
#include <ginkgo/core/base/batch_multi_vector.hpp>
```

Public Member Functions

- `std::unique_ptr< unbatch_type > create_view_for_item (size_type item_id)`
Creates a mutable view (of [matrix::Dense](#) type) of one item of the Batch [MultiVector](#) object.
- `std::unique_ptr< const unbatch_type > create_const_view_for_item (size_type item_id) const`
Creates a mutable view (of [matrix::Dense](#) type) of one item of the Batch [MultiVector](#) object.
- `batch_dim< 2 > get_size () const`
Returns the batch size.
- `size_type get_num_batch_items () const`
Returns the number of batch items.
- `dim< 2 > get_common_size () const`
Returns the common size of the batch items.
- `value_type * get_values () noexcept`
Returns a pointer to the array of values of the multi-vector.
- `const value_type * get_const_values () const noexcept`
Returns a pointer to the array of values of the multi-vector.
- `value_type * get_values_for_item (size_type batch_id) noexcept`
Returns a pointer to the array of values of the multi-vector for a specific batch item.
- `const value_type * get_const_values_for_item (size_type batch_id) const noexcept`
Returns a pointer to the array of values of the multi-vector for a specific batch item.
- `size_type get_num_stored_elements () const noexcept`
Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.
- `size_type get_cumulative_offset (size_type batch_id) const`
Get the cumulative storage size offset.
- `value_type & at (size_type batch_id, size_type row, size_type col)`
Returns a single element for a particular batch item.

- `value_type at (size_type batch_id, size_type row, size_type col) const`
Returns a single element for a particular batch item.
- `ValueType & at (size_type batch_id, size_type idx) noexcept`
Returns a single element for a particular batch item.
- `ValueType at (size_type batch_id, size_type idx) const noexcept`
Returns a single element for a particular batch item.
- `void scale (ptr_param< const MultiVector< ValueType >> alpha)`
Scales the vector with a scalar (aka: BLAS scal).
- `void add_scaled (ptr_param< const MultiVector< ValueType >> alpha, ptr_param< const MultiVector< ValueType >> b)`
Adds b scaled by alpha to the vector (aka: BLAS axpy).
- `void compute_dot (ptr_param< const MultiVector< ValueType >> b, ptr_param< MultiVector< ValueType >> result) const`
Computes the column-wise dot product of each multi-vector in this batch and its corresponding entry in b.
- `void compute_conj_dot (ptr_param< const MultiVector< ValueType >> b, ptr_param< MultiVector< ValueType >> result) const`
Computes the column-wise conjugate dot product of each multi-vector in this batch and its corresponding entry in b.
- `void compute_norm2 (ptr_param< MultiVector< remove_complex< ValueType >>> result) const`
Computes the Euclidean (L^2) norm of each multi-vector in this batch.
- `void fill (ValueType value)`
Fills the input MultiVector with a given value.

Static Public Member Functions

- `static std::unique_ptr< MultiVector > create_with_config_of (ptr_param< const MultiVector > other)`
Creates a MultiVector with the configuration of another MultiVector.
- `static std::unique_ptr< MultiVector > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size=batch_dim< 2 >{})`
Creates an uninitialized multi-vector of the specified size.
- `static std::unique_ptr< MultiVector > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, array< value_type > values)`
Creates a MultiVector from an already allocated (and initialized) array.
- `template<typename InputValueType >`
`static std::unique_ptr< MultiVector > create (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &size, std::initializer_list< InputValueType > values)`
create(std::shared_ptr<const
- `static std::unique_ptr< const MultiVector > create_const (std::shared_ptr< const Executor > exec, const batch_dim< 2 > &sizes, gko::detail::const_array_view< ValueType > &&values)`
Creates a constant (immutable) batch multi-vector from a constant array.

47.161.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::batch::MultiVector< ValueType >
```

`MultiVector` stores multiple vectors in a batched fashion and is useful for batched operations.

For example, if you want to store two batch items with multi-vectors of size (3 x 2) given below:

```
[1 2 ; 3 4 1 2 ; 3 4 1 2 ; 3 4]
```

In memory, they would be stored as a single array: [1 2 1 2 1 2 3 4 3 4 3 4].

Access functions `@at` can help access individual item if necessary.

The values of the different batch items are stored consecutively and in each batch item, the multi-vectors are stored in a row-major fashion.

Template Parameters

<i>ValueType</i>	precision of multi-vector elements
------------------	------------------------------------

47.161.2 Member Function Documentation

47.161.2.1 `add_scaled()`

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::add_scaled (
    ptr_param< const MultiVector< ValueType >> alpha,
    ptr_param< const MultiVector< ValueType >> b )
```

Adds `b` scaled by `alpha` to the vector (aka: BLAS axpy).

Parameters

<i>alpha</i>	the scalar
<i>b</i>	a multi-vector of the same dimension as this

Note

If `alpha` is 1x1 [MultiVector](#) matrix, the entire multi-vector (all batches) is scaled by `alpha`. If it is a [MultiVector](#) row vector of values, then *i*-th column of the vector is scaled with the *i*-th element of `alpha` (the number of columns of `alpha` has to match the number of columns of the multi-vector).

47.161.2.2 `at()` [1/4]

```
template<typename ValueType = default_precision>
ValueType gko::batch::MultiVector< ValueType >::at (
    size_type batch_id,
    size_type idx ) const [inline], [noexcept]
```

Returns a single element for a particular batch item.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

```

267     {
268         return values_.get_const_data() [linearize_index(batch_id, idx)];
269     }

```

References `gko::array< ValueType >::get_const_data()`.

47.161.2.3 at() [2/4]

```

template<typename ValueType = default_precision>
ValueType& gko::batch::MultiVector< ValueType >::at (
    size_type batch_id,
    size_type idx ) [inline], [noexcept]

```

Returns a single element for a particular batch item.

Useful for iterating across all elements of the vector. However, it is less efficient than the two-parameter variant of this method.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>idx</i>	a linear index of the requested element

Note

the method has to be called on the same [Executor](#) the vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

References `gko::array< ValueType >::get_data()`.

47.161.2.4 at() [3/4]

```

template<typename ValueType = default_precision>
value_type& gko::batch::MultiVector< ValueType >::at (
    size_type batch_id,
    size_type row,
    size_type col ) [inline]

```

Returns a single element for a particular batch item.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

References `gko::array< ValueType >::get_data()`, and `gko::batch::MultiVector< ValueType >::get_num_batch_items()`.

47.161.2.5 at() [4/4]

```
template<typename ValueType = default_precision>
value_type gko::batch::MultiVector< ValueType >::at (
    size_type batch_id,
    size_type row,
    size_type col ) const [inline]
```

Returns a single element for a particular batch item.

Parameters

<i>batch_id</i>	the batch item index to be queried
<i>row</i>	the row of the requested element
<i>col</i>	the column of the requested element

Note

the method has to be called on the same [Executor](#) the vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

References `gko::array< ValueType >::get_const_data()`, and `gko::batch::MultiVector< ValueType >::get_num_batch_items()`.

47.161.2.6 compute_conj_dot()

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::compute_conj_dot (
    ptr_param< const MultiVector< ValueType >> b,
    ptr_param< MultiVector< ValueType >> result ) const
```

Computes the column-wise conjugate dot product of each multi-vector in this batch and its corresponding entry in `b`.

If the vector has complex `value_type`, then the conjugate of this is taken.

Parameters

<i>b</i>	a MultiVector of same dimension as this
<i>result</i>	a MultiVector row vector, used to store the dot product (the number of column in the vector must match the number of columns of this)

47.161.2.7 compute_dot()

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::compute_dot (
    ptr_param< const MultiVector< ValueType >> b,
    ptr_param< MultiVector< ValueType >> result ) const
```

Computes the column-wise dot product of each multi-vector in this batch and its corresponding entry in *b*.

Parameters

<i>b</i>	a MultiVector of same dimension as this
<i>result</i>	a MultiVector row vector, used to store the dot product

47.161.2.8 compute_norm2()

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::compute_norm2 (
    ptr_param< MultiVector< remove_complex< ValueType >>> result ) const
```

Computes the Euclidean (L^2) norm of each multi-vector in this batch.

Parameters

<i>result</i>	a MultiVector , used to store the norm (the number of columns in the vector must match the number of columns of this)
---------------	---

47.161.2.9 create() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<MultiVector> gko::batch::MultiVector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    array< value_type > values ) [static]
```

Creates a [MultiVector](#) from an already allocated (and initialized) array.

Parameters

<i>exec</i>	Executor associated to the vector
<i>size</i>	sizes of the batch matrices in a batch_dim object
<i>values</i>	array of values

Note

If `values` is not an rvalue, not an array of `ValueType`, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the vector.

47.161.2.10 create() [2/3]

```
template<typename ValueType = default_precision>
template<typename InputValueType >
static std::unique_ptr<MultiVector> gko::batch::MultiVector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size,
    std::initializer_list< InputValueType > values ) [inline], [static]
```

`create(std::shared_ptr<const`

`create(std::shared_ptr<const executor>="", const batch_dim<2>&, array<value_type>)`

References `gko::batch::MultiVector< ValueType >::create()`.

47.161.2.11 create() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<MultiVector> gko::batch::MultiVector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & size = batch_dim< 2 >{} ) [static]
```

Creates an uninitialized multi-vector of the specified size.

Parameters

<i>exec</i>	<code>Executor</code> associated to the vector
<i>size</i>	size of the batch multi vector

Returns

A smart pointer to the newly created matrix.

Referenced by `gko::batch::MultiVector< ValueType >::create()`.

47.161.2.12 create_const()

```
template<typename ValueType = default_precision>
static std::unique_ptr<const MultiVector> gko::batch::MultiVector< ValueType >::create_const
```

```
(
    std::shared_ptr< const Executor > exec,
    const batch_dim< 2 > & sizes,
    gko::detail::const_array_view< ValueType > && values ) [static]
```

Creates a constant (immutable) batch multi-vector from a constant array.

Parameters

<i>exec</i>	the executor to create the vector on
<i>size</i>	the dimensions of the vector
<i>values</i>	the value array of the vector
<i>stride</i>	the row-stride of the vector

Returns

A smart pointer to the constant multi-vector wrapping the input array (if it resides on the same executor as the vector) or a copy of the array on the correct executor.

47.161.2.13 create_const_view_for_item()

```
template<typename ValueType = default_precision>
std::unique_ptr<const unbatch_type> gko::batch::MultiVector< ValueType >::create_const_view↵
_for_item (
    size_type item_id ) const
```

Creates a mutable view (of `matrix::Dense` type) of one item of the Batch `MultiVector` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<i>item↵ _id</i>	The index of the batch item
----------------------	-----------------------------

Returns

a `matrix::Dense` object with the data from the batch item at the given index.

47.161.2.14 create_view_for_item()

```
template<typename ValueType = default_precision>
std::unique_ptr<unbatch_type> gko::batch::MultiVector< ValueType >::create_view_for_item (
    size_type item_id )
```

Creates a mutable view (of `matrix::Dense` type) of one item of the Batch `MultiVector` object.

Does not perform any deep copies, but only returns a view of the data.

Parameters

<i>item</i> ↔ _id	The index of the batch item
----------------------	-----------------------------

Returns

a [matrix::Dense](#) object with the data from the batch item at the given index.

47.161.2.15 create_with_config_of()

```
template<typename ValueType = default_precision>
static std::unique_ptr<MultiVector> gko::batch::MultiVector< ValueType >::create_with_↔
config_of (
    ptr_param< const MultiVector< ValueType > > other ) [static]
```

Creates a [MultiVector](#) with the configuration of another [MultiVector](#).

Parameters

<i>other</i>	The other multi-vector whose configuration needs to copied.
--------------	---

47.161.2.16 fill()

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::fill (
    ValueType value )
```

Fills the input [MultiVector](#) with a given value.

Parameters

<i>value</i>	the value to be filled
--------------	------------------------

47.161.2.17 get_common_size()

```
template<typename ValueType = default_precision>
dim<2> gko::batch::MultiVector< ValueType >::get_common_size ( ) const [inline]
```

Returns the common size of the batch items.

Returns

the common size stored

References gko::batch_dim< Dimensionality, DimensionType >::get_common_size().

Referenced by gko::batch::MultiVector< ValueType >::get_cumulative_offset().

47.161.2.18 get_const_values()

```
template<typename ValueType = default_precision>
const value_type* gko::batch::MultiVector< ValueType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns a pointer to the array of values of the multi-vector.

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.161.2.19 get_const_values_for_item()

```
template<typename ValueType = default_precision>
const value_type* gko::batch::MultiVector< ValueType >::get_const_values_for_item (
    size_type batch_id ) const [inline], [noexcept]
```

Returns a pointer to the array of values of the multi-vector for a specific batch item.

Parameters

<i>batch_id</i>	the id of the batch item.
-----------------	---------------------------

Returns

the pointer to the array of values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`, `gko::batch::MultiVector< ValueType >::get_cumulative_offset()`, and `gko::batch::MultiVector< ValueType >::get_num_batch_items()`.

47.161.2.20 get_cumulative_offset()

```
template<typename ValueType = default_precision>
size_type gko::batch::MultiVector< ValueType >::get_cumulative_offset (
    size_type batch_id ) const [inline]
```

Get the cumulative storage size offset.

Parameters

<i>batch_id</i>	the batch id
-----------------	--------------

Returns

the cumulative offset

References `gko::batch::MultiVector< ValueType >::get_common_size()`.

Referenced by `gko::batch::MultiVector< ValueType >::get_const_values_for_item()`, and `gko::batch::MultiVector< ValueType >::get_values_for_item()`.

47.161.2.21 get_num_batch_items()

```
template<typename ValueType = default_precision>
size_type gko::batch::MultiVector< ValueType >::get_num_batch_items ( ) const [inline]
```

Returns the number of batch items.

Returns

the number of batch items

References `gko::batch_dim< Dimensionality, DimensionType >::get_num_batch_items()`.

Referenced by `gko::batch::MultiVector< ValueType >::at()`, `gko::batch::MultiVector< ValueType >::get_const_values_for_item()`, and `gko::batch::MultiVector< ValueType >::get_values_for_item()`.

47.161.2.22 get_num_stored_elements()

```
template<typename ValueType = default_precision>
size_type gko::batch::MultiVector< ValueType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the batch matrix, cumulative across all the batch items.

Returns

the number of elements explicitly stored in the vector, cumulative across all the batch items

References gko::array< ValueType >::get_size().

47.161.2.23 get_size()

```
template<typename ValueType = default_precision>
batch_dim<2> gko::batch::MultiVector< ValueType >::get_size ( ) const [inline]
```

Returns the batch size.

Returns

the batch size

47.161.2.24 get_values()

```
template<typename ValueType = default_precision>
value_type* gko::batch::MultiVector< ValueType >::get_values ( ) [inline], [noexcept]
```

Returns a pointer to the array of values of the multi-vector.

Returns

the pointer to the array of values

References gko::array< ValueType >::get_data().

47.161.2.25 get_values_for_item()

```
template<typename ValueType = default_precision>
value_type* gko::batch::MultiVector< ValueType >::get_values_for_item (
    size_type batch_id ) [inline], [noexcept]
```

Returns a pointer to the array of values of the multi-vector for a specific batch item.

Parameters

<i>batch</i> ↔ <i>_id</i>	the id of the batch item.
------------------------------	---------------------------

Returns

the pointer to the array of values

References `gko::batch::MultiVector< ValueType >::get_cumulative_offset()`, `gko::array< ValueType >::get_data()`, and `gko::batch::MultiVector< ValueType >::get_num_batch_items()`.

47.161.2.26 scale()

```
template<typename ValueType = default_precision>
void gko::batch::MultiVector< ValueType >::scale (
    ptr_param< const MultiVector< ValueType >> alpha )
```

Scales the vector with a scalar (aka: BLAS scal).

Parameters

<i>alpha</i>	the scalar
--------------	------------

Note

If alpha is 1x1 [MultiVector](#) matrix, the entire multi-vector (all batches) is scaled by alpha. If it is a [MultiVector](#) row vector of values, then i-th column of the vector is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of columns of the multi-vector). If it is a [MultiVector](#) of the same size as this, then an element wise scaling is performed.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/batch_multi_vector.hpp`

47.162 gko::log::ProfilerHook::NestedSummaryWriter Class Reference

Receives the results from [ProfilerHook::create_nested_summary\(\)](#).

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

- virtual void [write_nested](#) (const nested_summary_entry &root, std::chrono::nanoseconds overhead)=0
Callback to write out the summary results.

47.162.1 Detailed Description

Receives the results from [ProfilerHook::create_nested_summary\(\)](#).

47.162.2 Member Function Documentation

47.162.2.1 write_nested()

```
virtual void gko::log::ProfilerHook::NestedSummaryWriter::write_nested (
    const nested_summary_entry & root,
    std::chrono::nanoseconds overhead ) [pure virtual]
```

Callback to write out the summary results.

Parameters

<i>root</i>	the root range with runtime and count.
<i>overhead</i>	an estimate of the profiler overhead

Implemented in [gko::log::ProfilerHook::TableSummaryWriter](#).

The documentation for this class was generated from the following file:

- ginkgo/core/log/profiler_hook.hpp

47.163 gko::NotCompiled Class Reference

[NotCompiled](#) is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [NotCompiled](#) (const std::string &file, int line, const std::string &func, const std::string &module)
Initializes a [NotCompiled](#) error.

47.163.1 Detailed Description

[NotCompiled](#) is thrown when attempting to call an operation which is a part of a module that was not compiled on the system.

47.163.2 Constructor & Destructor Documentation

47.163.2.1 NotCompiled()

```
gko::NotCompiled::NotCompiled (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & module ) [inline]
```

Initializes a [NotCompiled](#) error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the function that has not been compiled
<i>module</i>	The name of the module which contains the function

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.164 gko::NotImplemented Class Reference

[NotImplemented](#) is thrown in case an operation has not yet been implemented (but will be implemented in the future).

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [NotImplemented](#) (const std::string &file, int line, const std::string &func)
Initializes a [NotImplemented](#) error.

47.164.1 Detailed Description

[NotImplemented](#) is thrown in case an operation has not yet been implemented (but will be implemented in the future).

47.164.2 Constructor & Destructor Documentation

47.164.2.1 NotImplemented()

```
gko::NotImplemented::NotImplemented (
    const std::string & file,
    int line,
    const std::string & func ) [inline]
```

Initializes a [NotImplemented](#) error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the not-yet implemented function

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.165 gko::NotSupported Class Reference

[NotSupported](#) is thrown in case it is not possible to perform the requested operation on the given object type.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [NotSupported](#) (const std::string &file, int line, const std::string &func, const std::string &obj_type)
Initializes a [NotSupported](#) error.

47.165.1 Detailed Description

[NotSupported](#) is thrown in case it is not possible to perform the requested operation on the given object type.

47.165.2 Constructor & Destructor Documentation

47.165.2.1 NotSupported()

```
gko::NotSupported::NotSupported (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & obj_type ) [inline]
```

Initializes a [NotSupported](#) error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the function where the error occurred
<i>obj_type</i>	The object type on which the requested operation cannot be performed.

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.166 gko::null_deleter< T > Class Template Reference

This is a deleter that does not delete the object.

```
#include <ginkgo/core/base/utils_helper.hpp>
```

Public Member Functions

- void [operator\(\)](#) (pointer) const noexcept
Deletes the object.

47.166.1 Detailed Description

```
template<typename T>
class gko::null_deleter< T >
```

This is a deleter that does not delete the object.

It is useful where the object has been allocated elsewhere and will be deleted manually.

47.166.2 Member Function Documentation

47.166.2.1 operator>()

```
template<typename T >
void gko::null_deleter< T >::operator() (
    pointer ) const [inline], [noexcept]
```

Deletes the object.

Parameters

<i>ptr</i>	pointer to the object being deleted
------------	-------------------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/base/utils_helper.hpp

47.167 gko::nvidia_device Class Reference

[nvidia_device](#) handles the number of executor on Nvidia devices and have the corresponding recursive_mutex.

```
#include <ginkgo/core/base/device.hpp>
```

47.167.1 Detailed Description

[nvidia_device](#) handles the number of executor on Nvidia devices and have the corresponding recursive_mutex.

The documentation for this class was generated from the following file:

- ginkgo/core/base/device.hpp

47.168 gko::OmpExecutor Class Reference

This is the [Executor](#) subclass which represents the OpenMP device (typically CPU).

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- `std::shared_ptr< Executor > get_master ()` noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `std::shared_ptr< const Executor > get_master ()` const noexcept override
Returns the master [OmpExecutor](#) of this [Executor](#).
- `void synchronize ()` const override
Synchronize the operations launched on the executor with its master.
- `std::string get_description ()` const override
- `virtual void run (const Operation &op)` const=0
Runs the specified [Operation](#) using this [Executor](#).
- `template<typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename ClosureDpcpp >`
`void run (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.
- `template<typename ClosureReference, typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename Closure↔`
`Dpcpp >`
`void run (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↔`
`Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const`
Runs one of the passed in functors, depending on the [Executor](#) type.

Static Public Member Functions

- static std::shared_ptr< [OmpExecutor](#) > [create](#) (std::shared_ptr< [CpuAllocatorBase](#) > [alloc](#)=std::make_shared< [CpuAllocator](#) >())
Creates a new [OmpExecutor](#).

47.168.1 Detailed Description

This is the [Executor](#) subclass which represents the OpenMP device (typically CPU).

47.168.2 Member Function Documentation

47.168.2.1 [get_description\(\)](#)

```
std::string gko::OmpExecutor::get_description ( ) const [override], [virtual]
```

Returns

a textual representation of the executor and its device.

Implements [gko::Executor](#).

Reimplemented in [gko::ReferenceExecutor](#).

47.168.2.2 [get_master\(\)](#) [1/2]

```
std::shared_ptr<const Executor> gko::OmpExecutor::get_master ( ) const [override], [virtual],  
[noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.168.2.3 `get_master()` [2/2]

```
std::shared_ptr<Executor> gko::OmpExecutor::get_master ( ) [override], [virtual], [noexcept]
```

Returns the master [OmpExecutor](#) of this [Executor](#).

Returns

the master [OmpExecutor](#) of this [Executor](#).

Implements [gko::Executor](#).

47.168.2.4 `run()` [1/3]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of <code>op_omp</code>
<i>ClosureCuda</i>	type of <code>op_cuda</code>
<i>ClosureHip</i>	type of <code>op_hip</code>
<i>ClosureDpcpp</i>	type of <code>op_dpcpp</code>

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

47.168.2.5 `run()` [2/3]

```
virtual void gko::Executor::run
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

47.168.2.6 run() [3/3]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureReference ,
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of op_ref
<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.169 gko::Operation Class Reference

Operations can be used to define functionalities whose implementations differ among devices.

```
#include <ginkgo/core/base/executor.hpp>
```


Public Member Functions

- virtual const char * [get_name](#) () const noexcept

Returns the operation's name.

47.169.1 Detailed Description

Operations can be used to define functionalities whose implementations differ among devices.

This is done by extending the [Operation](#) class and implementing the overloads of the `Operation::run()` method for all [Executor](#) types. When invoking the `Executor::run()` method with the [Operation](#) as input, the library will select the `Operation::run()` overload corresponding to the dynamic type of the [Executor](#) instance.

Consider an overload of `operator<<` for Executors, which prints some basic device information (e.g. device type and id) of the [Executor](#) to a C++ stream:

```
std::ostream& operator<<(std::ostream &os, const gko::Executor &exec);
```

One possible implementation would be to use RTTI to find the dynamic type of the [Executor](#). However, using the [Operation](#) feature of Ginkgo, there is a more elegant approach which utilizes polymorphism. The first step is to define an [Operation](#) that will print the desired information for each [Executor](#) type.

```
class DeviceInfoPrinter : public gko::Operation {
public:
    explicit DeviceInfoPrinter(std::ostream &os) : os_(os) {}
    void run(const gko::OmpExecutor *) const override { os_ << "OMP"; }
    void run(const gko::CudaExecutor *exec) const override
{ os_ << "CUDA(" << exec->get_device_id() << ")"; }
    void run(const gko::HipExecutor *exec) const override
{ os_ << "HIP(" << exec->get_device_id() << ")"; }
    void run(const gko::DpcppExecutor *exec) const override
{ os_ << "DPC++(" << exec->get_device_id() << ")"; }
    // This is optional, if not overloaded, defaults to OmpExecutor overload
    void run(const gko::ReferenceExecutor *) const override
{ os_ << "Reference CPU"; }
private:
    std::ostream &os_;
};
```

Using `DeviceInfoPrinter`, the implementation of `operator<<` is as simple as calling the `run()` method of the executor.

```
std::ostream& operator<<(std::ostream &os, const gko::Executor &exec)
{
    DeviceInfoPrinter printer(os);
    exec.run(printer);
    return os;
}
```

Now it is possible to write the following code:

```
auto omp = gko::OmpExecutor::create();
std::cout << *omp << std::endl
    << *gko::CudaExecutor::create(0, omp) << std::endl
    << *gko::HipExecutor::create(0, omp) << std::endl
    << *gko::DpcppExecutor::create(0, omp) << std::endl
    << *gko::ReferenceExecutor::create() << std::endl;
```

which produces the expected output:

```
OMP
CUDA(0)
HIP(0)
DPC++(0)
Reference CPU
```

One might feel that this code is too complicated for such a simple task. Luckily, there is an overload of the `Executor::run()` method, which is designed to facilitate writing simple operations like this one. The method takes four closures as input: one which is run for OMP, one for CUDA executors, one for HIP executors, and the last one for DPC++ executors. Using this method, there is no need to implement an [Operation](#) subclass:

```
std::ostream& operator<<(std::ostream &os, const gko::Executor &exec)
{
    exec.run(
        [&]() { os << "OMP"; }, // OMP closure
```

```

[&]() { os << "CUDA(" // CUDA closure
    << static_cast<gko::CudaExecutor&>(exec)
        .get_device_id()
    << ")"; },
[&]() { os << "HIP(" // HIP closure
    << static_cast<gko::HipExecutor&>(exec)
        .get_device_id()
    << ")"; });
[&]() { os << "DPC++(" // DPC++ closure
    << static_cast<gko::DpcppExecutor&>(exec)
        .get_device_id()
    << ")"; });
return os;
}

```

Using this approach, however, it is impossible to distinguish between a [OmpExecutor](#) and [ReferenceExecutor](#), as both of them call the OMP closure.

47.169.2 Member Function Documentation

47.169.2.1 get_name()

```
virtual const char* gko::Operation::get_name ( ) const [virtual], [noexcept]
```

Returns the operation's name.

Returns

the operation's name

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.170 gko::log::operation_data Struct Reference

Struct representing Operator related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.170.1 Detailed Description

Struct representing Operator related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.171 gko::OutOfBoundsError Class Reference

[OutOfBoundsError](#) is thrown if a memory access is detected to be out-of-bounds.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [OutOfBoundsError](#) (const std::string &file, int line, [size_type](#) index, [size_type](#) bound)
Initializes an [OutOfBoundsError](#).

47.171.1 Detailed Description

[OutOfBoundsError](#) is thrown if a memory access is detected to be out-of-bounds.

47.171.2 Constructor & Destructor Documentation

47.171.2.1 OutOfBoundsError()

```
gko::OutOfBoundsError::OutOfBoundsError (
    const std::string & file,
    int line,
    size\_type index,
    size\_type bound ) [inline]
```

Initializes an [OutOfBoundsError](#).

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>index</i>	The position that was accessed
<i>bound</i>	The first out-of-bound index

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.172 gko::OverflowError Class Reference

[OverflowError](#) is thrown when an index calculation for storage requirements overflows.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [OverflowError](#) (const std::string &file, const int line, const std::string &index_type)

47.172.1 Detailed Description

[OverflowError](#) is thrown when an index calculation for storage requirements overflows.

This most likely means that the index type is too small.

47.172.2 Constructor & Destructor Documentation

47.172.2.1 OverflowError()

```
gko::OverflowError::OverflowError (
    const std::string & file,
    const int line,
    const std::string & index_type ) [inline]
```

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>index_type</i>	The integer type that overflowed

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.173 gko::log::Papi< ValueType > Class Template Reference

[Papi](#) is a Logger which logs every event to the PAPI software.

```
#include <ginkgo/core/log/papi.hpp>
```

Public Member Functions

- const std::string [get_handle_name](#) () const
Returns the unique name of this logger, which can be used in the PAPI_read() call.
- const papi_handle_t [get_handle](#) () const
Returns the corresponding papi_handle_t for this logger.

Static Public Member Functions

- static std::shared_ptr< [Papi](#) > [create](#) (std::shared_ptr< const [gko::Executor](#) >, const Logger::mask_type &enabled_events=Logger::all_events_mask)
Creates a [Papi](#) Logger.
- static std::shared_ptr< [Papi](#) > [create](#) (const Logger::mask_type &enabled_events=Logger::all_events_mask)
Creates a [Papi](#) Logger.

47.173.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::log::Papi< ValueType >
```

[Papi](#) is a Logger which logs every event to the PAPI software.

Thanks to this logger, applications which interface with PAPI can access Ginkgo internal data through PAPI. For an example of usage, see `examples/papi_logging/papi_logging.cpp`

The logged values for each event are the following:

- all allocation events: number of bytes per executor
- all free events: number of calls per executor
- copy_started: number of bytes per executor from (to), in copy_started_from (respectively copy_started_to).
- copy_completed: number of bytes per executor from (to), in copy_completed_from (respectively copy_completed_to).
- all polymorphic objects and operation events: number of calls per executor
- all apply events: number of calls per LinOp (argument "A").
- all factory events: number of calls per factory
- criterion_check_completed event: the residual norm is stored in a record (per criterion)
- iteration_complete event: the number of iteration is counted (per solver)

Template Parameters

<i>ValueType</i>	the type of values stored in the class (e.g. residuals)
------------------	---

47.173.2 Member Function Documentation

47.173.2.1 [create\(\)](#) [1/2]

```
template<typename ValueType = default_precision>
static std::shared_ptr<Papi> gko::log::Papi< ValueType >::create (
```

```
const Logger::mask_type & enabled_events = Logger::all_events_mask ) [inline],
[static]
```

Creates a [Papi](#) Logger.

Parameters

<i>enabled_events</i>	the events enabled for this Logger
-----------------------	------------------------------------

```
195     {
196         return std::shared_ptr<Papi>(new Papi(enabled_events), [](auto logger) {
197             auto handle = logger->get_handle();
198             delete logger;
199             papi_sde_shutdown(handle);
200         });
201     }
```

References `gko::log::Papi< ValueType >::get_handle()`.

47.173.2.2 create() [2/2]

```
template<typename ValueType = default_precision>
static std::shared_ptr<Papi> gko::log::Papi< ValueType >::create (
    std::shared_ptr< const gko::Executor > ,
    const Logger::mask_type & enabled_events = Logger::all_events_mask ) [inline],
[static]
```

Creates a [Papi](#) Logger.

Parameters

<i>enabled_events</i>	the events enabled for this Logger
-----------------------	------------------------------------

47.173.2.3 get_handle()

```
template<typename ValueType = default_precision>
const papi_handle_t gko::log::Papi< ValueType >::get_handle ( ) const [inline]
```

Returns the corresponding `papi_handle_t` for this logger.

Returns

the corresponding `papi_handle_t` for this logger

Referenced by `gko::log::Papi< ValueType >::create()`.

47.173.2.4 get_handle_name()

```
template<typename ValueType = default_precision>
const std::string gko::log::Papi< ValueType >::get_handle_name ( ) const [inline]
```

Returns the unique name of this logger, which can be used in the PAPI_read() call.

Returns

the unique name of this logger

The documentation for this class was generated from the following file:

- ginkgo/core/log/papi.hpp

47.174 gko::factorization::Parlc< ValueType, IndexType > Class Template Reference

ParlC is an incomplete Cholesky factorization which is computed in parallel.

```
#include <ginkgo/core/factorization/par_ic.hpp>
```

Static Public Member Functions

- static parameters_type `parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

Additional Inherited Members

47.174.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::factorization::Parlc< ValueType, IndexType >
```

ParlC is an incomplete Cholesky factorization which is computed in parallel.

L is a lower triangular matrix, which approximates a given matrix A with $A \approx LL^H$. Here, $L + L^H$ has the same sparsity pattern as A , which is also called IC(0).

The ParlC algorithm generates the incomplete factors iteratively, using a fixed-point iteration of the form

$$F(L) = \begin{cases} \sqrt{a_{ii} - \sum_{k=1}^{i-1} |l_{ik}|^2}, & i == j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i < j \end{cases}$$

In general, the entries of L can be iterated in parallel and in asynchronous fashion, the algorithm asymptotically converges to the incomplete factors L and L^H fulfilling $(R = A - L \cdot L^H)|_{\mathcal{S}} = 0|_{\mathcal{S}}$ where $\mathcal{S}$ is the pre-defined sparsity pattern (in case of IC(0) the sparsity pattern of the system matrix A). The number of ParlC sweeps needed for convergence depends on the parallelism level: For sequential execution, a single sweep is sufficient, for fine-grained parallelism, the number of sweeps necessary to get a good approximation of the incomplete factors depends heavily on the problem. On the OpenMP executor, 3 sweeps usually give a decent approximation in our experiments, while GPU executors can take 10 or more iterations.

The ParlC algorithm in Ginkgo follows the design of E. Chow and A. Patel, Fine-grained Parallel Incomplete LU Factorization, SIAM Journal on Scientific Computing, 37, C169-C193 (2015).

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.174.2 Member Function Documentation

47.174.2.1 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::factorization::ParIc< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/par_ic.hpp

47.175 gko::factorization::ParIc< ValueType, IndexType > Class Template Reference

ParICT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

```
#include <ginkgo/core/factorization/par_ict.hpp>
```


Static Public Member Functions

- static parameters_type `parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

Additional Inherited Members

47.175.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::factorization::Parlct< ValueType, IndexType >
```

ParlCT is an incomplete threshold-based Cholesky factorization which is computed in parallel.

L is a lower triangular matrix which approximates a given symmetric positive definite matrix A with $A \approx LL^T$. Here, L has a sparsity pattern that is improved iteratively based on its element-wise magnitude. The initial sparsity pattern is chosen based on the lower triangle of A .

One iteration of the ParlCT algorithm consists of the following steps:

1. Calculating the residual $R = A - LL^T$
2. Adding new non-zero locations from R to L . The new non-zero locations are initialized based on the corresponding residual value.
3. Executing a fixed-point iteration on L according to $F(L) = \begin{cases} \frac{1}{l_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} l_{jk} \right), & i \neq j \\ \sqrt{a_{ij} - \sum_{k=1}^{j-1} l_{ik} l_{jk}}, & i = j \end{cases}$
4. Removing the smallest entries (by magnitude) from L
5. Executing a fixed-point iteration on the (now sparser) L

This ParlCT algorithm thus improves the sparsity pattern and the approximation of L simultaneously.

The implementation follows the design of H. Anzt et al., ParlLUT - A Parallel Threshold ILU for GPUs, 2019 IEEE International Parallel and Distributed Processing Symposium (IPDPS), pp. 231–241.

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.175.2 Member Function Documentation

47.175.2.1 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::factorization::ParIct< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/par_ict.hpp

47.176 gko::factorization::ParIlu< ValueType, IndexType > Class Template Reference

ParILU is an incomplete LU factorization which is computed in parallel.

```
#include <ginkgo/core/factorization/par_ilu.hpp>
```

Static Public Member Functions

- static parameters_type **parse** (const config::pnode &config, const config::registry &context, const config::type_descriptor &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

Additional Inherited Members

47.176.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::factorization::Parllu< ValueType, IndexType >
```

ParLLU is an incomplete LU factorization which is computed in parallel.

L is a lower unitriangular, while U is an upper triangular matrix, which approximate a given matrix A with $A \approx LU$. Here, L and U have the same sparsity pattern as A , which is also called ILU(0).

The ParLLU algorithm generates the incomplete factors iteratively, using a fixed-point iteration of the form

$$F(L, U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i > j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i \leq j \end{cases}$$

In general, the entries of L and U can be iterated in parallel and in asynchronous fashion, the algorithm asymptotically converges to the incomplete factors L and U fulfilling $(R = A - L \cdot U)|_S = 0|_S$ where S is the pre-defined sparsity pattern (in case of ILU(0) the sparsity pattern of the system matrix A). The number of ParLLU sweeps needed for convergence depends on the parallelism level: For sequential execution, a single sweep is sufficient, for fine-grained parallelism, the number of sweeps necessary to get a good approximation of the incomplete factors depends heavily on the problem. On the OpenMP executor, 3 sweeps usually give a decent approximation in our experiments, while GPU executors can take 10 or more iterations.

The ParLLU algorithm in Ginkgo follows the design of E. Chow and A. Patel, Fine-grained Parallel Incomplete LU Factorization, SIAM Journal on Scientific Computing, 37, C169-C193 (2015).

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.176.2 Member Function Documentation

47.176.2.1 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::factorization::Parllu< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/par_ilu.hpp

47.177 gko::factorization::Parllut< ValueType, IndexType > Class Template Reference

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

```
#include <ginkgo/core/factorization/par_ilut.hpp>
```

Static Public Member Functions

- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType, IndexType >())
Create the parameters from the property_tree.

Additional Inherited Members

47.177.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::factorization::Parllut< ValueType, IndexType >
```

ParILUT is an incomplete threshold-based LU factorization which is computed in parallel.

L is a lower unitriangular, while U is an upper triangular matrix, which approximate a given matrix A with $A \approx LU$. Here, L and U have a sparsity pattern that is improved iteratively based on their element-wise magnitude. The initial sparsity pattern is chosen based on the $ILU(0)$ factorization of A .

One iteration of the ParILUT algorithm consists of the following steps:

1. Calculating the residual $R = A - LU$
2. Adding new non-zero locations from R to L and U . The new non-zero locations are initialized based on the corresponding residual value.
3. Executing a fixed-point iteration on L and U according to $F(L, U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i > j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i \leq j \end{cases}$
For a more detailed description of the fixed-point iteration, see [Parllu](#).
4. Removing the smallest entries (by magnitude) from L and U
5. Executing a fixed-point iteration on the (now sparser) L and U

This ParILUT algorithm thus improves the sparsity pattern and the approximation of L and U simultaneously.

The implementation follows the design of H. Anzt et al., ParILUT - A Parallel Threshold ILU for GPUs, 2019 IEEE International Parallel and Distributed Processing Symposium (IPDPS), pp. 231–241.

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.177.2 Member Function Documentation

47.177.2.1 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::factorization::ParIlut< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↵ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/factorization/par_ilut.hpp

47.178 gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType > Class Template Reference

Represents a partition of a range of indices [0, size) into a disjoint set of parts.

```
#include <ginkgo/core/distributed/partition.hpp>
```

Public Member Functions

- `size_type get_size () const`
Returns the total number of elements represented by this partition.
- `size_type get_num_ranges () const noexcept`
Returns the number of ranges stored by this partition.
- `comm_index_type get_num_parts () const noexcept`
Returns the number of parts represented in this partition.
- `comm_index_type get_num_empty_parts () const noexcept`
Returns the number of empty parts within this partition.
- `const global_index_type * get_range_bounds () const noexcept`
Returns the ranges boundary array stored by this partition.
- `const comm_index_type * get_part_ids () const noexcept`
Returns the part IDs of the ranges in this partition.
- `const local_index_type * get_range_starting_indices () const noexcept`
Returns the part-local starting index for each range in this partition.
- `const local_index_type * get_part_sizes () const noexcept`
Returns the part size array.
- `local_index_type get_part_size (comm_index_type part) const`
Returns the size of a part given by its part ID.
- `bool has_connected_parts () const`
Checks if each part has no more than one contiguous range.
- `bool has_ordered_parts () const`
Checks if the ranges are ordered by their part index.

Static Public Member Functions

- `static std::unique_ptr< Partition > build_from_mapping (std::shared_ptr< const Executor > exec, const array< comm_index_type > &mapping, comm_index_type num_parts)`
Builds a partition from a given mapping global_index -> part_id.
- `static std::unique_ptr< Partition > build_from_contiguous (std::shared_ptr< const Executor > exec, const array< global_index_type > &ranges, const array< comm_index_type > &part_ids={})`
Builds a partition consisting of contiguous ranges, one for each part.
- `static std::unique_ptr< Partition > build_from_global_size_uniform (std::shared_ptr< const Executor > exec, comm_index_type num_parts, global_index_type global_size)`
Builds a partition by evenly distributing the global range.

47.178.1 Detailed Description

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
class gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >
```

Represents a partition of a range of indices [0, size) into a disjoint set of parts.

The partition is stored as a set of consecutive ranges [begin, end) with an associated part ID and local index (number of indices in this part before `begin`). Global indices are stored as 64 bit signed integers (int64), part-local indices use LocalIndexType, Part IDs use 32 bit signed integers (int).

For example, consider the interval [0, 13) that is partitioned into the following ranges:
[0, 3), [3, 6), [6, 8), [8, 10), [10, 13).

These ranges are distributed on three part with:

```
p_0 = [0, 3) + [6, 8) + [10, 13),
p_1 = [3, 6),
p_2 = [8, 10).
```

The part ids can be queried from the [get_part_ids](#) array, and the ranges are represented as offsets, accessed by [get_range_bounds](#), leading to the offset array:

```
r = [0, 3, 6, 8, 10, 13]
```

so that individual ranges are given by $[r[i], r[i + 1])$. Since each part may be associated with multiple ranges, it is possible to get the starting index for each range that is local to the owning part, see [get_range_starting_indices](#). These indices can be used to easily iterate over part local data. For example, the above partition has the following starting indices

```
starting_index[0] = 0,
starting_index[1] = 0,
starting_index[2] = 3, // second range of part 0
starting_index[3] = 0,
starting_index[4] = 5, // third range of part 0
```

which you can use to iterate only over the the second range of part 0 (the third global range) with

```
for(int i = 0; i < r[3] - r[2]; ++i){
    data[starting_index[2] + i] = val;
}
```

Template Parameters

<i>LocalIndexType</i>	The index type used for part-local indices. To prevent overflows, no single part's size may exceed this index type's maximum value.
<i>GlobalIndexType</i>	The index type used for the global indices. Needs to be at least as large a type as <i>LocalIndexType</i> .

47.178.2 Member Function Documentation

47.178.2.1 build_from_contiguous()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
static std::unique_ptr<Partition> gko::experimental::distributed::Partition< LocalIndexType,
GlobalIndexType >::build_from_contiguous (
    std::shared_ptr< const Executor > exec,
    const array< global_index_type > & ranges,
    const array< comm_index_type > & part_ids = {} ) [static]
```

Builds a partition consisting of contiguous ranges, one for each part.

Parameters

<i>exec</i>	the Executor on which the partition should be built
<i>ranges</i>	the boundaries of the ranges representing each part. Part <i>part_id[i]</i> contains the indices $[ranges[i], ranges[i + 1])$. Has to contain at least one element. The first element has to be 0.
<i>part_ids</i>	the part ids of the provided ranges. If empty, then it will assume range <i>i</i> belongs to part <i>i</i> .

Returns

a [Partition](#) representing the given contiguous partitioning.

47.178.2.2 build_from_global_size_uniform()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
static std::unique_ptr<Partition> gko::experimental::distributed::Partition< LocalIndexType,
GlobalIndexType >::build_from_global_size_uniform (
    std::shared_ptr< const Executor > exec,
    comm_index_type num_parts,
    global_index_type global_size ) [static]
```

Builds a partition by evenly distributing the global range.

Parameters

<i>exec</i>	the Executor on which the partition should be built
<i>num_parts</i>	the number of parst used in this partition
<i>global_size</i>	the global size of this partition

Returns

a [Partition](#) where each range has either `floor(global_size/num_parts)` or `floor(global_size/num_parts) + 1` indices.

47.178.2.3 build_from_mapping()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
static std::unique_ptr<Partition> gko::experimental::distributed::Partition< LocalIndexType,
GlobalIndexType >::build_from_mapping (
    std::shared_ptr< const Executor > exec,
    const array< comm_index_type > & mapping,
    comm_index_type num_parts ) [static]
```

Builds a partition from a given mapping `global_index -> part_id`.

Parameters

<i>exec</i>	the Executor on which the partition should be built
<i>mapping</i>	the mapping from global indices to part IDs.
<i>num_parts</i>	the number of parts used in the mapping.

Returns

a [Partition](#) representing the given mapping as a set of ranges

47.178.2.4 get_num_empty_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
comm_index_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >↵
::get_num_empty_parts ( ) const [inline], [noexcept]
```

Returns the number of empty parts within this partition.

Returns

number of empty parts.

```
126     {
127         return num_empty_parts_;
128     }
```

47.178.2.5 get_num_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
comm_index_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >↵
::get_num_parts ( ) const [inline], [noexcept]
```

Returns the number of parts represented in this partition.

Returns

number of parts.

47.178.2.6 get_num_ranges()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
size_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_↵
num_ranges ( ) const [inline], [noexcept]
```

Returns the number of ranges stored by this partition.

This size refers to the data returned by [get_range_bounds\(\)](#).

Returns

number of ranges.

References `gko::array< ValueType >::get_size()`.

47.178.2.7 `get_part_ids()`

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const comm_index_type* gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_ids ( ) const [inline], [noexcept]
```

Returns the part IDs of the ranges in this partition.

For each range from `get_range_bounds()`, it stores the part ID in the interval `[0, get_num_parts() - 1]`.

Returns

part ID array.

References `gko::array< ValueType >::get_const_data()`.

47.178.2.8 `get_part_size()`

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
local_index_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_size (
    comm_index_type part ) const
```

Returns the size of a part given by its part ID.

Warning

Triggers a copy from device to host.

Parameters

<i>part</i>	the part ID.
-------------	--------------

Returns

size of part.

47.178.2.9 `get_part_sizes()`

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const local_index_type* gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_part_sizes ( ) const [inline], [noexcept]
```

Returns the part size array.

`part_sizes[p]` stores the total number of indices in part `p`.

Returns

part size array.

References gko::array< ValueType >::get_const_data().

47.178.2.10 get_range_bounds()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const global_index_type* gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_range_bounds ( ) const [inline], [noexcept]
```

Returns the ranges boundary array stored by this partition.

range_bounds[i] is the beginning (inclusive) and range_bounds[i + 1] is the end (exclusive) of the ith range.

Returns

range boundaries array.

References gko::array< ValueType >::get_const_data().

47.178.2.11 get_range_starting_indices()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
const local_index_type* gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_range_starting_indices ( ) const [inline], [noexcept]
```

Returns the part-local starting index for each range in this partition.

Consider the partition on [0, 10) with

```
p_1 = [0-4) + [7-10),
p_2 = [4-7).
```

Then range_starting_indices[0] = 0, range_starting_indices[1] = 0, range_starting_indices[2] = 4.

Returns

part-local starting index array.

References gko::array< ValueType >::get_const_data().

47.178.2.12 get_size()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
size_type gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::get_↵
size ( ) const [inline]
```

Returns the total number of elements represented by this partition.

Returns

number elements.

47.178.2.13 has_connected_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
bool gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::has_↵
connected_parts ( ) const
```

Checks if each part has no more than one contiguous range.

Returns

true if each part has no more than one contiguous range.

47.178.2.14 has_ordered_parts()

```
template<typename LocalIndexType = int32, typename GlobalIndexType = int64>
bool gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >::has_↵
ordered_parts ( ) const
```

Checks if the ranges are ordered by their part index.

Implies that the partition is connected.

Returns

true if the ranges are ordered by their part index.

The documentation for this class was generated from the following files:

- ginkgo/core/distributed/assembly.hpp
- ginkgo/core/distributed/partition.hpp

47.179 gko::log::PerformanceHint Class Reference

[PerformanceHint](#) is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

```
#include <ginkgo/core/log/performance_hint.hpp>
```

Public Member Functions

- void [print_status](#) () const
Writes out the cross-executor writes and allocations that have been stored so far.

Static Public Member Functions

- static std::unique_ptr< [PerformanceHint](#) > [create](#) (std::ostream &os=std::cerr, [size_type](#) allocation_size_limit=16, [size_type](#) copy_size_limit=16, [size_type](#) histogram_max_size=1024)
Creates a [PerformanceHint](#) logger.

47.179.1 Detailed Description

[PerformanceHint](#) is a Logger which analyzes the performance of the application and outputs hints for unnecessary copies and allocations.

The specific patterns it checks for are:

- repeated cross-executor copies from or to the same pointer
- repeated allocation/free pairs of the same size

47.179.2 Member Function Documentation

47.179.2.1 create()

```
static std::unique_ptr<PerformanceHint> gko::log::PerformanceHint::create (
    std::ostream & os = std::cerr,
    size\_type allocation_size_limit = 16,
    size\_type copy_size_limit = 16,
    size\_type histogram_max_size = 1024 ) [inline], [static]
```

Creates a [PerformanceHint](#) logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

<code>os</code>	the stream used for this logger
<code>allocation_size_limit</code>	ignore allocations below this limit (bytes)
<code>copy_size_limit</code>	ignore copies below this limit (bytes)
<code>histogram_max_size</code>	how many allocation sizes and/or pointers to keep track of at most?

Returns

an `std::unique_ptr` to the the constructed object

```

64     {
65         return std::unique_ptr<PerformanceHint>(new PerformanceHint(
66             os, allocation_size_limit, copy_size_limit, histogram_max_size));
67     }

```

The documentation for this class was generated from the following file:

- `ginkgo/core/log/performance_hint.hpp`

47.180 `gko::Permutable< IndexType >` Class Template Reference

Linear operators which support permutation should implement the [Permutable](#) interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual `std::unique_ptr< LinOp >` [permute](#) (const `array< IndexType >` *permutation_indices) const
Returns a `LinOp` representing the symmetric row and column permutation of the [Permutable](#) object.
- virtual `std::unique_ptr< LinOp >` [inverse_permute](#) (const `array< IndexType >` *permutation_indices) const
Returns a `LinOp` representing the symmetric inverse row and column permutation of the [Permutable](#) object.
- virtual `std::unique_ptr< LinOp >` [row_permute](#) (const `array< IndexType >` *permutation_indices) const =0
Returns a `LinOp` representing the row permutation of the [Permutable](#) object.
- virtual `std::unique_ptr< LinOp >` [column_permute](#) (const `array< IndexType >` *permutation_indices) const =0
Returns a `LinOp` representing the column permutation of the [Permutable](#) object.
- virtual `std::unique_ptr< LinOp >` [inverse_row_permute](#) (const `array< IndexType >` *permutation_indices) const =0
Returns a `LinOp` representing the row permutation of the inverse permuted object.
- virtual `std::unique_ptr< LinOp >` [inverse_column_permute](#) (const `array< IndexType >` *permutation_↵indices) const =0
Returns a `LinOp` representing the row permutation of the inverse permuted object.

47.180.1 Detailed Description

```
template<typename IndexType>
class gko::Permutable< IndexType >
```

Linear operators which support permutation should implement the [Permutable](#) interface.

It provides functions to permute the rows and columns of a LinOp, independently or symmetrically, and with a regular or inverted permutation.

After a regular row permutation with permutation array `perm` the row `i` in the output LinOp contains the row `perm[i]` from the input LinOp. After an inverse row permutation, the row `perm[i]` in the output LinOp contains the row `i` from the input LinOp. Equivalently, after a column permutation, the output stores in column `i` the column `perm[i]` from the input, and an inverse column permutation stores in column `perm[i]` the column `i` from the input. A symmetric permutation is functionally equivalent to calling `as<Permutable>(A->row_permute(perm))->column_permute(perm)`, but the implementation can provide better performance due to kernel fusion.

47.180.1.1 Example: Permuting a Csr matrix:

```
{c++}
//Permuting an object of LinOp type.
//The object you want to permute.
auto op = matrix::Csr::create(exec);
//Permute the object by first converting it to a Permutable type.
auto perm = op->row_permute(permutation_indices);
```

47.180.2 Member Function Documentation

47.180.2.1 column_permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::column_permute (
    const array< IndexType > * permutation_indices ) const [pure virtual]
```

Returns a LinOp representing the column permutation of the [Permutable](#) object.

In the resulting LinOp, the column `i` contains the input column `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^T .

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order <code>perm</code> .
----------------------------------	---

Returns

a pointer to the new column permuted object

Implemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

47.180.2.2 inverse_column_permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::inverse_column_permute (
    const array< IndexType > * permutation_indices ) const [pure virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the column `perm[i]` contains the input column `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation AP^{-T} .

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order <code>perm</code> .
----------------------------------	---

Returns

a pointer to the new inverse permuted object

Implemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

47.180.2.3 inverse_permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::inverse_permute (
    const array< IndexType > * permutation_indices ) const [inline], [virtual]
```

Returns a LinOp representing the symmetric inverse row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location `(perm[i], perm[j])` contains the input value `(i, j)`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}AP^{-T}$.

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order.
----------------------------------	--

Returns

a pointer to the new permuted object

Reimplemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

47.180.2.4 inverse_row_permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::inverse_row_permute (
    const array< IndexType > * permutation_indices ) const [pure virtual]
```

Returns a LinOp representing the row permutation of the inverse permuted object.

In the resulting LinOp, the row `perm[i]` contains the input row `i`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation $P^{-1}A$.

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order <code>perm</code> .
----------------------------------	---

Returns

a pointer to the new inverse permuted object

Implemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

47.180.2.5 permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::permute (
    const array< IndexType > * permutation_indices ) const [inline], [virtual]
```

Returns a LinOp representing the symmetric row and column permutation of the [Permutable](#) object.

In the resulting LinOp, the entry at location `(i, j)` contains the input value `(perm[i], perm[j])`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PAP^T .

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order.
----------------------------------	--

Returns

a pointer to the new permuted object

Reimplemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

47.180.2.6 row_permute()

```
template<typename IndexType>
virtual std::unique_ptr<LinOp> gko::Permutable< IndexType >::row_permute (
    const array< IndexType > * permutation_indices ) const [pure virtual]
```

Returns a LinOp representing the row permutation of the [Permutable](#) object.

In the resulting LinOp, the row `i` contains the input row `perm[i]`.

From the linear algebra perspective, with $P_{ij} = \delta_{i\pi(j)}$, this represents the operation PA .

Parameters

<code>permutation_indices</code>	the array of indices containing the permutation order.
----------------------------------	--

Returns

a pointer to the new permuted object

Implemented in [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), and [gko::matrix::Csr< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/base/lin_op.hpp`

47.181 gko::matrix::Permutation< IndexType > Class Template Reference

[Permutation](#) is a matrix format that represents a permutation matrix, i.e.

```
#include <ginkgo/core/matrix/permutation.hpp>
```

Public Member Functions

- `index_type * get_permutation () noexcept`
Returns a pointer to the array of permutation.
- `const index_type * get_const_permutation () const noexcept`
Returns a pointer to the array of permutation.
- `size_type get_permutation_size () const noexcept`
Returns the number of elements explicitly stored in the permutation array.
- `std::unique_ptr< Permutation > compute_inverse () const`
Returns the inverse permutation.
- `std::unique_ptr< Permutation > compose (ptr_param< const Permutation > other) const`
Composes this permutation with another permutation.
- `void write (gko::matrix_data< value_type, index_type > &data) const override`
Writes a matrix to a [matrix_data](#) structure.

Static Public Member Functions

- static std::unique_ptr< [Permutation](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size=0)
Creates an uninitialized [Permutation](#) arrays on the specified executor.
- static std::unique_ptr< [Permutation](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, [array](#)< IndexType > permutation_indices)
Creates a [Permutation](#) matrix from an already allocated (and initialized) row and column permutation arrays.
- static std::unique_ptr< const [Permutation](#) > [create_const](#) (std::shared_ptr< const [Executor](#) > exec, [size_type](#) size, gko::detail::const_array_view< IndexType > &&perm_idx, mask_type enabled↵ permute=row_permute)
Creates a constant (immutable) [Permutation](#) matrix from a constant array.
- static std::unique_ptr< const [Permutation](#) > [create_const](#) (std::shared_ptr< const [Executor](#) > exec, gko↵::detail::const_array_view< IndexType > &&perm_idx)
Creates a constant (immutable) [Permutation](#) matrix from a constant array.

47.181.1 Detailed Description

```
template<typename IndexType = int32>
class gko::matrix::Permutation< IndexType >
```

[Permutation](#) is a matrix format that represents a permutation matrix, i.e.

a matrix where each row and column has exactly one entry. The matrix can only be applied to [Dense](#) inputs, where it represents a row permutation: $A' = PA$ means $A'(i, j) = A(p[i], j)$.

Template Parameters

<i>IndexType</i>	precision of permutation array indices.
------------------	---

47.181.2 Member Function Documentation

47.181.2.1 compose()

```
template<typename IndexType = int32>
std::unique_ptr<Permutation> gko::matrix::Permutation< IndexType >::compose (
    ptr\_param< const Permutation< IndexType > > other ) const
```

Composes this permutation with another permutation.

The resulting permutation fulfills `result[i] = this[other[i]]` or `result = other * this` from the matrix perspective, which is equivalent to first permuting by `this` and then by `other`: Combining permutations P_1 and P_2 with `P = P_1.combine(P_2)` performs the operation `permute(A, P) = permute(permute(A, P_1), P_2)`.

Parameters

<i>other</i>	the other permutation
--------------	-----------------------

Returns

the combined permutation

47.181.2.2 compute_inverse()

```
template<typename IndexType = int32>
std::unique_ptr<Permutation> gko::matrix::Permutation< IndexType >::compute_inverse ( ) const
```

Returns the inverse permutation.

Returns

a newly created [Permutation](#) object storing the inverse permutation of this [Permutation](#).

47.181.2.3 create() [1/2]

```
template<typename IndexType = int32>
static std::unique_ptr<Permutation> gko::matrix::Permutation< IndexType >::create (
    std::shared_ptr< const Executor > exec,
    array< IndexType > permutation_indices ) [static]
```

Creates a [Permutation](#) matrix from an already allocated (and initialized) row and column permutation arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the permutation array.
<i>permutation_indices</i>	array of permutation array
<i>enabled_permute</i>	mask for the type of permutation to apply.

Note

If `permutation_indices` is not an rvalue, not an array of `IndexType`, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.181.2.4 create() [2/2]

```
template<typename IndexType = int32>
static std::unique_ptr<Permutation> gko::matrix::Permutation< IndexType >::create (
    std::shared_ptr< const Executor > exec,
    size_type size = 0 ) [static]
```

Creates an uninitialized [Permutation](#) arrays on the specified executor.

Parameters

<i>exec</i>	Executor associated to the LinOp
-------------	--

Returns

A smart pointer to the newly created matrix.

47.181.2.5 create_const() [1/2]

```
template<typename IndexType = int32>
static std::unique_ptr<const Permutation> gko::matrix::Permutation< IndexType >::create_const
(
    std::shared_ptr< const Executor > exec,
    gko::detail::const_array_view< IndexType > && perm_idxes ) [static]
```

Creates a constant (immutable) [Permutation](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the size of the square matrix
<i>perm_idxes</i>	the permutation index array of the matrix
<i>enabled_permute</i>	the mask describing the type of permutation

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

47.181.2.6 create_const() [2/2]

```
template<typename IndexType = int32>
static std::unique_ptr<const Permutation> gko::matrix::Permutation< IndexType >::create_const
(
    std::shared_ptr< const Executor > exec,
```

```

size_type size,
gko::detail::const_array_view< IndexType > && perm_idx,
mask_type enabled_permute = row_permute ) [static]

```

Creates a constant (immutable) [Permutation](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the size of the square matrix
<i>perm_idx</i>	the permutation index array of the matrix
<i>enabled_permute</i>	the mask describing the type of permutation

Returns

A smart pointer to the constant matrix wrapping the input array (if it resides on the same executor as the matrix) or a copy of the array on the correct executor.

47.181.2.7 [get_const_permutation\(\)](#)

```

template<typename IndexType = int32>
const index_type* gko::matrix::Permutation< IndexType >::get_const_permutation ( ) const [inline],
[noexcept]

```

Returns a pointer to the array of permutation.

Returns

the pointer to the row permutation array.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References [gko::array< ValueType >::get_const_data\(\)](#).

47.181.2.8 [get_permutation\(\)](#)

```

template<typename IndexType = int32>
index_type* gko::matrix::Permutation< IndexType >::get_permutation ( ) [inline], [noexcept]

```

Returns a pointer to the array of permutation.

Returns

the pointer to the row permutation array.

References [gko::array< ValueType >::get_data\(\)](#).

47.181.2.9 get_permutation_size()

```
template<typename IndexType = int32>
size_type gko::matrix::Permutation< IndexType >::get_permutation_size ( ) const [noexcept]
```

Returns the number of elements explicitly stored in the permutation array.

Returns

the number of elements explicitly stored in the permutation array.

47.181.2.10 write()

```
template<typename IndexType = int32>
void gko::matrix::Permutation< IndexType >::write (
    gko::matrix_data< value_type, index_type > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< default_precision, IndexType >](#).

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/csr.hpp
- ginkgo/core/matrix/permutation.hpp

47.182 gko::matrix::Csr< ValueType, IndexType >::permuring_reuse_info Struct Reference

A struct describing a transformation of the matrix that reorders the values of the matrix into the transformed matrix.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [permuring_reuse_info](#) ()
Creates an empty reuse info.
- [permuring_reuse_info](#) (std::unique_ptr< [Permutation](#)< index_type >> value_permutation)
Creates a reuse info structure from its value permutation.
- void [update_values](#) (ptr_param< const Csr > input, ptr_param< Csr > output) const
Propagates the values from an input matrix to the transformed matrix.

47.182.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
struct gko::matrix::Csr< ValueType, IndexType >::permuting_reuse_info
```

A struct describing a transformation of the matrix that reorders the values of the matrix into the transformed matrix.

47.182.2 Member Function Documentation

47.182.2.1 update_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::permuting_reuse_info::update_values (
    ptr_param< const Csr > input,
    ptr_param< Csr > output ) const
```

Propagates the values from an input matrix to the transformed matrix.

The output matrix needs to have been computed using the transformation that was also used to generate this reuse data. Internally, this permutes the input value vector into the output value vector.

The documentation for this struct was generated from the following file:

- ginkgo/core/matrix/csr.hpp

47.183 gko::Perturbation< ValueType > Class Template Reference

The [Perturbation](#) class can be used to construct a LinOp to represent the operation (`identity + scalar * basis * projector`).

```
#include <ginkgo/core/base/perturbation.hpp>
```

Public Member Functions

- `const std::shared_ptr< const LinOp > get\_basis () const noexcept`
Returns the basis of the perturbation.
- `const std::shared_ptr< const LinOp > get\_projector () const noexcept`
Returns the projector of the perturbation.
- `const std::shared_ptr< const LinOp > get\_scalar () const noexcept`
Returns the scalar of the perturbation.

Static Public Member Functions

- static `std::unique_ptr< Perturbation > create` (`std::shared_ptr< const Executor > exec`)
Creates an empty perturbation operator (0x0 operator).
- static `std::unique_ptr< Perturbation > create` (`std::shared_ptr< const LinOp > scalar`, `std::shared_ptr< const LinOp > basis`)
Creates a perturbation with scalar and basis by setting projector to the conjugate transpose of basis.
- static `std::unique_ptr< Perturbation > create` (`std::shared_ptr< const LinOp > scalar`, `std::shared_ptr< const LinOp > basis`, `std::shared_ptr< const LinOp > projector`)
Creates a perturbation of scalar, basis and projector.

47.183.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::Perturbation< ValueType >
```

The [Perturbation](#) class can be used to construct a LinOp to represent the operation (`identity + scalar * basis * projector`).

This operator adds a movement along a direction constructed by `basis` and `projector` on the LinOp. `projector` gives the coefficient of `basis` to decide the direction.

For example, the Householder matrix can be represented with the [Perturbation](#) operator as follows. If `u` is the Householder factor then we can generate the [Householder transformation](#), $H = (I - 2 u u^*)$. In this case, the parameters of [Perturbation](#) class are `scalar = -2`, `basis = u`, and `projector = u*`.

Template Parameters

<i>ValueType</i>	precision of input and result vectors
------------------	---------------------------------------

Note

the apply operations of [Perturbation](#) class are not thread safe

47.183.2 Member Function Documentation

47.183.2.1 create() [1/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Perturbation> gko::Perturbation< ValueType >::create (
    std::shared_ptr< const Executor > exec ) [static]
```

Creates an empty perturbation operator (0x0 operator).

Parameters

<i>exec</i>	Executor associated to the perturbation
-------------	---

Returns

A smart pointer to the newly created perturbation.

47.183.2.2 create() [2/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Perturbation> gko::Perturbation< ValueType >::create (
    std::shared_ptr< const LinOp > scalar,
    std::shared_ptr< const LinOp > basis ) [static]
```

Creates a perturbation with scalar and basis by setting projector to the conjugate transpose of basis.

Basis must be transposable. [Perturbation](#) will throw [gko::NotSupported](#) if basis is not transposable.

Parameters

<i>scalar</i>	scaling of the movement
<i>basis</i>	the direction basis

Returns

A smart pointer to the newly created perturbation.

47.183.2.3 create() [3/3]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Perturbation> gko::Perturbation< ValueType >::create (
    std::shared_ptr< const LinOp > scalar,
    std::shared_ptr< const LinOp > basis,
    std::shared_ptr< const LinOp > projector ) [static]
```

Creates a perturbation of scalar, basis and projector.

Parameters

<i>scalar</i>	scaling of the movement
<i>basis</i>	the direction basis
<i>projector</i>	decides the coefficient of basis

Returns

A smart pointer to the newly created perturbation.

47.183.2.4 get_basis()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_basis ( ) const [inline],
[noexcept]
```

Returns the basis of the perturbation.

Returns

the basis of the perturbation

```
52     {
53         return basis_;
54     }
```

47.183.2.5 get_projector()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_projector ( ) const
[inline], [noexcept]
```

Returns the projector of the perturbation.

Returns

the projector of the perturbation

47.183.2.6 get_scalar()

```
template<typename ValueType = default_precision>
const std::shared_ptr<const LinOp> gko::Perturbation< ValueType >::get_scalar ( ) const [inline],
[noexcept]
```

Returns the scalar of the perturbation.

Returns

the scalar of the perturbation

The documentation for this class was generated from the following file:

- ginkgo/core/base/perturbation.hpp

47.184 gko::multigrid::Pgm< ValueType, IndexType > Class Template Reference

Parallel graph match ([Pgm](#)) is the aggregate method introduced in the paper M.

```
#include <ginkgo/core/multigrid/pgm.hpp>
```

Public Member Functions

- `std::shared_ptr< const LinOp > get\_system\_matrix () const`
Returns the system operator (matrix) of the linear system.
- `IndexType * get\_agg () noexcept`
Returns the aggregate group.
- `const IndexType * get\_const\_agg () const noexcept`
Returns the aggregate group.

Static Public Member Functions

- static `parameters_type parse (const config::pnode &config, const config::registry &context, const config::type\_descriptor &td_for_child=config::make_type_descriptor< ValueType, IndexType >())`
Create the parameters from the property_tree.

47.184.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::multigrid::Pgm< ValueType, IndexType >
```

Parallel graph match ([Pgm](#)) is the aggregate method introduced in the paper M.

Naumov et al., "AmgX: A Library for GPU Accelerated Algebraic Multigrid and Preconditioned Iterative Methods". Current implementation only contains size = 2 version.

[Pgm](#) creates the aggregate group according to the matrix value not the structure. [Pgm](#) gives two steps (one-phase handshaking) to group the elements. 1: get the strongest neighbor of each unaggregated element. 2: group the elements whose strongest neighbor is each other. repeating until reaching the given conditions. After that, the un-aggregated elements are assigned to an aggregated group or are left alone.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.184.2 Member Function Documentation

47.184.2.1 get_agg()

```
template<typename ValueType = default_precision, typename IndexType = int32>
IndexType* gko::multigrid::Pgm< ValueType, IndexType >::get_agg ( ) [inline], [noexcept]
```

Returns the aggregate group.

Aggregate group whose size is same as the number of rows. Stores the mapping information from row index to coarse row index. i.e., `agg[row_idx] = coarse_row_idx`.

Returns

the aggregate group.

```
79 { return agg_.get_data(); }
```

References `gko::array< ValueType >::get_data()`.

47.184.2.2 get_const_agg()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const IndexType* gko::multigrid::Pgm< ValueType, IndexType >::get_const_agg ( ) const [inline],
[noexcept]
```

Returns the aggregate group.

Aggregate group whose size is same as the number of rows. Stores the mapping information from row index to coarse row index. i.e., `agg[row_idx] = coarse_row_idx`.

Returns

the aggregate group.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.184.2.3 get_system_matrix()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const LinOp> gko::multigrid::Pgm< ValueType, IndexType >::get_system_matrix (
) const [inline]
```

Returns the system operator (matrix) of the linear system.

Returns

the system operator (matrix)

47.184.2.4 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::multigrid::Pgm< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >()
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

The documentation for this class was generated from the following files:

- ginkgo/core/distributed/matrix.hpp
- ginkgo/core/multigrid/pgm.hpp

47.185 gko::config::pnode Class Reference

pnode describes a tree of properties.

```
#include <ginkgo/core/config/property_tree.hpp>
```

Public Types

- enum [tag_t](#)

tag_t is the indicator for the current node storage.

Public Member Functions

- [pnode](#) ()
Default constructor: create an empty node.
- [pnode](#) (bool boolean)
Constructor for bool.
- [pnode](#) (T integer)
Constructor for integer with all integer type.
- [pnode](#) (const std::string &str)
Constructor for string.
- [pnode](#) (const char *str)
Constructor for char (otherwise, it will use bool)*
- [pnode](#) (double real)
Constructor for double (and also float)
- [pnode](#) (const array_type &array)
Constructor for array.
- [pnode](#) (const map_type &map)
Constructor for map.
- [operator bool](#) () const noexcept
bool conversion.
- [tag_t get_tag](#) () const
Get the current node tag.
- const array_type & [get_array](#) () const
Access the array stored in this property node.
- const map_type & [get_map](#) () const
Access the map stored in this property node.
- bool [get_boolean](#) () const
Access the boolean value stored in this property node.
- std::int64_t [get_integer](#) () const
- double [get_real](#) () const
Access the real floating point value stored in this property node.
- const std::string & [get_string](#) () const
Access the string stored in this property node.
- const [pnode](#) & [get](#) (const std::string &key) const
This function is to access the data under the map.
- const [pnode](#) & [get](#) (int index) const
This function is to access the data under the array.

47.185.1 Detailed Description

pnode describes a tree of properties.

A pnode can either be empty, hold a value (a string, integer, real, or bool), contain an array of pnode., or contain a mapping between strings and pnodes.

47.185.2 Constructor & Destructor Documentation

47.185.2.1 pnode() [1/7]

```
gko::config::pnode::pnode (
    bool boolean ) [explicit]
```

Constructor for bool.

Parameters

<i>boolean</i>	the bool type value
----------------	---------------------

47.185.2.2 pnode() [2/7]

```
template<typename T , std::enable_if_t< std::is_integral< T >::value > * >
gko::config::pnode::pnode (
    T integer ) [explicit]
```

Constructor for integer with all integer type.

Template Parameters

<i>T</i>	input type
----------	------------

Parameters

<i>integer</i>	the integer type value
----------------	------------------------

```
200             : tag_(tag_t::integer)
201 {
202     if (integer > std::numeric_limits<std::int64_t>::max() ||
203         (std::is_signed<T>::value &&
204          integer < std::numeric_limits<std::int64_t>::min())) {
205         throw std::runtime_error("The input is out of the range of int64_t.");
206     }
207     union_data_.integer_ = static_cast<std::int64_t>(integer);
208 }
```

47.185.2.3 pnode() [3/7]

```
gko::config::pnode::pnode (
    const std::string & str ) [explicit]
```

Constructor for string.

Parameters

<i>str</i>	string type value
------------	-------------------

47.185.2.4 pnode() [4/7]

```
gko::config::pnode::pnode (  
    const char * str ) [explicit]
```

Constructor for char* (otherwise, it will use bool)

Parameters

<i>str</i>	the string like "..."
------------	-----------------------

47.185.2.5 pnode() [5/7]

```
gko::config::pnode::pnode (  
    double real ) [explicit]
```

Constructor for double (and also float)

Parameters

<i>real</i>	the floating point type value
-------------	-------------------------------

47.185.2.6 pnode() [6/7]

```
gko::config::pnode::pnode (  
    const array_type & array ) [explicit]
```

Constructor for array.

Parameters

<i>array</i>	an pnode array
--------------	----------------

47.185.2.7 pnode() [7/7]

```
gko::config::pnode::pnode (  
    const map_type & map ) [explicit]
```

Constructor for map.

Parameters

<i>map</i>	a (string, pnode)-map
------------	-----------------------

47.185.3 Member Function Documentation

47.185.3.1 `get()` [1/2]

```
const pnode& gko::config::pnode::get (
    const std::string & key ) const
```

This function is to access the data under the map.

It will throw error when it does not hold a map. When access non-existent key in the map, it will return an empty node.

Parameters

<i>key</i>	the key for the node of the map
------------	---------------------------------

Returns

node. If the map does not have the key, return an empty node.

47.185.3.2 `get()` [2/2]

```
const pnode& gko::config::pnode::get (
    int index ) const
```

This function is to access the data under the array.

It will throw error when it does not hold an array or access out-of-bound index.

Parameters

<i>index</i>	the node index in array
--------------	-------------------------

Returns

node.

47.185.3.3 get_array()

```
const array_type& gko::config::pnode::get_array ( ) const
```

Access the array stored in this property node.

Throws [gko::InvalidStateError](#) if the property node does not store an array.

Returns

the array

47.185.3.4 get_boolean()

```
bool gko::config::pnode::get_boolean ( ) const
```

Access the boolean value stored in this property node.

Throws [gko::InvalidStateError](#) if the property node does not store a boolean value.

Returns

the boolean value

47.185.3.5 get_integer()

```
std::int64_t gko::config::pnode::get_integer ( ) const
```

- Access the integer value stored in this property node. Throws [gko::InvalidStateError](#) if the property node does not store an integer value.

Returns

the integer value

47.185.3.6 get_map()

```
const map_type& gko::config::pnode::get_map ( ) const
```

Access the map stored in this property node.

Throws [gko::InvalidStateError](#) if the property node does not store a map.

Returns

the map

47.185.3.7 `get_real()`

```
double gko::config::pnode::get_real ( ) const
```

Access the real floating point value stored in this property node.

Throws [gko::InvalidStateError](#) if the property node does not store a real value

Returns

the real floating point value

47.185.3.8 `get_string()`

```
const std::string& gko::config::pnode::get_string ( ) const
```

Access the string stored in this property node.

Throws [gko::InvalidStateError](#) if the property node does not store a string.

Returns

the string

47.185.3.9 `get_tag()`

```
tag_t gko::config::pnode::get_tag ( ) const
```

Get the current node tag.

Returns

the tag

47.185.3.10 `operator bool()`

```
gko::config::pnode::operator bool ( ) const [explicit], [noexcept]
```

bool conversion.

It's true if and only if it is not empty.

The documentation for this class was generated from the following file:

- `ginkgo/core/config/property_tree.hpp`

47.186 gko::log::polymorphic_object_data Struct Reference

Struct representing [PolymorphicObject](#) related data.

```
#include <ginkgo/core/log/record.hpp>
```

47.186.1 Detailed Description

Struct representing [PolymorphicObject](#) related data.

The documentation for this struct was generated from the following file:

- ginkgo/core/log/record.hpp

47.187 gko::PolymorphicObject Class Reference

A [PolymorphicObject](#) is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.

```
#include <ginkgo/core/base/polymorphic_object.hpp>
```

Public Member Functions

- `std::unique_ptr< PolymorphicObject > create_default (std::shared_ptr< const Executor > exec) const`
Creates a new "default" object of the same dynamic type as this object.
- `std::unique_ptr< PolymorphicObject > create_default () const`
Creates a new "default" object of the same dynamic type as this object.
- `std::unique_ptr< PolymorphicObject > clone (std::shared_ptr< const Executor > exec) const`
Creates a clone of the object.
- `std::unique_ptr< PolymorphicObject > clone () const`
Creates a clone of the object.
- `PolymorphicObject * copy_from (const PolymorphicObject *other)`
Copies another object into this object.
- `template<typename Derived , typename Deleter >
std::enable_if_t< std::is_base_of< PolymorphicObject, std::decay_t< Derived > >::value, PolymorphicObject > * copy_from (std::unique_ptr< Derived, Deleter > &&other)`
Moves another object into this object.
- `template<typename Derived , typename Deleter >
std::enable_if_t< std::is_base_of< PolymorphicObject, std::decay_t< Derived > >::value, PolymorphicObject > * copy_from (const std::unique_ptr< Derived, Deleter > &other)`
Copies another object into this object.
- `PolymorphicObject * copy_from (const std::shared_ptr< const PolymorphicObject > &other)`
Copies another object into this object.
- `PolymorphicObject * move_from (ptr_param< PolymorphicObject > other)`
Moves another object into this object.
- `PolymorphicObject * clear ()`
Transforms the object into its default state.
- `std::shared_ptr< const Executor > get_executor () const noexcept`
Returns the [Executor](#) of the object.

47.187.1 Detailed Description

A [PolymorphicObject](#) is the abstract base for all "heavy" objects in Ginkgo that behave polymorphically.

It defines the basic utilities (copying moving, cloning, clearing the objects) for all such objects. It takes into account that these objects are dynamically allocated, managed by smart pointers, and used polymorphically. Additionally, it assumes their data can be allocated on different executors, and that they can be copied between those executors.

Note

Most of the public methods of this class should not be overridden directly, and are thus not virtual. Instead, there are equivalent protected methods (ending in `<method_name>_impl`) that should be overridden instead. This allows polymorphic objects to implement default behavior around virtual methods (parameter checking, type casting).

See also

[EnablePolymorphicObject](#) if you wish to implement a concrete polymorphic object and have sensible defaults generated automatically. [EnableAbstractPolymorphicObject](#) if you wish to implement a new abstract polymorphic object, and have the return types of the methods updated to your type (instead of having them return [PolymorphicObject](#)).

47.187.2 Member Function Documentation

47.187.2.1 `clear()`

```
PolymorphicObject* gko::PolymorphicObject::clear ( ) [inline]
```

Transforms the object into its default state.

Equivalent to `this->copy_from(this->create\_default\(\))`.

See also

`clear_impl()` when implementing this method

Returns

this

47.187.2.2 clone() [1/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::clone ( ) const [inline]
```

Creates a clone of the object.

This is a shorthand for `clone(std::shared_ptr<const Executor>)` that uses the executor of this object to construct the new object.

Returns

A clone of the LinOp.

47.187.2.3 clone() [2/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::clone (
    std::shared_ptr< const Executor > exec ) const [inline]
```

Creates a clone of the object.

This is the polymorphic equivalent of the *executor copy constructor* `decltype(*this)(exec, this)`.

Parameters

<i>exec</i>	the executor where the clone will be created
-------------	--

Returns

A clone of the LinOp.

References `create_default()`.

47.187.2.4 copy_from() [1/4]

```
PolymorphicObject* gko::PolymorphicObject::copy_from (
    const PolymorphicObject * other ) [inline]
```

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

`copy_from_impl(const PolymorphicObject *)`

Parameters

<i>other</i>	the object to copy
--------------	--------------------

Returns

this

Referenced by `copy_from()`.

47.187.2.5 `copy_from()` [2/4]

```
PolymorphicObject* gko::PolymorphicObject::copy_from (
    const std::shared_ptr< const PolymorphicObject > & other ) [inline]
```

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

`copy_from_impl(const PolymorphicObject *)`

Parameters

<i>other</i>	the object to copy
--------------	--------------------

Returns

this

References `copy_from()`.

47.187.2.6 `copy_from()` [3/4]

```
template<typename Derived , typename Deleter >
std::enable_if_t< std::is_base_of<PolymorphicObject, std::decay_t<Derived> >::value, PolymorphicObject>*
gko::PolymorphicObject::copy_from (
    const std::unique_ptr< Derived, Deleter > & other ) [inline]
```

Copies another object into this object.

This is the polymorphic equivalent of the copy assignment operator.

See also

`copy_from_impl(const PolymorphicObject *)`

Parameters

<i>other</i>	the object to copy
--------------	--------------------

Returns

this

Template Parameters

<i>Derived</i>	the actual pointee type of the parameter, it needs to be derived from PolymorphicObject .
<i>Deleter</i>	the deleter of the <code>unique_ptr</code> parameter

References `copy_from()`.

47.187.2.7 `copy_from()` [4/4]

```
template<typename Derived , typename Deleter >
std::enable_if_t< std::is_base_of<PolymorphicObject, std::decay_t<Derived> >::value, PolymorphicObject>*
gko::PolymorphicObject::copy_from (
    std::unique_ptr< Derived, Deleter > && other ) [inline]
```

Moves another object into this object.

This is the polymorphic equivalent of the move assignment operator.

See also

`copy_from_impl(std::unique_ptr<PolymorphicObject>)`

Parameters

<i>other</i>	the object to move from
--------------	-------------------------

Returns

this

Template Parameters

<i>Derived</i>	the actual pointee type of the parameter, it needs to be derived from PolymorphicObject .
<i>Deleter</i>	the deleter of the <code>unique_ptr</code> parameter

47.187.2.8 create_default() [1/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::create_default ( ) const [inline]
```

Creates a new "default" object of the same dynamic type as this object.

This is a shorthand for `create_default(std::shared_ptr<const Executor>)` that uses the executor of this object to construct the new object.

Returns

a polymorphic object of the same type as this

Referenced by `clone()`.

47.187.2.9 create_default() [2/2]

```
std::unique_ptr<PolymorphicObject> gko::PolymorphicObject::create_default (
    std::shared_ptr< const Executor > exec ) const [inline]
```

Creates a new "default" object of the same dynamic type as this object.

This is the polymorphic equivalent of the *executor default constructor* `decltype(*this) (exec) ;`.

Parameters

<i>exec</i>	the executor where the object will be created
-------------	---

Returns

a polymorphic object of the same type as this

47.187.2.10 get_executor()

```
std::shared_ptr<const Executor> gko::PolymorphicObject::get_executor ( ) const [inline],
[noexcept]
```

Returns the [Executor](#) of the object.

Returns

[Executor](#) of the object

Referenced by `gko::preconditioner::lc< LSolverType, IndexType >::conj_transpose()`, `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::conj_transpose()`, `gko::preconditioner::lc< LSolverType, IndexType >::lc()`, `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::lc()`, `gko::matrix::Csr< ValueType, IndexType >::inv_scale()`, `gko::preconditioner::lc< LSolverType, IndexType >::operator=()`, `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::operator=()`, `gko::matrix::Csr< ValueType, IndexType >::scale()`, `gko::preconditioner::lc< LSolverType, IndexType >::transpose()`, and `gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >::transpose()`.

47.187.2.11 move_from()

```
PolymorphicObject* gko::PolymorphicObject::move_from (
    ptr_param< PolymorphicObject > other ) [inline]
```

Moves another object into this object.

This is the polymorphic equivalent of the move assignment operator.

See also

move_from_impl(PolymorphicObject *)

Parameters

<i>other</i>	the object to move from
--------------	-------------------------

Returns

this

References gko::ptr_param< T >::get().

The documentation for this class was generated from the following file:

- ginkgo/core/base/polymorphic_object.hpp

47.188 gko::precision_reduction Class Reference

This class is used to encode storage precisions of low precision algorithms.

```
#include <ginkgo/core/base/types.hpp>
```

Public Types

- using [storage_type](#) = uint8
The underlying datatype used to store the encoding.

Public Member Functions

- constexpr [precision_reduction](#) () noexcept
Creates a default [precision_reduction](#) encoding.
- constexpr [precision_reduction](#) ([storage_type](#) preserving, [storage_type](#) nonpreserving) noexcept
Creates a [precision_reduction](#) encoding with the specified number of conversions.
- constexpr [operator storage_type](#) () const noexcept
Extracts the raw data of the encoding.
- constexpr [storage_type get_preserving](#) () const noexcept
Returns the number of preserving conversions in the encoding.
- constexpr [storage_type get_nonpreserving](#) () const noexcept
Returns the number of non-preserving conversions in the encoding.

Static Public Member Functions

- constexpr static [precision_reduction autodetect](#) () noexcept
Returns a special encoding which instructs the algorithm to automatically detect the best precision.
- constexpr static [precision_reduction common](#) ([precision_reduction](#) x, [precision_reduction](#) y) noexcept
Returns the common encoding of input encodings.

47.188.1 Detailed Description

This class is used to encode storage precisions of low precision algorithms.

Some algorithms in Ginkgo can improve their performance by storing parts of the data in lower precision, while doing computation in full precision. This class is used to encode the precisions used to store the data. From the user's perspective, some algorithms can provide a parameter for fine-tuning the storage precision. Commonly, the special value returned by [precision_reduction::autodetect\(\)](#) should be used to allow the algorithm to automatically choose an appropriate value, though manually selected values can be used for fine-tuning.

In general, a lower precision floating point value can be obtained by either dropping some of the insignificant bits of the significand (keeping the same number of exponent bits, and thus preserving the range of representable values) or using one of the hardware or software supported conversions between IEEE formats, such as double to float or float to half (reducing both the number of exponent, as well as significand bits, and thus decreasing the range of representable values).

The [precision_reduction](#) class encodes the lower precision format relative to the base precision used and the algorithm in question. The encoding is done by specifying the amount of range non-preserving conversions and the amount of range preserving conversions that should be done on the base precision to obtain the lower precision format. For example, starting with a double precision value (11 exp, 52 sig. bits), the encoding specifying 1 non-preserving conversion and 1 preserving conversion would first use a hardware-supported non-preserving conversion to obtain a single precision value (8 exp, 23 sig. bits), followed by a preserving bit truncation to obtain a value with 8 exponent and 7 significand bits. Note that non-preserving conversion are always done first, as preserving conversions usually result in datatypes that are not supported by builtin conversions (thus, it is generally not possible to apply a non-preserving conversion to the result of a preserving conversion).

If the specified conversion is not supported by the algorithm, it will most likely fall back to using full precision for storing the data. Refer to the documentation of specific algorithms using this class for details about such special cases.

47.188.2 Constructor & Destructor Documentation

47.188.2.1 [precision_reduction\(\)](#) [1/2]

```
constexpr gko::precision_reduction::precision_reduction ( ) [inline], [constexpr], [noexcept]
```

Creates a default [precision_reduction](#) encoding.

This encoding represents the case where no conversions are performed.

Referenced by [common\(\)](#).

47.188.2.2 [precision_reduction\(\)](#) [2/2]

```
constexpr gko::precision_reduction::precision_reduction (
    storage_type preserving,
    storage_type nonpreserving ) [inline], [constexpr], [noexcept]
```

Creates a [precision_reduction](#) encoding with the specified number of conversions.

Parameters

<i>preserving</i>	the number of range preserving conversion
<i>nonpreserving</i>	the number of range non-preserving conversions

47.188.3 Member Function Documentation

47.188.3.1 autodetect()

```
constexpr static precision_reduction gko::precision_reduction::autodetect ( ) [inline], [static],
[constexpr], [noexcept]
```

Returns a special encoding which instructs the algorithm to automatically detect the best precision.

Returns

a special encoding instructing the algorithm to automatically detect the best precision.

47.188.3.2 common()

```
constexpr static precision_reduction gko::precision_reduction::common (
    precision_reduction x,
    precision_reduction y ) [inline], [static], [constexpr], [noexcept]
```

Returns the common encoding of input encodings.

The common encoding is defined as the encoding that does not have more preserving, nor non-preserving conversions than the input encodings.

Parameters

<i>x</i>	an encoding
<i>y</i>	an encoding

Returns

the common encoding of *x* and *y*

References `precision_reduction()`.

47.188.3.3 `get_nonpreserving()`

```
constexpr storage_type gko::precision_reduction::get_nonpreserving ( ) const [inline], [constexpr],  
[noexcept]
```

Returns the number of non-preserving conversions in the encoding.

Returns

the number of non-preserving conversions in the encoding.

47.188.3.4 `get_preserving()`

```
constexpr storage_type gko::precision_reduction::get_preserving ( ) const [inline], [constexpr],  
[noexcept]
```

Returns the number of preserving conversions in the encoding.

Returns

the number of preserving conversions in the encoding.

47.188.3.5 `operator storage_type()`

```
constexpr gko::precision_reduction::operator storage_type ( ) const [inline], [constexpr],  
[noexcept]
```

Extracts the raw data of the encoding.

Returns

the raw data of the encoding

The documentation for this class was generated from the following file:

- `ginkgo/core/base/types.hpp`

47.189 `gko::Preconditionable` Class Reference

A `LinOp` implementing this interface can be preconditioned.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual `std::shared_ptr< const LinOp >` [get_preconditioner](#) () const
Returns the preconditioner operator used by the [Preconditionable](#).
- virtual void [set_preconditioner](#) (std::shared_ptr< const LinOp > new_precond)
Sets the preconditioner operator used by the [Preconditionable](#).

47.189.1 Detailed Description

A LinOp implementing this interface can be preconditioned.

47.189.2 Member Function Documentation

47.189.2.1 [get_preconditioner](#)()

```
virtual std::shared_ptr<const LinOp> gko::Preconditionable::get_preconditioner ( ) const
[inline], [virtual]
```

Returns the preconditioner operator used by the [Preconditionable](#).

Returns

the preconditioner operator used by the [Preconditionable](#)

Referenced by `gko::solver::EnablePreconditionable< Bicg< ValueType > >::operator=()`.

47.189.2.2 [set_preconditioner](#)()

```
virtual void gko::Preconditionable::set_preconditioner (
    std::shared_ptr< const LinOp > new_precond ) [inline], [virtual]
```

Sets the preconditioner operator used by the [Preconditionable](#).

Parameters

<i>new_precond</i>	the new preconditioner operator used by the Preconditionable
--------------------	--

Reimplemented in `gko::solver::EnablePreconditionable< DerivedType >`, `gko::solver::EnablePreconditionable< Bicgstab< ValueType > >`, `gko::solver::EnablePreconditionable< Gmres< ValueType > >`, `gko::solver::EnablePreconditionable< CbGmres< ValueType > >`, `gko::solver::EnablePreconditionable< Fcg< ValueType > >`, `gko::solver::EnablePreconditionable< Ldr< ValueType > >`, `gko::solver::EnablePreconditionable< Cgs< ValueType > >`, `gko::solver::EnablePreconditionable< Gcr< ValueType > >`, `gko::solver::EnablePreconditionable< Cg< ValueType > >`, and `gko::solver::EnablePreconditionable< Bicg< ValueType > >`.

Referenced by `gko::solver::EnablePreconditionable< Bicg< ValueType > >::set_preconditioner()`.

The documentation for this class was generated from the following file:

- ginkgo/core/base/lin_op.hpp

47.190 gko::log::ProfilerHook Class Reference

This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Classes

- class [NestedSummaryWriter](#)
Receives the results from [ProfilerHook::create_nested_summary\(\)](#).
- class [SummaryWriter](#)
Receives the results from [ProfilerHook::create_summary\(\)](#).
- class [TableSummaryWriter](#)
Writes the results from [ProfilerHook::create_summary\(\)](#) and [ProfilerHook::create_nested_summary\(\)](#) to a ASCII table in Markdown format.

Public Member Functions

- void [set_object_name](#) (ptr_param< const [PolymorphicObject](#) > obj, std::string name)
Sets the name for an object to be profiled.
- void [set_synchronization](#) (bool synchronize)
Should the events call executor->synchronize on operations and copy/allocation? This leads to a certain overhead, but makes the execution timeline of kernels synchronous.
- [profiling_scope_guard](#) user_range (const char *name) const
Creates a scope guard for a user-defined range to be included in the profile.

Static Public Member Functions

- static std::shared_ptr< [ProfilerHook](#) > [create_tau](#) (bool initialize=true)
Creates a logger annotating Ginkgo events with TAU ranges via PerfStubs.
- static std::shared_ptr< [ProfilerHook](#) > [create_ytune](#) ()
Creates a logger annotating Ginkgo events with VTune ITT ranges.
- static std::shared_ptr< [ProfilerHook](#) > [create_nvtx](#) (uint32 color_argb=color_yellow_argb)
Creates a logger annotating Ginkgo events with NVTX ranges for CUDA.
- static std::shared_ptr< [ProfilerHook](#) > [create_roctx](#) ()
Creates a logger annotating Ginkgo events with ROCTX ranges for HIP.
- static std::shared_ptr< [ProfilerHook](#) > [create_for_executor](#) (std::shared_ptr< const [Executor](#) > exec)
Creates a logger annotating Ginkgo events with the most suitable backend for the given executor: NVTX for NSight Systems in CUDA, ROCTX for rocprof in HIP, TAU for everything else.
- static std::shared_ptr< [ProfilerHook](#) > [create_summary](#) (std::shared_ptr< [Timer](#) > timer=std::make_shared< [CpuTimer](#) >(), std::unique_ptr< [SummaryWriter](#) > writer=std::make_unique< [TableSummaryWriter](#) >(), bool debug_check_nesting=false)
Creates a logger measuring the runtime of Ginkgo events and printing a summary when it is destroyed.
- static std::shared_ptr< [ProfilerHook](#) > [create_nested_summary](#) (std::shared_ptr< [Timer](#) > timer=std::make_shared< [CpuTimer](#) >(), std::unique_ptr< [NestedSummaryWriter](#) > writer=std::make_unique< [TableSummaryWriter](#) >(), bool debug_check_nesting=false)
Creates a logger measuring the runtime of Ginkgo events in a nested fashion and printing a summary when it is destroyed.
- static std::shared_ptr< [ProfilerHook](#) > [create_custom](#) (hook_function begin, hook_function end)
Creates a logger annotating Ginkgo events with a custom set of functions for range begin and end.

Static Public Attributes

- constexpr static [uint32](#) [color_yellow_argb](#) = 0xFFFFCB05U
The Ginkgo yellow background color as packed 32 bit ARGB value.

47.190.1 Detailed Description

This Logger can be used to annotate the execution of Ginkgo functionality with profiler-specific ranges.

It currently supports TAU, VTune, NSightSystems (NVTX) and rocPROF(ROCTX) and custom profiler hooks.

The Logger should be attached to the [Executor](#) that is being used to run the application for a full, program-wide annotation, or to individual objects to only highlight events caused directly by them (not operations and memory allocations though)

47.190.2 Member Function Documentation

47.190.2.1 create_nested_summary()

```
static std::shared_ptr<ProfilerHook> gko::log::ProfilerHook::create_nested_summary (
    std::shared_ptr< Timer > timer = std::make_shared< CpuTimer > (),
    std::unique_ptr< NestedSummaryWriter > writer = std::make_unique< TableSummaryWriter > (),
    bool debug_check_nesting = false ) [static]
```

Creates a logger measuring the runtime of Ginkgo events in a nested fashion and printing a summary when it is destroyed.

Parameters

<i>timer</i>	The timer used to record time points.
<i>writer</i>	The NestedSummaryWriter to receive the performance results.
<i>debug_check_nesting</i>	Enable this flag if the output looks like it might contain incorrect nesting. This increases the overhead slightly, but recognizes mismatching push/pop pairs on the range stack.

Note

For this logger to provide reliable GPU timings, either use [Timer::create_for_executor](#) or enable synchronization via `set_synchronization(true)`.

47.190.2.2 create_nvtx()

```
static std::shared_ptr<ProfilerHook> gko::log::ProfilerHook::create_nvtx (
    uint32 color_argb = color_yellow_argb ) [static]
```

Creates a logger annotating Ginkgo events with NVTX ranges for CUDA.

Parameters

<i>color_argb</i>	The color of the NVTX ranges in the NSight Systems output. It has to be a 32 bit packed ARGB value.
-------------------	---

47.190.2.3 create_summary()

```
static std::shared_ptr<ProfilerHook> gko::log::ProfilerHook::create_summary (
    std::shared_ptr< Timer > timer = std::make_shared< CpuTimer > (),
    std::unique_ptr< SummaryWriter > writer = std::make_unique< TableSummaryWriter > (),
    bool debug_check_nesting = false ) [static]
```

Creates a logger measuring the runtime of Ginkgo events and printing a summary when it is destroyed.

Parameters

<i>timer</i>	The timer used to record time points.
<i>writer</i>	The SummaryWriter to receive the performance results.
<i>debug_check_nesting</i>	Enable this flag if the output looks like it might contain incorrect nesting. This increases the overhead slightly, but recognizes mismatching push/pop pairs on the range stack.

Note

For this logger to provide reliable GPU timings, either use [Timer::create_for_executor](#) or enable synchronization via `set_synchronization(true)`.

47.190.2.4 create_tau()

```
static std::shared_ptr<ProfilerHook> gko::log::ProfilerHook::create_tau (
    bool initialize = true ) [static]
```

Creates a logger annotating Ginkgo events with TAU ranges via PerfStubs.

Parameters

<i>initialize</i>	Should we call TAU's initialization and finalization functions, or does the application take care of it? The initialization will happen immediately, the finalization at program exit.
-------------------	--

47.190.2.5 set_object_name()

```
void gko::log::ProfilerHook::set_object_name (
    ptr_param< const PolymorphicObject > obj,
    std::string name )
```

Sets the name for an object to be profiled.

Every instance of that object in the profile will be replaced by the name instead of its runtime type.

Parameters

<i>obj</i>	the object
<i>name</i>	its name

47.190.2.6 user_range()

```
profiling_scope_guard gko::log::ProfilerHook::user_range (
    const char * name ) const
```

Creates a scope guard for a user-defined range to be included in the profile.

Parameters

<i>name</i>	the name of the range
-------------	-----------------------

Returns

the scope guard. It will begin a range immediately and end it at the end of its scope.

The documentation for this class was generated from the following file:

- ginkgo/core/log/profiler_hook.hpp

47.191 gko::log::profiling_scope_guard Class Reference

Scope guard that annotates its scope with the provided profiler hooks.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

- [profiling_scope_guard](#) ()
Creates an empty (moved-from) scope guard.
- [profiling_scope_guard](#) (const char *name, [profile_event_category](#) category, ProfilerHook::hook_function begin, ProfilerHook::hook_function end)
Creates the scope guard.
- [~profiling_scope_guard](#) ()
Calls the range end function if the scope guard was not moved from.

47.191.1 Detailed Description

Scope guard that annotates its scope with the provided profiler hooks.

47.191.2 Constructor & Destructor Documentation

47.191.2.1 profiling_scope_guard()

```
gko::log::profiling_scope_guard::profiling_scope_guard (
    const char * name,
    profile_event_category category,
    ProfilerHook::hook_function begin,
    ProfilerHook::hook_function end )
```

Creates the scope guard.

Parameters

<i>name</i>	the name of the profiler range
<i>category</i>	the category of the profiler range
<i>begin</i>	the hook function to begin a range
<i>end</i>	the hook function to end a range

The documentation for this class was generated from the following file:

- ginkgo/core/log/profiler_hook.hpp

47.192 gko::ptr_param< T > Class Template Reference

This class is used for function parameters in the place of raw pointers.

```
#include <ginkgo/core/base/utils_helper.hpp>
```

Public Member Functions

- [ptr_param](#) (T *ptr)
Initializes the [ptr_param](#) from a raw pointer.
- `template<typename U, std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr>`
[ptr_param](#) (const std::shared_ptr< U > &ptr)
Initializes the [ptr_param](#) from a shared_ptr.
- `template<typename U, typename Deleter, std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr>`
[ptr_param](#) (const std::unique_ptr< U, Deleter > &ptr)
Initializes the [ptr_param](#) from a unique_ptr.

- `template<typename U, std::enable_if_t< std::is_base_of< T, U >::value > * = nullptr>`
`ptr_param (const ptr_param< U > &ptr)`
Initializes the `ptr_param` from a `ptr_param` of a derived type.
- `T & operator* () const`
- `T * operator-> () const`
- `T * get () const`
- `operator bool () const`

47.192.1 Detailed Description

```
template<typename T>
class gko::ptr_param< T >
```

This class is used for function parameters in the place of raw pointers.

Pointer parameters should be used for everything that does not involve transfer of ownership. It can be converted to from raw pointers, shared pointers and unique pointers of the specified type or any derived type. This allows functions to be called without having to use `gko::lend` or calling `.get()` for every pointer argument. It probably has no use outside of function parameters, as it is immutable.

Template Parameters

<code>T</code>	the pointed-to type
----------------	---------------------

47.192.2 Member Function Documentation

47.192.2.1 get()

```
template<typename T>
T* gko::ptr_param< T >::get ( ) const [inline]
```

Returns

the underlying pointer.

Referenced by `gko::as()`, `gko::Executor::copy_from()`, `gko::matrix::Diagonal< ValueType >::inverse_apply()`, `gko::PolymorphicObject::move_from()`, `gko::ptr_param< T >::ptr_param()`, and `gko::matrix::Diagonal< ValueType >::rapply()`.

47.192.2.2 operator bool()

```
template<typename T>
gko::ptr_param< T >::operator bool ( ) const [inline], [explicit]
```

Returns

true iff the underlying pointer is non-null.

47.192.2.3 operator*()

```
template<typename T>
T& gko::ptr_param< T >::operator* ( ) const [inline]
```

Returns

a reference to the underlying pointee.

47.192.2.4 operator->()

```
template<typename T>
T* gko::ptr_param< T >::operator-> ( ) const [inline]
```

Returns

the underlying pointer.

The documentation for this class was generated from the following file:

- ginkgo/core/base/utils_helper.hpp

47.193 gko::range< Accessor > Class Template Reference

A range is a multidimensional view of the memory.

```
#include <ginkgo/core/base/range.hpp>
```

Public Types

- using `accessor` = Accessor
The type of the underlying accessor.

Public Member Functions

- `~range()`=default
Use the default destructor.
- template<typename... AccessorParams, typename = std::enable_if_t< sizeof...(AccessorParams) != 1 || !std::is_same< range, std::decay<detail::head_t<AccessorParams...>>>::value>>>
constexpr `range` (AccessorParams &&... params)
Creates a new range.
- template<typename... DimensionTypes>
constexpr auto `operator()` (DimensionTypes &&... dimensions) const -> decltype(std::declval< accessor >())(std::forward< DimensionTypes >(dimensions)...)
Returns a value (or a sub-range) with the specified indexes.
- template<typename OtherAccessor >
const `range` & `operator=` (const `range`< OtherAccessor > &other) const
const `range` & `operator=` (const `range` &other) const
Assigns another range to this range.
- constexpr `size_type length` (`size_type` dimension) const
Returns the length of the specified dimension of the range.
- constexpr const `accessor` * `operator->` () const noexcept
Returns a pointer to the accessor.
- constexpr const `accessor` & `get_accessor` () const noexcept
Returns a reference to the accessor.

Static Public Attributes

- static constexpr [size_type dimensionality](#) = accessor::dimensionality
The number of dimensions of the range.

47.193.1 Detailed Description

```
template<typename Accessor>
class gko::range< Accessor >
```

A range is a multidimensional view of the memory.

The range does not store any of its values by itself. Instead, it obtains the values through an accessor (e.g. [accessor::row_major](#)) which describes how the indexes of the range map to physical locations in memory.

There are several advantages of using ranges instead of plain memory pointers:

1. Code using ranges is easier to read and write, as there is no need for index linearizations.
2. Code using ranges is safer, as it is impossible to accidentally miscalculate an index or step out of bounds, since range accessors perform bounds checking in debug builds. For performance, this can be disabled in release builds by defining the `NDEBUG` flag.
3. Ranges enable generalized code, as algorithms can be written independent of the memory layout. This does not impede various optimizations based on memory layout, as it is always possible to specialize algorithms for ranges with specific memory layouts.
4. Ranges have various pointwise operations predefined, which reduces the amount of loops that need to be written.

47.193.1.1 Range operations

Ranges define a complete set of pointwise unary and binary operators which extend the basic arithmetic operators in C++, as well as a few pointwise operations and mathematical functions useful in ginkgo, and a couple of non-pointwise operations. Compound assignment (`+=`, `*=`, etc.) is not yet supported at this moment. Here is a complete list of operations:

- standard unary operations: `+`, `-`, `!`, `~`
- standard binary operations: `+`, `*` (this is pointwise, not matrix multiplication), `/`, `%`, `<`, `>`, `<=`, `>=`, `==`, `!=`, `||`, `&&`, `|`, `&`, `^`, `<<`, `>>`
- useful unary functions: `zero`, `one`, `abs`, `real`, `imag`, `conj`, `squared_norm`
- useful binary functions: `min`, `max`

All binary pointwise operations also work as expected if one of the operands is a scalar and the other is a range. The scalar operand will have the effect as if it was a range of the same size as the other operand, filled with the specified scalar.

Two "global" functions `transpose` and `mmul` are also supported. `transpose` transposes the first two dimensions of the range (i.e. `transpose(r)(i, j, ...) == r(j, i, ...)`). `mmul` performs a (batched) matrix multiply of the ranges - the first two dimensions represent the matrices, while the rest represent the batch. For example, given the ranges `r1` and `r2` of dimensions `(3, 2, 3)` and `(2, 4, 3)`, respectively, `mmul(r1, r2)` will return a range of dimensions `(3, 4, 3)`, obtained by multiplying the 3 frontal slices of the range, and stacking the result back vertically.

47.193.1.2 Compound operations

Multiple range operations can be combined into a single expression. For example, an "axpy" operation can be obtained using `y = alpha * x + y`, where `x` and `y` are ranges, and `alpha` is a scalar. Range operations are optimized for memory access, and the above code does not allocate additional storage for intermediate ranges `alpha * x` or `alpha * x + y`. In fact, the entire computation is done during the assignment, and the results of operations `+` and `*` only register the data, and the types of operations that will be computed once the results are needed.

It is possible to store and reuse these intermediate expressions. The following example will overwrite the range `x` with it's 4th power:

```
{c++}
auto square = x * x; // this is range constructor, not range assignment!
x = square; // overwrites x with x*x (this is range assignment)
x = square; // overwrites new x (x*x) with (x*x)*(x*x) (as is this)
```

47.193.1.3 Caveats

`__mmul` is not a highly-optimized BLAS-3 version of the matrix multiplication. The current design of ranges and accessors prevents that, so if you need a high-performance matrix multiplication, you should use one of the libraries that provide that, or implement your own (you can use pointwise range operations to help simplify that). However, range design might get improved in the future to allow efficient implementations of BLAS-3 kernels.

Aliasing the result range in `mmul` and `transpose` is not allowed. Constructs like `A = transpose(A)`, `A = mmul(A, A)`, or `A = mmul(A, A) + C` lead to undefined behavior. However, aliasing input arguments is allowed: `C = mmul(A, A)`, and even `C = mmul(A, A) + C` is valid code (in the last example, only pointwise operations are aliased). `C = mmul(A, A + C)` is not valid though.

47.193.1.4 Examples

The range unit tests in `core/test/base/range.cpp` contain lots of simple 1-line examples of range operations. The accessor unit tests in `core/test/base/range.cpp` show how to use ranges with concrete accessors, and how to use range slices using `spans` as arguments to range function call operator. Finally, `examples/range` contains a complete example where ranges are used to implement a simple version of the right-looking LU factorization.

Template Parameters

<i>Accessor</i>	underlying accessor of the range
-----------------	----------------------------------

47.193.2 Constructor & Destructor Documentation

47.193.2.1 range()

```
template<typename Accessor>
template<typename... AccessorParams, typename = std::enable_if_t< sizeof...(AccessorParams)
!= 1 || !std::is_same< range, std::decay<detail::head_t<AccessorParams...>>::value>>
constexpr gko::range< Accessor >::range (
    AccessorParams &&... params ) [inline], [explicit], [constexpr]
```

Creates a new range.

Template Parameters

<i>AccessorParam</i>	types of parameters forwarded to the accessor constructor
----------------------	---

Parameters

<i>params</i>	parameters forwarded to Accessor constructor.
---------------	---

```

329         : accessor_{std::forward<AccessorParams>(params)...}
330     {}

```

47.193.3 Member Function Documentation

47.193.3.1 get_accessor()

```

template<typename Accessor>
constexpr const accessor& gko::range< Accessor >::get_accessor ( ) const [inline], [constexpr],
[noexcept]

```

`Returns a reference to the accessor.

Returns

reference to the accessor

Referenced by gko::range< Accessor >::operator=().

47.193.3.2 length()

```

template<typename Accessor>
constexpr size_type gko::range< Accessor >::length (
    size_type dimension ) const [inline], [constexpr]

```

Returns the length of the specified dimension of the range.

Parameters

<i>dimension</i>	the dimensions whose length is returned
------------------	---

Returns

the length of the *dimension*-th dimension of the range

Referenced by gko::matrix_data< ValueType, IndexType >::matrix_data().

47.193.3.3 operator>()()

```
template<typename Accessor>
template<typename... DimensionTypes>
constexpr auto gko::range< Accessor >::operator() (
    DimensionTypes &&... dimensions ) const -> decltype(std::declval<accessor>()) (
    std::forward<DimensionTypes>(dimensions)...))    [inline], [constexpr]
```

Returns a value (or a sub-range) with the specified indexes.

Template Parameters

<i>DimensionTypes</i>	The types of indexes. Supported types depend on the underlying accessor, but are usually either integer types or spans. If at least one index is a span, the returned value will be a sub-range.
-----------------------	--

Parameters

<i>dimensions</i>	the indexes of the values.
-------------------	----------------------------

Returns

a value on position (*dimensions...*).

References `gko::range< Accessor >::dimensionality`.

47.193.3.4 operator->()

```
template<typename Accessor>
constexpr const accessor* gko::range< Accessor >::operator-> ( ) const    [inline], [constexpr],
[noexcept]
```

Returns a pointer to the accessor.

Can be used to access data and functions of a specific accessor.

Returns

pointer to the accessor

47.193.3.5 operator=() [1/2]

```
template<typename Accessor>
const range& gko::range< Accessor >::operator= (
    const range< Accessor > & other ) const    [inline]
```

Assigns another range to this range.

The order of assignment is defined by the accessor of this range, thus the memory access will be optimized for the resulting range, and not for the other range. If the sizes of two ranges do not match, the result is undefined. Sizes of the ranges are checked at runtime in debug builds.

Note

Temporary accessors are allowed to define the implementation of the assignment as deleted, so do not expect `r1 * r2 = r2` to work.

Parameters

<i>other</i>	the range to copy the data from
--------------	---------------------------------

References gko::range< Accessor >::get_accessor().

47.193.3.6 operator=() [2/2]

```
template<typename Accessor>
template<typename OtherAccessor >
const range& gko::range< Accessor >::operator= (
    const range< OtherAccessor > & other ) const [inline]
```

This is a version of the function which allows to copy between ranges of different accessors.

Template Parameters

<i>OtherAccessor</i>	accessor of the other range
----------------------	-----------------------------

The documentation for this class was generated from the following file:

- ginkgo/core/base/range.hpp

47.194 gko::syn::range< Start, End, Step > Struct Template Reference

range records start, end, step in template

```
#include <ginkgo/core/synthesizer/containers.hpp>
```

47.194.1 Detailed Description

```
template<int Start, int End, int Step = 1>
struct gko::syn::range< Start, End, Step >
```

range records start, end, step in template

Template Parameters

<i>Start</i>	start of range
<i>End</i>	end of range
<i>Step</i>	step of range. default is 1

The documentation for this struct was generated from the following file:

- `ginkgo/core/synthesizer/containers.hpp`

47.195 `gko::experimental::reorder::Rcm< IndexType >` Class Template Reference

`Rcm` (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.

```
#include <ginkgo/core/reorder/rcm.hpp>
```

Public Member Functions

- `const parameters_type & get_parameters ()`
Returns the parameters used to construct the factory.
- `std::unique_ptr< permutation_type > generate (std::shared_ptr< const LinOp > system_matrix) const`

Static Public Member Functions

- `static parameters_type build ()`
Creates a new parameter_type to set up the factory.

47.195.1 Detailed Description

```
template<typename IndexType = int32>
class gko::experimental::reorder::Rcm< IndexType >
```

`Rcm` (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.

Such a reordering typically also significantly reduces fill-in, though usually not as effective as more complex algorithms, specifically AMD and nested dissection schemes. The advantage of this algorithm is its low runtime.

The class is a `LinOpFactory` generating a Permutation matrix out of a Csr system matrix, to be used with `Csr<...>::permute(...)`.

There are two "starting strategies" currently available: minimum degree and pseudo-peripheral. These strategies control how a starting vertex for a connected component is chosen, which is then renumbered as first vertex in the component, starting the algorithm from there. In general, the bandwidths obtained by choosing a pseudo-peripheral vertex are slightly smaller than those obtained from choosing a vertex of minimum degree. On the other hand, this strategy is much more expensive, relatively. The algorithm for finding a pseudo-peripheral vertex as described in "Computer Solution of Sparse Linear Systems" (George, Liu, Ng, Oak Ridge National Laboratory, 1994) is implemented here.

Template Parameters

<i>IndexType</i>	Type of the indices of all matrices used in this class
------------------	--

47.195.2 Member Function Documentation

47.195.2.1 generate()

```
template<typename IndexType = int32>
std::unique_ptr<permutation_type> gko::experimental::reorder::Rcm< IndexType >::generate (
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a Permutation instead of a generic LinOp, which would need to be cast to Permutation again to access its indices. It is only necessary because smart pointers aren't covariant.

47.195.2.2 get_parameters()

```
template<typename IndexType = int32>
const parameters_type& gko::experimental::reorder::Rcm< IndexType >::get_parameters ( ) [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
205 { return parameters_; }
```

The documentation for this class was generated from the following file:

- ginkgo/core/reorder/rcm.hpp

47.196 gko::reorder::Rcm< ValueType, IndexType > Class Template Reference

Rcm (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.

```
#include <ginkgo/core/reorder/rcm.hpp>
```

Public Member Functions

- std::shared_ptr< const [PermutationMatrix](#) > [get_permutation](#) () const
Gets the permutation (permutation matrix, output of the algorithm) of the linear operator.
- std::shared_ptr< const [PermutationMatrix](#) > [get_inverse_permutation](#) () const
Gets the inverse permutation (permutation matrix, output of the algorithm) of the linear operator.

47.196.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::reorder::Rcm< ValueType, IndexType >
```

[Rcm](#) (Reverse Cuthill-McKee) is a reordering algorithm minimizing the bandwidth of a matrix.

Such a reordering typically also significantly reduces fill-in, though usually not as effective as more complex algorithms, specifically AMD and nested dissection schemes. The advantage of this algorithm is its low runtime. It requires the input matrix to be structurally symmetric.

Note

This class is derived from polymorphic object but is not a LinOp as it does not make sense for this class to implement the apply methods. The objective of this class is to generate a reordering/permutation vector (in the form of the Permutation matrix), which can be used to apply to reorder a matrix as required.

There are two "starting strategies" currently available: minimum degree and pseudo-peripheral. These strategies control how a starting vertex for a connected component is chosen, which is then renumbered as first vertex in the component, starting the algorithm from there. In general, the bandwidths obtained by choosing a pseudo-peripheral vertex are slightly smaller than those obtained from choosing a vertex of minimum degree. On the other hand, this strategy is much more expensive, relatively. The algorithm for finding a pseudo-peripheral vertex as described in "Computer Solution of Sparse Linear Systems" (George, Liu, Ng, Oak Ridge National Laboratory, 1994) is implemented here.

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

47.196.2 Member Function Documentation

47.196.2.1 `get_inverse_permutation()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const PermutationMatrix> gko::reorder::Rcm< ValueType, IndexType >::get_↵
inverse_permutation ( ) const [inline]
```

Gets the inverse permutation (permutation matrix, output of the algorithm) of the linear operator.

Returns

the inverse permutation (permutation matrix)

47.196.2.2 get_permutation()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<const PermutationMatrix> gko::reorder::Rcm< ValueType, IndexType >::get_↵
permutation ( ) const [inline]
```

Gets the permutation (permutation matrix, output of the algorithm) of the linear operator.

Returns

the permutation (permutation matrix)

The documentation for this class was generated from the following file:

- ginkgo/core/reorder/rcm.hpp

47.197 gko::ReadableFromMatrixData< ValueType, IndexType > Class Template Reference

A LinOp implementing this interface can read its data from a [matrix_data](#) structure.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual void [read](#) (const [matrix_data](#)< ValueType, IndexType > &data)=0
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [matrix_assembly_data](#)< ValueType, IndexType > &data)
Reads a matrix from a [matrix_assembly_data](#) structure.
- virtual void [read](#) (const [device_matrix_data](#)< ValueType, IndexType > &data)
Reads a matrix from a [device_matrix_data](#) structure.
- virtual void [read](#) ([device_matrix_data](#)< ValueType, IndexType > &&data)
Reads a matrix from a [device_matrix_data](#) structure.

47.197.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::ReadableFromMatrixData< ValueType, IndexType >
```

A LinOp implementing this interface can read its data from a [matrix_data](#) structure.

47.197.2 Member Function Documentation

47.197.2.1 read() [1/4]

```
template<typename ValueType, typename IndexType>
virtual void gko::ReadableFromMatrixData< ValueType, IndexType >::read (
    const device\_matrix\_data< ValueType, IndexType > & data ) [inline], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Ell< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), [gko::matrix::Sellp< ValueType, IndexType >](#), and [gko::matrix::SparsityCsr< ValueType, IndexType >](#).

47.197.2.2 read() [2/4]

```
template<typename ValueType, typename IndexType>
void gko::ReadableFromMatrixData< ValueType, IndexType >::read (
    const matrix_assembly_data< ValueType, IndexType > & data ) [inline]
```

Reads a matrix from a [matrix_assembly_data](#) structure.

Parameters

<i>data</i>	the matrix_assembly_data structure
-------------	--

47.197.2.3 read() [3/4]

```
template<typename ValueType, typename IndexType>
virtual void gko::ReadableFromMatrixData< ValueType, IndexType >::read (
    const matrix_data< ValueType, IndexType > & data ) [pure virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Ell< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), [gko::matrix::Sellp< ValueType, IndexType >](#), and [gko::matrix::SparsityCsr< ValueType, IndexType >](#).

47.197.2.4 read() [4/4]

```
template<typename ValueType, typename IndexType>
virtual void gko::ReadableFromMatrixData< ValueType, IndexType >::read (
    device_matrix_data< ValueType, IndexType > && data ) [inline], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Ell< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), [gko::matrix::Sellp< ValueType, IndexType >](#), and [gko::matrix::SparsityCsr< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/lin_op.hpp](#)

47.198 gko::log::Record Class Reference

[Record](#) is a Logger which logs every event to an object.

```
#include <ginkgo/core/log/record.hpp>
```

Classes

- struct [logged_data](#)
Struct storing the actually logged data.

Public Member Functions

- const [logged_data](#) & [get](#) () const noexcept
Returns the logged data.
- [logged_data](#) & [get](#) () noexcept

Static Public Member Functions

- static std::unique_ptr< [Record](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const mask_type &enabled_events=Logger::all_events_mask, [size_type](#) max_storage=1)
Creates a [Record](#) logger.
- static std::unique_ptr< [Record](#) > [create](#) (const mask_type &enabled_events=Logger::all_events_mask, [size_type](#) max_storage=1)
Creates a [Record](#) logger.

47.198.1 Detailed Description

[Record](#) is a Logger which logs every event to an object.

The object can then be accessed at any time by asking the logger to return it.

Note

Please note that this logger can have significant memory and performance overhead. In particular, when logging events such as the `check` events, all parameters are cloned. If it is sufficient to clone one parameter, consider implementing a specific logger for this. In addition, it is advised to tune the history size in order to control memory overhead.

47.198.2 Member Function Documentation

47.198.2.1 create() [1/2]

```
static std::unique_ptr<Record> gko::log::Record::create (
    const mask_type & enabled_events = Logger::all_events_mask,
    size_type max_storage = 1 ) [inline], [static]
```

Creates a [Record](#) logger.

This dynamically allocates the memory, constructs the object and returns an `std::unique_ptr` to this object.

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.
<i>max_storage</i>	the size of storage (i.e. history) wanted by the user. By default 0 is used, which means unlimited storage. It is advised to control this to reduce memory overhead of this logger.

Returns

an `std::unique_ptr` to the the constructed object

```
433 {
434     return std::unique_ptr<Record>(new Record(enabled_events, max_storage));
435 }
```

47.198.2.2 create() [2/2]

```
static std::unique_ptr<Record> gko::log::Record::create (
    std::shared_ptr< const Executor > exec,
    const mask_type & enabled_events = Logger::all_events_mask,
    size_type max_storage = 1 ) [inline], [static]
```

Creates a [Record](#) logger.

This dynamically allocates the memory, constructs the object and returns an `std::unique_ptr` to this object.

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.
<i>max_storage</i>	the size of storage (i.e. history) wanted by the user. By default 0 is used, which means unlimited storage. It is advised to control this to reduce memory overhead of this logger.

Returns

an std::unique_ptr to the the constructed object

47.198.2.3 get() [1/2]

```
const logged_data& gko::log::Record::get ( ) const [inline], [noexcept]
```

Returns the logged data.

Returns

the logged data

47.198.2.4 get() [2/2]

```
logged_data& gko::log::Record::get ( ) [inline], [noexcept]
```

The documentation for this class was generated from the following file:

- ginkgo/core/log/record.hpp

47.199 gko::ReferenceExecutor Class Reference

This is a specialization of the [OmpExecutor](#), which runs the reference implementations of the kernels used for debugging purposes.

```
#include <ginkgo/core/base/executor.hpp>
```

Public Member Functions

- std::string [get_description](#) () const override
- void [run](#) (const [Operation](#) &op) const override
Runs the specified [Operation](#) using this [Executor](#).
- virtual void [run](#) (const [Operation](#) &op) const=0
Runs the specified [Operation](#) using this [Executor](#).
- template<typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename ClosureDpcpp >
void [run](#) (const ClosureOmp &op_omp, const ClosureCuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const
Runs one of the passed in functors, depending on the [Executor](#) type.
- template<typename ClosureReference, typename ClosureOmp, typename ClosureCuda, typename ClosureHip, typename Closure↔Dpcpp >
void [run](#) (std::string name, const ClosureReference &op_ref, const ClosureOmp &op_omp, const Closure↔Cuda &op_cuda, const ClosureHip &op_hip, const ClosureDpcpp &op_dpcpp) const
Runs one of the passed in functors, depending on the [Executor](#) type.

Additional Inherited Members

47.199.1 Detailed Description

This is a specialization of the [OmpExecutor](#), which runs the reference implementations of the kernels used for debugging purposes.

47.199.2 Member Function Documentation

47.199.2.1 `get_description()`

```
std::string gko::ReferenceExecutor::get_description ( ) const [inline], [override], [virtual]
```

Returns

a textual representation of the executor and its device.

Reimplemented from [gko::OmpExecutor](#).

47.199.2.2 `run()` [1/4]

```
template<typename ClosureOmp , typename ClosureCuda , typename ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureOmp</i>	type of <code>op_omp</code>
<i>ClosureCuda</i>	type of <code>op_cuda</code>
<i>ClosureHip</i>	type of <code>op_hip</code>
<i>ClosureDpcpp</i>	type of <code>op_dpcpp</code>

Parameters

<i>op_omp</i>	functor to run in case of a OmpExecutor or ReferenceExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

47.199.2.3 run() [2/4]

```
void gko::ReferenceExecutor::run (
    const Operation & op ) const [inline], [override], [virtual]
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

Implements [gko::Executor](#).

47.199.2.4 run() [3/4]

```
virtual void gko::Executor::run
```

Runs the specified [Operation](#) using this [Executor](#).

Parameters

<i>op</i>	the operation to run
-----------	----------------------

47.199.2.5 run() [4/4]

```
template<typename ClosureReference , typename ClosureOmp , typename ClosureCuda , typename
ClosureHip , typename ClosureDpcpp >
void gko::Executor::run (
    typename ClosureReference ,
    typename ClosureOmp ,
    typename ClosureCuda ,
    typename ClosureHip ,
    typename ClosureDpcpp ) [inline]
```

Runs one of the passed in functors, depending on the [Executor](#) type.

Template Parameters

<i>ClosureReference</i>	type of op_ref
<i>ClosureOmp</i>	type of op_omp
<i>ClosureCuda</i>	type of op_cuda
<i>ClosureHip</i>	type of op_hip
<i>ClosureDpcpp</i>	type of op_dpcpp

Parameters

<i>name</i>	the name of the operation
<i>op_ref</i>	functor to run in case of a ReferenceExecutor
<i>op_omp</i>	functor to run in case of a OmpExecutor
<i>op_cuda</i>	functor to run in case of a CudaExecutor
<i>op_hip</i>	functor to run in case of a HipExecutor
<i>op_dpcpp</i>	functor to run in case of a DpcppExecutor

The documentation for this class was generated from the following file:

- ginkgo/core/base/executor.hpp

47.200 gko::config::registry Class Reference

This class stores additional context for creating Ginkgo objects from configuration files.

```
#include <ginkgo/core/config/registry.hpp>
```

Public Member Functions

- [registry](#) (const configuration_map &additional_map={})
registry constructor
- [registry](#) (const std::unordered_map< std::string, detail::allowed_ptr > &stored_map, const configuration_map &additional_map={})
registry constructor
- template<typename T >
bool [emplace](#) (std::string key, std::shared_ptr< T > data)
Store the data with the key.

47.200.1 Detailed Description

This class stores additional context for creating Ginkgo objects from configuration files.

The context can contain user-provided objects that derive from the following base types:

- LinOp
- [LinOpFactory](#)
- CriterionFactory

Additionally, users can provide mappings from a configuration (provided as a pnode) to user-defined types that are derived from [LinOpFactory](#)

47.200.2 Constructor & Destructor Documentation

47.200.2.1 registry() [1/2]

```
gko::config::registry::registry (
    const configuration_map & additional_map = {} )
```

registry constructor

Parameters

<i>additional_map</i>	the additional map to dispatch the class base. Users can extend the map to fit their own LinOpFactory . We suggest using "usr::" as the prefix in the key to simply avoid conflict with ginkgo's map.
-----------------------	---

47.200.2.2 registry() [2/2]

```
gko::config::registry::registry (
    const std::unordered_map< std::string, detail::allowed_ptr > & stored_map,
    const configuration_map & additional_map = {} )
```

registry constructor

Parameters

<i>stored_map</i>	the map stores the shared pointer of users' objects. It can hold any type whose base type is LinOp/LinOpFactory/CriterionFactory.
<i>additional_map</i>	the additional map to dispatch the class base. Users can extend the map to fit their own LinOpFactory . We suggest using "usr::" as the prefix in the key to simply avoid conflict with ginkgo's map.

47.200.3 Member Function Documentation**47.200.3.1 emplace()**

```
template<typename T >
bool gko::config::registry::emplace (
    std::string key,
    std::shared_ptr< T > data ) [inline]
```

Store the data with the key.

Template Parameters

<i>T</i>	the type
----------	----------

Parameters

<i>key</i>	the unique key string
<i>data</i>	the shared pointer of the object

```
236 {
237     auto it = stored_map.emplace(key, data);
```

```
238     return it.second;
239 }
```

The documentation for this class was generated from the following file:

- ginkgo/core/config/registry.hpp

47.201 gko::stop::RelativeResidualNorm< ValueType > Class Template Reference

The [RelativeResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.201.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::stop::RelativeResidualNorm< ValueType >
```

The [RelativeResidualNorm](#) class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the right-hand side, i.e.

when $\text{norm}(\text{residual}) / \text{norm}(\text{right_hand_side}) < \text{threshold}$. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on `b` in order to compute the norm of the right-hand side. If this is not correctly provided, an exception `::gko::NotSupported()` is thrown.

The documentation for this class was generated from the following file:

- ginkgo/core/stop/residual_norm.hpp

47.202 gko::reorder::ReorderingBase< IndexType > Class Template Reference

The [ReorderingBase](#) class is a base class for all the reordering algorithms.

```
#include <ginkgo/core/reorder/reordering_base.hpp>
```


Additional Inherited Members

47.202.1 Detailed Description

```
template<typename IndexType = int32>
class gko::reorder::ReorderingBase< IndexType >
```

The [ReorderingBase](#) class is a base class for all the reordering algorithms.

It contains a factory to instantiate the reorderings. It is up to each specific reordering to decide what to do with the data that is passed to it.

The documentation for this class was generated from the following file:

- ginkgo/core/reorder/reordering_base.hpp

47.203 gko::reorder::ReorderingBaseArgs Struct Reference

This struct is used to pass parameters to the `EnableDefaultReorderingBaseFactory::generate()` method.

```
#include <ginkgo/core/reorder/reordering_base.hpp>
```

47.203.1 Detailed Description

This struct is used to pass parameters to the `EnableDefaultReorderingBaseFactory::generate()` method.

It is the `ComponentsType` of `ReorderingBaseFactory`.

The documentation for this struct was generated from the following file:

- ginkgo/core/reorder/reordering_base.hpp

47.204 gko::experimental::mpi::request Class Reference

The request class is a light, move-only wrapper around the `MPI_Request` handle.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

- [request](#) ()
The default constructor.
- `MPI_Request *` [get](#) ()
Get a pointer to the underlying `MPI_Request` handle.
- [status wait](#) ()
Allows a rank to wait on a particular request handle.

47.204.1 Detailed Description

The request class is a light, move-only wrapper around the MPI_Request handle.

47.204.2 Constructor & Destructor Documentation

47.204.2.1 request()

```
gko::experimental::mpi::request::request ( ) [inline]
```

The default constructor.

It creates a null MPI_Request of MPI_REQUEST_NULL type.

47.204.3 Member Function Documentation

47.204.3.1 get()

```
MPI_Request* gko::experimental::mpi::request::get ( ) [inline]
```

Get a pointer to the underlying MPI_Request handle.

Returns

a pointer to MPI_Request handle

Referenced by gko::experimental::mpi::communicator::i_all_gather(), gko::experimental::mpi::communicator::i_all_reduce(), gko::experimental::mpi::communicator::i_all_to_all(), gko::experimental::mpi::communicator::i_all_to_all_v(), gko::experimental::mpi::communicator::i_broadcast(), gko::experimental::mpi::communicator::i_gather(), gko::experimental::mpi::communicator::i_gather_v(), gko::experimental::mpi::communicator::i_recv(), gko::experimental::mpi::communicator::i_reduce(), gko::experimental::mpi::communicator::i_scan(), gko::experimental::mpi::communicator::i_scatter(), gko::experimental::mpi::communicator::i_scatter_v(), gko::experimental::mpi::communicator::i_send(), gko::experimental::mpi::window< ValueType >::r_accumulate(), gko::experimental::mpi::window< ValueType >::r_get(), gko::experimental::mpi::window< ValueType >::r_get_accumulate(), and gko::experimental::mpi::window< ValueType >::r_put().

47.204.3.2 wait()

```
status gko::experimental::mpi::request::wait ( ) [inline]
```

Allows a rank to wait on a particular request handle.

Parameters

<i>req</i>	The request to wait on.
<i>status</i>	The status variable that can be queried.

References gko::experimental::mpi::status::get().

The documentation for this class was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.205 gko::stop::ResidualNorm< ValueType > Class Template Reference

The [ResidualNorm](#) class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.205.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::stop::ResidualNorm< ValueType >
```

The [ResidualNorm](#) class is a stopping criterion which stops the iteration process when the actual residual norm is below a certain threshold relative to.

1. the norm of the right-hand side, $\text{norm}(\text{residual}) \leq \text{threshold} * \text{norm}(\text{right_hand_side})$.
2. the initial residual, $\text{norm}(\text{residual}) \leq \text{threshold} * \text{norm}(\text{initial_residual})$.
3. one, $\text{norm}(\text{residual}) \leq \text{threshold}$.

For better performance, the checks are run on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on either `b` or the `initial_residual` in order to compute their norms. If this is not correctly provided, an exception `gko::NotSupported()` is thrown.

The documentation for this class was generated from the following file:

- ginkgo/core/stop/residual_norm.hpp

47.206 `gko::stop::ResidualNormBase< ValueType >` Class Template Reference

The `ResidualNormBase` class provides a framework for stopping criteria related to the residual norm.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.206.1 Detailed Description

```
template<typename ValueType>
class gko::stop::ResidualNormBase< ValueType >
```

The `ResidualNormBase` class provides a framework for stopping criteria related to the residual norm.

These criteria differ in the way they initialize `starting_tau_`, so in the value they compare the residual norm against. The provided `check_impl` uses the actual residual to check for convergence.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/residual_norm.hpp`

47.207 `gko::stop::ResidualNormReduction< ValueType >` Class Template Reference

The `ResidualNormReduction` class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

```
#include <ginkgo/core/stop/residual_norm.hpp>
```

47.207.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::stop::ResidualNormReduction< ValueType >
```

The `ResidualNormReduction` class is a stopping criterion which stops the iteration process when the residual norm is below a certain threshold relative to the norm of the initial residual, i.e.

when $\text{norm}(\text{residual}) / \text{norm}(\text{initial_residual}) < \text{threshold}$. For better performance, the checks are run thanks to kernels on the executor where the algorithm is executed.

Note

To use this stopping criterion there are some dependencies. The constructor depends on `initial_residual` in order to compute the first relative residual norm. The check method depends on either the `residual_norm` or the `residual` being set. When any of those is not correctly provided, an exception `::gko::NotSupported()` is thrown.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/residual_norm.hpp`

47.208 gko::accessor::row_major< ValueType, Dimensionality > Class Template Reference

A `row_major` accessor is a bridge between a range and the row-major memory layout.

```
#include <ginkgo/core/base/range_accessors.hpp>
```

Public Types

- using `value_type` = `ValueType`
Type of values returned by the accessor.
- using `data_type` = `value_type *`
Type of underlying data storage.

Public Member Functions

- constexpr `value_type & operator()` (`size_type` row, `size_type` col) const
Returns the data element at position (row, col)
- constexpr `range< row_major > operator()` (const `span` &rows, const `span` &cols) const
Returns the sub-range spanning the range (rows, cols)
- constexpr `size_type length` (`size_type` dimension) const
*Returns the length in dimension *dimension*.*
- template<typename OtherAccessor >
void `copy_from` (const OtherAccessor &other) const
Copies data from another accessor.

Public Attributes

- const `data_type data`
Reference to the underlying data.
- const `std::array< const size_type, dimensionality > lengths`
An array of dimension sizes.
- const `size_type stride`
Distance between consecutive rows.

Static Public Attributes

- static constexpr `size_type dimensionality` = 2
Number of dimensions of the accessor.

47.208.1 Detailed Description

```
template<typename ValueType, size_type Dimensionality>
class gko::accessor::row_major< ValueType, Dimensionality >
```

A `row_major` accessor is a bridge between a range and the row-major memory layout.

You should never try to explicitly create an instance of this accessor. Instead, supply it as a template parameter to a range, and pass the constructor parameters for this class to the range (it will forward it to this class).

Warning

The current implementation is incomplete, and only allows for 2-dimensional ranges.

Template Parameters

<i>ValueType</i>	type of values this accessor returns
<i>Dimensionality</i>	number of dimensions of this accessor (has to be 2)

47.208.2 Member Function Documentation

47.208.2.1 copy_from()

```
template<typename ValueType , size_type Dimensionality>
template<typename OtherAccessor >
void gko::accessor::row_major< ValueType, Dimensionality >::copy_from (
    const OtherAccessor & other ) const [inline]
```

Copies data from another accessor.

Warning

Do not use this function since it is not optimized for a specific executor. It will always be performed sequentially. Please write an optimized version (adjusted to the architecture) by iterating through the values yourself.

Template Parameters

<i>OtherAccessor</i>	type of the other accessor
----------------------	----------------------------

Parameters

<i>other</i>	other accessor
--------------	----------------

```
141     {
142         for (size_type i = 0; i < lengths[0]; ++i) {
143             for (size_type j = 0; j < lengths[1]; ++j) {
144                 (*this)(i, j) = other(i, j);
145             }
146         }
147     }
```

References `gko::accessor::row_major< ValueType, Dimensionality >::lengths`.

47.208.2.2 length()

```
template<typename ValueType , size_type Dimensionality>
constexpr size_type gko::accessor::row_major< ValueType, Dimensionality >::length (
    size_type dimension ) const [inline], [constexpr]
```

Returns the length in dimension `dimension`.

Parameters

<i>dimension</i>	a dimension index
------------------	-------------------

Returns

length in dimension *dimension*

References gko::accessor::row_major< ValueType, Dimensionality >::lengths.

47.208.2.3 operator>() [1/2]

```
template<typename ValueType , size_type Dimensionality>
constexpr range<row_major> gko::accessor::row_major< ValueType, Dimensionality >::operator()
(
    const span & rows,
    const span & cols ) const [inline], [constexpr]
```

Returns the sub-range spanning the range (rows, cols)

Parameters

<i>rows</i>	row span
<i>cols</i>	column span

Returns

sub-range spanning the range (rows, cols)

References gko::span::begin, gko::accessor::row_major< ValueType, Dimensionality >::data, gko::span::end, gko::span::is_valid(), gko::accessor::row_major< ValueType, Dimensionality >::lengths, and gko::accessor::row↵_major< ValueType, Dimensionality >::stride.

47.208.2.4 operator>() [2/2]

```
template<typename ValueType , size_type Dimensionality>
constexpr value_type& gko::accessor::row_major< ValueType, Dimensionality >::operator() (
    size_type row,
    size_type col ) const [inline], [constexpr]
```

Returns the data element at position (row, col)

Parameters

<i>row</i>	row index
<i>col</i>	column index

Returns

data element at (row, col)

References `gko::accessor::row_major< ValueType, Dimensionality >::data`, `gko::accessor::row_major< ValueType, Dimensionality >::lengths`, and `gko::accessor::row_major< ValueType, Dimensionality >::stride`.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/range_accessors.hpp`

47.209 `gko::matrix::RowGatherer< IndexType >` Class Template Reference

`RowGatherer` is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.

```
#include <ginkgo/core/matrix/row_gatherer.hpp>
```

Public Member Functions

- `index_type * get_row_idxes ()` noexcept
Returns a pointer to the row index array for gathering.
- `const index_type * get_const_row_idxes ()` const noexcept
Returns a pointer to the row index array for gathering.

Static Public Member Functions

- `static std::unique_ptr< RowGatherer > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size={})`
Creates uninitialized `RowGatherer` arrays of the specified size.
- `static std::unique_ptr< RowGatherer > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, array< index_type > row_idxes)`
Creates a `RowGatherer` matrix from an already allocated (and initialized) row gathering array.
- `static std::unique_ptr< const RowGatherer > create_const (std::shared_ptr< const Executor > exec, const dim< 2 > &size, gko::detail::const_array_view< IndexType > &&row_idxes)`
Creates a constant (immutable) `RowGatherer` matrix from a constant array.

47.209.1 Detailed Description

```
template<typename IndexType = int32>
class gko::matrix::RowGatherer< IndexType >
```

`RowGatherer` is a matrix "format" which stores the gather indices arrays which can be used to gather rows to another matrix.

Template Parameters

<i>IndexType</i>	precision of rowgatherer array indices.
------------------	---

Note

This format is used mainly to allow for an abstraction of the rowgatherer and provides the user with an apply method which calls the respective [Dense](#) rowgatherer operation. As such it only stores an array of the rowgatherer indices.

47.209.2 Member Function Documentation

47.209.2.1 create() [1/2]

```
template<typename IndexType = int32>
static std::unique_ptr<RowGatherer> gko::matrix::RowGatherer< IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< index_type > row_idxes ) [static]
```

Creates a [RowGatherer](#) matrix from an already allocated (and initialized) row gathering array.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the rowgatherer array.
<i>row_idxes</i>	array of rowgatherer array

Note

If *row_idxes* is not an rvalue, not an array of *IndexType*, or is on the wrong executor, an internal copy will be created, and the original array data will not be used in the matrix.

Returns

A smart pointer to the newly created matrix.

47.209.2.2 create() [2/2]

```
template<typename IndexType = int32>
static std::unique_ptr<RowGatherer> gko::matrix::RowGatherer< IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = {} ) [static]
```

Creates uninitialized [RowGatherer](#) arrays of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the RowGatherable matrix

Returns

A smart pointer to the newly created matrix.

47.209.2.3 create_const()

```
template<typename IndexType = int32>
static std::unique_ptr<const RowGatherer> gko::matrix::RowGatherer< IndexType >::create_const
(
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    gko::detail::const\_array\_view< IndexType > && row_idxes ) [static]
```

Creates a constant (immutable) [RowGatherer](#) matrix from a constant array.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>row_idxes</i>	the gathered row indices of the matrix

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of the arrays on the correct executor.

47.209.2.4 get_const_row_idxes()

```
template<typename IndexType = int32>
const index_type* gko::matrix::RowGatherer< IndexType >::get_const_row_idxes ( ) const [inline],
[noexcept]
```

Returns a pointer to the row index array for gathering.

Returns

the pointer to the row index array for gathering.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

```

64     {
65         return row_idx_.get_const_data();
66     }

```

References gko::array< ValueType >::get_const_data().

47.209.2.5 get_row_idxes()

```

template<typename IndexType = int32>
index_type* gko::matrix::RowGatherer< IndexType >::get_row_idxes ( ) [inline], [noexcept]

```

Returns a pointer to the row index array for gathering.

Returns

the pointer to the row index array for gathering.

References gko::array< ValueType >::get_data().

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/row_gatherer.hpp

47.210 gko::ScaledIdentityAddable Class Reference

Adds the operation $M \leftarrow aI + bM$ for matrix M , identity operator I and scalars a and b , where M is the calling object.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- void [add_scaled_identity](#) (ptr_param< const LinOp > const a, ptr_param< const LinOp > const b)
Scales this and adds another scalar times the identity to it.

47.210.1 Detailed Description

Adds the operation $M \leftarrow aI + bM$ for matrix M , identity operator I and scalars a and b , where M is the calling object.

47.210.2 Member Function Documentation**47.210.2.1 add_scaled_identity()**

```

void gko::ScaledIdentityAddable::add_scaled_identity (
    ptr_param< const LinOp > const a,
    ptr_param< const LinOp > const b ) [inline]

```

Scales this and adds another scalar times the identity to it.

Parameters

<i>a</i>	Scalar to multiply the identity operator before adding.
<i>b</i>	Scalar to multiply this before adding the scaled identity to it.

References `gko::make_temporary_clone()`.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/lin_op.hpp`

47.211 `gko::matrix::ScaledPermutation< ValueType, IndexType >` Class Template Reference

[`ScaledPermutation`](#) is a matrix combining a permutation with scaling factors.

```
#include <ginkgo/core/matrix/scaled_permutation.hpp>
```

Public Member Functions

- `value_type * get\_scaling\_factors ()` noexcept
Returns a pointer to the scaling factors.
- `const value_type * get\_const\_scaling\_factors ()` const noexcept
Returns a pointer to the scaling factors.
- `index_type * get\_permutation ()` noexcept
Returns a pointer to the permutation indices.
- `const index_type * get\_const\_permutation ()` const noexcept
Returns a pointer to the permutation indices.
- `std::unique_ptr< ScaledPermutation > compute\_inverse ()` const
Returns the inverse of this operator as a scaled permutation.
- `std::unique_ptr< ScaledPermutation > compose (ptr_param< const ScaledPermutation > other)` const
Composes this scaled permutation with another scaled permutation.
- `void write (gko::matrix_data< value_type, index_type > &data)` const override
Writes a matrix to a [matrix_data](#) structure.

Static Public Member Functions

- `static std::unique_ptr< ScaledPermutation > create (std::shared_ptr< const Executor > exec, size_type size=0)`
Creates an uninitialized [ScaledPermutation](#) matrix.
- `static std::unique_ptr< ScaledPermutation > create (ptr_param< const Permutation< IndexType >> permutation)`
Create a [ScaledPermutation](#) from a [Permutation](#).
- `static std::unique_ptr< ScaledPermutation > create (std::shared_ptr< const Executor > exec, array< value_type > scaling_factors, array< index_type > permutation_indices)`
Creates a [ScaledPermutation](#) matrix from already allocated arrays.
- `static std::unique_ptr< const ScaledPermutation > create\_const (std::shared_ptr< const Executor > exec, gko::detail::const_array_view< value_type > &&scale, gko::detail::const_array_view< index_type > &&perm_idx)`
Creates a constant (immutable) [ScaledPermutation](#) matrix from constant arrays.

47.211.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::ScaledPermutation< ValueType, IndexType >
```

[ScaledPermutation](#) is a matrix combining a permutation with scaling factors.

It is a combination of [Diagonal](#) and [Permutation](#), and can be read as $SP = P \cdot S$, i.e. the scaling gets applied before the permutation.

Template Parameters

<i>IndexType</i>	index type of permutation indices
<i>ValueType</i>	value type of the scaling factors

47.211.2 Member Function Documentation

47.211.2.1 compose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >↔
::compose (
    ptr_param< const ScaledPermutation< ValueType, IndexType > > other ) const
```

Composes this scaled permutation with another scaled permutation.

This means $\text{result} = \text{other} * \text{this}$ from the matrix perspective, which is equivalent to first scaling and permuting by *this* and then by *other*.

Parameters

<i>other</i>	the other permutation
--------------	-----------------------

Returns

the combined permutation

47.211.2.2 compute_inverse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >↔
::compute_inverse ( ) const
```

Returns the inverse of this operator as a scaled permutation.

It is computed via $(PS)^{-1} = P^{-1}(PSP^{-1})$.

Returns

a newly created [ScaledPermutation](#) object storing the inverse of the permutation and scaling factors of this [ScaledPermutation](#).

47.211.2.3 create() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >::create (
    ptr\_param< const Permutation< IndexType >> permutation ) [static]
```

Create a [ScaledPermutation](#) from a [Permutation](#).

The permutation will be copied, the scaling factors are all set to 1.0.

Parameters

<i>permutation</i>	the permutation
--------------------	-----------------

Returns

A smart pointer to the newly created matrix.

47.211.2.4 create() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    array< value_type > scaling_factors,
    array< index_type > permutation_indices ) [static]
```

Creates a [ScaledPermutation](#) matrix from already allocated arrays.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>permutation_indices</i>	array of permutation indices
<i>scaling_factors</i>	array of scaling factors

Returns

A smart pointer to the newly created matrix.

47.211.2.5 create() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    size_type size = 0 ) [static]
```

Creates an uninitialized [ScaledPermutation](#) matrix.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	dimensions of the (square) scaled permutation matrix

Returns

A smart pointer to the newly created matrix.

47.211.2.6 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const ScaledPermutation> gko::matrix::ScaledPermutation< ValueType, IndexType >::create_const (
    std::shared_ptr< const Executor > exec,
    gko::detail::const_array_view< value_type > && scale,
    gko::detail::const_array_view< index_type > && perm_idxs ) [static]
```

Creates a constant (immutable) [ScaledPermutation](#) matrix from constant arrays.

Parameters

<i>exec</i>	the executor to create the object on
<i>perm_idxs</i>	the permutation index array of the matrix
<i>scale</i>	the scaling factor array

Returns

A smart pointer to the constant matrix wrapping the input arrays (if it resides on the same executor as the matrix) or a copy of the arrays on the correct executor.

47.211.2.7 get_const_permutation()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::ScaledPermutation< ValueType, IndexType >::get_const_permutation
( ) const [inline], [noexcept]
```

Returns a pointer to the permutation indices.

Returns

the pointer to the permutation indices.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

```

79     {
80         return permutation_.get_const_data();
81     }

```

References `gko::array< ValueType >::get_const_data()`.

47.211.2.8 get_const_scaling_factors()

```

template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::ScaledPermutation< ValueType, IndexType >::get_const_scaling_↵
factors () const [inline], [noexcept]

```

Returns a pointer to the scaling factors.

Returns

the pointer to the scaling factors.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.211.2.9 get_permutation()

```

template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::ScaledPermutation< ValueType, IndexType >::get_permutation ( ) [inline],
[noexcept]

```

Returns a pointer to the permutation indices.

Returns

the pointer to the permutation indices.

References `gko::array< ValueType >::get_data()`.

47.211.2.10 get_scaling_factors()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::ScaledPermutation< ValueType, IndexType >::get_scaling_factors ( )
[inline], [noexcept]
```

Returns a pointer to the scaling factors.

Returns

the pointer to the scaling factors.

References gko::array< ValueType >::get_data().

47.211.2.11 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::ScaledPermutation< ValueType, IndexType >::write (
    gko::matrix_data< value_type, index_type > & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/scaled_permutation.hpp

47.212 gko::experimental::reorder::ScaledReordered< ValueType, IndexType > Class Template Reference

Provides an interface to wrap reorderings like [Rcm](#) and diagonal scaling like equilibration around a LinOp like e.g.

```
#include <ginkgo/core/reorder/scaled_reordered.hpp>
```

Additional Inherited Members

47.212.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::experimental::reorder::ScaledReordered< ValueType, IndexType >
```

Provides an interface to wrap reorderings like [Rcm](#) and diagonal scaling like equilibration around a LinOp like e.g.

a sparse direct solver.

Reorderings can be useful for reducing fill-in in the numerical factorization phase of direct solvers, diagonal scaling can help improve the numerical stability by reducing the condition number of the system matrix.

With a permutation matrix P , a row scaling R and a column scaling C , the inner operator is applied to the system matrix $P^*R^*A^*C^*P^T$ instead of A . Instead of $A*x = b$, the inner operator attempts to solve the equivalent linear system $P^*R^*A^*C^*P^T*y = P^*R^*b$ and retrieves the solution $x = C^*P^T*y$. Note: The inner system matrix is computed from a clone of A , so the original system matrix is not changed.

Template Parameters

<i>ValueType</i>	Type of the values of all matrices used in this class
<i>IndexType</i>	Type of the indices of all matrices used in this class

The documentation for this class was generated from the following file:

- `ginkgo/core/reorder/scaled_reordered.hpp`

47.213 gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType > Class Template Reference

A [Schwarz](#) preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

```
#include <ginkgo/core/distributed/preconditioner/schwarz.hpp>
```

Public Member Functions

- `bool apply_uses_initial_guess ()` const override
Return whether the local solvers use the data in x as an initial guess.

Static Public Member Functions

- static `parameters_type parse` (const `config::pnode` &config, const `config::registry` &context, const `config::type_descriptor` &td_for_child=config::make_type_descriptor< ValueType, LocalIndexType, GlobalIndexType >())
Create the parameters from the property_tree.

47.213.1 Detailed Description

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename GlobalIndexType = int64>
class gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >
```

A [Schwarz](#) preconditioner is a simple domain decomposition preconditioner that generalizes the Block Jacobi preconditioner, incorporating options for different local subdomain solvers and overlaps between the subdomains.

See Iterative Methods for Sparse Linear Systems (Y. Saad) for a general treatment and variations of the method.

Note

Currently overlap and coarse grid correction are not supported (TODO).

Template Parameters

<i>ValueType</i>	precision of matrix element
<i>LocalIndexType</i>	local integer type of the matrix
<i>GlobalIndexType</i>	global integer type of the matrix

47.213.2 Member Function Documentation

47.213.2.1 apply_uses_initial_guess()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
bool gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >::apply_uses_initial_guess ( ) const [override]
```

Return whether the local solvers use the data in x as an initial guess.

Returns

true when the local solvers use the data in x as an initial guess. otherwise, false.

Note

TODO: after adding refining step, need to revisit this.

47.213.2.2 parse()

```
template<typename ValueType = default_precision, typename LocalIndexType = int32, typename
GlobalIndexType = int64>
static parameters_type gko::experimental::distributed::preconditioner::Schwarz< ValueType,
LocalIndexType, GlobalIndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, LocalInd
) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children objects.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children objects. The default uses the value/local/global index type of this class.

Returns

parameters

The documentation for this class was generated from the following file:

- ginkgo/core/distributed/preconditioner/schwarz.hpp

47.214 gko::scoped_device_id_guard Class Reference

This move-only class uses RAII to set the device id within a scoped block, if necessary.

```
#include <ginkgo/core/base/scoped_device_id_guard.hpp>
```

Public Member Functions

- [scoped_device_id_guard](#) (const [ReferenceExecutor](#) *exec, int device_id)
Create a scoped device id from an *Reference*.
- [scoped_device_id_guard](#) (const [OmpExecutor](#) *exec, int device_id)
Create a scoped device id from an *OmpExecutor*.
- [scoped_device_id_guard](#) (const [CudaExecutor](#) *exec, int device_id)
Create a scoped device id from an *CudaExecutor*.
- [scoped_device_id_guard](#) (const [HipExecutor](#) *exec, int device_id)
Create a scoped device id from an *HipExecutor*.
- [scoped_device_id_guard](#) (const [DpcppExecutor](#) *exec, int device_id)
Create a scoped device id from an *DpcppExecutor*.

47.214.1 Detailed Description

This move-only class uses RAII to set the device id within a scoped block, if necessary.

The class behaves similar to `std::scoped_lock`. The scoped guard will make sure that the device code is run on the correct device within one scoped block, when run with multiple devices. Depending on the executor it will record the current device id and set the device id to the one being passed in. After the scope has been exited, the destructor sets the `device_id` back to the one before entering the scope. The [OmpExecutor](#) and [DpcppExecutor](#) don't require setting an device id, so in those cases, the class is a no-op.

The device id scope has to be constructed from a executor with concrete type (not plain [Executor](#)) and a device id. Only the type of the executor object is relevant, so the pointer will not be accessed, and may even be a `nullptr`. From the executor type the correct derived class of `detail::generic_scoped_device_id_guard` is picked. The following illustrates the usage of this class:

```
{
    scoped_device_id_guard g{static_cast<CudaExecutor>(nullptr), 1};
    // now the device id is set to 1
}
// now the device id is reverted again
```

47.214.2 Constructor & Destructor Documentation

47.214.2.1 `scoped_device_id_guard()` [1/5]

```
gko::scoped_device_id_guard::scoped_device_id_guard (
    const ReferenceExecutor * exec,
    int device_id )
```

Create a scoped device id from an [Reference](#).

The resulting object will be a noop.

Parameters

<i>exec</i>	Not used.
<i>device_id</i>	Not used.

47.214.2.2 `scoped_device_id_guard()` [2/5]

```
gko::scoped_device_id_guard::scoped_device_id_guard (
    const OmpExecutor * exec,
    int device_id )
```

Create a scoped device id from an [OmpExecutor](#).

The resulting object will be a noop.

Parameters

<i>exec</i>	Not used.
<i>device</i> ↔ <i>_id</i>	Not used.

47.214.2.3 `scoped_device_id_guard()` [3/5]

```
gko::scoped_device_id_guard::scoped_device_id_guard (
    const CudaExecutor * exec,
    int device_id )
```

Create a scoped device id from an [CudaExecutor](#).

The resulting object will set the cuda device id accordingly.

Parameters

<i>exec</i>	Not used.
<i>device</i> ↔ <i>_id</i>	The device id to use within the scope.

47.214.2.4 `scoped_device_id_guard()` [4/5]

```
gko::scoped_device_id_guard::scoped_device_id_guard (
    const HipExecutor * exec,
    int device_id )
```

Create a scoped device id from an [HipExecutor](#).

The resulting object will set the hip device id accordingly.

Parameters

<i>exec</i>	Not used.
<i>device</i> ↔ <i>_id</i>	The device id to use within the scope.

47.214.2.5 `scoped_device_id_guard()` [5/5]

```
gko::scoped_device_id_guard::scoped_device_id_guard (
    const DpcppExecutor * exec,
    int device_id )
```

Create a scoped device id from an [DpcppExecutor](#).

The resulting object will be a noop.

Parameters

<i>exec</i>	Not used.
<i>device↔ _id</i>	Not used.

The documentation for this class was generated from the following file:

- ginkgo/core/base/scoped_device_id_guard.hpp

47.215 gko::segmented_array< T > Struct Template Reference

A minimal interface for a segmented array.

```
#include <ginkgo/core/base/segmented_array.hpp>
```

Public Member Functions

- [segmented_array](#) (std::shared_ptr< const [Executor](#) > exec)
Create an empty segmented array.
- [segmented_array](#) (std::shared_ptr< const [Executor](#) > exec, const [segmented_array](#) &other)
Copies a segmented array to a different executor.
- [segmented_array](#) (std::shared_ptr< const [Executor](#) > exec, [segmented_array](#) &&other)
Moves a segmented array to a different executor.
- [size_type](#) [get_size](#) () const
Get the total size of the stored buffer.
- [size_type](#) [get_segment_count](#) () const
Get the number of segments.
- T * [get_flat_data](#) ()
Access to the flat buffer.
- const T * [get_const_flat_data](#) () const
Const-access to the flat buffer.
- const [gko::array](#)< int64 > & [get_offsets](#) () const
Access to the segment offsets.
- std::shared_ptr< const [Executor](#) > [get_executor](#) () const
Access the executor.

Static Public Member Functions

- static [segmented_array](#) [create_from_sizes](#) (const [gko::array](#)< int64 > &sizes)
Creates an uninitialized segmented array with predefined segment sizes.
- static [segmented_array](#) [create_from_sizes](#) ([gko::array](#)< T > buffer, const [gko::array](#)< int64 > &sizes)
Creates a segmented array from a flat buffer and segment sizes.
- static [segmented_array](#) [create_from_offsets](#) ([gko::array](#)< int64 > offsets)
Creates an uninitialized segmented array from offsets.
- static [segmented_array](#) [create_from_offsets](#) ([gko::array](#)< T > buffer, [gko::array](#)< int64 > offsets)
Creates a segmented array from a flat buffer and offsets.

47.215.1 Detailed Description

```
template<typename T>
struct gko::segmented_array< T >
```

A minimal interface for a segmented array.

The segmented array is stored as a flat buffer with an offsets array. The segment `i` contains the index range `[offset[i], offset[i + 1])` of the flat buffer.

Template Parameters

<code>T</code>	value type stored in the arrays
----------------	---------------------------------

47.215.2 Constructor & Destructor Documentation

47.215.2.1 segmented_array() [1/3]

```
template<typename T>
gko::segmented_array< T >::segmented_array (
    std::shared_ptr< const Executor > exec ) [explicit]
```

Create an empty segmented array.

Parameters

<code>exec</code>	executor for storage arrays
-------------------	-----------------------------

47.215.2.2 segmented_array() [2/3]

```
template<typename T>
gko::segmented_array< T >::segmented_array (
    std::shared_ptr< const Executor > exec,
    const segmented_array< T > & other )
```

Copies a segmented array to a different executor.

Parameters

<code>exec</code>	the executor to copy to
<code>other</code>	the segmented array to copy from

47.215.2.3 segmented_array() [3/3]

```
template<typename T>
gko::segmented_array< T >::segmented_array (
    std::shared_ptr< const Executor > exec,
    segmented_array< T > && other )
```

Moves a segmented array to a different executor.

Parameters

<i>exec</i>	the executor to move to
<i>other</i>	the segmented array to move from

47.215.3 Member Function Documentation**47.215.3.1 create_from_offsets()** [1/2]

```
template<typename T>
static segmented_array gko::segmented_array< T >::create_from_offsets (
    gko::array< int64 > offsets ) [static]
```

Creates an uninitialized segmented array from offsets.

Parameters

<i>offsets</i>	the index offsets for each segment, and the total size of the buffer as last element
----------------	--

47.215.3.2 create_from_offsets() [2/2]

```
template<typename T>
static segmented_array gko::segmented_array< T >::create_from_offsets (
    gko::array< T > buffer,
    gko::array< int64 > offsets ) [static]
```

Creates a segmented array from a flat buffer and offsets.

Parameters

<i>buffer</i>	the flat buffer whose size has to match the last element of offsets
<i>offsets</i>	the index offsets for each segment, and the total size of the buffer as last element

47.215.3.3 create_from_sizes() [1/2]

```
template<typename T>
static segmented_array gko::segmented_array< T >::create_from_sizes (
    const gko::array< int64 > & sizes ) [static]
```

Creates an uninitialized segmented array with predefined segment sizes.

Parameters

<i>exec</i>	executor for storage arrays
<i>sizes</i>	the sizes of each segment

47.215.3.4 create_from_sizes() [2/2]

```
template<typename T>
static segmented_array gko::segmented_array< T >::create_from_sizes (
    gko::array< T > buffer,
    const gko::array< int64 > & sizes ) [static]
```

Creates a segmented array from a flat buffer and segment sizes.

Parameters

<i>buffer</i>	the flat buffer whose size has to match the sum of sizes
<i>sizes</i>	the sizes of each segment

47.215.3.5 get_const_flat_data()

```
template<typename T>
const T* gko::segmented_array< T >::get_const_flat_data ( ) const
```

Const-access to the flat buffer.

Returns

the flat buffer

47.215.3.6 get_executor()

```
template<typename T>
std::shared_ptr<const Executor> gko::segmented_array< T >::get_executor ( ) const
```

Access the executor.

Returns

the executor

47.215.3.7 get_flat_data()

```
template<typename T>
T* gko::segmented_array< T >::get_flat_data ( )
```

Access to the flat buffer.

Returns

the flat buffer

47.215.3.8 get_offsets()

```
template<typename T>
const gko::array<int64>& gko::segmented_array< T >::get_offsets ( ) const
```

Access to the segment offsets.

Returns

the segment offsets

47.215.3.9 get_segment_count()

```
template<typename T>
size_type gko::segmented_array< T >::get_segment_count ( ) const
```

Get the number of segments.

Returns

the number of segments

47.215.3.10 get_size()

```
template<typename T>
size_type gko::segmented_array< T >::get_size ( ) const
```

Get the total size of the stored buffer.

Returns

the total size of the stored buffer.

The documentation for this struct was generated from the following file:

- ginkgo/core/base/segmented_array.hpp

47.216 gko::matrix::Sellp< ValueType, IndexType > Class Template Reference

SELL-P is a matrix format similar to ELL format.

```
#include <ginkgo/core/matrix/sellp.hpp>
```

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< [Diagonal](#)< ValueType > > [extract_diagonal](#) () const override
Extracts the diagonal entries of the matrix into a vector.
- std::unique_ptr< [absolute_type](#) > [compute_absolute](#) () const override
Gets the AbsoluteLinOp.
- void [compute_absolute_inplace](#) () override
Compute absolute inplace on each element.
- [value_type](#) * [get_values](#) () noexcept
Returns the values of the matrix.
- const [value_type](#) * [get_const_values](#) () const noexcept
Returns the values of the matrix.
- [index_type](#) * [get_col_idxs](#) () noexcept
Returns the column indexes of the matrix.
- const [index_type](#) * [get_const_col_idxs](#) () const noexcept
Returns the column indexes of the matrix.
- [size_type](#) * [get_slice_lengths](#) () noexcept
Returns the lengths(columns) of slices.
- const [size_type](#) * [get_const_slice_lengths](#) () const noexcept
Returns the lengths(columns) of slices.
- [size_type](#) * [get_slice_sets](#) () noexcept
Returns the offsets of slices.
- const [size_type](#) * [get_const_slice_sets](#) () const noexcept
Returns the offsets of slices.
- [size_type](#) [get_slice_size](#) () const noexcept
Returns the size of a slice.
- [size_type](#) [get_stride_factor](#) () const noexcept
Returns the stride factor(t) of SELL-P.
- [size_type](#) [get_total_cols](#) () const noexcept
Returns the total column number.
- [size_type](#) [get_num_stored_elements](#) () const noexcept
Returns the number of elements explicitly stored in the matrix.
- [value_type](#) & [val_at](#) ([size_type](#) row, [size_type](#) slice_set, [size_type](#) idx) noexcept
*Returns the *idx*-th non-zero element of the *row*-th row with *slice_set* slice set.*

- `value_type val_at (size_type row, size_type slice_set, size_type idx) const noexcept`
Returns the `idx`-th non-zero element of the `row`-th row with `slice_set` slice set.
- `index_type & col_at (size_type row, size_type slice_set, size_type idx) noexcept`
Returns the `idx`-th column index of the `row`-th row with `slice_set` slice set.
- `index_type col_at (size_type row, size_type slice_set, size_type idx) const noexcept`
Returns the `idx`-th column index of the `row`-th row with `slice_set` slice set.
- `Sellp & operator= (const Sellp &)`
Copy-assigns a `Sellp` matrix.
- `Sellp & operator= (Sellp &&)`
Move-assigns a `Sellp` matrix.
- `Sellp (const Sellp &)`
Copy-assigns a `Sellp` matrix.
- `Sellp (Sellp &&)`
Move-assigns a `Sellp` matrix.

Static Public Member Functions

- `static std::unique_ptr< Sellp > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size={}, size_type total_cols=0)`
Creates an uninitialized `Sellp` matrix of the specified size.
- `static std::unique_ptr< Sellp > create (std::shared_ptr< const Executor > exec, const dim< 2 > &size, size_type slice_size, size_type stride_factor, size_type total_cols)`
Creates an uninitialized `Sellp` matrix of the specified size.

47.216.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Sellp< ValueType, IndexType >
```

SELL-P is a matrix format similar to ELL format.

The difference is that SELL-P format divides rows into smaller slices and store each slice with ELL format.

This implementation uses the column index value `invalid_index<IndexType>()` to mark padding entries that are not part of the sparsity pattern.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indexes

47.216.2 Constructor & Destructor Documentation

47.216.2.1 Sellp() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Sellp< ValueType, IndexType >::Sellp (
    const Sellp< ValueType, IndexType > & )
```

Copy-assigns a [Sellp](#) matrix.

Inherits the executor, copies the data and parameters.

47.216.2.2 Sellp() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Sellp< ValueType, IndexType >::Sellp (
    Sellp< ValueType, IndexType > && )
```

Move-assigns a [Sellp](#) matrix.

Inherits the executor, moves the data and parameters. The moved-from object is empty (0x0 with valid slice_sets and unchanged parameters).

47.216.3 Member Function Documentation**47.216.3.1 col_at() [1/2]**

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type gko::matrix::Sellp< ValueType, IndexType >::col_at (
    size_type row,
    size_type slice_set,
    size_type idx ) const [inline], [noexcept]
```

Returns the `idx`-th column index of the `row`-th row with `slice_set` slice set.

Parameters

<i>row</i>	the row of the requested element in the slice
<i>slice_set</i>	the slice set of the slice
<i>idx</i>	the <code>idx</code> -th stored element of the row in the slice

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

```
287     {
288         return this
289             ->get_const_col_idxs()[this->linearize_index(row, slice_set, idx)];
290     }
```

47.216.3.2 col_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type& gko::matrix::Sellp< ValueType, IndexType >::col_at (
    size_type row,
    size_type slice_set,
    size_type idx ) [inline], [noexcept]
```

Returns the `idx`-th column index of the `row`-th row with `slice_set` slice set.

Parameters

<i>row</i>	the row of the requested element in the slice
<i>slice_set</i>	the slice set of the slice
<i>idx</i>	the <code>idx</code> -th stored element of the row in the slice

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References `gko::matrix::Sellp< ValueType, IndexType >::get_col_idxs()`.

47.216.3.3 compute_absolute()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<absolute_type> gko::matrix::Sellp< ValueType, IndexType >::compute_absolute (
) const [override], [virtual]
```

Gets the `AbsoluteLinOp`.

Returns

a pointer to the new absolute object

Implements `gko::EnableAbsoluteComputation< remove_complex< Sellp< ValueType, IndexType > > >`.

47.216.3.4 create() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Sellp> gko::matrix::Sellp< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    size_type slice_size,
    size_type stride_factor,
    size_type total_cols ) [static]
```

Creates an uninitialized [Sellp](#) matrix of the specified size.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>slice_size</i>	number of rows in each slice
<i>stride_factor</i>	factor for the stride in each slice (strides should be multiples of the stride_factor)
<i>total_cols</i>	number of the sum of all cols in every slice.

Returns

A smart pointer to the newly created matrix.

47.216.3.5 create() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<Sellp> gko::matrix::Sellp< ValueType, IndexType >::create (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = {},
    size\_type total_cols = 0 ) [static]
```

Creates an uninitialized [Sellp](#) matrix of the specified size.

(The slice_size and stride_factor are set to the default values.)

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>total_cols</i>	number of the sum of all cols in every slice.

Returns

A smart pointer to the newly created matrix.

47.216.3.6 extract_diagonal()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<Diagonal<ValueType> > gko::matrix::Sellp< ValueType, IndexType >::extract_↔
diagonal ( ) const [override], [virtual]
```

Extracts the diagonal entries of the matrix into a vector.

Parameters

<i>diag</i>	the vector into which the diagonal will be written
-------------	--

Implements [gko::DiagonalExtractable< ValueType >](#).

47.216.3.7 get_col_idxxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::Sellp< ValueType, IndexType >::get_col_idxxs ( ) [inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

References [gko::array< ValueType >::get_data\(\)](#).

Referenced by [gko::matrix::Sellp< ValueType, IndexType >::col_at\(\)](#).

47.216.3.8 get_const_col_idxxs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_col_idxxs ( ) const
[inline], [noexcept]
```

Returns the column indexes of the matrix.

Returns

the column indexes of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References [gko::array< ValueType >::get_const_data\(\)](#).

47.216.3.9 get_const_slice_lengths()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const size_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_lengths ( ) const
[inline], [noexcept]
```

Returns the lengths(columns) of slices.

Returns

the lengths(columns) of slices.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References [gko::array< ValueType >::get_const_data\(\)](#).

47.216.3.10 get_const_slice_sets()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const size_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_slice_sets ( ) const
[inline], [noexcept]
```

Returns the offsets of slices.

Returns

the offsets of slices.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.216.3.11 get_const_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::Sellp< ValueType, IndexType >::get_const_values ( ) const [inline],
[noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.216.3.12 get_num_stored_elements()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Sellp< ValueType, IndexType >::get_num_stored_elements ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References `gko::array< ValueType >::get_size()`.

47.216.3.13 get_slice_lengths()

```
template<typename ValueType = default_precision, typename IndexType = int32>  
size_type* gko::matrix::Sellp< ValueType, IndexType >::get_slice_lengths ( ) [inline], [noexcept]
```

Returns the lengths(columns) of slices.

Returns

the lengths(columns) of slices.

References gko::array< ValueType >::get_data().

47.216.3.14 get_slice_sets()

```
template<typename ValueType = default_precision, typename IndexType = int32>  
size_type* gko::matrix::Sellp< ValueType, IndexType >::get_slice_sets ( ) [inline], [noexcept]
```

Returns the offsets of slices.

Returns

the offsets of slices.

References gko::array< ValueType >::get_data().

47.216.3.15 get_slice_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>  
size_type gko::matrix::Sellp< ValueType, IndexType >::get_slice_size ( ) const [inline],  
[noexcept]
```

Returns the size of a slice.

Returns

the size of a slice.

47.216.3.16 get_stride_factor()

```
template<typename ValueType = default_precision, typename IndexType = int32>  
size_type gko::matrix::Sellp< ValueType, IndexType >::get_stride_factor ( ) const [inline],  
[noexcept]
```

Returns the stride factor(t) of SELL-P.

Returns

the stride factor(t) of SELL-P.

47.216.3.17 get_total_cols()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Selp< ValueType, IndexType >::get_total_cols ( ) const [inline],
[noexcept]
```

Returns the total column number.

Returns

the total column number.

References `gko::array< ValueType >::get_size()`.

47.216.3.18 get_values()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::Selp< ValueType, IndexType >::get_values ( ) [inline], [noexcept]
```

Returns the values of the matrix.

Returns

the values of the matrix.

References `gko::array< ValueType >::get_data()`.

47.216.3.19 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Selp& gko::matrix::Selp< ValueType, IndexType >::operator= (
    const Selp< ValueType, IndexType > & )
```

Copy-assigns a [Selp](#) matrix.

Preserves the executor, copies the data and parameters.

47.216.3.20 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
Selp& gko::matrix::Selp< ValueType, IndexType >::operator= (
    Selp< ValueType, IndexType > && )
```

Move-assigns a [Selp](#) matrix.

Preserves the executor, moves the data and parameters. The moved-from object is empty (0x0 with valid `slice_sets` and unchanged parameters).

47.216.3.21 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Selp< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.216.3.22 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Sellp< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.216.3.23 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Sellp< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.216.3.24 val_at() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type gko::matrix::Sellp< ValueType, IndexType >::val_at (
    size_type row,
    size_type slice_set,
    size_type idx ) const [inline], [noexcept]
```

Returns the *idx*-th non-zero element of the *row*-th row with *slice_set* slice set.

Parameters

<i>row</i>	the row of the requested element in the slice
<i>slice_set</i>	the slice set of the slice
<i>idx</i>	the idx-th stored element of the row in the slice

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References `gko::array< ValueType >::get_const_data()`.

47.216.3.25 val_at() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type& gko::matrix::Sellp< ValueType, IndexType >::val_at (
    size_type row,
    size_type slice_set,
    size_type idx ) [inline], [noexcept]
```

Returns the `idx`-th non-zero element of the `row`-th row with `slice_set` slice set.

Parameters

<i>row</i>	the row of the requested element in the slice
<i>slice_set</i>	the slice set of the slice
<i>idx</i>	the idx-th stored element of the row in the slice

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the CPU results in a runtime error)

References `gko::array< ValueType >::get_data()`.

47.216.3.26 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Sellp< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData](#)< [ValueType](#), [IndexType](#) >.

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/csr.hpp
- ginkgo/core/matrix/sellp.hpp

47.217 gko::log::SolverProgress Class Reference

This Logger outputs the value of all scalar values (and potentially vectors) stored internally by the solver after each iteration.

```
#include <ginkgo/core/log/solver_progress.hpp>
```

Static Public Member Functions

- static std::shared_ptr< [SolverProgress](#) > [create_scalar_table_writer](#) (std::ostream &output, int precision=6, int column_width=12)
Creates a logger printing the value for all scalar values in the solver after each iteration in an ASCII table.
- static std::shared_ptr< [SolverProgress](#) > [create_scalar_csv_writer](#) (std::ostream &output, int precision=6, char separator=',')
Creates a logger printing the value for all scalar values in the solver after each iteration in a CSV table.
- static std::shared_ptr< [SolverProgress](#) > [create_vector_storage](#) (std::string output_file_prefix="solver_↵", bool binary=false)
Creates a logger storing all vectors and scalar values in the solver after each iteration on disk.

47.217.1 Detailed Description

This Logger outputs the value of all scalar values (and potentially vectors) stored internally by the solver after each iteration.

It needs to be attached to the solver being inspected.

47.217.2 Member Function Documentation

47.217.2.1 create_scalar_csv_writer()

```
static std::shared_ptr<SolverProgress> gko::log::SolverProgress::create_scalar_csv_writer (
    std::ostream & output,
    int precision = 6,
    char separator = ',' ) [static]
```

Creates a logger printing the value for all scalar values in the solver after each iteration in a CSV table.

If the solver is applied to multiple right-hand sides, only the first right-hand side gets printed.

Parameters

<i>output</i>	the stream to write the output to.
<i>precision</i>	the number of digits of precision to print
<i>separator</i>	the character separating columns from each other

47.217.2.2 create_scalar_table_writer()

```
static std::shared_ptr<SolverProgress> gko::log::SolverProgress::create_scalar_table_writer (
    std::ostream & output,
    int precision = 6,
    int column_width = 12 ) [static]
```

Creates a logger printing the value for all scalar values in the solver after each iteration in an ASCII table.

If the solver is applied to multiple right-hand sides, only the first right-hand side gets printed.

Parameters

<i>output</i>	the stream to write the output to.
<i>precision</i>	the number of digits of precision to print
<i>column_width</i>	the number of characters an output column is wide

47.217.2.3 create_vector_storage()

```
static std::shared_ptr<SolverProgress> gko::log::SolverProgress::create_vector_storage (
    std::string output_file_prefix = "solver_",
    bool binary = false ) [static]
```

Creates a logger storing all vectors and scalar values in the solver after each iteration on disk.

This logger can handle multiple right-hand sides, in contrast to `create_scalar_table_writer` or `create_scalar_csv_writer`.

Parameters

<i>output</i>	the path and file name prefix used to generate the output file names.
<i>precision</i>	the number of digits of precision to print when outputting matrices in text format
<i>binary</i>	if true, write data in Ginkgo's own binary format (lossless), if false write data in the MatrixMarket format (potentially lossy)

The documentation for this class was generated from the following file:

- ginkgo/core/log/solver_progress.hpp

47.218 gko::preconditioner::Sor< ValueType, IndexType > Class Template Reference

This class generates the (S)SOR preconditioner.

```
#include <ginkgo/core/preconditioner/sor.hpp>
```

Public Member Functions

- const parameters_type & [get_parameters](#) ()
Returns the parameters used to construct the factory.
- const parameters_type & [get_parameters](#) () const
Returns the parameters used to construct the factory.
- std::unique_ptr< [composition_type](#) > [generate](#) (std::shared_ptr< const LinOp > system_matrix) const

Static Public Member Functions

- static parameters_type [build](#) ()
Creates a new parameter_type to set up the factory.

47.218.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::preconditioner::Sor< ValueType, IndexType >
```

This class generates the (S)SOR preconditioner.

The SOR preconditioner starts from a splitting of the the matrix A into $A = D + L + U$, where L contains all entries below the diagonal, and U contains all entries above the diagonal. The application of the preconditioner is then defined as solving $Mx = y$ with

$$M = \frac{1}{\omega}(D + \omega L), \quad 0 < \omega < 2.$$

ω is known as the relaxation factor. The preconditioner can be made symmetric, leading to the SSOR preconditioner. Here, M is defined as

$$M = \frac{1}{\omega(2 - \omega)}(D + \omega L)D^{-1}(D + \omega U), \quad 0 < \omega < 2.$$

A detailed description can be found in Iterative Methods for Sparse Linear Systems (Y. Saad) ch. 4.1.

This class is a factory, which will only generate the preconditioner. The resulting LinOp will represent the application of M^{-1} .

Template Parameters

<i>ValueType</i>	The value type of the internally used CSR matrix
<i>IndexType</i>	The index type of the internally used CSR matrix

47.218.2 Member Function Documentation

47.218.2.1 generate()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<composition_type> gko::preconditioner::Sor< ValueType, IndexType >::generate
(
    std::shared_ptr< const LinOp > system_matrix ) const
```

Note

This function overrides the default LinOpFactory::generate to return a Factorization instead of a generic LinOp, which would need to be cast to Factorization again to access its factors. It is only necessary because smart pointers aren't covariant.

47.218.2.2 get_parameters() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const parameters_type& gko::preconditioner::Sor< ValueType, IndexType >::get_parameters ( )
[inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

```
91 { return parameters_; }
```

47.218.2.3 get_parameters() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
const parameters_type& gko::preconditioner::Sor< ValueType, IndexType >::get_parameters ( )
const [inline]
```

Returns the parameters used to construct the factory.

Returns

the parameters used to construct the factory.

The documentation for this class was generated from the following file:

- ginkgo/core/preconditioner/sor.hpp

47.219 gko::span Struct Reference

A span is a lightweight structure used to create sub-ranges from other ranges.

```
#include <ginkgo/core/base/range.hpp>
```

Public Member Functions

- constexpr `span (size_type point)` noexcept
Creates a span representing a point `point`.
- constexpr `span (size_type begin, size_type end)` noexcept
Creates a span.
- constexpr bool `is_valid ()` const
Checks if a span is valid.
- constexpr `size_type length ()` const
Returns the length of a span.

Public Attributes

- const `size_type begin`
Beginning of the span.
- const `size_type end`
End of the span.

47.219.1 Detailed Description

A span is a lightweight structure used to create sub-ranges from other ranges.

A span `s` represents a contiguous set of indexes in one dimension of the range, starting on index `s.begin` (inclusive) and ending at index `s.end` (exclusive). A span is only valid if its starting index is smaller than its ending index.

Spans can be compared using the `==` and `!=` operators. Two spans are identical if both their `begin` and `end` values are identical.

Spans also have two distinct partial orders defined on them:

1. `x < y` (`y > x`) if and only if `x.end < y.begin`
2. `x <= y` (`y >= x`) if and only if `x.end <= y.begin`

Note that the orders are in fact partial - there are spans `x` and `y` for which none of the following inequalities holds: `x < y`, `x > y`, `x == y`, `x <= y`, `x >= y`. An example are spans `span{0, 2}` and `span{1, 3}`.

In addition, `<=` is a distinct order from `<`, and not just an extension of the strict order to its weak equivalent. Thus, `x <= y` is not equivalent to `x < y || x == y`.

47.219.2 Constructor & Destructor Documentation

47.219.2.1 span() [1/2]

```
constexpr gko::span::span (  
    size_type point ) [inline], [constexpr], [noexcept]
```

Creates a span representing a point `point`.

The `begin` of this span is set to `point`, and the `end` to `point + 1`.

Parameters

<i>point</i>	the point which the span represents
--------------	-------------------------------------

47.219.2.2 span() [2/2]

```
constexpr gko::span::span (
    size_type begin,
    size_type end ) [inline], [constexpr], [noexcept]
```

Creates a span.

Parameters

<i>begin</i>	the beginning of the span
<i>end</i>	the end of the span

References begin.

47.219.3 Member Function Documentation**47.219.3.1 is_valid()**

```
constexpr bool gko::span::is_valid ( ) const [inline], [constexpr]
```

Checks if a span is valid.

Returns

true if and only if `this->begin <= this->end`

References begin, and end.

Referenced by `gko::accessor::row_major< ValueType, Dimensionality >::operator()()`.

47.219.3.2 length()

```
constexpr size_type gko::span::length ( ) const [inline], [constexpr]
```

Returns the length of a span.

Returns

`this->end - this->begin`

References begin, and end.

The documentation for this struct was generated from the following file:

- `ginkgo/core/base/range.hpp`

47.220 gko::matrix::Csr< ValueType, IndexType >::sparselib Class Reference

sparselib is a [strategy_type](#) which uses the sparselib csr.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [sparselib](#) ()
Creates a sparselib strategy.
- void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow) override
Computes srow according to row pointers.
- int64_t [clac_size](#) (const int64_t nnz) override
Computes the srow size according to the number of nonzeros.
- std::shared_ptr< [strategy_type](#) > [copy](#) () override
Copy a strategy.

47.220.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::sparselib
```

sparselib is a [strategy_type](#) which uses the sparselib csr.

Note

Uses cusparse in cuda and hipsparse in hip.

47.220.2 Member Function Documentation

47.220.2.1 clac_size()

```
template<typename ValueType = default_precision, typename IndexType = int32>
int64_t gko::matrix::Csr< ValueType, IndexType >::sparselib::clac_size (
    const int64_t nnz ) [inline], [override], [virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.220.2.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::sparselib::copy ( )
[inline], [override], [virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

47.220.2.3 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Csr< ValueType, IndexType >::sparselib::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [inline], [override], [virtual]
```

Computes srow according to row pointers.

Parameters

<i>mtx_row_ptrs</i>	the row pointers of the matrix
<i>mtx_srow</i>	the srow of the matrix

Implements [gko::matrix::Csr< ValueType, IndexType >::strategy_type](#).

The documentation for this class was generated from the following file:

- `ginkgo/core/matrix/csr.hpp`

47.221 [gko::matrix::SparsityCsr< ValueType, IndexType >](#) Class Template Reference

[SparsityCsr](#) is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

```
#include <ginkgo/core/matrix/sparsity_csr.hpp>
```

Public Member Functions

- void [read](#) (const [mat_data](#) &data) override
Reads a matrix from a [matrix_data](#) structure.
- void [read](#) (const [device_mat_data](#) &data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [read](#) ([device_mat_data](#) &&data) override
Reads a matrix from a [device_matrix_data](#) structure.
- void [write](#) ([mat_data](#) &data) const override
Writes a matrix to a [matrix_data](#) structure.
- std::unique_ptr< LinOp > [transpose](#) () const override
Returns a LinOp representing the transpose of the [Transposable](#) object.
- std::unique_ptr< LinOp > [conj_transpose](#) () const override
Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.
- std::unique_ptr< [SparsityCsr](#) > [to_adjacency_matrix](#) () const
Transforms the sparsity matrix to an adjacency matrix.
- void [sort_by_column_index](#) ()
Sorts each row by column index.
- index_type * [get_col_idx](#)s () noexcept
Returns the column indices of the matrix.
- const index_type * [get_const_col_idx](#)s () const noexcept
Returns the column indices of the matrix.
- index_type * [get_row_ptr](#)s () noexcept
Returns the row pointers of the matrix.
- const index_type * [get_const_row_ptr](#)s () const noexcept
Returns the row pointers of the matrix.
- value_type * [get_value](#) () noexcept
Returns the value stored in the matrix.
- const value_type * [get_const_value](#) () const noexcept
Returns the value stored in the matrix.
- [size_type](#) [get_num_nonzeros](#) () const noexcept
Returns the number of elements explicitly stored in the matrix.
- [SparsityCsr](#) & [operator=](#) (const [SparsityCsr](#) &)
Copy-assigns a [SparsityCsr](#) matrix.
- [SparsityCsr](#) & [operator=](#) ([SparsityCsr](#) &&)
Move-assigns a [SparsityCsr](#) matrix.
- [SparsityCsr](#) (const [SparsityCsr](#) &)
Copy-constructs a [SparsityCsr](#) matrix.
- [SparsityCsr](#) ([SparsityCsr](#) &&)
Move-constructs a [SparsityCsr](#) matrix.

Static Public Member Functions

- static std::unique_ptr< [SparsityCsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size=[dim](#)< 2 > {}, [size_type](#) num_nonzeros={})
Creates an uninitialized [SparsityCsr](#) matrix of the specified size.
- static std::unique_ptr< [SparsityCsr](#) > [create](#) (std::shared_ptr< const [Executor](#) > exec, const [dim](#)< 2 > &size, [array](#)< index_type > col_idx, [array](#)< index_type > row_ptr, value_type value=[one](#)< ValueType >())
Creates a [SparsityCsr](#) matrix from already allocated (and initialized) row pointer and column index arrays.

- `template<typename ColIndexType, typename RowPtrType >`
`static std::unique_ptr< SparsityCsr > create (std::shared_ptr< const Executor > exec, const dim< 2 >`
`&size, std::initializer_list< ColIndexType > col_idxes, std::initializer_list< RowPtrType > row_ptrs, value_type`
`value=one< ValueType >())`
`create(std::shared_ptr<const`
- `static std::unique_ptr< SparsityCsr > create (std::shared_ptr< const Executor > exec, std::shared_ptr<`
`const LinOp > matrix)`
Creates a Sparsity matrix from an existing matrix.
- `static std::unique_ptr< const SparsityCsr > create_const (std::shared_ptr< const Executor > exec, const`
`dim< 2 > &size, gko::detail::const_array_view< IndexType > &&col_idxes, gko::detail::const_array_view<`
`IndexType > &&row_ptrs, ValueType value=one< ValueType >())`
Creates a constant (immutable) SparsityCsr matrix from constant arrays.

47.221.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::SparsityCsr< ValueType, IndexType >
```

[SparsityCsr](#) is a matrix format which stores only the sparsity pattern of a sparse matrix by compressing each row of the matrix (compressed sparse row format).

The values of the nonzero elements are stored as a value array of length 1. All the values in the matrix are equal to this value. By default, this value is set to 1.0. A row pointer array also stores the linearized starting index of each row. An additional column index array is used to identify the column where a nonzero is present.

Template Parameters

<i>ValueType</i>	precision of vectors in apply
<i>IndexType</i>	precision of matrix indexes

47.221.2 Constructor & Destructor Documentation

47.221.2.1 SparsityCsr() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::SparsityCsr< ValueType, IndexType >::SparsityCsr (
    const SparsityCsr< ValueType, IndexType > & )
```

Copy-constructs a [SparsityCsr](#) matrix.

Inherits executor, strategy and data.

47.221.2.2 SparsityCsr() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::SparsityCsr< ValueType, IndexType >::SparsityCsr (
    SparsityCsr< ValueType, IndexType > && )
```

Move-constructs a [SparsityCsr](#) matrix.

Inherits executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor, no nonzeros and valid row pointers).

47.221.3 Member Function Documentation

47.221.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::SparsityCsr< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.221.3.2 create() [1/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::create
(
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    array< index_type > col_idxes,
    array< index_type > row_ptrs,
    value_type value = one< ValueType >() ) [static]
```

Creates a [SparsityCsr](#) matrix from already allocated (and initialized) row pointer and column index arrays.

Template Parameters

<i>ColIdxesArray</i>	type of <code>col_idxes</code> array
<i>RowPtrsArray</i>	type of <code>row_ptrs</code> array

Parameters

<i>exec</i>	Executor associated to the matrix
<i>size</i>	size of the matrix
<i>col_idxes</i>	array of column indexes
<i>row_ptrs</i>	array of row pointers
<i>value</i>	value stored. (same value for all matrix elements)

Note

If one of `row_ptrs` or `col_idxes` is not an rvalue, not an array of `IndexType` and `IndexType` respectively, or is on the wrong executor, an internal copy of that array will be created, and the original array data will not be used in the matrix.

47.221.3.3 create() [2/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
template<typename ColIndexType , typename RowPtrType >
static std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::create
(
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    std::initializer_list< ColIndexType > col_idxes,
    std::initializer_list< RowPtrType > row_ptrs,
    value_type value = one<ValueType>() ) [inline], [static]
```

`create(std::shared_ptr<const`

```
create(std::shared_ptr<const executor>="", const dim<2>&, array<index_type>, array<index_type>, value←
_type)
233 {
234     return create(exec, size, array<index_type>{exec, std::move(col_idxes)},
235                  array<index_type>{exec, std::move(row_ptrs)}, value);
236 }
```

References `gko::matrix::SparsityCsr< ValueType, IndexType >::create()`.

47.221.3.4 create() [3/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::create
(
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size = dim< 2 >{},
    size_type num_nonzeros = {} ) [static]
```

Creates an uninitialized `SparsityCsr` matrix of the specified size.

Parameters

<i>exec</i>	<code>Executor</code> associated to the matrix
<i>size</i>	size of the matrix
<i>num_nonzeros</i>	number of nonzeros

Referenced by `gko::matrix::SparsityCsr< ValueType, IndexType >::create()`.

47.221.3.5 create() [4/4]

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::create
(
    std::shared_ptr< const Executor > exec,
    std::shared_ptr< const LinOp > matrix ) [static]
```

Creates a Sparsity matrix from an existing matrix.

Uses the `copy_and_convert_to` functionality.

Parameters

<i>exec</i>	Executor associated to the matrix
<i>matrix</i>	The input matrix

47.221.3.6 create_const()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static std::unique_ptr<const SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >↔
::create_const (
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & size,
    gko::detail::const_array_view< IndexType > && col_idxes,
    gko::detail::const_array_view< IndexType > && row_ptrs,
    ValueType value = one<ValueType>() ) [inline], [static]
```

Creates a constant (immutable) `SparsityCsr` matrix from constant arrays.

Parameters

<i>exec</i>	the executor to create the matrix on
<i>size</i>	the dimensions of the matrix
<i>values</i>	the value array of the matrix
<i>col_idxes</i>	the column index array of the matrix
<i>row_ptrs</i>	the row pointer array of the matrix
<i>strategy</i>	the strategy the matrix uses for SpMV operations

Returns

A smart pointer to the constant matrix wrapping the input arrays (if they reside on the same executor as the matrix) or a copy of these arrays on the correct executor.

47.221.3.7 get_col_idxes()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_col_idxes ( ) [inline],
[noexcept]
```

Returns the column indices of the matrix.

Returns

the column indices of the matrix.

References `gko::array< ValueType >::get_data()`.

47.221.3.8 `get_const_col_idxes()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_col_idxes ( )
const [inline], [noexcept]
```

Returns the column indices of the matrix.

Returns

the column indices of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.221.3.9 `get_const_row_ptrs()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
const index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_row_ptrs ( )
const [inline], [noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References `gko::array< ValueType >::get_const_data()`.

47.221.3.10 get_const_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
const value_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_const_value ( ) const
[inline], [noexcept]
```

Returns the value stored in the matrix.

Returns

the value of the matrix.

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

References gko::array< ValueType >::get_const_data().

47.221.3.11 get_num_nonzeros()

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::SparsityCsr< ValueType, IndexType >::get_num_nonzeros ( ) const [inline],
[noexcept]
```

Returns the number of elements explicitly stored in the matrix.

Returns

the number of elements explicitly stored in the matrix

References gko::array< ValueType >::get_size().

47.221.3.12 get_row_ptrs()

```
template<typename ValueType = default_precision, typename IndexType = int32>
index_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_row_ptrs ( ) [inline],
[noexcept]
```

Returns the row pointers of the matrix.

Returns

the row pointers of the matrix.

References gko::array< ValueType >::get_data().

47.221.3.13 get_value()

```
template<typename ValueType = default_precision, typename IndexType = int32>
value_type* gko::matrix::SparsityCsr< ValueType, IndexType >::get_value ( ) [inline], [noexcept]
```

Returns the value stored in the matrix.

Returns

the value of the matrix.

References `gko::array< ValueType >::get_data()`.

47.221.3.14 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
SparsityCsr& gko::matrix::SparsityCsr< ValueType, IndexType >::operator= (
    const SparsityCsr< ValueType, IndexType > & )
```

Copy-assigns a [SparsityCsr](#) matrix.

Preserves executor, copies everything else.

47.221.3.15 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
SparsityCsr& gko::matrix::SparsityCsr< ValueType, IndexType >::operator= (
    SparsityCsr< ValueType, IndexType > && )
```

Move-assigns a [SparsityCsr](#) matrix.

Preserves executor, moves the data and leaves the moved-from object in an empty state (0x0 LinOp with unchanged executor, no nonzeros and valid row pointers).

47.221.3.16 read() [1/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::SparsityCsr< ValueType, IndexType >::read (
    const device_mat_data & data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.221.3.17 read() [2/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::SparsityCsr< ValueType, IndexType >::read (
    const mat_data & data ) [override], [virtual]
```

Reads a matrix from a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.221.3.18 read() [3/3]

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::SparsityCsr< ValueType, IndexType >::read (
    device_mat_data && data ) [override], [virtual]
```

Reads a matrix from a [device_matrix_data](#) structure.

The structure may be emptied by this function.

Parameters

<i>data</i>	the device_matrix_data structure.
-------------	---

Reimplemented from [gko::ReadableFromMatrixData< ValueType, IndexType >](#).

47.221.3.19 to_adjacency_matrix()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<SparsityCsr> gko::matrix::SparsityCsr< ValueType, IndexType >::to_adjacency↵
_matrix ( ) const
```

Transforms the sparsity matrix to an adjacency matrix.

As the adjacency matrix has to be square, the input [SparsityCsr](#) matrix for this function to work has to be square.

Note

The adjacency matrix in this case is the sparsity pattern but with the diagonal ones removed. This is mainly used for the reordering/partitioning as taken in by graph libraries such as METIS.

47.221.3.20 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::matrix::SparsityCsr< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

47.221.3.21 write()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::SparsityCsr< ValueType, IndexType >::write (
    mat_data & data ) const [override], [virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implements [gko::WritableToMatrixData< ValueType, IndexType >](#).

The documentation for this class was generated from the following files:

- ginkgo/core/matrix/csr.hpp
- ginkgo/core/matrix/sparsity_csr.hpp

47.222 gko::experimental::mpi::status Struct Reference

The status struct is a light wrapper around the MPI_Status struct.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Member Functions

- [status](#) ()
The default constructor.
- MPI_Status * [get](#) ()
Get a pointer to the underlying MPI_Status object.
- template<typename T >
int [get_count](#) (const T *data) const
Get the count of the number of elements received by the communication call.

47.222.1 Detailed Description

The status struct is a light wrapper around the MPI_Status struct.

47.222.2 Constructor & Destructor Documentation

47.222.2.1 status()

```
gko::experimental::mpi::status::status ( ) [inline]
```

The default constructor.

It creates an empty MPI_Status

47.222.3 Member Function Documentation

47.222.3.1 get()

```
MPI_Status* gko::experimental::mpi::status::get ( ) [inline]
```

Get a pointer to the underlying MPI_Status object.

Returns

a pointer to MPI_Status object

Referenced by gko::experimental::mpi::communicator::recv(), and gko::experimental::mpi::request::wait().

47.222.3.2 get_count()

```
template<typename T >  
int gko::experimental::mpi::status::get_count (   
    const T * data ) const [inline]
```

Get the count of the number of elements received by the communication call.

Template Parameters

<i>T</i>	The datatype of the object that was received.
----------	---

Parameters

<i>data</i>	The data object of type T that was received.
-------------	--

Returns

the count

The documentation for this struct was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.223 gko::stopping_status Class Reference

This class is used to keep track of the stopping status of one vector.

```
#include <ginkgo/core/stop/stopping_status.hpp>
```

Public Member Functions

- bool [has_stopped](#) () const noexcept
Check if any stopping criteria was fulfilled.
- bool [has_converged](#) () const noexcept
Check if convergence was reached.
- bool [is_finalized](#) () const noexcept
Check if the corresponding vector stores the finalized result.
- [uint8 get_id](#) () const noexcept
Get the id of the stopping criterion which caused the stop.
- void [reset](#) () noexcept
Clear all flags.
- void [stop](#) (uint8 id, bool set_finalized=true) noexcept
Call if a stop occurred due to a hard limit (and convergence was not reached).
- void [converge](#) (uint8 id, bool set_finalized=true) noexcept
Call if convergence occurred.
- void [finalize](#) () noexcept
Set the result to be finalized (it needs to be stopped or converged first).

Friends

- bool [operator==](#) (const [stopping_status](#) &x, const [stopping_status](#) &y) noexcept
Checks if two stopping statuses are equivalent.
- bool [operator!=](#) (const [stopping_status](#) &x, const [stopping_status](#) &y) noexcept
Checks if two stopping statuses are different.

47.223.1 Detailed Description

This class is used to keep track of the stopping status of one vector.

47.223.2 Member Function Documentation

47.223.2.1 converge()

```
void gko::stopping_status::converge (
    uint8 id,
    bool set_finalized = true ) [inline], [noexcept]
```

Call if convergence occurred.

Parameters

<i>id</i>	id of the stopping criteria.
<i>set_finalized</i>	Controls if the current version should count as finalized (set to true) or not (set to false).

References `has_stopped()`.

47.223.2.2 get_id()

```
uint8 gko::stopping_status::get_id ( ) const [inline], [noexcept]
```

Get the id of the stopping criterion which caused the stop.

Returns

Returns the id of the stopping criterion which caused the stop.

Referenced by `has_stopped()`.

47.223.2.3 has_converged()

```
bool gko::stopping_status::has_converged ( ) const [inline], [noexcept]
```

Check if convergence was reached.

Returns

Returns true if convergence was reached.

47.223.2.4 has_stopped()

```
bool gko::stopping_status::has_stopped ( ) const [inline], [noexcept]
```

Check if any stopping criteria was fulfilled.

Returns

Returns true if any stopping criteria was fulfilled.

References `get_id()`.

Referenced by `converge()`, `finalize()`, and `stop()`.

47.223.2.5 is_finalized()

```
bool gko::stopping_status::is_finalized ( ) const [inline], [noexcept]
```

Check if the corresponding vector stores the finalized result.

Returns

Returns true if the corresponding vector stores the finalized result.

47.223.2.6 stop()

```
void gko::stopping_status::stop (
    uint8 id,
    bool set_finalized = true ) [inline], [noexcept]
```

Call if a stop occurred due to a hard limit (and convergence was not reached).

Parameters

<i>id</i>	id of the stopping criteria.
<i>set_finalized</i>	Controls if the current version should count as finalized (set to true) or not (set to false).

References `has_stopped()`.

47.223.3 Friends And Related Function Documentation

47.223.3.1 operator!=

```
bool operator!= (
    const stopping_status & x,
    const stopping_status & y ) [friend]
```

Checks if two stopping statuses are different.

Parameters

<i>x</i>	a stopping status
<i>y</i>	a stopping status

Returns

true if and only if ! (*x* == *y*)

47.223.3.2 operator==

```
bool operator== (
    const stopping_status & x,
    const stopping_status & y ) [friend]
```

Checks if two stopping statuses are equivalent.

Parameters

<i>x</i>	a stopping status
<i>y</i>	a stopping status

Returns

true if and only if both *x* and *y* have the same mask and converged and finalized state

The documentation for this class was generated from the following file:

- ginkgo/core/stop/stopping_status.hpp

47.224 gko::matrix::Csr< ValueType, IndexType >::strategy_type Class Reference

[strategy_type](#) is to decide how to set the csr algorithm.

```
#include <ginkgo/core/matrix/csr.hpp>
```

Public Member Functions

- [strategy_type](#) (std::string name)
Creates a [strategy_type](#).
- std::string [get_name](#) ()
Returns the name of strategy.
- virtual void [process](#) (const [array](#)< index_type > &mtx_row_ptrs, [array](#)< index_type > *mtx_srow)=0
Computes srow according to row pointers.
- virtual int64_t [clac_size](#) (const int64_t nnz)=0
Computes the srow size according to the number of nonzeros.
- virtual std::shared_ptr< [strategy_type](#) > [copy](#) ()=0
Copy a strategy.

47.224.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Csr< ValueType, IndexType >::strategy_type
```

[strategy_type](#) is to decide how to set the csr algorithm.

The practical strategy method should inherit [strategy_type](#) and implement its `process`, `clac_size` function and the corresponding device kernel.

47.224.2 Constructor & Destructor Documentation

47.224.2.1 [strategy_type](#)()

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::matrix::Csr< ValueType, IndexType >::strategy_type::strategy_type (
    std::string name ) [inline]
```

Creates a [strategy_type](#).

Parameters

<i>name</i>	the name of strategy
-------------	----------------------

47.224.3 Member Function Documentation

47.224.3.1 [clac_size](#)()

```
template<typename ValueType = default_precision, typename IndexType = int32>
virtual int64_t gko::matrix::Csr< ValueType, IndexType >::strategy_type::clac_size (
```

```
const int64_t nnz ) [pure virtual]
```

Computes the srow size according to the number of nonzeros.

Parameters

<i>nnz</i>	the number of nonzeros
------------	------------------------

Returns

the size of srow

Implemented in [gko::matrix::Csr< ValueType, IndexType >::load_balance](#), [gko::matrix::Csr< ValueType, IndexType >::sparselib](#), [gko::matrix::Csr< ValueType, IndexType >::cusparse](#), [gko::matrix::Csr< ValueType, IndexType >::merge_path](#), and [gko::matrix::Csr< ValueType, IndexType >::classical](#).

47.224.3.2 copy()

```
template<typename ValueType = default_precision, typename IndexType = int32>
virtual std::shared_ptr<strategy_type> gko::matrix::Csr< ValueType, IndexType >::strategy_type::copy ( ) [pure virtual]
```

Copy a strategy.

This is a workaround until strategies are revamped, since strategies like `automatical` do not work when actually shared.

Implemented in [gko::matrix::Csr< ValueType, IndexType >::load_balance](#), [gko::matrix::Csr< ValueType, IndexType >::sparselib](#), [gko::matrix::Csr< ValueType, IndexType >::cusparse](#), [gko::matrix::Csr< ValueType, IndexType >::merge_path](#), and [gko::matrix::Csr< ValueType, IndexType >::classical](#).

47.224.3.3 get_name()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::string gko::matrix::Csr< ValueType, IndexType >::strategy_type::get_name ( ) [inline]
```

Returns the name of strategy.

Returns

the name of strategy

47.224.3.4 process()

```
template<typename ValueType = default_precision, typename IndexType = int32>
virtual void gko::matrix::Csr< ValueType, IndexType >::strategy_type::process (
    const array< index_type > & mtx_row_ptrs,
    array< index_type > * mtx_srow ) [pure virtual]
```

Computes srow according to row pointers.

Parameters

<i>mtx_row_ptrs</i>	the row pointers of the matrix
<i>mtx_srow</i>	the srow of the matrix

Implemented in [gko::matrix::Csr< ValueType, IndexType >::load_balance](#), [gko::matrix::Csr< ValueType, IndexType >::sparselib](#), [gko::matrix::Csr< ValueType, IndexType >::cusparse](#), [gko::matrix::Csr< ValueType, IndexType >::merge_path](#), and [gko::matrix::Csr< ValueType, IndexType >::classical](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/matrix/csr.hpp](#)

47.225 gko::matrix::Hybrid< ValueType, IndexType >::strategy_type

Class Reference

[strategy_type](#) is to decide how to set the hybrid config.

```
#include <ginkgo/core/matrix/hybrid.hpp>
```

Public Member Functions

- [strategy_type](#) ()
Creates a [strategy_type](#).
- void [compute_hybrid_config](#) (const [array](#)< [size_type](#) > &row_nnz, [size_type](#) *ell_num_stored_elements_↵
per_row, [size_type](#) *coo_nnz)
Computes the config of the [Hybrid](#) matrix (ell_num_stored_elements_per_row and coo_nnz).
- [size_type](#) [get_ell_num_stored_elements_per_row](#) () const noexcept
Returns the number of stored elements per row of the ell part.
- [size_type](#) [get_coo_nnz](#) () const noexcept
Returns the number of nonzeros of the coo part.
- virtual [size_type](#) [compute_ell_num_stored_elements_per_row](#) ([array](#)< [size_type](#) > *row_nnz) const =0
Computes the number of stored elements per row of the ell part.

47.225.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::matrix::Hybrid< ValueType, IndexType >::strategy_type
```

[strategy_type](#) is to decide how to set the hybrid config.

It computes the number of stored elements per row of the ell part and then set the number of residual nonzeros as the number of nonzeros of the coo part.

The practical strategy method should inherit [strategy_type](#) and implement its [compute_ell_num_stored_↵
elements_per_row](#) function.

47.225.2 Member Function Documentation

47.225.2.1 compute_ell_num_stored_elements_per_row()

```
template<typename ValueType = default_precision, typename IndexType = int32>
virtual size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_ell_num_stored_elements_per_row (
    array< size_type > * row_nnz ) const [pure virtual]
```

Computes the number of stored elements per row of the ell part.

Parameters

<i>row_nnz</i>	the number of nonzeros of each row
----------------	------------------------------------

Returns

the number of stored elements per row of the ell part

Implemented in [gko::matrix::Hybrid< ValueType, IndexType >::automatic](#), [gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage](#), [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit](#), [gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bound](#), and [gko::matrix::Hybrid< ValueType, IndexType >::column_limit](#).

Referenced by [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config\(\)](#).

47.225.2.2 compute_hybrid_config()

```
template<typename ValueType = default_precision, typename IndexType = int32>
void gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_hybrid_config (
    const array< size_type > & row_nnz,
    size_type * ell_num_stored_elements_per_row,
    size_type * coo_nnz ) [inline]
```

Computes the config of the [Hybrid](#) matrix (ell_num_stored_elements_per_row and coo_nnz).

For now, it copies row_nnz to the reference executor and performs all operations on the reference executor.

Parameters

<i>row_nnz</i>	the number of nonzeros of each row
<i>ell_num_stored_elements_per_row</i>	the output number of stored elements per row of the ell part
<i>coo_nnz</i>	the output number of nonzeros of the coo part

References [gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::compute_ell_num_stored_elements_per_row\(\)](#), [gko::array< ValueType >::get_executor\(\)](#), and [gko::array< ValueType >::get_size\(\)](#).

47.225.2.3 `get_coo_nnz()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::get_coo_nnz ( ) const
[inline], [noexcept]
```

Returns the number of nonzeros of the coo part.

Returns

the number of nonzeros of the coo part

47.225.2.4 `get_ell_num_stored_elements_per_row()`

```
template<typename ValueType = default_precision, typename IndexType = int32>
size_type gko::matrix::Hybrid< ValueType, IndexType >::strategy_type::get_ell_num_stored_↵
elements_per_row ( ) const [inline], [noexcept]
```

Returns the number of stored elements per row of the ell part.

Returns

the number of stored elements per row of the ell part

The documentation for this class was generated from the following file:

- ginkgo/core/matrix/hybrid.hpp

47.226 `gko::log::Stream< ValueType >` Class Template Reference

`Stream` is a Logger which logs every event to a stream.

```
#include <ginkgo/core/log/stream.hpp>
```

Static Public Member Functions

- static std::unique_ptr< `Stream` > `create` (std::shared_ptr< const `Executor` > exec, const Logger::mask_type &enabled_events=Logger::all_events_mask, std::ostream &os=std::cout, bool verbose=false)
Creates a `Stream` logger.
- static std::unique_ptr< `Stream` > `create` (const Logger::mask_type &enabled_events=Logger::all_events_↵mask, std::ostream &os=std::cerr, bool verbose=false)
Creates a `Stream` logger.

47.226.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::log::Stream< ValueType >
```

`Stream` is a Logger which logs every event to a stream.

This can typically be used to log to a file or to the console.

Template Parameters

<i>ValueType</i>	the type of values stored in the class (i.e. ValueType template parameter of the concrete Loggable this class will log)
------------------	---

47.226.2 Member Function Documentation

47.226.2.1 create() [1/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Stream> gko::log::Stream< ValueType >::create (
    const Logger::mask_type & enabled_events = Logger::all_events_mask,
    std::ostream & os = std::cerr,
    bool verbose = false ) [inline], [static]
```

Creates a [Stream](#) logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.
<i>os</i>	the stream used for this logger
<i>verbose</i>	whether we want detailed information or not. This includes always printing residuals and other information which can give a large output.

Returns

an std::unique_ptr to the the constructed object

```
195     {
196         return std::unique_ptr<Stream>(new Stream(enabled_events, os, verbose));
197     }
```

47.226.2.2 create() [2/2]

```
template<typename ValueType = default_precision>
static std::unique_ptr<Stream> gko::log::Stream< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    const Logger::mask_type & enabled_events = Logger::all_events_mask,
    std::ostream & os = std::cout,
    bool verbose = false ) [inline], [static]
```

Creates a [Stream](#) logger.

This dynamically allocates the memory, constructs the object and returns an std::unique_ptr to this object.

Parameters

<i>exec</i>	the executor
<i>enabled_events</i>	the events enabled for this logger. By default all events.
<i>os</i>	the stream used for this logger
<i>verbose</i>	whether we want detailed information or not. This includes always printing residuals and other information which can give a large output.

Returns

an `std::unique_ptr` to the the constructed object

The documentation for this class was generated from the following file:

- `ginkgo/core/log/stream.hpp`

47.227 gko::StreamError Class Reference

[StreamError](#) is thrown if accessing a stream failed.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [StreamError](#) (const `std::string` &file, int line, const `std::string` &func, const `std::string` &message)
Initializes a file access error.

47.227.1 Detailed Description

[StreamError](#) is thrown if accessing a stream failed.

47.227.2 Constructor & Destructor Documentation

47.227.2.1 StreamError()

```
gko::StreamError::StreamError (
    const std::string & file,
    int line,
    const std::string & func,
    const std::string & message ) [inline]
```

Initializes a file access error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The name of the function that tried to access the file
<i>message</i>	The error message

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.228 gko::log::ProfilerHook::SummaryWriter Class Reference

Receives the results from [ProfilerHook::create_summary\(\)](#).

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

- virtual void [write](#) (const std::vector< summary_entry > &entries, std::chrono::nanoseconds overhead)=0
Callback to write out the summary results.

47.228.1 Detailed Description

Receives the results from [ProfilerHook::create_summary\(\)](#).

47.228.2 Member Function Documentation

47.228.2.1 write()

```
virtual void gko::log::ProfilerHook::SummaryWriter::write (
    const std::vector< summary_entry > & entries,
    std::chrono::nanoseconds overhead ) [pure virtual]
```

Callback to write out the summary results.

Parameters

<i>entries</i>	the vector of ranges with runtime and count.
<i>overhead</i>	an estimate of the profiler overhead

Implemented in [gko::log::ProfilerHook::TableSummaryWriter](#).

The documentation for this class was generated from the following file:

- ginkgo/core/log/profiler_hook.hpp

47.229 gko::log::ProfilerHook::TableSummaryWriter Class Reference

Writes the results from [ProfilerHook::create_summary\(\)](#) and [ProfilerHook::create_nested_summary\(\)](#) to a ASCII table in Markdown format.

```
#include <ginkgo/core/log/profiler_hook.hpp>
```

Public Member Functions

- [TableSummaryWriter](#) (std::ostream &output=std::cerr, std::string header="Runtime summary")
Constructs a writer on an output stream.
- void [write](#) (const std::vector< summary_entry > &entries, std::chrono::nanoseconds overhead) override
Callback to write out the summary results.
- void [write_nested](#) (const nested_summary_entry &root, std::chrono::nanoseconds overhead) override
Callback to write out the summary results.

47.229.1 Detailed Description

Writes the results from [ProfilerHook::create_summary\(\)](#) and [ProfilerHook::create_nested_summary\(\)](#) to a ASCII table in Markdown format.

47.229.2 Constructor & Destructor Documentation

47.229.2.1 TableSummaryWriter()

```
gko::log::ProfilerHook::TableSummaryWriter::TableSummaryWriter (
    std::ostream & output = std::cerr,
    std::string header = "Runtime summary" )
```

Constructs a writer on an output stream.

Parameters

<i>output</i>	the output stream to write the table to.
<i>header</i>	the header to write above the table.

47.229.3 Member Function Documentation

47.229.3.1 write()

```
void gko::log::ProfilerHook::TableSummaryWriter::write (
    const std::vector< summary_entry > & entries,
    std::chrono::nanoseconds overhead ) [override], [virtual]
```

Callback to write out the summary results.

Parameters

<i>entries</i>	the vector of ranges with runtime and count.
<i>overhead</i>	an estimate of the profiler overhead

Implements [gko::log::ProfilerHook::SummaryWriter](#).

47.229.3.2 write_nested()

```
void gko::log::ProfilerHook::TableSummaryWriter::write_nested (
    const nested_summary_entry & root,
    std::chrono::nanoseconds overhead ) [override], [virtual]
```

Callback to write out the summary results.

Parameters

<i>root</i>	the root range with runtime and count.
<i>overhead</i>	an estimate of the profiler overhead

Implements [gko::log::ProfilerHook::NestedSummaryWriter](#).

The documentation for this class was generated from the following file:

- ginkgo/core/log/profiler_hook.hpp

47.230 gko::stop::Time Class Reference

The [Time](#) class is a stopping criterion which stops the iteration process after a certain amount of time has passed.

```
#include <ginkgo/core/stop/time.hpp>
```

47.230.1 Detailed Description

The [Time](#) class is a stopping criterion which stops the iteration process after a certain amount of time has passed.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/time.hpp`

47.231 `gko::time_point` Class Reference

An opaque wrapper for a time point generated by a timer.

```
#include <ginkgo/core/base/timer.hpp>
```

47.231.1 Detailed Description

An opaque wrapper for a time point generated by a timer.

The documentation for this class was generated from the following file:

- `ginkgo/core/base/timer.hpp`

47.232 `gko::Timer` Class Reference

Represents a generic timer that can be used to record time points and measure time differences on host or device streams.

```
#include <ginkgo/core/base/timer.hpp>
```

Public Member Functions

- [time_point](#) `create_time_point` ()
Returns a newly created time point.
- virtual void [record](#) ([time_point](#) &time)=0
Records a time point at the current time.
- virtual void [wait](#) ([time_point](#) &time)=0
Waits until all kernels in-process when recording the time point are finished.
- `std::chrono::nanoseconds` [difference](#) ([time_point](#) &start, [time_point](#) &stop)
Computes the difference between the two time points in nanoseconds.
- virtual `std::chrono::nanoseconds` [difference_async](#) (const [time_point](#) &start, const [time_point](#) &stop)=0
Computes the difference between the two time points in nanoseconds.

Static Public Member Functions

- static `std::unique_ptr< Timer >` [create_for_executor](#) (`std::shared_ptr< const Executor >` exec)
Creates the timer type most suitable for recording accurate timings of kernels on the given executor.

47.232.1 Detailed Description

Represents a generic timer that can be used to record time points and measure time differences on host or device streams.

To keep the runtime overhead of timing minimal, time points need to be allocated beforehand using [Timer::create_time_point](#):

```
auto begin = timer->create_time_point();
auto end = timer->create_time_point();
// ...
timer->record(begin);
run_expensive_operation();
timer->record(end);
auto elapsed = timer->difference(begin, end);
```

47.232.2 Member Function Documentation

47.232.2.1 create_for_executor()

```
static std::unique_ptr<Timer> gko::Timer::create_for_executor (
    std::shared_ptr< const Executor > exec ) [static]
```

Creates the timer type most suitable for recording accurate timings of kernels on the given executor.

Parameters

<code>exec</code>	the executor to create a Timer for
-------------------	--

Returns

[CpuTimer](#) for [ReferenceExecutor](#) and [OmpExecutor](#), [CudaTimer](#) for [CudaExecutor](#), [HipTimer](#) for [HipExecutor](#) or [DpcppTimer](#) for [DpcppExecutor](#).

47.232.2.2 create_time_point()

```
time_point gko::Timer::create_time_point ( )
```

Returns a newly created time point.

Time points may only be used with the timer they were created with.

47.232.2.3 difference()

```
std::chrono::nanoseconds gko::Timer::difference (
    time_point & start,
    time_point & stop )
```

Computes the difference between the two time points in nanoseconds.

The function synchronizes with `stop` before computing the difference.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

47.232.2.4 difference_async()

```
virtual std::chrono::nanoseconds gko::Timer::difference_async (
    const time_point & start,
    const time_point & stop ) [pure virtual]
```

Computes the difference between the two time points in nanoseconds.

This asynchronous version does not synchronize itself, so the time points need to have been synchronized with, i.e. `timer->wait(stop)` needs to have been called. The version is intended for more advanced users who want to measure the overhead of timing functionality separately.

Parameters

<i>start</i>	the first time point (earlier)
<i>end</i>	the second time point (later)

Returns

the difference between the time points in nanoseconds.

Implemented in [gko::DpcppTimer](#), [gko::HipTimer](#), [gko::CudaTimer](#), and [gko::CpuTimer](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/timer.hpp](#)

47.233 gko::Transposable Class Reference

Linear operators which support transposition should implement the [Transposable](#) interface.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual `std::unique_ptr< LinOp > transpose ()` const =0
Returns a *LinOp* representing the transpose of the *Transposable* object.
- virtual `std::unique_ptr< LinOp > conj_transpose ()` const =0
Returns a *LinOp* representing the conjugate transpose of the *Transposable* object.

47.233.1 Detailed Description

Linear operators which support transposition should implement the [Transposable](#) interface.

It provides two functionalities, the normal transpose and the conjugate transpose.

The normal transpose returns the transpose of the linear operator without changing any of its elements representing the operation, $B = A^T$.

The conjugate transpose returns the conjugate of each of the elements and additionally transposes the linear operator representing the operation, $B = A^H$.

47.233.1.1 Example: Transposing a Csr matrix:

```
{c++}
//Transposing an object of LinOp type.
//The object you want to transpose.
auto op = matrix::Csr::create(exec);
//Transpose the object by first converting it to a transposable type.
auto trans = op->transpose();
```

47.233.2 Member Function Documentation

47.233.2.1 conj_transpose()

```
virtual std::unique_ptr<LinOp> gko::Transposable::conj_transpose ( ) const [pure virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >](#), [gko::preconditioner::Jacobi< ValueType, IndexType >](#), [gko::matrix::Fft3](#), [gko::preconditioner::lc< LSolverType, IndexType >](#), [gko::solver::UpperTrs< ValueType, IndexType >](#), [gko::preconditioner::lsai< IsaiType, ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Fft2](#), [gko::solver::lr< ValueType >](#), [gko::matrix::SparsityCsr< ValueType, IndexType >](#), [gko::solver::Gmres< ValueType >](#), [gko::matrix::Diagonal< ValueType >](#), [gko::solver::LowerTrs< ValueType, IndexType >](#), [gko::solver::ldr< ValueType >](#), [gko::solver::Bicg< ValueType >](#), [gko::Combination< ValueType >](#), [gko::solver::Bicgstab< ValueType >](#), [gko::solver::Fcg< ValueType >](#), [gko::matrix::Fft](#), [gko::Composition< ValueType >](#), [gko::solver::Gcr< ValueType >](#), [gko::solver::Cg< ValueType >](#), [gko::solver::Cgs< ValueType >](#), [gko::experimental::solver::Direct< ValueType, IndexType >](#), and [gko::matrix::Identity< ValueType >](#)

47.233.2.2 transpose()

```
virtual std::unique_ptr<LinOp> gko::Transposable::transpose ( ) const [pure virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implemented in [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Dense< ValueType >](#), [gko::matrix::Dense< value_type >](#), [gko::preconditioner::Ilu< LSolverType, USolverType, ReverseApply, IndexType >](#), [gko::preconditioner::Jacobi< ValueType, IndexType >](#), [gko::matrix::Fft3](#), [gko::solver::UpperTrs< ValueType, IndexType >](#), [gko::preconditioner::Ic< LSolverType, IndexType >](#), [gko::preconditioner::Isai< IsaiType, ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Fft2](#), [gko::solver::Ir< ValueType >](#), [gko::matrix::SparsityCsr< ValueType, IndexType >](#), [gko::solver::Gmres< ValueType >](#), [gko::matrix::Diagonal< ValueType >](#), [gko::solver::LowerTrs< ValueType, IndexType >](#), [gko::solver::Idr< ValueType >](#), [gko::solver::Bicg< ValueType >](#), [gko::Combination< ValueType >](#), [gko::solver::Bicgstab< ValueType >](#), [gko::solver::Fcg< ValueType >](#), [gko::matrix::Fft](#), [gko::Composition< ValueType >](#), [gko::solver::Gcr< ValueType >](#), [gko::solver::Cg< ValueType >](#), [gko::solver::Cgs< ValueType >](#), [gko::experimental::solver::Direct< ValueType, IndexType >](#), and [gko::matrix::Identity< ValueType >](#)

The documentation for this class was generated from the following file:

- [ginkgo/core/base/lin_op.hpp](#)

47.234 gko::config::type_descriptor Class Reference

This class describes the value and index types to be used when building a Ginkgo type from a configuration file.

```
#include <ginkgo/core/config/type_descriptor.hpp>
```

Public Member Functions

- [type_descriptor](#) (std::string value_typestr="float64", std::string index_typestr="int32", std::string global_index_typestr="int64")
type_descriptor constructor.
- const std::string & [get_value_typestr](#) () const
Get the value type string.
- const std::string & [get_index_typestr](#) () const
Get the index type string.
- const std::string & [get_local_index_typestr](#) () const
Get the local index type string, which gives the same result as [get_index_typestr\(\)](#)
- const std::string & [get_global_index_typestr](#) () const
Get the global index type string.

47.234.1 Detailed Description

This class describes the value and index types to be used when building a Ginkgo type from a configuration file.

A [type_descriptor](#) is passed in order to define the parse function defines which template parameters, in terms of `value_type` and/or `index_type`, the created object will have. For example, a CG solver created like this:

```
auto cg = parse(config, context, type_descriptor("float64", "int32"));
```

will have the value type `float64` and the index type `int32`. Any Ginkgo object that does not require one of these types will just ignore it. In `value_type`, one additional value `void` can be used to specify that no default type is provided. In this case, the configuration has to provide the necessary template types.

If the configuration specifies one field (only allow `value_type` now):

```
value_type: "some_value_type"
```

this type will take precedence over the [type_descriptor](#).

47.234.2 Constructor & Destructor Documentation

47.234.2.1 type_descriptor()

```
gko::config::type_descriptor::type_descriptor (
    std::string value_typestr = "float64",
    std::string index_typestr = "int32",
    std::string global_index_typestr = "int64" ) [explicit]
```

[type_descriptor](#) constructor.

There is free function `make_type_descriptor` to create the object by template.

Parameters

<i>value_typestr</i>	the value type string. "void" means no default.
<i>index_typestr</i>	the (local) index type string.
<i>global_index_typestr</i>	the global index type string.

Note

there is no way to call the constructor with explicit template, so we create another free function to handle it.

The documentation for this class was generated from the following file:

- `ginkgo/core/config/type_descriptor.hpp`

47.235 gko::experimental::mpi::type_impl< T > Struct Template Reference

A struct that is used to determine the `MPI_Datatype` of a specified type.

```
#include <ginkgo/core/base/mpi.hpp>
```

47.235.1 Detailed Description

```
template<typename T>
struct gko::experimental::mpi::type_impl< T >
```

A struct that is used to determine the MPI_Datatype of a specified type.

Template Parameters

<i>T</i>	type of which the MPI_Datatype should be inferred.
----------	--

Note

any specialization of this type has to provide a static function `get_type()` that returns an MPI_Datatype

The documentation for this struct was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.236 gko::syn::type_list< Types > Struct Template Reference

[type_list](#) records several types in template

```
#include <ginkgo/core/synthesizer/containers.hpp>
```

47.236.1 Detailed Description

```
template<typename... Types>
struct gko::syn::type_list< Types >
```

[type_list](#) records several types in template

Template Parameters

<i>Types</i>	the types in the list
--------------	-----------------------

The documentation for this struct was generated from the following file:

- ginkgo/core/synthesizer/containers.hpp

47.237 gko::UnsupportedMatrixProperty Class Reference

Exception throws if a matrix does not have a property required by a numerical method.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [UnsupportedMatrixProperty](#) (const std::string &file, const int line, const std::string &msg)
Initializes the [UnsupportedMatrixProperty](#) error.

47.237.1 Detailed Description

Exception throws if a matrix does not have a property required by a numerical method.

Currently, a message is specified at the call-site manually.

47.237.2 Constructor & Destructor Documentation

47.237.2.1 UnsupportedMatrixProperty()

```
gko::UnsupportedMatrixProperty::UnsupportedMatrixProperty (
    const std::string & file,
    const int line,
    const std::string & msg ) [inline]
```

Initializes the [UnsupportedMatrixProperty](#) error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>msg</i>	A message describing the property required.

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.238 gko::stop::Criterion::Updater Class Reference

The [Updater](#) class serves for convenient argument passing to the [Criterion](#)'s check function.

```
#include <ginkgo/core/stop/criterion.hpp>
```

Public Member Functions

- [Updater](#) (const [Updater](#) &)=delete
Prevent copying and moving the object This is to enforce the use of argument passing and calling check at the same time.
- bool [check](#) (uint8 stopping_id, bool set_finalized, array< [stopping_status](#) > *stop_status, bool *one_↔ changed) const
Calls the parent [Criterion](#) object's check method.

47.238.1 Detailed Description

The [Updater](#) class serves for convenient argument passing to the [Criterion](#)'s check function.

The pattern used is a Builder, except [Updater](#) builds a function's arguments before calling the function itself, and does not build an object. This allows calling a [Criterion](#)'s check in the form of: `stop_criterion->update().num_iterations(num_iterations) .ignore_residual_check(ignore_residual_check) .residual_norm(residual_norm) .implicit_sq_residual_norm(implicit_sq_residual_norm) .residual(residual) .solution(solution) .check(converged);`

If there is a need for a new form of data to pass to the [Criterion](#), it should be added here.

47.238.2 Member Function Documentation

47.238.2.1 check()

```
bool gko::stop::Criterion::Updater::check (
    uint8 stopping_id,
    bool set_finalized,
    array< stopping_status > * stop_status,
    bool * one_changed ) const [inline]
```

Calls the parent [Criterion](#) object's check method.

References `gko::stop::Criterion::check()`.

The documentation for this class was generated from the following file:

- `ginkgo/core/stop/criterion.hpp`

47.239 gko::solver::UpperTrs< ValueType, IndexType > Class Template Reference

[UpperTrs](#) is the triangular solver which solves the system $Ux = b$, when U is an upper triangular matrix.

```
#include <ginkgo/core/solver/triangular.hpp>
```

Public Member Functions

- `std::unique_ptr< LinOp > transpose ()` const override
Returns a `LinOp` representing the transpose of the [Transposable](#) object.
- `std::unique_ptr< LinOp > conj_transpose ()` const override
Returns a `LinOp` representing the conjugate transpose of the [Transposable](#) object.
- `UpperTrs (const UpperTrs &)`
Copy-assigns a triangular solver.
- `UpperTrs (UpperTrs &&)`
Move-assigns a triangular solver.
- `UpperTrs & operator= (const UpperTrs &)`
Copy-constructs a triangular solver.
- `UpperTrs & operator= (UpperTrs &&)`
Move-constructs a triangular solver.

Static Public Member Functions

- static parameters_type [parse](#) (const [config::pnode](#) &config, const [config::registry](#) &context, const [config::type_descriptor](#) &td_for_child=config::make_type_descriptor< ValueType, IndexType >())

Create the parameters from the property_tree.

47.239.1 Detailed Description

```
template<typename ValueType = default_precision, typename IndexType = int32>
class gko::solver::UpperTrs< ValueType, IndexType >
```

[UpperTrs](#) is the triangular solver which solves the system $Ux = b$, when U is an upper triangular matrix.

It works best when passing in a matrix in CSR format. If the matrix is not in CSR, then the generate step converts it into a CSR matrix. The generation fails if the matrix is not convertible to CSR.

Note

As the constructor uses the copy and convert functionality, it is not possible to create a empty solver or a solver with a matrix in any other format other than CSR, if none of the executor modules are being compiled with.

Template Parameters

<i>ValueType</i>	precision of matrix elements
<i>IndexType</i>	precision of matrix indices

47.239.2 Constructor & Destructor Documentation

47.239.2.1 UpperTrs() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::solver::UpperTrs< ValueType, IndexType >::UpperTrs (
    const UpperTrs< ValueType, IndexType > & )
```

Copy-assigns a triangular solver.

Preserves the executor, shallow-copies the system matrix. If the executors mismatch, clones system matrix onto this executor. Solver analysis information will be regenerated.

47.239.2.2 UpperTrs() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
gko::solver::UpperTrs< ValueType, IndexType >::UpperTrs (
    UpperTrs< ValueType, IndexType > && )
```

Move-assigns a triangular solver.

Preserves the executor, moves the system matrix. If the executors mismatch, clones system matrix onto this executor and regenerates solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

47.239.3 Member Function Documentation

47.239.3.1 conj_transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::UpperTrs< ValueType, IndexType >::conj_transpose ( )
const [override], [virtual]
```

Returns a LinOp representing the conjugate transpose of the [Transposable](#) object.

Returns

a pointer to the new conjugate transposed object

Implements [gko::Transposable](#).

47.239.3.2 operator=() [1/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
UpperTrs& gko::solver::UpperTrs< ValueType, IndexType >::operator= (
    const UpperTrs< ValueType, IndexType > & )
```

Copy-constructs a triangular solver.

Preserves the executor, shallow-copies the system matrix. Solver analysis information will be regenerated.

47.239.3.3 operator=() [2/2]

```
template<typename ValueType = default_precision, typename IndexType = int32>
UpperTrs& gko::solver::UpperTrs< ValueType, IndexType >::operator= (
    UpperTrs< ValueType, IndexType > && )
```

Move-constructs a triangular solver.

Preserves the executor, moves the system matrix and solver analysis information. Moved-from object is empty (0x0 and nullptr system matrix)

47.239.3.4 parse()

```
template<typename ValueType = default_precision, typename IndexType = int32>
static parameters_type gko::solver::UpperTrs< ValueType, IndexType >::parse (
    const config::pnode & config,
    const config::registry & context,
    const config::type_descriptor & td_for_child = config::make_type_descriptor< ValueType, IndexType >() ) [static]
```

Create the parameters from the property_tree.

Because this is directly tied to the specific type, the value/index type settings within config are ignored and type_↔ descriptor is only used for children configs.

Parameters

<i>config</i>	the property tree for setting
<i>context</i>	the registry
<i>td_for_child</i>	the type descriptor for children configs. The default uses the value/index type of this class.

Returns

parameters

47.239.3.5 transpose()

```
template<typename ValueType = default_precision, typename IndexType = int32>
std::unique_ptr<LinOp> gko::solver::UpperTrs< ValueType, IndexType >::transpose ( ) const
[override], [virtual]
```

Returns a LinOp representing the transpose of the [Transposable](#) object.

Returns

a pointer to the new transposed object

Implements [gko::Transposable](#).

The documentation for this class was generated from the following file:

- ginkgo/core/solver/triangular.hpp

47.240 gko::UseComposition< ValueType > Class Template Reference

The [UseComposition](#) class can be used to store the composition information in LinOp.

```
#include <ginkgo/core/base/composition.hpp>
```

Public Member Functions

- std::shared_ptr< [Composition](#)< ValueType > > [get_composition](#) () const
Returns the composition operators.
- std::shared_ptr< const LinOp > [get_operator_at](#) (size_type index) const
Returns the operator at index-th position of composition.

47.240.1 Detailed Description

```
template<typename ValueType = default_precision>
class gko::UseComposition< ValueType >
```

The [UseComposition](#) class can be used to store the composition information in LinOp.

Template Parameters

<i>ValueType</i>	precision of input and result vectors
------------------	---------------------------------------

47.240.2 Member Function Documentation**47.240.2.1 get_composition()**

```
template<typename ValueType = default_precision>
std::shared_ptr<Composition<ValueType> > gko::UseComposition< ValueType >::get_composition (
) const [inline]
```

Returns the composition operators.

Returns

composition

47.240.2.2 get_operator_at()

```
template<typename ValueType = default_precision>
std::shared_ptr<const LinOp> gko::UseComposition< ValueType >::get_operator_at (
    size_type index ) const [inline]
```

Returns the operator at index-th position of composition.

Returns

index-th operator

Note

when this composition is not set, this function always returns nullptr. However, when this composition is set, it will throw exception when exceeding index.

Exceptions

<i>std::out_of_range</i>	if index is out of bound when composition is existed.
--------------------------	---

Referenced by `gko::multigrid::EnableMultigridLevel< ValueType >::get_coarse_op()`, `gko::multigrid::EnableMultigridLevel< ValueType >::get_prolong_op()`, and `gko::multigrid::EnableMultigridLevel< ValueType >::get_restrict_op()`.

The documentation for this class was generated from the following file:

- ginkgo/core/base/composition.hpp

47.241 gko::syn::value_list< T, Values > Struct Template Reference

[value_list](#) records several values with the same type in template.

```
#include <ginkgo/core/synthesizer/containers.hpp>
```

47.241.1 Detailed Description

```
template<typename T, T... Values>
struct gko::syn::value_list< T, Values >
```

[value_list](#) records several values with the same type in template.

Template Parameters

<i>T</i>	the value type of the list
<i>Values</i>	the values in the list

The documentation for this struct was generated from the following file:

- ginkgo/core/synthesizer/containers.hpp

47.242 gko::ValueMismatch Class Reference

[ValueMismatch](#) is thrown if two values are not equal.

```
#include <ginkgo/core/base/exception.hpp>
```

Public Member Functions

- [ValueMismatch](#) (const std::string &file, int line, const std::string &func, [size_type](#) val1, [size_type](#) val2, const std::string &clarification)

Initializes a value mismatch error.

47.242.1 Detailed Description

[ValueMismatch](#) is thrown if two values are not equal.

47.242.2 Constructor & Destructor Documentation

47.242.2.1 ValueMismatch()

```
gko::ValueMismatch::ValueMismatch (
    const std::string & file,
    int line,
    const std::string & func,
    size_type val1,
    size_type val2,
    const std::string & clarification ) [inline]
```

Initializes a value mismatch error.

Parameters

<i>file</i>	The name of the offending source file
<i>line</i>	The source code line number where the error occurred
<i>func</i>	The function name where the error occurred
<i>val1</i>	The first value to be compared.
<i>val2</i>	The second value to be compared.
<i>clarification</i>	An additional message further describing the error

The documentation for this class was generated from the following file:

- ginkgo/core/base/exception.hpp

47.243 gko::experimental::distributed::Vector< ValueType > Class Template Reference

[Vector](#) is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

```
#include <ginkgo/core/distributed/vector.hpp>
```

Public Member Functions

- void [read_distributed](#) (const [device_matrix_data](#)< ValueType, [int64](#) > &data, [ptr_param](#)< const [Partition](#)< [int64](#), [int64](#) >> partition)
Reads a vector from the [device_matrix_data](#) structure and a global row partition.
- void [read_distributed](#) (const [matrix_data](#)< ValueType, [int64](#) > &data, [ptr_param](#)< const [Partition](#)< [int64](#), [int64](#) >> partition)
Reads a vector from the [matrix_data](#) structure and a global row partition.
- std::unique_ptr< absolute_type > [compute_absolute](#) () const override
Gets the AbsoluteLinOp.
- void [compute_absolute_inplace](#) () override

- Compute absolute inplace on each element.*

 - `std::unique_ptr< complex\_type > make\_complex ()` const

Creates a complex copy of the original vectors.

 - `void make\_complex (ptr_param< complex\_type > result)` const

Writes a complex copy of the original vectors to given complex vectors.

 - `std::unique_ptr< real\_type > get\_real ()` const

Creates new real vectors and extracts the real part of the original vectors into that.

 - `void get\_real (ptr_param< real\_type > result)` const

Extracts the real part of the original vectors into given real vectors.

 - `std::unique_ptr< real\_type > get\_imag ()` const

Creates new real vectors and extracts the imaginary part of the original vectors into that.

 - `void get\_imag (ptr_param< real\_type > result)` const

Extracts the imaginary part of the original vectors into given real vectors.

 - `void fill (ValueType value)`

Fill the distributed vectors with a given value.

 - `void scale (ptr_param< const LinOp > alpha)`

Scales the vectors with a scalar (aka: BLAS scal).

 - `void inv\_scale (ptr_param< const LinOp > alpha)`

Scales the vectors with the inverse of a scalar.

 - `void add\_scaled (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)`

Adds b scaled by alpha to the vectors (aka: BLAS axpy).

 - `void sub\_scaled (ptr_param< const LinOp > alpha, ptr_param< const LinOp > b)`

Subtracts b scaled by alpha from the vectors (aka: BLAS axpy).

 - `void compute\_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result)` const

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

 - `void compute\_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp)` const

Computes the column-wise dot product of this (multi-)vector and b using a global reduction.

 - `void compute\_conj\_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result)` const

Computes the column-wise dot product of this (multi-)vector and conj(b) using a global reduction.

 - `void compute\_conj\_dot (ptr_param< const LinOp > b, ptr_param< LinOp > result, array< char > &tmp)` const

Computes the column-wise dot product of this (multi-)vector and conj(b) using a global reduction.

 - `void compute\_squared\_norm2 (ptr_param< LinOp > result)` const

Computes the square of the column-wise Euclidean (L^2) norm of this (multi-)vector using a global reduction.

 - `void compute\_squared\_norm2 (ptr_param< LinOp > result, array< char > &tmp)` const

Computes the square of the column-wise Euclidean (L^2) norm of this (multi-)vector using a global reduction.

 - `void compute\_norm2 (ptr_param< LinOp > result)` const

Computes the Euclidean (L^2) norm of this (multi-)vector using a global reduction.

 - `void compute\_norm2 (ptr_param< LinOp > result, array< char > &tmp)` const

Computes the Euclidean (L^2) norm of this (multi-)vector using a global reduction.

 - `void compute\_norm1 (ptr_param< LinOp > result)` const

Computes the column-wise (L^1) norm of this (multi-)vector.

 - `void compute\_norm1 (ptr_param< LinOp > result, array< char > &tmp)` const

Computes the column-wise (L^1) norm of this (multi-)vector using a global reduction.

 - `void compute\_mean (ptr_param< LinOp > result)` const

Computes the column-wise mean of this (multi-)vector using a global reduction.

 - `void compute\_mean (ptr_param< LinOp > result, array< char > &tmp)` const

Computes the column-wise arithmetic mean of this (multi-)vector using a global reduction.

 - `value_type & at\_local (size_type row, size_type col)` noexcept

Returns a single element of the multi-vector.

 - `value_type at\_local (size_type row, size_type col)` const noexcept

- `ValueType & at_local (size_type idx)` noexcept
Returns a single element of the multi-vector.
- `ValueType at_local (size_type idx)` const noexcept
- `value_type * get_local_values ()`
Returns a pointer to the array of local values of the multi-vector.
- `const value_type * get_const_local_values ()` const
Returns a pointer to the array of local values of the multi-vector.
- `const local_vector_type * get_local_vector ()` const
Direct (read) access to the underlying local local_vector_type vectors.
- `std::unique_ptr< const real_type > create_real_view ()` const
Create a real view of the (potentially) complex original multi-vector.
- `std::unique_ptr< real_type > create_real_view ()`
Create a real view of the (potentially) complex original multi-vector.

Static Public Member Functions

- `static std::unique_ptr< Vector > create_with_config_of (ptr_param< const Vector > other)`
Creates a distributed Vector with the same size and stride as another Vector.
- `static std::unique_ptr< Vector > create_with_type_of (ptr_param< const Vector > other, std::shared_ptr< const Executor > exec)`
Creates an empty Vector with the same type as another Vector, but on a different executor.
- `static std::unique_ptr< Vector > create_with_type_of (ptr_param< const Vector > other, std::shared_ptr< const Executor > exec, const dim< 2 > &global_size, const dim< 2 > &local_size, size_type stride)`
Creates an Vector with the same type as another Vector, but on a different executor and with a different size.
- `static std::unique_ptr< Vector > create (std::shared_ptr< const Executor > exec, mpi::communicator comm, dim< 2 > global_size, dim< 2 > local_size, size_type stride)`
Creates an empty distributed vector with a specified size.
- `static std::unique_ptr< Vector > create (std::shared_ptr< const Executor > exec, mpi::communicator comm, dim< 2 > global_size={}, dim< 2 > local_size={})`
Creates an empty distributed vector with a specified size.
- `static std::unique_ptr< Vector > create (std::shared_ptr< const Executor > exec, mpi::communicator comm, dim< 2 > global_size, std::unique_ptr< local_vector_type > local_vector)`
Creates a distributed vector from local vectors with a specified size.
- `static std::unique_ptr< Vector > create (std::shared_ptr< const Executor > exec, mpi::communicator comm, std::unique_ptr< local_vector_type > local_vector)`
Creates a distributed vector from local vectors.
- `static std::unique_ptr< const Vector > create_const (std::shared_ptr< const Executor > exec, mpi::communicator comm, dim< 2 > global_size, std::unique_ptr< const local_vector_type > local_vector)`
Creates a constant (immutable) distributed Vector from a constant local vector.
- `static std::unique_ptr< const Vector > create_const (std::shared_ptr< const Executor > exec, mpi::communicator comm, std::unique_ptr< const local_vector_type > local_vector)`
Creates a constant (immutable) distributed Vector from a constant local vector.

47.243.1 Detailed Description

```
template<typename ValueType = double>
class gko::experimental::distributed::Vector< ValueType >
```

`Vector` is a format which explicitly stores (multiple) distributed column vectors in a dense storage format.

The (multi-)vector is distributed by row, which is described by a

See also

[Partition](#). The [local](#) vectors are stored using the

Dense format. The vector should be filled using the [read_distributed](#) method, e.g.

```
auto part = Partition<...>::build_from_mapping(...);
auto vector = Vector<...>::create(exec, comm);
vector->read_distributed(matrix_data, part);
```

Using this approach the size of the global vectors, [as well as](#) the size of the [local](#) vectors, will be automatically inferred. It is possible to [create](#) a vector with specified global and [local](#) sizes and [fill](#) the [local](#) vectors using the [accessor get_local_vector](#).

Note

Operations between two vectors (axpy, dot product, etc.) are only valid if both vectors were created using the same partition.

Template Parameters

<i>ValueType</i>	The precision of vector elements.
------------------	-----------------------------------

47.243.2 Member Function Documentation

47.243.2.1 add_scaled()

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::add_scaled (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b )
```

Adds *b* scaled by *alpha* to the vectors (aka: BLAS axpy).

Parameters

<i>alpha</i>	If <i>alpha</i> is 1x1 Dense matrix, the all vectors of <i>b</i> are scaled by <i>alpha</i> . If it is a Dense row vector of values, then <i>i</i> -th column vector of <i>b</i> is scaled with the <i>i</i> -th element of <i>alpha</i> (the number of columns of <i>alpha</i> has to match the number of vectors).
<i>b</i>	a (multi-)vector of the same dimension as this

47.243.2.2 at_local() [1/4]

```
template<typename ValueType = double>
ValueType gko::experimental::distributed::Vector< ValueType >::at_local (
    size_type idx ) const [noexcept]
```

47.243.2.3 at_local() [2/4]

```
template<typename ValueType = double>
ValueType& gko::experimental::distributed::Vector< ValueType >::at_local (
    size_type idx ) [noexcept]
```

Returns a single element of the multi-vector.

Useful for iterating across all elements of the multi-vector. However, it is less efficient than the two-parameter variant of this method.

Parameters

<i>idx</i>	a linear index of the requested element (ignoring the stride)
------------	---

Note

the method has to be called on the same [Executor](#) the matrix is stored at (e.g. trying to call this method on a GPU matrix from the OMP results in a runtime error)

47.243.2.4 at_local() [3/4]

```
template<typename ValueType = double>
value_type gko::experimental::distributed::Vector< ValueType >::at_local (
    size_type row,
    size_type col ) const [noexcept]
```

47.243.2.5 at_local() [4/4]

```
template<typename ValueType = double>
value_type& gko::experimental::distributed::Vector< ValueType >::at_local (
    size_type row,
    size_type col ) [noexcept]
```

Returns a single element of the multi-vector.

Parameters

<i>row</i>	the local row of the requested element
<i>col</i>	the local column of the requested element

Note

the method has to be called on the same [Executor](#) the multi-vector is stored at (e.g. trying to call this method on a GPU multi-vector from the OMP results in a runtime error)

47.243.2.6 compute_absolute()

```
template<typename ValueType = double>
std::unique_ptr<absolute_type> gko::experimental::distributed::Vector< ValueType >::compute←
_absolute ( ) const [override], [virtual]
```

Gets the AbsoluteLinOp.

Returns

a pointer to the new absolute object

Implements [gko::EnableAbsoluteComputation< remove_complex< Vector< ValueType > > >](#).

47.243.2.7 compute_conj_dot() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_conj_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result ) const
```

Computes the column-wise dot product of this (multi-)vector and `conj (b)` using a global reduction.

Parameters

<i>b</i>	a (multi-)vector of same dimension as this
<i>result</i>	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)

47.243.2.8 compute_conj_dot() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_conj_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise dot product of this (multi-)vector and `conj (b)` using a global reduction.

Parameters

<i>b</i>	a (multi-)vector of same dimension as this
<i>result</i>	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.9 compute_dot() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result ) const
```

Computes the column-wise dot product of this (multi-)vector and **b** using a global reduction.

Parameters

<i>b</i>	a (multi-)vector of same dimension as this
<i>result</i>	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)

47.243.2.10 compute_dot() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_dot (
    ptr_param< const LinOp > b,
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise dot product of this (multi-)vector and **b** using a global reduction.

Parameters

<i>b</i>	a (multi-)vector of same dimension as this
<i>result</i>	a Dense row matrix, used to store the dot product (the number of column in result must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.11 compute_mean() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_mean (
    ptr_param< LinOp > result ) const
```

Computes the column-wise mean of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row matrix, used to store the mean (the number of columns in result must match the number of columns of this)
---------------	---

47.243.2.12 compute_mean() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_mean (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise arithmetic mean of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row matrix, used to store the mean (the number of columns in result must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.13 compute_norm1() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_norm1 (
    ptr_param< LinOp > result ) const
```

Computes the column-wise (L^1) norm of this (multi-)vector.

Parameters

<i>result</i>	a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
---------------	---

47.243.2.14 compute_norm1() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_norm1 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the column-wise (L^1) norm of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.15 compute_norm2() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_norm2 (
    ptr_param< LinOp > result ) const
```

Computes the Euclidean (L^2) norm of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
---------------	---

47.243.2.16 compute_norm2() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_norm2 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the Euclidean (L^2) norm of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row matrix, used to store the norm (the number of columns in result must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.17 compute_squared_norm2() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_squared_norm2 (
    ptr_param< LinOp > result ) const
```

Computes the square of the column-wise Euclidean (L^2) norm of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
---------------	---

47.243.2.18 compute_squared_norm2() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::compute_squared_norm2 (
    ptr_param< LinOp > result,
    array< char > & tmp ) const
```

Computes the square of the column-wise Euclidean (L^2) norm of this (multi-)vector using a global reduction.

Parameters

<i>result</i>	a Dense row vector, used to store the norm (the number of columns in the vector must match the number of columns of this)
<i>tmp</i>	the temporary storage to use for partial sums during the reduction computation. It may be resized and/or reset to the correct executor.

47.243.2.19 create() [1/4]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    dim< 2 > global_size,
    dim< 2 > local_size,
    size_type stride ) [static]
```

Creates an empty distributed vector with a specified size.

Parameters

<i>exec</i>	Executor associated with vector
<i>comm</i>	Communicator associated with vector
<i>global_size</i>	Global size of the vector
<i>local_size</i>	Processor-local size of the vector
<i>stride</i>	Stride of the local vector.

Returns

A smart pointer to the newly created vector.

47.243.2.20 create() [2/4]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    dim< 2 > global_size,
    std::unique_ptr< local_vector_type > local_vector ) [static]
```

Creates a distributed vector from local vectors with a specified size.

Note

The data from the `local_vector` will be moved into the new distributed vector. You could either move in a `std::unique_ptr` directly, copy a local vector with `gko::clone`, or create a unique non-owning view of a given local vector with `gko::make_dense_view`.

Parameters

<i>exec</i>	<code>Executor</code> associated with this vector
<i>comm</i>	Communicator associated with this vector
<i>global_size</i>	The global size of the vector
<i>local_vector</i>	The underlying local vector, the data will be moved into this

Returns

A smart pointer to the newly created vector.

47.243.2.21 create() [3/4]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    dim< 2 > global_size = {},
    dim< 2 > local_size = {} ) [static]
```

Creates an empty distributed vector with a specified size.

Parameters

<i>exec</i>	<code>Executor</code> associated with vector
<i>comm</i>	Communicator associated with vector
<i>global_size</i>	Global size of the vector
<i>local_size</i>	Processor-local size of the vector, uses <code>local_size[1]</code> as the stride

Returns

A smart pointer to the newly created vector.

47.243.2.22 create() [4/4]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    std::unique_ptr< local_vector_type > local_vector ) [static]
```

Creates a distributed vector from local vectors.

The global size will be deduced from the local sizes, which will incur a collective communication.

Note

The data from the local_vector will be moved into the new distributed vector. You could either move in a std::unique_ptr directly, copy a local vector with [gko::clone](#), or create a unique non-owning view of a given local vector with [gko::make_dense_view](#).

Parameters

<i>exec</i>	Executor associated with this vector
<i>comm</i>	Communicator associated with this vector
<i>local_vector</i>	The underlying local vector, the data will be moved into this.

Returns

A smart pointer to the newly created vector.

47.243.2.23 create_const() [1/2]

```
template<typename ValueType = double>
static std::unique_ptr<const Vector> gko::experimental::distributed::Vector< ValueType >::create_const (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    dim< 2 > global_size,
    std::unique_ptr< const local_vector_type > local_vector ) [static]
```

Creates a constant (immutable) distributed [Vector](#) from a constant local vector.

Parameters

<i>exec</i>	Executor associated with this vector
<i>comm</i>	Communicator associated with this vector
<i>global_size</i>	The global size of the vector
<i>local_vector</i>	The underlying local vector, of which a view is created

Returns

A smart pointer to the newly created vector.

47.243.2.24 create_const() [2/2]

```
template<typename ValueType = double>
static std::unique_ptr<const Vector> gko::experimental::distributed::Vector< ValueType >↵
::create_const (
    std::shared_ptr< const Executor > exec,
    mpi::communicator comm,
    std::unique_ptr< const local_vector_type > local_vector ) [static]
```

Creates a constant (immutable) distributed [Vector](#) from a constant local vector.

The global size will be deduced from the local sizes, which will incur a collective communication.

Parameters

<i>exec</i>	Executor associated with this vector
<i>comm</i>	Communicator associated with this vector
<i>local_vector</i>	The underlying local vector, of which a view is created

Returns

A smart pointer to the newly created vector.

47.243.2.25 create_real_view() [1/2]

```
template<typename ValueType = double>
std::unique_ptr<real_type> gko::experimental::distributed::Vector< ValueType >::create_real↵
_view ( )
```

Create a real view of the (potentially) complex original multi-vector.

If the original vector is real, nothing changes. If the original vector is complex, the result is created by viewing the complex vector with as real with a `reinterpret_cast` with twice the number of columns and double the stride.

47.243.2.26 create_real_view() [2/2]

```
template<typename ValueType = double>
std::unique_ptr<const real_type> gko::experimental::distributed::Vector< ValueType >::create↵
_real_view ( ) const
```

Create a real view of the (potentially) complex original multi-vector.

If the original vector is real, nothing changes. If the original vector is complex, the result is created by viewing the complex vector with as real with a `reinterpret_cast` with twice the number of columns and double the stride.

47.243.2.27 create_with_config_of()

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create_↵
with_config_of (
    ptr_param< const Vector< ValueType > > other ) [static]
```

Creates a distributed [Vector](#) with the same size and stride as another [Vector](#).

Parameters

<i>other</i>	The other vector whose configuration needs to copied.
--------------	---

47.243.2.28 create_with_type_of() [1/2]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create_↵
with_type_of (
    ptr_param< const Vector< ValueType > > other,
    std::shared_ptr< const Executor > exec ) [static]
```

Creates an empty [Vector](#) with the same type as another [Vector](#), but on a different executor.

Parameters

<i>other</i>	The other multi-vector whose type we target.
<i>exec</i>	The executor of the new multi-vector.

Note

The new multi-vector uses the same communicator as other.

Returns

an empty [Vector](#) with the type of other.

47.243.2.29 create_with_type_of() [2/2]

```
template<typename ValueType = double>
static std::unique_ptr<Vector> gko::experimental::distributed::Vector< ValueType >::create_↵
with_type_of (
    ptr_param< const Vector< ValueType > > other,
    std::shared_ptr< const Executor > exec,
    const dim< 2 > & global_size,
    const dim< 2 > & local_size,
    size_type stride ) [static]
```

Creates an [Vector](#) with the same type as another [Vector](#), but on a different executor and with a different size.

Parameters

<i>other</i>	The other multi-vector whose type we target.
<i>exec</i>	The executor of the new multi-vector.
<i>global_size</i>	The global size of the multi-vector.
<i>local_size</i>	The local size of the multi-vector.
<i>stride</i>	The stride of the new multi-vector.

Returns

a [Vector](#) of specified size with the type of other.

47.243.2.30 fill()

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::fill (
    ValueType value )
```

Fill the distributed vectors with a given value.

Parameters

<i>value</i>	the value to be filled
--------------	------------------------

47.243.2.31 get_const_local_values()

```
template<typename ValueType = double>
const value_type* gko::experimental::distributed::Vector< ValueType >::get_const_local_values
( ) const
```

Returns a pointer to the array of local values of the multi-vector.

Returns

the pointer to the array of local values

Note

This is the constant version of the function, which can be significantly more memory efficient than the non-constant version, so always prefer this version.

47.243.2.32 get_local_values()

```
template<typename ValueType = double>
value_type* gko::experimental::distributed::Vector< ValueType >::get_local_values ( )
```

Returns a pointer to the array of local values of the multi-vector.

Returns

the pointer to the array of local values

47.243.2.33 get_local_vector()

```
template<typename ValueType = double>
const local_vector_type* gko::experimental::distributed::Vector< ValueType >::get_local_vector
( ) const
```

Direct (read) access to the underlying local local_vector_type vectors.

Returns

a constant pointer to the underlying local_vector_type vectors

47.243.2.34 inv_scale()

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::inv_scale (
    ptr_param< const LinOp > alpha )
```

Scales the vectors with the inverse of a scalar.

Parameters

<i>alpha</i>	If alpha is 1x1 Dense matrix, the all vectors are scaled by 1 / alpha. If it is a Dense row vector of values, then i-th column vector is scaled with the inverse of the i-th element of alpha (the number of columns of alpha has to match the number of vectors).
--------------	--

47.243.2.35 make_complex() [1/2]

```
template<typename ValueType = double>
std::unique_ptr<complex_type> gko::experimental::distributed::Vector< ValueType >::make_↵
complex ( ) const
```

Creates a complex copy of the original vectors.

If the original vectors were real, the imaginary part of the result will be zero.

47.243.2.36 make_complex() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::make_complex (
    ptr_param< complex_type > result ) const
```

Writes a complex copy of the original vectors to given complex vectors.

If the original vectors were real, the imaginary part of the result will be zero.

47.243.2.37 read_distributed() [1/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::read_distributed (
    const device_matrix_data< ValueType, int64 > & data,
    ptr_param< const Partition< int64, int64 >> partition )
```

Reads a vector from the [device_matrix_data](#) structure and a global row partition.

The number of rows of the matrix data is ignored, only its number of columns is relevant. Both the number of local and global rows are inferred from the row partition.

Note

The matrix data can contain entries for rows other than those owned by the process. Entries for those rows are discarded.

Parameters

<i>data</i>	The device_matrix_data structure
<i>partition</i>	The global row partition

47.243.2.38 read_distributed() [2/2]

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::read_distributed (
    const matrix_data< ValueType, int64 > & data,
    ptr_param< const Partition< int64, int64 >> partition )
```

Reads a vector from the [matrix_data](#) structure and a global row partition.

See [@read_distributed](#)

Note

For efficiency it is advised to use the [device_matrix_data](#) overload.

47.243.2.39 scale()

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::scale (
    ptr_param< const LinOp > alpha )
```

Scales the vectors with a scalar (aka: BLAS scal).

Parameters

<i>alpha</i>	If alpha is 1x1 Dense matrix, the all vectors are scaled by alpha. If it is a Dense row vector of values, then i-th column vector is scaled with the i-th element of alpha (the number of columns of alpha has to match the number of vectors).
--------------	---

47.243.2.40 sub_scaled()

```
template<typename ValueType = double>
void gko::experimental::distributed::Vector< ValueType >::sub_scaled (
    ptr_param< const LinOp > alpha,
    ptr_param< const LinOp > b )
```

Subtracts *b* scaled by *alpha* from the vectors (aka: BLAS axpy).

Parameters

<i>alpha</i>	If alpha is 1x1 Dense matrix, the all vectors of <i>b</i> are scaled by alpha. If it is a Dense row vector of values, then i-th column vector of <i>b</i> is scaled with the i-th element of alpha (the number of c
<i>b</i>	a (multi-)vector of the same dimension as this

The documentation for this class was generated from the following files:

- ginkgo/core/distributed/matrix.hpp
- ginkgo/core/distributed/vector.hpp

47.244 gko::version Struct Reference

This structure is used to represent versions of various Ginkgo modules.

```
#include <ginkgo/core/base/version.hpp>
```

Public Attributes

- const uint64 *major*
The major version number.
- const uint64 *minor*
The minor version number.
- const uint64 *patch*
The patch version number.
- const char *const *tag*
Addition tag string that describes the version in more detail.

47.244.1 Detailed Description

This structure is used to represent versions of various Ginkgo modules.

Version structures can be compared using the usual relational operators.

47.244.2 Member Data Documentation

47.244.2.1 tag

```
const char* const gko::version::tag
```

Addition tag string that describes the version in more detail.

It does not participate in comparisons.

Referenced by `gko::operator<<()`.

The documentation for this struct was generated from the following file:

- `ginkgo/core/base/version.hpp`

47.245 gko::version_info Class Reference

Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:

```
#include <ginkgo/core/base/version.hpp>
```

Static Public Member Functions

- static const [version_info](#) & `get ()`
Returns an instance of [version_info](#).

Public Attributes

- [version_header_version](#)
Contains version information of the header files.
- [version_core_version](#)
Contains version information of the core library.
- [version_reference_version](#)
Contains version information of the reference module.
- [version_omp_version](#)
Contains version information of the OMP module.
- [version_cuda_version](#)
Contains version information of the CUDA module.
- [version_hip_version](#)
Contains version information of the HIP module.
- [version_dpcpp_version](#)
Contains version information of the DPC++ module.

47.245.1 Detailed Description

Ginkgo uses version numbers to label new features and to communicate backward compatibility guarantees:

1. Versions with different major version number have incompatible interfaces (parts of the earlier interface may not be present anymore, and new interfaces can appear).
2. Versions with the same major number X, but different minor numbers Y1 and Y2 numbers keep the same interface as version X.0.0, but additions to the interface in X.0.0 present in X.Y1.0 may not be present in X.Y2.0 and vice versa.
3. Versions with the same major and minor version numbers, but different patch numbers have exactly the same interface, but the functionality may be different (something that is not implemented or has a bug in an earlier version may have this implemented or fixed in a later version).

This structure provides versions of different parts of Ginkgo: the headers, the core and the kernel modules (reference, OpenMP, CUDA, HIP, DPCPP). To obtain an instance of [version_info](#) filled with information about the current version of Ginkgo, call the [version_info::get\(\)](#) static method.

47.245.2 Member Function Documentation

47.245.2.1 get()

```
static const version\_info& gko::version_info::get ( ) [inline], [static]
```

Returns an instance of [version_info](#).

Returns

an instance of version info

47.245.3 Member Data Documentation

47.245.3.1 core_version

```
version gko::version_info::core_version
```

Contains version information of the core library.

This is the version of the static/shared library called "ginkgo".

47.245.3.2 cuda_version

```
version gko::version_info::cuda_version
```

Contains version information of the CUDA module.

This is the version of the static/shared library called "ginkgo_cuda".

47.245.3.3 dpcpp_version

```
version gko::version_info::dpcpp_version
```

Contains version information of the DPC++ module.

This is the version of the static/shared library called "ginkgo_dpcpp".

47.245.3.4 hip_version

```
version gko::version_info::hip_version
```

Contains version information of the HIP module.

This is the version of the static/shared library called "ginkgo_hip".

47.245.3.5 omp_version

```
version gko::version_info::omp_version
```

Contains version information of the OMP module.

This is the version of the static/shared library called "ginkgo_omp".

47.245.3.6 reference_version

```
version gko::version_info::reference_version
```

Contains version information of the reference module.

This is the version of the static/shared library called "ginkgo_reference".

The documentation for this class was generated from the following file:

- ginkgo/core/base/version.hpp

47.246 gko::experimental::mpi::window< ValueType > Class Template Reference

This class wraps the MPI_Window class with RAII functionality.

```
#include <ginkgo/core/base/mpi.hpp>
```

Public Types

- enum [create_type](#)
The create type for the window object.
- enum [lock_type](#)
The lock type for passive target synchronization of the windows.

Public Member Functions

- [window](#) ()
The default constructor.
- [window](#) ([window](#) &&other)
The move constructor.
- [window](#) & [operator=](#) ([window](#) &&other)
The move assignment operator.
- [window](#) (std::shared_ptr< const [Executor](#) > exec, ValueType *base, int num_elems, const [communicator](#) &comm, const int disp_unit=sizeof(ValueType), MPI_Info input_info=MPI_INFO_NULL, [create_type](#) c_↵ type=create_type::create)
Create a window object with a given data pointer and type.
- MPI_Win [get_window](#) () const
Get the underlying window object of MPI_Win type.
- void [fence](#) (int assert=0) const
The active target synchronization using MPI_Win_fence for the window object.
- void [lock](#) (int rank, [lock_type](#) lock_t=lock_type::shared, int assert=0) const
Create an epoch using MPI_Win_lock for the window object.
- void [unlock](#) (int rank) const
Close the epoch using MPI_Win_unlock for the window object.
- void [lock_all](#) (int assert=0) const
Create the epoch on all ranks using MPI_Win_lock_all for the window object.
- void [unlock_all](#) () const
Close the epoch on all ranks using MPI_Win_unlock_all for the window object.
- void [flush](#) (int rank) const
Flush the existing RDMA operations on the target rank for the calling process for the window object.
- void [flush_local](#) (int rank) const
Flush the existing RDMA operations on the calling rank from the target rank for the window object.
- void [flush_all](#) () const
Flush all the existing RDMA operations for the calling process for the window object.
- void [flush_all_local](#) () const
Flush all the local existing RDMA operations on the calling rank for the window object.
- void [sync](#) () const
Synchronize the public and private buffers for the window object.
- [~window](#) ()
The deleter which calls MPI_Win_free when the window leaves its scope.
- template<typename PutType >
void [put](#) (std::shared_ptr< const [Executor](#) > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const
Put data into the target window.
- template<typename PutType >
[request](#) [r_put](#) (std::shared_ptr< const [Executor](#) > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const
Put data into the target window.

- `template<typename PutType >`
`void accumulate (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const`
Accumulate data into the target window.
- `template<typename PutType >`
`request r_accumulate (std::shared_ptr< const Executor > exec, const PutType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const`
(Non-blocking) Accumulate data into the target window.
- `template<typename GetType >`
`void get (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const`
Get data from the target window.
- `template<typename GetType >`
`request r_get (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, const int target_rank, const unsigned int target_disp, const int target_count) const`
Get data (with handle) from the target window.
- `template<typename GetType >`
`void get_accumulate (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, GetType *result_buffer, const int result_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const`
Get Accumulate data from the target window.
- `template<typename GetType >`
`request r_get_accumulate (std::shared_ptr< const Executor > exec, GetType *origin_buffer, const int origin_count, GetType *result_buffer, const int result_count, const int target_rank, const unsigned int target_disp, const int target_count, MPI_Op operation) const`
(Non-blocking) Get Accumulate data (with handle) from the target window.
- `template<typename GetType >`
`void fetch_and_op (std::shared_ptr< const Executor > exec, GetType *origin_buffer, GetType *result_buffer, const int target_rank, const unsigned int target_disp, MPI_Op operation) const`
Fetch and operate on data from the target window (An optimized version of Get_accumulate).

47.246.1 Detailed Description

```
template<typename ValueType>
class gko::experimental::mpi::window< ValueType >
```

This class wraps the `MPI_Window` class with RAII functionality.

Different create and lock type methods are setup with enums.

`MPI_Window` is primarily used for one sided communication and this class provides functionalities to fence, lock, unlock and flush the communication buffers.

47.246.2 Constructor & Destructor Documentation

47.246.2.1 window() [1/3]

```
template<typename ValueType >
gko::experimental::mpi::window< ValueType >::window ( ) [inline]
```

The default constructor.

It creates a null window of MPI_WIN_NULL type.

47.246.2.2 window() [2/3]

```
template<typename ValueType >
gko::experimental::mpi::window< ValueType >::window (
    window< ValueType > && other ) [inline]
```

The move constructor.

Move the other object and replace it with MPI_WIN_NULL

Parameters

<i>other</i>	the window object to be moved.
--------------	--------------------------------

47.246.2.3 window() [3/3]

```
template<typename ValueType >
gko::experimental::mpi::window< ValueType >::window (
    std::shared_ptr< const Executor > exec,
    ValueType * base,
    int num_elems,
    const communicator & comm,
    const int disp_unit = sizeof(ValueType),
    MPI_Info input_info = MPI_INFO_NULL,
    create_type c_type = create_type::create ) [inline]
```

Create a window object with a given data pointer and type.

A collective operation.

Parameters

<i>exec</i>	The executor, on which the base pointer is located.
<i>base</i>	the base pointer for the window object.
<i>num_elems</i>	the num_elems of type ValueType the window points to.
<i>comm</i>	the communicator whose ranks will have windows created.
<i>disp_unit</i>	the displacement from base for the window object.
<i>input_info</i>	the MPI_Info object used to set certain properties.
<i>c_type</i>	the type of creation method to use to create the window.

References `gko::experimental::mpi::communicator::get()`.

47.246.3 Member Function Documentation

47.246.3.1 `accumulate()`

```
template<typename ValueType >
template<typename PutType >
void gko::experimental::mpi::window< ValueType >::accumulate (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

Accumulate data into the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to put
<i>target_rank</i>	the rank to put the data to
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call
<i>operation</i>	the reduce operation. See <code>@MPI_Op</code>

References `gko::experimental::mpi::window< ValueType >::get_window()`.

47.246.3.2 `fence()`

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::fence (
    int assert = 0 ) const [inline]
```

The active target synchronization using `MPI_Win_fence` for the window object.

This is called on all associated ranks.

Parameters

<i>assert</i>	the optimization level. 0 is always valid.
---------------	--

47.246.3.3 fetch_and_op()

```
template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::fetch_and_op (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    GetType * result_buffer,
    const int target_rank,
    const unsigned int target_disp,
    MPI_Op operation ) const [inline]
```

Fetch and operate on data from the target window (An optimized version of Get_accumulate).

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>target_rank</i>	the rank to get the data from
<i>target_disp</i>	the displacement at the target window
<i>operation</i>	the reduce operation. See @MPI_Op

References gko::experimental::mpi::window< ValueType >::get_window().

47.246.3.4 flush()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::flush (
    int rank ) const [inline]
```

Flush the existing RDMA operations on the target rank for the calling process for the window object.

Parameters

<i>rank</i>	the target rank.
-------------	------------------

47.246.3.5 flush_local()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::flush_local (
    int rank ) const [inline]
```

Flush the existing RDMA operations on the calling rank from the target rank for the window object.

Parameters

<i>rank</i>	the target rank.
-------------	------------------

47.246.3.6 get()

```

template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::get (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]

```

Get data from the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to get
<i>target_rank</i>	the rank to get the data from
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call

References `gko::experimental::mpi::window< ValueType >::get_window()`.

47.246.3.7 get_accumulate()

```

template<typename ValueType >
template<typename GetType >
void gko::experimental::mpi::window< ValueType >::get_accumulate (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    GetType * result_buffer,
    const int result_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]

```

Get Accumulate data from the target window.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to get
<i>result_buffer</i>	the buffer to receive the target data
<i>result_count</i>	the number of elements to get
<i>target_rank</i>	the rank to get the data from
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call
<i>operation</i>	the reduce operation. See @MPI_Op

References gko::experimental::mpi::window< ValueType >::get_window().

47.246.3.8 get_window()

```
template<typename ValueType >
MPI_Win gko::experimental::mpi::window< ValueType >::get_window ( ) const [inline]
```

Get the underlying window object of MPI_Win type.

Returns

the underlying window object.

Referenced by gko::experimental::mpi::window< ValueType >::accumulate(), gko::experimental::mpi::window< ValueType >::fetch_and_op(), gko::experimental::mpi::window< ValueType >::get(), gko::experimental::mpi::window< ValueType >::get_accumulate(), gko::experimental::mpi::window< ValueType >::put(), gko::experimental::mpi::window< ValueType >::r_accumulate(), gko::experimental::mpi::window< ValueType >::r_get(), gko::experimental::mpi::window< ValueType >::r_get_accumulate(), and gko::experimental::mpi::window< ValueType >::r_put().

47.246.3.9 lock()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::lock (
    int rank,
    lock_type lock_t = lock_type::shared,
    int assert = 0 ) const [inline]
```

Create an epoch using MPI_Win_lock for the window object.

Parameters

<i>rank</i>	the target rank.
<i>lock_t</i>	the type of the lock: shared or exclusive
<i>assert</i>	the optimization level. 0 is always valid.

47.246.3.10 lock_all()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::lock_all (
    int assert = 0 ) const [inline]
```

Create the epoch on all ranks using MPI_Win_lock_all for the window object.

Parameters

<i>assert</i>	the optimization level. 0 is always valid.
---------------	--

47.246.3.11 operator=()

```
template<typename ValueType >
window& gko::experimental::mpi::window< ValueType >::operator= (
    window< ValueType > && other ) [inline]
```

The move assignment operator.

Move the other object and replace it with MPI_WIN_NULL

Parameters

<i>other</i>	the window object to be moved.
--------------	--------------------------------

47.246.3.12 put()

```
template<typename ValueType >
template<typename PutType >
void gko::experimental::mpi::window< ValueType >::put (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]
```

Put data into the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send

Parameters

<i>origin_count</i>	the number of elements to put
<i>target_rank</i>	the rank to put the data to
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call

References `gko::experimental::mpi::window< ValueType >::get_window()`.

47.246.3.13 r_accumulate()

```
template<typename ValueType >
template<typename PutType >
request gko::experimental::mpi::window< ValueType >::r_accumulate (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

(Non-blocking) Accumulate data into the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to put
<i>target_rank</i>	the rank to put the data to
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call
<i>operation</i>	the reduce operation. See @MPI_Op

Returns

the request handle for the send call

References `gko::experimental::mpi::request::get()`, and `gko::experimental::mpi::window< ValueType >::get_↵window()`.

47.246.3.14 r_get()

```
template<typename ValueType >
template<typename GetType >
```

```
request gko::experimental::mpi::window< ValueType >::r_get (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]
```

Get data (with handle) from the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to get
<i>target_rank</i>	the rank to get the data from
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call

Returns

the request handle for the send call

References `gko::experimental::mpi::request::get()`, and `gko::experimental::mpi::window< ValueType >::get_↔window()`.

47.246.3.15 r_get_accumulate()

```
template<typename ValueType >
template<typename GetType >
request gko::experimental::mpi::window< ValueType >::r_get_accumulate (
    std::shared_ptr< const Executor > exec,
    GetType * origin_buffer,
    const int origin_count,
    GetType * result_buffer,
    const int result_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count,
    MPI_Op operation ) const [inline]
```

(Non-blocking) Get Accumulate data (with handle) from the target window.

Parameters

<i>exec</i>	The executor, on which the message buffers are located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to get
<i>result_buffer</i>	the buffer to receive the target data
<i>result_count</i>	the number of elements to get
<i>target_rank</i>	the rank to get the data from
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call
<i>operation</i>	the reduce operation. See <code>@MPI_Op</code>

Returns

the request handle for the send call

References `gko::experimental::mpi::request::get()`, and `gko::experimental::mpi::window< ValueType >::get_↵window()`.

47.246.3.16 r_put()

```
template<typename ValueType >
template<typename PutType >
request gko::experimental::mpi::window< ValueType >::r_put (
    std::shared_ptr< const Executor > exec,
    const PutType * origin_buffer,
    const int origin_count,
    const int target_rank,
    const unsigned int target_disp,
    const int target_count ) const [inline]
```

Put data into the target window.

Parameters

<i>exec</i>	The executor, on which the message buffer is located.
<i>origin_buffer</i>	the buffer to send
<i>origin_count</i>	the number of elements to put
<i>target_rank</i>	the rank to put the data to
<i>target_disp</i>	the displacement at the target window
<i>target_count</i>	the request handle for the send call

Returns

the request handle for the send call

References `gko::experimental::mpi::request::get()`, and `gko::experimental::mpi::window< ValueType >::get_↵window()`.

47.246.3.17 unlock()

```
template<typename ValueType >
void gko::experimental::mpi::window< ValueType >::unlock (
    int rank ) const [inline]
```

Close the epoch using `MPI_Win_unlock` for the window object.

Parameters

<i>rank</i>	the target rank.
-------------	------------------

The documentation for this class was generated from the following file:

- ginkgo/core/base/mpi.hpp

47.247 gko::solver::workspace_traits< Solver > Struct Template Reference

Traits class providing information on the type and location of workspace vectors inside a solver.

```
#include <ginkgo/core/solver/solver_base.hpp>
```

47.247.1 Detailed Description

```
template<typename Solver>
struct gko::solver::workspace_traits< Solver >
```

Traits class providing information on the type and location of workspace vectors inside a solver.

The documentation for this struct was generated from the following file:

- ginkgo/core/solver/solver_base.hpp

47.248 gko::WritableToMatrixData< ValueType, IndexType > Class Template Reference

A LinOp implementing this interface can write its data to a [matrix_data](#) structure.

```
#include <ginkgo/core/base/lin_op.hpp>
```

Public Member Functions

- virtual void [write](#) ([matrix_data](#)< ValueType, IndexType > &data) const =0
Writes a matrix to a [matrix_data](#) structure.

47.248.1 Detailed Description

```
template<typename ValueType, typename IndexType>
class gko::WritableToMatrixData< ValueType, IndexType >
```

A LinOp implementing this interface can write its data to a [matrix_data](#) structure.

47.248.2 Member Function Documentation

47.248.2.1 write()

```
template<typename ValueType, typename IndexType>
virtual void gko::WritableToMatrixData< ValueType, IndexType >::write (
    matrix_data< ValueType, IndexType > & data ) const [pure virtual]
```

Writes a matrix to a [matrix_data](#) structure.

Parameters

<i>data</i>	the matrix_data structure
-------------	---

Implemented in [gko::matrix::Fft3](#), [gko::matrix::Fft2](#), [gko::matrix::Fft](#), [gko::matrix::Fft3](#), [gko::matrix::Fft2](#), [gko::matrix::Fft](#), [gko::matrix::Fft3](#), [gko::matrix::Fft2](#), [gko::matrix::Fft](#), [gko::matrix::Csr< ValueType, IndexType >](#), [gko::matrix::Hybrid< ValueType, IndexType >](#), [gko::preconditioner::Jacobi< ValueType, IndexType >](#), [gko::matrix::Fbcsr< ValueType, IndexType >](#), [gko::matrix::Ell< ValueType, IndexType >](#), [gko::matrix::Coo< ValueType, IndexType >](#), [gko::matrix::Sellp< ValueType, IndexType >](#), [gko::matrix::SparsityCsr< ValueType, IndexType >](#), [gko::matrix::Permutation< IndexType >](#), and [gko::matrix::ScaledPermutation< IndexType >](#).

The documentation for this class was generated from the following file:

- [ginkgo/core/base/lin_op.hpp](#)

Index

- abs
 - gko, [336](#)
- accumulate
 - gko::experimental::mpi::window< ValueType >, [1112](#)
- add_logger
 - gko::Executor, [706](#)
 - gko::log::Loggable, [851](#)
- add_scaled
 - gko::batch::MultiVector< ValueType >, [904](#)
 - gko::experimental::distributed::Vector< ValueType >, [1091](#)
 - gko::matrix::Dense< ValueType >, [593](#)
- add_scaled_identity
 - gko::batch::matrix::Csr< ValueType, IndexType >, [525](#)
 - gko::batch::matrix::Dense< ValueType >, [577](#)
 - gko::batch::matrix::Ell< ValueType, IndexType >, [661](#)
 - gko::ScaledIdentityAddable, [1013](#)
- add_value
 - gko::matrix_assembly_data< ValueType, IndexType >, [878](#)
- all_gather
 - gko::experimental::mpi::communicator, [474](#)
- all_reduce
 - gko::experimental::mpi::communicator, [476](#)
- all_to_all
 - gko::experimental::mpi::communicator, [477](#), [478](#)
- all_to_all_v
 - gko::experimental::mpi::communicator, [478](#), [479](#)
- alloc
 - gko::Executor, [706](#)
- allocation_mode
 - gko, [335](#)
- AllocationError
 - gko::AllocationError, [410](#)
- apply
 - gko::batch::matrix::Csr< ValueType, IndexType >, [526](#), [527](#)
 - gko::batch::matrix::Dense< ValueType >, [577](#), [578](#)
 - gko::batch::matrix::Ell< ValueType, IndexType >, [661](#), [662](#)
 - gko::batch::matrix::Identity< ValueType >, [788](#), [789](#)
- apply2
 - gko::matrix::Coo< ValueType, IndexType >, [511](#), [512](#)
- apply_uses_initial_guess
 - gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >, [1021](#)
 - gko::solver::Bicg< ValueType >, [440](#)
 - gko::solver::Bicgstab< ValueType >, [443](#)
 - gko::solver::Cg< ValueType >, [456](#)
 - gko::solver::Cgs< ValueType >, [459](#)
 - gko::solver::Fcg< ValueType >, [730](#)
 - gko::solver::Gcr< ValueType >, [747](#)
 - gko::solver::Gmres< ValueType >, [750](#)
 - gko::solver::Idr< ValueType >, [795](#)
 - gko::solver::Ir< ValueType >, [826](#)
 - gko::solver::Multigrid, [897](#)
- array
 - gko, [336](#)
 - gko::array< ValueType >, [415–419](#)
- as
 - gko, [337–340](#)
- as_array
 - gko::syn, [405](#)
- as_const_view
 - gko::array< ValueType >, [420](#)
- as_list
 - gko::syn, [404](#)
- as_view
 - gko::array< ValueType >, [420](#)
- assemble_rows_from_neighbors
 - gko::experimental::distributed, [380](#)
- assembly_mode
 - gko::experimental::distributed, [379](#)
- at
 - gko::batch::matrix::Dense< ValueType >, [578](#), [579](#), [581](#)
 - gko::batch::MultiVector< ValueType >, [904–906](#)
 - gko::matrix::Dense< ValueType >, [593–595](#)
- at_local
 - gko::experimental::distributed::Vector< ValueType >, [1091](#), [1092](#)
- autodetect
 - gko::precision_reduction, [975](#)
- automatic
 - gko, [336](#)
- BadDimension
 - gko::BadDimension, [430](#)
- batch_dim
 - gko::batch_dim< Dimensionality, DimensionType >, [431](#)
- BatchLinOp, [318](#)
 - EnableDefaultBatchLinOpFactory, [320](#)

- GKO_ENABLE_BATCH_LIN_OP_FACTORY, 319
- bind_to_core
 - gko::machine_topology, 859
- bind_to_cores
 - gko::machine_topology, 860
- bind_to_pu
 - gko::machine_topology, 860
- bind_to_pus
 - gko::machine_topology, 860
- block_at
 - gko::BlockOperator, 449
- BlockOperator
 - gko::BlockOperator, 449
- BlockSizeError
 - gko::BlockSizeError< IndexType >, 452
- broadcast
 - gko::experimental::mpi::communicator, 480
- build_from_contiguous
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 937
- build_from_global_size_uniform
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 938
- build_from_mapping
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 938
- build_partition_from_local_range
 - gko::experimental::distributed, 380
- build_partition_from_local_size
 - gko::experimental::distributed, 381
- build_smoother
 - gko::solver, 399, 400
- ceildiv
 - gko, 340
- check
 - gko::stop::Criterion, 522
 - gko::stop::Criterion::Updater, 1082
- clac_size
 - gko::matrix::Csr< ValueType, IndexType >::classical, 463
 - gko::matrix::Csr< ValueType, IndexType >::cusparse, 569
 - gko::matrix::Csr< ValueType, IndexType >::load_balance, 850
 - gko::matrix::Csr< ValueType, IndexType >::merge_path, 892
 - gko::matrix::Csr< ValueType, IndexType >::sparselib, 1047
 - gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1064
- clear
 - gko::array< ValueType >, 420
 - gko::index_set< IndexType >, 817
 - gko::PolymorphicObject, 968
- clone
 - gko, 341
 - gko::PolymorphicObject, 968, 969
- col_at
 - gko::matrix::Ell< ValueType, IndexType >, 672, 673
 - gko::matrix::Sellp< ValueType, IndexType >, 1032
- col_scale
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 868
- column_limit
 - gko::matrix::Hybrid< ValueType, IndexType >::column_limit, 465
- column_permute
 - gko::matrix::Csr< ValueType, IndexType >, 538
 - gko::matrix::Dense< ValueType >, 595, 596
 - gko::Permutable< IndexType >, 945
- columns
 - gko::matrix, 393
- Combination
 - gko::Combination< ValueType >, 467
- combine
 - Stopping criteria, 316
- combined
 - gko::experimental::distributed, 380
- comm_index_type
 - gko::experimental::distributed, 379
- common
 - gko::precision_reduction, 975
- communicator
 - gko::experimental::mpi::communicator, 473, 474
- compose
 - gko::matrix::Permutation< IndexType >, 949
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1015
- Composition
 - gko::Composition< ValueType >, 499, 500
- compute_absolute
 - gko::EnableAbsoluteComputation< AbsoluteLinOp >, 683
 - gko::experimental::distributed::Vector< ValueType >, 1092
 - gko::matrix::Coo< ValueType, IndexType >, 512
 - gko::matrix::Csr< ValueType, IndexType >, 539
 - gko::matrix::Dense< ValueType >, 597
 - gko::matrix::Diagonal< ValueType >, 637
 - gko::matrix::Ell< ValueType, IndexType >, 673
 - gko::matrix::Fbcsr< ValueType, IndexType >, 720
 - gko::matrix::Hybrid< ValueType, IndexType >, 769
 - gko::matrix::Sellp< ValueType, IndexType >, 1033
- compute_absolute_linop
 - gko::AbsoluteComputable, 407
 - gko::EnableAbsoluteComputation< AbsoluteLinOp >, 683
- compute_conj_dot
 - gko::batch::MultiVector< ValueType >, 906
 - gko::experimental::distributed::Vector< ValueType >, 1093
 - gko::matrix::Dense< ValueType >, 597, 598
- compute_dot
 - gko::batch::MultiVector< ValueType >, 907

- gko::experimental::distributed::Vector< ValueType
>, 1094
- gko::matrix::Dense< ValueType >, 598
- compute_ell_num_stored_elements_per_row
 - gko::matrix::Hybrid< ValueType, IndexType
>::automatic, 429
 - gko::matrix::Hybrid< ValueType, IndexType
>::column_limit, 465
 - gko::matrix::Hybrid< ValueType, IndexType
>::imbalance_bounded_limit, 806
 - gko::matrix::Hybrid< ValueType, IndexType
>::imbalance_limit, 808
 - gko::matrix::Hybrid< ValueType, IndexType
>::minimal_storage_limit, 895
 - gko::matrix::Hybrid< ValueType, IndexType
>::strategy_type, 1067
- compute_hybrid_config
 - gko::matrix::Hybrid< ValueType, IndexType
>::strategy_type, 1067
- compute_inverse
 - gko::matrix::Permutation< IndexType >, 950
 - gko::matrix::ScaledPermutation< ValueType, IndexType
>, 1015
- compute_mean
 - gko::experimental::distributed::Vector< ValueType
>, 1094, 1095
 - gko::matrix::Dense< ValueType >, 599
- compute_norm1
 - gko::experimental::distributed::Vector< ValueType
>, 1095
 - gko::matrix::Dense< ValueType >, 599, 600
- compute_norm2
 - gko::batch::MultiVector< ValueType >, 907
 - gko::experimental::distributed::Vector< ValueType
>, 1096
 - gko::matrix::Dense< ValueType >, 600
- compute_squared_norm2
 - gko::experimental::distributed::Vector< ValueType
>, 1096, 1097
 - gko::matrix::Dense< ValueType >, 601
- compute_storage_space
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 445
- concatenate
 - gko::syn, 404
- cond
 - gko::matrix_data< ValueType, IndexType >, 885
- conj
 - gko, 342
- conj_transpose
 - gko::Combination< ValueType >, 468
 - gko::Composition< ValueType >, 500
 - gko::experimental::solver::Direct< ValueType, IndexType
>, 649
 - gko::matrix::Csr< ValueType, IndexType >, 539
 - gko::matrix::Dense< ValueType >, 601, 602
 - gko::matrix::Diagonal< ValueType >, 637
 - gko::matrix::Fbcsr< ValueType, IndexType >, 720
 - gko::matrix::Fft, 733
 - gko::matrix::Fft2, 737
 - gko::matrix::Fft3, 741
 - gko::matrix::Identity< ValueType >, 791
 - gko::matrix::SparsityCsr< ValueType, IndexType
>, 1051
 - gko::preconditioner::lc< LSolverType, IndexType
>, 784
 - gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 802
 - gko::preconditioner::lsai< IsaiType, ValueType, IndexType
>, 831
 - gko::preconditioner::Jacobi< ValueType, IndexType
>, 837
 - gko::solver::Bicg< ValueType >, 440
 - gko::solver::Bicgstab< ValueType >, 443
 - gko::solver::Cg< ValueType >, 457
 - gko::solver::Cgs< ValueType >, 459
 - gko::solver::Fcg< ValueType >, 731
 - gko::solver::Gcr< ValueType >, 747
 - gko::solver::Gmres< ValueType >, 750
 - gko::solver::ldr< ValueType >, 795
 - gko::solver::lr< ValueType >, 827
 - gko::solver::LowerTrs< ValueType, IndexType >, 854
 - gko::solver::UpperTrs< ValueType, IndexType >, 1084
 - gko::Transposable, 1077
- const_view
 - gko::array< ValueType >, 420
- contains
 - gko::index_set< IndexType >, 817, 818
 - gko::matrix_assembly_data< ValueType, IndexType
>, 878
- contiguous_type
 - gko::experimental::mpi::contiguous_type, 502
- converge
 - gko::stopping_status, 1061
- convert_to
 - gko::ConvertibleTo< ResultType >, 507
 - gko::EnablePolymorphicAssignment< ConcreteType, ResultType >, 696
 - gko::matrix::Fbcsr< ValueType, IndexType >, 721
 - gko::preconditioner::Jacobi< ValueType, IndexType
>, 837
- coordinate
 - gko, 336
- copy
 - gko::Executor, 707
 - gko::matrix::Csr< ValueType, IndexType >::classical, 464
 - gko::matrix::Csr< ValueType, IndexType >::cusparse, 570
 - gko::matrix::Csr< ValueType, IndexType >::load_balance, 850
 - gko::matrix::Csr< ValueType, IndexType >::merge_path, 892
 - gko::matrix::Csr< ValueType, IndexType >::sparselib,

- 1048
- gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1065
- copy_and_convert_to
 - gko, 342–344
- copy_from
 - gko::accessor::row_major< ValueType, Dimensionality >, 1008
 - gko::Executor, 707
 - gko::PolymorphicObject, 969–971
- copy_to_host
 - gko::device_matrix_data< ValueType, IndexType >, 630
- copy_val_to_host
 - gko::Executor, 708
- core_version
 - gko::version_info, 1107
- create
 - gko::batch::log::BatchConvergence< ValueType >, 434
 - gko::batch::matrix::Csr< ValueType, IndexType >, 527, 528
 - gko::batch::matrix::Dense< ValueType >, 581, 582
 - gko::batch::matrix::Ell< ValueType, IndexType >, 662, 663
 - gko::batch::matrix::Identity< ValueType >, 789
 - gko::batch::MultiVector< ValueType >, 907, 908
 - gko::BlockOperator, 450
 - gko::CudaExecutor, 560, 561
 - gko::DpcppExecutor, 653
 - gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >, 688
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 868–873
 - gko::experimental::distributed::Vector< ValueType >, 1097–1099
 - gko::HipExecutor, 758
 - gko::log::Convergence< ValueType >, 504, 505
 - gko::log::Papi< ValueType >, 927, 928
 - gko::log::PerformanceHint, 943
 - gko::log::Record, 996
 - gko::log::Stream< ValueType >, 1069
 - gko::matrix::Coo< ValueType, IndexType >, 512, 513
 - gko::matrix::Csr< ValueType, IndexType >, 539–541
 - gko::matrix::Dense< ValueType >, 602, 603
 - gko::matrix::Diagonal< ValueType >, 637, 638
 - gko::matrix::Ell< ValueType, IndexType >, 674, 675
 - gko::matrix::Fbcsr< ValueType, IndexType >, 721–723
 - gko::matrix::Fft, 733, 734
 - gko::matrix::Fft2, 737, 738
 - gko::matrix::Fft3, 741, 742
 - gko::matrix::Hybrid< ValueType, IndexType >, 769–772
 - gko::matrix::Identity< ValueType >, 791, 792
 - gko::matrix::IdentityFactory< ValueType >, 793
 - gko::matrix::Permutation< IndexType >, 950
 - gko::matrix::RowGatherer< IndexType >, 1011
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1016
 - gko::matrix::Sellp< ValueType, IndexType >, 1033, 1034
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1051, 1052
 - gko::Perturbation< ValueType >, 955, 956
- create_const
 - gko::batch::matrix::Csr< ValueType, IndexType >, 528
 - gko::batch::matrix::Dense< ValueType >, 583
 - gko::batch::matrix::Ell< ValueType, IndexType >, 664
 - gko::batch::MultiVector< ValueType >, 908
 - gko::experimental::distributed::Vector< ValueType >, 1099, 1100
 - gko::matrix::Coo< ValueType, IndexType >, 514
 - gko::matrix::Csr< ValueType, IndexType >, 541
 - gko::matrix::Dense< ValueType >, 603
 - gko::matrix::Diagonal< ValueType >, 638
 - gko::matrix::Ell< ValueType, IndexType >, 675
 - gko::matrix::Fbcsr< ValueType, IndexType >, 723
 - gko::matrix::Permutation< IndexType >, 951
 - gko::matrix::RowGatherer< IndexType >, 1012
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1017
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1053
- create_const_view_for_item
 - gko::batch::matrix::Csr< ValueType, IndexType >, 529
 - gko::batch::matrix::Dense< ValueType >, 583
 - gko::batch::matrix::Ell< ValueType, IndexType >, 664
 - gko::batch::MultiVector< ValueType >, 909
- create_const_view_of
 - gko::matrix::Dense< ValueType >, 604
- create_default
 - gko::PolymorphicObject, 971, 972
- create_for_executor
 - gko::Timer, 1075
- create_from_combined_lu
 - gko::experimental::factorization::Factorization< ValueType, IndexType >, 715
- create_from_composition
 - gko::experimental::factorization::Factorization< ValueType, IndexType >, 716
- create_from_host
 - gko::device_matrix_data< ValueType, IndexType >, 631
- create_from_offsets
 - gko::segmented_array< T >, 1027
- create_from_sizes
 - gko::segmented_array< T >, 1027, 1028

- create_from_symm_composition
 - gko::experimental::factorization::Factorization<
 - ValueType, IndexType >, 716
- create_nested_summary
 - gko::log::ProfilerHook, 979
- create_nvtx
 - gko::log::ProfilerHook, 979
- create_real_view
 - gko::experimental::distributed::Vector<
 - ValueType >, 1100
 - gko::matrix::Dense<
 - ValueType >, 604
- create_scalar_csv_writer
 - gko::log::SolverProgress, 1041
- create_scalar_table_writer
 - gko::log::SolverProgress, 1042
- create_submatrix
 - gko::matrix::Csr<
 - ValueType, IndexType >, 542
 - gko::matrix::Dense<
 - ValueType >, 605
- create_summary
 - gko::log::ProfilerHook, 980
- create_tau
 - gko::log::ProfilerHook, 980
- create_time_point
 - gko::Timer, 1075
- create_vector_storage
 - gko::log::SolverProgress, 1042
- create_view_for_item
 - gko::batch::matrix::Csr<
 - ValueType, IndexType >, 529
 - gko::batch::matrix::Dense<
 - ValueType >, 584
 - gko::batch::matrix::Ell<
 - ValueType, IndexType >, 665
 - gko::batch::MultiVector<
 - ValueType >, 909
- create_view_of
 - gko::matrix::Dense<
 - ValueType >, 605
- create_with_config_of
 - gko::batch::MultiVector<
 - ValueType >, 910
 - gko::experimental::distributed::Vector<
 - ValueType >, 1100
 - gko::matrix::Dense<
 - ValueType >, 606
- create_with_type_of
 - gko::experimental::distributed::Vector<
 - ValueType >, 1101
 - gko::matrix::Dense<
 - ValueType >, 606, 607
- criterion
 - gko::log, 390
- Csr
 - gko::matrix::Csr<
 - ValueType, IndexType >, 538
- CublasError
 - gko::CublasError, 556
- CUDA Executor, 287
- cuda_stream
 - gko::cuda_stream, 557
- cuda_version
 - gko::version_info, 1107
- CudaError
 - gko::CudaError, 558
- CufftError
 - gko::CufftError, 567
- CurandError
 - gko::CurandError, 568
- CusparsError
 - gko::CusparsError, 571
- cycle
 - gko::solver::multigrid, 401
- deferred_factory_parameter
 - gko::deferred_factory_parameter<
 - FactoryType >, 573
- Dense
 - gko::matrix::Dense<
 - ValueType >, 592, 593
- device_matrix_data
 - gko::device_matrix_data<
 - ValueType, IndexType >, 629, 630
- diag
 - gko::matrix_data<
 - ValueType, IndexType >, 886–888
- difference
 - gko::Timer, 1075
- difference_async
 - gko::CpuTimer, 520
 - gko::CudaTimer, 566
 - gko::DpcppTimer, 658
 - gko::HipTimer, 766
 - gko::Timer, 1076
- dim
 - gko::dim<
 - Dimensionality, DimensionType >, 644
- DimensionMismatch
 - gko::DimensionMismatch, 648
- DPC++ Executor, 288
- dpcpp_version
 - gko::version_info, 1108
- Ell
 - gko::matrix::Ell<
 - ValueType, IndexType >, 672
- ell_col_at
 - gko::matrix::Hybrid<
 - ValueType, IndexType >, 772
- ell_val_at
 - gko::matrix::Hybrid<
 - ValueType, IndexType >, 773
- emplace
 - gko::config::registry, 1001
- empty_out
 - gko::device_matrix_data<
 - ValueType, IndexType >, 631
- EnableDefaultBatchLinOpFactory
 - BatchLinOp, 320
- EnableDefaultCriterionFactory
 - gko::stop, 403
- EnableDefaultLinOpFactory
 - Linear Operators, 304
- EnableDefaultReorderingBaseFactory
 - gko::reorder, 396
- EnableIterativeBase
 - gko::solver::EnableIterativeBase<
 - DerivedType >, 690
- EnablePreconditionable

- gko::solver::EnablePreconditionable< DerivedType
>, 698
- EnableSolverBase
 - gko::solver::EnableSolverBase< DerivedType, Ma-
trixType >, 701
- environment
 - gko::experimental::mpi::environment, 702
- Error
 - gko::Error, 703
- executor_deleter
 - gko::executor_deleter< T >, 713
- Executors, 289
 - GKO_REGISTER_HOST_OPERATION, 290
 - GKO_REGISTER_OPERATION, 291
- extract_diagonal
 - gko::DiagonalExtractable< ValueType >, 641
 - gko::matrix::Coo< ValueType, IndexType >, 514
 - gko::matrix::Csr< ValueType, IndexType >, 543
 - gko::matrix::Dense< ValueType >, 607, 608
 - gko::matrix::Ell< ValueType, IndexType >, 676
 - gko::matrix::Fbcsr< ValueType, IndexType >, 724
 - gko::matrix::Hybrid< ValueType, IndexType >, 774
 - gko::matrix::Sellp< ValueType, IndexType >, 1034
- extract_diagonal_linop
 - gko::DiagonalExtractable< ValueType >, 641
 - gko::DiagonalLinOpExtractable, 642
- Factorizations, 293
- factory
 - gko::log, 390
- Fbcsr
 - gko::matrix::Fbcsr< ValueType, IndexType >, 720
- fence
 - gko::experimental::mpi::window< ValueType >,
1112
- fetch_and_op
 - gko::experimental::mpi::window< ValueType >,
1113
- fill
 - gko::array< ValueType >, 421
 - gko::batch::MultiVector< ValueType >, 910
 - gko::experimental::distributed::Vector< ValueType
>, 1102
 - gko::matrix::Dense< ValueType >, 608
- flush
 - gko::experimental::mpi::window< ValueType >,
1113
- flush_local
 - gko::experimental::mpi::window< ValueType >,
1113
- free
 - gko::Executor, 708
- gather
 - gko::experimental::mpi::communicator, 480
- gather_v
 - gko::experimental::mpi::communicator, 481
- generate
 - gko::AbstractFactory< AbstractProductType, Com-
ponentsType >, 409
 - gko::experimental::factorization::Cholesky< Value-
Type, IndexType >, 461
 - gko::experimental::factorization::Lu< ValueType,
IndexType >, 857
 - gko::experimental::reorder::Amd< IndexType >,
411
 - gko::experimental::reorder::Mc64< ValueType, In-
dexType >, 891
 - gko::experimental::reorder::Rcm< IndexType >,
991
 - gko::preconditioner::GaussSeidel< ValueType, In-
dexType >, 745
 - gko::preconditioner::Sor< ValueType, IndexType
>, 1044
- get
 - gko::config::pnode, 964
 - gko::cuda_stream, 557
 - gko::experimental::mpi::communicator, 481
 - gko::experimental::mpi::contiguous_type, 503
 - gko::experimental::mpi::request, 1004
 - gko::experimental::mpi::status, 1059
 - gko::experimental::mpi::window< ValueType >,
1114
 - gko::hip_stream, 754
 - gko::log::Record, 997
 - gko::ptr_param< T >, 983
 - gko::version_info, 1107
- get_accessor
 - gko::range< Accessor >, 987
- get_accumulate
 - gko::experimental::mpi::window< ValueType >,
1114
- get_agg
 - gko::multigrid::Pgm< ValueType, IndexType >, 958
- get_approximate_inverse
 - gko::preconditioner::Isai< IsaiType, ValueType, In-
dexType >, 831
- get_array
 - gko::config::pnode, 964
- get_basis
 - gko::Perturbation< ValueType >, 956
- get_blas_handle
 - gko::CudaExecutor, 561
 - gko::HipExecutor, 758
- get_block_offset
 - gko::preconditioner::block_interleaved_storage_scheme<
IndexType >, 446
- get_block_size
 - gko::BlockOperator, 450
 - gko::matrix::Fbcsr< ValueType, IndexType >, 724
- get_blocks
 - gko::preconditioner::Jacobi< ValueType, Index-
Type >, 838
- get_boolean
 - gko::config::pnode, 965
- get_closest_numa

- gko::CudaExecutor, [561](#)
- gko::HipExecutor, [759](#)
- get_closest_pus
 - gko::CudaExecutor, [561](#)
 - gko::HipExecutor, [759](#)
- get_coarse_op
 - gko::multigrid::EnableMultigridLevel< ValueType >, [693](#)
 - gko::multigrid::MultigridLevel, [901](#)
- get_coarsest_solver
 - gko::solver::Multigrid, [898](#)
- get_coefficients
 - gko::Combination< ValueType >, [468](#)
- get_col_idxs
 - gko::batch::matrix::Csr< ValueType, IndexType >, [530](#)
 - gko::batch::matrix::Ell< ValueType, IndexType >, [665](#)
 - gko::device_matrix_data< ValueType, IndexType >, [631](#)
 - gko::matrix::Coo< ValueType, IndexType >, [515](#)
 - gko::matrix::Csr< ValueType, IndexType >, [543](#)
 - gko::matrix::Ell< ValueType, IndexType >, [676](#)
 - gko::matrix::Fbcsr< ValueType, IndexType >, [724](#)
 - gko::matrix::Sellp< ValueType, IndexType >, [1035](#)
 - gko::matrix::SparsityCsr< ValueType, IndexType >, [1053](#)
- get_col_idxs_for_item
 - gko::batch::matrix::Ell< ValueType, IndexType >, [665](#)
- get_common_size
 - gko::batch::MultiVector< ValueType >, [910](#)
 - gko::batch_dim< Dimensionality, DimensionType >, [432](#)
- get_communicator
 - gko::experimental::distributed::DistributedBase, [651](#)
- get_complex_subspace
 - gko::solver::ldr< ValueType >, [795](#)
- get_composition
 - gko::UseComposition< ValueType >, [1086](#)
- get_conditioning
 - gko::preconditioner::Jacobi< ValueType, IndexType >, [838](#)
- get_const_agg
 - gko::multigrid::Pgm< ValueType, IndexType >, [959](#)
- get_const_block_pointers
 - gko::batch::preconditioner::Jacobi< ValueType, IndexType >, [842](#)
- get_const_blocks
 - gko::batch::preconditioner::Jacobi< ValueType, IndexType >, [843](#)
- get_const_blocks_cumulative_offsets
 - gko::batch::preconditioner::Jacobi< ValueType, IndexType >, [843](#)
- get_const_col_idxs
 - gko::batch::matrix::Csr< ValueType, IndexType >, [530](#)
- gko::batch::matrix::Ell< ValueType, IndexType >, [666](#)
- gko::device_matrix_data< ValueType, IndexType >, [631](#)
- gko::matrix::Coo< ValueType, IndexType >, [515](#)
- gko::matrix::Csr< ValueType, IndexType >, [543](#)
- gko::matrix::Ell< ValueType, IndexType >, [677](#)
- gko::matrix::Fbcsr< ValueType, IndexType >, [724](#)
- gko::matrix::Sellp< ValueType, IndexType >, [1035](#)
- gko::matrix::SparsityCsr< ValueType, IndexType >, [1054](#)
- get_const_col_idxs_for_item
 - gko::batch::matrix::Ell< ValueType, IndexType >, [666](#)
- get_const_coo_col_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, [774](#)
- get_const_coo_row_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, [774](#)
- get_const_coo_values
 - gko::matrix::Hybrid< ValueType, IndexType >, [775](#)
- get_const_data
 - gko::array< ValueType >, [421](#)
- get_const_ell_col_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, [775](#)
- get_const_ell_values
 - gko::matrix::Hybrid< ValueType, IndexType >, [775](#)
- get_const_flat_data
 - gko::segmented_array< T >, [1028](#)
- get_const_local_values
 - gko::experimental::distributed::Vector< ValueType >, [1102](#)
- get_const_map_block_to_row
 - gko::batch::preconditioner::Jacobi< ValueType, IndexType >, [843](#)
- get_const_permutation
 - gko::matrix::Permutation< IndexType >, [952](#)
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, [1017](#)
- get_const_row_idxs
 - gko::device_matrix_data< ValueType, IndexType >, [632](#)
 - gko::matrix::Coo< ValueType, IndexType >, [515](#)
 - gko::matrix::RowGatherer< IndexType >, [1012](#)
- get_const_row_ptrs
 - gko::batch::matrix::Csr< ValueType, IndexType >, [530](#)
 - gko::matrix::Csr< ValueType, IndexType >, [544](#)
 - gko::matrix::Fbcsr< ValueType, IndexType >, [725](#)
 - gko::matrix::SparsityCsr< ValueType, IndexType >, [1054](#)
- get_const_scaling_factors
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, [1018](#)
- get_const_slice_lengths
 - gko::matrix::Sellp< ValueType, IndexType >, [1035](#)
- get_const_slice_sets
 - gko::matrix::Sellp< ValueType, IndexType >, [1035](#)
- get_const_srow

- gko::matrix::Csr< ValueType, IndexType >, 544
- get_const_value
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1054
- get_const_values
 - gko::batch::matrix::Csr< ValueType, IndexType >, 531
 - gko::batch::matrix::Dense< ValueType >, 584
 - gko::batch::matrix::Ell< ValueType, IndexType >, 667
 - gko::batch::MultiVector< ValueType >, 911
 - gko::device_matrix_data< ValueType, IndexType >, 632
 - gko::matrix::Coo< ValueType, IndexType >, 516
 - gko::matrix::Csr< ValueType, IndexType >, 544
 - gko::matrix::Dense< ValueType >, 608
 - gko::matrix::Diagonal< ValueType >, 639
 - gko::matrix::Ell< ValueType, IndexType >, 677
 - gko::matrix::Fbcsr< ValueType, IndexType >, 725
 - gko::matrix::Sellp< ValueType, IndexType >, 1036
- get_const_values_for_item
 - gko::batch::matrix::Csr< ValueType, IndexType >, 531
 - gko::batch::matrix::Dense< ValueType >, 584
 - gko::batch::matrix::Ell< ValueType, IndexType >, 667
 - gko::batch::MultiVector< ValueType >, 911
- get_coo
 - gko::matrix::Hybrid< ValueType, IndexType >, 776
- get_coo_col_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, 776
- get_coo_nnz
 - gko::matrix::Hybrid< ValueType, IndexType >::strategy_type, 1067
- get_coo_num_stored_elements
 - gko::matrix::Hybrid< ValueType, IndexType >, 776
- get_coo_row_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, 777
- get_coo_values
 - gko::matrix::Hybrid< ValueType, IndexType >, 777
- get_core
 - gko::machine_topology, 862
- get_count
 - gko::experimental::mpi::status, 1059
- get_cublas_handle
 - gko::CudaExecutor, 562
- get_cumulative_offset
 - gko::batch::matrix::Dense< ValueType >, 585
 - gko::batch::MultiVector< ValueType >, 912
- get_cuspars_handle
 - gko::CudaExecutor, 562
- get_cycle
 - gko::solver::Multigrid, 898
- get_data
 - gko::array< ValueType >, 422
- get_description
 - gko::CudaExecutor, 562
 - gko::DpcppExecutor, 653
 - gko::Executor, 709
 - gko::HipExecutor, 759
 - gko::OmpExecutor, 920
 - gko::ReferenceExecutor, 998
- get_deterministic
 - gko::solver::ldr< ValueType >, 795
- get_device_id
 - gko::DpcppExecutor, 653
- get_device_type
 - gko::DpcppExecutor, 654
- get_dynamic_type
 - gko::name_demangling, 394
- get_ell
 - gko::matrix::Hybrid< ValueType, IndexType >, 777
- get_ell_col_idxs
 - gko::matrix::Hybrid< ValueType, IndexType >, 777
- get_ell_num_stored_elements
 - gko::matrix::Hybrid< ValueType, IndexType >, 778
- get_ell_num_stored_elements_per_row
 - gko::matrix::Hybrid< ValueType, IndexType >, 778
- get_ell_strategy_type
 - gko::matrix::Hybrid< ValueType, IndexType >::strategy_type, 1068
- get_ell_stride
 - gko::matrix::Hybrid< ValueType, IndexType >, 778
- get_ell_values
 - gko::matrix::Hybrid< ValueType, IndexType >, 778
- get_executor
 - gko::array< ValueType >, 423
 - gko::device_matrix_data< ValueType, IndexType >, 632
 - gko::index_set< IndexType >, 818
 - gko::PolymorphicObject, 972
 - gko::segmented_array< T >, 1028
- get_fine_op
 - gko::multigrid::EnableMultigridLevel< ValueType >, 693
 - gko::multigrid::MultigridLevel, 901
- get_flat_data
 - gko::segmented_array< T >, 1028
- get_global_block_offset
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 446
- get_global_index
 - gko::index_set< IndexType >, 818
- get_group_offset
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 446
- get_group_size
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 447
- get_handle
 - gko::log::Papi< ValueType >, 928
- get_handle_name
 - gko::log::Papi< ValueType >, 928
- get_hipblas_handle
 - gko::HipExecutor, 759
- get_hipspars_handle
 - gko::HipExecutor, 760

get_id
 gko::stopping_status, 1061
 get_implicit_sq_resnorm
 gko::log::Convergence< ValueType >, 505
 get_instance
 gko::machine_topology, 862
 get_integer
 gko::config::pnode, 965
 get_inverse_permutation
 gko::reorder::Rcm< ValueType, IndexType >, 992
 get_kappa
 gko::solver::ldr< ValueType >, 796
 get_krylov_dim
 gko::solver::CbGmres< ValueType >, 453
 gko::solver::Gcr< ValueType >, 748
 gko::solver::Gmres< ValueType >, 751
 get_l_solver
 gko::preconditioner::lc< LSolverType, IndexType >, 784
 gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 802
 get_lh_solver
 gko::preconditioner::lc< LSolverType, IndexType >, 784
 get_local_index
 gko::index_set< IndexType >, 819
 get_local_matrix
 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 873
 get_local_values
 gko::experimental::distributed::Vector< ValueType >, 1102
 get_local_vector
 gko::experimental::distributed::Vector< ValueType >, 1103
 get_loggers
 gko::log::Loggable, 852
 get_map
 gko::config::pnode, 965
 get_master
 gko::CudaExecutor, 562, 563
 gko::DpcppExecutor, 654
 gko::Executor, 709
 gko::HipExecutor, 760
 gko::OmpExecutor, 920
 get_max_block_size
 gko::batch::preconditioner::Jacobi< ValueType, IndexType >, 844
 get_max_iterations
 gko::batch::solver::BatchSolver, 437
 get_max_subgroup_size
 gko::DpcppExecutor, 654
 get_max_workgroup_size
 gko::DpcppExecutor, 655
 get_max_workitem_sizes
 gko::DpcppExecutor, 655
 get_mg_level_list
 gko::solver::Multigrid, 898
 get_mid_smoother_list
 gko::solver::Multigrid, 898
 get_name
 gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1065
 gko::Operation, 924
 get_non_local_matrix
 gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 873
 get_nonpreserving
 gko::precision_reduction, 975
 get_num_batch_items
 gko::batch::MultiVector< ValueType >, 912
 gko::batch_dim< Dimensionality, DimensionType >, 432
 get_num_block_cols
 gko::matrix::Fbcsr< ValueType, IndexType >, 725
 get_num_block_rows
 gko::matrix::Fbcsr< ValueType, IndexType >, 726
 get_num_blocks
 gko::batch::preconditioner::Jacobi< ValueType, IndexType >, 844
 gko::preconditioner::Jacobi< ValueType, IndexType >, 838
 get_num_columns
 gko::matrix::Hybrid< ValueType, IndexType >::column_limit, 466
 get_num_computing_units
 gko::DpcppExecutor, 655
 get_num_cores
 gko::machine_topology, 862
 get_num_devices
 gko::DpcppExecutor, 655
 get_num_elements_per_item
 gko::batch::matrix::Csr< ValueType, IndexType >, 532
 gko::batch::matrix::Dense< ValueType >, 585
 gko::batch::matrix::Ell< ValueType, IndexType >, 668
 get_num_elems
 gko::array< ValueType >, 423
 gko::device_matrix_data< ValueType, IndexType >, 633
 gko::index_set< IndexType >, 819
 get_num_empty_parts
 gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 939
 get_num_iterations
 gko::batch::log::BatchConvergence< ValueType >, 435
 gko::log::Convergence< ValueType >, 505
 get_num_nonzeros
 gko::matrix::SparsityCsr< ValueType, IndexType >, 1055
 get_num_numas
 gko::machine_topology, 862
 get_num_parts

- gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 939
- get_num_pci_devices
 - gko::machine_topology, 863
- get_num_pus
 - gko::machine_topology, 863
- get_num_ranges
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 939
- get_num_srow_elements
 - gko::matrix::Csr< ValueType, IndexType >, 545
- get_num_stored_blocks
 - gko::matrix::Fbcsr< ValueType, IndexType >, 726
- get_num_stored_elements
 - gko::batch::matrix::Csr< ValueType, IndexType >, 532
 - gko::batch::matrix::Dense< ValueType >, 585
 - gko::batch::matrix::Ell< ValueType, IndexType >, 668
 - gko::batch::MultiVector< ValueType >, 912
 - gko::batch::preconditioner::Jacobi< ValueType, IndexType >, 844
 - gko::device_matrix_data< ValueType, IndexType >, 633
 - gko::matrix::Coo< ValueType, IndexType >, 516
 - gko::matrix::Csr< ValueType, IndexType >, 545
 - gko::matrix::Dense< ValueType >, 609
 - gko::matrix::Ell< ValueType, IndexType >, 677
 - gko::matrix::Fbcsr< ValueType, IndexType >, 726
 - gko::matrix::Hybrid< ValueType, IndexType >, 779
 - gko::matrix::Sellp< ValueType, IndexType >, 1036
 - gko::matrix_assembly_data< ValueType, IndexType >, 879
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 838
- get_num_stored_elements_per_row
 - gko::batch::matrix::Ell< ValueType, IndexType >, 668
 - gko::matrix::Ell< ValueType, IndexType >, 678
- get_num_subsets
 - gko::index_set< IndexType >, 820
- get_offsets
 - gko::segmented_array< T >, 1029
- get_operator_at
 - gko::UseComposition< ValueType >, 1086
- get_operators
 - gko::Combination< ValueType >, 468
 - gko::Composition< ValueType >, 500
- get_ordered_data
 - gko::matrix_assembly_data< ValueType, IndexType >, 879
- get_parameters
 - gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >, 688
 - gko::experimental::factorization::Cholesky< ValueType, IndexType >, 461, 462
 - gko::experimental::factorization::Lu< ValueType, IndexType >, 857
- gko::experimental::reorder::Amd< IndexType >, 411
- gko::experimental::reorder::Mc64< ValueType, IndexType >, 891
- gko::experimental::reorder::Rcm< IndexType >, 991
- gko::preconditioner::GaussSeidel< ValueType, IndexType >, 746
- gko::preconditioner::Sor< ValueType, IndexType >, 1044
- get_part_ids
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 939
- get_part_size
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 940
- get_part_sizes
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 940
- get_pci_device
 - gko::machine_topology, 863, 864
- get_percentage
 - gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit, 806
 - gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit, 808
 - gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit, 895
- get_permutation
 - gko::matrix::Permutation< IndexType >, 952
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1018
 - gko::reorder::Rcm< ValueType, IndexType >, 992
- get_permutation_size
 - gko::matrix::Permutation< IndexType >, 952
- get_post_smoother_list
 - gko::solver::Multigrid, 899
- get_pre_smoother_list
 - gko::solver::Multigrid, 899
- get_preconditioner
 - gko::batch::solver::BatchSolver, 437
 - gko::Preconditionable, 977
- get_preserving
 - gko::precision_reduction, 976
- get_projector
 - gko::Perturbation< ValueType >, 957
- get_prolong_op
 - gko::multigrid::EnableMultigridLevel< ValueType >, 693
 - gko::multigrid::MultigridLevel, 901
- get_provided_thread_support
 - gko::experimental::mpi::environment, 702
- get_pu
 - gko::machine_topology, 864
- get_range_bounds
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 941

- get_range_starting_indices
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 941
- get_ratio
 - gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit, 806
- get_real
 - gko::config::pnode, 965
- get_remote_global_idxs
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 811
- get_remote_local_idxs
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 811
- get_remote_target_ids
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 812
- get_residual
 - gko::log::Convergence< ValueType >, 506
- get_residual_norm
 - gko::batch::log::BatchConvergence< ValueType >, 435
 - gko::log::Convergence< ValueType >, 506
- get_restrict_op
 - gko::multigrid::EnableMultigridLevel< ValueType >, 694
 - gko::multigrid::MultigridLevel, 901
- get_row_idxs
 - gko::device_matrix_data< ValueType, IndexType >, 633
 - gko::matrix::Coo< ValueType, IndexType >, 516
 - gko::matrix::RowGatherer< IndexType >, 1013
- get_row_ptrs
 - gko::batch::matrix::Csr< ValueType, IndexType >, 532
 - gko::matrix::Csr< ValueType, IndexType >, 545
 - gko::matrix::Fbcsr< ValueType, IndexType >, 726
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1055
- get_scalar
 - gko::Perturbation< ValueType >, 957
- get_scaling_factors
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1018
- get_segment_count
 - gko::segmented_array< T >, 1029
- get_significant_bit
 - gko, 345
- get_size
 - gko::array< ValueType >, 423
 - gko::batch::MultiVector< ValueType >, 913
 - gko::device_matrix_data< ValueType, IndexType >, 634
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 941
 - gko::index_set< IndexType >, 820
 - gko::matrix_assembly_data< ValueType, IndexType >, 879
 - gko::segmented_array< T >, 1029
- get_slice_lengths
 - gko::matrix::Sellp< ValueType, IndexType >, 1036
- get_slice_sets
 - gko::matrix::Sellp< ValueType, IndexType >, 1037
- get_slice_size
 - gko::matrix::Sellp< ValueType, IndexType >, 1037
- get_solver
 - gko::solver::lr< ValueType >, 827
- get_sparselib_handle
 - gko::CudaExecutor, 563
 - gko::HipExecutor, 760
- get_srow
 - gko::matrix::Csr< ValueType, IndexType >, 546
- get_static_type
 - gko::name_demangling, 394
- get_stop_criterion_factory
 - gko::solver::IterativeBase, 834
- get_storage_precision
 - gko::solver::CbGmres< ValueType >, 453
- get_storage_scheme
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 839
- get_strategy
 - gko::matrix::Csr< ValueType, IndexType >, 546
 - gko::matrix::Hybrid< ValueType, IndexType >, 779
- get_stream
 - gko::CudaExecutor, 563
- get_stride
 - gko::matrix::Dense< ValueType >, 609
 - gko::matrix::Ell< ValueType, IndexType >, 678
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 447
- get_stride_factor
 - gko::matrix::Sellp< ValueType, IndexType >, 1037
- get_string
 - gko::config::pnode, 966
- get_subgroup_sizes
 - gko::DpcppExecutor, 656
- get_subsets_begin
 - gko::index_set< IndexType >, 820
- get_subsets_end
 - gko::index_set< IndexType >, 820
- get_subspace_dim
 - gko::solver::ldr< ValueType >, 796
- get_superior_power
 - gko, 345
- get_superset_indices
 - gko::index_set< IndexType >, 821
- get_system_matrix
 - gko::batch::solver::BatchSolver, 437
 - gko::multigrid::FixedCoarsening< ValueType, IndexType >, 744
 - gko::multigrid::Pgm< ValueType, IndexType >, 959
- get_tag
 - gko::config::pnode, 966
- get_tolerance
 - gko::batch::solver::BatchSolver, 437

- get_tolerance_type
 - gko::batch::solver::BatchSolver, [438](#)
- get_total_cols
 - gko::matrix::Sellp< ValueType, IndexType >, [1037](#)
- get_u_solver
 - gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, [803](#)
- get_value
 - gko::matrix::SparsityCsr< ValueType, IndexType >, [1055](#)
 - gko::matrix::assembly_data< ValueType, IndexType >, [879](#)
- get_values
 - gko::batch::matrix::Csr< ValueType, IndexType >, [533](#)
 - gko::batch::matrix::Dense< ValueType >, [586](#)
 - gko::batch::matrix::Ell< ValueType, IndexType >, [669](#)
 - gko::batch::MultiVector< ValueType >, [913](#)
 - gko::device_matrix_data< ValueType, IndexType >, [634](#)
 - gko::matrix::Coo< ValueType, IndexType >, [517](#)
 - gko::matrix::Csr< ValueType, IndexType >, [546](#)
 - gko::matrix::Dense< ValueType >, [609](#)
 - gko::matrix::Diagonal< ValueType >, [639](#)
 - gko::matrix::Ell< ValueType, IndexType >, [678](#)
 - gko::matrix::Fbcsr< ValueType, IndexType >, [727](#)
 - gko::matrix::Sellp< ValueType, IndexType >, [1038](#)
- get_values_for_item
 - gko::batch::matrix::Csr< ValueType, IndexType >, [533](#)
 - gko::batch::matrix::Dense< ValueType >, [586](#)
 - gko::batch::matrix::Ell< ValueType, IndexType >, [669](#)
 - gko::batch::MultiVector< ValueType >, [913](#)
- get_walltime
 - gko::experimental::mpi, [386](#)
- get_window
 - gko::experimental::mpi::window< ValueType >, [1115](#)
- give
 - gko, [346](#)
- gko, [321](#)
 - abs, [336](#)
 - allocation_mode, [335](#)
 - array, [336](#)
 - as, [337–340](#)
 - automatic, [336](#)
 - ceildiv, [340](#)
 - clone, [341](#)
 - conj, [342](#)
 - coordinate, [336](#)
 - copy_and_convert_to, [342–344](#)
 - get_significant_bit, [345](#)
 - get_superior_power, [345](#)
 - give, [346](#)
 - highest_precision, [333](#)
 - imag, [346](#)
 - is_complex, [347](#)
 - is_complex_or_scalar, [347](#)
 - is_complex_or_scalar_s, [333](#)
 - is_complex_s, [334](#)
 - is_finite, [348](#)
 - is_nan, [349](#)
 - is_nonzero, [350](#)
 - is_zero, [350](#)
 - layout_type, [335](#)
 - lend, [351](#)
 - log_propagation_mode, [336](#)
 - make_array_view, [352](#)
 - make_const_array_view, [352](#)
 - make_const_dense_view, [353](#)
 - make_dense_view, [353](#)
 - make_temporary_clone, [354](#)
 - make_temporary_conversion, [354](#)
 - make_temporary_output_clone, [355](#)
 - max, [356](#)
 - min, [356](#)
 - mixed_precision_dispatch, [357](#)
 - mixed_precision_dispatch_real_complex, [357](#)
 - nan, [358](#)
 - never, [336](#)
 - one, [358, 359](#)
 - operator!=, [359, 360](#)
 - operator<<, [360, 361](#)
 - operator==, [361](#)
 - pi, [362](#)
 - precision_dispatch, [362](#)
 - precision_dispatch_real_complex, [362, 363](#)
 - previous_precision_base, [334](#)
 - read, [363](#)
 - read_binary, [364](#)
 - read_binary_raw, [365](#)
 - read_generic, [365](#)
 - read_generic_raw, [366](#)
 - read_raw, [367](#)
 - real, [367](#)
 - reduce_add, [368](#)
 - remove_complex, [334](#)
 - round_down, [369](#)
 - round_up, [369](#)
 - safe_divide, [370](#)
 - share, [370](#)
 - squared_norm, [371](#)
 - to_complex, [334](#)
 - to_real, [335](#)
 - transpose, [371, 372](#)
 - unit_root, [372](#)
 - with_matrix_type, [373](#)
 - write, [374](#)
 - write_binary, [374](#)
 - write_binary_raw, [375](#)
 - write_raw, [375](#)
 - zero, [376](#)
- gko::AbsoluteComputable, [407](#)
 - compute_absolute_linop, [407](#)

- gko::AbstractFactory< AbstractProductType, ComponentsType >, 408
 - generate, 409
- gko::accessor, 377
- gko::accessor::row_major< ValueType, Dimensionality >, 1007
 - copy_from, 1008
 - length, 1008
 - operator(), 1009
- gko::AllocationError, 409
 - AllocationError, 410
- gko::Allocator, 410
- gko::amd_device, 412
- gko::are_all_integral< Args >, 413
- gko::array< ValueType >, 413
 - array, 415–419
 - as_const_view, 420
 - as_view, 420
 - clear, 420
 - const_view, 420
 - fill, 421
 - get_const_data, 421
 - get_data, 422
 - get_executor, 423
 - get_num_elems, 423
 - get_size, 423
 - is_owning, 424
 - operator=, 424–426
 - resize_and_reset, 427
 - set_executor, 427
 - view, 428
- gko::BadDimension, 430
 - BadDimension, 430
- gko::batch::BatchLinOpFactory, 435
- gko::batch::EnableBatchLinOp< ConcreteBatchLinOp, PolymorphicBase >, 685
- gko::batch::log, 377
- gko::batch::log::BatchConvergence< ValueType >, 434
 - create, 434
 - get_num_iterations, 435
 - get_residual_norm, 435
- gko::batch::matrix::Csr< ValueType, IndexType >, 523
 - add_scaled_identity, 525
 - apply, 526, 527
 - create, 527, 528
 - create_const, 528
 - create_const_view_for_item, 529
 - create_view_for_item, 529
 - get_col_idxs, 530
 - get_const_col_idxs, 530
 - get_const_row_ptrs, 530
 - get_const_values, 531
 - get_const_values_for_item, 531
 - get_num_elements_per_item, 532
 - get_num_stored_elements, 532
 - get_row_ptrs, 532
 - get_values, 533
 - get_values_for_item, 533
 - scale, 533
- gko::batch::matrix::Dense< ValueType >, 574
 - add_scaled_identity, 577
 - apply, 577, 578
 - at, 578, 579, 581
 - create, 581, 582
 - create_const, 583
 - create_const_view_for_item, 583
 - create_view_for_item, 584
 - get_const_values, 584
 - get_const_values_for_item, 584
 - get_cumulative_offset, 585
 - get_num_elements_per_item, 585
 - get_num_stored_elements, 585
 - get_values, 586
 - get_values_for_item, 586
 - scale, 587
 - scale_add, 587
- gko::batch::matrix::Ell< ValueType, IndexType >, 659
 - add_scaled_identity, 661
 - apply, 661, 662
 - create, 662, 663
 - create_const, 664
 - create_const_view_for_item, 664
 - create_view_for_item, 665
 - get_col_idxs, 665
 - get_col_idxs_for_item, 665
 - get_const_col_idxs, 666
 - get_const_col_idxs_for_item, 666
 - get_const_values, 667
 - get_const_values_for_item, 667
 - get_num_elements_per_item, 668
 - get_num_stored_elements, 668
 - get_num_stored_elements_per_row, 668
 - get_values, 669
 - get_values_for_item, 669
 - scale, 669
- gko::batch::matrix::Identity< ValueType >, 787
 - apply, 788, 789
 - create, 789
- gko::batch::MultiVector< ValueType >, 902
 - add_scaled, 904
 - at, 904–906
 - compute_conj_dot, 906
 - compute_dot, 907
 - compute_norm2, 907
 - create, 907, 908
 - create_const, 908
 - create_const_view_for_item, 909
 - create_view_for_item, 909
 - create_with_config_of, 910
 - fill, 910
 - get_common_size, 910
 - get_const_values, 911
 - get_const_values_for_item, 911
 - get_cumulative_offset, 912
 - get_num_batch_items, 912
 - get_num_stored_elements, 912

- get_size, 913
- get_values, 913
- get_values_for_item, 913
- scale, 914
- gko::batch::preconditioner::Jacobi< ValueType, IndexType >, 841
 - get_const_block_pointers, 842
 - get_const_blocks, 843
 - get_const_blocks_cumulative_offsets, 843
 - get_const_map_block_to_row, 843
 - get_max_block_size, 844
 - get_num_blocks, 844
 - get_num_stored_elements, 844
- gko::batch::solver::BatchSolver, 436
 - get_max_iterations, 437
 - get_preconditioner, 437
 - get_system_matrix, 437
 - get_tolerance, 437
 - get_tolerance_type, 438
 - reset_max_iterations, 438
 - reset_tolerance, 438
 - reset_tolerance_type, 439
- gko::batch::solver::Bicgstab< ValueType >, 441
- gko::batch::solver::Cg< ValueType >, 455
- gko::batch::solver::EnableBatchSolver< ConcreteSolver, ValueType, PolymorphicBase >, 686
- gko::batch_dim< Dimensionality, DimensionType >, 431
 - batch_dim, 431
 - get_common_size, 432
 - get_num_batch_items, 432
 - operator!=, 433
 - operator==, 433
- gko::BlockOperator, 448
 - block_at, 449
 - BlockOperator, 449
 - create, 450
 - get_block_size, 450
 - operator=, 451
- gko::BlockSizeError< IndexType >, 451
 - BlockSizeError, 452
- gko::Combination< ValueType >, 466
 - Combination, 467
 - conj_transpose, 468
 - get_coefficients, 468
 - get_operators, 468
 - operator=, 469
 - transpose, 469
- gko::Composition< ValueType >, 498
 - Composition, 499, 500
 - conj_transpose, 500
 - get_operators, 500
 - operator=, 500, 501
 - transpose, 501
- gko::config::pnode, 960
 - get, 964
 - get_array, 964
 - get_boolean, 965
 - get_integer, 965
 - get_map, 965
 - get_real, 965
 - get_string, 966
 - get_tag, 966
 - operator bool, 966
 - pnode, 961–963
- gko::config::registry, 1000
 - emplace, 1001
 - registry, 1000, 1001
- gko::config::type_descriptor, 1078
 - type_descriptor, 1079
- gko::ConvertibleTo< ResultType >, 507
 - convert_to, 507
 - move_to, 508
- gko::CpuAllocator, 518
- gko::CpuAllocatorBase, 519
- gko::CpuTimer, 519
 - difference_async, 520
- gko::cpx_real_type< T >, 520
 - type, 521
- gko::CublasError, 555
 - CublasError, 556
- gko::cuda_stream, 556
 - cuda_stream, 557
 - get, 557
- gko::CudaAllocator, 557
- gko::CudaAllocatorBase, 558
- gko::CudaError, 558
 - CudaError, 558
- gko::CudaExecutor, 559
 - create, 560, 561
 - get_blas_handle, 561
 - get_closest_numa, 561
 - get_closest_pus, 561
 - get_cublas_handle, 562
 - get_cusparses_handle, 562
 - get_description, 562
 - get_master, 562, 563
 - get_sparselib_handle, 563
 - get_stream, 563
 - run, 563, 565
- gko::CudaTimer, 566
 - difference_async, 566
- gko::CufftError, 567
 - CufftError, 567
- gko::CurandError, 568
 - CurandError, 568
- gko::CusparsesError, 571
 - CusparsesError, 571
- gko::default_converter< S, R >, 571
 - operator(), 572
- gko::deferred_factory_parameter< FactoryType >, 572
 - deferred_factory_parameter, 573
 - on, 574
- gko::device_matrix_data< ValueType, IndexType >, 627
 - copy_to_host, 630
 - create_from_host, 631

- device_matrix_data, 629, 630
- empty_out, 631
- get_col_idx, 631
- get_const_col_idx, 631
- get_const_row_idx, 632
- get_const_values, 632
- get_executor, 632
- get_num_elems, 633
- get_num_stored_elements, 633
- get_row_idx, 633
- get_size, 634
- get_values, 634
- remove_zeros, 634
- resize_and_reset, 634, 635
- sum_duplicates, 635
- gko::device_matrix_data< ValueType, IndexType >::arrays, 428
- gko::DiagonalExtractable< ValueType >, 641
 - extract_diagonal, 641
 - extract_diagonal_linop, 641
- gko::DiagonalLinOpExtractable, 642
 - extract_diagonal_linop, 642
- gko::dim< Dimensionality, DimensionType >, 643
 - dim, 644
 - operator bool, 644
 - operator!=, 646
 - operator<=, 646
 - operator*, 646
 - operator==, 647
 - operator[], 645
- gko::DimensionMismatch, 647
 - DimensionMismatch, 648
- gko::DpcppExecutor, 652
 - create, 653
 - get_description, 653
 - get_device_id, 653
 - get_device_type, 654
 - get_master, 654
 - get_max_subgroup_size, 654
 - get_max_workgroup_size, 655
 - get_max_workitem_sizes, 655
 - get_num_computing_units, 655
 - get_num_devices, 655
 - get_subgroup_sizes, 656
 - run, 656, 657
- gko::DpcppTimer, 658
 - difference_async, 658
- gko::enable_parameters_type< ConcreteParameterType, Factory >, 681
 - on, 682
- gko::EnableAbsoluteComputation< AbsoluteLinOp >, 682
 - compute_absolute, 683
 - compute_absolute_linop, 683
- gko::EnableAbstractPolymorphicObject< AbstractObject, PolymorphicBase >, 684
- gko::EnableCreateMethod< ConcreteType >, 687
 - gko::EnableDefaultFactory< ConcreteFactory, ProductType, ParametersType, PolymorphicBase >, 687
 - create, 688
 - get_parameters, 688
 - gko::EnableLinOp< ConcreteLinOp, PolymorphicBase >, 691
 - gko::EnablePolymorphicAssignment< ConcreteType, ResultType >, 694
 - convert_to, 696
 - move_to, 696
 - gko::EnablePolymorphicObject< ConcreteObject, PolymorphicBase >, 697
 - gko::Error, 703
 - Error, 703
 - gko::Executor, 704
 - add_logger, 706
 - alloc, 706
 - copy, 707
 - copy_from, 707
 - copy_val_to_host, 708
 - free, 708
 - get_description, 709
 - get_master, 709
 - memory_accessible, 709
 - remove_logger, 710
 - run, 710, 711
 - set_log_propagation_mode, 712
 - should_propagate_log, 712
 - gko::executor_deleter< T >, 713
 - executor_deleter, 713
 - operator(), 714
 - gko::experimental::distributed, 377
 - assemble_rows_from_neighbors, 380
 - assembly_mode, 379
 - build_partition_from_local_range, 380
 - build_partition_from_local_size, 381
 - combined, 380
 - comm_index_type, 379
 - index_space, 379
 - local, 380
 - make_temporary_conversion, 382
 - non_local, 380
 - precision_dispatch, 383
 - precision_dispatch_real_complex, 383, 384
 - gko::experimental::distributed::DistributedBase, 650
 - get_communicator, 651
 - operator=, 651
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 809
 - get_remote_global_idx, 811
 - get_remote_local_idx, 811
 - get_remote_target_idx, 812
 - index_map, 811
 - map_to_local, 812
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 864
 - col_scale, 868

- create, [868–873](#)
- get_local_matrix, [873](#)
- get_non_local_matrix, [873](#)
- Matrix, [867](#), [868](#)
- operator=, [874](#)
- read_distributed, [875](#), [876](#)
- row_scale, [877](#)
- gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, [935](#)
 - build_from_contiguous, [937](#)
 - build_from_global_size_uniform, [938](#)
 - build_from_mapping, [938](#)
 - get_num_empty_parts, [939](#)
 - get_num_parts, [939](#)
 - get_num_ranges, [939](#)
 - get_part_ids, [939](#)
 - get_part_size, [940](#)
 - get_part_sizes, [940](#)
 - get_range_bounds, [941](#)
 - get_range_starting_indices, [941](#)
 - get_size, [941](#)
 - has_connected_parts, [942](#)
 - has_ordered_parts, [942](#)
- gko::experimental::distributed::preconditioner, [385](#)
- gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >, [1020](#)
 - apply_uses_initial_guess, [1021](#)
 - parse, [1021](#)
- gko::experimental::distributed::Vector< ValueType >, [1088](#)
 - add_scaled, [1091](#)
 - at_local, [1091](#), [1092](#)
 - compute_absolute, [1092](#)
 - compute_conj_dot, [1093](#)
 - compute_dot, [1094](#)
 - compute_mean, [1094](#), [1095](#)
 - compute_norm1, [1095](#)
 - compute_norm2, [1096](#)
 - compute_squared_norm2, [1096](#), [1097](#)
 - create, [1097–1099](#)
 - create_const, [1099](#), [1100](#)
 - create_real_view, [1100](#)
 - create_with_config_of, [1100](#)
 - create_with_type_of, [1101](#)
 - fill, [1102](#)
 - get_const_local_values, [1102](#)
 - get_local_values, [1102](#)
 - get_local_vector, [1103](#)
 - inv_scale, [1103](#)
 - make_complex, [1103](#)
 - read_distributed, [1104](#)
 - scale, [1104](#)
 - sub_scaled, [1105](#)
- gko::experimental::factorization::Cholesky< ValueType, IndexType >, [460](#)
 - generate, [461](#)
 - get_parameters, [461](#), [462](#)
 - parse, [462](#)
- gko::experimental::factorization::Factorization< ValueType, IndexType >, [714](#)
 - create_from_combined_lu, [715](#)
 - create_from_composition, [716](#)
 - create_from_symm_composition, [716](#)
 - unpack, [717](#)
- gko::experimental::factorization::Lu< ValueType, IndexType >, [856](#)
 - generate, [857](#)
 - get_parameters, [857](#)
 - parse, [857](#)
- gko::experimental::mpi, [385](#)
 - get_walltime, [386](#)
 - map_rank_to_device_id, [386](#)
 - wait_all, [387](#)
- gko::experimental::mpi::communicator, [470](#)
 - all_gather, [474](#)
 - all_reduce, [476](#)
 - all_to_all, [477](#), [478](#)
 - all_to_all_v, [478](#), [479](#)
 - broadcast, [480](#)
 - communicator, [473](#), [474](#)
 - gather, [480](#)
 - gather_v, [481](#)
 - get, [481](#)
 - i_all_gather, [482](#)
 - i_all_reduce, [482](#), [483](#)
 - i_all_to_all, [484](#)
 - i_all_to_all_v, [485](#), [486](#)
 - i_broadcast, [487](#)
 - i_gather, [487](#)
 - i_gather_v, [488](#)
 - i_recv, [489](#)
 - i_reduce, [489](#)
 - i_scan, [490](#)
 - i_scatter, [491](#)
 - i_scatter_v, [491](#)
 - i_send, [492](#)
 - node_local_rank, [493](#)
 - operator!=, [493](#)
 - operator==, [493](#)
 - rank, [494](#)
 - recv, [494](#)
 - reduce, [495](#)
 - scan, [495](#)
 - scatter, [496](#)
 - scatter_v, [496](#)
 - send, [497](#)
 - size, [498](#)
 - synchronize, [498](#)
- gko::experimental::mpi::contiguous_type, [501](#)
 - contiguous_type, [502](#)
 - get, [503](#)
 - operator=, [503](#)
- gko::experimental::mpi::environment, [701](#)
 - environment, [702](#)
 - get_provided_thread_support, [702](#)

- gko::experimental::mpi::request, 1003
 - get, 1004
 - request, 1004
 - wait, 1004
- gko::experimental::mpi::status, 1058
 - get, 1059
 - get_count, 1059
 - status, 1059
- gko::experimental::mpi::type_impl< T >, 1079
- gko::experimental::mpi::window< ValueType >, 1108
 - accumulate, 1112
 - fence, 1112
 - fetch_and_op, 1113
 - flush, 1113
 - flush_local, 1113
 - get, 1114
 - get_accumulate, 1114
 - get_window, 1115
 - lock, 1115
 - lock_all, 1116
 - operator=, 1116
 - put, 1116
 - r_accumulate, 1117
 - r_get, 1117
 - r_get_accumulate, 1118
 - r_put, 1119
 - unlock, 1119
 - window, 1110, 1111
- gko::experimental::reorder, 387
 - mc64_strategy, 388
- gko::experimental::reorder::Amd< IndexType >, 411
 - generate, 411
 - get_parameters, 411
- gko::experimental::reorder::Mc64< ValueType, IndexType >, 890
 - generate, 891
 - get_parameters, 891
- gko::experimental::reorder::Rcm< IndexType >, 990
 - generate, 991
 - get_parameters, 991
- gko::experimental::reorder::ScaledReordered< ValueType, IndexType >, 1019
- gko::experimental::solver::Direct< ValueType, IndexType >, 648
 - conj_transpose, 649
 - parse, 649
 - transpose, 650
- gko::factorization, 388
 - incomplete_algorithm, 388
- gko::factorization::lc< ValueType, IndexType >, 786
 - parse, 787
- gko::factorization::llu< ValueType, IndexType >, 799
 - parse, 799
- gko::factorization::Parlc< ValueType, IndexType >, 929
 - parse, 930
- gko::factorization::Parlct< ValueType, IndexType >, 930
 - parse, 931
- gko::factorization::Parllu< ValueType, IndexType >, 932
 - parse, 933
- gko::factorization::Parllut< ValueType, IndexType >, 934
 - parse, 935
- gko::half, 752
- gko::hip_stream, 753
 - get, 754
 - hip_stream, 754
- gko::HipAllocatorBase, 755
- gko::HipblasError, 755
 - HipblasError, 755
- gko::HipError, 756
 - HipError, 756
- gko::HipExecutor, 757
 - create, 758
 - get_blas_handle, 758
 - get_closest_numa, 759
 - get_closest_pus, 759
 - get_description, 759
 - get_hipblas_handle, 759
 - get_hipsparse_handle, 760
 - get_master, 760
 - get_sparselib_handle, 760
 - run, 761, 762
- gko::HipfftError, 762
 - HipfftError, 763
- gko::HiprandError, 763
 - HiprandError, 764
- gko::HipsparseError, 764
 - HipsparseError, 765
- gko::HipTimer, 765
 - difference_async, 766
- gko::index_set< IndexType >, 812
 - clear, 817
 - contains, 817, 818
 - get_executor, 818
 - get_global_index, 818
 - get_local_index, 819
 - get_num_elems, 819
 - get_num_subsets, 820
 - get_size, 820
 - get_subsets_begin, 820
 - get_subsets_end, 820
 - get_superset_indices, 821
 - index_set, 814–817
 - is_contiguous, 821
 - map_global_to_local, 821
 - map_local_to_global, 822
 - operator=, 822, 823
 - to_global_indices, 823
- gko::InvalidStateError, 824
 - InvalidStateError, 824
- gko::KernelNotFound, 845
 - KernelNotFound, 845
- gko::LinOpFactory, 846
- gko::log, 389
 - criterion, 390
 - factory, 390

- internal, 390
- linop, 390
- memory, 390
- object, 390
- operation, 390
- profile_event_category, 390
- solver, 390
- user, 390
- gko::log::Convergence< ValueType >, 503
 - create, 504, 505
 - get_implicit_sq_resnorm, 505
 - get_num_iterations, 505
 - get_residual, 506
 - get_residual_norm, 506
 - has_converged, 506
- gko::log::criterion_data, 523
- gko::log::EnableLogging< ConcreteLoggable, PolymorphicBase >, 692
- gko::log::executor_data, 713
- gko::log::iteration_complete_data, 833
- gko::log::linop_data, 846
- gko::log::linop_factory_data, 846
- gko::log::Loggable, 851
 - add_logger, 851
 - get_loggers, 852
 - remove_logger, 852
- gko::log::operation_data, 924
- gko::log::Papi< ValueType >, 926
 - create, 927, 928
 - get_handle, 928
 - get_handle_name, 928
- gko::log::PerformanceHint, 943
 - create, 943
- gko::log::polymorphic_object_data, 967
- gko::log::ProfilerHook, 978
 - create_nested_summary, 979
 - create_nvtx, 979
 - create_summary, 980
 - create_tau, 980
 - set_object_name, 980
 - user_range, 981
- gko::log::ProfilerHook::NestedSummaryWriter, 914
 - write_nested, 915
- gko::log::ProfilerHook::SummaryWriter, 1071
 - write, 1071
- gko::log::ProfilerHook::TableSummaryWriter, 1072
 - TableSummaryWriter, 1072
 - write, 1073
 - write_nested, 1073
- gko::log::profiling_scope_guard, 981
 - profiling_scope_guard, 982
- gko::log::Record, 995
 - create, 996
 - get, 997
- gko::log::Record::logged_data, 852
- gko::log::SolverProgress, 1041
 - create_scalar_csv_writer, 1041
 - create_scalar_table_writer, 1042
 - create_vector_storage, 1042
- gko::log::Stream< ValueType >, 1068
 - create, 1069
- gko::machine_topology, 858
 - bind_to_core, 859
 - bind_to_cores, 860
 - bind_to_pu, 860
 - bind_to_pus, 860
 - get_core, 862
 - get_instance, 862
 - get_num_cores, 862
 - get_num_numas, 862
 - get_num_pci_devices, 863
 - get_num_pus, 863
 - get_pci_device, 863, 864
 - get_pu, 864
- gko::matrix, 391
 - columns, 393
 - inverse, 393
 - inverse_columns, 393
 - inverse_rows, 393
 - inverse_symmetric, 393
 - none, 393
 - permute_mode, 392
 - rows, 393
 - symmetric, 393
- gko::matrix::Coo< ValueType, IndexType >, 509
 - apply2, 511, 512
 - compute_absolute, 512
 - create, 512, 513
 - create_const, 514
 - extract_diagonal, 514
 - get_col_idx, 515
 - get_const_col_idx, 515
 - get_const_row_idx, 515
 - get_const_values, 516
 - get_num_stored_elements, 516
 - get_row_idx, 516
 - get_values, 517
 - read, 517, 518
 - write, 518
- gko::matrix::Csr< ValueType, IndexType >, 534
 - column_permute, 538
 - compute_absolute, 539
 - conj_transpose, 539
 - create, 539–541
 - create_const, 541
 - create_submatrix, 542
 - Csr, 538
 - extract_diagonal, 543
 - get_col_idx, 543
 - get_const_col_idx, 543
 - get_const_row_ptr, 544
 - get_const_srow, 544
 - get_const_values, 544
 - get_num_srow_elements, 545
 - get_num_stored_elements, 545
 - get_row_ptr, 545

- get_srow, 546
- get_strategy, 546
- get_values, 546
- inv_scale, 546
- inverse_column_permute, 547
- inverse_permute, 547
- inverse_row_permute, 548
- operator=, 548
- permute, 548, 549
- permute_reuse, 550, 551
- read, 551, 552
- row_permute, 552
- scale, 553
- scale_permute, 553
- set_strategy, 554
- transpose, 554
- transpose_reuse, 554
- write, 555
- gko::matrix::Csr< ValueType, IndexType >::classical, 463
 - clac_size, 463
 - copy, 464
 - process, 464
- gko::matrix::Csr< ValueType, IndexType >::cuspars, 569
 - clac_size, 569
 - copy, 570
 - process, 570
- gko::matrix::Csr< ValueType, IndexType >::load_balance, 847
 - clac_size, 850
 - copy, 850
 - load_balance, 848, 849
 - process, 850
- gko::matrix::Csr< ValueType, IndexType >::merge_path, 891
 - clac_size, 892
 - copy, 892
 - process, 893
- gko::matrix::Csr< ValueType, IndexType >::permute_reuse_in, 953
 - update_values, 954
- gko::matrix::Csr< ValueType, IndexType >::sparselib, 1047
 - clac_size, 1047
 - copy, 1048
 - process, 1048
- gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1063
 - clac_size, 1064
 - copy, 1065
 - get_name, 1065
 - process, 1065
 - strategy_type, 1064
- gko::matrix::Dense< ValueType >, 587
 - add_scaled, 593
 - at, 593–595
 - column_permute, 595, 596
 - compute_absolute, 597
 - compute_conj_dot, 597, 598
 - compute_dot, 598
 - compute_mean, 599
 - compute_norm1, 599, 600
 - compute_norm2, 600
 - compute_squared_norm2, 601
 - conj_transpose, 601, 602
 - create, 602, 603
 - create_const, 603
 - create_const_view_of, 604
 - create_real_view, 604
 - create_submatrix, 605
 - create_view_of, 605
 - create_with_config_of, 606
 - create_with_type_of, 606, 607
 - Dense, 592, 593
 - extract_diagonal, 607, 608
 - fill, 608
 - get_const_values, 608
 - get_num_stored_elements, 609
 - get_stride, 609
 - get_values, 609
 - inv_scale, 610
 - inverse_column_permute, 610, 611
 - inverse_permute, 611–613
 - inverse_row_permute, 613, 614
 - make_complex, 614, 615
 - operator=, 615
 - permute, 615–619
 - row_gather, 619–621
 - row_permute, 621–623
 - scale, 623
 - scale_permute, 623–625
 - sub_scaled, 626
 - transpose, 626
- gko::matrix::Diagonal< ValueType >, 635
 - compute_absolute, 637
 - conj_transpose, 637
 - create, 637, 638
 - create_const, 638
 - get_const_values, 639
 - get_values, 639
 - inverse_apply, 639
 - rapply, 640
 - transpose, 640
- gko::matrix::Ell< ValueType, IndexType >, 670
 - col_at, 672, 673
 - compute_absolute, 673
 - create, 674, 675
 - create_const, 675
 - Ell, 672
 - extract_diagonal, 676
 - get_col_idxs, 676
 - get_const_col_idxs, 677
 - get_const_values, 677
 - get_num_stored_elements, 677
 - get_num_stored_elements_per_row, 678

- get_stride, 678
- get_values, 678
- operator=, 678, 679
- read, 679, 680
- val_at, 680
- write, 681
- gko::matrix::Fbcsr< ValueType, IndexType >, 717
 - compute_absolute, 720
 - conj_transpose, 720
 - convert_to, 721
 - create, 721–723
 - create_const, 723
 - extract_diagonal, 724
 - Fbcsr, 720
 - get_block_size, 724
 - get_col_idxxs, 724
 - get_const_col_idxxs, 724
 - get_const_row_ptrs, 725
 - get_const_values, 725
 - get_num_block_cols, 725
 - get_num_block_rows, 726
 - get_num_stored_blocks, 726
 - get_num_stored_elements, 726
 - get_row_ptrs, 726
 - get_values, 727
 - is_sorted_by_column_index, 727
 - operator=, 727
 - read, 728
 - transpose, 729
 - write, 729
- gko::matrix::Fft, 732
 - conj_transpose, 733
 - create, 733, 734
 - transpose, 734
 - write, 734, 735
- gko::matrix::Fft2, 736
 - conj_transpose, 737
 - create, 737, 738
 - transpose, 738
 - write, 738, 739
- gko::matrix::Fft3, 740
 - conj_transpose, 741
 - create, 741, 742
 - transpose, 742
 - write, 742, 743
- gko::matrix::Hybrid< ValueType, IndexType >, 766
 - compute_absolute, 769
 - create, 769–772
 - ell_col_at, 772
 - ell_val_at, 773
 - extract_diagonal, 774
 - get_const_coo_col_idxxs, 774
 - get_const_coo_row_idxxs, 774
 - get_const_coo_values, 775
 - get_const_ell_col_idxxs, 775
 - get_const_ell_values, 775
 - get_coo, 776
 - get_coo_col_idxxs, 776
 - get_coo_num_stored_elements, 776
 - get_coo_row_idxxs, 777
 - get_coo_values, 777
 - get_ell, 777
 - get_ell_col_idxxs, 777
 - get_ell_num_stored_elements, 778
 - get_ell_num_stored_elements_per_row, 778
 - get_ell_stride, 778
 - get_ell_values, 778
 - get_num_stored_elements, 779
 - get_strategy, 779
 - Hybrid, 769
 - operator=, 780
 - read, 780, 781
 - write, 781
- gko::matrix::Hybrid< ValueType, IndexType >::automatic, 429
 - compute_ell_num_stored_elements_per_row, 429
- gko::matrix::Hybrid< ValueType, IndexType >::column_limit, 465
 - column_limit, 465
 - compute_ell_num_stored_elements_per_row, 465
 - get_num_columns, 466
- gko::matrix::Hybrid< ValueType, IndexType >::imbalance_bounded_limit, 805
 - compute_ell_num_stored_elements_per_row, 806
 - get_percentage, 806
 - get_ratio, 806
- gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit, 807
 - compute_ell_num_stored_elements_per_row, 808
 - get_percentage, 808
 - imbalance_limit, 807
- gko::matrix::Hybrid< ValueType, IndexType >::minimal_storage_limit, 894
 - compute_ell_num_stored_elements_per_row, 895
 - get_percentage, 895
- gko::matrix::Hybrid< ValueType, IndexType >::strategy_type, 1066
 - compute_ell_num_stored_elements_per_row, 1067
 - compute_hybrid_config, 1067
 - get_coo_nnz, 1067
 - get_ell_num_stored_elements_per_row, 1068
- gko::matrix::Identity< ValueType >, 790
 - conj_transpose, 791
 - create, 791, 792
 - transpose, 792
- gko::matrix::IdentityFactory< ValueType >, 792
 - create, 793
- gko::matrix::Permutation< IndexType >, 948
 - compose, 949
 - compute_inverse, 950
 - create, 950
 - create_const, 951
 - get_const_permutation, 952
 - get_permutation, 952
 - get_permutation_size, 952

- write, [953](#)
- `gko::matrix::RowGatherer< IndexType >`, [1010](#)
 - create, [1011](#)
 - create_const, [1012](#)
 - get_const_row_idxs, [1012](#)
 - get_row_idxs, [1013](#)
- `gko::matrix::ScaledPermutation< ValueType, IndexType >`, [1014](#)
 - compose, [1015](#)
 - compute_inverse, [1015](#)
 - create, [1016](#)
 - create_const, [1017](#)
 - get_const_permutation, [1017](#)
 - get_const_scaling_factors, [1018](#)
 - get_permutation, [1018](#)
 - get_scaling_factors, [1018](#)
 - write, [1019](#)
- `gko::matrix::Sellp< ValueType, IndexType >`, [1030](#)
 - col_at, [1032](#)
 - compute_absolute, [1033](#)
 - create, [1033](#), [1034](#)
 - extract_diagonal, [1034](#)
 - get_col_idxs, [1035](#)
 - get_const_col_idxs, [1035](#)
 - get_const_slice_lengths, [1035](#)
 - get_const_slice_sets, [1035](#)
 - get_const_values, [1036](#)
 - get_num_stored_elements, [1036](#)
 - get_slice_lengths, [1036](#)
 - get_slice_sets, [1037](#)
 - get_slice_size, [1037](#)
 - get_stride_factor, [1037](#)
 - get_total_cols, [1037](#)
 - get_values, [1038](#)
 - operator=, [1038](#)
 - read, [1038](#), [1039](#)
 - Sellp, [1031](#), [1032](#)
 - val_at, [1039](#), [1040](#)
 - write, [1040](#)
- `gko::matrix::SparsityCsr< ValueType, IndexType >`, [1048](#)
 - conj_transpose, [1051](#)
 - create, [1051](#), [1052](#)
 - create_const, [1053](#)
 - get_col_idxs, [1053](#)
 - get_const_col_idxs, [1054](#)
 - get_const_row_ptrs, [1054](#)
 - get_const_value, [1054](#)
 - get_num_nonzeros, [1055](#)
 - get_row_ptrs, [1055](#)
 - get_value, [1055](#)
 - operator=, [1056](#)
 - read, [1056](#), [1057](#)
 - SparsityCsr, [1050](#)
 - to_adjacency_matrix, [1057](#)
 - transpose, [1057](#)
 - write, [1058](#)
- `gko::matrix_assembly_data< ValueType, IndexType >`, [877](#)
 - add_value, [878](#)
 - contains, [878](#)
 - get_num_stored_elements, [879](#)
 - get_ordered_data, [879](#)
 - get_size, [879](#)
 - get_value, [879](#)
 - set_value, [880](#)
- `gko::matrix_data< ValueType, IndexType >`, [880](#)
 - cond, [885](#)
 - diag, [886–888](#)
 - matrix_data, [882–884](#)
 - nonzeros, [889](#)
 - sum_duplicates, [889](#)
- `gko::matrix_data_entry< ValueType, IndexType >`, [889](#)
- `gko::MetisError`, [893](#)
 - MetisError, [893](#)
- `gko::MpiError`, [895](#)
 - MpiError, [896](#)
- `gko::multigrid`, [393](#)
- `gko::multigrid::EnableMultigridLevel< ValueType >`, [692](#)
 - get_coarse_op, [693](#)
 - get_fine_op, [693](#)
 - get_prolong_op, [693](#)
 - get_restrict_op, [694](#)
- `gko::multigrid::FixedCoarsening< ValueType, IndexType >`, [744](#)
 - get_system_matrix, [744](#)
- `gko::multigrid::MultigridLevel`, [900](#)
 - get_coarse_op, [901](#)
 - get_fine_op, [901](#)
 - get_prolong_op, [901](#)
 - get_restrict_op, [901](#)
- `gko::multigrid::Pgm< ValueType, IndexType >`, [958](#)
 - get_agg, [958](#)
 - get_const_agg, [959](#)
 - get_system_matrix, [959](#)
 - parse, [959](#)
- `gko::name_demangling`, [393](#)
 - get_dynamic_type, [394](#)
 - get_static_type, [394](#)
- `gko::NotCompiled`, [915](#)
 - NotCompiled, [916](#)
- `gko::NotImplemented`, [916](#)
 - NotImplemented, [916](#)
- `gko::NotSupported`, [917](#)
 - NotSupported, [917](#)
- `gko::null_deleter< T >`, [918](#)
 - operator(), [918](#)
- `gko::nvidia_device`, [919](#)
- `gko::OmpExecutor`, [919](#)
 - get_description, [920](#)
 - get_master, [920](#)
 - run, [921](#), [922](#)
- `gko::Operation`, [922](#)
 - get_name, [924](#)
- `gko::OutOfBoundsError`, [925](#)

- OutOfBoundsError, 925
- gko::OverflowError, 925
 - OverflowError, 926
- gko::Permutable< IndexType >, 944
 - column_permute, 945
 - inverse_column_permute, 946
 - inverse_permute, 946
 - inverse_row_permute, 946
 - permute, 947
 - row_permute, 947
- gko::Perturbation< ValueType >, 954
 - create, 955, 956
 - get_basis, 956
 - get_projector, 957
 - get_scalar, 957
- gko::PolymorphicObject, 967
 - clear, 968
 - clone, 968, 969
 - copy_from, 969–971
 - create_default, 971, 972
 - get_executor, 972
 - move_from, 972
- gko::precision_reduction, 973
 - autodetect, 975
 - common, 975
 - get_nonpreserving, 975
 - get_preserving, 976
 - operator storage_type, 976
 - precision_reduction, 974
- gko::Preconditionable, 976
 - get_preconditioner, 977
 - set_preconditioner, 977
- gko::preconditioner, 395
 - isai_type, 395
- gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 444
 - compute_storage_space, 445
 - get_block_offset, 446
 - get_global_block_offset, 446
 - get_group_offset, 446
 - get_group_size, 447
 - get_stride, 447
 - group_power, 448
- gko::preconditioner::GaussSeidel< ValueType, IndexType >, 745
 - generate, 745
 - get_parameters, 746
- gko::preconditioner::lc< LSolverType, IndexType >, 782
 - conj_transpose, 784
 - get_l_solver, 784
 - get_lh_solver, 784
 - lc, 783
 - operator=, 784, 785
 - parse, 785
 - transpose, 786
- gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 800
 - conj_transpose, 802
 - get_l_solver, 802
 - get_u_solver, 803
 - llu, 801, 802
 - operator=, 803
 - parse, 804
 - transpose, 804
- gko::preconditioner::lsai< IsaiType, ValueType, IndexType >, 829
 - conj_transpose, 831
 - get_approximate_inverse, 831
 - lsai, 830, 831
 - operator=, 831, 832
 - parse, 832
 - transpose, 832
- gko::preconditioner::Jacobi< ValueType, IndexType >, 835
 - conj_transpose, 837
 - convert_to, 837
 - get_blocks, 838
 - get_conditioning, 838
 - get_num_blocks, 838
 - get_num_stored_elements, 838
 - get_storage_scheme, 839
 - Jacobi, 836, 837
 - move_to, 839
 - operator=, 839, 840
 - parse, 840
 - transpose, 840
 - write, 841
- gko::preconditioner::Sor< ValueType, IndexType >, 1043
 - generate, 1044
 - get_parameters, 1044
- gko::ptr_param< T >, 982
 - get, 983
 - operator bool, 983
 - operator*, 983
 - operator->, 984
- gko::range< Accessor >, 984
 - get_accessor, 987
 - length, 987
 - operator(), 987
 - operator->, 988
 - operator=, 988, 989
 - range, 986
- gko::ReadableFromMatrixData< ValueType, IndexType >, 993
 - read, 993, 994
- gko::ReferenceExecutor, 997
 - get_description, 998
 - run, 998, 999
- gko::reorder, 396
 - EnableDefaultReorderingBaseFactory, 396
- gko::reorder::Rcm< ValueType, IndexType >, 991
 - get_inverse_permutation, 992
 - get_permutation, 992
- gko::reorder::ReorderingBase< IndexType >, 1002
- gko::reorder::ReorderingBaseArgs, 1003

- gko::ScaledIdentityAddable, 1013
 - add_scaled_identity, 1013
- gko::scoped_device_id_guard, 1022
 - scoped_device_id_guard, 1023, 1024
- gko::segmented_array< T >, 1025
 - create_from_offsets, 1027
 - create_from_sizes, 1027, 1028
 - get_const_flat_data, 1028
 - get_executor, 1028
 - get_flat_data, 1028
 - get_offsets, 1029
 - get_segment_count, 1029
 - get_size, 1029
 - segmented_array, 1026
- gko::solver, 397
 - build_smoother, 399, 400
 - initial_guess_mode, 399
 - provided, 399
 - rhs, 399
 - trisolve_algorithm, 399
 - zero, 399
- gko::solver::ApplyWithInitialGuess, 412
- gko::solver::Bicg< ValueType >, 439
 - apply_uses_initial_guess, 440
 - conj_transpose, 440
 - parse, 440
 - transpose, 441
- gko::solver::Bicgstab< ValueType >, 442
 - apply_uses_initial_guess, 443
 - conj_transpose, 443
 - parse, 443
 - transpose, 444
- gko::solver::CbGmres< ValueType >, 452
 - get_krylov_dim, 453
 - get_storage_precision, 453
 - parse, 454
 - set_krylov_dim, 454
- gko::solver::Cg< ValueType >, 456
 - apply_uses_initial_guess, 456
 - conj_transpose, 457
 - parse, 457
 - transpose, 457
- gko::solver::Cgs< ValueType >, 458
 - apply_uses_initial_guess, 459
 - conj_transpose, 459
 - parse, 459
 - transpose, 460
- gko::solver::EnableApplyWithInitialGuess< DerivedType >, 685
- gko::solver::EnableIterativeBase< DerivedType >, 689
 - EnableIterativeBase, 690
 - operator=, 690
 - set_stop_criterion_factory, 690
- gko::solver::EnablePreconditionable< DerivedType >, 697
 - EnablePreconditionable, 698
 - operator=, 698
 - set_preconditioner, 699
- gko::solver::EnablePreconditionedIterativeSolver< ValueType, DerivedType >, 699
- gko::solver::EnableSolverBase< DerivedType, MatrixType >, 700
 - EnableSolverBase, 701
 - operator=, 701
- gko::solver::Fcg< ValueType >, 729
 - apply_uses_initial_guess, 730
 - conj_transpose, 731
 - parse, 731
 - transpose, 731
- gko::solver::Gcr< ValueType >, 746
 - apply_uses_initial_guess, 747
 - conj_transpose, 747
 - get_krylov_dim, 748
 - parse, 748
 - set_krylov_dim, 749
 - transpose, 749
- gko::solver::Gmres< ValueType >, 749
 - apply_uses_initial_guess, 750
 - conj_transpose, 750
 - get_krylov_dim, 751
 - parse, 751
 - set_krylov_dim, 752
 - transpose, 752
- gko::solver::has_with_criteria< SolverType, typename >, 753
- gko::solver::ldr< ValueType >, 794
 - apply_uses_initial_guess, 795
 - conj_transpose, 795
 - get_complex_subspace, 795
 - get_deterministic, 795
 - get_kappa, 796
 - get_subspace_dim, 796
 - parse, 796
 - set_complex_subspace, 797
 - set_complex_subspace, 797
 - set_deterministic, 797
 - set_kappa, 798
 - set_subspace_dim, 798
 - transpose, 798
- gko::solver::lr< ValueType >, 824
 - apply_uses_initial_guess, 826
 - conj_transpose, 827
 - get_solver, 827
 - lr, 826
 - operator=, 827
 - parse, 828
 - set_solver, 828
 - transpose, 829
- gko::solver::IterativeBase, 834
 - get_stop_criterion_factory, 834
 - set_stop_criterion_factory, 834
- gko::solver::LowerTrs< ValueType, IndexType >, 853
 - conj_transpose, 854
 - LowerTrs, 854
 - operator=, 854, 855
 - parse, 855

- transpose, 855
- gko::solver::Multigrid, 896
 - apply_uses_initial_guess, 897
 - get_coarsest_solver, 898
 - get_cycle, 898
 - get_mg_level_list, 898
 - get_mid_smoother_list, 898
 - get_post_smoother_list, 899
 - get_pre_smoother_list, 899
 - parse, 899
 - set_cycle, 900
- gko::solver::multigrid, 400
 - cycle, 401
 - mid_smooth_type, 401
- gko::solver::UpperTrs< ValueType, IndexType >, 1082
 - conj_transpose, 1084
 - operator=, 1084
 - parse, 1084
 - transpose, 1085
 - UpperTrs, 1083
- gko::solver::workspace_traits< Solver >, 1120
- gko::span, 1045
 - is_valid, 1046
 - length, 1046
 - span, 1045, 1046
- gko::stop, 402
 - EnableDefaultCriterionFactory, 403
- gko::stop::AbsoluteResidualNorm< ValueType >, 408
- gko::stop::Combined, 470
- gko::stop::Criterion, 521
 - check, 522
 - update, 522
- gko::stop::Criterion::Updater, 1081
 - check, 1082
- gko::stop::CriterionArgs, 523
- gko::stop::ImplicitResidualNorm< ValueType >, 809
- gko::stop::Iteration, 833
- gko::stop::RelativeResidualNorm< ValueType >, 1002
- gko::stop::ResidualNorm< ValueType >, 1005
- gko::stop::ResidualNormBase< ValueType >, 1006
- gko::stop::ResidualNormReduction< ValueType >, 1006
- gko::stop::Time, 1073
- gko::stopping_status, 1060
 - converge, 1061
 - get_id, 1061
 - has_converged, 1061
 - has_stopped, 1061
 - is_finalized, 1062
 - operator!=, 1062
 - operator==, 1063
 - stop, 1062
- gko::StreamError, 1070
 - StreamError, 1070
- gko::syn, 403
 - as_array, 405
 - as_list, 404
 - concatenate, 404
 - gko::syn::range< Start, End, Step >, 989
 - gko::syn::type_list< Types >, 1080
 - gko::syn::value_list< T, Values >, 1087
 - gko::time_point, 1074
 - gko::Timer, 1074
 - create_for_executor, 1075
 - create_time_point, 1075
 - difference, 1075
 - difference_async, 1076
 - gko::Transposable, 1076
 - conj_transpose, 1077
 - transpose, 1077
 - gko::UnsupportedMatrixProperty, 1080
 - UnsupportedMatrixProperty, 1081
 - gko::UseComposition< ValueType >, 1085
 - get_composition, 1086
 - get_operator_at, 1086
 - gko::ValueMismatch, 1087
 - ValueMismatch, 1088
 - gko::version, 1105
 - tag, 1106
 - gko::version_info, 1106
 - core_version, 1107
 - cuda_version, 1107
 - dpcpp_version, 1108
 - get, 1107
 - hip_version, 1108
 - omp_version, 1108
 - reference_version, 1108
 - gko::WritableToMatrixData< ValueType, IndexType >, 1120
 - write, 1121
 - gko::xstd, 405
- GKO_CREATE_FACTORY_PARAMETERS
 - Linear Operators, 301
- GKO_ENABLE_BATCH_LIN_OP_FACTORY
 - BatchLinOp, 319
- GKO_ENABLE_BUILD_METHOD
 - Linear Operators, 301
- GKO_ENABLE_LIN_OP_FACTORY
 - Linear Operators, 302
- GKO_FACTORY_PARAMETER
 - Linear Operators, 303
- GKO_FACTORY_PARAMETER_SCALAR
 - Linear Operators, 303
- GKO_FACTORY_PARAMETER_VECTOR
 - Linear Operators, 304
- GKO_REGISTER_HOST_OPERATION
 - Executors, 290
- GKO_REGISTER_OPERATION
 - Executors, 291
- group_power
 - gko::preconditioner::block_interleaved_storage_scheme< IndexType >, 448
- has_connected_parts
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 942
- has_converged

- gko::log::Convergence< ValueType >, 506
- gko::stopping_status, 1061
- has_ordered_parts
 - gko::experimental::distributed::Partition< LocalIndexType, GlobalIndexType >, 942
- has_stopped
 - gko::stopping_status, 1061
- highest_precision
 - gko, 333
- HIP Executor, 294
- hip_stream
 - gko::hip_stream, 754
- hip_version
 - gko::version_info, 1108
- HipblasError
 - gko::HipblasError, 755
- HipError
 - gko::HipError, 756
- HipfftError
 - gko::HipfftError, 763
- HiprandError
 - gko::HiprandError, 764
- HipsparseError
 - gko::HipsparseError, 765
- Hybrid
 - gko::matrix::Hybrid< ValueType, IndexType >, 769
- i_all_gather
 - gko::experimental::mpi::communicator, 482
- i_all_reduce
 - gko::experimental::mpi::communicator, 482, 483
- i_all_to_all
 - gko::experimental::mpi::communicator, 484
- i_all_to_all_v
 - gko::experimental::mpi::communicator, 485, 486
- i_broadcast
 - gko::experimental::mpi::communicator, 487
- i_gather
 - gko::experimental::mpi::communicator, 487
- i_gather_v
 - gko::experimental::mpi::communicator, 488
- i_recv
 - gko::experimental::mpi::communicator, 489
- i_reduce
 - gko::experimental::mpi::communicator, 489
- i_scan
 - gko::experimental::mpi::communicator, 490
- i_scatter
 - gko::experimental::mpi::communicator, 491
- i_scatter_v
 - gko::experimental::mpi::communicator, 491
- i_send
 - gko::experimental::mpi::communicator, 492
- lc
 - gko::preconditioner::lc< LSolverType, IndexType >, 783
- llu
 - gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 801, 802
- imag
 - gko, 346
- imbalance_limit
 - gko::matrix::Hybrid< ValueType, IndexType >::imbalance_limit, 807
- incomplete_algorithm
 - gko::factorization, 388
- index_map
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 811
- index_set
 - gko::index_set< IndexType >, 814–817
- index_space
 - gko::experimental::distributed, 379
- initial_guess_mode
 - gko::solver, 399
- initialize
 - Linear Operators, 305–307
- internal
 - gko::log, 390
- inv_scale
 - gko::experimental::distributed::Vector< ValueType >, 1103
 - gko::matrix::Csr< ValueType, IndexType >, 546
 - gko::matrix::Dense< ValueType >, 610
- InvalidStateError
 - gko::InvalidStateError, 824
- inverse
 - gko::matrix, 393
- inverse_apply
 - gko::matrix::Diagonal< ValueType >, 639
- inverse_column_permute
 - gko::matrix::Csr< ValueType, IndexType >, 547
 - gko::matrix::Dense< ValueType >, 610, 611
 - gko::Permutable< IndexType >, 946
- inverse_columns
 - gko::matrix, 393
- inverse_permute
 - gko::matrix::Csr< ValueType, IndexType >, 547
 - gko::matrix::Dense< ValueType >, 611–613
 - gko::Permutable< IndexType >, 946
- inverse_row_permute
 - gko::matrix::Csr< ValueType, IndexType >, 548
 - gko::matrix::Dense< ValueType >, 613, 614
 - gko::Permutable< IndexType >, 946
- inverse_rows
 - gko::matrix, 393
- inverse_symmetric
 - gko::matrix, 393
- lr
 - gko::solver::lr< ValueType >, 826
- is_complex
 - gko, 347
- is_complex_or_scalar
 - gko, 347
- is_complex_or_scalar_s
 - gko, 333
- is_complex_s

- gko, 334
- is_contiguous
 - gko::index_set< IndexType >, 821
- is_finalized
 - gko::stopping_status, 1062
- is_finite
 - gko, 348
- is_nan
 - gko, 349
- is_nonzero
 - gko, 350
- is_owning
 - gko::array< ValueType >, 424
- is_sorted_by_column_index
 - gko::matrix::Fbcsr< ValueType, IndexType >, 727
- is_valid
 - gko::span, 1046
- is_zero
 - gko, 350
- Isai
 - gko::preconditioner::Isai< IsaiType, ValueType, IndexType >, 830, 831
- isai_type
 - gko::preconditioner, 395
- Jacobi
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 836, 837
- Jacobi Preconditioner, 295
- KernelNotFound
 - gko::KernelNotFound, 845
- layout_type
 - gko, 335
- lend
 - gko, 351
- length
 - gko::accessor::row_major< ValueType, Dimensionality >, 1008
 - gko::range< Accessor >, 987
 - gko::span, 1046
- Linear Operators, 296
 - EnableDefaultLinOpFactory, 304
 - GKO_CREATE_FACTORY_PARAMETERS, 301
 - GKO_ENABLE_BUILD_METHOD, 301
 - GKO_ENABLE_LIN_OP_FACTORY, 302
 - GKO_FACTORY_PARAMETER, 303
 - GKO_FACTORY_PARAMETER_SCALAR, 303
 - GKO_FACTORY_PARAMETER_VECTOR, 304
 - initialize, 305–307
- linop
 - gko::log, 390
- load_balance
 - gko::matrix::Csr< ValueType, IndexType >::load_balance, 848, 849
- local
 - gko::experimental::distributed, 380
- lock
 - gko::experimental::mpi::window< ValueType >, 1115
- lock_all
 - gko::experimental::mpi::window< ValueType >, 1116
- log_propagation_mode
 - gko, 336
- Logging, 308
- LowerTrs
 - gko::solver::LowerTrs< ValueType, IndexType >, 854
- make_array_view
 - gko, 352
- make_complex
 - gko::experimental::distributed::Vector< ValueType >, 1103
 - gko::matrix::Dense< ValueType >, 614, 615
- make_const_array_view
 - gko, 352
- make_const_dense_view
 - gko, 353
- make_dense_view
 - gko, 353
- make_temporary_clone
 - gko, 354
- make_temporary_conversion
 - gko, 354
 - gko::experimental::distributed, 382
- make_temporary_output_clone
 - gko, 355
- map_global_to_local
 - gko::index_set< IndexType >, 821
- map_local_to_global
 - gko::index_set< IndexType >, 822
- map_rank_to_device_id
 - gko::experimental::mpi, 386
- map_to_local
 - gko::experimental::distributed::index_map< LocalIndexType, GlobalIndexType >, 812
- Matrix
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 867, 868
- matrix_data
 - gko::matrix_data< ValueType, IndexType >, 882–884
- max
 - gko, 356
- mc64_strategy
 - gko::experimental::reorder, 388
- memory
 - gko::log, 390
- memory_accessible
 - gko::Executor, 709
- MetisError
 - gko::MetisError, 893
- mid_smooth_type
 - gko::solver::multigrid, 401
- min

- gko, 356
- mixed_precision_dispatch
 - gko, 357
- mixed_precision_dispatch_real_complex
 - gko, 357
- mode
 - Stopping criteria, 316
- move_from
 - gko::PolymorphicObject, 972
- move_to
 - gko::ConvertibleTo< ResultType >, 508
 - gko::EnablePolymorphicAssignment< ConcreteType, ResultType >, 696
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 839
- MpiError
 - gko::MpiError, 896
- nan
 - gko, 358
- never
 - gko, 336
- node_local_rank
 - gko::experimental::mpi::communicator, 493
- non_local
 - gko::experimental::distributed, 380
- none
 - gko::matrix, 393
- nonzeros
 - gko::matrix_data< ValueType, IndexType >, 889
- NotCompiled
 - gko::NotCompiled, 916
- NotImplemented
 - gko::NotImplemented, 916
- NotSupported
 - gko::NotSupported, 917
- object
 - gko::log, 390
- omp_version
 - gko::version_info, 1108
- on
 - gko::deferred_factory_parameter< FactoryType >, 574
 - gko::enable_parameters_type< ConcreteParametersType, Factory >, 682
- one
 - gko, 358, 359
- OpenMP Executor, 311
- operation
 - gko::log, 390
- operator bool
 - gko::config::pnode, 966
 - gko::dim< Dimensionality, DimensionType >, 644
 - gko::ptr_param< T >, 983
- operator storage_type
 - gko::precision_reduction, 976
- operator!=
 - gko, 359, 360
- gko::batch_dim< Dimensionality, DimensionType >, 433
- gko::dim< Dimensionality, DimensionType >, 646
- gko::experimental::mpi::communicator, 493
- gko::stopping_status, 1062
- operator<<
 - gko, 360, 361
 - gko::dim< Dimensionality, DimensionType >, 646
- operator*
 - gko::dim< Dimensionality, DimensionType >, 646
 - gko::ptr_param< T >, 983
- operator()
 - gko::accessor::row_major< ValueType, Dimensionality >, 1009
 - gko::default_converter< S, R >, 572
 - gko::executor_deleter< T >, 714
 - gko::null_deleter< T >, 918
 - gko::range< Accessor >, 987
- operator->
 - gko::ptr_param< T >, 984
 - gko::range< Accessor >, 988
- operator=
 - gko::array< ValueType >, 424–426
 - gko::BlockOperator, 451
 - gko::Combination< ValueType >, 469
 - gko::Composition< ValueType >, 500, 501
 - gko::experimental::distributed::DistributedBase, 651
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 874
 - gko::experimental::mpi::contiguous_type, 503
 - gko::experimental::mpi::window< ValueType >, 1116
 - gko::index_set< IndexType >, 822, 823
 - gko::matrix::Csr< ValueType, IndexType >, 548
 - gko::matrix::Dense< ValueType >, 615
 - gko::matrix::Ell< ValueType, IndexType >, 678, 679
 - gko::matrix::Fbcsr< ValueType, IndexType >, 727
 - gko::matrix::Hybrid< ValueType, IndexType >, 780
 - gko::matrix::Sellp< ValueType, IndexType >, 1038
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1056
 - gko::preconditioner::lc< LSolverType, IndexType >, 784, 785
 - gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 803
 - gko::preconditioner::lsai< IsaiType, ValueType, IndexType >, 831, 832
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 839, 840
 - gko::range< Accessor >, 988, 989
 - gko::solver::EnableIterativeBase< DerivedType >, 690
 - gko::solver::EnablePreconditionable< DerivedType >, 698
 - gko::solver::EnableSolverBase< DerivedType, MatrixType >, 701

- gko::solver::lr< ValueType >, 827
- gko::solver::LowerTrs< ValueType, IndexType >, 854, 855
- gko::solver::UpperTrs< ValueType, IndexType >, 1084
- operator==
 - gko, 361
 - gko::batch_dim< Dimensionality, DimensionType >, 433
 - gko::dim< Dimensionality, DimensionType >, 647
 - gko::experimental::mpi::communicator, 493
 - gko::stopping_status, 1063
- operator[]
 - gko::dim< Dimensionality, DimensionType >, 645
- OutOfBoundsError
 - gko::OutOfBoundsError, 925
- OverflowError
 - gko::OverflowError, 926
- parse
 - gko::experimental::distributed::preconditioner::Schwarz< ValueType, LocalIndexType, GlobalIndexType >, 1021
 - gko::experimental::factorization::Cholesky< ValueType, IndexType >, 462
 - gko::experimental::factorization::Lu< ValueType, IndexType >, 857
 - gko::experimental::solver::Direct< ValueType, IndexType >, 649
 - gko::factorization::lc< ValueType, IndexType >, 787
 - gko::factorization::llu< ValueType, IndexType >, 799
 - gko::factorization::Parlc< ValueType, IndexType >, 930
 - gko::factorization::Parlct< ValueType, IndexType >, 931
 - gko::factorization::Parllu< ValueType, IndexType >, 933
 - gko::factorization::Parllut< ValueType, IndexType >, 935
 - gko::multigrid::Pgm< ValueType, IndexType >, 959
 - gko::preconditioner::lc< LSolverType, IndexType >, 785
 - gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 804
 - gko::preconditioner::lsai< IsaiType, ValueType, IndexType >, 832
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 840
 - gko::solver::Bicg< ValueType >, 440
 - gko::solver::Bicgstab< ValueType >, 443
 - gko::solver::CbGmres< ValueType >, 454
 - gko::solver::Cg< ValueType >, 457
 - gko::solver::Cgs< ValueType >, 459
 - gko::solver::Fcgs< ValueType >, 731
 - gko::solver::Gcr< ValueType >, 748
 - gko::solver::Gmres< ValueType >, 751
 - gko::solver::ldr< ValueType >, 796
 - gko::solver::lr< ValueType >, 828
 - gko::solver::LowerTrs< ValueType, IndexType >, 855
 - gko::solver::Multigrid, 899
 - gko::solver::UpperTrs< ValueType, IndexType >, 1084
- permute
 - gko::matrix::Csr< ValueType, IndexType >, 548, 549
 - gko::matrix::Dense< ValueType >, 615–619
 - gko::Permutable< IndexType >, 947
- permute_mode
 - gko::matrix, 392
- permute_reuse
 - gko::matrix::Csr< ValueType, IndexType >, 550, 551
- pi
 - gko, 362
- pnode
 - gko::config::pnode, 961–963
- precision_dispatch
 - gko, 362
 - gko::experimental::distributed, 383
- precision_dispatch_real_complex
 - gko, 362, 363
 - gko::experimental::distributed, 383, 384
- precision_reduction
 - gko::precision_reduction, 974
- Preconditioners, 312
- previous_precision_base
 - gko, 334
- process
 - gko::matrix::Csr< ValueType, IndexType >::classical, 464
 - gko::matrix::Csr< ValueType, IndexType >::cusparse, 570
 - gko::matrix::Csr< ValueType, IndexType >::load_balance, 850
 - gko::matrix::Csr< ValueType, IndexType >::merge_path, 893
 - gko::matrix::Csr< ValueType, IndexType >::sparselib, 1048
 - gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1065
- profile_event_category
 - gko::log, 390
- profiling_scope_guard
 - gko::log::profiling_scope_guard, 982
- provided
 - gko::solver, 399
- put
 - gko::experimental::mpi::window< ValueType >, 1116
- r_accumulate
 - gko::experimental::mpi::window< ValueType >, 1117
- r_get

- gko::experimental::mpi::window< ValueType >, 1117
- r_get_accumulate
 - gko::experimental::mpi::window< ValueType >, 1118
- r_put
 - gko::experimental::mpi::window< ValueType >, 1119
- range
 - gko::range< Accessor >, 986
- rank
 - gko::experimental::mpi::communicator, 494
- rapply
 - gko::matrix::Diagonal< ValueType >, 640
- read
 - gko, 363
 - gko::matrix::Coo< ValueType, IndexType >, 517, 518
 - gko::matrix::Csr< ValueType, IndexType >, 551, 552
 - gko::matrix::Ell< ValueType, IndexType >, 679, 680
 - gko::matrix::Fbcsr< ValueType, IndexType >, 728
 - gko::matrix::Hybrid< ValueType, IndexType >, 780, 781
 - gko::matrix::Sellp< ValueType, IndexType >, 1038, 1039
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1056, 1057
 - gko::ReadableFromMatrixData< ValueType, IndexType >, 993, 994
- read_binary
 - gko, 364
- read_binary_raw
 - gko, 365
- read_distributed
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 875, 876
 - gko::experimental::distributed::Vector< ValueType >, 1104
- read_generic
 - gko, 365
- read_generic_raw
 - gko, 366
- read_raw
 - gko, 367
- real
 - gko, 367
- recv
 - gko::experimental::mpi::communicator, 494
- reduce
 - gko::experimental::mpi::communicator, 495
- reduce_add
 - gko, 368
- Reference Executor, 313
- reference_version
 - gko::version_info, 1108
- registry
 - gko::config::registry, 1000, 1001
- remove_complex
 - gko, 334
- remove_logger
 - gko::Executor, 710
 - gko::log::Loggable, 852
- remove_zeros
 - gko::device_matrix_data< ValueType, IndexType >, 634
- request
 - gko::experimental::mpi::request, 1004
- reset_max_iterations
 - gko::batch::solver::BatchSolver, 438
- reset_tolerance
 - gko::batch::solver::BatchSolver, 438
- reset_tolerance_type
 - gko::batch::solver::BatchSolver, 439
- resize_and_reset
 - gko::array< ValueType >, 427
 - gko::device_matrix_data< ValueType, IndexType >, 634, 635
- rhs
 - gko::solver, 399
- round_down
 - gko, 369
- round_up
 - gko, 369
- row_gather
 - gko::matrix::Dense< ValueType >, 619–621
- row_permute
 - gko::matrix::Csr< ValueType, IndexType >, 552
 - gko::matrix::Dense< ValueType >, 621–623
 - gko::Permutable< IndexType >, 947
- row_scale
 - gko::experimental::distributed::Matrix< ValueType, LocalIndexType, GlobalIndexType >, 877
- rows
 - gko::matrix, 393
- run
 - gko::CudaExecutor, 563, 565
 - gko::DpcppExecutor, 656, 657
 - gko::Executor, 710, 711
 - gko::HipExecutor, 761, 762
 - gko::OmpExecutor, 921, 922
 - gko::ReferenceExecutor, 998, 999
- safe_divide
 - gko, 370
- scale
 - gko::batch::matrix::Csr< ValueType, IndexType >, 533
 - gko::batch::matrix::Dense< ValueType >, 587
 - gko::batch::matrix::Ell< ValueType, IndexType >, 669
 - gko::batch::MultiVector< ValueType >, 914
 - gko::experimental::distributed::Vector< ValueType >, 1104
 - gko::matrix::Csr< ValueType, IndexType >, 553
 - gko::matrix::Dense< ValueType >, 623

- scale_add
 - gko::batch::matrix::Dense< ValueType >, 587
- scale_permute
 - gko::matrix::Csr< ValueType, IndexType >, 553
 - gko::matrix::Dense< ValueType >, 623–625
- scan
 - gko::experimental::mpi::communicator, 495
- scatter
 - gko::experimental::mpi::communicator, 496
- scatter_v
 - gko::experimental::mpi::communicator, 496
- scoped_device_id_guard
 - gko::scoped_device_id_guard, 1023, 1024
- segmented_array
 - gko::segmented_array< T >, 1026
- Sellp
 - gko::matrix::Sellp< ValueType, IndexType >, 1031, 1032
- send
 - gko::experimental::mpi::communicator, 497
- set_complex_subspace
 - gko::solver::ldr< ValueType >, 797
- set_complex_subspace
 - gko::solver::ldr< ValueType >, 797
- set_cycle
 - gko::solver::Multigrid, 900
- set_deterministic
 - gko::solver::ldr< ValueType >, 797
- set_executor
 - gko::array< ValueType >, 427
- set_kappa
 - gko::solver::ldr< ValueType >, 798
- set_krylov_dim
 - gko::solver::CbGmres< ValueType >, 454
 - gko::solver::Gcr< ValueType >, 749
 - gko::solver::Gmres< ValueType >, 752
- set_log_propagation_mode
 - gko::Executor, 712
- set_object_name
 - gko::log::ProfilerHook, 980
- set_preconditioner
 - gko::Preconditionable, 977
 - gko::solver::EnablePreconditionable< DerivedType >, 699
- set_solver
 - gko::solver::lr< ValueType >, 828
- set_stop_criterion_factory
 - gko::solver::EnableIterativeBase< DerivedType >, 690
 - gko::solver::IterativeBase, 834
- set_strategy
 - gko::matrix::Csr< ValueType, IndexType >, 554
- set_subspace_dim
 - gko::solver::ldr< ValueType >, 798
- set_value
 - gko::matrix_assembly_data< ValueType, IndexType >, 880
- share
 - gko, 370
- should_propagate_log
 - gko::Executor, 712
- size
 - gko::experimental::mpi::communicator, 498
- solver
 - gko::log, 390
- Solvers, 314
- span
 - gko::span, 1045, 1046
- SparsityCsr
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1050
- SpMV employing different Matrix formats, 309
- squared_norm
 - gko, 371
- status
 - gko::experimental::mpi::status, 1059
- stop
 - gko::stopping_status, 1062
- Stopping criteria, 315
 - combine, 316
 - mode, 316
- strategy_type
 - gko::matrix::Csr< ValueType, IndexType >::strategy_type, 1064
- StreamError
 - gko::StreamError, 1070
- sub_scaled
 - gko::experimental::distributed::Vector< ValueType >, 1105
 - gko::matrix::Dense< ValueType >, 626
- sum_duplicates
 - gko::device_matrix_data< ValueType, IndexType >, 635
 - gko::matrix_data< ValueType, IndexType >, 889
- symmetric
 - gko::matrix, 393
- synchronize
 - gko::experimental::mpi::communicator, 498
- TableSummaryWriter
 - gko::log::ProfilerHook::TableSummaryWriter, 1072
- tag
 - gko::version, 1106
- to_adjacency_matrix
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1057
- to_complex
 - gko, 334
- to_global_indices
 - gko::index_set< IndexType >, 823
- to_real
 - gko, 335
- transpose
 - gko, 371, 372
 - gko::Combination< ValueType >, 469
 - gko::Composition< ValueType >, 501

- gko::experimental::solver::Direct< ValueType, IndexType >, 650
- gko::matrix::Csr< ValueType, IndexType >, 554
- gko::matrix::Dense< ValueType >, 626
- gko::matrix::Diagonal< ValueType >, 640
- gko::matrix::Fbcsr< ValueType, IndexType >, 729
- gko::matrix::Fft, 734
- gko::matrix::Fft2, 738
- gko::matrix::Fft3, 742
- gko::matrix::Identity< ValueType >, 792
- gko::matrix::SparsityCsr< ValueType, IndexType >, 1057
- gko::preconditioner::lc< LSolverType, IndexType >, 786
- gko::preconditioner::llu< LSolverType, USolverType, ReverseApply, IndexType >, 804
- gko::preconditioner::lsai< IsaiType, ValueType, IndexType >, 832
- gko::preconditioner::Jacobi< ValueType, IndexType >, 840
- gko::solver::Bicg< ValueType >, 441
- gko::solver::Bicgstab< ValueType >, 444
- gko::solver::Cg< ValueType >, 457
- gko::solver::Cgs< ValueType >, 460
- gko::solver::Fcg< ValueType >, 731
- gko::solver::Gcr< ValueType >, 749
- gko::solver::Gmres< ValueType >, 752
- gko::solver::ldr< ValueType >, 798
- gko::solver::lr< ValueType >, 829
- gko::solver::LowerTrs< ValueType, IndexType >, 855
- gko::solver::UpperTrs< ValueType, IndexType >, 1085
- gko::Transposable, 1077
- transpose_reuse
 - gko::matrix::Csr< ValueType, IndexType >, 554
- trisolve_algorithm
 - gko::solver, 399
- type
 - gko::cpx_real_type< T >, 521
- type_descriptor
 - gko::config::type_descriptor, 1079
- unit_root
 - gko, 372
- unlock
 - gko::experimental::mpi::window< ValueType >, 1119
- unpack
 - gko::experimental::factorization::Factorization< ValueType, IndexType >, 717
- UnsupportedMatrixProperty
 - gko::UnsupportedMatrixProperty, 1081
- update
 - gko::stop::Criterion, 522
- update_values
 - gko::matrix::Csr< ValueType, IndexType >::permuting_reuse_info, 954
- UpperTrs
 - gko::solver::UpperTrs< ValueType, IndexType >, 1083
- user
 - gko::log, 390
- user_range
 - gko::log::ProfilerHook, 981
- val_at
 - gko::matrix::Ell< ValueType, IndexType >, 680
 - gko::matrix::Sellp< ValueType, IndexType >, 1039, 1040
- ValueMismatch
 - gko::ValueMismatch, 1088
- view
 - gko::array< ValueType >, 428
- wait
 - gko::experimental::mpi::request, 1004
- wait_all
 - gko::experimental::mpi, 387
- window
 - gko::experimental::mpi::window< ValueType >, 1110, 1111
- with_matrix_type
 - gko, 373
- write
 - gko, 374
 - gko::log::ProfilerHook::SummaryWriter, 1071
 - gko::log::ProfilerHook::TableSummaryWriter, 1073
 - gko::matrix::Coo< ValueType, IndexType >, 518
 - gko::matrix::Csr< ValueType, IndexType >, 555
 - gko::matrix::Ell< ValueType, IndexType >, 681
 - gko::matrix::Fbcsr< ValueType, IndexType >, 729
 - gko::matrix::Fft, 734, 735
 - gko::matrix::Fft2, 738, 739
 - gko::matrix::Fft3, 742, 743
 - gko::matrix::Hybrid< ValueType, IndexType >, 781
 - gko::matrix::Permutation< IndexType >, 953
 - gko::matrix::ScaledPermutation< ValueType, IndexType >, 1019
 - gko::matrix::Sellp< ValueType, IndexType >, 1040
 - gko::matrix::SparsityCsr< ValueType, IndexType >, 1058
 - gko::preconditioner::Jacobi< ValueType, IndexType >, 841
 - gko::WritableToMatrixData< ValueType, IndexType >, 1121
- write_binary
 - gko, 374
- write_binary_raw
 - gko, 375
- write_nested
 - gko::log::ProfilerHook::NestedSummaryWriter, 915
 - gko::log::ProfilerHook::TableSummaryWriter, 1073
- write_raw
 - gko, 375
- zero
 - gko, 376

gko::solver, [399](#)